

SECOND EDITION

Programming iOS 5

Matt Neuburg

O'REILLY®

Beijing • Cambridge • Farnham • Köln • Sebastopol • Tokyo

Programming iOS 5, Second Edition

by Matt Neuburg

Revision History for the :

Dec 23, 2011

See <http://oreilly.com/catalog/errata.csp?isbn=9781449319342> for release details.

ISBN: 978-1-449-31934-2
1326317995

Table of Contents

Preface	xvii
---------------	------

Part I. Language

1. Just Enough C	3
Compilation, Statements, and Comments	4
Variable Declaration, Initialization, and Data Types	6
Structs	8
Pointers	10
Arrays	12
Operators	13
Flow Control and Conditions	15
Functions	20
Pointer Parameters and the Address Operator	22
Files	24
The Standard Library	27
More Preprocessor Directives	28
Data Type Qualifiers	29
2. Object-Based Programming	31
Objects	31
Messages and Methods	32
Classes and Instances	33
Class Methods	36
Instance Variables	37
The Object-Based Philosophy	38
3. Objective-C Objects and Messages	43
An Instance Reference Is a Pointer	43
Instance References, Initialization, and nil	44

Instance References and Assignment	47
Instance References and Memory Management	48
Messages and Methods	50
Sending a Message	50
Declaring a Method	51
Nesting Method Calls	52
No Overloading	53
Parameter Lists	54
Unrecognized Selectors	54
Typecasting and the id Type	56
Messages as Data Type	60
C Functions	61
C Struct Pointers	62
Blocks	63
4. Objective-C Classes	67
Class and Superclass	67
Interface and Implementation	69
Header File and Implementation File	71
Class Methods	73
The Secret Life of Classes	74
5. Objective-C Instances	77
How Instances Are Created	77
Ready-Made Instances	77
Instantiation from Scratch	78
Nib-Based Instantiation	81
Polymorphism	82
The Keyword self	84
The Keyword super	87
Instance Variables and Accessors	89
Key–Value Coding	91
Properties	92
How to Write an Initializer	94

Part II. IDE

6. Anatomy of an Xcode Project	99
New Project	99
The Project Window	101
The Navigator Pane	103
The Utilities Pane	107

The Editor	109
The Project File and Its Dependents	111
The Target	114
Build Phases	114
Build Settings	115
Configurations	117
Schemes and Destinations	118
From Project to App	120
Build Settings	122
Property List Settings	122
Nib Files and Storyboard Files	123
Other Resources	124
Code	126
Frameworks and SDKs	128
7. Nib Management	133
A Tour of the Nib-Editing Interface	134
The Dock	135
Canvas	136
Inspectors and Libraries	138
Nib Loading and File's Owner	140
Making and Loading a Nib	142
Outlet Connections	143
More Ways to Create Outlets	148
More About Outlets	150
Action Connections	151
Additional Initialization of Nib-Based Instances	155
8. Documentation	157
The Documentation Window	158
Class Documentation Pages	159
Sample Code	163
Other Resources	164
Quick Help	164
Symbols	165
Header Files	165
Internet Resources	166
9. Life Cycle of a Project	169
Choosing a Device Architecture	169
Localization	173
Editing Your Code	174
Autocompletion	175

Snippets	176
Live Syntax Checking and Fix-it	177
Navigating Your Code	178
Debugging	180
Caveman Debugging	180
The Xcode Debugger	183
Static Analyzer	188
Clean	189
Running in the Simulator	190
Running on a Device	191
Device Management	195
Version Control	196
Instruments	198
Distribution	202
Ad Hoc Distribution	204
Final App Preparations	205
Icons in the App	206
Other Icons	207
Launch Images	207
Screenshots	208
Property List Settings	209
Submission to the App Store	210

Part III. Cocoa

10. Cocoa Classes	215
Subclassing	215
Categories	218
Splitting a Class	219
Private Method Declarations	220
Protocols	221
Optional Methods	226
Some Foundation Classes	227
Useful Structs and Constants	227
NSString and Friends	228
NSDate and Friends	230
NSNumber	230
NSValue	231
NSData	231
Equality and Comparison	232
NSIndexSet	232
NSArray and NSMutableArray	233

NSSet and Friends	234
NSDictionary and NSMutableDictionary	235
NSNull	237
Immutable and Mutable	237
Property Lists	238
The Secret Life of NSObject	238
11. Cocoa Events	243
Reasons for Events	244
Subclassing	244
Notifications	246
Receiving a Built-In Notification	247
Unregistering	249
NSTimer	251
Delegation	251
Data Sources	255
Actions	256
The Responder Chain	261
Deferring Responsibility	261
Nil-Targeted Actions	262
Application Lifetime Events	263
Swamped by Events	268
12. Accessors and Memory Management	273
Accessors	273
Key–Value Coding	275
Memory Management	279
Principles of Cocoa Memory Management	279
The Golden Rules of Memory Management	281
What ARC Is and What It Does	283
How Cocoa Objects Manage Memory	286
Autorelease	288
Memory Management of Instance Variables (Non-ARC)	291
Memory Management of Instance Variables (ARC)	295
Retain Cycles and Weak References	297
Nib Loading and Memory Management	303
Memory Management of Global Variables	305
Memory Management of Pointer-to-Void Context Info	306
Memory Management of CFTypeRefs	308
Properties	310
13. Data Communication	317
Model–View–Controller	317

Instance Visibility	319
Visibility by Instantiation	320
Visibility by Relationship	321
Global Visibility	322
Notifications	323
Key–Value Observing	324

Part IV. Views

14. Views	333
The Window	333
Subview andSuperview	336
Frame	339
Bounds and Center	341
Layout	344
Transform	347
Visibility and Opacity	351
15. Drawing	353
UIImage and UIImageView	353
Graphics Contexts	357
UIImage Drawing	361
CGImage Drawing	362
CIFilter and CIImage	365
Drawing a UIView	368
Graphics Context Settings	370
Paths and Drawing	371
Clipping	375
Gradients	376
Colors and Patterns	378
Graphics Context Transforms	380
Shadows	382
Points and Pixels	383
Content Mode	383
16. Layers	387
View and Layer	388
Layers and Sublayers	390
Manipulating the Layer Hierarchy	391
Positioning a Sublayer	392
CAScrollLayer	393
Layout of Sublayers	394

Drawing in a Layer	394
Content Resizing and Positioning	397
Layers that Draw Themselves	399
Transforms	400
Depth	404
Shadows, Borders, and More	406
Layers and Key–Value Coding	409
17. Animation	411
Drawing, Animation, and Threading	412
UIImageView and UIImageView Animation	415
View Animation	417
Animation Blocks	417
Modifying an Animation Block	418
Transition Animations	422
Block-Based View Animation	423
Implicit Layer Animation	428
Animation Transactions	429
Media Timing Functions	430
Core Animation	432
CABasicAnimation and Its Inheritance	432
Using a CABasicAnimation	434
Keyframe Animation	437
Making a Property Animatable	438
Grouped Animations	439
Transitions	443
The Animations List	445
Actions	447
What an Action Is	447
The Action Search	448
Hooking Into the Action Search	449
Nonproperty Actions	452
Emitter Layers	453
18. Touches	461
Touch Events and Views	462
Receiving Touches	464
Restricting Touches	465
Interpreting Touches	466
Gesture Recognizers	471
Gesture Recognizer Classes	471
Multiple Gesture Recognizers	476
Subclassing Gesture Recognizers	477

Gesture Recognizer Delegate	479
Touch Delivery	481
Hit-Testing	482
Initial Touch Event Delivery	487
Gesture Recognizer and View	488
Touch Exclusion Logic	489
Recognition	490
Touches and the Responder Chain	491

Part V. Interface

19. View Controllers	495
The View Controller Hierarchy	498
View Controller and View Creation	502
Manual View	505
Generic Automatic View	507
View in a Separate Nib	509
Nib-Instantiated View Controller	512
Storyboard-Instantiated View Controller	515
Rotation	517
Initial Orientation	519
Rotation Events	522
Presented View Controller	523
Presented View Animation	528
Presentation Styles	529
Presented Views and Rotation	532
Tab Bar Controllers	534
Tab Bar Items	535
Configuring a Tab Bar Controller	536
Navigation Controllers	538
Bar Button Items	542
Navigation Items	544
Toolbar Items	546
Configuring a Navigation Controller	547
Page View Controller	549
Container View Controllers	552
Storyboards	554
View Controller Lifetime Events	560
View Controller Memory Management	562
20. Scroll Views	567
Creating a Scroll View	568

Scrolling	571
Paging	574
Tiling	575
Zooming	577
Zooming Programmatically	579
Zooming with Detail	579
ScrollView Delegate	582
ScrollView Touches	584
ScrollView Performance	589
21. Table Views	591
Table View Cells	594
Built-In Cell Styles	595
Custom Cells	601
Table View Data	609
The Three Big Questions	610
Table View Sections	614
Refreshing Table View Data	618
Variable Row Heights	619
Table View Selection	621
Table View Scrolling and Layout	627
Table View Searching	628
Table View Editing	634
Deleting Table Items	637
Editable Content in Table Items	638
Inserting Table Items	640
Rearranging Table Items	642
Dynamic Table Content	643
Table View Menus	644
22. Popovers and Split Views	647
Configuring and Displaying a Popover	649
Managing a Popover	654
Dismissing a Popover	654
Popover Segues	658
Automatic Popovers	659
Split Views	661
23. Text	667
UILabel	668
UITextField	669
Editing and the Keyboard	672
Configuring the Keyboard	676

Text Field Delegate and Control Event Messages	677
The Text Field Menu	679
UITextView	681
Core Text	684
24. Web Views	693
Loading Content	694
Communicating with a Web View	700
25. Controls and Other Views	703
UIActivityIndicatorView	703
UIProgressView	704
UIPickerView	707
UISearchBar	709
UIControl	712
UISwitch	714
UIStepper	714
UIPageControl	715
UIDatePicker	716
UISlider	718
UISegmentedControl	721
UIButton	723
Custom Controls	727
Bars	730
UINavigationBar	730
UIToolbar	734
UITabBar	735
Appearance Proxy	739
26. Modal Dialogs	743
Alert View	744
Action Sheet	746
Dialog Alternatives	750
Local Notifications	751

Part VI. Some Frameworks

27. Audio	759
System Sounds	759
Audio Session	760
Audio Player	764
Remote Control of Your Sound	766

Playing Sound in the Background	768
Further Topics in Sound	769
28. Video	771
MPMoviePlayerController	772
MPMoviePlayerViewController	777
UIVideoEditorController	778
Further Topics in Video	779
29. Music Library	783
Exploring the Music Library	783
The Music Player	787
The Music Picker	791
30. Photo Library	795
UIImagePickerController	795
Choosing from the Photo Library	796
Using the Camera	797
The Assets Library Framework	800
31. Address Book	803
Address Book Database	803
Address Book Interface	805
ABPeoplePickerNavigationController	806
ABPersonViewController	808
ABNewPersonViewController	808
ABUnknownPersonViewController	809
32. Calendar	811
Calendar Database	811
Calendar Interface	816
33. Mail	819
Mail Message	819
SMS Message	820
34. Maps	821
Presenting a Map	821
Annotations	823
Overlays	828

35. Sensors	833
Location	833
Heading and Course	836
Acceleration	837
Shake Events	837
UIAccelerometer	838
Core Motion	841
<hr/>	
Part VII. Final Topics	
36. Persistent Storage	845
The Sandbox	845
Basic File Operations	846
Saving and Reading Files	847
User Defaults	849
File Sharing	851
Document Types	851
Handing Off a Document	853
XML	856
SQLite	862
Image File Formats	863
37. Basic Networking	867
HTTP Requests	867
Bonjour	873
Push Notifications	875
Beyond Basic Networking	876
38. Threads	877
The Main Thread	877
Why Threading Is Hard	880
Three Ways of Threading	881
Manual Threads	883
NSOperation	884
Grand Central Dispatch	889
Threads and App Backgrounding	891
39. Undo	895
The Undo Manager	895
The Undo Interface	898
The Undo Architecture	901

40. Epilogue	903
--------------------	-----

Preface

Aut lego vel scribo; doceo scrutorve sophian.

—Sedulius Scottus

With the advent of version 2 of the iPhone system, Apple proved they could do a remarkable thing — adapt their existing Cocoa computer application programming framework to make applications for a touch-based device with limited memory and speed and a dauntingly tiny display. The resulting Cocoa Touch framework, in fact, turned out to be in many ways better than the original Cocoa.

A programming framework has a kind of personality, an overall flavor that provides an insight into the goals and mindset of those who created it. When I first encountered Cocoa Touch, my assessment of its personality was: “Wow, the people who wrote this are really clever!” On the one hand, the number of built-in interface widgets was severely and deliberately limited; on the other hand, the power and flexibility of some of those widgets, especially such things as UITableView, was greatly enhanced over their Mac OS X counterparts. Even more important, Apple created a particularly brilliant way (UIViewController) to help the programmer make entire blocks of interface come and go and supplant one another in a controlled, hierarchical manner, thus allowing that tiny iPhone display to unfold virtually into multiple interface worlds within a single app without the user becoming lost or confused.

Even more impressive, Apple took the opportunity to recreate and rationalize Cocoa from the ground up as Cocoa Touch. Cocoa itself is very old, having begun life as NeXTStep before Mac OS X even existed. It has grown by accretion and with a certain conservatism in order to maintain something like backward compatibility. With Cocoa Touch, on the other hand, Apple had the opportunity to throw out the baby with the bath water, and they seized this opportunity with both hands.

So, although Cocoa Touch is conceptually based on Mac OS X Cocoa, it is very clearly *not* Mac OS X Cocoa, nor is it limited or defined by Mac OS X Cocoa. It’s an independent creature, a leaner, meaner, smarter Cocoa. I could praise Cocoa Touch’s deliberate use of systematization (and its healthy respect for Occam’s Razor) through numerous examples. Where Mac OS X’s animation layers are glommed onto views as a kind of afterthought, a Cocoa Touch view always has an animation layer counterpart.

Memory management policies, such as how top-level objects are managed when a nib loads, are simplified and clarified. And so on.

At the same time, Cocoa Touch is still a form of Cocoa. It still requires a knowledge of Objective-C. It is not a scripting language; it is certainly not aimed at nonprogrammers, like HyperCard’s HyperTalk or Apple’s AppleScript. It is still huge and complicated. In fact, it’s rather difficult.

The popularity of the iPhone, with its largely free or very inexpensive apps, and the subsequent popularity of the iPad, have brought and will continue to bring into the fold many new programmers who see programming for these devices as worthwhile and doable, even though they may not have felt the same way about Mac OS X. Apple’s own annual WWDC developer conventions have reflected this trend, with their emphasis shifted from Mac OS X to iOS instruction.

The widespread eagerness to program iOS, however, though delightful on the one hand, has also fostered a certain tendency to try to run without first learning to walk. iOS gives the programmer mighty powers that can seem as limitless as imagination itself, but it also has fundamentals. I often see questions online from programmers who are evidently deep into the creation of some interesting app, but who are stymied in a way that reveals quite clearly that they are unfamiliar with the basics of the very world in which they are so happily cavorting.

It is this state of affairs that has motivated me to write this book, which is intended to ground the reader in the fundamentals of iOS. I love Cocoa and have long wished to write about it, but it is iOS and its popularity that has given me a proximate excuse to do so. Indeed, my working title was “Fundamentals of Cocoa Touch Programming.” Here I have attempted to marshal and expound, in what I hope is a pedagogically helpful and instructive yet ruthlessly Euclidean and logical order, the principles on which sound iOS programming rests, including a good basic knowledge of Objective-C (starting with C itself) and the nature of object-oriented programming, advice on the use of the tools, the full story on how Cocoa objects are instantiated, referred to, put in communication with one another, and managed over their lifetimes, and a survey of the primary interface widgets and other common tasks. My hope, as with my previous books, is that you will both read this book cover to cover (learning something new often enough to keep you turning the pages) and keep it by you as a handy reference.

This book is not intended to disparage Apple’s own documentation and example projects. They are wonderful resources and have become more wonderful as time goes on. I have depended heavily on them in the preparation of this book. But I also find that they don’t fulfill the same function as a reasoned, ordered presentation of the facts. The online documentation must make assumptions as to how much you already know; it can’t guarantee that you’ll approach it in a given order. And online documentation is more suitable to reference than to instruction. A fully written example, no matter how well commented, is difficult to follow; it demonstrates, but it does not teach.

A book, on the other hand, has numbered chapters and sequential pages; I can assume you know C before you know Objective-C for the simple reason that Chapter 1 precedes Chapter 2. And along with facts, I also bring to the table a degree of experience, which I try to communicate to you. Throughout this book you'll see me referring to "common beginner mistakes"; in most cases, these are mistakes that I have made myself, in addition to seeing others make them. I try to tell you what the pitfalls are because I assume that, in the course of things, you will otherwise fall into them just as naturally as I did as I was learning. You'll also see me construct many examples piece by piece or extract and explain just one tiny portion of a larger app. It is not a massive finished program that teaches programming, but an exposition of the thought process that developed that program. It is this thought process, more than anything else, that I hope you will gain from reading this book.

iOS is huge, massive, immense. It's far too big to be encompassed in a book even of this size. And in any case, that would be inappropriate and unnecessary. There are entire areas of Cocoa Touch that I have ruthlessly avoided discussing. Some of them would require an entire book of their own. Others you can pick up well enough, when the time comes, from the documentation. This book is only a beginning — the fundamentals. But I hope that it will be the firm foundation that will make it easier for you to tackle whatever lies beyond, in your own fun and rewarding iOS programming future.

Acknowledgments for the First Edition

It's a poor craftsman who blames his tools. No blame attaches to the really great tools by which I have been assisted in the writing of this book. I am particularly grateful to the Unicomp Model M keyboard (<http://pckeyboard.com>), without which I could not have produced so large a book so painlessly. I was also aided by wonderful software, including TextMate (<http://macromates.com>) and AsciiDoc (<http://www.methods.co.nz/asciidoc>). BBEdit (<http://www.barebones.com>) helped with its diff display. Screenshots were created with Snapz Pro X (<http://www.ambrosiasw.com>) and GraphicConverter (<http://www.lemkesoft.com>); diagrams were drawn with OmniGraffle (<http://www.omnigroup.com>).

The splendid O'Reilly production process converted my AsciiDoc text files into PDF while I worked, allowing me to proofread in simulated book format. Were it not for this, and the Early Release program that permitted me to provide my readers with periodic updates of the book as it grew, I would never have agreed to undertake this project in the first place. I would like particularly to thank Tools maven Abby Fox for her constant assistance.

I have taken advice from two tech reviewers, Dave Smith and David Rowland, and have been assisted materially and spiritually by many readers who submitted errata and encouragement. I was particularly fortunate in having Brian Jepson as editor; he provided enthusiasm for the O'Reilly tools and the electronic book formats, a watchful

eye, and a trusting attitude; he also endured the role of communications pipeline when I needed to prod various parts of the O'Reilly machine. I have never written an O'Reilly book without the help of Nancy Kotary, and I didn't intend to start now; her sharp eye has smoothed the bristles of my punctuation-laden style. For errors that remain, I take responsibility, of course.

Notes on the Second Printing

For the second printing of this book, screen shots have been rendered more legible, and a major technical error in the presentation of key-value coding in [Chapter 5](#) has been corrected. In addition, numerous small emendations have been made; many of these have resulted from errata submissions by my readers, whom I should like to thank once again for their continued assistance and kind support. Please note that these changes have altered the pagination of the printed and PDF editions of the book.

Notes on the Second Edition

In order to describe the relationship of the second edition of this book with the first edition, it will help if I first recap the recent history of iOS and Xcode versions.

At the time I started writing the first edition this book, system versions 3.1.3 (on the iPhone) and 3.2 (on the iPad) were current. As I was working on the book, iOS 4 and the iPhone 4 came into being, but iOS 4 didn't yet run on the iPad. Subsequently iOS 4.2 emerged; this was the first system able to run on both the iPhone and the iPad. At the same time, Xcode was improved up to 3.2.5. iOS 4 was the first version of the system to support multitasking, which necessitated much scurrying about on the part of developers, to adapt their apps to the new world order.

Just in time for my final revisions of the first edition, Xcode 3.2.6 and iOS 4.3 were released, along with the first public version of the long-awaited Xcode 4. Xcode 4 was a thorough overhaul of the IDE: menus, windows, and preferences are quite different from Xcode 3.2.x. Both Xcode 4 and Xcode 3.2.x can coexist on the same machine (up through Snow Leopard) and can be used to work on the same project; moreover, Xcode 3.2.x has some specialized capabilities that Xcode 4 lacks, so some long-standing developers may well continue to use it. This situation presents a dilemma for an author describing the development process. However, for iOS programming, I recommend adoption of Xcode 4, and the first edition of this book assumed that the reader had adopted it.

Such was the situation in May 2011, when the first edition was formally released, describing how to program iOS 4.

Less than five months later, in October 2011, Apple released iOS 5. Some of the features that are new in iOS 5 are dramatic and pervasive, and it is this fact which has necessitated a full revision of this book. At the same time, Apple also released Xcode 4.2, and

this book assumes that you are using that version of Xcode (or later), since it is the earliest version of Xcode on which iOS 5 development is officially possible. (It may be that, by deep trickery, one can develop for iOS 5 using an earlier version of Xcode, but that would constitute unsupported behavior.) The first edition had a few mentions of menu commands and other interface in Xcode 3.2.x, but they have been excised from this edition. Xcode 4.2 comes in two flavors, depending whether you’re running Snow Leopard (Mac OS X 10.6) or Lion (Mac OS X 10.7) on your development machine; they are supposed to behave more or less identically, but in fact each has its own bugs, so feel free to try both.

The chief purpose of this new edition, then, is to bring the book up to date for iOS 5. You, the reader, might be coming to iOS programming for the first time, so this edition assumes no prior knowledge of iOS 4 or any previous version. On the other hand, you, like me, could be making the transition from iOS 4 to iOS 5, so this edition lays some special emphasis on features that are new in iOS 5. This emphasis could also be useful to new iOS programmers who are thinking of writing apps that can also run under iOS 4. My goal, however, is not to burden the reader with outdated information. The vast majority of devices that could run iOS 4 have probably been updated to iOS 5, and you will probably be right in assuming that there will plenty of iOS 5 users out there, without your having to bother to target earlier systems. And from a pedagogical point of view, it seems counterproductive to describe how things used to be — especially as, if you’re really interested, you can always consult the previous edition of this book! For this reason, some references to the state of things before iOS 4.2 have been excised from this edition.

Here is a case in point, showing my attitude and pedagogical approach with regard to new iOS 5 features in this edition. iOS 5 introduces ARC (automatic reference counting), which changes the way in which Objective-C programmers manage object memory so profoundly as to render Objective-C a different language. Use of ARC is optional in programming iOS, but it is extraordinarily helpful to have it turned on, and in this book I therefore assume throughout that you do have it turned on. In [Chapter 12](#), where I discuss memory management, I still describe what life is like without ARC, as I did in the previous edition; but, outside that chapter, all code examples, unless specifically state otherwise, are supposed to be running under ARC. If you start a new Xcode project with File → New Project and pick any iOS application template, then if “Use Automatic Reference Counting” is checked in the second screen, you’re using ARC.

iOS 5 also introduces storyboards. A storyboard file is similar to a nib file: it’s a place where Xcode lets you “draw” parts of the interface. The main difference is that a single storyboard file can do the work of multiple nib files. Nib files and storyboard files are not identical, nor are they used identically, but because of their similarity, when I speak of a nib file generically, in this book, I mean a nib or storyboard file, indifferently. I’ll try to indicate this at the time, but the reader will forgive me if I don’t keep saying “nib or storyboard” all the time.

In closing, I should like to say a few words to the people who have, in my opinion, gratuitously criticized the previous edition of this book on one or more of the following grounds:

- a. It isn't a "cookbook" (a book full of step-by-step instructions for creating full working applications).
- b. It devotes hundreds of pages to fundamentals.
- c. It doesn't get the reader started early on with hands-on programming; there isn't even a "Hello, World" tutorial.

All of that is perfectly true. It is also quite deliberate. As both the table of contents and this preface are at pains to make clear, this is not that type of book. To paraphrase Butler's Law, this book is the type of book it is, and not some other type. That's why I wrote this book in the first place. The books of the type that these critics seem to want this book to be exist by the score; books of the type that this book is, however, seemed to me not to exist at all. As with all my other books, so with this one: when I couldn't find the book I wanted, I wrote it myself. I expect this book to be useful to those who need this type of book. People who prefer some other type of book should get some other type of book, and not mar my book's Web page by criticizing it for not being what it was never intended to be.

The purpose of this book is to proceed by focusing on the underlying knowledge needed for an actual understanding of iOS programming. That is precisely the opposite of a cookbook. This book has no simple recipes ready for you to drop into your own code and come up with an app. I don't give you some fish; I teach you what a fish is and what you need to know to obtain one. The number of books out there that skip blithely past the fundamentals, leaving the reader to pick up Objective-C somehow independently, is painfully huge. The result is that numerous learners are encouraged to try their hand at programming when, to judge from the nature of the questions they ask and the confusion they exhibit, they don't understand what they are doing. This book acts as a corrective, and this requires that space be devoted to fundamentals. The book does not hold a gun to your head and force you to read all about all of those fundamentals; if *you* already know everything there is to know about C, about Objective-C, about Xcode, about Cocoa, about views and drawing or whatever (but *do you?* do you *really?*), then by all means, skip those opening chapters. But don't begrudge to the people who need them the explanations that this book contains, as those are the people at whom they are aimed.

That explains why there's no attempt, in this book, to rush the reader into hands-on programming. My book does not pander to a desire for the false, cheap gratification of making believe that one is a programmer merely because one can parrot some instructions. My book is about *knowledge* — hard-won, rigorously gained knowledge. It's about gaining an understanding of what you are doing when you program iOS 5. It calls for a substantial investment of time and thought, and many pages elapse before any practical programming is demonstrated.

However, the complaint that the reader of my book doesn't get to run any code is factually false. The book is crammed full of substantial code examples — all of which, moreover, are available for download from my GitHub site (<https://github.com/mattneub>), so you can download them, run them in Xcode, and play with them to your heart's content. So you can and should, in fact, be running code right from the outset. Nevertheless, the purpose of the code in this book is not for the fun of running it. All of my code is to support your understanding of whatever *concepts* I'm explaining at that point in the book.

In any case, perfectly good hands-on “Hello, World” tutorials are a dime a dozen; they're plastered all over the Internet, including at Apple's own site (<http://developer.apple.com/library/ios/#documentation/iPhone/Conceptual/iPhone101/Articles/>). You don't need *me* to show you that the process of writing a trivial iPhone application is fun and easy.

Still, for those who feel strongly that I haven't done my job unless I supply a “Hello, World” example, here is one, complete with step-by-step instructions:

1. Install Xcode 4.2, and launch the Xcode application.
2. Choose File → New → New Project.
3. In the “Choose a template” dialog, under “iOS” on the left (not “Mac OS X”), click Application. On the right, click Empty Application. Click Next.
4. For Product Name, type `Hello`. Enter a company identifier if there isn't one already, such as `com.yourLastName.yourFirstName`. Choose Device Family: iPhone. Uncheck all three checkboxes. Click Next.
5. Navigate to the Desktop. Uncheck “Create local git repository.” Click Create.
6. The project window opens. Press Command-1. At the left, click `AppDelegate.m`.
7. Work in the editor in the middle of the window. Look for the words “Override point for customization after application launch.” Place the cursor to the right of those words and hit Return a few times, to make some white space. Click in that white space and type these lines of code:

```
UILabel* label = [[UILabel alloc] init];
label.text = @"Hello, world!";
[label sizeToFit];
CGRect f = label.frame;
f.origin = CGPointMake(100,100);
label.frame = f;
[self.window addSubview:label];
```

8. Press Command-R. If you see a dialog asking whether you want to save, accept.
9. After a while, the iOS Simulator application appears, containing a white window with “Hello, world!” in it.

Congratulations. You've made a “Hello, world” example. Wasn't that easy? Wasn't it boring? Wasn't it pointless? And are you any the wiser as to what, in fact, you just did?

To find out—and, even more important, to know enough to be able to progress further on your own—read this book.

Language

Apple has provided a vast toolbox for programming iOS to make an app come to life and behave the way you want it to. That toolbox is the *API* (application programming interface). To use the API, you must speak the API's language. That language, for the most part, is Objective-C, which itself is built on top of C; some pieces of the API use C itself. This part of the book instructs you in the basics of these languages:

- [Chapter 1](#) explains C. In general, you will probably not need to know all the ins and outs of C, so this chapter restricts itself to those aspects of C that you need to know in order to use both Objective-C and the C-based areas of the API.
- Objective-C adds object-based programming features to C. [Chapter 2](#) discusses object-based programming in general architectural terms. It also explains some extremely important terms that will be used throughout the book, along with the concepts that lie behind them.
- [Chapter 3](#) introduces the basic syntax of Objective-C.
- [Chapter 4](#) continues the explanation of Objective-C, discussing the nature of Objective-C classes, with emphasis on how to create a class in code.
- [Chapter 5](#) completes the introduction to Objective-C, discussing how instances are created and initialized, along with an explanation of such related topics as polymorphism, instance variables, accessors, `self` and `super`, key-value coding, and properties.

Even at this point, our discussion of Objective-C is not yet complete; we'll return in [Part III](#) to a description of further aspects of the language: those that are particularly bound up with the Cocoa frameworks.

Just Enough C

Do you believe in C? Do you believe in anything that has to do with me?

—Leonard Bernstein and Stephen Schwartz, *Mass*

To program for iOS, you need some knowledge of the C programming language, for two reasons:

- Most of your iOS programming will be in the Objective-C language, and Objective-C is a superset of C. This means that Objective-C presupposes C; everything that is true of C trickles up to Objective-C. A common mistake is to forget that “Objective-C is C” and to neglect a basic understanding of C.
- Some of the iOS API involves C rather than Objective-C. Even in Objective-C code, you often need to use C data structures and C function calls. For example, a rectangle is represented as a CGRect, which is a C *struct*, and to create a CGRect from four numbers you call CGRectMake, which is a C *function*. The iOS API documentation will very often show you C expressions and expect you to understand them.

The best way to learn C is to read *The C Programming Language* (PTR Prentice Hall, 1988) by Brian W. Kernighan and Dennis M. Ritchie, commonly called K&R (Ritchie was the creator of C). It is one of the best computer books ever written: brief, dense, and stunningly precise and clear. K&R is so important for effective iOS (and Mac OS X) programming that I keep a physical copy beside me at all times while coding, and I recommend that you do the same. Another useful manual is *The C Book*, by Mike Banahan, Declan Brady and Mark Doran, available online at http://publications.gbdirect.co.uk/c_book/.

You don’t have to know *all* about C in order to use Objective-C effectively, though; and that’s a good thing. C is not a large or difficult language, but it has some tricky corners and can be extremely subtle, powerful, and low-level. Also, it would be impossible, and unnecessary, for me to describe all of C in a single chapter. C is described far more fully and correctly in K&R, *The C Book*, and elsewhere than I could possibly

do it. Sooner or later, you’re probably going to have technical questions about C that this chapter doesn’t (and shouldn’t) make any attempt to answer. So I emphasize that you really, really ought to have K&R or something similar at hand and resort to it as needed.

What I can do, and what this chapter will attempt to do, is tell you what aspects of C are important to understand at the outset, before you even start using Objective-C for iOS programming. That’s why this chapter is “Just Enough C”: it’s just enough to get you going, comfortably and safely.

If you know no C at all, I suggest that, as an accompaniment to this chapter, you also read parts of K&R (think of this as “C: The Good Parts Version”). Here’s my proposed K&R syllabus:

- Quickly skim K&R Chapter 1, the tutorial.
- Carefully read K&R Chapters 2 through 4.
- Read the first three sections of K&R Chapter 5 on pointers and arrays. You don’t need to read the rest of Chapter 5 because you won’t typically be doing any pointer arithmetic, but you do need to understand clearly what a pointer is, as Objective-C is all about objects and every reference to an object is a pointer; you’ll be seeing and using that `*` character constantly.
- Read also the first section of K&R Chapter 6, on structures (structs); as a beginner, you probably won’t define any structs, but you will use them quite a lot, so you’ll need to know the notation (for example, as I’ve already said, a `CGRect` is a struct).
- Glance over K&R Appendix B, which covers the standard library, because you may find yourself making certain standard library calls, such as the mathematical functions; forgetting that the library exists is a typical beginner mistake.

Just to make things a little more confusing, the C defined in K&R is not precisely the C that forms the basis of Objective-C. Developments subsequent to K&R have resulted in further C standards (ANSI C, C89, C99), and the Xcode compiler extends the C language in its own ways. By default, Xcode projects are treated as GNU99, which is itself an extension of C99 (though you could specify another C standard if you really wanted to). Fortunately, the most important differences between K&R’s C and Xcode’s C are small, convenient improvements that are easily remembered, so K&R remains the best and most reliable C reference.

Compilation, Statements, and Comments

C is a compiled language. You write your program as text; to run the program, things proceed in two stages. First your text is compiled into machine instructions; then those machine instructions are executed. Thus, as with any compiled language, you can make two kinds of mistake:

- Any purely syntactic errors (meaning that you spoke the C language incorrectly) will be caught by the compiler, and the program won’t even begin to run.
- If your program gets past the compiler, then it will run, but there is no guarantee that you haven’t made some other sort of mistake, which can be detected only by noticing that the program doesn’t behave as intended.

The C compiler is fussy, but you should accept its interference with good grace. The compiler is your friend: learn to love it. It may emit what looks like an irrelevant or incomprehensible error message, but when it does, the fact is that you’ve done something wrong and the compiler has helpfully caught it for you. Also, the compiler can warn you if something seems like a possible mistake, even though it isn’t strictly illegal; these warnings, which differ from outright errors, are also helpful and should not be ignored.

I have said that running a program requires a preceding stage: compilation. But in fact there is a third stage that precedes compilation: preprocessing. (It doesn’t really matter whether you think of preprocessing as a stage preceding compilation or as the first stage of compilation.) Preprocessing modifies your text, so when your text is handed to the compiler, it is not identical to the text you wrote. Preprocessing might sound tricky and intrusive, but in fact it proceeds only according to your instructions and is helpful for making your code clearer and more compact.

Xcode allows you to view the effects of preprocessing on your program text (choose Product → Generate Output → Generate Preprocessed File), so if you think you’ve made a mistake in instructing the preprocessor, you can track it down. I’ll talk more later about some of the things you’re likely to say to the preprocessor.

C is a statement-based language; every statement ends in a semicolon. (Forgetting the semicolon is a common beginner’s mistake.) For readability, programs are mostly written with one statement per line, but this is by no means a hard and fast rule: long statements (which, unfortunately, arise very commonly because of Objective-C’s verbosity) are commonly split over multiple lines, and extremely short statements are sometimes written two or three to a line. You cannot split a line just anywhere, however; for example, a literal string can’t contain a return character. Indentation is linguistically meaningless and is purely a matter of convention (and C programmers argue over those conventions with near-religious fervor); Xcode helps “intelligently” by indenting automatically, and you can use its automatic indentation both to keep your code readable and to confirm that you’re not making any basic syntactic mistakes.

Comments are delimited in K&R C by /* ... */; the material between the delimiters can consist of multiple lines (K&R 1.2). In modern versions of C, a comment also can be denoted by two slashes (//); the rule is that if two slashes appear, they and everything after them on the same line are ignored:

```
int lower = 0; // lower limit of temperature table
```

Choosing a Compiler

The compiler situation in Xcode is rather complicated. Originally, Xcode's compiler was the free open source GCC (<http://gcc.gnu.org>). More recently, Xcode has phased in use of another free open source compiler, LLVM (<http://llvm.org>). Changing compilers is scary, so Apple has proceeded in stages, as follows:

- A hybrid compiler, LLVM-GCC, provides the advantages of LLVM compilation, but the code is parsed with GCC for maximum backward compatibility.
- A pure LLVM compiler (also referred to as Clang) does its own parsing and provides more intelligent and helpful error messages and warnings.

As Xcode 3.2.x evolved, LLVM-GCC was eventually considered the best choice, but Apple was hesitant to make it the default compiler, so GCC remained the default. When Xcode 4 emerged, LLVM-GCC was the default compiler, but GCC remained available. Finally, in Xcode 4.2, LLVM 3.0 became the default compiler, and pure GCC was withdrawn. (The choice of compiler is a project-level build setting; see [Chapter 6](#).)

These are sometimes called C++-style comments and are much more convenient for brief comments than the K&R comment syntax.

Throughout the C language (and therefore, throughout Objective-C as well), capitalization matters. All names are case-sensitive. There is no such data type as Int; it's lowercase "int." If you declare an int called `lower` and then try to speak of the same variable as `Lower`, the compiler will complain. By convention, variable names tend to start with a lowercase letter.

Variable Declaration, Initialization, and Data Types

C is a strongly typed language. Every variable must be declared, indicating its data type, before it can be used. Declaration can also involve explicit initialization; a variable that is declared but not explicitly initialized is of uncertain value (and should be regarded as dangerous until it *is* initialized). In K&R C, declarations must precede all other statements, but in modern versions of C, this rule is relaxed so that you don't have to declare a variable until just before you start using it. The usual convention is thus to declare a variable and assign it a value as it makes its first appearance on the scene:

```
int height = 2;
int width = height * 2;
height = height + 1;
int area = height * width;
```

The basic built-in C data types are all numeric: char (one byte), int (four bytes), float and double (floating-point numbers), and varieties such as short (short integer), long (long integer), unsigned short, and so on. iOS makes use of some further numeric types

derived from the C numeric types (by way of the `typedef` statement, K&R 6.7); the most important of these are `NSInteger` (along with `NSUInteger`) and `CGFloat`. You don't need to use these explicitly unless an API tells you to, and even when you do, just think of `NSInteger` as `int` and `CGFloat` as `float`, and you'll be fine.

To *cast* (or *typecast*) a variable's value explicitly to another type, precede the variable's name with the other type's name in parentheses:

```
int height = 2;
float fheight = (float)height;
```

In that particular example, the explicit cast is unnecessary because the integer value will be cast to a float implicitly as it is assigned to a float variable, but it illustrates the notation. You'll find yourself typecasting quite a bit in Objective-C, mostly in order to subdue the worries of the compiler (examples appear in [Chapter 3](#)).

Another form of numeric initialization is the enum (K&R 2.3). It's a way of assigning names to a sequence of numeric values and is useful when a value represents one of several possible options. The Cocoa API uses this device a lot. For example, the three possible types of status bar animation are defined like this:

```
typedef enum {
    UIStatusBarAnimationNone,
    UIStatusBarAnimationFade,
    UIStatusBarAnimationSlide,
} UIBarStatusAnimation;
```

That definition assigns the value 0 to the name `UIStatusBarAnimationNone`, the value 1 to the name `UIStatusBarAnimationFade`, and the value 2 to the name `UIStatusBarAnimationSlide`. The upshot is, however, that you can use the suggestively meaningful names without caring about, or even knowing, the arbitrary numeric values they represent. It's a useful idiom, and you may well have reason to define enums in your own code.

There appears to be a native text type (a string) in C, but this is something of an illusion; behind the scenes, it is actually a null-terminated array of char. For example, in C you can write a string literal like this:

```
"string"
```

But in fact this is stored as 7 bytes, the numeric (ASCII) equivalents of each letter followed by a byte consisting of 0 to signal the end of the string. This data structure, called a C string, is rather tricky, and if you're lucky you'll rarely or never encounter one while programming iOS. In general, when working with strings, you'll use an Objective-C object type called `NSString`. An `NSString` is totally different from a C string; it happens, however, that Objective-C lets you write a literal `NSString` in a way that looks very like a C string:

```
@"string"
```

Notice the at-sign! This expression is actually a directive to the Objective-C compiler to form an `NSString` object. A common mistake is forgetting the at-sign, thus causing your expression to be interpreted as a C string, which is a completely different animal.

Because the notation for literal `NSStrings` is modeled on the notation for C strings, it is worth knowing something about C strings, even though you won't generally encounter them. For example, K&R lists a number of escaped characters (K&R 2.3), which you can also use in a literal `NSString`, including the following:

\n	A Unix newline character
\t	A tab character
\"	A quotation mark (escaped to show that this is not the end of the string literal)
\\	A backslash



NSStrings are natively Unicode-based, but because Objective-C is C, including non-ASCII characters in a literal `NSString` was, until quite recently, remarkably tricky, and you needed to know about such things as the `\x` and `\u` escape sequences. Now, however, it is perfectly legal to type a bullet or any other non-ASCII character directly into an `NSString` literal, and you should ignore old Internet postings (and even an occasional sentence in Apple's own documentation) warning that it is not.

K&R also mention a notation for concatenating string literals, in which multiple string literals separated only by white space are automatically concatenated and treated as a single string literal. This notation is useful for splitting a long string into multiple lines for legibility, and Objective-C copies this convention for literal `NSStrings` as well, except that you have to remember the at-sign:

```
@"This is a big long literal string "
@"which I have broken over two lines of code.;"
```

Structs

C offers few simple native data types, so how are more complex data types made? There are three ways: structures, pointers, and arrays. Both structures and pointers are going to be crucial when you're programming iOS. You're less likely to need a C array, because Objective-C has its own `NSArray` object type, but it will arise in a couple of examples later in this book.

A C structure, usually called a struct (K&R 6.1), is a compound data type: it combines multiple data types into a single type, which can be passed around as a single entity. Moreover, the elements constituting the compound entity have names and can be accessed by those names through the compound entity, using dot-notation. The iOS API has many commonly used structs, typically accompanied by convenience functions for working with them.

For example, the iOS documentation tells you that a CGPoint is defined as follows:

```
struct CGPoint {  
    CGFloat x;  
    CGFloat y;  
};  
typedef struct CGPoint CGPoint;
```

Recall that a CGFloat is basically a float, so this is a compound data type made up of two simple native data types; in effect, a CGPoint has two CGFloat parts, and their names are `x` and `y`. (The rather odd-looking last line merely asserts that one can use the term `CGPoint` instead of the more verbose `struct CGPoint`.) So we can write:

```
CGPoint myPoint;  
myPoint.x = 4.3;  
myPoint.y = 7.1;
```

Just as we can assign to `myPoint.x` in order to *set* this part of the struct, we can say `myPoint.x` to *get* this part of the struct. It's as if `myPoint.x` were the name of a variable. Moreover, an element of a struct can itself be a struct, and the dot-notation can be chained. To illustrate, first note the existence of another iOS struct, CGSize:

```
struct CGSize {  
    CGFloat width;  
    CGFloat height;  
};  
typedef struct CGSize CGSize;
```

Put a CGPoint and a CGSize together and you've got a CGRect:

```
struct CGRect {  
    CGPoint origin;  
    CGSize size;  
};  
typedef struct CGRect CGRect;
```

So suppose we've got a CGRect variable called `myRect`, already initialized. Then `myRect.origin` is a CGPoint, and `myRect.origin.x` is a CGFloat. Similarly, `myRect.size` is a CGSize, and `myRect.size.width` is a CGFloat. You could change just the `width` part of our CGRect directly, like this:

```
myRect.size.width = 8.6;
```

Instead of initializing a struct by assigning to each of its elements, you can initialize it at declaration time by assigning values for all its elements at once, in curly braces, separated by commas, like this:

```
CGPoint myPoint = { 4.3, 7.1 };
CGRect myRect = { myPoint, {10, 20} };
```

You don't actually have to be assigning to a struct-typed variable to use a struct initializer; you can use an initializer anywhere the given struct type is expected, but you might also have to cast to that struct type in order to explain to the compiler what your curly braces mean, like this:

```
CGContextFillRect(con, (CGRect){myPoint, {10, 20}});
```

In that example, `CGContextFillRect` is a function. I'll talk about functions later in this chapter, but the upshot of the example is that what comes after the first comma has to be a `CGRect`, and can therefore be a `CGRect` initializer provided this is accompanied by a `CGRect` cast.

Pointers

The other big way that C extends its range of data types is by means of pointers (K&R 5.1). A pointer is an integer (of some size or other) with a meaning: it designates the location in memory where the real data is to be found. Knowing the structure of that data and how to work with it, as well as allocating a block of memory of the required size beforehand and disposing of that block of memory when it's no longer needed, is a very complicated business. Luckily, this is exactly the sort of complicated business that Objective-C is going to take care of for us. So all you really have to know in order to use pointers is what they are and what notation is used to refer to them.

Let's start with a simple declaration. If we wanted to declare an integer in C, we could say:

```
int i;
```

That line says, "i is an integer." Now let's instead declare a *pointer* to an integer:

```
int* intPtr;
```

That line says, "intPtr is a pointer to an integer." Never mind how we know there really is going to be an integer at the address designated by this point; here, I'm concerned only with the notation. It is permitted to place the asterisk in the declaration before the name rather than after the type:

```
int *intPtr;
```

You could even put a space on both sides of the asterisk (though this is rarely done):

```
int * intPtr;
```

I prefer the first form, but I do occasionally use the second form, and Apple quite often uses it, so be sure you understand that these are all ways of saying the same thing. No matter how the spaces are inserted, the name of the type is still `int*`. If you are asked what type is `intPtr` is, the answer is `int*` (a pointer to an `int`); the asterisk is part of the name of the type of this variable. If you needed to cast a variable `p` to this type, you'd

cast like this: `(int*)p`. Once again, it is possible that you'll see code where there's a space before the asterisk, like this: `(int *)p`.

Pointers are very important in Objective-C, because Objective-C is all about objects ([Chapter 2](#)), and every variable referring to an object is itself a pointer. For example, I've already mentioned that the Objective-C string type is called `NSString`. So the way to declare an `NSString` variable is as a pointer to an `NSString`:

```
NSString* s;
```

An `NSString` literal is an `NSString` value, so we can even declare and initialize this `NSString` object, thus writing a seriously useful line of Objective-C code:

```
NSString* s = @"Hello, world!";
```

In pure C, having declared a pointer-to-integer called `IntPtr`, you are liable to speak later in your code of `*IntPtr`. This notation, outside of a declaration, means “the thing pointed to by the pointer `IntPtr`.” You speak of `*IntPtr` because you wish to access the integer at the far end of the pointer.

But in Objective-C, this is generally *not* the case. In your code, you'll be treating the pointer to an object as the object. So, for example, having declared `s` as a pointer to an `NSString`, you will *not* then proceed to speak of `*s`; rather, you will speak simply of `s`, as if it *were* the string. All the Objective-C stuff you'll want to do with an object will expect the pointer, not the object at the far end of the pointer; behind the scenes, Objective-C itself will take care of the messy business of following the pointer to its block of memory and doing whatever needs to be done in that block of memory. This fact is extremely convenient for you as a programmer, but it does cause Objective-C users to speak a little loosely; we tend to say that “`s` is an `NSString`,” when of course it is actually a pointer to an `NSString`.

You must never let this convenience lull you into forgetting the crucial fact that a pointer is a pointer. The logic of how pointers work is different from the logic of how simple data types work. The difference is particularly evident with assignment. Assignment to a simple data type changes the data value. Assignment to a pointer repoints the pointer. Suppose `ptr1` and `ptr2` are both pointers, and you say:

```
ptr1 = ptr2;
```

Now `ptr1` and `ptr2` are pointing at the same thing. Any change to the thing pointed to by `ptr1` will also change the thing pointed to by `ptr2`, because they are the same thing. Meanwhile, whatever `ptr1` was pointing to before the assignment is now not being pointed to by `ptr1`; it might, indeed, be pointed to by nothing (which could be bad). A firm understanding of these facts is crucial when working in Objective-C ([Figure 1-1](#)).

The most general type of pointer is *pointer-to-void* (`void*`), the *generic pointer*. It is legal to use a generic pointer wherever a specific type of pointer is expected. In effect, pointer-to-void casts away type checking as to what's at the far end of the pointer. Thus, the following is legal:

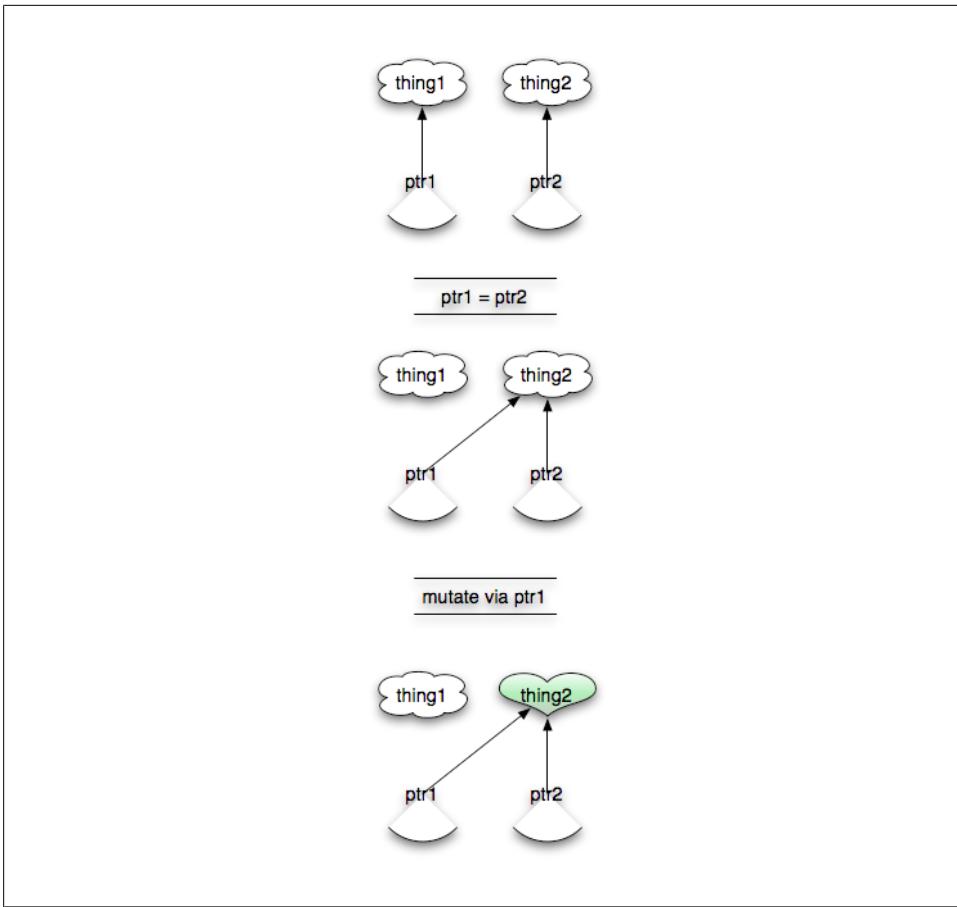


Figure 1-1. Pointers and assignment

```
int* p1; // and pretend p1 has a value
void* p2;
p2 = p1;
p1 = p2;
```

Arrays

A C array (K&R 5.3) consists of multiple elements of the same data type. An array declaration states the data type of the elements, followed by the name of the array, along with square brackets containing the number of elements:

```
int arr[3]; // means: arr is an array consisting of 3 ints
```

To refer to an element of an array, use the array's name followed by the element number in square brackets. The first element of an array is numbered 0. So we can initialize an array by assigning values to each element in turn:

```
int arr[3];
arr[0] = 123;
arr[1] = 456;
arr[2] = 789;
```

Alternatively, you can initialize an array at declaration time by assigning a list of values in curly braces, just as with a struct. In this case, the size of the array can be omitted from the declaration, because it is implicit in the initialization (K&R 4.9):

```
int arr[] = {123, 456, 789};
```

Curiously, the name of an array is the name of a pointer (to the first element of the array). Thus, for example, having declared `arr` as in the preceding examples, you can use `arr` wherever a value of type `int*` (a pointer to an int) is expected. This fact is the basis of some highly sophisticated C idioms that you almost certainly won't need to know about (which is why I don't recommend that you read any of K&R Chapter 5 beyond section 3).

C arrays rarely arise in practice when programming iOS, because you'll work mostly with the `NSArray` object type instead. But here's a case where they do. The function `CGContextStrokeLineSegments` is declared like this:

```
void CGContextStrokeLineSegments (
    CGContextRef c,
    const CGPoint points[],
    size_t count
);
```

The second parameter is an array (meaning a C array) of `CGPoint`s. That's what the square brackets tell you. So to call this function, you'd need to know at least how to make an array of `CGPoint`s. You might do it like this:

```
CGPoint arr[] = {{4,5}, {6,7}, {8,9}, {10,11}};
```

Having done that, you can pass `arr` as the second argument in a call to `CGContextStrokeLineSegments`.

Also, a C string, as I've already mentioned, is actually an array. For example, the `NSString` method `stringWithUTF8String:` takes (according to the documentation) "a NULL-terminated C array of bytes in UTF8 encoding," but the parameter is declared not as an array, but as a `char*`. Those are the same thing, and are both ways of saying that this method takes a C string.

Operators

Arithmetic operators are straightforward (K&R 2.5), but watch out for the rule that "integer division truncates any fractional part." This rule is the cause of much novice error in C. If you have two integers and you want to divide them in such a way as to get a fractional result, you must represent at least one of them as a float:

```
int i = 3;
float f = i/2; // beware! not 1.5
```

To get 1.5, you should have written `i/2.0` or `(float)i/2`.

The integer increment and decrement operators (K&R 2.8), `++` and `--`, work differently depending on whether they precede or follow their variable. The expression `++i` replaces the value of `i` by 1 more than its current value and then uses the resulting value; the expression `i++` uses the current value of `i` and then replaces it with 1 more than its current value. This is one of C's coolest features.

C also provides bitwise operators (K&R 2.9), such as bitwise-and (`&`) and bitwise-or (`|`); they operate on the individual binary bits that constitute integers. Of these, the one you are most likely to need is bitwise-or, because the Cocoa API often uses bits as switches when multiple options are to be specified simultaneously. For example, there are various ways in which a `UIView` can be resized automatically as its superview is resized, and you're supposed to provide one or more of these when setting a `UIView`'s `autoresizingMask` property. The autoresizing options are listed in the documentation as follows:

```
enum {
    UIViewAutoresizingNone          = 0,
    UIViewAutoresizingFlexibleLeftMargin = 1 << 0,
    UIViewAutoresizingFlexibleWidth   = 1 << 1,
    UIViewAutoresizingFlexibleRightMargin = 1 << 2,
    UIViewAutoresizingFlexibleTopMargin = 1 << 3,
    UIViewAutoresizingFlexibleHeight   = 1 << 4,
    UIViewAutoresizingFlexibleBottomMargin = 1 << 5
};
typedef NSUInteger UIViewAutoresizing;
```

The `<<` symbol is the left shift operator; the right operand says how many bits to shift the left operand. So pretend that an `NSUInteger` is 8 bits (it isn't, but let's keep things simple and short). Then this enum means that the following name-value pairs are defined (using binary notation for the values):

```
UIViewAutoresizingNone
00000000

UIViewAutoresizingFlexibleLeftMargin
00000001

UIViewAutoresizingFlexibleWidth
00000010

UIViewAutoresizingFlexibleRightMargin
00000100

UIViewAutoresizingFlexibleTopMargin
00001000
```

and so on. The reason for this bit-based representation is that these values can be combined into a single value (a *bitmask*) that you pass to set the `autoresizingMask`. All Cocoa has to do in order to understand your intentions is to look to see which bits in the value that you pass are set to 1. So, for example, `00001010` would mean that `UIViewAutoresizingFlexibleTopMargin` and `UIViewAutoresizingFlexibleWidth` are true (and that the others, by implication, are all false).

The question is how to form the value `00001010` in order to pass it. You could just do the math, figure out that binary `00001010` is decimal 10, and set the `autoresizingMask` property to 10, but that's not what you're supposed to do, and it's not a very good idea, because it's error-prone and makes your code incomprehensible. Instead, use the bitwise-or operator to combine the desired options:

```
myView.autoresizingMask =
    UIViewAutoresizingFlexibleTopMargin | UIViewAutoresizingFlexibleWidth;
```

This notation works because the bitwise-or operator combines its operands by setting in the result any bits that are set in either of the operands, so `00001000 | 00000010` is `00001010`, which is just the value we're trying to convey.

Simple assignment (K&R 2.10) is by the equal sign. But there are also compound assignment operators that combine assignment with some other operation. For example:

```
height *= 2; // same as saying: height = height * 2;
```

The ternary operator (`?:`) is a way of specifying one of two values depending on a condition (K&R 2.11). The scheme is as follows:

```
(condition) ? exp1 : exp2
```

If the condition is true (see the next section for what that means), the expression `exp1` is evaluated and the result is used; otherwise, the expression `exp2` is evaluated and the result is used. For example, you might use the ternary operator while performing an assignment, using this schema:

```
myVariable = (condition) ? exp1 : exp2;
```

What gets assigned to `myVariable` depends on the truth value of the condition. There's nothing happening here that couldn't be accomplished more verbosely with flow control (see the next section), but the ternary operator can greatly improve clarity, and I use it a lot.

Flow Control and Conditions

Basic flow control is fairly simple and usually involves a condition in parentheses and a block of conditionally executed code in curly braces. These curly braces constitute a new scope, into which new variables can be introduced. So, for example:

```
if (x == 7) {
    int i = 0;
    i += 1;
}
```

After the closing curly brace in the fourth line, the `i` introduced in the second line has ceased to exist, because its scope is the inside of the curly braces. If the contents of the curly braces consist of a single statement, the curly braces can be omitted, but I would advise beginners against this shorthand, as you can confuse yourself. A common beginner mistake (which will be caught by the compiler) is forgetting the parentheses around the condition. The full set of flow control statements is given in K&R Chapter 3, and I'll just summarize them schematically here ([Example 1-1](#)).

Example 1-1. The C flow control constructs

```
if (condition) {
    statements;
}

if (condition) {
    statements;
} else {
    statements;
}

if (condition) {
    statements;
} else if (condition) {
    statements;
} else {
    statements;
}

while (condition) {
    statements;
}

do {
    statements;
} while (condition);

for (before-all; condition; after-each) {
    statements;
}
```

The `if...else if...else` structure can have as many `else if` blocks as needed, and the `else` block is optional. Instead of an extended `if...else if...else if...else` structure, when the conditions would consist of comparing various values against a single value, you can use the `switch` statement; be careful, though, as it is rather confusing and can easily go wrong (see K&R 3.4 for full details). The main trick is to remember to end every case with a `break` statement, unless you want it to “fall through” to the next case ([Example 1-2](#)).

Example 1-2. A switch statement

```
NSString* key;
switch (tag) {
    case 1: { // i.e. if tag == 1
        key = @"lesson";
        break;
    }
    case 2: { // i.e. if tag == 2
        key = @"lessonSection";
        break;
    }
    case 3: { // i.e. if tag == 3
        key = @"lessonSectionPartFirstWord";
        break;
    }
}
```

The C for loop needs some elaboration for beginners ([Example 1-1](#)). The *before-all* statement is executed once as the for loop is first encountered and is usually used for initialization of the counter. The condition is then tested, and if true, the block is executed; the condition is usually used to test whether the counter has reached its limit. The *after-each* statement is then executed, and is usually used to increment or decrement the counter; the condition is then immediately tested again. Thus, to execute a block using integer values 1, 2, 3, 4, and 5 for *i*, the notation is:

```
int i;
for (i = 1; i < 6; i++) {
    // ... statements ...
}
```

The need for a counter intended to exist solely within the for loop is so common that C99 permits the declaration of the counter as part of the *before-all* statement; the declared variable's scope is then inside the curly braces:

```
for (int i = 1; i < 6; i++) {
    // ... statements ...
}
```

The for loop is one of the few areas in which Objective-C extends C's flow-control syntax. Certain Objective-C objects represent enumerable collections of other objects; "enumerable" basically means that you can cycle through the collection, and cycling through a collection is called *enumerating* the collection. To make enumerating easy, Objective-C provides a **for...in** operator, which works like a for loop:

```
SomeType* oneItem;
for (oneItem in myCollection) {
    // ... statements ....
}
```

On each pass through the loop, the variable **oneItem** (or whatever you call it) takes on the next value from within the collection. As with the C99 for loop, **oneItem** can be declared in the **for** statement, limiting its scope to the curly braces:

```
for (SomeType* oneItem in myCollection) {  
    // ... statements ....  
}
```

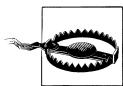
To abort a loop from inside the curly braces, use the `break` statement. To abort the current iteration from within the curly braces and proceed to the next iteration, use the `continue` statement. In the case of `while` and `do`, `continue` means to perform immediately the conditional test; in the case of a for loop, `continue` means to perform immediately the *after-each* statement and then the conditional test.

C also has a `goto` statement that allows you to jump to a named (labeled) line in your code (K&R 3.8); even though `goto` is notoriously “considered harmful,” there are situations in which it is pretty much necessary, especially because C’s flow control is otherwise so primitive.



It is permissible for a C statement to be compounded of multiple statements, separated by commas, to be executed sequentially. The last of the multiple statements is the value of the compound statement as a whole. This construct, for instance, lets you perform some secondary action before each test of a condition or perform more than one *after-each* action (an example appears in [Chapter 17](#)).

We can now turn to the question of what a condition consists of. C has no separate boolean type; a condition either evaluates to 0, in which case it is considered false, or it doesn’t, in which case it is true. Comparisons are performed using the equality and relational operators (K&R 2.6); for example, `==` compares for equality, and `<` compares for whether the first operand is less than the second. Logical expressions can be combined using the logical-and operator (`&&`) and the logical-or operator (`||`); using these along with parentheses and the not operator (`!`) you can form complex conditions. Evaluation of logical-and and logical-or expressions is short-circuited, meaning that if the left condition settles the question, the right condition is never even evaluated.



Don’t confuse the logical-and operator (`&&`) and the logical-or operator (`||`) with the bitwise-and operator (`&`) and the bitwise-or operator (`|`) discussed earlier. Writing `&` when you mean `&&` (or *vice versa*) can result in surprising behavior.

The operator for testing basic equality, `==`, is not a simple equal sign; forgetting the difference is a common novice mistake. The problem is that such code is legal: simple assignment, which is what the equal sign means, has a value, and any value is legal in a condition. So consider this piece of (nonsense) code:

```
int i = 0;  
while (i = 1) {  
    i = 0;  
}
```

You might think that the while condition tests whether `i` is 1. You might then think: `i` is 0, so the while body will never be performed. Right? Wrong. The while condition does not test whether `i` is 1; it assigns 1 to `i`. The value of that assignment is also 1, so the condition evaluates to 1, which means true. So the while body *is* performed. Moreover, even though the while body assigns 0 to `i`, the condition is then evaluated again and assigns 1 to `i` a second time, which means true yet again. And so on, forever; we've written an endless loop, and the program will hang. (And, depending on what compiler and settings you're using, you might not even get a warning of trouble ahead.)

C programmers actually revel in the fact that testing for zero and testing for false are the same thing and use it to create compact conditional expressions, which are considered elegant and idiomatic. I don't recommend that you make use of such idioms, as they can be confusing, but I must admit that even I do occasionally resort to this sort of thing:

```
NSString* s = nil;  
// ...  
if (s) {  
    // ...  
}
```

The idea of that code is to test whether the `NSString` object `s`, between the time it was declared and the start of the if-block, has been set to an actual string. Because `nil` is a form of 0, the condition is asking whether `s` is non-nil. Some Objective-C programmers would take me to task for this style of writing code; if I want to test whether `s` is `nil`, they would say, I should test it explicitly:

```
if (s == nil)
```

In fact, some would say, it is even better to write the terms of the comparison in the opposite order:

```
if (nil == s)
```

Why? Because if I were to omit accidentally the second equal sign, thus turning the equality comparison into an assignment, the first expression would compile (and misbehave, because I am now assigning `nil` to `s`), but the second expression would certainly be caught by the compiler as an error, because assigning a value to `nil` is illegal.

Objective-C introduces a `BOOL` type, which you should use if you need to capture or maintain a condition's value as a variable, along with constants `YES` and `NO` (actually representing 1 and 0), which you should use when setting a boolean value. Don't compare anything against a `BOOL`, not even `YES` or `NO`, because a value like 2 is true in a condition but is not equal to `YES` or `NO`. Just use the `BOOL` directly as a condition, or as part of a complex condition, and all will be well. For example:

```
BOOL snil = (nil == s);  
// ...  
if (snil) // ... not: if (snil == YES)
```

Functions

C is a function-based language (K&R 4.1). A *function* is a block of code defining what should happen; when other code *calls* (invokes) that function, the function's code does happen. A function returns a value, which is substituted for the call to that function.

Here's a definition of a function that accepts an integer and returns its square:

```
int square(int i) {  
    return i * i;  
}
```

Now I'll call that function:

```
int i = square(3);
```

Because of the way `square` is defined, that is exactly like saying:

```
int i = 9;
```

That example is extremely simple, but it illustrates many key aspects of functions.

Let's analyze how a function is defined:

```
int❶ square❷(❸int i) {❹  
    return i * i;  
}
```

- ❶ We start with the type of value that the function returns; here, it returns an int.
- ❷ Then we have the name of the function, which is `square`.
- ❸ Then we have parentheses, and here we place the data type and name of any values that this function expects to receive. Here, `square` expects to receive one value, an int, which we are calling `i`. The name `i` (along with its expected data type) is a *parameter*; when the function is called, its value will be supplied as an *argument*. If a function expects to receive more than one value, multiple parameters in its definition are separated by a comma (and when the function is called, the arguments supplied are likewise separated by a comma).
- ❹ Finally, we have curly braces containing the statements that are to be executed when the function is called.

Those curly braces constitute a scope; variables declared within them are local to the function. The names used for the parameters in the function definition are also local to the function; in other words, the `i` in the first line of the function definition is the same as the `i` in the second line of the function definition, but it has nothing to do with any `i` used outside the function definition (as when the result of the function call is assigned to a variable called `i`). The value of the `i` parameter in the function definition is assigned from the corresponding argument when the function is actually called; in the previous example, it is 3, which is why the function result is 9. Supplying a function call with arguments is thus a form of assignment. Suppose a function is defined like this:

```
int myfunction(int i, int j) { // ...
```

And suppose we call that function:

```
int result = myfunction(3, 4);
```

That function call effectively assigns 3 to the function's *i* parameter and 4 to the function's *j* parameter.

When a `return` statement is encountered, the value accompanying it is handed back as the result of the function call, and the function terminates. It is legal for a function to return no value; in such a case, the `return` statement has no accompanying value, and the definition states the type of value returned by the function as `void`. It is also legal to call a function and ignore its return value even if it has one. For example, we could say:

```
square(3);
```

That would be a somewhat silly thing to say, because we have gone to all the trouble of calling the function and having it generate the square of 3 — namely 9 — but we have done nothing to *capture* that 9. It is exactly as if we had said:

```
9;
```

You're allowed to say that, but it doesn't seem to serve much purpose. On the other hand, the point of a function might be not so much the value it returns as other things it does as it is executing, so then it might make perfect sense to ignore its result.

The parentheses in a function's syntax are crucial. Parentheses are how C knows there's a function. Parentheses after the function name in the function *definition* are how C knows this is a function definition, and they are needed even if this function takes no parameters. Parentheses after the function name in the function *call* are how C knows this is a function call, and they are needed even if this function call supplies no arguments. Using the bare name of a function is possible, because the name is effectively a kind of variable (and I'll talk later about why you might want to do that), but it doesn't call the function.

Let's return to the simple C function definition and call that I used as my example earlier. Suppose we combine that function definition and the call to that function into a single program:

```
int square(int i) {  
    return i * i;  
}  
int i = square(3);
```

That is a legal program, but only because the definition of the `square` function precedes the call to that function. If we wanted to place the definition of the `square` function elsewhere, such as after the call to it, we would need at least to precede the call with a declaration of the `square` function ([Example 1-3](#)). The declaration looks just like the first line of the definition, but it is a statement, ending with a semicolon, rather than a left curly brace.

Example 1-3. Declaring, calling, and defining a function

```
int square(int i);
int i = square(3);
int square(int i) {
    return i * i;
}
```

The parameter names in the declaration do not have to match the parameter names in the definition, but all the types (and, of course, the name of the function) must match. The types constitute the *signature* of this function. In other words, it does not matter if the first line, the declaration, is rewritten thus:

```
int square(int j);
```

What does matter is that, both in the declaration and in the definition, `square` is a function taking one int parameter and returning an int.

In Objective-C, when you’re sending a message to an object ([Chapter 2](#)), you won’t use a function call; you’ll use a method call ([Chapter 3](#)). But you will most definitely use plenty of C function calls as well. For example, earlier we initialized a `CGPoint` by setting its `x` element and its `y` element and by assigning its elements values in curly braces. But what you’ll usually do to make a new `CGPoint` is to call `CGPointMake`, which is declared like this:

```
CGPoint CGPointMake (
    CGFloat x,
    CGFloat y
);
```

Despite its multiple lines and its indentations, this is indeed a C function declaration, just like the declaration for our simple `square` function. It says that `CGPointMake` is a C function that takes two `CGFloat` parameters and returns a `CGPoint`. So now you know (I hope) that it would be legal (and typical) to write this sort of thing:

```
CGPoint myPoint = CGPointMake(4.3, 7.1);
```

Pointer Parameters and the Address Operator

I’ve mentioned several times that your variables referring to Objective-C objects are going to be pointers:

```
NSString* s = @"Hello, world!";
```

Although it is common to speak loosely of `s` as an `NSString` (or just as a string), it is actually an `NSString*` — a pointer to an `NSString`. Therefore, when a C function or an Objective-C method expects an `NSString*` parameter, there’s no problem, because that’s exactly what you’ve got. For example, one way to concatenate two `NSStrings` is to call the `NSString` method `stringByAppendingString`: (that’s not a misprint; the colon is part of the name), which the documentation tells you is declared as follows:

```
- (NSString *)stringByAppendingString:(NSString *)aString
```

The space between the class name and the asterisk is optional, so this declaration is telling you (after you allow for the Objective-C syntax) that this method expects one `NSString*` parameter and returns an `NSString*`. That's splendid because those kinds of pointers are just what you've got and just what you want. So this code would be legal:

```
NSString* s1 = @"Hello, ";
NSString* s2 = @"World!";
NSString* s3 = [s1 stringByAppendingString: s2];
```

The idea, then, is that although Objective-C is chock-a-block with pointers and asterisks, they don't make things more complicated, as long as you remember that they *are* pointers.

Sometimes, however, a function expects as a parameter a pointer to something, but what you've got is not a pointer but the thing itself. Thus, you need a way to create a pointer to that thing. The solution is the address operator (K&R 5.1), which is an ampersand before the name of the thing.

For example, there's an `NSString` method for reading from a file into an `NSString`, which is declared like this:

```
+ (id)stringWithContentsOfFile:(NSString *)path
                           encoding:(NSStringEncoding)enc
                             error:(NSError **)error
```

Now, never mind what an `id` is, and don't worry about the Objective-C method declaration syntax. Just consider the types of the parameters. The first one is an `NSString*`; that's no problem, as every reference to an `NSString` is actually a pointer to an `NSString`. An `NSStringEncoding` turns out to be merely an alias to a primitive data type, an `NSUInteger`, so that's no problem either. But what on earth is an `NSError**`?

By all logic, it looks like an `NSError**` should be a pointer to a pointer to an `NSError`. And that's exactly what it is. This method is asking to be passed a pointer to a pointer to an `NSError`. Well, it's easy to declare a pointer to an `NSError`:

```
NSError* myError;
```

But how can we obtain a pointer to that? With the address operator! So our code might look, schematically, like this:

```
NSString* myPath = // something or other;
NSStringEncoding myEnc = // something or other;
NSError* myError = nil;
NSString* result = [NSString stringWithContentsOfFile: myPath
                                             encoding: myEnc
                                               error: &myError];
```

The important thing to notice is the ampersand. Because `myError` is a pointer to an `NSError`, `&myError` is a pointer to a pointer to an `NSError`, which is just what we're expected to provide. Thus, everything goes swimmingly.

This device lets Cocoa effectively return *two* results from this method call. It returns a real result, which we have captured by assigning it to the `NSString` pointer we're calling

`result`. But if there's an error, it also wants to set the value of another object, an `NSError` object; the idea is that you can then study that `NSError` object to find out what went wrong. (Perhaps the file wasn't where you said it was, or it wasn't stored in the encoding you claimed it was.) By passing a pointer to a pointer to an `NSError`, you give the method free rein to do that. Before the call to `stringWithContentsOfFile:`, `myError` was uninitialized; during the call to `stringWithContentsOfFile:`, Cocoa can, if it likes, repoint the pointer, thus giving `myError` an actual value.

So the idea is that you first check `result` to see whether it's nil. If it isn't, fine; it's the string you asked for. If it is, you then study the `NSError` that `myError` is now pointing to, to learn what went wrong. This pattern is frequently used in Cocoa.

You can use the address operator to create a pointer to any named variable. A C function is technically a kind of named variable, so you can even create a pointer to a function! This is an example of when you'd use the name of the function without the parentheses: you aren't calling the function, you're talking about it. For example, `&square` is a pointer to the `square` function. In [Chapter 9](#), I describe a situation in which this is a useful thing to do.

Another operator used in connection with pointers, or when memory must be allocated dynamically, is `sizeof`. It may be followed by a type name in parentheses or by a variable name; a variable name needn't be in parentheses, but it *can* be, so most programmers ignore the distinction and use parentheses routinely, as if `sizeof` were a function.

For example, the documentation shows the declaration for `AudioSessionSetProperty` like this:

```
OSStatus AudioSessionSetProperty (
    AudioSessionPropertyID    inID,
    UInt32                  inDataSize,
    const void*              *inData
);
```

Never mind what an `AudioSessionPropertyID` is; it's merely a value that you obtain and pass on. `UInt32` is one of those derived numeric types I mentioned earlier. The discussion has already dealt with pointer-to-void and how to derive a pointer using the address operator. But look at the *name* of the second parameter; the function is asking for the size of the thing pointed to by the third parameter. Here's an actual call to this function (from [Chapter 27](#)):

```
UInt32 ambi = kAudioSessionCategory_AmbientSound;
AudioSessionSetProperty(kAudioSessionProperty_AudioCategory, sizeof(ambi), &ambi);
```

Files

The little dance of declaring a function before calling it ([Example 1-3](#)) may seem rather absurd, but it is of tremendous importance in the C language, because it is what allows a C program to be arbitrarily large and complex.

As your program grows, you can divide and organize it into multiple files. This kind of organization can make a large program much more maintainable — easier to read, easier to understand, easier to change without accidentally breaking things. A large C program therefore usually consists of two kinds of file: code files, whose filename extension is `.c`, and header files, whose filename extension is `.h`. The build system will automatically “see” all the files and will know that together they constitute a single program, but there is also a rule in C that code inside one file cannot “see” another file unless it is explicitly told to do so. Thus, a file itself constitutes a scope; this is a deliberate and valuable feature of C, because it helps you keep things nicely pigeonholed.

The way you tell a C file to “see” another file is with the `#include` directive. The hash sign in the term `#include` is a signal that this line is an instruction to the preprocessor. In this case, the word `#include` is followed by the name of another file, and the directive means that the preprocessor should simply replace the directive by the entire contents of the file that’s named.

So the strategy for constructing a large C program is something like this:

- In each `.c` file, put the code that only this file needs to know about; typically, each file’s code consists of related functionality.
- In each `.h` file, put the function declarations that multiple `.c` files might need to know about.
- Have each `.c` file include those `.h` files containing the declarations it needs to know about.

So, for example, if `function1` is defined in `file1.c`, but `file2.c` might need to call `function1`, the declaration for `function1` can go in `file1.h`. Now `file1.c` can include `file1.h`, so all of its functions, regardless of order, can call `function1`, and `file2.c` can also include `file1.h`, so all of its functions can call `function1` ([Figure 1-2](#)). In short, header files are a way of letting code files share knowledge about one another without actually sharing code (because, if they did share code, that would violate the entire point of keeping the code in separate files).

But how does the compiler know where, among all these multiple `.c` files, to begin execution? Every real C program contains, somewhere, exactly one function called `main`, and this is always the entry point for the program as a whole: the compiler sets things up so that when the program executes, `main` is called.

The organization for large C programs that I’ve just described will also be, in effect, the organization for your iOS programs. (The chief difference will be that instead of `.c` files, you’ll use `.m` files, because `.m` is the conventional filename extension for telling Xcode that your files are written in Objective-C, not pure C.) Moreover, if you look at any iOS Xcode project, you’ll discover that it contains a file called `main.m`; and if you look at that file, you’ll find that it contains a function called `main`. That’s the entry point to your application’s code when it runs.

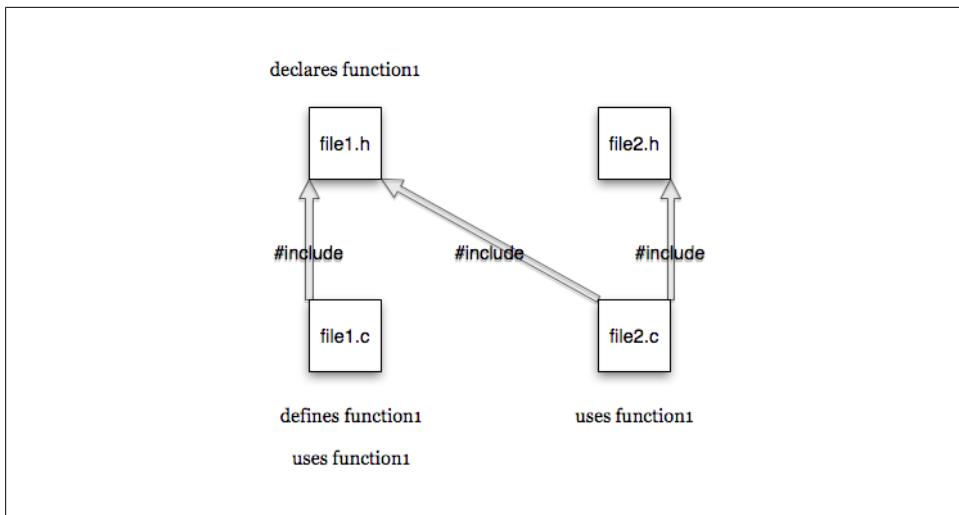


Figure 1-2. How a large C program is divided into files

The big difference between your Objective-C code files and the C code files I've been discussing is that instead of saying `#include`, your files will say `#import`. The `#import` preprocessor directive is not mentioned in K&R. It's an Objective-C addition to the language. It's based on `#include`, but it is used instead of `#include` because it (`#import`) contains some logic for making sure that the same material is not included more than once. Such repeated inclusion is a danger whenever there are many cross-dependent header files; use of `#import` solves the problem neatly.

Furthermore, your iOS programs consist not only of *your* code files and their corresponding `.h` files, but also of Apple's code files and *their* corresponding `.h` files. The difference is that Apple's code files (which are what constitutes Cocoa) have already been compiled. But your code must still `#import` Apple's `.h` files so as to be able to see Apple's declarations. If you look at an iOS Xcode project, you'll find that any `.m` files it contains by default, as well as its `main.m` file, contain a line of this form:

```
#import <UIKit/UIKit.h>
```

That line is essentially a single massive `#import` that copies into your program the declarations for the entire basic iOS API. Moreover, each of your `.m` files `#imports` its corresponding `.h` file, including whatever the `.h` file `#imports`. Thus, all your code files include the basic iOS declarations.

For example, earlier I said that `CGPoint` was defined like this:

```

struct CGPoint {
    CGFloat x;
    CGFloat y;
};
typedef struct CGPoint CGPoint;

```

After the preprocessor operates on all your files, your `.m` files actually *contain* that definition of `CGPoint`. (You can even choose Product → Generate Output → Generate Preprocessed File, as I mentioned earlier, to confirm that this is true.) And that is why your code is able to use a `CGPoint`!

The `#import` directive, like the `#include` directive (K&R 4.11), can specify a file in angle brackets or in quotation marks:

```
#import <UIKit/UIKit.h>
#import "MyHeader.h"
```

The quotation marks form means “look for the named file in the same folder as this file” (the `.m` file in which the `#import` line occurs). The angle brackets form means to look among the various header search paths supplied in the build settings; these search paths are set for you automatically, and you normally won’t need to modify them. In general, the distinction means that you’ll use angle brackets to refer to a header file owned by the Cocoa API and quotation marks to refer to a header file that you wrote. If you’re curious as to what an `#import` directive imports, select it (in Xcode) and choose File → Open Quickly to display the contents of the designated header file.

The Standard Library

You also have at your disposal a large collection of built-in C library files. A library file is a centrally located collection of C functions, along with a `.h` file that you can include in order to make those functions available to your code.

For example, suppose you want to round a float up to the next highest integer. The way to do this is to call some variety of the `ceil` function. You can read the `ceil` man page by typing `man ceil` in the Terminal. The documentation tells you what `#include` to use to incorporate the correct header and also shows you the function declarations and tells you what those functions do. A small pure C program might thus look like this:

```
#include <math.h>
float f = 4.5;
int i = ceil(f); // now i is 5
```

In your iOS programs, `math.h` is included for you as part of the massive UIKit `#import`, so there’s no need to include it again. But some library functions might require an explicit `#import`.

The standard library is discussed in K&R Appendix B. But the modern standard library has evolved since K&R; it is a superset of K&R’s library. The `ceil` function, for example, is listed in K&R appendix B, but the `ceilf` function is not. Similarly, if you wanted to generate a random number (which is likely if you’re writing a game program that needs to incorporate some unpredictable behavior), you probably wouldn’t use the `rand` function listed in K&R; you’d use the `random` function, which supersedes it.

Forgetting that Objective-C is C and that the C library functions are available to your code is a common beginner mistake.

More Preprocessor Directives

Of the many other available preprocessor directives, the one you'll use most often is `#define`. It is followed by a name and a value; at preprocess time, the value is substituted for the name down through this code file. As K&R very well explain (K&R 1.4), this is a good way to prevent “magic numbers” from being hidden and hard-coded into your program in a way that makes the program difficult to understand and maintain.

For example, in an iOS app that lays out some text fields vertically, I might want them all to have the same space between them. Let's say this space is 3.0. I shouldn't write 3.0 repeatedly throughout my code as I calculate the layout; instead, I write:

```
#define MIDSPACE 3.0
```

Now instead of the “magic number” 3.0, my code uses a meaningful name, `MIDSPACE`; at preprocessor time, the text `MIDSPACE` is replaced with the text `3.0`. So it amounts to the same thing, but if I decide to change this value and try a different one, all I have to do change is the `#define` line, not every occurrence of the number 3.0.

A `#define` simply performs text substitution, so any expression can be used as the value. Sometimes you'll want that expression to be an `NSString` literal. In Cocoa, `NSString` literals can be used as a key to a dictionary or the name of a notification. (Never mind for now what a dictionary or a notification is.) This situation is an invitation to error. If you have a dictionary containing a key `@“mykey”` and you mistype this elsewhere in your code as `@“myKey”` or `@“mikey”`, the compiler won't complain, but your program will misbehave. The solution is to define a name for this literal string:

```
#define MYKEY @“mykey”
```

Now use `MYKEY` throughout your code instead of `@“mykey”`, and if you mistype `MYKEY` the preprocess substitution won't be performed and the compiler *will* complain, catching the mistake for you.

The `#define` directive can also be used to create a macro (K&R 4.11.2), a more elaborate form of text substitution. You'll encounter a few Cocoa macros in the course of this book, but they will appear indistinguishable from functions; their secret identity as macros won't concern you.

The `#warning` directive deliberately triggers a warning in Xcode at compile time; this can be a way to remind yourself of some impending task or requirement:

```
#warning Don't forget to fix this bit of code
```

There is also a `#pragma mark` directive that's useful with Xcode; I talk about it when discussing the Xcode programming environment ([Chapter 9](#)).

Data Type Qualifiers

A variable's data type can be declared with a qualifier before the name of the type, modifying something about how that variable is to be used. For example, the declaration can be preceded by the term `const`, which means (K&R 2.4) that it is illegal to change the variable's value; the variable must be initialized in the same line as the declaration, and that's the only value it can ever have.

You can use a `const` variable as an alternative way (instead of `#define`) to prevent “magic numbers” and similar expressions. For example:

```
const NSString* MYKEY = @"Howdy";
```

The Cocoa API itself makes heavy use of this device. For example, in some circumstances Cocoa will pass a dictionary of information to your code. The documentation tells you what keys this dictionary contains. But instead of telling you a key as a string, the documentation tells you the key as a `const NSString` variable name:

```
UIKIT_EXTERN NSString *const UIApplicationStatusBarOrientationUserInfoKey;
```

(Never mind what `UIKIT_EXTERN` means.) This declaration tells you that `UIApplicationStatusBarOrientationUserInfoKey` is the name of an `NSString`, and you are to trust that its value is set for you. You are to go ahead and use this name whenever you want to speak of this particular key, secure in the knowledge that the actual key name string will be substituted. You do not have to know what that actual key name string is. In this way, if you make a mistake in typing the variable name, the compiler will catch the mistake because you'll be using the name of an undefined variable.

Another commonly used qualifier is `static`. This term is unfortunately used in two rather different ways in C; the way I commonly use it is inside a function. Inside a function, `static` indicates that the memory set aside for a variable should not be released after the function returns; rather, the variable remains and maintains its value for the next time the function is called. A static variable is useful, for example, when you want to call a function many times without the overhead of calculating the result each time (after the first time). First test to see whether the static value has already been calculated: if it hasn't, this must be the first time the function is being called, so you calculate it; if it has, you just return it. Here's a schematic version:

```
int myfunction() {
    static int result = 0; // 0 means we haven't done the calculation yet
    if (result == 0) {
        // calculate result and set it
    }
    return result;
}
```

A very common use of a static variable in Objective-C is to implement a singleton instance returned by a class factory method. If that sounds complicated, don't worry; it isn't. Here's an example from my own code, which you can grasp even though we haven't discussed Objective-C yet:

```
+ (CardPainter*) sharedPainter {
    static CardPainter* sp = nil;
    if (nil == sp)
        sp = [[CardPainter alloc] init];
    return sp;
}
```

That code says: If the CardPainter instance `sp` has never been created, create it, and in any case, now return it. Thus, no matter how many times this method is called, the instance will be created just once and that same instance will be returned every time.

Object-Based Programming

My object all sublime.

—W. S. Gilbert, *The Mikado*

Objective-C, the native language for programming the Cocoa API, is an object-oriented language; in order to use it, the programmer must have an appreciation of the nature of objects and object-based programming. There's little point in learning the syntax of Objective-C message sending or instantiation without a clear understanding of what a message or an instance is. That is what this chapter is about.

Objects

An object, in programming, is based on the concept of an object in the real world. It's an independent, self-contained thing. These objects, unlike purely inert objects in the real world, have abilities. So an object in programming is more like a clock than a rock; it doesn't just sit there, but actually does something. Perhaps one could compare an object in programming more to the animate objects of the real world, as opposed to the inanimate objects, except that — unlike real-world animate things — a programming object is supposed to be predictable: in particular, it does what you tell it. In the real world, you tell a dog to sit and anything can happen; in the programming world, you tell a dog to sit and it sits. (This is why so many of us prefer programming to dealing with the real world.)

In object-based programming, a program is organized into many discrete objects. This organization can make life much easier for the programmer. Each object has abilities that are specialized for that object. You can think of this as being a little like how an automobile assembly line works. Each worker or station along the line does one thing (screw on the bumpers, or paint the door, or whatever) and does it well. You can see immediately how this organization helps the programmer. If the car is coming off the assembly line with the door badly painted, it is very likely that the blame lies with the door-painting object, so we know where to look for the bug in our code. Or, if we decide

to change the color that the door is to be painted, we have but to make a small change in the door-painting object. Meanwhile, other objects just go on doing what they do. They neither know nor care what the door-painting object does or how it works.

Objects, then, are an organizational tool, a set of boxes for encapsulating the code that accomplishes a particular task. They are also a conceptual tool. The programmer, being forced to think in terms of discrete objects, must divide the goals and behaviors of the program into discrete tasks, each task being assigned to an appropriate object. Of course, objects can cooperate with one another, and the ways in which this cooperation can be arranged are innumerable. The assembly-line analogy illustrates one such arrangement — first, object 1 operates upon the end-product; then it hands it off to object 2, and object 2 operates upon the end-product, and so on — but that arrangement won't be appropriate to most tasks. Coming up with an appropriate arrangement — an *architecture* — for the cooperative and orderly relationship between objects is one of the most challenging aspects of object-based programming.

Messages and Methods

Nothing in a computer program happens unless it is instructed to happen. In a C program, all code belongs to a function and doesn't run unless that function is called. In an object-based program, all code belongs to an object, and doesn't run unless that object is told to run that code. All the action in an object-based program happens because an object was told to act. What does it mean to tell an object something?

An object, in object-based programming, has a well-defined set of abilities — things it knows how to do. For example, imagine an object that is to represent a dog. We can design a highly simplified, schematic dog that knows how to do an extremely limited range of things: eat, come for a walk, bark, sit, lie down, sleep. The purpose of these abilities is so that the object can be told, as appropriate, to exercise them. So, again, we can imagine our schematic dog, rather like some child's toy robot, responding to simple commands: Eat! Come for a walk! Bark!

In object-based programming, a command directed to an object is called a *message*. To make the dog object eat, we send the eat message to the dog object. This mechanism of message sending is the basis of all activity in the program. The program consists entirely of objects, so its activity consists entirely of objects sending messages to one another.

For objects to send messages to one another, objects must know about one another in some appropriate way at some appropriate time. Ensuring such mutual knowledge is part of the architectural design process I spoke of earlier. Returning for a moment to the assembly-line architecture, it's no use saying that object 1 operates on the end-product and then object 2 operates on the end-product; that isn't going to happen all by itself. It has to be arranged somehow. We can imagine various architectures for arranging it. Perhaps we will set things up so that object 1 knows about object 2, and

as the last step in its own operation, sends a message to object 2, handing it the end-product and telling it to commence its own operation. Or perhaps we will have a conveyor-belt object, which will hand the end-product to object 1 and tell it to commence its operation, wait until object 1 finishes with it, and then hand the end-product to object 2 and tell *it* to commence its operation. Each of these is a perfectly reasonable architectural pattern, and many others are possible; it is the programmer's job to implement an architecture that not only makes the program work appropriately, but also makes the program itself clear and easy for the programmer to work on. But the problem of making sure that within that architecture, each object knows about — technically, has a *reference* to — any other object to which it might need to send a message can be quite tricky (so much so, indeed, that an entire chapter of this book, [Chapter 13](#), is devoted to it).

A moment ago, I said that in a C program, all code belongs to a function. The object-based analogue to a function is called a *method*. So, for example, a dog object might have an `eat` method. When the dog object is sent the `eat` message, it responds by calling the `eat` method.

It may sound as if I'm not drawing any clear distinction between a message and a method. But there is a difference. A message is what one object says to another. A method is a bundle of code that gets called. The connection between the two is not perfectly direct. You might send a message to an object that corresponds to no method of that object. For example, you might tell the dog to recite the soliloquy from Hamlet. I'm not sure what will happen if you do that; the details are implementation-dependent. (The dog might just sit there silently. Or it might get annoyed and bite you. Or, I suppose, it might nip off, read Hamlet, memorize the soliloquy, and recite it.) But that implementation-dependence is exactly the point of the distinction between message and method.

Nevertheless, in general the distinction between sending a message and calling a method won't usually be important in real life. Most of the time, when you're using Objective-C, your reason for sending a message to an object will be that that object implements the corresponding method and you are expecting to call that method. So sending a message to an object and calling a method of an object will appear to be the same act.

Classes and Instances

We come now to an extremely characteristic and profound feature of object-based programming. Just like in the real world, every object in the object-based programming world is of some type. This type, called a *class*, is the object-based analogy to the data type in C. Just as a simple variable in C might be an `int` or a `float`, an object in the object-based programming world might be a `Dog` (or an `NSString`). In the object-based programming world, the idea of this arrangement is to ensure that more than one individual object can be relied upon to act the same way.

There can, for example, be more than one dog. You might have a dog called Fido and I might have a dog called Rover. But both dogs know how to eat, come for a walk, and bark. In object-based programming, they know that because they both belong to the Dog class. The knowledge of how to eat, come for a walk, and bark is part of the Dog class. Your dog Fido and my dog Rover possess this knowledge solely by virtue of being Dog objects.

From the programmer's point of view, what this means is simple: all the code you write is put into a class. All the methods you write will be part of some class or other. You don't program an individual dog object: you program the Dog class.

But I just got through saying that an object-based program works through the sending of messages to individual objects. So even though the programmer does not write the code for an individual dog object, there still needs to *be* an individual dog object in order for there to be something to send a message to. It is the Dog class that knows how to bark, but it is an individual dog object that is told to bark, and that actually does bark. So the question is: if all Dog code lives in a Dog class, where do individual dogs come from?

The answer is that they have to be created in the course of the program as it runs. When the program starts out, it contains code for a Dog class, but no individual dog objects. If any barking by any dogs is to be done, the program must first create an individual dog object. This object will belong to the Dog class, so it can be sent the `bark` message. An individual object belonging to the Dog class (or any class) is an *instance* of that class. To manufacture, from a class, an actual individual object that is an instance of that class, is to *instantiate* that class.

So every individual object, such as I talked about in the preceding sections — every individual object, that is, to which a message can be sent — is an instance of some class. Classes exist from the get-go, as part of the fact that the program exists in the first place; they are where the code is. Instances are manufactured, deliberately and individually, as the program runs. Each instance is manufactured *from* a class, it is an instance of that class, and it has methods by virtue of the fact that the class has those methods. The instance can then be sent a message; what it will do in response depends on what code the class contains in its methods. The instance is the individual thing that can be sent messages; the class, with its methods, is the locus of the thing's ability to respond to messages (Figure 2-1).

This relationship between instance and class begins to sound rather ethereal or metaphysical. Instances and classes seem to be programming-language analogies to what a philosopher would call particulars and universals. Indeed, the whole setup reminds one of nothing so much as Plato's theory of Forms. For Plato, this world of ours is the world of individual things, but those things derive their natures by virtue of archetypal Forms that live off in another world. I'm not the only person ever to make this comparison to Platonic Forms — it is, indeed, implicit in the design of object-based languages and has been evoked explicitly in discussions of such languages ever since

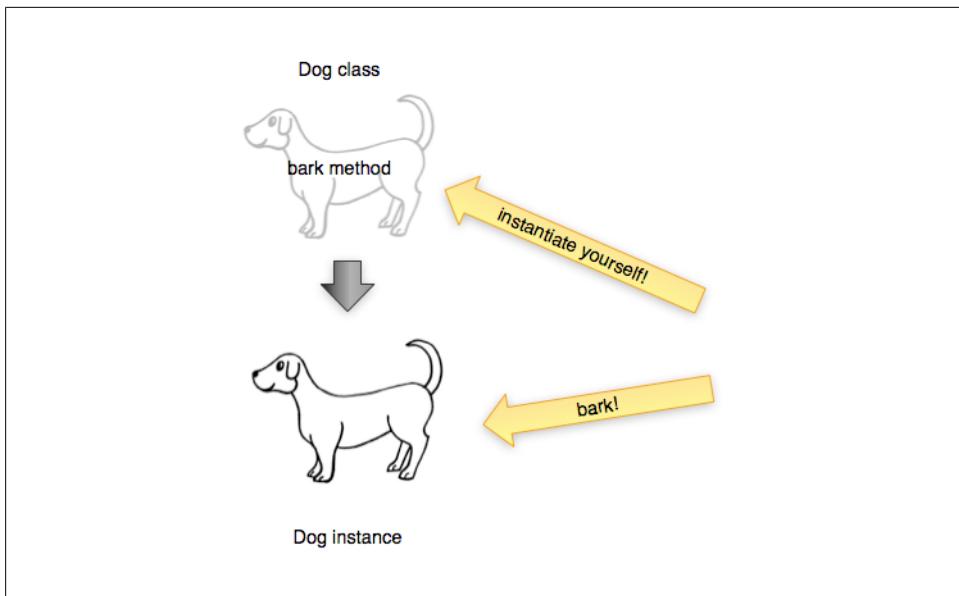


Figure 2-1. Class and instance

Smalltalk. But the comparison is still an apt one. As I said many years ago in my book [REALbasic: The Definitive Guide](#):

Indeed, object-oriented programming seems to fulfill Plato's philosophical program announced in the *Euthyphro* (6e, my translation):

SOCRATES. Now, you recall that I asked you to explain to me, not this or that particular pious thing, but that Form Itself through which all pious things are pious? You did say, I believe, that it was through one Form that impious things are impious and pious things are pious; don't you remember?

EUTHYPHRO. Yes, I do.

SOCRATES. All right, then; so, explain to me what is this Form Itself, so that by keeping my eyes upon it and using it as a model, I may declare that whatever you or anyone else does that is of this sort, is pious, and that whatever is not, is not.

The problems with Plato's characterization are well known: the Form seems to be a "thing" separate from the particular things of the world around us, the notion "through" is crucial but slippery, and Plato seems to equivocate rather glibly between the Form's being responsible for a thing's being such and such and our ability to know that a thing is such and such; thus, his program is almost certainly doomed to failure as an explanation of how the world works. But he is perfectly accurate about how an object-oriented program works! If an instance is of the Pious type, there really is a separate Pious class that really is responsible for the instance being such as it is.

Because every individual object is an instance of a class, to know what messages you can officially send to that object, you need to know at least what methods its class has

endowed it with. The public knowledge of this information is that class's API. (A class may also have methods that you're not really supposed to call from outside that object; these would not be public and other objects couldn't officially send those messages to an instance of that class.) That's why Apple's own Cocoa documentation consists largely of pages listing and describing the methods supplied by some class. For example, to know what messages you can send to an NSString object (instance), you'd start by studying the NSString class documentation. That page is really just a big list of methods, so it tells you what an NSString object can do. That isn't everything in the world there is to know about an NSString, but it's a big percentage of it.

Class Methods

Up to now I've been keeping something back, and if you've been paying close attention, you may have caught me at it, because it looks as though I've contradicted myself. I said that nothing happens in a program unless a message is sent to an object. But I also said that there are no instances until they are created as the program runs. The contradiction is that if messages can be sent only to instances, it appears that no instances can ever be created (because, when the program starts up, there are no instances to which you can send the message asking for an instance to be created).

The truth that I've been keeping back, which complicates things only a little, is that classes are themselves objects and can be sent messages. This revelation solves the contradiction completely. No instances exist as the program starts up, but the classes do. The classes may live off in a world of Platonic Forms, but they can still be sent messages. And one of the most important things you can ask a class to do by sending it a message is to instantiate itself.

You cannot, however, ask an instance to instantiate itself. It thus begins to look as if there must be two kinds of message: messages that you are allowed to send to a class (such as telling the Dog class to instantiate itself) and messages that you are allowed to send to an instance (such as telling an individual dog to bark). That is exactly true. More precisely, all code lives as a method in a class, but methods are of two kinds: class methods and instance methods. If a method is a class method, you can send that message to the class. If a method is an instance method, you can send that message to an instance of the class.

In Objective-C syntax, class methods and instance methods are distinguished by the use of a plus sign or a minus sign. For example, Apple's NSString class documentation page listing the methods of the NSString class starts out like this:

```
+ string  
- init
```

The `string` method is a class method. The `init` method is an instance method.

In general, though not exclusively, class methods tend to be factory methods — that is, methods for generating an instance. This makes sense, because making an instance

of itself is one of the main things you’re likely to want to ask a class to do. You might think that a class really needs only *one* class method for generating an instance of itself, and that is rigorously true, but classes tend to provide multiple factory methods purely as a convenience to the programmer. For example, here are three NSString class methods:

```
+ string  
+ stringWithFormat:  
+ stringWithContentsOfFile:encoding:error:
```

They all make instances. The first class method, `string`, generates an empty NSString instance (a string with no text). The second class method, `stringWithFormat:`, generates an NSString instance based on text that you provide, which can include transforming other values into text; for example, you might use it to start with an integer 9 and generate an NSString instance @"9". The third class method reads the contents of a file and generates an NSString instance from those contents. When you come to write your own classes, you too might well create multiple class methods that act as instance factories for your own future programming convenience.

Instance Variables

Now that I’ve revealed that classes are objects and can be sent messages, you might be wondering why there need to be instances at all. Why doesn’t the mere existence of classes as objects suffice for object-based programming? Why would you ever bother to instantiate any of the classes? Why wouldn’t you write all your code as class methods, have the program send messages from one class object to another, and be done with it?

The answer is that instances have a feature that classes do not: instance variables. An instance variable is just what the name suggests: it’s a variable belonging to an instance. Like instance methods, instance variables are defined as part of the class. But the *value* of an instance variable is set as the program runs and belongs to one instance alone. In other words, different instances can have different values for the same instance variable.

For example, suppose we have a Dog class and we decide that it might be a good idea for every dog to have a name. Just as you can learn a real-world dog’s name by reading the tag on its collar, we want to be able to assign every dog instance a name and, subsequently, to learn what that name is. So, in designing the Dog class, we declare that this class has an instance variable called `name`, whose value is a string (probably an NSString, as we’re using Objective-C). Now when our program runs we can instantiate Dog and assign the resulting dog instance a name (that is, we can assign its `name` instance variable a value). We can also instantiate Dog again and assign *that* resulting dog instance a name. Let’s say these are two different names: one is @"Rover" and one is @"Fido". Then we’ve got two instances of Dog, and they are significantly different; they differ in the value of their `name` instance variables ([Figure 2-2](#)).

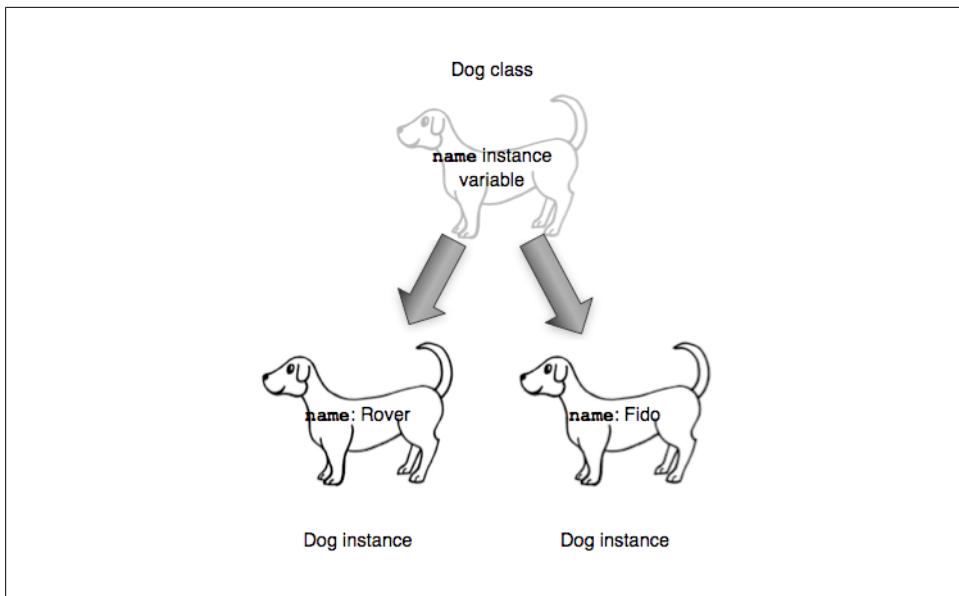


Figure 2-2. Instance variables

So an instance is a reflection of the instance methods of its class, but that isn't *all* it is; it's also a collection of instance variables. The class is responsible for what instance variables the instance has, but not for the values of those variables. The values can change as the program runs and apply only to a particular instance. An instance is a cluster of particular instance variable values.

In short, an instance is both code and data. The code it gets from its class and in a sense is shared with all other instances of that class, but the data belong to it alone. The data can persist as long as the instance persists. The instance has, at every moment, a state — the complete collection of its own personal instance variable values. An instance is a device for maintaining state. It's a box for storage of data.

The Object-Based Philosophy

In my REALbasic book, I summarized the nature of objects in two phrases: encapsulation of functionality, and maintenance of state:

Encapsulation of functionality

Each object does its own job, and presents to the rest of the world — to other objects, and indeed in a sense to the programmer — an opaque wall whose only entrances are the methods to which it promises to respond and the actions it promises to perform when the corresponding messages are sent to it. The details of how, behind the scenes, it actually implements those actions are secreted within itself; no other object needs to know them.

Maintenance of state

Each individual instance is a bundle of data that it maintains. Typically that data is private, which means that it's encapsulated as well; no other object knows what that data is or in what form it is kept. The only way to discover from outside what data an object is maintaining is if there's a method that reveals it.

As an example, imagine an object whose job is to implement a stack — it might be an instance of a *Stack* class. A *stack* is a data structure that maintains a set of data in LIFO order (last in, first out). It responds to just two messages: *push* and *pop*. Push means to add a given piece of data to the set. Pop means to remove from the set the piece of data that was most recently pushed and hand it out. It's like a stack of plates: plates are placed onto the top of the stack or removed from the top of the stack one by one, so the first plate to go onto the stack can't be retrieved until all other subsequently added plates have been removed ([Figure 2-3](#)).

The stack object illustrates encapsulation of functionality because the outside world knows nothing of how the stack is actually implemented. It might be an array, it might be a linked list, it might be any of a number of other implementations. But a client object — an object that actually sends a *push* or *pop* message to the stack object — knows nothing of this and cares less, provided the stack object adheres to its contract of behaving like a stack. This is also good for the programmer, who can, as the program develops, safely substitute one implementation for another without harming the vast machinery of the program as a whole. And just the other way round, the stack object knows nothing and cares less about who is telling it to push or to pop, and why. It just hums along and does its job in its reliable little way.

The stack object illustrates maintenance of state because it isn't just the gateway to the stack data — it *is* the stack data. Every object that has a reference to the stack object has the same access to its data, the same ability to push or to pop. (And that's all it can do. The stack data is effectively inside the stack object; no one else can see it. All that another object can do is push or pop.) If a certain object is at the top of our stack object's stack right now, then whatever object sends the *pop* message to this stack object will receive that object in return. If no object sends the *pop* message to this stack object, then the object at the top of the stack will just sit there, waiting.

As a second example of the philosophy and nature of object-based programming at work, I'll revert to another imaginary scenario I used in my *REALbasic* book. Pretend we're writing an arcade game where the user is to "shoot" at moving "targets," and the score increases every time a target is hit. We immediately have a sense of how we might organize our code using object-based programming and can see how object-based programming will fulfill its nature and purpose:

- There will be a Target class. Every target object will be an instance of this class. This decision makes sense because we want every target to behave the same way. A target will need to know how to draw itself; that knowledge will be part of the

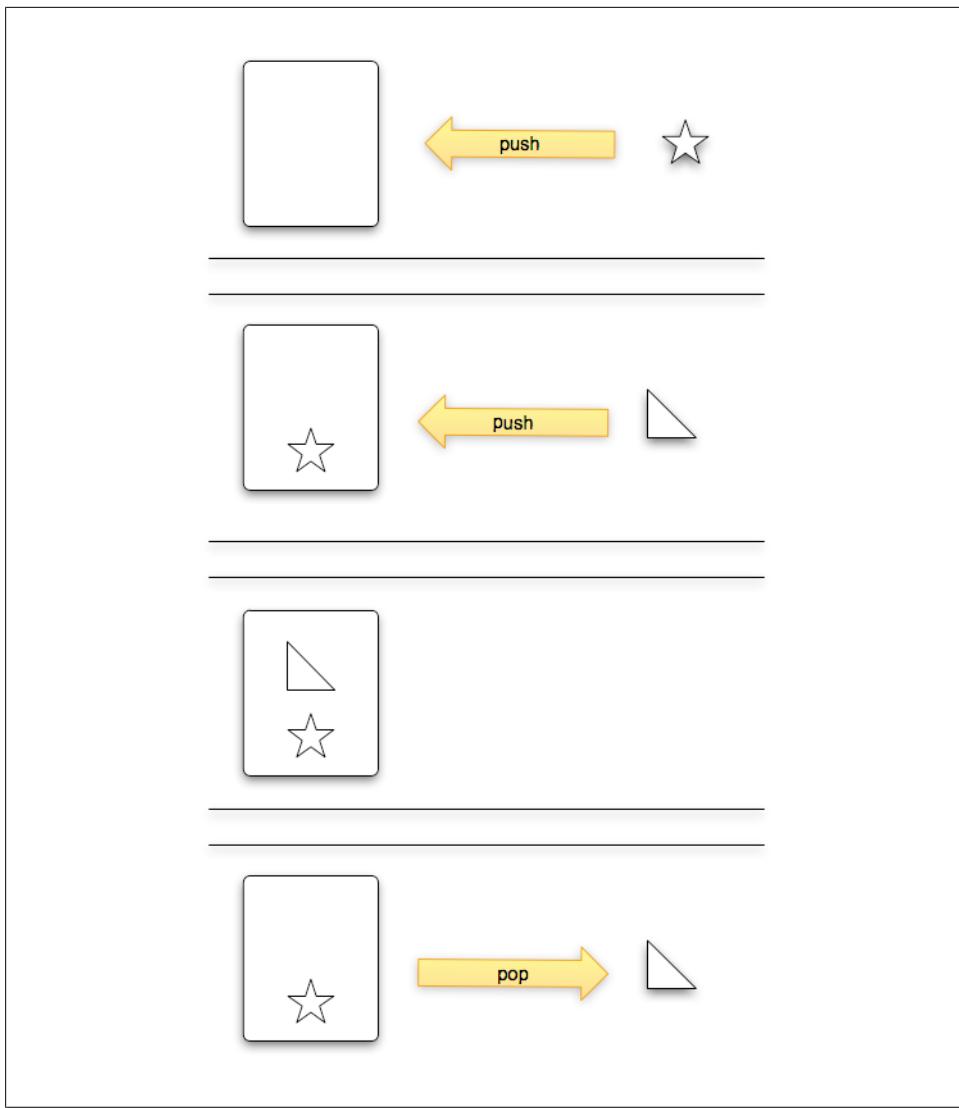


Figure 2-3. A stack

Target class, which makes sense because all targets will draw themselves in the same way. Thus we have the relationship between class and instance.

- Targets may draw themselves the same way, but they may also differ in appearance. Perhaps some targets are blue, others are red, and so on. This difference between individual targets can be expressed as an instance variable. Call it `color`. Every time we instantiate a target, we'll assign it a color. The Target class's code for drawing an individual target will look at that target's `color` instance variable and use it when filling in the target's shape. Clearly, we could extend this individualization as much

as we like: targets could have different sizes, different shapes, and so on, and all of these parametric distinctions could be made on an individual basis through the use of instance variables. Thus we have both encapsulation of functionality and maintenance of state. A target has a state, the parameters that describe how it should look, and also has the ability to draw itself, expressing that state visually.

- When a target is hit by the user, it will explode. So perhaps the Target class will have an `explode` instance method; thus, every target knows how to explode. One thing that should happen whenever a target explodes is that the user's score should increase. So let's imagine a score object — an instance of the Score class. Give every target object a reference to this score object so that it can send a message to it. When a target explodes, one of the things its `explode` instance method will do is send an `increase` message to the score object. Thus we have both encapsulation of functionality and maintenance of state. The score object responds indifferently to any object that sends it the `increase` message; it doesn't need to know why it's being sent that message. Nor does the score object even need to know that targets exist, or indeed that it's part of a game. It just sits there maintaining the score, and when it receives the `increase` message, it increases it.

This chapter has described only the rudiments of object-based philosophy — enough to communicate the correct mind-set. Using object-based programming effectively to make a program clear and maintainable is something of an art; your abilities will improve with experience. Eventually, you may want to do some further reading on how to construct an object-based program most effectively. I recommend in particular two classic, favorite books. *Refactoring*, by Martin Fowler (Addison-Wesley, 1999), describes how you can get a sense that you might need to rearrange what methods belong to what classes (and how to conquer your fear of doing so). *Design Patterns*, by Erich Gamma, Richard Helm, Ralph Johnson, and John Vlissides (also known as “the Gang of Four”), is the bible on architecting object-based programs, listing all the ways you can arrange objects with the right powers and the right knowledge of one another (Addison-Wesley, 1994).

Objective-C Objects and Messages

One of the first object-based programming languages to achieve maturity and widespread dissemination was Smalltalk. It was developed during the 1970s at Xerox PARC under the leadership of Alan Kay and started becoming widely known in 1980. The purpose of Objective-C, created by Brad Cox and Tom Love in 1986, was to build Smalltalk-like syntax and behavior on top of C. Objective-C was licensed by NeXT in 1988 and was the basis for its application framework API, NeXTStep. Eventually, NeXT and Apple merged, and the NeXT application framework evolved into Cocoa, the framework for Mac OS X applications, still revolving around Objective-C. That history explains why Objective-C is the base language for iOS programming. (It also explains why Cocoa class names often begin with “NS” — it stands for “NeXTStep.”)

Having learned the basics of C ([Chapter 1](#)) and the nature of object-based programming ([Chapter 2](#)), you are ready to meet Objective-C. This chapter describes Objective-C structural fundamentals; the next two chapters provide more detail about how Objective-C classes and instances work. (A few additional features of the language are discussed in [Chapter 10](#).) As with the C language, my intention is not to describe the Objective-C language completely, but to provide a practical linguistic grounding, founded on my own experience of those aspects of the language that need to be firmly understood as a basis for iOS programming.

An Instance Reference Is a Pointer

In C, every variable must be declared to be of some type. In an object-based language such as Objective-C, an instance’s type is its class. The C language includes very few basic data types. To facilitate the multiplicity of class types required by its object-based nature, Objective-C takes advantage of C pointers. So, in Objective-C, if a variable is an instance of the class `MyClass`, that variable is of type `MyClass*` — a pointer to a `MyClass`. In general, in Objective-C, a reference to an instance is a pointer and the name of the data type of what’s at the far end of that pointer is the name of the instance’s class.



Note the convention for capitalization. Variable names tend to start with a lowercase letter; class names tend to start with an uppercase letter.

As I mentioned in [Chapter 1](#), the fact that a reference to an instance is a pointer in Objective-C will generally not cause you any difficulties, because pointers are used consistently throughout the language. For example, a message to an instance is directed at the pointer, so there is no need to dereference the pointer. Indeed, having established that a variable representing an instance is a pointer, you’re likely to forget that this variable even *is* a pointer and just work directly with that variable:

```
NSString* s = @"Hello, world!";
NSString* s2 = [s uppercaseString];
```

Having established that `s` is an `NSString*`, you would never dereference `s` (that is, you would never speak of `*s`) to access the “real” `NSString`. So it feels as if the pointer *is* the real `NSString`. Thus, in the previous example, once the variable `s` is declared as a pointer to an `NSString`, the `uppercaseString` message is sent directly to the variable `s`. (The `uppercaseString` message asks an `NSString` to generate and return an uppercase version of itself; so, after that code, `s2` is `@"HELLO, WORLD!"`)

The tie between a pointer, an instance, and the class of that instance is so close that it is natural to speak of an expression like `MyClass*` as meaning “a `MyClass` instance,” and of a `MyClass*` value as “a `MyClass`.” A Objective-C programmer will say simply that, in the previous example, `s` *is* an `NSString`, that `uppercaseString` returns “an `NSString`,” and so forth. It is fine to speak like that, and I do it myself (and will do it in this book) — provided you remember that this is a shorthand. Such an expression means “an `NSString` instance,” and because an instance is represented as a C pointer, it means an `NSString*`, a pointer to an `NSString`.

Although the fact that instance references in Objective-C are pointers does not cause any special difficulty, you must still be conscious of what pointers are and how they work. As I emphasized in [Chapter 1](#), when you’re working with pointers, you must keep in mind the special meaning of your actions. So here are some basic facts about pointers that you should keep in mind when working with instance references in Objective-C.



Forgetting the asterisk in an instance declaration is a common beginner mistake, and will net you a mysterious compiler error message, such as “Interface type cannot be statically allocated.”

Instance References, Initialization, and nil

Merely declaring an instance reference’s type doesn’t bring any instance into existence. For example:

```
NSString* s; // only a declaration; no instance is pointed to
```

After that declaration, `s` is *typed* as a pointer to an `NSString`, but it is not *in fact* pointing to an `NSString`. You have created a pointer, but you haven't supplied an `NSString` for it to point to. It's just sitting there, waiting for you to point it at an `NSString`, typically by assignment (as we did with `@"Hello, world!"` earlier). Such assignment *initializes* the variable, giving it an actual meaningful value of the proper type.

You can declare a variable as an instance reference in one line of code and initialize it later, like this:

```
NSString* s;  
// ... time passes ...  
s = @"Hello, world!";
```

But this is not common. It is much more common, wherever possible, to declare and initialize a variable all in one line of code:

```
NSString* s = @"Hello, world!";
```

Declaration *without* initialization, before the advent of iOS 5 and ARC ([Chapter 12](#)), created a dangerous situation:

```
NSString* s;
```

What *is* `s` after a mere declaration like that? It could be anything. But it is *claiming* to be a pointer to an `NSString`, and so your code might proceed to *treat* it as a pointer to an `NSString`. But it is pointing at garbage. A pointer pointing at garbage is liable to cause serious trouble down the road when you accidentally try to use it as an instance. Sending a message to a garbage pointer, or otherwise treating it as a meaningful instance, can crash your program. Even worse, it might *not* crash your program: it might cause your program to behave very, very oddly instead — and figuring out why can be difficult.

For this reason, if you *aren't* going to initialize an instance reference pointer at the moment you declare it by assigning it a real value, it's a good idea to assign it `nil`:

```
NSString* s = nil;
```

A small but delightful bonus feature of using ARC is that this assignment is performed for you, implicitly and invisibly, as soon as you declare a variable without initializing it:

```
NSString* s; // under ARC, s is immediately set to nil for you
```

This prevents the existence of a garbage pointer, and could save you from yourself by preventing a crash when you accidentally use `s` as an instance without initializing it. Nevertheless, long years of habit have trained me to initialize or explicitly set to `nil` an instance pointer as soon as I declare it, and you'll see that I continue to do so in examples in this book.

What is `nil`? It's simply a form of zero — the form of zero appropriate to an instance reference. The `nil` value simply means: "This instance reference isn't pointing to any

instance.” Indeed, you can test an instance reference against nil as a way of finding out whether it is in fact pointing to a real instance. This is an extremely common thing to do:

```
if (nil == s) // ...
```

As I mentioned in [Chapter 1](#), the explicit comparison with nil isn’t strictly necessary; because nil is a form of zero, and because zero means false in a condition, you can perform the same test like this:

```
if (!s) // ...
```

I do in fact write nil tests in that second form all the time, but some programmers would take me to task for bad style. The first form has the advantage that its real meaning is made explicit, rather than relying on a cute implicit feature of C. The first form places nil first in the comparison so that if the programmer accidentally omits an equal sign, performing an assignment instead of a comparison, the compiler will catch the error (because assignment to nil is illegal).

Many Cocoa methods use a return value of nil, instead of an expected instance, to signify that something went wrong. You are supposed to capture this return value and test it for nil in order to discover whether something *did* go wrong. For example, the documentation for the `NSString` class method `stringWithContentsOfFile:encoding:error:` says that it returns “a string created by reading data from the file named by `path` using the encoding, `enc`. If the file can’t be opened or there is an encoding error, returns nil.” So, as I described in [Chapter 1](#), your next move after calling this method and capturing the result should be to test that result against nil, just to make sure you’ve really got an instance now:

```
NSString* path = // ... whatever;
NSStringEncoding enc = // ... whatever;
NSError* err = nil;
NSString* s = [NSString stringWithContentsOfFile:path encoding:enc error:&err];
if (nil == s) // oops! something went wrong...
```

You should now be wondering about the implications of a nil-value pointer for sending a message to a noninstance. For example, you can send a message to an `NSString` instance like this:

```
NSString* s2 = [s uppercaseString];
```

That code sends the `uppercaseString` message to `s`. So `s` is supposedly an `NSString` instance. But what if `s` is nil? With some object-based programming languages, sending a message to nil constitutes a runtime error and will cause your program to terminate prematurely (REALbasic and Ruby are examples). But Objective-C doesn’t work like that. In Objective-C, sending a message to nil is legal and does not interrupt execution. Moreover, if you capture the result of the method call, it will be a form of zero — which means that if you assign that result to an instance reference pointer, it too will be nil:

```
NSString* s = nil; // now s is nil
NSString* s2 = [s uppercaseString]; // now s2 is nil
```

Whether this behavior of Objective-C is a good thing is a quasi-religious issue and a subject of vociferous debate among programmers. It is useful, but it also extremely easy to be tricked by it. The usual scenario is that you accidentally send a message to a nil reference without realizing it, and then later your program doesn't behave as expected. Because the point where the unexpected behavior occurs is later than the moment when the nil pointer arose in the first place, the genesis of the nil pointer can be difficult to track down (indeed, it often fails to occur to the programmer that a nil pointer is the cause of the trouble in the first place).

Short of peppering your code with tests to ascertain that your instance reference pointers are not accidentally nil, which is not generally a good idea, there isn't much you can do about this. This behavior is strongly built into the language and is not going to change. It's just something you need to be aware of.

If, on the other hand, a method call can return nil, be conscious of that fact. Don't assume that everything will go well and that it won't return nil. On the contrary, if something can go wrong, it probably will. For example, to omit the nil test after calling `stringWithContentsOfFile:encoding:error:` is just stupid. I don't care if you know perfectly well that the file exists and the encoding is what you say it is — test the result for nil!

Instance References and Assignment

As I said in [Chapter 1](#), assigning to a pointer does not mutate the value at the far end of the pointer; rather, it repoints the pointer. Moreover, assigning one pointer to another repoints the pointer in such a way that both pointers are now pointing to the very same thing. Failure to keep these simple facts firmly in mind can have results that range from surprising to disastrous.

For example, instances in general are usually mutable: they typically have instance variables that can change. If two references are pointing at one and the same instance, then when the instance is mutated by way of one reference, that mutation also affects the instance as seen by the other reference. To illustrate, pretend that we've implemented the Stack class described in the previous chapter:

```
Stack* myStack1 = // ... create Stack instance and initialize myStack1 ... ;
Stack* myStack2 = myStack1;
[myStack1 push: @"Hello"];
[myStack1 push: @"World"];
NSString* s = [myStack2 pop];
```

After we pop `myStack2`, `s` is @"World" even though nothing was ever pushed onto `myStack2` (and the stack `myStack1` contains only @"Hello" even though nothing was ever popped off of `myStack1`). That's because we did push two strings onto `myStack1` and then pop one string off `myStack2`, and `myStack1` is `myStack2` — in the sense that they are both pointers to the very same stack instance. That's perfectly fine, as long as you understand and intend this behavior.

In real life, you’re likely to pass an instance off to some other object, or to receive it from some other object:

```
Stack* myStack = // ... create Stack instance and initialize myStack ... ;  
// ... more code might go here ...  
[myObject doSomethingWithThis: myStack]; // pass myStack to myObject
```

After that code, `myObject` has a pointer to the very same instance we’re already pointing to as `myStack`. So we must be careful and thoughtful. The object `myObject` might mutate `myStack` right under our very noses. Even more, the object `myObject` might *keep* its reference to the stack instance and mutate it *later* — possibly much later, in a way that could surprise us. This is possible because instances can have instance variables that point to other objects, and those pointers can persist as long as the instances themselves do. This kind of shared referent situation can be intentional, but it is also something to watch out for and be conscious of ([Figure 3-1](#)).

Another possible misunderstanding is to imagine that the assignment `myStack2 = myStack1` somehow makes a new, separate instance that duplicates `myStack1`. That’s not at all the case. It doesn’t make a new instance; it just points `myStack2` at the very same instance that `myStack1` is pointing at. It may be possible to make a new instance that duplicates a given instance, but the ability to do so is not a given and it is not going to happen through mere assignment. (For how a separate duplicate instance might be generated, see the `NSCopying` protocol and the `copy` method mentioned in [Chapter 10](#).)

Instance References and Memory Management

The pointer nature of instance references in Objective-C also has implications for management of memory. The scope, and in particular the lifetime, of variables in pure C is typically quite straightforward: if you bring a piece of variable storage into existence by declaring that variable within a certain scope, then when that scope ceases to exist, the variable storage ceases to exist. That sort of variable is called *automatic* (K&R 1.10). So, for example:

```
void myFunction() {  
    int i; // storage for an int is set aside  
    i = 7; // 7 is placed in that storage  
} // the scope ends, so the int storage and its contents vanish
```

But in the case of a pointer, there are two pieces of memory to worry about: the pointer itself, which is an integer signifying an address in memory, and whatever is at the far end of that pointer. Nothing about the C language causes the destruction of what a pointer points to when the pointer itself is automatically destroyed as it goes out of scope:

```
void myFunction() {  
    NSString* s = @"Hello, world!"; // storage for a pointer is set aside  
    NSString* s2 = [s uppercaseString]; // storage for another pointer is set aside  
} // the two pointers go out of existence...  
// ... but what about the two NSStrings they point to?
```

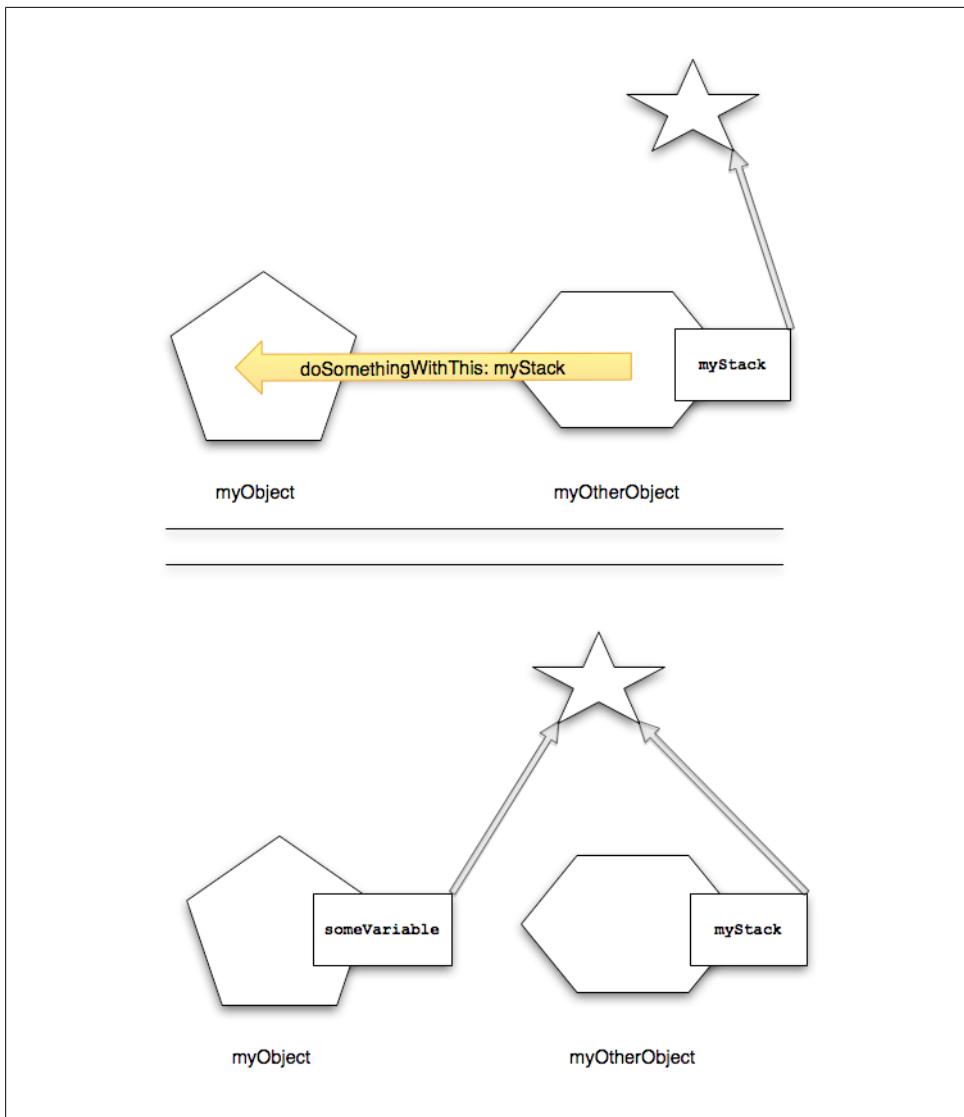


Figure 3-1. Two instances end up with pointers to the same third instance

Some object-based programming languages in which a reference to an instance is a pointer do manage automatically the memory pointed to by instance references (REALbasic and Ruby are examples). But Objective-C, at least the way it's implemented when you're programming for iOS, is not one of those languages. Because the C language has nothing to say about the automatic destruction of what is pointed to by a reference to an instance, Objective-C implements an explicit mechanism for the management of memory. I'll talk in a later chapter ([Chapter 12](#)) about what that mechanism

is and what responsibilities for the programmer it entails. (Fortunately, under ARC, those responsibilities are fewer than they used to be.)

Messages and Methods

An Objective-C method is defined as part of a class. It has three aspects:

Whether it's a class method or an instance method

If it's a class method, you call it by sending a message to the class itself. If it's an instance method, you call it by sending a message to an instance of the class.

Its parameters and return value

As with a C function, an Objective-C method takes some number of parameters; each parameter is of some specified type. And, as with a C function, it may return a value, which is also of some specified type; if the method returns nothing, its return type is declared as `void`.

Its name

An Objective-C method's name must contain as many colons as it takes parameters. The name is split after each colon in a method call or declaration, so it is usual for the part of the name preceding each colon to describe the corresponding parameter.

Sending a Message

As you've doubtless gathered, the syntax for sending a message to an object involves square brackets. The first thing in the square brackets is the object to which the message is to be sent; this object is the message's *receiver*. Then follows the message:

```
NSString* s2 = [s uppercaseString]; // send "uppercaseString" message to s ...
// ... (and assign result to s2)
```

If the message is a method that takes parameters, each corresponding argument value comes after a colon:

```
[myStack1 push: @"Hello"]; // send "push:" message to myStack1 ...
// ...with one argument, the NSString @"Hello"
```

To send a message to a class (calling a class method), you can represent the class by the literal name of the class:

```
NSString* s = [NSString string]; // send "string" message to NSString class
```

To send a message to an instance (calling an instance method), you'll need a reference to an instance, which (as you know) is a pointer:

```
NSString* s = @"Hello, world!"; // and now s is initialized as an NSString instance
NSString* s2 = [s uppercaseString]; // send "uppercaseString" message to s
```

You can send a class method to a class, and an instance method to an instance, no matter how you got hold of and represent the class or the instance. For example, @"Hello, world!" is itself an `NSString` instance, so it's legal to say:

```
NSString* s2 =[@"Hello, world!" uppercaseString];
```

If a method takes no parameters, then its name contains no colons, like the `NSString` instance method `uppercaseString`. If a method takes one parameter, then its name contains one colon, which is the final character of the method name, like the hypothetical Stack instance method `push:`. If a method takes two or more parameters, its name contains that number of colons. In the minimal case, its name ends with that number of colons. For example, a method taking three parameters might be called `here-AreThreeStrings:::`. To call it, we split the name after each colon and follow each colon with an argument, which looks like this:

```
[someObject hereAreThreeStrings:@"string1":@"string2":@"string3"];
```

That's a legal way to name a method, but it isn't very common, mostly because it isn't very informative. Usually the name will have more text; in particular, the part before each colon will describe the parameter that follows that colon.

For example, there's a `UIColor` class method for generating an instance of a `UIColor` from four `CGFloat` numbers representing its red, green, blue, and alpha (transparency) components, and it's called `colorWithRed:green:blue:alpha:`. Notice the clever construction of this name. The `colorWith` part tells something about the method's purpose: it generates a *color*, starting *with* some set of information. All the rest of the name, `Red:green:blue:alpha:`, describes the meaning of each parameter. And you call it like this:

```
UIColor* c =[UIColor colorWithRed: 0.0 green: 0.5 blue: 0.25 alpha: 1.0];
```

The space after each colon in the method call is optional. (Space before a colon is also legal, though in practice one rarely sees this.)

The rules for naming an Objective-C method, along with the conventions governing such names (like trying to make the name informative about the method's purpose and the meanings of its parameters), lead to some rather long and unwieldy method names, such as `getBytes:maxLength:usedLength:encoding:options:range:remainingRange:`. Such verbosity of nomenclature is characteristic of Objective-C. Method calls, and even method declarations, are often split across multiple lines to prevent a single line of code from becoming so long that it wraps within the editor, as well as for clarity.

Declaring a Method

The declaration for a method has three parts:

- Either `+` or `-`, meaning that the method is a class method or an instance method, respectively.
- The data type of the return value, in parentheses.

- The name of the method, split after each colon. Following each colon is the corresponding parameter, expressed as the data type of the parameter, in parentheses, followed by a placeholder name for the parameter.

So, for example, Apple's documentation tells us that the declaration for the UIColor class method `colorWithRed:green:blue:alpha:` is:

```
+ (UIColor*) colorWithRed: (CGFloat) red green: (CGFloat) green
                      blue: (CGFloat) blue alpha: (CGFloat) alpha
```

(Note that I've split the declaration into two lines, for legibility and to fit onto this page. The documentation puts it all on a single line.)

Make very sure you can read this declaration! You should be able to look at it and say to yourself instantly, "The name of this method is `colorWithRed:green:blue:alpha:`. It's a class method that returns a UIColor and takes four CGFloat parameters."

It is not uncommon, outside of code, to write a method's name along with the plus sign or the minus sign, to make it clear whether this is a class method or an instance method. So you might speak informally of "-uppercaseString," just as a way of reminding yourself or a reader that this is an instance method. Again outside of code, it is not uncommon, especially when communicating with other Objective-C programmers, to speak of a method's name along with the class in which this method is defined. So you might say "NSString's -uppercaseString," or even something like "-[NSString uppercaseString]." Notice that that isn't code, or even pseudo-code, because you are not actually speaking of a method call, and in any case you could never send the `uppercaseString` message to the NSString class; it's just a compact way of saying, "I'm talking about the `uppercaseString` that's an instance method of NSString."

Nesting Method Calls

Wherever in a method call an object of a certain type is supposed to appear, you can put another method call that returns that type. Thus you can nest method calls. A method call can appear as the message's receiver:

```
NSString* s = [[NSString string] uppercaseString]; // silly but legal
```

That's legal because NSString's class method `string` returns an NSString instance (formally, an `NSString*` value, remember), so we can send an NSString instance method to that result. Similarly, a method call can appear as an argument in a method call:

```
[myStack push: [NSString string]]; // ok if push: expects an NSString* parameter
```

However, I must caution you against overdoing that sort of thing. Code with a lot of nested square brackets is very difficult to read (and to write). Furthermore, if one of the nested method calls happens to return nil unexpectedly, you have no way to detect this fact. It is often better, then, to be even more verbose and declare a temporary variable for each piece of the method call. Just to take an example from my own code, instead of writing this:

```
NSArray* arr = [[MPMediaQuery albumsQuery] collections];
```

I might write this:

```
MPMediaQuery* query = [MPMediaQuery albumsQuery];
NSArray* arr = [query collections];
```

Even though the first version is quite short and legible, and even though the variable `query` will never be used again — it exists solely in order to be the receiver of the `collections` message in the second line — it is worth creating it as a separate variable. For one thing, it makes this code far easier to step through in the debugger later on, when I want to pause after the `albumsQuery` call and see whether the expected sort of result is being returned.



Incorrect number or pairing of nested square brackets can net you some curious messages from the compiler. For example, too many pairs of square brackets (e.g. `[[query collections]]`) or an unbalanced left square bracket (`[[query collections]`) is reported as “Expected identifier.”

No Overloading

The data type returned by a method, together with the data types of each of its parameters in order, constitute that method’s *signature*. It is illegal for two methods of the same type (class method or instance method) to exist in the same class with the same name but different signatures.

So, for example, you could not have two `MyClass` instance methods called `myMethod`, one of which returns `void` and one of which returns an `NSString`. Similarly, you could not have two `MyClass` instance methods called `myMethod:`, both returning `void`, one taking a `CGFloat` parameter and one taking an `NSString` parameter. An attempt to violate this rule will be stopped dead in its tracks by the compiler, which will announce a “duplicate declaration” error. The reason for this rule is that if two such conflicting methods were allowed to exist, there would be no way to determine from a method call to one of them *which* method was being called.

You might think that the issue could be decided by looking at the types involved in the call. If one `myMethod:` takes a `CGFloat` parameter and the other `myMethod:` takes an `NSString` parameter, you might think that when `myMethod:` is called, Objective-C could look at the actual argument and realize that the former method is meant if the argument is a `CGFloat` and the latter if the argument is an `NSString`. But Objective-C doesn’t work that way. There are languages that permit this feature, called *overloading*, but Objective-C is not one of them.

Parameter Lists

It isn't uncommon for an Objective-C method to require an unknown number of parameters. A good example is the NSArray class method `arrayWithObjects:`, which looks from the name as if it takes one parameter but in fact takes any number of parameters, separated by comma. The parameters are the objects of which the NSArray is to consist. The trick here, however, which you must discover by reading the documentation, is that the list must end with nil. The nil is not one of the objects to go into the NSArray (nil isn't an object, so an NSArray can't contain nil); it's to show where the list ends.

So, here's a correct way to call the `arrayWithObjects:` method:

```
NSArray* pep = [NSArray arrayWithObjects:@"Manny", @"Moe", @"Jack", nil];
```

The declaration for `arrayWithObjects:` uses three dots to show that a comma-separated list is legal:

```
+ (id)arrayWithObjects:(id)firstObj, ... ;
```

Without the nil terminator, the program will not know where the list ends, and bad things will happen when the program runs, as it goes hunting off into the weeds of memory, incorporating all sorts of garbage into the NSArray that you never meant to have incorporated. Forgetting the nil terminator is a common beginner error, but not as common as it used to be: by a bit of deep-C voodoo, the Objective-C compiler now notices if you've forgotten the nil, and warns you ("missing sentinel in method dispatch"). Even though it's just a warning, don't run that code.

The C language has explicit provision for argument lists of unspecified length, which Objective-C methods such as `arrayWithObjects:` are using behind the scenes. I'm not going to explain the C mechanism, because I don't expect you'll ever write a method or function that requires it; see K&R 7.3 if you need the gory details.

Unrecognized Selectors

Objective-C messaging is dynamic, meaning that the compiler takes no formal responsibility for whether a particular object is a legal recipient of a given message. That's because whether an object can deal with a message sent to it isn't decided until the program actually runs and the message actually arrives. Objective-C has various devices for dealing at runtime with a message that doesn't correspond directly to a method, and for all the compiler knows, one of them might come into play in this case. For example, at the time the program runs, the recipient of the message might be nil — and it's harmless to send *any* message to nil.

Thus, it is theoretically legal to direct a message at an object with no corresponding method. The only guardian against this possibility is the compiler. Before ARC, the compiler was not a very strong guardian in this respect. For example:

```
NSString* s = @"Hello, world!";
[s rockTheCasbah]; // without ARC, compiler warns
```

An `NSString` has no method `rockTheCasbah`. But the compiler will not stop you from running a program containing this code; it's legal. The compiler will *warn* you, but it won't stop you. There are actually two possible warnings:

- If no `rockTheCasbah` method is defined *anywhere* in your code, the compiler will say: "Instance method '-rockTheCasbah' not found (return type defaults to 'id')." Without going into the details, what the compiler means is: "I know of no instance method `rockTheCasbah`, so I can't check its signature against the return type and arguments you're actually using, so I'll just make some loose assumptions and let it pass."
- If a `rockTheCasbah` method *is* defined somewhere in your code, the compiler will say: "'`NSString`' may not respond to 'rockTheCasbah'." This means: "There's a `rockTheCasbah` method, all right, but you seem to be sending the `rockTheCasbah` method to an instance of a class that doesn't have it as an instance method."

This is a good example of what I meant in [Chapter 2](#) when I said that sending a message and calling a method were not the same thing. The compiler is saying that `NSString` has no `rockTheCasbah` instance method, but that it isn't going to stop you from sending an `NSString` a `rockTheCasbah` message. At runtime, the object that receives the `rockTheCasbah` message might be able to deal with it, for all the compiler knows.

With ARC, however, the compiler is much stricter. The example above won't compile at all under ARC! The compiler declares a fatal compilation error: "Receiver type '`NSString`' for instance message does not declare a method with selector 'rockTheCasbah'." There is no `NSString` method `rockTheCasbah`, and by golly the compiler isn't going to let you send the `rockTheCasbah` message to an `NSString`, and that's final.

This is another of those delightful secondary benefits of using ARC. In order to do what it primarily does (manage memory), ARC must insist on more information about classes and their methods than the Objective-C standard calls for. Here, ARC is demanding that you *prove* that an `NSString` can respond to `rockTheCasbah`, or it won't let you run this code at all. (Nevertheless, if you really want to, you can slip past even ARC's stringent guardianship; I'll explain how in the next section.)

Let us assume for a moment, however, that we are compiling without ARC, or that we have somehow tricked even ARC into letting us compile successfully. Warning or no warning, we are now ready to run a program that sends the `rockTheCasbah` message to an `NSString`, and damn the consequences. What might those consequences be? Quite simply, if you send a message to an object that can't deal with it, your program will crash at that moment. So, for example, our attempt to send an `NSString` the `rockTheCasbah` message will crash our program, with a message (in the console log) of this form: "-[NSCFConstantString rockTheCasbah]: unrecognized selector sent to instance 0x3048."

The important thing here is the phrase *unrecognized selector*. The term "selector" is roughly equivalent to "message," so this is a way of saying a certain instance was sent

a message it couldn't deal with. The console message also tries to tell us *what* instance this was. `0x3048` is the value of the instance pointer; it is the address in memory to which our `NSString*` variable `s` was actually pointing. (Never mind why the `NSString` is described as an `NSCFConstantString`; this has to do with `NSString`'s implementation behind the scenes.)

(Strictly speaking, I should not say that a situation like this will "crash our program." What it will actually do is to generate an *exception*, an internal message as the program runs signifying that something bad has happened. It is possible for Objective-C code to "catch" an exception, in which case the program will not crash. The reason the program crashes, technically, is not that a message was sent to an object that couldn't handle it, but that the exception generated in response wasn't caught. That's why the crash log may also say, "Terminating app due to uncaught exception.")

Typecasting and the `id` Type

One way to silence the compiler when it warns in the way I've just described is by typecasting. A typecast, however, is not a viable way of fixing the problem unless it also tells the truth. It is perfectly possible to lie to the compiler by typecasting; this is not nice, and is not likely to yield nice consequences.

For example, suppose we've defined a class `MyClass` that does contain an instance method `rockTheCasbah`. As a result, it is fine with the compiler if you send the `rockTheCasbah` message to a `MyClass`, although it is not fine to send the `rockTheCasbah` message to an `NSString`. So you can silence the compiler by claiming that an `NSString` instance *is* a `MyClass` instance:

```
NSString* s = @"Hello, world!";
[(MyClass*)s rockTheCasbah];
```

The typecast silences the compiler; there is no warning. Notice that the typecast is not a value conversion; it's merely a claim about what the type will turn out to be at runtime. You're saying that when the program runs, `s` will magically turn out to be a `MyClass` instance. Because `MyClass` has a `rockTheCasbah` instance method, that silences the compiler. Of course, you've lied to the compiler, so when the program runs it will crash anyway, in exactly the same way as before! You're still sending an `NSString` a message it can't deal with, so the very same exception about sending an unrecognized selector to an `NSCFConstantString` instance will result. So don't do that!

Sometimes, however, typecasting to silence the compiler is exactly what you do want to do. This situation quite often arises in connection with class inheritance. We haven't discussed class inheritance yet, but I'll give an example anyway. Let's take the built-in Cocoa class `UINavigationController`. Its `topViewController` method is declared to return a `UIViewController` instance. In real life, though, it is likely to return an instance of some class you've created. So in order to call a method of the class you've created on the instance returned by `topViewController` without upsetting the compiler, you

have to reassure the compiler that this instance really will be an instance of the class you've created. That's what I'm doing in this line from one of my own apps:

```
[((RootViewController*)[navigationController topViewController]) setAlbums: arr];
```

The expression `(RootViewController*)` is a typecast in which I'm assuring the compiler that at this moment in the program, the value returned by the `topViewController` method call will in fact be an instance of `RootViewController`, which is my own defined class. The typecast silences the compiler when I send this instance the `setAlbums:` message, because my `RootViewController` class has a `setAlbums:` instance method and the compiler knows this. And the program doesn't crash, because I'm not lying: this `topViewController` method call really *will* return a `RootViewController` instance.

Objective-C also provides a special type designed to silence the compiler's worries about object data types altogether. This is the `id` type. An `id` is a pointer, so you don't say `id*`. It is defined to mean "an object pointer," plain and simple, with no further specification. Thus, every instance reference is also an `id`.

Use of the `id` type causes the compiler to stop worrying about the relationship between object types and messages. The compiler can't know anything about what the object will really be, so it throws up its hands and doesn't warn about anything. Moreover, any object value can be assigned or typecast to an `id`, and *vice versa*. The notion of assignment includes parameter passing. So you can pass a value typed as an `id` as an argument where a parameter of some particular object type is expected, and you can pass any object as an argument where a parameter of type `id` is expected. (I like to think of an `id` as analogous to both type AB blood and type O blood: it is both a universal recipient and a universal donor.) So, for example:

```
NSString* s = @"Hello, world!";
id unk = s;
[unk rockTheCasbah];
```

The second line is legal, because any object value can be assigned to an `id`. The third line doesn't generate any compiler warning, because any message can be sent to an `id`. (Of course the program will *still* crash when it actually runs and `unk` turns out to be an `NSString` and incapable of receiving of the `rockTheCasbah` message.)

That trick works even under ARC, with one caveat. ARC is willing to let that code compile — but only if a matching `rockTheCasbah` method is defined *somewhere* in your code (even if it isn't an `NSString` method). If there's no such method, ARC will stop you with a different error: "No known instance method for selector 'rockTheCasbah'." This is another way of saying the same thing the non-ARC compiler said earlier: "I know of no instance method `rockTheCasbah`, so I can't check its signature against the return type and arguments you're actually using." But instead of implicitly adding, "So I'll just make some loose assumptions and let it pass," ARC is stricter. After all, even without knowing what class `unk` will turn out to be when the program runs, ARC can be pretty sure that that class won't have a `rockTheCasbah` method, because *no* known

class has a `rockTheCasbah` method. So ARC, like a good guardian, continues to bar the way.

If, however, a matching `rockTheCasbah` method *is* defined somewhere in your code, even though it isn't an `NSString` method, ARC now takes its hands off the tiller entirely, and permits the program to compile and run without warning. You are now sending a message to an `id`, and an `id` can legally receive any message. If you crash at runtime, that's your problem; ARC can't save you from yourself.

If an `id`'s ability to receive any message reminds you of nil, it should. I have already said that nil is a form of zero; I can now specify what form of zero it is. It's zero cast as an `id`. Of course, it still makes a difference at runtime whether an `id` is nil or something else; sending a message to nil won't crash the program, but sending an unknown message to an actual object will.

Thus, `id` is a device for turning off the compiler's type checking altogether. Concerns about what type an object is are postponed until the program is actually running. All the compiler can do is intelligently analyze your code to see if you might be making a mistake that could matter at runtime. Using `id` turns off this part of the compiler's intelligence and leaves you to your own devices.

I do not recommend that you make extensive use of `id` to live in a world of pure dynamism. The compiler is your friend; you should let it use what intelligence it has to catch mistakes in your code. Thus, I almost never declare a variable or parameter as an `id`. I want my object types to be specific, so that the compiler can help check my code.

On the other hand, the Cocoa API does make frequent use of `id`, because it has to. For example, consider the `NSArray` class, which is the object-based version of an array. In pure C, you have to declare what type of thing lives in an array; for example, you could have "an array of int." In Objective-C, using an `NSArray`, you can't do that. Every `NSArray` is an array of `id`, meaning that every element of the array can be of any object type. You can put a specific type of object into an `NSArray` because any specific type of object can be assigned to an `id` (`id` is the universal recipient). You can get any specific type of object back out of an `NSArray` because an `id` can be assigned to any specific type of object (`id` is the universal donor).

So, for example, `NSArray`'s `lastObject` method is defined as returning an `id`. So, given an `NSArray arr`, I can fetch its last element like this:

```
id unk = [arr lastObject];
```

However, after that code, `unk` can now be sent any message at all, and we are dispensing with the compiler's type checking. Therefore, if I happen to *know* what type of object an array element is, I always assign or cast it to that type. For example, let's say I happen to know that `arr` contains nothing but `NSString` instances (because I put them there in the first place). Then I will say:

```
NSString* s = [arr lastObject];
```

The compiler doesn't complain, because an `id` can be assigned to any specific type of object (`id` is the universal donor). Moreover, from here on in, the compiler regards `s` as an `NSString`, and uses its type checking abilities to make sure I don't send `s` any non-`NSString` messages, which is just what I wanted. And I didn't lie to the compiler; at runtime, `s` really is an `NSString`, so everything is fine.

The compiler's type checking is called *static typing*, as opposed to the dynamic behavior that takes place when the program actually runs. What I'm saying here, then, is that I prefer to take advantage of static typing as much as possible.

The Cocoa API will sometimes return an `id` from a method call where you might not expect it. It's good to be conscious of this, because otherwise the compiler can mislead you into thinking you're doing something safe when you're not. For example, consider this code:

```
UIColor* c = [NSString string];
```

This is clearly a mistake — you're assigning an `NSString` to a `UIColor` variable, which is likely to lead to a crash later on — but the compiler is silent. Why doesn't the compiler warn here? It's because the `NSString` `string` class method is declared like this:

```
+ (id)string
```

The `string` method returns an `NSString`, but its return value is typed as an `id`. An `id` can be assigned where any object type is expected, so the compiler doesn't complain when it's assigned to a `UIColor` variable. This fact is a common source of programmer mistakes (especially if the programmer is me).

Earlier, I said that it is illegal for the same class to define methods of the same type (class method or instance method) with the same name but different signatures. But I did not say what happens when two *different* classes declare conflicting signatures for the same method name. This is another case in which it matters whether you're using static or dynamic typing. If you're using static typing — that is, the type of the object receiving the message is specified — there's no problem, because there's no doubt which method is being called (it's the one in that object's class). But if you're using dynamic typing, where the object receiving the message is an `id`, you might get a warning from the compiler; and if you're using ARC, you'll get a downright error: "Multiple methods named 'rockTheCasbah' found with mismatched result, parameter type or attributes." This is another reason why method names are so verbose: it's in order to make each method name unique, preventing two different classes from declaring conflicting signatures for the same method.



Accidentally defining your own method with the same name as an existing Cocoa method can cause mysterious problems. For example, in a recent online query, a programmer was confused because the compiler complained that his call to `initWithObjects:` lacked a nil terminator, even though his `initWithObjects:` didn't need a nil terminator. No, *his* `initWithObjects:` didn't, but *Cocoa's* did, and the compiler couldn't distinguish them because this message was being sent to an `id`. He should have picked a different name.

Messages as Data Type

Objective-C is so dynamic that it doesn't have to know until runtime what message to send to an object or what object to send it to. Certain important methods actually accept both pieces of information as parameters. For example, consider this method declaration from Cocoa's `NSNotificationCenter` class:

```
- (void)addObserver:(id)notificationObserver selector:(SEL)notificationSelector  
name:(NSString *)notificationName object:(id)notificationSender
```

We'll discuss later what this method does (when we talk about notifications in [Chapter 11](#)), but the important thing to understand here is that it constitutes an instruction to send a certain message to a certain object at some later, appropriate time. For example, our purpose in calling this method might be to arrange to have the message `tickleMeElmo:` sent at some later, appropriate time to the object `myObject`.

So let's consider how we might actually make this method call. The object to which the message will be sent is here called `notificationObserver`, and is typed as an `id` (making it possible to specify any type of object to send the message to). So, for the `notificationObserver` parameter, we're going to pass `myObject`. The message itself is the `notificationSelector` parameter, which has a special data type, `SEL` (for "selector," the technical term for a message name). The question now is how to express the message name `tickleMeElmo:`.

You can't just put `tickleMeElmo:` as a bare term; that doesn't work syntactically. You might think you could express it as an `NSString`, `@"tickleMeElmo:"`, but surprisingly, that doesn't work either. It turns out that the correct way to do it is like this:

```
@selector(tickleMeElmo:)
```

The term `@selector()` is a directive to the compiler, telling it that what's in parentheses is a message name. Notice that what's in parentheses is not an `NSString`; it's the bare message name. And because it is the name, it must have no spaces and must include any colons that are part of the message name.

So the rule is extremely easy: when a `SEL` is expected, you'll usually pass a `@selector()` expression. Failure to get this syntax right, however, is a common beginner error. Notice also that this syntax is an invitation to make a typing mistake, especially because there is no checking by the compiler. If `myObject` implements a `tickleMe-`

`Elmo`: method and I accidentally type `@selector(tickleMeElmo)`, forgetting the colon or making any other mistake in specifying the message name, there is no compiler error; the problem won't be discovered until the program runs and something bad happens. (In this case, if the `tickleMeElmo` message without the colon is ever sent to `myObject`, the app will probably crash with an unrecognized selector exception.)

C Functions

Although your code will certainly call many Objective-C methods, it will also probably call quite a few C functions. For example, I mentioned in [Chapter 1](#) that the usual way of initializing a `CGPoint` based on its `x` and `y` values is to call `CGPointMake`, which is declared like this:

```
CGPoint CGPointMake (
    CGFloat x,
    CGFloat y
);
```

Make certain that you can see at a glance that this *is* a C function, not an Objective-C method, and be sure you understand the difference in the calling syntax. To call an Objective-C method, you send a message to an object, in square brackets, with each argument following a colon in the method's name; to call a C function, you use the function's name followed by parentheses containing the arguments.

You might even have reason to write your own C functions as part of a class, instead of writing a method. A C function has lower overhead than a full-fledged method; so even though it lacks the object-oriented abilities of a method, it is sometimes useful to write one, as when some utility calculation must be called rapidly and frequently. Also, once in a while you might encounter a Cocoa method or function that requires you to supply a C function as a “callback.”

An example is the `NSArray` method `sortedArrayUsingFunction:context:`. The first parameter is typed like this:

```
NSInteger (^)(id, id, void *)
```

That expression denotes, in the rather tricky C syntax used for these things, a pointer to a function that takes three parameters and returns an `NSInteger`. The three parameters of the function are an `id`, an `id`, and a pointer-to-void (which means any C pointer). The address operator (see [Chapter 1](#)) can be used to obtain a pointer to a C function. So to call `sortedArrayUsingFunction:context:` you'd need to write a C function that meets this description, and use its name, preceded by an ampersand, as the first argument.

To illustrate, I'll write a “callback” function to sort an `NSArray` of `NSStrings` on the last character of each string. (This would be an odd thing to do, but it's only an example!) The `NSInteger` returned by the function has a special meaning: it indicates whether the first parameter is to be considered less than, equal to, or larger than the second. I'll

obtain it by calling the `NSString` `compare:` method, which returns an `NSInteger` with that same meaning. Here's the function:

```
NSInteger sortByLastCharacter(id string1, id string2, void* context) {
    NSString* s1 = (NSString*) string1;
    NSString* s2 = (NSString*) string2;
    NSString* string1end = [s1 substringFromIndex:[s1 length] - 1];
    NSString* string2end = [s2 substringFromIndex:[s2 length] - 1];
    return [string1end compare:string2end];
}
```

And here's how we'd call `sortedArrayUsingFunction:context:` with that function as our callback (assume that `arr` is an `NSArray` of strings):

```
NSArray* arr2 = [arr sortedArrayUsingFunction:&sortByLastCharacter context:NULL];
```

(`NULL` is the C equivalent of `nil`, a pointer to nothing; where `nil` is zero typed as `id`, `NULL` is zero typed as `void*`. The `context` argument is expected to be a `void*`, which is a C value, so I've used `NULL` rather than `nil`. In fact, `nil` would have worked — the two are to some extent implicitly interchangeable — but they are not identical, so it's a good idea to try to keep them straight. For one thing, you can't send a message to `NULL`.)

C Struct Pointers

Many Objective-C objects have lower-level C counterparts, along with C functions for manipulating them. For example, besides the Objective-C `NSString`, there is also something called a `CFString`; the “CF” stands for “Core Foundation,” which is a lower-level C-based API. A `CFString` is an opaque C struct (“opaque” means that the elements constituting this struct are kept secret, and that you should operate on a `CFString` only by means of appropriate functions). As with an `NSString` or any other object, in your code you'll typically refer to a `CFString` by way of a C pointer; the pointer to a `CFString` has a type name, `CFStringRef` (a “reference to a `CFString`,” evidently). You work with a `CFString` in pure C, by calling functions.

You might, on occasion, actually have to work with a Core Foundation type even when a corresponding object type exists. For example, you might find that `NSString`, for all its power, fails to implement a needed piece of functionality, which is in fact available for a `CFString`. Luckily, an `NSString` (a value typed as `NSString*`) and a `CFString` (a value typed as `CFStringRef`) are interchangeable: you can use one where the other is expected, though you will have to typecast in order to quiet the worries of the compiler. The documentation describes this interchangeability by saying that `NSString` and `CFString` are “toll-free bridged” to one another.

To illustrate, I'll use a `CFString` to convert an `NSString` representing an integer to that integer. (This use of `CFString` is unnecessary, and is just by way of demonstrating the syntax; `NSString` has an `intValue` method.)

```
NSString* answer = @"42";
int ans = CFStringGetIntValue((CFStringRef)answer); // non-ARC
```

The typecast prevents the compiler from complaining, and works because NSString is toll-free bridged to CFString — in effect, behind the scenes, an NSString is a CFString.

Under ARC, however, that code won't compile unless you supply a little more information. ARC, as we'll see in [Chapter 12](#), is about memory management; but ARC manages only Objective-C objects, not their C counterparts. So although ARC manages the memory for an NSString, it leaves memory management for a CFStringRef up to you; and in order to compile that code, it needs you to show it that you understand the memory management status of this value as it crosses the toll-free bridge. You do so like this:

```
NSString* answer = @"42";
int ans = CFStringGetIntValue((__bridge CFStringRef)answer); // ARC
```

The extra qualifier `__bridge` means: "Don't worry, ARC, I know I'm crossing the toll-free bridge, and I assure you that this has no implications for memory management." On the other hand, there are situations where crossing the toll-free bridge *does* have implications for memory management, and you may rest assured that I'll discuss them in [Chapter 12](#).

The pointer-to-struct C data types, whose name typically ends in "Ref", may be referred to collectively as CFTyperef, which is actually just the generic pointer-to-void.. Thus, crossing the toll-free bridge may usefully be thought of as a cast between an object pointer and a generic pointer — that is, in general terms, from `id` to `void*` or from `void*` to `id`. Even where there is no toll-free bridging between *specific* types (as there is with NSString and CFString), there is always bridging at the top of the hierarchy, so to speak, between NSObject (the base object class, as explained in [Chapter 4](#)) and CFTyperef.



It is sometimes necessary to assign a CFTyperef to an `id` variable or parameter. For example, a CALayer's `setContents:` method expects an `id` parameter, but the actual value must be a CGImageRef. This is legal, because a pointer is just a pointer, but the compiler will complain unless you also typecast to an `id`.

Blocks

A *block* is an extension to the C language, introduced in Mac OS X 10.6 and available in iOS 4.0 or later. It's a way of bundling up some code and handing off that entire bundle as an argument to a C function or Objective-C method. This is similar to what we did at the end of the preceding section, handing off a pointer to a function as an argument, but instead we're handing off the code *itself*. The latter has some major advantages over the former, which I'll discuss in a moment.

As an example, I'll rewrite the earlier array-sorting example to use a block instead of a function pointer. Instead of calling `sortedArrayUsingFunction:context:`, I'll call `sortedArrayUsingComparator:`, which takes a block as its parameter. The block is typed like this:

```
NSComparisonResult (^)(id obj1, id obj2)
```

That's similar to the syntax for specifying the type of a pointer to a function, but a caret character is used instead of an asterisk character. So this means a block that takes two `id` parameters and returns an `NSComparisonResult` (which is merely an `NSInteger`, with just the same meaning as in the previous example). We can define the block and hand it off as the argument to `sortedArrayUsingComparator:` all in a single move, like this:

```
NSArray* arr2 = [arr sortedArrayUsingComparator: ^(id obj1, id obj2) {
    NSString* s1 = (NSString*) obj1;
    NSString* s2 = (NSString*) obj2;
    NSString* string1end = [s1 substringFromIndex:[s1 length] - 1];
    NSString* string2end = [s2 substringFromIndex:[s2 length] - 1];
    return [string1end compare:string2end];
}];
```

The syntax of the inline block definition is:

```
^①(id obj1, id obj2)② {③}
```

- ① First, the caret character.
- ② Then, parentheses containing the parameters.
- ③ Finally, curly braces containing the block's content.

Thanks to the block, as you can see, we've combined the definition of the callback function with its use. Of course, you might object that this means the callback isn't reusable; if we had *two* calls to `sortedArrayUsingComparator:` using the same callback, we'd have to write out the callback in full twice. To avoid such repetition, or simply for clarity, a block can be assigned to a variable:

```
NSComparisonResult (^sortByLastCharacter)(id, id) = ^(id obj1, id obj2) {
    NSString* s1 = (NSString*) obj1;
    NSString* s2 = (NSString*) obj2;
    NSString* string1end = [s1 substringFromIndex:[s1 length] - 1];
    NSString* string2end = [s2 substringFromIndex:[s2 length] - 1];
    return [string1end compare:string2end];
};
NSArray* arr2 = [arr sortedArrayUsingComparator: sortByLastCharacter];
NSArray* arr4 = [arr3 sortedArrayUsingComparator: sortByLastCharacter];
```



The return type in an inline block definition is usually omitted. If included, it *follows* the caret character, *not* in parentheses. If omitted, you may have to use typecasting in the `return` line to make the returned type match the expected type. For a complete technical syntax specification for blocks, see <http://clang.llvm.org/docs/BlockLanguageSpec.txt>.

The power of blocks really starts to emerge when they are used instead of a selector name. In an example earlier in this chapter, we talked about how you could pass `@selector(tickleMeElmo:)` as the second argument to `addObserver:selector:name:object:` as a way of saying, “When the time comes, please call my `tickleMeElmo:` method.” We also talked about how error-prone this syntax was: make a typing error, and your `tickleMeElmo:` method mysteriously won’t be called. Moreover, such code is hard to maintain; there’s the `tickleMeElmo:` method sitting there, completely separate from the code that calls `addObserver:selector:name:object:,` yet existing only to specify what should happen at the later time when our message arrives. (I’ll talk about this problem again in [Chapter 11](#).) It might well be clearer and more compact to call `addObserverForName:object:queue:usingBlock:` and specify there and then as a block what should happen at message time, with no separate method callback.

Variables in scope at the point where a block is defined keep their meaning within the block at that moment, even though the block may be executed at some later moment. (Technically, we say that a block is a *closure*.) It is this aspect of blocks that makes them useful for specifying functionality to be executed at some later time, or even in some other thread.

Variables in scope whose meaning is captured by the closure are protected from direct assignment from within the block, unless you deliberately turn off this protection. Thus, if code inside a block tries to assign directly to a variable whose meaning comes from outside the block, the compiler will prevent it. To turn off this protection, declare the variable using the `_block` qualifier. But of course if such a variable is an object reference, messages can be sent to it and the object may be mutated (because message sending is not assignment) even without the `_block` qualifier.

Objective-C Classes

This chapter describes some linguistic and structural features of Objective-C having to do with classes; in the next chapter, we'll do the same for instances.

Class and Superclass

In Objective-C, as in many other object-oriented languages, a mechanism is provided for specifying a relationship between two classes: they can be *subclass* and *superclass* of one another. For example, we might have a class Quadruped and a class Dog and make Quadruped the superclass of Dog. A class may have many subclasses, but a class can have only one immediate superclass. (I say “immediate” because that superclass might itself have a superclass, and so on in a rising chain, until we get to the ultimate superclass, called the *base class*, or *root class*.)

Because a class can have many subclasses but only one superclass, we can imagine all classes in a program as being arranged in a tree that splits into branches, such that each branch splits into smaller branches, each smaller branch splits into even smaller branches, and so on. Or we can imagine all the classes arranged in a hierarchy, such as might be displayed in an outline, with a single ultimate superclass, then all of its immediate subclasses in the next level below that, then each of *their* immediate subclasses in the next level below that, and so on. Indeed, before you write a line of your own code, Cocoa already consists of exactly such a vast repertoire of classes arranged in exactly such a hierarchical relationship. Xcode will actually display this relationship for you: choose View → Navigators → Show Symbol Navigator and click Hierarchical, with the first and third icons in the filter bar darkened ([Figure 4-1](#)).

The reason for the class–subclass relationship is to allow related classes to share functionality. Suppose, for example, we have a Dog class and a Cat class, and we are considering defining a `walk` method for both of them. We might reason that both a dog and a cat walk in pretty much the same way, by virtue of both being quadrupeds. So it might make sense to define `walk` as a method of the Quadruped class, and make both Dog and Cat subclasses of Quadruped. The result is that both Dog and Cat can be sent

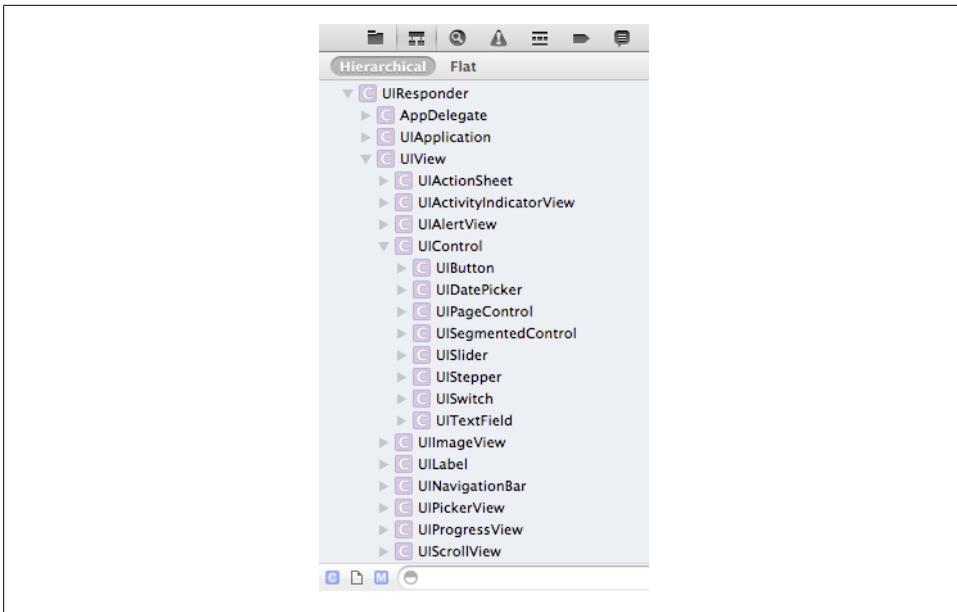


Figure 4-1. Browsing the built-in class hierarchy in Xcode 4

the `walk` message, even if neither of them has a `walk` method, because each of them has a superclass that *does* have a `walk` method. We say that a subclass *inherits* the methods of its superclass.

The purpose of subclassing is not merely so that a class can inherit another class's methods; it's so that it can define methods of its own. Typically, a subclass consists of the methods inherited from its superclass *and then some*. If `Dog` has no methods of its own, it is hard to see why it should exist separately from `Quadruped`. But if a `Dog` knows how to do something that not every `Quadruped` knows how to do — let's say, `bark` — then it makes sense as a separate class. If we define `bark` in the `Dog` class, and `walk` in the `Quadruped` class, and make `Dog` a subclass of `Quadruped`, then `Dog` inherits the ability to walk from the `Quadruped` class and also knows how to bark.

It is also permitted for a subclass to redefine a method inherited from its superclass. For example, perhaps some dogs bark differently from other dogs. We might have a class `NoisyDog`, for instance, that is a subclass of `Dog`. `Dog` defines `bark`, but `NoisyDog` also defines `bark`, and defines it differently from how `Dog` defines it. This is called *overriding*. The very natural rule is that if a subclass overrides a method inherited from its superclass, then when the corresponding message is sent to an instance of that subclass, it is the subclass's version of that method that is called.

Interface and Implementation

As you already know from [Chapter 2](#), all your code is going to go into some class or other. So the first thing we must do is specify what is meant by putting code “into a class” in Objective-C. How does Objective-C say, linguistically and structurally, “This is the code for such-and-such a class”?

To write the code for class, you must actually provide two chunks or sections of code, called the *interface* and the *implementation*. Here’s the complete minimum code required to define a class called MyClass. This class is so minimal that it doesn’t even have any methods of its own:

```
@interface MyClass  
@end  
@implementation MyClass  
@end
```

The `@interface` and `@implementation` compiler directives show the compiler where the interface and implementation sections begin for the class that’s being defined, MyClass; the corresponding `@end` lines show where each of those sections end.

In real life, the implementation section is where any methods for MyClass would be defined. So here’s a class that’s actually defined to do something:

```
@interface MyClass  
@end  
@implementation MyClass  
- (NSString*) sayGoodnightGracie {  
    return @"Good night, Gracie!";  
}  
@end
```

Observe how a method is defined. The first line is just like the method declaration, stating the type of method (class or instance), the type of value returned, and the name of the method along with the types of any parameters and local names for those parameters (see [Chapter 3](#)). Then come curly braces containing the code to be executed when the method is called, just as with a C function (see [Chapter 1](#)).

However, this class is still pretty much useless, because it can’t be instantiated. In Cocoa, knowledge of how to be instantiated, plus how to do a number of other things that any class should know how to do, resides in the base class, which is the NSObject class. Therefore, all Cocoa classes must be based ultimately on the NSObject class, by declaring as the superclass for your class either NSObject or some other class that inherits from NSObject (as just about any other Cocoa class does). The syntax for this declaration is a colon followed by the superclass name in the `@interface` line, like this:

```
@interface MyClass : NSObject
@end
@implementation MyClass
- (NSString*) sayGoodnightGracie {
    return @"Good night, Gracie!";
}
@end
```



NSObject is not the only Cocoa base class. It used to be, but there is now another, NSProxy. NSProxy is used only in very special circumstances and is not discussed in this book. If you have no reason for your class to inherit from any other class, make it inherit from NSObject.

In its fullest form, the interface section might contain some more material. In particular, if we want to declare our methods, those method declarations go into the interface section. Method declarations are not required, but without a method declaration, a method cannot be “seen” by other methods defined before it in the same class. A method declaration matches the name and signature for the method definition and ends with a semicolon (required):

```
@interface MyClass : NSObject
- (NSString*) sayGoodnightGracie;
@end
@implementation MyClass
- (NSString*) sayGoodnightGracie {
    return @"Good night, Gracie!";
}
@end
```

There are also instance variables to be considered. If our class is to have any instance variables (other than those inherited from its superclass), they must be declared (although, as we shall see in [Chapter 5](#) and [Chapter 12](#), it is possible in some circumstances to skip explicit declaration of a publicly accessible instance variable and declare it implicitly instead). Before iOS 5.0, such explicit declaration had to take place in curly braces at the start of the interface section:

```
@interface MyClass : NSObject {
    // instance variable declarations go here
}
- (NSString*) sayGoodnightGracie;
@end
@implementation MyClass
- (NSString*) sayGoodnightGracie {
    return @"Good night, Gracie!";
}
@end
```

However, starting in iOS 5.0, it is permitted to put instance variable declarations in curly braces at the start of the implementation section instead. This is a more logical place for variable declarations to go, because, as I'll explain in the next section, the

interface section is usually visible to other classes, but there is no reason why instance variables need to be visible to other classes, as they are usually private. Therefore, I prefer the new style, and will use it consistently throughout this book:

```
@interface MyClass : NSObject
- (NSString*) sayGoodnightGracie;
@end
@implementation MyClass {
    // instance variable declarations go here (starting in iOS 5)
}
- (NSString*) sayGoodnightGracie {
    return @"Good night, Gracie!";
}
@end
```

I'll go into more detail about instance variables in the next chapter.

Header File and Implementation File

It's perfectly possible for the interface and implementation of a class to appear in the same file, or for multiple classes to be defined in a single file, but this is not the usual convention. The usual convention is one class, two files: one file containing the interface section, the other file containing the implementation section. For example, let's suppose you are defining a class `MyClass`. Then you have two files, `MyClass.h` and `MyClass.m`. (The file naming is not magical or necessary; it's just part of the convention. The file extensions are pretty much necessary, though, because the build process and Xcode itself rely on them.) The interface section goes into `MyClass.h`, which is called the *header file*. The implementation section goes into `MyClass.m`, which is called the *implementation file*. This separation into two files is not inconvenient, because Xcode, expecting you to follow this convention, makes it easy to jump from editing a `.h` file to the corresponding `.m` file and *vice versa* (Navigate → Jump to Next Counterpart). Finally, the implementation file imports the header file (see [Chapter 1](#) on the `#import` directive); this effectively unites the full class definition, making the definition legal even though it is split between two files, and allowing the implementation section to "see" any method declarations in the interface section.

With this arrangement in place, further imports become easy to configure. The header file imports the basic header file for the entire Cocoa framework; in the case of an iOS program, that's `UIKit.h` (again, see [Chapter 1](#)). There is no need for the implementation file to import `UIKit.h`, because the header file imports it, and the implementation file imports the header file. If a class needs to know about another class that isn't already imported in this way, its implementation file imports that class's header file. [Example 4-1](#) summarizes this conventional schema.

Example 4-1. Conventional schema for defining a class

```
// [MyClass.h]

#import <UIKit/UIKit.h>

@interface MyClass : NSObject
- (NSString*) sayGoodnightGracie;
@end

// [MyClass.m]

#import "MyClass.h"
#import "OtherClass.h"

@implementation MyClass {
    // instance variable declarations go here
}
- (NSString*) sayGoodnightGracie {
    return @"Good night, Gracie!";
}
@end
```

The result of this arrangement is that everything has the right visibility. No file ever imports an implementation file; that way, what's inside a class's implementation file is private to that class. If something about a class needs to be public, such as a method that you want other classes to be able to call, it is declared in the header file, and other classes import that header file in their implementation files (as I do with *OtherClass.h* in [Example 4-1](#)); this keeps the chain of imports clear and simple.

A header file is also an appropriate place to define constants. In [Chapter 1](#), for example, I talked about the problem of mistyping the name of a notification or dictionary key, which is a literal NSString, and how you could solve this problem by defining a name for such a string:

```
#define MYKEY @"mykey"
```

The question then arises of where to put that definition. If only one class needs to know about it, the definition can go near the start of its implementation file (it doesn't need to be inside the implementation section). But if multiple classes need to know about this name, then a header file is an appropriate location; every implementation file that imports this header file will acquire the definition, and you can use the name **MYKEY** in that implementation file.

A slight problem arises when a header file needs to mention one of your other classes. Suppose, for example, that *MyClass* has a public method that takes or returns an instance of *MyOtherClass*, or that *MyClass* has an instance variable whose type is *MyOtherClass*. So *MyClass.h* needs to speak of *MyOtherClass**. But *MyClass.h* does not import *MyOtherClass.h*, so *MyClass.h* doesn't know about *MyOtherClass*, and the compiler will complain. To silence the compiler without violating the arrangement of imports (by importing *MyOtherClass.h* in the header file *MyClass.h*), use the **@class**

directive. The word `@class` is followed by a comma-separated list of class names, ending with a semicolon. So `MyClass.h` might start out like this:

```
#import <UIKit/UIKit.h>
@class MyOtherClass;
```

Then the interface section would follow, as before. The `@class` directive simply tells the compiler, “Don’t worry, `MyOtherClass` really is the name of a class.” That’s all the compiler needs to know in order to permit the mention of the type `MyOtherClass*` in the header file.

If, on the other hand, `MyClass` is to be a subclass of some other class, then `MyClass`’s header file must import that superclass’s header file (or some other header file that imports that superclass’s header file). Thus, for example, in [Example 4-1](#), `MyClass.h` imports `UIKit.h`; thus it knows about `NSObject`, so that `MyClass` can declare `NSObject` as its superclass.

A question that may occur to you at this point is how to declare a method without making it public. For example, let’s say that many methods in `MyClass` need to call `myCoolMethod`, which is also a `MyClass` method. To make `myCoolMethod` visible to all `MyClass` methods, regardless of the order in which they are defined in the implementation section, you can declare `myCoolMethod` in the interface section. But this effectively “publishes” `myCoolMethod`, because any other class that imports `MyClass.h` will now know about it. If that isn’t something you want to do, there’s a trick for creating a second interface section that only `MyClass`’s implementation section can see ([Chapter 10](#)).



The Global Namespace

When defining classes, choose your class names wisely to prevent name collisions. Objective-C has no namespaces; there’s a single vast namespace containing all names. You don’t want your own class name (or, for that matter, any other top-level constant name) to match a name defined in Cocoa. Instead of namespaces, there’s a convention: each Cocoa framework prefixes its names with a particular pair of capital letters (`NSString` and `NSArray`, `CGFloat` and `CGRect`, and so on). Apple suggests that you use a prefix of your own as well; in fact, when you create a new project in Xcode, you’re offered an opportunity to specify a prefix, which will appear before the automatically created class names. Don’t use any of Apple’s prefixes. Nothing limits your prefix to two letters, or requires that both letters be uppercase. In fact, because all of Apple’s own prefixes *are* two uppercase letters, “`My`” as a prefix is safe.

Class Methods

Class methods are useful in general for two main purposes:

Cocoa's Own Header Files

The Cocoa classes themselves also follow the convention described in [Example 4-1](#): each class is separated into a header file (containing the interface) and an implementation file. However, the Cocoa class implementation files are not visible to you. This is one of the major limitations of Cocoa; unlike many programming frameworks, you can't see the source code for Cocoa — it's secret. To figure out how Cocoa works, you have to rely purely on the documentation (and experimentation). You can, however, see the Cocoa header files, and indeed you are expected to look at them, as they can be a useful form of documentation (see [Chapter 8](#)).

Factory methods

A factory method is a method that dispenses an instance of that class. For example, the `UIFont` class has a class method `fontWithName:Size:`. You supply a name and a size, and the `UIFont` class hands you back a `UIFont` instance corresponding to a font with that name and size.

Global utility methods

Classes are global (visible from all code), so a class is a good place to put a utility method that anyone might need to call and that doesn't require the overhead of an instance. For example, the `UIFont` class has a class method `familyNames`. It returns an array of strings (that is, an `NSArray` of `NSString` instances) consisting of the names of the font families installed on this device. Because this method has to do with fonts, the `UIFont` class is as good a place as any to put it.

Most methods that you write will be instance methods, but now and then you might write a class method. When you do, your purpose will probably be similar to those examples.

The Secret Life of Classes

A class method may be called by sending a message directly to the name of a class. For example, the `familyNames` class method of `UIFont` that I mentioned a moment ago might be called like this:

```
NSArray* fams = [UIFont familyNames];
```

Clearly, this is possible because a class is an object ([Chapter 2](#)), and the name of the class here represents that object.

You don't have to do anything to create a class object. One class object for every class your program defines is created for you automatically as the program starts up. (This includes the classes your program imports, so there's a `MyClass` class object because you defined `MyClass`, and there's an `NSString` class object because you imported `UI-`

Kit.h and the whole Cocoa framework.) It is to this class object that you’re referring when you send a message to the name of the class.

Your ability to send a message directly to the bare name of a class is due to a kind of syntactic shorthand. You can use the bare class name only in two ways (and we already know about both of them):

To send a message to

In the expression `[UIFont familyNames]`, the bare name `UIFont` is sent the `familyNames` message.

To specify an instance type

In the expression `NSString*`, the bare name `NSString` is followed by an asterisk to signify a pointer to an instance of this class.

Otherwise, to speak of a class object, you need to obtain that object formally. One way to do this is to send the `class` message to a class or instance. For example, `[MyClass class]` returns the actual class object. Some built-in Cocoa methods expect a class object parameter (whose type is described as `Class`). To supply this as an argument, you’d need to obtain a class object formally. Take, for example, introspection on an object to inquire what its class is. The `isKindOfClass:` instance method is declared like this:

```
- (BOOL)isKindOfClass:(Class)aClass
```

So that means you could call it like this:

```
if ([someObject isKindOfClass: [MyClass class]]) // ...
```

A class object is not an instance, but it is definitely a full-fledged object. Therefore, a class object can be used wherever an object can be used. For example, it can be assigned to a variable of type `id`:

```
id classObject = [MyClass class];
```

You could then call a class method by sending a message to that object, because it is the class object:

```
id classObject = [MyClass class];
[classObject someClassMethod];
```

All class objects are also members of the `Class` class, so you could say this:

```
Class classObject = [MyClass class];
[classObject someClassMethod];
```


Objective-C Instances

Instances are the heart of the action in an Objective-C program. Most of the methods you'll define when creating your own classes will be instance methods; most of the messages you'll send in your code will call instance methods. This chapter describes how instances come into existence and how they work.

How Instances Are Created

Your class objects are created for you automatically as your program starts up, but instances must be created deliberately as the program runs. The entire question of where instances come from is thus crucial. Ultimately, every instance comes into existence in just one way: someone turns to a class and asks that class to instantiate itself. But there are three different ways in which this can occur: ready-made instances, instantiation from scratch, and nib-based instantiation.

Ready-Made Instances

One way to create an instance is indirectly, by calling code that does the instantiation for you. You can think of an instance obtained in this indirect manner as a “ready-made instance.” (That’s my made-up phrase, not an official technical term.) For example, consider this simple code:

```
NSString* s2 = [s uppercaseString];
```

The documentation for the `NSString` instance method `uppercaseString` says that it returns an `NSString*` that is “an uppercased representation of the receiver.” In other words, you send the `uppercaseString` message to an `NSString`, and you get back a *different*, newly created `NSString`. After that line of code, `s2` points to an `NSString` instance that didn’t exist beforehand.

The `NSString` produced by the `uppercaseString` method is a ready-made `NSString` instance. Your code didn’t say anything about instantiation; it just sent the `uppercaseString` message. But clearly *someone* said something about instantiation, because in-

stantiation took place; this is a newly minted `NSString` instance. That someone is presumably some code inside the `NSString` class. But we don't have to worry about the details. We are guaranteed of receiving a complete ready-made ready-to-roll `NSString`, and that's all we care about.

Similarly, any class factory method instantiates the class and dispenses the resulting instance as a ready-made instance. So, for example, the `NSString` class method `stringWithContentsOfFile:encoding:error:` reads a file and produces an `NSString` representing its contents. All the work of instantiation has been done for you. You just accept the resulting string and away you go.



A Cocoa class factory method is likely to have its return value typed as `id`. As I mentioned in [Chapter 3](#), this can lead to trouble later if you mistakenly assign the resulting instance where a different class of object is expected; the compiler doesn't complain (because `id` is the universal donor) but you can mysteriously crash later when the wrong message is sent to the instance.

Not every method that returns an instance returns a new instance, of course. For example, this is how you ask an array (an `NSArray`) for its last element:

```
id last = [myArray lastObject];
```

The `NSArray` `myArray` didn't *create* the object that it hands you. That object already existed; `myArray` was merely containing it, as it were — it was holding the object, pointing to it. Now it's sharing that object with you, that's all.

Similarly, many classes dispense one particular object. For example, your app has exactly one instance of the `UIApplication` class (we call this the *singleton* `UIApplication` instance); to access it, you send the `sharedApplication` class method to the `UIApplication` class:

```
UIApplication* theApp = [UIApplication sharedApplication];
```

This singleton instance existed before you asked for it; indeed, it existed before any code of yours could possibly run. You don't care how it was brought into being; all you care is that you can get hold of it when you want it. I'll talk more about globally available singleton objects of this kind in [Chapter 13](#).

Instantiation from Scratch

The alternative to requesting a ready-made instance is to tell a class, yourself, directly, to instantiate itself. There is basically one way to do this: you send a class the `alloc` message. The `alloc` class method is implemented by the `NSObject` class, the root class from which all other classes inherit. It causes memory to be set aside for the instance so that an instance pointer can point to it. (Management of that memory is a separate issue, discussed in [Chapter 12](#).)

You must never, never, *never* call `alloc` by itself. You must *immediately* call another method, an instance method that *initializes* the newly created instance, placing it into a known valid state so that it can be sent other messages. Such a method is called an *initializer*. Moreover, an initializer returns an instance — usually the same instance, initialized. Therefore you can, and always should, call `alloc` and the initializer in the same line of code. The minimal initializer is `init`. So the basic pattern, known informally as “alloc-init,” looks like [Example 5-1](#).

Example 5-1. The basic pattern for instantiation from scratch

```
SomeClass* aVariable = [[SomeClass alloc] init];
```

You cannot instantiate from scratch if you do not also know how to initialize, so we turn immediately to a discussion of initialization.

Initialization

Every class defines (or inherits) at least one initializer. This is an instance method; the instance has just been created (by calling `alloc` on the class), and it is to this newly minted instance that the initializer message must be sent. An initialization message must be sent to an instance immediately after that instance is created by means of the `alloc` message, and it must not be sent to an instance at any other time.

The basic initialization pattern, as shown in [Example 5-1](#), is to nest the `alloc` call in the initialization call, assigning the result of the *initialization* (not the `alloc!`) to a variable. One reason for this is that if something goes wrong and the instance can't be created or initialized, the initializer will return `nil`; therefore it's important to capture the result of the initializer and treat that, not the result of `alloc`, as the pointer to the instance.

To help you identify initializers, all initializers are named in a conventional manner. The convention is that all initializers, and only initializers, begin with the word `init`. The ultimate bare-bones initializer is called simply `init`, and takes no parameters. Other initializers do take parameters, and usually begin with the phrase `initWith` followed by descriptions of their parameters. For example, the `NSArray` class documentation lists these methods:

- `initWithArray:`
- `initWithArray:copyItems:`
- `initWithContentsOfFile:`
- `initWithContentsOfURL:`
- `initWithObjects:`
- `initWithObjects:count:`

Let's try a real example. A particularly easy and generally useful initializer for `NSArray` is `initWithObjects:`. It takes a list of objects; the list must be terminated by `nil`. In [Chapter 3](#), we illustrated this by creating an `NSArray` from three strings, by means of a class factory method that returned a ready-made instance:

```
NSArray* pep = [NSArray arrayWithObjects:@"Manny", @"Moe", @"Jack", nil];
```

Now we'll do what amounts to exactly the same thing, except that we'll create the instance ourselves, from scratch:

```
NSArray* pep = [[NSArray alloc] initWithObjects:@"Manny", @"Moe", @"Jack", nil];
```

In that particular case, there exist both a factory method and an initializer that work from the same set of data. Ultimately, it makes no difference which you use; given the same arguments, both approaches result in `NSArray*` instances that are indistinguishable from one another. It will turn out in the discussion of memory management ([Chapter 12](#)) that there might be a reason to choose instantiation from scratch over ready-made instances.

In looking for an initializer, don't forget to look upward through the class hierarchy. For example, the class documentation for `UIWebView` lists no initializers, but `UIWebView` inherits from `UIView`, and in `UIView`'s class documentation you'll discover `initWithFrame:`. Moreover, the `init` method is defined as an instance method of the `NSObject` class, so every class inherits it and every newly minted instance can be sent the `init` message. Thus it is a given that if a class defines no initializers of its own, you can initialize an instance of it with `init`. For example, the `UIResponder` class documentation lists no initializers at all (and no factory methods). So to create a `UIResponder` instance from scratch, you'd call `alloc` and `init`.



In just the single case where `init` is the initializer you want to call, you can collapse the successive calls to `alloc` and `init` with a call to `new`. In other words, `[MyClass new]` is a synonym for `[[MyClass alloc] init]`.

The designated initializer

If a class does define initializers, one of them may be described in the documentation as the *designated initializer*. (There's nothing about a method's name that tells you it's the designated initializer; you must peruse the documentation to find out.) For example, in the `UIView` class documentation, the `initWithFrame:` method is described as the designated initializer. A class that does not define a designated initializer inherits its designated initializer; the ultimate designated initializer, inherited by all classes without any other designated initializer anywhere in their superclass chain, is `init`.

The designated initializer is the initializer on which any other initializers depend, in this class or any subclasses: ultimately, they *must* call it. The designated initializer might have the most parameters, allowing the most instance variables to be set explicitly, with the other initializers supplying default values for some instance variables, for convenience. Or it might just be the most basic form of initialization. But in any case, it is a bottleneck through which all other initializers pass. Here are some examples:

- The `NSDate` class documentation says that `initWithTimeIntervalSinceReferenceDate:` is the designated initializer, and that `initWithTimeIntervalSinceNow:` calls it.

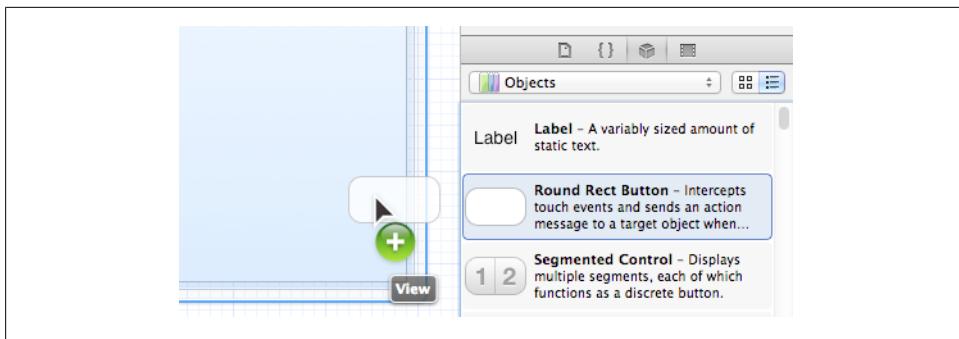


Figure 5-1. Dragging a button into a view

- The UIView class documentation says that `initWithFrame:` is the designated initializer. UIView contains no other initializers, but some of its subclasses do. UIWebView, a UIView subclass, has no initializer, so `initWithFrame:` is its inherited designated initializer. UIImageView, a UIView subclass, has initializers such as `initWithImage:`, but none of them is a designated initializer; so `initWithFrame:` is its inherited designated initializer as well, and `initWithImage:` must call `initWithFrame:`.

Moreover, a class that implements a designated initializer will override the designated initializer inherited from its superclass. The idea is typically that even the inherited designated initializer, if called, will call this class's designated initializer. For example, UIView overrides the inherited `init` to call its own designated initializer, `initWithFrame:`, with a value of `(CGRect){{0, 0}, {0, 0}}`.

Nib-Based Instantiation

The third means of instantiation is through a nib file (or storyboard file). A nib file is where Xcode lets you “draw” parts of the user interface. Most Xcode projects will include at least one nib file, which will be built into the app bundle, and will then be loaded as the app runs. A nib file consists, in a sense, of the names of classes along with instructions for instantiating and initializing them. When the app runs and a nib file is loaded, those instructions are carried out — those classes *are* instantiated and initialized.

For example, suppose you'd like the user to be presented with a view containing a button whose title is “Howdy.” Xcode lets you arrange this graphically by editing a nib file: you drag a button from the Object library into the view, place it at a certain position in the view, and then set its title to “Howdy” (Figure 5-1). In effect, you create a drawing of what you want the view and its contents to look like.

When the app runs, the nib file loads, and that drawing is turned into reality. To do this, the drawing is treated as a set of instructions for instantiating objects. The button that you dragged into the view is treated as a representative of the UIButton class. The

`UIButton` class is told to instantiate itself, and that instance is then initialized, giving it the same position you gave it in the drawing (the instance's `frame`), the same title you gave it in the drawing (the instance's `title`), and putting it into the window. In effect, the loading of your nib file is equivalent to this code (assuming that `view` is a reference to the view object):

```
UIButton* b =
    [UIButton buttonWithType:UIButtonTypeRoundedRect]; // factory method, instantiate
[b setTitle:@"Howdy!" forState:UIControlStateNormal]; // set up title
[b setFrame: CGRectMake(100,100,100,35)]; // set up frame
[view addSubview:b]; // place button in view
```

The fact that nib files are a source of instances, and that those instances are brought into existence as the nib file is loaded, is a source of confusion to beginners. I'll discuss nib files and how they are used to generate instances in much more detail in [Chapter 7](#).

Polymorphism

The compiler, even in the world of static typing, is perfectly happy for you to supply a subclass instance where a superclass type is declared. To see this, let's start with the first line of the previous example:

```
UIButton* b = [UIButton buttonWithType:UIButtonTypeRoundedRect];
```

`UIButton` is a subclass of `UIControl`, which is a subclass of `UIView`. So it would be perfectly legal and acceptable to say this:

```
UIButton* b = [UIButton buttonWithType:UIButtonTypeRoundedRect];
UIView* v = b;
```

The variable `b` is a `UIButton` instance, but I'm assigning it to a variable declared as a `UIView`. That's legal and acceptable because `UIView` is an ancestor (up the superclass chain) of `UIButton`. Putting it another way, I'm behaving as if a `UIButton` were a `UIView`, and the compiler accepts this because a `UIButton` *is* a `UIView`.

What's important when the app runs, however, is not the declared class of a variable, but the actual class of the object to which that variable points. Even if I assign the `UIButton` instance `b` to a `UIView` variable `v`, the object to which the variable `v` points is still a `UIButton`. So I can send it messages appropriate to a `UIButton`. For example:

```
UIButton* b = [UIButton buttonWithType:UIButtonTypeRoundedRect];
UIView* v = b;
[v setTitle:@"Howdy!" forState:UIControlStateNormal];
```

That code will cause the compiler to complain, because `UIView` doesn't implement `setTitle:forState:`; under ARC, in fact, that code won't even compile. So I'll calm the compiler's fears by typecasting:

```
UIButton* b = [UIButton buttonWithType:UIButtonTypeRoundedRect];
UIView* v = b;
[(UIButton*)v setTitle:@"Howdy!" forState:UIControlStateNormal];
```

The typecast calms the compiler's fears, but the important thing is what happens when the program runs. What happens is that this code works just fine! It works fine not because I typecast `v` to a `UIButton` (typecasting doesn't magically convert anything to anything else; it's just a hint to the compiler), but because `v` really *is* a `UIButton`. So when the message `setTitle:forState:` arrives at the object pointed to by `v`, everything is fine. (If `v` had been a `UIView` but not a `UIButton`, on the other hand, the program would have crashed at that moment.)

An object, then, responds to a message sent to it on the basis of what it really is, not on the basis of anything said about what it is — and what it really is cannot be known until the program actually runs and the message is actually sent to that object.

Now let's turn the tables. We called a `UIButton` a `UIView` and sent it a `UIButton` message. Now we're going to call a `UIButton` a `UIButton` and send it a `UIView` message.

What an object really is depends not just upon its class but also upon that class's inheritance. A message is acceptable even if an object's own class doesn't implement a corresponding method, provided that the method is implemented somewhere up the superclass chain. For example, returning again to the same code:

```
UIButton* b = [UIButton buttonWithType:UIButtonTypeRoundedRect];
[b setFrame: CGRectMake(100,100,100,35)];
```

This code works fine, too. But you won't find `setFrame:` in the documentation for the `UIButton` class. That's because you're looking in the wrong place. A `UIButton` is a `UIControl`, and a `UIControl` is a `UIView`. To find out about `setFrame:`, look in the `UIView` class's documentation. (Okay, it's more complicated than that; you won't find `setFrame:` there either. But you will find a term `frame` which is called a "property," and this amounts to the same thing, as I'll explain later in this chapter.) So the `setFrame:` message is sent to a `UIButton`, but it corresponds to a method defined on a `UIView`. Yet it works fine, because a `UIButton` *is* a `UIView`.



A common beginner mistake is to consult the documentation without following the superclass chain. If you want to know what you can say to a `UIButton`, don't just look in the `UIButton` class documentation: also look in the `UIControl` class documentation, the `UIView` class documentation, and so on.

To sum up: we treated a `UIButton` object as a `UIView`, yet we were still able to send it a `UIButton` message. We treated a `UIButton` as a `UIButton`, yet we were still able to send it a `UIView` message. What matters when a message is sent to an object is not how the variable pointing to that object is declared but what class the object really is. What an object really is depends upon its class, along with that class's inheritance from the superclass chain; these facts are innate to the object and are independent of how the variable pointing to the object presents itself to the world. This independent maintenance of object type integrity is the basis of what is called *polymorphism*.

But it is not quite the whole of polymorphism. To understand the whole of polymorphism, we must go further into the dynamics of message sending.

The Keyword self

A common situation is that code in an instance method defined in a class must call another instance method defined within the same class. We have not yet discussed how to do this. A method is called by sending a message to an object; in this situation, what object would that be? The answer is supplied by a special keyword, `self`. Here's a simple example:

```
@implementation MyClass  
- (NSString*) greeting {  
    return @"Goodnight, Gracie!";  
}  
  
- (NSString*) sayGoodnightGracie {  
    return [self greeting];  
}  
  
@end
```

When the `sayGoodnightGracie` message is sent to a `MyClass` instance, the `sayGoodnightGracie` method runs. It sends the `greeting` message to `self`. As a result, the `greeting` instance method is called; it returns the string `@"Goodnight, Gracie!"`, and this same string is then returned from the `sayGoodnightGracie` method.

The example seems straightforward enough, and it is. In real life, your code when you define a class will sometimes consist of a few public instance methods along with lots of other instance methods on which they rely. The instance methods within this class will be calling each other constantly. They do this by sending messages to `self`.

Behind this simple example, though, is a subtle and important mechanism having to do with the real meaning of the keyword `self`. The keyword `self` does not actually mean “in the same class.” It’s an instance, after all, not a class. What instance? It’s this same instance. The same as what? The same instance to which the message was sent that resulted in the keyword `self` being encountered in the first place.

So let’s consider in more detail what happens when we instantiate `MyClass` and send the `sayGoodnightGracie` message to that instance:

```
MyClass* thing = [[MyClass alloc] init];  
NSString* s = [thing sayGoodnightGracie];
```

We instantiate `MyClass` and assign the instance to a variable `thing`. We then send the `sayGoodnightGracie` message to `thing`, the instance we just created. The message arrives, and it turns out this instance is a `MyClass`. Sure enough, `MyClass` implements a `sayGoodnightGracie` method, and this method is called. As it runs, the keyword `self` is encountered. It means “the instance to which the original message was sent in the first

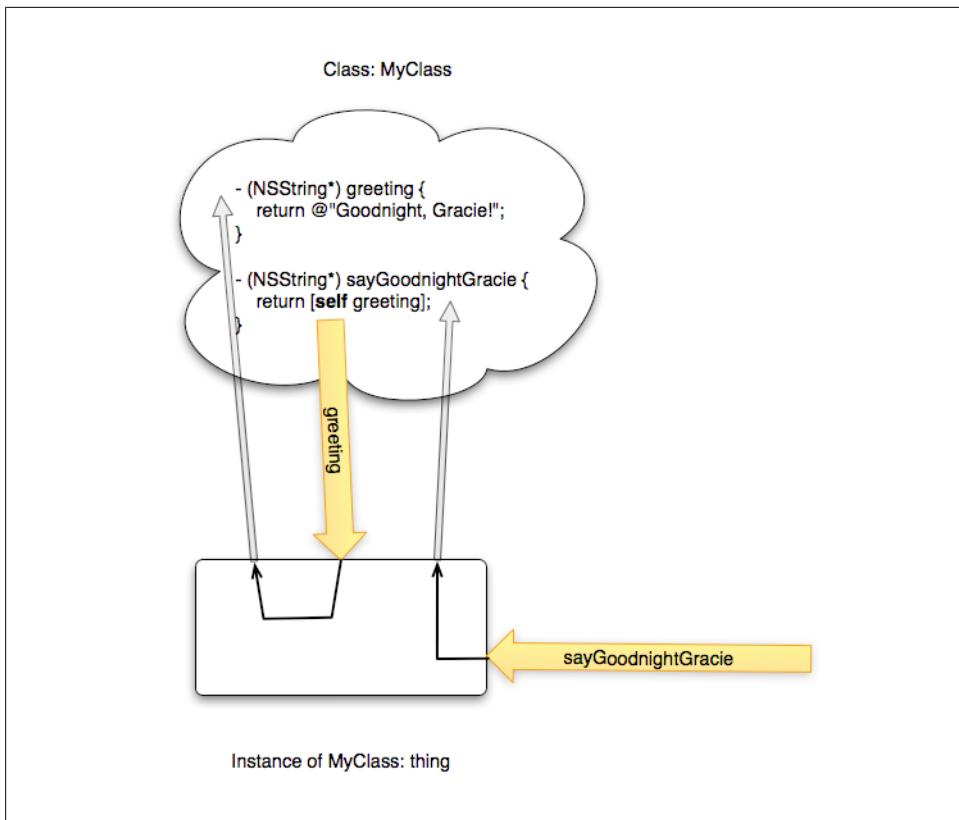


Figure 5-2. The meaning of self

place.” That, as it happens, is the instance pointed to by the variable `thing`. So now the `greeting` message is sent to that instance ([Figure 5-2](#)).

This mechanism may seem rather elaborate, considering that the outcome is just what you’d intuitively expect. But the mechanism *needs* to be elaborate in order to get the right outcome. This is particularly evident when superclasses are involved and a class overrides a method of its superclass. To illustrate, suppose we have a class `Dog` with an instance method `bark`. And suppose `Dog` also has an instance method `speak`, which simply calls `bark`. Now suppose we subclass `Dog` with a class `Basenji`, which overrides `bark` (because Basenjis can’t bark). What happens when we send the `speak` message to a `Basenji` instance, as in [Example 5-2](#)?

Example 5-2. Polymorphism in action

`@implementation Dog`

```

- (NSString*) bark {
    return @"Woof!";
}

```

```

- (NSString*) speak {
    return [self bark];
}

@end

@implementation Basenji : Dog

- (NSString*) bark {
    return @""; // empty string, Basenjis can't bark
}

@end

// [so, in some other class...]

Basenji* b = [[Basenji alloc] init];
NSString* s = [b speak];

```

If the keyword `self` meant “the same class where this keyword appears,” then when we send the `speak` message to a Basenji instance, we would arrive at the implementation of `speak` in the `Dog` class, and the `Dog` class’s `bark` method would be called. This would be terrible, because it would make nonsense of the notion of overriding; we’d return `@"Woof!"`, which is wrong for a Basenji. But that is *not* what the keyword `self` means. It has to do with the instance, not the class.

So here’s what happens. The `speak` message is sent to our Basenji instance, `b`. The Basenji class doesn’t implement a `speak` method, so we look upward in the class hierarchy and discover that `speak` is implemented in the superclass, `Dog`. We call `Dog`’s instance method `speak`, the `speak` method runs, and the keyword `self` is encountered. It means “the instance to which the original message was sent in the first place.” That instance is still our Basenji instance `b`. So we send the `bark` message to the Basenji instance `b`. The Basenji class implements a `bark` instance method, so this method is found and called, and the empty string is returned ([Figure 5-3](#)).

Of course, if the Basenji class had *not* overridden `bark`, then when the `bark` message was sent to the Basenji instance, we would have looked upward in the class hierarchy *again* and found the `bark` method implemented in the `Dog` class and called that. Thus, thanks to the way the keyword `self` works, inheritance works correctly both when there is overriding and when there is not.

If you understand that example, you understand polymorphism. The mechanism I’ve just described is crucial to polymorphism and is the basis of object-oriented programming. (Observe that I now speak of object-oriented programming, not just object-based programming as in [Chapter 2](#). That’s because, in my view, the addition of polymorphism is what turns object-based programming into object-oriented programming.)

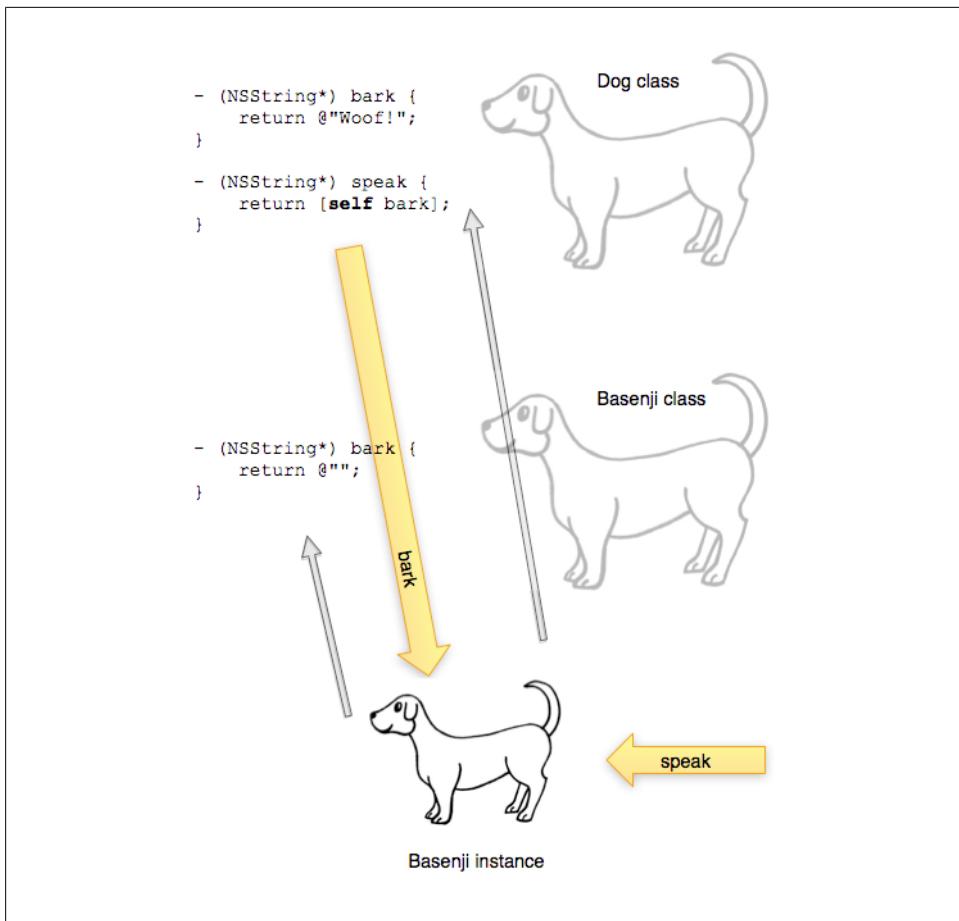


Figure 5-3. Class inheritance, overriding, self, and polymorphism

The Keyword super

Sometimes (quite often, in Cocoa programming) you want to override an inherited method but still access the overridden functionality. To do so, you'll use the keyword `super`. Like `self`, the keyword `super` is something you send a message to. But its meaning has nothing to do with “this instance” or any other instance. The keyword `super` is class-based, and it means: “Start the search for messages I receive in the superclass of this class” (where “this class” is the class where the keyword `super` appears).

You can do anything you like with `super`, but its primary purpose, as I've already said, is to access overridden functionality — typically from within the very functionality that does the overriding, so as to get both the overridden functionality and some additional functionality.

For example, suppose we define a class `NoisyDog`, a subclass of `Dog`. When told to bark, it barks twice:

```
@implementation NoisyDog : Dog  
- (NSString*) bark {  
    return [NSString stringWithFormat: @"%@ %@", [super bark], [super bark]];  
}  
  
@end
```

That code calls `super`'s implementation of `bark`, twice; it assembles the two resulting strings into a single string with a space between, and returns that (using the `stringWithFormat:` method). Because `Dog`'s `bark` method returns `@"Woof!"`, `NoisyDog`'s `bark` method returns `@"Woof! Woof!"`. Notice that there is no circularity or recursion here: `NoisyDog`'s `bark` method will never call itself.

A nice feature of this architecture is that by sending a message to the keyword `super`, rather than hard-coding `@"Woof!"` into `NoisyDog`'s `bark` method, we ensure maintainability: if `Dog`'s `bark` method is changed, the result of `NoisyDog`'s `bark` method will change to match. For example, if we later go back and change `Dog`'s `bark` method to return `@"Arf!"`, `NoisyDog`'s `bark` method will return `@"Arf! Arf!"` with no further change on our part.

In real Cocoa programming, it will very often be Cocoa's own methods that you're overriding. For example, the `UIViewController` class, which is built into Cocoa, implements a method `viewDidAppear:`, defined as follows:

```
- (void)viewDidAppear:(BOOL)animated
```

The documentation says that `UIViewController` is a class for which you are very likely to define a subclass (so as to get all of `UIViewController`'s mighty powers — we'll find out what they are in [Chapter 19](#) — along with your own custom behavior). The documentation proceeds to suggest that in your subclass of `UIViewController` you might want to override this method, but cautions that if you do, “you must call `super` at some point in your implementation.” The phrase “call `super`” is a kind of shorthand, meaning “pass on to `super` the very same call and arguments that were sent to you.” So your own implementation might look like this:

```
@implementation MyViewController : UIViewController  
// ...  
- (void)viewDidAppear: (BOOL) animated {  
    [super viewDidAppear: animated];  
    // ... do more stuff here ...  
}
```

The result is that when `viewDidAppear:` is called in a `MyViewController` instance, we do both the standard stuff that its superclass `UIViewController` does in response to `viewDidAppear:` and the custom stuff pertaining our own class `MyViewController`. In this particular case, we don't even know exactly what the `UIViewController` stuff is,

and we don't care. When the documentation tells you to call `super` when overriding, call `super` when overriding!

Instance Variables and Accessors

In [Chapter 3](#), I explained that one of the main reasons there are instances and not just classes is that instances can have instance variables. Instance variables, you remember, are declared when you define the class, and in [Chapter 4](#) I said that these declarations go into the curly-braces part of the class's interface section or, starting in iOS 5, its implementation section. But the value of an instance variable differs for each instance.



The term "instance variable" arises so often that it is often abbreviated to *ivar*. I'll use both terms indiscriminately from now on.

Let's write a class that uses an instance variable. Suppose we have a Dog class and we want every Dog instance to have a number, which should be an int. (For example, this number might correspond to the dog's license number, or something like that.) So the interface section for the Dog class might look like this:

```
@interface Dog : NSObject {  
    int number;  
}  
// public method declarations go here  
@end
```

Or, if we're using iOS 5's new convention (and I propose to do so henceforward), we could declare `number` in the implementation section for the Dog class, like this:

```
@implementation Dog {  
    // ivars can now be declared in the implementation section  
    int number;  
}  
// method implementations go here  
@end
```

(You might ask why, for this example, I don't use instead the concept of giving the dog a name. The reason is that a name would be an `NSString` instance, which is an object; instance variables that are pointers to objects raise some additional issues I don't want to discuss just now. But instance variables that are simple C data types raise no such issues. We'll return to this matter in [Chapter 12](#).)

By default, instance variables are *protected*, meaning that other classes (except for subclasses) can't see them. So if, somewhere else, I instantiate a Dog, I won't be able to access that Dog instance's `number` instance variable. This is a deliberate feature of Objective-C; you can work around it if you like, but in general you should not. Instead,

if you want to provide public access to an instance variable, write an accessor method and make the method declaration public.

Within a class, on the other hand, that class’s own instance variables are global. Any Dog method can just use the variable name `number` and access this instance variable, just like any other variable. But global variables can be confusing when you’re reading code; suddenly there’s a variable called `number` and you don’t understand what it is, because there’s no declaration for it (the declaration is stashed away in the interface section, which is in a different file). So I often use a different notation, like this: `self->ivarName`. The “arrow” operator, formed by a minus sign and a greater-than sign, is called the *structure pointer* operator, because of its original use in C (K&R 6.2).

So let’s write, in Dog’s implementation section, a method that allows setting a value for the `number` ivar:

```
- (void) setNumber: (int) n {
    self->number = n;
}
```

Of course, we must also declare `setNumber:` in Dog’s interface section:

```
@interface Dog : NSObject
- (void) setNumber: (int) n;
@end
```

We can now instantiate a Dog and assign that instance a number:

```
Dog* fido = [[Dog alloc] init];
[fido setNumber: 42];
```

We can now set a Dog’s number, but we can’t get it (from outside that Dog instance). To correct this problem, we’ll write a second accessor method, one that allows for getting the value of the `number` ivar:

```
- (int) number {
    return self->number;
}
```

Again, we declare the `number` method in Dog’s interface section. (You’re not going to be confused, are you, by the fact that Dog has both a `number` method and a `number` instance variable? This doesn’t confuse the compiler, because they are used in completely different ways in code, so it shouldn’t confuse you either.) Now we can both set and get a Dog instance’s `number`:

```
Dog* fido = [[Dog alloc] init];
[fido setNumber: 42];
int n = [fido number];
// sure enough, n is now 42!
```

This architecture is very typical. Your class can have as many ivars as you like, but if you want them to be publicly accessible, you must provide accessor methods. Luckily, Objective-C 2.0 — which is what you’re using to program for iOS — provides a mechanism for generating accessor methods automatically (discussed in [Chapter 12](#)), so you

won't have to go through the tedium of writing them by hand every time you want to make an ivar publicly accessible. (Though, to be honest, I don't see why you shouldn't have to go through that tedium; before Objective-C 2.0, we all had to, so why shouldn't you? We also had to clean the roads with our tongues on the way to school. And we liked it! You kids today, you don't know what real programming is.)

Key–Value Coding

Objective-C provides a means for translating from a string to an instance variable accessor, called *key–value coding*. Such translation is useful, for example, when the name of the desired instance variable will not be known until runtime. So instead of calling `[fido number]`, we might have a string `@"number"` that tells us what accessor to call. This string is the “key.” The key–value coding equivalent of calling a getter is `valueForKey:`; the equivalent of calling a setter is `setValue:forKey:`.

Thus, for example, suppose we wish to get the value of the `number` instance variable from the `fido` instance. We can do this by sending `valueForKey:` to `fido`. However, even though the `number` instance variable is an `int`, the value returned by `valueForKey:` is an object — in this case, an `NSNumber`, the object equivalent of a number (see [Chapter 10](#)). If we want the actual `int`, `NSNumber` provides an instance method, `intValue`, that lets us extract it:

```
NSNumber* num = [fido valueForKey: @"number"];
int n = [num intValue];
```

Similarly, to use key–value coding to set the value of the `number` instance variable in the `fido` instance, we would say:

```
NSNumber* num = [NSNumber numberWithInt:42];
[fido setValue: num forKey: @"number"];
```

In this case there is no advantage to using key–value coding over just calling the accessors. But suppose we had received the value `@"number"` in a variable (as the result of a method call, perhaps). Suppose that variable is called `something`. Then we could say:

```
id result = [fido valueForKey: something];
```

Thus we could access a different instance variable under different circumstances. This powerful flexibility is possible because Objective-C is such a dynamic language that a message to be sent to an object does not have to be formed until the program is already running.

When you call `valueForKey:` or `setValue:forKey:`, the correct accessor method is called if there is one. Thus, when we use `@"number"` as the key, a `number` method and a `setNumber:` method are called if they exist. (This is one reason why your accessors should be properly named.) On the other hand, if there isn't an accessor method, the instance variable is accessed directly. Such direct access violates the privacy of instance variables, so there's a way to turn off this feature for a particular class if you don't like it. (I'll explain what it is, with more about key–value coding, in [Chapter 12](#).)

Properties

A *property* is a syntactical feature of Objective-C 2.0 designed to provide an alternative to the standard syntax for calling an instance variable accessor. In other words, a property is merely syntactic sugar for calling an accessor method. I'll use the Dog class as an example. If the Dog class has a getter method called `number` and a setter method called `setNumber:`, then the Dog class might also declare a `number` property. If it does, then instead of saying things like this:

```
[fido setNumber: 42];
int n = [fido number];
```

You can talk like this:

```
fido.number = 42;
int n = fido.number;
```

As you can see, this is a very pleasant syntax. You use dot-notation to chain the property name to the instance, and you can use the resulting expression either on the left side of an equal sign (to set the instance variable's value) or elsewhere (to fetch the instance variable's value). Remember, though, that you can do this only if the class you're talking to has declared a property corresponding to the accessor methods in question. Remember also that your use of property syntax is not compulsory. If Dog has a `number` property, it has getter and setter methods `number` and `setNumber:`, and you are free to call them directly if you like. When you use a property in code, it is translated behind the scenes into a call to the corresponding getter or setter method, so it's all the same if you call the corresponding getter or setter method explicitly.

To use a property within the class that declares that property, you must use `self` explicitly. So, for example:

```
self.number = 42;
```



Do not confuse a property with an instance variable. An expression like `self->number = n`, or even simply `number = n`, sets the instance variable directly (and is possible only within the class, because instance variables are protected by default). An expression like `fido.number` or `self.number` involves a property and is equivalent to calling a getter or setter method. That getter or setter method may access an instance variable, and that instance variable may have the same name as the property, but that doesn't make them the same thing.

I have not yet told you *how* to declare a property corresponding to an instance variable. Plus, there are many options when declaring a property that affect how it can be used and what it means. All of that will be taken up in [Chapter 12](#). But I'm telling you about properties now because they are so widely used in Cocoa and because you'll see them so frequently in the documentation. For example, in [Chapter 1](#), I talked about setting a `UIView`'s `autoresizingMask` property:

```
myView.autoresizingMask =
    UIViewAutoresizingFlexibleTopMargin | UIViewAutoresizingFlexibleWidth;
```

How did I know I could talk that way? Because the `UIView` documentation says that `UIView` declares an `autoresizingMask` property. Near the top of the documentation page, we see this line:

`autoresizingMask` *property*

And further down, we get the details:

autoresizingMask

An integer bit mask that determines how the receiver resizes itself when its bounds change.

```
@property(nonatomic) UIViewAutoresizing autoresizingMask
```

That last line *is* the property declaration. Never mind for now what `nonatomic` means; the point is that `autoresizingMask` is a property. That's how I knew I could use property syntax as a way of calling a setter method; alternatively, I could have called the `setAutoresizingMask:` method explicitly.

Similarly, earlier in this chapter I called `UIView`'s `setFrame:` method, even though no such method is mentioned in the `UIView` documentation. What the `UIView` documentation does say is this:

frame

The frame rectangle, which describes the view's location and size in its superview's coordinate system.

```
@property(nonatomic) CGRect frame
```

The documentation is telling me that I can call a `UIView` setter method either by assigning to a `frame` property using dot-notation or by calling `setFrame:` explicitly.

Objective-C uses dot-notation for properties, and C uses dot-notation for structs; these can be chained. So, for example, `UIView`'s `frame` is a property whose value is a struct (a `CGRect`); thus, you can say `myView.frame.size.height`, where `frame` is a property that returns a struct, `size` is a component of that struct, and `height` is a component of *that* struct. But a struct is not a pointer, so you cannot (for example) *set* a frame's height directly through a chain starting with the `UIView`, like this:

```
myView.frame.size.height = 36.0; // compile error, "Expression is not assignable"
```

Instead, if you want to change a component of a struct property, you must fetch the property value into a struct variable, change the struct variable's value, and set the entire property value from the struct variable:

```
CGRect f = myView.frame;
f.size.height = 0;
myView.frame = f;
```

How to Write an Initializer

Now that you know about `self` and `super` and instance variables, we can return to a topic that I blithely skipped over earlier. I described how to initialize a newly minted instance by calling an initializer, and emphasized that you must always do so, but I said nothing about how to write an initializer in your own classes. You will wish to do so only when you want your class to provide a convenient initializer that goes beyond the functionality of the inherited initializers. Often your purpose will be to accept some parameters and use them to set the initial values of some instance variables.

For example, in our example of a Dog with a number, let's say we don't want any Dog instances to come into existence without a number; every Dog *must* have one. So having a value for its `number` ivar is a *sine qua non* of a Dog being instantiated in the first place. An initializer publicizes this rule and helps to enforce it — especially if it is the class's designated initializer. So let's decide that this initializer will be Dog's designated initializer.

Moreover, let's say that a Dog's number should not be changed. Once the Dog has come into existence, along with a number, that number should remain attached to that Dog instance for as long as that Dog instance persists.

So delete the `setNumber:` method and its declaration, thus destroying any ability of other classes to set a Dog instance's `number` after it has been initialized. Instead, we're going to set a Dog's `number` as it is initialized, using a method we'll declare like this:

```
- (id) initWithNumber: (int) n
```

Our return value is typed as `id`, not as a pointer to a Dog, even though in fact we will return a Dog object. This is a convention that we should obey. The name is conventional as well; as you know, the `init` beginning tells the world this is an initializer.

Now I'm just going to show you the actual code for the initializer ([Example 5-3](#)). Much of this code is conventional — a dance you are required to do. You should not question this dance: just do it. I'll describe the meaning of the code, but I'm not going to try to justify all the parts of the convention.

Example 5-3. Conventional schema for an initializer

```
- (id) initWithNumber: (int) n {
    self = [super init]; ❶ ❷
    if (self) {
        self->number = n; ❸
    }
    return self; ❹
}
```

The parts of the convention are:

- ❶ We send some sort of initialization message, calling a designated initializer. If this is our class's designated initializer, this message is sent to `super` and calls the super-

class's designated initializer. Otherwise, it is sent to `self` and calls either this class's designated initializer or another initializer that calls this class's designated initializer. In this case, this is our class's designated initializer, and the superclass's designated initializer is `init`.

- ❷ We capture the result of the initialization message to `super`, and assign that result to `self`. It comes as a surprise to many beginners (and not-so-beginners) that one can assign to `self` at all or that it would make sense to do so. But one can assign to `self` (because of how Objective-C messaging works behind the scenes), and it makes sense to do so because in certain cases the instance returned from the call to `super` might not be same as the `self` we started with.
- ❸ If `self` is not nil, we initialize any instance variables we care to. This part of the code is typically the only part you'll customize; the rest will be according to the pattern. Observe that I don't use any setter methods (or properties); in initializing an instance variable not inherited from the superclass, you should assign directly to the instance variable.
- ❹ We return `self`.



All instance variables are set to a form of zero by `alloc`. Therefore, any instance variables not initialized explicitly in an initializer remain 0. This means, among other things, that by default a BOOL instance variable is NO and an object reference instance variable is nil. It is common practice to take advantage of these defaults in your program; if the default values are satisfactory initial values, you won't bother to set them in your designated initializer.

But we are not finished. Recall from earlier in this chapter that a class that defines a designated initializer should also override the inherited designated initializer (in this case, `init`). And you can see why: if we don't, someone could say `[[Dog alloc] init]` and create a dog without a number — the very thing our initializer is trying to prevent. Just for the sake of the example, I'll make the overridden `init` assign a negative number as a signal that there's a problem. Notice that we're still obeying the rules: this initializer is not the designated initializer, so it calls this class's designated initializer.

```
- (id) init {
    return [self initWithNumber: -9999];
}
```

Just to complete the story, here's some code showing how we now would instantiate a Dog:

```
Dog* fido = [[Dog alloc] initWithNumber:42];
int n = [fido number];
// n is now 42; our initialization worked!
```


PART II

IDE

By now, you’re doubtless anxious to jump in and start writing an app. To do that, you need a solid grounding in the tools you’ll be using. The heart and soul of those tools can be summed up in one word: Xcode. In this part of the book we explore Xcode, the *IDE* (integrated development environment) in which you’ll be programming iOS. Xcode is a big program, and writing an app involves coordinating a lot of pieces; this part of the book will help you become comfortable with Xcode. Along the way, we’ll generate a simple working app through some hands-on tutorials.

- [Chapter 6](#) tours Xcode and explains the architecture of the *project*, the collection of files from which an app is generated.
- [Chapter 7](#) is about nibs. A *nib* is a file containing a drawing of your interface. Understanding nibs — knowing how they work and how they relate to your code — is crucial to your use of Xcode and to proper development of just about any app.
- [Chapter 8](#) pauses to discuss the Xcode documentation and other sources of information on the API.
- [Chapter 9](#) explains editing your code, testing and debugging your code, and the various steps you’ll take on the way to submitting your app to the App Store.

Anatomy of an Xcode Project

Xcode is the application used to develop an iOS app. An Xcode *project* is the entire collection of files and settings needed in order to construct an app. The source for an app is an Xcode project. To develop and maintain an app, you must know how to manipulate an Xcode project. That means you must know your way around a project, as displayed by Xcode. By the same token, you must know your way around Xcode sufficiently to manipulate a project.



The term “Xcode” is actually used in two ways. It’s the name for the entire suite of developer tools — the Xcode tools — and it’s the name of one application within that suite, the application in which you edit and build your app. This ambiguity should generally present little difficulty.

Xcode is a powerful, complex, and extremely large program. My approach when introducing Xcode to new users is to suggest that they adopt a kind of deliberate tunnel vision: if you don’t understand something, don’t worry about it, and don’t even look at it (and don’t touch it, because you might change something important). That’s the approach I’ll take here. This and subsequent chapters will undertake a simplified survey of Xcode, charting a somewhat restricted path, focusing on aspects of Xcode that you most need to understand immediately and resolutely ignoring those that you don’t.

For full information, study Apple’s own documentation (choose Help → Xcode Help); it may seem overwhelming at first, but what you need to know is probably in there somewhere. There are also entire books devoted to describing and explaining Xcode.

New Project

Even before you’ve written any code, an Xcode project is quite elaborate. To see this, let’s make a new, essentially “empty” project; you’ll see instantly that it isn’t empty at all.

1. Start up Xcode and choose File → New → New Project.
2. The “Choose a template” dialog appears. The template is your project’s initial set of files and settings. When you pick a template, you’re really picking an existing folder full of files; it will be one of the folders at some depth inside */Developer/Platforms/iPhoneOS.platform/Developer/Library/Xcode/Templates/Project Templates/Application* (I use the folder name */Developer* on the assumption that when you installed the Xcode tools you accepted the default and installed into a top-level */Developer* folder). This folder will essentially be copied, and a few values will be filled in, in order to create your project.

So, in this case, on the left, under iOS (not Mac OS X!), choose Application. On the right, select Single View Application. Click Next.

3. You are now asked to provide a name for your project (Product Name). Let’s call our new project *Empty Window*.

In a real project, you should give some thought to the project’s name, as you’re going to be living in close quarters with it. As Xcode copies the template folder, it’s going to use the project’s name to “fill in the blank” in several places, including some filenames and some settings, such as the name of the app. Thus, whatever you type at this moment is something you’ll be seeing in a lot of places throughout your project, for as long as you work with this project. So use a name that is either your app’s final name or at least approximates it.

It’s fine to use spaces in a project name. Spaces are legal in the folder name, the project name, the app name, and the various names of files that Xcode will generate automatically; and in the few places where spaces are problematic (such as the bundle identifier, discussed in the next paragraph), the name you type as the Product Name will have its spaces converted to hyphens.

4. Just below the Product Name field is the Company Identifier field. The first time you create a project, this field will be blank, and you should fill it in. The goal here is to create a unique string; your app’s bundle identifier, which is shown in gray below the company identifier, will consist of the company identifier plus a version of the project’s name, and because every project should have a unique name, the bundle identifier will also be unique and will thus uniquely identify this project along with the app that it produces and everything else connected with it. The convention is to start the company identifier with `com.` and to follow it with a string (possibly with multiple dot-components) that no one else is likely to use. For example, I use `com.neuburg.matt.`
5. Make sure the Device Family pop-up menu is set to iPhone. Make sure that Use Automatic Reference Counting is checked but that the other two checkboxes are unchecked. Ignore the Class Prefix field for now; it should be empty, with its default value “XYZ” shown in grey. Click Next.
6. You’ve now told Xcode how to construct your project. Basically, it’s going to copy the *Single View Application.xctemplate* folder from within the *Project Templates/*

Application folder I mentioned earlier. But you need to tell it where to copy this folder to. That's why Xcode is now presenting a Save As dialog. You are to specify the location of a folder that is about to be created — a folder that will be the *project folder* for this project.

The project folder can go just about anywhere, and you can move it after creating it. So the location doesn't matter much; I usually create new projects on the Desktop.

7. Xcode 4 also offers to create a git repository for your project. In real life, this can be a great convenience, but for now, uncheck that checkbox. Click Create.
8. The *Empty Window* project folder is created on disk (on the Desktop, if that's the location you just specified), and the project window for the *Empty Window* project opens in Xcode.

The project we've just created is a working project; it really does build an iOS app called *Empty Window*. To see this, make sure that the Scheme pop-up menu in the project window's toolbar reads *Empty Window* → iPhone 5.0 Simulator (though the exact system version number might be different), and choose Product → Run. After a while, the iOS Simulator application opens and displays your app running — an empty grey screen.



To *build* a project is to compile its code and assemble the compiled code, together with various resources, into the actual app. Typically, if you want to know whether your code compiles and your project is consistently and correctly constructed, you'll build the project (Product → Build). To *run* a project is to launch the built app, in the Simulator or on a connected device; if you want to know whether your code works as expected, you'll run the project (Product → Run), which automatically builds first if necessary.

The Project Window

An Xcode project must embody a lot of information about what files constitute the project and how they are to be used when building the app, such as:

- The source files (your code) that are to be compiled
- Any resources, such as icons, images, or sound files, as well as nib and storyboard files, that are to be part of the app
- Any frameworks to which the code must be linked as the app is built
- All settings (instructions to the compiler, to the linker, and so on) that are to be obeyed as the app is built

Xcode presents this information in graphical form, and this is one reason why a project window is so elaborate, and why learning to navigate and understand it takes time.

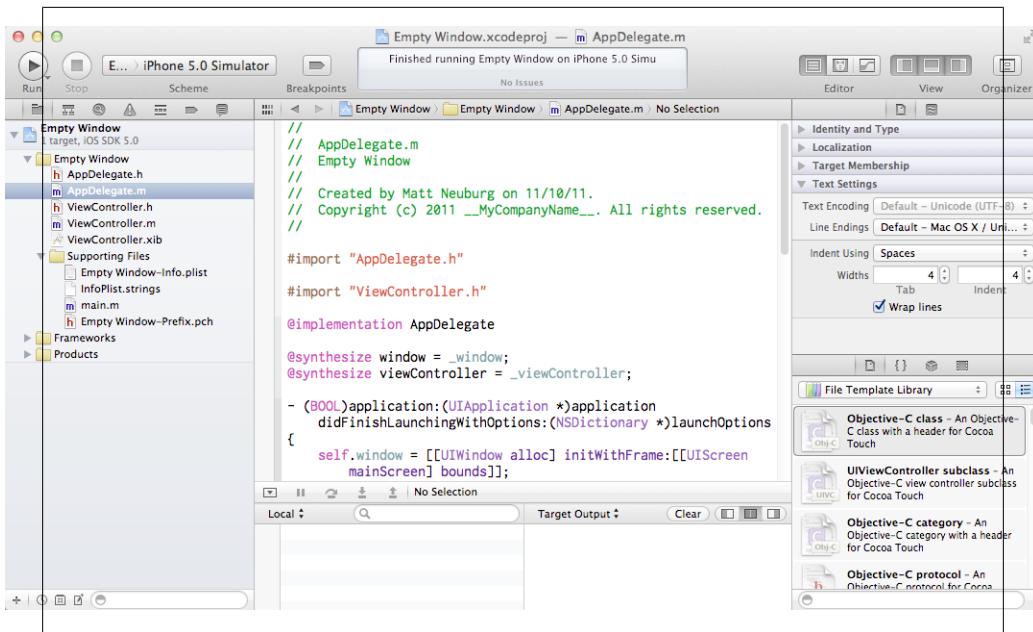


Figure 6-1. The project window, on steroids

Also, this single window must let you access, edit, and navigate your code, as well as reporting the progress and results of such procedures as building or debugging an app. In short, the single project window displays a lot of information and embodies a lot of functionality. You won't lose your way, however, if you just take a moment to explore this window and see how it is constructed.

[Figure 6-1](#) shows the project window, configured in rather an extreme manner, in order to display as many parts of the window as possible. In real life, you'd probably never show all these parts of the window at the same time, except very briefly, unless you had a really big monitor.

1. On the left is the Navigator pane. Show and hide it with [View → Navigators → Show/Hide Navigator](#) (Command-0) or with the first button in the View segmented control in the toolbar.
2. In the middle is the Editor pane (or simply “editor”). A project window always contains at least one Editor pane. I could have displayed this window with multiple Editor panes, but I was afraid that might make you run screaming from the room.
3. On the right is the Utilities pane. Show and hide it with [View → Utilities → Show/Hide Utilities](#) (Command-Option-0) or with the third button in the View segmented control in the toolbar.

4. At the bottom is the Debugger pane. Show and hide it with View → Show/Hide Debug Area (Shift-Command-Y) or with the second button in the View segmented control in the toolbar.



All Xcode keyboard shortcuts can be customized; see the Key Bindings pane of the Preferences window. Keyboard shortcuts that I cite are the defaults.

The Navigator Pane

All navigation of the project window begins ultimately with the Navigator pane, the column of information at the left of the window. It is possible to toggle the visibility of the Navigator pane (View → Navigators → Hide/Show Navigator, or Command-0); for example, once you've used the Navigator pane to reach the item you want to see or work on, you might hide the Navigator pane temporarily to maximize your screen real estate (especially on a smaller monitor). You can change the Navigator pane's width by dragging the vertical line at its right edge.

The Navigator pane itself can display seven different sets of information; thus, there are actually seven navigators. These are represented by the seven icons across its top; to switch among them, use these icons or their keyboard shortcuts (Command-1, Command-2, and so on). You will quickly become adept at switching to the navigator you want; their keyboard shortcuts will become second nature. If the Navigator pane is hidden, pressing a navigator's keyboard shortcut both shows the Navigator pane and switches to that navigator.

Depending on your settings in the Behaviors pane of Xcode's preferences, a navigator might show itself automatically when you perform a certain action. For example, by default, when you build your project, if warning messages or error messages are generated, the Issue navigator will appear. This automatic behavior will not prove troublesome, because it is generally precisely the behavior you want, and if it isn't, you can change it; plus you can easily switch to a different navigator at any time.

The most important general use pattern for the Navigator pane is: you select something in the Navigator pane, and that thing is displayed in the main area of the project window. Let's begin experimenting immediately with the various navigators:

Project navigator (Command-1)

Click here for basic navigation through the files that constitute your project. For example, in the Empty Window folder (these folder-like things in the Project navigator are actually called *groups*) click *AppDelegate.m* to view its code ([Figure 6-2](#)).

At the top level of the Project navigator, with a blue Xcode icon, is the Empty Window project itself; click it to view the settings associated with your project and

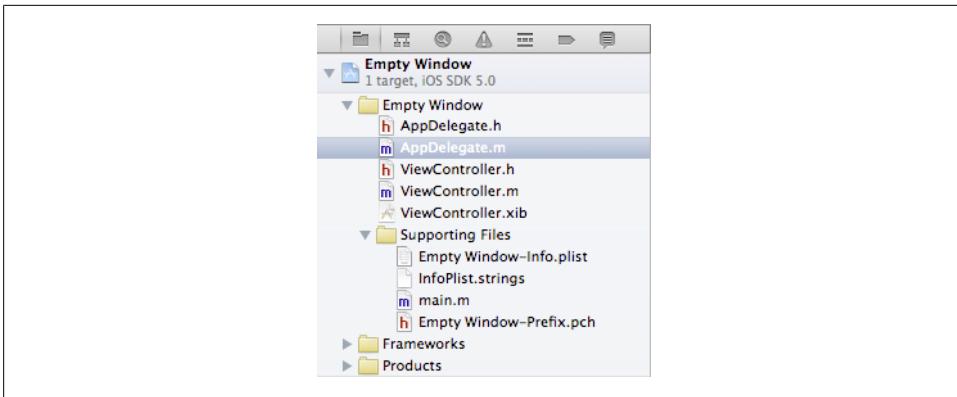


Figure 6-2. The Project navigator

its targets. Don't change anything here without knowing what you're doing! I'll talk later in this chapter about what these settings are for.

Symbol navigator (Command-2)

A *symbol* is a name, typically the name of a class or method. Depending on which of the three icons in the filter bar at the bottom of the Symbol navigator you highlight, you can view Cocoa's built-in symbols or the symbols defined in your project. The former can be a useful form of documentation; the latter can be helpful for navigating your code. For example, highlight the first two icons in the filter bar (the first two are dark-colored, the third is light), and see how quickly you can reach the definition of AppDelegate's `applicationDidBecomeActive:` method.

Feel free to highlight the filter bar icons in various ways to see how the contents of the Symbol navigator change. Note too that you can type in the search field in the filter bar to limit what appears in the Symbol navigator; for example, try typing "active" in the search field, and see what happens.

Search navigator (Command-3)

This is a powerful search facility for finding text globally in your project, and even in the headers of Cocoa frameworks. You can also summon the Search navigator with **Edit** → **Find** → **Find in Workspace** (Shift-Command-F). To access the full set of options, click the magnifying glass and choose **Show Find Options**. For example, try searching for "delegate" (Figure 6-3). Click a search result to jump to it in your code.

Issue navigator (Command-4)

You'll need this navigator primarily when your code has issues. This doesn't refer to emotional instability; it's Xcode's term for warning and error messages emitted when you build your project.

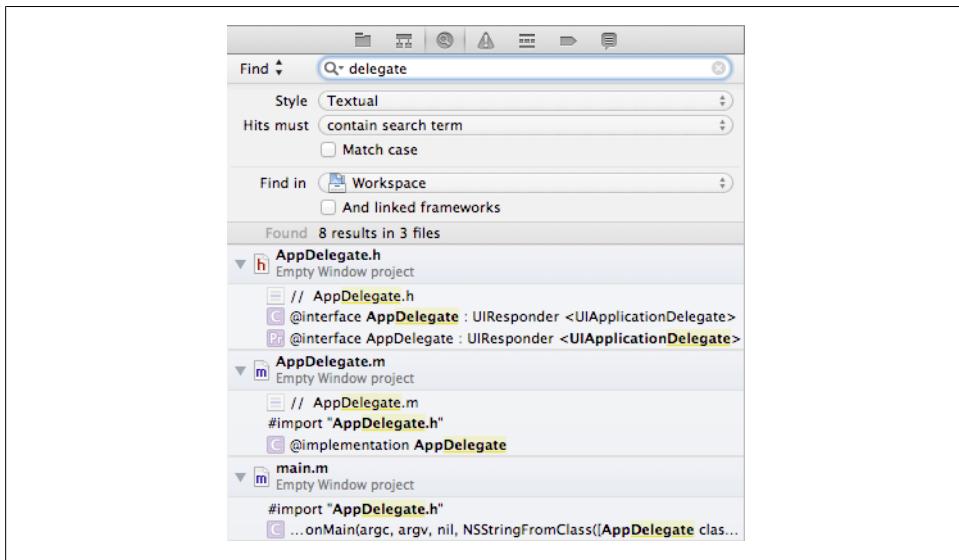


Figure 6-3. The Search navigator

To see the Issue navigator in action, you'll need to give your code an issue. For example, navigate (as you already know how to do, in at least three different ways) to the file *AppDelegate.m*, and in the blank line after the last comment at the top of the file, above the `#import` line, type `howdy`. Build (Command-B), saving if you're prompted to. The Issue navigator will display numerous error messages, showing that the compiler is totally unable to cope with this illegal word appearing in an illegal place. Click an issue to see it within its file. In your code, issue "balloons" may appear to the right of lines containing issues; if you're distracted or hampered by these, toggle their visibility with Editor → Issues → Hide/Show All Issues. (Now that you've made Xcode miserable, select "howdy" and delete it; build again, and your issues will be gone. If only real life were this easy!)

Debug navigator (Command-5)

By default, this navigator will appear when your code is paused while you're debugging it. There is not a strong distinction in Xcode between running and debugging; the milieu is the same. (The difference is mostly a matter of whether breakpoints are obeyed; more about that, and about debugging in general, in [Chapter 9](#).) However, if your code runs and doesn't pause, the Debug navigator by default won't come into play.

To see the Debug navigator in action, you'll need to give your code a breakpoint. Navigate once more to the file *AppDelegate.m*, select in the line that says `return YES`, and choose Product → Debug → Add Breakpoint at Current Line to make a blue breakpoint arrow appear on that line. Run the project. (If the project is already running, the Stop dialog may appear; click Stop to terminate the current

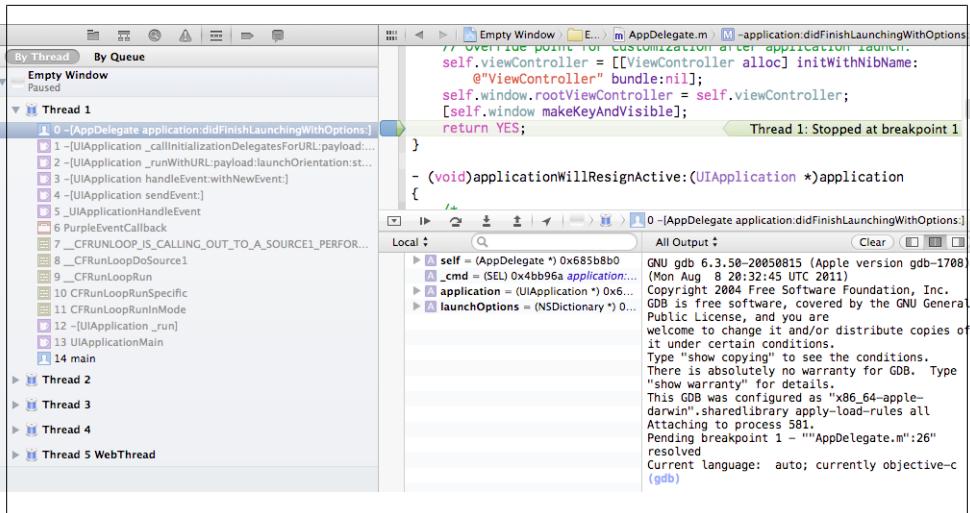


Figure 6-4. The Debug layout

run and begin a new one.) By default, as the breakpoint is encountered, the Navigator pane switches to the Debug navigator, and the Debug pane appears at the bottom of the window.

This overall layout (Figure 6-4) will rapidly become familiar as you debug your projects. The Debug navigator displays the call stack, with the names of the nested methods in which the pause occurs; as you would expect, you can click on a method name to navigate to it. You can shorten or lengthen the list with the slider at the bottom of the pane. The Debug pane, which can be shown or hidden at will (View → Hide/Show Debug Area, or Shift-Command-Y) consists of two subpanes, either of which can be hidden using the segmented control at the top right of the pane.

- On the left, the variables list is populated with the variables in scope for the selected method in the call stack (and you can optionally display processor registers as well).
- On the right is the console, where the debugger displays text messages; that's how you learn of exceptions thrown by your running app. Exceptions are extremely important to know about, and this is your only way to know about them, so keep an eye on the console as your app runs. Note also that View → Debug Area → Activate Console shows the console pane if only the variables list is displayed.

You can also use the console to communicate via text with the debugger. This can often be a better way to explore variable values during a pause than the variables list.

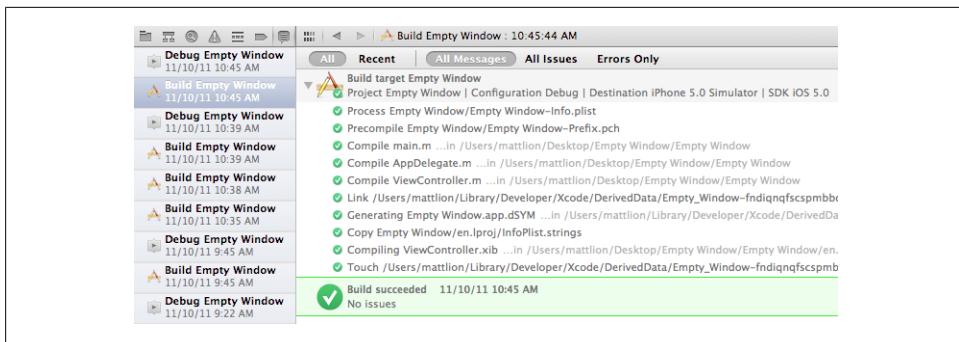


Figure 6-5. Viewing a log

Breakpoint navigator (Command-6)

This navigator lists all your breakpoints. At the moment you've only one, but when you're actively debugging a large project, you'll be glad of this navigator. Also, this is where you create special breakpoints (such as symbolic breakpoints), and in general it's your center for managing existing breakpoints. We'll return to this topic in [Chapter 9](#).

Log navigator (Command-7)

This navigator lists your recent major actions, such as building or running (debugging) your project. Click on a listing to see the log file generated when you performed that action. The log file might contain information that isn't displayed in any other way, and also it lets you dredge up messages from the recent past ("What was that exception I got while debugging a moment ago?").

For example, by clicking on the listing for a successful build, and by choosing to display All and All Messages using the filter switches at the top of the log, we can see the steps by which a build takes place ([Figure 6-5](#)). To reveal the full text of a step, click on that step and then click the Expand Transcript button that appears at the far right (and see also the menu items in the Editor menu).

When navigating by clicking in the Navigator pane, modifications to your click can determine where navigation takes place. For the settings that govern these click modifications, see the General pane of Xcode's preferences. For example, if you haven't changed the original settings, Option-click navigates in an assistant pane (discussed later in this chapter), and double-click navigates by opening a new window.

The Utilities Pane

The Utilities pane, the column at the right of the project window, consists partly of inspectors that provide information about, and in some cases let you change the specifications of, the current selection, and partly of libraries that function as a source of objects you may need while editing your project. Its importance emerges mostly when you're working in the nib editor ([Chapter 7](#)), and you'll probably keep it hidden the

rest of the time. But if you have sufficient screen real estate, you might like to keep it open while editing code, because Quick Help, a form of documentation ([Chapter 8](#)), is displayed here as well; plus, the Utilities pane is the source of code snippets ([Chapter 9](#)). To toggle the visibility of the Utilities pane, choose **View** → **Utilities** → **Hide/Show Utilities** (Command-Option-0). You can change the Utilities pane’s width by dragging the vertical line at its left edge.

Many individual inspectors and libraries are discussed in subsequent chapters. Here, I’ll just describe the overall physical characteristics of the Utilities pane.

The Utilities pane consists of a set of palettes. Actually, there are so many of these palettes that they are clumped into multiple sets, divided into two major groups: the top half of the pane and the bottom half of the pane. You can change the relative heights of these two halves by dragging the horizontal line that separates them.

The top half

What appears in the top half of the Utilities pane depends on what’s selected in the current editor. There are two main cases:

A code file is being edited

The top half of the Utilities pane shows either the File inspector or Quick Help. Toggle between them with the icons at the top of this half of the Utilities pane, or with their keyboard shortcuts (Command-Option-1, Command-Option-2). The File inspector is rarely needed, but Quick Help can be useful as documentation. The File inspector consists of multiple sections, each of which can be expanded or collapsed by clicking its header.

A nib or storyboard file is being edited

The top half of the Utilities pane shows, in addition to the File inspector and Quick Help, the Identity inspector (Command-Option-3), the Attributes inspector (Command-Option-4), the Size inspector (Command-Option-5), and the Connections inspector (Command-Option-6). Like the File inspector, these can consist of multiple sections, each of which can be expanded or collapsed by clicking its header.

The bottom half

The bottom half of the Utilities pane shows one of four libraries. Toggle between them with the icons at the top of this half of the Utilities pane, or with their keyboard shortcuts. They are the File Template library (Command-Option-Control-1), the Code Snippet library (Command-Option-Control-2), the Object library (Command-Option-Control-3), and the Media library (Command-Option-Control-4). The Object library is the most important; you’ll use it heavily when editing a nib or storyboard.

To see a help pop-up describing the currently selected item in a library, press Spacebar.

The Editor

In the middle of the project window is the *editor*. This is where you get actual work done, reading and writing your code ([Chapter 9](#)), or designing your interface in a nib or storyboard file ([Chapter 7](#)). The editor is the core of the project window. You can eliminate the Navigator pane, the Utilities pane, and the Debug pane, but there is no such thing as a project window without an editor (though you can cover the editor completely with the Debug pane).

The editor provides its own form of navigation, the *jump bar* across the top. I'll talk more later about the jump bar, but for now, observe that not only does it show you hierarchically what file is currently being edited, but also it allows you to switch to a different file. In particular, each path component in the jump bar is also a pop-up menu. These pop-up menus can be summoned by clicking on a path component, or by using keyboard shortcuts (shown in the second section of the View → Standard Editor submenu). For example, Control-4 summons a hierarchical pop-up menu, which can be navigated entirely with the keyboard, allowing you to choose a different file in your project to edit. Thus you can navigate your project even if the Project navigator isn't showing.

It is extremely likely, as you develop a project, that you'll want to edit more than one file simultaneously, or obtain multiple views of a single file so that you can edit two areas of it simultaneously. This can be achieved in three ways: assistants, tabs, and secondary windows.

Assistants

You can split the editor into multiple editors by summoning an *assistant pane*. To do so, click the second button in the Editor segmented control in the toolbar, or choose View → Assistant Editor → Show Assistant Editor (Command-Option-Return). Also, by default, adding the Option key to navigation opens an assistant pane; for example, Option-click in the Navigator pane, or Option-choose in the jump bar, to navigate by opening an assistant pane (or to navigate in an existing assistant pane if there is one). To remove the assistant pane, click the first button in the Editor segmented control in the toolbar, or choose View → Standard Editor → Show Standard Editor (Command-Return), or click the “x” button at the assistant pane's top right.

Your first task will be to decide how you want multiple editor panes arranged with respect to one another. To do so, choose from the View → Assistant Editor submenu. I usually prefer All Editors Stacked Vertically, but it's purely a matter of personal taste and convenience.

Once you've summoned an assistant pane, you can split it further into additional assistant panes. To do so, click the “+” button at the top right of an assistant pane. To dismiss an assistant pane, click the “x” button at its top right.

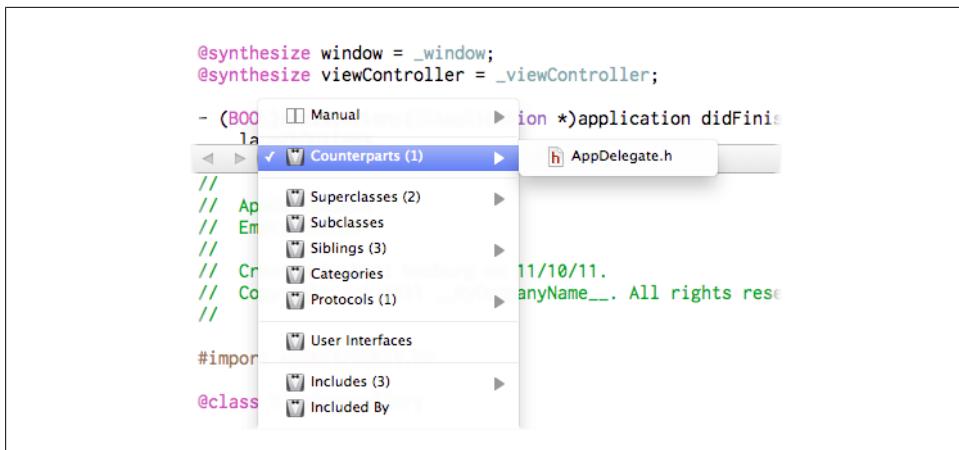


Figure 6-6. Telling an assistant pane to display counterparts

An assistant pane is not merely split-pane editing; an assistant pane can bear a special relationship to the primary editor pane. The primary editor pane is the one whose contents, by default, are determined by what you click on in the Navigator pane; an assistant pane, meanwhile, can respond to what file is being edited in the primary editor pane by changing intelligently what file it (the assistant pane) is editing. This is called *tracking*.

To see tracking in action, open a single assistant pane and set the first component in its jump bar to Counterparts (Figure 6-6). Now use the Project navigator to select *AppDelegate.m*; the primary editor pane displays this file, and the assistant automatically displays *AppDelegate.h*. Next, use the Project navigator to select *AppDelegate.h*; the primary editor pane displays this file, and the assistant automatically displays *AppDelegate.m*. There's a lot of convenience and power lurking here, which you'll explore as you need it.

Tabs

You can embody the entire project window interface as a tab. To do so, choose File → New → New Tab (Command-T), revealing the tab bar (just below the toolbar) if it wasn't showing already. Use of a tabbed interface will likely be familiar from applications such as Safari. You can switch between tabs by clicking on a tab, or with Command-Shift-}. At first, your new tab will look largely identical to the original window from which it was spawned. But now you can make changes in a tab — change what panes are showing or what file is being edited, for example — without affecting any other tabs. Thus you can get multiple views of your project.

Secondary windows

A secondary project window is similar to a tab, but it appears as a separate window instead of a tab in the same window. To create one, choose File → New → New Window (Command-Shift-T). Alternatively, you can promote a tab to be a window

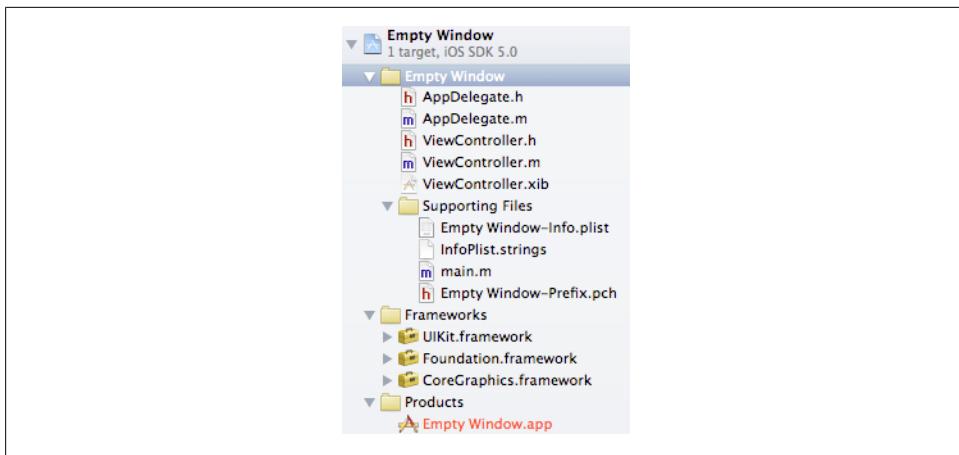


Figure 6-7. The Project navigator again

by dragging it right out of its current window. Yet another way to make a secondary project window is to choose Navigate → Open In and navigate left in the resulting dialog until the dialog offers to make a new window.

There isn't a strong difference between a tab and a secondary window; which you use, and for what, will be a matter of taste and convenience. I find that the advantage of a secondary window is that you can see it at the same time as the main window, and that it can be small. Thus, when I have a file I frequently want to refer to, I often spawn into a secondary window an editor displaying that file, making it fairly small and without any additional panes.

The Project File and Its Dependents

The first item in the Project navigator (Command-1) represents the project file on disk (in our new project, this is called *Empty Window*). Hierarchically dependent upon it are items that contribute to the building of the project (Figure 6-7).

Many of these items, including the project file itself, correspond to items on disk in the project folder. To survey this correspondence, let's examine the project folder in the Finder simultaneously with the Xcode project window. Select the project file listing in the Project navigator and choose File → Show in Finder.

The Finder displays the contents of your project folder (Figure 6-8). The most important of these is *Empty Window.xcodeproj*. This is the project file. All Xcode's knowledge about your project — what files it consists of and how to build the project — is stored in this file.

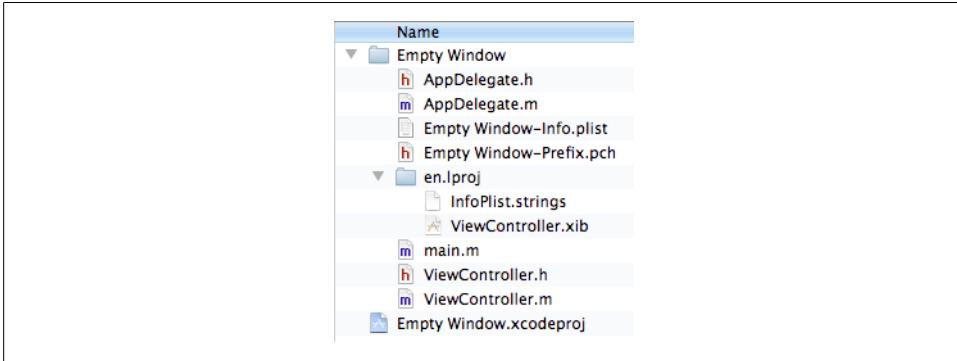
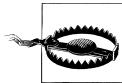


Figure 6-8. The project folder



To open a project from the Finder, double-click the project file. This will launch Xcode if it isn't already running.



Never, never, *never* touch anything in a project folder by way of the Finder, except for double-clicking the project file to open the project. Don't put anything directly into a project folder. Don't remove anything from a project folder. Don't rename anything in a project folder. Don't touch anything in a project folder! Do all your interaction with the project through the project window in Xcode.

The reason for the foregoing warning is that in order to work properly, the project expects things in the project folder to be a certain way. If you make any alterations to the project folder directly in the Finder, behind the project's back, you can upset those expectations and break the project. When you work in the project window, it is Xcode itself that makes any necessary changes in the project folder, and all will be well. (When you're an Xcode power user, you'll know when you can disobey this rule. Until then, just obey it blindly and rigorously.)

Consider now the groups and files shown in the Project navigator as hierarchically dependent upon the project file, and how they correspond to reality on disk as portrayed in the Finder. (Recall that *group* is the technical term for the folder-like objects shown in the Project navigator.)

The first thing you'll notice is that groups in the Project navigator don't necessarily correspond to folders on disk in the Finder, and folders on disk in the Finder don't necessarily correspond to groups in the Project navigator.

- The Empty Window group is, to some extent, real; it corresponds directly to the *Empty Window* folder on disk. If you were to create additional files (which, in real life, you would almost certainly do in the course of developing your project), you

would likely put them in the Empty Window group in the Project navigator so that they'd be in the *Empty Window* folder on disk. (Doing so, however, is not a requirement; your files can live anywhere and your project will still work fine.)

- The Supporting Files group, on the other hand, corresponds to nothing on disk; it's just a way of clumping some items together in the Project navigator, so that they can be located easily and can be shown or hidden together. The things *inside* this group are real, however; you can see that the four files *Empty Window-Info.plist*, *InfoPlist.strings*, *Empty Window-Prefix.pch*, and *main.m* do exist on disk — they're just not inside anything called *Supporting Files*. Rather, they're at the top level of the *Empty Window* folder.
- Two files, *InfoPlist.strings* and *ViewController.xib*, appear in the Finder inside a folder called *en.lproj*, which doesn't appear in the Project navigator. The folder *en.lproj* has to do with *localization*, which I'll discuss in [Chapter 9](#).

You may be tempted to find all this confusing. Don't! Remember what I said about not involving yourself with the project folder on disk in the Finder. Keep your attention on the Project navigator, make your modifications to the project there, and all will be well.

By convention, as you add other files to your project that are not code but need to be copied into the app as it is built, such as sound and image files, you would usually put them into yet another group — probably, though not necessarily, a group inside the Empty Window group. You might call this group Resources. (I usually do.) And as your project grows further, you should feel free to create even more groups to help organize your files. To make a new group, choose File → New → New Group. To rename a group, select it in the Project navigator and press Return to make the name editable.

When I say “feel free,” I mean it. You want navigating your project to be easy and intuitive. That's what groups are for. They are just ways of making the Project navigator work well for you. As we've seen, they don't necessarily affect how the actual files are stored on disk. Even more important, they don't affect how the app is built. It is not the placement of files in groups or in the Finder that causes them to be built into the app; it's their inclusion in the appropriate target build phase, as I'll explain later in this chapter.

The things in the Frameworks group and the Products group don't correspond to anything in the project folder, but they do correspond to real things that the project needs to know about in order to build and run:

Frameworks

This group, by convention, lists frameworks (Cocoa code) that your code calls. Frameworks exist on disk, but they are not built into your app when it is constructed; they don't have to be, because they are present on the target device (an iPhone, iPod touch, or iPad). Instead, the frameworks are *linked* to the app, meaning that the app knows about them and expects to find them on the device when it runs. Thus, all the framework code is omitted from the app itself, saving considerable space.

Products

This group, by convention, holds an automatically generated reference to the built app.

The Target

A *target* is a collection of parts along with rules and settings for how to build a product from them. It is a major determinant of how an app is built. Whenever you build, what you're really building is a target.

Select the Empty Window project at the top of the Project navigator, and you'll see two things on the left side of the editor: the project itself, and a list of your targets. In this case, there is only one target, called Empty Window (just like the project itself). But there *could* be more than one target, under certain circumstances. For example, you might want to write an app that can be built as an iPhone app or as an iPad app — two different apps that share a lot of the same code. So you might want one project containing two targets.

If you select the project in the left side of the editor, you *edit the project*. If you select the target in the left side of the editor, you *edit the target*. I'll use those expressions a lot in later instructions.

Build Phases

Edit the target and click Build Phases at the top of the editor (Figure 6-9). These are the stages by which your app is built. By default, there are three of them with content — Compile Sources, Link Binary With Libraries, and Copy Bundle Resources — and those are the only stages you'll usually need, though you can add others. The build phases are both a report to you on how the target will be built and a set of instructions to Xcode on how to build the target; if you change the build phases, you change the build process.

The meanings of the three build phases are pretty straightforward:

Compile Sources

Certain files (your code) are compiled, and the resulting compiled code (a single file called the *binary*) is copied into the app.

Link Binary With Libraries

Certain libraries, usually frameworks, are linked to the compiled code, so that it will expect them to be present on the device when the app runs.

Copy Bundle Resources

Certain files are copied into the app, so that your code or the system can find them there when the app runs. For example, if your app had an icon, it would need to be copied into the app so the device could find and display it.

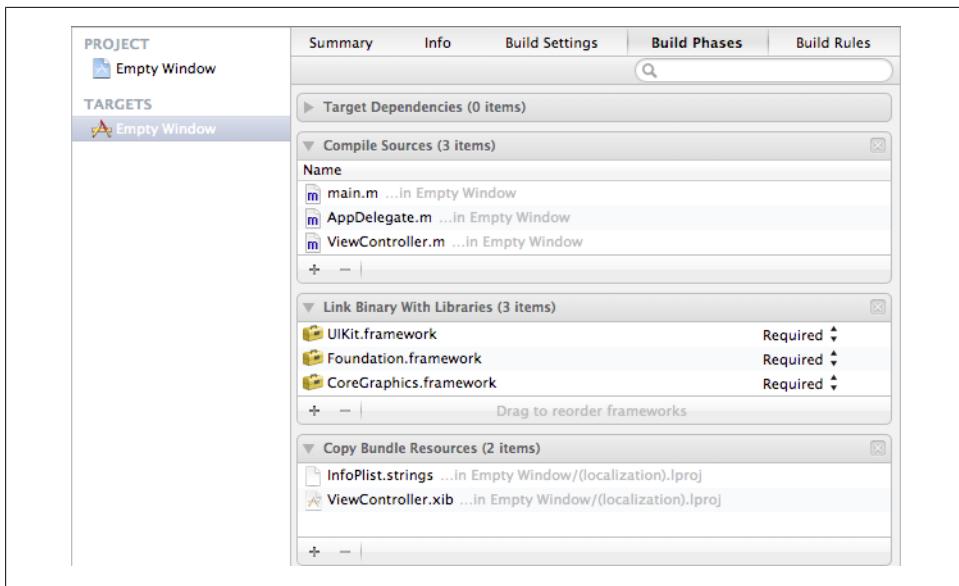


Figure 6-9. Build phases

By opening the build phases in the editor, you can see the files to which each phase applies. The first phase, Compile Sources, presently compiles two files (*main.m* and *AppDelegate.m*). The second phase, Link Binary With Libraries, presently links three libraries (frameworks). The third phase, Copy Bundle Resources, presently copies two files (*InfoPlist.strings*, along with *MainWindow.xib*, the nib file).

You can alter these lists. If something in your project was not in Copy Bundle Resources and you wanted it copied into the app during the build process, you could drag it from the Project navigator into the Copy Bundle Resources list, or (easier) click the “+” button beneath the Copy Bundle Resources list to get a helpful dialog listing everything in your project. If something in your project was in Copy Bundle Resources and you didn’t want it copied in the app, you would delete it from the list; this would not delete it from your project, from the Project navigator, or from the Finder, but only from the list of things to be copied into your app.

Build Settings

Build phases are only one aspect of how a target knows how to build the app. The other aspect is build settings. To see them, edit the target and click Build Settings at the top of the editor (Figure 6-10). Here you’ll find a long list of settings, most of which you’ll never touch. But Xcode examines this list in order to know what to do at various stages of the build process. Build settings are the reason your project compiles and builds the way it does.

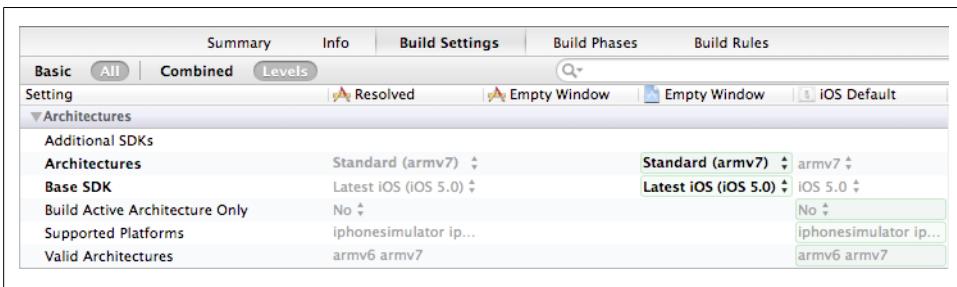


Figure 6-10. Target build settings

You can determine what build settings are displayed by clicking Basic or All. The settings are combined into categories, and you can close or open each category heading to save room. If you know something about a setting you want to see, such as its name, you can use the search field at the top right to filter what settings are shown.

You can determine how build settings are displayed by clicking Combined or Levels; in [Figure 6-10](#), I've clicked Levels, in order to discuss what levels are. It turns out that not only does a *target* contain values for the build settings, but the *project* also contains values for the same build settings; furthermore, Xcode has certain built-in default build setting values. The Levels display shows all of these levels at once, so you can understand the derivation of the actual values used for every build setting.

To understand the chart, read from right to left. For example, the Base SDK build setting (whose meaning I'll discuss in [Chapter 9](#)) is set to be iOS 5.0 by the built-in Xcode default (the rightmost column). Then, however, the project comes along with a different value for this build setting, namely Latest iOS (second column from the right). The target does not override this value (third column from the right). Therefore the actual value used will be Latest iOS (fourth column from the right, "Resolved"). It happens that the latest iOS is iOS 5.0 at the moment, so the end result is just the same as if the project had not overridden Xcode's built-in default value for this setting; but in theory an iOS 5.1 SDK could someday be installed on this machine, and now the project would use iOS 5.1 as its Base SDK even though Xcode's default setting continues to specify iOS 5.0.

If you wanted to change this value, you could, here and now. You could change the value at the project level or at the target level. I'm not suggesting that you should do so; indeed, you will rarely have occasion to manipulate build settings directly, as the defaults are usually acceptable. Nevertheless, you *can* change build setting values, and this is where you would do so. For details on what the various build settings are, consult Apple's documentation, especially the *Xcode Build Setting Reference*. Also, you can select a build setting and show Quick Help in the Utilities pane to learn more about it.

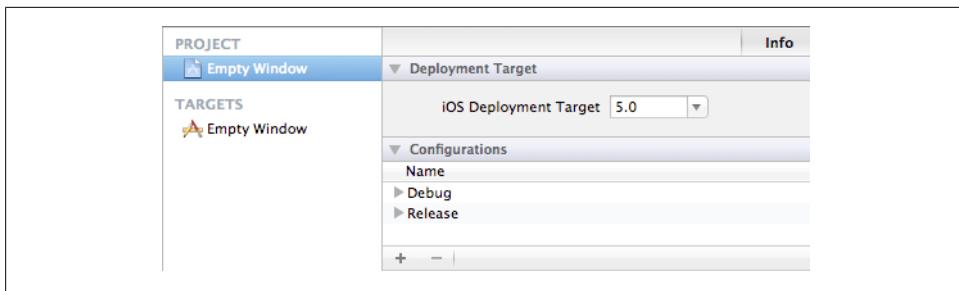


Figure 6-11. Configurations

Configurations

There are actually multiple lists of build setting values — though only one such list applies when a build is performed. Each such list is called a *configuration*. Multiple configurations are needed because you build in different ways at different times for different purposes, and thus you'll want certain build settings to take on different values under different circumstances.

By default, there are two configurations:

Debug

This configuration is used throughout the development process, as you write and run your app.

Release

This configuration is used for late-stage testing, when you want to check performance on a device.

Configurations exist at all because the project says so. To see where the project says so, edit the project and click Info at the top of the editor ([Figure 6-11](#)). Note that these configurations are just names. You can make additional configurations, and when you do, you're just adding to a list of names. The importance of configurations emerges only when those names are coupled with build setting values. Configurations can affect build setting values both at the project level and at the target level.

For example, return to the target build settings ([Figure 6-10](#)) and type “Optim” into the search field. Now you can look at the Optimization Level build setting. The Debug configuration value for Optimization Level is None: while you’re developing your app, you build with the Debug configuration, so your code is just compiled line by line in a straightforward way. The Release configuration value for Optimization Level is Fastest, Smallest; when your app is ready to ship, you build it with the Release configuration, so that the resulting binary is faster and smaller, which is great for your users installing and running the app on a device, but would be no good while you’re developing the app because breakpoints and stepping in the debugger wouldn’t work properly. (A not uncommon beginner error is building with the Release configuration and then wondering why the debugger isn’t pausing at breakpoints any more.)

Schemes and Destinations

So far, I have said that there are configurations, and I have explained that you may need to switch between configurations in order to get the build setting values appropriate for your current purpose. But I have not said how the configuration is determined as you actually build. It's determined by a scheme.

A *scheme* unites a target (or multiple targets) with a build configuration, with respect to the purpose for which you're building. A new project, such as Empty Window, comes by default with a single scheme, named after the project's single target. Thus this project's single scheme is called, by default, Empty Window. To see it, choose Product → Edit Scheme. The scheme editor dialog opens. Make sure that Info at the top of the dialog is selected.

On the left side of the scheme editor are listed various actions you might perform from the Product menu. Click an action to see its corresponding settings in this scheme. The first action, the Build action, is different from the other actions, because it is common to all of them (the other actions all implicitly involve building); thus the Build action merely determines what target(s) will be built when each of the other actions is performed, and for our simple project this is trivial, because we've only one target and we always need it built. So, now consider the Run action.

When you click the Run action at the left, the editor displays the settings that will be used when you build and run (Figure 6-12). As you can see, the Build Configuration pop-up menu is set to Debug. That explains where the current build configuration comes from. At the moment, whenever you build and run, you're using the Debug build configuration and the build setting values that correspond to it, because you're using this scheme, and that's what this scheme says to do when you build and run.

Now dismiss the scheme editor, and consider how you might proceed if you wanted to build and run using the Release build configuration. (The Debug build configuration settings may affect the behavior of the built app, so you want to test the app as an actual user would experience it.) One way would be to return to the scheme editor and change the build configuration for the Run action for this scheme. Xcode makes this convenient: hold the Option key as you choose Product → Run (or as you click the Run button in the toolbar). The scheme editor appears, containing a Run button. So now you can make any changes you like, such as setting the Build Configuration pop-up menu to Release for the Run action, and proceed directly to build and run the app by clicking Run.

(If you're following along and you did make this change, open the scheme editor again and set the Build Configuration pop-up for the Run action in our Empty Window scheme back to Debug.)

On the other hand, if you were to find yourself often wanting to switch between building and running with the Debug configuration and building and running with the Release configuration, you might create a distinct, additional scheme that uses the Release

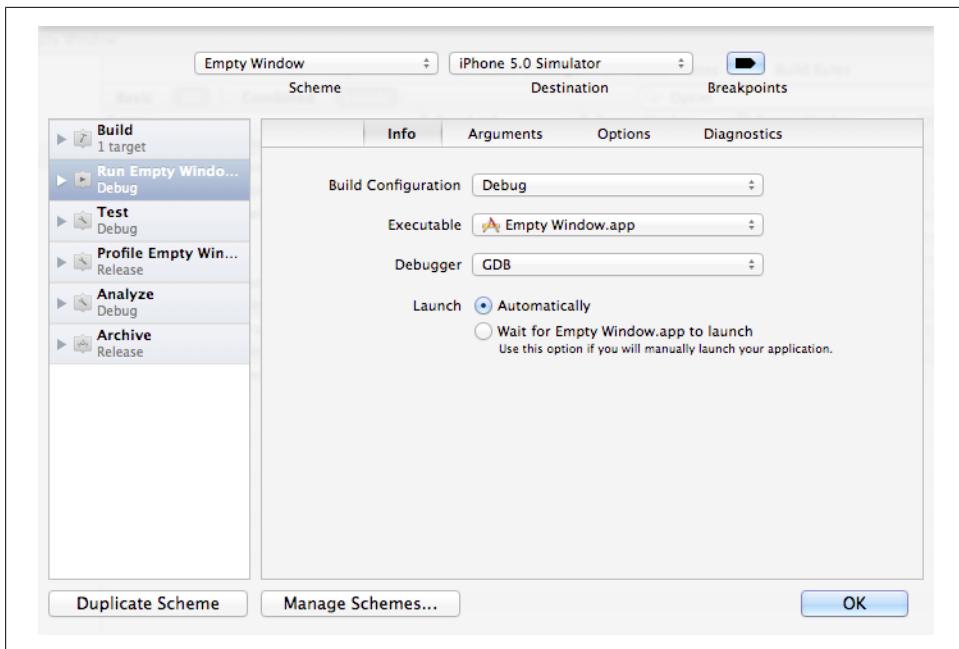


Figure 6-12. The scheme editor

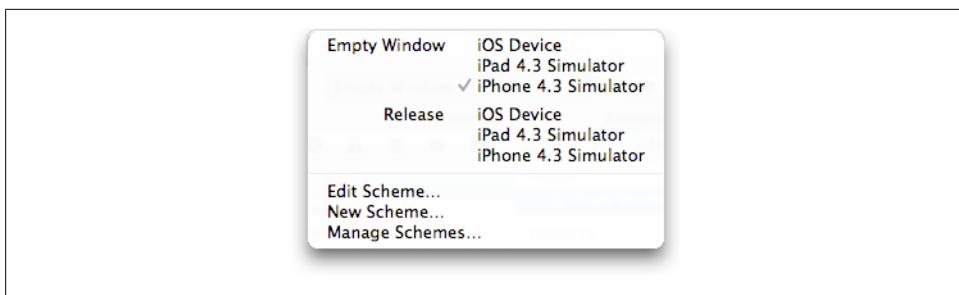


Figure 6-13. The Scheme pop-up menu

debug configuration for the Run action. This is easy to do: in the scheme editor, click Duplicate Scheme. The name of the new scheme is editable; let's call it Release. Change the Build Configuration pop-up for the Run action in our new scheme to Release, and dismiss the scheme editor.

Now you have two schemes, Empty Window (whose build configuration for running is Debug) and Release (whose build configuration for running is Release). To switch between them easily, you can use the Scheme pop-up menu in the project window toolbar (Figure 6-13) before you build and run.

The Scheme pop-up menu lists each scheme, along with each destination on which you might run your built app. A *destination* is effectively a machine that can run your app. For example, you might want to run the app in the Simulator or on a physical device. There is no configuration of destinations; you are automatically assigned destinations, depending on what system your project is set to run on and what devices are connected to your computer.

Destinations and schemes have nothing to do with one another; your app is built the same way regardless of your chosen destination. The presence of destinations in the Scheme pop-up menu is intended as a convenience, allowing you to use the pop-up menu to choose either a scheme or a destination, or both, in a single move. To switch easily among destinations without changing schemes, click near the right end of the Scheme pop-up menu. To switch among schemes, possibly also determining the destination (as shown in [Figure 6-13](#)), click near the left end of the Scheme pop-up menu. You can also switch among schemes or among destinations by using the scheme editor.

From Project to App

An app file is really a special kind of folder called a *package* (and a special kind of package called a *bundle*). The Finder normally disguises a package as a file and does not dive into it to reveal its contents to the user, but you can bypass this protection and investigate an app bundle with the Show Package Contents command. By doing so, you can study the structure of your built app bundle.

We'll use the Empty Window app that we built earlier as a sample minimal app to investigate. You'll have to locate it in the Finder; by default, it should be somewhere in your user *Library/Developer/Xcode/DerivedData* folder, as shown in [Figure 6-14](#). (In theory, you should be able to select the app under Products in the Navigation pane and choose File → Show in Finder, but there seems to be a long-standing bug preventing this.)

In the Finder, Control-click the Empty Window app (if you're using Lion, I presume you know how to show the user *Library* directory), and choose Show Package Contents from the contextual menu.

Looking inside our minimal app bundle ([Figure 6-15](#)), we see that it contains just five files:

PkgInfo

A tiny text file reading **APPL????**, signifying the type and creator codes for this app. The *PkgInfo* file something of a dinosaur; it isn't really necessary for the functioning of an iOS app and is generated automatically. You'll never need to touch it.

InfoPlist.strings

A text file intended for text appearing in our app that might need to be translated into different languages. It is copied directly from *InfoPlist.strings* in the project.

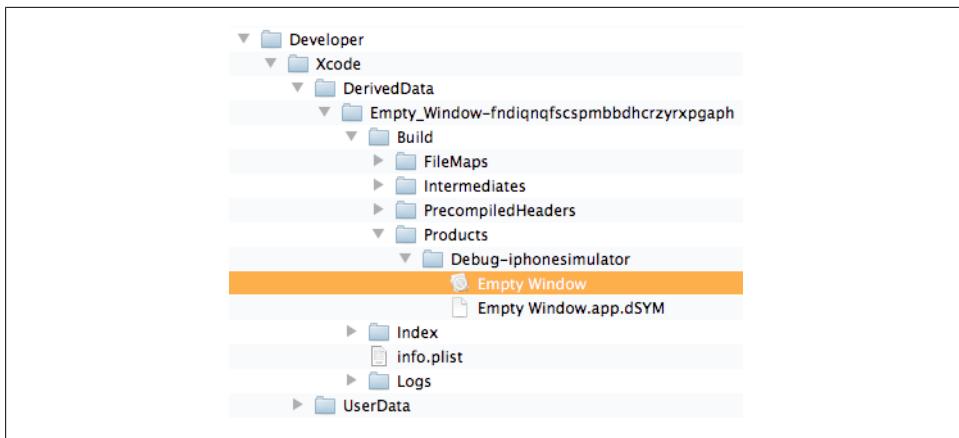


Figure 6-14. The built app, in the Finder

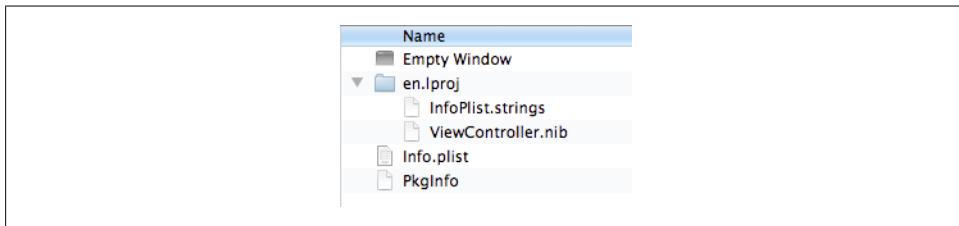


Figure 6-15. Contents of the app package

We haven't edited this file, and our app currently appears only in English, so this file is of no interest at the moment.

ViewController.nib

Currently, our app's only nib file. It contains instructions for generating the initial contents of our app's main window (currently just a grey rectangle). It is created ("compiled") from the *ViewController.xib* file in the project; a *.xib* file and a *.nib* file are different forms of the same thing.

Empty Window

Our app's compiled code (the binary). When the app is launched, the binary is linked to the various frameworks, and the code begins to run (starting with the entry point in the `main` function).

Info.plist

A configuration file in a strict text format (a *property list* file). It is derived from the project file *Empty Window-Info.plist*. It contains instructions to the system about how to treat and launch the app. For example, if our app had an icon, *Info.plist* would tell the system its name, so that the system could dive into the app bundle, find it, and display it. It also tells the system things like the name of the binary, so that the system can find it and launch the app correctly.

In real life, an app bundle will contain more files, but the difference will mostly be one of degree, not kind. For example, our project might have additional nib files, icon image files, and image or sound files. All of these would make their way into the app bundle.

You are now in a position to appreciate, in a general sense, how the components of our project are treated and assembled into an app, and what responsibilities accrue to you, the programmer, in order to ensure that the app is built correctly. The rest of this chapter outlines what goes into the building of an app from a project.

Build Settings

We have already talked about how build settings are determined. Xcode itself, the project, and the target all contribute to the resolved build setting values, some of which may differ depending on the build configuration. Before building, you, the programmer, will have already specified a scheme; the scheme determines the build configuration, the specific set of build setting values that will apply as the build proceeds.

Property List Settings

Your project contains a property list file that will be used to generate the built app's *Info.plist* file. The target knows what file it is because it is named in the Info.plist File build setting. For example, in our project, the value of the Info.plist File build setting has been set automatically to *Empty Window/Empty Window-Info.plist*. (Take a look at the build settings and see!)



Because the name of the file in your project from which the built app's *Info.plist* file is generated will vary, depending on the name of the project, I'll refer to it generically as the project's *Info.plist*.

The property list file is a collection of key–value pairs. You can edit it, and you may well need to do so. There are two main ways to edit your project's *Info.plist*:

- Select the file in the Project navigator and edit in the editor. By default, the key names (and some of the values) are displayed descriptively, in terms of their functionality; for example, it says "Bundle name" instead of the actual key, which is `CFBundleName`. But you can view the actual keys by choosing Editor → Show Raw Keys & Values (you might have to click in the editor to enable this menu item).
- Edit the target, and click Info at the top of the editor. This pane shows effectively the same information as editing the *Info.plist* in the editor.

I'm not going to enumerate all the key–value pairs you might want to edit in your project's property list file, but I'll just call attention to a few that you will almost certainly want to edit (and I'll talk about others in [Chapter 9](#) and elsewhere):

Bundle display name (CFBundleDisplayName)

The name that appears under your app’s icon on the device screen; this name needs to be short in order to avoid truncation.

Bundle identifier (CFBundleIdentifier)

Your app’s unique identifier, used throughout the development process and when submitting to the App Store. I talked earlier in this chapter about how this is derived from your company name when you create a project.

For a complete list of the possible keys and their meanings, see Apple’s document *Information Property List Key Reference*.

Nib Files and Storyboard Files

You edit a nib file (technically, this will probably be a *.xib* file) to describe graphically some objects that you want instantiated when the nib file loads ([Chapter 5](#)). Your app is likely to have at least one nib file. By breaking your interface into multiple nib files, you simplify the relationship between each nib file and your code; also, if nibs that aren’t needed when your app launches aren’t loaded until they *are* needed, you speed up your app’s launch time, and you streamline your app’s memory usage (because nib objects are not instantiated until the nib is loaded, and can then be destroyed when they are no longer needed).

Your app might have also one or more storyboard files (a *.storyboard* file). A storyboard file is like many nib files in one: in it, you describe graphically the various interfaces (called *scenes*) that you want to appear as the user works with your app. Just as with multiple nib files, a storyboard scene is transformed into actual interface only when it is needed for display, and the memory needed to maintain that interface can be given back when that interface is no longer showing. A single storyboard file may in fact replace all the nib files in your app; but there are cases where you might use both nib files and one or more storyboard files.

The target knows about your nib files because they appear in its Copy Bundle Resources build phase. In the case of a nib file in *.xib* format, the file is not merely copied into the app bundle; Xcode also translates (compiles) it into a smaller *.nib* file (using the *ibtool* tool). Similarly, Xcode translates (compiles) a *.storyboard* file into a smaller *.storyboardc* file in the built app (again, using the *ibtool* tool).

Nib files located inside your app bundle are typically loaded when they are needed as the app runs, usually because code tells them to load. If you elect to use a storyboard as the basis of your main interface, however, it will need to load before any code has a chance to do so. Such a storyboard file is called the *main storyboard file*. This situation is handled through the *Info.plist* file; it contains a key “Main storyboard file base name” (*UIStoryboardFile*), and the system sees this and loads the designated storyboard file automatically as the app launches. (Instead of a main storyboard file, it is possible to have a main nib file that loads automatically when the app launches; this was the

standard approach for apps created with Xcode 3.2.x and Xcode 4.0, but none of the Xcode 4.2 templates support this approach, so I don't discuss it in this edition of the book.)



A universal app — that is, an app that runs both on the iPad and on the iPhone — typically has nib files or storyboard files in pairs, one to be loaded on the iPad and the other to be loaded on the iPhone. Thus the app can have different basic interfaces on the two different types of device. Naming conventions and *Info.plist* keys allow the runtime to know which nib or storyboard to load depending on the device type. For example, a second *Info.plist* key, “Main storyboard file base name (iPad)” (`UIStoryboardFile~ipad`), specifies the storyboard file to be loaded at launch time on the iPad.

See [Chapter 7](#) for more details about nib files; both nib files and storyboard files, and how they are loaded and why, are discussed in detail in [Chapter 19](#).

Other Resources

Our app doesn't currently have any additional resources — not even an icon file. But if it did, the target would know about them because they appear in its Copy Bundle Resources build phase. In general, such resources would be copied unchanged into the app bundle.

With the exception of the app's icon and some images with standardized names, all of which are found and used by the system, additional resources are present because you want your running app to be able to fetch them out of its bundle. For example, if your app needs to display a certain image, you'd add the image to your project and make sure it appears in the Copy Bundle Resources build phase. When the app runs, your code (or possibly the code implied by a loaded nib file) reaches into the app bundle, locates the image, and displays it ([Chapter 15](#)).

To add a resource to your project, start in the Project navigator and choose File → Add Files to Empty Window (or whatever the name of the project is). Alternatively, drag the resource from the Finder into the Project navigator. Either way, a dialog appears ([Figure 6-16](#)) containing a pane in which you make the following settings:

Copy items into destination group's folder (if needed)

You should almost certainly check this checkbox. Doing so causes the resource to be copied into the project folder. If you leave this checkbox unchecked, your project will be relying on a file that's outside the project folder and that you might delete or change unintentionally. Keeping everything your project needs inside the project folder is far safer.

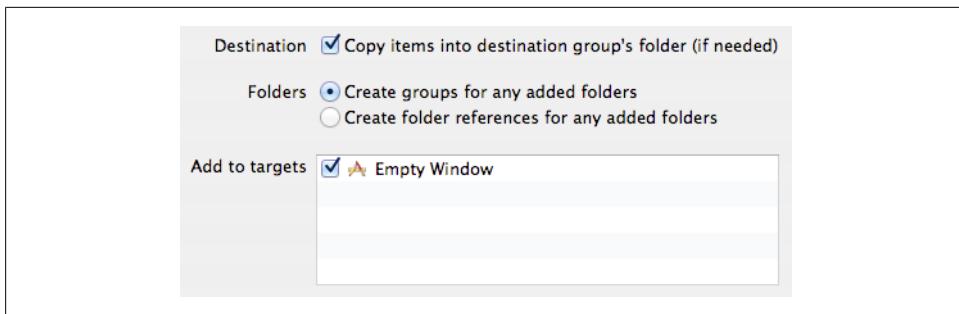


Figure 6-16. Options when adding a resource to a project

Folders

This choice matters only if what you’re adding to the project is a folder. In both cases, whether the folder is copied into the project folder depends on whether you checked the checkbox discussed in the previous paragraph; the difference is in how the project references the folder contents:

Create groups for any added folders

The folder is expressed as a group within the Project navigator, but its contents all appear individually in the Copy Bundle Resources build phase, so they will all be copied individually into the app bundle.

Create folder references for any added folders

The folder itself is shown in blue in the Project navigator and appears as a folder in the Copy Bundle Resources build phase; thus, the build process will copy the entire folder and its contents into the app bundle. This means that the resources inside the folder won’t be at the top level of the bundle, but in a subfolder of it; your code might have to specify the folder name when loading such a resource. Such an arrangement can be valuable if you have many resources and you want to separate them into categories (rather than clumping them all at the top level of the app bundle) or if the folder hierarchy among resources is meaningful to your app.

Add to Targets

Checking this checkbox causes the resource to be added to the target’s Copy Bundle Resources build phase. Thus you will almost certainly want to check it; why else would you be adding this resource to the project? But if this checkbox is unchecked and you realize later that a resource listed in the Project navigator needs to be added to the Copy Bundle Resources build phase, you can add it manually, as I described earlier.

An alternative way to copy resources from your project into the app bundle is through a custom Copy Files build phase that you add to your target. To do so, edit the target, switch to Build Phases, and click Add Build Phase (at the lower right) and choose Add Copy Files. A Copy Files build phase appears; open its triangle, and you’ll find you can

specify a custom path within the app bundle. For example, if you leave the Destination pop-up menu set to Resources and type “Pix” in the Subpath field, then any resources you add to this build phase will be copied into a folder called *Pix* in the app bundle.

A custom Copy Files build phase of this sort can be a good way of keeping resources organized by folder inside your app bundle; I frequently use it for this purpose. Bear in mind, however, that it is entirely up to you to make sure that the desired resources are placed inside the appropriate Copy Files build phase (and that they are not placed in the normal Copy Bundle Resources build phase, because if they are, you’ll end up with two copies of the resource in your app bundle).



If you copy resources into a subfolder of your app bundle, either with a folder reference or a custom Copy Files build phase, your code may have to specify that subfolder in order to fetch the resource from inside the app bundle.

Code

Code declaring two classes, *AppDelegate* and *ViewController*, was created for you when the project was created; the implementation files for these classes (*AppDelegate.m* and *ViewController.m*) appear in the target’s Compile Sources build phase. If you create any further class files, you’ll specify that they should be added to the target, and they too will then have their implementation files listed in the Compile Sources build phase. This (the contents of the Compile Sources build phase) is how your target knows what files to compile to create the app’s binary.

The binary that results from compilation of these files is your project’s *executable*, and is placed into the app bundle, with its name being by default the same as the name of the target. The app bundle’s *Info.plist* file has an “Executable file” (*CFBundle-Executable*) key whose value is the name of the binary; this is how the system knows how to locate the executable and launch the app.

Besides the class code files you create (or that Xcode creates for you), your project contains a *main.m* file. This too is in the Compile Sources build phase; it had better be, because this file contains the all-important *main* function, the entry point to your app’s code! Here are the *main* function’s contents:

```
int main(int argc, char *argv[])
{
    @autoreleasepool {
        return UIApplicationMain(argc, argv, nil,
                               NSStringFromClass([AppDelegate class]));
    }
}
```

The *main* function is very simple, but it’s crucial. It calls *UIApplicationMain*, which sets everything else in motion. *UIApplicationMain* is responsible for solving some tricky problems. As your app starts up, how will any of its code ever run? And how will its

starting repertoire of instances ever be generated? `UIApplicationMain` takes care of these issues. First, your app is a C program, and a C program's `main` function is always called, to start the program running; thus, `UIApplicationMain` will in fact be called. Then, `UIApplicationMain` does the following things:

- It creates your very first instance — the shared application instance (later accessible in code by calling `[UIApplication sharedApplication]`). The third argument to `UIApplicationMain` specifies, as a string, what class the shared application instance should be an instance of. If nil, which will usually be the case, the default class is `UIApplication`; but you can make a subclass of `UIApplication` and specify that subclass here by substituting something like this (depending on what the subclass is called) as the third argument:
`NSStringFromClass([MyUIApplicationSubclass class])`

- Optionally, it also creates your second instance — the application instance's *delegate*. Delegation is an important and pervasive Cocoa pattern, described in detail in [Chapter 10](#), but for now let's just say that it is crucial that every app you write have an app delegate instance. The fourth argument to `UIApplicationMain` specifies, as a string, what class the app delegate instance should be. If this class is specified, as here, `UIApplicationMain` instantiates that class and ties that instance to the shared application instance as the latter's delegate. If this class is *not* specified (the fourth argument is nil), it is up to you to provide a delegate instance in some other way; since you cannot do this in code, you would have to do it through the loading of the main nib file. (Before iOS 5 and Xcode 4.2, this was in fact the usual way in which the app delegate was instantiated; but Apple has now changed the default pattern so that the app delegate is generated in code by the call to `UIApplicationMain`.)
- If the `Info.plist` file specifies a main storyboard file or main nib file, `UIApplicationMain` loads it.
- An app delegate instance has now been generated, either because `UIApplicationMain` instantiated it directly in response to the value of its fourth argument, or because `UIApplicationMain` loaded a main nib file which instantiated it. `UIApplicationMain` now turns to this app delegate instance and starts calling some of its code — in particular, it calls `application:didFinishLaunchingWithOptions:`, which is typically responsible, in turn, for displaying your app's initial interface.
- The app is now launched and visible to the user. `UIApplicationMain` is still running (like Charlie on the M.T.A., `UIApplicationMain` never returns), and is now just sitting there, watching for the user to do something, maintaining the *event loop*, which will respond to user actions as they occur.

The call to `UIApplicationMain` is wrapped in some memory management functionality (the `@autoreleasepool` curly braces) that I'll explain in [Chapter 12](#).

Finally, notice the file `Empty Window-Prefix.pch` in the Project navigator. This is your project's *precompiled header* file. It isn't listed in the Compile Sources build phase be-

cause it is actually compiled *before* that build phase; the target knows about it because it is pointed to by the Prefix Header build setting.

The precompiled header is a device for making compilation go faster. It's a header file; it is compiled once (or at least, very infrequently) and the results are cached (off in `/var/folders/`) and are implicitly imported by all your code files. So the precompiled header should consist primarily of `#import` directives for headers that never change (such as the built-in Cocoa headers); it is also a reasonable place to put `#defines` that will never change and that are to be shared by all your code.

The default precompiled header file imports `<Foundation/Foundation.h>` (the Core Foundation framework header) and `<UIKit/UIKit.h>` (the Cocoa framework).

Frameworks and SDKs

A *framework* is a library of compiled code used by your code. Most of the frameworks you are likely to use will be Apple's built-in frameworks; they are built-in in the sense that they are part of the system on the device where your app will run — they live in `/System/Library/Frameworks`, though you can't tell that on an iPhone or iPad because there's no way (normally) to view the filesystem directly.

However, your code needs to use these frameworks not only when running on a device but also when building and when running in the Simulator. To make this possible, part of the device's system — in particular, the part containing its frameworks — is duplicated on your computer, in the `/Developer` folder. This duplicated subset of the device's system is called an *SDK* (for "software development kit") and is something you can see directly in the Finder. For example, look at `/Developer/Platforms/iPhoneOS.platform/Developer/SDKs/iPhoneOS5.0.sdk/System/Library/Frameworks`; behold, there are the frameworks included on a device running iOS 5.0.

To use a framework in your code, you must do two things:

Import the framework's header

A framework has a header file, which provides (usually by importing other header files within the framework) the interface information about classes in that framework. Your code needs this information in order to *compile* successfully. You import the header with an appropriate `#import` directive.

Link to the framework

A framework is a package; you must instruct the build system to associate this package with your app's executable binary, so that your binary's calls to code within that framework can be routed into the framework's compiled code. This is necessary in order for your app to *run* successfully. Such an association is called *linking* the binary with the framework, and you instruct the build system to do this by including the framework in the target's Link Binary With Libraries build phase.

You might think that linkage is impossible because the framework to which we ultimately want to link is off on a target device somewhere. But linkage is path-based, and the path is determined relative to the current SDK. Thus, the linkage to the UIKit framework uses the path *System/Library/Frameworks/UIKit.framework*. This path is relative to the current SDK, so if you’re using the iOS 5.0 SDK, the path during development will be */Developer/Platforms/iPhoneOS.platform/Developer/SDKs/iPhoneOS5.0.sdk/System/Library/Frameworks/UIKit.framework*.

But when the app runs on the device, there is no SDK, and the path becomes absolute, starting at the top level of the device. Thus, when the app runs in the Simulator, the framework is found successfully on your computer, and when the app runs on a device, the framework is found successfully on the device.

By default, three frameworks are linked into your target:

Foundation

Many basic Cocoa classes, such as `NSString` and `NSArray` and others whose names begin with “NS,” are part of the Foundation framework. The Foundation framework is imported in the precompiled header file (and, by default, in the headers of new classes that you create). In turn, it imports the Core Foundation headers and loads the Core Foundation framework as a subframework; thus, there is no need for you to import or link explicitly to the Core Foundation framework (which is full of functions and pointer types whose names begin with “CF,” such as `CFStringRef`).

UIKit

Cocoa classes that are specialized for iOS, whose names begin with “UI,” are part of the UIKit framework. The UIKit framework is imported in the precompiled header file (and by templated class code files such as `AppDelegate.h`).

Core Graphics

The Core Graphics framework defines many structs and functions connected with drawing, whose names begin with “CG.” It is imported by many UIKit headers, so you won’t need to import it explicitly.

You might find that the three default frameworks are sufficient to your needs, or you might find that you need other frameworks to provide additional functionality. How will you know that a class or function you want to use resides outside the three default frameworks? You might get a clue from its name, which won’t begin with “NS,” “UI,” or “CG”, but more often, if you’re like me, you’ll be alerted by banging up against the compiler.

For example, let’s say you’ve just found out about animation ([Chapter 17](#)) and you’re raring to try it in your app. So, in your code, you create a `CABasicAnimation`:

```
CABasicAnimation* anim = [CABasicAnimation animation];
```

The next time you try to build your app, the compiler complains that `CABasicAnimation` is undeclared (and that it therefore can’t make sense of `anim` either). That’s when

you realize you need to import a framework header. Near the start of the CABasicAnimation class documentation is a line announcing that it's in *QuartzCore.framework*. You might guess (correctly) that the way to import the main Quartz Core framework header is to put this line near the start of your implementation file:

```
#import <QuartzCore/QuartzCore.h>
```

This works to quiet the compiler. Remember, though, that I said that using a framework requires two things; we've done only one of them. So your code *still* doesn't build. This time, you get a build error during the link phase of the build process complaining about `_OBJC_CLASS_$_CABasicAnimation` and saying, "Symbol(s) not found." That mysterious-sounding error merely means that you've forgotten to link your target to the Quartz Core framework.

To link your target to a framework, edit the target, click Summary at the top of the editor, and scroll down to the Linked Frameworks and Libraries section. (This is the same information that appears which you click Build Phases at the top of the editor and open the Link Binary with Libraries build phase.) Click the "+" button at the left just below the frameworks. A dialog appears nicely listing the existing frameworks that are part of the active SDK. Select *QuartzCore.framework* and click Add. The Quartz Core framework is added to the target's Link Binary With Libraries build phase. (It also appears in the Project navigator; you might like to drag it manually into the Frameworks group, for the sake of neatness.) Now you can build (and run) your app.

You might wonder why the project isn't linked by default to *all* the frameworks, so that you don't have to go through this process every time you stray beyond the default three frameworks. It's just a matter of time and resources. Importing headers increases the size of your code; linking to frameworks slows down your app's launch time. You should link to only the frameworks needed for your code to run.

Where you import a framework header depends on how you intend to use it. It's simply a matter of scope. If a framework's classes are to be mentioned only within a single implementation file, then you can import it at the start of that implementation file. If you want to subclass one of the framework's classes, you'll need to import it at the start of the interface file that declares the subclass; in that case, every implementation file that imports this interface file imports the framework header, and there's no need to import the framework header separately in the implementation file. Of course, for maximum scope, you can simply import the framework header in the precompiled header file, making that framework available throughout your code.

Renaming Things

The name assigned to your project at creation time is used in many places throughout the project, leading beginners to worry that they can never rename a project without breaking something. But fear not! To rename a project, select the project listing at the top of the Project navigator, press Return to make its name editable, type the new name, and press Return again. Xcode presents a dialog proposing to change some other names to match, including the target, the built app, the precompiled header file, and the `Info.plist` — and, by implication, the build settings specifying these. You can check or uncheck any name, and click Rename; your project will continue to work correctly.

You can freely change the target name independently of the project name. It is the target name, not the project name, that is used to derive the name of the product and thus the bundle name, bundle display name, and bundle identifier mentioned earlier in this chapter. Thus, when you settle on a real name for your app, it might be sufficient to set the target name.

Changing the project name (or target name) does not automatically change the scheme name to match. There is no particular need to do so, but you can change a scheme name freely; choose Product → Manage Schemes and click on the scheme name to make it editable.

Changing the project name (or target name) does not automatically change the main group name to match. There is no particular need to do so, but you can freely change the name of a group in the Project navigator, because these names are arbitrary; they have no effect on the build settings or the build process. However, the main group is special, because (as I've already said) it corresponds to a real folder on disk, the folder that sits beside your project file at the top level of the project folder. You can change the group's name (changing the project name does not do this for you automatically), but you should not delete it, and beginners should not change the name of the folder on disk to which it corresponds, as that folder name is hard-coded into several build settings.

You can change the name of the project folder in the Finder at any time, and you can move the project folder in the Finder at will, because all build setting references to file and folder items in the project folder are relative.

If you want to change the name of a class or variable, Xcode can assist you with its Refactoring and Edit All In Scope features ([Chapter 9](#)).

Nib Management

A *nib file*, or simply *nib*, is a file containing a drawing of a piece of your interface. The term *nib* is not really an English word (it has nothing to do with fountain pens, for example); it is based on the file extension *.nib* that is used to signify this type of file, an extension that originated as an acronym (for “NeXTStep Interface Builder”). Nowadays, you will usually develop your interface using a file format whose extension is *.xib*; when your app is built, your target’s *.xib* files are translated (“compiled”) into *.nib* format ([Chapter 6](#)). But a *.xib* file is still referred to as a nib file. I will speak of the same nib file as having either a *.xib* extension (if you’re editing it) or a *.nib* extension (if it’s in the built app).

You construct your program in two ways — writing code, and drawing the interface. But these are really two ways of accomplishing the same ends; drawing the interface *is* a way of writing code. When the app runs and your drawing of the interface in a nib file is loaded, it is translated into instructions for instantiating and initializing the objects in the nib file. You could equally have instantiated and initialized those same objects in code. (This point is crucial; see “[Nib-Based Instantiation](#)” on page 81.) Indeed, deciding whether to create an interface object in code or through a nib file is not always easy; each approach has its advantages. The important thing is to understand how interface objects drawn in a nib file are instantiated and connected to your code when the app runs.

(This chapter applies in almost all details equally to storyboards. So do not, under any circumstances, skip this chapter on the grounds that you intend to use storyboards instead of nibs! Storyboards do not relieve you of the need to understand nib management thoroughly; and in any case, an Xcode programming life without nibs is still extremely improbable. I’ll address the use of storyboards in [Chapter 19](#).)



Up through Xcode 3.2.x, nib editing was performed in a separate application, Interface Builder. Starting in Xcode 4, the functionality of Interface Builder was rolled into Xcode itself. Nevertheless, the Xcode interface for nib editing is still commonly referred to as Interface Builder.

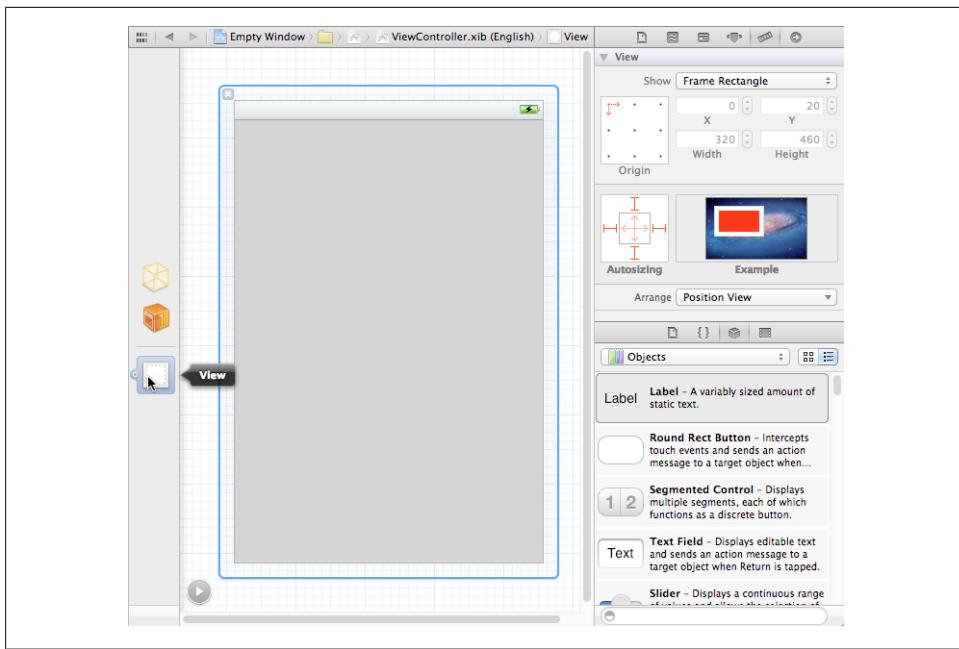


Figure 7-1. Editing a nib file

A Tour of the Nib-Editing Interface

Let's use an actual nib file to explore the Xcode nib-editing interface. In [Chapter 6](#), we created a simple Xcode project, Empty Window; it contains a nib file, so we'll use that. In Xcode, open the Empty Window project, locate the *ViewController.xib* listing in the Project navigator, and click it to edit it.

[Figure 7-1](#) shows the project window after selecting *ViewController.xib* and making some additional adjustments. The Navigator pane is hidden; the Utilities pane is showing. Within the Utilities pane, the Size inspector and the Object library are showing. The interface may be considered in four pieces:

1. At the left of the editor is the *dock*, showing the nib's top-level objects. The dock can be expanded by dragging its right edge or by clicking the right-pointing triangle-in-a-circle at its lower right; then it shows *all* of the nib's objects hierarchically.
2. The remainder of the editor is devoted to the *canvas*, where you physically design your app's interface. The canvas portrays views in your app's interface and things that can contain views. (A *view* is an interface object, which draws itself into a rectangular area. The phrase "things that can contain views" is my way of including view controllers, which are represented in the canvas even though they are not drawn in your app's interface.)

3. The inspectors in the Utilities pane are where you view and edit details of the currently selected object.
4. The libraries in the Utilities pane, especially the Object library, are your source for interface objects to be added to the nib.

The Dock

The dock, as I've already said, shows the nib's *top-level objects*. To see what this means, you need first to envision the nib as containing objects. Some of these objects — those that represent views — are arranged in a hierarchy of containment. Objects that are contained by no other object are top-level objects.

A view can contain other views (its *subviews*) and can be contained by another view (its *superview*); for example, a button might be a subview of a window, and that window would be that button's superview. One view can contain many subviews, which might themselves contain subviews. But each view can have only one immediate superview. Thus there is a hierarchical tree of subviews contained by their superviews with a single object at the top. The highest superview of any such hierarchy in the nib is a top-level object and appears in the dock. That's why the view object (labeled View in [Figure 7-1](#)) appears in this nib's dock: it is a view contained by no other view.

A nib file can actually contain two types of top-level object:

Placeholders (proxy objects)

A placeholder, or *proxy object*, represents an object that already exists in your app's code at the time the nib is loaded. Proxy objects appear in a nib file chiefly so that you can provide communication between objects in your app's code and objects instantiated from the nib. You can't create or delete a proxy object; the dock is populated automatically with them. Proxy objects are shown above the dividing line in the dock.

Nib objects

A nib object is an object that is instantiated by the nib — that is, the instance it represents will be created when your code runs and the nib loads. You can create new nib objects. Top-level nib objects are shown below the dividing line in the dock.

The dock can be expanded (by clicking the right-pointing triangle-in-a-circle at its lower right); it then portrays objects by name (label), and shows as an outline the full hierarchy of objects in this nib ([Figure 7-2](#)). At present, expanding the dock may seem silly, because there is no hierarchy; all objects in this nib are top-level objects. But when a nib contains many levels of hierarchically arranged objects, you're going to be very glad of the ability to survey them all in a nice outline, and to select the one you're after, thanks to the expanded dock. You can also rearrange the hierarchy here; for example, if you've made an object a subview of the wrong view, you can drag it onto the view it should be a subview of within this outline.



Figure 7-2. The dock, expanded

You can also select objects using the jump bar at the top of the editor. First, click on the canvas background so that no object is selected; the entire hierarchy of the objects in your nib is then shown as a set of hierarchical menus off the rightmost jump bar path component (Control-6). Again, this may seem like small potatoes now, when your nib contains just four top-level objects and nothing more, but it will be valuable when you've many nib objects in a hierarchy.

The names (labels) by which nib objects are designated are meaningful only while editing a nib file; they have no relationship to your code. When the dock is expanded, each object is portrayed by its label, as shown in [Figure 7-2](#). When the dock is collapsed, you can see a top-level object's label by hovering the mouse over it, as shown in [Figure 7-1](#). If you find an object's label unhelpful, you can change it: select the object and edit the Label field (whose placeholder reads “Xcode Specific Label”) in the Identity section of the Identity inspector (Command-Option-3).

Canvas

The canvas presents a graphical representation of a top-level nib object along with its subviews, similar to what you're probably accustomed to in any drawing program. If a top-level nib object has a graphical representation (not every top-level nib object has one), you can click on it in the dock to display that representation in the canvas. A little dot to the left of a top-level object in the collapsed dock indicates that it is currently being displayed graphically in the canvas.

To remove the canvas representation of a top-level nib object, click the “x” at its upper left; this merely clears the representation from the canvas — it does not remove the top-level nib object from the dock (or from the nib), and of course you can always bring back the graphical representation by clicking that nib object in the dock again. On the other hand, the canvas is scrollable and automatically accommodates all graphical representations within it, so you can keep as many graphical representations open in the canvas as you like, side by side, and scroll to see each one, regardless of the size of your monitor; thus you might never need to remove the canvas representation of a top-level nib object at all.

Our simple Empty Window project's *ViewController.xib* contains just one top-level nib object that has a graphical representation — the root view of the app's window, called View. (The term “root” here implies that the view occupies the entire window.) Because

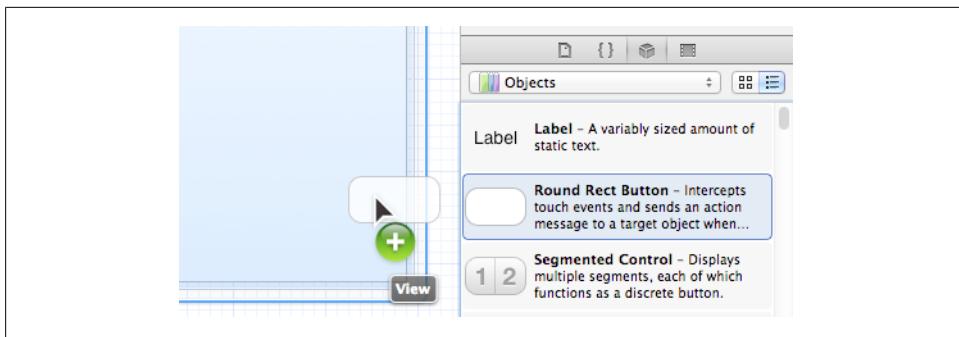


Figure 7-3. Dragging a button into a window

this view appears in our app’s window, any changes you make here will be reflected in the app’s user interface when you run it. To see this, we’re going to add a subview to it:

1. Ensure that the View in the dock is being displayed in the canvas.
2. Look at the Object library (Control-Option-Command-3). Click the second button in the segmented control to put the Object library into list view, if it isn’t in list view already. Locate the Round Rect Button (you can type “button” into the filter bar at the bottom of the library as a shortcut).
3. Drag the Round Rect Button from the Object library into the View in the canvas (Figure 7-3). Don’t accidentally drop the button onto the canvas background, outside of the window! This would cause the button to become a top-level object, which is not what you want. If that happens, select the button in the dock and press Delete, and try again.

A button now appears in the window in the canvas. The move we’ve just performed —dragging from the Object library into the canvas— is extremely characteristic; you’ll do it often as you design your interface. Here are two alternative ways to do the same thing:

- Double-click an object in the Object library; if a view (such as our window) is already selected in the canvas, a copy of that object becomes a subview of it.
- Type some part of an object’s name in the filter bar; you can then use arrow keys to select the correct object, if needed, and finally press Return to copy the object into the canvas. You can switch to the Object library with Control-Option-Command-3, and this also puts focus in the filter bar, so the whole operation can be performed with the keyboard.

Next, play around with the button in the window. Much as in a drawing program, the nib editor provides features to aid you in designing your interface. Here are some things to try:

- Select it: resizing handles appear.
- Resize it to make it wider: dimension information appears.

- Drag it near the edge of the window: a guideline appears, showing a standard margin space between the edge of the button and the edge of the window.
- With the button selected, hold down the Option key and hover the mouse outside the button: arrows and numbers appear showing the pixel distance between the button and the edges of the window. (If you accidentally clicked and dragged while you were holding Option, you'll now have two buttons. That's because Option-dragging an object duplicates it. Select the unwanted button and press Delete to remove it.)
- Shift-Control-click on the button: a menu appears, letting you select the button or whatever's behind it (in this case, the window).

Let's prove that we really are designing our app's interface. We'll run the app to see that its interface has changed.

1. Make sure that the Breakpoints button in the project window toolbar is *not* selected, as we don't want to pause at any breakpoints you may have created while reading the previous chapter.
2. Make sure the destination in the Scheme pop-up menu is the iPhone Simulator.
3. Choose Product → Run (or click the Run button in the toolbar).

After a heart-stopping pause, the iOS Simulator opens, and presto, our empty window is empty no longer ([Figure 7-4](#)); it contains a round rect button! You can tap this button with the mouse, emulating what the user would do with a finger; the button highlights as you tap it.

Inspectors and Libraries

There are four inspectors that appear only when you're editing a nib and apply to whatever object is selected in the dock or canvas:

Identity inspector (Command-Option-3)

Far and away the most important section of this inspector is the first one, the Custom Class. The selected object's Class setting tells you the object's class, and you can use it to change the object's class. Some situations in which you'll need to change the class of an object in the nib appear later in this chapter.

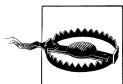
Attributes inspector (Command-Option-4)

Settings here correspond to properties and methods that you might use to configure the object in code. For example, changing the setting in the Background pop-up menu in the Attributes inspector for our view corresponds to setting the `backgroundColor` property for the view in code. Similarly, typing a value in the Title field in the Attributes inspector for our button is like calling the button's `setTitle:forState:` method.



Figure 7-4. The Empty Window app's window is empty no longer

The Attributes inspector has sections corresponding to the selected object's class inheritance. For example, the UIButton Attributes inspector has three sections, because a UIButton is also a UIControl ("Control" in the inspector) and a UIView ("View" in the inspector).



The correspondence between Attributes inspector settings and Objective-C methods is mostly a matter of guesswork. The Attributes inspector doesn't always tell you, and there's no way to see the code generated when the nib actually loads.

Size inspector (Command-Option-5)

The X, Y, Width, and Height fields determine the object's frame (its position and size within its superview), corresponding to its `frame` property in code; you can equally do this in the design window by dragging and resizing, but numeric precision can be desirable. The Autosizing box corresponds to the `autoresizingMask` property, determining how the object will be repositioned and resized when its superview is resized; a delightful animation demonstrates visually the implications of your settings. The Arrange pop-up menu contains useful commands for positioning the selected object.

Connections inspector (Command-Option-6)

I'll discuss this later in this chapter.

There are two libraries that are of particular importance when you're editing a nib:

Object library (Control-Option-Command-3)

This library, as we've already seen, is your source for types of object that you want to copy into the nib.

Media library (Control-Option-Command-4)

This library lists media in your project, such as images that you might want to drag into a UIImageView or directly into your interface (in which case a UIImageView is created for you).

Nib Loading and File's Owner

A nib file is useless until your app runs and the nib file is *loaded*. If a nib is designated by the *Info.plist* key “Main nib file base name” (`NSMainNibFile`, see [Chapter 6](#)), it is loaded automatically as the app launches; but this is an exceptional case, and in Xcode 4.2 it has fallen out of favor — there are no automatically loaded main nib files in the project templates. In general, nibs are loaded explicitly as needed while the app runs. In our Empty Window application, you can actually see where this happens, in *AppDelegate.m*:

```
self.viewController =  
    [[ViewController alloc] initWithNibName:@"ViewController" bundle:nil];
```

That line of code does several things, one of which is that (for reasons to be explained more fully in [Chapter 19](#)) it causes the nib named @"ViewController" (i.e. the nib file compiled from *ViewController.xib*, the nib file we've been editing) to be loaded, and the resulting views to be put into our app's interface — which is how we were able to obtain the outcome shown in [Figure 7-4](#).

So a nib is not loaded until the app runs and our code decides, at some point in the life of the app, that that nib is needed. This architecture is a source of great efficiency. For example, imagine our app has two complete sets of interface, and the user might never ask to see the second one. It makes obvious sense not to load a nib containing the second set of interface until the user *does* ask to see it. By this strategy, a nib is loaded when its instances are needed, and those instances are destroyed when they are no longer needed. Thus memory usage is kept to a minimum, which is important because memory is at a premium in a mobile device. Also, loading a nib takes time, so loading fewer nibs at launch time makes launching faster.

When a nib loads, some already existing instance is designated its owner. A nib cannot load without an owner, and the owner must exist before the nib can load. The owner will often be a UIViewController instance, because a UIViewController already knows how to load a nib and manage a view that it contains ([Chapter 19](#)); and indeed, this is

the case when our *ViewController.xib* file is loaded: an instance of the *ViewController* class, which is a *UIViewController* subclass, is created precisely to act as its owner. But a nib owner can be an instance of any class, and it is important to be conscious of that fact.

The File’s Owner top-level object in a nib file is a proxy for the instance that will be the nib’s owner when the nib loads, and its class should be set to that instance’s class. In the case of our Empty Window project’s *ViewController.xib*, the File’s Owner’s class has been correctly set in advance: its class is *ViewController* (do you see how to confirm this in the nib editor’s Identify inspector?), corresponding to the fact that a *ViewController* instance will be the nib’s owner when it loads. For nibs that you create, the File’s Owner’s class might not be set correctly, and you’ll have to set it yourself using the Identity inspector.

Let’s look at that line of code from *AppDelegate.m* once again:

```
self.viewController =  
    [[ViewController alloc] initWithNibName:@"ViewController" bundle:nil];
```

I mentioned earlier that it does several things. The first thing it does (by performing the “alloc-init” dance) is to instantiate *ViewController*. The second thing it does, by using *initWithNibName:bundle:* as the initializer for that new *ViewController* instance, is to tell that instance to load the nib named @"*ViewController*" *with itself as owner*. The class of the actual owner of the nib at the moment it loads thus corresponds to the class of the File’s Owner proxy in that nib.

When a nib loads, its nib objects are instantiated, meaning its top-level nib objects and all deeper-level nib objects hierarchically dependent on them. (Proxy objects, by definition, exist before the nib loads; nib loading does not instantiate them.) For example, in our nib, the view is instantiated when the nib loads, bringing with it the button inside it. (Again, see “[Nib-Based Instantiation](#)” on page 81; make very sure you understand this point!) This is what nibs are for — to instantiate objects when they load. To put it another way, that is what nib loading is — it is the instantiation of the nib objects described in the nib. At that point, having loaded, the nib’s work is done; the nib does not, for example, have to be “unloaded.”



The same nib can be loaded multiple times, generating an entirely new set of instances each time. A common beginner question is, “I have a view in a nib; how do I make multiple copies of this view?” The simple solution is to load that nib multiple times. This is common practice. For example, consider table view cells. Every “row” of a table view is a table view cell. Let’s say there’s a certain look and behavior you want each “row” to have. You design the cell in a nib of its own as a *UITableViewCell*. If the table has to display ten rows, you load that nib ten times ([Chapter 21](#)).

Making and Loading a Nib

Let's create our own nib-loading code, illustrating at the same time the fact that any instance can be a nib's owner. To do so, we'll need a second nib file in our project.

First, we'll make the nib:

1. Choose File → New → New File.
2. At the left of the dialog, under iOS (not Mac OS X!), choose User Interface, and select View in the main part of the dialog. Click Next.
3. For the Device Family, specify iPhone. Click Next.
4. Name the file *MyNib*; make sure you're saving into the Empty Window project folder, that the group is Empty Window, and that the target is Empty Window (and checked). Click Save.

We've now created a nib file, *MyNib.xib*, containing a single top-level nib object, a `UIView`. Look at *MyNib.xib* in the editor to see that this is true.

We'll also need an instance to act as the nib's owner. By the time our code will run, we will already have at least one instance we could use (the `AppDelegate` instance), but to illustrate the procedure fully, we'll create our own class whose sole purpose is to be instantiated so that this instance can act as the owner of the nib file as it loads:

1. In the Empty Window project in Xcode, choose File → New → New File. The "Choose a template" dialog for files appears.
2. At the left of the dialog, under iOS (not Mac OS X!), select Cocoa Touch, and select Objective-C Class in the main part of the dialog. Click Next.
3. The dialog now offers you a chance to specify what superclass the new class should be a subclass of. Make sure this is `NSObject`. Click Next.
4. Name the file *MyClass*; make sure you're saving into the Empty Window project folder, that the group is Empty Window, and that the target is Empty Window (and checked). Click Save.

We've now created files *MyClass.h* and *MyClass.m* declaring a class called *MyClass*.

Next, we'll write code that will load our new nib when the app runs. We need a place in our little app where our code is guaranteed to run: we'll use the `AppDelegate` instance method `application:didFinishLaunchingWithOptions:` (in the file *AppDelegate.m*). Just before or after the call to `makeKeyAndVisible`, insert this code to instantiate *MyClass* and load *MyNib.nib* with that instance as its owner:

```
 MyClass* mc = [[MyClass alloc] init];
 [[NSBundle mainBundle] loadNibNamed:@"MyNib" owner:mc options:nil];
```

Xcode will complain about this, because you can't speak of *MyClass* without importing its declaration, so after the existing `#import` at the start of this file, add this line:

```
#import "MyClass.h"
```

Now build and run the project. Our new *MyNib.nib* file loads, and its `UIView` top-level nib object is instantiated. Unfortunately, you can't *see* that this is true! The next section explains how to obtain visible proof that our nib is loading and that its top-level nib objects are being instantiated.

Outlet Connections

You know how to load a nib file, thus instantiating its top-level nib objects. But those instances are useless to you if you don't know how to get a *reference* to any of them in your code! Doing things with an object such as a label or a button or a text field or whatever (such as setting or getting the text it displays) is easy; but you have to be able to talk *to* the object in the first place, meaning that you need a reference to it, a variable that points to that instance ([Chapter 3](#)). Getting a reference to an instance that you created in code is trivial, because you assigned it to a variable at the time you created it ([Chapter 5](#)). But there's no such assignment when you load a nib; you just load it and that's the end of that:

```
[[NSBundle mainBundle] loadNibNamed:@"MyNib" owner:mc options:nil];
// no assignment??!! dude, where are my nib-created instances?
```

To refer in code to instances generated from nib objects when the nib loads, you need to have previously set up an outlet connection from a proxy object in the same nib.

A *connection* is a named unidirectional linkage from one object in a nib file (the connection's source) to another object in the same nib file (the connection's target). An *outlet* is a connection whose name corresponds to an instance variable in the source object. When the nib loads, and the target object is instantiated, the value of the instance variable is set to the target object. Thus the source object winds up with a reference to the target object as the value of one of its instance variables.

Connections can link any two objects in a nib file, but a proxy object as the source of a connection is special because it represents an object that exists before the nib loads. Thus an outlet from a proxy object causes an object that exists *before* the nib loads to end up with a reference to an object that doesn't exist until *after* the nib loads — an object that is in fact instantiated by the loading of the nib.

In the most typical configuration, the proxy object will be the File's Owner. The idea is that the instance that owns the nib has an instance variable, and the File's Owner in the nib has a corresponding outlet to a nib object; the nib loads, and the owner instance ends up with an instance variable that refers to the instance generated from the nib object ([Figure 7-5](#)).

To demonstrate, we'll implement exactly the schema illustrated in [Figure 7-5](#), by making an outlet from the File's Owner to a nib object in *MyNib.xib*. First, we need a nib object in *MyNib.xib* to make an outlet to. For visual impact, we'll replace the existing top-level view with a top-level label, which will draw some text:

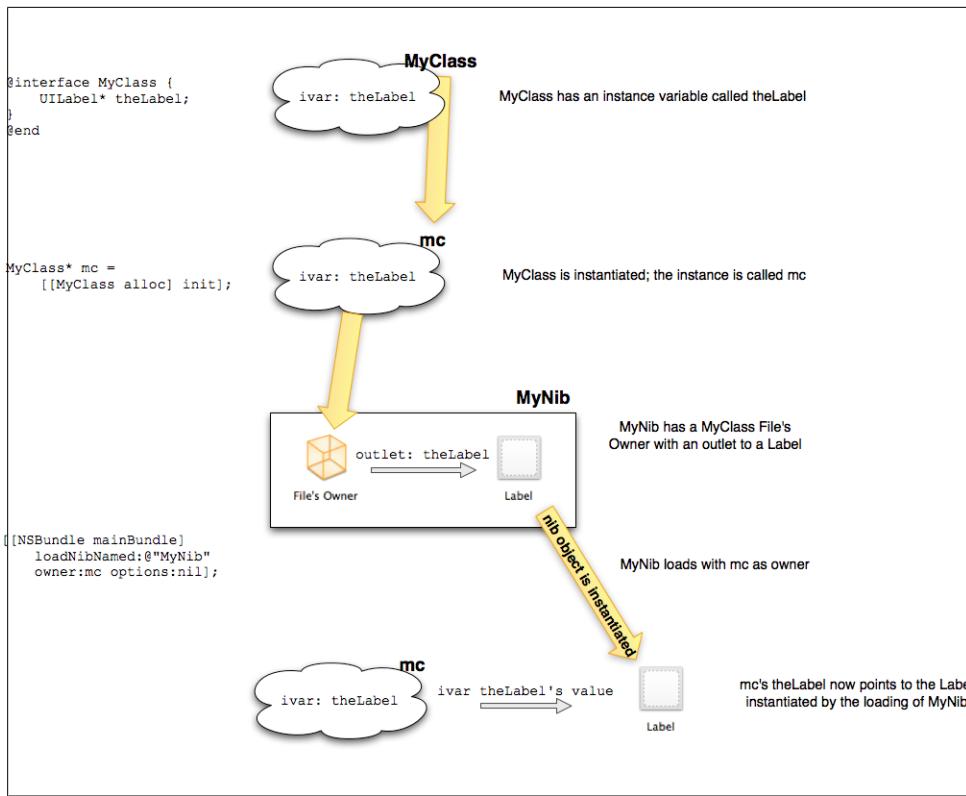


Figure 7-5. How an outlet provides a reference to a nib-instantiated object

1. In Xcode, click `MyNib.xib` to edit it.
2. In the dock, select the View object and delete it.
3. Drag a Label object (`UILabel`) from the Object library into the dock or the canvas to become a new top-level object. Its graphical representation appears in the canvas.
4. Double-click the word “Label” in the label’s graphical representation in the canvas and type “Hello, world!” Hit Return to stop editing and to make the label the size of its text.

The object that will own the nib file when it loads is a `MyClass` instance. But the nib doesn’t know this; we need to tell it:

1. Select the File’s Owner proxy object and look at the Identity inspector.
2. The Class, under Custom Class, is `NSObject`. Change this to `MyClass`. (If you type “My,” the word “`MyClass`” should just appear, as it’s the only class Xcode knows about whose name starts with “My.” Accept this by pressing Return.)

Now comes the really crucial part. We need two things, in two different places:

The instance variable

In its code, MyClass needs an instance variable whose value will be the label.

The outlet

In the nib, the File’s Owner proxy, representing a MyClass instance, needs an outlet pointing at the label — an outlet with the same name as the instance variable.

When the app runs and *MyNib.nib* is loaded with a MyClass instance as its owner, as we arranged in the preceding section, those two pairs of things will be effectively equated:

- The MyClass instance will be equated with the File’s Owner proxy in the nib, because it *will* be the nib’s owner as it loads.
- MyClass’s instance variable will be equated with the File’s Owner outlet pointing at the label, because *they have the same name*.



I’m oversimplifying. It isn’t really the identity of an instance variable’s name with that of the outlet that makes the match. It’s more complicated than that; the match is made using key-value coding. The rigorous details appear in [Chapter 12](#).

You thus need to work in two places at once: the nib, and MyClass’s code. Before Xcode 4, this required working separately in two *different* places, Xcode (where the code was edited) and Interface Builder (where the nib was edited). But in Xcode 4, the same program edits both the code and the nib, and furthermore you can see the code and the nib *at the same time*, all of which will make creating this pair of things, the instance variable and the outlet, much easier than it once was.

I want you now to arrange to see two things at once: *MyClass.m* (the MyClass implementation file, where we’ll declare the instance variable) and *MyNib.xib* (where we’ll create the outlet). You could use two project windows if you wanted, but for simplicity, let’s use an assistant: while editing *MyNib.xib*, switch to Assistant view (View → Assistant Editor → Show Assistant Editor) as in [Figure 7-6](#). If, when you showed the assistant pane, it didn’t appear with MyClass’s header file showing, use the jump bar in the assistant pane to make the assistant pane show *MyClass.m*.

In *MyClass.m* (in the assistant pane), at the start of the implementation section, create curly braces and declare a UILabel instance variable:

```
@implementation MyClass
{
    IBOutlet UILabel* theLabel;
}
@end
```

The term **IBOutlet** is linguistically meaningless; it is #defined as an empty string, so it is deleted before the compiler ever sees it. It’s purely a hint to Xcode to make it easy for you to create the outlet. Xcode responds by displaying an empty circle in the gutter

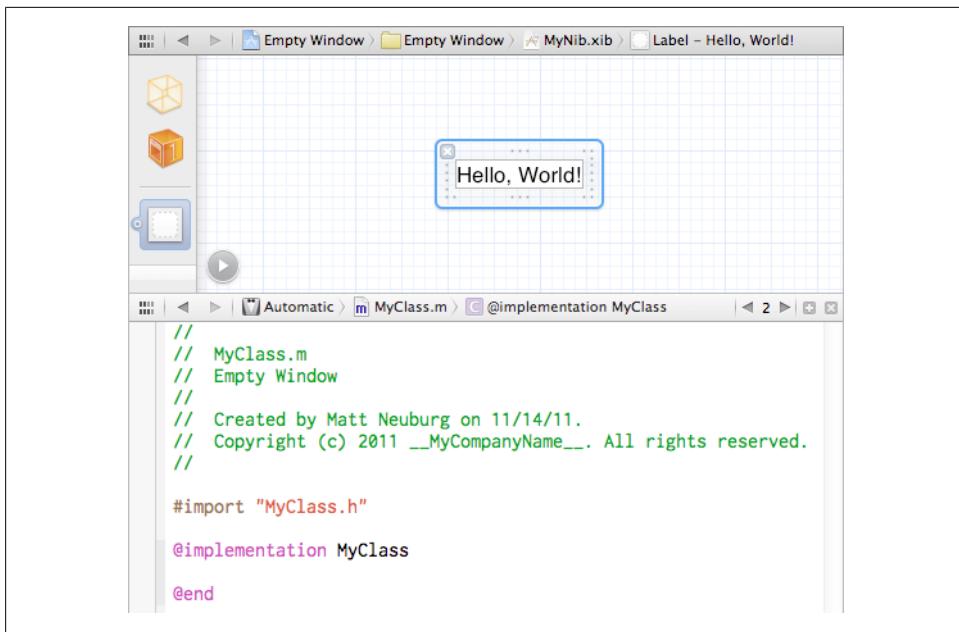


Figure 7-6. Editing a nib, with code in the assistant pane

to the left of the `IBOutlet` line; this indicates that although we're speaking of an outlet in our code, no corresponding outlet connection yet exists in a nib. We'll fix that in a moment.

We have typed the instance variable as a `UILabel*`, because we happen to know that this is the type of object that this instance variable will be pointing to; we could also use `id`, or any superclass of `UILabel`. If we do not use one of these alternatives (`id`, `UILabel`, or a superclass of `UILabel`), we will not be able to form the connection to a `UILabel` in the nib.

We have accomplished half our task: we've made the instance variable. Now we're ready for the other half, namely, to make the outlet connection. There are several ways to do this, so I'll just pick one for now and demonstrate the others later:

1. Select File's Owner in the nib (which, you remember, represents a `MyClass` instance) and switch to the Connections inspector. Lo and behold, the name of our instance variable, `theLabel`, is listed here! This is the work of the `IBOutlet` hint we typed earlier.
2. Click in the empty circle to the right of `theLabel` in the Connections inspector, drag to the Label object in the canvas (Figure 7-7), and release the mouse. (A kind of elastic line follows the mouse as you drag from the circle to show that you're creating a connection.)

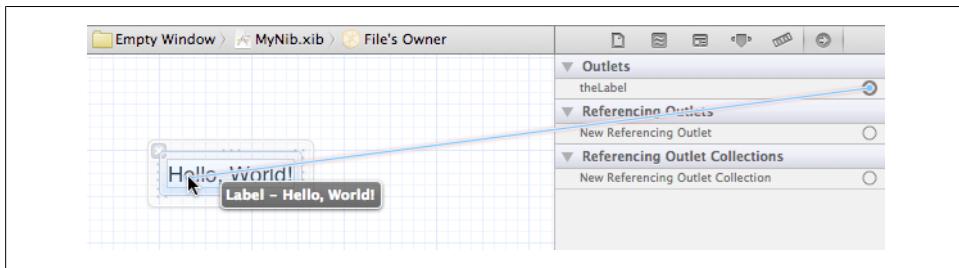


Figure 7-7. Connecting an outlet from the Connections inspector

With the File's Owner object selected, look again at the Connections inspector; it shows that `theLabel` is connected to the Label nib object, and if you hover the mouse over the filled circle at the right, the label object in the nib is highlighted. And look at the `IBOutlet` line in `MyClass.m`; the circle in the gutter is now filled in, and if you click that filled circle, the label is specified in a pop-up menu next to the circle, and the label object in the nib is highlighted. Mission accomplished! We have made an outlet connection in the nib from the File's Owner proxy (representing a `MyClass` instance) to the Label object, and this outlet connection has the same name as the instance variable `theLabel` in `MyClass`'s code.

Therefore, when the nib loads and a `MyClass` instance is the nib's owner, its `theLabel` instance variable will be set to the `UILabel` object that will be instantiated through the loading of the nib. To prove that this is indeed the case, we'll do something with that instance variable in our code. In particular, we'll stick the `UILabel` into our main window interface, thus making it visible. Its visibility will prove that the nib is loading and that the instance variable is being set by the outlet.

Return to `AppDelegate.m` and modify the nib-loading code like this (you added the first two lines earlier):

```
 MyClass* mc = [[MyClass alloc] init];
 [[NSBundle mainBundle] loadNibNamed:@"MyNib" owner:mc options:nil];
 UILabel* lab = [mc valueForKey: @"theLabel"];
 [self.window.rootViewController.view addSubview: lab];
 lab.center = CGPointMake(100,100);
 lab.frame = CGRectMakeIntegral(lab.frame);
```

(We haven't written an accessor method in `MyClass` for `theLabel`, so to save time I used key-value coding.) Build and run the app. The words "Hello, world!" appear in the window! This proves that our outlet worked. We loaded a nib and, using an outlet, we obtained a reference to a nib object and were able to manipulate that object, but putting it in our window.



Figure 7-8. Connecting an outlet by Control-dragging from the source object



Making an instance variable and giving it an `IBOutlet` hint, but forgetting to connect the outlet to anything in the nib, is an unbelievably common beginner (and not-so-beginner) mistake. Had we made this mistake, our code would have run without error, but “Hello, world!” would not appear in the window because `lab` would be `nil`. The unfilled circle that appears in the gutter next to an `IBOutlet` line for which no corresponding nib connection exists is your only clue that something’s amiss, so watch for it.

More Ways to Create Outlets

I said a moment ago that there were other ways to create the outlet. Let’s try some of them. Return to our assistant-panned nib editor, select the File’s Owner, switch to the Connections inspector, and delete the outlet by clicking the little “x” to its left. We’re going to make this outlet again, a different way:

1. Select the File’s Owner in the dock.
2. Hold down the Control key and drag from the File’s Owner to the label. An elastic line follows the mouse.
3. A little window (called a HUD, for “heads-up display”) appears, titled Outlets, listing `theLabel` as a possibility (Figure 7-8). Click `theLabel`.

Once again, look at the Connections inspector with the File’s Owner selected to confirm that this worked. You can even build and run the project again, to prove it to yourself if you’re in any doubt. Now delete the outlet *again*; we’re going to make this outlet in yet a *different* way:

1. Select the File’s Owner in the dock.
2. Control-click the File’s Owner in the dock. A HUD appears, looking a lot like the Connections inspector.
3. Drag from the circle to the right of `theLabel` to the label (Figure 7-9).

Now delete the outlet *again*; we’re going to make this outlet in *another* way. This time, we’re going to operate from the point of view of the label. The Connections inspector shows all connections emanating *from* the selected object; it also shows all connections linking *to* the selected object. So, select the label and look at the Connections inspector. It lists “New Referencing Outlet.” This means an outlet *from* something else *to* the thing we’re inspecting, the label. So:

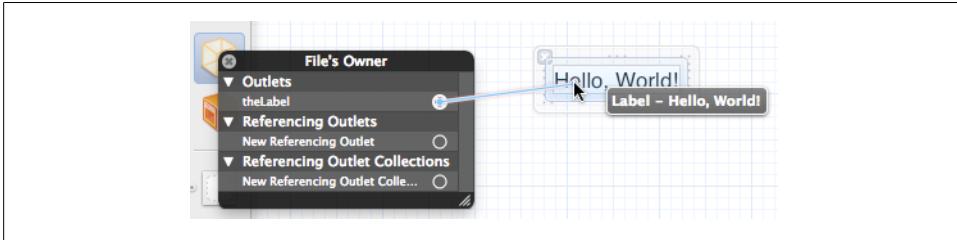


Figure 7-9. Connecting an outlet by dragging from the Connections HUD

1. From the circle at the right of “New Referencing Outlet,” drag to the File’s Owner. An elastic line follows the mouse.
2. A HUD saying `theLabel` appears. Click it.

Confirm that, once again, we’ve made an outlet from the File’s Owner to the label. (And we could also have done the same thing by Control-clicking the label to start with, to show its Connections HUD.) Now delete the outlet *again*; we’re going to make this outlet in *another* way. This time, we’re going to start with the label, but we’re going to connect directly to the *code* which is sitting in the assistant pane:

1. Select the label.
2. Make sure that `MyClass.m` is showing in the assistant pane and that you can see the `IBOutlet` line declaring the instance variable `theLabel`.
3. Hold down the Control key and drag from the label to that line of code. An elastic line follows the mouse.
4. When you’ve got the mouse positioned correctly, the words Connect Outlet will appear. Release the mouse.

Yet again, confirm that we’ve successfully made the desired outlet. And you could also have done the same thing in reverse; starting with the circle at the left of the `IBOutlet` line, you can drag (without holding Control) to the label in the nib.

Now delete the outlet one last time, and (get this) *delete the line of code* declaring the instance variable (but leave the curly braces). We’re going to create the outlet and the instance variable declaration, all in a single amazing move:

1. Select the label.
2. Make sure `MyClass.m` is showing in the assistant pane.
3. Hold down the Control key and drag from the label to the area within the curly braces. An elastic line follows the mouse. The words Insert Outlet or Outlet Collection appear. Release the mouse.
4. A little HUD appears, asking for the name of the instance variable that’s about to be created. Call it `theLabel` (and make sure the type is `UILabel`), and press Return.

Connecting to Code is an Illusion

Making an outlet by connecting directly between code and an interface object in the nib is extremely cool and convenient, but coolness and convenience do not relieve you of the necessity of understanding what an outlet is and how it works. No matter what physical gesture you make in Xcode, and no matter how much Xcode may give you the illusion that you are somehow connecting an object in the nib to the code, that's not really the case. There is *no* connection between the instance variable in the code and the object in the nib; there is *no* identity between the instance variable in the code and the outlet connection in the nib. There are always, if an outlet is to work properly, two distinct and separate things: an instance variable in a class, and an outlet in the nib, *with the same name*, and *coming from an instance of that class*. It is the identity of the names that allows the two to be matched at runtime when the nib loads, so that the instance variable is properly set at that moment.

The `IBOutlet` line declaring the instance variable is created, and the outlet is formed to match it.

More About Outlets

At the risk of seeming to repeat myself, let me emphasize an important thing to remember about outlets (and nib connections generally) that often confuses beginners: they apply to specific instances. Outlets appear in a nib, but a nib is just a template for specific instances. At the moment a nib loads, then and only then, the one specific instance which is the nib's owner (represented by the File's Owner in the nib) and the specific instances generated from the nib objects are all in existence together and are hooked together by their outlets.

All our examples so far have involved a proxy object, but an outlet connection can connect any two objects in the nib. The only requirement is that the source object be of a class that has an instance variable whose type matches the class of the target object.



Nothing in the documentation for a built-in Cocoa class tells you which of its instance variables are available as outlets. In general, the only way to learn what outlets a built-in class provides is to examine a representative of that class in a nib.

It is also possible to create an *outlet collection*. This is an `NSArray` instance variable matched by multiple connections to objects of the same type. For example, suppose a class contains this instance variable declaration:

```
IBOutletCollection(UILabel) NSArray* labels;
```

Connections Between Nibs

You cannot draw a connection from an object in one nib to an object in another nib. If you expect to be able to do this, you haven't understood what a nib is! An object in a nib is only a potential object, becoming a real object when the nib is loaded and the object is instantiated. This potentiality can be realized never, once, or many times. Two objects in the same nib will be instantiated together, so it's clear what a connection means. But a connection from an object in one nib to an object in another nib would be meaningless, because there's no way to say what actual future instances the connection is supposed to connect. The problem of communicating between an instance instantiated from one nib and an instance instantiated from another nib is just a special case of the more general problem of how to communicate between instances in a program and is discussed in [Chapter 13](#).

Then it is possible to form multiple `labels` outlets from an instance of that class in a nib, each one to a different `UILabel` in that nib. When the nib loads, those `UILabel` instances become the elements of the `NSArray labels`. The order in which the outlets are formed is the order of the elements in the array. This is a fairly new feature and I haven't written any code that uses it.

Action Connections

An *action* is a message emitted automatically by a Cocoa `UIControl` interface object (a *control*) when the user does something to it, such as tapping the control. The various user behaviors that will cause a control to emit an action message are called *events*. To see a list of possible events, look at the `UIControl` class documentation, under "Control Events." For example, in the case of a `UIButton`, the user tapping the button corresponds to the `UIControlEventTouchUpInside` event. In the case of a `UITextField`, the user typing or deleting or cutting or pasting corresponds to the `UIControlEventEditingChanged` event. A complete list of `UIControls` and what events they report is provided in [Chapter 11](#).

An action message, then, is a way for your code to respond when the user does something to a control in the interface, such as tapping a button. But your code will not receive an action message from a control unless you explicitly make prior arrangements with that control. You must tell the control what event should trigger an action message, what instance to send the action message to, and what the action message's name should be. There are two ways to make this arrangement: in code, or in a nib.

Either way, we're going to need a method for the action message to call. There are three standard signatures for a method that is to be called through an action message; the most commonly used one takes a single parameter, which will be a reference to the object that emitted the action message. (For full details, see [Chapter 11](#).) So, for ex-

ample, you could have a method like this (let's agree to put it in the implementation section for ViewController, in *ViewController.m*):

```
- (void) buttonPressed: (id) sender {
    UIAlertView* av = [[UIAlertView alloc] initWithTitle:@"Howdy!"
                                                message:@"You tapped me."
                                               delegate:nil
                                              cancelButtonTitle:@"Cool"
                                              otherButtonTitles:nil];
    [av show];
}
```

Now, as I mentioned a moment ago, it is possible to arrange *in code* for **buttonPressed:** to be called when the user taps a button. In particular, if **b** is a reference to the button, then some ViewController code could say:

```
[b addTarget:self action:@selector(buttonPressed:)
forControlEvents:UIControlEventTouchUpInside];
```

That last line means: “Hey there, button! When the user taps on you (**UIControlEventTouchUpInside**), send me (**self**) a **buttonPressed:** message.” (See [Chapter 3](#) if you’ve forgotten about the **@selector** directive.) Of course, such an instruction assumes that this object (**self**) really does implement a **buttonPressed:** method. (If it doesn’t, then when the user taps the button, the app will crash.)

However, instead of doing that, we’re going to use the existing button in *ViewController.xib* and arrange *in the nib* for its action message to be **buttonPressed:** and to be sent to a ViewController instance. We’re going to form an *action connection* in the nib. We can do this because, as I’ve already mentioned, the File’s Owner proxy in *ViewController.xib* is of the ViewController class.

As with outlets, there are several ways to do this; I’ll just show you the main ones and leave you to discover the rest. (They are all directly comparable to the many ways of creating an outlet connection.)

1. We need a hint, in our code, that a method with the expected signature exists. This hint involves substituting **IBAction** for the method’s **void** return type. (The substitution is legal because **IBAction** is **#defined** as **void**; Xcode can see the hint in your code, but the preprocessor will turn **IBAction** back to **void** before the compiler ever sees it.) So, in *ViewController.m*, change the first line of our **buttonPressed:** method implementation to look like this (and save the file):

```
- (IBAction) buttonPressed: (id) sender {
```

This causes an empty circle to appear in the gutter next to the **IBAction** line.

2. Now edit *ViewController.xib*, select the button in the window, and look at the Connections inspector. The event for which we’d like to send the action message is Touch Up Inside. Drag from its circle to the File’s Owner nib object in the dock, which is to receive the message ([Figure 7-10](#)).

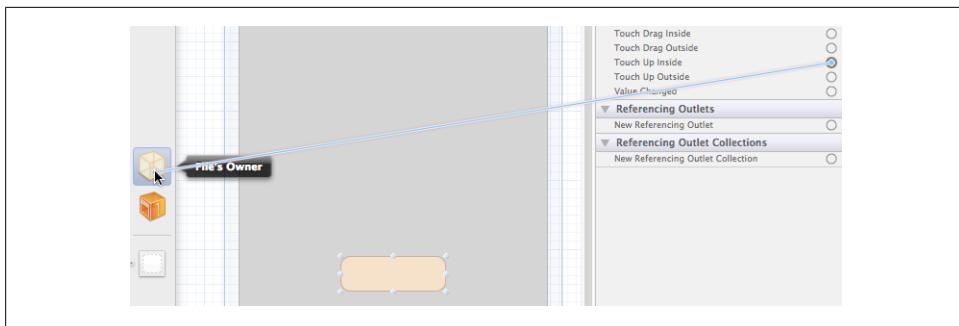


Figure 7-10. Connecting an action from the Connections inspector

3. A little window listing possible Empty_WindowAppDelegate action methods appears; in this case, it lists only `buttonPressed:`. Click on `buttonPressed:` to form the connection.

To see that the action connection has been formed, look at the Connections inspector. If you select the button, the Connections inspector reports that the button's Touch Up Inside event is connected to the Empty Window App Delegate's `buttonPressed:` method. If you select the File's Owner object, the Connections inspector reports a Received Action where `buttonPressed:` is called by the Rounded Rect Button's Touch Up Inside event. Finally, look at the code in `ViewController.m`; the circle next to the `IBAction` line is filled, and you can click it to reveal that the connection is from the button.

Finally, to make assurance doubly sure, you can also build and run the project to confirm that the action connection is working. In the running app, the button inside the window now actually does something when the user taps it! It summons an alert.

As with outlets, we could have formed the action connection by Control-dragging from the button directly to the File's Owner, instead of involving the Connections inspector. If you *just* Control-drag, Interface Builder assumes a default event for you (in this case, it would assume Touch Up Inside). If that isn't what you want, start by Control-clicking on the button to summon a HUD version of the Connections inspector, and drag from the desired event's circle just as you would do from the real Connections inspector.

As with outlets, you can also form the action connection directly to code. (But please reread “[Connecting to Code is an Illusion](#)” on page 150; that warning applies equally to action connections.) In Figure 7-11, we've Control-clicked the button to summon its Connections HUD, and dragged from the Touch Up Inside circle to the `buttonPressed:` implementation. And we could equally have gone the other way, dragging from the unfilled circle next to the `IBAction` line to the button.

But wait, there's more! Instead of writing the action method ahead of time, you can ask Xcode to stub it out for you. To do so, Control-drag from the nib to an empty spot in the implementation; the words Insert Action appear, and when you release the mouse, a dialog appears, letting you specify the name of the action method, the number of

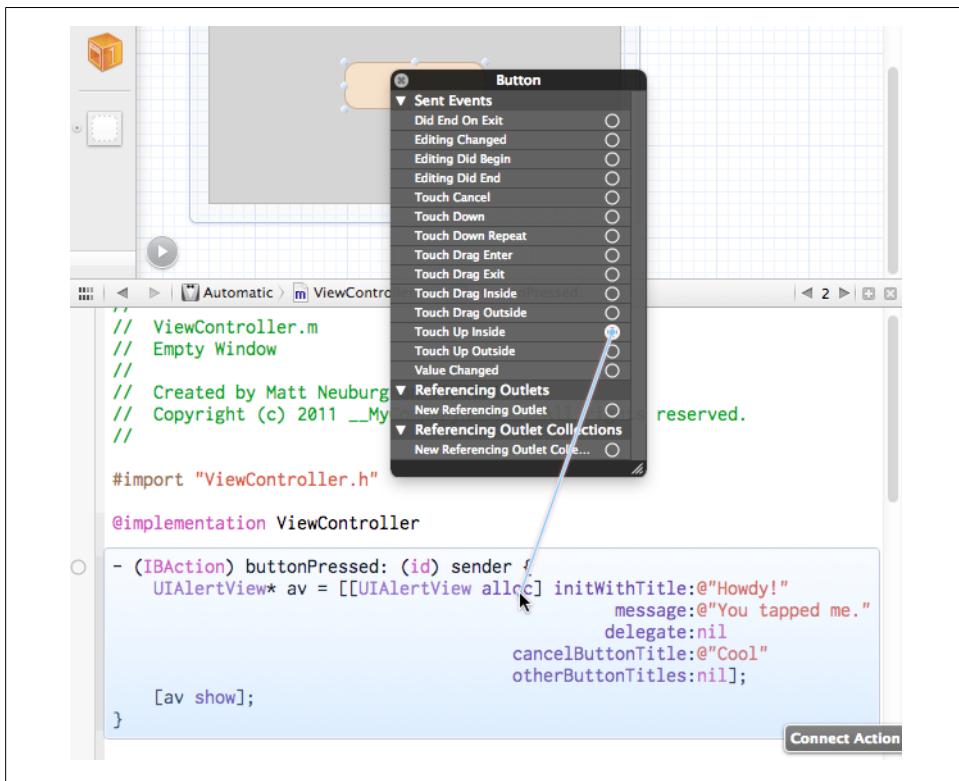


Figure 7-11. Connecting an action to a method implementation

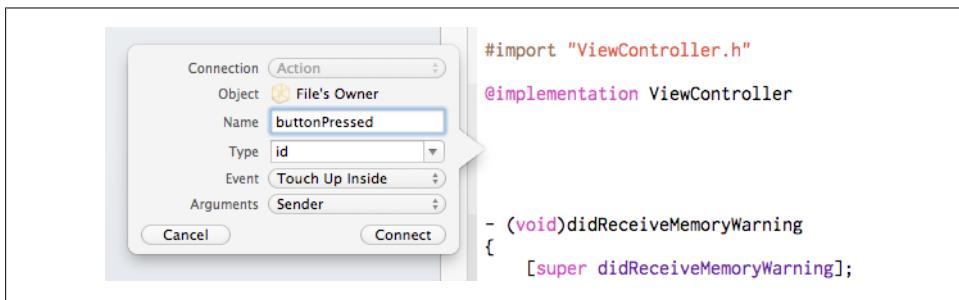


Figure 7-12. Connecting an action and creating a method implementation

arguments it should take, and the control event to be used as a trigger (Figure 7-12). Xcode inserts the method implementation, but doesn't put any code between the curly braces; it's smart, but not smart enough to guess what you want the method to do!

Additional Initialization of Nib-Based Instances

By the time a nib finishes loading, its instances are fully fledged; they have been initialized and configured with all the attributes dictated through the Attributes and Size inspectors, and their outlets have been used to set the values of the corresponding instance variables. Nevertheless, you might want to append your own code to the initialization process as an object is instantiated from a loading nib. Most commonly, to do this, you'll implement `awakeFromNib`. The `awakeFromNib` message is sent to all nib-instantiated objects just after they are instantiated by the loading of the nib: at the point where this happens, the object has been initialized and configured and its connections are operational.

For example, our Empty Window app is loading `MyNib.xib`, extracting a `UILabel` from it, and inserting that label into our interface; the result is that the words “Hello, world!” appear in our window. Let's modify the behavior of this `UILabel` so that it does some additional self-initialization in code. To do that, we will need a class of our own to which our `UILabel` will belong. Clearly, this needs to be a `UILabel` subclass. So:

1. In Xcode, choose File → New → New File and specify that you want a Cocoa Touch Objective-C class. Click Next.
2. Make the new class a subclass of `UILabel`. Click Next.
3. Call it `MyLabel`. Make sure you're saving into the project folder; set the Empty Window group and the Empty Window target. Click Save.
4. In `MyLabel.m`, somewhere in the implementation section, implement `awakeFromNib`:

```
- (void) awakeFromNib {
    [super awakeFromNib];
    self.text = @"I initialized myself!";
    [self sizeToFit];
}
```

5. That code won't apply to the label in `MyNib.xib` unless that label is a `MyLabel`, so edit `MyNib.xib` and change the label's class to `MyLabel` (in the Identity inspector).

Now build and run the project. Instead of “Hello, world!” we now see “I initialized myself!” in the window.



Mac OS X Programmer Alert

If you're an experienced Mac OS X programmer, you may be accustomed to rarely or never calling `super` from `awakeFromNib`; doing so used to raise an exception, in fact. In iOS, you must always call `super` in `awakeFromNib`. Another major difference is that in Mac OS X, a nib owner's `awakeFromNib` is called when the nib loads, so it's possible for an object to be sent `awakeFromNib` multiple times; in iOS, `awakeFromNib` is sent to an object only when that object is itself instantiated from a nib, so it can be sent to an object a maximum of once.

Sometimes, you might need to interfere with a nib object's initialization at an even earlier stage. If this object is a UIView or UIViewController (or a subclass of either), you can implement `initWithCoder:`. In your implementation, be sure to call `super` and return `self` as you would do in any initializer. Your purpose here would typically be to initialize additional instance variables that your subclass has declared, as with any initializer.

Here, for example, is an implementation of MyLabel that declares an instance variable that is an int called `num` and manipulates it first in `initWithCoder:` and then in `awakeFromNib`, thus proving that the two are called in that order:

```
@implementation MyLabel
{
    int num;
}

- (id) initWithCoder:(NSCoder *)aDecoder {
    self = [super initWithCoder:aDecoder];
    if (self) {
        self->num = 42;
    }
    return self;
}

- (void) awakeFromNib {
    [super awakeFromNib];
    self.text = [NSString stringWithFormat: @"The answer is %i", self->num];
    [self sizeToFit];
}
@end
```

That's trivial and unnecessary, but it illustrates the principle.

Documentation

Knowledge is of two kinds. We know a subject ourselves, or we know where we can find information upon it.

—Samuel Johnson, *Boswell's Life of Johnson*

You don't remember Cocoa; you look it up!

—Anonymous programmer, cited by Beam and Davidson, *Cocoa in a Nutshell*

No aspect of Cocoa programming is more important than a fluid and nimble relationship with the documentation. There is a huge number of built-in classes, with many methods and properties and other details. Apple’s documentation, whatever its flaws, is the definitive official word on how you can expect Cocoa to behave and on the contractual rules incumbent upon you in working with this massive framework whose inner workings you cannot see directly.

The Xcode documentation installed on your machine comes in large chunks called *documentation sets* (or *doc sets*, also called *libraries*). You do not merely install a documentation set; you subscribe to it, so that when Apple releases a documentation update (because a new version of iOS has been released, or because there has been an incremental revision of the documentation), you can obtain the updated version.

When you first install the Xcode tools, the bulk of the documentation is *not* installed on your machine; viewing the documentation in the documentation window (discussed in the next section) requires an Internet connection, so that you can see the online docs at Apple’s site. However, assuming that you checked Documentation in the installer, the documentation *will* be installed on your machine; you should start up Xcode immediately after installation to let it download and install your initial documentation sets. The process can be monitored, to some extent, in the Downloads pane of the Preferences window (under Documentation); you can also specify here whether you want updates installed automatically or whether you want to press Check and Install Now manually from time to time. This is also where you specify which doc sets you

want; I believe that in Xcode 4.2 the iOS 5.0 Library and the Xcode 4.2 Developer Library are all you need for iOS development. (But at the time of this writing there are some bugs in the documentation, such that the Mac OS X Core Library doc set is also needed: for example, the *Information Property List Key Reference* appears only in the Mac OS X Core Library; and although Core Image is now partly implemented in iOS 5, some important information is documented only for Mac OS X.) You may have to provide your machine's admin password when a doc set is first installed.

The Documentation Window

Your primary access to the documentation is in Xcode, through the Documentation tab of the Organizer window (Window → Organizer and then click Documentation, or Help → Documentation and API Reference). I'll refer to this as the *documentation window*, even though it's really an aspect of the Organizer window.

The documentation window behaves basically as a glorified web browser, because the documentation consists essentially of web pages. Indeed, most of the same pages can be accessed at Apple's developer site, <http://developer.apple.com>. And any page open in the documentation window can be opened instead in your web browser: Control-click for the contextual menu and choose Open Page in Browser. Notice too the contextual menu for links within a documentation window, such as Copy Link and Open Link in Browser. When you're trying to figure something out, the ability to spawn off a page as a secondary window in a browser while you go on searching in the Xcode documentation window can be very useful.

Each doc set has a home page, which you access from the Browse navigator (Editor → Explore Documentation) or from the first component of the jump bar (Control-4). A typical home page presents a full list of documents, which can be sorted by column and filtered by keyword. Some home pages, such as the iOS Library home page, also have a broad categorical list down the left side, which can similarly be used to filter the document list. In practice I rarely use these home pages, though they can come in handy when you're looking for broad topic introductions (click Guides on the left). The Browse navigator (and the jump bar) can also be used to explore a doc set by category.

When you encounter a documentation page to which you're likely to want to return, make it a bookmark (Editor → Add Bookmark). Bookmarks are accessed through the Bookmarks navigator (Editor → Documentation Bookmarks). Documentation bookmark management is simple but effective: you can rearrange bookmarks or delete a bookmark, and that's all.

My chief way into the documentation — and, I suspect, most users' chief way — is by searching (Editor → Search Documentation). Type a term into the search field (Shift-Option-Command-?). Click the magnifying glass to choose Show Find Options. It's important to set these options correctly:

Match Type

Your choices are Contains, Prefix, and Exact. These determine how your search terms are understood (as middles of words, starts of words, or whole words, respectively), and you'll probably want to switch among them fairly often, depending on what you're searching for. For example, if you are typing the start of the name of a class you want to search for, do a Prefix search, not a Contains search.

Doc Sets

Check only those doc sets that interest you; if you're doing iOS development, for example, uncheck any Mac OS X libraries to eliminate inapplicable and duplicate results.

Languages

Check only those languages you're likely to be interested in (probably Objective-C and C).

In Xcode 4, the search doesn't take place until you press Return. Search results are displayed in categories, in relevance order, in the navigation pane; click a result to see that page.

Alternatively, if you're editing code, select a term in the editor and choose Help → Search Documentation for Selected Text (Control-Option-Command- $/$). This command switches to the documentation window, enters the selected term into the search field, and performs the search using the current find options, in a single move.

Don't confuse searching the documentation with finding within the current page. To find within the current documentation page, make sure the focus is within the page itself (probably by clicking in the page), and then use the Edit → Find menu commands. Command-F summons a find bar, as in Safari.

A major difference between the display of a documentation page in Xcode and its display in Safari is that the latter often shows a Table of Contents column at the left side. In Xcode, this Table of Contents column is suppressed, which saves space, but makes it harder to get a sense for where you are in a document or a set of related documents. The intention is presumably that you should use the jump bar both to get your bearings and to navigate. The last component in the jump bar may show headings within the current document; the next-to-last component may show related documents in the same collection.

Class Documentation Pages

In the vast majority of cases, your target documentation page will be the documentation for a class. I have frequently spoken already of the importance of class documentation pages. A common move on your part will be to search on a class name in the documentation window. If you search on, say, `NSString`, the search result whose title is `NSString Class Reference` is the class documentation for `NSString`.

UIButton Class Reference

Inherits from	UIControl : UIView : UIResponder : NSObject
Conforms to	NSCoding NSCoding (UIView) UIAppearance (UIView) UIAppearanceContainer (UIView) NSObject (NSObject)
Framework	/System/Library/Frameworks/UIKit.framework
Availability	Available in iOS 2.0 and later.
Declared in	UIButton.h
Related sample code	AddMusic BubbleLevel iPhoneMixerEQGraphTest SimpleNetworkStreams TheElements

Overview

An instance of the `UIButton` class implements a button on the touch screen. A button intercepts touch events and sends an action message to a target object when tapped. Methods for setting the target and action are inherited from `UIControl`. This class provides methods for setting the title, image, and other appearance properties of a button. By using these accessors, you can specify a different appearance for each button state.

Tasks

Creating Buttons

+ `buttonWithType:`

Figure 8-1. The start of a typical class documentation page

Let's pause to notice the key features of a class documentation page. I'll use `UIButton` as an example (Figure 8-1):

Inherits from

Lists, and links to, the chain of superclasses. One of the biggest beginner mistakes is failing to read the documentation up the superclass chain. A class inherits from its superclasses, so the functionality or information you're looking for may be in a superclass. You won't find out about `addTarget:action:forControlEvents:` from the `UIButton` class page; that information is in the `UIControl` class page. You won't find out that a `UIButton` has a `frame` property from the `UIButton` class page; that information is in the `UIView` class page.

Conforms to

Lists, and links to, the protocols implemented by this class. Protocols are discussed in [Chapter 10](#). Fortunately, a class that conforms to a formal protocol usually lists that protocol's required methods as links (though the methods themselves are documented on the protocol's documentation page).



Methods injected into a class by a category ([Chapter 10](#)) are often *not* listed on that class's documentation page and can be very difficult to discover. This is a major weakness in Apple's organization and display of the documentation. A third-party documentation display application such as AppKiDo can be helpful here (<http://homepage.mac.com/aglee/downloads/appkido.html>).

Framework

Tells what framework this class is part of. Your code must link to this framework in order to use this class (see [Chapter 6](#)).

Availability

States the earliest version of the operating system where this class is implemented. For example, EKEventViewController, along with the whole EventKit framework (consisting of classes for querying the user's calendar; see [Chapter 32](#)) wasn't invented until iOS 4.0. So if you want to use this feature in your app, you must make sure either that your app targets only iOS 4.0 or later or that you take precautions not to call into this framework on earlier versions of the operating system. The availability information also confirms that you're looking at the right documentation page; if you're doing iOS programming and this class is available only on Mac OS X, reading this page is pointless. Note that individual methods also have availability information.

Companion guide

If a class documentation page lists a companion guide, you might want to click that link and read that guide. Guides are broad surveys of a topic; they provide important information (including, often, useful code examples), and they can serve to orient your thinking and make you aware of your options. (See the UIView class page for an example.)

Related sample code

If a class documentation page links to sample code, you might want to examine that code. (But see my remarks on sample code in the next section of this chapter.)

Overview

Some class pages provide extremely important introductory information in the Overview section, including links to related guides and further information. (See the UIView class page for an example.)

Tasks

This section lists in categorical order, and links to, the properties and methods that appear later on the page. (Recall from [Chapter 5](#) that a property is a syntactic shortcut for calling an accessor method; the documentation lists the property rather than the accessor.) Often, just looking over this list can give you the hint you’re looking for.

Properties, Class Methods, Instance Methods

These sections provide the full documentation for this class’s methods. In recent years, this part of the documentation has become quite splendid, with good hyperlinks. Note the following subsections:

The property or method name

This name is suitable for copying and pasting into your code (if, for example, you need to enter the name of a selector).

The property or method’s purpose

A short summary of what it does.

The formal declaration for the property or method

Read this to learn things like the method’s parameters and return type. ([Chapter 12](#) explains how to read a property declaration.) Suitable for copying and pasting into your code in order to enter a call to this method, though you are more likely to use Xcode’s code completion feature where possible (see [Chapter 9](#)).

Parameters and return value

Precise information on the meaning and purpose of these.

Discussion

Often contains extremely important further details about how this method behaves. Always pay attention to this section!

Availability

An old class can acquire new methods as the operating system advances; if a newer method is crucial to your app, you might want to exclude your app from running on older operating systems that don’t implement the method.

See also

Lists and links to related methods. Very helpful for giving you a larger perspective on how this method fits into the overall behavior of this class.

Related sample code

It can sometimes be worth consulting the sample code to see an example of how this particular method is used.

Declared in

The relevant header file. It can sometimes be worth looking at the header file, which may contain helpful comments or other details. Unfortunately, the list-

ing of a header file in the documentation window is not a clickable link. You can open the header file from your project window, as explained later in this chapter.

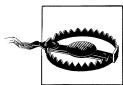
Constants

Many classes define constants that accompany particular methods. For example, to create a `UIButton` instance in code, you call the `buttonWithType:` class method; the argument value will be a constant, listed under `UIButtonType` in the Constants section. (To help you get there, there's a link from the `buttonWithType:` method to the `UIButtonType` section in Constants.) There's a formal definition of the constant; you won't usually care about this (but do see [Chapter 1](#) if you don't know how to read it). Then each value is explained, and the value name is suitable for copying and pasting into your code.

Sample Code

Apple provides plenty of sample code projects. You can view the code directly in the documentation window; sometimes this will be sufficient, but you can see only one class implementation or header file at a time, so it's difficult to get an overview. The alternative is to open the sample code project in Xcode.

When you look at a sample code page from your browser, there's a button that reads Download Sample Code. In fact, the sample code may already be on your computer. When you look at the same sample code page in the documentation window, the same button will read Open Project. The sample code on your hard disk is zipped, so even if the code is already on your computer, you are first asked to specify a "download folder" in which to save the unzipped project folder. This policy of keeping the sample code projects zipped on your hard disk is a good one, as it prevents you from accidentally altering the original, and you are free to experiment with the unzipped copy.



If a sample code project was linked against the frameworks of an older SDK that isn't installed on your computer, the project will be described in the Project navigator with the words "missing base SDK." In earlier versions of Xcode, this situation could prevent you from building and running the project, and features that depend on indexing might not work. In Xcode 4.2, however, the project should build and run regardless. To remove the "missing base SDK" annotation, edit the target, switch to Build Settings, and change the outdated Base SDK setting to Latest iOS.

As a form of documentation, sample code is both good and bad. It can be a superb source of working code that you can often copy and paste and use with very little alteration in your own projects. It is usually heavily commented, because the Apple folks are aware, as they write the code, that it is intended for instructional purposes. Sample code also illustrates concepts that users have difficulty extracting from the

documentation. (Users who have not grasped UITouch handling, for instance, often find that the lightbulb goes on when they discover the MoveMe example.) But the logic of a project is often spread over multiple files, and nothing is more difficult to understand than someone else's code (except, perhaps, your own code). Moreover, what learners most need is not the fait accompli of a fully written project but the reasoning process that constructed the project, which no amount of commentary can provide.

My own assessment is that Apple's sample code is generally very thoughtful and instructive and definitely a major component of the documentation, and that it deserves more appreciation and usage than it seems to get. But it is most useful, I think, after you've reached a certain level of competence and comfort.

Other Resources

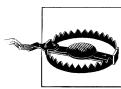
Here is a survey of other useful resources that supplement the documentation.

Quick Help

Quick Help is a condensed rendering of the documentation on some single topic, usually a symbol name (a class or method). It appears with regard to the current selection or insertion point automatically in the Quick Help inspector (Option-Command-2) if the inspector is showing. Thus, for example, if you're editing code and the insertion point or selection is within the term `CGPointMake`, documentation for `CGPointMake` appears in the Quick Help inspector if it is visible.

A slightly reduced version of the same Quick Help documentation can be displayed as a small floating window, without the Quick Help inspector, by Option-clicking on a term in code. Alternatively, select a term and choose Help → Quick Help for Selected Item (Shift-Control-Command-?). In the Quick Help window, click the “book” icon to open the full documentation in the documentation window; click the “H” icon to open the appropriate header file.

Both the Quick Help inspector and the Quick Help window may also contain links. Some of these may be to various other documentation aids, such as sample code. The most important link will probably be the first one, the name of the symbol being documented; this links to the appropriate spot in the full documentation in the documentation window.



Xcode 4 provides no *direct* path from a symbol in code to its documentation in the documentation window. You must pass through Quick Help to get there. You can select a term and choose Help → Search Documentation for Selected Text (Control-Option-Command-"/"), but this is hardly the same thing, as it doesn't jump to the actual API linked from Quick Help.

If you hold down Option and hover the mouse over code, the term that Quick Help would document if you were to click at that point appears with a dotted underline.

Quick Help is also available during code completion ([Chapter 9](#)), concerning the term currently being proposed as a completion; the question-mark icon at the right side of the code completion pop-up menu summons the Quick Help window. Plus, Quick Help is available in the Quick Help inspector for interface objects selected while editing a nib, for build settings while editing a project or target, and so forth.

Symbols

A *symbol* is a nonlocally defined term, such as the name of a class, method, or instance variable. If you can see the name of a symbol in your code in an editor in Xcode, Command-click it to jump to the definition for that symbol. Alternatively, select text and choose Navigate → Jump to Definition (Control-Command-J). If there are multiple definitions for a term, you'll get a little pop-up window where you can pick which one to jump to. If you hold down Command and hover the mouse over code, the symbol whose definition would be shown if you were to click at that point appears with a solid underline.

If the symbol is defined in a Cocoa framework, you jump to the header file. If the symbol is defined in your code, you jump to the class or method definition; this can be very helpful not only for understanding your code but also for navigating it.

The precise meaning of the notion “jump” depends upon the modifier keys you use in addition to the Command key, and on your settings in the General pane of Xcode’s preferences. For example, if you haven’t changed these settings from the default, Command-click jumps in the same editor, Command-Option-click jumps in an assistant pane, and Command-double-click jumps in a new window. Similarly, Control-Option-Command-J jumps in an assistant pane to the definition of the selected term.

Another way to see a list of your project’s symbols, and navigate to a symbol definition, is with the Symbol navigator ([Chapter 6](#)).

Header Files

Sometimes a header file can be a useful form of documentation. It compactly summarizes a class’s instance variables and methods and may contain comments and other helpful information — information that may be documented nowhere else. A single header file can contain declarations for multiple class interfaces and protocols. So it can be an excellent quick reference.

There are various ways to see a header file from an Xcode editor:

- If the class is your own and you’re in the implementation file, choose Navigate → Jump to Next Counterpart (Control-Command-Up).

- Click the Related Files button at the left of the jump bar (Control-1). The menu lets you jump to any header files imported in the current file (as well as any files that import the current file) and to the header files of the current class file's superclasses and subclasses and so forth. Hold Option to jump in an assistant pane.
- Select text and choose File → Open Quickly (Shift-Command-O). This command brings up a dialog listing all source and header files containing a given symbol.
- Command-click a symbol, choose Navigate → Jump to Definition, or pass through Quick Help, as described in the previous sections.
- Use the Symbol navigator ([Chapter 6](#)).

All of these approaches require that a project window be open; File → Open Quickly requires an active SDK for effective operation, and the others all operate on specific windows or words in an open project. An alternative that works under all circumstances is to switch to the Terminal and use the `open -h` command to open a header file in Xcode. The argument may represent part of a header file's name. The command is interactive if there's an ambiguity; for example, `open -h NSString` proposes to open `NSString.h` or `NSStringDrawing.h` (or both, or neither). I wish this command were built into Xcode itself.

Internet Resources

Programming has become a lot easier since the Internet came along and Google started indexing it. It's amazing what you can find out with a Google search. Your problem is very likely a problem someone else has faced, solved, and written about on the Internet. Often you'll find sample code that you can paste into your project and adapt.

Apple's documentation resources are available at <http://developer.apple.com>. These resources are updated before the changes are rolled into your doc sets for download. There are also some materials here that aren't part of the Xcode documentation on your computer. As a registered iOS developer, you have access to iTunes videos, including the videos for all WWDC 2011 sessions, and to Apple's developer forums (<https://devforums.apple.com>).

Apple maintains some public mailing lists (<http://lists.apple.com/mailman/listinfo>). I have long subscribed to the Xcode-users group (for questions about use of the Xcode tools) and the Cocoa-dev group (for questions about programming Cocoa). Cocoa-dev does now permit iOS questions, but it is not heavily used for these. The lists are searchable, but Apple's own search doesn't work very well; you're better off using Google with a `site:lists.apple.com` term, or <http://www.cocoabuilder.com>, which archives the lists. Apple has not added a mailing list devoted to iOS programming; that's what the developer forums are supposed to be for, though the interface for these is extraordinarily clunky in my view, and this — plus the lack of openness (to Google and to the world in general) — has limited their usefulness.

Other online resources, such as forums, have sprung up spontaneously as iOS programming has become more popular, and lots of iOS and Cocoa programmers blog about their experiences. I am particularly fond of Stack Overflow (<http://www.stackoverflow.com>); of course it isn't devoted exclusively to iOS programming, but lots of iOS programmers hang out there, questions are answered succinctly and correctly, and the interface lets you focus on the right answer quickly and easily.

Life Cycle of a Project

This chapter surveys some of the main stages in the life cycle of a project, from inception to submission at the App Store. This survey will provide an opportunity to discuss some additional features of the Xcode development environment. You already know how to create a project, define a class, and link to a framework ([Chapter 6](#)), as well as how to create and edit a nib ([Chapter 7](#)) and how to use the documentation ([Chapter 8](#)).

Choosing a Device Architecture

As you create a project, after you pick a project template, in the part of the dialog where you name your project, the Device Family pop-up menu offers a choice of iPhone, iPad, or Universal (meaning both iPhone and iPad).



A *universal app* is a single app built from a single target, containing some resources that are loaded only on one type of device or the other.

You are not tied forever to your initial decision, but your life will be easier if you decide correctly from the outset. The iPhone and iPad differ in their physical characteristics as well as their programming interfaces. The iPad has a larger screen size, along with some built-in interface features that don't exist on the iPhone, such as split views and popovers ([Chapter 22](#)); thus an iPad project's nib files and some other resources will differ from those of an iPhone project.

Historically, different types of device also ran different versions of the operating system: iOS 3.1.3 and before, plus iOS 4.0 and 4.1, were iPhone only, while iOS 3.2.x was iPad only. This made life very complicated for the programmer wishing to target both types of device; universal apps were particularly difficult to write. Fortunately, starting with iOS 4.2, Apple has unified the system versions, so that the same system now runs on both device types, so that if you write a universal app, you probably won't concern

yourself with possible system differences, although you will still be concerned about device differences.

Your choice in the Device Family pop-up menu affects what template your new project will be based on. It also affects your target's Targeted Device Family build setting:

iPad

The app will run only on an iPad.

iPhone

The app will run on an iPhone or iPod touch; it can also run on an iPad, but not as a native iPad app (it runs in a reduced enlargeable window, which I call the *iPhone Emulator*; Apple sometimes refers to this as “compatibility mode”).

iPhone/iPad

The app will run natively on both kinds of device, and should be structured as a universal app.

Two additional build settings work together and in conjunction with the Targeted Device Family to determine what systems your device will run on:

Base SDK

The *latest* system your app can run on: in Xcode 4.2, you have just two choices, iOS 5.0 and Latest iOS. As of this writing, Latest iOS means iOS 5.0, so what's the difference? It's that, in the latter case, if you update Xcode to develop for a subsequent system, your existing projects will use that newer system's SDK as their Base SDK automatically, without you also having to update their Base SDK setting. Latest iOS is the default when you create a new project.

iOS Deployment Target

The *earliest* system your app can run on: this can be any iOS system number from the current 5.0 all the way back to 3.0. (iOS 3.0 is also the earliest system on which a universal app will run.) You can change the iOS Deployment Target setting easily by editing your project or your target; the project's Info tab has an iOS Deployment Target pop-up menu, and the target's Summary tab has a Deployment Target pop-up menu. These both represent the iOS Deployment Target build setting; you will probably want to edit the target, because if you edit the project only, the target setting will override it.

Writing an app whose Deployment Target differs from its Base SDK is something of a challenge. The problem is that Xcode will happily allow you to compile using any features of the Base SDK, but an actual system, whether it's a Simulator SDK or a device, will crash your app if it uses any features not supported by that system.

For example, if you were to create a new iPad project using the Empty Application template and set the iOS Deployment Target to 3.2, and run it on an iPad with iOS 3.2 installed, the app would crash on launch, because the template contains this line, which is encountered as the app starts up:

Additional Simulator SDKs

When the Deployment Target is set to a system earlier than 5.0, earlier SDK versions of the Simulator may become available in the Scheme pop-up menu. Exactly what versions appear depends on the contents of */Developer/Platforms/iPhoneSimulator.platform/Developer/SDKs/*, which may in turn depends on what Mac OS X system you're using and the history of your Xcode installations. For example, in my Snow Leopard installation of Xcode 4.2, with a project's Deployment Target set to 3.0, I can see iPhone Simulator SDKs back as far as 4.0, and iPad Simulator SDKs back to 4.2 (along with a listing for iPad 3.2 Simulator, but this turns out to be unusable — trying to launch an app into the iPad Simulator for 3.2 fails on Xcode 4.2, though it works in Xcode 4.0). In my Lion installation of Xcode 4.2, on the other hand, with a project's Deployment Target set to 3.0, I see iPhone and iPad Simulator SDKs for iOS 5.0 and 4.3 and that's all; however, I was able to copy the Simulator SDKs from my Snow Leopard installation to my Lion installation, and they were all usable (except the iPad 3.2 Simulator).

```
self.window.rootViewController = self.viewController;
```

The problem is that the window `rootViewController` property wasn't invented until iOS 4.0. Here's an example that should be easier for you to test:

```
[UIButton appearance];
```

If that line of code is encountered while running in a 5.0 Simulator, all is well; if is encountered while running in a 4.3 Simulator, you'll crash, because the `appearance` method wasn't invented until iOS 5.0.

How can you guard against such problems? I would recommend that you not even attempt backwards compatibility with a device and system on which you cannot test directly. If you don't own an iPad running iOS 3.2, it would surely be unwise to set your Deployment Target to iOS 3.2; the prospect that a compatibility issue might not be discovered until the app has been let loose upon a world of users is highly unsettling. Earlier SDKs can help, to be sure; for this particular example, you might discover the crash by trying to run the project with the iPad 3.2 Simulator under Xcode 4.0. But there's more to testing an app that uses the Simulator; some apps, or the discovery of some bugs, might require a device. The fact is that ensuring backward compatibility is hard, and you might reasonably decide it isn't worth the effort.

Writing a universal app presents challenges of its own, because of the physical and system differences between the iPhone and the iPad. As you develop, you must juggle two versions of many files, such as nibs. Moreover, although you'll probably want to share some code between the iPhone and the iPad version of the app, to reduce duplication, some code will have to be kept separate, because your app will behave differently on the different types of device. As I already mentioned, you can't summon a popover on an iPhone; but the complexities can run considerably deeper, because the interfaces might behave very differently — tapping a table cell on the iPhone might summon an

entire new screenful of stuff, whereas on the larger iPad, it might only alter what appears in one region of the screen.

There are various programming devices to govern dynamically what code is encountered, based on what system or device type the app is running on; thus you can avoid executing code that will cause a crash in a particular environment, or otherwise make your app behave differently depending on the runtime circumstances (see also [Example 29-1](#)):

- The `UIDevice` class lets you query the current device as `[UIDevice currentDevice]` to learn its system version (`systemVersion`) and type (`userInterfaceIdiom`, either `UIUserInterfaceIdiomPhone` or `UIUserInterfaceIdiomPad`):

```
if ([UIDevice currentDevice].userInterfaceIdiom == UIUserInterfaceIdiomPhone) {  
    // do things appropriate to iPhone  
} else {  
    // do things appropriate to iPad  
}
```

For an actual example, make a new project from the Master–Detail Application template (with no storyboard) and look in `AppDelegate.m`. You’ll see how the code configures the initial interface differently, including loading a different nib, depending on the device type we’re running on.

- If your app is linked to a framework and tries to run on a system that lacks that framework, it will crash at launch time. The solution is to link to that framework optionally, by changing the Required pop-up menu item in its listing in the target to Optional (this is technically referred to as *weak-linking* the framework).
- You can test for the existence of a method using `respondsToSelector:` and related `NSObject` calls:

```
if ([UIButton respondsToSelector:@selector(appearance)]) {  
    // ok to call appearance method  
} else {  
    // don't call appearance method  
}
```

- You can test for the existence of a class using the `NSClassFromString` function, which yields nil if the class doesn’t exist. Also, if the Base SDK is 5.0 or later, and if the class’s framework is present or weak-linked, you can send the class any message (such as `[CIFilter class]`) and test the result for nil; this works because classes are weak-linked starting in iOS 5.

```
// assume Core Image framework is weak-linked  
if ([CIFilter class]) { // ok to do things with CIFilter
```

- You can test for the existence of a constant name, including the name of a C function, by taking the name’s address and testing against zero. For example:

```
if (&UIApplicationWillEnterForegroundNotification) {  
    // OK to refer to UIApplicationWillEnterForegroundNotification
```

Many calls that load resources by name from your app’s bundle will automatically select an alternative resource whose name (before the extension) ends with `~iphone` or `~ipad` as appropriate to the device type, thus relieving your code from using conditionals. For example, `UIImage`’s `imageNamed:` method, if you specify the image name as `@"Linen.png"`, will load an image called `Linen~ipad.png` if it finds one and if we’re running on an iPad. (We’ll see in [Chapter 15](#) that the same sort of naming convention will help you automatically load a double-resolution image on a device with a double-resolution screen.)

Localization

A device can be set by the user to prefer a certain language as its primary language. You might like the text in your app’s interface to respond to this situation by appearing in that language. This is achieved by *localizing* the app for that language.

Localization works through localization folders in your project and in the built app bundle. Every resource in one of these localization folders has a counterpart in the other localization folders. Then, when your app goes to load such a resource, it automatically loads the one appropriate to the user’s preferred language. For example, if there’s a copy of `ViewController.nib` in the English localization folder and a copy of `ViewController.nib` in the French localization folder, the latter will be loaded as the app launches on a device on which French is the preferred language. So the two copies of `ViewController.nib` should be identical except that all the text the user will see in the interface should be in French in the French version.

This approach solves the problem for resources that are physically loaded, such as nib files and images and sound files, but it doesn’t deal with strings generated from within your code, such as the text of an alert message. Surely you don’t want your code to consist of a bunch of massive `if` clauses every time there’s text to display. The problem is solved through the use of a strings file. A *strings file* is a specially formatted text file whose file extension is `.strings`; by default the name of the file is `Localizable.strings` (that is, this file will be sought by default, if no filename is specified), but you can use another name if you like. As with other localized resources, the strings file exists in multiple copies, one for each language. The strings file consists of key–value pairs; the keys are the same in all copies, but the values differ, depending on the target language. So instead of entering a string directly in your code, you tell your code to fetch the correct value from the appropriate strings file, based on the key:

```
NSString* myAlertText = NSLocalizedString(@"alertTextKey", nil);
```

Another specially named `.strings` file, `InfoPlist.strings`, stores localized versions of `Info.plist` key values. So, for example, the value of the `CFBundleDisplayName` key, as set in your project’s `Info.plist` file, appears under your app’s icon on the user’s device ([Chapter 6](#)); to change this name depending on the user’s primary language setting, you’d include appropriate key–value pairs in `InfoPlist.strings` files.



Figure 9-1. How a localized strings file is represented in Xcode

Localization explains the `en.lproj` folder seen in the Finder in our Empty Window project folder (Figure 6-8). That's an English localization folder; its contents, `InfoPlist.strings` and `ViewController.xib`, are localized for English. In Xcode, however, nothing seems to indicate this; you wouldn't know, from looking at the Project navigator, that there's anything special about these two files. That's because there's only one localization. As soon as a file has more than one localization, it's shown in the Project navigator as a kind of folder, inverted from how it's shown in the Finder: the file name contains hierarchically the names of the localizations (Figure 9-1). This makes it easy to find and edit the correct copy of the file.

To get started with localization in your project, select in the Project navigator a file that you want to localize and examine it in the Localization section of the File inspector (Command-Option-1). It is obvious how to add and remove localization languages here.

For full discussion, see Apple's *Internationalization Programming Topics*.

Editing Your Code

Many aspects of Xcode's editing environment can be modified to suit your tastes. Your first step should be to pick a font face and size you like in the Fonts & Colors preference pane. Nothing is so important as being able to read and write code comfortably! I like a largish size (14 or even 16) and a pleasant monospaced font such as Monaco, Menlo, or Consolas (or the freeware Inconsolata).

Xcode has some formatting, autotyping, and text selection features adapted for Objective-C. Exactly how these behave depends upon your settings in the Editing and Indentation tabs of Xcode's Text Editing preference pane. I'm not going to describe these settings in detail, but I urge you to take advantage of them. Under Editing, I like to check just about everything, including Line Numbers; visible line numbers are useful when debugging. Under Indentation, I like to have just about everything checked too; I find the way Xcode lays out Objective-C code to be excellent with these settings. A sound approach might be to check everything initially and then, when you've some experience editing with Xcode, switch off features you don't prefer.

If you like Xcode's smart syntax-aware indenting, but you find that once in a while a line of code isn't indenting itself correctly, try choosing Editor → Structure → Re-indent (Control-I), which autoindents the current line. (Autoindent problems can also be caused by incorrect syntax earlier in the file, so hunt for that too.)

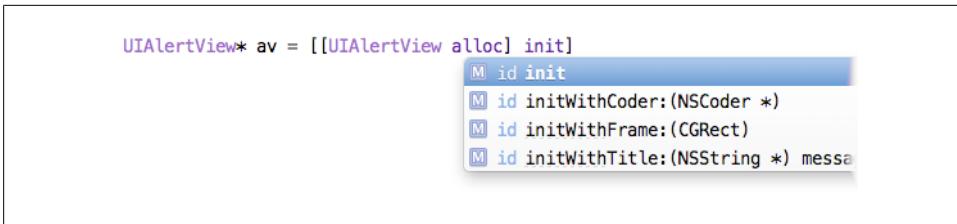


Figure 9-2. The completion menu

Under Editing, notice “Balance brackets in Objective-C method calls.” If this option is checked, then when you type a closing square bracket after some text, Xcode intelligently inserts the opening square bracket before the text. I like this feature, as it allows me to type nested square brackets without planning ahead. For example, I type this:

```
UIalertView* av = [UIalertView alloc]
```

I now type the right square bracket *twice*. The first right square bracket closes the open left square bracket (which highlights to indicate this). The second right square bracket also inserts a space before itself, plus the missing left square bracket:

```
UIalertView* av = [[UIalertView alloc] ]  
// insertion point is here: ^
```

The insertion point is positioned before the second right square bracket, ready for me to type `init`.

Autocompletion

As you write code, you’ll take advantage of Xcode’s completion feature. Objective-C is a verbose language, and whatever reduces your time and effort typing will be a relief. However, I personally do *not* check “Suggest completions while typing” under Editing; instead, I check “Escape key shows code completions”, and when I want autocompletion to happen, I ask for it manually, by pressing Esc.

For example, suppose my code is as displayed in the previous example, with the insertion point before the second right square bracket. I now type `init` and then press Esc, and a little menu pops up, listing the four `init` methods appropriate to a `UIAlertView` (Figure 9-2). You can navigate this menu, dismiss it, or accept the selection, using only the keyboard. So I would navigate to `initWithTitle:....` (no need, actually, as it is selected by default) and press Return to accept the selected choice.

Alternatively, I might press Control-Period instead of Esc. Pressing Control-Period repeatedly cycles through the alternatives. Again, press Return to accept the selected choice.

The template for the correct method call is now entered in my code (I’ve broken it manually into multiple lines to show it here):

```
[[UIAlertView alloc] initWithTitle:<#(NSString *)#>
                           message:<#(NSString *)#>
                           delegate:<#(id)#>
                          cancelButtonTitle:<#(NSString *)#>
                          otherButtonTitles:<#(NSString *), ...#>, nil]
```

The expressions in <#...#> are *placeholders*, showing the type of each parameter; you can select the next placeholder with Tab (if the insertion point precedes a placeholder) or by choosing Navigate → Jump to Next Placeholder (Control-Slash). Thus I can select a placeholder and type in its place the actual value I wish to pass, select the next placeholder and type its value, and so forth.



Placeholders are delimited by <#...#> behind the scenes, but they appear as “text tokens” to prevent them from being edited accidentally. To convert a placeholder to a normal string without the delimiters, select it and press Return, or double-click it.

Autocompletion also works for method declarations. You don’t have to know or enter a method’s return type beforehand. Just type the initial - or + (to indicate an instance method or a class method) followed by the first few letters of the method’s name. For example, in my app delegate I might type:

```
- appli
```

If I then press Esc, I see a list of methods such as `application:didChangeStatusBarFrame:`; these are methods that might be sent to my app delegate (by virtue of its being the app delegate, as discussed in [Chapter 11](#)). When I choose one, the declaration is filled in for me, including the return type and the parameter names:

```
- (void)application:(UIApplication *)application
didChangeStatusBarFrame:(CGRect)oldStatusBarFrame
```

At this point I’m ready to type the left curly brace, followed by a Return character; this causes the matching right curly brace to appear, with the insertion point positioned between them, ready for me to start typing the body of this method.

Snippets

Code autocompletion is supplemented by code snippets, which are text constructs with an abbreviation. Code snippets are kept in the Code Snippet library (Control-Option-Command-2), but you can use code snippets without showing the library. You type the abbreviation and the macro is included among the possible completions. For example, to enter an `if` block, I would type `if` and press Esc and then Return; the construct appears in my code, and the condition area (between the parentheses) and statements area (between the curly braces) are placeholders.

To learn a snippet’s abbreviation, you must open its editing window (select the snippet in the Code Snippet library and press Spacebar) and click Edit. You can add your own

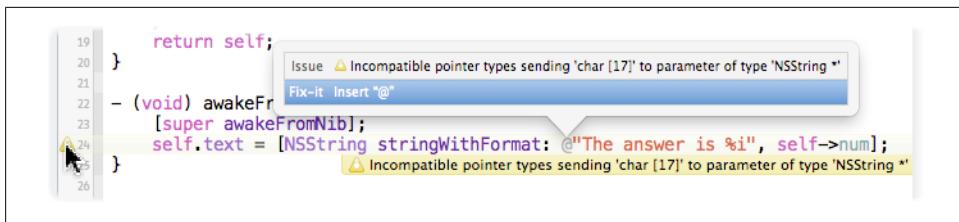


Figure 9-3. A warning with a Fix-it suggestion

snippets, which will be categorized as User snippets; the easiest way is to drag text into the Code Snippet library. Edit to suit your taste, providing a name, a description, and an abbreviation; use the `<#...#>` construct to form any desired placeholders.

If learning a snippet's abbreviation is too much trouble, simply drag it from the Code Snippet library into your text.

Live Syntax Checking and Fix-it

Xcode 4 can perform live syntax checking as you type. This feature can save you from mistakes; in addition, the extremely cool “Fix-it” feature can actually make *and implement* positive suggestions on how to avert a problem.

For instance, in [Figure 9-3](#) I’ve accidentally omitted the @ before an Objective-C NSString literal, and the compiler is warning (because what I’ve typed is a C string literal, a very different thing). By clicking on the warning symbol in the gutter, I’ve summoned a little dialog that not only describes the mistake but tells me how to fix it. Not only that: it has tentatively inserted the missing @ into my code. (Note that @ is a faded gray color. It’s not part of what I typed; Xcode has added it.) Not only *that*: if I press Return, or double-click the “Fix-it” button in the dialog, Xcode *really* inserts the missing @ into my code — and the warning vanishes, because the problem is solved. If I’m confident that Xcode will do the right thing, I can choose Editor → Fix All in Scope (Control-Option-Command-F), and Xcode will implement *all* nearby Fix-it suggestions without my even having to show the dialog.

Live syntax checking can be toggled on or off using the Enable Live Issues In Editors checkbox in the General preference pane. Personally, I keep it turned off, as I find it intrusive. My code is almost never valid while I’m typing, because the terms and parentheses are always half-finished; that’s what it means to be typing. For example, merely typing a left parenthesis will instantly cause the syntax checker to complain of a parse error (until I type the corresponding right parenthesis).

The good news is that, in Xcode 4.2, turning off live syntax checking doesn’t eliminate Fix-it. Even if live syntax checking is turned off, when I compile the code shown in [Figure 9-3](#), I am still shown the same warning and I can still click on the warning symbol in the gutter to get the same Fix-it suggestion dialog, and I can still accept its suggestion by pressing Return, by double-clicking, or by choosing Fix All in Scope.

Navigating Your Code

Developing an Xcode project involves editing code in many files at once. Xcode provides numerous ways to navigate your code. Many of these have been mentioned in previous chapters.

The Project navigator

If you know something about the name of a file, you can find it quickly in the Project navigator (Command-1) by typing into the search field in the filter bar at the bottom of the navigator (Edit → Filter → Filter in Navigator, Command-Option-J). For example, type `xib` to see just your nib files. Moreover, after using the filter bar, you can press Tab and then the Up or Down arrow key to navigate the Project navigator. Thus you can reach the desired nib file with the keyboard alone.

The Symbol navigator

If you highlight the first two icons (the first two are dark, the third is light), the Symbol navigator lists your project’s classes and their methods, making navigation to a desired method easy. As with the Project navigator, the filter bar can quickly get you where you want to go. For example, to see the `applicationDidBecomeActive:` method, type `active` in the search field.

The jump bar

Every path component of the jump bar is a menu:

The bottom level

At the bottom level (farthest right) in the jump bar is a list of your file’s method and function declarations and definitions, in the order in which they appear (hold Command while choosing the menu to see them in alphabetical order); choose one to navigate to it.

You can add section titles to this bottom-level menu using the `#pragma mark` directive. To see an example, make a project based on the Single View Application template; you’ll find this code in `ViewController.m`:

```
#pragma mark - View lifecycle  
  
- (void)viewDidLoad  
{  
    [super viewDidLoad];  
    // Do any additional setup after loading the view, typically from a nib.  
}
```

The result is that the “viewDidLoad” item in the bottom-level menu falls within a “View lifecycle” section. To make a section divider line in the menu, type a `#pragma mark` directive whose value is a hyphen; in this example, both a hyphen (to make a section divider line) and a title (to make a bold section title) are used. Similarly, comments outside of any method and starting with `TODO:, ???:,` or `!!!:` will appear in the bottom-level menu.

Higher levels

Higher-level path components are hierarchical menus; thus you can use any of them to work your way down the file hierarchy.

History

Each editor pane remembers the names of files you've edited in it. The Back and Forward triangles are both buttons and pop-up menus (or choose Navigate → Go Back and Navigate → Go Forward, Control-Command-Left and Control-Command-Right).

Related items

The leftmost button in the jump bar summons a hierarchical menu of files related to the current file, such as counterparts, superclasses, and included files.

The Assistant pane

The Assistant allows you to be in two places at once. Hold Option while navigating to open something in an Assistant pane instead of the main or current editor pane.

The first path component in an Assistant pane's jump bar sets its automatic relationship to the main pane (*tracking*). If that relationship involves multiple files, triangle buttons appear at the right end of the jump bar, letting you navigate between them; or choose from the second path component's pop-up menu (Control-5). For example, show *AppDelegate.m* in the main pane and switch the assistant pane's related items pop-up menu to Includes; the triangle buttons at the right end of the jump bar then navigation between different files #imported by *AppDelegate.m*.

You can also be in two places at once by opening a tab or a separate window.

Jump to definition

Navigate → Jump to Definition (Control-Command-J) lets you jump to the definition or implementation of the symbol already selected in your code.

Open quickly

File → Open Quickly (Shift-Command-O) searches in a dialog for a symbol in your code and the Cocoa headers. You can type the symbol in the search field, or, if a symbol is selected when you summon the dialog, it will be entered in the search field for you (and you can then navigate the dialog entirely with the keyboard).

Breakpoints

The Breakpoint navigator lists all breakpoints in your code. Xcode 4 lacks code bookmarks, but you can misuse a disabled breakpoint as a bookmark.

Finding

Finding is a form of navigation. Xcode has both a global find (Edit → Find → Find in Workspace, Shift-Command-F, which is the same as using the Search navigator) and an editor-level find (Edit → Find → Find, Command-F); don't confuse them.

Find options are all-important. Both sorts of find have options settings that you can summon by clicking the magnifying glass. The global find options ([Figure 6-3](#)) allow you to specify the scope of a search (which files will be searched) in sophisticated ways: choose Custom in the “Find in” pop-up menu to create a scope. The global find search bar also pops down a menu automatically as you type, letting you switch among the most important options. You can also find using regular expressions. There’s a lot of power lurking here.

To replace text, click on the word Find next to the search bar to summon the pop-up menu, and choose Replace. (It may be necessary to perform a global find first, before a global replace on the same search term will work.) You can replace all occurrences, or select particular find results in the Search navigator and replace only those (click Replace instead of Replace All). Even better, click Preview; it summons a dialog that shows you the effect of each possible replacement, and lets you check or uncheck particular replacements in advance of performing the replacement.

A sophisticated form of editor-level find is Editor → Edit All In Scope, which finds simultaneously all occurrences of the currently selected term (usually a variable name) within the current set of curly braces; you can use this to change the variable’s name throughout its scope, or just to survey how the name is used. To change a symbol’s name throughout your code, use Xcode’s Refactoring feature (see “[Making Global Changes to Your Code](#)” in the *Xcode 4 User Guide*).

Debugging

Debugging is the art of figuring out what’s wrong with the behavior of your app as it runs. I divide this art into two main techniques: caveman debugging and pausing your running app.

Caveman Debugging

Caveman debugging consists of altering your code, usually temporarily, typically by adding code to dump informative messages into the console.



To see the console as a full window, open a second project window or tab, show the Debug pane (View → Show Debug Area), and slide the top of the Debug pane all the way up to cover the editor. Eliminate the Navigator and Organizer panes, and the variables list. Now this window or tab contains nothing but the console. Switch to this window or tab when you want to read the console, but don’t run or stop while viewing it, as doing so may cause the Debug pane to close or change size.

The standard command for sending a message to the console is `NSLog`. It's a C function, and it takes an `NSString` which operates as a format string, followed by the format arguments.

A *format string* is a string (here, an `NSString`) containing symbols called *format specifiers*, for which values (the format arguments) will be substituted at runtime. All format specifiers begin with a percent sign (%), so the only way to enter a literal percent sign in a format string is as a double percent sign (%%). The character(s) following the percent sign specify the type of value that will be supplied at runtime. The most common format specifiers are %@ (an object reference), %i (an integer), %f (a float), and %p (a pointer, usually an object reference, shown as the address in memory pointed to, useful for making certain that two references refer to the same instance). For example:

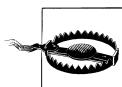
```
NSLog(@"the window: %@", self.window);
```

In that example, `self.window` is the first (and only) format argument, so its value will be substituted for the first (and only) format specifier, %@, when the format string is printed in the console. Thus the console output looks something like this:

```
the window: <UIWindow: 0x6a08140; frame = (0 0; 320 480);
layer = <UIWindowLayer: 0x6a08230>
```

This nice display of information is due to `UIWindow`'s implementation of the `description` method: an object's `description` method is called when that object is used with the %@ format specifier. For this reason, you will probably want to implement `description` in your own classes, so that you can investigate an instance with a simple `NSLog` call.

For the complete repertory of format specifiers available in a format string, read Apple's document *String Format Specifiers*. The format specifiers are largely based on those of the C `printf` standard library function; see K&R B1.2, the `sprintf` man page, and the IEEE `printf` specification linked from the documentation.



If an object reference has been set to `nil`, `NSLog` will report it as `(null)`. But if an object reference is uninitialized, an `NSLog` call referring to it will probably fail silently, or even crash the debugger. This is very frustrating; indeed, the fact that this object reference is uninitialized is probably just what you were trying to debug. Fortunately, if you use ARC ([Chapter 12](#)), an object reference is autoinitialized to `nil` if you don't initialize it explicitly.

The main ways to go wrong with `NSLog` (or any format string) are to supply a different number of format arguments from the number of format specifiers in the string, or to supply an argument value different from the type declared by the corresponding format specifier. These mistakes can send your app off into the weeds, or at least give misleading results. I often see beginners claim that logging shows a certain value to be nonsense, when in fact it is their `NSLog` call that is nonsense; for example, a format specifier was `%i` but the value of the corresponding argument was a float.

C structs are not objects, so to see a struct's value with NSLog you must somehow deconstruct or translate the struct. Common Cocoa structs usually supply convenience functions for this purpose. For example:

```
NSLog(@"%@", NSStringFromCGRect(self.window.frame)); // {{0, 0}, {320, 480}}
```

Purists may scoff at caveman debugging, but I use it heavily: it's easy, informative, and lightweight. And sometimes it's the only way. Unlike the debugger, NSLog works with any build configuration (Debug or Release) and wherever your app runs (in the Simulator or on a device). It even works on someone else's device, such as a tester to whom you've distributed your app. It's a little tricky for a tester to get a look at the console so as to be able to report back to you, but it can be done: the tester can connect the device to a computer and view its log in Xcode's Organizer window or with Apple's iPhone Configuration Utility; there's also a free utility app called Console that displays the log right on the device.

Remember to remove or comment out NSLog calls before shipping your app, as you probably don't want your finished app to dump lots of messages into the console. A useful trick (shamelessly stolen from Jens Alfke) is to call MyLog instead of NSLog, and define MyLog like this in your precompiled header:

```
#define MyLog if(0); else NSLog
```

When it's time to stop logging, change the 0 to 1.

A useful fact when logging is that the variable name `_cmd` holds the selector for the current method. Thus a single form of statement can signal where you are:

```
NSLog(@"Logging %@ in %@", NSStringFromSelector(_cmd), self);
```

(Similarly, in a C function, `NSLog(@"%@", __FUNCTION__)` logs the name of the function.)

Another sort of call with which you can pepper your code is *asserts*. Asserts are conditions that you claim (`assert`) are true at that moment — and you feel so strongly about this that you want your app to crash if you're wrong. Asserts are a very good way to confirm that the situation matches your expectations, not just now as you write your code, but in the future as the app develops.

The simplest form of assert is the C function (actually it's a macro) `assert()`, to which you pass one argument, a condition — something that can be evaluated as false (0) or true (some other value). If it's false, your app will crash when this line is encountered, along with a nice explanation in the log. For example, suppose we assert `NO`, which is false and will certainly cause a crash. Then when this line is encountered we crash with this log message:

```
Assertion failed: (NO),
function -[AppDelegate application:didFinishLaunchingWithOptions:],
file /Users/mattleopard/Desktop/testing/testing/AppDelegate.m, line 20.
```

That's plenty for us to track down the assertion failure: we know the assertion condition, the method in which the assertion occurred, the file containing that method, and the line number.

For higher-level asserts, look at `NSAssert` (used in Objective-C methods) and `NSCAssert` (used in C functions); they allow you to form your own log message, to appear in the console in addition to the native `assert` logging, and the log message can be a format string followed by values corresponding to the format specifiers, as with `NSLog`.

Some developers even think that asserts should be allowed to remain in your code when your app is finished. By default, however, they are disabled in a Release build. This is done through the Other C Flags build setting, which is set to `-DNS_BLOCK_ASSERTIONS=1` in Apple's project templates; the effect is to `#define` the pre-processor macro `NS_BLOCK_ASSERTIONS`, which in turn is the signal for asserts to be effectively neutered at precompile time. To keep asserts working in a Release build, clear that value from the Other C Flags build setting in your target.



As of this writing, assertions don't log properly in Xcode 4.2's console when running in the iOS 5.0 Simulator. The good news, however, is that they do work when running on a device; I believe this is a consequence of my having updated my device from iOS 5 to iOS 5.0.1.

The Xcode Debugger

When you're building and running in Xcode, you can pause in the debugger and use Xcode's debugging facilities. There isn't a strong difference between running and debugging in Xcode 4; the main distinction is whether breakpoints are effective or ignored. The effectiveness of breakpoints can be toggled at two levels:

Globally

Breakpoints as a whole are either active or inactive. If breakpoints are inactive, we won't pause at any breakpoints.

Individually

A given breakpoint is either enabled or disabled. Even if breakpoints are active, we won't pause at this one if it is disabled. Disabling a breakpoint allows you to leave in place a breakpoint that you might need later without pausing at it every time it's encountered.

A breakpoint, then, is ignored if it is disabled or if breakpoints as a whole are inactive.

The important thing, if you want to use the debugger, is that the app should be built with the Debug build configuration. The debugger is not very helpful against an app built with the Release build configuration, not least because compiler optimizations can destroy the correspondence between steps in the compiled code and lines in your code. Trying to debug a Release build is a common beginner error (though it's less likely

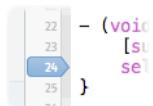


Figure 9-4. A breakpoint

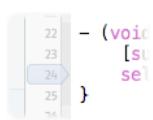


Figure 9-5. A disabled breakpoint

to occur accidentally in Xcode 4, in which by default a scheme's Run action uses the Debug build configuration).

To create a breakpoint ([Figure 9-4](#)), select in the editor the line where you want to pause, and choose Product → Debug → Add Breakpoint at Current Line (Command-Backslash). This keyboard shortcut toggles between adding and removing a breakpoint for the current line. The breakpoint is symbolized by an arrow in the gutter. Alternatively, a simple click in the gutter adds a breakpoint; to remove a breakpoint gesturally, drag it out of the gutter.

To disable a breakpoint at the current line, click on the breakpoint in the gutter to toggle its enabled status. Alternatively, Control-click on the breakpoint and choose Disable Breakpoint in the contextual menu. A dark breakpoint is enabled; a light breakpoint is disabled ([Figure 9-5](#)).

Once you have some breakpoints in your code, you'll want to survey and manage them. That's what the Breakpoint navigator is for. Here you can navigate to a breakpoint, enable or disable a breakpoint by clicking on its arrow in the navigator, and delete a breakpoint.

You can also edit a breakpoint's behavior. Control-click on the breakpoint, in the gutter or in the Breakpoint navigator, and choose Edit Breakpoint. This is a very powerful facility: you can have a breakpoint pause only under a certain condition or after it has been encountered a certain number of times, and you can have a breakpoint perform a certain action when it is encountered, such as logging or running a script.

A breakpoint can be configured to continue automatically after performing its action when it is encountered. This can be an excellent alternative to caveman debugging: Instead of inserting an NSLog call, which must be compiled into your code and later removed when the app is released, you can set a breakpoint that logs and continues, which operates only when you're debugging.

In the Breakpoint navigator, you can create two kinds of breakpoint that you can't create in a code editor: exception breakpoints and symbolic breakpoints. Click the "+" button at the bottom of the navigator and choose from its pop-up menu.

Exception breakpoint

An exception breakpoint causes your app to pause at the time an exception is thrown or caught, without regard to whether the exception would crash your app later. I recommend that you create an exception breakpoint to catch all exceptions when they are thrown, because this gives the best view of the call stack and variable values at the moment of the exception (rather than later when the crash actually occurs); you can see where you are in your code, and you can examine variable values, which may help you understand the cause of the problem. If you do create such an exception breakpoint, I also suggest that you use the contextual menu to say Move Breakpoint To → User, which makes this breakpoint permanent and global to all your projects.

Symbolic breakpoint

A symbolic breakpoint causes your app to pause when a certain method is called, regardless of what object called it or to what object the message is sent. The method name is entered in a special way — the instance method or class method symbol (- or +) followed by square brackets containing the class name and the method name. For example, to learn where in my app the `beginReceivingRemoteControlEvents` message was being sent to my shared application instance, I configured a symbolic breakpoint like this:

```
-[UIApplication beginReceivingRemoteControlEvents]
```

To toggle the active status of breakpoints as a whole, click the Breakpoints button in the project window toolbar, or choose Product → Debug → Activate/Deactivate Breakpoints (Command-Y). The active status of breakpoints as a whole doesn't affect the enabled or disabled status of any breakpoints; if breakpoints are inactive, they are simply ignored en masse, and no pausing at breakpoints takes place. Breakpoint arrows are blue if breakpoints are active, gray if they are inactive.

When the app runs with breakpoints active and an enabled breakpoint is encountered (and assuming its conditions are met, and so on), the app pauses. In the active project window, the editor shows the file containing the point of execution, which will usually be the file containing the breakpoint. The point of execution is shown as a green arrow; this is the line that is *about* to be executed (Figure 9-6). Depending on the settings for "Run pauses" in the Behaviors preference pane, the Debug navigator and the Debug pane will also appear.

Here are some things you might like to do while paused at a breakpoint:

See where you are

One common reason for setting a breakpoint is to make sure that the path of execution is passing through a certain line. You can see where you are in any of your

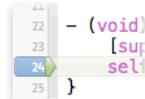


Figure 9-6. Paused at a breakpoint

methods by clicking on the method name in the call stack, shown in the Debug navigator.

Methods listed in the call stack with a User icon, with the text in black, are yours; click one to see where you are paused in that method. Other methods, with the text in gray, are methods for which you have no source code, so there would be little point clicking one unless you know something about assembly language. The slider at the bottom of the navigator hides chunks of the call chain, to save space, starting with the methods for which you have no source.



You can also navigate the call stack using the jump bar at the top of the Debug pane.

Study variable values

This is a very common reason for pausing. In the Debug pane, variable values for the current scope (corresponding to what's selected in the call stack) are visible in the variables list. You can see additional object features, such as collection elements, instance variables, and even some private information, by opening triangles.

Switch the pop-up menu above the variables list to Auto to see only those variables that Xcode thinks will interest you (because their value has been recently changed, for instance); if you're after completeness, Local will probably be the best setting. You can use the search field to filter variables by name or value.

In some cases, toggling Show Formatted Summaries in the contextual menu can give a faster or more reliable display of variables. Even with formatted summaries turned off, you can send `description` to an object variable and view the output in the console by selecting Print Description from the contextual menu.

Set a watchpoint

A watchpoint is like a breakpoint, but instead of depending on a certain line of code it depends on a variable's value: the debugger pauses whenever the variable's value changes. You can set a watchpoint only while paused in the debugger. Control-click on the variable in the variables list and choose Watch [Variable]. Watchpoints, once created, are listed and managed in the Breakpoint navigator.

Manage expressions

An expression is code to be added to the variables list and evaluated every time we pause. Choose Add Expression from the contextual menu in the variables list.

Talk to the debugger

You can communicate verbally with the debugger in the console. The most common command is `po` (for “print object”) followed by an object variable’s name or a method call that returns an object; it calls the object’s `description` method, just like `NSLog`.



Xcode’s debugger is a front end to an open source third-party command-line debugger tool. Thus, by talking directly to that command-line tool you can do everything that you can do through the Xcode debugger interface, and more. Throughout the history of Xcode up through Xcode 3.2.x, the debugger tool has been GDB; see *Debugging with GDB*. Starting in Xcode 4.2, the debugger LLDB is available as an alternative (<http://lldb.llvm.org>). LLDB is said to be the way of the future, and it has some powerful features, but it lacks watchpoints and there are some wrinkles that still need ironing out; for example, LLDB cannot see instance variables declared in a class’s implementation section. Still, feel free to try it. To do so, use the Scheme editor; in the Info pane of the Run action, switch the Debugger pop-up menu to LLDB.

Fiddle with breakpoints

You are free to create, destroy, enable and disable, and otherwise manage breakpoints dynamically even though your app is running, which is useful because where you’d like to pause next might depend on what you learn while you’re paused here.

Step or continue

To proceed with your app, you can either resume running (Product → Debug → Continue) until the next breakpoint is encountered or take one step and pause again. Also, if you hover the mouse over the gutter, a green Continue to Here button appears; pressing this, or alternatively choosing Product > Debug > Continue to Current Line (or Continue to Here in the contextual menu), effectively sets a breakpoint at the chosen line, continues, and removes the breakpoint.

The stepping commands (under Product → Debug) are:

Step Over

Pause at the next line.

Step Into

Pause in your method that the current line calls, if there is one; otherwise, pause at the next line.

Step Out

Pause when we return from the current method.

You can access these commands through convenient buttons at the top of the Debug pane. Even if the Debug pane is collapsed, the part containing the buttons appears while running.

You can also float the project window over everything else on your computer while debugging by choosing Product → Debug Workflow → Xcode Always In Front; after you then switch to the Simulator, you can interact with the Xcode window without giving it focus. If you do want to give it focus, to type in a filter bar for instance, click Focus in the toolbar. This mode of working could be useful while you're interacting with the Simulator, so as not to have to keep switching between the Simulator and Xcode. To return the project window to its normal state, choose Normal from the Debugging pop-up menu in the window toolbar (or click Stop in the toolbar to kill the running app).



Step Over and Step Into have advanced forms where you hold Control to step by machine-level instruction, and Control-Shift to step while blocking all other threads.

Start over, or abort

To kill the running app, click Stop in the toolbar (Product → Stop, Command-Period). To kill the running app and relaunch it without rebuilding it, Control-click Run in the toolbar (Product → Perform Action → Run Without Building, Control-Command-R). You can make changes to your code while the app is running or paused, but they are not magically communicated to the running app; you must run in the normal way (which includes building) to see your changes in action.



Clicking the Home button in the Simulator or on the device does *not* stop the running app in the multitasking world of iOS 4 and iOS 5.

Static Analyzer

From time to time, you should use the static analyzer to look for possible sources of error in your code; choose Product → Analyze (Shift-Command-B). This command causes your code to be compiled, and the static analyzer studies it and reports its findings in the Issue navigator and in your code.

The static analyzer is static (it's analyzing your code, not debugging in real time), but it is remarkably intelligent and may well alert you to potential problems that could otherwise escape your notice. You might think that the compiler — including ARC, if you're using it — knows all there is to know about your code; and it is certainly true that one of the main reasons for using the static analyzer, namely, to assist with manual

memory management of Objective-C instances, is essentially gone if you’re using ARC. Still, the static analyzer actually studies the possible values and paths of execution in your code, and can detect potential sources of trouble in your program’s *logic* that no mere compiler would worry about. For example, in this code, the static analyzer knows that `i` in the second line is uninitialized:

```
int i;
if (i) NSLog(@"here");
```

In this code, the static analyzer knows that the second line throws away the existing value of `i` without that value ever having been read:

```
int i=0;
i=1;
```

Those are tiny problems, but they illustrate how, in a complex program, the static analyzer is capable of noticing possible sources of trouble and bringing them to your attention. For more about the static analyzer, see <http://clang-analyzer.llvm.org>.

Clean

From time to time, during repeated testing and debugging, and before making a different sort of build (switching from Debug to Release, or running on a device instead of the Simulator), it is a good idea to *clean* your target. This means that existing builds will be removed and caches will be cleared, so that all code will be considered to be in need of compilation and the next build will build your app from scratch.

The first build of your app after you clean will take longer than usual. But it’s worth it, because cleaning removes the cruft, quite literally. For example, suppose you have been including a certain resource in your app, and you decide it is no longer needed. You can remove it from the Copy Bundle Resources build phase (or from your project as a whole), but that doesn’t remove it from your built app. Only cleaning will do that, because it removes the built app completely.

To clean, choose Product → Clean. For more complete cleaning, hold Option to get Product → Clean Build Folder.

In addition, Xcode 4 stores builds and project indexes in `~/Library/Developer/Xcode/DerivedData`. From time to time, with Xcode not running, I like to move the contents of that folder to the trash. This is effectively a massive and even more complete clean of every project that you’ve opened recently. Alternatively, to trash the folder in *DerivedData* for a single project from within Xcode, switch to the Projects tab of the Organizer window (Window → Organizer), select the project at the left, and click the Delete button next to the Derived Data listing at the top of the window. A project will take longer to open for the first time afterward, because its index must be rebuilt, and it will take longer to build, because its build information has been removed. But the space savings on your hard disk can be significant, and forcing the index to rebuild itself can actually ward off certain problems.

You should also from time to time remove all versions of your built app from the Simulator cache. Choose iOS Simulator → Reset Content and Settings. Alternatively, you can clean the Simulator cache by hand. To do so, first quit the Simulator if it's running. Then find the cache in `~/Library/Application Support/iPhone Simulator`, followed by the system version of the SDK (for example, there might be a folder called `5.0`); within this, find the *Applications* folder, and move the contents of that folder to the trash. If there are multiple system version folders here, you might want to jettison the contents of the *Applications* folders of all of them.

Running in the Simulator

When you build and run with Simulator as the destination, you run in the iOS Simulator application. The Simulator window represents a device. If your app runs on either iPhone or iPad (natively or in the iPhone emulator), you can choose which device is simulated as you choose your destination; similarly, if your app runs on multiple system versions, you can choose the system version of the simulated device as you choose your destination. (See [Chapter 6](#) on destinations, and the first section of this chapter on device architectures and the Deployment Target build setting.)

You can also switch device types by choosing Hardware → Device in the Simulator. This quits your app running in the Simulator; you can relaunch it by building and running in Xcode again, or by clicking your app's icon in the Simulator's springboard. In the latter case there is no longer any connection to Xcode (you aren't using the debugger, so you won't stop at breakpoints, and log messages won't be relayed to the Xcode console); still, you might do this just to check quickly on how your app looks or behaves on a different device.

The one key choice you can make using Hardware → Device in the Simulator that you *can't* make by choosing a destination in Xcode is between iPhone and iPhone (Retina). The latter has a double-resolution screen, so it is displayed at double size, with each pixel of the Retina display corresponding to a pixel of your computer's monitor. Changing this setting quits your app running in the Simulator, but your choice will stick if you return to Xcode and build and run on iPhone again.

The iPad device can be displayed at half, three-quarter, or full size (choose from Window → Scale). This is a matter of display merely, comparable to zooming a window, so your app running in the Simulator does not quit when you change this setting.

You can interact with the Simulator in some of the same basic ways as you would a device. Using the mouse, you can click the Home button and tap on the device's screen; hold Option to make the mouse represent two fingers and Option-Shift to move those fingers in parallel. Menu items let you perform hardware gestures such as rotating the device, shaking it, and locking its screen; you can also test your app by simulating certain rare events, such as a low-memory situation (and this is a useful thing to do from time to time; I'll talk more about it in [Chapter 19](#)).

New in Xcode 4.2 is the Debug menu in the Simulator. This is a valuable addition, useful for detecting problems with animations and drawing. You can choose from this menu while your app is running in the Simulator, without causing the app to quit. Toggle Slow Animations is unique to the Simulator; it makes animations unfold in slow motion so that you can see just what's happening (animation is discussed in [Chapter 17](#)). The other four menu items represent features that were previously available only when running on a device using Instruments (mentioned later in this chapter), under the Core Animation instrument; now they are rolled directly into the Simulator as well.

I'll return to the specifics of what these menu items do when discussing drawing ([Chapter 15](#)) and layers ([Chapter 16](#)); but here's an example you can try immediately. Return to the Empty Window project developed in [Chapter 7](#). In *AppDelegate.m*, we are pulling a label out of a nib and putting it into our interface by setting its center, like this:

```
[self.window.rootViewController.view addSubview: lab];
lab.center = CGPointMake(100,100);
lab.frame = CGRectMakeIntegral(lab.frame);
```

If you comment out that last line, run the project in the Simulator with the device set to iPhone — not iPhone (Retina) — and toggle on Debug → Color Misaligned Images, you may see the label painted with a magenta overlay. That's because, without the call to `CGRectIntegral`, the label is misaligned; by default, the label is 21 points high, which is an odd number, so setting its center to an integral point value has caused its vertical position to be halfway between two integer pixel values on the device. The effect of this misalignment is actually visible to the naked eye if you know what to look for: the text looks fuzzy or bold. Using Debug → Color Misaligned Images alerts you to the issue; calling `CGRectIntegral` fixes it.

Running on a Device

Sooner or later, you're going to want to switch from running and testing and debugging in the Simulator to running and testing and debugging on a real device. The Simulator is nice, but it's only a simulation; there are many differences between the Simulator and a real device. The Simulator is really your computer, which is fast and has lots of memory, so problems with memory management and speed won't be exposed until you run on a device. User interaction with the Simulator is limited to what can be done with a mouse: you can click, you can drag, you can hold Option to simulate use of two fingers, but more elaborate gestures can be performed only on an actual device. And many iOS facilities, such as the accelerometer and access to the music library, are not present on the Simulator at all, so that testing an app that uses them is possible *only* on a device.



Don't even think of developing an app without testing it on a device. You have no idea how your app *really* looks and behaves until you run it on a device. Submitting to the App Store an app that you have not run on a device is asking for trouble.

Before you can run your app on a device, even just to test, you must join the iOS Developer Program by paying the annual fee. (Yes, this is infuriating. Now get over it.) Only in this way can you obtain and provide to Xcode the credentials for running on a device. Once you have joined the iOS Developer Program, obtaining these credentials involves use of the iOS Provisioning Portal, which is accessed online, through your web browser (or, for certain actions, through Xcode itself).



To reach the iOS Provisioning Portal in your browser (once you're an iOS Developer Program member), go to http://developer.apple.com/dev_center/ios. Click Log In to log in, and then click iOS Provisioning Portal at the upper right.

You will need to perform the following steps just once:

1. Join the iOS Developer Program (<http://developer.apple.com/programs/ios>). This requires filling out a form and paying the annual fee. Unless you have multiple developers, all of whom might need to build and run on their own devices, the Individual program is sufficient. The Company program costs no more, but adds the ability to privilege additional developers in various roles. (You do *not* need the Company program just in order to distribute your built app to other users for testing.)
2. Obtain a *development certificate* that identifies and authorizes your computer. This is the computer to which you'll be attaching the device so you can run on it. Basically, this certificate matches the person who uses your computer to the person interacting with the iOS Provisioning Portal. The certificate will be stored in your computer's keychain, where Xcode will be able to see it automatically.

The certificate depends upon a private–public key pair. The private key will live in your keychain; the public key will be handed over to the iOS Provisioning Portal, to be built into the certificate. The way you give the Portal your public key is through a *request* for the certificate. So, you generate the private–public key pair; your keychain keeps the private key; the public key goes into the certificate request; you submit the request, containing the public key, to the Portal; and the Portal sends back the certificate, also containing the public key, which also goes into your keychain, where it is matched with the private key, thus ensuring that you are you.

Detailed instructions for generating the private–public key pair and the certificate request are available once you've joined the iOS Developer Program and have logged in at Apple's developer site. (The process is described at <http://developer.apple.com/ios/>

manage/certificates/team/howto.action. A video review of the steps involved is available to anyone at <http://developer.apple.com/ios/videos/popupcerts.action>.) Basically, you start up Keychain Access and choose Keychain Access → Certificate Assistant → Request a Certificate from a Certificate Authority. Using your name and email address as identifiers, you generate and save to disk a 2048-bit RSA certificate request file. Your private key is stored in your keychain then and there; the certificate request contains your public key.

You then go to the iOS Provisioning Portal in your browser. At the Portal, upload the certificate request file using the Development (not Distribution!) tab of the Certificates section. You may have to approve your own request.



If this is your very, very first time obtaining any certificate from the Portal, you will need *another* certificate: the WWDR Intermediate Certificate. This is the certificate that certifies that certificates issued by WWDR (the Apple Worldwide Developer Relations Certification Authority) are to be trusted. (You can't make this stuff up.) You'll see a link for this intermediate certificate; click it to download the intermediate certificate. Double-click the intermediate certificate file; it is imported by your keychain. You can then throw the file away.

When the development certificate itself is ready, you download it and double-click it; Keychain Access automatically imports the certificate and stores it in your keychain. You do not need to keep the certificate request file or the development certificate file; your keychain now contains all the needed credentials. If this has worked, you can see the certificate in your keychain, read its details, and observe that it is valid and linked to your private key ([Figure 9-7](#)). After you've done this once, your development certificate is good for all your app development from now on. (However, your development certificate expires when your year of iOS Developer Program membership expires; if you renew your membership, you'll have to revoke your current development certificate at the Portal, delete it from your keychain, and repeat the process of obtaining a new one.)

With your development certificate in place, you need to register a device for development use, meaning that you'll be able to build and run from Xcode onto that device rather than the Simulator. This can be done entirely from within Xcode. Open the Organizer window (Window → Organizer) and switch to the Devices tab. Select Provisioning Profiles under Library at the left, and make sure Automatic Device Provisioning is checked at the bottom of the window. Attach your device to the computer; the device name appears at the left under Devices. Select it, and click Use For Development. You'll be asked for your Portal username and password. Xcode connects to the Portal via the Internet and does two things:

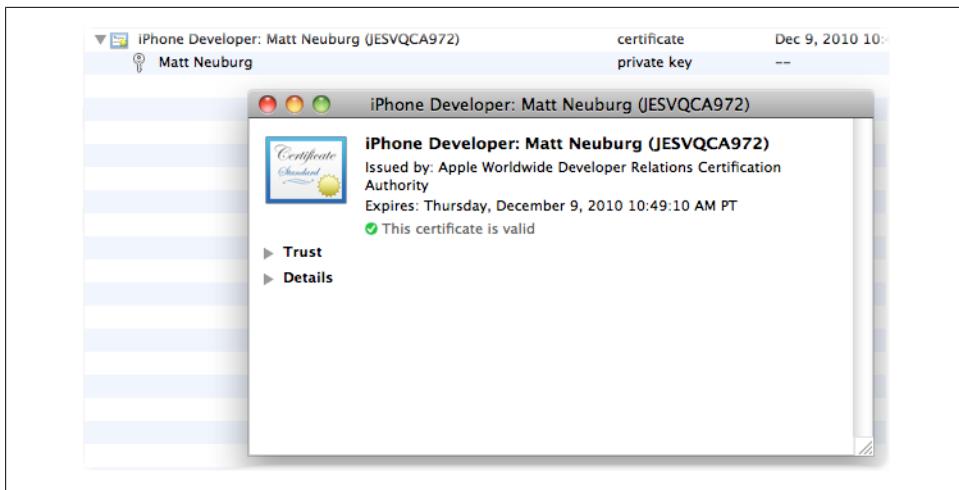


Figure 9-7. A valid development certificate, as shown in Keychain Access

- It registers your device at the Portal by its name and unique identifier number. You could have done this yourself in your browser (at the Portal, under Devices), but this way it is done for you.)
- It creates and downloads from the Portal a universal development provisioning profile (referred to as a Team Provisioning Profile) for development on this device. This is something you can't do at the Portal yourself. A development provisioning profile created manually at the Portal applies to a single app; in the past, when the Portal was the only way to obtain a development provisioning profile, you had to generate a new development provisioning profile for each app you wanted to test on a device, which was very inconvenient. But the development provisioning profile generated by Xcode applies to all apps, now and in the future (until it expires, at which time it can easily be regenerated). The universal development provisioning profile appears in the Organizer, under Provisioning Profiles; you can identify it because it is called iOS Team Provisioning Profile and has an app identifier consisting of just a key and an asterisk, like this: B398E68A3D.*.

If your device is already registered at the Portal but Use For Development doesn't appear in Xcode and you've no team provisioning profile, go back to Provisioning Profiles under Library and click Refresh at the bottom of the window. The team provisioning profile will be regenerated.



If you develop an app that uses certain specialized features, such as push notifications ([Chapter 37](#)) or in-app purchases, you must generate a development provisioning profile the old way, manually at the Portal. To do so, first enter your app by name and bundle id in the App IDs section of the Portal. Now go to Provisioning and the Development section and generate a new provisioning profile, specifying that app and your device(s). You can then download the provisioning profile in the Organizer window, under Provisioning Profiles, by clicking the Refresh button.

You can install the provisioning profile onto your device manually in the Organizer window by dragging its listing (under Provisioning Profiles) onto the device's name (under Devices). Alternatively, you can just start building and running on the device. Start with a project window. With the device attached to the computer, pick the destination in the Scheme pop-up menu corresponding to your device; then build and run. If Xcode complains that your device doesn't contain a copy of the provisioning profile, and offers to install it for you, accept that offer.

The app is built, loaded onto your device, and runs. As long as you launch the app from Xcode, everything is just as it was before: you can run, or you can debug, and the running app is in communication with Xcode, so that you can stop at breakpoints, read messages in the console, and so on. The outward difference is that to interact physically with the app, you use the device, not the Simulator.

Starting in Xcode 4.2, an option in the General preference pane offer to Support Wirelessly Connected Devices. In theory this should allow Xcode to install and debug an app on your device *without* that device being physically connected to your computer. For this to work, the device must also be set up in iTunes for wireless syncing. The indications are, however, that this feature is not yet ready for prime time (and you may be reluctant to enable iTunes wireless syncing for some other reason), so I don't recommend it.

Device Management

Your central location for management of identities (certificates), provisioning profiles, and devices is the Devices tab of the Organizer window (Window → Organizer). Under Library, select Developer Profile to see your identities and provisioning profiles. Select Provisioning Profiles for another list of profiles, along with a button that lets you easily renew an expired profile. Both these panes offer the ability to export and import identities and profiles, useful if you move development to a different machine.

When your device is attached to the computer, it is listed with a green dot under Devices. Click its name to access information on the device. You can see the device's unique identifier. You can see provisioning profiles that have been installed on the device. You can view the device's console log in real time, just as if you were running

the Console application to view your computer’s logs. You can see log reports for crashes that took place on the device. And you can take screenshots that image your device’s screen; you’ll need to do this for your app when you submit it to the App Store. Crash reports and screenshots are also available under Library.

Version Control

Various systems of *version control* exist for taking periodic snapshots (technically called *commits*) of your project. The value of such a system to you will depend on what system you use and how you use it; for example, you might use version control because it lets you store your commits in a repository offsite, so that your code isn’t lost in case of a local computer glitch or some equivalent “hit by a bus” scenario, or because it allows multiple developers to access the same code.

To me, personally, the chief value of version control is *freedom from fear*. Having version control actually changes the way I program. A project is a complicated thing, consisting of numerous files. Often, changes must be made in many files before a new feature can be tested. Thus it is all too easy to start down some virtual road involving creating or editing multiple files, only to find yourself at the end of a blind alley and needing to retrace your steps. Version control means that I can easily retrace my steps; I have but to say, in the language of some version control system I’ve been using, “Forget everything I just did and return the whole project to where it was at such-and-such a commit.” I rarely, if ever, *in fact* retrace my steps, but the knowledge that I *could* do so gives me the courage to try some programming strategy whose outcome may not be apparent until after many days of effort. Also, I can ask a version control system, “What the heck are all the changes I’ve made since the last commit?” In short, without version control I’d be lost, confused, hesitant, rooted to the spot, paralyzed with uncertainty; with it, I forge boldly ahead and get things done. For this reason, my current personal favorite version control system is git (<http://git-scm.com>), whose agile facilities for managing branches give me tremendous license to experiment.

Xcode provides various version control facilities. Starting with Xcode 4, those facilities concentrate on git and Subversion (<http://subversion.apache.org>). This doesn’t mean you can’t use any other version control system with your projects! It means only that you can’t use any other version control system in an integrated fashion from inside Xcode. Personally, I don’t find that to be any kind of restriction. For years I’ve used Subversion, and more recently git, on my Xcode projects from the command line in Terminal, or using other third-party GUI front ends (such as svnX for Subversion, <http://www.lachoseinteractive.net/en/products>). I’m comfortable and nimble at the command line, and access to version control from within Xcode itself is not a priority for me.

At the same time, version control integration in Xcode 4 is greatly improved and far more extensive than previously:

```

20 // we create the button in code and configure its
21 // action in code
22 UIButton* b = [UIButton buttonWithType:
23     UIButtonTypeRoundedRect];
24 [b setTitle:@"Howdy!" forState:UIControlStateNormal]
25 [b setFrame:CGRectMake(100,100,100,35)];
26 [self.window.rootViewController.view addSubview:b];
27 [b addTarget:self action:@selector(buttonPressed:)
28     forControlEvents:UIControlEventTouchUpInside];
29
30 // we create the button in code and configure its
31 // action in code
32 UIButton* b = [UIButton buttonWithType:
33     UIButtonTypeRoundedRect];
34 [b setTitle:@"Howdy!" forState:UIControlStateNormal]
35 [b setFrame:CGRectMake(100,100,100,35)];
36 [self.window addSubview:b];
37 [b addTarget:self action:@selector(buttonPressed:)
38     forControlEvents:UIControlEventTouchUpInside];
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74

```

Figure 9-8. Version comparison

Automatic git repository creation

When you create a new project in Xcode 4, the Save dialog includes a checkbox that offers to place a git repository into your project folder from the outset.

Automatic repository detection

When you open an existing project in Xcode 4, if that project is already managed with Subversion or git, Xcode detects this and is ready instantly to display version control information in its interface.

Version comparison

The Version editor (View → Version Editor → Show Version Editor, or click the third button in the Editor segmented control in the project window toolbar) includes a view similar to that of the File Merge utility, graphically displaying the differences between versions of a file. For example, in [Figure 9-8](#), I can see that in the more recent version of this file (on the left) I've changed `self.window` to `self.window.rootViewController.view`. (This change was necessary for the second edition of this book, because iOS 5 now complains if the window has no root view controller.) The Version editor also includes various ways to survey and navigate versions and commit logs. (For example, if I switch to Blame view I can see my own commit message associated with the change I just mentioned: "Giving window a root view controller and putting subviews inside that, so that no views ever go directly into the window any more even in simple examples.")

Without minimizing these features, I don't rely on them exclusively or even primarily (although I certainly take advantage of them where convenient). I find version control management through the command line far easier and clearer for many purposes, and Xcode doesn't come close to the command line's power, especially for managing branches; and Xcode has nothing like the visual branch representation of git's own `gitk` tool. (In fact, as of this writing, Xcode's git integration is fundamentally flawed: if I create a new file in a git-controlled project, Xcode adds it to the staging area rather than letting me compose the commit in my own way, so that I typically have to switch to the command line and say `git reset` to undo Xcode's action.)

Version control in general is a large and complicated topic. Use and configuration of any version control system can be tricky and scary at first and always requires some care. So I'm deliberately not going to say anything specific about it; I'm mentioning it at all only because version control of some sort is in fact likely, sooner or later, to play a role in the life cycle of your projects. When it does, you'll want to read up on the use of your chosen version control system, along with "Managing Versions of Your Project"

in the *Xcode 4 User Guide*. You'll find Xcode 4's integrated version control facilities in three chief locations:

The File menu

The relevant menu items are all under File → Source Control.

The Version editor

Choose View → Version Editor → Show Version Editor, or click the third button in the Editor segmented control in the project window toolbar.

The Organizer

The Repositories tab of the Organizer window lists known repositories and branches for each project, along with their commit logs. Also, use the “+” button at the bottom of the navigator to enter data about a remote repository, so that you can obtain a copy of its contents.

Xcode also contains its own way of taking and storing a snapshot of your project as a whole; this is done using File → Create Snapshot (and, according to your settings, some mass operations such as find-and-replace or refactoring may offer to take a snapshot first). Snapshots themselves are managed in the Projects tab of the Organizer window. Although these snapshots are not to be treated as full-fledged version control, they are in fact maintained as git repositories, and can certainly serve the purpose of giving confidence in advance of performing some change that might subsequently engender regret. You can manage snapshots in the Projects tab of the Organizer window; here you export a snapshot, thus resurrecting an earlier state of your project folder.

Instruments

As your app approaches completion, you may wish to fine-tune it for memory usage, speed, and other real-time behavior. Xcode provides a sophisticated and powerful utility application, Instruments, that lets you collect profiling data on your app as it runs. The graphical display and detailed data provided by Instruments may give you the clues you need to optimize your app.

You can use Instruments on the Simulator or the device. The device is where you'll do your ultimate testing, and certain instruments (such as Core Animation) are available only for the device; on the other hand, certain other instruments (such as Zombies) are available only in the Simulator.

To get started with Instruments, set the desired destination in the Scheme pop-up menu in the project window toolbar, and choose Product → Profile. (For memory usage, your destination can be the Simulator, but for most other forms of analysis, you'll want to run on a device.) Your app builds using the Profile action for your scheme; by default, this uses the Release build configuration, which is probably what you want. Instruments launches; if your scheme's Instrument pop-up menu for the Profile action is set to Ask on Launch, Instruments presents a dialog where you choose a trace template.

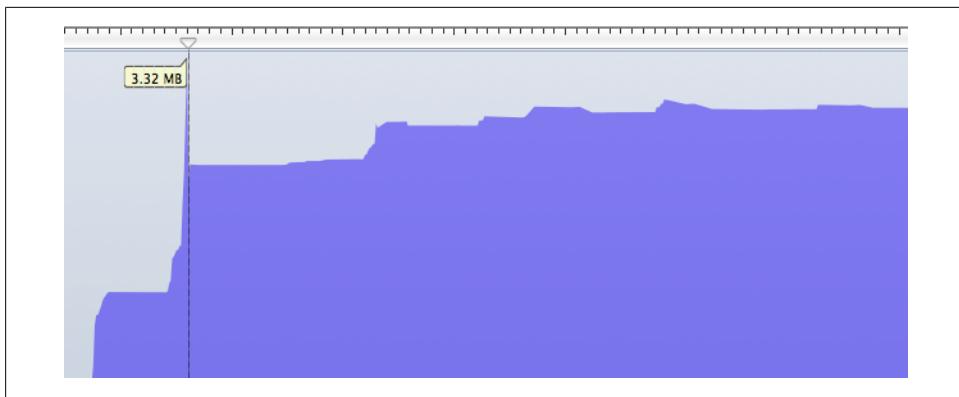


Figure 9-9. Instruments graphs memory usage over time

With Instruments running, you should interact with your app like a user; Instruments will record its statistics. Once Instruments is running, it can be further customized to profile the kind of data that particularly interests you, and you can save the structure of the Instruments window as a custom template.

Use of Instruments is an advanced topic and beyond the scope of this book. Indeed, an entire book could (and really should) be written about Instruments alone. For proper information, you should read Apple's documents, *Instruments User Guide* and *Instruments New Features User Guide*; also, many WWDC 2010 and 2011 videos are about Instruments; look for sessions with "Instruments" or "Performance" in their names. Here, I'll just demonstrate, without much explanation, the sort of thing Instruments can do.

I'll start by charting the memory usage of my TidBITS News app as it starts up and the user proceeds to work with it. Memory is a scarce resource on a mobile device, so it's important to be certain that we're not hogging too much of it. I'll set the destination to the Simulator and choose Product → Profile; Instruments launches, and I'll choose the Allocations trace template and click Profile. My app starts running in the Simulator, and I work with it for a while and then pause Instruments, which meanwhile has charted my memory usage (Figure 9-9). Examining the chart, I find there's a spike early on, as the app downloads and parses an RSS feed; but it's only 3.32 MB at its maximum, and the app settles down to use less than 3 MB pretty steadily thereafter. These are very gentle memory usage figures, so I'm happy.

Another field of Instruments expertise is the ability to detect memory leaks. Memory leaks, discussed further in Chapter 12, remain possible even under ARC. In this trivial example, I have two classes, MyClass1 and MyClass2; MyClass1 has an `ivar` property which is a MyClass2 instance, and MyClass2 has an `ivar` property which is a MyClass1 instance. The app runs this code:

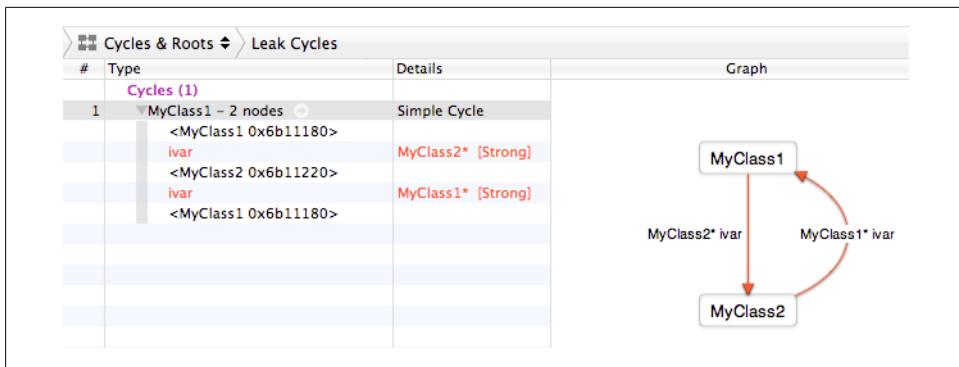


Figure 9-10. Instruments describes a leak

```
MyClass1* m1 = [MyClass1 new];
MyClass2* m2 = [MyClass2 new];
m1.ivar = m2;
m2.ivar = m1;
```

There are steps I could have taken to prevent this from being a memory leak, as I'll explain in [Chapter 12](#); but I haven't taken those steps, so it *is* a memory leak. I'll set the destination to the Simulator and choose Product → Profile; Instruments launches, and I'll choose the Leaks trace template and click Profile. My app starts running in the Simulator, and after about 10 seconds (the default interval at which Instruments runs its leak analysis), a leak is detected. I'll pause Instruments and, with some appropriate button-pushing, am actually shown a diagram of the mistake that's causing this leak ([Figure 9-10](#))!

In this final example, I'm concerned with whether I might be able to speed up the launch time of my Albumen app. I'll set the destination to a device, because that's where speed matters and needs to be measured, and choose Product → Profile; Instruments launches, and I'll choose the Time Profiler trace template and click Profile. I'll watch the device until Albumen has brought up its initial interface; then I'll pause Instruments. To get a better look at the resulting graph of CPU activity over time, I'll choose View → Increase Deck Size (twice) and View → Snap Track to Fit.

The result is [Figure 9-11](#), which is not of itself terribly helpful. However, it is evident from the graph that CPU usage during the app launch process is divided into three distinct phases. For reasons that I won't explain here, I'll focus on the third phase. I'll Option-drag in the graph to select the third phase, and alter some of the checkboxes at the left of the window.

Now we're getting somewhere ([Figure 9-12](#)). Of the time spent in calls within my code, nearly half is spent my RootViewController class's `tableView:cellForRowAtIndexPath:`. If I double-click that line in the table, I'm shown my own code, time-profiled ([Figure 9-13](#)).

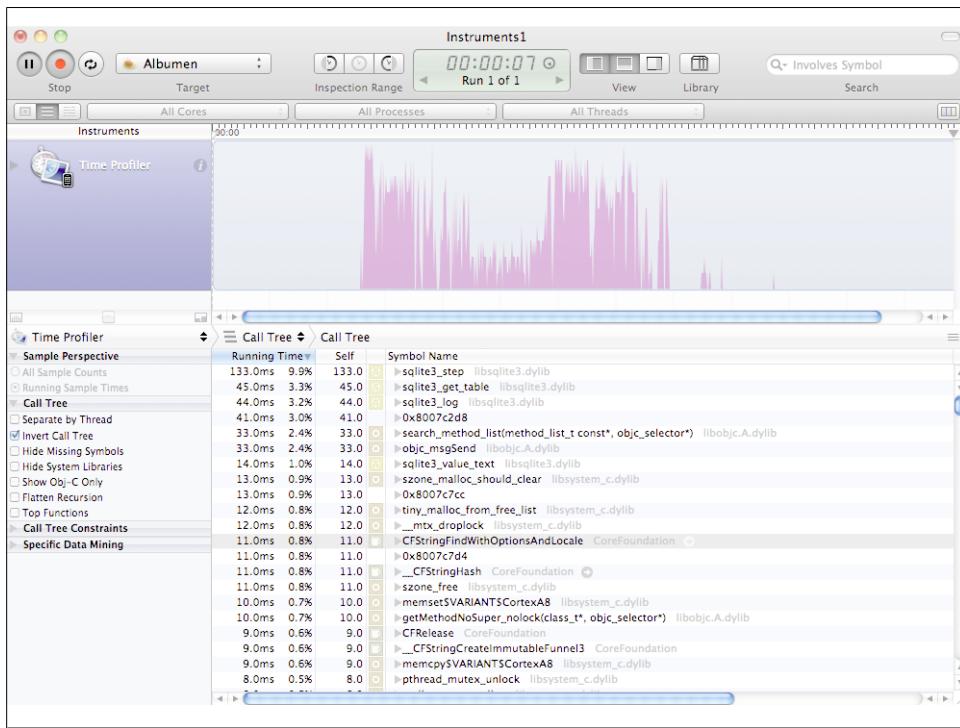


Figure 9-11. A time profile in Instruments

This is really useful information. For example, this line of code turns out to be relatively expensive:

```
cell.textLabel.font = [UIFont fontWithName:@"Georgia-Bold" size:15.0];
```

Part of the expense is that that line is being called ten times (because this code runs every time a table row is displayed, and there are initially ten rows visible in the table). But that's quite unnecessary, since this is a constant; it would be sufficient to make this call once, earlier in the app's lifetime, and store the result for subsequent use. A few small changes of this sort can shave about 50% off the time spent in each call to `tableView:cellForRowAtIndexPath:`.

It is a programming proverb that one should not optimize prematurely; equally, one should not spend time optimizing when the resulting savings will make no perceptible difference to the user. So I can't say, without further measurement, analysis, and experimentation, whether these changes would matter, either during app launch or over the lifetime of the app. But then, that's exactly the value of Instruments: for guesswork and impressions, it substitutes actual numbers and facts.

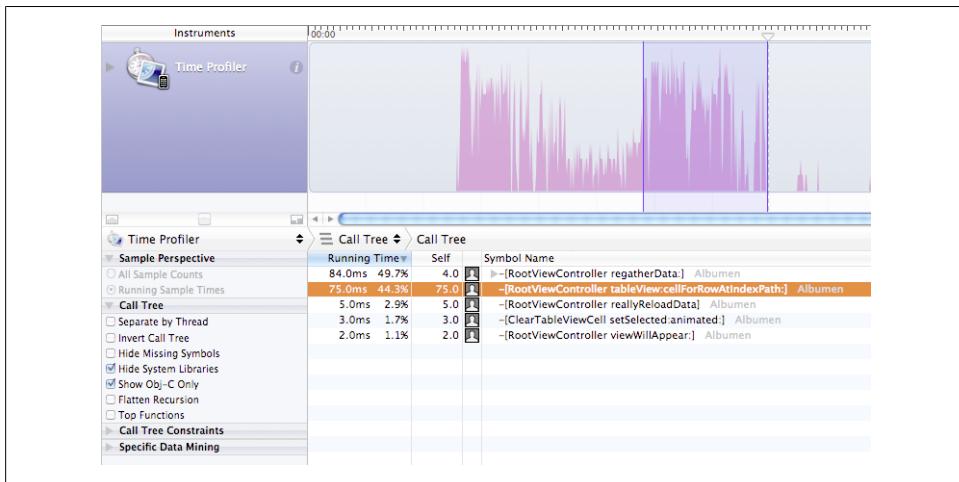


Figure 9-12. The time profile starts to make sense

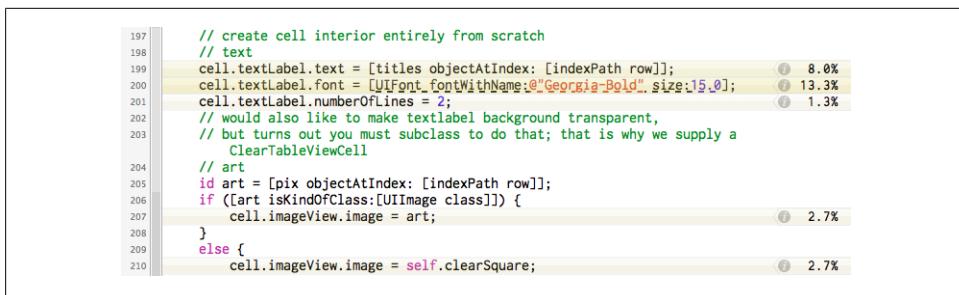


Figure 9-13. My code, time-profiled

Distribution

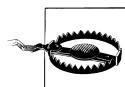
By *distribution* is meant providing your app to others who are not developers on your team. There are two kinds of distribution:

Ad Hoc distribution

You are providing a copy of your app to a limited set of known users so that they can try it on their devices and report bugs, make suggestions, and so forth.

App Store distribution

You are providing the app to the App Store so that anyone can download it (possibly for a fee) and run it.



The Portal imposes a registration limit of 100 devices per year *per developer* (not per app), which limits your number of Ad Hoc testers. Your own devices used for development are counted against this limit.

In order to perform any kind of distribution, you will need a *distribution certificate*, which is different from the development certificate discussed earlier in this chapter. Like the development certificate, you need only one distribution certificate; it identifies you as you. Obtaining a distribution certificate is exactly like obtaining a development certificate, except that, at the iOS Provisioning Portal, under Certificates, you use the Distribution tab instead of the Development tab. (And, like the development certificate, it expires when your year of iOS Developer Program membership expires; if you renew, you'll have to revoke your distribution certificate at the Portal, delete it from your keychain, and obtain a new distribution certificate.)

You will also need a *distribution profile* specifically for this app, which is different from the development profile you obtained earlier. You can't obtain a distribution profile from within Xcode; you must get it at the Portal in your browser. You might need two distribution profiles, because the profile for an Ad Hoc distribution is different from the profile for an App Store distribution. Remember, you will need a separate set of distribution profiles for each app you plan to distribute.

When you build for distribution, you'll use the Product → Archive command. Indeed, you can think of *archive* as meaning “build for distribution.” (Product → Archive isn't enabled unless your destination in the Scheme pop-up menu is a device.) However, you have some preparation to do before you can archive in such a way as to make an app you can actually distribute. If you look at the Archive action in your default scheme, you'll discover that it is set to use the Release distribution configuration. But if you examine the Code Signing Identity build setting for your project, you'll see that by default it uses a development project (most likely the team development profile). This won't do; when you intend to distribute a copy of your app, you want to use a distribution profile. The solution is to create a Distribution build configuration; you can then set the Archive action in your scheme to use the Distribution build configuration, and set the Code Signing Identity build setting to use the distribution profile when the Distribution build configuration is in force.

First, here are the steps for obtaining a distribution profile:

1. To obtain an Ad Hoc distribution profile, collect the unique identifiers of all the devices where this build is to run, and add each of the device identifiers at the Portal under Devices. (For an App Store profile, omit this step.)
2. In the Portal, in the Distribution (not Development!) tab of the Provisioning section, ask for a New Profile. In the New Profile form, ask for an Ad Hoc profile or an App Store profile, depending on which you're after.
3. Describe the profile, giving it a name, and specifying your distribution certificate and this app. For an Ad Hoc profile, also specify all the devices you want the app to run on. Be careful about the profile's name; I suggest that this name should contain both the name of the app and the term “adhoc” or “appstore,” so that you can identify it later easily from within Xcode.

4. Click Submit to generate the profile; you might then have to refresh the browser window to see the Download button next to your new profile. Download the profile and drag it onto Xcode's icon in the Dock. You can now throw the profile away in the Finder; Xcode has kept a copy (which should appear in the Organizer window).

A distribution profile expires together with your annual developer membership. If a distribution profile expires and you want to create a new distribution build, you can renew the profile by returning to the Distribution tab of the Provisioning section at the Portal and clicking Modify for the expired profile. In the next screen, you might need to make some change — any change — merely to bring the Submit button to life. (Apple, Masters of Interface!) Then you can change it back and click Submit.

You can view distribution profiles in the Organizer, in the Devices tab, under Developer Profile or Provisioning Profiles in the Library section.

Now here are the steps to create a separate build configuration for your project:

1. Edit the project. In the Info tab, click the Plus button below the list of configurations, and duplicate the Release configuration. Name the new configuration Distribution.
2. Edit the project (still). In the Build Settings tab, locate the Code Signing Identity entry. The Distribution build setting is now listed here. For the subentry Any iOS Device, set the value of this to a distribution profile *for this app*. I believe it won't matter which distribution profile you choose; the important thing here is that you're specifying a profile that is tied to your distribution certificate. Choosing iPhone Distribution under Automatic Profile Selector should suffice.
3. Edit the scheme, and switch to the Archive action. Change the build configuration to Distribution, and click OK.

Ad Hoc Distribution

To create and distribute an Ad Hoc distribution build, first switch to the iOS Device destination in the Scheme pop-up menu in the project window toolbar. Until you do this, the Product → Archive menu item will be disabled. You do *not* have to have a device connected; you are not building to run on a particular device, but to save an archive.



Apple's docs say that an Ad Hoc distribution build should include an icon that will appear in iTunes, but my experience is that this step is optional. If you want to include this icon, it should be a PNG or JPEG file, 512×512 pixels in size, and its name should be *iTunesArtwork*, with *no file extension*. Make sure the icon is included in the build, being present in the Copy Bundle Resources build phase.

Now choose Product → Archive. The build is created and copied into a date folder within `~/Library/Developer/Xcode/Archives`; it also appears in the Organizer window in the Archives tab. Locate the archive in the Organizer window. You can add a comment here; you can also change the archive’s name (this won’t affect the name of the app).



If Xcode refuses to build because of a Code Sign error, it may help to close the project window and delete this project’s entire build folder from the DerivedData directory, as described earlier in this chapter. When you reopen the project, after Xcode has finished reindexing it, check the build settings for your Distribution configuration; it may help to set once again the value of Any iOS Device under Code Signing Identity for the Distribution configuration once again. See also Apple’s tech note TN2250, “Understanding and Resolving Code Signing Issues.”

Select the archive and press the Share button at the upper right of the window. A dialog appears. Here, you are to specify a Contents type; choose iOS App Store Package (the default). You must also choose an Identity; specify the identity associated with the Ad Hoc distribution profile for this app. (This is the step in which it matters which of the app’s distribution profiles you specify.) Click Next.

After a while, a Save dialog appears. Give the file a useful name (again, this won’t affect the name of the app). Save the file to disk. It will have the suffix `.ipa` (“iPhone app”).

Locate in the Finder the file you just saved. Provide this file to your users with instructions. A user should launch iTunes and drag the `.ipa` file onto the iTunes icon in the Dock. Then the user should connect the device to the computer, make certain the app is present *and checked* in the list of apps for this device, and sync the device to cause the app to be copied to it.

If you listed your own device as one of the devices for which this Ad Hoc distribution profile was to be enabled, you can obey these instructions yourself to make sure the Ad Hoc distribution is working as expected. First, remove from your device any previous copies of this app (such as development copies). Then copy the app onto your device by syncing with iTunes as just described. The app should run on your device, and you should see the Ad Hoc distribution profile on your device (in the Settings app, under General → Profiles). Because you are not privileged over your other Ad Hoc testers, what works for you should work for them.

Final App Preparations

As the day approaches when you’re thinking of submitting your app to the App Store, don’t let the prospect of huge fame or big profits hasten you past the all-important final stages of app preparation. Apple has a lot of requirements for your app, such as icons and launch images, and failure to meet them can cause your app to be rejected. Take

your time. Make a checklist and go through it carefully. See the *iOS Application Programming Guide* for full details.

Xcode 4 makes it easier than in the past for you to fulfill these requirements, by providing interface for doing so. Edit the target, and switch to the Summary tab; there are spaces where for an iPhone app you can drag-and-drop normal and double-resolution icons, and normal and double-resolution launch images, and for an iPad app you can drag-and-drop an iPad icon along with portrait and landscape launch images.

At various stages, you can also obtain validation of your app to confirm that you haven't omitted certain requirements. For example, by default, a new project's Release build configuration has the Validate Build Product build setting set to Yes. Thus, when I do a build of the Empty Window app we've developed in previous chapters, if that build uses the Release build configuration (or the Distribution build configuration duplicated from it), Xcode warns that the app has no icon. When you submit your app to the App Store, it will be subjected to even more rigorous validation.

Icons in the App

An icon file must be a PNG file, without any alpha transparency, with an exact pixel size. It should be a full square, without shading (the "shine" effect that you see in the upper part of icons on your device); the rounding of the corners and shine will be added for you. You can prevent the shine effect from being added to the icon for your App Store build by defining and checking the "Icon already includes gloss and bevel effects" (`UIPrerenderedIcon`) key in your *Info.plist*. Make sure that the icon is copied into the built app by inclusion in the Copy Bundle Resources build phase. The required size is as follows (note that information about icons for apps that are to run on systems before iOS 3.2 is omitted):

- For an iPad app, the icon file should be 72×72 pixels in size.
- For an iPhone app, there should be two primary app icons, one 57×57 pixels, the other 114×114 pixels (for use on the double-resolution Retina display). A double-resolution variant of an icon should have the same name as the single-resolution variant, except for the addition of @2x to its name.

As I mentioned earlier, in Xcode 4 you can drag-and-drop the required icons into the appropriate spaces in the Summary tab when you're editing the target, and Xcode itself will incorporate them into the project and the target and configure the *Info.plist* for you. For example, if I drop a 57×57 PNG file called *myDumbIcon.png* onto the first icon space, it is copied into the project and added to the target, and the first element in the "Icon files" key in my *Info.plist* becomes *myDumbIcon.png*.

Alternatively, you can specify the icon file(s) manually using the *Info.plist*. In this case, you will have to add the relevant key-value pairs manually. To do so, set the "Icon files" key's value (`CFBundleIconFiles`, and notice the plural!); this value is an array, so you can list multiple icons.

You may also optionally include smaller versions of your icon to appear when the user does a search on the device (and in the Settings app, if you include a settings bundle, [Chapter 36](#)). The smaller icon sizes are 29×29 pixels (for an iPhone app, and for the Settings bundle for an iPad app), 50×50 pixels (for an iPad app), and 58×58 pixels (for an iPhone 4 app, on the double-resolution display). List the icons in the “Icon files” key.

The system determines which icon listed under the “Icon files” key to use under what circumstances by examining their sizes. That’s one reason why the sizes must be exactly correct.

For more information, see Apple’s tech note QA1686, “App Icons on iPad and iPhone,” and the “App-Related Resources” chapter of the *iOS App Programming Guide*.

Other Icons

When you submit the app to the App Store, you will be asked to supply a 512×512 PNG, JPEG, or TIFF icon to be displayed at the App Store. Have this icon ready before submission. Apple’s guidelines say that it should not merely be a scaled-up version of your app’s icon, but it must not differ perceptibly from your app’s icon, either, or your app will be rejected (I know this from experience).

The App Store icon does *not* need to be built into your app; indeed, it should not be, as it will merely swell the built app’s size unnecessarily (remember that space is at a premium on a device, and that your app must be downloaded from the App Store, so users appreciate your keeping your app as small as possible). On the other hand, you will probably want to keep it in your project (and in your project folder) so that you can find and maintain it easily. So create it and import it into your project, but do *not* add it to any target.



If you created a 512×512 icon file for Ad Hoc distribution, you may wish to delete it from the Copy Bundle Resources build phase now so that it doesn’t swell the final app’s size unnecessarily.

Launch Images

There may be a delay between the moment when the user taps your app’s icon to launch it and the moment when your app is up and running and displaying its initial window. To cover this delay and give the user the sense that something is happening, you should provide a launch image to be displayed during that interval.

The launch image might be just a blank depiction of the main elements or regions of the app’s interface, so that when the actual window appears, those elements or regions will seem to be filled in. The best way to create such a launch image is to start with a screenshot of your app’s actual initial interface. That way, all you have to do is blank

out the details. You don't need to blank out the status bar area; it will be covered by the real status bar. Taking screenshots is covered in the next section.

For an iPhone app, the launch image should be a PNG image, 320×480 pixels in size. It should be named *Default.png*. Create the launch image, import it into the project, and make sure it is built into the app, being present in the Copy Bundle Resources build phase. For the double-resolution Retina display, provide a second version of the launch image, 640×960 pixels, called *Default@2x.png*.

For an iPad app, you will probably provide at least *two* launch images. Here's why. On the iPhone, you get to dictate the orientation in which the app should launch (landscape or portrait), so you can be certain that the default image matches this. But on the iPad, you're not supposed to do that; your app should be prepared to launch in whatever orientation the device happens to be. Thus, you need a launch image for landscape orientation and a launch image for portrait orientation. The landscape image should be called *Default-Landscape.png*; it should be 748 pixels high and 1024 pixels wide. The portrait image should be called *Default-Portrait.png*; it should be 1004 pixels high and 768 pixels wide. Observe that these sizes omit the status bar area, unlike the iPhone launch image.

You can use the orientation suffixes in the names of launch images on the iPhone as well. To distinguish between a launch image to be used on the iPhone and a launch image with the same orientation suffix to be used on the iPad, use additional suffixes *~ipad* and *~iphone*. Thus you can end up with file names like *Default-Portrait@2x~iphone.png*. To make things even more confusing, you can replace *Default* with some other base name by creating the "Launch image" key in your *Info.plist* file and setting its value appropriately.

As I mentioned earlier, in Xcode 4 you can drag-and-drop the required icons into the appropriate spaces in the Summary tab when you're editing the target. For example, if I drop a 320×480 PNG file called *myLaunchImage.png* onto the first launch image space, it is copied into the project and added to the target, and the copy is renamed *Default.png*.

For further details, see the "App-Related Resources" chapter of the *iOS App Programming Guide*.

Screenshots

When you submit the app to the App Store, you will be asked for one or more screenshots of your app in action to be displayed at the App Store. You should take these screenshots beforehand and be prepared to provide them during the app submission process.

You can obtain screenshots either from the Simulator or from a device connected to the computer. (The Simulator is a particularly good approach when you need a double-

resolution launch image and you don't have a double-resolution device.) Run the app and get it into the desired state. Then:

Simulator

In the Simulator application, choose File > Save Screen Shot.

Device

In Xcode, in the Organizer window, locate your device under Devices and click Screenshots. Click New Screenshot, at the lower right of the window.

To make screenshots available for upload, select each one in the left side of the window and click Export to save it with a nice filename into the Finder. Apple asks that if the status bar is visible in a screenshot, you remove it. And you may need to rotate a screenshot to get it into the correct orientation.

To make a screenshot listed in the Organizer window a launch image (see the previous section), select it in the list and click Save as Launch Image. A dialog will ask you what name to assign to it and what open project to add it to.

Property List Settings

A number of settings in *Info.plist* are crucial to the proper behavior of your app. You should peruse Apple's *Information Property List Key Reference* for full information. Most of the required keys are created as part of the template, and are given reasonable default values, but you should check them anyway. In addition to those already mentioned in [Chapter 6](#), the following are particularly worthy of attention:

Status bar style (UIStatusBarStyle)

On the iPhone and iPod touch, the look of the status bar. (On the iPad, the status bar is always black opaque.) Your choices are "Gray style" (`UIStatusBarStyleDefault`), "Opaque black style" (`UIStatusBarStyleBlackOpaque`), and "Transparent black style" (`UIStatusBarStyleBlackTranslucent`). This setting will be used in conjunction with your launch image, even before the app is actually running. If the status bar is to be hidden initially, set "Status bar is initially hidden" (`UIStatusBarHidden`) instead.

Supported interface orientations (UISupportedInterfaceOrientations)

The initial orientation(s) in which the app is permitted to launch. (The app may support additional orientations later as it runs.) In Xcode 4, you can perform this setting graphically in the Summary tab when editing your target.

Required device capabilities (UIRequiredDeviceCapabilities)

You should set this key if the app requires capabilities that are not present on all devices. Be sure to look over the list of possible values. Don't use this key unless it makes no sense for your app to run *at all* on a device lacking the specified capabilities.

Note that a new Xcode 4.2 project includes `armv7` in the list of required device capabilities. This is a new restriction, and you might like to delete it:

- If you are interested in backwards compatibility, this restriction could prove unnecessarily restrictive, excluding many older devices.
- When I rewrote an existing app from scratch starting with a new Xcode 4.2 application template, my update wouldn't validate at the App Store initially, apparently because you aren't permitted to add to the list of required device capabilities in an update to an existing app.

Property list settings can adopt different values depending on what device type you're running on. To specify that a property list setting applies only on a particular type of device, you add to its key the suffix `~iphone`, `~ipod`, or `~ipad`. This feature is typically useful in a universal app, in which the distinction will be between iPhone and iPod, on the one hand, and iPad on the other. The general setting is used unless there is an applicable specific case, so in a universal app you might have one setting with no suffix and a second setting with the `~ipad` suffix. Thus, for example, you could have a different set of supported initial interface orientations on the iPad by adding settings for "Supported interface orientations (iPad)" (`UISupportedInterfaceOrientations~ipad`).

You app also needs a version number. Edit the Target and click Summary at the top of the page; enter a version string in the Version text field. Things are a little confusing here because there is both a Version field and a Bundle field; the former corresponds to "Bundle versions string, short" (`CFBundleShortVersionString`), while the latter corresponds to "Bundle version" (`CFBundleVersion`). As far as I can determine, Apple will pay attention to the former if it is set, and otherwise will fall back on the latter. In general I play it safe and set both to the same value when submitting to the App Store. The value needs to be a version string, such as "1.0". This version number will appear at the App Store, and you should increment it when you develop and submit an update to an existing app. Failure to increment the version string when submitting an update will cause the update to be rejected.

Submission to the App Store

When you're satisfied that your app works well, and you've installed or collected all the necessary resources, you're ready to submit your app to the App Store for distribution. The primary way to submit your app is through a website called iTunes Connect. You can find a link to it on the iOS developer pages when you've logged in at Apple's site. You can go directly to <http://itunesconnect.apple.com>, but you'll still need to log in with your iOS Developer username and password.

The first thing you should do at iTunes Connect is download the *iTunes Connect Developer Guide*. It's a PDF that gives you a good idea what to expect when you submit your app, as well as later when you return for financial and other reports, and when you update your app.

You should also go to the Contracts section at iTunes Connect and complete submission of your contract if you haven't already done so. You can't offer any apps for sale until you do, and even free apps require completion of a contractual form.

When you submit an app to iTunes Connect, you will have to supply a description of fewer than 4,000 characters; Apple recommends fewer than 580 characters, and the first paragraph is the most important, because this may be all that users see at first when they visit the App Store. It must be pure text, without HTML and without character styling.

You will also be asked for a list of keywords: a comma-separated list shorter than 100 characters. These keywords will be used, in addition to your app's name, to help users discover your app through the Search feature of the App Store.

iTunes Connect will also expect you to provide a website where users can find more information about your app; it's good to have that ready in advance.

Now build the app. The procedure is exactly as for an Ad Hoc build. Set the destination to iOS Device in the Scheme pop-up menu in the project window toolbar, and choose Product → Archive.

The archived build that appears in the Organizer window can be used to generate either an Ad Hoc build or an App Store build. You can't test an App Store build. So if you want to test one last time, use this archived build to generate an Ad Hoc build and test with that. When you generate the App Store build, you use the exact same binary, so you are guaranteed that its behavior will be exactly the same as the build you tested. (That is one of the purposes of archiving.)

Enter your app's information at the iTunes Connect website. I'm not going to recite all the steps you have to go through, as these are described thoroughly in the *iTunes Connect Developer Guide*. But I'll just mention a few possible pitfalls:

Your app's name

This is the name that will appear at the App Store; it need not be identical to the short name that will appear under the app's icon on the device, dictated by the Bundle Display Name setting in your *Info.plist* file. This name can be up to 255 characters long, though Apple recommends that you limit it to fewer than 70 and ideally to fewer than 35 characters. You can get a rude shock when you submit your app's information to iTunes Connect and discover that the name you wanted is already taken. There is no reliable way to learn this in advance, and such a discovery can require a certain amount of scrambling on your part: you might have to Build and Archive your app yet again with a new name and possibly other last-minute changes.

Copyright

Do not include a copyright symbol in this string; it will be added for you at the App Store.

SKU number

This is unimportant, so don't get nervous about it. It's just a unique identifier, unique within the world of your own apps. It's convenient if it has something to do with your app's name. It needn't be a number; it can actually be any string.

Price

You don't get to make up a price. You have to choose from a list of pricing "tiers."

Availability Date

Apple suggests that you set the availability date to the date you submit the app; that way, when it is approved, it will be available immediately.

When you've submitted the information for your app, you can do a final validation check: return to the Organizer window, select the archived build, and click Validate. (This feature has not worked well for me in the past, however.)

Finally, when you're ready to upload the app for which you've already submitted the information at iTunes Connect, and when the iTunes Connect status for your app is "Waiting for Upload," you can perform the upload using Xcode. Select the archived build in the Organizer and click Submit, specifying the App Store distribution profile.

Alternatively, you can use Application Loader, an application located in */Developer/Applications/Utilities/*, to upload the app. Application Loader first checks with iTunes Connect to see what apps are awaiting upload; pick the right one. In the Organizer, click Share and specify the App Store distribution profile (if you're asked to specify a distribution profile). The result will be an *.ipa* file in the Finder; this is what you will hand to Application Loader. Application Loader will perform many useful validation checks as it communicates with iTunes Connect and uploads your app.

You will subsequently receive emails from Apple informing you of your app's status as it passes through various stages: "Waiting For Review," "In Review," and finally, if all has gone well, "Ready For Sale" (even if it's a free app). Your app will then appear at the App Store.

PART III

Cocoa

When you program for iOS, you take advantage of a suite of frameworks provided by Apple. These frameworks, taken together, constitute *Cocoa*; the brand of Cocoa that provides the API for programming iOS is *Cocoa Touch*. Cocoa thus plays an important and fundamental role in iOS programming; your code will ultimately be almost entirely about communicating with Cocoa — interacting with the frameworks provided by Apple, in order to make an app that does what you want it to do.

The Cocoa Touch frameworks are a huge boon to you, the programmer, because they provide the underlying functionality that any iOS app needs to have. Your app can put up a window, show the interface containing a button, respond to that button being tapped by the user, and so forth, because Cocoa knows how to do those things. But with the great advantages of working with a framework come great responsibilities. You have to think the way the framework thinks, put your code where the framework expects it, and fulfill many obligations imposed on you by the framework.

- [Chapter 10](#) picks up where [Chapter 5](#) left off, describing some Objective-C linguistic features used by Cocoa, such as categories and protocols; it also surveys some important fundamental classes.
- An *event* is a message sent by Cocoa to your code. Cocoa is event-based; if Cocoa doesn't send your code an event, your code doesn't run. Getting your code to run at the appropriate moment is all about knowing what events you can expect Cocoa to send you and when. [Chapter 11](#) describes Cocoa's event-driven model, along with its major design patterns.
- [Chapter 12](#) describes your responsibilities for making your instances nicely encapsulated and good memory-management citizens in the world of Cocoa objects.
- [Chapter 13](#) surveys some answers to the question of how your objects are going to see and communicate with one another within the Cocoa-based world.

Cocoa Classes

Using the Cocoa frameworks requires an understanding of how those frameworks organize their classes. Cocoa class organization depends upon certain Objective-C language features that are introduced in this chapter. The chapter also surveys some commonly used Cocoa utility classes, along with a discussion of the Cocoa root class.

Subclassing

Cocoa effectively hands you a large repertory of objects that already know how to behave in certain desirable ways. A `UIButton`, for example, knows how to draw itself and how to respond when the user taps it; a `UITextField` knows how to summon the keyboard when the user taps in it, how to accept keyboard input, and how to respond when the user finishes inputting text.

Often, the default behavior or appearance of an object supplied by Cocoa won't be quite what you're after, and you'll want to customize it. Cocoa classes are heavily endowed with methods and properties for precisely this purpose, and these will be your first resort. Always study the documentation for a Cocoa class to see whether instances can already be made to do what you want. For example, the class documentation for `UILabel` ([Chapter 27](#)) shows that you can set the font, size, color, line-breaking behavior, and horizontal alignment of its text, among other things.

Nevertheless, sometimes setting properties and calling methods won't suffice to customize an instance the way you want to. In such cases, Cocoa may provide methods that are called internally as an instance does its thing, and whose behavior you can customize by subclassing and overriding. You don't have the code to any of Cocoa's built-in classes, but you can still subclass them, creating a new class that acts just like a built-in class except for the modifications you provide.

Oddly enough, however (and you might be particularly surprised by this if you've used another object-oriented application framework), subclassing is probably one of the less important ways in which your code will relate to Cocoa. Knowing or deciding when to

subclass can be somewhat tricky, but the general rule is that you probably shouldn't subclass unless you're invited to.

A common case involves custom drawing into a UIView. You don't actually draw *into* a UIView; rather, when a UIView needs drawing, its `drawRect:` method is called so that the view can draw itself. So the way to make a UIView that is drawn in some completely custom manner is to subclass UIView and implement `drawRect:` in the subclass. As the documentation says, "Subclasses override this method if they actually draw their views." That's a pretty strong hint that you *need* to subclass UIView in order to do custom drawing into a UIView.

For example, suppose we want our window to contain a horizontal line. There is no horizontal line interface widget, so we'll just have to roll our own — a UIView that draws itself as a horizontal line. Let's try it. First we'll code the class:

1. In our Empty Window example project, choose File → New → New File and specify a Cocoa Touch Objective-C class, and in particular a subclass of UIView. Call it `MyHorizLine`. Xcode creates `MyHorizLine.m` and `MyHorizLine.h`.
2. In `MyHorizLine.m`, replace the contents of the implementation section with this (without further explanation; you'll know all about this after you read [Chapter 15](#)):

```
- (id)initWithCoder:(NSCoder *)decoder {
    self = [super initWithCoder:decoder];
    if (self) {
        self.backgroundColor = [UIColor clearColor];
    }
    return self;
}

- (void)drawRect:(CGRect)rect {
    CGContextRef c = UIGraphicsGetCurrentContext();
    CGContextMoveToPoint(c, 0, 0);
    CGContextAddLineToPoint(c, self.bounds.size.width, 0);
    CGContextStrokePath(c);
}
```

3. Edit `ViewController.xib`. Find UIView in the Object library, and drag it into the View object in the canvas. You may resize it to be less tall.
4. With the UIView that you just dragged into the canvas still selected, use the Identity inspector to change its class to `MyHorizLine`.

Build and run the app in the Simulator. You'll see a horizontal line corresponding to the location of the top of the `MyHorizLine` instance in the window. Our view has drawn itself as a horizontal line, because we subclassed it to do so.

In that example, we started with a bare UIView that had no drawing functionality of its own. (That's why there was no need to call `super`; the default implementation of UIView's `drawRect:` does nothing.) But you might also be able to subclass a built-in UIView subclass to modify the way it already draws itself. Again using UILabel as an example, the documentation shows that two methods are present for exactly this pur-

pose. Both `drawTextInRect:` and `textRectForBounds:limitedToNumberOfLines:` explicitly tell us: “You should not call this method directly. This method should only be overridden by subclasses.” The implication is that these are methods that will be called for us, automatically, by Cocoa, as a label draws itself; thus, we can subclass `UILabel` and implement them in our subclass to modify how a particular type of label draws itself.

Here’s an example from one of my own apps, in which I subclass `UILabel` to make a label that draws its own rectangular border and has its content inset somewhat from that border, by overriding `drawTextInRect:`. As the documentation tells us: “In your overridden method, you can configure the current [graphics] context further and then invoke `super` to do the actual drawing [of the text].” Let’s try it:

1. In the Empty Window project, make a new class file, a `UILabel` subclass this time; call it `MyBoundedLabel`.
2. In `MyBoundedLabel.m`, insert this code into the implementation section:

```
- (void)drawTextInRect:(CGRect)rect {
    CGContextRef context = UIGraphicsGetCurrentContext();
    CGContextStrokeRect(context, CGRectMakeInset(self.bounds, 1.0, 1.0));
    [super drawTextInRect:CGRectMakeInset(rect, 5.0, 5.0)];
}
```
3. In `MyNib.xib`, select the `UILabel` and change its class in the Identity inspector to `MyBoundedLabel`.

Build and run the app, and you’ll see how the rectangle is drawn and the label’s text is inset within it.

Similarly, in a table view (a `UITableView`) you might very well be able to avoid subclassing the table view cell (`UITableViewCell`), because it provides so many properties through which you can customize its appearance. If you want text to appear in the cell using a certain font, the built-in cell styles and the ability to access and modify the cell’s labels might be quite sufficient. You can directly replace a cell’s background or put a checkmark at the right end of the cell. All of that is simply a matter of setting the cell’s built-in properties. But if you want a table view cell that doesn’t look or behave like any of the built-in cell styles, then you’ll have to subclass `UITableViewCell`. (We’ll go deeply into this in [Chapter 21](#).)

You wouldn’t subclass `UIApplication` (the class of the singleton shared application instance) just in order to respond when the application has finished launching, because the delegate mechanism ([Chapter 11](#)) provides a way to do that (`application:didFinishLaunchingWithOptions:`). On the other hand, if you need to perform certain tricky customizations of your app’s fundamental event messaging behavior, you’d have to subclass `UIApplication` in order to override `sendEvent:`. The documentation does tell you this, and it also tells you, rightly, that needing to do this would be fairly rare (though I have had occasion to do it).



If you do subclass `UIApplication`, you'll need to change the third argument in the call to `UIApplicationMain` in `main.m` from `nil` to the `NSString` name of your subclass. Otherwise your `UIApplication` subclass won't be instantiated as the shared application instance.

Another class that's commonly subclassed is `UIViewController` ([Chapter 19](#)). And any class you write will naturally need to be a subclass of `NSObject`, if nothing else. You definitely want your class to inherit all of `NSObject`'s yummy goodness, including `alloc` and `init`, which make it possible to instantiate your class in the first place. For more information, see "The Secret Life of `NSObject`," later in this chapter.

Categories

A *category* is an Objective-C language feature that allows you to reach right into an existing class and define additional methods. You can do this even if you don't have the code for the class, as with Cocoa's classes. Your instance methods can refer to `self`, and this will mean the instance to which the message was originally sent, as usual. A category, unlike a subclass, cannot define additional instance variables; it can override methods, but you should probably not take advantage of this ability.

Defining a category is just like defining the class on which the category is being defined: you need an interface section and an implementation section, and you'll typically distribute them into the standard `.h` and `.m` class file pair. At the start of both the interface section and the implementation section, where you give the class's name, you add a category name in parentheses. The `.h` file will probably need to import the header for the original class (or the header of the framework that defines it), and the `.m` file will, as usual, import the corresponding header file.

For example, in one of my apps I found myself performing a bunch of string transformations in order to derive the path to various resource files inside the app bundle based on the resource's name and purpose. I ended up with half a dozen utility methods. Given that these methods all operated on an `NSString`, it was appropriate to implement them as a category of `NSString`, thus allowing *any* `NSString`, anywhere in my code, to respond to them.

The code was structured like this (I'll show just one of the methods):

```
// [StringCategories.h]
#import <Foundation/Foundation.h>

@interface NSString (MyStringCategories)
- (NSString*) basePictureName;
@end

// [StringCategories.m]
#import "StringCategories.h"
```

```

@implementation NSString (MyStringCategories)
- (NSString*) basePictureName {
    return [self stringByAppendingString:@"IO"];
}
@end

```

If we had written a utility method within some other class, we'd have to pass an `NSString` to the method that operates on it (and if this were an instance method, we'd need to go to the extra trouble of obtaining a reference to an instance of that class). A category is neater and more compact. We've extended `NSString` *itself* to have `basePictureName` as an instance method, so we can send the `basePictureName` message directly to any `NSString` we want to transform:

```
NSString* aName = [someString basePictureName];
```

A category is particularly appropriate in the case of a class like `NSString`, because the documentation warns us that subclassing `NSString` is a bad idea. That's because `NSString` is part of a complex of classes called a *class cluster*, which means that an `NSString` object's real class might actually be some other class. A category is a much better way to modify a class within a class cluster than subclassing.

Keep in mind, too, that a method defined through a category can equally be a class method. Thus you can inject utility methods into any appropriate class and call those methods without the overhead of instantiating anything at all.

Splitting a Class

A category can be used to split a class over multiple `.h/.m` file pairs. If a class becomes long and unwieldy, yet it clearly needs to be a single class, you can define the basic part of it (including instance variables) in one file pair, and then add another file pair defining a category on your own class to provide further methods.

Cocoa itself does this. A good example is `NSString`. `NSString` is defined as part of the Foundation framework, and its basic methods are declared in `NSString.h`. Here we find that `NSString` itself, with no category, has just two methods, `length` and `characterAtIndex:`, because these are regarded as the minimum that a string needs to do in order to be a string. Additional methods (those that create a string, deal with a string's encoding, split a string, search in a string, and so on) are clumped into a category. A string may serve as a file pathname, so we also find a category on `NSString` in `NSPathUtilities.h`, where methods are declared for splitting a pathname string into its constituents and the like. Then, in `NSURL.h`, there's another `NSString` category, declaring a couple of methods for dealing with percent-escaping in a URL string. Finally, off in a completely different framework (UIKit), `UIStringDrawing.h` adds yet another `NSString` category, with methods for drawing a string in a graphics context.

This organization won't matter to you as a programmer, because an `NSString` is an `NSString`, no matter how it acquires its methods, but it can matter when you consult the documentation. The `NSString` methods declared in `NSString.h`, `NSPathUtilities.h`,

and `NSURL.h` are documented in the `NSString` class documentation page, but the `NSString` methods declared in `UIStringDrawing.h` are not, presumably because they originate in a different framework. Instead, they appear in a separate document, *NSString UIKit Additions Reference*. As a result, the string drawing methods can be difficult to discover, especially as the `NSString` class documentation doesn't link to the other document. I regard this as a major flaw in the structure of the Cocoa documentation. A third-party utility such as AppKiDo can be helpful here.

Private Method Declarations

A problem (already mentioned in [Chapter 4](#)) arises when you'd like to declare a method in such a way that all other methods in the same class can see the declaration (and can thus call the method) without putting that declaration in the class's interface section where any *other* class that imports the header file will *also* be able to see and call it. The solution is to put an interface section for a category on your own class *in the implementation file*, which no one imports ([Example 10-1](#)).

Example 10-1. Declaring a method privately

```
// [in MyClass.m]
#import "MyClass.h"

@interface MyClass (Tricky)
- (void) myMethod;
@end

@implementation MyClass
// all methods here can call myMethod
@end
```

This trick cannot completely prevent some other class from calling this class's `myMethod` — Objective-C is too dynamic for that — but at least a normal call to `myMethod` from some other class will get its hand slapped by the compiler.

In [Example 10-1](#), the compiler will not warn if the methods declared in the category interface section (such as `myMethod`) are not defined in the implementation section. If this worries you, there are two solutions. One is to provide a named category implementation section, corresponding to the named category interface section; if you fail to implement the category-declared methods in this category implementation section, the compiler *will* warn:

```
@implementation MyClass (Tricky)
// must implement myMethod here, or compiler will warn
@end
```

The other approach is just the opposite, namely to remove the category name altogether from the category implementation section; if you then fail to implement the category-declared methods in the *normal* implementation section, the compiler will warn:

```
@interface MyClass ()  
- (void) myMethod;  
@end
```

This nameless type of category, which exists only as an interface section (corresponding to the normal implementation section), is called a *class extension*. This is the approach that I prefer; in fact, use of a class extension to declare methods privately is such a common and valuable thing to do that I routinely insert a class extension immediately at the start of the code for any class that I create. (A class extension has another remarkable feature: you can declare instance variables and properties in a class extension interface section, which you can't do in a named category. I'll discuss a use of that feature in [Chapter 12](#), when we talk about Objective-C properties and the `@synthesize` directive.)

Protocols

Every reasonably sophisticated object-oriented language must face the fact that the hierarchy of subclasses and superclasses is insufficient to express the desired relationships between classes. For example, a Bee object and a Bird object might need to have certain features in common by virtue of the fact that both a bee and a bird can fly. But Bee might inherit from Insect, and not every insect can fly, so how can Bee acquire the aspects of a Flier in a way that isn't completely independent of how Bird acquires them?

Some object-oriented languages solve this problem through *mixin* classes. For example, in Ruby you could define a Flier module, complete with method definitions, and incorporate it into both Bee and Bird. Objective-C uses a simpler, lighter-weight approach — the *protocol*. Cocoa makes heavy use of protocols.

A protocol is just a named list of method declarations, with no implementation. A class may formally declare that it *conforms* to (or *adopts*) a protocol; such conformance is inherited by subclasses. This declaration satisfies the compiler when you try to send a corresponding message. If a protocol declares an instance method `myCoolMethod`, and if `MyClass` declares conformance to that protocol, then you can send the `myCoolMethod` message to a `MyClass` instance and the compiler won't complain.

Actually implementing the methods declared in a protocol is up to the class that conforms to it. A protocol method may be required or optional. If a protocol method is required, then if a class conforms to that protocol, the compiler will warn if that class fails to implement that method. Implementing optional methods, on the other hand, is optional. (Of course, that's just the compiler's point of view; at runtime, if a message is sent to an object with no implementation for the corresponding method, a crash can result.)

Here's an example of how Cocoa uses a protocol. Some objects can be copied; some can't. This has nothing to do with an object's class heritage. Yet we would like a uniform method to which any object that *can* be copied will respond. So Cocoa defines a pro-

tocol named NSCopying, which declares just one method, `copyWithZone:` (required). A class that explicitly conforms to NSCopying is promising that it implements `copyWithZone:`.

Here's how the NSCopying protocol is defined (in `NSObject.h`, where your code can see it):

```
@protocol NSCopying
- (id)copyWithZone:(NSZone *)zone;
@end
```

That's all there is to defining a protocol. The definition uses the `@protocol` compiler directive; it states the name of the protocol; it consists entirely of method declarations; and it is terminated by the `@end` compiler directive. A protocol definition will typically appear in a header file, so that classes that need to know about it, in order to call its methods, can import it. A `@protocol` section of a header file is not inside any other section (such as an `@interface` section). Any optional methods must be preceded by the `@optional` directive. A protocol definition may state that the protocol incorporates other protocols; these constitute a comma-separated list in angle brackets after the protocol's name, like this example from Apple's own code (`UIAlertView.h`):

```
@protocol UIAlertViewDelegate <NSObject>
@optional
// optional method declarations go here
@end
```

The NSCopying protocol definition in `NSObject.h` is just a definition; it is not a statement that `NSObject` conforms to NSCopying. Indeed, `NSObject` does *not* conform to NSCopying. To see this, try sending the `copyWithZone:` method to your own subclass of `NSObject`:

```
MyClass* mc = [[MyClass alloc] init];
MyClass* mc2 = [mc copyWithZone: [mc zone]];
```

The compiler warns that a `MyClass` instance may not respond to `copyWithZone:`; under ARC, this code won't compile at all.

To conform formally to a protocol, a class's `@interface` section appends the name of the protocol, in angle brackets, after the name of the superclass (or, if this is a category declaration, after the parentheses naming the category). This will necessitate importing the header file that declares the protocol (or some header file that imports that header file). To state that a class conforms to multiple protocols, put multiple protocol names in the angle brackets, separated by comma.

Let's see what happens if you conform formally to the NSCopying protocol. Modify the first line of the `@interface` section of your class as follows:

```
@interface MyClass : NSObject <NSCopying>
```

Now the compiler warns that `MyClass` fails to implement `copyWithZone:` and thus does not fully implement the NSCopying protocol (because `copyWithZone:` is a required method of the NSCopying protocol).

The name of a protocol may also be used when specifying an object type. Most often, the object will be typed as an `id`, but with the accompanying proviso that it conforms to a protocol, whose name appears in angle brackets.

To illustrate, let's look at another typical example of how Cocoa uses protocols, namely in connection with a table (`UITableView`). A `UITableView` has a `dataSource` property, declared like this:

```
@property (nonatomic, assign) id<UITableViewDataSource> dataSource
```

This property represents an instance variable whose type is `id <UITableViewDataSource>`. This means "I don't care what class my data source belongs to, but whatever it is, it should conform to the `UITableViewDataSource` protocol." Such conformance constitutes a promise that the data source will implement at least the required instance methods `tableView:numberOfRowsInSection:` and `tableView:cellForRowAtIndexPath:`, which the table view will call when it needs to know what data to display.

If you attempt to set a table view's `dataSource` property to an object that does *not* conform to `UITableViewDataSource`, you'll get a warning from the compiler. So, for example:

```
MyClass* mc = [[MyClass alloc] init];
UITableView* tv = [[UITableView alloc] init];
tv.dataSource = mc; // compiler warns
```

Under ARC, this warning is couched in rather confusing terms, along these lines: "Passing 'MyClass *const __strong' to parameter of incompatible type 'id<UITableViewDataSource>'."

To quiet the compiler, `MyClass`'s declaration should state that it conforms to `UITableViewDataSource`. Now `MyClass` *is* an `id <UITableViewDataSource>`, and the third line no longer generates a warning. Of course, you must also supply implementations of `tableView:numberOfRowsInSection:` and `tableView:cellForRowAtIndexPath:` in `MyClass` to avoid the other warning, namely that you're not fully implementing a protocol you've claimed to conform to.

A prevalent use of protocols in Cocoa is in connection with delegate objects. We'll talk in detail about delegates in [Chapter 11](#), but you can readily see that many classes have a `delegate` property and that the class of this property is often `id <SomeProtocol>`. For example, in our Empty Window project, the `AppDelegate` class provided by the project template is declared like this:

```
@interface AppDelegate : UIResponder <UIApplicationDelegate>
```

The reason is that `AppDelegate`'s purpose on earth is to serve as the shared application's delegate. The shared application object is a `UIApplication`, and a `UIApplication`'s `delegate` property is typed as an `id <UIApplicationDelegate>`. So `AppDelegate` announces its role by explicitly conforming to `UIApplicationDelegate`.

As a programmer, Cocoa's use of protocols will matter to you in two ways. First, when an object value that you wish to supply is typed as `id <SomeProtocol>`, you will need to

make sure that that object's class does indeed conform to SomeProtocol (and implements any methods required by that protocol).

Second, you must understand about protocols in order to use the documentation. A protocol has its own documentation page. When the UIApplication class documentation tells you that a UIApplication's `delegate` property is typed as an `id <UIApplicationDelegateDelegate>`, it's implicitly telling you that if you want to know what messages a UIApplication's delegate might receive, you need to look in the UIApplicationDelegate protocol documentation.

Similarly, when a class's documentation mentions that the class conforms to a protocol, don't forget to examine that protocol's documentation, because the latter might contain important information about how the class behaves. To learn what messages can be sent to an object, you need to look upward through the inheritance chain (the superclass); you also need to look at any protocols that this object's class conforms to.

Might you ever have cause to define a protocol yourself? Unless you're writing a framework, it probably wouldn't be necessary; but protocols can make your code neater and cleaner, because they effectively allow one class to declare a method that another class is to implement, which is sometimes appropriate architecturally.

For example, in one of my apps I present a view where the user can move three sliders to choose a color. Appropriately, its code is in a class called ColorPickerController. When the user taps Done or Cancel, the view should be dismissed; but first, the code that presented the view needs to hear about what color the user chose. So I need to send a message from the ColorPickerController instance back to the instance that presented it. Here is the declaration for that message:

```
- (void) colorPicker:(ColorPickerController*)picker  
didSetColorNamed:(NSString*)theName  
toColor:(UIColor*)theColor;
```

The question is: where should this declaration go? Now, it happens that in my app I know the class of the instance that will present the ColorPickerController's view: it is a SettingsController. So I could simply declare this method in the interface section of SettingsController's header file. But this feels wrong, for three reasons:

- It should not be up to SettingsController to declare a method that it is implementing only in deference to ColorPickerController.
- If SettingsController declares this message in its header file, ColorPickerController will have to import that header file in order to send the message. But this means that ColorPickerController now knows all about SettingsController, whereas the *only* thing it needs to know about SettingsController is that it implements this one method.
- It is merely a contingent fact that the instance being sent this message *is* a SettingsController; it should be open to *any* class to present and dismiss a ColorPick-

Informal Protocols

You may occasionally see, online or in the documentation, a reference to an *informal protocol*. An informal protocol isn't really a protocol at all; it's just a way of providing the compiler with a method signature so that it will allow a message to be sent without complaining. There are two complementary ways to implement an informal protocol. One is to define a category on NSObject; this makes any object eligible to receive the messages listed in the category. The other is to define a protocol without formally conforming to it; instead, send any message listed in the protocol only to objects typed as `id`, thus suppressing any possible objections from the compiler. These techniques were widespread before protocols could declare methods as optional; now they are largely unnecessary. (They are still used, but decreasingly so; in iOS 5 very few informal protocols remain.) They are also mildly dangerous, because you might accidentally define a method with the same name as an existing method but a different signature, with unpredictable results.

erController's view, and thus to be eligible to receive the `colorPickerController:didSetColor...` message.

Therefore we want ColorPickerController *itself* to declare the method that it is going to send, and we want it to send the message blindly to some receiver, without regard to the class of that receiver. Thus there needs to be a linkage, as it were, between the declaration of this method in ColorPickerController and the implementation of this method in the receiver. That linkage is precisely what a protocol creates! The solution, therefore, is for ColorPickerController to define a protocol in its header file, with this method as part of that protocol, and for the class that presents a ColorPickerController's view to conform to that protocol.

If you look at the project created by Xcode's own Utility Application project template, you will see that this is exactly the architecture it exemplifies. We start with a MainViewController. It eventually creates a FlipsideViewController. When the FlipsideViewController is ready to go out of existence, it is going to want to send the `flipsideViewControllerDidFinish:` message back to whoever created it. So FlipsideViewController defines a FlipsideViewControllerDelegate protocol requiring the `flipsideViewControllerDidFinish:` method, along with a `delegate` property typed as `id <FlipsideViewControllerDelegate>`. When a MainViewController instance creates a FlipsideViewController instance, it specifies that it itself, the MainViewController instance, is the FlipsideViewController's `delegate`; and it can do this, because MainViewController does in fact adopt the FlipsideViewControllerDelegate protocol! Problem solved, mission accomplished.

Optional Methods

The careful reader may have noticed that earlier sections of this chapter have listed two ways in which a method can be publicly declared without necessarily being implemented, and without the compiler complaining if it isn't:

- By defining a named category interface section with no corresponding named category implementation section.
- By defining a protocol in which some methods are explicitly designated as optional.

The question thus arises: How, in practice, is such an optional method feasible? We know that if a message is sent to an object and the object can't handle that message, an exception is raised (and your app will likely crash). But a method declaration is a contract suggesting that the object *can* handle that message. If we subvert that contract by declaring a method that might or might not be implemented, aren't we inviting crashes?

The answer is that Objective-C is not only dynamic but also introspective. You can ask an object whether it can deal with a message without actually sending it that message. This makes optional methods quite safe, provided you know that a method is optional.

The key method here is `NSObject`respondsToSelector:`, which takes a selector parameter and returns a `BOOL`. With it, you can send a message to an object only if it would be safe to do so:

```
MyClass* mc = [[MyClass alloc] init];
if ([mc respondsToSelector:@selector(woohoo)]) {
    [mc woohoo];
}
```

You wouldn't want to do this before sending just any old message, because it isn't necessary except for optional methods, and it slows things down a little. But Cocoa does in fact call `respondsToSelector:` on your objects as a matter of course. To see that this is true, implement `respondsToSelector:` on `AppDelegate` in our Empty Window project in such a way as to instrument it with logging:

```
- (BOOL) respondsToSelector: (SEL) sel {
    NSLog(@"%@", NSStringFromSelector(sel));
    return [super respondsToSelector:(sel)];
}
```

Here's the output on my machine, as the Empty Window app launches:

```
application:handleOpenURL:
application:openURL:sourceApplication:annotation:
applicationDidReceiveMemoryWarning:
applicationWillTerminate:
applicationSignificantTimeChange:
application:willChangeStatusBarOrientation:duration:
application:didChangeStatusBarOrientation:
```

```
application:willChangeStatusBarFrame:  
application:didChangeStatusBarFrame:  
application:deviceAccelerated:  
application:deviceChangedOrientation:  
applicationDidBecomeActive:  
applicationWillResignActive:  
applicationDidEnterBackground:  
applicationWillEnterForeground:  
applicationWillSuspend:  
application:didResumeWithOptions:  
application: didFinishLaunchingWithOptions:
```

That's Cocoa, checking to see which of the optional `UIApplicationDelegate` protocol methods (including a couple of undocumented methods) are actually implemented by our `AppDelegate` instance — which, because it is the `UIApplication` object's delegate and formally conforms to the `UIApplicationDelegate` protocol, has explicitly agreed that it *might* be willing to respond to any of those messages. The entire delegate pattern ([Chapter 11](#)) depends upon this technique. Observe the policy followed here by Cocoa: it checks all the optional protocol methods once, when it first meets the object in question, and presumably stores the results; thus, the app is slowed a tiny bit by this one-time initial bombardment of `respondsToSelector:` calls, but now Cocoa knows all the answers and won't have to perform any of these same checks on the same object later on.

Some Foundation Classes

The Foundation classes of Cocoa provide basic data types and utilities that will form the basis of much that you do in Cocoa. Obviously I can't list all of them, let alone describe them fully, but I can survey a few that I use frequently and that you'll probably want to look into before writing even the simplest Cocoa program. For more information, start with Apple's list of the Foundation classes in the *Foundation Framework Reference*.

Useful Structs and Constants

`NSRange` is a struct of importance in dealing with some of the classes I'm about to discuss. Its components are integers (`NSUInteger`), `location` and `length`. So a range whose `location` is 1 starts at the second element of something (because element counting is always zero-based), and if its `length` is 2 it designates this element and the next. Cocoa also supplies various convenience methods for dealing with a range; you'll use `NSMakeRange` frequently. (Note that the name, `NSMakeRange`, is backward compared to names like `CGPointMake` and `CGRectMake`.)

`NSNotFound` is a constant integer indicating that some requested element was not found. For example, if you ask for the index of a certain object in an `NSArray` and the object isn't present in the array, the result is `NSNotFound`. The result could not be 0

to indicate the absence of the object, because `0` would indicate the first element of the array. Nor could it be `nil`, because `nil` is `0` (and in any case is not appropriate when an integer is expected). The true numeric value of `NSNotFound` is of no concern to you; always compare against `NSNotFound` itself, to learn whether a result is a meaningful index.

If a search returns a range and the thing sought is not present, the `location` component of the resulting `NSRange` will be `NSNotFound`.

NSString and Friends

`NSString`, which has already been used rather liberally in examples earlier in this book, is the Cocoa object version of a string. You can create an `NSString` through a number of class methods and initializers, or by using the `NSString` literal notation `@"..."`, which is really a compiler directive. Particularly important is `stringWithFormat:`, which lets you convert numbers to strings and combine strings; see [Chapter 9](#), where I discussed format strings in connection with `NSLog`.

```
int x = 5;
NSString* s = @"widgets";
NSString* s2 = [NSString stringWithFormat:@"You have %i %@", x, s];
```

`NSString` has a modern, Unicode-based idea of what a string can consist of. A string’s “elements” are its characters, whose count is its `length`. These are not bytes, because the numeric representation of a Unicode character could be multiple bytes, depending on the encoding. Nor are they glyphs, because a composed character sequence that prints as a single “letter” can consist of multiple characters. Thus the length of an `NSRange` indicating a single “character” might be greater than 1.

An `NSString` can be searched using various `rangeOf...` methods, which return an `NSRange`. In addition, `NSScanner` lets you walk through a string looking for pieces that fit certain criteria; for example, with `NSScanner` (and `NSCharacterSet`) you can skip past everything in a string that precedes a number and then extract the number. The `rangeOfString:` family of search methods can specify an option `NSRegularExpressionSearch`, which lets you search using a regular expression; regular expressions are also supported as a separate class, `NSRegularExpression` (which uses `NSTextCheckingResult` to describe match results).

In this example from one of my apps, the user has tapped a button whose title is something like “5 by 4” or “4 by 3”. I want to know both numbers; one tells me how many rows the layout is to have, the other how many columns. I use an `NSScanner` to locate the two numbers:

```

NSString* s = [as buttonTitleAtIndex:ix];
NSScanner* sc = [NSScanner scannerWithString:s];
int rows, cols;
[sc scanInt:&rows];
[sc scanUpToCharactersFromSet:[NSCharacterSet decimalDigitCharacterSet]
    intoString:nil];
[sc scanInt:&cols];

```

Here's how I might do the same thing using a regular expression:

```

NSString* s = [as buttonTitleAtIndex:ix];
int rowcol[2]; int* prowcol = rowcol;
NSError* err = nil;
NSRegularExpression* r = [NSRegularExpression regularExpressionWithPattern:@"\\d"
    options:0
    error:&err];
// error-checking omitted
for (NSTextCheckingResult* match in [r matchesInString:s
    options:0
    range:NSMakeRange(0, [s length])])
    *prowcol++ = [[s substringWithRange: [match range]] intValue];

```

The syntax seems oddly tortured, though, because we must convert each match from an NSTextCheckingResult to a range, then to a substring of our original string, and finally to an integer.

More sophisticated automated textual analysis is supported by some additional classes, such as NSDataDetector, an NSRegularExpression subclass that efficiently finds certain types of string expression such as a URL or a phone number, and NSLinguisticTagger (new in iOS 5), which actually attempts to analyze text into its grammatical parts of speech.

An NSString object's string is immutable. You can use a string to generate another string in various ways, such as by appending another string or by extracting a substring, but you can't alter the string *itself*. For that, you need NSString's subclass, NSMutableString.

An NSString carries no font and size information. In iOS programming, interface objects that display strings (such as UILabel) have a font property that is a UIFont, which is used to determine the single font and size in which the string will display. String drawing in a graphics context can be performed simply with methods provided through the NSStringDrawing category on NSString (see the *String UIKit Additions Reference*). Complex string layout in a graphics context, including use of styled text, requires Core Text and is a separate topic ([Chapter 23](#)).

NSString has convenience utilities for working with a file path string, and is often used in conjunction with NSURL, which is another Foundation class worth looking into. NSString and some other classes discussed in this section provide methods for writing out to a file's contents or reading in a file's contents; when they do, the file can be specified either as an NSString file path or as an NSURL ([Chapter 36](#)).

NSDate and Friends

An NSDate is a date and time, represented internally as a number of seconds (NSTimeInterval) since some reference date. Calling [`NSDate date`] gives you a date object for the current date and time; other date operations may involve NSDateComponents and NSCalendar and can be a bit tricky because calendars are complicated (see the *Date and Time Programming Guide*).

You will also likely be concerned with dates represented as strings. Creation and parsing of date strings involves NSDateFormatter, which uses a format string similar to NSString's `stringWithFormat`. A complication is added by the fact that the exact string representation of a date component or format can depend upon the user's locale, consisting of language, region format, and calendar settings. (Actually, locale considerations can also play a role in NSString format strings.)

In this example from one of my apps, I prepare the content of a UILabel reporting the date and time when our data was last updated. The app is not localized — the word “at” appearing in the string is always going to be in English — so I want complete control of the presentation of the date and time components as well. To get it, I have to insist upon a particular locale:

```
NSDateFormatter *df = [[NSDateFormatter alloc] init];
if ([[NSLocale availableLocaleIdentifiers] indexOfObject:@"en_US"] != NSNotFound) {
    NSLocale* loc =
        [[NSLocale alloc] initWithLocaleIdentifier:@"en_US"];
    [df setLocale:loc]; // force English month name and time zone name if possible
}
[df setDateFormat:@"d MMMM yyyy 'at' h:mm a z"];
NSString* lastUpdated = [df stringFromDate: [NSDate date]];
```

Locales are an interesting and complicated topic; to learn more, consult the ICU (International Components for Unicode), from which the iOS support for creating and parsing date strings is derived. To study what locales exist, use the locale explorer at <http://demo.icu-project.org/icu-bin/locexp>.

NSNumber

An NSNumber is an object that wraps a numeric value (including BOOL). Thus, you can use it to store and pass a number where an object is expected. An NSNumber is formed from an actual number with a method that specifies the numeric type; for example, you can call `numberWithInt:` to form a number from an int:

```
[[NSUserDefaults standardUserDefaults] registerDefaults:
    [NSDictionary dictionaryWithObjectsAndKeys:
        [NSNumber numberWithInt: 4],
        @"cardMatrixRows",
        [NSNumber numberWithInt: 3],
        @"cardMatrixColumns",
        nil]];
```

An NSNumber is not itself a number, so you can't use it in calculations or where an actual number is expected. Instead, you must extract the number from its NSNumber wrapper using the inverse of the method that wrapped the number to begin with. So, for example, if an NSNumber wraps an int, you can call `intValue` to extract the int:

```
NSUserDefaults* ud = [NSUserDefaults standardUserDefaults];
int therows = [[ud objectForKey:@"cardMatrixRows"] intValue];
int thecols = [[ud objectForKey:@"cardMatrixColumns"] intValue];
```

Actually, this is such a common transformation when communicating with NSUserDefaults that it provides convenience methods. So I could have written the same thing this way:

```
NSUserDefaults* ud = [NSUserDefaults standardUserDefaults];
int therows = [ud integerForKey:@"cardMatrixRows"];
int thecols = [ud integerForKey:@"cardMatrixColumns"];
```

NSValue

NSValue is NSNumber's superclass. Use it for wrapping nonnumeric C values such as structs. Convenience methods provided through the NSValueUIKitGeometryExtensions category on NSValue (see the *NSValue UIKit Additions Reference*) allow easy wrapping and unwrapping of CGPoint, CGSize, CGRect, CGAffineTransform, UIEdgeInsets, and UIOffset; additional categories allow easy wrapping and unwrapping of CATransform3D, CMTime, CMTimeMapping, and CMTimeRange.

You are unlikely to need to store any other kind of C value in an NSValue, but you can if you need to.

NSData

NSData is a general sequence of bytes. It is immutable; the mutable version is its subclass NSMutableData.

In practice, NSData tends to arise in two main ways:

- When downloading data from the Internet. For example, the `NSURLConnection` class supplies whatever it retrieves from the Internet as `NSData`. Transforming it from there into (let's say) a string, specifying the correct encoding, would then be up to you.
- When storing an object as a file or in user preferences. For example, you can't store a `UIColor` value directly into user preferences. So if the user has made a color choice and you need to save it, you transform the `UIColor` into an `NSData` (using `NSKeyedArchiver`) and save that:

```
[[NSUserDefaults standardUserDefaults] registerDefaults:
    [NSDictionary dictionaryWithObjectsAndKeys:
        [NSKeyedArchiver archivedDataWithRootObject:[UIColor blueColor]],
        @"myColor",
        nil]];
```

The use of NSKeyedArchiver, and its reversal with NSKeyedUnarchiver, is a separate topic ([Chapter 36](#)).

Equality and Comparison

The foregoing types will quickly come to seem to you like basic data types, but of course they are actually object types. Therefore you cannot compare them with the C operators for testing equality as you would with actual numbers. That's because, in the case of object types, the C operators compare the pointers, not the object content of the instances. For example:

```
NSString* s1 = [NSString stringWithFormat:@"%@, %@", @"Hello", @"world"];
NSString* s2 = [NSString stringWithFormat:@"%@, %@", @"Hello", @"world"];
if (s1 == s2) // false
// ...
```

The two strings are equivalent (@"Hello, world") but are not the same object. (The example is deliberately elaborate because Cocoa's efficient management of string literals sees to it that two strings initialized directly as @"Hello, world" *are* the same object, which wouldn't illustrate the point I'm making.) It is up to individual classes to implement a test for equality. The general test, `isEqual:`, is inherited from `NSObject` and overridden, but some classes also define more specific and efficient tests. Thus, the correct way to perform the above test is like this:

```
if ([s1 isEqualToString: s2])
```

Similarly, it is up to individual classes to supply ordered comparison methods. The standard method is called `compare:`, and returns one of three constants: `NSOrderedAscending` (the receiver is less than the parameter), `NSOrderedSame` (the receiver is equal to the parameter), or `NSOrderedDescending` (the receiver is greater than the parameter).

NSIndexSet

`NSIndexSet` expresses a collection of unique whole numbers; its purpose is to express element numbers of an ordered collection, such as an `NSArray`. Thus, for instance, to retrieve multiple objects simultaneously from an array, you specify the desired indexes as an `NSIndexSet`. It is also used with other things that are array-like; for example, you pass an `NSIndexSet` to a `UITableView` to indicate what sections to insert or delete.

To take a specific example, let's say you want to speak of elements 1, 2, 3, 4, 8, 9, and 10 of an `NSArray`. `NSIndexSet` expresses this notion in some compact implementation that can be readily queried. The actual implementation is opaque, but you can imagine

that in this case the set might consist of two NSRange structs, {1,4} and {8,3}, and NSIndexSet's methods actually invite you to think of as NSIndexSet as composed of ranges.

An NSIndexSet is immutable; its mutable subclass is NSMutableIndexSet. You can form a simple NSIndexSet consisting of just one contiguous range directly, by passing an NSRange to `indexSetWithIndexesInRange:`; but to form a more complex index set you'll need to use NSMutableIndexSet so that you can append additional ranges.

Walking through (enumerating) the index values specified by an NSIndexSet is easy starting in iOS 4.0, which provides `enumerateIndexesUsingBlock:`. and iOS 5.0 adds ways of enumerating ranges as well. (As of this writing, these are not documented on the NSIndexRange class documentation page; but they do appear in the header. Look for the three methods whose names begin with `enumerateRanges....`) But if your code is to run on earlier systems, you can't use blocks, and no enumerator is provided, so you must resort to a rather clumsy construct ([Example 10-2](#)).

Example 10-2. Enumerating an NSIndexSet before iOS 4.0

```
NSIndexSet* ixen = //...;
NSUInteger ix = [ixen firstIndex];
do {
    // ... do something with ix ...
} while ((ix = [ixen indexGreaterThanOrEqualTo:ix]) != NSNotFound);
```

NSArray and NSMutableArray

An NSArray is an ordered collection of objects. Its length is its `count`, and a particular object can be obtained by index number using `objectAtIndex:`. The index of the first object is zero, so the index of the last object is `count` minus one. You can form an NSArray in various ways, but typically you'll start by supplying a list of the objects it is to contain (see [Chapter 3](#)).

An NSArray is immutable. This doesn't mean you can't mutate any of the objects it contains; it means that once the NSArray is formed you can't remove an object from it, insert an object into it, or replace an object at a given index. To do those things, you can derive a new array consisting of the original array plus or minus some objects, or use NSArray's subclass, NSMutableArray.

You can walk through (enumerate) every object in an array with the `for...in` construct described in [Chapter 1](#). (You'll get an exception if you try to mutate an array while enumerating it.)

You can seek an object within an array with `indexOfObject:` or `indexOfObjectIdenticalTo:`; the former's idea of equality is to call `isEqual:`, whereas the latter uses pointer equality.

Those familiar with other languages may miss such utility array functions as `map`, which builds a new array of the results of calling a method on each object in the array. (`make-`

`ObjectsPerformSelector`: requires a selector that returns no value, and `enumerateObjectsUsingBlock`: requires a block function that returns no value.) The usual workaround is to make an empty mutable array and then enumerate the original array, calling a method and appending each result to the mutable array ([Example 10-3](#)). It is also sometimes possible to use key-value coding as a `map` substitute (see [Chapter 12](#)).

Example 10-3. Building an array by enumerating another array

```
NSMutableArray* marr = [NSMutableArray array];
for (id obj in myArray) {
    id result = [obj doSomething];
    [marr addObject: result];
}
```

You can filter an array to produce a new array consisting of just those objects meeting a test that can be described as an `NSPredicate`:

```
NSArray* pep = [NSArray arrayWithObjects: @"Manny", @"Moe", @"Jack", nil];
NSPredicate* p = [NSPredicate predicateWithFormat:@"self BEGINSWITH[cd] 'm'"];
NSArray* ems = [pep filteredArrayUsingPredicate:p];
```

To search or filter an array on a more customized test, you can walk through the array applying the test and adding those that meet it to an `NSMutableArray` (similar to [Example 10-3](#)). And there are many methods that give you the ability to search or filter an array using a block:

```
NSArray* pep = [NSArray arrayWithObjects: @"Manny", @"Moe", @"Jack", nil];
NSArray* ems =
    [pep objectsAtIndexes: [pep indexesOfObjectsPassingTest:
        ^BOOL(id obj, NSUInteger idx, BOOL *stop) {
            return ((NSString*)obj rangeOfString:@"m"
                options:NSCaseInsensitiveSearch].location == 0);
    }]];
```

You can derive a sorted version of the array, supplying the sorting rules in various ways, or if it's a mutable array, you can sort it directly.



Forming a new array from some or all of the elements of an existing array is *not* an expensive operation. The objects constituting the elements of the first array are not copied; the new array consists merely of a new set of pointers to the already existing objects. The same is true for the other collection types I'm about to discuss.

NSSet and Friends

An `NSSet` is an unordered collection of distinct objects. “Distinct” means that no two objects in a set can return YES when they are compared using `isEqual:`. Learning whether an object is present in a set is much more efficient than seeking it in an array, and you can ask whether one set is a subset of, or intersects, another set. You can walk through (enumerate) a set with the `for...in` construct, though the order is of course

undefined. You can filter a set, as you can an array. Indeed, much of what you can do with a set is parallel to what you can do with an array, except that of course you can't do anything with a set that involves the notion of ordering. And starting in iOS 5, that restriction is lifted, through the introduction of a new class, `NSOrderedSet`.

An ordered set (`NSOrderedSet`) is *very* like an array, and the methods for working with it are very similar to the methods for working with an array, except that an ordered set's elements must be distinct. Handing an array over to an ordered set *uniques* the array, meaning that order is maintained but only the first occurrence of an equal object is moved to the set. An ordered set provides many of the advantages of sets: for example, as with an `NSSet`, learning whether an object is present in an ordered set is much more efficient than for an array, and you can readily take the union, intersection, or difference with another set. Since the distinctness restriction will often prove no restriction at all (because the elements were going to be distinct anyway), it is likely that programmers will want to get into the habit of using `NSOrderedSet` instead of `NSArray` wherever possible.

An `NSSet` is immutable. You can derive one `NSSet` from another by adding or removing elements, or you can use its subclass, `NSMutableSet`. Similarly, `NSOrderedSet` has its mutable counterpart, `NSMutableOrderedSet`. There is no penalty for adding to, or inserting into, a mutable set an object that the set already contains; nothing is added (and so the distinctness rule is enforced), but there's no error.

`NSCountedSet`, a subclass of `NSMutableSet`, is a mutable unordered collection of objects that are *not* necessarily distinct (this concept is usually referred to as a *bag*). It is implemented as a set plus a count of how many times each element has been added.

NSDictionary and NSMutableDictionary

An `NSDictionary` is an unordered collection of key–value pairs. The key is usually an `NSString`, though it doesn't have to be. The value can be any object. An `NSDictionary` is immutable; its mutable subclass is `NSMutableDictionary`.

The keys of a dictionary are distinct (using `isEqual:` for comparison). If you add a key–value pair to an `NSMutableDictionary`, then if that key is not already present, the key–value pair is simply added, but if the key is already present, then the corresponding value is replaced.

The fundamental use of an `NSDictionary` is to request an entry's value by key (using `objectForKey:`); if no such key exists, the result is nil, so this is also the way to find out whether a key is present. A dictionary is thus an easy, flexible data storage device, an object-based analogue to a struct. Cocoa often uses a dictionary to provide you with an extra packet of named values, as in the `userInfo` of an `NSNotification`, the `options` parameter of `application:didFinishLaunchingWithOptions:`, and so on.

Data structures such as an array of dictionaries, a dictionary of dictionaries, and so forth, are extremely common, and will often lie at the heart of an app's functionality.

Here's an example from one of my own apps. The app bundle contains a text file laid out like this:

```
chapterNumber [tab] pictureName [return]
chapterNumber [tab] pictureName [return]
```

As the app launches, I load this text file and parse it into a dictionary, each entry of which has the following structure:

```
key: (chapterNumber, as an NSNumber)
value: [Mutable Array]
    (pictureName)
    (pictureName)
...
...
```

Thus, as we walk the text file, we end up with all pictures for a chapter collected under the number of that chapter. This makes it easy for me later to present all the pictures for a given chapter. For each line of the text file, if the dictionary entry for that chapter number doesn't exist, we create it, with an empty mutable array as its value. Whether that dictionary entry existed or not, it does now, and its value is a mutable array, so we append the picture name to that mutable array. Observe how this single typical example ([Example 10-4](#)) brings together many of the Foundation classes discussed in this section.

Example 10-4. Parsing a file with Foundation classes

```
NSString* f = [[NSBundle mainBundle] pathForResource:@"index" ofType:@"txt"];
NSError* err = nil;
NSString* s = [NSString stringWithContentsOfFile:f
                                         encoding:NSUTF8StringEncoding
                                         error:&err];
// error-checking omitted
NSMutableDictionary* d = [NSMutableDictionary dictionary];
for (NSString* line in [s componentsSeparatedByString:@"\n"]) {
    NSArray* items = [line componentsSeparatedByString:@"\t"];
    NSInteger chnum = [[items objectAtIndex: 0] integerValue];
    NSNumber* key = [NSNumber numberWithInteger:chnum];
    NSMutableArray* marr = [d objectForKey: key];
    if (!marr) { // no such key, create key-value pair
        marr = [NSMutableArray array];
        [d setObject: marr forKey: key];
    }
    // marr is now a mutable array, empty or otherwise
    NSString* picname = [items objectAtIndex: 1];
    [marr addObject: picname];
}
```

You can get from an NSDictionary a list of keys, a sorted list of keys, or a list of values. You can walk through (enumerate) a dictionary by its keys with the `for...in` construct, though the order is of course undefined. A dictionary also supplies an `object-Enumerator`, which you can use with the `for...in` construct to walk through just the

values. Starting in iOS 4.0, you can also walk through the key–value pairs together using a block, and you can even filter an NSDictionary by some test against its values.

NSNull

NSNull does nothing but supply a pointer to a singleton object, `[NSNull null]`. Use this singleton object to stand for nil in situations where an actual object is required and nil is not permitted. For example, you can't use nil as the value of an element of a collection (such as NSArray, NSSet, or NSDictionary), so you'd use `[NSNull null]` instead.

Despite what I said earlier about equality, you can test an object against `[NSNull null]` using the C equality operator, because this is a singleton instance and therefore pointer comparison works.

Immutable and Mutable

Beginners sometimes have difficulty with the Foundation's immutable/mutable class pairs, so here are some hints.

The documentation may not make it completely obvious that the mutable classes obey and, if appropriate, override the methods of the immutable classes. Thus, for example, `[NSArray array]` generates an immutable array, but `[NSMutableArray array]` generates a mutable array. (You will look in vain for the expected `[NSMutableArray mutableArray]`.) The same is true of all the initializers and convenience class methods for instantiation: they may all have “array” in their name, but when sent to NSMutableArray, they yield a mutable array.

That fact also answers the question of how to make an immutable array mutable, and *vice versa*. If `arrayWithArray:`, sent to the NSArray class, yields a new immutable array containing the same objects in the same order as the original array, then the same method, `arrayWithArray:`, sent to the NSMutableArray class, yields a *mutable* array containing the same objects in the same order as the original. Thus this single method can transform an array between immutable and mutable in either direction. You can also use `copy` (produces an immutable copy) and `mutableCopy` (produces a mutable copy).

All of the above applies equally, of course, to the other immutable/mutable class pairs. You will often want to work internally and temporarily with a mutable instance but then store (and possibly vend, as an instance variable) an immutable instance, thus protecting the value from being changed accidentally or behind your own back. What matters is not a variable's declared class but what class the instance really is (polymorphism; see [Chapter 5](#)), so it's good that you can easily switch between an immutable and a mutable version of the same data.

To test whether an instance is mutable or immutable, do *not* ask for its `class`. These immutable/mutable class pairs are all implemented as *class clusters*, which means that Cocoa uses a secret class, different from the documented class you work with. This secret class is subject to change without notice, because it's none of your business and you should never have looked at it in the first place. Thus, code of this form is subject to breakage:

```
if ([NSStringFromClass([n class]) isEqualToString: @"NSCFArray"]) // wrong!
```

Instead, to learn whether an object is mutable, ask it whether it responds to a mutability method:

```
if ([n respondsToSelector:@selector(addObject:)]) // right
```

Bear in mind also that just because a collection class is immutable doesn't mean that the objects it collects are immutable. They are still objects and do not lose any of their normal behavior merely because they are pointed to by an immutable collection.

Property Lists

A *property list* is a string (XML) representation of data. The Foundation classes `NSString`, `NSData`, `NSArray`, and `NSDictionary` are the only classes that can be converted into a property list. Moreover, an `NSArray` or `NSDictionary` can be converted into a property list only if the only classes it collects are these classes, along with `NSDate` and `NSNumber`. (This is why, as mentioned earlier, you must convert a `UIColor` into an `NSData` in order to store it in user defaults; the user defaults is a property list.)

The primary use of a property list is to store data as a file. `NSArray` and `NSDictionary` provide convenience methods `writeToFile:atomically:` and `writeToURL:atomically:` that generate property list files given a pathname or file URL, respectively; they also provide inverse convenience methods that initialize an `NSArray` object or an `NSDictionary` object based on the property list contents of a given file. For this very reason, you are likely to start with one of these classes when you want to create a property list. (`NSString` and `NSData`, with their methods `writeToFile:...` and `writeToURL:...`, just write the data out as a file directly, not as a property list.)

When you initialize an `NSArray` or `NSDictionary` from a property list file in this way, the objects in the collection are all immutable. If you want them to be mutable, or if you want to convert an instance of one of the other property list classes to a property list, you'll use the `NSPropertyListSerialization` class (see the *Property List Programming Guide*).

The Secret Life of NSObject

Because every class inherits from `NSObject`, it's worth taking some time to investigate and understand `NSObject`. `NSObject` is constructed in a rather elaborate way:

- It defines some native class methods and instance methods having mostly to do with the basics of instantiation and of method sending and resolution. (See the *NSObject Class Reference*.)
- It adopts the *NSObject* protocol. This protocol declares instance methods having mostly to do with memory management, the relationship between an instance and its class, and introspection. Because all the *NSObject* protocol methods are required, the *NSObject* class implements them all. (See the *NSObject Protocol Reference*.) This architecture is what permits *NSProxy* to be a root class; it, too, adopts the *NSObject* protocol.
- It implements convenience methods related to the *NSCopying*, *NSMutableCopying*, and *NSCoding* protocols, without formally adopting those protocols. *NSObject* intentionally doesn't adopt these protocols because this would cause all other classes to adopt them, which would be wrong. But thanks to this architecture, if a class does adopt one of these protocols, you can call the corresponding convenience method. For example, *NSObject* implements the `copy` instance method, so you can call `copy` on any instance, but you'll crash unless the instance's class adopts the *NSCopying* protocol and implements `copyWithZone:`.
- A large number of methods are injected into *NSObject* by more than two dozen categories on *NSObject*, scattered among various header files. For example, `awakeFromNibFromNib` (see [Chapter 7](#)) comes from the *UINibLoadingAdditions* category on *NSObject*, declared in *UINibLoading.h*. And `performSelector:withObject:afterDelay:`, discussed later in this chapter and used in several examples later in this book, comes from the *NSDelayedPerforming* category on *NSObject*, declared in *NSRunLoop.h*.
- A class object, as explained in [Chapter 4](#), is an object. Therefore all classes, which are objects of type *Class*, inherit from *NSObject*. Therefore, *any method defined as an instance method by NSObject can be called on a class object as a class method!* For example, `respondsToSelector:` is defined as an instance method by *NSObject*, but it can (therefore) be treated also as a class method and sent to a class object.

The problem for the programmer is that Apple's documentation is rather rigid about classification. When you're trying to work out what you can say to an object, you don't care where that object's methods come from; you just care what you can say. But Apple differentiates methods by where they come from. Even though *NSObject* is the root class, the most important class, from which all other classes inherit, *no single page of the documentation provides a conspectus of all its methods*. Instead, you have to look at both the *NSObject Class Reference* and the *NSObject Protocol Reference* simultaneously, plus the pages documenting the *NSCopying*, *NSMutableCopying*, and *NSCoding* protocols (in order to understand how they interact with methods defined by *NSObject*), plus you have to supply mentally a class method version of every *NSObject* instance method!

Of the methods injected into `NSObject` by categories, many are delegate methods used in restricted situations (so that these are really informal protocols), and do not need centralized documentation; for example, `animationDidStart:` is documented under the `CAAnimation` class, quite rightly, because you need to know about it only and exactly when you’re working with `CAAnimation`. Others that are general in nature are documented on the `NSObject` class documentation page itself; for example, `cancelPreviousPerformRequestsWithTarget:` comes from a category declared in `NSRunLoop.h`, but it is documented under `NSObject`, quite rightly, since this is a class method, and therefore effectively a global method, that you might want to send at any time. However, every object responds to `awakeFromNib`, and it’s likely to be crucial to every app you write; yet you must learn about it outside of the `NSObject` documentation, sitting all by itself on the *NSObject UIKit Additions Reference*, where you’re extremely unlikely to discover it! The same goes, it might be argued, for all the key-value coding methods ([Chapter 12](#)) and key-value observing methods ([Chapter 13](#)).

Once you’ve collected all the `NSObject` methods, you can see that they fall into a certain natural classification, much as outlined in Apple’s documentation (see also “The Root Class” in the “Cocoa Objects” section of the *Cocoa Fundamentals Guide*):

Creation, destruction, and memory management

Methods for creating an instance, such as `alloc` and `copy`, along with methods that you might override in order to learn when something is happening in the lifetime of an object, such as `initialize` (see [Chapter 11](#)) and `dealloc` (see [Chapter 12](#)), plus methods that manage memory (see [Chapter 12](#)).

Class relationships

Methods for learning an object’s class and inheritance, such as `class`, `superclass`, `isKindOfClass:`, and `isMemberOfClass:`.

To check the class of an instance (or class), use methods such as `isKindOfClass:` and `isMemberOfClass:`. Direct comparison of two class objects, as in `[someObject class] == [otherObject class]`, is rarely advisable, especially because a Cocoa instance’s class might be a private, undocumented subclass of the class you expect. I mentioned this already in connection with class clusters, and it can happen in other cases.

Object introspection and comparison

Methods for asking what would happen if an object were sent a certain message, such as `respondsToSelector:`; for representing an object as a string (`description`, used in debugging; see [Chapter 9](#)); and for comparing objects (`isEqual:`).

Message response

Methods for meddling with what does happen when an object is sent a certain message, such as `doesNotRecognizeSelector:`. If you’re curious, see the *Objective-C Runtime Programming Guide*. An example appears in [Chapter 25](#).

Message sending

Methods for sending a message indirectly. For example, `performSelector:` takes a selector as parameter, and sending it to an object tells that object to perform that selector. This might seem identical to just sending that message to that object, but what if you don't know what message to send until runtime? Moreover, variants on `performSelector:` allow you send a message on a specified thread, or send a message after a certain amount of time has passed (`performSelector:withObject:afterDelay:` and similar); this is called *delayed performance*.

Delayed performance is a valuable technique. You often need to let Cocoa finish doing something, such as laying out interface, before proceeding to a further step; delayed performance with a very short delay (even as short as zero seconds) is enough to postpone a method call until after Cocoa has finished whatever it's in the middle of doing. Technically, it allows the current run loop to finish, completing and unwinding the entire current method call stack, before sending the specified selector. It can also be used for simple timing, such as when you want to do something different depending whether the user taps twice in quick succession or only once; basically, when the user first taps, you respond using delayed performance, to give the user time to tap again if two taps are intended. Examples of both uses appear in later chapters.



Another way to implement delayed performance is by calling the Grand Central Dispatch function `dispatch_after`. It has the advantage that what to do after the delay can be expressed as a block; for this reason, I've become quite fond of it. The syntax looks forbidding at first, but if you type `dispatch_after` and ask for code completion you get a nice bit of boilerplate code, all ready for you to enter the contents of the block. Be sure to change the delay, as the default in the boilerplate, two seconds, is unlikely to be what you want.

Cocoa Events

None of your code runs until Cocoa calls it. The art of Cocoa programming consists largely of knowing when and why Cocoa will call your code. If you know this, you can put your code in the correct place, with the correct method name, so that your code runs at the correct moment, and your app behaves the way you intend.

In [Chapter 7](#), for example, we wrote a method to be called when the user taps a certain button in our interface, and we also arranged things so that that method *would* be called when the user taps that button:

```
- (void) buttonPressed: (id) sender {  
    // ... react to the button being pressed  
}
```

This architecture typifies the underpinnings of a Cocoa program. Your code itself is like a panel of buttons, waiting for Cocoa to press one. If something happens that Cocoa feels your code needs to know about and respond to, it presses the right button — if the right button is there. You organize your code with Cocoa’s behavior in mind. Cocoa makes certain promises about how and when it will dispatch messages to your code. These are Cocoa’s *events*. You know what these events are, and you arrange your code to be ready when Cocoa delivers them.

Thus, to program for Cocoa, you must, in a sense, surrender control. Your code never gets to run just whenever it feels like it. It can run *only* in response to some kind of event. Something happens, such as the user making a gesture on the screen, or some specific stage arriving in the lifetime of your app, and Cocoa dispatches an event to your code — if your code is prepared to receive it. So you don’t write just any old code you want to and put it in any old place. You use the framework, by letting the framework use you. You submit to Cocoa’s rules and promises and expectations, so that your code will be called at the right time and in the right way.

The specific events that you can receive are listed in the documentation. The overall architecture of how and when events are dispatched and the ways in which your code arranges to receive them is the subject of this chapter.

Reasons for Events

Broadly speaking, the reasons you might receive an event may be divided informally into four categories. These categories are not official; I made them up. Often it isn't completely clear which of these categories an event fits into; an event may well appear to fit two categories. But they are still generally useful for visualizing how and why Cocoa interacts with your code.

User events

The user does something interactive, and an event is triggered directly. Obvious examples are events that you get when the user taps or swipes the screen, or types a key on the keyboard.

Lifetime events

These are events notifying you of the arrival of a stage in the life of the app, such as the fact that the app is starting up or is about to go into the background, or of a component of the app, such as the fact that a `UIViewController`'s view has just loaded or is about to be removed from the screen.

Functional events

Cocoa is about to do something, and turns to you in case you want to supply (additional) functionality. I would put into this category things like `UIView`'s `drawRect:` (your chance to have a view draw itself) and `UILabel`'s `drawTextInRect:` (your chance to modify the look of a label), with which we experimented in [Chapter 10](#).

Query events

Cocoa turns to you to ask a question; its behavior will depend upon your answer. For example, the way data appears in a table (a `UITableView`) is that whenever Cocoa needs a cell for a row of the table, it turns to you and asks for the cell.

Subclassing

A built-in Cocoa class may define methods that Cocoa itself will call and that you are invited (or required) to override in a subclass. Sometimes you know when the method will be called; at other times you don't know or care exactly when the method is called, but you know that you must override it so that whenever it *is* called, your behavior, and not (merely) the default behavior, will take place.

An example I gave in [Chapter 10](#) was `UIView`'s `drawRect:`. The built-in `UIView` implementation does nothing, so overriding `drawRect:` in a subclass is your only chance to dictate the full procedure by which a view draws itself. You don't know exactly when this method will be called, and you don't care; when it is, you draw, and this guarantees that the view will always appear the way you want it to. (You never call `drawRect:` yourself; if some underlying condition has changed and you want the view to be redrawn, you call `setNeedsDisplay` and let Cocoa call `drawRect:` in response.)

In addition to `UIView`, particular built-in `UIView` subclasses may have methods you'll want to customize through subclassing. Typically this will be in order to change the way the view is drawn. In [Chapter 10](#) I gave an example involving `UILabel` and its `drawTextInRect:`. Another example is `UISlider`, which lets you customize the position and size of the slider's "thumb" by overriding `thumbRectForBounds:trackRect:value:` ([Chapter 25](#)).

`UIViewController` ([Chapter 19](#)) is a good example of a class meant for subclassing. Of the methods listed in the `UIViewController` class documentation, just about all are methods you might have reason to override. If you create a `UIViewController` subclass in Xcode, you'll see that the template already includes about half a dozen methods for you to uncomment and override if desired.

For example, you must override `loadView` if your `UIViewController` creates its view in code, and you must create and assign it to this instance's `view` property at this moment. (I'd probably call that a functional event, because your code has a specific job to do, namely, supply the view.) You may override `viewDidLoad` to perform additional initializations as your view is first loaded, whether it comes from a nib or you created it in `loadView`. Methods like `viewWillAppear:` and `viewDidDisappear:` are called as your `UIViewController`'s view takes over the screen or is replaced on the screen by some other view; thus, `viewWillAppear:` is a moment to make sure that whatever happened while your view was offscreen is reflected in how it looks as it comes back onscreen. (Those are obviously lifetime events.)

A `UIViewController` method like `shouldAutorotateToInterfaceOrientation:` is what I call a query event. It is passed an orientation parameter; your job is to return a `BOOL` telling Cocoa whether your view can appear in that orientation. The default, if you don't implement it, is that your view can appear only in portrait orientation. If you want this `UIViewController`'s view to appear in some other orientation, you'll return `YES` for that orientation. If you return `YES` for more than one orientation and the user rotates the device, you might then receive messages like `willRotateToInterfaceOrientation:duration:` and `willAnimateRotationToInterfaceOrientation:duration:`, where you can customize what happens to the view as the orientation changes.

When looking for events that you can receive through subclassing, be sure to look upward through the inheritance hierarchy. For example, if you're wondering how to be notified when your custom `UILabel` subclass is embedded into another view, you won't find the answer in the `UILabel` class documentation; a `UILabel` receives the appropriate event by virtue of being a `UIView`. In the `UIView` class documentation, you'll learn that you can override `didMoveToSuperview` to be informed when this happens.

Even further up the inheritance hierarchy, you'll find things like `NSObject`'s `initialize` class method. Every class that is actually sent a class method message (including instantiation) is first sent the `initialize` message, once. Thus, `initialize` can be overridden in order to run code extremely early in a class's lifetime (before it even has an instance). Your project's application delegate class (such as `AppDelegate` in our

Empty Window project) is instantiated very early in the app’s lifetime, so its `initialize` can be a good place to perform very early app initializations, such as setting default values for any user preferences. When implementing `initialize`, we must test, as a matter of course, whether `self` really is the class in question (and this is one of the few situations in which we will compare two classes directly against one another); otherwise there is a chance that `initialize` will be called again (and our code will run again) if a subclass of this class is used. Here’s an example from Apple’s own sample code (the BubbleLevel example, in `LevelAppDelegate.m`):

```
+ (void)initialize {
    if ([self class] == [LevelAppDelegate class]) {
        // Register a default value for the instrument calibration.
        // This will be used if the user hasn't calibrated the instrument.
        NSNumber *defaultCalibrationOffset = [NSNumber numberWithFloat:0.0];
        NSDictionary *resourceDict =
            [NSDictionary dictionaryWithObject:defaultCalibrationOffset
                                         forKey:BubbleLevelCalibrationOffsetKey];
        [[NSUserDefaults standardUserDefaults] registerDefaults:resourceDict];
    }
}
```

Notifications

Cocoa provides your app with a single instance of `NSNotificationCenter`, informally called the *notification center*. This instance is the basis of a mechanism for sending messages called *notifications*. A notification includes an instance of `NSNotification` (a *notification object*). The idea is that any object can be registered with the notification center to receive certain notifications. Another object can hand the notification center a notification object to send out (this is called *posting* the notification). The notification center will then send that notification object, in a notification, to all objects that are registered to receive it.

The notification mechanism is often described as a dispatching or broadcasting mechanism, and with good reason. It lets an object send a message without knowing or caring what object or how many objects receive it. This relieves your app’s architecture from the formal responsibility of somehow hooking up instances just so a message can pass from one to the other. When objects are conceptually “distant” from one another, notifications can be a fairly lightweight way of permitting one to message the other.

An `NSNotification` object has three pieces of information associated with it, which can be retrieved by instance methods:

Its name

An `NSString` which specifies the notification’s meaning.

Its object

An instance associated with the notification; typically, the instance that posted it.

Its userInfo

Not every notification has a `userInfo`; it is an `NSDictionary`, and can contain additional information associated with the notification. What information this `NSDictionary` will contain, and under what keys, depends on the particular notification; you have to consult the documentation. For example, the documentation tells us that `UIApplication`'s `UIApplicationDidChangeStatusBarFrameNotification` includes a `userInfo` dictionary with a key `UIApplicationStatusBarFrameUserInfoKey` whose value is the status bar's frame. When you post a notification yourself, you can put anything you like into the `userInfo` for the notification's recipient(s) to retrieve.

Receiving a Built-In Notification

Cocoa itself posts notifications through the notification center, and your code can register to receive them. You'll find a separate Notifications section in the documentation for a class that provides them.

To register for a notification, you have two choices. One is to use the `addObserver:selector:name:object:` method, usually sent to the app's default notification center, `[NSNotificationCenter defaultCenter]`. The parameters are as follows:

`addObserver:`

The instance to which the notification is to be sent. This will typically be `self`; it isn't usual for one instance to register a different instance as the receiver of a notification.

`selector:`

The message to be sent to the observer instance when the notification occurs. The designated method should return `void` and should take one parameter, which will be the `NSNotification` object (so the parameter should be typed as `NSNotification*` or `id`). Don't forget to implement this method! If the notification center sends a notification by sending the message specified as the `selector:` here, and there is no method implemented to receive this message, your app will crash. Failing to specify the `selector:` accurately is a common beginner mistake.

`name:`

The `NSString` `name` of the notification you'd like to receive. If this parameter is `nil`, you're asking to receive *all* notifications sent by the object designated in the `object` parameter. A built-in Cocoa notification's name is usually a constant. As I explained in [Chapter 1](#), this is helpful, because if you flub the name of a constant, the compiler will complain, whereas if you enter the name of the notification directly as an `NSString` literal and you get it wrong, the compiler won't complain but you will mysteriously fail to get any notifications (because no notification has the name you actually entered) — a very difficult sort of mistake to track down.

object:

The `object` of the notification you’re interested in, which will usually be the object that posted it. If this is `nil`, you’re asking to receive *all* notifications with the name designated in the `name` parameter. (If both the `name` and `object` parameters are `nil`, you’re asking to receive all notifications!)

For example, in one of my apps I need to respond, by changing my interface, if the device’s music player starts playing a different song. The API for the device’s built-in music player is the `MPMusicPlayerController` class; it provides a notification to tell me when the built-in music player changes its playing state, listed under Notifications in the `MPMusicPlayerController`’s class documentation as `MPMusicPlayerControllerNowPlayingItemDidChangeNotification`.

It turns out, looking at the documentation, that this notification won’t be posted at all unless I call `MPMusicPlayerController`’s `beginGeneratingPlaybackNotifications` instance method. This architecture is not uncommon; Cocoa saves itself some time and effort by not sending out certain notifications unless they are switched on, as it were. So my first job is to get an instance of `MPMusicPlayerController` and call this method:

```
MPMusicPlayerController* mp = [MPMusicPlayerController iPodMusicPlayer];
[mp beginGeneratingPlaybackNotifications];
```

Now I register myself to receive the desired playback notification:

```
[[NSNotificationCenter defaultCenter] addObserver:self
    selector:@selector(nowPlayingItemChanged:)
    name:MPMusicPlayerControllerNowPlayingItemDidChangeNotification
    object:nil];
```

So now, whenever an `MPMusicPlayerControllerNowPlayingItemDidChangeNotification` is posted, my `nowPlayingItemChanged:` method will be called:

```
- (void) nowPlayingItemChanged: (NSNotification*) n {
    MPMusicPlayerController* mp = [MPMusicPlayerController iPodMusicPlayer];
    self->nowPlayingItem = mp.nowPlayingItem;
    // ... and so on ...
}
```

The other way to register to receive a notification is by calling `addObserverForName:object:queue:usingBlock:`. Its `name:` and `object:` parameters are just like those of `addObserver:selector:name:object:`, but it doesn’t specify an observer and it doesn’t specify a selector. Instead, you provide a block — the actual code to be executed when the notification fires. There is no need to specify an observer, because of the nature of blocks: `self` inside the block is the `self` who defined the block in the first place. (The `queue:` argument will almost always be `nil`, so I’m not going to discuss it further. There is no `userInfo:` parameter, because it isn’t needed: there can be no difficulty providing the block’s code with whatever values it needs.)

This way of registering for a notification has some tremendous advantages. Look at all the cautions I’m issuing in the preceding paragraphs when you’re working with `addObserver:selector:name:object:` — get the selector right, make sure you implement

the corresponding method. With a block, there is no selector and there is no separate method; everything happens right there in the call to `addObserverForName:object:queue:usingBlock:`, like this:

```
MPMusicPlayerController* mp = [MPMusicPlayerController iPodMusicPlayer];
[mp beginGeneratingPlaybackNotifications];
id ob = [[NSNotificationCenter defaultCenter]
addObserverForName:MPMusicPlayerControllerNowPlayingItemDidChangeNotification
object:nil queue:nil usingBlock:^(NSNotification *n) {
    self->nowPlayingItem = mp.nowPlayingItem;
    // ... and so on ...
}]
};
```

In addition to the convenience and reliability using a block as you register for a notification, there is also the matter of maintaining and understanding your code. Heavy use of `addObserver:selector:name:object:` means that your code ends up peppered with methods that exist solely in order to be called by the notification center. But there is nothing about these methods that tells you exactly *why* these methods are being called — you will probably want to use explicit comments in order to remind yourself — and the methods are separate from the registration call, all of which makes your code very method-heavy and confusing. With a block, on the other hand, the whole purpose of the registration is crystal-clear. You can read the preceding code almost as an English sentence expressing it: “When the now-playing item changes, store it in my `nowPlayingItem` instance variable...”. And notice how, in the block, I don’t have to re-define `mp` as I did in the separate method `nowPlayingItemChanged:`; it is still in scope from where it was defined a couple of lines earlier. Blocks are so convenient!

Indeed, when I wrote the first edition of this book, iOS 4, which introduced blocks, was relatively new, and I had code that still needed to run on earlier systems, where blocks didn’t exist. So, at that time, I was using `addObserver:selector:name:object:` exclusively. Now, however, I’ve given up all thought of backwards compatibility stretching that far, and I use `addObserverForName:object:queue:usingBlock:` almost exclusively. I can’t overstate how much clearer and more maintainable this change has made my code.

Unregistering

It is up to you, for every object that you register as a recipient of notifications, to unregister that object before it goes out of existence. If you fail to do this, and if the object does go out of existence, and if a notification for which that object is registered is posted, the notification center will attempt to send the appropriate message to that object, which is now missing in action. The result will be a crash at best, and chaos at worst.

To unregister an object as a recipient of notifications, send the notification center the `removeObserver:` message. (Alternatively, you can unregister an object for just a specific set of notifications with `removeObserver:name:object:..`) The object passed as the

observer: argument is the object that is no longer to receive notifications. What object that is depends on how you registered in the first place:

You called addObserver:....

You supplied an observer originally; that is the observer you must now unregister.

You called addObserverForName:....

The call returned an observer token object, which you captured as an **id** variable (its real class and nature are no concern of yours); that is the observer you must now unregister.

The trick is finding the right moment to do this. In most cases, the easiest solution is the registered instance's `dealloc` method, this being the last event an instance is sent before it goes out of existence ([Chapter 12](#)). Some would say, indeed, that even if you think there's another, better place to unregister, you should *still* unregister in `dealloc` as well, just to be on the safe side. (This use of `dealloc` is not affected by the fact that you may be using ARC.)

Keep it simple, because complicated logic for registering and unregistering for notifications can be difficult to debug, especially as `NSNotificationCenter` provides no kind of introspection: you cannot ask an `NSNotificationCenter` what objects are registered with it as notification recipients. I once had a devil of a time understanding why one of my instances was not receiving a notification for which it was registered. Caveman debugging didn't help. Eventually I realized that some code I'd forgotten about was unregistering my instance.

If you're calling `addObserverForName:....`, you're going to end up receiving from the notification center one or more observer tokens, which you need to preserve so that you can unregister by handing them back to the notification center. If you're going to unregister everything at once, it is easiest to handle this situation an instance variable that is a mutable collection. So, for example, I might have an `NSMutableSet` instance variable called `observers`. Early on, I initialize it to an empty set:

```
self->observers = [NSMutableSet set];
```

Each time I register for a notification using a block, I capture the result and add it to the set:

```
id ob = [[NSNotificationCenter defaultCenter]
          addObserverForName:@"whatever" object:nil queue:nil
          usingBlock:^(NSNotification *note) {
    // ... whatever ...
  }];
[self->observers addObject:ob];
```

Ultimately, I unregister by enumerating the set:

```
for (id ob in self->observers)
  [[NSNotificationCenter defaultCenter] removeObserver:ob];
self->observers = nil;
```

This sounds tedious, and indeed it is a slight price that one must pay in order to take advantage of blocks when using notifications. However, under ARC it's much less tedious than it used to be. (I'll explain in [Chapter 12](#) about that last line, setting the `observers` instance variable to nil.)



I am skipping over some other aspects of notifications that you probably won't need to know about. Read Apple's *Notification Programming Topics for Cocoa* if you want the gory details.

NSTimer

An timer (NSTimer) is not, strictly speaking, a notification; but it behaves very similarly. It is an object that gives off a signal (*fires*) after the lapse of a certain time interval. The signal is a message to one of your instances. Thus you can arrange to be notified when a certain time has elapsed. The timing is not perfectly accurate, but it's pretty good.

Timer management is not exactly tricky, but it is a little unusual. A timer that is actively watching the clock is said to be *scheduled*. A timer may fire once, or it may be a *repeating* timer. To make a timer go out of existence, it must be *invalidated*. A timer that is set to fire once is invalidated automatically after it fires; a repeating timer repeats until *you* invalidate it (by sending it the `invalidate` message). An invalidated timer should be regarded as off-limits: you cannot revive it or use it for anything further, and you should probably not send any messages to it.

The straightforward way to create a timer is with the NSTimer class method `scheduledWithTimeInterval:target:selector:userInfo:repeats:`. This creates the timer and schedules it, so that it begins watching the clock immediately. The target and selector determine what message will be sent to what object when the timer fires; the method in question should take one parameter, which will be a reference to the timer. The `userInfo` is just like the `userInfo` of a notification. (You can see why I categorize timers as being similar to notifications.)

For example, one of my apps is a game with a score; I want to penalize the user, by diminishing the score, for not making a move within ten seconds of the previous move. So each time the user makes a move, I create a repeating timer whose time interval is ten seconds (and I also invalidate any existing timer); in the method that the timer calls, I diminish the score.

Delegation

Delegation is an object-oriented design pattern, a relationship between two objects, in which the first object's behavior is customized or assisted by the second. The second

object is the first object's *delegate*. No subclassing is involved, and indeed the first object is agnostic about the second object's class.

As implemented by Cocoa, here's how delegation works. A built-in Cocoa class has an instance variable, usually called `delegate` (it will certainly have `delegate` in its name). For some instance of that Cocoa class, you set the value of this instance variable to an instance of one of *your* classes. At certain moments in its activity, the Cocoa class promises to turn to its delegate for instructions by sending it a certain message: if the Cocoa instance finds that its delegate is not `nil`, and that its delegate is prepared to receive that message (see [Chapter 10](#) on `respondsToSelector:`), the Cocoa instance sends the message to the delegate.

In the old days, delegate methods were listed in the Cocoa class's documentation, and their method signatures were made known to the compiler through an informal protocol (a category on `NSObject`). Now, though, a class's delegate methods are usually listed in a genuine protocol with its own documentation. There are over 70 Cocoa delegate protocols, showing how heavily Cocoa relies on delegation. Most delegate methods are optional, but in a few cases you'll discover some that are required.

To customize a Cocoa instance's behavior through delegation, you start with one of your classes, which, if necessary, declares conformance to the relevant delegate protocol. When the app runs, you set the Cocoa instance's `delegate` ivar (or whatever its name is) to an instance of your class. You might do this in code, usually through a property; or might do it, by connecting an instance's `delegate` outlet (or whatever it's called) to an appropriate instance that is to serve as delegate. Your delegate class will probably do other things besides serving as this instance's delegate. Indeed, one of the nice things about delegation is that it leaves you free to slot delegate code into your class architecture however you like. For example, if a view has a controller (a `UIViewController`), it will often make sense for the controller to serve also as the view's delegate.

Here's a simple example, involving `UIAlertView`. If a `UIAlertView` has no delegate, then when its Cancel button is tapped, the alert view is dismissed. But if you want to *do* something in response to the alert view being dismissed, you need to give it a delegate so that you can receive an event telling you that the alert view *was* dismissed. It's so common to give a `UIAlertView` a delegate that its designated initializer allows you to supply one; typically, the delegate will be the instance that summoned the alert view in the first place. Moreover, an alert view with a delegate is so common that the delegate is typed as a pure `id`, so you don't even have to bother conforming formally to the `UIAlertViewDelegate` protocol:

```
- (void) gameWon {
    UIAlertView* av =
        [[UIAlertView alloc] initWithTitle:@"Congratulations!"
                                    message:@"You won the game. Another game?"
                                   delegate:self
                          cancelButtonTitle:@"No, thanks."
                          otherButtonTitles:@"Sure!", nil];
```

```

        [av show];
    }

    - (void) alertView:(UIAlertView*) av didDismissWithButtonIndex: (NSInteger) ix {
        if (ix == 1) { // user said "Sure!"
            [self newGame];
        }
    }
}

```

The delegation mechanism is the last piece of the puzzle needed to explain the built-in bootstrapping procedure of a minimal app like our Empty Window project. Recall, from [Chapter 6](#), that the fourth argument to `UIApplicationMain` is the shared application instance's delegate. So, having instantiated the class of the shared application (usually `UIApplication`), `UIApplicationMain` instantiates the class nominated in its fourth argument — in this case, `AppDelegate`; there is now, therefore, before much of anything has happened, an `AppDelegate` instance. Moreover, `UIApplicationMain` sets the shared application's `delegate` property to this `AppDelegate` instance. `UIApplication`'s `delegate` is typed as `id <UIApplicationDelegate>` — and, by golly, `AppDelegate` is declared as conforming to the `UIApplicationDelegate` protocol.

This, as we saw in [Chapter 10](#), causes the shared application instance to bombard this `AppDelegate` instance with `respondsToSelector:` messages, to find out exactly which `UIApplicationDelegate` methods `AppDelegate` actually implements. One such method that it does implement is `application:didFinishLaunchingWithOptions:`. So the `UIApplication` instance now actually sends `application:didFinishLaunchingWithOptions:` to its delegate, the `AppDelegate` instance. That is why, in some earlier examples, we put code into the implementation of that method; this is code that is guaranteed, thanks to the `UIApplication`'s contract with its delegate, to run very early in the app's lifetime. We know the `application:didFinishLaunchingWithOptions:` message will be sent, we know *when* it will be sent (early in our app's lifetime), and we know to *whom* it will be sent (the application's delegate object); so we've put appropriate code there, waiting to be called. Moreover, there is already boilerplate code in the project template's implementation of this method that does some very important things, such as making our app's window and showing it:

```

self.window = [[UIWindow alloc] initWithFrame:[[UIScreen mainScreen] bounds]];
[self.window makeKeyAndVisible];

```



The `UIApplicationDelegate` methods are also provided as notifications. This lets an instance other than the app delegate hear conveniently about application lifetime events, by registering for them. A few other classes provide duplicate events similarly; for example, `UITableView`'s delegate method `tableView:didSelectRowAtIndexPath:` is matched by a notification `UITableViewSelectionDidChangeNotification`.

By convention, many Cocoa delegate method names contain the modal verbs `should`, `will`, or `did`. A `will` message is sent to the delegate just before something happens; a

`did` message is sent to the delegate just after something happens. A `should` method is special: it returns a BOOL, and you are expected to respond with YES to permit something or NO to prevent it. The documentation tells you what the default response is; you don't have to implement a `should` method if the default is acceptable. In many cases, a property will control the overall behavior; the delegate message lets you pick and choose the behavior based on circumstances at runtime.

For example, by default the user can tap the status bar to make a scroll view scroll quickly to the top. Even if the scroll view's `scrollsToTop` property is YES, you can prevent this behavior for a particular tap by returning NO from the delegate's `scrollViewShouldScrollToTop:`.

When you're searching the documentation for how you can be notified of a certain event, be sure to consult the corresponding delegate protocol, if there is one. (And don't forget to consult the class's superclasses to see if one of *them* has a corresponding delegate protocol.) You'd like to know when the user taps in a UITextField to start editing it? You won't find anything relevant in the UITextField documentation; what you're after is `textFieldDidBeginEditing:` in the UITextFieldDelegate protocol. You want to respond when the user rearranges items on your tab bar? Look in UITabBarControllerDelegate. You want to know how to make a UITextView zoomable (through the user making a pinch gesture)? A UITextView is a UIScrollView; a scroll view is not zoomable unless its delegate returns a view from `viewForZoomingInScrollView:`, documented under UIScrollViewDelegate.

You can implement the delegation pattern yourself if you like. I already mentioned, in [Chapter 10](#) in connection with protocols, how my ColorPickerController class's header file defines a ColorPickerDelegate protocol; to guarantee (and to assure the compiler) that it will be able to send the ColorPickerDelegate method `colorPicker:didSetColorNamed:toColor:` to the appropriate object, it defines a `delegate` property typed as an `id <ColorPickerDelegate>` (I'll explain the details of the property declaration in [Chapter 12](#)):

```
@protocol ColorPickerDelegate;
@interface ColorPickerController : UIViewController
@property (nonatomic, weak) id <ColorPickerDelegate> delegate;
@end

@protocol ColorPickerDelegate
// color == nil on cancel
- (void) colorPicker:(ColorPickerController *)picker
    didSetColorNamed:(NSString *)theName
        toColor:(UIColor *)theColor;
@end
```

The `@protocol` compiler directive in the first line is a *forward reference*; it quiets the compiler by asserting that ColorPickerDelegate *is* a protocol, defined later in the same file. In the implementation for ColorPickerController, I send the protocol message to the delegate, and all is well:

```

- (void) dismissColorPicker: (id) sender { // user has tapped our Done button
    [self.delegate colorPicker:self
        didSetColorNamed:self.colorName
        toColor:self.color];
}

```

The compiler is happy with that code, because this message is defined in a protocol that the delegate adopts. I'm happy for the same reason; assuming the delegate isn't nil, it must adopt this protocol and respond to this message, and so the right things will happen (and if the delegate *is* nil, there's no harm done, because sending a message to nil does nothing). And I don't bother to check whether the delegate implements this method, because I happen to know that it does — it's a required method. (Compare Xcode's own Utility Application project template, which uses exactly the same architecture, couched in all but identical language.)

Data Sources

A *data source* is like a delegate, except that its methods supply the data for another object to display. The only Cocoa classes with data sources are UITableView and UIPickerView. A table view displays data in rows; a picker view displays selectable choices using a rotating drum metaphor. In each case, the data source must formally conform to a protocol with required methods (UITableViewDataSource and UIPickerViewDataSource, respectively).

It comes as a surprise to some beginners that a data source is necessary at all. Why isn't a table's data just part of the table? Or why isn't there at least some fixed data structure that contains the data? The reason is that such policies would violate generality. Use of a data source separates the object that displays the data from the object that manages the data, and leaves the latter free to store and obtain that data however it likes (see on model–view–controller in [Chapter 12](#)). The only requirement is that the data source must be able to supply information quickly, because it will be asked for it in real time when the data needs displaying.

Another surprise is that the data source is different from the delegate. But this again is only for generality; it's an option, not a requirement. There is no reason why the data source and the delegate should not be the same object, and most of the time they probably will be.

In this simple example, we implement a UIPickerView that allows the user to select by name a day of the week (the Gregorian week, using English day names). The first two methods are UIPickerView data source methods; the third method is a UIPickerView delegate method:

```

- (NSInteger) numberOfComponentsInPickerView: (UIPickerView*) pickerView {
    return 1;
}

- (NSInteger) pickerView: (UIPickerView*) pickerView

```

```

    numberOfRowsInComponent: (NSInteger) component {
        return 7;
    }

    - (NSString*) pickerView:(UIPickerView*)pickerView
        titleForRow:(NSInteger)row
        forComponent:(NSInteger)component {
        NSArray* arr = [NSArray arrayWithObjects:
            @"Sunday",
            @"Monday",
            @"Tuesday",
            @"Wednesday",
            @"Thursday",
            @"Friday",
            @"Saturday",
            nil];
        return [arr objectAtIndex: row];
    }
}

```

Actions

An *action* is a message emitted by an instance of a UIControl subclass (a *control*) reporting that a significant user event taking place in that control. The UIControl subclasses ([Chapter 25](#)) are all simple interface objects that the user can interact with directly, like a button (UIButton), a switch (UISwitch), a segmented control (UISegmentedControl), a slider (UISlider), or a text field (UITextField).

The significant user events (*control events*) are listed under UIControlEvents in the Constants section of the UIControl class documentation:

- UIControlEventTouchDown
- UIControlEventTouchDownRepeat
- UIControlEventTouchDragInside
- UIControlEventTouchDragOutside
- UIControlEventTouchDragEnter
- UIControlEventTouchDragExit
- UIControlEventTouchUpInside
- UIControlEventTouchUpInside
- UIControlEventTouchCancel
- UIControlEventValueChanged
- UIControlEventEditingDidBegin
- UIControlEventEditingChanged
- UIControlEventEditingDidEnd
- UIControlEventEditingDidEndOnExit

- `UIControlEventAllTouchEvents`
- `UIControlEventAllEditingEvents`
- `UIControlEventAllEvents`

The control events also have informal names that are visible in the Connections inspector when you’re editing a nib. I’ll mostly use the informal names in what follows. Control events fall roughly into three groups: the user has touched the screen (Touch Down, Touch Drag Inside, Touch Up Inside, etc.), edited text (Editing Did Begin, Editing Changed, etc.), or changed the control’s value (Value Changed).

Apple’s documentation is rather coy about which controls normally emit actions for which control events, so here’s a list obtained through experimentation. Keep in mind that Apple’s silence on this matter may mean that the details are subject to change:

UIButton

All “Touch” events.

UIDatePicker

Value Changed.

UIPageControl

All “Touch” events, Value Changed.

UISegmentedControl

Value Changed.

UISlider

All “Touch” events, Value Changed.

UISwitch

All “Touch” events, Value Changed.

UITextField

All “Touch” events except the “Up” events, and all “Editing” events. The text field is either in touch mode or in edit mode; as it switches from the former to the latter (and the keyboard appears, and Editing Did Begin is triggered), a Touch Cancel event is triggered. If the user stops editing by tapping Return in the keyboard, Did End on Exit is triggered along with Editing Did End.

The way you hear about a control event is through an action message. A control maintains an internal dispatch table: for each control event, there is some number of target–action pairs, of which the *action* is a selector (the name of a method) and the *target* is the object to which that message is to be sent. When a control event occurs, the control consults its dispatch table, finds all the target–action pairs associated with that control event, and sends each action message to the corresponding target. This architecture is reminiscent of a notification ([Figure 11-1](#)).

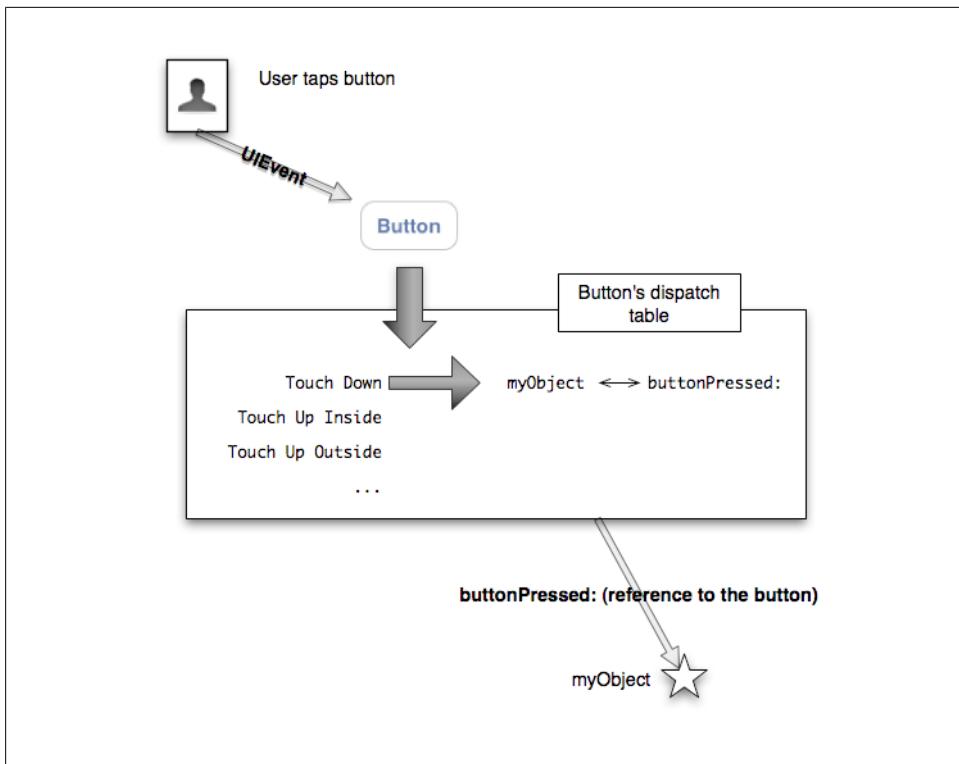


Figure 11-1. The target–action architecture



The action messaging mechanism is actually more complex than I've stated. The UIControl does not really send the action message directly; rather, it tells the shared application to send it. When a control wants to send an action message reporting a control event, it calls its own `sendAction:to:forEvent:` method. This in turn calls the shared application instance's `sendAction:to:from:forEvent:`, which actually sends the specified action message to the specified target. In theory, you could call or override either of these methods to customize this aspect of the message-sending architecture, but it is extremely unlikely that you would do so.

There are two ways to manipulate a control's action dispatch table: you can configure an action connection in a nib (as explained in [Chapter 7](#)), or you can use code. To use code, you send the control the message `addTarget:action:forControlEvents:`, where the `target:` is an object, the `action:` is a selector, and the `controlEvents:` are designated by a bitmask (see [Chapter 1](#) if you've forgotten how to construct a bitmask). Unlike a notification center, a control has methods for introspecting the dispatch table. Recall the example from [Chapter 7](#) (where `b` is a reference to a `UIButton`):

Touch Inside and Touch Outside

There is no explicit “Touch Down Inside” event, because *any* sequence of “Touch” events begins with “Touch Down,” which *must* be inside the control. (If it weren’t, this sequence of touches would not “belong” to the control, and there would be no control events at all.)

When the user taps within a control and starts dragging, the “Inside” events are triggered even after the drag moves outside the control’s bounds. But after a certain distance from the control is exceeded, an invisible boundary is crossed, Touch Drag Exit is triggered, and now “Outside” events are reported until the drag crosses back within the invisible boundary, at which point Touch Drag Enter is triggered and the “Inside” events are reported again. In the case of a UIButton, the crossing of this invisible boundary is exactly when the button automatically unhighlights (as the drag exits). Thus, to catch a legitimate button press, you probably want to consider only Touch Up Inside.

For other controls, there may be some slight complications. For example, a UISwitch will unhighlight when a drag reaches a certain distance from it, but the touch is still considered legitimate and can still change the UISwitch’s value; therefore, when the user’s finger leaves the screen, the UISwitch reports a Touch Up Inside event, even while reporting Touch Drag Outside events.

```
[b addTarget:self action:@selector(buttonPressed:)
forControlEvents:UIControlEventTouchUpInside];
```

After that, whenever the user taps in the button, our `buttonPressed:` method will be called. It might, again as in [Chapter 7](#), look like this:

```
- (void)buttonPressed: (id) sender {
    UIAlertView* av = [[UIAlertView alloc] initWithTitle:@"Howdy!"
                                                message:@"You tapped me."
                                               delegate:nil
                                              cancelButtonTitle:@"Cool"
                                              otherButtonTitles:nil];
    [av show];
}
```

The signature for the action selector can be in any of three forms. The fullest form takes two parameters:

- The control, usually typed as `id`.
- The UIEvent that generated the control event.

A shorter form (the most commonly used form) omits the second parameter; a still shorter form omits both parameters. If you’re perfectly certain what control is sending you the message, and you don’t need a reference to it, you might not bother with either of the action message’s parameters.



Mac OS X Programmer Alert

If you're an experienced Mac OS X Cocoa developer, you'll note that there are some major differences between the Mac OS X implementation of actions and the iOS implementation. In Mac OS X, a control has just one action; in iOS, a control may respond to multiple control events. In Mac OS X, an action has just one target; in iOS, a single event can trigger multiple action messages to multiple targets. In Mac OS X, an action message selector comes in just one form; in iOS, there are three possible forms.

What is the UIEvent, and what is it for? Well, a *touch event* is generated whenever the user does something with a finger (sets it down on the screen, moves it, raises it from the screen). UIEvents are the lowest-level objects charged with communication of touch events to your app. A UIEvent is basically a timestamp (a double) along with a collection (NSSet) of touch events (UITouch). The action mechanism deliberately shields you from the complexities of touch events, but by electing to receive the UIEvent, you can still deal with those complexities if you want to. (See [Chapter 18](#) for full details.)

In this example, I take advantage of the UIEvent's timestamp to do one thing if the user releases a UIButton after holding down a finger for a short time, but a different thing if the user releases the UIButton after holding it down for a longer time. Assume that the UIButton's dispatch table is configured so that its Touch Down control event calls my `buttonDown:event:` method, and its Touch Up Inside control event calls my `buttonUp:event:` method:

```
- (void) buttonDown: (id) sender event: (UIEvent*) e {
    self.downtime = [e timestamp]; // downtime is a property and ivar
}

- (void) buttonUp: (id) sender event: (UIEvent*) e {
    if ([e timestamp] - self.downtime < 0.3) {
        // respond to short tap
    } else {
        // respond to longer hold and release
    }
}
```



Curiously, none of the action selector parameters provide any way to learn *which* control event triggered the current action selector call! Thus, for example, to distinguish a Touch Up Inside control event from a Touch Up Outside control event, their corresponding target-action pairs must specify two different action handlers; if you dispatch them to the same action handler, that handler cannot discover which control event occurred.

The Responder Chain

A *responder* is an object that knows how to receive UIEvents directly (see the previous section). It knows this because it is an instance of UIResponder or a UIResponder subclass. If you examine the Cocoa class hierarchy, you'll find that just about any class that has anything to do with display on the screen is a responder. A UIView is a responder. A UIWindow is a responder. A UIViewController is a responder. Even a UIApplication is a responder. And starting in iOS 5 and the Xcode 4.2 templates, the app delegate is a responder.

If you look in the documentation for the UIResponder class, you'll find that it implements four low-level methods for receiving touch-related UIEvents: `touchesBegan:withEvent:`, `touchesMoved:withEvent:`, `touchesEnded:withEvent:` and `touchesCancelled:withEvent:`. These are called to notify a responder of a touch event. No matter how your code ultimately hears about a user-related touch event — indeed, even if your code *never* hears about a touch event (because Cocoa reacted in some automatic way to the touch, without your code's intervention) — the touch was initially communicated to a responder through one of these methods.

The mechanism for this communication starts by deciding which responder the user touched. The UIView methods `hitTest:withEvent:` and `pointInside:withEvent:` are called until the correct view (the *hit-test view*) is located. Then UIApplication's `sendEvent:` method is called, which calls UIWindow's `sendEvent:`, which calls the correct method of the hit-test view (a responder). I'll cover all this again in full detail in [Chapter 18](#).

The responders in your app participate in a *responder chain*, which essentially links them up through the view hierarchy. A UIView can sit inside another UIView, its *superview*, and so on until we reach the app's UIWindow (a UIView that has no superview). The responder chain, from bottom to top, looks like this:

1. The UIView that we start with (here, the hit-test view).
2. The UIViewController that controls that UIView, if there is one.
3. The UIView's superview, and then *its* UIViewController if there is one. Repeat this step, moving up the superview hierarchy one superview at a time, until we reach...
4. The UIWindow.
5. The UIApplication.
6. The UIApplication's delegate. (This is new in iOS 5, and is not well documented as of this writing.)

Deferring Responsibility

The responder chain can be used to let a responder defer responsibility for handling a touch event. If a responder receives a touch event and can't handle it, the event can be

passed up the responder chain to look for a responder that *can* handle it. This can happen in two main ways: (1) the responder doesn't implement the relevant method; (2) the responder implements the relevant method to call `super`.

For example, a plain vanilla `UIView` has no native implementation of the touch event methods. Thus, by default, even if a `UIView` is the hit-test view, the touch event effectively falls through the `UIView` and travels up the responder chain, looking for someone to respond to it. If this `UIView` is an instance of your own subclass, you might implement the touch event methods in that subclass to catch touch events in the `UIView` itself; but if the `UIView` is controlled by a `UIViewController`, you have already subclassed `UIViewController`, and that subclass is probably where the interface behavior logic for this `UIView` is already situated, so you might well prefer to implement the touch event methods there instead. You are thus taking advantage of the responder chain to defer responsibility for handling touch events from the `UIView` to its `UIViewController`, in a natural and completely automatic way.

Again, I'll come back to this in [Chapter 18](#); don't worry about it for now. I'm actually telling you about the responder chain chiefly in order to discuss nil-targeted actions, which is the subject of the next section.

Nil-Targeted Actions

A *nil-targeted action* is a target–action pair in which the target is `nil`. There is no designated target object, so the following rule is used: starting with the hit-test view (the view with which the user is interacting), Cocoa looks up the responder chain for an object that can respond to the action message.

Suppose, for example, we have a `UIButton` inside a `UIView`. And suppose we run this code early in the button's lifetime, where `b` is the button:

```
[b addTarget:nil action:@selector(doButton:)
    forControlEvents:UIControlEventTouchUpInside];
```

That's a nil-targeted action. So what happens when the user taps the button? First, Cocoa looks in the `UIButton` itself to see whether it responds to `doButton:`. If not, then it looks in the `UIView` that is its superview. And so on, up the responder chain. If a responder is found that handles `doButton:`, the action message is sent to that object; otherwise, the message goes unhandled.

Thus, suppose the `UIView` containing the `UIButton` is an instance of your own `UIView` subclass. Let's call it `MyView`. If `MyView` implements `doButton:`, then when the user taps the button, it is `MyView`'s `doButton:` that will be called.

To create a nil-targeted action in a nib, you form a connection to the First Responder proxy object (in the dock). This is what the First Responder proxy object is for! The First Responder isn't a real object with a known class, so before you can connect an action to it, you have to define the action message within the First Responder proxy object, like this:

The Term “First Responder”

Cocoa uses the term *first responder* in a confusing way. An arbitrary responder object can be assigned formal first responder status (by sending it `becomeFirstResponder`, provided that this responder returns YES from `canBecomeFirstResponder`). But this does not make the object first responder for purposes of handling nil-targeted actions! Cocoa’s hunt for a responder that can handle a nil-targeted action still starts with the control that the user is interacting with (the hit-test view) and goes up the responder chain from there.

Typical legitimate uses of `becomeFirstResponder` are:

- To put a UITextField into editing mode, as if the user had tapped in it ([Chapter 23](#)).
- To specify the object that should initially be sent remote events ([Chapter 27](#)).
- To designate an object that is to present a menu ([Chapter 39](#)).

1. Select the First Responder proxy in the nib, and switch to the Attributes inspector.
2. You’ll see a table (probably empty) of user-defined nil-targeted First Responder actions. Click the Plus button and give the new action a signature; it must take a single parameter (so that its name will end with a colon).
3. Now you can Control-drag from a control, such as a UIButton, to the First Responder proxy to specify a nil-targeted action with the signature you specified.

Application Lifetime Events

When you create an object, either obtaining an instance in code or instantiating it in a nib, you’ll concern yourself with the events you can arrange to receive in connection with that object. This book will do the same. We’ll talk about UIView events (including touches) when we talk about views; we’ll talk about the many lifetime events connected with UIViewController when we talk about view controllers; we’ll talk about UITextField events when we talk about text, UIControl events when we talk about controls, and so on. Indeed, it is not too much to say that the art of using Cocoa consists largely of mastering a two-way street, knowing what you can do to some built-in type of object (to configure it) and know what that object can do to you (what events you can receive in connection with it, and how to arrange to receive them).

Your application as a whole, on the other hand, is something, you never deliberately create. It was created for you as the shared application object by `UIApplicationMain` before any code of yours had a chance to run. So there won’t come any later point in this book any better for discussion of the various events you can receive in connection with your application’s lifetime than right now. As I’ve already mentioned, these events, marking certain important stages in the overall lifetime of your application, can arrive

either as messages to the app delegate (if you implement the appropriate methods) or as notifications to any object (if that object registers for those notifications).

What lifetime events you can receive depends on whether or not your app participates in *multitasking*. In the old days, before iOS 4, there was no multitasking. If the user pressed the Home button while running your app, your app was terminated. The next time the user launched your app by tapping its icon in the springboard, your app launched from scratch. Even under iOS 4 and later, your app can opt out of multitasking and behave like a pre-iOS 4 app. Also, if you set the “Application does not run in background” key (`UIApplicationExitsOnSuspend`) in your `Info.plist`. (For some apps, such as certain games, this might be a reasonable thing to do.)

The suite of events received by a nonmultitasking app is pretty simple:

`application:didFinishLaunchingWithOptions:`

The app has started up. This, as we have already seen, is the earliest opportunity for your code to configure the interface by creating showing the app’s window, and it’s fair to assume that every app you ever write will do exactly that. If you don’t show your app’s interface, the user won’t see anything! (The exceptional case, in iOS 5 and the Xcode 4.2 templates, is an app with a main storyboard; by default, the storyboard creates and shows the window for you as it loads.) Of course you can and will perform other early initializations here.

Alternatively, your app delegate can implement `application:didFinishLaunching:` instead, but there are no advantages to doing so, and there are many disadvantages — your app would miss out on the options dictionary, which can contain useful information, as we’ll see later in the book.

`applicationDidBecomeActive:`

The app has started up; received after `application:didFinishLaunchingWithOptions:`. Also received after the end of the situation that caused the app delegate to receive `applicationWillResignActive:`.

`applicationWillResignActive:`

Something has blocked the app’s interface. The most common cause is that the screen has been locked. An alert from outside your app could also cause this event. When this situation ends, the app delegate will receive `applicationDidBecomeActive:`.

`applicationWillTerminate:`

The app is about to quit. This is your last signal to preserve state (typically, by storing information with `NSUserDefaults`) and perform other final cleanup tasks.

In the multitasking environment of iOS 4 and later, however (unless you deliberately opt out, as I’ve already mentioned), the Home button doesn’t terminate your app; it backgrounds and suspends it. This means that your app is essentially freeze-dried in the background; its process still exists, but it isn’t getting any events (though notifications can be stored by the system for later delivery if your app comes to the front once

again). You'll probably *never* get `applicationWillTerminate:`, because when your app is terminated by the system, it will already have been suspended.



Under highly specialized circumstances (discussed, for instance, in [Chapter 27](#) and [Chapter 35](#)), your app can be backgrounded without being suspended. Nevertheless, throughout this section I'll speak as if backgrounding and suspension are one and the same.

Thus, you have to worry about what will happen when the app is suspended and when it returns from being suspended (`applicationDidEnterBackground:` and `applicationWillEnterForeground:`, and their corresponding notifications), and the notion of the application becoming inactive or active also takes on increased importance (`applicationWillResignActive:` and `applicationDidBecomeActive:`, and their notifications). These notifications all take on a wide range of meaning—indeed, in my opinion, the information your app is given is unfortunately too coarse-grained — so they are best understood by examining some typical scenarios:

The app launches freshly

Your app delegate receives these messages (just as in the premultitasking world):

- `application:didFinishLaunchingWithOptions:`
- `applicationDidBecomeActive:`

The user clicks the Home button

If your app is frontmost, it is suspended, and your app delegate receives these messages:

- `applicationWillResignActive:`
- `applicationDidEnterBackground:`

The user summons your suspended app to the front

Your app delegate receives these messages:

- `applicationWillEnterForeground:`
- `applicationDidBecomeActive:`

The user double-clicks the Home button

The user can now work in the app switcher. If your app is frontmost, your app delegate receives this message:

- `applicationWillResignActive:`

The user, in the app switcher, taps on your app's window

Your app delegate receives this message:

- `applicationDidBecomeActive:`

The user, in the app switcher, chooses another app

If your app is frontmost, your app delegate receives this message:

- `applicationDidEnterBackground:`

The screen is locked

If your app is frontmost, your app delegate receives this message:

- `applicationWillResignActive:`

The screen is unlocked

If your app is frontmost, your app delegate receives this message:

- `applicationDidBecomeActive:`

The user holds the screen-lock button down

The device offers to shut itself down. If your app is frontmost, your app delegate receives this message:

- `applicationWillResignActive:`

The user, as the device offers to shut itself down, cancels

If your app is frontmost, your app delegate receives this message:

- `applicationDidBecomeActive:`

The user, as the device offers to shut itself down, accepts

If your app is frontmost, your app delegate receives these messages:

- `applicationDidEnterBackground:`

- `applicationWillTerminate:` (probably the only way a normal app will receive this message in a multitasking world)

Juggling these events and meeting your responsibilities under all circumstances can be quite tricky. You have, as it were, a lot of bases to cover if you want to make certain that your app is in a known state as it passes through each of its possible lifetime stages; and, as I've already suggested, you aren't really given enough information to cover them coherently. To give an obvious example, `applicationWillResignActive:` could mean that the user is summoning the app switcher or that the user is locking the screen, or that the user has single-clicked the Home button to leave your app altogether (in which case you are about to receive `applicationDidEnterBackground:` as well). This makes it hard (not to say impossible) for you to do different things in response to these three different user actions.

In most cases, however, you can make the right thing happen at the right time, provided you are careful and conscious of the various scenarios I've just outlined. For example, let's say there are things my app wants to do when the user resumes it from a suspended state (`applicationWillEnterForeground:`), and there are things my app wants to do when the user unlocks the screen (`applicationDidBecomeActive:`). Some of these things are the same in both circumstances, but others are not. But the former event is always followed by the latter. So I have to be careful how I respond:

`applicationWillEnterForeground:`

Do things appropriate to resuming from suspension. But do *not* do things that are also appropriate to the user unlocking the screen, because `applicationDidBecomeActive:` is about to be called, so if I do those things here, I'll end up doing them *twice*.

`applicationDidBecomeActive:`

Do things appropriate both when resuming from suspension and when the user unlocks the screen.

Fine, but what about things that are appropriate when the user unlocks the screen but *not* when we resume from suspension? Apparently, Apple doesn't think there are any such things, since they haven't given you a sufficiently fine-grained repertoire of events to implement them. You can't simply put such things in `applicationDidBecomeActive:`, which is called when the user unlocks the screen, because it is *also* called when resuming from suspension. A possible workaround is to do a little dance with a flag (a BOOL instance variable) to detect when the app becomes active without having been suspended:

```
- (void)activeButNotForeground {
    // do things for when the screen unlocks
}

- (void)applicationWillEnterForeground:(UIApplication *)application {
    // do things for when we resume from suspension
    self->didForeground = YES;
}

- (void)applicationDidBecomeActive:(UIApplication *)application {
    // do things for when we resume from suspension or the screen unlocks
    if (!self->didForeground)
        [self activeButNotForeground];
    self->didForeground = NO;
}
```

Of all the tasks with which you'll concern yourself over your app's lifetime, probably the most important is saving state. If the user has done or specified things that the app needs to preserve, it must do so before it is terminated; since a multitasking app is probably never going to be notified by `applicationWillTerminate:`, you'll need to find an appropriate earlier moment. Moreover, it is a mark of most good iOS apps that they try to present themselves to the user on launch looking as they did when they were last terminated. This is easy, of course, if the app wasn't really terminated but only suspended; the app comes back to life automatically in exactly the same state as when it was suspended. But if the app is terminated while suspended, it will launch from scratch the next time, and will have to use its own internal logic to determine that it was in a particular state when it was suspended and to restore that state. An obvious example is a board game, such as chess; if the user was in the middle of a game when the app was terminated, the board should reappear on launch just as it was, with the same game history and with the same player about to move. Clearly, at some moment, all the

information about what pieces are on the board and whose turn it is needs to be saved. But what moment?

One obvious answer is: at every moment. In a game such as chess, there might be no reason why the app shouldn't save state after every move. And this is certainly part of the iOS state-saving philosophy: always be saving state. But some state might be too large or too convoluted to save repeatedly, or the app might not consist of discrete simple states the way a board game does. So what's the last best moment to save state? Looking over the list of scenarios, you can see that `applicationWillResignActive`: is the broadest signal that *something* might now happen such that your app will subsequently be terminated without further notice; but it is also possible that your app will next become active again. On the other hand, it is very unlikely that your app will be terminated without your first receiving `applicationDidEnterBackground`:, and this might well be the last event your app ever receives, so this is the default signal that you should save state. (You are given a little time to do this before your app is actually suspended; [Chapter 38](#) discusses what to do if you think a little time might not enough.)



Your app stands a better chance of not being terminated while suspended, the less memory resources it uses; you should therefore take `applicationDidEnterBackground`: as a signal to release whatever large memory resources you can. For example, you might write such resources out to disk, and recover them when your app receives `applicationWillEnterForeground`:. In addition, while running in the foreground, your app will be sent the delegate message `applicationDidReceiveMemoryWarning`: to inform you of a low memory situation. I'll return to this topic in [Chapter 19](#).

Swamped by Events

In your mind, assemble everything I've said in this chapter about events, and you can readily imagine that, as you develop your own iOS app, you're likely to wind up with classes consisting of numerous methods, the vast majority of which are *entry points* — methods that you have placed there so that they can be called by Cocoa through events. For example, suppose (as is extremely likely) that your iPhone app contains a screen that effectively consists entirely of a table view. You'll probably have a corresponding `UITableViewController` subclass; `UITableViewController` is a built-in `UIViewController` subclass, plus you'll probably use this same class as the table view's data source and delegate. I'll talk in detail about these matters in [Chapter 19](#) and [Chapter 21](#), but for now, just believe when I say that you're extremely likely, in this single class, to implement *at a minimum* such methods as these:

`initWithCoder`: or `initWithNibName:bundle`:

`UIViewController` lifetime method, where you perform custom instance initializations.

viewDidLoad:

UIViewController lifetime method, where you perform view-related initializations.

viewDidUnload:

UIViewController lifetime method, where you release any memory resources acquired in `viewDidLoad:`.

viewWillAppear:

UIViewController lifetime method, where you set up states that need to apply only while your view is onscreen. For example, if you’re going to register for a notification or set up a timer, this is a likely place to do it.

viewWillDisappear:

UIViewController lifetime method, where you reverse what you did in `viewWillAppear:`. For example, this would be a likely place to unregister for a notification or invalidate a repeating timer that you set up in `viewWillAppear:`.

shouldAutorotateToInterfaceOrientation:

UIViewController query method, where you specify what device orientations are allowed for this view.

numberOfSectionsInTableView:

UITableView data source query method, where you specify the contents of the table.

tableView:numberOfRowsInSection:

UITableView data source query method, where you specify the contents of the table.

tableView:cellForRowAtIndexPath:

UITableView data source query method, where you specify the contents of the table.

tableView:didSelectRowAtIndexPath:

UITableView delegate user action method, where you respond when the user taps a row of the table.

dealloc

NSObject lifetime method, where you perform end-of-life cleanup; if, for example, you registered for a notification or set up a timer in `viewWillAppear:`, you will unregister for the notification and invalidate a repeating timer here, as a backstop in case `viewWillDisappear:` isn’t called. (Also, if you’re not using ARC, you will have memory management tasks to perform here.)

Suppose, further, that you did in fact use `viewWillAppear:` to register for a notification and to set up a timer. Then that notification has a selector (unless you used a block), and the timer has a selector; you must implement the methods described by those selectors.

We already have, then, a dozen methods whose presence is effectively a matter of boilerplate. These are not *your* methods; *you* are never going to call them. They are *Cocoa's* methods, which you have placed here so that each can be called at the appropriate moment in the life story of your app. The logic of a program laid out in this fashion is by no means easy to understand! I'm not criticizing Cocoa here — indeed, it's hard to imagine how else an application framework could work — but, purely as an objective matter of fact, the result is that a Cocoa program, even your own program, even *while you're developing it*, is hard to read, because it consists of numerous disconnected entry points, each with its own meaning, each called at its own set moment which is not in any way obvious from looking at the program. To understand what our hypothetical class does, you have to know *already* such things as when `viewWillAppear:` is called and how it is typically used; otherwise, you don't even know where to look to find the program's logic and behavior, let alone how to interpret what you see when you do look there. And this difficulty is greatly compounded when you try to read someone else's code (this is one reason why, as I mentioned in [Chapter 8](#), sample code is not all that helpful to a beginner).

Looking at the code of an iOS program — even your own code — your eyes can easily glaze over at the sight of all these methods called automatically by Cocoa under various circumstances. To be sure, experience will teach you about such things as the overridden `UIViewController` methods, which form a standard sequence of moments in the lifetime of the view controller, and about the table view delegate and data source methods, which — although they are called at unpredictable moments — work together in a reliable and fairly simple way to describe the table and govern its behavior. On the other hand, no amount of experience will tell you that a certain method is called as a button's action or through a notification. Under these circumstances, comments really help, and I strongly advise you, as you develop any iOS app, to comment every method, quite heavily if need be, saying what that method does and under what circumstances you expect it to be called — especially if it is an entry point, where it is Cocoa itself that will do the calling.

In addition, since Cocoa is doing the calling and since you have no access to the Cocoa source code, the actual moment when entry points are called can be unpredictable. Your own code can trigger unintended events. The documentation might not make it clear just when a notification will be sent. There could even be a bug in Cocoa such that events are called in a way that seems to contradict the documentation. Therefore I also recommend also that as you develop your app, you instrument your code heavily with caveman debugging (`NSLog`; see [Chapter 9](#)). As you test your code, keep an eye on the console output and check whether the messages make sense. You may be surprised at what you discover.

For example, it has occasionally happened to me that Cocoa will mysteriously call an event too often — and without logging, I would never have been aware of this fact. For example, I mentioned earlier in this chapter an app of mine that registers for the `MPMusicPlayerControllerNowPlayingItemDidChangeNotification`, so that it can change its in-

terface in response when the device’s music player starts playing a different song. Unfortunately, it very often (though not always!) happens that this notification is delivered *twice* in quick succession when there’s a transition from one song to another. I would never have known this if I hadn’t been logging in my code triggered by the arrival of this notification.

You may be wondering what I did when I discovered this issue. My first response to wonder whether this might be a bug in Cocoa; but I quickly realized that such worries were counterproductive. The cause of the issue isn’t important, because there’s nothing I can do about it; the framework is in charge, and the framework is opaque. All that matters is how to guard against it. I then asked myself whether this issue was worth guarding against. The notification’s registered selector was my method `reloadData:`. If what I do in `reloadData:` isn’t expensive or sequential in nature, it might not matter if I do it twice in quick succession. It turns out that what I do in `reloadData:` *is* expensive. What, then, could I do to prevent my expensive code from running twice in quick succession? Two possible approaches presented themselves:

- Rearchitect my code so that what I do in `reloadData:` is no longer expensive. Unfortunately, at the time, I couldn’t find a convenient way to do that; I would effectively have to unwind the entire operation of this class and approach my interface in a completely different way.
- Use delayed performance (see [Chapter 10](#), on `NSObject`) in such a way that if `reloadData:` is called twice within a sufficiently short span of time, the first call is effectively thrown away. This is the solution to which I ultimately resorted:

```
- (void) reloadData: (NSNotification*) n {
    [UIApplication cancelPreviousPerformRequestsWithTarget:self
                                                selector:
                                                 @selector(reallyReloadData)
                                                object:nil];
    [self performSelector:@selector(reallyReloadData)
                   withObject:nil
                     afterDelay:0.3];
}

- (void) reallyReloadData {
    // do expensive thing here
}
```

In addition, your own code can indirectly trigger events in an unexpected way. This happened to me in developing a different app, where I suddenly found that in a `UIViewController` subclass, `viewDidLoad` was being called twice as the app started up, which should be impossible. Again, I wouldn’t even have noticed this mistake without `NSLog`. Not only that, but — as I discovered after adding even more `NSLog` messages — `viewDidLoad` was called while I was still in the middle of executing `awakeFromNib`, which should *really* be impossible. The reason turned out to be that I was making the mistake of mentioning my class’s `view` property during `awakeFromNib`; this actually causes `viewDidLoad` to be called. The problem went away when I corrected my mistake.

There is, in sum, a certain indeterminacy and uncertainty in the nature of events. They bombard your code; they control it. Their mutual relationship can be tricky, and in most cases you shouldn't rely on things happening in a definite order (although you probably can if the documentation says so). Your posture will therefore be, to some extent, one of uncertainty and defensiveness. You're like a tennis student being pelted by a ball-serving machine. Without comments, the tennis student doesn't know where the balls are coming from. Without logging, the tennis student is blindfolded.

Accessors and Memory Management

Even when you've understood how to slot in your code so as to get the messages that you want from Cocoa ([Chapter 11](#)), your obligations to Cocoa and your interactions with the framework are not over. You have additional responsibilities that emerge as you write the code for any class that will be instantiated. There are guidelines for how a well-behaved instance should be structured and how it should act—in fact, I seriously thought of calling this chapter “The Well-Behaved Instance.” If you don’t follow those guidelines, things can go wrong: outlets aren’t set, the wrong methods are called, memory gets used up, your app crashes. This chapter is about those guidelines.

Accessors

An *accessor* is a method for getting or setting the value of an instance variable. An accessor that gets the instance variable’s value is called a *getter*; an accessor that sets the instance variable’s value is called a *setter*.

There are naming conventions for accessors, and you should obey them. The conventions are simple:

- A setter’s name should start with `set` and be followed by a capitalized version of the instance variable’s name. If the instance variable is named `myVar`, the setter should be named `setMyVar:`. The setter should take one parameter: the new value to be assigned to the instance variable.
- A getter should have the same name as the instance variable. If the instance variable is named `myVar`, the getter should be named `myVar`. (This will not cause you or the compiler any confusion, because variable names and method names are used in completely different contexts.)

You can optionally start the getter’s name with `get` (e.g. `getMyVar`), though in fact I never do this. If the instance variable’s value is a BOOL, you may optionally start the getter’s name with `is` (for example, an ivar `showing` can have a getter `isShowing`), though in fact I never do this.

Accessors are important in part because instance variables are protected, whereas declared methods are public; without accessor methods, a protected instance variable can't be accessed by any object whose class (or superclass) isn't the one that declares the instance variable. You might be tempted to conclude from this that you needn't bother making an accessor for an instance variable that isn't intended for public access, and to some extent this is a reasonable conclusion (especially if your app uses ARC). But here are some counterarguments:

- Making accessors isn't really all that much bother, especially because you usually won't have to write the code for an accessor; Objective-C can write the code for you (more about that later in this chapter).
- If your app doesn't use ARC, and if your instance variable is an object, there are going to be memory management tasks to worry about every time you get and (especially) set that value; the best way to ensure that you're carrying out those tasks reliably and consistently is to pass through an accessor.
- Even your app does use ARC, there may be additional tasks that need to be performed every time the instance variable's value is touched; an accessor, acting as a gateway to the instance variable, ensures that these tasks are performed consistently.
- Cocoa often uses the string name of an instance variable to derive the name of the accessor and call it if it exists. (This conversion is called key-value coding, and is the subject of the next section.) Under these circumstances, there needs to be an accessor so that Cocoa can find it.

Outlets, in particular, use key-value coding to match the name of an outlet (in the nib) with the name of an instance variable (in your code). This means that your code should have accessors and that their names should obey the conventions; otherwise, key-value coding can't find the accessors. Suppose you have a class with an instance variable called `myVar` and you've drawn a `myVar` outlet from that class's representative in the nib to a Thing nib object. When the nib loads, the outlet name `myVar` is translated to the method name `setMyVar:`, and your instance's `setMyVar:` method, if it exists, is called with the Thing instance as its parameter, thus setting the value of your instance variable to the Thing (Figure 7-5). (I told you in Chapter 7 that there was more to the name-matching mechanism than I was letting on!)

It is important, therefore, to use accessor naming conventions correctly and consistently. By the same token, you should *not* use accessor names for methods that aren't accessors! For example, you probably would *not* want to have a method called `setMyVar:` if it is *not* the accessor for the `myVar` instance variable. If you did have such a method, it would be called when the nib loads, the Thing instance would be passed to it, and the Thing instance would *not* be assigned to the `myVar` instance variable! As a result, references in your code to `myVar` would be references to nil.

This example is not at all far-fetched; I very often see beginners complain that they are telling some part of their interface to do something and it isn't doing it. This is frequently

because they are accessing the object through an instance variable that is still nil, because it was never set properly through an outlet when the nib loaded, because they misused the name of the setter for some other purpose. (Of course it could also be because they forgot to draw the outlet in the nib in the first place.)



Although I keep saying that the names of the accessor methods use the name of the instance variable, there is no law requiring that they use the name of a *real* instance variable. Quite the contrary: you might deliberately have methods `myVar` and `setMyVar:` when in fact there is no `myVar` instance variable. Perhaps the accessors are masking the real name of the instance variable, or perhaps there is no instance variable at all, and these accessors are really doing something quite different behind the scenes. That, indeed, is one of the main reasons for using accessors; they effectively present a façade, as if there were a certain instance variable, shielding the caller from any knowledge of the underlying details.

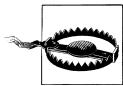
Key–Value Coding

The way Cocoa derives the name of an accessor from the name of an instance variable is through a mechanism called *key–value coding*, or simply *KVC*. (See also [Chapter 5](#), where I introduced key–value coding.) A *key* is a string (an `NSString`) that names the value to be accessed. The basis for key–value coding is the `NSKeyValueCoding` protocol, an informal protocol (it is actually a category) to which `NSObject` (and therefore every object) conforms.

The fundamental key–value coding methods are `valueForKey:` and `setValue:forKey:`. When one of these methods is called on an object, the object is introspected. In simplified terms, first the appropriate accessor is sought; if it doesn't exist, the instance variable is accessed directly. So, for example, suppose the call is this:

```
[myObject setValue:@"Hello" forKey:@"greeting"];
```

First, a method `setGreeting:` is sought in `myObject`; if it exists, it is called, passing `@"Hello"` as its argument. If that fails, but if `myObject` has an instance variable called `greeting`, the value `@"Hello"` is assigned directly to `myObject`'s `greeting` ivar.



The key–value coding mechanism can bypass completely the privacy of an instance variable! Cocoa knows that you might not want to allow that, so a class method `accessInstanceVariablesDirectly` is supplied, which you can override to return `NO` (the default is `YES`).

Both `valueForKey:` and `setValue:forKey:` require an object as the value. Your accessor's signature (or, if there is no accessor, the instance variable itself) might not use an object as the value, so the key–value coding mechanism converts for you. Numeric types (including `BOOL`) are expressed as an `NSNumber`; other types (such as `CGRect` and `CGPoint`) are expressed as an `NSValue`.



A class is *key-value coding compliant* on a given key if it implements the methods, or possesses the instance variable, required for access via that key.

Another useful pair of methods is `dictionaryWithValuesForKeys:` and `setValuesForKeysWithDictionary:`, which allow you to get and set multiple key–value pairs by way of an `NSDictionary` with a single command.

KVC is extremely dynamic. It allows you, in effect, to decide at runtime what instance variable to access; you obtain the instance variable’s name as an `NSString` and pass that to `valueForKey:` or `setValue:forKey:`. Thus, by using an `NSString` instead of an instance variable or method name, you’re throwing away compile-time checking as to the message you’re sending. Moreover, key–value coding is agnostic about the actual class of the object you’re talking to; you can send `valueForKey:` to *any* object and successfully get a result, provided the class of that object is key–value coding compliant for that key, so you’re throwing away compile-time checking as to the object you’re sending the message to. These are both strong advantages of key–value coding, and I often find myself using it because of them.

Here’s an example of key–value coding used in my own code on my own object. In a flashcard app, I have a class `Term`, representing a Latin term, that defines many instance variables. If the user taps any of three text fields, I want the interface to change from the term that’s current showing to the next term whose value is different for that text field. Thus this code is the same for all three text fields; the only difference is what instance variable to consider as we hunt for the term to be displayed. By far the simplest way to express this is through key–value coding (and, as I’ll point out later in this chapter, by defining properties for these instance variables, I get key–value coding compliance for free):

```
NSInteger tag = g.view.tag; // the tag tells us what text field was tapped
NSString* key = @"lesson";
switch (tag) {
    case 2: key = @"lessonSection"; break;
    case 3: key = @"lessonSectionPartFirstWord"; break;
}
// get current value for this instance variable
NSString* curValue = [[self currentCardController].term valueForKey: key];
// ... and so on
```

On the other hand, an attempt to access a nonexistent key through key–value coding will result, by default, in a crash at runtime, with an error message of this form: “This class is not key value coding-compliant for the key `myKey`.” The lack of quotation marks around the word after “the key” has misled many a beginner, so remember: the last word in that error message is the name of the key that gave Cocoa trouble. A common way to encounter this error message is to change the name of an instance variable so that the name of an outlet in a nib no longer matches it; at runtime, when the nib loads, Cocoa will attempt to use key–value coding to set a value in your object based on the

name of the outlet, will fail (because there is no longer an instance variable or accessor by that name), and will generate this error.

A number of built-in Cocoa classes permit you to use key–value coding in a special way. For example:

- If you send `valueForKey:` to an `NSArray`, it sends `valueForKey:` to each of its elements and returns a new array consisting of the results, an elegant shorthand (and a kind of poor man’s `map`). `NSSet` behaves similarly.
- `NSSortDescriptor` sorts an `NSArray` by sending `valueForKey:` to each of its elements.
- `NSDictionary` implements `valueForKey:` as an alternative to `objectForKey:` (useful particularly if you have an array of dictionaries); so does `NSUserDefaults`.
- `CALayer` ([Chapter 16](#)) and `CAAnimation` ([Chapter 17](#)) permit you to use key–value coding to define and retrieve the values for arbitrary keys, as if they were a kind of dictionary; this is useful for attaching identifying and configuration information to one of these instances.

As we have seen, Cocoa can access your instances via key–value coding (as it does when a nib containing an outlet is loaded).

There is also something called a *key path* that allows you to chain keys in a single expression. If an object is key–value coding compliant for a certain key, and if the value of that key is itself an object that is key–value coding compliant for another key, you can chain those keys by calling `valueForKeyPath:` and `setValue:forKeyPath:`. A key path string looks like a succession of key names joined with a dot (.). For example, `valueForKeyPath:@"key1.key2"` effectively calls `valueForKey:` on the message receiver, with `@"key1"` as the key, and then takes the object returned from that call and calls `valueForKey:` on that object, with `@"key2"` as the key.

To illustrate this shorthand, imagine that our object `myObject` has an instance variable `theData` which is an array of dictionaries such that each dictionary has a `name` key and a `description` key. I’ll show you the actual value of `theData` as displayed by `NSLog`:

```
(  
 {  
     description = "The one with glasses.>";  
     name = Manny;  
 },  
 {  
     description = "Looks a little like Governor Dewey.>";  
     name = Moe;  
 },  
 {  
     description = "The one without a mustache.>";  
     name = Jack;  
 } )
```

Then `[myObject valueForKeyPath: @"theData.name"]` returns an array consisting of the strings @"Manny", @"Moe", and @"Jack". If you don't understand why, review what I said a few paragraphs ago about how NSArray and NSDictionary implement `valueForKey:`.

Another feature of key–value coding is that it allows an object to implement a key as if its value were an array (or a set), even if it isn't. This is similar to what I said earlier about how accessors function as a façade, putting an instance variable name in front of hidden complexities. To illustrate, I'll add these methods to the class of our object `myObject`:

```
- (NSUInteger) countOfPepBoys {
    return [self.theData count];
}

- (id) objectInPepBoysAtIndex: (NSUInteger) ix {
    return [self.theData objectAtIndex: ix];
}
```

By implementing `countOf...` and `objectIn...AtIndex:`, I'm telling the key–value coding system to act as if the given key (@"pepBoys" in this case) existed and were an array. An attempt to fetch the value of the key @"pepBoys" by way of key–value coding will succeed, and will return an object that can be treated as an array, though in fact it is a proxy object (an `NSArray`). Thus we can now say `[myObject valueForKey: @"pepBoys"]` to obtain this array proxy, and we can say `[myObject valueForKeyPath: @"pepBoys.name"]` to get the same array of strings as before. This particular example may seem a little silly because the underlying implementation is already an array instance variable, but you can imagine an implementation whereby the result of `objectInPepBoysAtIndex:` is obtained through some completely different sort of operation.

The proxy object returned through this sort of façade behaves like an `NSArray`, not like an `NSMutableArray`. If you want the caller to be able to manipulate the proxy object provided by a KVC façade as if it were a mutable array, you must implement two more methods, and you must obtain a different proxy object by calling `mutableArrayValueForKey:`. So, for example:

```
- (void) insertObject: (id) val inPepBoysAtIndex: (NSUInteger) ix {
    [self.theData insertObject:val atIndex:ix];
}

- (void) removeObjectFromPepBoysAtIndex: (NSUInteger) ix {
    [self.theData removeObjectAtIndex: ix];
}
```

Now you can call `[myObject mutableArrayValueForKey: @"pepBoys"]` to obtain something that acts like a mutable array. (The true usefulness of `mutableArrayValueForKey:`, however, will be clearer when we talk about key–value observing, later on.)

A complication for the programmer is that none of these method names can be looked up directly in the documentation, because they involve key names that are specific to your object. You can't find out from the documentation what `removeObjectFromPep-`

`BoysAtIndex`: is for; you have to know, in some other way, that it is part of the implementation of key–value coding compliance for a key @"`pepBoys`" that can be obtained as a mutable array. Be sure to comment your code so that you'll be able to understand it later. Another complication, of course, is that getting a method name wrong can cause your object *not* to be key–value coding compliant. Figuring out why things aren't working as expected in a case like that can be tricky.

There is much more to key–value coding; see the *Key-Value Coding Programming Guide* for full information.

Memory Management

It comes as a surprise to many beginning Cocoa coders that the programmer has an important role to play in the explicit management of memory. What's more, managing memory incorrectly is probably the most frequent cause of crashes — or, inversely, of memory leakage, whereby your app's use of memory increases relentlessly until, in the worst-case scenario, there's no memory left.

Under iOS 5, if your app uses ARC, your explicit memory management responsibilities can be greatly reduced, which is a tremendous relief, as you are far less likely to make a mistake, and more of your time is liberated to concentrate on what your app actually does instead of dealing with memory management concerns; but even with ARC it is still possible to make a memory management mistake (I speak from personal experience), so you still need to understand Cocoa memory management, so that you know what ARC is doing for you, and so that you know how to interface with ARC in situations where it needs your assistance. Do not, therefore, suppose that you don't need to read this section on the grounds that you're going to be using ARC.

Principles of Cocoa Memory Management

The reason why memory must be managed at all is that object references are pointers. As I explained in [Chapter 1](#), the pointers themselves are simple C values (basically they are just integers) and are managed automatically, whereas what an object pointer points to is a hunk of memory that must explicitly be set aside when the object is brought into existence and that must explicitly be freed up when the object goes out of existence. We already know how the memory is set aside — that is what `alloc` does. But how is this memory to be freed up, and when should it happen?

At the very least, an object should certainly go out of existence when no other objects exist that have a pointer to it. An object without a pointer to it is useless; it is occupying memory, but no other object has, or can ever get, a reference to it. This is a *memory leak*. Many computer languages solve this problem through a policy called *garbage collection*. Simply put, the language prevents memory leaks by periodically sweeping through a central list of all objects and destroying those to which no pointer exists. But affixing a form of garbage collection to Objective-C would be an inappropriately ex-

pensive strategy on an iOS device, where memory is strictly limited and the processor is relatively slow (and may have only a single core). Thus, memory in iOS must be managed more or less manually.

But manual memory management is no piece of cake, because an object must go out of existence neither too late nor too soon. Suppose we endow the language with the ability for one object to command that another object go out of existence now, this instant. But multiple objects can have a pointer (a reference) to the very same object. If both the object Manny and the object Moe have a pointer to the object Jack, and if Manny tells Jack to go out of existence now, poor old Moe is left with a pointer to nothing (or worse, to garbage). A pointer whose object has been destroyed behind the pointer's back is a *dangling pointer*. If Moe subsequently uses that dangling pointer to send a message to the object that it thinks is there, the app will crash.

To prevent both dangling pointers and memory leakage, Objective-C and Cocoa implement a policy of manual memory management based on a number, maintained by every object, called its *retain count*. Other objects can increment or decrement an object's retain count. As long as an object's retain count is positive, the object will persist. No object has the direct power to tell another object to be destroyed; rather, as soon as an object's retain count is decremented to zero, it is destroyed automatically.

By this policy, every object that needs Jack to persist should increment Jack's retain count, and should decrement it once again when it no longer needs Jack to persist. As long as all objects are well-behaved in accordance with this policy, the problem of manual memory management is effectively solved:

- There cannot be any dangling pointers, because any object that has a pointer to Jack has incremented Jack's retain count, thus ensuring that Jack persists.
- There cannot be any memory leaks, because any object that no longer needs Jack decrements Jack's retain count, thus ensuring that eventually Jack will go out of existence (when the retain count reaches zero, indicating that no object needs Jack any longer).

Obviously, all of this depends upon all objects cooperating in obedience to this memory management policy. Cocoa's objects (objects that are instances of built-in Cocoa classes) are well-behaved in this regard, but *you* must make sure *your* objects are well-behaved. Before ARC, ensuring that your objects were well-behaved was entirely up to you and your explicit code; under ARC, your objects will be well-behaved more or less automatically, provided you understand how to cooperate with ARC's automated behavior.

Debugging Memory Management Mistakes

Memory management mistakes are among the most common pitfalls for beginners and even for experienced Cocoa programmers. Though far less likely to occur under ARC, they still *can* occur under ARC, especially because a programmer using ARC is prone to suppose (wrongly) that they can't. What experience really teaches is to use every tool at your disposal to ferret them out. Here are some:

- The static analyzer (Product → Analyze) knows a lot about memory management and can help call potential memory management mistakes to your attention.
- Instruments has excellent tools for noticing leaks and tracking memory management of individual objects (Product → Profile).
- Good old caveman debugging ([Chapter 9](#)) can help confirm that your objects are behaving as you want them to. I recently discovered that one of my apps had a memory leak by implementing `dealloc` (the `NSObject` method that is called as an object goes out of existence) with an `NSLog` call in one of its objects and finding that `dealloc` was never being called, even when I thought it should be — the object was leaking. Neither Instruments nor the static analyzer alerted me to the problem, and this app was using ARC, so I had been assuming (wrongly) that no such problem could arise in the first place.
- Dangling pointers are particularly difficult to track down, but they can often be located by “turning on zombies.” This is easy in Instruments with the Zombies template, but unfortunately it doesn’t work on a device. For a device, edit the Run action in your scheme, switch to the Diagnostics tab, and check Enable Zombie Objects. The result is that no object ever goes out of existence; instead, it is replaced by a “zombie” that will report to the console if a message is sent to it (“message sent to deallocated instance”). Be sure to turn zombies back off when you’ve finished tracking down your dangling pointers.

The Golden Rules of Memory Management

An object is well-behaved with respect to memory management as long as it adheres to certain very simple rules in conformity with the basic concepts of memory management outlined in the previous section.

Before I tell you the rules, it may help if I remind you (because this is confusing to beginners) that a variable name, including an instance variable, is just a pointer. When you send a message to that pointer, you are really sending a message *through* that pointer, to the object to which it points. The rules for memory management are rules about objects, not names, references, or pointers. You cannot increment or decrement the retain count of a pointer; there is no such thing. The memory occupied by the pointer is managed automatically (and is tiny). Memory management is concerned with the object to which the pointer points.

(That is why I've referred to my example objects by proper names — Manny, Moe, and Jack — and not by variable names. The question of who has retained Jack has nothing to do with the *name* by which any particular object *refers* to Jack.)

The two things are easily confused, especially because — as I've often pointed out in earlier chapters — the variable name pointing to an object is so often treated as the object that there is a tendency to think that it *is* the object, and to speak as if it were the object. It's clumsy, in fact, to separate talking about the name from talking about the object it points to. But in talking about memory management, I'll try to make that separation, for clarity and correctness, and to prevent confusion.

Here, then, are the golden rules of Cocoa memory management:

- To increment the retain count of any object, send it the `retain` message. This is called *retaining* the object. The object is now guaranteed to persist at least until its retain count is decremented once more. To make this a little more convenient, a `retain` call returns as its value the retained object — that is, `[myObject retain]` returns the object pointed to by `myObject`, but with its retain count incremented.
- When you (meaning a certain object) say `alloc` to a class, the resulting instance comes into the world with its retain count already incremented. You do *not* need to retain an object you've just instantiated by saying `alloc` (and you should not). Similarly, when you say `copy` to an instance, the resulting new object (the copy) comes into the world with its retain count already incremented. You do *not* need to retain an object you've just instantiated by saying `copy` (and you should not).
- To decrement the retain count of any object, send it the `release` message. This is called *releasing* the object. If you (meaning a certain object) obtained an object by saying `alloc` or `copy`, or if you said `retain` to an object, you (meaning the same object) *must* balance this eventually by saying `release` to that object, *once*. You should assume that thereafter the object may no longer exist.

A general way of understanding the golden rules of Cocoa memory management is to think in terms of *ownership*. If Manny has said `alloc`, `retain`, or `copy` with regard to Jack, Manny has asserted ownership of Jack. More than one object can own Jack at once, but each such object is responsible only for managing its own ownership of Jack correctly. It is the responsibility of an owner of Jack eventually to release Jack, and a nonowner of Jack must *never* release Jack. As long as all objects that ever take ownership of Jack behave this way, Jack will not leak nor will any pointer to Jack be left dangling.

Now, under ARC, as I shall explain presently in more detail, these rules remain exactly the same, but they are obeyed for you in an automated fashion by the compiler. In an ARC-based app, you *never* say `retain` or `release` — in fact, you're not allowed to. Instead, the compiler says `retain` or `release` for you, using exactly the principles you would have had to use if you *had* said them (the golden rules of Cocoa memory management)! Since the compiler is smarter (or at least more ruthlessly tenacious) than you

are about this sort of nit-picky rule-based behavior, it won't make any of the mistakes you might have made due to carelessness or confusion.

The moment an object is released, there is a chance it will be destroyed. Before ARC, this fact was a big worry for programmers. In a non-ARC program, you must take care not to send any messages subsequently through any pointer to an object that has been destroyed — including the pointer you just used to release the object. In effect, you've just turned your *own* pointer into a possible dangling pointer! If there is any danger that you might accidentally attempt to use this dangling pointer, a wise policy is to *nilify* the pointer — that is, to set the pointer itself to nil. A message to nil has no effect, so if you do send a message through that pointer, it won't do any good, but at least it won't do any harm (kind of like chicken soup).

In an ARC-based program, this policy, too, is strictly followed: ARC will nilify for you any pointer to whose object it has just sent the last balancing `release` message (meaning that the object might now have gone out of existence). Since, as I mentioned in [Chapter 3](#), ARC also sets an instance pointer to nil when you declare it (if you don't initialize it yourself, there and then, to point to an actual instance), there follows as the night the day the following delightful corollary: *under ARC, every instance pointer either points either to an actual instance or is nil*. This fact alone should send you rushing to convert all your existing non-ARC apps to ARC if you possibly can.

What ARC Is and What It Does

When you create a new Xcode project and choose an application template, a checkbox in the second dialog lets you elect to Use Automatic Reference Counting. Automatic Reference Counting is ARC. If this checkbox is checked, then (among other things):

- The LLVM compiler build setting Objective-C Automatic Reference Counting (`CLANG_ENABLE_OBJC_ARC`) for your project is set to YES.
- Any `retain` or `release` statements that would have been present in the non-ARC version of any of the project template's `.m` files are stripped out.
- Any code that Xcode subsequently inserts automatically, such as a property generated by Control-dragging from a nib into code, will conform to ARC conventions.

It is also possible to convert an existing non-ARC project to ARC; choose `Edit → Refactor → Convert to Objective-C ARC` for assistance with the necessary code changes. (For full details, see the appropriate WWDC 2011 videos.)



ARC is actually a feature of the LLVM 3.0 compiler, and is one of the main purposes for which the LLVM compiler was developed. For full technical information, see <http://clang.llvm.org/docs/AutomaticReferenceCounting.html>. — You do not have to adopt ARC for an entire project; if you have old non-ARC code, possibly written by someone else, you may wish to incorporate that code into your ARC-based project without substantially altering the non-ARC code. To do so, confine all non-ARC code to its own files, and for each of those files, edit the target, switch to the Build Phases tab, and in the Compile Sources section, double-click the non-ARC file's listing and type `-fno-objc-arc` in the box (to enter it in the Compiler Flags column).

When you write code for an ARC-based project, as I've already mentioned, you never say `retain` or `release`. When you compile an ARC-based project, the compiler will treat any `retain` or `release` commands as an error, and will instead, behind the scenes, insert its own commands that effectively do the exact same work as `retain` and `release` commands. Your code is thus manually memory-managed, in conformity with the principles and golden rules of manual memory management that I've already described, but the author of the manual memory-management code is the compiler (and the memory-management code itself is invisible, unless you feel like reading assembly language).

ARC does its work of inserting `retain` and `release` commands in two stages:

1. It behaves very, very conservatively; basically, if in doubt, it retains (and of course later releases). In effect, ARC retains at every juncture that might have the slightest implications for memory management: it retains when an object is received as an argument, it retains when an object is assigned to a variable, and so forth. It may even insert temporary variables to enable it to refer sufficiently early to an object so that it can retain it. But of course it also releases to match. This means that at the end of the first stage, memory management is technically correct; there may be far more retains and releases on a single object than you would have put if you were writing those commands yourself, but at least you can be confident that no pointer will dangle and no object will leak.
2. It optimizes, removing as many `retain` and `release` pairs from each object as it possibly can while still ensuring safety with regard to the program's actual behavior. This means that at the end of the second stage, memory management is still technically correct, and it is also efficient.

So, for example, consider the following code:

```
- (void) myMethod {
    NSArray* myArray = [NSArray array];
    NSArray* myOtherArray = myArray;
}
```

Now, in actual fact, no additional memory management code is needed here (for reasons that I'll clarify in the next section). But in its first pass, we may imagine that ARC

will behave very, very conservatively: it will ensure that every variable is nil or points to an object, and it will retain on every assignment, releasing any existing value pointed to by the variable being assigned to. So we may imagine (though this is unlikely to be exactly correct) a scenario where ARC compiles that code at first into the equivalent of [Example 12-1](#).

Example 12-1. Imaginary scenario of ARC's conservative memory management

```
- (void) myMethod {
    // create all new object pointers as nil
    NSArray* myArray = nil;
    // retain as you assign
    id temp1 = myArray;
    myArray = [NSArray array];
    [myArray retain];
    [temp1 release]; // no effect, it's nil
    // create all new object pointers as nil
    NSArray* myOtherArray = nil;
    // retain and release as you assign
    id temp2 = myOtherArray;
    myOtherArray = myArray;
    [myOtherArray retain];
    [temp2 release]; // no effect, it's nil
    // method is ending, balance out retains on local variables
    [myArray release];
    myArray = nil;
    [myOtherArray release];
    myOtherArray = nil;
}
```

The ARC optimizer will come along and reduce the amount of work being done here. For example, it may observe that `myArray` and `myOtherArray` turn out to be pointers to the same object, so it may therefore remove some of the intermediate retains and releases. And it may observe that there's no need to send `release` to nil. But retains and releases are so efficient under ARC that it wouldn't much matter if the optimizer didn't remove any of the intermediate retains and releases.

However, the keen-eyed reader may have noticed that there is more to the manual memory management balancing act than matching `retain` and `release`: in particular, I said earlier that `alloc` and `copy` yielded objects whose retain count had already been incremented, so that they, too, must be balanced by `release`. In order to obey this part of the golden rules of Cocoa memory management, ARC resorts to *assumptions about how methods are named*. This means that *you* had better conform, in your code, to the same assumptions about how methods are named, or you can accidentally cause ARC to do the wrong things (although, as it turns out, there are ways out of this predicament if you have a wrongly-named method whose name you absolutely can't change).

In particular, when your code receives an object as the returned value of a method call, ARC looks at the opening word (or words) of the camelCased method name. (The term *camelCased* describes a compound word whose individual words are demarcated by

internal capitalization, like the words “camel” and “Cased” in the word “camelCased.”) If the opening word of the name of that method is `alloc`, `init`, `new`, `copy`, or `mutableCopy`, ARC assumes that this object has an incremented retain count that will need to be balanced with a corresponding `retain`.

So, in the preceding example, if the array had been received from a call to `[NSArray new]` instead of `[NSArray array]`, ARC would know that an extra `release` will eventually be needed, to balance the incremented retain count of the object returned from a method whose name begins with `new`.

Your own responsibility in this regard, then, is *not* to name any of your methods inappropriately in such a way as to set off that sort of alarm bell in ARC’s head. The easiest approach is not to start any of your own method names with `alloc`, `init` (unless you’re writing an initializer, of course), `new`, `copy`, or `mutableCopy`. Doing so might not cause any damage, but it is better not to take the risk: obey the ARC naming conventions if you possibly can.



All code examples in preceding chapters (and subsequent chapters, for that matter) represent ARC code and are therefore correctly memory-managed. In a non-ARC project, however, some code examples from preceding chapters require further explicit memory management or objects will leak. In the first edition of this book, I included at this point a section rewriting those earlier examples to show how they would look with memory properly managed. In this edition there is no need, since I’m using ARC throughout and I’m assuming you’re using it as well, so that section is omitted.

How Cocoa Objects Manage Memory

Built-in Cocoa objects will take ownership of objects you hand them, by retaining them, if it makes sense for them to do so. (Indeed, this is so generally true that if a Cocoa object is *not* going to retain an object you hand it, there will be a note to that effect in the documentation.) Thus, you don’t need to worry about managing memory for an object if the only thing you’re going to do with it is hand it over to a Cocoa object.

A good example is an `NSArray`. Consider the following minimal example:

```
NSString* s = [[NSDate date] description];
NSArray* arr = [NSArray arrayWithObject: s];
```

When you hand the string to the array, the array retains the string. As long as the array exists and the string is in the array, the string will exist. When the array goes out of existence, it will also release the string; if no other object is retaining the string, the string will then go out of existence in good order, without leaking, and all will be well. All of this is right and proper; the array could hardly “contain” the string without taking ownership of it.

An NSMutableArray works the same way, with additions. When you add an object to an NSMutableArray, the array retains it. When you remove an object from an NSMutableArray, the array releases it. Again, the array is always doing the right thing.

Thus you should stay out of, and not worry yourself about, memory management for objects you don't own; the right thing will happen all by itself. For instance, look back at [Example 10-4](#). Here it is again:

```
NSString* f = [[NSBundle mainBundle] pathForResource:@"index" ofType:@"txt"];
NSError* err = nil;
NSString* s = [NSString stringWithContentsOfFile:f
                                         encoding:NSUTF8StringEncoding
                                         error:&err];
// error-checking omitted
NSMutableDictionary* d = [NSMutableDictionary dictionary];
for (NSString* line in [s componentsSeparatedByString:@"\n"]) {
    NSArray* items = [line componentsSeparatedByString:@"\t"];
    NSInteger chnum = [[items objectAtIndex: 0] integerValue];
    NSNumber* key = [NSNumber numberWithInteger:chnum];
    NSMutableArray* marr = [d objectForKey: key];
    if (!marr) { // no such key, create key-value pair
        marr = [NSMutableArray array];
        [d setObject: marr forKey: key];
    }
    // marr is now a mutable array, empty or otherwise
    NSString* picname = [items objectAtIndex: 1];
    [marr addObject: picname];
}
```

No explicit memory management is happening here, and no additional memory management needs to happen (even if you aren't using ARC). We're generating a lot of objects, but never do we say `alloc` (or `copy`), so we have no ownership, and memory management is therefore not our concern. Moreover, no bad thing is going to happen between one line and the next while this code is running. The mutable dictionary `d`, for example, generated by calling `[NSMutableDictionary dictionary]`, is not going to vanish mysteriously before we can finish adding objects to it. (I'll say a bit more, later in this chapter, about why I'm so confident of this.)

On the other hand, it is possible (if you aren't using ARC) to be tripped up by how Cocoa objects manage memory. Consider the following:

```
NSString* s = [myMutableArray objectAtIndex: 0];
[myMutableArray removeObjectAtIndex: 0]; // Bad idea! (But just fine under ARC)
```

Here we remove a string from an array, keeping a reference to it ourselves as `s`. But, as I just said, when you remove an object from an NSMutableArray, the array releases it. So the commented line of code in the previous example involves an implicit release of the string in question, and if this reduces the string's retain count to zero, it will be destroyed. In effect, we've just done the thing I warned you about at the end of the previous section: we've turned our own pointer `s` into a possible dangling pointer, and a crash may be in our future when we try to use it as if it were a string.

The way to ensure against such possible destruction in non-ARC code is to retain the object before doing anything that might destroy it ([Example 12-2](#)).

Example 12-2. How non-ARC code ensures a collection element's persistence

```
NSString* s = [myMutableArray objectAtIndex: 0];
[s retain]; // this is non-ARC code
[myMutableArray removeObjectAtIndex: 0];
```

Of course, now you have made management of this object your business; you have asserted ownership of it, and must make sure that this `retain` is eventually balanced by a subsequent `release`, or the string object may leak.

However, the very same code works perfectly under ARC:

```
NSString* s = [myMutableArray objectAtIndex: 0];
[myMutableArray removeObjectAtIndex: 0]; // Just fine under ARC
```

The reason is that, as I mentioned earlier, ARC is insanely conservative at the outset. Just as in [Example 12-1](#), ARC retains on assignment, so we may imagine that ARC will operate according to something like the following scenario:

```
NSString* s = nil;
id temp = s;
s = [myMutableArray objectAtIndex: 0];
[s retain];
[temp release];
[myMutableArray removeObjectAtIndex: 0];
// ... and later ...
[s release];
s = nil;
```

This turns out to be exactly the right thing to do! When the call to `removeObjectAtIndex:` comes along, the retain count of the object received from the array is still incremented, exactly as in our non-ARC [Example 12-2](#).

Autorelease

When you call a method and receive as a result what [Chapter 5](#) calls a ready-made instance, how does memory management work? Consider, for example, this code:

```
NSArray* myArray = [NSArray array];
```

According to the golden rules of memory management, the object now pointed to by `myArray` doesn't need memory management. You didn't say `alloc` in order to get it, so you haven't claimed ownership of it and you don't need to release it (and shouldn't do so). But how is this possible? How is the `NSArray` class able to vend an array that you don't have to release without also leaking that object?

If you don't see why this is mysterious, pretend that *you* are `NSArray`. How would you implement the `array` method so as to generate an array that the caller doesn't have to memory-manage? Don't say that you'd just call some *other* `NSArray` method that vends

a ready-made instance; that merely pushes the same problem back one level. You *are* NSArray. Sooner or later, *you* must somehow supply this magical instance. Ultimately you will have to generate the instance from scratch, and then how will you manage its memory? You can't do it like this:

```
- (NSArray*) array {
    NSArray* arr = [[NSArray alloc] init];
    return arr; // hmmm, not so fast...
}
```

This, it appears, can't work. On the one hand, we generated `arr`'s value by saying `alloc`. This means we must release the object pointed to by `arr`. On the other hand, when are we going to do this? If we do it just *before* returning `arr`, `arr` will be pointing to garbage and we will be vending garbage. We cannot do it just *after* returning `arr`, because our method exits when we say `return`. This is a puzzle. It is our job, if we are to be a good Cocoa citizen and follow the golden rules of memory management, to decrement the retain count of this object. We need a way to vend this object without decrementing its retain count *now* (so that it stays in existence long enough for the caller to receive and work with it), yet ensure that we *will* decrement its retain count (to balance our `alloc` call and fulfill our own management of this object's memory).

The solution, which is explicit in pre-ARC code, is `autorelease`:

```
- (NSArray*) array {
    NSArray* arr = [[NSArray alloc] init];
    [arr autorelease];
    return arr;
}
```

Or, because `autorelease` returns `self`, we can condense that:

```
- (NSArray*) array {
    NSArray* arr = [[NSArray alloc] init];
    return [arr autorelease];
}
```

Here's how `autorelease` works. Your code runs in the presence of something called an *autorelease pool*. (If you look in `main.m`, you can actually see an autorelease pool being created.) When you send `autorelease` to an object, that object is placed in the autorelease pool, and a number is incremented saying how many times this object has been placed in this autorelease pool. From time to time, when nothing else is going on, the autorelease pool is automatically *drained*. This means that the autorelease pool sends `release` to each of its objects, the same number of times as that object was placed in this autorelease pool, and empties itself of all objects. If that causes an object's retain count to be zero, fine; the object is destroyed in the usual way. So `autorelease` is just like `release` — effectively, it *is* a form of `release` — but with a proviso, “later, not right this second.”

You don't need to know exactly when the current autorelease pool will be drained; indeed, you can't know (unless you force it, as we shall see). The important thing is

that in a case like our method `array`, there will be plenty of time for whoever called `array` to retain the vended object if desired.

The vended object in a case like our method `array` is called an *autoreleased object*. The object that is doing the vending has in fact completed its memory management of the vended object. The vended object thus potentially has a zero retain count. But it doesn't have a zero retain count just yet. The vended object is not going to vanish right this second, because your code is running and so the autorelease pool is not going to be drained right this second. The recipient of such an object needs to bear in mind that this object may be autoreleased. It won't vanish while the code that called the method that vended the object is running, but if the receiving object wants to be sure that the vended object will persist later on, it should retain it.

This explains why there's no explicit memory management in [Example 10-4](#) (cited earlier in this chapter): we don't madly retain every object we obtain in that code, even in non-ARC code, because those objects will all persist long enough for our code to finish. This fits with the golden rules of memory management. An object you receive by means *other* than those listed among the golden rules as asserting ownership (`alloc` or `copy`) isn't under your ownership. The object will either be owned and retained by some other persistent object, in which case it won't vanish while the other object persists, or it will be independent but autoreleased, in which case it will at least persist while your code continues to run. If you want it to persist and you're afraid it might not, you should take ownership of it by retaining it.

Under ARC, as you might expect, all the right things happen of their own accord. You don't have to say `autorelease`, and indeed you cannot. Instead, ARC will say it for you. And it says it in accordance with the method naming rule I described earlier. A method called `array`, for example, does not start with a camelCase unit `new`, `init`, `alloc`, `copy`, or `mutableCopy`. Therefore it must return an object whose memory management is balanced, using `autorelease` for the last release. ARC will see to it that this is indeed the case. On the other side of the ledger, the method that called `array` and received an array in return must assume that this object is autoreleased and could go out of existence if we don't retain it. That's exactly what ARC does assume.

Sometimes you may wish to drain the autorelease pool immediately. Consider the following:

```
for (NSString* aWord in myArray) {
    NSString* lowerAndShorter = [[aWord lowercaseString] substringFromIndex:1];
    [myMutableArray addObject: lowerAndShorter];
}
```

Every time through that loop, two objects are added to the autorelease pool: the lowercase version of the string we start with, and the shortened version of that. The first object, the lowercase version of the string, is purely an *intermediate object*: as the current iteration of the loop ends, no one except the autorelease pool has a pointer to it. If this loop had very many repetitions, or if these intermediate objects were themselves very large in size, this could add up to a lot of memory. These intermediate objects will all

be released when the autorelease pool drains, so they are not leaking; nevertheless, they are accumulating in memory, and in certain cases there could be a danger that we will run out of memory before the autorelease pool drains. The problem can be even more acute than you know, because you might repeatedly call a built-in Cocoa method that itself accumulates a lot of intermediate objects.

The solution is to intervene in the autorelease pool mechanism by supplying your own autorelease pool. This works because the autorelease pool used to store an autoreleased object is the most recently created pool. So you can just create an autorelease pool at the top of the loop and drain it at the bottom of the loop, each time through the loop. In iOS 5, the notation for doing this is to surround the code that is to run under its own autorelease pool with the directive `@autoreleasepool{}`, like this:

```
for (NSString* aWord in myArray) {  
    @autoreleasepool {  
        NSString* lowerAndShorter = [[aWord lowercaseString] substringFromIndex:1];  
        [myMutableArray addObject: lowerAndShorter];  
    }  
}
```

(Prior to iOS 5, the notation was to instantiate an `NSAutoreleasePool` object at the top of the loop and call `release` or `drain` on that object at the bottom of the loop. Under ARC this notation is not available, and even in non-ARC code `@autoreleasepool{}` is considered preferable; the truth is that an `NSAutoreleasePool` instance was never a real instance, so the newer notation is more honest.)

Many classes provide the programmer with two equivalent ways to obtain an object: either an autoreleased object or an object that you create yourself with `alloc` and some form of `init`. So, for example, `NSArray` supplies both the class method `arrayWithObjects:` and the instance method `initWithObjects:..` Which should you use? Before ARC, Apple has stated that they would prefer you to lean toward `initWithObjects:..` In general, where you can generate an object with `alloc` and some form of `init`, they'd like you to do so. That way, you are in charge of releasing the object. This policy prevents your objects from hanging around in the autorelease pool and keeps your use of memory as low as possible. Under ARC, I still tend to adhere to this policy from force of habit, but in fact the ARC autorelease pool architecture is so efficient that the old policy may no longer provide any advantage.

Memory Management of Instance Variables (Non-ARC)

Before ARC, the main place for the programmer to make a memory management mistake was with respect to instance variables. Memory management of temporary variable within a single method is pretty easy; you can see the whole method at once, so now just follow the golden rules of memory management, balancing every `retain`, `alloc`, or `copy` with a `release` (or, if you're returning an object with an incremented retain count, `autorelease`). But instance variables make things complicated, for many reasons:

- Instance variables are persistent. Your own instance variables will persist when this method is over and your code has stopped running and the autorelease pool has been drained. So if you want an object value pointed to by an instance variable not to vanish in a puff of smoke, you'd better retain it as you assign it to the instance variable.
- Instance variables are managed from different places in your code. This means that memory management can be spread out over several different methods, making it difficult to get right and difficult to debug if you get it wrong. For example, if you retained a value assigned to an instance variable, you'll later need to release it in order to conform to the golden rules of memory management — but in some other method.
- Instance variables might not belong to you. You will often assign to or get a value from an instance variable belonging to another object. You are now sharing access to a value with some other object. If that object were to go out of existence and release its instance variables, and you have a pointer to the instance variable value coming from that other object and you haven't asserted your own ownership by retaining that value, you can wind up with a dangling pointer.

To see what I mean, return once more to [Example 10-4](#). As I have already explained, there was no need to worry about memory management during this code, even without ARC. We have a mutable dictionary `d`, which we acquired as a ready-made instance by calling `[NSMutableDictionary dictionary]`, and it isn't going to vanish while we're working with it. Now, however, suppose that in the *next* line we propose to assign `d` to an instance variable of ours:

```
self->theData = d; // in non-ARC code this would be a bad idea!
```

Before ARC, that code constituted a serious potential mistake. If our code now comes to a stop, we're left with a persistent pointer to an object over which we have never asserted ownership; it might vanish, leaving us with a dangling pointer. The solution, obviously, is to retain this object as we assign it to our instance variable. You could do it like this:

```
[d retain];
self->theData = d;
```

Or you could do it like this:

```
self->theData = d;
[self->theData retain];
```

Or, because `retain` returns `self`, you could actually do it like this:

```
self->theData = [d retain];
```

(Make sure you understand why those are all equivalent. It's because `d` and `self->theData` are just names; they are pointers. What you're retaining is the object pointed to. How you refer to that object, under what name, is neither here nor there.)

So which should you use? Probably none of them. Consider what a lot of trouble it will be if you ever want to assign a *different* value to `self->theData`. You’re going to have to remember to release the object already pointed to (to balance the retain you’ve used here), and you’re going to have to remember to retain the next value as well. It would be much better to encapsulate memory management for this instance variable in an accessor (a setter). That way, as long as you always pass through the accessor, memory will be managed correctly. A standard template for such an accessor might look like [Example 12-3](#).

Example 12-3. A simple retaining setter

```
- (void) setTheData: (NSMutableArray*) value {
    if (self->theData != value) {
        [self->theData release];
        self->theData = [value retain];
    }
}
```

In [Example 12-3](#), we release the object currently pointed to by our instance variable (and if that object is nil, no harm done) and retain the incoming value before assigning it to our instance variable (and if that value is nil, no harm done either). The test for whether the incoming value is the object already pointed to by our instance variable is not just to save a step; it’s because if we were to release that object, it could vanish then and there, turning not only `self->theData` but also `value` (which points to the same thing) into a dangling pointer.

The setter accessor now manages memory correctly for us; provided we always use it to set our instance variable, all will be well. This is one of the main reasons why accessors are so important! So the assignment to the instance variable in our original code should now look like this:

```
[self setTheData: d];
```

Observe that we can also use this setter subsequently to release the value of the instance variable and nilify the instance variable itself, thus preventing a dangling pointer, all in a single easy step:

```
[self setTheData: nil];
```

So there’s yet another benefit of using an accessor to manage memory.

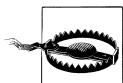
Our memory management for this instance variable is still incomplete, however. We (meaning the object whose instance variable this is) must also remember to release the object pointed to by this instance variable at the last minute before we ourselves go out of existence. Otherwise, if this instance variable points to a retained object, there will be a memory leak. The “last minute” is typically `dealloc`, the `NSObject` method ([Chapter 10](#)) that is called as an object goes out of existence.

In `dealloc`, there is no need to use accessors to refer to an instance variable, and in fact it’s not a good idea to do so, because you never know what other side effects an accessor

might have. And (under non-ARC code) you must always call `super` last of all. So here's our implementation of this object's `dealloc`:

```
- (void) dealloc {
    [self->theData release];
    [super dealloc];
}
```

That completes the memory management for one instance variable. In general, if you are not using ARC, you will need to make sure that *every* object of yours has a `dealloc` that releases *every* instance variable whose value has been retained. This, obviously, is one more very good opportunity for you to make a mistake.



Never, never call `dealloc` in your code, except to call `super` last of all in your override of `dealloc`. Under ARC, you *can't* call `dealloc` — yet another example of how ARC saves you from yourself.

Just as it's not a good idea to use your own accessors to refer to your own instance variable in `dealloc`, so you should not use your own accessors to refer to your own instance variables in an initializer (see [Chapter 5](#)). The reason is in part that the object is not yet fully formed, and in part that an accessor can have other side effects. Instead, you will set your instance variables directly, but you must also remember to manage memory.

To illustrate, I'll rewrite the example initializer from [Chapter 5 \(Example 5-3\)](#). This time I'll allow our object (a Dog) to be initialized with a name. The reason I didn't discuss this possibility in [Chapter 5](#) is that a string is an object whose memory must be managed! So, imagine now that we have an instance variable `name` whose value is an `NSString`, and we want an initializer that allows the caller to pass in a value for this instance variable. It might look like [Example 12-4](#).

Example 12-4. A simple initializer that retains an ivar

```
- (id) initWithName: (NSString*) s {
    self = [super init];
    if (self) {
        self->name = [s retain];
    }
    return self;
}
```

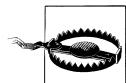
Actually, it is more likely in the case of an `NSString` that you would copy it rather than merely retain it. The reason is that `NSString` has a mutable subclass `NSMutableString`, so some other object might call `initWithName:` and hand you a mutable string to which it still holds a reference — and then mutate it, thus changing this Dog's name behind your back. So the initializer would look like [Example 12-5](#).

Example 12-5. A simple initializer that copies an ivar

```
- (id) initWithName: (NSString*) s {
    self = [super init];
    if (self) {
        self->name = [s copy];
    }
    return self;
}
```

In [Example 12-5](#), we don't bother to release the existing value of `name`; it is certainly not pointing to any *previous* value (because there is no previous value), so there's no point.

Thus, memory management for an instance variable may take place in as many as three places: the initializer, the setter, and `dealloc`. This is a common architecture. It is a lot of work, and a common source of error, having to look in multiple places to check that you are managing memory consistently and correctly, but that's what you must do if you aren't using ARC (though, as I'll point out later in this chapter, Objective-C has the ability to write your accessors for you).



Earlier, I mentioned that KVC will set an instance variable directly if it can't find a setter corresponding to the key. When it does this, *it retains the incoming value*. This fact is little-known and poorly documented — and scary. The last thing you want, in non-ARC code, is implicit memory management. This is one more reason to provide accessors. On the other hand, if you're using ARC, this is not such a worry, since ARC is already providing implicit memory management.

Memory Management of Instance Variables (ARC)

If you're using ARC, ARC will manage your instance variable memory for you; you needn't (and, by and large, you can't) do it for yourself. By default, ARC will treat an instance variable the same way it treats any variable: on assignment to that instance variable, it creates a temporary variable, retains the assigned value in it, releases the current value of the instance variable, and performs the assignment. Thus, you write this code:

```
self->theData = d; // an NSMutableDictionary
```

ARC, in effect, in accordance with its rule that it retains and releases on assignment, substitutes something like this scenario:

```
id temp = self->theData;
self->theData = d;
[self->theData retain];
[temp release];
```

This is exactly the right thing to have happened; in fact, it will not have escaped your attention that it is virtually the same code you would have written for a formal accessor

such as [Example 12-3](#). So much for worrying about release and retain on assignment! If you did want to write a setter, it might consist of no more than a direct assignment:

```
- (void) setTheData: (NSMutableArray*) value {
    self->theData = value;
}
```

Moreover, when your object goes out of existence, ARC releases its retained instance variable values. So much for worrying about releasing in `dealloc`! You may still need, under ARC, to implement `dealloc` for other reasons — for example, it could still be the right place to unregister for a notification ([Chapter 11](#)) — but you won’t call `release` on any instance variables there, and you won’t call `super`. At the time `dealloc` is called, your instance variables have not yet been released, so it’s fine to refer to them in `dealloc`.

At this point you may be imagining that, under ARC, you might be able to live without any accessors at all: instead, you can just assign directly to your instance variables and all the right memory-management things will happen, so who needs a formal setter? However, a formal accessor, as I’ll explain later, can do things above and beyond ARC’s automated insertion of release-and-retain, such as copying instead of retaining, dealing with multithreading, and adding your own custom behaviors. Also, obviously, accessors can be made public and so available to other objects, whereas an instance variable is not public. Nevertheless, it is certainly true that in my own code I’ve found myself using more private instance variables and assigning to them directly, simply because ARC makes it so easy (and memory-safe) to do so.



In the absence of a `release` call, which is forbidden under ARC, what happens if you want to release an instance variable’s value manually?

The solution is simple: set the instance variable to nil. When you nilify a variable, ARC releases its existing value for you by default.

You may be wondering about ARC’s implications for the way you’ll write an initializer that involves setting object instance variable values, as in [Example 12-4](#) and [Example 12-5](#). The code for these initializers will be just the same under ARC as under non-ARC, except that you needn’t (and can’t) say `retain`. So [Example 12-4](#) under ARC would look like this ([Example 12-6](#)):

Example 12-6. A simple initializer that retains an ivar under ARC

```
- (id) initWithName: (NSString*) s {
    self = [super init];
    if (self) {
        self->name = s;
    }
    return self;
}
```

[Example 12-5](#) under ARC will be unchanged; you can still say `copy` under ARC, and ARC understands how to manage the memory of an object returned from a method whose camelCased name starts with (or simply is) `copy` ([Example 12-7](#)):

Example 12-7. A simple initializer that copies an ivar under ARC

```
- (id) initWithName: (NSString*) s {
    self = [super init];
    if (self) {
        self->name = [s copy];
    }
    return self;
}
```

Retain Cycles and Weak References

ARC's behavior is automatic and mindless; it knows nothing of the logic of the relationships between objects in your app. Sometimes, you have to provide ARC with further instructions to prevent it from doing something detrimental. Typically, this detrimental thing will be the creation of a retain cycle.

A *retain cycle* is a situation in which object A and object B are each retaining one another. This can arise quite innocently, because relationships in an object graph can run both ways. For example, in a system of orders and items, an order needs to know what its items are and an item might need to know what orders it is a part of, so you might be tempted to let be the case *both* that an order retains its items *and* that an item retains its orders. That's a retain cycle, with object A (an order) retaining object B (an item) and *vice versa*. A moment's thought will reveal that such a situation, if allowed to persist, will result in a leak of both objects. Object B's retain count can never decrement to zero so long as object A continues to retain it, and object A's retain count can never decrement to zero so long as object B continues to retain it. If either object's retain count were to decrement to zero, ARC would then release all the objects pointed to by its instance variables, and the other object's retain count might thus decrement to zero as well; both objects would go out of existence together. But this can never happen; instead, both objects keep each other in existence forever.

In order to illustrate the problem simply, without yet getting into properties and their memory management policies — I'll talk about that later in this chapter — I'll suppose a rather unusual class `MyClass` with a public instance variable:

```
@interface MyClass : NSObject {
    @public
    id thing;
}
```

We now run this code:

```
 MyClass* m1 = [MyClass new];
 MyClass* m2 = [MyClass new];
 m1->thing = m2;
 m2->thing = m1;
```

In non-ARC code, where you are responsible for memory management, direct assignment has no memory management implications, and this is not a retain cycle. But under ARC, unless you take steps to the contrary, this will be a retain cycle; by default, `m1` and `m2` are now retaining one another, because by default, ARC retains on assignment. To see this, implement `dealloc` with an `NSLog` message in `MyClass`; when the code runs, `dealloc` is never called for either of our `MyClass` instances. They have leaked.

You can prevent an instance variable from retaining the object assigned to it by specifying that the instance variable should be a *weak reference*. You can do this with the `_weak` qualifier in the instance variable's declaration:

```
@interface MyClass : NSObject {
    @public
    _weak id thing;
}
```

Now there is no retain cycle. In our particular example, since both `m1` and `m2` exist only as automatic variables in the scope of the code that creates them, they will both go out of existence instantly when that code comes to an end and ARC releases them both (to balance the `new` calls that created them).



In ARC, a reference not explicitly declared weak is a *strong reference*. Thus, a strong reference is one where ARC retains as it assigns. There is in fact a `_strong` qualifier, but in practice you'll never use it, as it is the default. (There are also two additional qualifiers, `_unsafe_unretained` and `_autoreleasing`, but they are rarely needed and I don't talk about them in this book.) Observe that although, by a sort of trick, an ARC-based app can be made to work backwards-compatibly in iOS 4, ARC weak references cannot.

In real life, a weak reference is most likely to arise with regard to something like an object and its delegate ([Chapter 11](#)). A delegate is an independent entity; there is usually no reason why an object needs to claim ownership of its delegate. The object should have no role in the persistence of its delegate; and it could even be that the delegate might for some reason retain the object, causing a retain cycle. Therefore, most delegates should be declared as weak references. For example, in an ARC project created from Xcode's Utility Application project template, you'll find this line:

```
@property (weak, nonatomic) id <FlipsideViewControllerDelegate> delegate;
```

(The delegate may also be tagged as an `IBOutlet`, but you can ignore that.) The keyword `weak` in the property declaration, as I'll explain more fully later in this chapter, is equivalent to declaring the `delegate` instance variable as `_weak`.

In non-ARC code, a reference can be prevented from causing a retain cycle merely by not retaining when assigning to that reference; the reference isn't memory-managed at all. You will see this referred to a weak reference; it is not, however, quite the same thing as an ARC weak reference. A non-ARC weak reference risks turning into a dangling pointer when the instance to which it points is released in such a way as to reduce its retain count to zero and vanishes in a puff of smoke behind its back, as it were. This means that it is possible for the reference to be non-nil and pointing to garbage, so that a message sent to it can have mysteriously disastrous consequences. Amazingly, however, this cannot happen with an ARC weak reference: the instance to which it points can be released and have its retain count reach zero and vanish, but when it does, any ARC weak reference that was pointing to it is set to nil! This amazing feat is accomplished by some behind-the-scenes bookkeeping: when an object is assigned to a weak reference, ARC in effect notes this fact on a scratchpad list. When the object is released, ARC consults the scratchpad list and discovers the existence of the weak reference to it, and assigns nil to that weak reference. This is yet another reason for preferring to use ARC wherever possible. ARC sometimes refers to non-ARC weak references, disdainfully but accurately, as "unsafe." (Non-ARC weak references are in fact the `__unsafe_unretained` references I mentioned a moment ago.)

Most properties of built-in Cocoa classes that keep weak references are, however, non-ARC weak references (because they are old and backwards-compatible, whereas ARC is new). Such properties are declared using the keyword `assign`. For example, `UINavigationController`'s `delegate` property is declared like this:

```
@property(nonatomic, assign) id<UINavigationControllerDelegate> delegate
```

This means that if you (some object) are assigned to a `UINavigationController` as its `delegate`, and if you are about to go out of existence at a time when this `UINavigationController` still exists, you have a duty to set that `UINavigationController`'s `delegate` property to nil; otherwise, it might try to send a message to its `delegate` at some future time, when you no longer exist and its `delegate` property is a dangling pointer, and the app will then crash — and, since this happens at some future time, figuring out the cause of the crash can be quite difficult. (This is the sort of situation in which you might need to turn on zombies in order to debug, as described earlier in this chapter.)

The same consideration applies to delegates in general, as well as to such non-retained non-ARC weak references as data sources and action targets. *Be sure to nilify any non-ARC weak references belonging a persistent object if the object referred to is about to go out of existence while the persistent object will continue to exist.* The `dealloc` method is the usual place for fulfilling such duties. The situation does not often arise — typically, the thing whose delegate or data source or action target you are will go out of existence before you do — but it can, and when it does, you can crash in a mysterious way later on if you have failed to do your memory-management duty in this regard. I stress that this is every bit as much an issue in an ARC project as in a non-ARC project.

This is also the reason why, when you go out of existence, you must deregister yourself from the shared notification center if you are registered there to receive any notifications. If you registered using `addObserver:selector:name:object:`, you handed the notification center a reference to yourself as the first argument; the notification center's reference to you is a weak reference, as well it might be, since the notification center would have no business keeping you in existence, but this is a non-ARC weak reference, and there is a danger that the notification will try to send a notification to whatever is referred to. By deregistering yourself, you remove the notification center's reference to you.

If you registered with the notification center using `addObserverForName:object:queue:usingBlock:`, memory management can be a little tricky, under ARC in particular, for two reasons:

- As long as you are registered with the notification center, the notification center may be retaining you.
- The observer token returned from the call to `addObserverForName:object:queue:usingBlock:` may also be retaining you.

The result can be a retain cycle. Here's an example (this is all ARC code). We have a class `MyClass` with an instance method `start` which causes the object to register itself with the notification center for a fictitious notification:

```
- (void) start {
    // observer is an id instance variable
    self->observer = [[NSNotificationCenter defaultCenter]
        addObserverForName:@"heyho"
        object:nil queue:nil usingBlock:^(NSNotification *note) {
            NSLog(@"%@", self);
    }];
}
```

We create a `MyClass` instance and tell it to `start`, and retain the instance:

```
self->thing = [MyClass new];
[self->thing start];
```

At some later time, we release the `MyClass` instance by nilifying the instance variable that retains it:

```
self->thing = nil;
```

The `MyClass` instance does not go out of existence, as we can confirm by implementing `dealloc` with logging. We are leaking. This must be because someone else is still retaining the `MyClass` instance. But who? Well, it could be the notification center. The `MyClass` instance has behaved badly here by not unregistering itself from the notification center. It is for this purpose that we stored the observer token in our `observer` instance variable; we are failing to use it for the purpose for which we stored it. Note, however, that it would be pointless to try to solve this problem by unregistering in `MyClass`'s `dealloc` method, because we already know that `MyClass`'s `dealloc` is *never*

being called. That's the very problem we're trying to solve. When you have a retain cycle, `dealloc` is too late to break the cycle; that's what makes it a retain cycle.

So let's introduce a further `MyClass` method `stop`, which unregisters the `MyClass` instance from the notification center:

```
- (void) stop {
    [[NSNotificationCenter defaultCenter] removeObserver:self->observer];
}
```

We call that method before releasing the `MyClass` instance:

```
[self->thing stop];
self->thing = nil;
```

That should solve the problem, right? Wrong! The `MyClass` instance's `dealloc` method is *still* not being called; it is still leaking. The reason is that the `MyClass` instance is *retaining itself*, by retaining the observer token. We must break this retain cycle.

The retain cycle can be broken by using any of three techniques, each of which could be useful depending on the circumstances:

- Make the instance variable that holds the reference to the observer token a weak reference:

```
@implementation MyClass {
    __weak id observer;
}
```

This is probably the simplest and best solution in this case. When the `MyClass` instance unregisters itself from the notification center, the notification center releases the observer token. The `MyClass` instance never retained the observer token, because the instance variable is declared a weak reference. The observer token thus goes out of existence, releasing the `MyClass` instance, which is now free to go out of existence in good order. If the `MyClass` instance persists, its `observer` instance variable is now nil, and not a dangling pointer, thanks to the magic of ARC weak references.

- The `observer` instance variable retains the observer token, but when the `MyClass` instance unregisters itself from the notification center, it also nilifies the `observer` instance variable. This releases the `observer` instance variable and breaks the retain cycle. One might say that whereas the previous solution automatically avoided the retain cycle, this solution breaks the retain cycle manually. It has the disadvantage that you must remember to do it, but sometimes manual memory management makes things clearer:

```
- (void) stop {
    [[NSNotificationCenter defaultCenter] removeObserver:self->observer];
    self->observer = nil;
}
```

This solution is useful also when a class registers for multiple notifications and must unregister for all of them at once. Instead of retaining individual `observer`

instance variables, we retain an NSMutableSet into which we place each `observer` instance variable as it is received. When it's time to unregister for all of them, we cycle through the NSMutableSet. I demonstrated this approach in [Chapter 11](#); but that code will leak under ARC, because the NSMutableSet, which must be retained by the MyClass instance, itself retains its elements, each of which retains the MyClass instance. Manually nilifying the reference to the mutable set breaks the retain cycle.

- The `observer` instance variable retains the observer token, and we concentrate instead on how the MyClass instance itself got retained. This happened because the block refers to `self`. Normally there is no complication when a block refers to `self`, but the NSNotificationCenter class documentation on `addObserverForName:object:queue:usingBlock:` says: “The block is copied by the notification center and (the copy) held until the observer registration is removed.” Under ARC, this should cause memory management alarm bells to go off; a copied block retains `self` if `self` is referred to, even indirectly (that is, even if what is referred to is an instance variable). We can prevent this by making the block’s reference to `self` a weak reference. This is done by a technique demonstrated in Apple’s WWDC 2011 videos, commonly called “the weak-strong dance” ([Example 12-8](#)).

Example 12-8. The weak-strong dance prevents a copied block from retaining self

```
- (void) start {
    __weak MyClass* wself = self; ❶
    self->observer = [[NSNotificationCenter defaultCenter]
        addObserverForName:@"heyho"
        object:nil queue:nil usingBlock:^(NSNotification *note) {
            MyClass* sself = wself; ❷
            if (sself) {
                NSLog(@"%@", sself); ❸
            }
        }];
}
```

The weak-strong dance works like this:

- ❶ We form a *local* weak reference to `self`, outside the block but where the block can see it.
- ❷ Inside the block, we form from that weak reference a normal reference.
- ❸ We use that normal reference in place of any references to `self` inside the block.

This breaks the retain cycle and is safe. Observe that this technique is needed only in cases where the block is copied; in the vast majority of situations in which you’ll use blocks, there is no need for the weak-strong dance.

The weak-strong dance may seem elaborate, but it has one huge advantage: it is the *only* one of the three proposed solutions that allows `dealloc` to be called *before* you unregister the observer. Thus, it is the only solution that allows you to unregister the

observer *in* your `dealloc` implementation — and that's important, because this is often the very place where you *want* to unregister the observer.

In using `addObserverForName:object:queue:usingBlock:`, it was up to us to pay attention to the `NSNotificationCenter` class documentation and notice that this method could have unusual memory management implications. In general, you must be on the lookout for Cocoa objects with unusual memory management behavior. Such behavior will be called out clearly in the documentation. For example, the `UIWebView` documentation warns: “Before releasing an instance of `UIWebView` for which you have set a delegate, you must first set its delegate property to `nil`.” And a `CAAnimation` object *retains its delegate*; this is exceptional and can cause trouble if you’re not conscious of it.

Another example is `NSTimer` ([Chapter 10](#)). The `NSTimer` class documentation says that “run loops retain their timers”; it then says of `scheduledTimerWithTimeInterval:target:selector:userInfo:repeats:` that “The target object is retained by the timer and released when the timer is invalidated.” This means that as long as a repeating timer has not been invalidating, the target is being retained by the run loop; the only way to stop this is to send the `invalidate` message to the timer. (With a non-repeating timer, the problem doesn’t arise, because the timer invalidates itself immediately after firing.)

When you called `scheduledTimerWithTimeInterval:target:selector:userInfo:repeats:`, you probably supplied `self` as the `target:` argument. This means that you (`self`) are being retained, and cannot go out of existence until you invalidate the timer. You can’t do this in your `dealloc` implementation, because as long as the timer is repeating and has not been sent the `invalidate` message, `dealloc` won’t be called. You therefore need to find another appropriate moment for sending `invalidate` to the timer. In order to send `invalidate` to the timer later, you will have kept a reference to the timer when you created it (by calling `scheduledTimerWithTimeInterval:target:selector:userInfo:repeats:`); that reference is retained by default under ARC, but that should not be a difficulty — this is not a retain cycle but a case of special memory management (the timer doesn’t release the target until it is invalidated).

There are also situations where the documentation fails to warn of any special memory management considerations, but ARC itself will warn of a possible retain cycle due to the use of `self` in a block. Again, the weak-strong dance is likely to be your best defense. (An example is the completion handler of `UIPageViewController`’s instance method `setViewControllers:direction:animated:completion:`, where the compiler will warn, “Capturing ‘`self`’ strongly in this block is likely to lead to a retain cycle.” Using the weak-strong dance, you capture `self` weakly instead.)

Nib Loading and Memory Management

On iOS, when a nib loads, the top-level nib objects that it instantiates are autoreleased. So if someone doesn’t retain them, they’ll eventually vanish in a puff of smoke. There are two primary strategies for preventing that from happening:

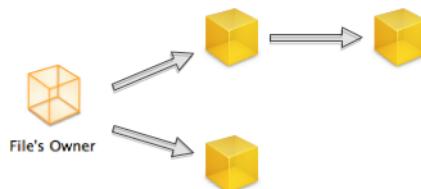


Figure 12-1. An outlet graph with retain

Outlet graph with retain

A memory management graph is formed: every top-level object is retained by another top-level object (without retain cycles, of course), with the File’s Owner as the start of the graph. So, the File’s Owner proxy has an outlet to a top-level object; when the nib loads and the top-level object is assigned to the corresponding instance variable of the actual nib owner instance ([Chapter 7](#)), it is retained. Under ARC this would happen automatically by default, but the usual solution, under both ARC and non-ARC code, is for this instance variable to be backed by a setter accessor method; key-value coding sees to it that this setter is called, and the setter makes sure that the incoming value is retained as it is assigned to the instance variable. Under ARC, mere assignment accomplishes this; under non-ARC code, the incoming value is explicitly retained. Thus, when the nib loads, the nib owner retains this top-level object (and must, of course, release it before it itself goes out of existence). And so on, for every top-level object ([Figure 12-1](#)). This is the strategy you’ll typically use when loading a nib. Fortunately, as I’ll explain later in this chapter, Objective-C can write the accessor for you.

You can see this strategy being used, for example, in a project made from the Single View Application template. The ViewController class is a UIViewController subclass; UIViewController has a `view` property which retains the value assigned to it. Inside the nib, an outlet called `view` runs from the File’s Owner, which is a ViewController, to the top-level UIView (called View) in the nib. Thus this view is retained when the nib loads and doesn’t vanish in a puff of smoke.

Mass retain

The call to `loadNibNamed:owner:options:` ([Chapter 7](#)) returns an NSArray of the nib-instantiated objects; retain this NSArray. (This is the strategy used by `UIApplicationMain` when it loads the app’s main nib, if there is one.)

Objects in the nib that are *not* top-level objects are already part of a memory management object graph, so there’s no need for you to retain them directly. For example, if you have a top-level UIView in the nib, and it contains a UIButton, the UIButton is the UIView’s subview — and a view retains its subviews and takes ownership of them. Thus, it is sufficient to manage the UIView’s memory and to let the UIView manage

the UIButton. If you have an outlet to this button, you typically don't have to retain the button, because it is retained by the UIView as long as the UIButton is inside it (though you would want to retain the button in the rare case where you are planning at some point in your code to remove the button from its superview while keeping it on hand for later use).

In ARC code, when you draw a connection in Xcode between a nib object and your code and Xcode offers to create an instance variable or property declaration for you, it also offers a pop-up menu allowing you to specify whether you want a strong (retaining) or weak reference.

In non-ARC code, you should always use an accessor, even if the setter does nothing but assign the nib object directly to the instance variable. The reason is that, as I warned earlier in this chapter, if there is no accessor, key-value coding itself will assign the nib object to the instance variable *and will retain it*. This implicit, behind-the-scenes memory management is the last thing you want when you're managing memory manually, as it can easily cause a leak. As I'll show later in this chapter, it is easy to arrange for Objective-C to write the accessor for you.



Mac OS X Programmer Alert

Memory management for nib-loaded instances is different on iOS than on Mac OS X. On Mac OS X, nib-loaded instances are not autoreleased, so they don't have to be retained, and memory management is usually automatic in any case because the file's owner is usually an NSWindowController, which takes care of these things for you. On iOS, memory management of top-level nib objects is up to you. On Mac OS X, an outlet to a non-top-level object does not cause an extra retain if there is no accessor for the corresponding ivar; on iOS, it does.

Memory Management of Global Variables

In C, and therefore Objective-C, it is permitted to declare a variable outside of any method. K&R ([Chapter 1](#)) calls this an *external variable* (see K&R 4.3); I call it a *global variable*. It is common practice, though not strictly required, to qualify such a variable's declaration as `static`; such qualification is a technical matter of scope and has no effect the variable's persistence or its global nature.

In Objective-C code, a global variable is not uncommonly used as a lightweight alternative to an instance variable, especially when its value is constant. Also, in a class method, there are no instance variables, but there can be global variables.

Before ARC, where an instance variable had to be backed by explicit memory management in any case, a global variable could be simpler; with ARC, on the other hand, an instance variable is usually *less* trouble to manage than a global variable. The reason is that under ARC, if a global variable is to point to an instance, it must be memory-managed manually. ARC will of course retain on assignment, but when your class's

instance is destroyed, ARC will not automatically release global variables as it will instance variables. Therefore it is up to you to nilify your global variables, thus releasing them, as your instance goes out of existence (probably in `dealloc`). With an instance variable, there's no need to do that.

On the other hand, a global variable does have one advantage an instance variable does not: you can initialize it as you declare it:

```
static NSString* g_myString = @"my string";
```

Under ARC, that's actually safe in a way that it wasn't previously.

Memory Management of Pointer-to-Void Context Info

A number of Cocoa methods take an optional parameter typed as `void*`, and often called `context`:. You might think that `void*`, the universal pointer type, would be the same as `id`, the universal object type, because a reference to an object is a pointer. But an `id` is a universal *object* type; `void*` is just a C pointer. This means that Cocoa won't treat this value as an object. So the use of the `void*` type is a clue to you that Cocoa won't do any memory management on this value. Thus, making sure that it persists long enough to be useful is up to you.

The big change wrought by ARC in this regard has to do, not with memory management, but with "crossing the bridge" between an `id` and a `void*`. Before ARC, these two types were treated as equivalent, in the sense that you could supply one where the other was expected. But ARC is not so sanguine. ARC manages memory for objects only. Thus, it manages memory for something typed as `id`, but *not* for something typed as `void*`. Therefore, if you want to use an object where a `void*` is expected, or a `void*` where an object is expected, you must reassure ARC that you know what you're doing.

When an object comes into existence by instantiation under ARC, it is memory-managed by ARC from birth to death, as it were. But when an object is cast to a `void*`, it passes out of the realm of ARC memory management, and ARC will not let go without more information, because it doesn't understand what its memory management responsibilities should be at this moment. Similarly, when a non-object (a `void*`) is cast to an object type, it passes ready-made into the realm of ARC memory management, and ARC will not accept it without more information, again because it doesn't understand what its memory management responsibilities should be at this moment.

In general, when you supply an object as a `context`: argument, it will probably be easiest if you keep a reference to that object and manage its memory through that reference. Therefore, as the `context`: argument passes out of ARC's purview (and later passes back in again), ARC has *no* memory management responsibilities: you just want ARC to permit the cast. The way to indicate this is to cast the value explicitly, with a `_bridge` qualifier.

As an example, I'll use `beginAnimations:context:`. You call this on a `UIView` before changing one or more of its property values, such as its size, position, or opacity, to

make those changes appear animated. Whatever you pass as the `context:` argument comes back to you later in two delegate messages as the animation proceeds, indicating that the animation is about to start and that the animation has ended; basically, the `context:` is a kind of envelope in which information can be carried from stage to stage during the animation, for any purpose you desire. The `context:` is a `void*`. So, ignoring for a moment the question of memory management, you might say this:

```
NSDictionary* d = [NSDictionary dictionaryWithObject: @"object" forKey: @"key"];
[UIView beginAnimations:@"shrinkImage" context:d];
```

Before ARC, that code was legal. Under ARC, it isn't; the compiler will stop you in your tracks with this complaint: "Implicit conversion of an Objective-C pointer to 'void *' is disallowed with ARC." The solution is to cast `d` as the `context:` argument with a `_bridge` qualifier, like this:

```
NSDictionary* d = [NSDictionary dictionaryWithObject: @"object" forKey: @"key"];
[UIView beginAnimations:@"shrinkImage" context:(__bridge void*)d];
```

In the same way, when the context comes back to you in a delegate message, you'll cast it back to a `_bridge id` so that ARC will accept it and you can continue treating it as an object.

Now let's talk about how to manage the `context:` argument's memory (assuming all the time, of course, that it is an object). Remember, Cocoa is not going to do any memory management behind the scenes. So the `context:` argument must be retained when it is passed as the `context:` argument, or it will vanish in a puff of smoke and garbage will be supplied in its place in the delegate message. And this retain must be balanced by a release at some time after the delegate message is received, or the object will leak.

Clearly, then, you want to retain the `context:` argument as you hand it to `beginAnimations:context:` and release it later, when you're done with it. But when exactly is "later," and when are you actually "done with it?" A good answer would appear to be, "When the animation is over," which is when the `animationDidStop:finished:context:` delegate message arrives. So you could release the context object when you receive it in `animationDidStop:finished:context:`. But this solution is not very maintainable. You're balancing memory-management calls in two very different places, so you can't easily keep an eye on them both. Moreover, there are two different possible delegate messages. And, to top it all off, the intrusion of ARC's automatic retain and release behavior actually makes such a solution more difficult to implement; in effect, you're fighting the framework.

A better approach is to make your `context:` argument persistent as an instance variable. As I've already said, this is a lightweight approach under ARC, and we are now taking advantage of ARC's automatic retain and release behavior to manage its object's memory correctly with no effort on our part. The object may thus persist even after it is no longer needed, but it won't actually leak. Here's a complete example:

```

- (void) animate {
    // animcontext is an ivar typed as id
    // under ARC, this assignment retains the context and ensures that it persists
    animcontext = [NSDictionary dictionaryWithObject: @"object" forKey: @"key"];
    // prepare animation
    [UIView beginAnimations:@"shrinkImage" context:(__bridge void*)animcontext];
    [UIView setAnimationDelegate:self];
    [UIView setAnimationDidStopSelector:
     @selector(animationDidStop:finished:context:)];
    [imv setAlpha: 0];
    // request animation to start
    [UIView commitAnimations];
}

- (void) animationDidStop:(NSString*)anim
                     finished:(NSNumber *)f context:(void *)c {
    NSDictionary* d = (__bridge id)c; // cast back to dictionary, use as desired
    // ...
}

- (void) dealloc {
    // nothing to do here!
    // under ARC, animcontext will be released as we go out of existence
}

```

Considerations of this sort do not apply to parameters that are typed as objects. For instance, when you call `postNotificationName:object:userInfo:`, the `userInfo` is typed as an `NSDictionary` and is retained for you (and released after the notification is posted); its memory management behind the scenes is not your concern.

Memory Management of CFTypeRefs

A `CFTyperef` (see [Chapter 3](#)) is a value obtained through a C function that is a pointer to a struct; its type name will usually end in “Ref”. It is a kind of object, even though it isn’t a full-fledged Cocoa Objective-C object, and it must be managed in much the same way as a Cocoa object. *ARC is irrelevant to this fact.* ARC manages Objective-C objects; it has no concern with `CFTyprefs`. You must manage the memory of `CFTyprefs` manually, even if you’re using ARC. In fact, as I shall explain, the fact that you are using ARC actually *increases* the degree of your memory management responsibility.

I will divide the discussion into two halves: memory management of `CFTyprefs` on their own, and what happens when you “cross the bridge” between a `CFTyperef` and a full-fledged Objective-C object type.

Just as, in the Objective-C world of objects, certain method names (`alloc`, `copy`, and `retain`) alert you to your memory management responsibilities, so too in the world of `CFTyprefs`. The golden rule here is that if you obtained such an object through a function whose name contains the word `Create` or `Copy`, you are responsible for releasing it. In the case of a Core Foundation object (its type name begins with `CF`), you’ll release

it with the `CFRelease` function; other object creation functions are paired with their own object release functions.



An Objective-C object can be sent messages even if it is nil. But `CFRelease` cannot take a NULL argument. Be sure that a `CFTyperef` variable is not NULL before releasing it.

The matter is not a complicated one; it's much simpler than memory management of Cocoa objects, and the documentation will usually give you a hint about your memory management responsibilities. As an example, here (without further explanation) is some actual code from one of my apps, strongly modeled on Apple's own example code, in which I set up a base pattern color space (for drawing with a pattern):

```
- (void) addPattern: (CGContextRef) context color: (CGColorRef) incolor {
    CGColorSpaceRef baseSpace;
    CGColorSpaceRef patternSpace;
    baseSpace = CGColorSpaceCreateDeviceRGB ();
    patternSpace = CGColorSpaceCreatePattern (baseSpace);
    CGContextSetFillColorSpace (context, patternSpace);
    CGColorSpaceRelease (patternSpace);
    CGColorSpaceRelease (baseSpace);
    // ...
}
```

Never mind exactly what that code does; the important thing here is that the values for `baseSpace` and `patternSpace` are a `CFTyperef` (in particular, a `CGColorSpaceRef`) obtained through functions with `Create` in their name, so after we're done using them, we release them with the corresponding release function (here, `CGColorSpaceRelease`).

Similarly, you can retain a Core Foundation object, if you are afraid that it might go out of existence while you still need it, with the `CFRetain` function, and you are then, once again, responsible for releasing it with the `CFRelease` function.

We now come to the business of “crossing the bridge.” As I explained in [Chapter 3](#), many Core Foundation object types are toll-free bridged with a corresponding Cocoa object type. Now, from a theoretical point of view, memory management is memory management: it makes no difference whether you use Core Foundation memory management or Cocoa memory management. Thus, if you obtain a `CFStringRef` through a `Create` or `Copy` function and assign it to an `NSString` variable, sending `release` to it through the `NSString` variable is just as good as calling `CFRelease` on it. And before ARC, that was the end of the matter.

Under ARC, however, we face the same problem I described in the preceding section. ARC manages memory for Objective-C objects; it knows nothing of `CFTyprefs`. Therefore, ARC is not going to let you hand an object into or out of its memory-management purview without explicit information about how to manage its memory. This means a little extra thought for you, but it's a good thing, because it means you can tell ARC to do automatically what you would have done manually.

For example, a moment ago I said that, before ARC, you might obtain a CFStringRef through a `Create` or `Copy` function, cast it to an NSString, and later send `release` to it through the NSString. Under ARC, you can't say `release`, but you can arrange for ARC to do exactly the same thing: as you "cross the bridge", you pass the CFString through the `CFBridgingRelease()` function. The result is an `id` that can be assigned to an NSString variable, and that ARC will release to balance out the incremented retain count generated by the original `Create` or `Copy` function.

You have three choices as you cross the toll-free bridge:

`_bridge cast`

As illustrated in the previous section, you cast explicitly to the across-the-bridge type and qualify the cast with `_bridge`. This means that memory management responsibilities are independent on either side of the bridge. You're telling ARC that you're going to be performing complete and correct memory management on the CFTyperef side of the bridge.

`CFBridgingRelease()` function

You're crossing the bridge from the CFTyperef side to the object side. You're telling ARC that memory management for this object is incomplete: it has a raised retain count on the CFTyperef side (probably because you generated it with a `Create` or `Copy` function, or called `CFRetain` on it), and it will be up to ARC to perform the corresponding release on the object side. (Alternatively, you can do a `_bridge_transfer` cast.) Here's an artificial but correct example:

```
CFArrayRef arr_ref = CFLocaleCopyISOCountryCodes(); // note "copy"  
NSArray* arr = CFBridgingRelease(arr_ref); // memory management is now complete  
// ARC will manage this array correctly from here on
```

`CFBridgingRetain()` function

You're crossing the bridge from the object side to the CFTyperef side. You're telling ARC that it should leave memory management for this object incomplete: you're aware of the raised retain count on the object side, and you intend to call `CFRelease()` on it yourself on the CFTyperef side. (Alternatively, you can do a `_bridge_retained` cast.)

You may see `_bridge_transfer` and `_bridge_retained` in code written by other people, but I strongly recommend that you stick to `CFBridgingRelease()` and `CFBridgingRetain()` in your own code, as they are eminently clearer (and better named). Note that it is perfectly possible to pass an object out of the object world with `CFBridgingRetain()` and back into it later with `CFBridgingRelease()`.

Properties

A *property* (see [Chapter 5](#)) is syntactic sugar for calling an accessor by using dot-notation. An object does not have a property unless its class (or that class's superclass, of

course) declares it. For instance, in a earlier example we had an object with an `NSMutableArray` instance variable and a setter, which we called like this:

```
[self setTheData: d];
```

If this object's class code declares a property `theData`, we could instead say:

```
self.theData = d;
```

The effect would be exactly the same, because setting a property is just a shorthand for calling the setter. Similarly, suppose we were to say this:

```
NSMutableArray* arr = self.theData;
```

That is exactly the same as calling the getter.

Properties offer certain advantages that accessors, of themselves, do not:

- It is simpler to declare one property than to declare two accessor methods (a getter and a setter).
- With a property declaration, you can ask Cocoa to construct the accessors for you, automatically. Such an automatically constructed accessor is called a *synthesized accessor*. Writing accessors is boring and error-prone; with a property, you write two lines of code (the property declaration, and the request that Cocoa construct its accessors) and that's all: the accessors now exist, without you bothering to write them. Moreover, your class is now key-value coding compliant for the accessor name, with no effort on your part.
- With a property declaration, you don't have to declare an instance variable; if you don't supply an explicit instance variable declaration, and if you ask for a synthesized accessor, the instance variable will be synthesized automatically.
- A property declaration includes a statement of the setter's *memory management policy*. This statement has two purposes:
 - It lets you, the programmer, know easily, just by glancing at a property declaration, how the incoming value will be treated. You could find this out otherwise only by looking at the setter's code — which, if this is a built-in Cocoa type, you cannot do (and even in the case of your own code, it's a pain having to locate and consult the setter directly).
 - It tells the compiler what to do if you also ask for a synthesized setter (as well as if you let the instance variable be synthesized as well). The synthesized setter (and instance variable) will participate in managing memory for you, correctly, according to the policy you state in the property declaration.

The possible memory management policies correspond simply to what has already been said in this chapter about how a setter might behave, as well as the ARC reference types:

`strong, retain`

Under ARC, the instance variable will be a normal (strong) reference, so when ARC assigns the incoming value to it, it will retain the incoming value and release the

existing value of the instance variable. Under non-ARC, the setter’s own explicit code (or the code of the synthesized setter) will retain the incoming value and release the existing value of the instance variable. The terms are pure synonyms of one another and can be used in ARC or non-ARC code; `retain` is the term inherited, as it were, from pre-ARC days.

`copy`

The same as `strong` or `retain`, except that the incoming value is copied (by sending `copy` to it) and the copy, which has an increased retain count already, is assigned to the instance variable. This is appropriate particularly when a nonmutable class has a mutable subclass (such as `NSString` and `NSMutableString`, or `NSArray` and `NSMutableArray`), to prevent the setter’s caller from passing in an object of the mutable subclass; it is legal for the setter’s caller to do so, because (in accordance with polymorphism, [Chapter 5](#)) where an instance of a class is expected, an instance of its subclass can be passed, but the `copy` call creates an instance of the nonmutable class ([Chapter 10](#)).

`weak`

Under ARC, the instance variable will be a weak reference. ARC will assign the incoming value to it without retaining it. ARC will also magically nilify the instance variable if the instance to which it points goes out of existence. This is useful, as already explained earlier in this chapter, for breaking a potential retain cycle and for declining to retain inappropriately, and to reduce overhead where it is known that no memory management is needed, as with an interface object that is already retained by its superview. The setter can be synthesized only under ARC; using `weak` in non-ARC code is not strictly impossible but probably makes no sense.

`assign (the default)`

This policy is inherited from pre-ARC days; it is used in the same ways as `weak`. The setter does not manage memory; the incoming value is assigned directly to the instance variable. The instance variable, however, is *not* an ARC weak reference and will *not* be nilified automatically if the instance to which it points goes out of existence; it is a non-ARC weak reference (`__unsafe_unretained`) and can become a dangling pointer.

As I’ve already said, a property’s declared memory management policy is an instruction to the compiler if the setter is synthesized. If the setter is *not* synthesized, the declared memory management policy is “purely conventional” (as the LLVM documentation puts it), meaning that if you write your own setter, you’d better make that setter behave the way you declared you would, but nothing is going to force you to do so.

We come now to the formal syntax for declaring a property. A property is declared in the same part of a class’s interface section where you would declare methods. Its syntax schema is as follows:

```
@property (attribute, attribute, ...) type name;
```

Here's a real example, for the NSMutableArray instance variable we were talking about a moment ago:

```
@property (nonatomic, retain) NSMutableArray* theData;
```

The *type* and *name* will usually match the type and name of an instance variable, but what you're really indicating here are the name of the property (as used in dot-notation) and therefore the names of the setter (here, `setTheData:`) and getter (here, `theData`), and the type of value to be passed to the setter and obtained from the getter.

If this property will be represented by an outlet in a nib, you can say `IBOutlet` before the type. This is a hint to Xcode and has no formal meaning.

The *type* doesn't have to be an object type; it can be a simple type such as `BOOL`, `CGFloat`, or `CGSize`. Of course in that case no memory management is performed (as none is needed), and no memory management policy should be declared; but the advantages of using a property remain — the caller can use dot-notation, and the accessors can be synthesized.

The possible *attribute* values are:

A memory management policy

I listed the names of these a few paragraphs ago. You will supply exactly one; under ARC this will usually be `strong`. Do not omit it, as the default is `assign`; under ARC you'd be most unlikely to want that.

`nonatomic`

If omitted, the synthesized accessors will use locking to ensure correct operation if your app is multithreaded. This will rarely be a concern, and locking slows down the operation of the accessors, so you'll probably specify `nonatomic` most of the time. It's a pity that `nonatomic` isn't the default, but such is life.

`readonly` or `readonly`

If omitted, the default is `readwrite`. If you say `readonly`, any attempt to use the property as a setter will cause a compiler error (a useful feature), and if the accessors are to be synthesized, no setter is synthesized.

`getter=gname, setter=sname:`

By default, the property name is used to derive the names of the getter and setter methods that will be called when the property is used. If the property is named `myProp`, the default getter method name is `myProp` and the default setter name is `setMyProp:`. You can use either or both of these attributes to change that. If you say `getter=getALife`, you're saying that the getter method corresponding to this property is called `getALife` (and if the accessors are synthesized, the getter will be given this name).

To request that the accessors be synthesized for you, use the `@synthesize` directive. It appears anywhere inside the class's implementation section, any number of times, and takes a comma-separated list of property names. The behavior and names of the syn-

thesized accessors will accord with the property declaration attributes I've just talked about. You can state that the synthesized accessors should access an instance variable whose name differs from the property name by using the syntax *propertyName=iVarName* in the property name list; otherwise, the instance variable will have the same name as the property. If no such instance variable is declared, it will be synthesized as well.



A synthesized instance variable is strictly private, meaning that it is not inherited by subclasses. This fact will rarely prove troublesome, but if it does, simply declare the instance variable explicitly.

Thus, using our NSMutableArray instance variable `theData` as an example, the full code would look like this:

```
// [In the header file]
@interface MyClass : NSObject
@property (nonatomic, retain) NSMutableArray* theData;
@end

// [In the implementation file]
@implementation MyClass {
    // instance variables go here, but no need to declare theData
}
@synthesize theData
// other code goes here
@end
```

In the Xcode 4.2 application templates, a synthesized setter commonly takes advantage of the *propertyName=iVarName* syntax. For example, we find this at the top of the App-Delegate class's implementation section:

```
@synthesize window = _window;
```

The intention, apparently, is that we can refer in our code to the property explicitly as `self.window`, but if we were accidentally to refer to the instance variable directly as `window`, we'd get a compilation error, because there is no instance variable `window` (it's called `_window`). Thus the template adopts a policy designed both to prevent accidental direct access to the instance variable without passing through the accessors and to distinguish clearly in code which names are instance variables — they're the ones starting with an underscore. Moreover, this policy frees up the name `window` to be used as a local variable in a method, without getting a warning from the compiler that we're overshadowing the name of an instance variable. This can be a useful convention, and I have found myself using it increasingly in my own code.

To make a *property declaration private*, when its accessors are to be synthesized, put it in an anonymous category (a class extension), an interface section at the top of the implementation file. In this way, this class can access the instance variable through its accessors using dot-notation but other classes cannot ([Example 12-9](#)).

Example 12-9. A private property

```
// [In the implementation file]
@interface MyClass ()
@property (nonatomic, retain) NSMutableArray* theData; // private
@end

@implementation MyClass
@synthesize theData
// other code goes here
@end
```

Being able to declare private properties is so useful that I find myself almost routinely adding a class extension to the top of any new class files, to make it easy to add private properties later if I need to.

Another use of an anonymous category is to *redeclare* a property. For example, we might want our property to be `readonly` as far as the rest of the world knows, but `readwrite` for code within our class. To implement this, declare the property `readonly` in the interface section in the header file, which the rest of the world sees, and then redeclare it as `readwrite` in the anonymous category interface section in the implementation file, which only this class sees. All other attributes must match between both declarations.

If you do not ask explicitly that a declared property's accessors be synthesized, then you must supply them explicitly or the compiler will complain. This is somewhat annoying; one wishes that synthesis of accessors were the default, to save a step when writing and maintaining code. You can turn off this complaint by using `@dynamic` instead of `@synthesize`, but this is a promise to generate the accessors in some other way, at runtime, and is rarely used except in connection with Core Animation and Core Data. (An example of `@dynamic` with Core Animation appears in [Chapter 17](#).)

A useful trick is to take advantage of the `@synthesize` syntax `propertyName=ivarName` to override the synthesized accessor without losing any of its functionality. What I mean is this. Suppose you want the setter for `myIvar` to do more than just set `myIvar`. One possibility is to write your own setter; however, a setter from scratch is tedious and error-prone, whereas a synthesized setter does the job correctly and writing it is no work at all. The solution is to declare a property `myIvar` along with a corresponding private property ([Example 12-9](#))—let's call it `myIvarAlias`—and synthesize the private property `myIvarAlias` to access the `myIvar` instance variable. You must then write the accessors for `myIvar` by hand, but all they need to do, at a minimum, is use the `myIvarAlias` properties to set and get the value of `myIvar` respectively. The key point is that you can also do *other* stuff in those accessors ([Example 12-10](#)); whoever gets or sets the property `myIvar` will be doing that other stuff.

Example 12-10. Overriding synthesized accessors

```
// [In the header file]

@interface MyClass : NSObject
@property (nonatomic, retain) NSNumber* myIvar;
@end

// [In the implementation file]

@interface MyClass ()
@property (nonatomic, retain) NSNumber* myIvarAlias;
@end

@implementation MyClass
@synthesize myIvarAlias=myIvar;

- (void) setMyIvar: (NSNumber*) num {
    // do other stuff here
    self.myIvarAlias = num;
}

- (NSNumber*) myIvar {
    // do other stuff here
    return self.myIvarAlias;
}
@end
```

A property declaration can also appear in a protocol or category declaration. This makes sense because, with a property declaration, you're really just declaring accessor methods, and these are places where method declarations can go.



In Xcode 4.2, when you create an outlet by dragging from a nib object to code in such a way as to generate an instance variable declaration or a property declaration, Xcode creates some extra code for you. For example, when Xcode generates a property declaration in an interface section, it also generates the `@synthesize` directive in the implementation section, and if (as is often the case) this code belongs to a UIViewController, it will insert a `viewDidUnload` method (if need) that releases the nib object (by nilifying the instance variable through the setter). This is cool, but (1) it does make one wonder why Xcode 4 can't do this for *any* property or instance variable you create, and (2) it can be annoying if you change your mind later and have to backtrack to find and delete all that generated code.

Data Communication

As soon as an app grows to more than a few objects, things can become confusing. Beginners are sometimes puzzled about how to communicate data between one piece of code (one object, really) and another. The problem is essentially one of architecture. Constructing your code so that all the pieces fit together and key information can be shared is something of an art. But it isn't difficult. This chapter presents some general considerations that may provide the needed clue.

Model–View–Controller

In Apple's documentation and online, you will find references to the term *model–view–controller*, or *MVC*. This refers to an architectural goal of maintaining a distinction between three functional aspects of a program that displays information to the user and permits the user to alter that information. The whole notion goes back to the days of Smalltalk, and much has been written about it since then, but informally, here's what the terms mean:

Model

The data and its management (often referred to as the program's “business logic,” the hard-core stuff that the program is really all about).

View

What the user sees and interacts with.

Controller

The mediation between the model and the view.

Consider, for example, a game where the current score is displayed to the user:

- A `UILabel` that shows the user the current score for the game in progress is *view*; it is effectively nothing but a pixel-maker, and its business is to know how to draw itself. The knowledge of *what* it should draw — the score, and the fact that this *is* a score — lies elsewhere. A rookie programmer might try to use the score displayed by the `UILabel` as the actual score: to increment the score, read the `UILabel`'s

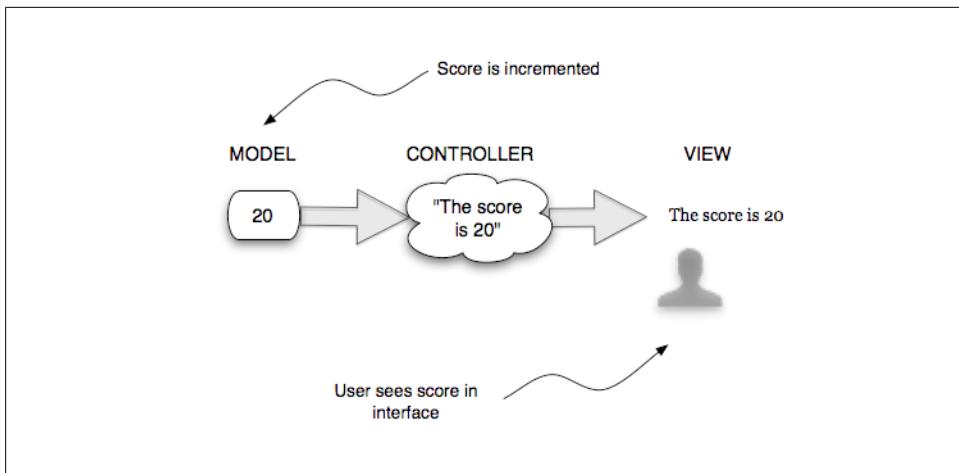


Figure 13-1. Model–view–controller

string, turn that string into a number, increment the number, turn the number back into a string, and present that string in place of the previous string. That is a gross violation of the MVC philosophy. The view presented to the user should *reflect* the score; it should not *store* the score.

- The score is data being maintained internally; it is *model*. It could be as simple as an instance variable along with a public `increment` method or as complicated as a Score object with a raft of methods. The score is numeric, whereas a `UILabel` displays a string; this alone is enough to show that the view and the model are naturally different.
- Telling the score when to change, and seeing that this fact is reflected in the user interface, is the work of the *controller*. This will be particularly clear if we imagine that the model's numeric score needs to be transformed in some way for presentation to the user. For example, suppose the `UILabel` that presents the score reads: “Your current score is 20”. The model is presumably storing and providing the number 20, so what’s the source of the phrase “Your current score is...”? Whoever is deciding that this phrase should precede the score in the presentation of the score to the user, and is making it so, is a controller.

Even this simplistic example (Figure 13-1) illustrates very well the advantages of MVC. By separating powers in this way, we allow the aspects of the program to evolve with a great degree of independence. Do you want a different font and size in the presentation of the score? Change the view; the model and controller need know nothing about it, but will just go on working exactly as they did before. Do you want to change the phrase that precedes the score? Change the controller; the model and view are unchanged.

Adherence to MVC is particularly appropriate in a Cocoa app, because Cocoa itself adheres to it. The very names of Cocoa classes reveal the MVC philosophy that underlies them. A `UIView` is a view. A `UIViewController` is a controller; its purpose is to

embody the logic that tells the view what to display. In [Chapter 11](#) we saw that a UIPickerView does not hold the data it displays; it gets that data from a data source. So the UIPickerView is a view; the data source is model.

Apple's documentation also points out this telling distinction: true model material and true view material should be quite reusable, in the sense that they can be transferred wholesale into some other app; controller material is generally not reusable, because it is concerned with how *this* app mediates between the model and the view.

In one of my own apps, for example, we download an XML (RSS) news feed and present the article titles to the user as a table. The storage and parsing of the XML are pure model material, and are so reusable that I didn't even write this part of the code (I used some code called FeedParser, by Kevin Ballard). The table is a UITableView, which is obviously reusable, seeing as I obtained it directly from Cocoa. But when the UITableView turns to me and asks what I'd like to display in this cell, and I turn to the XML and ask for the title of the article corresponding to this row of the table, that's controller logic.

By keeping the MVC architectural philosophy in mind as you develop your app, you'll implicitly solve one data communication problem. The data will live in the model, the view will be purely presentational in nature, and the communication between them will be handled by your own deliberately written controller code. You'll be communicating between the view and the model because controller code is *about* communicating between the view and model.

Instance Visibility

The problem of communication often comes down to one object being able to see another: object Manny needs to be able to find object Jack repeatedly and reliably over the long term so as to be able to send Jack messages. (This is the same problem I spoke of in [Chapter 2](#) as getting a *reference* to an object.)

An obvious solution is an instance variable of Manny whose value is Jack. This is appropriate particularly when Manny and Jack share certain responsibilities or supplement one another's functionality, and when they will both persist, especially when they will both persist together. A controller whose job is to configure and direct a certain view will need to exist just as long as the view does; they go together. The application object and its delegate, a table view and its data source, a UIViewController and its UIView — these are cases where the former must have an instance variable pointing at the latter.

With instance variables comes the question of memory management policy ([Chapter 12](#)). Should Manny, which has an instance variable pointing to Jack, also retain Jack? Basically, it depends on how closely allied the objects are. An object does not typically retain its delegate or its data source; it can exist without a delegate or a data source, and the delegate and data source have lives of their own — it is none of this

object's business to say whether the delegate or data source should be allowed to go out of existence. This object is therefore always prepared for the possibility that its delegate or data source may be nil. Similarly, an object that implements the target-action pattern, such as a UIControl, does not retain its target. On the other hand, a UIViewController is useless without a UIView to control; its very job is to be coterminous with its view, and to release its view when it itself goes out of existence. Similarly, an object that owns a nib as it loads rules the lifetimes of that nib's top-level objects.

Even when two objects go together closely, it will not necessarily be the case that each holds an instance variable pointing at the other. When each *does* point to the other, you must of course be careful not to let each *retain* the other; that's a retain cycle, and will cause both objects to leak. But if one object is the constant instigator of communication between the two, the first object can simply pass along a reference to itself as a method argument, if it thinks the second object might need this.

This behavior is conventional in a delegate message, for example. The parameter of the delegate message `textFieldShouldBeginEditing:` is a reference to the UITextField that sent the message. The same policy is followed by target-action messages in their fuller forms; the first parameter is a reference to the sender. You can follow a similar policy.

Visibility by Instantiation

The real question is how one object is to be introduced to the other in the first place. Much of the art of Cocoa programming (and of object-oriented programming generally) lies in getting a reference to a desired object. Every case is different and must be solved separately, but a major clue comes from the fact that every instance comes from somewhere. This means that some object commanded this instance to come into existence in the first place. That object therefore has a reference to the instance at that moment. That fact is always the starting point for establishment of future communication.

When Manny instantiates Jack, if Manny knows that Jack is going to need a reference to itself (Manny) or to some piece of data, Manny can hand Jack that reference early in Jack's lifetime. Indeed, you might write Jack with an initializer that will take this reference as a parameter, so that Jack will possess it from the moment it comes into existence. (Compare the approach taken, for example, by UIActionSheet and UIAlertView, where the delegate is one of the initializer's parameters, or by UIBarButtonItem, where the target is one of the initializer's parameters.)

This example, from one of my apps, is from a table view controller. The user has tapped a row of the table. We create a secondary table view controller, handing it the data it will need, and display the secondary table view. I deliberately devised TrackViewController to have a designated initializer `initWithMediaItemCollection:` to make it virtually obligatory for a TrackViewController to have access to the data it needs:

```

- (void)showItemsForRow: (NSIndexPath*) indexPath {
    // create subtable of tracks and go there
    TracksViewController *t =
        [[TracksViewController alloc] initWithMediaItemCollection:
            [self.albums objectAtIndex: [indexPath row]]];
    [self.navigationController pushViewController:t animated:YES];
}

```

The loading of a nib is also a case of visibility by instantiation. As I explained at length in [Chapter 7](#), you will want to prepare your nib objects, including some proxy object such as the File’s Owner, with outlets corresponding to instance variables or accessors in the corresponding class; as the nib loads, the object pointed to by each outlet will be handed to the corresponding instance variable as its value, or the setter as its argument, so that each instance that will need a reference to an object instantiated from the nib will in fact now have it and will be able to communicate as needed with the nib-instantiated object.

But what if two objects are conceptually distant from each other? A common case in point is when objects are going to be instantiated from *different* nibs. How can an instance from one nib get a reference to an instance from another nib? True, you can’t draw a connection between an object in nib A and an object in nib B. But someone (Manny) is going to be the file’s owner when nib A loads, and someone (Jack) is going to be the file’s owner when nib B loads. Those two file’s owners might be able to see each other; if so, the problem is solved. Perhaps they are the same object. Perhaps Manny instantiated Jack in the first place. Perhaps they are both instantiated by some third object, which provides a communication path for them.

Visibility by Relationship

Objects may acquire the ability to see one another automatically by virtue of their position in a built-in structure. Before worrying about how to supply one object with a reference to another, consider whether there may already be a chain of references leading from one to another.

For example, a subview can see its superview, through its `Superview` property. A superview can see all its subviews, through its `Subviews` property, and can pick out a specific subview through that subview’s `tag` property, by calling `viewWithTag:`. A subview in a window can see its window, through its `window` property. (There will be more about all that in [Chapter 14](#).) A responder ([Chapter 11](#)) can see the next responder in the responder chain, through the `nextResponder` method — which also means, because of the structure of the responder chain, that a `UIView` can see the `UIViewController` that manages it.

Similarly, if a `UIViewController` is currently presenting a view through a second `UIViewController`, the latter is the former’s `presentedViewController`, and the former is the latter’s `presentingViewController`. If a `UIViewController` is contained by a `UINavigationController`, the latter is its `navigationController`. A `UINavigationController`’s

visible view is controlled by its `visibleViewController`. And from any of these, you can reach the view controller's view through its `view` property, and so forth.

All of these relationships are public. So if you can get a reference to just one object within any of these structures or a similar structure, you can effectively navigate the whole structure through a chain of references and lay your hands on any other object within the structure.

Global Visibility

Some objects are globally visible (that is, visible to all other objects). In general, these are singletons vended by a class method. Some of these objects have properties pointing to other objects, making those other objects likewise globally visible.

For example, any object can see the singleton `UIApplication` instance by calling `[UIApplication sharedApplication]`. So any object can also see the app's primary window, because that is its `keyWindow` property, and any object can see the app delegate, because that is its `delegate` property. Thus, for example, in our Empty Window project, every object can see the `AppDelegate` instance created as the app launches. This means that any additional object can be made globally visible by designating a globally visible object, such as the app delegate, to hold a reference to it.

Another globally visible object is the shared defaults object obtained by calling `[NSUserDefaults standardUserDefaults]`. This object is the gateway to storage and retrieval of user defaults, which is similar to a dictionary (a collection of values named by keys). The user defaults are automatically saved when your application quits and are automatically available when your application is launched again later, so they are one of the main ways in which your app maintains state between launches. But, being globally visible, they are also a conduit for communicating values within your app.

For example, in one of my apps there's a setting I call `@"hazyStripy"`. This determines whether a certain visible interface object is drawn with a hazy fill or a stripy fill. This is a setting that the user can change, so there is a preferences interface allowing the user to make this change. When the user displays this preferences interface, I examine the `@"hazyStripy"` setting in the user defaults to configure the interface to reflect it; if the user interacts with the preferences interface to change the `@"hazyStripy"` setting, I respond by changing the actual `@"hazyStripy"` setting in the user defaults.

But the preferences interface is not the only object that uses the `@"hazyStripy"` setting in the user defaults; the drawing code that actually draws the hazy-or-stripy-filled object also uses it, so as to know which way to draw itself. Thus there no need for the object that draws the hazy-or-stripy-filled object and the object that manages the preferences interface to be able to see one another! They can both see this common object, the `@"hazyStripy"` user default ([Figure 13-2](#)). Indeed, it is not uncommon to “misuse” the user defaults storage to hold information that is *not* used to maintain state between

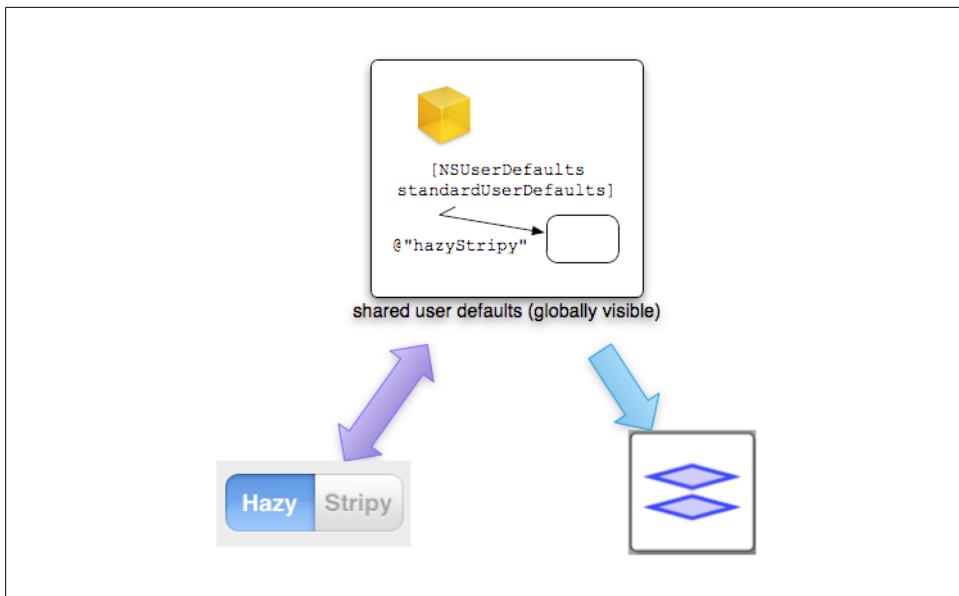


Figure 13-2. The global visibility of user defaults

runs of the app, but is placed there merely because this *is* a location globally visible to all objects.

Notifications

Notifications ([Chapter 11](#)) can be a way to communicate between objects that are conceptually distant from one another without bothering to provide *any* way for one to see the other. Using a notification in this way may seem lazy, an evasion of your responsibility to architect your objects sensibly. But sometimes one object doesn't need to know, and indeed shouldn't know, what object it is sending a message to.

I'll give a specific example. One of my apps consists of a bunch of flashcards. Only one card is showing at any one time, but the cards are actually embedded in a scroll view, so the user can move from one card to the next by swiping the screen. I've supplied classes to manage this interface: each card is managed by a CardController, and the scroll view as a whole is managed by a single ScrollViewController.

The flashcards all have the same layout: a foreign term, along with an English translation. To facilitate learning, the ScrollViewController displays a toolbar with a button that the user can tap to toggle visibility of the English translation *on all cards*. Thus, the user could hide the English translation and move from card to card, showing the English translation only occasionally to discover a forgotten translation or to confirm a remembered one.

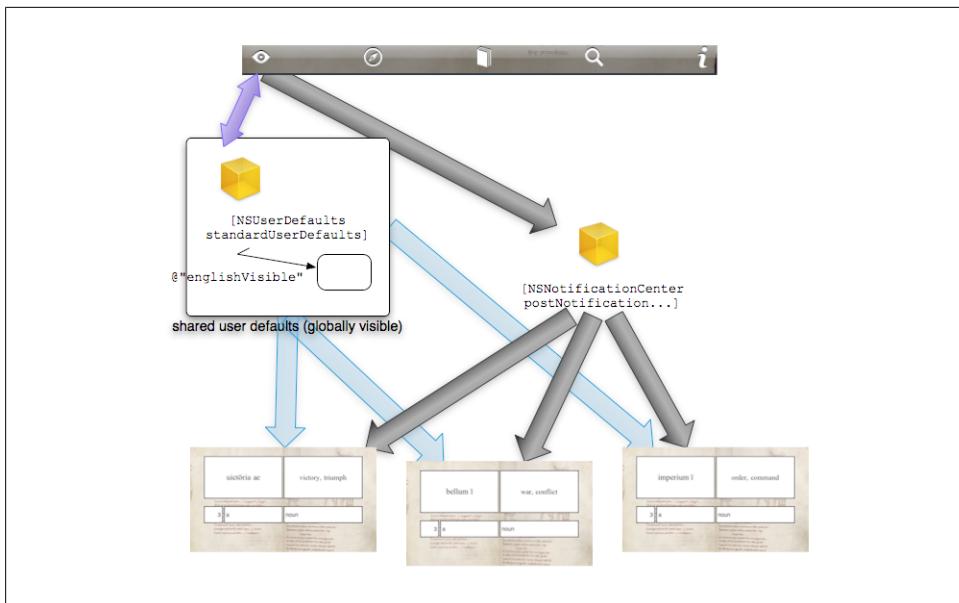


Figure 13-3. Posting a notification

What should happen, exactly, when the user taps the toolbar button? This seems a perfect use of a notification. I could cycle through all the existing `CardController` instances and tell each one to show or hide the English translation label, but this seems bulky and error-prone. How much simpler to have every `CardController`, as it comes into existence, register for the `@"toggleEnglish"` notification. Now the `ScrollViewController` can post a single notification and all the `CardController` instances will just hear about it, automatically.

This example also involves global storage in user defaults, discussed in the previous section. Just before posting the notification, `ScrollViewController` stores the desired state of the English translation's visibility in a user preference. Thus, to know what to do in response to the notification, each `CardController` just consults that preference and obeys it (Figure 13-3). Alternatively, `ScrollViewController` could have called `postNotificationName:object:userInfo:` and put this information into the `userInfo`, but there's a reason why I didn't do things that way: not every `CardController` that will ever exist is in existence now. I want each *new* `CardController` that *later* comes into existence to set the visibility of its English translation label correctly. To do so, it too has but to consult the globally available stored user preference.

Key–Value Observing

Key–value observing, or KVO, is a mechanism somewhat similar to the target–action mechanism, except that it is not limited to controls. (The KVO mechanism is provided

through an informal protocol, `NSKeyValueObserving`, which is actually a set of categories on `NSObject` and other classes.) The similarity is that objects register with a particular object to be notified when something happens. The “something” is that a certain value in that object is changed.

KVO can be broken down into three stages:

Registration

To hear about a change in a value belonging to object A, object B must be registered with object A.

Change

The change takes place in the value belonging to object A, and it must take place in a special way — a KVO compliant way.

Notification

Object B is notified that the value in object A has changed and can respond as desired.

Here’s a simple complete example — a rather silly example, but sufficient to demonstrate the KVO mechanism in action. We have a class `MyClass1`; this will be the class of object A. We also have a class `MyClass2`; this will be the class of object B. Finally, we have code that creates a `MyClass1` instance called `objectA` and a `MyClass2` instance called `objectB`; this code registers `objectB` to hear about changes in an instance variable of `objectA` called `value`, and then changes `value`, and sure enough, `objectB` is notified of the change:

```
// [In MyClass1.h]

@interface MyClass1 : NSObject
@property (nonatomic, copy) NSString* value;
@end

// [In MyClass1.m]

@implementation MyClass1
@synthesize value;
@end

// [In MyClass2.m (in its implementation section)]

- (void) observeValueForKeyPath:(NSString *)keyPath
    ofObject:(id)object
    change:(NSDictionary *)change
    context:(void *)context {
    NSLog(@"I heard about the change!");
}

// [somewhere else entirely]

MyClass1* objectA = [[MyClass1 alloc] init];
MyClass2* objectB = [[MyClass2 alloc] init];
```

```

// register for KVO
[objectA addObserver:objectB forKeyPath:@"value" options:0 context:NULL]; ①
// change the value in a KVO compliant way
objectA.value = @"Hello, world!"; ②
// result: objectB's observeValueForKeyPath:... is called

```

- ➊ We call `addObserver:forKeyPath:options:context:` to register `objectB` to hear about changes in `objectA`'s value. We didn't use any `options` or `context`; I'll talk about the `options` in a moment. (The `context` is for handing in a value that will be provided as part of the notification.)
- ➋ We change `objectA`'s `value`, and we do it in a KVO compliant way, namely, by passing through the setter (because setting a property is equivalent to passing through the setter). This is another reason why, as I said in [Chapter 12](#), accessors (and properties) are a good thing: they help you guarantee KVO compliance when changing a value.

When we change `objectA`'s `value`, the third stage takes place automatically: a call is made to `objectB`'s `observeValueForKeyPath:....`. We have implemented this method in `MyClass2` in order to receive the notification. In this simple example, we expect to receive only one notification, so we just log to indicate that we did indeed receive it. In real life, where a single object might be registered to receive more than one KVO notification, you'd use the incoming parameters to distinguish between different notifications and decide what to do.

At the very least, you'll probably want to know, when `observeValueForKeyPath:....` is called, what the new value is. We can find that out easily, because we are handed a reference to the object that changed, along with the key path for the value within that object. Thus we can use KVC to query the changed object in the most general way:

```

- (void) observeValueForKeyPath:(NSString *)keyPath
    ofObject:(id)object
    change:(NSDictionary *)change
    context:(void *)context {
    id newValue = [object valueForKeyPath:keyPath];
    NSLog(@"The key path %@ changed to %@", keyPath, newValue);
}

```

But it is also possible to request that the new value be included as part of the notification. This depends upon the `options` passed with the original registration. Here, we'll request that both the old and new values be included with the notification:

```

objectA.value = @"Hello";
[objectA addObserver:objectB forKeyPath:@"value"
    options: NSKeyValueObservingOptionNew | NSKeyValueObservingOptionOld
    context:NULL];
objectA.value = @"Goodbye"; // notification is triggered

```

When we receive the notification, we fetch the old and new values out of the `change` dictionary:

```

- (void) observeValueForKeyPath:(NSString *)keyPath
    ofObject:(id)object
    change:(NSDictionary *)change
    context:(void *)context {
    id newValue = [change objectForKey: NSKeyValueChangeNewKey];
    id oldValue = [change objectForKey: NSKeyValueChangeOldKey];
    NSLog(@"The key path %@ changed from %@ to %@", keyPath, oldValue, newValue);
}

```

No memory management happens as part of the registration process, so it is incumbent upon you to unregister object B before it is destroyed. Otherwise, object A may later attempt to send a notification to a dangling pointer. This is done by sending object A the `removeObserver:forKeyPath:` message.



The runtime will send you a nice warning in the log if an object *being observed* under KVO goes out of existence. But you get no warning if an *observer* object goes out of existence. It is crucial that the observer should be unregistered with the observed as the observer goes out of existence, to prevent the observed from trying to send it any notifications later. In real life, where an object will typically register itself as an observer, it will usually then unregister itself in its `dealloc` implementation (and this could be yet another reason to implement `dealloc` under ARC).

Beginners are often confused about how to use KVO to observe changes to a mutable array, to be notified when an object is added to, removed from, or replaced within the array. You can't add an observer to an array itself; you have to observe through an object that has a key path to the array (through accessors, for example). The simple-minded solution is then to access the array using `mutableArrayValueForKey:`, which provides an observable proxy object.

For example, recall how in [Chapter 12](#) we posited an object with an instance variable `theData` which is an array of dictionaries:

```

(
{
    description = "The one with glasses.";
    name = Manny;
},
{
    description = "Looks a little like Governor Dewey.";
    name = Moe;
},
{
    description = "The one without a mustache.";
    name = Jack;
}
)

```

Suppose this is an `NSMutableArray`. Then we can register with our object to observe the key path `@"theData"`:

```
[objectA addObserver:objectB forKeyPath:@"theData" options:0 context:NULL];
```

Now object B will be notified of changes to this mutable array, but only if those changes are performed through the `mutableArrayValueForKey:` proxy object:

```
[[objectA mutableArrayValueForKeyPath:@"theData"] removeObjectAtIndex:0];
// notification is triggered
```

But it seems onerous to require clients to know that they must call `mutableArrayValueForKey:`. The simple solution is for our object A itself to provide a getter that calls `mutableArrayValueForKey:`. Here's a possible implementation:

```
// [In MyClass1, in the header file]
```

```
@interface MyClass1 : NSObject
@property (nonatomic, strong, getter=theDataGetter) NSMutableArray* theData;
@end

// [In MyClass1, in the implementation section]

@synthesize theData;

- (NSMutableArray*) theDataGetter {
    return [self mutableArrayValueForKey:@"theData"];
}
```

The result is that, as far as any client knows, this object has a key @"theData" and a property `theData`, and we can register to observe with the key and then access the mutable array through the property:

```
[objectA addObserver:objectB forKeyPath:@"theData"
    options: NSKeyValueObservingOptionNew | NSKeyValueObservingOptionOld
    context:NULL];
[objectA.theData removeObjectAtIndex:0]; // notification is triggered
```

If you're going to take this approach, you should really also implement (in `MyClass1`) the four KVC compliance methods for a mutable array façade (see [Chapter 12](#)). Although things will appear to work just fine without them, and although they appear trivial (they are merely delegating to `self->theData` the equivalent calls), they will be called by the vended proxy object, which increases its efficiency (and, some would argue, its safety). Without these methods, the proxy object resorts to setting the instance variable directly, replacing the entire mutable array, every time a client changes the mutable array:

```
- (NSUInteger) countOfTheData {
    return [self->theData count];
}

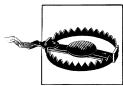
- (id) objectInTheDataAtIndex: (NSUInteger) ix {
    return [self->theData objectAtIndex: ix];
}

- (void) insertObject: (id) val inTheDataAtIndex: (NSUInteger) ix {
    [self->theData insertObject:val atIndex:ix];
```

```
}
```

```
- (void) removeObjectFromTheDataAtIndex: (NSUInteger) ix {
    [self->theData removeObjectAtIndex: ix];
}
```

If what you want to observe are mutations within an individual element of an array, things are more complicated. Suppose our array of dictionaries is an array of mutable dictionaries. To observe changes to the value of the @"`description`" key of any dictionary in the array, you'd need to register for that key with *each* dictionary in the array, separately. You can do that efficiently with `NSArray`'s instance method `addObserver:toObjectsAtIndexes:forKeyPath:options:context:`, but if the array *itself* is mutable then you're also going to have to register for that key with any *new* dictionaries that are subsequently added to the array (and unregister when a dictionary is removed from the array). This is doable but daunting, and I'm not going to go into the details here.



The properties of Apple's built-in classes are typically KVO compliant. Indeed, so are many classes that don't use properties per se; for example, `NSUserDefaults` is KVO compliant. Unfortunately, Apple warns that undocumented KVO compliance can't be counted on.

Key-value observing is a deep mechanism; consult Apple's *Key-Value Observing Guide* for full information. It does have some unfortunate shortcomings — for one thing, it's a pity that all notifications arrive by calling the same bottleneck method, `observeValueForKeyPath:...` — but in general KVO is useful for keeping values coordinated in different objects.



Mac OS X Programmer Alert

Mac OS X bindings don't exist on iOS, but you can sometimes use KVO to achieve similar aims.

PART IV

Views

This part of the book is about the things that appear in an app’s interface. All such things are, ultimately, *views*. A view is a unit of your app that knows how to draw itself. A view also knows how to sense that the user has touched it. Views are what your user sees on the screen, and what your user interacts with by touching the screen. Thus, views are the primary constituent of an app’s visible, touchable manifestation. They *are* your app’s interface. So it’s going to be crucial to know how views work.

- [Chapter 14](#) discusses views in their most general aspect — their hierarchy, position, and visibility.
- A view knows how to draw itself. [Chapter 15](#) is about drawing; it explains how to tell a view what you want it to draw, from simply displaying an already existing image to constructing a drawing line by line.
- The drawing power of a view comes ultimately from its *layer*. To put it another way, a layer is effectively the aspect of a view that knows how to draw — with even more power. [Chapter 16](#) explains about layers.
- A iOS app’s interface isn’t generally static; it’s lively. Much of that liveliness comes from animation. iOS gives you great power to animate your interface with remarkable ease; that power resides ultimately in layers. Now that you know about layers, [Chapter 17](#) tells about animation.
- A view knows how to draw itself; it also knows how to sense that the user is touching it. [Chapter 18](#) is about touches. It explains the iOS view-based mechanisms for sensing and responding to touches, with details on how touches are routed to the appropriate view and how you can customize that routing.

CHAPTER 14

Views

A *view* (an object whose class is a subclass of `UIView`) knows how to draw itself into a rectangular area of the interface. Your app has a visible interface thanks to views. Creating and configuring a view can be extremely simple: “Set it and forget it.” You’ve already seen that you can drag an interface widget, such as a `UIButton`, into a view in the nib; when the app runs, the button appears, and works properly. But you can also manipulate views in powerful ways, in real time. Your code can do some or all of the view’s drawing of itself; it can make the view appear and disappear, move, resize itself, and display many other physical changes, possibly with animation.

A view is also a responder (`UIView` is a subclass of `UIResponder`). This means that a view is subject to user interactions, such as taps and swipes. Thus, views are the basis not only of the interface that the user sees, but also of the interface that the user touches. Organizing your views so that the correct view reacts to a given touch allows you to allocate your code neatly and efficiently.

The *view hierarchy* is the chief mode of view organization. A view can have subviews; a subview has exactly one immediate superview. Thus there is a tree of views. This hierarchy allows views to come and go together. If a view is removed from the interface, its subviews are removed; if a view is hidden (made invisible), its subviews are hidden; if a view is moved, its subviews move with it; and other changes in a view are likewise shared with its subviews. The view hierarchy is also the basis of, though it is not identical to, the responder chain ([Chapter 11](#)).

A view may come from a nib, or you can create it in code. On balance, neither approach is to be preferred over the other; it depends on your needs and inclinations and on the overall architecture of your app.

The Window

The top of the view hierarchy is the app’s window. It is an instance of `UIWindow` (or your own subclass thereof), which is a `UIView` subclass. Your app should have exactly one main window. It occupies the entire screen and forms the background to, and the

ultimate superview of, all your other visible views. Other views are visible by virtue of being subviews, at some depth, of your app’s window. (If your app can display views on an external screen, you’ll create an additional UIWindow to contain those views; but in this chapter I’ll behave as if there were just one screen, the device’s own screen, and just one window.)

The project templates all generate your app’s window for you. In Xcode 4.2, the technique used, in a nonstoryboard app, is to create the window in code, in the app delegate’s `application:didFinishLaunchingWithOptions:`. The window must persist for the lifetime of the app, so the app delegate has a `window` property that retains it:

```
self.window = [[UIWindow alloc] initWithFrame:[[UIScreen mainScreen] bounds]];
```

The window’s designated initializer is `initWithFrame:`; I’ll explain in a moment what “frame” and “bounds” are, but the effect to make the window the same size as the screen. In the template, the comment, “Override point for customization after application launch,” comes *after* that line of code, because any code customizing what’s in the window will need the window to exist first. The various templates adopt various strategies for giving the window some content; in Xcode 4.2, this generally involves setting the window’s `rootViewController` to a UIViewController, whose `view` thus automatically becomes the window’s single primary subview. (I will refer to this as the window’s *root view*.)

For example, in the Single View Application project template, which we used for our earlier Empty Window example because it is the simplest complete template with a nib, the ViewController class is instantiated; the resulting instance is set to the window’s `rootViewController`, and the UIView that serves as the value of the ViewController instance’s `view` property, which was set from the File’s Owner’s `view` outlet in *ViewController.xib*, becomes the window’s single primary subview. (I’ll discuss this mechanism in detail in [Chapter 19](#).) Finally, the template code sends the window instance the `makeKeyAndVisible` message in order to make your app’s interface appear.

On the other hand, if you choose the Storyboard option as you specify a template, the app is given a main storyboard, pointed to by the *Info.plist* key “Main storyboard file base name” (`UIStoryboardName`). In that case, after `UIApplicationMain` instantiates the app delegate class ([Chapter 6](#)), it asks the app delegate for the value of its `window` property; if that value is nil, the window is created and assigned to the app delegate’s `window` property. The storyboard’s initial view controller is then instantiated and assigned to the window’s `rootViewController` property, with the result that its view is placed in the window as its root view. The window is then sent the `makeKeyAndVisible` message. All of that is done behind the scenes by `UIApplicationMain`, with no visible code whatever. That is why, in a storyboard template, the `application:didFinishLaunchingWithOptions:` implementation is empty.

It is improbable that you would want to subclass UIWindow and substitute an instance of your subclass as the app’s main window, but you can certainly do so. If the window is generated explicitly in code, you would obviously substitute the name of your win-

dow subclass as the class to be instantiated and assigned to the app delegate's `window` property in `application:didFinishLaunchingWithOptions:`, in the template code I quoted a moment ago. If you're using a main storyboard, however, `application:didFinishLaunchingWithOptions:` is too late; you'll have to perform the substitution when `UIApplicationMain` asks for the app delegate's `window` property, by implementing its getter so as to create the window and set the `window` property exactly once:

```
- (UIWindow*) window {
    UIWindow* w = self->_window;
    if (!w) {
        w = [[MyWindow alloc] initWithFrame:[[UIScreen mainScreen] bounds]];
        self->_window = w;
    }
    return w;
}
```

Once the app is and running, the app delegate forever points to the window as the value of its `window` property; so any code in the app delegate class can refer to the window as `self.window`. Code elsewhere can also get a reference to the app's window; here, I do it by way of the synthesized accessor:

```
UIWindow* theWindow = [[[UIApplication sharedApplication] delegate] window];
```

That code is unusual, though, and may require typecasting to quiet the compiler (because the class of the application's `delegate` property is otherwise unknown). You'd be more likely to use the application's `keyWindow` property:

```
UIWindow* theWindow = [[UIApplication sharedApplication] keyWindow];
```

An even more likely way to get a reference to your app's window would be through a subview of the window, at any depth of the hierarchy. You are very likely to have a reference to at least one such subview, and its `window` property points to the window that contains it, which is the app's window. You can also use a `UIView`'s `window` property as a way of asking whether it is ultimately embedded in a window; if it isn't, its `window` property is nil. A `UIView` whose `window` property is nil cannot be visible to the user.

Although your app will have exactly one primary window, it may generate other windows of which you are not conscious. For example, if you put up an alert view (`UIAlertView`), it is displayed in a secondary window that lies on top of your app's window; at that moment, this secondary window is the application's `keyWindow`. You would not be conscious of this fact, however, unless you needed a reference to your app's window while an alert was showing, which is unlikely.

The window's `backgroundColor` property, which it inherits from `UIView`, affects the appearance of the app if the window is visible behind its subviews. In many cases, though, you are likely to give your window a primary subview that occupies the entire window and blocks it from sight; the window's `backgroundColor` would then make no visible difference. The window would function solely as a container for the app's visible views.

As you work your way through this and subsequent chapters, you may be tempted to try out the example code, or code of your own, to experiment with views and how they are drawn. In iOS 4 and before, it was legal to put a subview directly into the window, so you could start with the Empty Application project template and, in `application:didFinishLaunchingWithOptions:`, create the view and add it as a subview to the window:

```
self.window = [[UIWindow alloc] initWithFrame:[[UIScreen mainScreen] bounds]];
// Override point for customization after application launch.
UIView* v = [[UIView alloc] initWithFrame:CGRectMake(100,100,50,50)];
v.backgroundColor = [UIColor redColor]; // small red square
[self.window addSubview: v];
self.window.backgroundColor = [UIColor whiteColor];
[self.window makeKeyAndVisible];
return YES;
```

This still works, but in iOS 5, the runtime complains if the window does not have a `rootViewController`. The minimal solution, permitting you to experiment with views created in code in an app drawn from the Empty Application project template, but without getting that warning from the runtime, is to instantiate a root view controller in code, assign it to the `rootViewController` property, and draw your experimental view into the root view controller's view instead, like this:

```
self.window = [[UIWindow alloc] initWithFrame:[[UIScreen mainScreen] bounds]];
// Override point for customization after application launch.
self.window.rootViewController = [UIViewController new];
UIView* v = [[UIView alloc] initWithFrame:CGRectMake(100,100,50,50)];
v.backgroundColor = [UIColor redColor]; // small red square
[self.window.rootViewController.view addSubview: v];
self.window.backgroundColor = [UIColor whiteColor];
[self.window makeKeyAndVisible];
return YES;
```

The result looks exactly the same, but there is now an invisible (transparent) root view occupying the entire window, and the small red square is actually a subview of that.

On the other hand, you might also want to experiment with views created in a nib. In Xcode 4.2, unlike earlier versions of Xcode, the Empty Application project template has no nib, so it will be simplest to start with the Single View Application project template, as we did with our earlier Empty Window example. The view supplied in the nib will become the window's root view, and whatever you drag into it in the nib will appear in the window when the app runs.

Subview and Superview

Once upon a time, and not so very long ago, a view owned precisely its rectangular area. No part of any view that was not a subview of this view could appear inside it, because when this view redrew its rectangle, it would erase the overlapping portion of

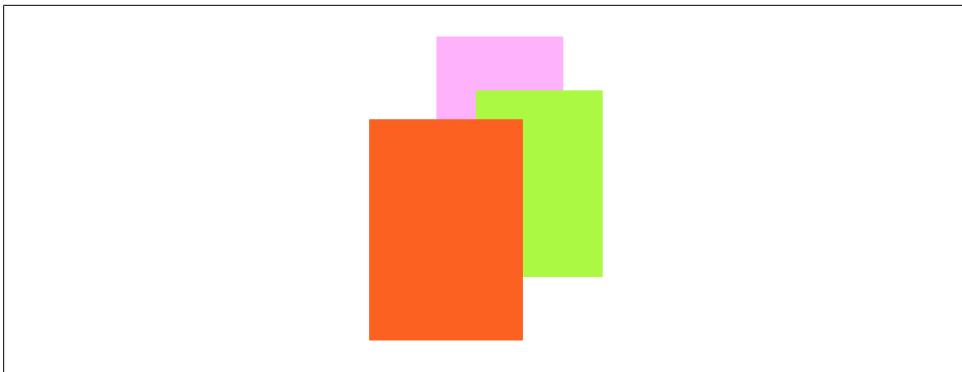


Figure 14-1. Overlapping views

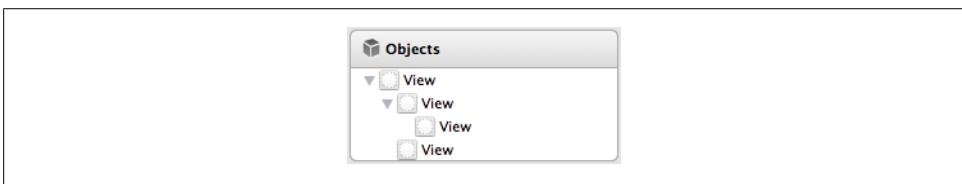


Figure 14-2. A view hierarchy as displayed in the nib

the other view. No part of anySubview of this view could appear outside it, because the view took responsibility for its own rectangle and no more.

Those rules, however, were gradually relaxed, and starting in Mac OS X 10.5 Apple introduced an entirely new architecture for view drawing that lifted those restrictions completely. iOS view drawing is based on this revised architecture. So now some or all of a subview can appear outside its superview, and a view can overlap another view and be drawn partially or totally in front of it without being its subview.

So, for example, [Figure 14-1](#) shows three overlapping views. All three views have a background color, so each is completely represented by a colored rectangle. You have no way of knowing, from this visual representation, exactly how the views are related within the view hierarchy. In actual fact, the view in the middle (horizontally) is a sibling view of the view on the left (they are both direct subviews of the root view), and the view on the right is a subview of the middle view.

When views are created in the nib, you can examine the view hierarchy in the expanded dock to learn their actual relationship ([Figure 14-2](#)). When views are created in code, you know their hierarchical relationship because you created that hierarchy. But the visible interface doesn't tell you, because view overlapping is so flexible.

Nevertheless, a view's position in the view hierarchy does affect how it is drawn. Most important, a view's position in the view hierarchy dictates the *order* in which it is drawn. Sibling subviews of the same superview have a layering order: one is "further back" than the other. This will make no visible difference if there is no overlap, but the subview

that is “further back” is drawn first, so if there *is* overlap, it will appear to be behind its sibling. Similarly, a superview is “further back” than its subviews; the superview is drawn first, so it will appear to be behind its subviews.

You can see this illustrated in [Figure 14-1](#). The view on the right is a subview of the view in the middle and is drawn on top of it. The view on the left is a sibling of the view in the middle, but it is a later sibling, so it is drawn on top of the view in the middle and on top of the view on the right. The view on the left *cannot* appear behind the view on the right but in front of the view in the middle, because those views are subview and superview and are drawn together — both are drawn either before or after the view on the left, depending on the “further back” ordering of the siblings.

This layering order can be governed in the nib by arranging the views in the expanded dock. (If you click in the canvas, you may be able to use the menu items of the Editor → Arrange menu instead — Send to Front, Send to Back, Send Forward, Send Backward.) In code, there are methods for arranging the sibling order of views, which we’ll come to in a moment.

Here are some other effects of the view hierarchy:

- If a view is removed from or moved within its superview, its subviews go with it.
- If a view’s size is changed, its subviews can be resized automatically.
- A view’s degree of transparency is inherited by its subviews.
- A view can optionally limit the drawing of its subviews so that any parts of them outside the view are not shown. This is called *clipping* and is set with the view’s `clipsToBounds` property.
- A superview *owns* its subviews, in the memory-management sense, much as an `NSArray` owns its elements; it retains them and is responsible for releasing a subview when that subview ceases to be its subview (it is removed from the collection of this view’s subviews) or when it itself goes out of existence.

A `UIView` has a `Superview` property (a `UIView`) and a `subviews` property (an `NSArray` of `UIViews`, in back-to-front order), allowing you to trace the view hierarchy in code. There is also a method `isDescendantOfView:` letting you check whether one view is a subview of another at any depth. If you need a reference to a particular view, you will probably arrange this beforehand as an instance variable, perhaps through an outlet. Alternatively, a view can have a numeric tag (its `tag` property), and can then be referred to by sending any view higher up the view hierarchy the `viewWithTag:` message. Seeing that all tags of interest are unique within their region of the hierarchy is up to you.

Manipulating the view hierarchy in code is easy. This is part of what gives iOS apps their dynamic quality, and it compensates for the fact that there is basically just a single window. It is perfectly reasonable for your code to rip an entire hierarchy of views out of the superview and substitute another. Such behavior can be implemented elegantly by using a `UIViewController`, a subject to which we’ll return later. But you can do it directly, too. The method `addSubview:` makes one view a subview of another; `remove-`

`FromSuperview` takes a subview out of its superview's view hierarchy. In both cases, if the superview is part of the visible interface, the subview will appear or disappear; and of course this view may itself have subviews that accompany it. Just remember that removing a subview from its superview releases it; if you intend to reuse that subview later on, you will wish to retain it first. This is often taken care of through a property with a retain policy.

Events inform a view of these dynamic changes. To respond to these events requires subclassing. Then you'll be able to override any of `didAddSubview:` and `willRemoveSubview:`, `didMoveToSuperview` and `willMoveToSuperview:`, `didMoveToWindow` and `willMoveToWindow:`.

When `addSubview:` is called, the view is placed last among its superview's subviews; thus it is drawn last, meaning that it appears frontmost. A view's subviews are indexed, starting at 0, which is rearmost. There are additional methods for inserting a subview at a given index (`insertSubview:atIndex:`), or below (behind) or above (in front of) a specific view (`insertSubview:belowSubview:`, `insertSubview:aboveSubview:`); for swapping two sibling views by index (`exchangeSubviewAtIndex:withSubviewAtIndex:`); and for moving a subview all the way to the front or back among its siblings (`bringSubviewToFront:`, `sendSubviewToBack:`).

Oddly, there is no command for removing all of a view's subviews at once. However, a view's `subviews` array is an immutable copy of the internal list of subviews, so it is legal to cycle through it and remove each subview one at a time:

```
for (UIView* v in view.subviews)
    [v removeFromSuperview];
```

Frame

A view's `frame` property, a `CGRect`, is the position of its rectangle within its superview, *in the superview's coordinate system*. By default, the superview's coordinate system will have the origin at its top left, with the x-coordinate growing positively rightward and the y-coordinate growing positively downward.

Setting a view's frame to a different `CGRect` value repositions the view, or resizes it, or both. If the view is visible, this change will be visibly reflected in the interface. On the other hand, you can also set a view's frame when the view is not visible — for example, when you create the view in code. In that case, the frame describes where the frame *will* be positioned within its superview when it is given a superview. `UIView`'s designated initializer is `initWithFrame:`, and you'll often assign a frame this way, especially because the default frame might otherwise be `{0,0},{0,0}`, which is rarely what you want.



Forgetting to assign a view a frame when creating it in code, and then wondering why it isn't appearing when added to a superview, is a common beginner mistake. A view with a zero-size frame is effectively invisible. If a view has a standard size that you want it to adopt, especially in relation to its contents (like a `UIButton` in relation to its title), an alternative is to send it the `sizeToFit` message.

Knowing this, we can generate programmatically the interface displayed in [Figure 14-1](#). This code might appear in the `application:didFinishLaunchingWithOptions:` method of the app delegate in an Empty Application template-based app (as I suggested earlier):

```
UIView* v1 = [[UIView alloc] initWithFrame:CGRectMake(113, 111, 132, 194)];  
v1.backgroundColor = [UIColor colorWithRed:1 green:.4 blue:1 alpha:1];  
UIView* v2 = [[UIView alloc] initWithFrame:CGRectMake(41, 56, 132, 194)];  
v2.backgroundColor = [UIColor colorWithRed:.5 green:1 blue:0 alpha:1];  
UIView* v3 = [[UIView alloc] initWithFrame:CGRectMake(43, 197, 160, 230)];  
v3.backgroundColor = [UIColor colorWithRed:1 green:0 blue:0 alpha:1];  
[self.window.rootViewController.view addSubview: v1];  
[v1 addSubview: v2];  
[self.window.rootViewController.view addSubview: v3];
```

In that code, we determined the layering order of `v1` and `v3` (the middle and left views, which are sibling subviews of the window) by the order in which we inserted them into the view hierarchy with `addSubview:`.

Part of the app's window may be covered by the status bar, which is actually another window, supplied by the system. This may affect where you want to draw in the window. A view centered within the window will be centered on the screen, but it may not look centered because it isn't centered in the visible part of the window exclusive of the status bar. Similarly, material drawn in the window at the point `{0,0}` (in the window's coordinates) may not be visible, because that point may be covered by the status bar. You can determine the rectangle currently not covered by the status bar as follows:

```
CGRect f = [[UIScreen mainScreen] applicationFrame];
```

By default, the `applicationFrame` is the frame of the window's root view, so a subview of the root view with a frame origin of `{0,0}` is *not* covered by the status bar (because the subview is positioned using the root view's coordinates, not the window's coordinates).

Complications are introduced by the possibility of the user rotating the device. This does not change anything about the window's coordinate system, so the window's `{0,0}` point might be in any corner (from the user's perspective). One of the main benefits of covering the window with a root view managed by a `UIViewController` is that the `UIViewController` deals with rotation. I'll discuss that in [Chapter 19](#); for now, I'll assume that the device is not rotated.

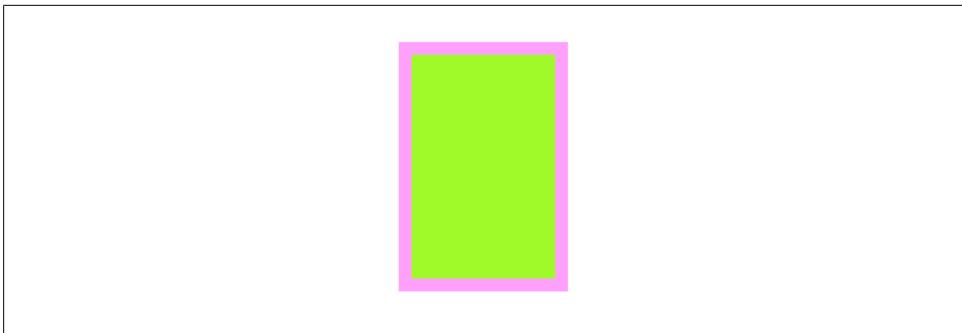


Figure 14-3. A subview inset from its superview

Bounds and Center

Suppose we wish to give a view a subview inset by 10 pixels, as in [Figure 14-3](#). The utility function `CGRectInset` makes it easy to derive one rectangle as an inset from another, but *what* rectangle should we use as a basis? Not the superview's frame; the frame represents a view's position within *its* superview, and in that superview's coordinates. What we're after is a `CGRect` describing our superview's rectangle in its *own* coordinates, because those are the coordinates in which the subview's frame is to be expressed. That `CGRect` is the view's `bounds` property.

So, the code to generate [Figure 14-3](#) looks like this:

```
UIView* v1 = [[UIView alloc] initWithFrame:CGRectMake(113, 111, 132, 194)];  
v1.backgroundColor = [UIColor colorWithRed:1 green:.4 blue:1 alpha:1];  
UIView* v2 = [[UIView alloc] initWithFrame:CGRectInset(v1.bounds, 10, 10)];  
v2.backgroundColor = [UIColor colorWithRed:.5 green:1 blue:0 alpha:1];  
[self.window.rootViewController.view addSubview: v1];  
[v1 addSubview: v2];
```

You'll very often use a view's `bounds` in this way. When you need coordinates for drawing inside a view, whether drawing manually or placing a subview, you'll often refer to the view's `bounds`.



The screen also has bounds, and functions in that sense as the window's superview, even though a `UIScreen` isn't a view. Moreover, the window's frame is always set to the screen's bounds (see the example earlier in this chapter of creating a window in code). Thus, *window coordinates are screen coordinates*. For example, when asking the screen for the `applicationFrame`, the answer comes back in screen coordinates, which are also window coordinates, and can thus be used for positioning something within the window.

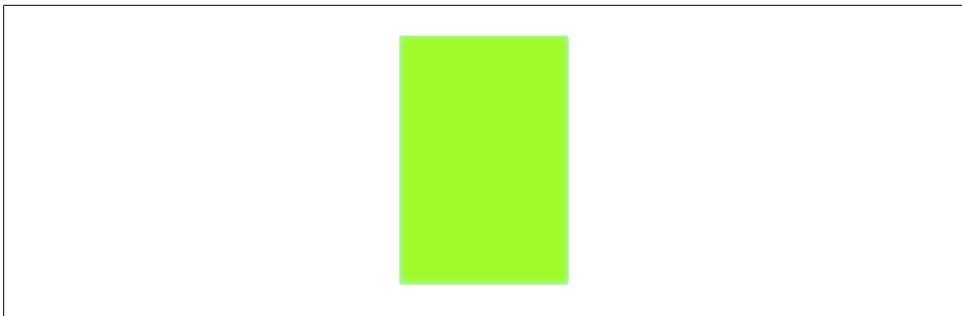


Figure 14-4. A subview exactly covering its superview

Interesting things happen when you set a view’s bounds. If you change a view’s bounds *size*, you change its *frame*. The change in the view’s frame takes place around its *center*, which remains unchanged. So, for example:

```
UIView* v1 = [[UIView alloc] initWithFrame:CGRectMake(113, 111, 132, 194)];  
v1.backgroundColor = [UIColor colorWithRed:1 green:.4 blue:1 alpha:1];  
UIView* v2 = [[UIView alloc] initWithFrame:CGRectMakeInset(v1.bounds, 10, 10)];  
v2.backgroundColor = [UIColor colorWithRed:0.5 green:1 blue:0 alpha:1];  
[self.window.rootViewController.view addSubview: v1];  
[v1 addSubview: v2];  
CGRect f = v2.bounds;  
f.size.height += 20;  
f.size.width += 20;  
v2.bounds = f;
```

What appears is a single rectangle; the subview completely and exactly covers its superview, its frame being the same as the superview’s bounds. The call to `CGRectInset` started with the superview’s bounds and shaved 10 points off the left, right, top, and bottom to set the subview’s frame ([Figure 14-3](#)). But then we added 20 points to the subview’s bounds height and width, and thus added 20 points to the subview’s frame height and width as well ([Figure 14-4](#)). The center didn’t move, so we effectively put the 10 points back onto the left, right, top, and bottom of the subview’s frame.

When you create a `UIView`, its bounds coordinate system’s $\{0,0\}$ point is at its top left. If you change a view’s bounds *origin*, you move the *origin of its internal coordinate system*. Because a subview is positioned in its superview with respect its superview’s coordinate system, such a change will change the apparent position of a subview relative to its superview. To illustrate, we start with our subview inset evenly within its superview, and then change the bounds origin of the superview:

```
UIView* v1 = [[UIView alloc] initWithFrame:CGRectMake(113, 111, 132, 194)];  
v1.backgroundColor = [UIColor colorWithRed:1 green:.4 blue:1 alpha:1];  
UIView* v2 = [[UIView alloc] initWithFrame:CGRectMakeInset(v1.bounds, 10, 10)];  
v2.backgroundColor = [UIColor colorWithRed:0.5 green:1 blue:0 alpha:1];  
[self.window.rootViewController.view addSubview: v1];  
[v1 addSubview: v2];
```

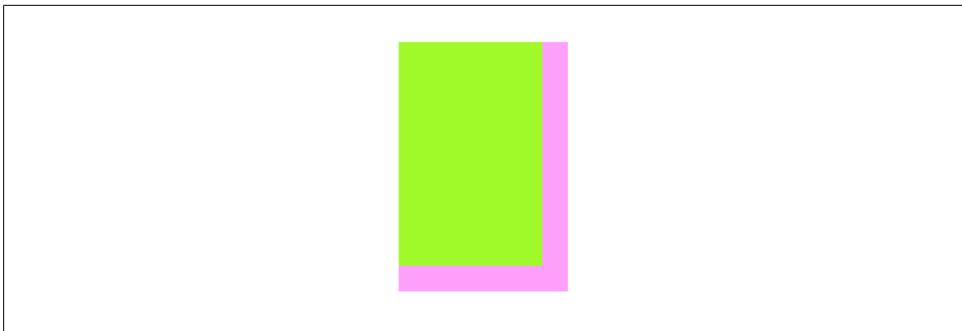


Figure 14-5. The superview's bounds origin has been shifted

```
CGRect f = v1.bounds;
f.origin.x += 10;
f.origin.y += 10;
v1.bounds = f;
```

Nothing happens to the superview's size or position. But the subview has moved up and to the left so that it is flush with its superview's top left corner ([Figure 14-5](#)). Basically, what we've done is to say to the superview, "Instead of calling the point at your upper left {0,0}, call that point {10,10}." Because the subview's frame origin is itself at {10,10}, the subview now touches the superview's top left corner. The effect of changing a view's bounds origin may seem directionally backward — we increased the superview's origin in the positive direction, but the subview moved in the negative direction — but think of it this way: a view's bounds origin point coincides with its frame's top left.

We have seen that changing a view's bounds size affects its frame size. The converse is also true: changing a view's frame size affects its bounds size. What is *not* affected by changing a view's bounds size is the view's **center**. This property, like the **frame** property, represents the view's position within its superview, in the superview's coordinates, but it is the position of the bounds center, the point derived from the bounds like this:

```
CGPoint c = CGPointMake(CGRectGetMidX(theView.bounds),
                        CGRectGetMidY(theView.bounds));
```

A view's center is thus a single point establishing the positional relationship between a view's bounds and its superview's bounds. Changing a view's bounds does not change its center (we already saw that when we increased a view's bounds size, its frame expanded around a stationary center); changing a view's center does not change its bounds.

Thus, a view's bounds and center are orthogonal (independent), and describe (among other things) both the view's size and its position within its superview. The view's frame is therefore superfluous! In fact, the **frame** property is merely a convenient expression of the **center** and **bounds** values. In most cases, this won't matter to you; you'll use the **frame** property anyway. When you first create a view from scratch, the designated initializer is **initWithFrame:**. You can change the frame, and the bounds size and center

will change to match. You can change the bounds size or the center, and the frame will change to match. Nevertheless, the proper and most reliable way to position and size a view within its superview is to use its bounds and center, not its frame; there are some situations in which the frame is meaningless (or will at least behave very oddly), but the bounds and center will always work.

We have seen that every view has its own coordinate system, expressed by its `bounds`, and that a view's coordinate system has a clear relationship to its superview's coordinate system, expressed by its `center`. This is true of every view in a window, so it is possible to convert between the coordinates of any two views in the same window. Convenience methods are supplied to perform this conversion both for a `CGPoint` and for a `CGRect`: `convertPoint:fromView:`, `convertPoint:toView:`, `convertRect:fromView:`, and `convertRect:toView:`. If the second parameter is nil, it is taken to be the window.

For example, if `v2` is a subview of `v1`, then to center `v2` within `v1` you could say:

```
v2.center = [v1 convertPoint:v1.center fromView:v1.superview];
```



When setting a view's position by setting its center, if the height or width of the view is not an even integer, the view can end up *misaligned* (on a normal-resolution screen): its point values in one or both dimensions are located between the screen pixels. This can cause the view to be displayed incorrectly; for example, if the view contains text, the text may be blurry. You can detect this situation in the Simulator by checking `Debug → Color Misaligned Images`. A simple solution is to set the view's frame, after positioning it, to the `CGRectIntegral` of its frame.

Layout

We have seen that a subview moves when its superview's bounds *origin* is changed. But what happens to a subview when its superview's bounds *size* is changed? (And remember, this includes changing the superview's frame size.)

Of its own accord, nothing happens. The subview's bounds and center haven't changed, and the superview's bounds origin hasn't moved, so the subview stays in the same position relative to the top left of its superview. In real life, however, that often won't be what you want. You'll want subviews to be resized and repositioned when their superview's bounds size is changed. This is called *layout*.

The need for layout is obvious in a context such as Mac OS X, where the user can freely resize a window, potentially disturbing your interface. For example, you'd want an OK button near the lower right corner to stay in the lower right corner as the window grows, while a text field at the top of the window should stay at the top of the window, but perhaps should widen as the window widens.

There are no user-resizable windows on an iOS device, but still, a superview might be resized dynamically. For example, you might respond to the user rotating the device 90 degrees by swapping the width and height values of a view; now its subviews should

shift to compensate. Or you might want to provide a reusable complex view, such as a table view cell containing several subviews, without knowing its precise final dimensions in advance.

Layout is performed in two primary ways, which can be combined:

Automatic layout

Automatic resizing of subviews depends on the superview's `autoresizesSubviews` property. To turn off a view's automatic resizing altogether, set this property to NO. If it is YES, then a subview will respond automatically to its superview's being resized, in accordance with the rules prescribed by the subview's `autoresizingMask` property value.

Manual layout

The superview is sent the `layoutSubviews` message whenever it is resized; so, to lay out subviews manually, provide your own subclass and override `layoutSubviews`. If you're going to use both approaches, automatic resizing is performed before `layoutSubviews` is called.

You should never call `layoutSubviews` yourself. Instead, if you wish to trigger layout, send `setNeedsLayout` to the view. This will cause the layout procedures to be followed at the next appropriate moment. Alternatively, if you really need layout to occur right this moment, send the view the `layoutIfNeeded` message; this may cause the layout of the entire view tree, not only below but also above this view, and is probably not a very common thing to do.

Automatic resizing is a matter of conceptually assigning a subview "springs and struts." A spring can stretch; a strut can't. Springs and struts can be assigned internally or externally. Thus you can specify, using internal springs and struts, whether and how the view can be resized, and, using external springs and struts, whether and how the view can be repositioned. For example:

- Imagine a subview that is centered in its superview and is to stay centered, but is to resize itself as the superview is resized. It would have struts externally and springs internally.
- Imagine a subview that is centered in its superview and is to stay centered, and is *not* to resize itself as the superview is resized. It would have springs externally and struts internally.
- Imagine an OK button that is to stay in the lower right of its superview. It would have struts internally, struts externally to its right and bottom, and springs externally to its top and left.
- Imagine a text field that is to stay at the top of its superview. It is to widen as the superview widens. It would have struts externally; internally it would have a vertical strut and a horizontal spring.

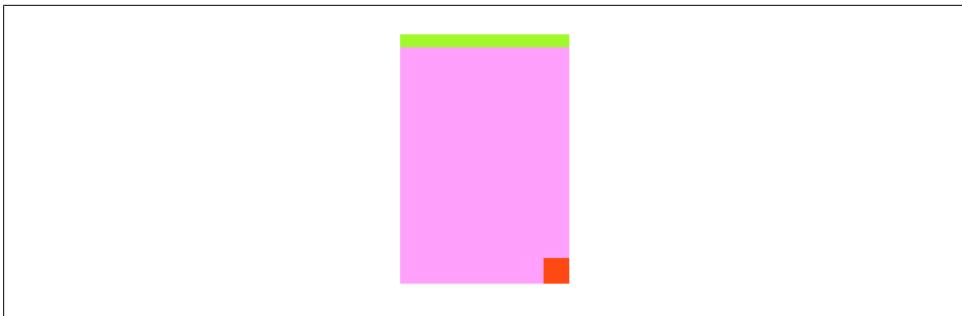


Figure 14-6. Before autoresizing

When editing a nib file, you can experiment with assigning a view springs and struts in the Size inspector (Autosizing). A solid line externally represents a strut; a solid line internally represents a spring. A helpful animation shows you the effect on your view's position as its superview is resized.

In code, a combination of springs and struts is set through a view's `autoresizingMask` property. It's a bitmask, so you use logical-or to combine options. The options, with names that start with "UIViewAutoresizingFlexible", represent springs; whatever isn't specified is a strut. The default is `UIViewAutoresizingNone`, meaning all struts.

To demonstrate autoresizing, I'll start with a view and two subviews, one stretched across the top, the other confined to the lower right ([Figure 14-6](#)):

```
UIView* v1 = [[UIView alloc] initWithFrame:CGRectMake(100, 111, 132, 194)];  
v1.backgroundColor = [UIColor colorWithRed:1 green:.4 blue:1 alpha:1];  
UIView* v2 = [[UIView alloc] initWithFrame:CGRectMake(0, 0, 132, 10)];  
v2.backgroundColor = [UIColor colorWithRed:.5 green:1 blue:0 alpha:1];  
UIView* v3 = [[UIView alloc] initWithFrame:v1.bounds.size.width-20,  
                                v1.bounds.size.height-20,  
                                20, 20)];  
v3.backgroundColor = [UIColor colorWithRed:1 green:0 blue:0 alpha:1];  
[self.window.rootViewController.view addSubview: v1];  
[v1 addSubview: v2];  
[v1 addSubview: v3];
```

To that example, I'll add code applying strings and struts to the two subviews to make them behave like the text field and the OK button I was hypothesizing earlier:

```
v2.autoresizingMask = UIViewAutoresizingFlexibleWidth;  
v3.autoresizingMask = UIViewAutoresizingFlexibleTopMargin |  
                    UIViewAutoresizingFlexibleLeftMargin;
```

Now I'll resize the superview, thus bringing autoresizing into play; as you can see ([Figure 14-7](#)), the subviews remain pinned in their correct relative positions:

```
CGRect f = v1.bounds;  
f.size.width += 40;  
f.size.height -= 50;  
v1.bounds = f;
```

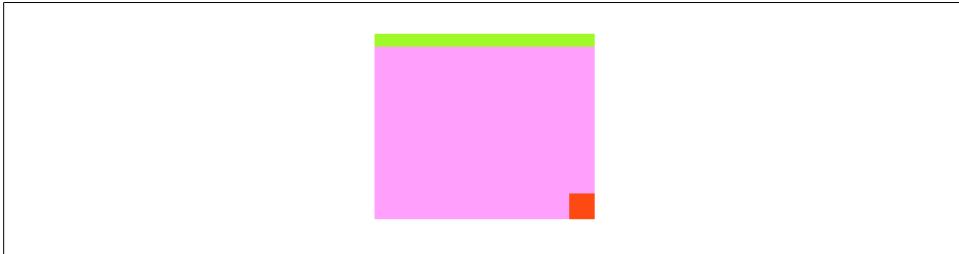


Figure 14-7. After autoresizing

Transform

A view’s `transform` property alters how the view is drawn — it may, for example, change the view’s perceived size and orientation — without affecting its bounds and center. A transformed view continues to behave correctly: a rotated button, for example, is still a button, and can be tapped in its apparent location and orientation.

A transform value is a `CGAffineTransform`, which is a struct representing six of the nine values of a 3×3 transformation matrix (the other three values are constants, so there’s no point representing them in the struct). You may have forgotten your high-school linear algebra, so you may not recall what a transformation matrix is. For the details, which are quite simple really, see the “Transforms” chapter of Apple’s *Quartz 2D Programming Guide*, especially the section called “The Math Behind the Matrices.” But you don’t really need to know those details, because convenience functions, whose names start with `CGAffineTransformMake...`, are provided for creating three of the basic types of transform: rotation, scaling, and translation (i.e., changing the view’s apparent position). A fourth basic transform type, skewing or shearing, has no convenience function.

By default, a view’s transformation matrix is `CGAffineTransformIdentity`, the identity transform. It has no visible effect, so you’re unaware of it. Any transform that you do apply takes place around the view’s center, which is held constant.

Here’s some code to illustrate use of a transform:

```
UIView* v1 = [[UIView alloc] initWithFrame:CGRectMake(113, 111, 132, 194)];  
v1.backgroundColor = [UIColor colorWithRed:1 green:.4 blue:1 alpha:1];  
UIView* v2 = [[UIView alloc] initWithFrame:CGRectInset(v1.bounds, 10, 10)];  
v2.backgroundColor = [UIColor colorWithRed:.5 green:1 blue:0 alpha:1];  
[self.window.rootViewController.view addSubview: v1];  
[v1 addSubview: v2];  
v1.transform = CGAffineTransformMakeRotation(45 * M_PI/180.0);
```

The `transform` property of the view `v1` is set to a rotation transform. The result (Figure 14-8) is that the view appears to be rocked 45 degrees clockwise. (I think in degrees, but Core Graphics thinks in radians, so my code has to convert.) Observe that the view’s `center` property is unaffected, so that the rotation seems to have occurred around the

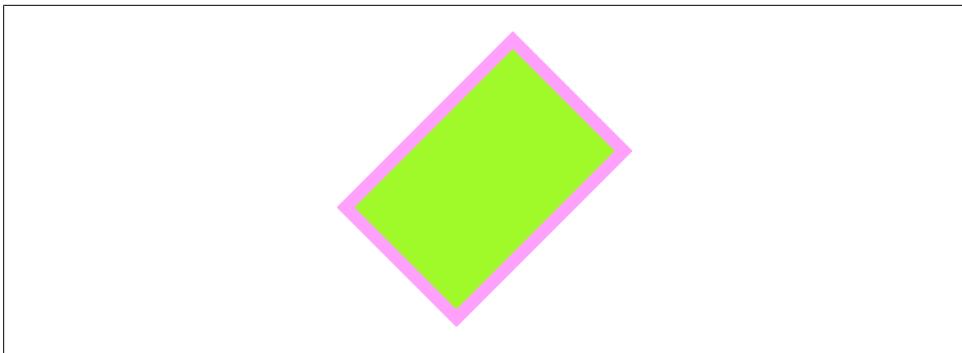


Figure 14-8. A rotation transform

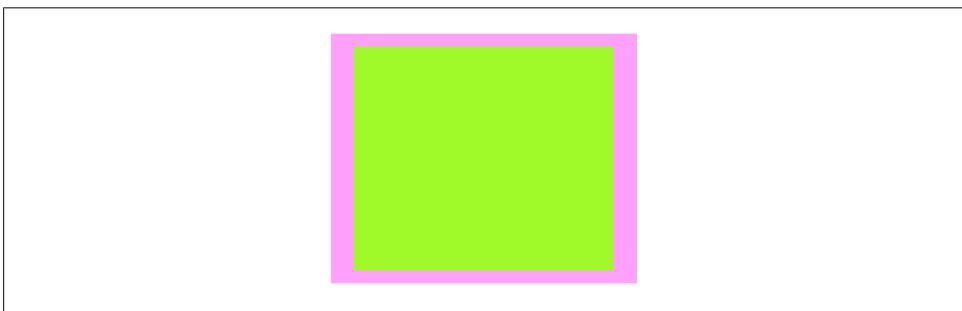


Figure 14-9. A scale transform

view's center. Moreover, the view's `bounds` property is unaffected; the internal coordinate system is unchanged, so the subview is drawn in the same place relative to its superview. The view's `frame`, however, is now useless, as no mere rectangle can describe the region of the superview apparently occupied by the view, its actual value, `\{\{63.7416, 92.7416\}, {230.517, 230.517}\}`, describes the minimal bounding rectangle surrounding view's apparent position. The rule is that if a view's `transform` is not the identity transform, you should not set its `frame`; also, automatic resizing of a subview requires that the superview's transform be the identity transform.

Suppose, instead of `CGAffineTransformMakeRotation`, we call `CGAffineTransformMakeScale`, like this:

```
v1.transform = CGAffineTransformMakeScale(1.8, 1);
```

The `bounds` property of the view `v1` is still unaffected, so the subview is still drawn in the same place relative to its superview; this means that the two views seem to have stretched horizontally together (Figure 14-9). No bounds or centers were harmed by the application of this transform!

Transformation matrices can be chained. There are convenience functions for applying one transform to another. Their names do *not* contain "Make." These functions are

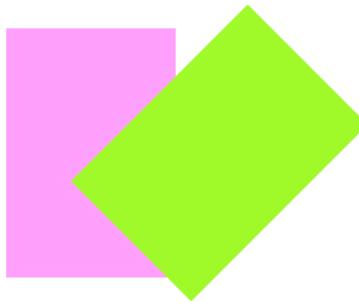


Figure 14-10. Translation, then rotation

not commutative; that is, order matters. If you start with a transform that translates a view to the right and then apply a rotation of 45 degrees, the rotated view appears to the right of its original position; on the other hand, if you start with a transform that rotates a view 45 degrees and then apply a translation to the right, the meaning of “right” has changed, so the rotated view appears 45 degrees down from its original position. To demonstrate the difference, I’ll start with a subview that exactly overlaps its superview:

```
UIView* v1 = [[UIView alloc] initWithFrame:CGRectMake(20, 111, 132, 194)];  
v1.backgroundColor = [UIColor colorWithRed:1 green:.4 blue:1 alpha:1];  
UIView* v2 = [[UIView alloc] initWithFrame:v1.bounds];  
v2.backgroundColor = [UIColor colorWithRed:.5 green:1 blue:0 alpha:1];  
[self.window.rootViewController.view addSubview: v1];  
[v1 addSubview: v2];
```

Then I’ll apply two successive transforms to the subview, leaving the superview to show where the subview was originally. In this example, I translate and then rotate ([Figure 14-10](#)):

```
v2.transform = CGAffineTransformMakeTranslation(100, 0);  
v2.transform = CGAffineTransformRotate(v2.transform, 45 * M_PI/180.0);
```

In this example, I rotate and then translate ([Figure 14-11](#)):

```
v2.transform = CGAffineTransformRotation(45 * M_PI/180.0);  
v2.transform = CGAffineTransformTranslate(v2.transform, 100, 0);
```

The function `CGAffineTransformConcat` concatenates two transform matrices using matrix multiplication. Again, this operation is not commutative. The order is the *opposite* of the order when using convenience functions for applying one transform to another. For example, this gives the same result as [Figure 14-11](#):

```
CGAffineTransform r = CGAffineTransformMakeRotation(45 * M_PI/180.0);  
CGAffineTransform t = CGAffineTransformMakeTranslation(100, 0);  
v2.transform = CGAffineTransformConcat(t,r); // not r,t
```

To remove a transform from a combination of transforms, apply its inverse. A convenience function lets you obtain the inverse of a given affine transform. Again, order

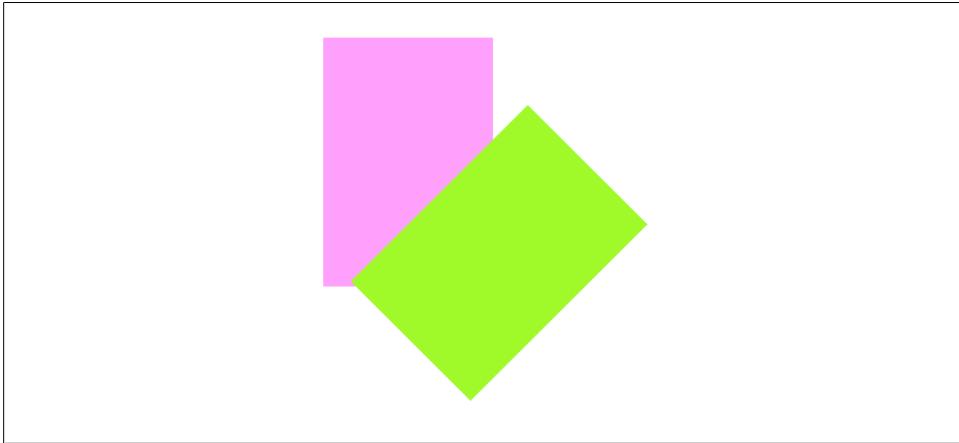


Figure 14-11. Rotation, then translation

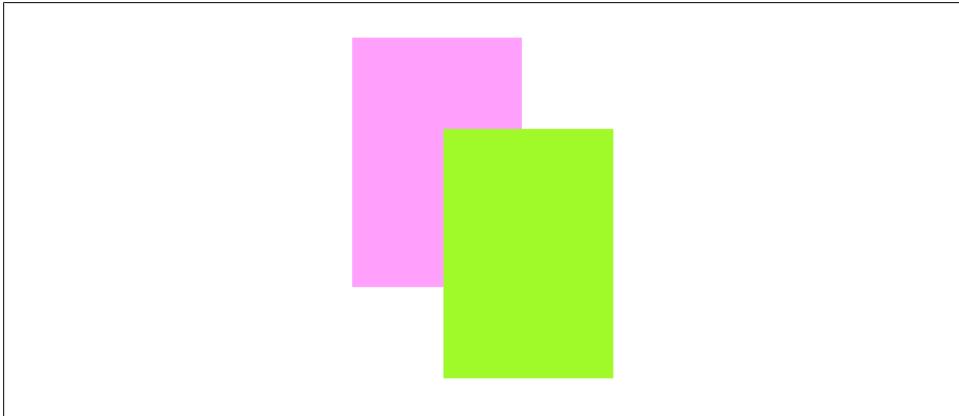


Figure 14-12. Rotation, then translation, then inversion of the rotation

matters. In this example, I rotate the subview and shift it to its “right,” and then remove the rotation ([Figure 14-12](#)):

```
CGAffineTransform r = CGAffineTransformMakeRotation(45 * M_PI/180.0);
CGAffineTransform t = CGAffineTransformMakeTranslation(100, 0);
v2.transform = CGAffineTransformConcat(t,r);
v2.transform = CGAffineTransformConcat(CGAffineTransformInvert(r), v2.transform);
```

Finally, as there are no convenience methods for creating a skew (shear) transform, I’ll illustrate by creating one manually, without further explanation ([Figure 14-13](#)):

```
v1.transform = CGAffineTransformMake(1, 0, -0.2, 1, 0, 0);
```

Transforms are useful particularly as temporary visual indicators. For example, you might call attention to a view by applying a transform that scales it up slightly, and then

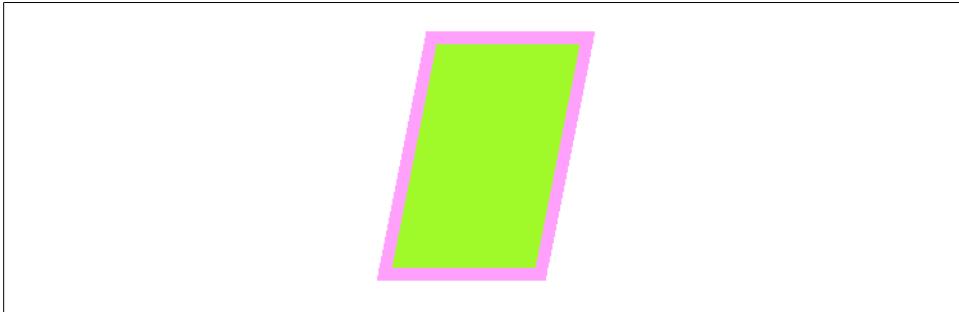


Figure 14-13. Skew (shear)

applying the identity transform to restore it to its original size, and animating those changes ([Chapter 17](#)).

Visibility and Opacity

A view can be made invisible by setting its `hidden` property to YES, and visible again by setting it to NO. This takes it (and its subviews, of course) out of the visible interface without the overhead of actually removing it from the view hierarchy. A hidden view does not (normally) receive touch events, so to the user it really is as if the view weren't there. But it is there, so it can still participate in layout and can be manipulated in other ways.

A view can be assigned a background color through its `backgroundColor` property, as we've been doing in the examples so far in this chapter—indeed, having a background color is the only thing that made our example views visible and distinguishable. A color is a `UIColor`; this is not a difficult class to use, and I'm not going to go into details. A view whose background color is `nil` (the default) has a transparent background. It is perfectly reasonable for a view to have a transparent background and to do no additional drawing of its own, just so that it can act as a convenient superview to other views, making them behave together.

A view can be made partially or completely transparent through its `alpha` property: `1.0` means opaque, `0.0` means transparent, and a value may be anywhere between them, inclusive. This affects subviews, as I've already mentioned; if a superview has an `alpha` of `0.5`, none of its subviews can have an *apparent* opacity of more than `0.5`, because whatever `alpha` value they have will be drawn relative to `0.5`. (Just to make matters more complicated, colors have an `alpha` value as well. So, for example, a view can have an `alpha` of `1.0` but still have a transparent background because its `backgroundColor` has an `alpha` less than `1.0`.) A view that is completely transparent (or very close to it) is like a view whose `hidden` is YES: it is invisible, along with its subviews, and cannot (normally) be touched.

A view's `alpha` property value affects the apparent transparency of its background color and the apparent transparency of its contents separately. For example, if a view displays an image and has a background color and its `alpha` is less than 1, the background color will seep through the image (and whatever is behind the view will seep through both).

A view's `opaque` property, on the other hand, is a horse of a different color; changing it has no effect on the view's appearance. Rather, this property is a hint to the drawing system. If a view completely fills its bounds with ultimately opaque material and its `alpha` is `1.0`, so that the view has no effective transparency, then it can be drawn more efficiently (with less drag on performance) if you inform the drawing system of this fact by setting its `opaque` to YES. Otherwise, you should set its `opaque` to NO. The `opaque` value is *not* changed for you when you set a view's `backgroundColor` or `alpha!` Setting it correctly is entirely up to you; the default, perhaps surprisingly, is YES.

CHAPTER 15

Drawing

Many UIView subclasses, such as a UIButton or a UITextField, know how to draw themselves; sooner or later, though, you’re going to want to do some drawing of your own. A class like UIImageView will display a static image; you can generate that image dynamically by drawing it in code. And a pure UIView does little or no drawing of its own; you can draw its appearance.

Drawing is not difficult, but it is a very large topic. This chapter will make you comfortable with the basic principles, so that you can consult and understand Apple’s documentation when you need further details.

UIImage and UIImageView

The basic general UIKit image class is UIImage. UIImage can read a file from disk, so if an image does not need to be created dynamically, but has already been created before your app runs, then drawing may be as simple as providing an image file as a resource in your app’s bundle. The system knows how to work with many standard image file types, such as TIFF, JPEG, GIF, and PNG. You can also obtain image data in some other way, such as by downloading it, and transform this into a UIImage. Conversely, you can draw your own image for display in your interface or for saving to disk (image file output is discussed in [Chapter 36](#)).

In the very simplest case, an image file in your app’s bundle can be obtained through the UIImage class method `imageNamed:`. This method looks at the top level of your app’s bundle for an image file with the supplied name, including the file extension, and reads it as a UIImage instance. A nice thing about this approach is that memory management is handled for you: the image data may be cached in memory, and if you ask for the same image by calling `imageNamed:` again later, the cached data may be supplied immediately. You can also read an image file from anywhere in your app’s bundle using the class method `imageWithContentsOfFile:` or the instance method `initWithContentsOfFile:`, both of which expect a pathname string; you can get a reference to your app’s

bundle with `[NSBundle mainBundle]`, and `NSBundle` then provides instance methods for getting the pathname of a file within the bundle, such as `pathForResource ofType:`.

Many built-in Cocoa interface objects will accept a `UIImage` as part of how they draw themselves; for example, a `UIButton` can display an image, and (starting in iOS 5) a `UINavigationBar` or a `UITabBar` can have a background image. I'll discuss those in [Chapter 25](#). But when you simply want an image to appear in your interface, you'll probably hand it to a `UIImageView`, which has the most knowledge and flexibility with regard to displaying images and is intended for this purpose. If a `UIImageView` instance begins life in a nib and is to display a `UIImage` from a file in your app's bundle, you won't even need any code; the `UIImageView` can be set to that file directly in the nib. (This mechanism works most easily if the file will be at the top level of the app's bundle.)

A `UIImageView` can actually have *two* images, one assigned to its `image` property and the other assigned to its `highlightedImage` property; the value of the `UIImageView`'s `highlighted` property dictates which of the two is displayed. A `UIImageView` does not automatically highlight itself, the way a button does, for example, merely because the user taps it. However, there are certain situations where a `UIImageView` will respond to the highlighting of its surroundings; for example, within a table view cell, a `UIImageView` will show its highlighted image when the cell is highlighted. You can, of course, also use the notion of `UIImageView` highlighting yourself however you like.

When an image is obtained by name from the bundle, as with `imageNamed` or the name you enter in the nib for a `UIImageView`'s `image`, a file with the same name extended by `~ipad` will automatically be used if the app is running on an iPad. You can use this in a universal app to supply different images automatically depending on whether the app runs on an iPhone or iPod touch, on the one hand, or on an iPad, on the other. This is true not just for images but for *any* resource obtained by name from the bundle. See Apple's *Resource Programming Guide*.

Similarly, on a device with a double-resolution screen (such as the iPhone 4 with Retina display), when an image is obtained by name from the bundle, a file with the same name extended by `@2x`, if there is one, will be used automatically, with the resulting `UIImage` marked as double-resolution by assigning it a `scale` property value of `2.0`. In this way, your app can contain both a single-resolution and a double-resolution version of an image file; on the double-resolution display device, the double-resolution version of the image is used, and is drawn at the same size as the single-resolution image. Thus, on the double-resolution screen, your code continues to work without change, but your images look sharper.

The documentation warns that if a `UIImageView` is to be assigned multiple images (such as an `image` and a `highlightedImage`), they must have the same `scale` property value. This is because the `UIImageView` gets its own internal scaling information from an image's scale at the time it is assigned to it; it does not change its internal scale merely because you switch the value of its `highlighted` property.



Figure 15-1. Mars appears in my interface

A UIImageView is a UIView, so it can have a background color in addition to its image, it can have an alpha (transparency) value, and so forth. A UIImageView without a background color is invisible except for its image, so the image simply appears in the interface, without the user being aware that it resides in a rectangular host. An image may have areas that are transparent, and a UIImageView will respect this; thus an image of any shape can appear. A UIImageView without an image and without a background color is invisible, so you could start with an empty UIImageView in the place where you will later need an image and assign the image in code as needed. You can assign a new image to substitute one image for another.

How a UIImageView draws its image depends upon the setting of its `contentMode` property. (The `contentMode` property is inherited from UIView; I'll discuss its more general purpose later in this chapter.) For example, `UIViewContentModeScaleToFill` means the image's width and height are set to the width and height of the view, thus filling the view completely even if this alters the image's aspect ratio; `UIViewContentModeCenter` means the image is drawn centered in the view without altering its size. The best way to get a feel for the meanings of the various `contentMode` settings is to assign a UIImageView a small image in a nib and then, in the Attributes inspector, change the Mode pop-up menu, and see where and how the image draws itself.

When creating a UIImageView in code, you can take advantage of a convenience initializer, `initWithImage:` (or `initWithImage:highlightedImage:`). The default `contentMode` is `UIViewContentModeScaleToFill`, but the image is not initially scaled; rather, the view itself is sized to match to the image. You will still probably need to set the UIImageView's frame or center in order to place it correctly in its superview. In this example, I'll put a picture of the planet Mars in the center of the window ([Figure 15-1](#)):

```
UIImageView* iv =
    [[UIImageView alloc] initWithImage:[UIImage imageNamed:@"Mars.png"]];
[self.window.rootViewController.view addSubview: iv];
iv.center = self.window.center;
```

If we have a second image file called *Mars@2x.png*, it will be used on a double-resolution device.

Starting in iOS 5, a UIImage can be transformed into a *resizable image*, by sending it the `resizableImageWithCapInsets:` message. (The notion of a resizable image supersedes the notion of a stretchable image from previous system versions.) The `capInsets:` argument is a `UIEdgeInsets`, a struct consisting of four floats representing inset



Figure 15-2. A tiled image of Mars

values starting at the top and moving *counterclockwise* — top, left, bottom, right. They represent distances inwards from the edges of the image. In a context (such as a UIImageView) larger than the image, a resizable image can behave in one of two ways:

- If the cap insets are all zero (UIEdgeInsetsZero), the image will be tiled (repeated).
- If the cap insets are nonzero, the cap inset regions (from the edges inwards) are drawn normally, but the interior of the image is tiled to fill the resulting area. Typically, you'll use cap insets that leave a tiling area of just one or two pixels; otherwise, you can get some very strange results.

First, I'll illustrate tiling:

```
UIImage* mars = [UIImage imageNamed:@"Mars.png"];
UIImage* marsTiled = [mars resizableImageWithCapInsets:UIEdgeInsetsZero];
UIImageView* iv = [[UIImageView alloc] initWithFrame:
    CGRectMake(20, 5, mars.size.width*2, mars.size.height*4)];
iv.image = marsTiled;
[self.window.rootViewController.view addSubview:iv];
```

The image view is eight times the size of the Mars image, and because we used a resizable version of the Mars image, we see eight tiled copies of the Mars image ([Figure 15-2](#)).

Next, I'll illustrate stretching. As I suggested earlier, I'll make almost half the original Mars image serve as a cap inset, leaving just a few pixels in the center to be repeated:

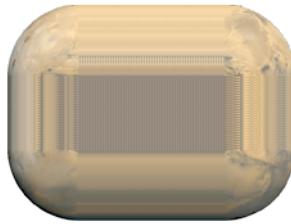


Figure 15-3. A stretched image of Mars

```
UIImage* mars = [UIImage imageNamed:@"Mars.png"];
CGFloat capw = mars.size.width / 2.0 - 1;
CGFloat caph = mars.size.height / 2.0 - 1;
UIImage* marsTiled = [mars resizableImageWithCapInsets:
    UIEdgeInsetsMake(caph, capw, caph, capw)];
UIImageView* iv = [[UIImageView alloc] initWithFrame:
    CGRectMake(20, 5, mars.size.width*2, mars.size.height*1.5)];
iv.image = marsTiled;
[self.window.rootViewController.view addSubview:iv];
```

The result is shown in [Figure 15-3](#). The resulting lozenge may appear somewhat alien at first, but if you look carefully you'll see that its corners are the four quarters of the original Mars image, with the remainder filled in by repeating the central pixels.

Both my examples of stretchable images are deliberately rather weird, but stretchable images are extremely useful. If you start with an image that tiles nicely, such as a fabric texture, you can easily fill a region (perhaps the entire background of your window) with that texture. Similarly, an appropriately stretched image can serve as a button background or a “frame” to set off an image or a region of your interface.

Graphics Contexts

`UIImageView` draws an image for you and takes care of all the details; in many cases, it will be all you'll need. Eventually, though, you may want to do some drawing yourself, directly, in code. To do so, you will always need a *graphics context*.

A graphics context is basically a place you can draw. Conversely, you can't draw in code unless you've got a graphics context. There are several ways in which you might obtain a graphics context; in this chapter I will concentrate on two, which have proven in my experience to be far and away the most common:

You create an image context

The function `UIGraphicsBeginImageContextWithOptions` creates a graphics context suitable for use as an image. You then draw into this context to generate the image. When you've done that, you call `UIGraphicsGetImageFromCurrentImageContext` to turn the context into a `UIImage`, and then `UIGraphicsEndImageContext` to dismiss

the context. Now you have a `UIImage` that you can display in your interface or draw into some other graphics context or save as a file.

Cocoa hands you a graphics context

You subclass `UIView` and implement `drawRect:`. At the time your `drawRect:` implementation is called, Cocoa has already created a graphics context and is asking you to draw into it, right now; whatever you draw is what the `UIView` will display. (A slight variant of this situation is that you subclass a `CALayer` and implement `drawInContext:`, or make some object the delegate of a layer and implement `drawLayer:inContext:`; layers are discussed in [Chapter 16](#).)

Moreover, at any given moment there either is or is not a *current graphics context*:

- `UIGraphicsBeginImageContextWithOptions` not only creates an image context, it also makes that context the current graphics context.
- When `drawRect` is called, the `UIView`'s drawing context is already the current graphics context.
- Callbacks with a `context:` argument have *not* made any context the current graphics context; rather, that argument is a reference to a graphics context.

What beginners find most confusing about drawing is there are two separate sets of tools with which you can draw, and they take different attitudes towards the context in which they will draw:

UIKit

Various Objective-C classes know how to draw themselves; these include `UIImage`, `NSString` (for drawing text), `UIBezierPath` (for drawing shapes), and `UIColor`. Some of these classes provide convenience methods with limited abilities; others are extremely powerful. In many cases, `UIKit` will be all you'll need.

With `UIKit`, you can draw *only into the current context*. So if you're in a `UIGraphicsBeginImageContextWithOptions` or `drawRect` situation, you can use the `UIKit` convenience methods directly; there is a current context and it's the one you want to draw into. If you've been handed a `context:` argument, on the other hand, then if you want to use the `UIKit` convenience methods, you'll have to make that context the current context; you do this by calling `UIGraphicsPushContext` (and be sure to restore things with `UIGraphicsPopContext` later).

Core Graphics

This is the full drawing API. Core Graphics, often referred to as Quartz, or Quartz 2D, is the drawing system that underlies all iOS drawing — `UIKit` drawing is built on top of it — so it is low-level and consists of C functions. There are a lot of them! This chapter will familiarize you with the fundamentals; for complete information, you'll want to study Apple's *Quartz 2D Programming Guide*.

With Core Graphics, you must *specify a graphics context* to draw into, explicitly, in every function call. If you've been handed a `context:` argument, then, hey presto,

you have a graphics context, and it's probably the graphics context you want to draw into. But in a `UIGraphicsBeginImageContextWithOptions` or `drawRect` situation, you have no reference to a context; to use Core Graphics, you need to get such a reference. Since the context you want to draw into is the current graphics context, you call `UIGraphicsGetCurrentContext` to get the needed reference.

So we have two sets of tools and three ways in which a context might be supplied; that makes six ways of drawing, and in case you're confused, I'll now demonstrate all six of them. Without worrying just yet about the actual drawing commands, focus your attention on how the context is specified and on whether we're using UIKit or Core Graphics. First, I'll draw a blue circle by implementing a `UIView` subclass's `drawRect:`, using UIKit to draw into the current context, which Cocoa has already prepared for me:

```
- (void) drawRect: (CGRect) rect {  
    UIBezierPath* p =  
        [UIBezierPath bezierPathWithOvalInRect:CGRectMake(0,0,100,100)];  
    [[UIColor blueColor] setFill];  
    [p fill];  
}
```

Now I'll do the same thing with Core Graphics; this will require that I first get a reference to the current context:

```
- (void) drawRect: (CGRect) rect {  
    CGContextRef con = UIGraphicsGetCurrentContext();  
    CGContextAddEllipseInRect(con, CGRectMake(0,0,100,100));  
    CGContextSetFillColorWithColor(con, [UIColor blueColor].CGColor);  
    CGContextFillPath(con);  
}
```

Next, I'll implement a `UIView` subclass's `drawLayer:inContext:`. In this case, we're handed a reference to a context, but it isn't the current context. So I have to make it the current context in order to use UIKit:

```
- (void)drawLayer:(CALayer*)lay inContext:(CGContextRef)con {  
    UIGraphicsPushContext(con);  
    UIBezierPath* p =  
        [UIBezierPath bezierPathWithOvalInRect:CGRectMake(0,0,100,100)];  
    [[UIColor blueColor] setFill];  
    [p fill];  
    UIGraphicsPopContext();  
}
```

To use Core Graphics in `drawLayer:inContext:`, I simply keep referring to the context I was handed:

```
- (void)drawLayer:(CALayer*)lay inContext:(CGContextRef)con {  
    CGContextAddEllipseInRect(con, CGRectMake(0,0,100,100));  
    CGContextSetFillColorWithColor(con, [UIColor blueColor].CGColor);  
    CGContextFillPath(con);  
}
```

Finally, for the sake of completeness, let's make a `UIImage` of a blue circle. We can do this at any time (we don't need to wait for some particular method to be called) and in any class (we don't need to be in a `UIView` subclass). First, I'll use UIKit:

```
UIGraphicsBeginImageContextWithOptions(CGSizeMake(100,100), NO, 0);
UIBezierPath* p = [UIBezierPath bezierPathWithOvalInRect:CGRectMake(0,0,100,100)];
[[UIColor blueColor] setFill];
[p fill];
UIImage* im = UIGraphicsGetImageFromCurrentImageContext();
UIGraphicsEndImageContext();
// im is the blue circle image, do something with it here ...
```

Here's the same thing using Core Graphics:

```
UIGraphicsBeginImageContextWithOptions(CGSizeMake(100,100), NO, 0);
CGContextRef con = UIGraphicsGetCurrentContext();
CGContextAddEllipseInRect(con, CGRectMake(0,0,100,100));
CGContextSetFillColorWithColor(con, [UIColor blueColor].CGColor);
CGContextFillPath(con);
UIImage* im = UIGraphicsGetImageFromCurrentImageContext();
UIGraphicsEndImageContext();
// im is the blue circle image, do something with it here ...
```

You may be wondering about the arguments to `UIGraphicsBeginImageContextWithOptions`. The first argument is obviously the size of the image to be created. The second argument declares whether the image should be opaque; if I had passed YES instead of NO here, my image would have a black background, which I don't want. The third argument specifies the image scale, corresponding to the `UIImage scale` property I discussed earlier; by passing 0, I'm telling the system to set the scale for me in accordance with the main screen resolution, so my image will look good on both single-resolution and double-resolution devices.

You don't have to use UIKit or Core Graphics exclusively; on the contrary, you can intermingle UIKit calls and Core Graphics calls to operate on the same graphics context. They simply have two different *ways* of talking about the same graphics context.



In the past, there was a caveat that accessing the current context in `drawRect` was not thread-safe, so there were certain situations where use of UIKit (and `UIGraphicsGetCurrentContext`) was forbidden; however, Apple's *Technical Q&A QA1637*, which originally issued this caveat, now says: "Beginning with iOS 4.0, drawing to a graphics context in UIKit is thread-safe. This includes accessing and manipulating the current graphics stack, drawing images and strings, and usage of color and font objects from secondary threads." This edition of the book doesn't concern itself with backwards compatibility beyond iOS 4.0, so I'll blithely assume that UIKit and Core Graphics are equally available.



Figure 15-4. Two images of Mars combined side by side

UIImage Drawing

UIImage provides methods for drawing itself into the current context. We now know how to obtain an image context and make it the current context, so we can experiment with these methods. Here, I'll make a UIImage consisting of two pictures of Mars side by side:

```
UIImage* mars = [UIImage imageNamed:@"Mars.png"];
CGSize sz = [mars size];
UIGraphicsBeginImageContextWithOptions(CGSizeMake(sz.width*2, sz.height), NO, 0);
[mars drawAtPoint:CGPointMake(0,0)];
[mars drawAtPoint:CGPointMake(sz.width,0)];
UIImage* im = UIGraphicsGetImageFromCurrentImageContext();
UIGraphicsEndImageContext();
```

If I now hand this image `im` over to a visible UIImageView, the image appears onscreen ([Figure 15-4](#)). I could do this, for example, by creating the UIImageView in code, as before:

```
UIImageView* iv = [[UIImageView alloc] initWithImage:im];
[self.window.rootViewController.view addSubview: iv];
iv.center = self.window.center;
```

Additional UIImage methods let you scale an image into a desired rectangle as you draw, and specify the compositing (blend) mode whereby the image should combine with whatever is already present. To illustrate, I'll create an image showing Mars centered in another image of Mars that's twice as large, using the Multiply blend mode ([Figure 15-5](#)):

```
UIImage* mars = [UIImage imageNamed:@"Mars.png"];
CGSize sz = [mars size];
UIGraphicsBeginImageContextWithOptions(CGSizeMake(sz.width*2, sz.height*2), NO, 0);
[mars drawInRect:CGRectMake(0,0,sz.width*2,sz.height*2)];
[mars drawInRect:CGRectMake(sz.width/2.0, sz.height/2.0, sz.width, sz.height)
blendMode:kCGBlendModeMultiply alpha:1.0];
UIImage* im = UIGraphicsGetImageFromCurrentImageContext();
UIGraphicsEndImageContext();
```

There is no UIImage drawing method for specifying the source rectangle — that is, for specifying that you want to extract a smaller region of the original image. You can work around this by specifying a smaller graphics context and positioning the image drawing so that the desired region falls into it. For example, to obtain an image of the right half



Figure 15-5. Two images of Mars in different sizes, composited



Figure 15-6. Half the original image of Mars

of Mars, you'd make a graphics context half the width of the `mars` image, and then draw `mars` shifted left, so that only its right half intersects the graphics context. There is no harm in doing this, and it's a perfectly standard device; the left half of `mars` simply isn't drawn (Figure 15-6):

```
UIImage* mars = [UIImage imageNamed:@"Mars.png"];
CGSize sz = [mars size];
UIGraphicsBeginImageContextWithOptions(CGSizeMake(sz.width/2.0, sz.height), NO, 0);
[mars drawAtPoint:CGPointMake(-sz.width/2.0, 0)];
UIImage* im = UIGraphicsGetImageFromCurrentImageContext();
UIGraphicsEndImageContext();
```

CGImage Drawing

The Core Graphics version of `UIImage` is `CGImage` (actually a `CGImageRef`). They are easily converted to one another: a `UIImage` has a `CGImage` property that accesses its Quartz image data, and you can make a `UIImage` from a `CGImage` using `imageWithCGImage:` or `initWithCGImage:`.

A `CGImage` lets you create a new image directly from a rectangular region of the original image. (It also lets you apply an image mask, which you can't do with `UIImage`.) I'll demonstrate by splitting the image of Mars in half and drawing the two halves sepa-

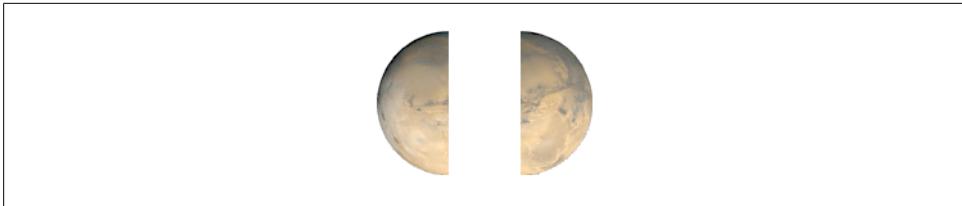


Figure 15-7. Image of Mars split in half

rately ([Figure 15-7](#)). Observe we are now operating in the CFTyperef world and must take care to manage memory manually; ARC won't help us here ([Chapter 12](#)):

```
UIImage* mars = [UIImage imageNamed:@"Mars.png"];
// extract each half as a CGImage
CGSize sz = [mars size];
CGImageRef marsLeft = CGImageCreateWithImageInRect([mars CGImage],
    CGRectMake(0,0,sz.width/2.0,sz.height));
CGImageRef marsRight = CGImageCreateWithImageInRect([mars CGImage],
    CGRectMake(sz.width/2.0,0,sz.width/2.0,sz.height));
// draw each CGImage into an image context
UIGraphicsBeginImageContextWithOptions(CGSizeMake(sz.width*1.5, sz.height), NO, 0);
CGContextRef con = UIGraphicsGetCurrentContext();
CGContextDrawImage(con, CGRectMake(0,0,sz.width/2.0,sz.height), marsLeft);
CGContextDrawImage(con, CGRectMake(sz.width/2.0,0,sz.width/2.0,sz.height), marsRight);
UIImage* im = UIGraphicsGetImageFromCurrentImageContext();
UIGraphicsEndImageContext();
CGImageRelease(marsLeft); CGImageRelease(marsRight);
```

But there's a problem with that example: the drawing is upside-down! It isn't rotated; it's mirrored top to bottom, or, to use the technical term, *flipped*. This phenomenon can arise when you create a CGImage and then draw it with CGContextDrawImage and is due to a mismatch in the native coordinate systems of the source and target contexts.

There are various ways of compensating for this mismatch between the coordinate systems. One is to draw the CGImage into an intermediate UIImage and extract *another* CGImage from that. [Example 15-1](#) presents a utility function for doing this.

Example 15-1. Utility for flipping an image drawing

```
CGImageRef flip (CGImageRef im) {
    CGSize sz = CGSizeMake(CGImageGetWidth(im), CGImageGetHeight(im));
    UIGraphicsBeginImageContextWithOptions(sz, NO, 0);
    CGContextDrawImage(UIGraphicsGetCurrentContext(),
        CGRectMake(0, 0, sz.width, sz.height), im);
    CGImageRef result = [UIGraphicsGetImageFromCurrentImageContext() CGImage];
    UIGraphicsEndImageContext();
    return result;
}
```

Armed with the utility function from [Example 15-1](#), we can now draw the halves of Mars the right way up in the previous example:

```

CGContextDrawImage(con, CGRectMake(0,0,sz.width/2.0,sz.height),
    flip(marsLeft));
CGContextDrawImage(con, CGRectMake(sz.width,0,sz.width/2.0,sz.height),
    flip(marsRight));

```

However, we've *still* got a problem: on a double-resolution device, if there is a high-resolution (@2x) version of our image file, the drawing comes out all wrong. The reason is that we are loading our original Mars image using `imageNamed:`, which automatically substitutes the high-resolution version of the image on the high-resolution device. The `UIImage` compensates for the doubled size of the image by setting its own `scale` property to match. But a `CGImage` doesn't have a `scale` property, and knows nothing of the fact that the image dimensions are doubled!

When you call a `UIImage`'s `CGImage` method, therefore, you can't assume that the resulting `CGImage` is the same size as the original `UIImage`; a `UIImage`'s `size` property is the same for a single-resolution image and its double-resolution counterpart, because the `scale` tells it how to compensate, but the `CGImage` of a double-resolution `UIImage` is twice as large in both dimensions as the `CGImage` of the corresponding single-resolution image.

So, in extracting a desired piece of the `CGImage`, we must either multiply all appropriate values by the `scale` or express ourselves in terms of the `CGImage`'s dimensions. In this case, as we are extracting the left and right halves of the image, the latter is obviously the simpler course. So here's a version of our original code that draws correctly on either a single-resolution or a double-resolution device, and compensates for flipping:

```

UIImage* mars = [UIImage imageNamed:@"Mars.png"];
CGSize sz = [mars size];
// Derive CGImage and use its dimensions to extract its halves
CGImageRef marsCG = [mars CGImage];
CGSize szCG = CGSizeMake(CGImageGetWidth(marsCG), CGImageGetHeight(marsCG));
CGImageRef marsLeft = CGImageCreateWithImageInRect(marsCG,
    CGRectMake(0,0,szCG.width/2.0,0,szCG.height));
CGImageRef marsRight = CGImageCreateWithImageInRect(marsCG,
    CGRectMake(szCG.width/2.0,0,0,szCG.width/2.0,szCG.height));
UIGraphicsBeginImageContextWithOptions(CGSizeMake(sz.width*1.5, sz.height), NO, 0);
// The rest is as before, calling flip() to compensate for flipping
CGContextRef con = UIGraphicsGetCurrentContext();
CGContextDrawImage(con, CGRectMake(0,0,sz.width/2.0,sz.height),
    flip(marsLeft));
CGContextDrawImage(con, CGRectMake(sz.width,0,sz.width/2.0,sz.height),
    flip(marsRight));
UIImage* im = UIGraphicsGetImageFromCurrentImageContext();
UIGraphicsEndImageContext();
CGImageRelease(marsLeft); CGImageRelease(marsRight);

```

If this is starting to look rather clumsy and involved, don't worry; I have up my sleeve another flipping solution that simplifies things considerably. Instead of calling our `flip` utility, you can wrap your `CGImage` in a `UIImage` before drawing. This has two big advantages:

- The `UIImage` compensates for flipping automatically as it draws.

Why Flipping Happens

The ultimate source of accidental flipping is that Core Graphics comes from the Mac OS X world, where the coordinate system's origin is located by default at the bottom left and the positive y-direction is upward, whereas on iOS the origin is located by default at the top left and the positive y-direction is downward. In most drawing situations, no problem arises, because the coordinate system of the graphics context is adjusted to compensate. Thus, the default coordinate system for drawing in a Core Graphics context on iOS has the origin at the top left, just as you expect. But creating and drawing a CGImage exposes the issue.

- The UIImage can be formed in such a way as to compensate for scale: call `imageWithCGImage:scale:orientation:` as you form the UIImage from the CGImage.

So here's a self-contained approach that deals with both flipping and scale:

```
UIImage* mars = [UIImage imageNamed:@"Mars.png"];
CGSize sz = [mars size];
// Derive CGImage and use its dimensions to extract its halves
CGImageRef marsCG = [mars CGImage];
CGSize szCG = CGSizeMake(CGImageGetWidth(marsCG), CGImageGetHeight(marsCG));
CGImageRef marsLeft = CGImageCreateWithImageInRect(marsCG,
    CGRectMake(0,0,szCG.width/2.0,szCG.height));
CGImageRef marsRight = CGImageCreateWithImageInRect(marsCG,
    CGRectMake(szCG.width/2.0,0,szCG.width/2.0,szCG.height));
UIGraphicsBeginImageContextWithOptions(CGSizeMake(sz.width*1.5, sz.height), NO, 0);
[[UIImage imageWithCGImage:marsLeft
    scale:[mars scale]
    orientation:UIImageOrientationUp]
 drawAtPoint:CGPointMake(0,0)];
[[UIImage imageWithCGImage:marsRight
    scale:[mars scale]
    orientation:UIImageOrientationUp]
 drawAtPoint:CGPointMake(sz.width,0)];
UIImage* im = UIGraphicsGetImageFromCurrentImageContext();
UIGraphicsEndImageContext();
CGImageRelease(marsLeft); CGImageRelease(marsRight);
```

Yet another solution to flipping is to apply a transform to the graphics context before drawing the CGImage, effectively flipping the context's internal coordinate system. This is elegant, but can be confusing if there are other transforms in play. I'll talk more about graphics context transforms later in this chapter.

CIFilter and CIIImage

CIFilter and CIIImage are new in iOS 5, though they have been present in Mac OS X for many years. The “CI” stands for Core Image, a technology for transforming images through mathematical filters. But don't start imagining blurry buttons or an iPad ver-

sion of Photoshop; iOS devices are not powerful enough for Apple to have provided any mathematically intensive filters. To use Core Image, you'll have to link your target to *CoreImage.framework*.

A filter is a CIFilter. The available filters fall naturally into several categories:

Patterns and gradients

These filters create CIImages that can then be combined with other CIImages, such as a single color, a checkerboard, stripes, or a gradient.

Compositing

These filters combine one image with another, using compositing blend modes familiar from image processing programs such as Photoshop.

Color

These filters adjust or otherwise modify the colors of an image. Thus you can alter an image's saturation, hue, brightness, contrast, gamma and white point, exposure, shadows and highlights, and so on.

Geometric

These filters perform basic geometric transformations on an image, such as scaling, rotation, and cropping.

Use of a CIFilter is quite simple; it basically works as if a filter were a kind of dictionary consisting of keys and values. You create the filter by supplying the string name of a filter; to learn what these names are, consult Apple's *Core Image Filter Reference*, or call the CIFilter class method `filterNamesInCategories`: with a nil argument. Each filter has a small number of keys and values that determine its behavior; for each key that you're interested in, you supply a key–value pair, either by calling `setValue:forKey:` or by supplying all the keys and values as you specify the filter name. In supplying values, a number will be an NSNumber, generated with `numberWithFloat:`, and there are a few supporting classes such as CIVector (like CGPoint and CGRect combined) and CI-Color, whose use is easy to grasp.

These keys include any image or images on which the filter is to operate; such an image must be a CIImage. You can obtain a CIImage from a CGImage with `initWithCGImage:`; we already know how to obtain a CGImage from a UIImage. You can also obtain a CGImage as the output of a filter; thus filters can be chained together.

The interesting thing is that as you build a chain of filters, nothing actually happens. The only calculation-intensive move comes at the very end, when you produce the result of the entire chain as a CGImage. You do this by creating a CIContext (by calling `contextWithOptions:`) and calling `createCGImage:fromRect:`. The only mildly tricky thing here is that a CIImage doesn't have a frame or bounds; it has an extent. You will often use this as the second argument to `createCGImage:fromRect:`. The final output CGImage is ready for any purpose, such as for display in your app, for transformation into a UIImage, or for use in further drawing.



Figure 15-8. A photo of me, filtered

To illustrate, I'll start with an ordinary photo of myself (it's true I'm wearing a motorcycle helmet, but it's still ordinary) and black out everything except my head by compositing a circular white-to-black gradient onto it using Darken blend mode:

```
UIImage* moi = [UIImage imageNamed:@"moi.jpg"];
CIIImage* moi2 = [[CIIImage alloc] initWithCGImage:moi.CGImage];

CIFilter* grad = [CIFilter filterWithName:@"CIRadialGradient"];
CIVector* center = [CIVector vectorWithX:moi.size.width/2.0 Y:moi.size.height/2.0];
[grad setValue:center forKey:@"inputCenter"];
CIFilter* dark = [CIFilter filterWithName:@"CIDarkenBlendMode"
    keysAndValues:
        @{@"inputImage", grad.outputImage,
        @"inputBackgroundImage", moi2,
        nil};

CIContext* con = [CIContext contextWithOptions:nil];
CGImageRef moi3 = [con createCGImage:dark.outputImage
    fromRect:moi2.extent];
UIImage* moi4 = [UIImage imageWithCGImage:moi3];
CGImageRelease(moi3);
```

The UIImage moi4 can now be displayed in our app's interface ([Figure 15-8](#)).

The example may not seem overly compelling, since everything I've just done (except, perhaps, Core Image's use of the device's GPU) can effectively be done using Core Graphics alone; Core Graphics can make a circular gradient and can composite images using blend modes. However, Core Image is much simpler, especially when you want to reuse a chain of filters on multiple input images, and its color tweaking features are beyond anything Core Graphics supplies.

It is also possible to draw a filter's output directly into an OpenGL context, but OpenGL is outside the scope of this book. Core Image can also perform automatic face detection in an image.

Drawing a UIView

The most flexible way to draw a UIView is to draw it yourself. As I've already said, you don't actually draw a UIView; you subclass UIView and endow the subclass with the ability to draw itself. When a UIView needs drawing, its `drawRect:` method is called. Overriding that method is your chance to draw. At the time that `drawRect:` is called, the current graphics context has already been set to the view's own graphics context. You can use Core Graphics functions or UIKit convenience methods to draw into that context. Thus, everything I did in the previous two sections, generating a UIImage and displaying it somehow in the interface, could have been done instead by putting into my interface a UIView subclass that knows how to display itself as desired.



You should *never* call `drawRect:` yourself! If a view needs updating and you want its `drawRect:` called, send the view the `setNeedsDisplay` message. This will cause `drawRect:` to be called at the next proper moment. Also, don't override `drawRect:` unless you are assured that this is legal. For example, it is not legal to override `drawRect:` in a subclass of UIImageView; you cannot combine your drawing with that of the UIImageView.

So, for example, let's say we have a UIView subclass called MyView. How this class gets instantiated, and how the instance gets into our view hierarchy, isn't important. One possibility would be to drag a UIView into a view in the nib and set its class to MyView in the identity inspector. Another would be to create the MyView instance and put it into the interface in code:

```
MyView* mv = [[MyView alloc] initWithFrame:  
    CGRectMake(0, 0, self.window.bounds.size.width - 50, 150)];  
mv.center = self.window.center;  
mv.opaque = NO;  
[self.window.rootViewController.view addSubview: mv];
```

Let's suppose that MyView's job is to draw the two halves of Mars, one at each end of the view. Then we can readily adapt the last example from the previous section. There is no need for an image context; we just draw directly into the current context, which is the view's own graphics context:

```
- (void)drawRect:(CGRect)rect {  
    CGRect b = self.bounds;  
    UIImage* mars = [UIImage imageNamed:@"Mars.png"];  
    CGSize sz = [mars size];  
    CGImageRef marsCG = [mars CGImage];  
    CGSize szCG = CGSizeMake(CGImageGetWidth(marsCG), CGImageGetHeight(marsCG));  
    CGImageRef marsLeft = CGImageCreateWithImageInRect(marsCG,  
        CGRectMake(0,0,szCG.width/2.0,szCG.height));  
    CGImageRef marsRight = CGImageCreateWithImageInRect(marsCG,  
        CGRectMake(szCG.width/2.0,0,szCG.width/2.0,szCG.height));  
    [[UIImage imageWithCGImage:marsLeft  
        scale:[mars scale]
```

```

        orientation:UIImageOrientationUp]
    drawAtPoint:CGPointMake(0,0)];
    [[UIImage imageNamed:marsRight
        scale:[mars scale]
        orientation:UIImageOrientationUp]
    drawAtPoint:CGPointMake(b.size.width-sz.width/2.0,0)];
    CGImageRelease(marsLeft); CGImageRelease(marsRight);
}

```

There is no need to call `super`, because the superclass here is `UIView`, whose `drawRect:` does nothing.

The need to draw in real time, on demand, surprises some beginners, who worry that drawing may be a time-consuming operation. Where drawing is extensive and can be compartmentalized into sections, you may be able to gain some efficiency by paying attention to the `rect` parameter passed into `drawRect:`. It designates the region of the view's bounds that needs refreshing. The system knows this either because this is the area that has just been exposed by the removal of some covering view or because you called `setNeedsDisplayInRect:`, specifying it. Thus, you could call `setNeedsDisplayInRect:` to tell your `drawRect:` to redraw a subregion of the view; the rest of the view will be left alone.

In general, however, you should not optimize prematurely. The code for a drawing operation may appear verbose and yet be extremely fast. Moreover, the iOS drawing system is efficient; it doesn't call `drawRect:` unless it has to (or is told to, through a call to `setNeedsDisplay`), and once a view has drawn itself, the result is cached so that the cached drawing can be reused instead of repeating the drawing operation from scratch. You can readily satisfy yourself of this fact with some caveman debugging, logging in your `drawRect:` implementation; you may be amazed to discover that your code is called only once in the entire lifetime of the app! In fact, moving code to `drawRect:` is a common way to *increase* efficiency. This is because it is more efficient for the drawing engine to render directly onto the screen than for it render once offscreen and then copy those pixels onto the screen.

You may have noticed, oh keen-eyed reader, that when I created the `MyView` instance in code I set its `opaque` property to NO. Without this, the view would have been drawn with a black background. Of course, if a view fills its rectangle with opaque drawing or has an opaque background color, you can leave `opaque` set to YES and gain some drawing efficiency (see [Chapter 14](#)).

Obtaining a transparent `UIView` background is a source of considerable confusion among beginners. The black background arises when two things are true:

- The view's `backgroundColor` is nil.
- The view's `opaque` is YES.

Unfortunately, when creating a `UIView` in code, both those things *are* true by default! So if you don't want the black background, you must do something about one or the other of them (or both). For example, I could have eliminated the black background

by setting the view's `backgroundColor` to `[UIColor clearColor]`. But then I should *still* set its `opaque` to `NO`, because the view isn't opaque, and it's up to me to tell the drawing system this.

With a `UIView` created in the nib, on the other hand, the black background problem doesn't arise. This is because such a `UIView`'s `backgroundColor` is not `nil`. The nib assigns it *some* actual background color, even if that color is `[UIColor clearColor]`.

Graphics Context Settings

When you draw in a graphics context, the drawing obeys the context's current settings. Thus, the procedure is always to configure the context's settings first, and then draw. For example, to draw a red line followed by a blue line, you would first set the context's line color to red, and then draw the first line; then you'd set the context's line color to blue, and then draw the second line. To the eye, it appears that the redness and blueness are properties of the individual lines, but in fact, at the time you draw each line, line color is a feature of the entire graphics context. This is true regardless of whether you use `UIKit` methods or `Core Graphics` functions.

A graphics context thus has, at every moment, a *state*, which is the sum total of all its settings; the way a piece of drawing looks is the result of what the graphics context's state was at the moment that piece of drawing was performed. To help you manipulate entire states, the graphics context provides a *stack* for holding states. Every time you call `CGContextSaveGState`, the context pushes the entire current state onto the stack; every time you call `CGContextRestoreGState`, the context retrieves the state from the top of the stack (the state that was most recently pushed) and sets itself to that state.

Thus, a common pattern is: call `CGContextSaveGState`; manipulate the context's settings, thus changing its state; draw; call `CGContextRestoreGState` to restore the state and the settings to what they were before you manipulated them. You do not have to do this before *every* manipulation of a context's settings, however, because settings don't necessarily conflict with one another or with past settings. You can set the context's line color to red and then later to blue without any difficulty. But in certain situations you do want your manipulation of settings to be undoable, and I'll point out several such situations later in this chapter.

Many of the settings that constitute a graphics context's state, and that determine the behavior and appearance of drawing performed at that moment, are similar to those of any drawing application. Here are some of them, along with some of the commands that determine them; I provide `Core Graphics` functions here, but keep in mind that `UIKit` commands are actually calling these same functions and manipulating the context's state in the same ways:

Line thickness and dash style

`CGContextSetLineWidth`, `CGContextSetLineDash`

Line end-cap style and join style

`CGContextSetLineCap, CGContextSetLineJoin, CGContextSetMiterLimit`

Line color or pattern

`CGContextSetRGBStrokeColor, CGContextSetGrayStrokeColor, CGContextSetStroke-
ColorWithColor, CGContextSetStrokePattern`

Fill color or pattern

`CGContextSetRGBFillColor, CGContextSetGrayFillColor, CGContextSetFillColor-
WithColor, CGContextSetFillPattern`

Shadow

`CGContextSetShadow, CGContextSetShadowWithColor`

Blend mode

`CGContextSetBlendMode` (this determines how drawing that you do now will be
composited with drawing already present)

Overall transparency

`CGContextSetAlpha` (individual colors also have an alpha component)

Text features

`CGContextSelectFont, CGContextSetFont, CGContextSetFontSize, CGContextSetText-
DrawingMode, CGContextSetCharacterSpacing`

Whether anti-aliasing and font smoothing are in effect

`CGContextSetShouldAntialias, CGContextSetShouldSmoothFonts`

Additional settings include:

Clipping area

Drawing outside the clipping area is not physically drawn.

Transform (or “CTM,” for “current transform matrix”)

Changes how points that you specify in subsequent drawing commands are mapped onto the physical space of the canvas.

Many (but not all) of these settings will be illustrated by examples later in this chapter.

Paths and Drawing

By issuing a series of instructions for moving an imaginary pen, you trace out a *path*. Such a path does *not* constitute drawing! First you provide a path; *then* you draw. Drawing can mean stroking the path or filling the path, or both. Again, this should be a familiar notion from certain drawing applications.

A path is constructed by tracing it out from point to point. Think of the drawing system as holding a pen. Then you must first tell that pen where to position itself, setting the current point; after that, you issue a series of commands telling it how to trace out each

subsequent piece of the path. Each additional piece of the path starts at the current point; its end becomes the new current point.

Here are some path-drawing commands you’re likely to give:

Position the current point

`CGContextMoveToPoint`

Trace a line

`CGContextAddLineToPoint`, `CGContextAddLines`

Trace a rectangle

`CGContextAddRect`, `CGContextAddRects`

Trace an ellipse or circle

`CGContextAddEllipseInRect`

Trace an arc

`CGContextAddArcToPoint`, `CGContextAddArc`

Trace a Bezier curve with one or two control points

`CGContextAddQuadCurveToPoint`, `CGContextAddCurveToPoint`

Close the current path

`CGContextClosePath`. This appends a line from the last point of the path to the first point. There’s no need to do this if you’re about to fill the path, since it’s done for you.

Stroke or fill the current path

`CGContextStrokePath`, `CGContextFillPath`, `CGContextEOFillPath`, `CGContextDrawPath`. Stroking or filling the current path *clears the path*. Use `CGContextDrawPath` if you want both to fill and to stroke the path in a single command, because if you merely stroke it first with `CGContextStrokePath`, the path is cleared and you can no longer fill it.

There are also a lot of convenience functions that create a path and stroke or fill it all in a single move: `CGContextStrokeLineSegments`, `CGContextStrokeRect`, `CGContextStrokeRectWithWidth`, `CGContextFillRect`, `CGContextFillRects`, `CGContextStrokeEllipseInRect`, `CGContextFillEllipseInRect`.

A path can be compound, meaning that it consists of multiple independent pieces. For example, a single path might consist of two separate closed shapes: a rectangle and a circle. When you call `CGContextMoveToPoint` in the middle of constructing a path (that is, after tracing out a path and without clearing it by filling, stroking, or calling `CGContextBeginPath`), you pick up the imaginary pen and move it to a new location without tracing a segment, thus preparing to start an independent piece of the same path. If you’re worried, as you begin to trace out a path, that there might be an existing path and that your new path might be seen as a compound part of that existing path, you can call `CGContextBeginPath` to specify that this is a different path; many of Apple’s examples do this, but in practice I usually do not find it necessary.

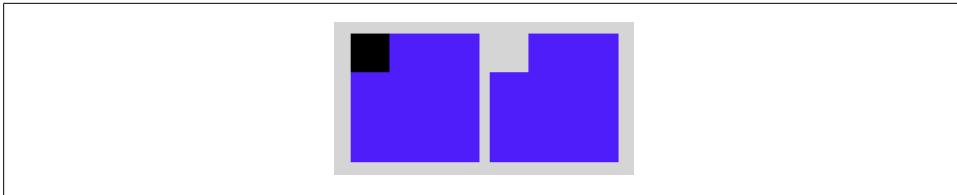


Figure 15-9. The very strange behavior of `CGContextClearRect`

There is also a function for erasing an area: `CGContextClearRect`. This erases all existing drawing in a rectangle; combined with clipping, though, it can erase an area of any shape. The result can “punch a hole” through all existing drawing.

The behavior of `CGContextClearRect` depends on whether the context is transparent or opaque. This is particularly obvious and intuitive when drawing into an image context. If the image context is transparent — the second argument to `UIGraphicsBeginImageContextWithOptions` is NO — `CGContextClearRect` erases to transparent; otherwise it erases to black.

When drawing directly into a view (as with `drawRect:` or `drawLayer:inContext:`), if the view’s background color is nil or a color with even a tiny bit of transparency, the result of `CGContextClearRect` will appear to be transparent, punching a hole right through the view including its background color; if the background color is completely opaque, the result of `CGContextClearRect` will be black. This is because *the view’s background color determines whether the view’s graphics context is transparent or opaque*; thus, this is essentially the same behavior that I described in the preceding paragraph.

[Figure 15-9](#) illustrates; the blue square on the left has been partly cut away to black, while the blue square on the right has been partly cut away to transparency. Yet these are instances of the same `UIView` subclass, drawn with exactly the same code! The difference between the views is that the `backgroundColor` of the first view is set in the nib to solid red with an alpha of 1, while the `backgroundColor` of the second view is set in the nib to solid red with an alpha of 0.99. This difference is utterly imperceptible to the eye (not to mention that the red color never appears, as it is covered with a blue fill), but it completely changes the effect of `CGContextClearRect`. The `UIView` subclass’s `drawRect:` looks like this:

```
CGContextRef con = UIGraphicsGetCurrentContext();
CGContextSetFillColorWithColor(con, [UIColor blueColor].CGColor);
CGContextFillRect(con, rect);
CGContextClearRect(con, CGRectMake(0,0,30,30));
```

To illustrate the typical use of path-drawing commands, I’ll generate the up-pointing arrow shown in [Figure 15-10](#). This might not be the best way to create the arrow, and I’m deliberately avoiding use of the convenience functions, but it’s clear and shows a nice basic variety of typical commands:

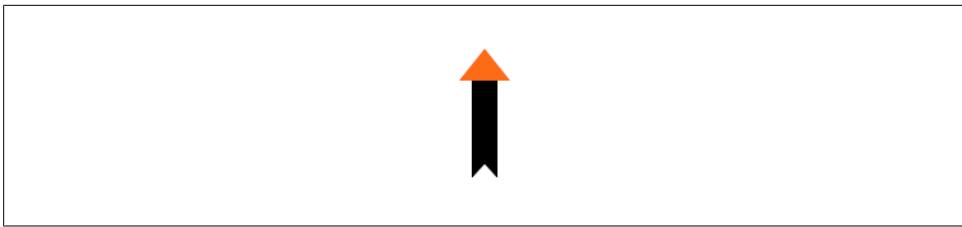


Figure 15-10. A simple path drawing

```
// obtain the current graphics context
CGContextRef con = UIGraphicsGetCurrentContext();

// draw a black (by default) vertical line, the shaft of the arrow
CGContextMoveToPoint(con, 100, 100);
CGContextAddLineToPoint(con, 100, 19);
CGContextSetLineWidth(con, 20);
CGContextStrokePath(con);

// draw a red triangle, the point of the arrow
CGContextSetFillColorWithColor(con, [[UIColor redColor] CGColor]);
CGContextMoveToPoint(con, 80, 25);
CGContextAddLineToPoint(con, 100, 0);
CGContextAddLineToPoint(con, 120, 25);
CGContextFillPath(con);

// snip a triangle out of the shaft by drawing in Clear blend mode
CGContextMoveToPoint(con, 90, 101);
CGContextAddLineToPoint(con, 100, 90);
CGContextAddLineToPoint(con, 110, 101);
CGContextSetBlendMode(con, kCGBlendModeClear);
CGContextFillPath(con);
```

Properly speaking, we should probably surround our drawing code with calls to `CGContextSaveGState` and `CGContextRestoreGState`, just in case. It probably wouldn't make any difference in this particular example, as the context does not persist between calls to `drawRect:`, but it can't hurt.

If a path needs to be reused or shared, you can encapsulate it as a `CGPath`, which is actually a `CGPathRef`. You can either create a new `CGMutablePathRef` and construct the path using various `CGPath` functions that parallel the graphics path-construction functions, or you can copy the graphics context's current path using `CGContextCopyPath`. There are also a number of `CGPath` functions for creating a path based on simple geometry (`CGPathCreateWithRect`, `CGPathCreateWithEllipseInRect`) or based on an existing path (`CGPathCreateCopyBy strokingPath`, `CGPathCreateCopyByDashingPath`, `CGPathCreateCopyByTransformingPath`).

A UIKit class, `UIBezierPath`, wraps `CGPath`. It provides methods for drawing certain path shapes, as well as for stroking, filling, and for accessing certain settings of the current graphics context state. Similarly, `UIColor` provides methods for setting the

current graphics context's stroke and fill colors. Thus we could rewrite our arrow-drawing routine like this:

```
UIBezierPath* p = [UIBezierPath bezierPath];
[p moveToPoint:CGPointMake(100,100)];
[p addLineToPoint:CGPointMake(100, 19)];
[p setLineWidth:20];
[p stroke];

[[UIColor redColor] set];
[p removeAllPoints];
[p moveToPoint:CGPointMake(80,25)];
[p addLineToPoint:CGPointMake(100, 0)];
[p addLineToPoint:CGPointMake(120, 25)];
[p fill];

[p removeAllPoints];
[p moveToPoint:CGPointMake(90,101)];
[p addLineToPoint:CGPointMake(100, 90)];
[p addLineToPoint:CGPointMake(110, 101)];
[p fillWithBlendMode:kCGBlendModeClear alpha:1.0];
```

There's no savings of code in this particular case, but `UIBezierPath` still might be useful if you need object features, and it does offer one convenience method, `bezierPathWithRoundedRect:cornerRadius:`, that is particularly attractive; drawing a rectangle with rounded corners using only Core Graphics functions is rather tedious.

Clipping

Another use of a path is to mask out areas, protecting them from future drawing. This is called *clipping*. By default, a graphics context's clipping region is the entire graphics context: you can draw anywhere within the context.

The clipping area is a feature of the context as a whole, and any new clipping area is applied by intersecting it with the existing clipping area; so if you apply your own clipping region, the way to remove it from the graphics context later is to plan ahead and wrap things with calls to `CGContextSaveGState` and `CGContextRestoreGState`.

To illustrate, I'll rewrite the code that generated our original arrow ([Figure 15-10](#)) to use clipping instead of a blend mode to "punch out" the triangular notch in the tail of the arrow. This is a little tricky, because what we want to clip to is not the region inside the triangle but the region outside it. To express this, we'll use a compound path consisting of more than one closed area — the triangle, and the drawing area as a whole (which we can obtain with `CGContextGetClipBoundingBox`).

Both when filling a compound path and when using it to express a clipping region, the system follows one of two rules:

How Big Is My Context?

At first blush, it appears that there's no way to learn a graphics context's size. Typically, this doesn't matter, because either you created the graphics context or it's the graphics context of some object whose size you know, such as a `UIView`. But in fact, because the default clipping region of a graphics context is the entire context, you can use `CGContextGetClipBoundingBox` to learn the context's "bounds" (before changing the clipping region, of course).

Winding rule

The fill or clipping area is denoted by an alternation in the direction (clockwise or counterclockwise) of the path demarcating each region.

Even-odd rule (EO)

The fill or clipping area is denoted by a simple count of the paths demarcating each region.

Our situation is extremely simple, so it's easier to use the even-odd rule. So we set up the clipping area using `CGContextEOClip` and then draw the arrow:

```
// obtain the current graphics context
CGContextRef con = UIGraphicsGetCurrentContext();

// punch triangular hole in context clipping region
CGContextMoveToPoint(con, 90, 100);
CGContextAddLineToPoint(con, 100, 90);
CGContextAddLineToPoint(con, 110, 100);
CGContextClosePath(con);
CGContextAddRect(con, CGContextGetClipBoundingBox(con));
CGContextEOClip(con);

// draw the vertical line
CGContextMoveToPoint(con, 100, 100);
CGContextAddLineToPoint(con, 100, 19);
CGContextSetLineWidth(con, 20);
CGContextStrokePath(con);

// draw the red triangle, the point of the arrow
CGContextSetFillColorWithColor(con, [[UIColor redColor] CGColor]);
CGContextMoveToPoint(con, 80, 25);
CGContextAddLineToPoint(con, 100, 0);
CGContextAddLineToPoint(con, 120, 25);
CGContextFillPath(con);
```

Gradients

Gradients can range from the simple to the complex. A simple gradient (which is all I'll describe here) is determined by a color at one endpoint along with a color at the other

endpoint, plus (optionally) colors at intermediate points; the gradient is then painted either linearly between two points in the context or radially between two circles in the context.

You can't use a gradient as a path's fill color, but you can restrict a gradient to a path's shape by clipping, which amounts to the same thing.

To illustrate, I'll redraw our arrow, using a linear gradient as the "shaft" of the arrow ([Figure 15-11](#)):

```
// obtain the current graphics context
CGContextRef con = UIGraphicsGetCurrentContext();
CGContextSaveGState(con);

// punch triangular hole in context clipping region
CGContextMoveToPoint(con, 90, 100);
CGContextAddLineToPoint(con, 100, 90);
CGContextAddLineToPoint(con, 110, 100);
CGContextClosePath(con);
CGContextAddRect(con, CGContextGetClipBoundingBox(con));
CGContextEOClip(con);

// draw the vertical line, add its shape to the clipping region
CGContextMoveToPoint(con, 100, 100);
CGContextAddLineToPoint(con, 100, 19);
CGContextSetLineWidth(con, 20);
CGContextReplacePathWithStrokedPath(con);
CGContextClip(con);

// draw the gradient
CGFloat locs[3] = { 0.0, 0.5, 1.0 };
CGFloat colors[12] = {
    0.3,0.3,0.3,0.8, // starting color, transparent gray
    0.0,0.0,0.0,1.0, // intermediate color, black
    0.3,0.3,0.3,0.8 // ending color, transparent gray
};
CGColorSpaceRef sp = CGColorSpaceCreateDeviceGray();
CGGradientRef grad = CGGradientCreateWithColorComponents (sp, colors, locs, 3);
CGContextDrawLinearGradient (con, grad, CGPointMake(89,0), CGPointMake(111,0), 0);
CGColorSpaceRelease(sp);
CGGradientRelease(grad);

CGContextRestoreGState(con); // done clipping

// draw the red triangle, the point of the arrow
CGContextSetFillColorWithColor(con, [[UIColor redColor] CGColor]);
CGContextMoveToPoint(con, 80, 25);
CGContextAddLineToPoint(con, 100, 0);
CGContextAddLineToPoint(con, 120, 25);
CGContextFillPath(con);
```

The call to `CGContextReplacePathWithStrokedPath` pretends to stroke the current path, using the current line width and other line-related context state settings, but then cre-

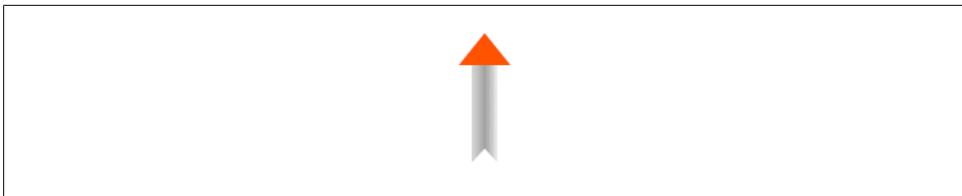


Figure 15-11. Drawing with a gradient

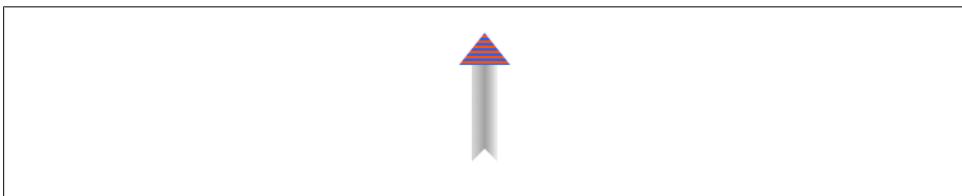


Figure 15-12. A patterned fill

ates a new path representing the outside of that stroked path. Thus, instead of a thick line we have a rectangular region that we can use as the clip region.

We then create the gradient and paint it. The procedure is verbose but simple; everything is boilerplate. We describe the gradient as a set of locations on the continuum between one endpoint (0.0) and the other endpoint (1.0), along with the colors corresponding to each location; in this case, I want the gradient to be lighter at the edges and darker in the middle, so I use three locations, with the dark one at 0.5. We must also supply a color space in order to create the gradient. Finally, we create the gradient, paint it into place, and release the color space and the gradient.

Colors and Patterns

A color is a CGColor (actually a CGColorRef). CGColor is not difficult to work with, and is bridged to UIColor through UIColor's `colorWithCGColor:` and `CGColor` methods.

A pattern, on the other hand, is a CGPattern (actually a CGPatternRef). You can create a pattern and stroke or fill with it. The process is rather elaborate. As an extremely simple example, I'll replace the red triangular arrowhead with a red-and-blue striped triangle ([Figure 15-12](#)). To do so, remove this line:

```
CGContextSetFillColorWithColor(con, [[UIColor redColor] CGColor]);
```

In its place, put the following:

```
CGColorSpaceRef sp2 = CGColorSpaceCreatePattern(NULL);
CGContextSetFillColorSpace (con, sp2);
CGColorSpaceRelease (sp2);
CGPatternCallbacks callback = {
    0, &drawStripes, NULL
};
```

```

CGAffineTransform tr = CGAffineTransformIdentity;
CGPatternRef patt = CGPatternCreate(NULL,
    CGRectMake(0,0,4,4),
    tr,
    4, 4,
    kCGPatternTilingConstantSpacingMinimalDistortion,
    true,
    &callback);

CGFloat alph = 1.0;
CGContextSetFillPattern(con, patt, &alph);
CGPatternRelease(patt);

```

That code is verbose, but it is almost entirely boilerplate. To understand it, it almost helps to read it backward. What we’re leading up to is the call to `CGContextSetFillPattern`; instead of setting a fill color, we’re setting a fill pattern, to be used the next time we fill a path (in this case, the triangular arrowhead). The third parameter to `CGContextSetFillPattern` is a pointer to a `CGFloat`, so we have to set up the `CGFloat` itself beforehand. The second parameter to `CGContextSetFillPattern` is a `CGPatternRef`, so we have to create that `CGPatternRef` beforehand (and release it afterward).

So now let’s talk about the call to `CGPatternCreate`. A pattern is a drawing in a rectangular “cell”; we have to state both the size of the cell (the second argument) and the spacing between origin points of cells (the fourth and fifth arguments). In this case, the cell is 4×4 , and every cell exactly touches its neighbors both horizontally and vertically. We have to supply a transform to be applied to the cell (the third argument); in this case, we’re not doing anything with this transform, so we supply the identity transform. We supply a tiling rule (the sixth argument). We have to state whether this is a color pattern or a stencil pattern; it’s a color pattern, so the seventh argument is `true`. And we have to supply a pointer to a callback function that actually draws the pattern into its cell (the eighth argument).

Except that that’s *not* what we have to supply as the eighth argument. To make matters more complicated, what we actually have to supply here is a pointer to a `CGPatternCallbacks` struct. This struct consists of the number `0` and pointers to *two* functions, one called to draw the pattern into its cell, the other called when the pattern is released. We’re not specifying the second function, however; it is for memory management, and we don’t need it in this simple example.

We have almost worked our way backward to the start of the code. It turns out that before you can call `CGContextSetFillPattern` with a colored pattern, you have to set the context’s fill color space to a pattern color space. If you neglect to do this, you’ll get an error when you call `CGContextSetFillPattern`. So we create the color space, set it as the context’s fill color space, and release it.

But we are *still* not finished, because I haven’t shown you the function that actually draws the pattern cell! This is the function whose address is taken as `&drawStripes` in our code. Here it is:

```

void drawStripes (void *info, CGContextRef con) {
    // assume 4 x 4 cell
    CGContextSetFillColorWithColor(con, [[UIColor redColor] CGColor]);
    CGContextFillRect(con, CGRectMake(0,0,4,4));
    CGContextSetFillColorWithColor(con, [[UIColor blueColor] CGColor]);
    CGContextFillRect(con, CGRectMake(0,0,4,2));
}

```

As you can see, the actual pattern-drawing code is very simple. The only tricky issue is that the call to `CGPatternCreate` must be in agreement with the pattern-drawing function as to the size of a cell, or the pattern won't come out the way you expect. We know in this case that the cell is 4×4. So we fill it with red, and then fill its lower half with blue. When these cells are tiled touching each other horizontally and vertically, we get the stripes that you see in [Figure 15-12](#).

Note, finally, that the code as presented has left the graphics context in an undesirable state, with its fill color space set to a pattern color space. This would cause trouble if we were later to try to set the fill color to a normal color. The solution, as usual, is to wrap the code in calls to `CGContextSaveGState` and `CGContextRestoreGState`.

You may have observed in [Figure 15-12](#) that the stripes do not fit neatly inside the triangle of the arrow-head: the bottommost stripe is something like half a blue stripe. This is because a pattern is positioned not with respect to the shape you are filling (or stroking), but with respect to the graphics context as a whole. We could shift the pattern position by calling `CGContextSetPatternPhase` before drawing.

Graphics Context Transforms

Just as a `UIView` can have a transform, so can a graphics context. However, applying a transform to a graphics context has no effect on the drawing that's already in it; it affects only the drawing that takes place after it is applied, altering the way the coordinates you provide are mapped onto the graphics context's area. A graphics context's transform is called its CTM, for “current transformation matrix.”

It is quite usual to take full advantage of a graphics context's CTM to save yourself from performing even simple calculations. You can multiply the current transform by any `CGAffineTransform` using `CGContextConcatCTM`; there are also convenience functions for applying a translate, scale, or rotate transform to the current transform.

The base transform for a graphics context is already set for you when you obtain the context; this is how the system is able to map context drawing coordinates onto screen coordinates. Whatever transforms you apply are applied to the current transform, so the base transform remains in effect and drawing continues to work. You can return to the base transform after applying your own transforms by wrapping your code in calls to `CGContextSaveGState` and `CGContextRestoreGState`.

For example, we have hitherto been drawing our upward-pointing arrow with code that knows how to place that arrow at only one location: the top left of its rectangle is



Figure 15-13. Drawing rotated with a CTM

hard-coded at `{80,0}`. This is silly. It makes the code hard to understand, as well as inflexible and difficult to reuse. Surely the sensible thing would be to draw the arrow at `{0,0}`, by subtracting 80 from all the x-values in our existing code. Now it is easy to draw the arrow at *any* position, simply by applying a translation transform beforehand, mapping `{0,0}` to the desired top left corner of the arrow. So, to draw it at `{80,0}`, we would say:

```
CGContextTranslateCTM(con, 80, 0);
// now draw the arrow at (0,0)
```

A rotate transform is particularly useful, allowing you to draw in a rotated orientation without any nasty trigonometry. However, it's a bit tricky because the point around which the rotation takes place is the origin. This is rarely what you want, so you have apply a translate transform first, to map the origin to the point around which you really want to rotate. But then, after rotating, in order to figure out where to draw you will probably have to reverse your translate transform.

To illustrate, here's code to draw our arrow repeatedly at several angles, pivoting around the end of its tail (Figure 15-13). First, we'll encapsulate the drawing of the arrow as a `UIImage`. Then we simply draw that `UIImage` repeatedly:

```
UIGraphicsBeginImageContextWithOptions(CGSizeMake(40,100), NO, 0.0);
CGContextRef con = UIGraphicsGetCurrentContext();

// draw the arrow into the image context
// draw it at (0,0)! adjust all x-values by subtracting 80
// ... actual code omitted ...

UIImage* im = UIGraphicsGetImageFromCurrentImageContext();
UIGraphicsEndImageContext();

con = UIGraphicsGetCurrentContext();

[im drawAtPoint:CGPointMake(0,0)];
for (int i=0; i<3; i++) {
    CGContextTranslateCTM(con, 20, 100);
    CGContextRotateCTM(con, 30 * M_PI/180.0);
    CGContextTranslateCTM(con, -20, -100);
    [im drawAtPoint:CGPointMake(0,0)];
}
```

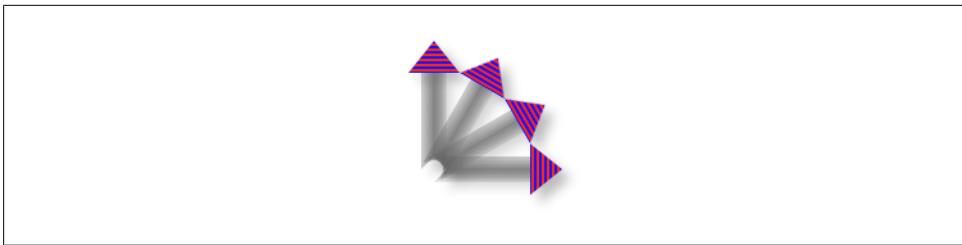


Figure 15-14. Drawing with a shadow

A transform is also one more solution for the “flip” problem we encountered earlier with `CGContextDrawImage`. Instead of reversing the drawing, we can reverse the context into which we draw it. Essentially, we apply a “flip” transform to the context’s coordinate system. You move the context’s top downward, and then reverse the direction of the y-coordinate by applying a scale transform whose y-multiplier is -1:

```
CGContextTranslateCTM(con, 0, theHeight);
CGContextScaleCTM(con, 1.0, -1.0);
```

How far down you move the context’s top depends on how you intend to draw the image. So, for example, we could draw the two halves of Mars (from the example earlier in this chapter) without flipping, like this:

```
CGContextTranslateCTM(con, 0, sz.height); // sz is [mars size]
CGContextScaleCTM(con, 1.0, -1.0);
CGContextDrawImage(con,
    CGRectMake(0,0,sz.width/2.0,sz.height),
    marsLeft);
CGContextDrawImage(con,
    CGRectMake(b.size.width-sz.width/2.0, 0, sz.width/2.0, sz.height),
    marsRight);
```

Shadows

To add a shadow to a drawing, give the context a shadow value before drawing. The shadow position is expressed as a `CGSize`, where the positive direction for both values indicates down and to the right. The blur value is an open-ended positive number; Apple doesn’t explain how the scale works, but experimentation shows that 12 is nice and blurry, 99 is so blurry as to be shapeless, and higher values become problematic.

Figure 15-14 shows the result of the same code that generated Figure 15-13, except that before we start drawing the arrow into the real context, we give the real context a shadow with this line:

```
CGContextSetShadow(con, CGSizeMake(7, 7), 12);
```

However, there’s a subtle cosmetic problem with this approach. It may not be evident from Figure 15-14, but we are adding a shadow each time we draw. Thus the arrows are able to cast shadows on one another. What we want, however, is for all the arrows to cast a single shadow collectively. The way to achieve this is with a *transparency*

layer; this is basically a subcontext that accumulates all drawing and then adds the shadow. Our code for drawing the shadowed arrows would thus look like this:

```
CGContextSetShadow(con, CGSizeMake(7, 7), 12);
CGContextBeginTransparencyLayer(con, NULL);
[im drawAtPoint:CGPointMake(0,0)];
for (int i=0; i<3; i++) {
    CGContextTranslateCTM(con, 20, 100);
    CGContextRotateCTM(con, 30 * M_PI/180.0);
    CGContextTranslateCTM(con, -20, -100);
    [im drawAtPoint:CGPointMake(0,0)];
}
CGContextEndTransparencyLayer(con);
```

Points and Pixels

A point is a dimensionless location described by an x-coordinate and a y-coordinate. When you draw in a graphics context, you specify the points at which to draw, and this works regardless of the device's resolution, because Core Graphics maps your drawing nicely onto the physical output (using the base CTM, along with any anti-aliasing and smoothing). Therefore, throughout this chapter I've concerned myself with graphics context points, disregarding their relationship to screen pixels.

However, pixels do exist. A pixel is a physical, integral, dimensioned unit of display in the real world. Whole-numbered points effectively lie between pixels, and this can matter if you're fussy, especially on a single-resolution device. For example, if a vertical path with whole-number coordinates is stroked with a line width of 1, half the line falls on each side of the path, and the drawn line on the screen of a single-resolution device will seem to be 2 pixels wide (because the device can't illuminate half a pixel).

You will sometimes encounter advice suggesting that if this effect is objectionable, you should try shifting the line's position by 0.5, to center it in its pixels. This advice may appear to work, but it makes some simple-minded assumptions. A more sophisticated approach is to obtain the `UIView`'s `contentScaleFactor` property. This value will be either `1.0` or `2.0`, so you can divide by it to convert from pixels to points. Consider also that the most accurate way to draw a vertical or horizontal line is not to stroke a path but to fill a rectangle. So this `UIView` subclass code will draw a perfect 1-pixel-wide vertical line on any device:

```
CGContextFillRect(con, CGRectMake(100,0,1.0/self.contentScaleFactor,100));
```

Content Mode

A view that draws something within itself, as opposed to merely having a background color and subviews (as in the previous chapter), has *content*. This means that its `contentMode` property becomes important whenever the view is resized. As I mentioned earlier, the drawing system will avoid asking a view to redraw itself from scratch if possible;

instead, it will use the cached result of the previous drawing operation. So, if the view is resized, the system may simply stretch or shrink or reposition the cached drawing, if your `contentMode` setting instructs it to do so.

It's a little tricky to illustrate this point, because I have to arrange for the view to be resized without also causing it to be redrawn (that is, without triggering a call to `drawRect:`). Here's how I'll do that. As the app starts up, I'll create the `MyView` instance in code and put it in the window, much as before. Then I'll use delayed performance to resize the `MyView` instance after the window has shown and the interface has been initially displayed:

```
- (BOOL)application:(UIApplication *)application  
didFinishLaunchingWithOptions:(NSDictionary *)launchOptions  
{  
    self.window = [[UIWindow alloc] initWithFrame:[[UIScreen mainScreen] bounds]];  
    self.window.rootViewController = [UIViewController new];  
  
    MyView* mv =  
        [[MyView alloc] initWithFrame:  
            CGRectMake(0, 0, self.window.bounds.size.width - 50, 150)];  
    mv.center = self.window.center;  
    [self.window.rootViewController.view addSubview: mv];  
    mv.opaque = NO;  
    mv.tag = 111; // so I can get a reference to this view later  
  
    [self performSelector:@selector(resize:) withObject:nil afterDelay:0.1];  
  
    self.window.backgroundColor = [UIColor whiteColor];  
    [self.window makeKeyAndVisible];  
    return YES;  
}  
  
- (void) resize: (id) dummy {  
    UIView* mv = [self.window viewWithTag:111];  
    CGRect f = mv.bounds;  
    f.size.height *= 2;  
    mv.bounds = f;  
}
```

We double the height of the view without causing `drawRect:` to be called. The result is that the view's drawing appears at double its correct height. For example, if our view's `drawRect:` code is the same as the code that generated [Figure 15-11](#), we get [Figure 15-15](#).

This, however, is almost certainly not what we want. Sooner or later `drawRect:` will be called, and the drawing will be refreshed in accordance with our code. Our code doesn't say to draw the arrow at a height that is relative to the height of the view's bounds; it draws the arrow at a fixed height. Thus, not only has the arrow stretched, but at some future time, it will snap back to its original size.

The moral is that our view's `contentMode` property needs to be in agreement with how the view draws itself. For example, our `drawRect:` code dictates the size and position of the arrow relative to the view's bounds origin, its top left. So we could set its `content-`

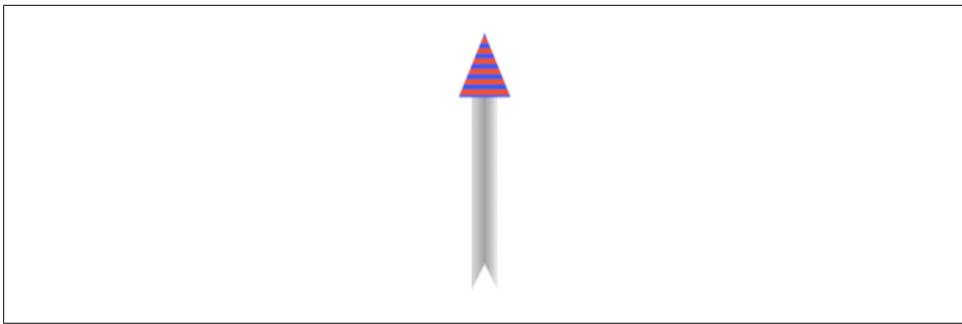


Figure 15-15. Automatic stretching of content

Mode to `UIViewContentModeTopLeft`. Alternatively, and more likely, we could set it to `UIViewContentModeRedraw`; this will cause automatic scaling and repositioning of the cached content to be turned off, and instead the view's `setNeedsDisplay` method will be called, ultimately triggering `drawRect:` to redraw the content.

On the other hand, if a view might be resized only *momentarily* — say, as part of an animation — then stretching behavior might be exactly what you want. Suppose we're going to animate the view by making it get a little larger for a moment and then returning it to its original size, perhaps as a way of attracting the user's attention. Then presumably we do want the view's content to stretch and shrink as the view stretches and shrinks; that's the whole point of the animation. This is precisely what the default `contentMode` value, `UIViewContentModeScaleToFill`, does for us. And remember, it does it efficiently; what's being stretched and shrunk is just a cached image of our view's content.

CHAPTER 16

Layers

The tale told in [Chapter 14](#) and [Chapter 15](#) of how a `UIView` works and how it draws itself is only half the story. A `UIView` has a partner called its *layer*, a `CALayer`. A `UIView` does not actually draw itself onto the screen; it draws itself into its layer, and it is the layer that appears on the screen. As I've already mentioned, a view is not redrawn frequently; instead, its drawing is cached, and the cached version of the drawing is used where possible. The cached version is, in fact, the layer. What I spoke of in [Chapter 15](#) as the view's graphics context is actually the layer's graphics context.

This might seem like a mere implementation detail, but layers are important and interesting. To understand layers is to understand views more deeply; layers extend the power of views. In particular:

Layers have properties that affect drawing.

Layers have drawing-related properties beyond those of a `UIView`. Because a layer is the recipient and presenter of a view's drawing, you can modify how a view is drawn on the screen by accessing the layer's properties. In other words, by reaching down to the level of its layer, you can make a view do things you can't do through `UIView` methods alone.

Layers can be combined within a single view.

A `UIView`'s partner layer can contain additional layers. Since the purpose of layers is to draw, portraying visible material on the screen, this allows a `UIView`'s drawing to be composed of multiple distinct pieces. This can make drawing easier, with the constituents of a drawing being treated as objects.

Layers are the basis of animation.

Animation allows you to add clarity, emphasis, and just plain coolness to your interface. Layers are made to be animated (the "CA" in "CALayer" stands for "Core Animation"). Animation is the subject of [Chapter 17](#); understanding layers is a prerequisite for reading that chapter.

For example, suppose we want to add a compass indicator to our app's interface. [Figure 16-1](#) portrays a simple version of such a compass. It takes advantage of the arrow

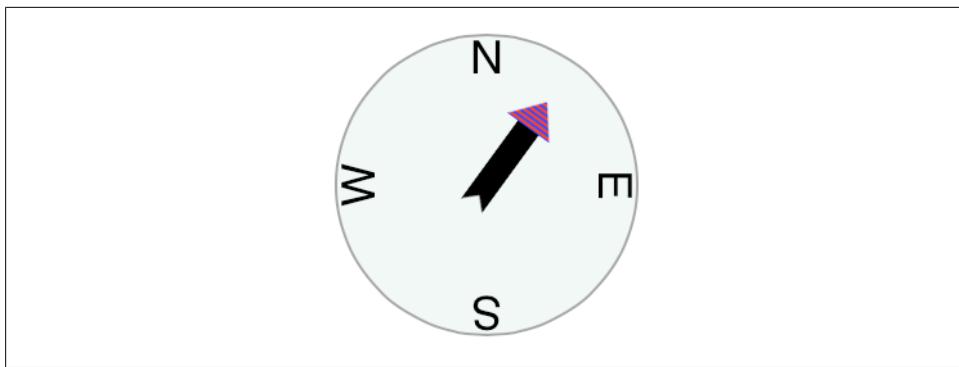


Figure 16-1. A compass, composed of layers

that we figured out how to draw in [Chapter 15](#); the arrow is drawn into a layer of its own. The other parts of the compass are layers too: the circle is a layer, and each of the cardinal point letters is a layer. The drawing is thus easy to composite in code (and later in this chapter, that's exactly what we'll do); even more intriguing, the pieces can be repositioned and animated separately, so it's easy to rotate the arrow without moving the circle (and in [Chapter 17](#), that's exactly what we'll do).

The documentation discusses layers chiefly in connection with animation (in particular, in the *Core Animation Programming Guide*). This categorization gives the impression that layers are of interest only if you intend to animate. That's misleading. Layers are the basis of animation, but they are also the basis of view drawing, and are useful and important even if you don't use them for animation.

CALayer is not part of UIKit. It's part of the Quartz Core framework, which is not linked by default into the project template. If your app contains code that refers to CALayer or related classes, you must link your target to *QuartzCore.framework*, and you must import `<QuartzCore/QuartzCore.h>` into any file containing such code.

View and Layer

A UIView instance has an accompanying CALayer instance, accessible as the view's `layer` property. This layer has a special status: it is partnered with this view to embody all of the view's drawing. The layer has no corresponding `view` property, but the view is the layer's `delegate`. The documentation sometimes speaks of this layer as the view's "underlying layer."

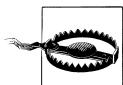
By default, when a UIView is instantiated, its layer is an instance of CALayer. But if you subclass UIView and you want your subclass's underlying layer to be an instance of a CALayer subclass (built-in or your own), implement the UIView subclass's `layerClass` class method.

That, for example, is how the compass in [Figure 16-1](#) is created. We have a UIView subclass, CompassView, and a CALayer subclass, CompassLayer. CompassView contains these lines:

```
+ (Class) layerClass {
    return [CompassLayer class];
}
```

Thus, when CompassView is instantiated, its underlying layer is a CompassLayer. In this example, there is no drawing in CompassView; its job is to give CompassLayer a place in the visible interface, because a layer cannot appear without a view.

Because every view has an underlying layer, there is a tight integration between the two. The layer is on the screen and portrays all the view's drawing; if the view draws, it does so by contributing to the layer's drawing. The view is the layer's delegate. And the view's properties are often merely a convenience for accessing the layer's properties. For example, when you set the view's `backgroundColor`, you are really setting the layer's `backgroundColor`, and if you set the layer's `backgroundColor` directly, the view's `backgroundColor` is set to match. Similarly, the view's `frame` is really the layer's `frame` and *vice versa*.



A CALayer's `delegate` property is settable, but you must never set the `delegate` property of a UIView's underlying layer. To do so would be to break the integration between them, thereby causing drawing to stop working correctly. A UIView *must* be the delegate of its underlying layer; moreover, it must *not* be the delegate of any *other* layer. *Don't do anything to mess this up.*

The view draws into its layer, and the layer caches that drawing; the layer can then be manipulated, changing the view's appearance, without necessarily asking the view to redraw itself. This is a source of great efficiency in the drawing system. It also explains such phenomena as the content stretching that we encountered in the last section of [Chapter 15](#): when the view's bounds size changes, the drawing system, by default, simply stretches or repositions the cached layer image, until such time as the view is told to generate a new drawing of itself (`drawRect:`) to replace the layer's content.



Mac OS X Programmer Alert

On Mac OS X, NSView existed long before CALayer was introduced, so today a view might have no layer, or, if it does have a layer, it might relate to it in various ways. You may be accustomed to terms like *layer-backed view* or *layer-hosting view*. On iOS, layers were incorporated from the outset: every UIView has an underlying layer and relates to it in the same way.

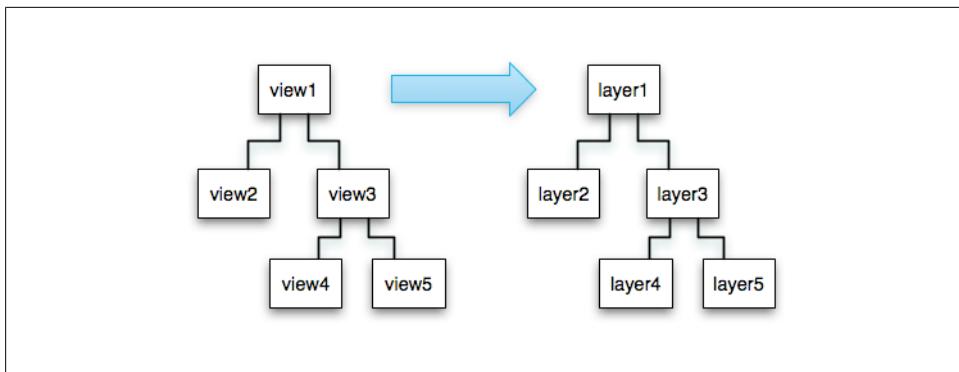


Figure 16-2. A hierarchy of views and the hierarchy of layers underlying it

Layers and Sublayers

A layer can have sublayers, and a layer has at most one superlayer. Thus there is a tree of layers. This is similar and parallel to the tree of views ([Chapter 14](#)). In fact, so tight is the integration between a view and its underlying layer, that these hierarchies are effectively the same hierarchy. Given a view and its underlying layer, that layer's superlayer is the view's superview's underlying layer, and that layer has as sublayers all the underlying layers of all the view's subviews. Indeed, because the layers are how the views actually get drawn, one might say that the view hierarchy really *is* a layer hierarchy ([Figure 16-2](#)).

At the same time, the layer hierarchy can go beyond the view hierarchy. A view has exactly one underlying layer, but a layer can have sublayers that are not the underlying layers of any view. So the hierarchy of layers that underlie views exactly matches the hierarchy of views ([Figure 16-2](#)), but the total layer tree may be a superset of that hierarchy. In [Figure 16-3](#), we see the same view-and-layer hierarchy as in [Figure 16-2](#), but two of the layers have additional sublayers that are theirs alone (that is, sublayers that are not any view's underlying layers).

From a visual standpoint, there may be nothing to distinguish a hierarchy of views from a hierarchy of layers. For example, in [Chapter 14](#) we drew three overlapping rectangles using a hierarchy of views ([Figure 14-1](#)). This code gives exactly the same visible display by manipulating layers:

```
UIView* v = self.window.rootViewController.view;
CALayer* lay1 = [[CALayer alloc] init];
lay1.frame = CGRectMake(113, 111, 132, 194);
lay1.backgroundColor = [[UIColor colorWithRed:1 green:.4 blue:1 alpha:1] CGColor];
[v.layer addSublayer:lay1];
CALayer* lay2 = [[CALayer alloc] init];
lay2.backgroundColor = [[UIColor colorWithRed:.5 green:1 blue:0 alpha:1] CGColor];
lay2.frame = CGRectMake(41, 56, 132, 194);
[lay1 addSublayer:lay2];
```

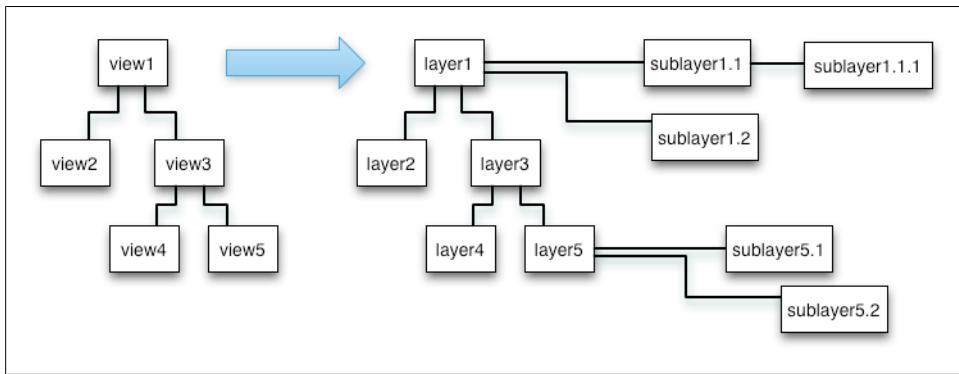


Figure 16-3. Layers that have sublayers of their own

```
CALayer* lay3 = [[CALayer alloc] init];
lay3.backgroundColor = [[UIColor colorWithRed:1 green:0 blue:0 alpha:1] CGColor];
lay3.frame = CGRectMake(43, 197, 160, 230);
[v.layer addSublayer:lay3];
```

There are, indeed, situations in which it is not clear whether a piece of interface should be constructed as a view hierarchy or a layer hierarchy. Several of my apps have an interface that is a rectangular grid of objects of the same type; in some cases, I implement these as layers, in some cases I implement them as views, and sometimes it isn't clear to me that my choice is much more than arbitrary. A layer on its own is more lightweight than a view; on the other hand, a view is a UIResponder, so it can respond to touches, and layers lack automatic layout (as implemented through a UIView's `autoresizingMask`).

Manipulating the Layer Hierarchy

Layers come with a full set of methods for reading and manipulating the layer hierarchy, parallel to the methods for reading and manipulating the view hierarchy. A layer has a `superlayer` property and a `sublayers` property; there are methods `addSublayer:`, `insertSublayer:atIndex:`, `insertSublayer:below:`, `insertSublayer:above:`, `replaceSublayer:with:`, and `removeFromSuperlayer`.

Unlike a view's `subviews` property, a layer's `sublayers` property is writable; thus, you can give a layer multiple sublayers in a single move, by assigning to its `sublayers` property. To remove all of a layer's sublayers, set its `sublayers` property to nil.

Although a layer's sublayers have an order, reflected in the `sublayers` order and regulated with the methods I've just mentioned, this is not necessarily the same as their back-to-front drawing order. By default, it is, but a layer also has a `zPosition` property, a `CGFloat`, and this also determines drawing order. The rule is that all sublayers with the same `zPosition` are drawn in the order they are listed among their `sublayers` siblings,

but lower `zPosition` siblings are drawn before higher `zPosition` siblings. (The default `zPosition` is 0.)

Sometimes, the `zPosition` property is a more convenient way of dictating drawing order than sibling order is. For example, if layers represent playing cards laid out in a solitaire game, it will likely be a lot easier and more flexible to determine how the cards overlap by setting their `zPosition` than by rearranging their sibling order.

Methods are also provided for converting between the coordinate systems of layers within the same layer hierarchy: `convertPoint:fromLayer:`, `convertPoint:toLayer:`, `convertRect:fromLayer:`, and `convertRect:toLayer:`.

Positioning a Sublayer

Layer coordinate systems and positioning are similar to those of views. A layer's own internal coordinate system is expressed by its `bounds`, just like a view; its size is its `bounds` size, and its `bounds` origin is the internal coordinate at its top left.

However, a sublayer's position within its superlayer is not described by its `center`, like a view; a layer does not have a `center`. Instead, a sublayer's position within its superlayer is defined by a combination of two properties, its `position` and its `anchorPoint`. Think of the sublayer as pinned to its superlayer; then you have to say both where the pin passes through the sublayer and where it passes through the superlayer. (I didn't make up that analogy, but it's pretty apt.)

`position`

A point expressed in the superlayer's coordinate system.

`anchorPoint`

Where the `position` point is with respect to the layer's own bounds. It is a pair of floating-point numbers (a `CGPoint`) describing a fraction (or multiple) of the layer's own bounds width and bounds height. Thus, for example, `{0,0}` is the layer's top left, and `{1,1}` is its bottom right.

If the `anchorPoint` is `{0.5,0.5}` (the default), the `position` property works like a view's `center` property. A view's `center` is thus a special case of a layer's `position`. This is quite typical of the relationship between view properties and layer properties; the view properties are often a simpler, more convenient, and less powerful version of the layer properties.

A layer's `position` and `anchorPoint` are orthogonal (independent); changing one does not change the other. Therefore, changing either of them without changing the other changes where the layer is drawn within its superlayer.

For example, in [Figure 16-1](#), the most important point in the circle is its center; all the other objects need to be positioned with respect to it. Therefore they all have the same `position`: the center of the circle. But they differ in their `anchorPoint`. For example, the arrow's `anchorPoint` is `{0.5,0.8}`, the middle of the shaft, near the end. On the other

hand, the `anchorPoint` of a cardinal point letter is more like `{0.5,3.8}`, well outside the letter's bounds, so as to place the letter near the edge of the circle.

A layer's `frame` is a purely derived property. When you get the `frame`, it is calculated from the bounds size along with the `position` and `anchorPoint`. When you set the `frame`, you set the bounds size and `position`. In general, you should regard the `frame` as a convenient façade and no more. Nevertheless, it is convenient! For example, to position a sublayer so that it exactly overlaps its superlayer, you can just set the sublayer's `frame` to the superlayer's `bounds`.



A layer created in code (as opposed to a view's underlying layer) has a `frame` and `bounds` of `{0,0},{0,0}` and will not be visible on the screen even when you add it to a superlayer that is on the screen. Be sure to give your layer a nonzero width and height if you want to be able to see it.

CAScrollLayer

If you're going to be moving a layer's bounds origin as a way of repositioning its sublayers *en masse*, you might like to make the layer a `CAScrollLayer`, a `CALayer` subclass that provides convenience methods for this sort of thing. (Despite the name, a `CAScrollLayer` provides no scrolling interface; the user can't scroll it by dragging, for example.) By default, a `CAScrollLayer`'s `masksToBounds` property is YES; thus, the `CAScrollLayer` acts like a window through which you see can only what is within its bounds. (You can set its `masksToBounds` to NO, but this would be an odd thing to do, as it somewhat defeats the purpose.)

To move the `CAScrollLayer`'s bounds, you can talk either to it or to a sublayer (at any depth):

Talking to the CAScrollLayer

`scrollToPoint:`

Changes the `CAScrollLayer`'s bounds origin to that point.

`scrollToRect:`

Changes the `CAScrollLayer`'s bounds origin minimally so that the given portion of the bounds rect is visible.

Talking to a sublayer

`scrollPoint:`

Changes the `CAScrollLayer`'s bounds origin so that the given point of *the sublayer* is at the top left of the `CAScrollLayer`.

`scrollRectVisible:`

Changes the `CAScrollLayer`'s bounds origin so that the given rect of *the sublayer's bounds* is within the `CAScrollLayer`'s bounds area. You can also ask the

sublayer for its `visibleRect`, the part of this sublayer now within the CAS-`scrollLayer`'s bounds.

Layout of Sublayers

The only option for layout of sublayers on iOS is manual layout. When a layer needs layout, either because its bounds have changed or because you called `setNeedsLayout`, you can respond in either of two ways:

- The layer's `layoutSublayers` method is called; to respond, override `layoutSublayers` in your `CALayer` subclass.
- Alternatively, implement `layoutSublayersOfLayer:` in the layer's delegate. (Remember, if the layer is a view's underlying layer, the view is its delegate.)

To do effective manual layout of sublayers, you'll probably need a way to identify or refer to the sublayers. There is no layer equivalent of `viewWithTag:`, so such identification and reference is entirely up to you. Key-value coding can be helpful here; layers implement key-value coding in a special way, discussed at the end of this chapter.



Mac OS X Programmer Alert

On Mac OS X, layers have extensive layout support, including both “springs and struts” (*constraints*) and custom layout managers. But iOS lacks all of this.

Drawing in a Layer

There are various ways to make a layer display something (apart from having a partnered view draw into it, in the ways discussed in [Chapter 15](#)).

The simplest way to make something appear in a layer is through its `contents` property. This is parallel to the `image` in a `UIImageView` ([Chapter 15](#)). It is expected to be a `CGImageRef` (or `nil`, signifying no image). A `CGImageRef` is not an object type, but the `contents` property is typed as an `id`; in order to quiet the compiler, you'll have to typecast your `CGImageRef` to an `id` as you assign it, like this:

```
layer.contents = (id)[im CGImage];
```

You may be wondering why, under ARC, we don't also have to “cross the bridge” from the `CFTyperef` world of a `CGImageRef` to the object world of an `id` by supplying a `__bridge` cast, as discussed in [Chapter 12](#). It's because the `CGImage` method is a Cocoa method and supplies ARC with the memory management information it needs. Coming back the other way, though, we would need an explicit `__bridge` cast:

```
CGImageRef imref = (__bridge CGImageRef)layer.contents;
```



Setting a layer's `contents` to a `UIImage`, rather than a `CGImage`, will fail silently—the image doesn't appear, but there is no error either. This is absolutely maddening, and I wish I had a nickel for every time I've done it and then wasted hours figuring out why my layer isn't appearing.

There are also four methods that can be implemented to provide or draw a layer's content on demand, similar to a `UIView`'s `drawRect:`. A layer is very conservative about calling these methods (and you must not call any of them directly). When a layer *does* call these methods, I will say that the layer *redisplays itself*. Here is how a layer can be caused to redisplay itself:

- If the layer's `needsDisplayOnBoundsChange` property is NO (the default), then the only way to cause the layer to redisplay itself is by calling `setNeedsDisplay` (or `setNeedsDisplayInRect:`). Even this might not cause these methods to be called right away; if that's crucial, then you will also call `displayIfNeeded`.
- If the layer's `needsDisplayOnBoundsChange` property is YES, then the layer will also redisplay itself when the layer's bounds change (rather like a `UIView`'s `UIViewContentModeRedraw`).

Here are the four methods that can be called when a layer redisplays itself; pick one to implement (don't try to combine them, you'll just confuse things):

`display` in a subclass

Your `CALayer` subclass can override `display`. There's no graphics context at this point, so `display` is pretty much limited to setting the `contents` image.

`displayLayer:` in the delegate

You can set the `CALayer`'s `delegate` property and implement `displayLayer:` in the delegate. As with `display`, there's no graphics context, so you'll just be setting the `contents` image.

`drawInContext:` in a subclass

Your `CALayer` subclass can override `drawInContext:`. The parameter is a graphics context into which you can draw directly; it is *not* automatically made the current context.

`+drawLayer:inContext:` in the delegate

You can set the `CALayer`'s `delegate` property and implement `drawLayer:inContext:`. The second parameter is a graphics context into which you can draw directly; it is *not* automatically made the current context.

Remember, you must not set the `delegate` property of a view's underlying layer! The view is its delegate and must remain its delegate. This restriction is not as onerous as it seems; there's generally an easy architectural way to draw into a layer by way of some other delegate if that's what you want to do.

For example, in one of my apps there's an overlay view, sitting on top of everything else on the screen; the user is unaware of this, because the view is transparent and

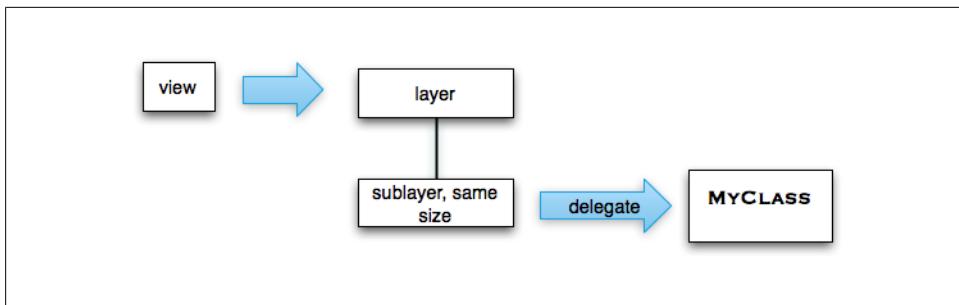


Figure 16-4. A view and a layer delegate that draws into it

usually does no drawing, and the view ignores touches, which fall through to the visible views, as if the overlay were not there at all. But every once in a while I want the overlay view to display something (this is its purpose). I don't want the overhead of making an image, and my app has a main controller, which already knows what needs drawing, so I want to draw *using this controller as a layer delegate*. But it can't be the delegate of the overlay view's underlying layer, so I give that layer a sublayer and make my main controller *that* sublayer's delegate. Thus we have a view and its underlying layer that do nothing, except to serve as a host for this sublayer (Figure 16-4) — and there's nothing wrong with that.

Assigning a layer a `contents` image and drawing directly into the layer are, in effect, mutually exclusive. So:

- If a layer's `contents` is assigned an image, this image is shown immediately and replaces whatever drawing may have been displayed in the layer.
- If a layer redisplays itself and `drawInContext:` or `drawLayer:inContext:` draws into the layer, the drawing replaces whatever image may have been displayed in the layer.
- If a layer redisplays itself and none of the `display` methods provides content (perhaps because you didn't override any of them), the layer will be empty of content.

Three layer properties strongly affect what the layer displays, in ways that can be baffling to beginners: its `backgroundColor` property, its `opaque` property, and its `opacity` property. Here's what you need to know:

- Think of the `backgroundColor` as separate from the layer's own drawing, and as painted *behind* the layer's own drawing. It is equivalent to a view's `backgroundColor` (and if this layer is a view's underlying layer, it *is* the view's `backgroundColor`). Changing the `backgroundColor` takes effect immediately.
- The `opaque` property determines whether the layer's *graphics context is opaque*. An opaque graphics context is black; you can draw on top of that blackness, but the blackness is still there. A non-opaque graphics context is clear; where no drawing is, it is completely transparent. Changing the `opaque` property has no effect until the layer redraws itself.

Automatically Redisplaying a View's Underlying Layer

A layer is not told automatically to redisplay itself (unless its bounds are resized when `needsDisplayOnBoundsChange` is YES), but a view is. For example, a view is told to redraw itself when it first appears; basically, it is sent `setNeedsDisplay`, much as if you had sent it explicitly. When a view is sent `setNeedsDisplay`, the view's underlying layer is also sent `setNeedsDisplay` — unless the view has no `drawRect:` implementation (because in that case, it is assumed that the view never needs redrawing). So, if you're drawing a view entirely by drawing its underlying layer directly, and if you want the underlying layer to be redrawn automatically when the view is told to redraw itself, you should implement `drawRect:, even if it does nothing.` (This technique has no effect on sublayers of the underlying layer.)

If a layer is a view's underlying layer, then setting the view's `backgroundColor` to an opaque color (alpha component of 1) sets the layer's `opaque`, though not the view's `opaque`, to YES. I regard this as extremely weird. (It is the reason behind the strange behavior of `CGContextClearRect` described in [Chapter 15](#).)

- The `opacity` property affects the overall apparent transparency of the layer. It is equivalent to a view's `alpha` (and if this layer is a view's underlying layer, it is the view's `alpha`). It affects the apparent transparency of the layer's sublayers as well. It affects the apparent transparency of the background color and the apparent transparency of the layer's content separately (just as with a view's `alpha`). Changing the `opacity` property takes effect immediately.

When drawing directly into a layer, the behavior of `CGContextClearRect` differs from what was described in [Chapter 15](#): instead of punching a hole through the background color, it effectively paints with the layer's background color. (This can have curious side effects.)

Content Resizing and Positioning

Once a layer has content, regardless of whether this content came from an image (setting the `contents` property) or from direct drawing into its context (`drawInContext:`, `drawLayer:inContext:`), various properties dictate how the content should be drawn in relation to the layer's bounds. It is as if the cached content is itself treated as an image, which can then be resized, repositioned, cropped, and so on. These properties are:

`contentsGravity`

This property, a string, is parallel to a `UIView`'s `contentMode` property, and describes how the content should be positioned or stretched in relation to the bounds. For example, `kCAGravityCenter` means the content is centered in the bounds without resizing; `kCAGravityResize` (the default) means the content is sized to fit the bounds, even if this means distorting its aspect; and so forth.



For historical reasons, the terms “bottom” and “top” in the names of the `contentsGravity` settings have the opposite of their expected meanings.

`contentsRect`

A `CGRect` expressing the proportion of the content that is to be displayed. The default is `{ {0,0}, {1,1} }`, meaning the entire content is displayed. The specified part of the content is sized and positioned in relation to the bounds in accordance with the `contentsGravity`. Thus, for example, you can scale up part of the content to fill the bounds, or slide part of a larger image into view without redrawing or changing the `contents` image.

You can also use the `contentsRect` to scale down the content, by specifying a larger `contentsRect` such as `{ {-.5, -.5}, {1.5, 1.5} }`; but any content pixels that touch the edge of the `contentsRect` will then be extended outwards to the edge of the layer (to prevent this, make sure that the outmost pixels of the content are all empty).

`contentsCenter`

A `CGRect` expressing the central region of nine rectangular regions of the `contentsRect` that are variously allowed to stretch if the `contentsGravity` calls for stretching. The central region (the actual value of the `contentsCenter`) stretches in both directions. Of the other eight regions (inferred from the value you provide), the four corner regions don’t stretch, and the four side regions stretch in one direction. (`UIView` has a somewhat similar property, `contentStretch`, which I didn’t discuss.)

If you’re drawing directly into the layer’s graphics context (e.g., with `drawLayer:in-Context:`), and the `contentsRect` is the entire content, then if the layer redraws itself, the `contentsGravity` won’t matter, because the graphics context fills the layer. But if the layer’s bounds are resized when `needsDisplayOnBoundsChange` is `NO`, then its cached content from the last time you drew are treated as an image. By a judicious combination of settings, you can attain some fairly sophisticated automatic behavior, with no need to redraw the content yourself. For example, [Figure 16-5](#) shows the result of the following settings:

```
arrow.needsDisplayOnBoundsChange = NO;
arrow.contentsCenter = CGRectMake(0.0, 0.4, 1.0, 0.6);
arrow.contentsGravity = kCAGravityResizeAspect;
arrow.bounds = CGRectInset(arrow.bounds, -20, -20);
```

Because `needsDisplayOnBoundsChange` is `NO`, the content is not redrawn when the arrow’s bounds are increased; instead, the cached content is used. The `contentsGravity` setting tells us to resize proportionally; therefore, the arrow is both longer and wider than in [Figure 16-1](#), but not in such a way as to distort its proportions. However, notice that although the triangular arrowhead is wider, it is not longer; the increase in

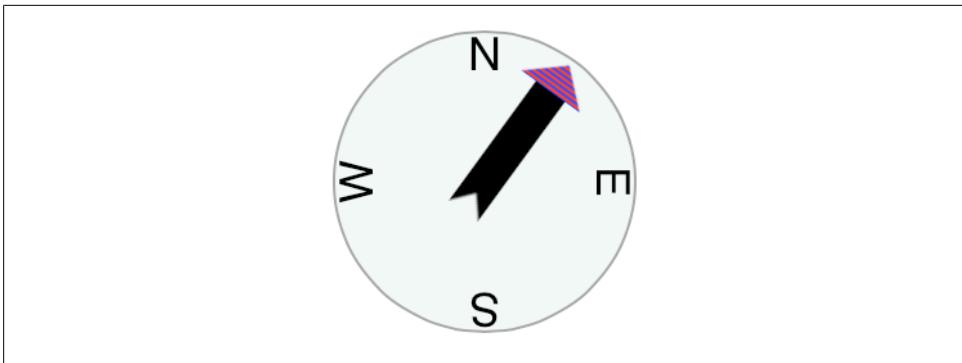


Figure 16-5. One way of resizing the compass arrow

length is due entirely to the stretching of the shaft. That's because the `contentsCenter` region is restricted to the shaft of the arrow.

If the content is larger than the bounds of the layer (which can easily happen if you're assigning a `contents` image), and if the `contentsGravity` and `contentsRect` do not resize the content to fit the bounds, then by default the content will be drawn larger than the layer; the layer does not automatically clip its content to its bounds (just as it does not automatically clip its sublayers to its bounds). To get such clipping, for both content and sublayers, set the layer's `masksToBounds` property to YES.



The value of a layer's bounds origin does not affect where its content is drawn. It affects only where its sublayers are drawn.

Layers that Draw Themselves

A few built-in CALayer subclasses provide some basic but extremely helpful self-drawing ability:

`CATextLayer`

A `CATextLayer` has a `string` property, which can be an `NSString` or `NSAttributedString`, along with other text formatting properties; it draws its `string`. The default text color, the `foregroundColor` property, is white, which is unlikely to be what you want. The text is different from the `contents` and is mutually exclusive with it: either the contents image or the text will be drawn, but not both, so in general you should not give a `CATextLayer` any contents image. In [Figure 16-1](#), the cardinal point letters are `CATextLayer` instances.

The fact that a `CATextLayer`'s `string` can be an `NSAttributedString` gives it a power that `UILabel` lacks, namely, to display text in multiple font, sizes, and styles. For example, using a `CATextLayer`, you could underline one word of the text; you can't do that with a `UILabel`. An example appears in [Chapter 23](#).

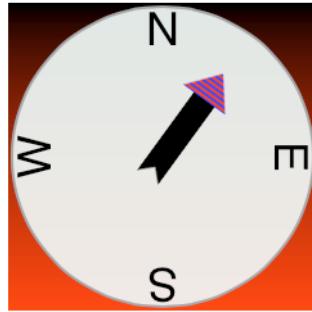


Figure 16-6. A gradient drawn behind the compass

CAShapeLayer

A CAShapeLayer has a `path` property, which is a CGPath. It fills or strokes this path, or both, depending on its `fillColor` and `strokeColor` values, and displays the result; the default is a `fillColor` of black and no `strokeColor`. A CAShapeLayer may also have `contents`; the shape is displayed on top of the contents image, but there is no property permitting you to specify a compositing mode. In [Figure 16-1](#), the background circle is a CAShapeLayer instance, stroked with gray and filled with a lighter, slightly transparent gray.

CAGradientLayer

A CAGradientLayer covers its background with a simple linear gradient; thus, it's an easy way to composite a gradient into your interface (and if you need something more elaborate you can always draw with Core Graphics instead). The gradient is defined much as in the Core Graphics gradient example in [Chapter 15](#), an array of locations and an array of corresponding colors (except that these are NSArrays, of course, not C arrays), along with a start and end point. To clip the gradient, you can add a mask to the CAGradientLayer (masks are discussed later in this chapter). A CAGradientLayer's `contents` are not displayed.

The `colors` array requires CGColors, not UIColors. But a CGColorRef is not an object type, whereas NSArray expects objects, so to quiet the compiler you'll probably need to typecast at least the first item of the array (to `id`).

[Figure 16-6](#) shows our compass drawn with an extra CAGradientLayer behind it.

Transforms

The way a layer is drawn on the screen can be modified through a transform. This is not surprising, because a view can have a transform (see [Chapter 14](#)), and a view is drawn on the screen by its layer. As with the bounds and other properties, a view and its underlying layer are tightly linked; when you change the transform of one, you are

changing the transform of the other. But, as so often happens, the layer's transform is more powerful than the view's transform. Thus, you can use the transform of the underlying layer to accomplish things with a view that you can't accomplish with the view's transform alone.

In the simplest cases, when a transform is two-dimensional, you can use the `setAffineTransform:` and `affineTransform` methods. The value is a `CGAffineTransform`, familiar from [Chapter 14](#) and [Chapter 15](#). The transform is applied around the `anchorPoint`. Thus, the `anchorPoint` has a second purpose that I didn't tell you about when discussing it earlier.

You now know everything you need to know in order to understand the code that generated [Figure 16-6](#), so here is that code. Notice how the four cardinal point letters are drawn by a `CATextLayer` and placed using a transform. They are drawn at the same coordinates, but they have different rotation transforms. Moreover, even though the `CATextLayers` are small (just 40 by 30) and appear near the perimeter of the circle, they are anchored, and so their rotation is centered, at the center of the circle. In this code, `self` is the `CompassLayer`; it does no drawing of its own, but merely assembles and configures its sublayers. To generate the arrow, we make ourselves the arrow layer's delegate and call `setNeedsDisplay`; this causes `drawLayer:inContext:` to be called in `CompassLayer` (that code is just the same code we developed for drawing the arrow into a context in [Chapter 15](#), and is not repeated here):

```
// the gradient
CAGradientLayer* g = [[CAGradientLayer alloc] init];
g.frame = self.bounds;
g.colors = [NSArray arrayWithObjects:
            (id)[[UIColor blackColor] CGColor],
            [[UIColor redColor] CGColor],
            nil];
g.locations = [NSArray arrayWithObjects:
                [NSNumber numberWithFloat: 0.0],
                [NSNumber numberWithFloat: 1.0],
                nil];
[self addSublayer:g];

// the circle
CAShapeLayer* circle = [[CAShapeLayer alloc] init];
circle.lineWidth = 2.0;
circle.fillColor =
[[UIColor colorWithRed:0.9 green:0.95 blue:0.93 alpha:0.9] CGColor];
circle.strokeColor = [[UIColor grayColor] CGColor];
CGMutablePathRef p = CGPathCreateMutable();
CGPathAddEllipseInRect(p, NULL, CGRectMakeInset(self.bounds, 3, 3));
circle.path = p;
[self addSublayer:circle];
circle.bounds = self.bounds;
circle.position = CGPointMake(CGRectGetMidX(self.bounds),
                            CGRectGetMidY(self.bounds));

// the four cardinal points
```

```

NSArray* pts = [NSArray arrayWithObjects: @"N", @"E", @"S", @"W", nil];
for (int i = 0; i < 4; i++) {
    CATextLayer* t = [[CATextLayer alloc] init];
    t.string = [pts objectAtIndex: i];
    t.bounds = CGRectMake(0, 0, 40, 30);
    t.position = CGPointMake(CGRectGetMidX(circle.bounds),
                           CGRectGetMidY(circle.bounds));
    CGFloat vert = (CGRectGetMidY(circle.bounds) - 5) / CGRectGetHeight(t.bounds);
    t.anchorPoint = CGPointMake(0.5, vert);
    t.alignmentMode = kCAAlignmentCenter;
    t.foregroundColor = [[UIColor blackColor] CGColor];
    [t setAffineTransform:CGAffineTransformMakeRotation(i*M_PI/2.0)];
    [circle addSublayer:t];
}

// the arrow
CALayer* arrow = [[CALayer alloc] init];
arrow.bounds = CGRectMake(0, 0, 40, 100);
arrow.position = CGPointMake(CGRectGetMidX(self.bounds),
                           CGRectGetMidY(self.bounds));
arrow.anchorPoint = CGPointMake(0.5, 0.8);
arrow.delegate = self;
[arrow setAffineTransform:CGAffineTransformMakeRotation(M_PI/5.0)];
[self addSublayer:arrow];
[arrow setNeedsDisplay];

```

A full-fledged layer transform, the value of the `transform` property, takes place in three-dimensional space; its description includes a z-axis, perpendicular to both the x-axis and y-axis. (By default, the positive z-axis points out of the screen, toward the viewer's face.) Layers do not magically give you realistic three-dimensional rendering — for that you would use OpenGL, which is beyond the scope of this discussion. Layers are two-dimensional objects, and they are designed for speed and simplicity. Nevertheless, they do operate in three dimensions, quite sufficiently to give a cartoonish but effective sense of reality, especially when performing an animation. We've all seen the screen image flip like turning over a piece of paper to reveal what's on the back; that's a rotation in three dimensions.

A three-dimensional transform takes place around a three-dimensional extension of the `anchorPoint`, whose z-component is supplied by the `anchorPointZ` property. Thus, in the reduced default case where `anchorPointZ` is 0, the `anchorPoint` is sufficient, as we've already seen in using `CGAffineTransform`.

The transform itself is described mathematically by a struct called a `CATransform3D`. The *Core Animation Function Reference* lists the functions for working with these transforms. They are a lot like the `CGAffineTransform` functions, except they've got a third dimension. For example, here's the declaration of the function for making a 2D scale transform:

```

CGAffineTransform CGAffineTransformMakeScale (
    CGFloat sx,
    CGFloat sy
);

```

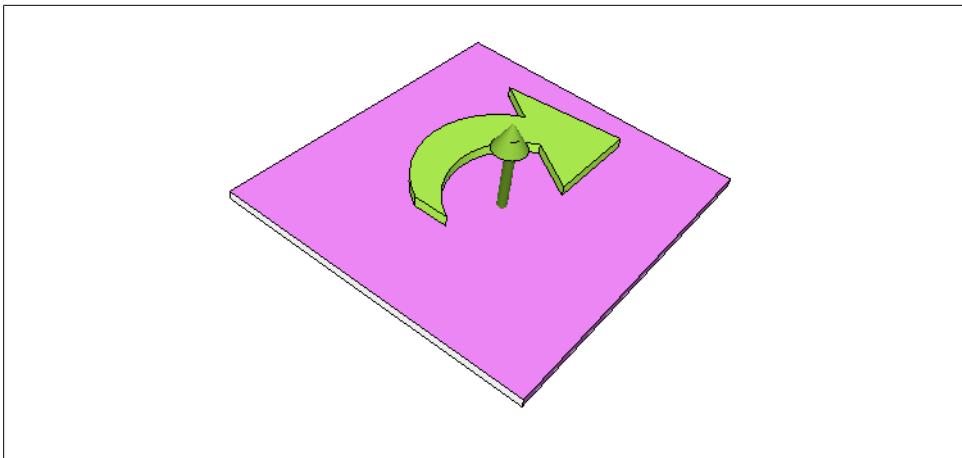


Figure 16-7. An anchor point plus a vector defines a rotation plane

And here's the declaration of the function for making a 3D scale transform:

```
CATransform3D CATransform3DMakeScale (
    CGFloat sx,
    CGFloat sy,
    CGFloat sz
);
```

The rotation 3D transform is a little more complicated. In addition to the angle, you also have to supply three coordinates describing the vector around which the rotation takes place. Perhaps you've forgotten from your high-school math what a vector is, or perhaps trying to visualize three dimensions boggles your mind, so think of it this way.

Pretend for purposes of discussion that the anchor point is the origin, $\{0,0,0\}$. Now imagine an arrow emanating from the anchor point; its other end, the pointy end, is described by the three coordinates you provide. Now imagine a plane that intersects the anchor point, perpendicular to the arrow. That is the plane in which the rotation will take place; a positive angle is a clockwise rotation, as seen from the side of the plane with the arrow (Figure 16-7). In effect, the three points you supply describe, relative to the anchor point, where your eye would have to be to see this rotation as an old-fashioned two-dimensional rotation.

The three values you give specify a direction, not a point. Thus it makes no difference on what scale you give them: $\{1,1,1\}$ means the same thing as $\{10,10,10\}$. If the three values are $\{0,0,1\}$, with all other things being equal, the case is collapsed to a simple CGAffineTransform, because the rotational plane is the screen. On the other hand, if the three values are $\{0,0,-1\}$, it's a backward CGAffineTransform, so that a positive angle looks counterclockwise (because we are looking at the "back side" of the rotational plane).

A layer can itself be rotated in such a way that its "back" is showing. For example, the following rotation flips a layer around its y-axis:

```
someLayer.transform = CATransform3DMakeRotation(M_PI, 0, 1, 0);
```

By default, the layer is considered double-sided, so when it is flipped to show its “back,” what’s drawn is an appropriately reversed version of the content of the layer (along with its sublayers). But if the layer’s `doubleSided` property is NO, then when it is flipped to show its “back,” the layer disappears (along with its sublayers); its “back” is transparent and empty.

Depth

There are two ways to place layers at different nominal depths with respect to their siblings. One is through the z-component of their `position`, which is the `zPosition` property. Thus the `zPosition`, too, has a second purpose that I didn’t tell you about earlier. The other is to apply a transform that translates the layer’s position in the z-direction. These two values (the z-component of a layer’s position and the z-component of its translation transform) are related; in some sense, the `zPosition` is a shorthand for a translation transform in the z-direction. (If you provide both a `zPosition` and a z-direction translation, you can rapidly confuse yourself.)

In the real world, changing an object’s `zPosition` would make it appear larger or smaller, as it is positioned closer or further away; but this, by default, is not the case in the world of layer drawing. There is no attempt to portray perspective; the layer planes are drawn at their actual size and flattened onto one another, with no illusion of distance. (This is called *orthographic projection*, and is the way blueprints are often drawn to display an object from one side.)

However, there’s a widely used trick for introducing a quality of perspective into the way layers are drawn: make them sublayers of a layer whose `sublayerTransform` property maps all points onto a “distant” plane. (This is probably just about the only thing the `sublayerTransform` property is ever used for.) Combined with orthographic projection, the effect is to apply one-point perspective to the drawing, so that things do get perceptibly smaller in the negative z-direction.

For example, let’s try applying a sort of “page-turn” rotation to our compass: we’ll anchor it at its right side and then rotate it around the y-axis. For purposes of the example, the sublayer we’re actually rotating is accessed through a property, `rotationLayer`:

```
self->rotationLayer.anchorPoint =
    CGPointMake(1,0.5);
self->rotationLayer.position =
    CGPointMake(CGRectGetMaxX(self.bounds), CGRectGetMidY(self.bounds));
self->rotationLayer.transform =
    CATransform3DMakeRotation(M_PI/4.0, 0, 1, 0);
```

The results are disappointing ([Figure 16-8](#)); the compass looks more squashed than rotated. Now, however, we’ll also apply the distance-mapping transform. The super-layer here is `self`:



Figure 16-8. A disappointing page-turn rotation

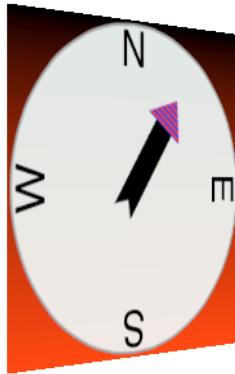


Figure 16-9. A dramatic page-turn rotation

```
CATransform3D transform = CATransform3DIdentity;
transform.m34 = -1.0/1000.0;
self.sublayerTransform = transform;
```

The results (shown in Figure 16-9) are better, and you can experiment with values to replace `1000.0`; for example, `500.0` gives an even more exaggerated effect. Also, the `z-Position` of the `rotationLayer` will now affect how large it is.

Another way to draw layers with depth is to use `CATransformLayer`. This `CALayer` subclass doesn't do any drawing of its own; it is intended solely as a host for other layers. It has the remarkable feature that you can apply a transform to it and it will maintain the depth relationships among its sublayers.

Figure 16-10 shows our page-turn rotation yet again, still with the `sublayerTransform` applied to `self`, but this time the only sublayer of `self` is a `CATransformLayer`. The `CATransformLayer`, to which the page-turn transform is applied, holds the gradient layer, the circle layer, and the arrow layer. Those three layers are at different depths



Figure 16-10. Page-turn rotation applied to a CATransformLayer

(using different `zPosition` settings), and you can see that the circle layer floats in front of the gradient layer. (This is clear from its apparent offset, but I wish you could see this page-turn as an animation, which makes the circle jump right out from the gradient as the rotation proceeds.) I've also tried to emphasize the arrow's separation from the circle by adding a shadow.

Even more remarkable, note the little white peg sticking through the arrow and running into the circle. It is a `CAShapeLayer`, rotated to be perpendicular to the `CATransformLayer`. Normally, it runs straight out of the circle toward the viewer, so it is seen end-on, and because a layer has no thickness, it is invisible. But as the `CATransformLayer` pivots forward in our page-turn rotation, the peg maintains its orientation relative to the circle, and comes into view.

There is, I think, a slight additional gain in realism if the same `sublayerTransform` is applied also to the `CATransformLayer`, but I have not done so here.

Shadows, Borders, and More

A `CALayer` has many additional properties that affect details of how it is drawn. Once again, all of these drawing details can, of course, be applied equally to a `UIView`; changing these properties of the `UIView`'s underlying layer changes how the view is drawn. Thus, these are effectively view features as well.

A `CALayer` can have a shadow, defined by its `shadowColor`, `shadowOpacity`, `shadowRadius`, and `shadowOffset` properties. To make the layer draw a shadow, set the `shadowOpacity` to a nonzero value. The shadow is normally based on the shape of the layer's nontransparent region, but deriving this shape can be calculation-intensive (so much so that in early versions of iOS, layer shadows weren't implemented). You can vastly



Figure 16-11. A layer with a mask

improve performance by defining the shape yourself and assigning this shape as a `CGPath` to the `shadowPath` property.

A CALayer can have a border (`borderWidth`, `borderColor`); the `borderWidth` is drawn inward from the bounds, potentially covering some of the content unless you compensate.

A CALayer can be bounded by a rounded rectangle, by giving it a `cornerRadius` greater than zero. If the layer has a border, the border has rounded corners too. If the layer has a `backgroundColor`, that background is clipped to the shape of the rounded rectangle. If the layer's `masksToBounds` is YES, the layer's content and its sublayers are clipped by the rounded corners.

Like a UIView, a CALayer has a `hidden` property that can be set to take it and its sublayers out of the visible interface without actually removing it from its superlayer.

A CALayer can have a `mask`. This is itself a layer, whose content must be provided somehow. The transparency of the mask's content in a particular spot becomes (all other things being equal) the transparency of the layer at that spot. For example, [Figure 16-11](#) shows our arrow layer, with the gray circle layer behind it, and a mask applied to the arrow layer. The mask is silly, but it illustrates very well how masks work: it's an ellipse, with an opaque fill and a thick, semitransparent stroke. Here's the code that generates and applies the mask:

```
CAShapeLayer* mask = [[CAShapeLayer alloc] init];
mask.frame = arrow.bounds;
CGMutablePathRef p2 = CGPathCreateMutable();
CGPathAddEllipseInRect(p2, NULL, CGRectMakeInset(mask.bounds, 10, 10));
mask.strokeColor = [[UIColor colorWithWhite:0.0 alpha:0.5] CGColor];
mask.lineWidth = 20;
mask.path = p2;
arrow.mask = mask;
CGPathRelease(p2);
```

To position the mask, pretend it's a sublayer. The hues in the mask's colors are irrelevant; only transparency matters.

Using a mask, we can do manually and in a more general way what the `cornerRadius` and `masksToBounds` properties do, clipping all our content drawing and our sublayers (and, if this layer is a view's underlying layer, the view's subviews) to a desired path. Here's a utility method that generates a CALayer suitable for use as a rounded rectangle mask:

```

- (CALayer*) maskOfSize:(CGSize)sz roundingCorners:(CGFloat)rad {
    CGRect r = (CGRect){CGPointZero, sz};
    UIGraphicsBeginImageContextWithOptions(r.size, NO, 0);
    CGContextRef con = UIGraphicsGetCurrentContext();
    CGContextSetFillColorWithColor(con,[UIColor colorWithWhite:0 alpha:0].CGColor);
    CGContextFillRect(con, r);
    CGContextSetFillColorWithColor(con,[UIColor colorWithWhite:0 alpha:1].CGColor);
    UIBezierPath* p = [UIBezierPath bezierPathWithRoundedRect:r cornerRadius:rad];
    [p fill];
    UIImage* im = UIGraphicsGetImageFromCurrentImageContext();
    UIGraphicsEndImageContext();
    CALayer* mask = [CALayer layer];
    mask.frame = r;
    mask.contents = (id)im.CGImage;
    return mask;
}

```

By now, you’re probably envisioning all sorts of compositing fun, with layers masking sublayers and laid semitransparently over other layers. There’s nothing wrong with that, but when an iOS device is asked to animate the movement of its drawing from place to place, the animation may stutter because the device lacks the necessary computing power to composite repeatedly and rapidly. This sort of issue is likely to emerge particularly when your code performs an animation ([Chapter 17](#)) or when the user is able to animate drawing through touch, as when scrolling a table view ([Chapter 21](#)). You may be able to detect these problems by eye, and you can quantify them on a device by using the Core Animation template, which shows the frame rate achieved during animation. Also, both the Core Animation template and the Simulator’s Debug menu let you summon colored overlays that provide clues as to possible sources of inefficient drawing which can lead to such problems.

In general, opaque drawing is most efficient. If a layer will always be shown over a background consisting of a single color, you can give the layer its own background of that same color; when additional layer content is supplied, the visual effect will be the same as if that additional layer content were composited over a transparent background. For example, instead of an image masked to a rounded rectangle (with a layer’s `cornerRadius` or `mask` property), you could use Core Graphics to clip the drawing of that image into the graphics context of a layer whose background color is the same as that of the destination in front of which the drawing will be shown. Here’s an example from a view’s `drawRect:` in one of my own apps:

```

// clip to rounded rect
CGRect r = CGRectMakeInset(rect, 1, 1);
[[UIBezierPath bezierPathWithRoundedRect:r cornerRadius:6] addClip];
// draw image
UIImage* im = [UIImage imageNamed: @"linen.jpg"];
// simulate UIViewContentModeScaleAspectFill
// make widths the same, image height will be too tall
CGFloat scale = im.size.width/rect.size.width;

```

```
CGFloat y = (im.size.height/scale - rect.size.height) / 2.0;
CGRect r2 = CGRectMake(0,-y,im.size.width/scale, im.size.height/scale);
r2 = CGRectIntegral(r2); // looks a whole lot better if we do this
[im drawInRect:r2];
```

Another way to gain some efficiency is by “freezing” the entirety of the layer’s drawing as a bitmap. In effect, you’re drawing everything in the layer to a secondary cache and using the cache to draw to the screen. To do this, set the layer’s `shouldRasterize` to YES and its `rasterizationScale` to some sensible value (probably `[UIScreen mainScreen].scale`). You can always turn rasterization off again by setting `shouldRasterize` to NO, so it’s easy to rasterize just before some massive or sluggish rearrangement of the screen and then unrasterize afterward. (In addition, you can get some cool “out of focus” effects by setting the `rasterizationScale` to around 0.3.)

Layers and Key–Value Coding

All of a layer’s properties are accessible through key–value coding by way of keys with the same name as the property. Thus, to apply a mask to a layer, instead of saying this:

```
layer.mask = mask;
```

we could have said:

```
[layer setValue: mask forKey: @"mask"];
```

In addition, `CATransform3D` and `CGAffineTransform` values can be expressed through key–value coding and key paths. For example, instead of writing this earlier:

```
self->rotationLayer.transform =
    CATransform3DMakeRotation(M_PI/4.0, 0, 1, 0);
```

we could have written this:

```
[self->rotationLayer setValue:[NSNumber numberWithFloat:M_PI/4.0]
    forKeyPath:@"transform.rotation.y"];
```

This notation is possible because `CATransform3D` is key–value coding compliant for a repertoire of keys and key paths. These are not properties, however; a `CATransform3D` doesn’t have a `rotation` property. It doesn’t have *any* properties, because it isn’t even an object. You cannot say:

```
self->rotationLayer.transform.rotation.y = //... No, sorry
```

The transform key paths you’ll use most often are `rotation.x`, `rotation.y`, `rotation.z`, `rotation` (same as `rotation.z`), `scale.x`, `scale.y`, `scale.z`, `translation.x`, `translation.y`, `translation.z`, and `translation` (two-dimensional, a `CGSize`).

The Quartz Core framework also injects KVC compliance into `CGPoint`, `CGSize`, and `CGRect`, allowing you to use keys and key paths matching their struct component names. For a complete list of KVC compliant classes related to `CALayer`, along with the keys and key paths they implement, plus rules for how to wrap nonobject values

as objects, see “Core Animation Extensions to Key-Value Coding” in the *Core Animation Programming Guide*.

Moreover, you can treat a CALayer as a kind of NSDictionary, and get and set the value for *any* key. This is tremendously useful, because it means you can attach arbitrary information to an individual layer instance and retrieve it later. For example, earlier I mentioned that to apply manual layout to a layer’s sublayers, you will need a way of identifying those sublayers. This feature could provide a way of doing that.

Also, CALayer has a `defaultValueForKey:` class method; to implement it, you’ll need to subclass and override. In the case of keys whose value you want to provide a default for, return that value; otherwise, return the value that comes from calling `super`. Thus, even if a value for a particular key has never been explicitly provided, it can have a non-nil value.

The truth is that this feature, though delightful (and I often wish that all classes behaved like this), is not put there for your convenience and enjoyment. It’s there to serve as the basis for animation, which is the subject of the next chapter.

Animation

Animation is the visible change of an attribute over time. The changing attribute might be positional, but not necessarily. For example, a view’s background color might change from red to green, not instantly, but perceptibly fading from one to the other. Or a view’s opacity might change from opaque to transparent, not instantly, but perceptibly fading away.

Without help, most of us would find animation beyond our reach. There are just too many complications — complications of calculation, of timing, of screen refresh, of threading, and many more. Fortunately, help is provided. You don’t perform an animation yourself; you describe it, you order it, and it is performed for you. You get *animation on demand*.

Asking for an animation can be as simple as setting a property value; under some circumstances, a single line of code will result in animation:

```
myLayer.backgroundColor = [[UIColor redColor] CGColor]; // animate change to red
```

And this is no coincidence. Apple wants to facilitate your use of animation. Animation is crucial to the character of the iOS interface. It isn’t just cool and fun; it clarifies that something is changing or responding. For example, one of my first apps was based on a Mac OS X game in which the user clicks cards to select them. In the Mac OS X version, a card was highlighted to show it was selected, and the computer would beep to indicate a click on an ineligible card. On iOS, these indications were insufficient: the highlighting felt weak, and you can’t use a sound warning in an environment where the user might have the volume turned off or be listening to music. So in the iOS version, animation is the indicator for card selection (a selected card wiggles eagerly) and for tapping on an ineligible card (the whole interface shudders, as if to shrug off the tap).

Recall from [Chapter 16](#) that CALayer requires the Quartz Core framework; so do the other “CA” (Core Animation) classes discussed here, such as CAAnimation. You’ll link your target to *QuartzCore.framework* and import *<QuartzCore/QuartzCore.h>*.



As you experiment with animation code, bear in mind that the Simulator's Debug → Toggle Slow Animations menu item helps you inspect animations by making them run more slowly.

Drawing, Animation, and Threading

When you change a visible view property *without* animation, that change does *not* visibly take place there and then. Rather, the system records that this is a change you would like to make, and marks the view as needing to be redrawn. You can change many visible view properties, but these changes are all just accumulated for later. Later, when all your code has run to completion and the system has, as it were, a free moment, then it redraws all views that need redrawing, applying their new visible property features. I call this the *redraw moment*. (The documentation calls it “when the [current] thread’s run-loop next iterates.”)

You can see that this is true simply by changing some visible aspect of a view and changing it back again, in the same code: on the screen, nothing happens. For example, suppose a view’s background color is green. Suppose your code changes it to red, and then later changes it back to green:

```
// view starts out green
view.backgroundColor = [UIColor redColor];
// ... time-consuming code goes here ...
view.backgroundColor = [UIColor greenColor];
// code ends, redraw moment arrives
```

The system accumulates all the desired changes until the redraw moment happens, and the redraw moment doesn’t happen until after your code has finished, so when the redraw moment does happen, the last accumulated change in the view’s color is to green — which is its color already. Thus, no matter how much time-consuming code lies between the change from green to red and the change from red to green, the user won’t see any color change at all.

(That’s why you don’t order a view to be redrawn; rather, you tell it that it *needs* redrawing — `setNeedsDisplay` — at the next redraw moment. It’s also why I used delayed performance in the `contentMode` example in [Chapter 15](#): by calling `performSelector:withObject:afterDelay:`, I give the redraw moment a chance to happen, thus giving the view some content, *before* resizing the view. This use of delayed performance to let a redraw moment happen is quite common, though later in this chapter I’ll suggest another way of accomplishing the same goal.)

Similarly, when you ask for an animation to be performed, the animation doesn’t start happening on the screen until the next redraw moment. (You can force an animation to be performed immediately, but this is unusual.)

While the animation lasts, it is effectively in charge of the screen. Imagine that the animation is a kind of movie, a cartoon, interposed between the user and the “real”

screen. When the animation is finished, this movie is removed, revealing the state of the “real” screen behind it. The user is unaware of this, because (if you’ve done things correctly) at the time that it starts, the movie’s first frame looks just like the state of the “real” screen at that moment, and at the time that it ends, the movie’s last frame looks just like the state of the “real” screen at *that* moment.

So, when you reposition a view from position 1 to position 2 with animation, you can envision a typical sequence of events like this:

1. The view is set to position 2, but there has been no redraw moment, so it is still portrayed at position 1.
2. The rest of your code runs to completion.
3. The redraw moment arrives. If there were no animation, the view would now be portrayed at position 2. But there *is* an animation, and it (the “animation movie”) starts with the view portrayed at position 1, so that is still what the user sees.
4. The animation proceeds, portraying the view at intermediate positions between position 1 and position 2. The documentation describes the animation as now *in-flight*.
5. The animation ends, portraying the view ending up at position 2.
6. The “animation movie” is removed, revealing the view indeed at position 2.

Animation takes place on an independent thread. Multithreading is generally rather tricky and complicated, but the system makes it easy in this case. Nevertheless, you can’t completely ignore the threaded nature of animation. Awareness of threading issues, and having a mental picture of how animation is performed, will help you to ask yourself the right questions and thus to avoid confusion and surprises. For example:

1. *The time when an animation starts is somewhat indefinite (because you don’t know exactly when the next redraw moment will be). The time when an animation ends is also somewhat indefinite (because the animation happens on another thread, so your code cannot just wait for it to end). So what if your code needs to do something in response to an animation beginning or ending?*

An animation can have a delegate; there is a delegate message that is sent when an animation starts, and another when it ends. Thus, you can arrange to receive an event at these crucial moments. On iOS 4 and later, you can also supply a block to be run after an animation ends.

2. *Since animation happens on its own thread, something might cause code of yours to start running while an animation is still in-flight. What happens if your code now changes a property that is currently being animated? What happens if your code asks for another animation?*

If you change a property while it is being animated, it won’t tie the system in knots, but the end result may look odd, if the value you set differs from the final value in the animation. If a property is being animated from value 1 to value 2 and mean-

while you set it to value 3, then the property may appear very suddenly to take on value 3 (because, in effect, the animation movie is removed, and its final frame shown doesn't agree with the state of things revealed behind it). If that isn't what you intend, don't do that; on the other hand, this can be a useful feature, as it provides a coherent way of effectively canceling an in-flight animation.

If you ask for an animation when an animation is already scheduled for the next redraw moment or already in-flight, there might be no problem; both animations can take place simultaneously. But that's impossible if both animations attempt to animate *the same property*. In that case, the first animation may be forced to end instantly; that is, the change it represents ceases to be animated and is portrayed as happening suddenly instead. This is typically not what's intended. But there are many alternative approaches. If you want to chain animations, you can wait until one animation ends (using the delegate message to learn when that is) before ordering the next one. Or you can create a single animation combining multiple changes; these changes needn't start at the same moment or be the same length. And a simple call (such as `setAnimationBeginsFromCurrentState:`) will "blend" the second animation with the first.

3. *While an animation is in-flight, if your code is not running, the interface is responsive to the user. What happens if the user tries to tap a view whose position is currently being animated?*

The problem is a very real one: the view might not really be where it appears to be on the screen, so the user might try to tap it and miss, or might tap elsewhere and accidentally tap it, because where the user tapped is where the view really is. The usual way of coping is to turn off responsiveness in your app's interface. (I'll talk more in [Chapter 18](#) about how to let the user tap a view while it's in animated motion.)

To prevent the interface as a whole from responding while an animation is in-flight, you can call the `UIApplication` instance method `beginIgnoringInteractionEvents` when the animation starts and call `endIgnoringInteractionEvents` when the animation is over (possibly using the delegate messages to learn when those things happen). If that's too broad, you can block responsiveness to touches at the level of individual views; for example, you can turn off a view's `userInteractionEnabled` until the animation is over. But all of this is up to you; the system has, generally speaking, no policy of automatically disabling touch responsiveness. (But there's one exception: if you use block-based view animation, the system *does* turn off user interaction for an animated view, while it is being animated, by default.)

4. *In a multitasking world, the user can suspend my app without quitting it. What happens if an animation is in-flight at that moment?*

If your app is suspended ([Chapter 11](#)) during animation, the animation is removed. This simply means that the "animation movie" is cancelled. Any animation, whether in-flight or scheduled, is simply a slowed-down visualization of a property

Presentation Layer

There isn't really an "animation movie" in front of the screen — though the effect is much the same. In reality, it is not a layer that draws itself on the screen; it's a derived layer called the *presentation layer*. Thus, when you animate the change of a view's position or a layer's position from position 1 to position 2, its nominal position changes immediately; meanwhile, the presentation layer's position remains unchanged until the redraw moment, and then changes over time, and because that's what's actually drawn on the screen, that's what the user sees.

A layer's presentation layer can be accessed through its `presentationLayer` property (and the layer itself is the presentation layer's `modelLayer`). It is typed as an `id`, so in order to work with it as a layer, you will probably want to typecast it to a `CALayer`*. Accessing the `presentationLayer` is not a common thing to do, but it might come in handy if your code needs to learn the current state of an in-flight animation.

change; that property is still changed, and indeed was probably changed before the animation even started. If your app is resumed, therefore, no animations will be running, and properties that were changed remain changed, and are shown as changed.

UIImageView and UIImage Animation

`UIImageView` provides a form of animation that is so simple and crude as to be scarcely deserving of the name. Nevertheless, sometimes this form of animation is all you need — a trivial solution to what might otherwise be a tricky problem. Supply the `UIImageView` with an array of `UIImage`s, as the value of its `animationImages` or `highlightedAnimationImages` property; this causes the `image` or `highlightedImage` to be hidden. This array represents the "frames" of a simple cartoon; when you send the `startAnimating` message, the images are displayed in turn, at a frame rate determined by the `animationDuration` property, repeating as many times as specified by the `animationRepeatCount` property (the default is 0, meaning to repeat forever, or until the `stopAnimating` message is received).

For example, suppose we want an image of Mars to appear out of nowhere and flash three times on the screen. This might seem to require some sort of `NSTimer`-based solution (see [Chapter 11](#)), but it's far simpler to use an animating `UIImageView`:

```
UIImage* mars = [UIImage imageNamed: @"mars.png"];
UIGraphicsBeginImageContextWithOptions(mars.size, NO, 0);
UIImage* empty = UIGraphicsGetImageFromCurrentImageContext();
UIGraphicsEndImageContext();
NSArray* arr = [NSArray arrayWithObjects: mars, empty, mars, empty, mars, nil];
UIImageView* iv = [[UIImageView alloc] initWithFrame:CGRectMake(56, 63, 208, 208)];
[self.window.rootViewController.view addSubview: iv];
```

```
iv.animationImages = arr;
iv.animationDuration = 2;
iv.animationRepeatCount = 1;
[iv startAnimating];
```

You can combine UIImageView animation with other kinds of animation. For example, you could flash the image of Mars while at the same time sliding the UIImageView rightward, using view animation as described in the next section.

In addition, starting in iOS 5, UIImage itself supplies a simple form of animation parallel to what UIImageView provides. An image can itself be an *animated image*. Just as with UIImageView, this really means that you've multiple images that form a sequence serving as the "frames" of a simple cartoon. You can designate an image as an animated image with one of two UIImage class methods:

`animatedImageWithImages:duration:`

As with UIImageView's `animationImages`, you supply an array of UIImages. You also supply the duration for the whole animation.

`animatedImageNamed:duration:`

You supply the name of a single image file at the top level of your app bundle, as with `imageNamed:` — except that you omit the file extension, and the system does not look for an image file by this name. Instead, it appends @"0" to the name you supply (and then, I presume, several different possible file extensions) and looks for *that* image file, and makes it the first image in the animation sequence. Then it appends @"1" to the name you supply and looks for *that* image file. And so on, until it fails to find any more image files with any of these constructed names, up through @"1024". It is fine if image files for some constructed names don't exist; for example, `animatedImageNamed:@"moi"` works even if the only "moi" image files are `moi101.png` and `moi293.png`.

A third method, `animatedResizableImageNamed:capInsets:duration:` combines an animated image with a resizable image ([Chapter 15](#)).

You do not tell an animated image to start animating, nor are you able to tell it how long you want the animation to repeat. Rather, an animated image is *always animating*, so long as it appears in your interface; to control the animation, add the image to your interface or remove it from the interface, possibly exchanging it for a similar image that isn't animated. Moreover, an animated image can appear in the interface *anywhere* a UIImage can appear as a property of some interface object. So, it can appear in a UIImageView, but it can also appear as the background of a UIButton or a UINavigationBar, and so forth.

In this example, I construct a sequence of red circles of different sizes, in code, and build an animated image which I then display in a UIButton to the left of its title:

```
NSMutableArray* arr = [NSMutableArray array];
float w = 18;
for (int i = 0; i < 6; i++) {
    UIGraphicsBeginImageContextWithOptions(CGSizeMake(w,w), NO, 0);
```

```

CGContextRef con = UIGraphicsGetCurrentContext();
CGContextSetFillColorWithColor(con, [UIColor redColor].CGColor);
CGContextAddEllipseInRect(con, CGRectMake(0+i,0+i,w-i*2,w-i*2));
CGContextFillPath(con);
UIImage* im = UIGraphicsGetImageFromCurrentImageContext();
UIGraphicsEndImageContext();
[arr addObject:im];
}
UIImage* im = [UIImage animatedImageWithImages:arr duration:0.5];
UIButton* b = [UIButton buttonWithType:UIButtonTypeRoundedRect];
[b setTitle:@"Howdy" forState:UIControlStateNormal];
[b setImage:im forState:UIControlStateNormal];
b.center = CGPointMake(100,100);
[b sizeToFit];
[self.window.rootViewController.view addSubview:b];

```

But please, use your mighty powers for good, not evil. Remember the early days of web browsers, when people discovered that HTML could flash text and overused this ability, to horrible effect? I do *not* want to see any apps that overuse animated images like that!

View Animation

Animation is ultimately layer animation. However, for a limited range of attributes, you can animate a `UIView` directly: these are its `alpha`, `backgroundColor`, `bounds`, `center`, `frame`, and `transform`. You can also animate a `UIView`'s change of contents. Despite the brevity of the list, `UIView` animation is an excellent way to become acquainted with animation and to experiment with the various parameters you can use to determine how an animation behaves; in many cases it will prove quite sufficient.

There are actually two ways to ask for `UIView` animation: the old way using animation blocks (before iOS 4.0, and still available), and the new way using Objective-C blocks ([Chapter 3](#)). I'll describe the old way first; the documentation describes this approach as "discouraged," but it is not officially deprecated and it does still work, and it may offer some advantages over the newer block-based animation.

Animation Blocks

To animate a change to an animatable `UIView` property the old way, wrap the change in calls to the `UIView` class methods `beginAnimations:context:` and `commitAnimations`. The region between these calls is referred to as an *animation block*, even though it is *not* a block in the syntactical Objective-C sense.

So, animating a change to a view's background color could be as simple as this:

```

[UIView beginAnimations:nil context:NULL];
v.backgroundColor = [UIColor yellowColor];
[UIView commitAnimations];

```

Any animatable change made within an animation block will be animated, so we can animate a change both to the view's color and its position simultaneously:

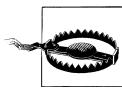
```
[UIView beginAnimations:nil context:NULL];
v.backgroundColor = [UIColor yellowColor];
CGPoint p = v.center;
p.y -= 100;
v.center = p;
[UIView commitAnimations];
```

We can also animate changes to multiple views. For example, suppose we want to make one view dissolve into another. We start with the second view present in the view hierarchy, but with an `alpha` of 0, so that it is invisible. Then we animate the change of the first view's `alpha` to 0 and the second view's `alpha` to 1, simultaneously. This might be a way, for example, to make the text of a label or the title of a button appear to dissolve while changing.

The two parameters to `beginAnimations:context:` are an `NSString` and a pointer-to-void that are completely up to you; the idea is that an animation can have a delegate (so that you can be notified when the animation starts and ends), and you can supply values here that will be passed along in the delegate messages, helping you identify the animation and so forth.

Modifying an Animation Block

An animation has various characteristics that you can modify, and an animation block provides a way to make such modifications: within the animation block, you call a `UIView` class method whose name begins with `setAnimation....`



Some of the `setAnimation....` method calls are oddly picky as to whether they precede or follow the actual property value changes within the animation block. If a call seems to be having no effect, try moving it to the beginning or end of the animation block. I find that in general these calls work best if they *precede* the value changes.

Animation blocks can be nested. The result is a single animation, whose description is not complete until the outermost animation block is terminated with `commitAnimations`. Therefore, by using `setAnimation...` method calls in the different nested animation blocks, you can give the parts of the animation different characteristics. Within each animation block, the animation for any property changes will have the default characteristics unless you change them.



Nested animation blocks are different from successive top-level animation blocks; successive top-level animation blocks are different animations, which, as I mentioned earlier, can have undesirable effects, possibly causing the earlier animation to be cancelled abruptly.

Here are the `setAnimation...` `UIView` class methods:

`setAnimationDuration:`

Sets the “speed” of the animation, by dictating (in seconds) how long it takes to run from start to finish. Obviously, if two views are told to move different distances in the same time, the one that must move further must move faster.

`setAnimationRepeatAutoreverses:`

If YES, the animation will run from start to finish (in the given duration time), and will then run from finish to start (also in the given duration time).

`setAnimationRepeatCount:`

Sets how many times the animation should be repeated. Unless the animation also autoreverses, the animation will “jump” from its end to its start to begin the next repetition. The value is a float, so it is possible to end the repetition at some midpoint of the animation.

`setAnimationCurve:`

Describes how the animation changes speed during its course. Your options are:

- `UIViewControllerAnimatedCurveEaseInOut` (the default)
- `UIViewControllerAnimatedCurveEaseIn`
- `UIViewControllerAnimatedCurveEaseOut`
- `UIViewControllerAnimatedCurveLinear`

The term “ease” means that there is a gradual acceleration or deceleration between the animation’s central speed and the zero speed at its start or end.

`setAnimationStartDate:, setAnimationDelay:`

These are both ways of postponing the start of the animation; in my experience, the former is and always has been broken, so you should use `setAnimationDelay:` exclusively.

`setAnimationDelegate:`

Arranges for your code to be notified as the animation starts or ends; the methods to be called on the delegate are specified as follows:

`setAnimationWillStartSelector:`

The `start` method must take two parameters; these are the values passed into `beginAnimations:context:`, namely an identifying `NSString` and a pointer-to-void. This method is not called unless something within the animation block triggers an actual animation.

`setAnimationDidStopSelector:`

The `stop` method must take three parameters: the second parameter is a `BOOL` wrapped as an `NSNumber`, indicating whether the animation completed successfully, and the other two are like the “start” method parameters. This method is called, with the second parameter representing YES, even if nothing within the animation block triggers any animations.

setAnimationsEnabled:

Call this with a NO argument to perform subsequent animatable property changes within the animation block without making them part of the animation.

setAnimationBeginsFromCurrentState:

If YES, and if this animation animates a property already being animated by an animation that is previously ordered or in-flight, then instead of canceling the previous animation (completing the requested change instantly), this animation will use the presentation layer to decide where to start, and will “blend” its animation with the previous animation if possible.

If an animation autoreverses, and if, when the animation ends, the view’s actual property is still at the finish value, the view will appear to “jump” from start to finish as the “animation movie” is removed. So, for example, suppose we want a view to animate its position to the right and then back to its original position. This code causes the view to animate right, animate left, and then (unfortunately) jump right:

```
[UIView beginAnimations:nil context:NULL];
[UIView setAnimationRepeatAutoreverses:YES];
CGPoint p = v.center;
p.x += 100;
v.center = p;
[UIView commitAnimations];
```

How can we prevent this? We want the view to stay at the start value after the animation reverses and ends. If we try to eliminate the jump at the end by setting the view’s position back to its starting point after the animation block, there is no animation at all (because when the redraw moment arrives, there is no property change):

```
[UIView beginAnimations:nil context:NULL];
[UIView setAnimationRepeatAutoreverses:YES];
CGPoint p = v.center;
p.x += 100;
v.center = p;
[UIView commitAnimations];
p = v.center;
p.x -= 100;
v.center = p;
```

The coherent solution is to use the “stop” delegate method to set the view’s position back to its starting point when the animation ends:

```
- (void) someMethod {
    [UIView beginAnimations:nil context:NULL];
    [UIView setAnimationRepeatAutoreverses:YES];
    [UIView setAnimationDelegate:self];
    [UIView setAnimationDidStopSelector:@selector(stopped:fin:context:)];
    CGPoint p = v.center;
    p.x += 100;
    v.center = p;
    [UIView commitAnimations];
}
```

```

- (void) stopped:(NSString *)anim fin:(NSNumber*)fin context:(void *)context {
    CGPoint p = v.center;
    p.x -= 100;
    v.center = p;
}

```

In that example, we happened to know how the animation had changed the view's position, so we could hard-code the instructions for reversing the change. To be more general, we could take advantage of our ability to pass information into the animation block and retrieve this same information in the delegate method. Or, we could store the view's original position in its layer (recall that a CALayer is a dictionary-like container):

```

- (void) someMethod {
    [UIView beginAnimations:nil context:NULL];
    [UIView setAnimationRepeatAutoreverses:YES];
    [UIView setAnimationDelegate:self];
    [UIView setAnimationDidStopSelector:@selector(stopped:fin:context:)];
    CGPoint p = v.center;
    [v.layer setValue:[NSValue valueWithCGPoint:p] forKey:@"origCenter"];
    p.x += 100;
    v.center = p;
    [UIView commitAnimations];
}

- (void) stopped:(NSString *)anim fin:(NSNumber*)fin context:(void *)context {
    v.center = [[v.layer valueForKey:@"origCenter"] CGPointValue];
}

```

To illustrate `setAnimationBeginsFromCurrentState:`, consider the following:

```

[UIView beginAnimations:nil context:NULL];
[UIView setAnimationDuration:1];
CGPoint p = v.center;
p.x += 100;
v.center = p;
[UIView commitAnimations];

[UIView beginAnimations:nil context:NULL];
// uncomment the next line to fix the problem
//[UIView setAnimationBeginsFromCurrentState:YES];
[UIView setAnimationDuration:1];
CGPoint p2 = v.center;
p2.x = 0;
v.center = p2;
[UIView commitAnimations];

```

The result is that the view jumps 100 points rightward, and then animates leftward. That's because the second animation caused the first animation to be thrown away; the move 100 points rightward was performed instantly, instead of being animated. But if we uncomment the call to `setAnimationBeginsFromCurrentState:`, the result is that the view animates leftward from its current position, with no jump.

Even more interesting is what happens when we change `x` to `y` in the second animation. If we uncomment the call to `setAnimationBeginsFromCurrentState:`, both the `x`-component and the `y`-component of the view's position are animated together, as if we had ordered one animation instead of two.

Transition Animations

A *transition* is a sort of animated redrawing of a view. The usual reason for a transition animation is that you are making some change in the view's appearance, and you want to emphasize this by animating the view. To order a transition animation using an animation block, call `setAnimationTransition:forView:cache:`.

- The first parameter describes how the animation should behave; your choices are:
 - `UIViewControllerAnimatedFlipFromLeft`
 - `UIViewControllerAnimatedFlipFromRight`
 - `UIViewControllerAnimatedCurlUp`
 - `UIViewControllerAnimatedCurlDown`
- The second parameter is the view.
- The third parameter is whether to cache the view's contents right now, effectively taking a “snapshot” of those contents at the moment and as they will be after the contents change, and using these snapshots throughout the transition. The alternative is to redraw the contents repeatedly throughout the transition. You'll want to say YES wherever possible.

Here's a simple example that flips a `UIImageView` while changing its image. The result is that the `UIImageView` appears to flip over, like a piece of paper being rotated to show its reverse side — a piece of paper with Mars on its front and Saturn on its back:

```
[UIView beginAnimations:nil context:NULL];
[UIView setAnimationTransition:UIViewControllerAnimatedFlipFromLeft
    forView:iv cache:YES];
// iv is a UIImageView whose image is Mars.png
iv.image = [UIImage imageNamed:@"Saturn.gif"];
[UIView commitAnimations];
```

The example is a little misleading, because the change in the image does not necessarily have to be inside the animation block. The animation described by `setAnimationTransition:...` will be performed in any case. The change of image will be performed in any case as well. They will both happen at the redraw moment, so they are performed together. Thus, we could have written the same example this way:

```
iv.image = [UIImage imageNamed:@"Saturn.gif"];
[UIView beginAnimations:nil context:NULL];
[UIView setAnimationTransition:UIViewControllerAnimatedFlipFromLeft
    forView:iv cache:YES];
[UIView commitAnimations];
```

Nevertheless, it is customary to order the changes in the view from inside the animation block, and I'll continue to do so in subsequent examples.

You can do the same sort of thing with any built-in view subclass. Here's a button that seems to be labeled "Start" on one side and "Stop" on the other:

```
[UIView beginAnimations:nil context:NULL];
// "b" is a UIButton; "stopped" is presumably a BOOL variable or ivar
[UIView setAnimationTransition:UIViewAnimationTransitionFlipFromLeft
    forView:b cache:YES];
[b setTitle:(stopped ? @"Start" : @"Stop") forState:UIControlStateNormal];
[UIView commitAnimations];
```

To do the same thing with a custom UIView subclass that knows how to draw itself in its `drawRect:`, call `setNeedsDisplay` to cause a redraw. For example, imagine a UIView subclass with a `reverse` BOOL property, which draws an ellipse if `reverse` is YES and a square if `reverse` is NO. Then we can animate the square flipping over and becoming an ellipse (or *vice versa*):

```
v.reverse = !v.reverse;
[UIView beginAnimations:nil context:NULL];
[UIView setAnimationTransition:UIViewAnimationTransitionFlipFromLeft
    forView:v cache:YES];
[v setNeedsDisplay];
[UIView commitAnimations];
```

Of course you can also animate a view while doing such things to it as removing or adding a subview.

Block-Based View Animation

A UIView can also be animated using a syntax involving Objective-C blocks. This is intended to replace the old animation block syntax I've just been describing. In the new syntax:

- The behavior to be animated is a block.
- The code to be run when the animation ends is also a block. Thus, there is no need for the two-part structure involving an animation block and a separate delegate method.
- Options describing the animation are part of the original animation method call, not separate calls as with an animation block.
- Transition animations have more options than with animation blocks.
- User touch interactions with an animated view are disabled, by default. This is not the case with an animation block (or with layer-based animation, described later in this chapter). The option `UIViewControllerAnimatedOptionAllowUserInteraction` permits user touch interaction with the animated view.



Under iOS 4, the rule was more comprehensive: *all* user touch interactions were disabled during a block-based view animation, unless you used `UIViewAnimationOptionAllowUserInteraction`. *This is a major difference in behavior between iOS 4 and iOS 5.*

The basis of the new syntax is the `UIView` class method `animateWithDuration:delay:options:animations:completion:`. There are also two reduced calls, the first letting you omit the `delay` and `options` parameters and the second letting you also omit the `completion` parameter. The parameters of the full form are:

duration

The duration of the animation.

delay

The delay before the animation starts. The default, in the reduced forms, is no delay.

options

A bitmask stating additional options. The default, in the reduced forms, is `UIViewControllerAnimatedOptionCurveEaseInOut` (which is also the default animation curve for animation blocks). For an ordinary animation (not a transition), the chief options are:

Animation curve

Your choices are:

- `UIViewControllerAnimatedOptionCurveEaseInOut`
- `UIViewControllerAnimatedOptionCurveEaseIn`
- `UIViewControllerAnimatedOptionCurveEaseOut`
- `UIViewControllerAnimatedOptionCurveLinear`

Repetition and autoreverse

Your options are:

- `UIViewControllerAnimatedOptionRepeat`
- `UIViewControllerAnimatedOptionAutoreverse`

There is no way to specify a certain number of repetitions; you either repeat forever or not at all. This feels like an oversight (a serious oversight); I'll suggest a workaround in a moment. The documentation's claim that you can autoreverse only if you also repeat is incorrect.

animations

The block containing view property changes to be animated.

completion

The block to run when the animation ends. It takes one `BOOL` parameter indicating whether the animation ran to completion. (There is no way to specify a notification when the animation starts, but this should not be needed, as the animation

code is itself a block.) It's fine for this block to order a further animation. This block is called, with a parameter indicating YES, even if nothing in the `animations` block triggers any animations.

Here's an example, recasting an earlier example to use Objective-C blocks instead of animation blocks. We move a view rightward and reverse it back into place. With animation blocks, we used a delegate so that we could set the view back to its original position, and we stored that position in the layer so as to be able to retrieve it in the delegate method. With blocks, however, the original position can live in a variable that remains in scope, so things are much simpler (to increase readability, I've expressed the blocks and the options as named variables):

```
CGPoint p = v.center;
CGPoint pOrig = p;
p.x += 100;
void (^anim) (void) = ^{
    v.center = p;
};
void (^after) (BOOL) = ^(BOOL f) {
    v.center = pOrig;
};
NSUInteger opts = UIViewAnimationOptionAutoreverse;
[UIView animateWithDuration:1 delay:0 options:opts
    animations:anim completion:after];
```

Working around the inability to specify a finite number of repetitions is not easy. Here's one approach using recursion:

```
- (void) animate: (int) count {
    CGPoint p = v.center;
    CGPoint pOrig = p;
    p.x += 100;
    void (^anim) (void) = ^{
        v.center = p;
    };
    void (^after) (BOOL) = ^(BOOL f) {
        v.center = pOrig;
        if (count)
            [self animate: count-1];
    };
    NSUInteger opts = UIViewAnimationOptionAutoreverse;
    [UIView animateWithDuration:1 delay:0 options:opts
        animations:anim completion:after];
}
```

If we call the `animate` method with an argument of 2, our animation takes place three times and stops. There is always a danger, with recursion, of filling up the stack and running out of memory, but I think we're safe if we start with a small `count` value.

In addition to the options I've already listed, there are some options saying what should happen if an animation is already ordered or in-flight.

`UIViewControllerAnimatedOptionBeginFromCurrentState`

Similar to `setAnimationBeginsFromCurrentState:`.

`UIViewControllerAnimatedOptionOverrideInheritedDuration`

Prevents inheriting duration from an already ordered or in-flight animation (the default is to inherit it).

`UIViewControllerAnimatedOptionOverrideInheritedCurve`

Prevents inheriting the animation curve from an already ordered or in-flight animation (the default is to inherit it).

A widely used technique for canceling a repeating animation is to start another animation of the same property of the same view. (Reread the first section of this chapter if you don't understand why that would work.) This example builds on our previous examples; we have an autoreversing animation of our view's `center` that repeats, nominally infinitely. To stop it, we animate the same `center` property of the same view back to its original position. This is a great opportunity to use `UIViewControllerAnimatedOptionBeginFromCurrentState`; without it, the animation ends abruptly. Two different methods need access to the view's original center, so I've put it in a static global:

```
static CGPoint pOrig;

- (void) animate {
    CGPoint p = v.center;
    pOrig = p;
    p.x += 100;
    void (^anim) (void) = ^{
        v.center = p;
    };
    void (^after) (BOOL) = ^(BOOL f) {
        v.center = pOrig;
    };
   NSUInteger opts = UIViewAnimationOptionAutoreverse |
                      UIViewAnimationOptionRepeat;
    [UIView animateWithDuration:1 delay:0 options:opts
                          animations:anim completion:after];
}

- (void) cancelAnimation {
    void (^anim) (void) = ^{
        v.center = pOrig;
    };
    NSUInteger opts = UIViewAnimationOptionBeginFromCurrentState;
    [UIView animateWithDuration:0.2 delay:0 options:opts
                          animations:anim completion:nil];
}
```

There is also a layout option, `UIViewControllerAnimatedOptionLayoutSubviews`. This is useful if the view that you are about to animate does its layout manually through an override of `layoutSubviews` ([Chapter 14](#)). In that case, if you supply this option, `layoutSubviews` is called while we are still within the animation block; the changes ordered by your `layoutSubviews` implementation will thus be part of the animation. If you *don't* supply this

option, the changes ordered by your `layoutSubviews` implementation will appear as a sudden jump as the animation begins.

Transitions are ordered using one of two methods. The one that's parallel to `setAnimationTransition:forView:cache:`, described earlier in connection with animation blocks, is `transitionWithView:duration:options:animations:completion:`. The transition animation types are expressed as part of the `options` bitmask; the last three are new in iOS 5, and have no parallel in the older animation block syntax:

- `UIViewControllerAnimatedOptionTransitionFlipFromLeft`
- `UIViewControllerAnimatedOptionTransitionFlipFromRight`
- `UIViewControllerAnimatedOptionTransitionCurlUp`
- `UIViewControllerAnimatedOptionTransitionCurlDown`
- `UIViewControllerAnimatedOptionTransitionCrossDissolve`
- `UIViewControllerAnimatedOptionTransitionFlipFromBottom`
- `UIViewControllerAnimatedOptionTransitionFlipFromTop`

Here's a recasting, using `transitionWithView...`, of the earlier example where we flip a rectangle into an ellipse by means of a custom `UIView` subclass whose `drawRect:` behavior depends on its `reverse` property:

```
v.reverse = !v.reverse;
void (^anim) (void) = ^{
    [v setNeedsDisplay];
};
NSUInteger opts = UIViewAnimationOptionTransitionFlipFromLeft;
[UIView transitionWithView:v duration:1 options:opts
    animations:anim completion:nil];
```

During a transition, by default, a snapshot of the view's final appearance is used; this is parallel to what happens when you supply YES as the `cache:` argument to `setAnimationTransition:forView:cache:`. If that isn't what you want, use `UIViewControllerAnimatedOptionAllowAnimatedContent` in the `options` bitmask. For example, suppose `v` is the view to be animated using a transition, and `v2` is a subview of `v` that currently occupies part of its width. In this block, to be used as the animation during the transition, we increase the width of `v2` to occupy the entire width of `v`:

```
void (^anim) (void) = ^{
    CGRect f = v2.frame;
    f.size.width = v.frame.size.width;
    f.origin.x = 0;
    v2.frame = f;
};
```

Without `UIViewControllerAnimatedOptionAllowAnimatedContent`, that change in the frame of `v2` takes place in a jump after the transition is over. With `UIViewControllerAnimatedOptionAllowAnimatedContent`, it is seen to happen smoothly as part of the transition animation.

The second transition method is `transitionFromView:toView:duration:options:completion:`. It names two views; the first is replaced by the second, while their superview undergoes the transition animation. This has no parallel in the older animation block syntax. There are actually two possible configurations, depending on the options you provide:

Remove one subview, add the other

If `UIViewControllerAnimatedOptionShowHideTransitionViews` is *not* one of the options, then the second subview is not in the view hierarchy when we start; the first subview is removed from its superview and the second subview is added to that same superview.

Hide one subview, show the other

If `UIViewControllerAnimatedOptionShowHideTransitionViews` is one of the options, then both subviews are in the view hierarchy when we start; the `hidden` of the first is NO, the `hidden` of the second is YES, and these values are reversed.

So, for example, this code causes the superview of `v1` to rotate like a piece of paper being turned over, while at the same `v1` is removed from it and `v2` is added to it:

```
NSUInteger opts = UIViewAnimationOptionTransitionFlipFromLeft;
[UIView transitionFromView:v1 toView:v2 duration:1 options:opts completion:nil];
```

It's up to you to make sure beforehand that `v2` has the desired position, so that it will appear in the right place in its superview.

Implicit Layer Animation

If a layer is not a view's underlying layer, animating it can be as simple as setting a property. A change in what the documentation calls an *animatable property* is automatically interpreted as a request to animate that change. In other words, animation of layer property changes is the default! Multiple property changes are considered part of the same animation. This mechanism is called *implicit animation*.



You cannot use implicit animation on the underlying layer of a `UIView`. You can animate a `UIView`'s underlying layer directly, but you must use explicit layer animation (discussed later in this chapter).

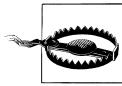
For example, in [Chapter 16](#) we constructed a compass out of layers. The compass itself is a `CompassView` that does no drawing of its own; its underlying layer is a `CompassLayer` that also does no drawing, serving only as a superlayer for the layers that constitute the drawing. None of the layers that constitute the actual drawing is the underlying layer of a view, so a property change to any of them is animated automatically.

So, presume that we have a reference to the arrow layer, a property `arrow` of the `CompassLayer`, and also a reference to the `CompassView`, a property `compass` of the app delegate, which is `self`. If we rotate the arrow by changing its `transform` property, that rotation is animated:

```
CompassLayer* c = (CompassLayer*)self.compass.layer;
// the next line is an implicit animation
c.arrow.transform = CATransform3DRotate(c.arrow.transform, M_PI/4.0, 0, 0, 1);
```

`CALayer` properties listed in the documentation as animatable in this way are `anchorPoint` and `anchorPointZ`, `backgroundColor`, `borderColor`, `borderWidth`, `bounds`, `contents`, `contentsCenter`, `contentsRect`, `cornerRadius`, `doubleSided`, `hidden`, `masksToBounds`, `opacity`, `position` and `zPosition`, `rasterizationScale` and `shouldRasterize`, `shadowColor`, `shadowOffset`, `shadowOpacity`, `shadowRadius`, and `sublayerTransform` and `transform`. In addition, a `CAShapeLayer`'s `path`, `fillColor`, `strokeColor`, `lineWidth`, `lineDashPhase`, and `miterLimit` are animatable; so are a `CATextLayer`'s `fontSize` and `foregroundColor`. (See [Chapter 16](#) for discussion of what all those properties do.)

Basically, a property is animatable because there's some sensible way to interpolate the intermediate values between one value and another. The nature of the animation attached to each property is therefore just what you would intuitively expect. When you change a layer's `hidden` property, it fades out of view (or into view). When you change a layer's `contents`, the old contents are dissolved into the new contents. And so forth.



Observe that I didn't say `frame` was an animatable property. That's because it isn't an animatable property! A `CALayer`'s frame is a purely derived value. To animate the changing of a layer's frame, you'll change other properties such as its `bounds` and `position`.

Animation Transactions

Implicit animation operates with respect to a *transaction* (a `CATransaction`), which groups animation requests into a single animation. Every animation request takes place in the context of a transaction. You can make this explicit by wrapping your animation requests in calls to the `CATransaction` class methods `begin` and `commit`; the result is a *transaction block*. But additionally there is already an *implicit transaction* surrounding all your code, and you can operate on this implicit transaction without any `begin` and `commit`.

To modify the characteristics of an implicit animation, you modify its transaction. Typically, you'll use these `CATransaction` class methods:

`setAnimationDuration`:

The duration of the animation.

`setAnimationTimingFunction`:

A `CAMediaTimingFunction`; timing functions are discussed in the next section.

`setCompletionBlock:`

A block to be called when the animation ends. The block takes no parameters. The block is called even if no animation is triggered during this transaction.

By nesting transaction blocks, you can apply different animation characteristics to different elements of an animation. But you can also use transaction commands outside of any transaction block to modify the implicit transaction.

So, in our previous example, we could slow down the animation of the arrow like this:

```
CompassLayer* c = (CompassLayer*)self.compass.layer;
[CATransaction setAnimationDuration:0.8];
c.arrow.transform = CATransform3DRotate(c.arrow.transform, M_PI/4.0, 0, 0, 1);
```

Another useful feature of animation transactions is to turn implicit animation *off*. It's important to be able to do this, because implicit animation is the default, and can be unwanted (and a performance drag). To do so, call the CATransaction class method `setDisableActions:` with argument YES. There are other ways to turn off implicit animation (discussed later in this chapter), but this is the simplest.

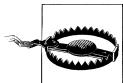
Observe that, since `setCompletionBlock:` is called even if there is no animation, it constitutes a way of saying what to do after the next redraw moment. Previously ([Chapter 15](#)) I did that using delayed performance:

```
[self performSelector:@selector(resize:) withObject:nil afterDelay:0.1];
```

A completion block on the implicit transaction might be a better way:

```
[CATransaction setCompletionBlock:^{[self resize:nil]}];
```

CATransaction implements KVC to allow you set and retrieve a value for an arbitrary key, similar to CALayer. An example appears later in this chapter.



An explicit transaction block that orders an animation to a layer, if the block is *not preceded by any other changes to the layer*, can cause animation to begin immediately when the CATransaction class method `commit` is called, without waiting for the redraw moment, while your code continues running. In my experience, this can cause confusion (for example, animation delegate messages cannot arrive, and the presentation layer can't be queried properly) and should be avoided.

Media Timing Functions

The CATransaction class method `setAnimationTimingFunction:` takes as its parameter a media timing function (CAMediaTimingFunction). This class is the general expression of the animation curves we have already met (ease-in-out, ease-in, ease-out, and linear); in fact, you are most likely to use it with those very same predefined curves, by calling the CAMediaTimingFunction class method `functionWithName:` with one of these parameters:

- `kCAMediaTimingFunctionLinear`

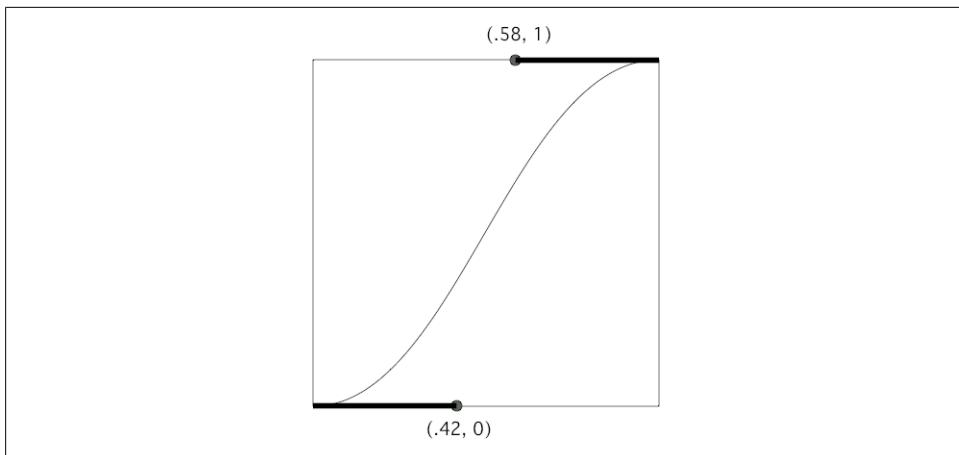


Figure 17-1. An ease-in-out Bézier curve

- `kCAMediaTimingFunctionEaseIn`
- `kCAMediaTimingFunctionEaseOut`
- `kCAMediaTimingFunctionEaseInEaseOut`
- `kCAMediaTimingFunctionDefault`

In reality, a media timing function is a Bézier curve defined by two points. The curve graphs the fraction of the animation’s time that has elapsed (the x-axis) against the fraction of the animation’s change that has occurred (the y-axis); its endpoints are therefore at $\{0,0\}$ and $\{1,1\}$, because at the beginning of the animation there has been no elapsed time and no change, and at the end of the animation all the time has elapsed and all the change has occurred.

Because the curve’s defining points are its endpoints, each needs only one Bézier control point to define the tangent to the curve. And because the curve’s endpoints are known, defining the two control points is sufficient to describe the entire curve. And because a point is a pair of floating point values, a media timing function can be expressed as four floating-point values. That is, in fact, how it is expressed.

So, for example, the ease-in-out timing function is expressed as the four values `0.42`, `0.0`, `0.58`, `1.0`. That defines a Bézier curve with one endpoint at $\{0,0\}$, whose control point is $\{0.42,0\}$, and the other endpoint at $\{1,1\}$, whose control point is $\{0.58,1\}$ ([Figure 17-1](#)).

If you want to define your own media timing function, you can supply the coordinates of the two control points by calling `functionWithControlPoints::::` or `initWithControlPoints:::::`; this is one of those rare cases, mentioned in [Chapter 3](#), where the parameters of an Objective-C method have no name. (It helps to design the curve in a standard drawing program first so that you can visualize how the placement of the control points shapes the curve.) For example, here’s a media timing function that starts out quite

slowly and then whips quickly into place after about two-thirds of the time has elapsed. I call this the “clunk” timing function, and it looks great with the compass arrow:

```
CAMediaTimingFunction* clunk =
    [CAMediaTimingFunction functionWithControlPoints:.9 :.1 :.7 :.9];
[CATransaction setAnimationTimingFunction: clunk];
c.arrow.transform = CATransform3DRotate(c.arrow.transform, M_PI/4.0, 0, 0, 1);
```

Core Animation

Core Animation is the fundamental underlying iOS animation technology. View animation and implicit layer animation are merely convenient façades for Core Animation. Core Animation is *explicit layer animation*, and revolves primarily around the CAAnimation class and its subclasses, which allow you to create far more elaborate specifications of an animation than anything we’ve encountered so far.

You may never program at the level of Core Animation, but you should read this section anyway, if only to learn how animation really works and to get a sense of the mighty powers you would acquire if you *did* elect to use Core Animation directly. In particular, Core Animation:

- Works even on a view’s underlying layer. Thus, Core Animation is the *only* way to apply full-on layer property animation to a view.
- Provides fine control over the intermediate values and timing of an animation.
- Allows animations to be grouped into complex combinations.
- Adds transition animation effects that aren’t available otherwise, such as new content “pushing” the previous content out of a layer.

CABasicAnimation and Its Inheritance

The simplest way to animate a property with Core Animation is with a CABasicAnimation object. CABasicAnimation derives much of its power through its inheritance, so I’m going to describe that inheritance as well as CABasicAnimation itself. You will readily see that all the property animation features we have met so far are embodied in a CABasicAnimation instance.

CAAnimation

CAAnimation is an abstract class, meaning that you’ll only ever use a subclass of it. Some of CAAnimation’s powers come from its implementation of the CAMediaTiming protocol.

animation

A class method, a convenient way of creating an animation object.

delegate

The delegate messages are `animationDidStart:` and `animationDidStop:finished:`, which should sound familiar from the analogous UIView animation

delegate messages. A CAAnimation instance *retains its delegate*; this is very unusual behavior and can cause trouble if you’re not conscious of it (I’m speaking from experience).

Alternatively, don’t set a delegate; use the CATransaction class method `setCompletionBlock:` instead, to run code after the animation ends. Besides avoiding the delegate retention problem, the code structure may be clearer.

`duration, timingFunction`

The length of the animation, and its timing function (a CAMediaTimingFunction). A duration of 0 (the default) means .25 seconds unless overridden by the transaction.

`autoreverses, repeatCount, repeatDuration, cumulative`

The first two are familiar from UIView animation. The `repeatDuration` property is a different way to govern repetition, specifying how long the repetition should continue rather than how many repetitions should occur; don’t specify both a `repeatCount` and a `repeatDuration`. If `cumulative` is YES, a repeating animation starts each repetition where the previous repetition ended (rather than jumping back to the start value).

`beginTime`

The delay before the animation starts. To delay an animation with respect to now, call `CACurrentMediaTime` and add the desired delay in seconds. The delay does not eat into the animation’s duration.

`timeOffset`

A shift in the animation’s overall timing; looked at another way, specifies the starting frame of the “animation movie,” which is treated as a loop. For example, an animation with a duration of 8 and a time offset of 4 plays its second half followed by its first half.

CAAnimation, along with all its subclasses, implements KVC to allow you set and retrieve a value for an arbitrary key, similar to CALayer ([Chapter 16](#)) and CATransaction.

CAPropertyAnimation

CAPropertyAnimation is a subclass of CAAnimation. It too is abstract, and adds the following:

`keyPath`

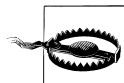
The all-important string specifying the CALayer key that is to be animated. Recall from [Chapter 16](#) that CALayer properties are accessible through KVC keys; now we are using those keys! A CAPropertyAnimation convenience class method `animationWithKeyPath:` creates the instance and assigns it a `keyPath`.

`additive`

If YES, the values supplied by the animation are added to the current presentation-layer value.

valueFunction

Converts a simple scalar value that you supply into a transform.



There is no animatable CALayer key called @"frame" (because `frame` is not an animatable property).

CABasicAnimation

`CABasicAnimation` is a subclass (not abstract!) of `CAPropertyAnimation`. It adds the following:

fromValue, toValue

The starting and ending values for the animation. These values must be objects, so numbers and structs will have to be wrapped accordingly, using `NSNumber` and `NSValue`. If neither `fromValue` nor `toValue` is provided, the former and current values of the property are used. If just one of `fromValue` or `toValue` is provided, the other uses the current value of the property.

byValue

Expresses one of the endpoint values as a *difference* from the other rather than in absolute terms. So you would supply a `byValue` instead of a `fromValue` or instead of a `toValue`, and the actual `fromValue` or `toValue` would be calculated for you by subtraction or addition with respect to the other value. If you supply *only* a `byValue`, the `fromValue` is the property's current value.

Using a `CABasicAnimation`

Having constructed and configured a `CABasicAnimation`, the way you order it to be performed is to *add it to a layer*. This is done with the `CALayer` instance method `addAnimation:forKey:`. (I'll discuss the purpose of the `forKey:` parameter later; it's fine to ignore it and use `nil`, as I do in the examples that follow.)

However, there's a slight twist. A `CAAnimation` is *merely* an animation; all it does is describe the hoops that the presentation layer is to jump through, the "animation movie" that is to be presented. It has no effect on the layer *itself*. Thus, if you naively create a `CABasicAnimation` and add it to a layer with `addAnimation:forKey:`, the animation happens and then the "animation movie" is whipped away to reveal the layer sitting there in exactly the same state as before. It is up to *you* to change the layer to match what the animation will ultimately portray.

This requirement may seem odd, but keep in mind that we are now in a much more fundamental, flexible world than the automatic, convenient worlds of view animation and implicit layer animation. Using explicit animation is more work, but you get more power. The converse, as we shall see, is that you *don't* have to change the layer if it *doesn't* change as a result of the animation.

To assure good results, we'll start by taking a plodding, formulaic approach to the use of CABasicAnimation, like this:

1. Capture the start and end values for the layer property you're going to change, because you're likely to need these values in what follows.
2. Change the layer property to its end value, first calling `setDisableActions`: to prevent implicit animation.
3. Construct the explicit animation, using the start and end values you captured earlier, and with its `keyPath` corresponding to the layer property you just changed.
4. Add the explicit animation to the layer.

Here's how you'd use this approach to animate our compass arrow rotation:

```
CompassLayer* c = (CompassLayer*)self.compass.layer;
// capture the start and end values
CATransform3D startValue = c.arrow.transform;
CATransform3D endValue = CATransform3DRotate(startValue, M_PI/4.0, 0, 0, 1);
// change the layer, without implicit animation
[CATransaction setDisableActions:YES];
c.arrow.transform = endValue;
// construct the explicit animation
CABasicAnimation* anim = [CABasicAnimation animationWithKeyPath:@"transform"];
anim.duration = 0.8;
CAMediaTimingFunction* clunk =
    [CAMediaTimingFunction functionWithControlPoints:.9 :.1 :.7 :.9];
anim.timingFunction = clunk;
anim.fromValue = [NSValue valueWithCATransform3D:startValue];
anim.toValue = [NSValue valueWithCATransform3D:endValue];
// ask for the explicit animation
[c.arrow addAnimation:anim forKey:nil];
```

Once you know the full form, you will find that in many cases it can be condensed. For example, when `fromValue` and `toValue` are not set, the former and current values of the property are used automatically. (This magic is possible because the presentation layer still has the former value of the property, while the layer itself has the new value.) Thus, in this case there was no need to set them, and so there was no need to capture the start and end values beforehand either. Here's the condensed version:

```
CompassLayer* c = (CompassLayer*)self.compass.layer;
[CATransaction setDisableActions:YES];
c.arrow.transform = CATransform3DRotate(c.arrow.transform, M_PI/4.0, 0, 0, 1);
CABasicAnimation* anim = [CABasicAnimation animationWithKeyPath:@"transform"];
anim.duration = 0.8;
CAMediaTimingFunction* clunk =
    [CAMediaTimingFunction functionWithControlPoints:.9 :.1 :.7 :.9];
anim.timingFunction = clunk;
[c.arrow addAnimation:anim forKey:nil];
```

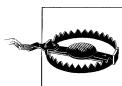
As I mentioned earlier, you will omit changing the layer if it doesn't change as a result of the animation. For example, let's make the compass arrow appear to vibrate rapidly, without ultimately changing its current orientation. To do this, we'll waggle it back

and forth, using a repeated animation, between slightly clockwise from its current position and slightly counterclockwise from its current position. The “animation movie” neither starts nor stops at the current position of the arrow, but for this animation it doesn’t matter, because it all happens so quickly as to appear perfectly natural:

```
CompassLayer* c = (CompassLayer*)self.compass.layer;
// capture the start and end values
CATransform3D nowValue = c.arrow.transform;
CATransform3D startValue = CATransform3DRotate(nowValue, M_PI/40.0, 0, 0, 1);
CATransform3D endValue = CATransform3DRotate(nowValue, -M_PI/40.0, 0, 0, 1);
// construct the explicit animation
CABasicAnimation* anim = [CABasicAnimation animationWithKeyPath:@"transform"];
anim.duration = 0.05;
anim.timingFunction =
    [CAMediaTimingFunction functionWithName:kCAMediaTimingFunctionLinear];
anim.repeatCount = 3;
anim.autoreverses = YES;
anim.fromValue = [NSValue valueWithCATransform3D:startValue];
anim.toValue = [NSValue valueWithCATransform3D:endValue];
// ask for the explicit animation
[c.arrow addAnimation:anim forKey:nil];
```

That code, too, can be shortened considerably from its full form. We can eliminate the need to calculate the new rotation values based on the arrow’s current transform by setting our animation’s `additive` property to YES; this means that the animation’s property values are added to the existing property value for us, so that they are relative, not absolute. For a transform, “added” means “matrix-multiplied,” so we can describe the waggle without any dependence on the arrow’s current rotation. Moreover, because our rotation is so simple (around a cardinal axis), we can take advantage of `CAPropertyAnimation`’s `valueFunction`; the animation’s property values can then be simple scalars (in this case, angles), because the `valueFunction` tells the animation to interpret these as rotations around the z-axis:

```
CompassLayer* c = (CompassLayer*)self.compass.layer;
CABasicAnimation* anim = [CABasicAnimation animationWithKeyPath:@"transform"];
anim.duration = 0.05;
anim.timingFunction =
    [CAMediaTimingFunction functionWithName:kCAMediaTimingFunctionLinear];
anim.repeatCount = 3;
anim.autoreverses = YES;
anim.additive = YES;
anim.valueFunction = [CAValueFunction functionWithName:kCAValueFunctionRotateZ];
anim.fromValue = [NSNumber numberWithFloat:M_PI/40];
anim.toValue = [NSNumber numberWithFloat:-M_PI/40];
[c.arrow addAnimation:anim forKey:nil];
```



Instead of using a `valueFunction`, we could have achieved the same effect by setting the animation’s key path to `@"transform.rotation.z"`. However, Apple advises against this, as it can result in mathematical trouble when there is more than one rotation.

Remember that there is no @"frame" key. To animate a layer's frame, if both its **position** and **bounds** are to change, you must animate both. Recall this earlier example, using block-based animation (where **v** is **v2**'s superview):

```
void (^anim) (void) = ^{
    CGRect f = v2.frame;
    f.size.width = v.frame.size.width;
    f.origin.x = 0;
    v2.frame = f;
};
```

Here's how to do that with Core Animation:

```
CABasicAnimation* anim1 = [CABasicAnimation animationWithKeyPath:@"bounds"];
CGRect f = v2.layer.bounds;
f.size.width = v.layer.bounds.size.width;
v2.layer.bounds = f;
[v2.layer addAnimation:anim1 forKey: nil];
CABasicAnimation* anim2 = [CABasicAnimation animationWithKeyPath:@"position"];
CGPoint p = v2.layer.position;
p.x = CGRectGetMidX(v.layer.bounds);
v2.layer.position = p;
[v2.layer addAnimation:anim2 forKey: nil];
```

Keyframe Animation

Keyframe animation (**CAKeyframeAnimation**) is an alternative to basic animation (**CABasicAnimation**); they are both subclasses of **CAPropertyAnimation** and they are used in identical ways. The difference is that a keyframe animation, in addition to specifying a starting and ending value, also specifies multiple values through which the animation should pass on the way, the stages (*frames*) of the animation. This can be as simple as setting the animation's **values** property (an **NSArray**).

Here's a more sophisticated version of our animation for wagging the compass arrow: the animation includes both the start and end states, and the degree of waggle gets progressively smaller:

```
CompassLayer* c = (CompassLayer*)self.compass.layer;
NSMutableArray* values = [NSMutableArray array];
[values addObject: [NSNumber numberWithFloat:0]];
int direction = 1;
for (int i = 20; i < 60; i += 5, direction *= -1) { // reverse direction each time
    [values addObject: [NSNumber numberWithFloat: direction*M_PI/(float)i]];
}
[values addObject: [NSNumber numberWithFloat:0]];
CAKeyframeAnimation* anim = [CAKeyframeAnimation animationWithKeyPath:@"transform"];
anim.values = values;
anim.additive = YES;
anim.valueFunction = [CAValueFunction functionWithName: kCAValueFunctionRotateZ];
[c.arrow addAnimation:anim forKey:nil];
```

Here are some **CAKeyframeAnimation** properties:

values

The array of values the animation is to adopt, including the starting and ending value.

timingFunctions

An array of timing functions, one for each stage of the animation (so that this array will be one element shorter than the **values** array).

keyTimes

An array of times to accompany the array of values, defining when each value should be reached. The times start at 0 and are expressed as increasing fractions of 1, ending at 1.

calculationMode

Describes how the **values** are treated to create *all* the values through which the animation must pass.

- The default is **kCAAnimationLinear**, a simple straight-line interpolation from value to value.
- **kCAAnimationCubic** constructs a single smooth curve passing through all the values (and additional advanced properties, **tensionValues**, **continuityValues**, and **biasValues**, allow you to refine the curve).
- **kCAAnimationPaced** and **kCAAnimationCubicPaced** means the timing functions and key times are ignored, and the velocity is made constant through the whole animation.
- **kCAAnimationDiscrete** means no interpolation: we jump directly to each value at the corresponding key time.

path

When you’re animating a property whose values are pairs of floats (CGPoints), this is an alternative way of describing the values; instead of a **values** array, which must be interpolated to arrive at the intermediate values along the way, you supply the entire interpolation as a single CGPathRef. The points used to draw the path are the keyframe values, so you can still apply timing functions and key times. If you’re animating a position, the **rotationMode** property lets you ask the animated object to rotate so as to remain perpendicular to the path.

Making a Property Animatable

So far, we’ve been animating built-in animatable properties. If you define your own property on a CALayer subclass, you can make that property animatable through a CAPropertyAnimation (a CABasicAnimation or a CAKeyframeAnimation). You do this by declaring the property **@dynamic** (so that Core Animation can create its accessors) and returning YES from the class method **needsDisplayForKey:**, where the key is the string name of the property.

For example, here we'll start writing the code for a layer class `MyLayer` with an animatable `thickness` property:

```
// the interface section
@interface MyLayer : CALayer
@property (nonatomic, assign) CGFloat thickness;
@end

// the implementation section
@implementation MyLayer
@dynamic thickness;

+ (BOOL) needsDisplayForKey:(NSString *)key {
    if ([key isEqualToString: @"thickness"])
        return YES;
    return [super needsDisplayForKey:key];
}

@end
```

Returning YES from `needsDisplayForKey:` causes this layer to be redisplayed repeatedly as the `thickness` property changes. So if we want to *see* the animation, this layer also needs to draw itself in some way that depends on the `thickness` property. Here, I'll implement the layer's `drawInContext:` to make `thickness` the thickness of the black border around a red rectangle:

```
- (void) drawInContext:(CGContextRef)con {
    CGRect r = CGRectMakeInset(self.bounds, 20, 20);
    CGContextSetFillColorWithColor(con, [UIColor redColor].CGColor);
    CGContextFillRect(con, r);
    CGContextSetLineWidth(con, self.thickness);
    CGContextStrokeRect(con, r);
}
```

Now we can animate the rectangle's thickness using explicit animation (`lay` points to a `MyLayer` instance):

```
CABasicAnimation* ba = [CABasicAnimation animationWithKeyPath:@"thickness"];
ba.toValue = [NSNumber numberWithFloat: 10.0];
ba.autoreverses = YES;
[lay addAnimation:ba forKey:nil];
```

At every step of the animation, `drawInContext:` is called, and because the `thickness` value differs at each step, it is animated.

Grouped Animations

A grouped animation (`CAAnimationGroup`) combines multiple animations into one, by means of its `animations` property (an `NSArray` of animations). By delaying and timing the various component animations, complex effects can be created.

A `CAAnimationGroup` is itself an animation; it is a `CAAnimation` subclass, so it has a `duration` and other animation features. Think of the `CAAnimationGroup` as the parent

and its `animations` as its children. Then *the children inherit default values from their parent*. Thus, for example, if you don't set a child's duration explicitly, it will inherit the parent's duration. Also, make sure the parent's duration is sufficient to include all parts of the child animations that you want displayed.

For example, we can form a sequence where the compass arrow rotates and then wiggles. Very little change is required. We express the first animation in its full form, with explicit `fromValue` and `toValue`. We postpone the second animation using its `beginTime` property; notice that we express this in relative terms, as a number of seconds into the parent's duration, not with respect to `CACurrentMediaTime`. Finally, we set the overall parent duration to the sum of the child durations, so that it can embrace both of them:

```
CompassLayer* c = (CompassLayer*)self.compass.layer;
// capture current value, set final value
CGFloat rot = M_PI/4.0;
[CATransaction setDisableActions:YES];
CGFloat current = [[c.arrow valueForKeyPath:@"transform.rotation.z"] floatValue];
[c.arrow setValue: [NSNumber numberWithFloat: current + rot]
    forKeyPath:@"transform.rotation.z"];

// first animation (rotate and clunk)
CABasicAnimation* anim1 = [CABasicAnimation animationWithKeyPath:@"transform"];
anim1.duration = 0.8;
CAMediaTimingFunction* clunk =
    [CAMediaTimingFunction functionWithControlPoints:.9 :.1 :.7 :.9];
anim1.timingFunction = clunk;
anim1.fromValue = [NSNumber numberWithFloat: current];
anim1.toValue = [NSNumber numberWithFloat: current + rot];
anim1.valueFunction = [CAValueFunction functionWithName:kCAValueFunctionRotateZ];

// second animation (waggle)
NSMutableArray* values = [NSMutableArray array];
[values addObject: [NSNumber numberWithFloat:0]];
int direction = 1;
for (int i = 20; i < 60; i += 5, direction *= -1) { // reverse direction each time
    [values addObject: [NSNumber numberWithFloat: direction*M_PI/(float)i]];
}
[values addObject: [NSNumber numberWithFloat:0]];
CAKeyframeAnimation* anim2 =
    [CAKeyframeAnimation animationWithKeyPath:@"transform"];
anim2.values = values;
anim2.duration = 0.25;
anim2.beginTime = anim1.duration;
anim2.additive = YES;
anim2.valueFunction = [CAValueFunction functionWithName:kCAValueFunctionRotateZ];

// group
CAAnimationGroup* group = [CAAnimationGroup animation];
group.animations = [NSArray arrayWithObjects: anim1, anim2, nil];
group.duration = anim1.duration + anim2.duration;
[c.arrow addAnimation:group forKey:nil];
```

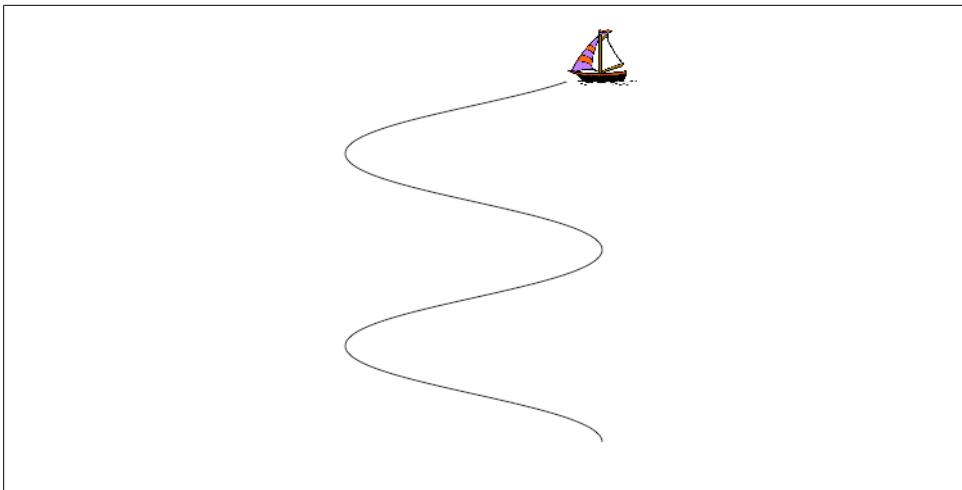


Figure 17-2. A boat and the course she'll sail

In that example, I grouped two animations that animated the same property sequentially. Now let's go to the other extreme and group some animations that animate different properties simultaneously. I have a small view (about 56×38), located near the top right corner of the screen, whose layer contents are a picture of a sailboat facing to the left. I'll "sail" the boat in a curving path, both down the screen and left and right across the screen, like an extended letter "S" (Figure 17-2). Each time the boat comes to a vertex of the curve, changing direction across the screen, I'll turn the boat picture so that it faces the way it's about to move. At the same time, I'll constantly rock the boat, so that it always appears to be pitching a little on the waves.

Here's the first animation, the movement of the boat along its curving path. It illustrates the use of a CAKeyframeAnimation with a CGPath; the `calculationMode` of `kCAAnimationPaced` ensures an even speed over the whole path. We don't set an explicit duration because we want to adopt the duration of the group:

```
CGFloat h = 200;
CGFloat v = 75;
CGMutablePathRef path = CGPathCreateMutable();
int leftright = 1;
CGPoint next = self.view.layer.position;
CGPoint pos;
CGPathMoveToPoint(path, NULL, next.x, next.y);
for (int i = 0; i < 4; i++) {
    pos = next;
    leftright *= -1;
    next = CGPointMake(pos.x+h*leftright, pos.y+v);
    CGPathAddCurveToPoint(path, NULL, pos.x, pos.y+30, next.x, next.y-30,
                          next.x, next.y);
}
CAKeyframeAnimation* anim1 = [CAKeyframeAnimation animationWithKeyPath:@"position"];
anim1.path = path;
anim1.calculationMode = kCAAnimationPaced;
```

Here's the second animation, the reversal of the direction the boat is facing. This is simply a rotation around the y-axis. We make no attempt at visually animating this reversal, so we set the `calculationMode` to `kCAAnimationDiscrete` (the boat image reversal is a sudden change). There is one fewer value than the number of points in our first animation's path, and the first animation has an even speed, so the reversals take place at each curve apex with no further effort on our part. (If the pacing were more complicated, we could give both the first and the second animation identical `keyTimes` arrays, to coordinate them.) Once again, we don't set an explicit duration:

```
NSArray* revs = [NSArray arrayWithObjects:
    [NSNumber numberWithFloat:0],
    [NSNumber numberWithFloat:M_PI],
    [NSNumber numberWithFloat:0],
    [NSNumber numberWithFloat:M_PI],
    nil];
CAKeyframeAnimation* anim2 =
    [CAKeyframeAnimation animationWithKeyPath:@"transform"];
anim2.values = revs;
anim2.valueFunction = [CAValueFunction functionWithName:kCAValueFunctionRotateY];
anim2.calculationMode = kCAAnimationDiscrete;
```

Here's the third animation, the rocking of the boat. It has a short duration, and repeats indefinitely (by giving its `repeatCount` an immense value):

```
NSArray* pitches = [NSArray arrayWithObjects:
    [NSNumber numberWithFloat:0],
    [NSNumber numberWithFloat:M_PI/60.0],
    [NSNumber numberWithFloat:0],
    [NSNumber numberWithFloat:-M_PI/60.0],
    [NSNumber numberWithFloat:0],
    nil];
CAKeyframeAnimation* anim3 =
    [CAKeyframeAnimation animationWithKeyPath:@"transform"];
anim3.values = pitches;
anim3.repeatCount = HUGE_VALF;
anim3.duration = 0.5;
anim3.additive = YES;
anim3.valueFunction = [CAValueFunction functionWithName:kCAValueFunctionRotateZ];
```

Finally, we combine the three animations, assigning the group an explicit duration that will be adopted by the first two animations. As we hand the animation over to the layer displaying the boat, we also change the layer's position to match the final position from the first animation, so that the boat won't jump back to its original position afterward:

```
CAAnimationGroup* group = [CAAnimationGroup animation];
group.animations = [NSArray arrayWithObjects: anim1, anim2, anim3, nil];
group.duration = 8;
[view.layer addAnimation:group forKey:nil];
[CATransaction setDisableActions:YES];
view.layer.position = next;
```

Here are some further CAAnimation properties (from the CAMediaTiming protocol) that come into play especially when animations are grouped:

speed

The ratio between a child’s timescale and the parent’s timescale. For example, if a parent and child have the same duration, but the child’s **speed** is **1.5**, its animation runs one-and-a-half times as fast as the parent.

fillMode

Suppose the child animation begins after the parent animation, or ends before the parent animation, or both. What should happen to the appearance of the property being animated, outside the child animation’s boundaries? The answer depends on the child’s **fillMode**:

- **kCAFillModeRemoved** means the child animation is removed, revealing the layer property at its actual current value whenever the child is not running.
- **kCAFillModeForwards** means the final presentation layer value of the child animation remains afterward.
- **kCAFillModeBackwards** means the initial presentation layer value of the child animation appears right from the start.
- **kCAFillModeBoth** combines the previous two.



CALayer adopts the CAMediaTiming protocol, in the sense that a layer can have a **speed**. This will affect any animation attached to it. A CALayer with a speed of **2** will play a 10-second animation in 5 seconds.

Transitions

A layer transition is an animation involving two “copies” of a single layer, in which the second “copy” appears to replace the first. It is described by an instance of CATransition (a CAAnimation subclass), which has these chief properties describing the animation:

type

Your choices are:

- **kCATransitionFade**
- **kCATransitionMoveIn**
- **kCATransitionPush**
- **kCATransitionReveal**

subtype

If the type is not **kCATransitionFade**, your choices are:

- **kCATransitionFromRight**
- **kCATransitionFromLeft**
- **kCATransitionFromTop**

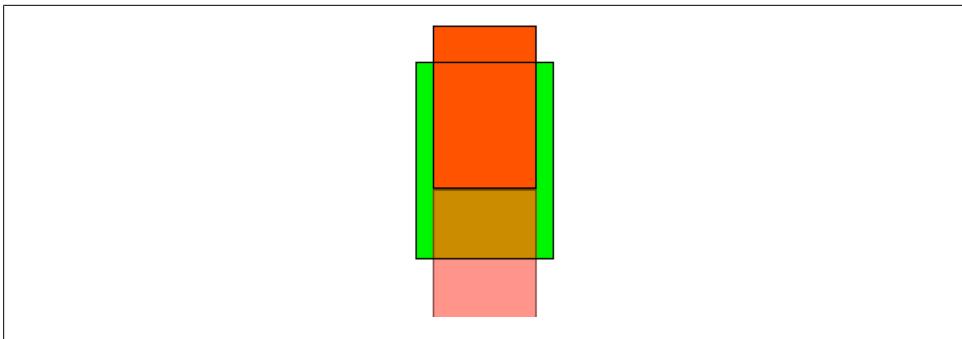


Figure 17-3. A push transition

- `kCATransitionFromBottom`



For historical reasons, the terms “bottom” and “top” in the names of the `subtype` settings have the opposite of their expected meanings.

To understand the nature of a transition animation, the best approach is to try one, without doing anything else. For example:

```
CATransition* t = [CATransition animation];
t.type = kCATransitionPush;
t.subtype = kCATransitionFromBottom;
[layer addAnimation: t forKey: nil];
```

It will help if the layer’s frame is visible (give it a `borderWidth`, perhaps). What you’ll see, then, is that the entire layer exits moving down from its original place, and another “copy” of the same layer enters moving down from above. In [Figure 17-3](#), the green layer (the wider rectangle) is the superlayer of the red layer (the narrower rectangle, which appears twice). The red layer is normally centered in the green layer, but I’ve managed to freeze the red layer in the middle of a transition.

You can use a layer’s superlayer to help restrict the visible part of the layer’s transition. If the superlayer’s `masksToBounds` is NO, the user can see the entire transition; its movements will have the whole screen as their visible boundaries. But if the superlayer’s `masksToBounds` is YES, then the visible part of the transition movement is restricted to the superlayer’s bounds: it’s as if you’re seeing the movements through a window that is the superlayer. In [Figure 17-3](#), for example, if the green layer’s `masksToBounds` were YES, we wouldn’t see any of the part of the transition animation outside its boundaries. A common device is to have the layer that is to be transitioned live inside a superlayer that is exactly the same size and whose `masksToBounds` is YES. This confines the visible transition to the bounds of the layer itself.

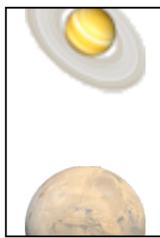


Figure 17-4. Another push transition

Our example appears silly, because there was no motivation for this animation; the two “copies” of the layer are identical. A typical motivation would be that you’re changing the contents of a layer and you want to dramatize this. Here, we change the example so that an image of Saturn replaces an image of Mars by pushing it away from above (Figure 17-4). We get a slide effect, as if one layer were being replaced by another; but in fact there is just one layer that holds first one picture, then the other:

```
CATransition* t = [CATransition animation];
t.type = kCATTransitionPush;
t.subtype = kCATransitionFromBottom;
[CATransaction setDisableActions:YES];
layer.contents = (id)[[UIImage imageNamed: @"Saturn.gif"] CGImage];
[layer addAnimation: t forKey: nil];
```

A transition on a superlayer can happen simultaneously with animation of a sublayer. The animation will be seen to occur on the second “copy” of the layer as it moves into position. This is analogous to what we achieved earlier with the `UIViewAnimationOption-AllowAnimatedContent` option using block-based view animation.

The Animations List

The method that asks for an explicit animation to happen is CALayer’s `addAnimation:forKey:`. To understand how this method actually works (and what the “key” is), you need to know about a layer’s *animations list*.

An animation is an object (a CAAnimation) that modifies how a layer is drawn. It does this merely by being attached to the layer; the layer’s drawing mechanism does the rest. A layer maintains a list of animations that are currently in force. To add an animation to this list, you call `addAnimation:forKey:`. When the time comes to draw itself, the layer looks through its animations list and draws itself in accordance with any animations it finds there. (The list of things the layer must do in order to draw itself is sometimes referred to by the documentation as the *render tree*.)

The animations list is maintained in a curious way. The list is not exactly a dictionary, but it behaves somewhat like a dictionary. An animation has a key — the `forKey:` parameter in `addAnimation:forKey:`. If an animation with a certain key is added to the list, and an animation with that key is already in the list, the one that is already in the

list is removed. Thus a rule is maintained that *only one animation with a given key* can be in the list at a time (the *exclusivity rule*). This explains why sometimes ordering an animation can cancel an animation already ordered or in-flight: the two animations had the same key, so the first one was removed. It is also possible to add an animation with *no key* (the key is nil); it is then *not* subject to the exclusivity rule (that is, there can be more than one animation in the list with no key). The order in which animations were added to the list is the order in which they are applied.

The `forKey:` parameter in `addAnimation:forKey:` is thus *not a property name*. It *could* be a property name, but it can be any arbitrary value. Its purpose is to enforce the exclusivity rule. It does *not* have any meaning with regard to what property a `CAPrropertyAnimation` animates; that is the job of the animation's `keyPath`. (Apple's use of the term "key" in `addAnimation:forKey:` is thus unfortunate and misleading; I wish they had named this method `addAnimation:withIdentifier:` or something like that.)



Actually, there *is* a relationship between the "key" in `addAnimation:forKey:` and a `CAPrropertyAnimation`'s `keyPath`: if a `CAPrropertyAnimation`'s `keyPath` is nil at the time that it is added to a layer with `addAnimation:forKey:`, *that keyPath is set to the call's forKey: parameter value*. Thus, you can *misuse* the `forKey:` parameter in `addAnimation:forKey:` as a way of specifying what `keyPath` an animation animates. (This fact is not documented, so far as I know, but it's easily verified experimentally, and it should remain reliably true, as implicit animation crucially depends on it.) I have seen many prominent but misleading examples that use this technique, apparently in the mistaken belief that the "key" in `addAnimation:forKey:` is the way you are *supposed* to specify what property to animate. *This is wrong*. Set the `CAPrropertyAnimation`'s `keyPath` explicitly (as do all my examples); that's what it's for.

You can use the exclusivity rule to your own advantage, to keep your code from stepping on its own feet. Some code of yours might add an animation to the list using a certain key; then later, some other code might come along and correct this, removing that animation and replacing it with another. By using the same key, the second code is easily able to override the first: "You may have been given some other animation with this key, but throw it away; play this one instead."

In some cases, the key you supply is ignored and a different key is substituted. In particular, the key with which a `CATransition` is added to the list is always `kCATransition` (which happens to be @"`transition`"); thus there can be only one transition animation in the list.

You can't access the entire animations list directly. You can access the key names of the animations in the list, with `animationKeys`; and you can obtain or remove an animation with a certain key, with `animationForKey:` and `removeAnimationForKey:`; but animations with a nil key are inaccessible. You can, however, remove all animations, including animations with a nil key, using `removeAllAnimations`.



In the multitasking world, when the app is suspended ([Chapter 11](#)), `removeAllAnimations` is called on all layers for you.

You can think of an animation in a layer's animations list as being the “animation movie” I spoke of at the start of this chapter. As long as an animation is in the list, the movie is present, either waiting to be played or actually playing. An animation that has finished playing is, in general, pointless; the animation should now be removed from the list. Therefore, an animation has a `removedOnCompletion` property, which defaults to YES: when the “movie” is over, the animation removes itself from the list.

You can, if desired, set `removedOnCompletion` to NO. However, even the presence in the list of an animation that has already played might make no difference to the layer's appearance, because an animation's `fillMode` is `kCAFillModeRemoved`, which removes the animation from the layer's drawing when the movie is over. Thus, it can usually do no harm to leave an animation in the list after it has played, but it's not a great idea either, because this is just one more thing for the drawing system to worry about. Typically, you'll leave `removedOnCompletion` set at YES.



You may encounter examples that set `removedOnCompletion` to NO and set the animation's `fillMode` to `kCAFillModeForwards` or `kCAFillModeBoth`, as a way of causing the layer to keep the appearance of the last frame of the “animation movie” even after the animation is over, and preventing a property from apparently jumping back to its initial value when the animation ends. *This is wrong*. The correct approach, as I have explained, is to change the property value to match the final frame of the animation. The chief use of `kCAFillModeForwards` is in connection with grouped animations, as explained earlier.

Actions

For the sake of completeness, I will now explain how implicit animation works — that is, how implicit animation is turned into explicit animation behind the scenes. The basis of implicit animation is the *action mechanism*.

What an Action Is

An *action* is an object that adopts the `CAAction` protocol. This means simply that it implements `runActionForKey:object:arguments:`.

The action object could do *anything* in response to this message. The notion of an action is completely general. However, in real life, the only class that adopts the `CAAction` protocol is `CAAnimation`. So, an animation is a special case of an action, but in fact it is also the *only* case of an action.

What an animation does when it receives `runActionForKey:object:arguments:` is to assume that the second parameter, the `object:`, is a layer, and to add itself to that layer's animations list. Thus, for an animation, receiving the `runActionForKey:object:arguments:` message is like being told: "Play yourself!"

You would never send `runActionForKey:object:arguments:` to an animation directly. Rather, this message is sent to an animation for you, as the basis of implicit animation.

The Action Search

When you set a property of a layer and trigger an implicit animation, you are actually triggering the *action search*. This basically means that the layer searches for an *action object* to which it can send the `runActionForKey:object:arguments:` message; because that action object will be an animation, and because it will respond to this message by adding itself to the layer's animations list, this is the same as saying that the layer searches for an animation to play itself with respect to the layer. The procedure by which the layer searches for this animation is quite elaborate.

The search for an action object begins because you do something that causes the layer to be sent the `actionForKey:` message. Let us presume that what you do is to change the value of an animatable property. (Other things can cause the `actionForKey:` message to be sent, as I'll show later.) The action mechanism then treats the name of the property as a key, and the layer receives `actionForKey:` with that key — and the action search begins.

At each stage of the action search, the following rules are obeyed regarding what is returned from that stage of the search:

An action object

If an action object (an animation) is produced, that is the end of the search. The action mechanism sends that animation the `runActionForKey:object:arguments:` message; the animation responds by adding itself to the layer's animations list.

[`NSNull null`]

If [`NSNull null`] is produced, this means, "Do nothing and stop searching." The search comes to an end; there will be no implicit animation.

`nil`

If nil is produced, the search continues to the next stage.

The action search proceeds by stages, as follows:

1. The layer might terminate the search before it even starts. For example, the layer will do this if it is the underlying layer of a view, or if a property is set to the same value it already has. In such a case, there should be no implicit animation, so the whole mechanism is nipped in the bud.

2. If the layer has a delegate that implements `actionForLayer:forKey:`, that message is sent to the delegate, with this layer as the layer and the property name as the key. If an animation or `[NSNull null]` is returned, the search ends.
3. The layer has a property called `actions`, which is a dictionary. If there is an entry in this dictionary with the given key, that value is used, and the search ends.
4. The layer has a property called `style`, which is a dictionary. If there is an entry in this dictionary with the key `actions`, it is assumed to be a dictionary; if this `actions` dictionary has an entry with the given key, that value is used, and the search ends. Otherwise, if there is an entry in the `style` dictionary called `style`, the same search is performed within it, and so on recursively until either an `actions` entry with the given key is found (the search ends) or there are no more `style` entries (the search continues).



If the `style` dictionary sounds profoundly weird, that's because it is profoundly weird. It is actually a special case of a larger, separate mechanism, which is also profoundly weird, having to do not with `actions`, but with a CALayer's implementation of KVC. When you call `valueForKey:` on a layer, if the key is undefined by the layer itself, the `style` dictionary is consulted. I have never written or seen code that uses this mechanism for anything, and I'll say no more about it.

5. The layer's class is sent `defaultActionForKey:`, with the property name as the key. If an animation or `[NSNull null]` is returned, the search ends.
6. If the search reaches this last stage, a default animation is supplied, as appropriate. For a property animation, this is a plain vanilla CABasicAnimation.

Both the delegate's `actionForLayer:forKey:` and the subclass's `defaultActionForKey:` are declared as returning an `id<CAAction>`. To return `[NSNull null]`, therefore, you'll need to typecast it to `id<CAAction>` to quiet the compiler; you're lying (`NSNull` does not adopt the `CAAction` protocol), but it doesn't matter.

Hooking Into the Action Search

You can affect the action search at various stages to modify what happens when the search is triggered. For example, you could cause some stage of the search to produce an animation; that animation will then be used. Assuming that the search is triggered by setting an animatable layer property, you would then be affecting how implicit animation behaves.

You will probably want your animation to be fairly minimal. You may have no way of knowing the former and current values of the property that is being changed, so it would then be pointless (and very strange) to set a CABasicAnimation's `fromValue` or `toValue`. Moreover, although animation properties that you don't set can be set through

CATransaction, in the usual manner for implicit property animation, animation properties that you *do* set can *not* be overridden through CATransaction. For example, if you set the duration of the animation that you produce at some stage of the action search, a call to CATransaction’s `setAnimationDuration:` class method cannot change it.

Let’s say we want a certain layer’s duration for an implicit `position` animation to be 5 seconds. We can achieve this with a minimally configured animation, like this:

```
CABasicAnimation* ba = [CABasicAnimation animation];
ba.duration = 5;
```

The idea now is to situate this animation, `ba`, where it will be produced by the action search when implicit animation is triggered on the `position` property of our layer. We could, for instance, put it into the layer’s `actions` dictionary:

```
layer.actions = [NSDictionary dictionaryWithObject: ba forKey: @"position"];
```

The result is that when we set that layer’s `position`, if an implicit animation results, its duration is 5 seconds, even if we try to change it through CATransaction:

```
[CATransaction setAnimationDuration:1];
layer.position = CGPointMake(100,200); // animation takes 5 seconds
```

Let’s use that example to tease apart how the action mechanism makes implicit animation work:

1. You set the value of the layer’s `position` property.
2. If your setting does not represent a change in the `position` value, or if this layer is a view’s underlying layer, that’s the end of the story; there is no implicit property animation.
3. Otherwise, the action search begins. There is no delegate in this case, so the search proceeds to the next stage, the `actions` dictionary.
4. There is an entry under the key `@"position"` in the `actions` dictionary (because we put it there), and it is an animation. That animation is the action, and that is the end of the search.
5. The animation is sent `runActionForKey:object:arguments:`.
6. The animation responds by calling `[object addAnimation:self forKey: @"position"]`. The animation’s `keyPath` was `nil`, *so this call also sets the keyPath to the same key!* Thus, there is now an animation in the layer’s `animations` list that animates its `position`, because its `keyPath` is `@"position"`. Moreover, we didn’t set the `fromValue` or `toValue`, so the property’s previous and new values are used. The animation therefore shows the layer moving from its current position to `{100,200}`.

Using the layer’s `actions` dictionary to set default animations is a somewhat inflexible way to hook into the action search, however. It’s a good way to disable implicit animation for specific properties; just set the value for that key to `[NSNull null]`. But it has the disadvantage in general that you must write your animation beforehand.

By contrast, if you set the layer's delegate to a instance that responds to `actionForLayer:forKey:`, your code runs at the time the animation is needed, and you have access to the layer that is to be animated. So you can create the animation on the fly, possibly modifying it in response to current circumstances.

Recall also that `CATransaction` implements KVC to allow you to set and retrieve the value of arbitrary keys. We can take advantage of this fact to pass an additional message from the code that sets the property value, and triggers the action search, to the code that supplies the action. This works because they both take place within the same transaction.

In this example, we use the layer delegate to change the default `position` animation so that instead of being a straight line, the path has a slight waggle. To do this, the delegate constructs a keyframe animation. The animation depends on the old `position` value and the new `position` value; the delegate can get the former direct from the layer, but the latter must be handed to the delegate somehow. Here, a `CATransaction` key @"newP" is used to communicate this information. When we set the layer's `position`, we put its future value where the delegate can retrieve it, like this:

```
CGPoint newP = CGPointMake(200,300);
[CATransaction setValue: [NSValue valueWithCGPoint: newP] forKey: @"newP"];
layer.position = newP; // the delegate will waggle the layer into place
```

The delegate is called by the action search and constructs the animation:

```
- (id < CAAction >)actionForLayer:(CALayer *)layer forKey:(NSString *)key {
    if ([key isEqualToString: @"position"]) {
        CGPoint oldP = layer.position;
        CGPoint newP = [[CATransaction valueForKey: @"newP"] CGPointValue];
        CGFloat d = sqrt(pow(oldP.x - newP.x, 2) + pow(oldP.y - newP.y, 2));
        CGFloat r = d/3.0;
        CGFloat theta = atan2(newP.y - oldP.y, newP.x - oldP.x);
        CGFloat wag = 10*M_PI/180.0;
        CGPoint p1 = CGPointMake(oldP.x + r*cos(theta+wag),
                               oldP.y + r*sin(theta+wag));
        CGPoint p2 = CGPointMake(oldP.x + r*2*cos(theta-wag),
                               oldP.y + r*2*sin(theta-wag));
        CAKeyframeAnimation* anim = [CAKeyframeAnimation animation];
        anim.values = [NSArray arrayWithObjects:
                      [NSValue valueWithCGPoint:oldP],
                      [NSValue valueWithCGPoint:p1],
                      [NSValue valueWithCGPoint:p2],
                      [NSValue valueWithCGPoint:newP],
                      nil];
        anim.calculationMode = kCAAnimationCubic;
        return anim;
    }
    return nil;
}
```

Finally, for the sake of completeness, I'll demonstrate overriding `defaultActionForKey:`. This code would go into a `CALayer` subclass where setting its `contents` is to trigger a push transition from the left:

```

+ (id < CAAction >)defaultActionForKey:(NSString *)aKey {
    if ([aKey isEqualToString:@"contents"]) {
        CATransition* tr = [CATransition animation];
        tr.type = kCATransitionPush;
        tr.subtype = kCATransitionFromLeft;
        return tr;
    }
    return [super defaultActionForKey: aKey];
}

```

Nonproperty Actions

Changing a property is not the only way to trigger a search for an action; an action search is also triggered when a layer is added to a superlayer (key `kCAOnOrderIn`) and when a layer's sublayers are changed by adding or removing a sublayer (key `@"sublayers"`). We can watch for these keys in the delegate and return an animation.



These triggers and their keys are incorrectly described in Apple's documentation, and there are additional triggers and keys that are not mentioned there.

In this example, we use our layer's delegate so that when our layer is added to a superlayer, it will "pop" into view. We implement this by fading the layer quickly in from an opacity of 0 and at the same time scaling the layer's transform to make it momentarily appear a little larger:

```

- (id < CAAction >)actionForLayer:(CALayer *)lay forKey:(NSString *)key {
    if ([key isEqualToString:kCAOnOrderIn]) {
        CABasicAnimation* anim1 =
            [CABasicAnimation animationWithKeyPath:@"opacity"];
        anim1.fromValue = [NSNumber numberWithFloat: 0.0];
        anim1.toValue = [NSNumber numberWithFloat: lay.opacity];
        CABasicAnimation* anim2 =
            [CABasicAnimation animationWithKeyPath:@"transform"];
        anim2.toValue = [NSValue valueWithCATransform3D:
            CATransform3DScale(lay.transform, 1.1, 1.1, 1.0)];
        anim2.autoreverses = YES;
        anim2.duration = 0.1;
        CAAimationGroup* group = [CAAnimationGroup animation];
        group.animations = [NSArray arrayWithObjects: anim1, anim2, nil];
        group.duration = 0.2;
        return group;
    }
}

```

The documentation says that when a layer is removed from a superlayer, an action is sought under the key `kCAOnOrderOut`. This is true but useless, because by the time the action is sought, the layer has already been removed from the superlayer, so returning an animation has no visible effect. Similarly, an animation returned as an action when a layer's `hidden` is set to YES is never played. Apple has admitted that this is a bug. A

possible workaround is to trigger the animation via the `opacity` property, perhaps in conjunction with a `CATransaction` key, and remove the layer afterward:

```
[CATransaction setCompletionBlock: ^{
    [layer removeFromSuperlayer];
}];

[CATransaction setValue:@"" forKey:@"byebye"];
layer.opacity = 0;
```

Now `actionForLayer:forKey:` can test for the incoming key @"`opacity`" and the `CATransaction` key @"`byebye`", and return the animation appropriate to removal from the superlayer. Here's a possible implementation:

```
if ([key isEqualToString:@"opacity"]) {
    if ([CATransaction valueForKey:@"byebye"]) {
        CABasicAnimation* anim1 =
        [CABasicAnimation animationWithKeyPath:@"opacity"];
        anim1.fromValue = [NSNumber numberWithFloat: layer.opacity];
        anim1.toValue = [NSNumber numberWithFloat: 0.0];
        CABasicAnimation* anim2 =
        [CABasicAnimation animationWithKeyPath:@"transform"];
        anim2.toValue = [NSValue valueWithCATransform3D:
                         CATransform3DScale(layer.transform, 0.1, 0.1, 1.0)];
        CAAnimationGroup* group = [CAAnimationGroup animation];
        group.animations = [NSArray arrayWithObjects: anim1, anim2, nil];
        group.duration = 0.2;
        return group;
    }
}
```

Emitter Layers

Emitter layers (`CAEmitterLayer`) are new to iOS 5; they've been around on the desktop since Mac OS X 10.6. To some extent, they are on a par with animated images: once you've set up an emitter layer, it just sits there animating all by itself. The nature of this animation is rather narrow: an emitter layer emits particles, which are `CAEmitterCell` instances. However, by clever setting of the properties of an emitter layer and its emitter cells, you can achieve some astonishing effects. Moreover, the animation is itself animatable using Core Animation.

It is easiest to understand emitter layers and emitter cells if you start with some stupid settings to achieve a boring effect. Let's start with the emitter cells. Here are some useful basic properties of a `CAEmitterCell`:

`contents`, `contentsRect`

These are modeled after the eponymous `CALayer` properties, although `CAEmitterLayer` is not a `CALayer` subclass; so, respectively, an image (a `CGImageRef`) and a `CGRect` defining a region of that image. They define the image that a cell will portray.

birthrate, lifetime

How many cells per second should be emitted, and how many seconds each cell should live before vanishing, respectively.

velocity

The speed at which a cell moves. The meaning of the actual value isn't documented; perhaps it's points per second.

emissionLatitude, emissionLongitude

The angle at which the cell is emitted from the emitter, as a variation from the perpendicular. Longitude is an angle within the plane; latitude is an angle out of the plane.

So, here's code to create a very elementary emitter cell:

```
UIGraphicsBeginImageContextWithOptions(CGSizeMake(10,10), NO, 0);
CGContextRef con = UIGraphicsGetCurrentContext();
CGContextAddEllipseInRect(con, CGRectMake(0,0,10,10));
CGContextSetFillColorWithColor(con, [UIColor blackColor].CGColor);
CGContextFillPath(con);
UIImage* im = UIGraphicsGetImageFromCurrentImageContext();
UIGraphicsEndImageContext();

CAEmitterCell* cell = [CAEmitterCell emitterCell];
emit.emitterCells = [NSArray arrayWithObject:cell];
cell.birthRate = 5;
cell.lifetime = 1;
cell.velocity = 100;
cell.contents = (id)im.CGImage;
```

The result is that little black circles should be emitted slowly and steadily, five per second, each one vanishing in five seconds. Now we need an emitter layer from which these circles are to be emitted. Here are some basic CAEmitterLayer properties (beyond those it inherits from CALayer); these define an imaginary object, an emitter, that will be producing the emitter cells:

emitterPosition

The point at which the emitter should located, in superlayer coordinates. You can optionally add a third dimension to this point, `emitterZPosition`.

emitterSize

The size of the emitter.

emitterShape

The shape of the emitter. The dimensions of the shape depend on the emitter's size; the cuboid shape depends also on a third size dimension, `emitterDepth`. Your choices are:

- `kCAEmitterLayerPoint`
- `kCAEmitterLayerLine`
- `kCAEmitterLayerRectangle`

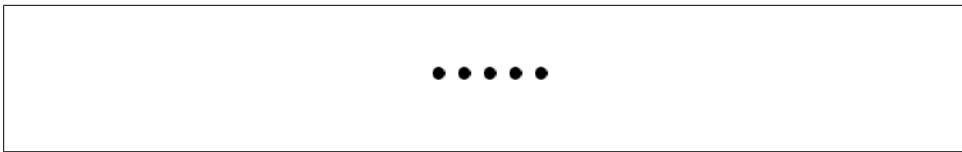


Figure 17-5. A really boring emitter layer

- `kCAEmitterLayerCuboid`
- `kCAEmitterLayerCircle`
- `kCAEmitterLayerSphere`

`emitterMode`

The region of the shape from which cells should be emitted. Your choices are:

- `kCAEmitterLayerPoints`
- `kCAEmitterLayerOutline`
- `kCAEmitterLayerSurface`
- `kCAEmitterLayerVolume`

Let's start with the simplest possible case, a single point emitter:

```
CAEmitterLayer* emit = [[CAEmitterLayer alloc] init];
emit.emitterPosition = CGPointMake(30,100);
emit.emitterShape = kCAEmitterLayerPoint;
emit.emitterMode = kCAEmitterLayerPoints;
```

We tell the emitter what types of cell to emit by assigning to its `emitterCells` property (an array of `CAEmitterCell`). We have only one type of cell. We then add the emitter to our interface, and presto, it starts emitting:

```
emit.emitterCells = [NSArray arrayWithObject:cell];
[self.window.rootViewController.view.layer addSublayer:emit];
```

The result is a constant stream of black circles emitted from the point `{30,100}`, each circle marching steadily to the right and vanishing after one second ([Figure 17-5](#)).

Now that we've succeeded in creating a boring emitter layer, we can start to vary some parameters. The `emissionRange` defines a cone in which cells will be emitted; if we increase the `birthRate` and widen the `emissionRange`, we get something that looks like a stream coming from a water hose:

```
cell.birthRate = 100;
cell.lifetime = 1;
cell.velocity = 100;
cell.emissionRange = M_PI/10;
```

As the cell moves, it can be made to accelerate (or decelerate) in each dimension, using its `xAcceleration`, `yAcceleration`, and `zAcceleration` properties. Here, we turn the stream into a falling cascade, like a waterfall coming from the left:

```
cell.birthRate = 100;
cell.lifetime = 1.5;
cell.velocity = 100;
cell.emissionRange = M_PI/10;
cell.xAcceleration = -40;
cell.yAcceleration = 200;
```

All aspects of cell behavior can be made to vary, using the following CAEmitterCell properties:

lifetimeRange, velocityRange

How much the lifetime and velocity values are allowed to vary randomly for different cells.

scale

scaleRange, scaleSpeed

The scale alters the size of the cell; the range and speed determine how far and how rapidly this size alteration is allowed to change over the lifetime of each cell.

color

redRange, greenRange, blueRange, alphaRange

redSpeed, greenSpeed, blueSpeed, alphaSpeed

The color is painted in accordance with the opacity of the cell's contents image; it combines with the image's color, so if we want the color stated here to appear in full purity, our contents image should use only white. The range and speed determine how far and how rapidly each color component is to change.

spin, spinRange

The spin is a rotational speed (in radians per second); its range determines how far this speed is allowed to change over the lifetime of each cell.

Here we apply some variation so that the circles behave a little more independently of one another. Some live longer than others, some come out of the emitter faster than others. And they all start out a shade of blue, but change to a shade of green about halfway through the stream ([Figure 17-6](#)):

```
cell.birthRate = 100;
cell.lifetime = 1.5;
cell.lifetimeRange = .4;
cell.velocity = 100;
cell.velocityRange = 20;
cell.emissionRange = M_PI/5;
cell.scale = 1;
cell.scaleRange = .2;
cell.scaleSpeed = .2;
cell.xAcceleration = -40;
cell.yAcceleration = 200;
cell.color = [UIColor blueColor].CGColor;
cell.greenRange = .5;
cell.greenSpeed = .75;
```

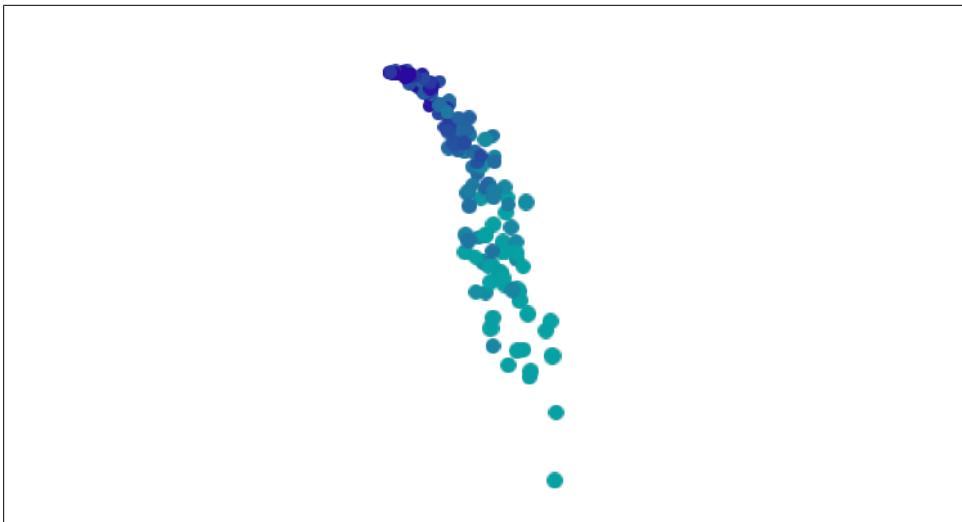


Figure 17-6. An emitter layer that makes a sort of waterfall

But wait, there's more! Once the emitter layer is in place and animating, you can change its parameters and the parameters of its emitter cells. To do so, use KVC on the emitter layer. You can access the emitter cells through the emitter layer's @"emitterCells" key path; to specify a cell type, use its `name` property (which you'll have to have assigned earlier) the next piece of the key path. For example, suppose we've set `cell.name` to @"circle"; now we'll change the cell's `greenSpeed` so that each cell changes from blue to green much earlier in its lifetime:

```
[emit setValue:[NSNumber numberWithFloat:3]
          forKeyPath:@"emitterCells.circle.greenSpeed"];
```

But wait, there's *still* more: such changes can themselves be animated! Here, we'll attach to the emitter layer a repeating animation that causes our cell's `greenSpeed` to move back and forth between two values. The result is that the stream is sometimes mostly blue and sometimes mostly green:

```
CABasicAnimation* ba =
    [CABasicAnimation animationWithKeyPath:@"emitterCells.circle.greenSpeed"];
ba.fromValue = [NSNumber numberWithFloat:-1];
ba.toValue = [NSNumber numberWithFloat:3];
ba.duration = 4;
ba.autoreverses = YES;
ba.repeatCount = HUGE_VALF;
[emit addAnimation:ba forKey:nil];
```

But wait, there's *still* still more! A CAEmitterCell can itself function as an emitter — that is, it can have cells of its own. Both CAEmitterLayer and CAEmitterCell conform to the CAMediaTiming protocol, and their `beginTime` and `duration` properties can be used to govern their times of operation, much as in a grouped animation. For example,

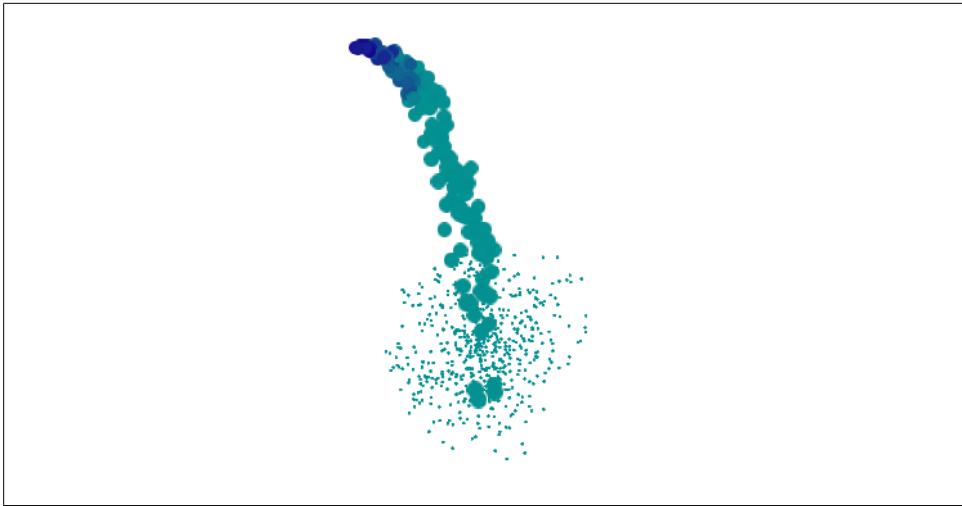


Figure 17-7. The waterfall makes a kind of splash

this code causes our existing waterfall to spray tiny droplets in the region of the “nozzle” (the emitter):

```
CAEmitterCell* cell2 = [CAEmitterCell emitterCell];
cell.emitterCells = [NSArray arrayWithObject: cell2];
cell2.contents = (id)im.CGImage;
cell2.emissionRange = M_PI;
cell2.birthRate = 200;
cell2.lifetime = 0.4;
cell2.velocity = 200;
cell2.scale = 0.2;
cell2.beginTime = .04;
cell2.duration = .2;
```

But if we change the `beginTime` to be larger (hence later), the tiny droplets happen near the bottom of the cascade. We must also increase the `duration`, or stop setting it altogether, since if the `duration` is less than the `beginTime`, no emission takes place at all ([Figure 17-7](#)):

```
cell2.beginTime = .7;
cell2.duration = .8;
```

Of course we can also completely change the picture by changing the behavior of the emitter itself. This change turns the emitter into a line, so that our cascade becomes broader:

```
emit.emitterPosition = CGPointMake(100,25);
emit.emitterSize = CGSizeMake(100,100);
emit.emitterShape = kCAEmitterLayerLine;
emit.emitterMode = kCAEmitterLayerOutline;
cell.emissionLongitude = 3*M_PI/4;
```

There remains more to know about emitter layers and emitter cells, but at this point you know enough to understand Apple's sample code, one portraying fire and smoke, and the other simulating fireworks, and you can explore further on your own.

CHAPTER 18

Touches

*[Winifred the Woebegone illustrates hit-testing:] Hey nonny nonny, is it you? — Hey nonny nonny nonny no!
— Hey nonny nonny, is it you? — Hey nonny nonny nonny no!*

—Marshall Barer, *Once Upon a Mattress*

A *touch* is an instance of the user putting a finger on the screen. The system and the hardware, working together, know *when* a finger contacts the screen and *where* it is. (Fingers are fat, but the system and the hardware cleverly reduce the finger's location to a single appropriate point.)

A UIView, by virtue of being a UIResponder, is the visible locus of touches. There are other UIResponder subclasses, but none of them is visible on the screen. What the user sees are views; what the user is touching are views. (The user may also see layers, but a layer is not a UIResponder and is not involved with touches. I'll talk later about how to make it seem as if the user can touch a layer.)

It would make sense, therefore, if every touch were reported directly to the view in which it occurred. However, what the system “sees” is not particular views but an app as a whole. So a touch is represented as an object (a UITouch instance) which is bundled up in an envelope (a UIEvent) which the system delivers to your app. It is then up to your app to deliver the envelope to an appropriate UIView. In the vast majority of cases, this will happen automatically the way you expect, and you will respond to a touch by way of the view in which the touch occurred.

In fact, usually you won't concern yourself with UIEvents and UITouches at all. Most built-in interface views deal with these low-level touch reports themselves, and notify your code at a higher level. When a UIButton emits an action message to report a control event such as Touch Up Inside (Chapter 11), it has already performed a reduction of a complex sequence of touches (“the user put a finger down inside me and then, possibly with some dragging hither and yon, raised it when it was still reasonably close to me”). A UITextField reports touches on the keyboard as changes in its own text. A UITableView

View reports that the user selected a cell. A UIScrollView, when dragged, reports that it scrolled; when pinched outward, it reports that it zoomed. Some interface views respond to touches internally without reporting to your code at all; for example, a UIWebView, when dragged, just scrolls.

Nevertheless, it is useful to know how to respond to touches directly, so that you can implement your own touchable views, and so that you understand what Cocoa's built-in views are actually doing. This chapter discusses touch detection and response by views (and other UIResponders) at their lowest level, along with a slightly higher-level mechanism, gesture recognizers, that categorizes touches into gesture types for you; then it deconstructs the touch-delivery architecture by which touches are reported to your views in the first place.

Touch Events and Views

Imagine a screen that the user is not touching at all: the screen is “finger-free.” Now the user touches the screen with one or more fingers. From that moment to the time the screen is once again finger-free, all touches and finger movements together constitute what Apple calls a single *multitouch sequence*.

The system reports to your app, during a given multitouch sequence, every change in finger configuration, so that your app can figure out what the user is doing. Every such report is a UIEvent. In fact, every report having to do with the same multitouch sequence is *the same UIEvent instance*, arriving repeatedly, each time there's a change in finger configuration.

Every UIEvent reporting a change in the user's finger configuration contains one or more UITouch objects. Each UITouch object corresponds to a single finger; conversely, every finger touching the screen is represented in the UIEvent by a UITouch object. Once a certain UITouch instance has been created to represent a finger that has touched the screen, *the same UITouch instance* is used to represent that finger throughout this multitouch sequence until the finger leaves the screen.

Now, it might sound as if the system has to bombard the app with huge numbers of reports constantly during a multitouch sequence. But that's not really true. The system needs to report only *changes* in the finger configuration. For a given UITouch object (representing, remember, a specific finger), only four things can happen. These are called *touch phases*, and are described by a UITouch instance's *phase* property:

UITouchPhaseBegan

The finger touched the screen for the first time; this UITouch instance has just been created. This is always the first phase, and arrives only once.

UITouchPhaseMoved

The finger moved upon the screen.

UITouchPhaseStationary

The finger remained on the screen without moving. Why is it necessary to report this? Well, remember, once a UITouch instance has been created, it must be present every time the UIEvent arrives. So if the UIEvent arrives because something *else* happened (e.g., a new finger touched the screen), we must report what *this* finger has been doing, even if it has been doing nothing.

UITouchPhaseEnded

The finger left the screen. Like `UITouchPhaseBegan`, this phase arrives only once. The UITouch instance will now be destroyed and will no longer appear in UIEvents for this multitouch sequence.

Those four phases are sufficient to describe everything that a finger can do. Actually, there is one more possible phase:

UITouchPhaseCancelled

The system has aborted this multitouch sequence because something interrupted it.

What might interrupt a multitouch sequence? There are many possibilities. Perhaps the user clicked the Home button or the screen lock button in the middle of the sequence. A local notification alert may have appeared ([Chapter 26](#)); on an actual iPhone, a call might have come in. (As we shall see, a gesture recognizer recognizing its gesture may also trigger touch cancellation.) The point is, if you’re dealing with touches yourself, you cannot afford to ignore touch cancellation; they are your opportunity to get things into a coherent state when the sequence is interrupted.

When a UITouch first appears (`UITouchPhaseBegan`), your app works out which UIView it is associated with. (I’ll give full details, later in this chapter, as to how it does that.) This view is then set as the touch’s `view` property; from then on, this UITouch is *always* associated with this view. In other words, *a touch’s view is that touch’s view forever* (until that finger leaves the screen).

The same UIEvent containing the same UITouches can be sent to multiple views; these are programmatic objects, not real-world envelopes containing actual fingers. Accordingly, a UIEvent is distributed to *all the views of all the UITouches it contains*. Conversely, if a view is sent a UIEvent, it’s because that UIEvent contains at least one UITouch whose `view` is this view.

If every UITouch in a UIEvent associated with a certain UIView has the phase `UITouchPhaseStationary`, that UIEvent is *not* sent to that UIView. There’s no point, because as far as that view is concerned, nothing happened.

Receiving Touches

A UIResponder, and therefore a UIView, has four methods corresponding to the four UITouch phases that require UIEvent delivery. A UIEvent is delivered to a view by calling one or more of these four methods (the *touches... methods*):

`touchesBegan:withEvent:`

A finger touched the screen, creating a UITouch.

`touchesMoved:withEvent:`

A finger previously reported to this view with `touchesBegan:withEvent:` has moved.

`touchesEnded:withEvent:`

A finger previously reported to this view with `touchesBegan:withEvent:` has left the screen.

`touchesCancelled:withEvent:`

We are bailing out on a finger previously reported to this view with `touchesBegan:withEvent:`.

The parameters of these methods are:

The relevant touches

These are the event's touches whose phase corresponds to the name of the method and (normally) whose view is this view. They arrive as an NSSet ([Chapter 10](#)). If you know for a fact that there is only one touch in the set, or that any touch in the set will do, you can retrieve it with `anyObject` (an NSSet doesn't implement `lastObject` because a set is unordered).

The event

This is the UIEvent instance. It contains its touches as an NSSet, which you can retrieve with the `allTouches` message. This means *all* the event's touches, including but not necessarily limited to those in the first parameter; there might be touches in a different phase or intended for some other view. You can call `touchesForView:` or `touchesForWindow:` to ask for the set of touches associated with a particular view or window.

A UITouch has some useful methods and properties:

`locationInView:, previousLocationInView:`

The current and previous location of this touch with respect to the coordinate system of a given view. The view you'll be interested in will often be `self` or `self.superview`; supply `nil` to get the location with respect to the window. The previous location will be of interest only if the phase is `UITouchPhaseMoved`.

`timestamp`

When the touch last changed. A touch is timestamped when it is created (`UITouchPhaseBegan`) and each time it moves (`UITouchPhaseMoved`).

tapCount

If two touches are in roughly the same place in quick succession, and the first one is brief, the second one may be characterized as a repeat of the first. They are different touch objects, but the second will be assigned a `tapCount` one larger than the previous one. The default is 1, so if (for example) a touch's `tapCount` is 3 then this is the third tap in quick succession in roughly the same spot.

view

The view with which this touch is associated.

Here are some additional UIEvent properties:

type

This will be `UIEventTypeTouches`. There are other event types, but you're not going to receive any of them this way.

timestamp

When the event occurred.

So, when we say that a certain view *is receiving a touch*, that is a shorthand expression meaning that it is being sent a UIEvent containing this UITouch, over and over, by calling one of its `touches...` methods, corresponding to the phase this touch is in, from the time the touch is created until the time it is destroyed.

Restricting Touches

Touch events can be turned off entirely at the application level with `UIApplication`'s `beginIgnoringInteractionEvents`. It is quite common to do this during animations and other lengthy operations during which responding to a touch could cause undesirable results. This call should be balanced by `endIgnoringInteractionEvents`. Pairs can be nested, in which case interactivity won't be restored until the outermost `endIgnoringInteractionEvents` has been reached.

A number of high-level `UIView` properties also restrict the delivery of touches to particular views:

userInteractionEnabled

If set to NO, this view (along with its subviews) is excluded from receiving touches. Touches on this view or one of its subviews "fall through" to a view behind it.

opacity

If set to 0.0 (or extremely close to it), this view (along with its subviews) is excluded from receiving touches. Touches on this view or one of its subviews "fall through" to a view behind it.

`hidden`

If set to YES, this view (along with its subviews) is excluded from receiving touches. This makes sense, since from the user's standpoint, the view and its subviews are not even present.

`multipleTouchEnabled`

If set to NO, this view never receives more than one touch simultaneously; once it receives a touch, it doesn't receive any other touches until that first touch has ended.

`exclusiveTouch`

This is the only one of these properties that can't be set in the nib. An `exclusiveTouch` view receives a touch only if no other views in the same window have touches associated with them; once an `exclusiveTouch` view has received a touch, then while that touch exists no other view in the same window receives any touches.



A UIWindow ignores `multipleTouchEnabled`; it always receives multiple touches. Moreover, a UIWindow's behavior with respect to `exclusiveTouch` is unreliable, presumably because it is not itself a view in the window. In general this should not be an issue, since you'll always have a root view covering the window anyway.

Interpreting Touches

Thanks to the existence of gesture recognizers (discussed later in this chapter), in most cases you won't have to interpret touches at all; you'll let a gesture recognizer do most of that work. Even so, it is beneficial to be conversant with the nature of touch interpretation; this will help you interact with a gesture recognizer, write your own gesture recognizer, or subclass an existing one. Furthermore, not every touch sequence can be codified through a gesture recognizer; sometimes, directly interpreting touches is the best approach.

To figure out what's going on as touches are received by a view, your code must essentially function as a kind of state machine. You'll receive various `touches...` method calls, and your response will partly depend upon what happened previously, so you'll have to record somehow, such as in instance variables, the information that you'll need in order to decide what to do when the next `touches...` method is called. Such an architecture can make writing and maintaining touch-analysis code quite tricky. Moreover, although you can distinguish a particular UITouch or UIEvent object over time by keeping a reference to it, you mustn't retain that reference; it doesn't belong to you.

To illustrate the business of interpreting touches, we'll start with a view that can be dragged with the user's finger. For simplicity, I'll assume that this view receives only a single touch at a time. (This assumption is easy to enforce by setting the view's `multipleTouchEnabled` to NO, which is the default.)

The trick to making a view follow the user's finger is to realize that a view is positioned by its `center`, which is in superview coordinates, but the user's finger might not be at the center of the view. So at every stage of the drag we must change the view's center by the change in the user's finger position in superview coordinates:

```
- (void) touchesMoved:(NSSet *)touches withEvent:(UIEvent *)event {
    CGPoint loc = [[touches anyObject] locationInView: self.superview];
    CGPoint oldP = [[touches anyObject] previousLocationInView: self.superview];
    CGFloat deltaX = loc.x - oldP.x;
    CGFloat deltaY = loc.y - oldP.y;
    CGPoint c = self.center;
    c.x += deltaX;
    c.y += deltaY;
    self.center = c;
}
```

Next, let's add a restriction that the view can be dragged only vertically or horizontally. All we have to do is hold one coordinate steady; but which coordinate? Everything seems to depend on what the user does initially. So we'll do a one-time test the first time we receive `touchesMoved:withEvent:`. Now we're maintaining two state variables, `decided` and `horiz`:

```
- (void) touchesBegan:(NSSet *)touches withEvent:(UIEvent *)event {
    self->decided = NO;
}

- (void) touchesMoved:(NSSet *)touches withEvent:(UIEvent *)event {
    if (!self->decided) {
        self->decided = YES;
        CGPoint then = [[touches anyObject] previousLocationInView: self];
        CGPoint now = [[touches anyObject] locationInView: self];
        CGFloat deltaX = fabs(then.x - now.x);
        CGFloat deltaY = fabs(then.y - now.y);
        self->horiz = (deltaX >= deltaY);
    }
    CGPoint loc = [[touches anyObject] locationInView: self.superview];
    CGPoint oldP = [[touches anyObject] previousLocationInView: self.superview];
    CGFloat deltaX = loc.x - oldP.x;
    CGFloat deltaY = loc.y - oldP.y;
    CGPoint c = self.center;
    if (self->horiz)
        c.x += deltaX;
    else
        c.y += deltaY;
    self.center = c;
}
```

Look at how things are trending. We are maintaining state variables, which we are managing across multiple methods, and we are subdividing a `touches...` method implementation into tests depending on the state of our state machine. Our state machine is very simple, involving just two state variables, but already our code is becoming difficult to read and to maintain. Things only become more messy as we try to make our view's behavior more sophisticated. (In the first edition of this book, I illustrated

by proceeding to show how to allow the user to “fling” the view horizontally or vertically and how to deal with multiple simultaneous touches, but in this edition I’ve decided it would be best to leave that code to your imagination.)

Another area in which manual touch handling can rapidly prove overwhelming is when it comes to distinguishing between different gestures that the user is to be permitted to perform on a view. Imagine, for example, a view that distinguishes between a finger tapping briefly and a finger remaining down for a longer time. We can’t know how long a tap is until it’s over, so one approach might be to wait until then before deciding:

```
- (void) touchesBegan:(NSSet *)touches withEvent:(UIEvent *)event {
    self->time = [[touches anyObject] timestamp];
}

- (void) touchesEnded:(NSSet *)touches withEvent:(UIEvent *)event {
    NSTimeInterval diff = event.timestamp - self->time;
    if (diff < 0.4)
        NSLog(@"short");
    else
        NSLog(@"long");
}
```

On the other hand, one might argue that if a tap hasn’t ended after some set time (here, 0.4 seconds), we know that it is long, and so we could begin responding to it without waiting for it to end. The problem is that we don’t automatically get an event after 0.4 seconds. So we’ll create one, using delayed performance:

```
- (void) touchesBegan:(NSSet *)touches withEvent:(UIEvent *)event {
    self->time = [[touches anyObject] timestamp];
    [self performSelector:@selector(touchWasLong) withObject:nil afterDelay:0.4];
}

- (void) touchesEnded:(NSSet *)touches withEvent:(UIEvent *)event {
    NSTimeInterval diff = event.timestamp - self->time;
    if (diff < 0.4)
        NSLog(@"short");
}

- (void) touchWasLong {
    NSLog(@"long");
}
```

But there’s a bug. If the tap is short, we report that it was short, but we *also* report that it was long. That’s because the delayed call to `touchWasLong` arrives anyway. We could use some sort of boolean flag to tell us when to ignore that call, but there’s a better way: `NSObject` has a class method that lets us cancel any pending delayed performance calls. So:

```
- (void) touchesBegan:(NSSet *)touches withEvent:(UIEvent *)event {
    self->time = [[touches anyObject] timestamp];
    [self performSelector:@selector(touchWasLong) withObject:nil afterDelay:0.4];
}
```

```

- (void) touchesEnded:(NSSet *)touches withEvent:(UIEvent *)event {
    NSTimeInterval diff = event.timestamp - self->time;
    if (diff < 0.4) {
        NSLog(@"short");
        [NSObject cancelPreviousPerformRequestsWithTarget:self
                                                    selector:@selector(touchWasLong)
                                                    object:nil];
    }
}

- (void) touchWasLong {
    NSLog(@"long");
}

```

Here's another use of the same technique. We'll distinguish between a single tap and a double tap. The `UITouch` `tapCount` property already makes this distinction, but that, by itself, is not enough to help us react differently to the two. What we must do, having received a tap whose `tapCount` is 1, is to delay responding to it long enough to give a second tap a chance to arrive. This is unfortunate, because it means that if the user intends a single tap, some time will elapse before anything happens in response to it; however, there's nothing we can easily do about that.

Distributing our various tasks correctly is a bit tricky. We *know* when we have a double tap as early as `touchesBegan:withEvent:`, so that's when we cancel our delayed response to a single tap, but we *respond* to the double tap in `touchesEnded:withEvent:`. We don't start our delayed response to a single tap until `touchesEnded:withEvent:`, because what matters is the time between the taps as a whole, not between the starts of the taps. This code is adapted from Apple's own example:

```

- (void) touchesBegan:(NSSet *)touches withEvent:(UIEvent *)event {
    int ct = [[touches anyObject] tapCount];
    if (ct == 2) {
        [NSObject cancelPreviousPerformRequestsWithTarget:self
                                                    selector:@selector(singleTap)
                                                    object:nil];
    }
}

- (void) touchesEnded:(NSSet *)touches withEvent:(UIEvent *)event {
    int ct = [[touches anyObject] tapCount];
    if (ct == 1)
        [self performSelector:@selector(singleTap) withObject:nil afterDelay:0.3];
    if (ct == 2)
        NSLog(@"double tap");
}

- (void) singleTap {
    NSLog(@"single tap");
}

```

Now let's consider combining our detection for a single or double tap with our earlier code for dragging a view horizontally or vertically. This is to be a view that can detect four kinds of gesture: a single tap, a double tap, a horizontal drag, and a vertical drag.

We must include the code for all possibilities and make sure they don't interfere with each other. The result is rather horrifying, a forced join between two already complicated sets of code, along with an additional pair of state variables to track the decision between the tap gestures on the one hand and the drag gestures on the other:

```
- (void) touchesBegan:(NSSet *)touches withEvent:(UIEvent *)event {
    // be undecided
    self->decidedTapOrDrag = NO;
    // prepare for a tap
    int ct = [[touches anyObject] tapCount];
    if (ct == 2) {
        [NSObject cancelPreviousPerformRequestsWithTarget:self
                                                    selector:@selector(singleTap)
                                                    object:nil];
        self->decidedTapOrDrag = YES;
        self->drag = NO;
        return;
    }
    // prepare for a drag
    self->decidedDirection = NO;
}

- (void) touchesMoved:(NSSet *)touches withEvent:(UIEvent *)event {
    if (self->decidedTapOrDrag && !self->drag)
        return;
    self->decidedTapOrDrag = YES;
    self->drag = YES;
    if (!self->decidedDirection) {
        self->decidedDirection = YES;
        CGPoint then = [[touches anyObject] previousLocationInView: self];
        CGPoint now = [[touches anyObject] locationInView: self];
        CGFloat deltaX = fabs(then.x - now.x);
        CGFloat deltaY = fabs(then.y - now.y);
        self->horiz = (deltaX >= deltaY);
    }
    CGPoint loc = [[touches anyObject] locationInView: self.superview];
    CGPoint oldP = [[touches anyObject] previousLocationInView: self.superview];
    CGFloat deltaX = loc.x - oldP.x;
    CGFloat deltaY = loc.y - oldP.y;
    CGPoint c = self.center;
    if (self->horiz)
        c.x += deltaX;
    else
        c.y += deltaY;
    self.center = c;
}

- (void) touchesEnded:(NSSet *)touches withEvent:(UIEvent *)event {
    if (!self->decidedTapOrDrag || !self->drag) {
        // end for a tap
        int ct = [[touches anyObject] tapCount];
        if (ct == 1)
            [self performSelector:@selector(singleTap) withObject:nil
                           afterDelay:0.3];
        if (ct == 2)
```

```

        NSLog(@"double tap");
    return;
}
}

- (void) singleTap {
    NSLog(@"single tap");
}

```

That code seems to work, but it's hard to say whether it covers all possibilities coherently; it's barely legible and the logic borders on the mysterious. This is the kind of situation for which gesture recognizers were devised.

Gesture Recognizers

Writing and maintaining a state machine that interprets touches across a combination of three or four `touches...` methods is hard enough when a view confines itself to expecting only one kind of gesture, such as dragging. It becomes even more involved when a view wants to accept and respond differently to different kinds of gesture. Furthermore, many types of gesture are conventional and standard; it seems insane to require developers to implement independently the elements that constitute what is, in effect, a universal vocabulary.

The solution is gesture recognizers, which standardize common gestures and allow the code for different gestures to be separated and encapsulated into different objects.

Gesture Recognizer Classes

A gesture recognizer (a subclass of `UIGestureRecognizer`) is an object attached to a `UIView`, which has for this purpose methods `addGestureRecognizer:` and `removeGestureRecognizer:`, and a `gestureRecognizers` property. A `UIGestureRecognizer` implements the four `touches...` handlers, but it is not a responder (a `UIResponder`), so it does not participate in the responder chain.

If a new touch is going to be delivered to a view, it is also associated with and delivered to that view's gesture recognizers if it has any, and to that view's superview's gesture recognizers if it has any, and so on up the view hierarchy. Thus, the place of a gesture recognizer in the view hierarchy matters, even though it isn't part of the responder chain.

`UITouch` and `UIEvent` provide complementary ways of learning how touches and gesture recognizers are associated. `UITouch`'s `gestureRecognizers` lists the gesture recognizers that are currently handling this touch. `UIEvent`'s `touchesForGestureRecognizer:` lists the touches that are currently being handled by a particular gesture recognizer.

Each gesture recognizer maintains its own state as touch events arrive, building up evidence as to what kind of gesture this is. When one of them decides that it has rec-

ognized its own type of gesture, it emits either a single message (to indicate, for example, that the user tapped this view) or a series of messages (to indicate, for example, that the user is dragging this view); the distinction here is between a *discrete* and a *continuous* gesture. What message a gesture recognizer emits, and to what object it sends it, is set through a target–action dispatch table attached to the gesture recognizer; a gesture recognizer is rather like a UIControl ([Chapter 11](#)) in this regard. Indeed, one might say that a gesture recognizer simplifies the touch handling of *any* view to be like that of a control. The difference is that one control may report several different control events, whereas each gesture recognizer reports only one gesture type, with different gestures being reported by different gesture recognizers. This architecture implies that it is unnecessary to subclass UIView merely in order to implement touch analysis.

UIGestureRecognizer itself is abstract, providing methods and properties to its subclasses. Among these are:

`initWithTarget:action:`

The designated initializer. Each message emitted by a UIGestureRecognizer is simply a matter of sending the action message to the target. Further target–action pairs may be added with `addTarget:action:`, and removed with `removeTarget:action:`.

Two forms of selector are possible: either there is no parameter, or there is a single parameter which will be the gesture recognizer. Most commonly, you’ll use the second form, so that the target can identify and query the gesture recognizer; moreover, using the second form also gives the target a reference to the view, because the gesture recognizer provides a reference to its view as the `view` property.

`locationOfTouch:inView:`

The touch is specified by an index number. The `numberOfTouches` property provides a count of current touches; the touches themselves are inaccessible from outside the gesture recognizer.

`enabled`

A convenient way to turn a gesture recognizer off without having to remove it from its view.

`state, view`

I’ll discuss state later on. The view is the view to which this gesture recognizer is attached.

Built-in UIGestureRecognizer subclasses are provided for six common gesture types: tap, pinch, pan (drag), swipe, rotate, and long press. These embody properties and methods likely to be needed for each type of gesture, either in order to configure the gesture recognizer beforehand or in order to query it as to the state of an ongoing gesture:

`UITapGestureRecognizer (discrete)`

Configuration: `numberOfTapsRequired`, `numberOfTouchesRequired` (“touches” means simultaneous fingers).

UIPinchGestureRecognizer (*continuous*)

State: `scale`, `velocity`.

UIRotationGestureRecognizer (*continuous*)

State: `rotation`, `velocity`.

UISwipeGestureRecognizer (*discrete*)

Configuration: `direction` (meaning permitted directions, a bitmask), `numberOfTouchesRequired`.

UIPanGestureRecognizer (*continuous*)

Configuration: `minimumNumberOfTouches`, `maximumNumberOfTouches`.

State: `translationInView:`, `setTranslation:inView:`, and `velocityInView:`; the coordinate system of the specified view is used, so to follow a finger you'll use the superview of the view being dragged, just as we did in the examples earlier.

UILongPressGestureRecognizer (*continuous*)

Configuration: `numberOfTapsRequired`, `numberOfTouchesRequired`, `minimumPressDuration`, `allowableMovement`. The `numberOfTapsRequired` is the count of taps *before* the tap that stays down; so it can be 0 (the default). The `allowableMovement` setting lets you compensate for the fact that the user's finger is unlikely to remain steady during an extended press; thus we need to provide some limit before deciding that this gesture is, say, a drag, and not a long press after all. On the other hand, once the long press is recognized, the finger is permitted to drag.

UIGestureRecognizer also provides a `locationInView:` method. This is a single point, even if there are multiple touches. The subclasses implement this variously. For example, for *UIPanGestureRecognizer*, the location is where the touch is if there's a single touch, but it's a sort of midpoint ("centroid") if there are multiple touches.

We already know enough to implement, using a gesture recognizer, a view that responds to a single tap, or a view that responds to a double tap. We don't yet know quite enough to implement a view that lets itself be dragged around, or a view that can respond to more than one gesture; we'll come to that. Meanwhile, here's code that implements a view that responds to a single tap:

```
UITapGestureRecognizer* t = [[UITapGestureRecognizer alloc]
                           initWithTarget:self
                           action:@selector(singleTap)];
[v addGestureRecognizer:t];
// ...
- (void) singleTap {
    NSLog(@"single");
}
```

And here's code that implements a view that responds to a double tap:

```
UITapGestureRecognizer* t = [[UITapGestureRecognizer alloc]
                           initWithTarget:self
                           action:@selector(doubleTap)];
t.numberOfTapsRequired = 2;
```

```

[v addGestureRecognizer:t];
// ...
- (void) doubleTap {
    NSLog(@"double");
}

```

For a continuous gesture like dragging, we need to know both when the gesture is in progress and when the gesture ends. This brings us to the subject of a gesture recognizer's state.

A gesture recognizer implements a notion of *states* (the `state` property); it passes through these states in a definite progression. The gesture recognizer remains in the Possible state until it can make a decision one way or the other as to whether this is in fact the correct gesture. The documentation neatly lays out the possible progressions:

Wrong gesture

Possible → Failed. No action message is sent.

Discrete gesture (like a tap), recognized

Possible → Ended. One action message is sent, when the state changes to Ended.

Continuous gesture (like a drag), recognized

Possible → Began → Changed (repeatedly) → Ended. Action messages are sent once for Began, as many times as necessary for Changed, and once for Ended.

Continuous gesture, recognized but later cancelled

Possible → Began → Changed (repeatedly) → Cancelled. Action messages are sent once for Began, as many times as necessary for Changed, and once for Cancelled.

The actual phase names are `UIGestureRecognizerStatePossible` and so forth. The name `UIGestureRecognizerStateRecognized` is actually a synonym for the Ended state; I find this unnecessary and confusing and I'll ignore it in my discussion.

We now know enough to implement, using a gesture recognizer, a view that lets itself be dragged around in any direction by a single finger. Our maintenance of state is greatly simplified, because a UIPanGestureRecognizer maintains a delta (translation) for us. This delta, available using `translationInView:`, is reckoned from the touch's initial position. So we need to store our center only once:

```

UIPanGestureRecognizer* p = [[UIPanGestureRecognizer alloc]
                           initWithTarget:self
                           action:@selector(dragging:)];
[v addGestureRecognizer:p];
// ...
- (void) dragging: (UIPanGestureRecognizer*) p {
    UIView* vv = p.view;
    if (p.state == UIGestureRecognizerStateBegan)
        self->origC = vv.center;
    CGPoint delta = [p translationInView: vv.superview];

```

```

    CGPoint c = self->origC;
    c.x += delta.x; c.y += delta.y;
    vv.center = c;
}

```

Actually, it's possible to write that code without maintaining any state at all, because we are allowed to reset the UIPanGestureRecognizer's delta, using `setTranslation:inView:`. So:

```

- (void) dragging: (UIPanGestureRecognizer*) p {
    UIView* vv = p.view;
    if (p.state == UIGestureRecognizerStateBegan ||
        p.state == UIGestureRecognizerStateChanged) {
        CGPoint delta = [p translationInView: vv.superview];
        CGPoint c = vv.center;
        c.x += delta.x; c.y += delta.y;
        vv.center = c;
        [p setTranslation: CGPointMakeZero inView: vv.superview];
    }
}

```

A gesture recognizer also works, as I've already mentioned, if it is attached to the superview (or further up the hierarchy) of the view in which the user gestures. For example, if a tap gesture recognizer is attached to the window, the user can tap on any view within that window, and the tap will be recognized; the view's presence does not "block" the window from recognizing the gesture, even if it is a UIControl that responds autonomously to touches.

This behavior comes as a surprise to beginners, but it makes sense, because if it were not the case, certain gestures would be impossible. Imagine, for example, a pair of views on each of which the user can tap individually, but which the user can also touch simultaneously (one finger on each view) and rotate together around their mutual centroid. Neither view can detect the rotation *qua* rotation, because neither view receives both touches; only the superview can detect it, so the fact that the views themselves respond to touches must not prevent the superview's gesture recognizer from operating.

Suppose, then, that your window's root view has a UITapGestureRecognizer attached to it (perhaps because you want to be able to recognize taps on the background), but there is also a UIButton within it. How is the root view to ignore a tap on the button? The best way, I think, is through the gesture recognizer's delegate, which can tell the gesture recognizer to ignore touches intended for another view. (I'll talk about gesture recognizer delegates later in this chapter.) And if the gesture recognizer attached to the root view is for some gesture other than a tap, then the problem never arises, because a tap on the button won't trigger the gesture recognizer's action handler in the first place. Nevertheless, the view hierarchy does complicate the use of gesture recognizers; fortunately, as I shall explain, gesture recognizers usually take care of these complications themselves, and when they don't, they provide ways for you to resolve them.

Multiple Gesture Recognizers

The question naturally arises of what happens when multiple gesture recognizers are in play. This isn't a matter merely of multiple recognizers attached to a single view, because, as I have just said, if a view is touched, not only its own gesture recognizers but any gesture recognizers attached to views further up the view hierarchy are also in play, simultaneously. I like to think of a view as surrounded by a swarm of gesture recognizers — its own and those of its superview (and so on). In reality, it is a touch that has a swarm of gesture recognizers; that's why a `UITouch` has a `gestureRecognizers` property, in the plural.

In general, once a gesture recognizer succeeds in recognizing its gesture, any *other* gesture recognizers associated with its touches are *forced into the Failed state*, and whatever touches were associated with those gesture recognizers are no longer sent to them; in effect, the first gesture recognizer in a swarm that recognizes its gesture owns the gesture, and those touches, from then on.

In many cases, this behavior alone will correctly eliminate conflicts. For example, we can add *both* our `UITapGestureRecognizer` for a single tap *and* our `UIPanGestureRecognizer` to a view and everything will just work.

What happens if we also add the `UITapGestureRecognizer` for a double tap? Dragging works, and single tap works; double tap works too, but without preventing the single tap from working. So, on a double tap, both the single tap action handler and the double tap action handler are called.

If that isn't what we want, we don't have to use delayed performance, as we did earlier. Instead, we can create a *dependency* between one gesture recognizer and another, telling the first to suspend judgement until the second has decided whether this is its gesture, by sending the first the `requireGestureRecognizerToFail:` message. This message doesn't mean "force this recognizer to fail"; it means, "you can't succeed until this recognizer fails."

So our view is now configured as follows:

```
UITapGestureRecognizer* t2 = [[UITapGestureRecognizer alloc]
                           initWithTarget:self
                           action:@selector(doubleTap)];
t2.numberOfTapsRequired = 2;
[v addGestureRecognizer:t2];

UITapGestureRecognizer* t1 = [[UITapGestureRecognizer alloc]
                           initWithTarget:self
                           action:@selector(singleTap)];
[t1 requireGestureRecognizerToFail:t2];
[v addGestureRecognizer:t1];

UIPanGestureRecognizer* p = [[UIPanGestureRecognizer alloc]
                           initWithTarget:self
                           action:@selector(dragging:)];
[v addGestureRecognizer:p];
```



Apple would prefer, if you’re going to have a view respond both to single tap and double tap, that you *not* make the former wait upon the latter (because this delays your response after the user single taps). Rather, they would like you to arrange things so that it doesn’t matter that you respond to a single tap that is the first tap of a double tap. This isn’t always feasible, of course; Apple’s own Mobile Safari is a clear counterexample.

Subclassing Gesture Recognizers

To subclass a built-in gesture recognizer subclass, you must do the following things:

- At the start of the implementation file, import `<UIKit UIGestureRecognizerSubclass.h>`. This file contains a category on `UIGestureRecognizer` that allows you to set the gesture recognizer’s state (which is otherwise read-only), along with declarations for the methods you may need to override.
- Override any `touches...` methods you need to (as if the gesture recognizer were a `UIResponder`); you will almost certainly call `super` so as to take advantage of the built-in behavior. In overriding a `touches...` method, you need to think like a gesture recognizer. As these methods are called, a gesture recognizer is setting its state; you must interact with that process.

To illustrate, we will subclass `UIPanGestureRecognizer` so as to implement a view that can be moved only horizontally or vertically. Our strategy will be to make *two* `UIPanGestureRecognizer` subclasses — one that allows only horizontal movement, and another that allows only vertical movement. They will make their recognition decisions in a mutually exclusive manner, so we can attach an instance of each to our view. This separates the decision-making logic in a gorgeously encapsulated object-oriented manner — a far cry from the spaghetti code we wrote earlier to do this same task.

I will show only the code for the horizontal drag gesture recognizer, because the vertical recognizer is symmetrically identical. We maintain just one instance variable, `origLoc`, which we will use once to determine whether the user’s initial movement is horizontal. We override `touchesBegan:withEvent:` to set our instance variable with the first touch’s location:

```
- (void) touchesBegan:(NSSet *)touches withEvent:(UIEvent *)event {
    self->origLoc = [[touches anyObject] locationInView:self.view.superview];
    [super touchesBegan: touches withEvent: event];
}
```

We then override `touchesMoved:withEvent:`; all the recognition logic is here. This method will be called for the first time with the state still at `Possible`. At that moment, we look to see if the user’s movement is more horizontal than vertical. If it isn’t, we set the state to `Failed`. But if it is, we just step back and let the superclass do its thing:

```

- (void) touchesMoved:(NSSet *)touches withEvent:(UIEvent *)event {
    if (self.state == UIGestureRecognizerStatePossible) {
        CGPoint loc = [[touches anyObject] locationInView:self.view.superview];
        CGFloat deltaX = fabs(loc.x - origLoc.x);
        CGFloat deltaY = fabs(loc.y - origLoc.y);
        if (deltaY >= deltaX)
            self.state = UIGestureRecognizerStateFailed;
    }
    [super touchesMoved: touches withEvent:event];
}

```

We now have a view that moves only if the user's initial gesture is horizontal. But that isn't the entirety of what we want; we want a view that, itself, moves horizontally only. To implement this, we'll simply lie to our client about where the user's finger is, by overriding `translationInView:`:

```

- (CGPoint)translationInView:(UIView *)v {
    CGPoint proposedTranslation = [super translationInView:v];
    proposedTranslation.y = 0;
    return proposedTranslation;
}

```

That example was simple, because we subclassed a fully functional built-in `UIGestureRecognizer` subclass. If you were to write your own `UIGestureRecognizer` subclass entirely from scratch, there would be more work to do:

- You should definitely implement all four `touches...` handlers. Their job, at a minimum, is to advance the gesture recognizer through the canonical progression of its states. When the first touch arrives at a gesture recognizer, its state will be Possible; you never explicitly set the recognizer's state to Possible yourself. As soon as you know this can't be our gesture, you set the state to Failed (Apple says that a gesture recognizer should "fail early, fail often"). If the gesture gets past all the failure tests, you set the state instead either to Ended (for a discrete gesture) or to Began (for a continuous gesture); if Began, then you might set it to Changed, and ultimately you must set it to Ended. Action messages will be sent automatically at the appropriate moments.
- You should probably implement `reset`. This is called after you reach the end of the progression of states to notify you that the gesture recognizer's state is about to be set back to Possible; it is your chance to return your state machine to its starting configuration (resetting instance variables, for example).

Keep in mind that your gesture recognizer might stop receiving touches without notice. Just because it gets a `touchesBegan:withEvent:` call for a particular touch doesn't mean it will ever get `touchesEnded:withEvent:` for that touch. If your gesture recognizer fails to recognize its gesture, either because it declares failure or because it is still in the Possible state when another gesture recognizer recognizes, it won't get any more `touches...` calls for any of the touches that were being sent to it. This is why `reset` is so important; it's the one reliable signal that it's time to clean up and get ready to receive the beginning of another possible gesture.

Gesture Recognizer Delegate

A gesture recognizer can have a delegate, which can perform two types of task:

Block a gesture recognizer's operation

`gestureRecognizerShouldBegin`: is sent to the delegate before the gesture recognizer passes out of the Possible state; return NO to force the gesture recognizer to proceed to the Failed state.

`gestureRecognizer:shouldReceiveTouch`: is sent to the delegate before a touch is sent to the gesture recognizer's `touchesBegan:...` method; return NO to prevent that touch from ever being sent to the gesture recognizer.

Mediate simultaneous gesture recognition

When a gesture recognizer is about to declare that it recognizes its gesture, `gestureRecognizer:shouldRecognizeSimultaneouslyWithGestureRecognizer`: is sent to the delegate of that gesture recognizer, if this declaration would force the failure of another gesture recognizer, and to the delegate of a gesture recognizer whose failure would be forced. Return YES to prevent that failure, thus allowing both gesture recognizers to operate simultaneously. For example, a view could respond to both a two-fingered pinch and a two-fingered pan, the one applying a scale transform, the other changing the view's center.

As an artificial example, we will use delegate messages to combine a `UILongPressGestureRecognizer` and a `UIPanGestureRecognizer`, as follows: the user must perform a tap-and-a-half (tap and hold) to "get the view's attention," which we will indicate by a pulsing animation on the view; then (and only then) the user can drag the view. The example is artificial because a long press gesture recognizer can do all the work itself: its `Changed` state indicates a drag. However, the example illustrates the principle of using a gesture recognizer delegate.

As we create our gesture recognizers, we'll keep a reference to the `UILongPressGestureRecognizer`, and we'll make ourself the `UIPanGestureRecognizer`'s delegate:

```
UIPanGestureRecognizer* p = [[UIPanGestureRecognizer alloc]
    initWithTarget:self
    action:@selector(panning:)];
UILongPressGestureRecognizer* lp = [[UILongPressGestureRecognizer alloc]
    initWithTarget:self
    action:@selector(longPress:)];
lp.numberOfTapsRequired = 1;
[view addGestureRecognizer:p];
[view addGestureRecognizer:lp];
self.longPresser = lp;
p.delegate = self;
```

In keeping with encapsulation, the `UILongPressGestureRecognizer`'s handler will take care of starting and stopping the animation, and the `UIPanGestureRecognizer`'s handler will take care of the drag in the familiar manner:

```

- (void) longPress: (UILongPressGestureRecognizer*) lp {
    if (lp.state == UIGestureRecognizerStateBegan) {
        CABasicAnimation* anim =
            [CABasicAnimation animationWithKeyPath: @"transform"];
        anim.toValue =
            [NSValue valueWithCATransform3D:CATransform3DMakeScale(1.1, 1.1, 1)];
        anim.fromValue =
            [NSValue valueWithCATransform3D:CATransform3DIdentity];
        anim.repeatCount = HUGE_VALF;
        anim.autoreverses = YES;
        [lp.view.layer addAnimation:anim forKey:nil];
    }
    if (lp.state == UIGestureRecognizerStateChanged || 
        lp.state == UIGestureRecognizerStateCancelled) {
        [lp.view.layer removeAllAnimations];
    }
}

- (void) panning: (UIPanGestureRecognizer*) p {
    UIView* vv = p.view;
    if (p.state == UIGestureRecognizerStateBegan)
        self->origC = vv.center;
    CGPoint delta = [p translationInView: v.superview];
    CGPoint c = self->origC;
    c.x += delta.x; c.y += delta.y;
    vv.center = c;
}

```

Now for the delegate methods. We are the UIPanGestureRecognizer's delegate. If the UIPanGestureRecognizer tries to declare success while the UILongPressGestureRecognizer's state is Failed or still at Possible, we prevent it. If the UILongPressGestureRecognizer succeeds, we permit the UIPanGestureRecognizer to operate as well:

```

- (BOOL) gestureRecognizerShouldBegin: (UIGestureRecognizer*) g {
    if (self.longPresser.state == UIGestureRecognizerStatePossible ||
        self.longPresser.state == UIGestureRecognizerStateFailed)
        return NO;
    return YES;
}

- (BOOL)gestureRecognizer: (UIGestureRecognizer*) g1
    shouldRecognizeSimultaneouslyWithGestureRecognizer: (UIGestureRecognizer*) g2 {
    return YES;
}

```

The result is that the view can be dragged only if it is pulsing; in effect, what we've done is to compensate, using delegate methods, for the fact that UIGestureRecognizer has no `requireGestureRecognizerToSucceed:` method.

As I said earlier, the example is a bit artificial, because a UILongPressGestureRecognizer can implement draggability all on its own. It lacks the convenient `translationInView:` method, but we know how to work around that. So here, for completeness, is the same behavior implemented using a single gesture recognizer and a single handler:

```

- (void) longPress: (UILongPressGestureRecognizer*) lp {
    UIView* vv = lp.view;
    if (lp.state == UIGestureRecognizerStateBegan) {
        CABasicAnimation* anim =
            [CABasicAnimation animationWithKeyPath: @"transform"];
        anim.toValue =
            [NSValue valueWithCATransform3D:CATransform3DMakeScale(1.1, 1.1, 1)];
        anim.fromValue =
            [NSValue valueWithCATransform3D:CATransform3DIdentity];
        anim.repeatCount = HUGE_VALF;
        anim.autoreverses = YES;
        [vv.layer addAnimation:anim forKey:nil];
        self->origOffset =
            CGPointMake(CGRectGetMidX(vv.bounds) - [lp locationInView:vv].x,
                       CGRectGetMidY(vv.bounds) - [lp locationInView:vv].y);
    }
    if (lp.state == UIGestureRecognizerStateChanged) {
        CGPoint c = [lp locationInView: vv.superview];
        c.x += self->origOffset.x;
        c.y += self->origOffset.y;
        vv.center = c;
    }
    if (lp.state == UIGestureRecognizerStateChanged ||
        lp.state == UIGestureRecognizerStateCancelled) {
        [vv.layer removeAllAnimations];
    }
}

```

If you are subclassing a gesture recognizer class, you can incorporate delegate-like behavior into the subclass. By overriding `canPreventGestureRecognizer:` and `canBePreventedByGestureRecognizer:`, you can mediate simultaneous gesture recognition at the class level. The built-in gesture recognizer subclasses already do this; that is why, for example, a UITapGestureRecognizer whose `numberOfTapsRequired` is 1 does not, by recognizing its gesture, cause the failure of a UITapGestureRecognizer whose `numberOfTapsRequired` is 2.

You can also, in a gesture recognizer subclass, send `ignoreTouch:forEvent:` directly to a gesture recognizer (typically, to `self`). This has the same effect as the delegate method `gestureRecognizer:shouldReceiveTouch:` returning NO, blocking delivery of that touch to the gesture recognizer for as long as it exists. For example, if you're in the middle of an already recognized gesture and a new touch arrives, you might well elect to ignore it.

Touch Delivery

Let's now return to the very beginning of the touch reporting process, when the system sends the app a UIEvent containing touches, and tease apart in full detail the entire procedure by which a touch is delivered to views and gesture recognizers:

1. Whenever a new touch appears, the application calls the UIView instance method `hitTest:withEvent:` on the window, which returns the view (called, appropriately,

the *hit-test view*) that will be permanently associated with this touch. This method uses the `UIView` instance method `pointInside:withEvent:` along with `hitTest:withEvent:` recursively down the view hierarchy to find the frontmost view containing the touch's location and capable of receiving a touch. The logic of how a view's `userInteractionEnabled`, `hidden`, and `opacity` affect its touchability is implemented at this stage.

2. Each time the touch situation changes, the application calls its own `sendEvent:`, which in turn calls the window's `sendEvent:`. The window delivers each of an event's touches by calling the appropriate `touches...` method(s), as follows:
 - a. As a touch first appears, it is initially delivered to the hit-test view's swarm of gesture recognizers. It is then also delivered to that view. The logic of withholding touches in obedience to `multipleTouchEnabled` and `exclusiveTouch` is also implemented at this stage. For example, additional touches won't be delivered to a view if that view currently has a touch and has `multipleTouchEnabled` set to NO.
 - b. If a gesture is recognized by a gesture recognizer, then for any touch associated with this gesture recognizer:
 - i. `touchesCancelled:forEvent:` is sent to the touch's view, and the touch is no longer delivered to its view.
 - ii. If that touch was associated with any other gesture recognizer, that gesture recognizer is forced to fail.
 - c. If a gesture recognizer fails, either because it declares failure or because it is forced to fail, its touches are no longer delivered to it, but (except as already specified) they continue to be delivered to their view.
 - d. If a touch would be delivered to a view, but that view does not respond to the appropriate `touches...` method, a responder further up the responder chain ([Chapter 11](#)) is sought that does respond to it, and the touch is delivered there.

The rest of this chapter elaborates on each stage of this standard procedure, nearly every bit of which can be customized to some extent.

Hit-Testing

Hit-testing is the determination of what view the user touched. View hit-testing uses the `UIView` instance method `hitTest:withEvent:`, which returns either a view (the hit-test view) or nil. The idea is to find the frontmost view containing the touch point. This method uses an elegant recursive algorithm, as follows:

1. A view's `hitTest:withEvent:` first calls the same method on its own subviews, if it has any, because a subview is considered to be in front of its superview. The subviews are queried in reverse order, because that's front-to-back order ([Chapter 14](#)): thus, if two sibling views overlap, the one in front reports the hit first.

2. If, as a view hit-tests its subviews, any of those subviews responds by returning a view, it stops querying its subviews and immediately returns the view that was returned to it. Thus, the very first view to declare itself the hit-test view immediately percolates all the way to the top of call chain and *is* the hit-test view.
3. If, on the other hand, a view has no subviews, or if all of its subviews return nil (indicating that neither they nor their subviews was hit), then the view calls `pointInside:withEvent:` on itself. If this call reveals that the touch was inside this view, the view returns itself, declaring itself the hit-test view; otherwise it returns nil.

No problem arises if a view has a transform, because `pointInside:withEvent:` takes the transform into account. That's why a rotated button continues to work correctly.

It is also up to `hitTest:withEvent:` to implement the logic of touch restrictions exclusive to a view. If a view's `userInteractionEnabled` is NO, or its `hidden` is YES, or its `opacity` is close to 0.0, it returns nil without hit-testing any of its subviews and without calling `pointInside:withEvent:`. Thus these restrictions do not, of themselves, exclude a view from being hit-tested; on the contrary, they operate precisely by modifying a view's hit-test result.

However, hit-testing knows nothing about `multipleTouchEnabled` (because its behavior involves multiple touches) or `exclusiveTouch` (because its behavior involves multiple views). The logic of obedience to these properties is implemented at a later stage of the story.

You can use hit-testing yourself at any moment where it might prove useful. In calling `hitTest:withEvent:`, supply a point *in the coordinates of the view to which the message is sent*. The second parameter can be nil if you have no event.

For example, suppose we have a `UIView` with two `UIImageView` subviews. We want to detect a tap in either `UIImageView`, but we want to handle this at the level of the `UIView`. We can attach a `UITapGestureRecognizer` to the `UIView`, but how will we know which subview, if any, the tap was in?

Our first step must be to set `userInteractionEnabled` to YES for both `UIImageView`s. (This step is crucial; `UIImageView` is one of the few built-in view classes where this is NO by default, and a view whose `userInteractionEnabled` is NO won't normally be the result of a call to `hitTest:withEvent:`.) Now, when our gesture recognizer's action handler is called, the view can use hit-testing to determine where the tap was:

```
CGPoint p = [g locationOfTouch:0 inView:self]; // g is the gesture recognizer  
UIView* v = [self hitTest:p withEvent:nil];
```

You can also override `hitTest:withEvent:` in a view subclass, to alter its results during touch delivery, thus customizing the touch delivery mechanism. I call this *hit-test munging*. Hit-test munging can be used selectively as a way of turning user interaction on or off in an area of the interface. In this way, some unusual effects can be produced.

For example, an important use of hit-test munging is to permit the touching of parts of subviews outside the bounds of their superview. If a view's `clipsToBounds` is NO, a paradox arises: the user can *see* the regions of its subviews that are outside its bounds, but the user can't *touch* them. This can be confusing and seems wrong. The solution is for the view to override `hitTest:withEvent:` as follows:

```
- (UIView *)hitTest:(CGPoint)point withEvent:(UIEvent *)event {
    UIView* result = [super hitTest:point withEvent:event];
    if (result)
        return result;
    for (UIView* sub in [self.subviews reverseObjectEnumerator]) {
        CGPoint pt = [self convertPoint:point toView:sub];
        result = [sub hitTest:pt withEvent:event];
        if (result)
            return result;
    }
    return nil;
}
```

Here are some further possible uses of hit-test munging, just to stimulate your imagination:

- If a superview contains a UIButton but doesn't return that UIButton from `hitTest:withEvent:`, that button can't be tapped.
- You might override `hitTest:withEvent:` to return the result from `super` most of the time, but to return `self` under certain conditions, effectively making all subviews untouchable without making the superview itself untouchable (as setting its `userInteractionEnabled` to NO would do).
- A view whose `userInteractionEnabled` is NO can break the normal rules and return itself from hit-testing and can thus end up as the hit-test view.

Hit-testing for layers

There is also hit-testing for layers. It doesn't happen automatically, as part of `sendEvent:` or anything else; it's up to you. It's just a convenient way of finding out which layer would receive a touch at a point, if layers received touches. To hit-test layers, call `hitTest:` on a layer, with a point *in superlayer coordinates*.

Keep in mind, though, that layers do *not* receive touches. A touch is reported to a view, not a layer. A layer, except insofar as it is a view's underlying layer and gets touch reporting because of its view, is completely untouchable; from the point of view of touches and touch reporting, it's as if the layer weren't on the screen at all. No matter where a layer may appear to be, a touch falls right through the layer to whatever view is behind it.

In the case of the layer that is a view's underlying layer, you don't need hit-testing. It is the view's drawing; where it appears is where the view is. So a touch in that layer is equivalent to a touch in its view. Indeed, one might say that this is what views are actually for: to provide layers with touchability.

The only layers on which you'd need special hit-testing, then, would presumably be layers that are not themselves any view's underlying layer, because those are the only ones you don't find out about by normal view hit-testing. However, all layers, including a layer that is its view's underlying layer, are part of the layer hierarchy, and can participate in layer hit-testing. So the most comprehensive way to hit-test layers is to start with the topmost layer, the window's layer. In this example, we subclass UIWindow and override its `hitTest:withEvent:` so as to get layer hit-testing every time there is view hit-testing:

```
- (UIView*) hitTest:(CGPoint)point withEvent:(UIEvent *)event {
    CALayer* lay = [self.layer hitTest:point];
    // ... possibly do something with that information ...
    return [super hitTest:point withEvent:event];
}
```

Because this is the window, the view hit-test point works as the layer hit-test point; window bounds are screen bounds. But usually you'll have to convert to superlayer coordinates. In this example, we return to the CompassView developed in [Chapter 16](#), in which all the parts of the compass are layers; we want to know whether the user tapped on the arrow layer. For simplicity, we've given the CompassView a UITapGestureRecognizer, and this is its action handler, in the CompassView itself. We convert to our superview's coordinates, because these are also our layer's superlayer coordinates:

```
// self is the CompassView
CGPoint p = [t locationOfTouch: 0 inView: self.superview];
CALayer* hit = [self.layer hitTest:p];
if (hit == ((CompassLayer*)self.layer).arrow) // ...
```

Layer hit-testing works by calling `containsPoint:`. However, `containsPoint:` takes a point in the layer's coordinates, so to hand it a point that arrives through `hitTest:` you must first convert from superlayer coordinates:

```
BOOL hit = [lay containsPoint: [lay convertPoint:point fromLayer:lay.superlayer]];
```

Layer hit-testing knows nothing of the restrictions on touch delivery; it just reports on every sublayer, even those whose view has `userInteractionEnabled` set to NO.



The documentation warns that `hitTest:` must not be called on a CA-TransformLayer.

Hit-testing for drawings

The preceding example (letting the user tap on the compass arrow) worked, but we might complain that it is reporting a hit on the arrow even if the hit misses the *drawing* of the arrow. That's true for view hit-testing as well. A hit is reported if we are within the view or layer as a whole; hit-testing knows nothing of drawing, transparent areas, and so forth.

If you know how the region is drawn and can reproduce the edge of that drawing as a CGPath, you can test whether a point is inside it with `CGPathContainsPoint`. So, for a layer, you could override `hitTest` along these lines:

```
- (CALayer*) hitTest:(CGPoint)p {
    CGPoint pt = [self convertPoint:p fromLayer:self.superlayer];
    CGMutablePathRef path = CGPathCreateMutable();
    // ... draw path here ...
    CALayer* result = CGPathContainsPoint(path, NULL, pt, true) ? self : nil;
    CGPathRelease(path);
    return result;
}
```

Alternatively, it might be the case that if a pixel of the drawing is transparent, it's outside the drawn region. Unfortunately there's no way to ask a drawing (or a view, or a layer) for the color of a pixel; you have to make a bitmap and copy the drawing into it, and then ask the bitmap for the color of a pixel. If you can reproduce the content as an image, and all you care about is transparency, you can make a one-pixel alpha-only bitmap, draw the image in such a way that the pixel we want to test is the pixel drawn into the bitmap, and examine the transparency of the resulting pixel:

```
// assume im is a UIImage, point is the CGPoint to test
unsigned char pixel[1] = {0};
CGContextRef context = CGBitmapContextCreate(pixel,
                                              1, 1, 8, 1, NULL,
                                              kCGImageAlphaOnly);

UIGraphicsPushContext(context);
[im drawAtPoint:CGPointMake(-point.x, -point.y)];
UIGraphicsPopContext();
CGContextRelease(context);
CGFloat alpha = pixel[0]/255.0;
BOOL transparent = alpha < 0.01;
```

However, there can be complications; for example, there may not be a one-to-one relationship between the pixels of the underlying drawing and the points of the drawing as portrayed on the screen (because the drawing is stretched, for example). It's a tricky problem, but in many cases, the CALayer method `renderInContext:` can be helpful here. This method allows you to copy a layer's actual drawing into a context of your choice. If that context is, say, an image context created with `UIGraphicsBeginImageContextWithOptions`, you can now use the resulting image as `im` in the code above.

Hit-testing during animation

If user interaction is allowed during an animation that moves a view from one place to another, then if the user taps on the animated view, the tap might mysteriously fail because the view in the model layer is elsewhere; conversely, the user might accidentally tap where the view actually is in the model layer, and the tap will hit the animated view even though it appears to be elsewhere. If the position of a view or layer is being animated and you want the user to be able to tap on it, therefore, you'll need to hit-test the presentation layer (see [Chapter 17](#)).

In this simple example, we have a superview containing a subview. To allow the user to tap on the subview even when it is being animated, we interfere with hit-testing in the superview:

```
- (UIView*) hitTest:(CGPoint)point withEvent:(UIEvent *)event {
    // v is the animated subview
    CALayer* lay = [v.layer presentationLayer];
    CALayer* hitLayer = [lay hitTest: point];
    if (hitLayer == lay)
        return v;
    UIView* hitView = [super hitTest:point withEvent:event];
    if (hitView == v)
        return self;
    return hitView;
}
```

If the user taps outside the presentation layer, we cannot simply call `super`, because the user might tap at the spot to which the subview has in reality already moved (behind the “animation movie”), in which case `super` will report that it hit the subview. So if `super` does report this, we return `self` (assuming that we are what’s behind the animated subview at its new location).

However, as Apple puts it in the WWDC 2011 videos, the animated view “swallows the touch.” For example, suppose the view in motion is a button. Although our hit-test munging makes it possible for the user to tap the button as it is being animated, and although the user sees the button highlight in response, the button’s action message is not sent in response to this highlighting until the animation is over (or nearly over). This behavior seems unfortunate, but it’s generally possible to work around it (for instance, with a gesture recognizer).

Initial Touch Event Delivery

When the touch situation changes, an event containing all touches is handed to the `UIApplication` instance by calling its `sendEvent:`, and the `UIApplication` in turn hands it to the relevant `UIWindow` by calling *its* `sendEvent:`. The `UIWindow` then performs the complicated logic of examining, for every touch, the hit-test view and its superviews and their gesture recognizers and deciding which of them should be sent a `touches...` message, and does so.

These are delicate and crucial maneuvers, and you wouldn’t want to lame your application by interfering with them. Nevertheless, you can override `sendEvent:` in a subclass, and there are situations where you might wish to do so. This is just about the *only* case in which you might subclass `UIApplication`; if you do, remember to change the third argument in the call to `UIApplicationMain` in your `main.m` file to the string name of your `UIApplication` subclass so that your subclass is used to generate the app’s singleton `UIApplication` instance. If you subclass `UIWindow`, remember to change the window’s class in the app delegate code that instantiates the window.

Now that gesture recognizers exist, it is unlikely that you will need to resort to such measures. A typical case, in the past, was that you needed to detect touches directed to an object of some built-in interface class in a way that subclassing it wouldn't permit. For example, you want to know when the user swipes a UIWebView; you're not allowed to subclass UIWebView, and in any case it eats the touch. The solution used to be to subclass UIWindow and override `sendEvent:`; you would then work out whether this was a swipe on the UIWebView and respond accordingly, or else call `super`. Now, however, you can attach a UISwipeGestureRecognizer to the UIWebView.

Gesture Recognizer and View

When a touch first appears and is delivered to a gesture recognizer, it is also delivered to its hit-test view, the same `touches...` method being called on both. This comes as a surprise to beginners, but it is the most reasonable approach, as it means that touch interpretation by a view isn't jettisoned just because gesture recognizers are in the picture. Later on in the multitouch sequence, if all the gesture recognizers in a view's swarm declare failure to recognize their gesture, that view's internal touch interpretation just proceeds as if gesture recognizers had never been invented.

However, if a gesture recognizer in a view's swarm recognizes its gesture, that view is sent `touchesCancelled:withEvent:` for any touches that went to that gesture recognizer and were hit-tested to that view, and subsequently the view no longer receives those touches.

This behavior can be changed by setting a gesture recognizer's `cancelsTouchesInView` property to NO. If this is the case for every gesture recognizer in a view's swarm, the view will receive touch events more or less as if no gesture recognizers were in the picture. Making this change, however, alters delivery logic rather drastically; it seems unlikely that you'd want to do that.

If a gesture recognizer happens to be ignoring a touch (because it was told to do so by `ignoreTouch:forEvent:`), then `touchesCancelled:withEvent:` won't be sent to the view for that touch when the gesture recognizer recognizes its gesture. Thus, a gesture recognizer's ignoring a touch is the same as simply letting it fall through to the view, as if the gesture recognizer weren't there.

Gesture recognizers can also *delay* the delivery of touches to a view, and by default they do. The UIGestureRecognizer property `delaiesTouchesEnded` is YES by default, meaning that when a touch reaches `UITouchPhaseEnded` and the gesture recognizer's `touchesEnded:withEvent:` is called, if the gesture recognizer is still allowing touches to be delivered to the view because its state is still Possible, it doesn't deliver this touch until it has resolved the gesture. When it does, either it will recognize the gesture, in which case the view will have `touchesCancelled:withEvent:` called instead (as already explained), or it will declare failure and *now* the view will have `touchesEnded:withEvent:` called.

The reason for this behavior is most obvious with a gesture where multiple taps are required. The first tap ends, but this is insufficient for the gesture recognizer to declare success or failure, so it withholds that touch from the view. In this way, the gesture recognizer gets the proper priority. In particular, if there is a second tap, the gesture recognizer should succeed and send `touchesCancelled:withEvent:` to the view — but it can't do that if the view has already been sent `touchesEnded:withEvent:`.

It is also possible to delay the entire suite of `touches...` methods from being called on a view, by setting a gesture recognizer's `delaysTouchesBegan` property to YES. Again, this delay would be until the gesture recognizer can resolve the gesture: either it will recognize it, in which case the view will have `touchesCancelled:withEvent:` called, or it will declare failure, in which case the view will receive `touchesBegan:withEvent:` plus any further `touches...` calls that were withheld — except that it will receive *at most* one `touchesMoved:withEvent:` call, the last one, because if a lot of these were withheld, to queue them all up and send them all at once now would be simply insane.

It is unlikely that you'll change a gesture recognizer's `delaysTouchesBegan` property to YES, however. You might do so, for example, if you have an elaborate touch analysis within a view that simply cannot operate simultaneously with a gesture recognizer, but this is improbable, and the latency involved may look strange to your user.

When touches are delayed and then delivered, what's delivered is the original touch with the original event, which still have their original timestamps. Because of the delay, these timestamps may differ significantly from now. For this reason (and many others), Apple warns that touch analysis that is concerned with timing should always look at the timestamp, not the clock.

Touch Exclusion Logic

It is up to the `UIWindow`'s `sendEvent:` to implement the logic of `multipleTouchEnabled` and `exclusiveTouch`.

If a new touch is hit-tested to a view whose `multipleTouchEnabled` is NO and which already has an existing touch hit-tested to it, then `sendEvent:` never delivers the new touch to that view. However, that touch *is* delivered to the view's swarm of gesture recognizers.

Similarly, if there's an `exclusiveTouch` view in the window, then `sendEvent:` must decide whether a particular touch should be delivered, as already described. If a touch is not delivered to a view because of `exclusiveTouch` restrictions, it is not delivered to its swarm of gesture recognizers either. (This behavior with regard to gesture recognizers has changed in a confusing and possibly buggy way from system to system, but I believe I'm describing it correctly for iOS 5. The statement in Apple's *SimpleGestureRecognizers* sample code that "Recognizers ignore the exclusive touch setting for views" now appears to be false.)

Recognition

When a gesture recognizer recognizes its gesture, everything changes. As we've already seen, the touches for this gesture recognizer are sent to their hit-test views as a `touchesCancelled:forEvent:` message, and then no longer arrive at those views (unless the gesture recognizer's `cancelsTouchesInView` is NO). Moreover, all other gesture recognizers pending with regard to these touches are made to fail, and then are no longer sent the touches they were receiving either.

If the very same event would cause more than one gesture recognizer to recognize, there's an algorithm for picking the one that will succeed and make the others fail: a gesture recognizer lower down the view hierarchy (closer to the hit-test view) prevails over one higher up the hierarchy, and a gesture recognizer more recently added to its view prevails over one less recently added.

There are various means for modifying this "first past the post, winner takes all" behavior. One is by telling a gesture recognizer, in effect, that being first isn't good enough:

- `requireGestureRecognizerToFail`: constitutes a dependency order, possibly causing the gesture recognizer to which it is sent to be put on hold when it tries to transition from the Possible state to the Began (continuous) or Ended (discrete) state; only if a certain other gesture recognizer fails is this one permitted to perform that transition. (So, "require to fail" means "*you* cannot succeed without this other's failure.")

Apple says that in a dependency like this, the gesture recognizer that fails first is not sent `reset` (and won't receive any touches) until the second finishes its state sequence and is sent `reset`, so that they resume recognizing together.

- The delegate method `gestureRecognizerShouldBegin`: , by returning NO, turns success into failure; at the moment when the gesture recognizer is about to declare that it recognizes its gesture, the delegate is telling it to fail instead.

Another approach is to permit simultaneous recognition; a gesture recognizer succeeds, but some other gesture recognizer is *not* forced to fail. There are two ways to achieve this:

- A subclass can implement `canPreventGestureRecognizer`: or `canBePreventedByGestureRecognizer`: (or both). Here, "prevent" means "by succeeding, you force failure upon this other," and "be prevented" means "by succeeding, this other forces failure upon you."

These two methods work together as follows. `canPreventGestureRecognizer`: is called first; if it returns NO, that's the end of the story for that gesture recognizer, and `canPreventGestureRecognizer`: is called on the other gesture recognizer. But if `canPreventGestureRecognizer`: returns YES when it is first called, the other gesture recognizer is sent `canBePreventedByGestureRecognizer`: . If it returns YES, that's the end of the story; if it returns NO, the process starts over the other way around,

sending `canPreventGestureRecognizer`: to the second gesture recognizer, and so forth. In this way, conflicting answers are resolved without the device exploding: prevention is regarded as exceptional (even though it is in fact the norm) and will happen only if it is acquiesced to by everyone involved.

- The delegate method `gestureRecognizer:shouldRecognizeSimultaneouslyWithGestureRecognizer`: can return YES to permit one gesture recognizer to succeed without forcing the other to fail.

Touches and the Responder Chain

A UIView is a responder, and participates in the responder chain ([Chapter 11](#)). In particular, if a touch is to be delivered to a UIView (because, for example, it's the hit-test view) and that view doesn't implement the relevant `touches...` method, a walk up the responder chain is performed, looking for a responder that *does* implement it; if such a responder is found, the touch is delivered to that responder. Moreover, the default implementation of the `touches...` methods — the behavior that you get if you call `super` — is to perform the same walk up the responder chain, starting with the next responder in the chain.

The relationship between touch delivery and the responder chain can be useful, but you must be careful not to allow it to develop into an incoherency. For example, if `touchesBegan:withEvent:` is implemented in a superview but not in a subview, then a touch to the subview will result in the superview's `touchesBegan:withEvent:` being called, with the first parameter (the `touches`) containing a touch whose `view` is the subview. But most UIView implementations of the `touches...` methods rely upon the assumption that the first parameter consists of all and only touches whose `view` is `self`; built-in UIView subclasses certainly assume this.

Again, if `touchesBegan:withEvent:` is implemented in both a superview and a subview, and you call `super` in the subview's implementation, passing along the same arguments that came in, then the same touch delivered to the subview will trigger both the subview's `touchesBegan:withEvent:` and the superview's `touchesBegan:withEvent:` (and once again the first parameter to the superview's `touchesBegan:withEvent:` will contain a touch whose `view` is the subview).

The solution is to behave rationally, as follows:

- If all the responders in the affected part of the responder chain are instances of your own subclass of UIView itself or of your own subclass of UIViewController, you will generally want to follow the simplest possible rule: implement *all* the `touches...` events together in one class, so that touches arrive at an instance either because it was the hit-test view or because it is up the responder chain from the hit-test view, and do *not* call `super` in any of them. In this way, “the buck stops here” — the touch handling for this object or for objects below it in the responder chain is bottlenecked into one well-defined place.

- If you subclass a built-in UIView subclass and you override its touch handling, you don't have to override every single `touches...` event, but you *do* need to call `super` so that the built-in touch handling can occur.
- Don't allow touches to arrive from lower down the responder chain at an instance of a built-in UIView subclass that implements built-in touch handling, because such a class is completely unprepared for the first parameter of a `touches...` method containing a touch not intended for itself. Judicious use of `userInteractionEnabled` or hit-test munging can be a big help here.

I'm not saying, however, that you have to block all touches from percolating up the responder chain; it's normal for unhandled touches to arrive at the UIWindow or UIApplication, for example, because these classes do not (by default) do any touch handling — so those touches will remain unhandled and will percolate right off the end of the responder chain, which is perfectly fine.

- Never call a `touches...` method directly (except to call `super`).



Apple's documentation has some discussion of a technique called *event forwarding* where you *do* call `touches...` methods directly. But you are far less likely to need this now that gesture recognizers exist, and it can be extremely tricky and even downright dangerous to implement, so I won't give an example here, and I suggest that you not use it.

PART V

Interface

The previous part of the book introduced views. This part of the book is about the particular kinds of view provided by the Cocoa framework — the built-in “widgets” with which you’ll construct an app’s interface. These are surprisingly few, but impressively powerful.

- [Chapter 19](#) is about view controllers. View controllers are a brilliant mechanism for allowing an entire interface to be replaced by another; this ability is especially crucial on the iPhone’s small screen. They are also the basis of an app’s ability to compensate when the user rotates the device. In real life, every app you write will probably have its interface managed by view controllers.
- [Chapter 20](#) is about scroll views, the iOS mechanism for letting the user scroll and zoom the interface.
- [Chapter 21](#) explains table views, an extremely important and powerful type of scroll view that lets the user navigate through any amount of data.
- [Chapter 22](#) is about two forms of interface unique to, and characteristic of, the iPad — popovers and split views.
- [Chapter 23](#) describes several ways of presenting text in an app’s interface — labels, text fields, text views, and text drawn manually with Core Text.
- [Chapter 24](#) discusses web views. A web view is an easy-to-use interface widget backed by the power of a full-fledged web browser. It can also be used to present a PDF and various other forms of data.
- [Chapter 25](#) describes all the remaining built-in iOS (UIKit) interface widgets.
- [Chapter 26](#) is about the forms of modal dialog that can appear in front of an app’s interface.

CHAPTER 19

View Controllers

An iOS app's interface is dynamic in a way that a Mac OS X application is not. In Mac OS X, an application's windows can be big, and there can be more than one of them, so there's room for lots of interface. With iOS, everything needs to fit on a single display consisting of a single window, which in the case of the iPhone is almost forbiddingly tiny. The iOS solution to this is to swap out interface and replace it with other interface, as needed. Thus, entire regions of interface material — often the entire contents of the screen — must come and go in an agile fashion that is understandable to the user. Animation is often used to emphasize and clarify the replacement of one view by another.

Management of this task resides in a *view controller*, an instance of `UIViewController`. Actually, a view controller is most likely to be an instance of a `UIViewController` subclass. The `UIViewController` class is designed to be subclassed. You are very unlikely to use a plain vanilla `UIViewController` object. You might write your own `UIViewController` subclass, or you might use a built-in `UIViewController` subclass such as `UINavigationController` or `UITabBarController`. On the other hand, you will generally *not* subclass a built-in `UIViewController` subclass such as `UINavigationController` or `UITabBarController`, unless specifically instructed by the documentation to do so. You'll either use a built-in `UIViewController` subclass as is, or you'll subclass `UIViewController` yourself. (The major exception is `UITableViewController`, a `UIViewController` subclass that you'll probably always subclass, as I'll explain in [Chapter 21](#).)

A view controller manages a single view (which can, of course, have subviews); its `view` property points to the view it manages. The view has no explicit pointer to the view controller that manages it, but a view controller is a `UIResponder` and is in the responder chain just above its view ([Chapter 11](#)), so it is the view's `nextResponder`.

The chief concepts involved in the use of view controllers are as follows:

Rotation

The user can rotate the device, and you might like the interface to rotate in response, to compensate. A window is effectively pinned to the physical display (window bounds are screen bounds and do not change), but a view can be given a

transform so that its top moves to the current top of the display. A `UIViewController` responds to device rotation by applying this transform.

Root view controller

Every real-life iOS app should have a single view controller that acts as the root view controller for the whole app. Its job is to supply the view that covers the entire window and acts as the superview for all other interface ([Chapter 7](#), [Chapter 14](#)). The user may never see or be conscious of the root view, as it may be completely covered by its subviews, but it still has an important function with regard to rotation of the device: it allows the entire interface to rotate in response to device rotation. If the user does see the root view, it is the highest view of the hierarchy that the user will ever see; there is nothing behind it except the window, and the user won't see the window except, perhaps, in a glimpse as view controllers are swapped along with animation of their views.

Prior to iOS 5 it was theoretically possible for an iOS app to lack a root view controller. In iOS 5, it's still theoretically possible, but it's strongly discouraged (the runtime issues a warning if the app launches without a root view controller: "Applications are expected to have a root view controller at the end of application launch"). That is why our Empty Window project ([Chapter 6](#) and following) was based on the Single View Application project template: this is the minimal Xcode 4.2 template that supplies a root view controller along with a nib containing its view. All code in this edition of the book that puts interface into a view assumes there is a root view controller, and puts that interface into the root view controller's view.

Parentage

A view controller can *contain* another view controller. The containing view controller is the *parent* of the contained view controller; the contained view controller is a *child* of the containing view controller. Replacement of one view with another often involves a parent view controller managing its children. For example, [Figure 19-1](#) shows the TidBITS News app displaying a conventional and typical iPhone interface, consisting of a list of story headlines and blurbs; if the user taps an entry in the list, the whole list will slide away to the left and the text of the actual story will slide in from the right. This is done by a parent view controller (a `UINavigationController`) adding a new child view controller.

In iOS 4 and before, only built-in view controllers such as `UITabBarController`, `UINavigationController`, and `UISplitViewController` could act as parent view controllers. In iOS 5, you are free to write your own view controller subclasses that act as parent view controllers.

Presentation

In iOS 4 and before, there was a notion of a *modal view controller*, whose view effectively replaced the entire interface by appearing on top of it. In iOS 5 this has evolved into a notion of a *presented* view controller. One view controller *presents*

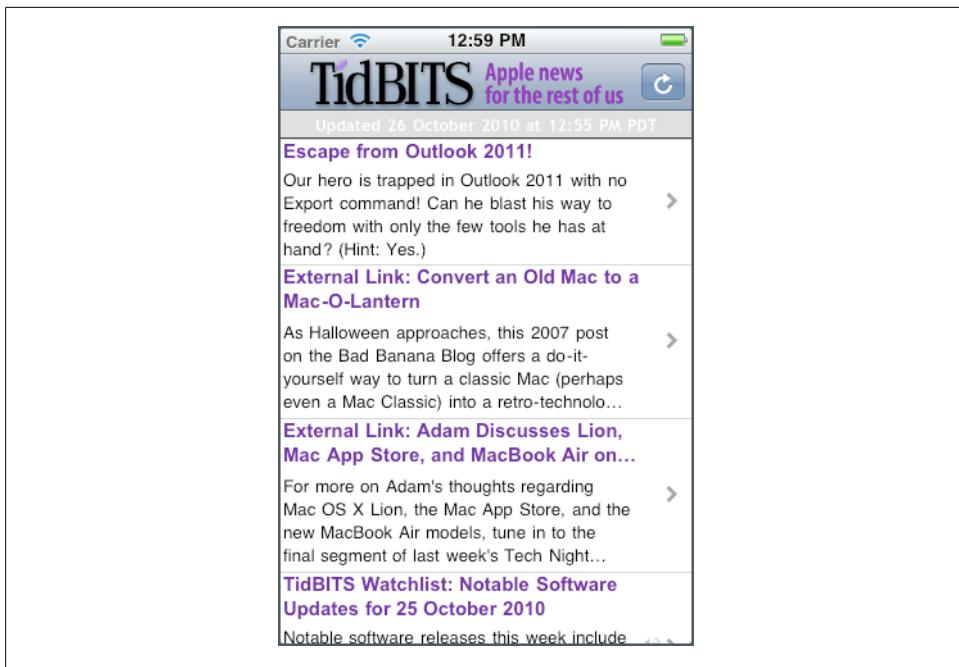


Figure 19-1. The TidBITS News app

another view controller; this means that the first view controller, the *presenting* view controller, remains in place, but the presented view controller's view is in front of it.

This relationship between view controllers is different from the parent–child relationship. In iOS 4 and before, a modal view controller was treated as a child of the presenting view controller, but in iOS 5 it is not. A presenting view controller is *not* the parent view controller of the view controller it presents — it is its presenting view controller.

Animation

The act of swapping views by manipulating child view controllers or presenting a view controller is very frequently accompanied by animation. It is the job of the parent view controller or presenting view controller to perform this animation. Certain animation types are built-in and conventional. For example, as I mentioned a moment ago, in Figure 19-1, tapping a story listing will cause the list to slide out to the left while the new view slides in from the right; this is the default behavior of a UINavigationController when it adds a new view controller and makes its view appear. Similarly, as a view controller is presented, the new view can slide in from below, flip into place like a piece of paper being flipped over, and so forth.

Customization

A view controller has properties and methods that are used to customize the interface and its behavior when its view is showing. For example, when a UINavigationController substitutes another view controller's view into its interface (by adding the view controller as a child), it looks for that view controller's `navigationItem.titleView` property — and puts it into the navigation bar at the top of the interface. That is how the TidBITS logo in [Figure 19-1](#) appears in the navigation bar — it's because it is a view controller's `navigationItem.titleView`. Similarly, if a view controller is to be presented, it has properties that allow it to dictate the style of animation that should take place as its view appears.

The View Controller Hierarchy

Because of containment and presentation, there is a hierarchy of view controllers. In a properly constructed iOS app, there should be exactly one root view controller, and it is the *only* view controller that has neither a parent view controller nor a presenting view controller. Any other view controller, if its view is to appear in the interface, must be a child view controller (of some parent view controller) or a presented view controller (of some presenting view controller).

At the same time, at any given moment, the actual views of the interface form a hierarchy dictated by and parallel to some portion of the view controller hierarchy. Every view visible in the interface owes its presence either to the fact that it is a view controller's view or to the fact that it is, at some depth, a subview of a view controller's view. Moreover, a child view controller's view is, at some depth, its parent view controller's view's subview.

Most important, the place of a view controller's view in the view hierarchy will usually be *automatic*, by virtue of the view controller's place in the view controller hierarchy. You will almost never need to put a UIViewController's view into the view hierarchy manually (although there is a clear-cut type of exception to that rule, which I'll talk about later).

For example, in [Figure 19-1](#), we see three interface elements (from top to bottom):

- The navigation bar, containing the TidBITS logo.
- Some text, which is actually a UILabel, stating when the list of stories was last updated.
- The list of stories, which is actually a UITableView.

I will describe how all of this comes to appear on the screen through the view controller hierarchy and the view hierarchy ([Figure 19-2](#)). The app's root view controller is a UINavigationController; the UINavigationController's view, which is never seen in isolation, is the window's sole immediate subview (the root view), and the navigation bar is a subview of that view. The UINavigationController contains a second UIView-

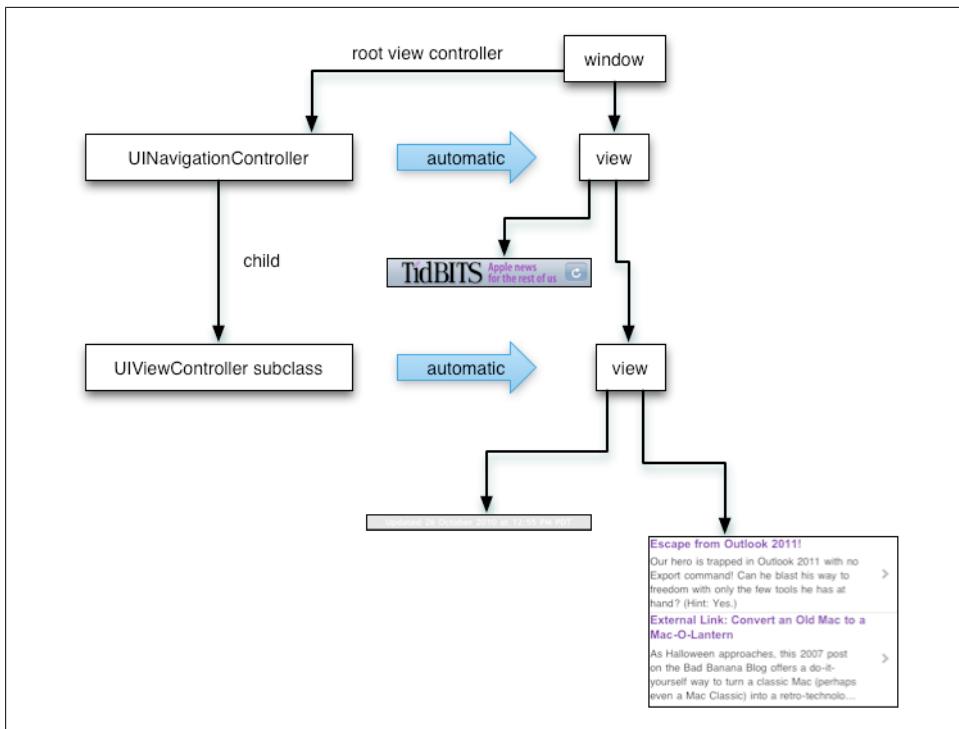


Figure 19-2. The TidBITS News app’s initial view controller and view hierarchy

Controller — a parent–child relationship. The child is a custom UIViewController subclass; its view is what occupies the rest of the window, as anotherSubview of the UINavigationController’s view. That view contains the UILabel and the UITableView as subviews. This architecture means that when the user taps a story listing in the UITableView, the whole label-and-table complex will slide out, to be replaced by the view of a different UIViewController, while the navigation bar stays.

In Figure 19-2, notice the word “automatic” in the two large right-pointing arrows associating a view controller with its view. This is intended to tell you *how the view controller’s view became part of the view hierarchy*. The UINavigationController’s view became the window’s subview *automatically*, by virtue of the UINavigationController being the window’s rootViewController. The UIViewController’s view became the UINavigationController’s view’s second subview *automatically*, by virtue of the UIViewController being the UINavigationController’s child.

Now, as I said a moment ago, there is an exception to this rule about views taking their place in the view hierarchy automatically — namely, when you write *your own parent view controller class*. In that case, you will need to put a child view controller’s view into the interface *manually*, as a subview (at some level) of the parent view controller’s view, if you want it to appear in the interface. (Conversely, you should *not* put a view controller’s view into the interface manually under any *other* circumstances.)

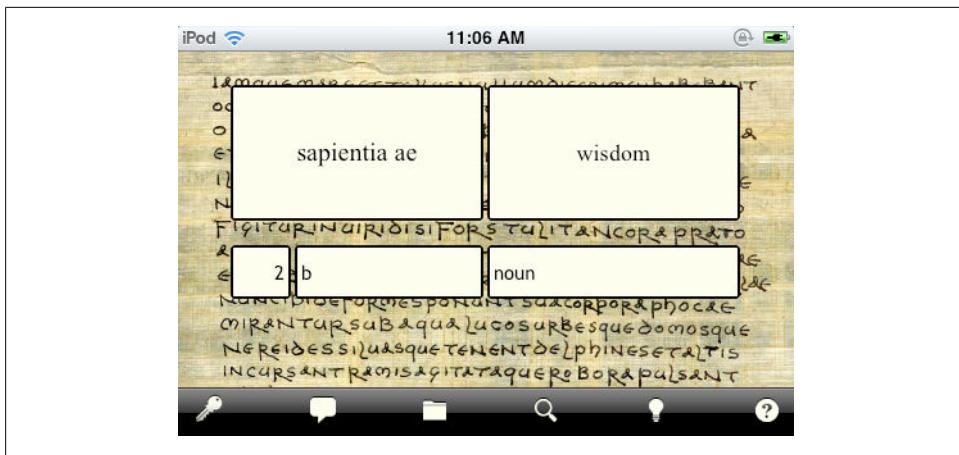


Figure 19-3. A Latin flashcard app

I'll illustrate with another app of mine (Figure 19-3). The interface displays a flashcard containing information about a Latin word, along with a toolbar (the black area at the bottom) where the user can tap an icon to choose additional functionality.

Again, I will describe how the interface shown in Figure 19-3 comes to appear on the screen through the view controller hierarchy and the view hierarchy (Figure 19-4). The app actually contains over a thousand of these Latin words, and I want the user to be able to navigate to the next or previous word; iOS 5 provides an excellent view controller for this purpose, the `UIPageViewController`. However, that's just for the card; the toolbar at the bottom just stays there, so it can't be inside the `UIPageViewController`'s view. Therefore the app's root view controller is my own `UIViewController` subclass, `RootViewController`; its view contains the toolbar and the `UIPageViewController`'s view.

In Figure 19-4, my `RootViewController`'s view becomes the window's subview (the root view) automatically, by virtue of the `RootViewController`'s being the window's `rootViewController`. But then, because I want to put a `UIPageViewController`'s view into my `RootViewController`'s view, it is up to me to make `RootViewController` function as a parent view controller; I must make the `UIPageViewController` the `RootViewController`'s child, and I must put the `UIPageViewController`'s view *manually* into my `RootViewController`'s view. Finally, the way `UIPageViewController` works as it replaces one interface by another is by swapping out a child view controller; so I hand the `UIPageViewController` an instance of my `CardController` class (another `UIViewController` subclass) as its child, and the `UIPageViewController` displays the `CardController`'s view automatically.

Finally, here's an example of a presented view controller. My Latin flashcard app has a second mode, where the user is drilled on a small subset of the cards in random order;

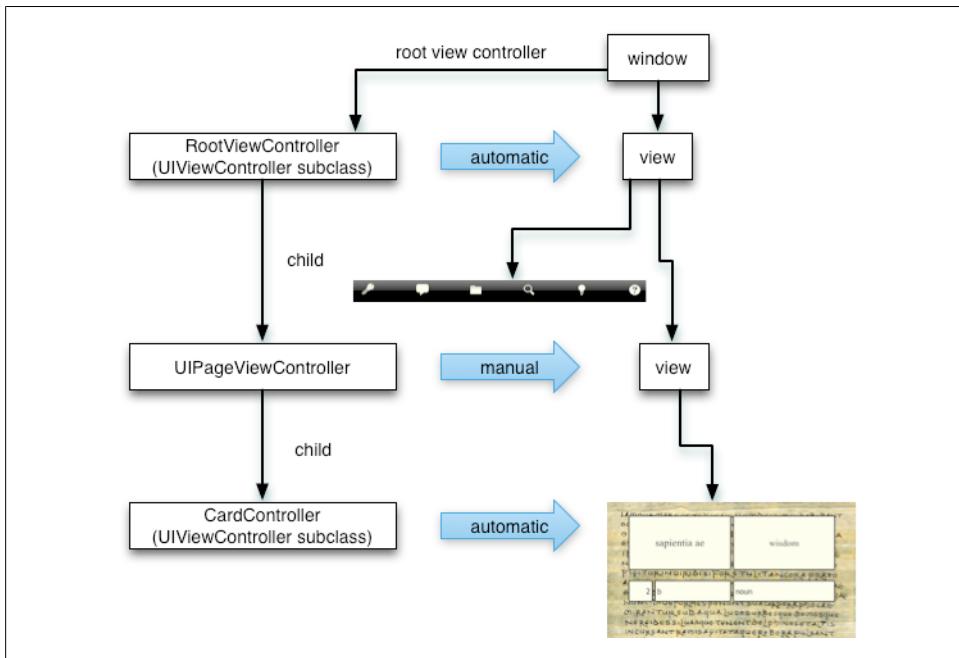


Figure 19-4. The Latin flashcard app’s initial view controller and view hierarchy

the interface looks very much like the first mode’s interface (Figure 19-5), but it behaves completely differently.

To implement this, I have another UIViewController subclass, DrillViewController; it is structured very much like RootViewController. When the user is in drill mode, a DrillViewController is being *presented* by the RootViewController, meaning that the DrillViewController’s interface takes over the screen automatically: the DrillViewController’s view, and its whole subview hierarchy, replaces the RootViewController’s view and *its* whole subview hierarchy. The RootViewController and its hierarchy of child view controllers remains in place, but the corresponding view hierarchy is not in the interface; it will be returned to the interface automatically when we leave drill mode (because the presented DrillViewController is dismissed), and the situation will look like Figure 19-4 once again.

For any app that you write, you should be able to construct a diagram similar to Figure 19-2 and the others, showing the hierarchy of view controllers and charting how each view controller’s view fits into the view hierarchy; and each view controller’s view should fit into the hierarchy as shown in the diagrams, namely, in parallel with the view controller hierarchy. If view controller B is view controller A’s child, then view controller B’s view should be view controller A’s view’s subview (at some level); there should be no crossed wires or orphan views.

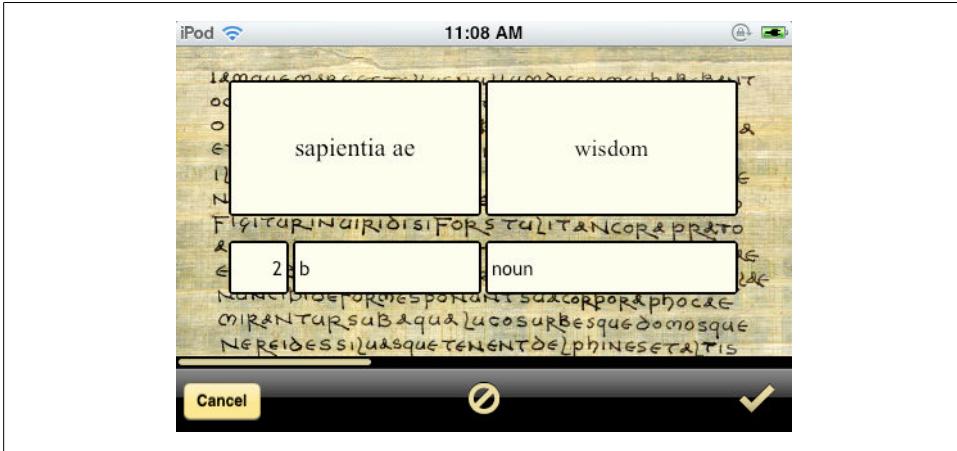


Figure 19-5. The Latin flashcard app, in drill mode

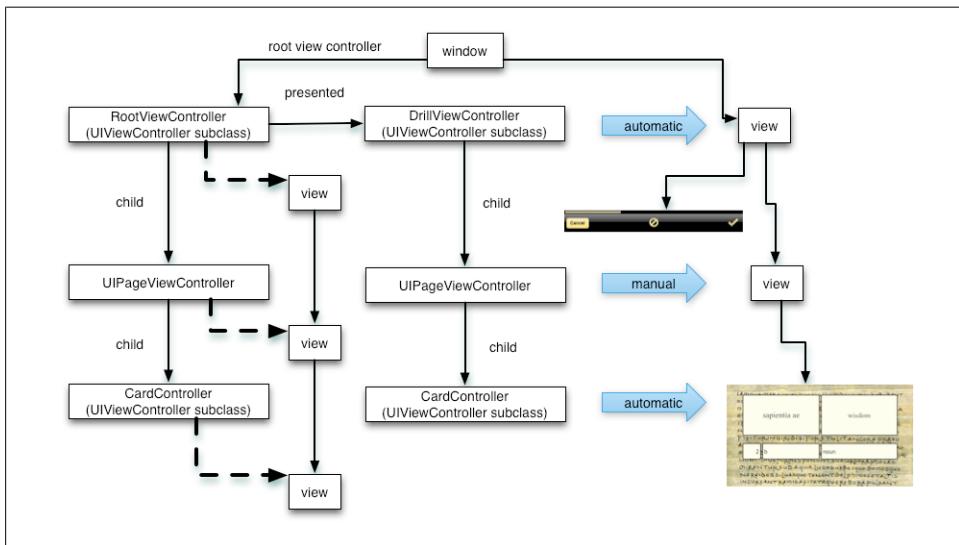


Figure 19-6. The Latin flashcard app's drill mode view controller and view hierarchy

View Controller and View Creation

On the whole, a view controller is created exactly like any other object. An view controller instance comes into existence because you instantiate a view controller class, either in code or by loading a nib ([Chapter 5](#)). But the instantiation of a view controller introduces some additional considerations:

- How will the view controller persist?

- How will the view controller’s view get into the interface?
- Where will the view controller’s view come from?

We begin with the issue of *persistence*. Even if you’re using ARC, memory must be managed somehow ([Chapter 12](#)). A view controller instance, once brought into existence, will eventually go right back out of existence if it is not retained somehow. The distinction between a view controller and its view can add to the confusion. It is possible, if things are mismanaged, for a view controller’s view to get into the interface while the view controller itself is allowed to go out of existence. This must not be permitted. If it does, at the very least the view will apparently misbehave, failing to perform its intended functionality, because that functionality is embodied by the view controller, which no longer exists. (I’ve made this mistake, so I speak from experience here.)

Fortunately, Cocoa follows a simple rule: if you hand a view controller to some other object whose job is to use that view controller somehow, the other object retains the view controller. For example, assigning a view controller to a window’s `rootViewController` property retains the view controller. Making a view controller another view controller’s child, or presenting a view controller from another view controller, retains it. Passing a view controller as the argument to `UIPopoverController`’s `initWithContentViewController:` retains the view controller. (There is then the problem of who will retain the `UIPopoverController`; this will cause much gnashing of teeth in [Chapter 22](#).) And so on.

This means that if you construct the view controller hierarchy correctly, the persistence problem will be solved. It also means that this little dance in the Single View Application template’s app delegate code is unnecessary:

```
self.viewController =
    [[ViewController alloc] initWithNibName:@"ViewController" bundle:nil];
self.window.rootViewController = self.viewController;
```

That code instantiates the `ViewController` class. It first assigns the instance to a `viewController` property with a `strong` retain policy; it then assigns the same instance to the window’s `rootViewController` property. The first of those steps is pointless: the `viewController` property is performing no useful function, as it is not needed in order to retain the view controller, and is never referred to again. This code would have done exactly the same thing:

```
self.window.rootViewController =
    [[ViewController alloc] initWithNibName:@"ViewController" bundle:nil];
```

If the same code or any other code ever does need to refer to this view controller again, it can get a reference to it through the window’s `rootViewController` property:

```
UIViewController* rvc =
    [[[UIApplication sharedApplication] keyWindow] rootViewController];
```

Now let’s talk about how the view controller’s view will get *into the interface*. As I’ve already said in the preceding section, and emphasized in the diagrams there, this will

nearly always happen automatically. Before iOS 4, which introduced the `rootViewController` property of the `UIWindow` class, the root view often had to be placed into the window manually, and the first edition of this book spilled a goodly measure of ink discussing how to do it properly; this edition, however, assumes that the `rootViewController` property exists and that you'll use it. When a view controller is assigned to the window's `rootViewController` property, the view controller's view is made the window's subview (the root view), with a correct frame, automatically.

Similarly, built-in view controllers are responsible for displaying the views of their child view controllers in their own views, as in [Figure 19-2](#) where the `UINavigationController` puts its child view controller's view into its own view, thus displaying that view's subviews, the label and the table view. And a presented view controller's view automatically replaces in the interface the view of the presenting view controller, as in [Figure 19-6](#).

The exceptional case, as I've already mentioned, is when your custom `UIViewController` subclass is acting as a parent view controller. In that case, the custom `UIViewController` subclass will need to contain some code that puts a child view controller's view into its own view, as appropriate. I'll return to this issue and demonstrate with actual code later in this chapter.

Finally, we have the issue of *where a view controller's view comes from*. For a built-in view controller class that you don't subclass, this is not a problem; in fact, you may not even be particularly conscious of the view controller's view. In [Figure 19-1](#) and [Figure 19-2](#), the `UINavigationController`'s view is barely a player. Even though it is in fact the app's root view, it is never seen in the interface as a distinct entity, and there is never any need to speak of it in code. You assign the `UINavigationController` to the window's `rootViewController` property, and you assign a child view controller to the `UINavigationController`, and the child view controller's view appears in the interface — and that's the end of that. The `UINavigationController` created its own view automatically and put both the navigation bar and its child's view into it automatically, and the window put the `UINavigationController`'s view into the interface automatically; the `UINavigationController`'s view functions as a kind of intermediary that you aren't concerned with, containing the interface that you *are* concerned with.

When you write a `UIViewController` subclass, however, the question of where its view is to come from is an extremely important question. It is *crucial* that you understand the answer to this question, which quite possibly causes more confusion to beginners than any other matter connected with iOS programming. The answer is rather involved, though, because there are several different options. The rest of this section treats those options one by one. To anticipate, the alternatives are as follows:

- The view may be created in code, manually.
- The view may be created as an empty generic view, automatically.
- The view may be created in its own separate nib.

- The view may be created in a nib, which is the same nib from which the view controller itself is instantiated.

Before we proceed, here's a caveat: distinguish between *creating* a view and *populating* that view. With a view controller, these are very clearly two different operations. Once the view controller has its view, your UIViewController subclass code will get plenty of further opportunities to customize what's in that view. I'll talk about that, of course, but keep in mind that the primary question with which we're concerned is how the UIViewController instance obtains its actual view at all, the view that can be accessed as its `view` property.

Manual View

To provide a UIViewController with a view manually, implement its `loadView` method. Your job here is to obtain an instance of UIView (or a subclass of UIView) and *assign it to self.view*. You must *not* call `super` (for reasons that I'll make clear later on).

Let's try it. Start with a project made from the Empty Application project template:

- We need a UIViewController subclass, so choose File → New → New File; specify an iOS Cocoa Touch UIViewController subclass. Click Next.
- Name the file `RootViewController`. Make sure this is a UIViewController subclass. Uncheck both checkboxes. Click Next.
- Save into the appropriate folder and group, as part of the target.

We now have a `RootViewController` class, and we proceed to edit its code. In `RootViewController.m`, the UIViewController subclass template has already provided a `loadView` method, commented out. So uncomment it and let's implement it. To convince ourselves that the example is working correctly, we'll give the view an identifiable color, and we'll put some interface inside that view, in this case a "Hello, World" label:

```
- (void) loadView {
    UIView* v = [[UIView alloc] init];
    v.backgroundColor = [UIColor greenColor];
    self.view = v;
    UILabel* label = [[UILabel alloc] init];
    [v addSubview:label];
    label.text = @"Hello, World!";
    label.autoresizingMask = (
        UIViewAutoresizingFlexibleTopMargin |
        UIViewAutoresizingFlexibleLeftMargin |
        UIViewAutoresizingFlexibleBottomMargin |
        UIViewAutoresizingFlexibleRightMargin
    );
    [label sizeToFit];
    label.center = CGPointMake(CGRectGetMidX(v.bounds), CGRectGetMidY(v.bounds));
    label.frame = CGRectIntegral(label.frame); // prevent fuzzies
}
```

We have not yet given any `RootViewController` instance a place in our view controller hierarchy — in fact, we have no view controller hierarchy. Let’s make one. To do so, we turn to `AppDelegate.m`. (It’s a little frustrating having to set things up in two different places before our labors can bear any visible fruit, but such is life.)

In `AppDelegate.m`, add the line `#import "RootViewController.h"` at the start, so that our code can speak of the `RootViewController` class. Then modify the start of the implementation of `application:didFinishLaunchingWithOptions:` to create a `RootViewController` instance and make it the window’s `rootViewController`. Observe that we must do this *after* our `window` property actually has a `UIWindow` as its value! That is why the template’s comment, “Override point for customization after application launch,” comes *after* the line that creates the `UIWindow`:

```
self.window = [[[UIWindow alloc] initWithFrame:[[UIScreen mainScreen] bounds]]];  
// Override point for customization after application launch.  
RootViewController* theRVC = [[RootViewController alloc] init];  
self.window.rootViewController = theRVC; // retains, places view into interface  
// ... and the rest is as in the template
```

Build and run the app. Sure enough, there’s our green background and our “Hello, world” label!

We have proved that we can create a view controller and get its view into the interface. But perhaps you’re not persuaded that the view controller is *managing* that view in an interesting way. To prove this, let’s make it possible to rotate our interface. Return to `RootViewController.m` and modify the template’s implementation of `shouldAutorotateToInterfaceOrientation:` to look like this:

```
- (BOOL)shouldAutorotateToInterfaceOrientation:(UIInterfaceOrientation)io  
{  
    return YES;  
}
```

Build and run the app, and rotate the device; in the simulator, you can do this by choosing Hardware → Rotate Left or Hardware → Rotate Right. Observe that the view, as indicated by the orientation of the “Hello, World” label, automatically rotates to compensate; that’s the work of the view controller. Observe also that we were careful to give the label an appropriate `autoresizingMask`, such as to keep it centered in the view even when the view’s bounds are changed to fit the rotated window. (I’ll talk more about rotation later in this chapter.)

We are not bothering to give our view a reasonable frame. This is because we are relying on someone else to frame the view appropriately. In this case, the “someone else” is the window, which responds to having its `rootViewController` property set to a view controller by framing the view controller’s view appropriately as the root view before putting it into the window as a subview. To be precise, the root view’s frame as it goes into the window in an iPhone app is `{0, 20}, {320, 460}` — that is, the root view fills the part of the window not covered by the status bar. This magic is easily accomplished by setting the view’s frame to `[[UIScreen mainScreen] applicationFrame]`,

which is exactly what the window does to the view controller’s view. (If you’re concerned about your view having a reasonable frame in the later part of your `loadView` code, you can certainly set its frame to `[[UIScreen mainScreen] applicationFrame]` yourself in the earlier part of your `loadView` code.)

If there is no status bar — for example, if “Status bar is initially hidden” is YES in our `Info.plist`, a possibility that I mentioned in [Chapter 9](#) — the call to `[[UIScreen mainScreen] applicationFrame]` will return the entire bounds of the window, and our view will fill the screen, which is still a correct result.

On the other hand, if the status bar is present but its status bar style is set to “Transparent black style,” then our root view’s frame fills the part of the window not covered by the status bar once again. This may or may not be what you want. If it is, fine. But if you want the root view to underlap the transparent status bar, you’ll need to set the view controller’s `wantsFullScreenLayout` to YES. You could do that in the app delegate:

```
RootViewController* theRVC = [[RootViewController alloc] init];
theRVC.wantsFullScreenLayout = YES;
self.window.rootViewController = theRVC;
```

Alternatively, you might feel that it is the view controller’s job to know that its view should underlap the status bar, in which case you could do it at some early point in the life of the view controller, such as `loadView`:

```
- (void) loadView {
    self.wantsFullScreenLayout = YES;
    // ... and so on ...
```

Generic Automatic View

Earlier, I said that we must draw a distinction between creating a view and populating it. It could be argued that the example in the preceding section fails to draw this distinction. The lines that create our `RootViewController`’s view are these:

```
UIView* v = [[UIView alloc] init];
self.view = v;
```

Everything else merely configures and populates the view, turning it green and putting a label in it. A more appropriate place to populate a view controller’s view is in its `viewDidLoad` implementation, which is called after the view exists (so that it can be referred to as `self.view`). We could therefore rewrite the preceding example like this:

```
- (void) loadView {
    UIView* v = [[UIView alloc] init];
    self.view = v;
}

- (void)viewDidLoad
{
    [super viewDidLoad];
    UIView* v = self.view;
    v.backgroundColor = [UIColor greenColor];
```

```

UILabel* label = [[UILabel alloc] init];
[v addSubview:label];
label.text = @"Hello, World!";
label.autoresizingMask = (
    UIViewAutoresizingFlexibleTopMargin |
    UIViewAutoresizingFlexibleLeftMargin |
    UIViewAutoresizingFlexibleBottomMargin |
    UIViewAutoresizingFlexibleRightMargin
);
[label sizeToFit];
label.center = CGPointMake(CGRectGetMidX(v.bounds), CGRectGetMidY(v.bounds));
label.frame = CGRectIntegral(label.frame); // prevent fuzzies
}

```

But if we're going to do that, we can go even further and remove our implementation of `loadView` altogether. If you don't implement `loadView`, and if no view is supplied in any other way, then `UIViewController`'s implementation of `loadView` will do exactly what we are already doing in code: it creates a generic `UIView` object and assigns it to `self.view`. If we needed our view controller's view to be a particular `UIView` subclass, that wouldn't be acceptable; but in this case, our view controller's view *is* a generic `UIView` object, so it *is* acceptable. Comment out or delete the `loadView` implementation, and our example still works:

```

/*
- (void)loadView {
    UIView* v = [[UIView alloc] init];
    self.view = v;
}

- (void)viewDidLoad
{
    [super viewDidLoad];
    UIView* v = self.view;
    v.backgroundColor = [UIColor greenColor];
    UILabel* label = [[UILabel alloc] init];
    [v addSubview:label];
    label.text = @"Hello, World!";
    label.autoresizingMask = (
        UIViewAutoresizingFlexibleTopMargin |
        UIViewAutoresizingFlexibleLeftMargin |
        UIViewAutoresizingFlexibleBottomMargin |
        UIViewAutoresizingFlexibleRightMargin
    );
    [label sizeToFit];
    label.center = CGPointMake(CGRectGetMidX(v.bounds), CGRectGetMidY(v.bounds));
    label.frame = CGRectIntegral(label.frame); // prevent fuzzies
}

```

View in a Separate Nib

A view controller's view can be provided from a nib file. This approach gives you the convenience of configuring and populating the view using the nib editor interface ([Chapter 7](#)).

(Since a nib-based approach is possible, you might wonder whether anyone in his right mind would ever use the techniques demonstrated in the preceding examples; why would someone create and populate a view in code when it is possible to do it in the nib? Actually, I'm quite fond of the code-based approach. It's verbose, to be sure, but it is clear and straightforward and gives me complete control, whereas loading a view from a nib always leaves me scratching my head a bit, wondering whether I checked all the right checkboxes and whether they will do the right thing.)

Our next technique for providing a view controller with a view, then, will be to make that view an object in a nib file. For this to work, it is necessary to prepare the nib file, as follows:

1. The File's Owner class must be set to a view controller class (usually, you'll want this to be our specific view controller subclass).
2. Performing the preceding step causes the File's Owner proxy in the nib to have a `view` outlet, corresponding to a `UIViewController`'s `view` property. This outlet must be connected to the view.

Do you see where this is heading? We will then load the nib file *with the view controller instance as its owner*. The view controller's class matches the File's Owner class, the view controller's `view` property is set via the `view` outlet in the nib to the view object, and presto, our view controller has a view. *If you don't understand what I just said, reread Chapter 7.* It is crucial that you comprehend how this technique works.

Now let's try it. We can start with the example we've already developed, with our `RootViewController` class. However, for now we must delete the implementation of `loadView` and `viewDidLoad` from `RootViewController.m`, because we want the view to come from a nib and we're going to populate it in the nib. Then:

1. Choose File → New → New File and specify an iOS User Interface View nib file. Click Next.
2. In the Device Family pop-up menu, choose iPhone. Click Next.
3. Name the file `MyNib`, and specify the appropriate folder and group, and make it part of the target. Click Create.
4. Edit `MyNib.xib` in the way I described a moment ago: set the File's Owner class to `RootViewController` (in the Identity inspector), and connect the File's Owner `view` outlet to the View object.
5. Design the view. To make it clear that this is not the same view we were creating previously, perhaps you should give the view a red background color (in the At-

tributes inspector). Drag a UILabel into the middle of the view and give it some text, such as “Hello, World!”

Back in *AppDelegate.m*, where we create our RootViewController instance, we must load *MyNib.xib* with the RootViewController instance as its owner. It is, in fact, possible to do this using the technique described back in [Chapter 7](#), though one shouldn’t:

```
// shouldn't do this!
RootViewController* theRVC = [[RootViewController alloc] init];
[[NSBundle mainBundle] loadNibNamed:@"MyNib" owner:theRVC options:nil];
self.window.rootViewController = theRVC; // retains, places view into interface
```

The correct approach is to instantiate the view controller and tell it what nib it is to load as owner, but let *it* load the nib when it needs to. The view controller then manages the loading of the nib and all the associated housekeeping correctly. This technique involves initializing the view controller using `initWithNibName:bundle:` (which is actually UIViewController’s designated initializer), like this:

```
RootViewController* theRVC =
    [[RootViewController alloc] initWithNibName:@"MyNib" bundle:nil];
self.window.rootViewController = theRVC; // retains, places view into interface
```

That works (and you can run the project to prove it), but now I’m going to show you a shortcut. It turns out that if the nib file has the same name as the view controller’s class, we can pass nil as the nib name. This means, in effect, that we can return to using `init` to initialize the view controller; the designated initializer is `initWithNibName:bundle:`, so UIViewController’s `init` actually calls `initWithNibName:bundle:`, passing nil for both arguments. The rule is, then, that if the nib name passed to `initWithNibName:bundle:` is nil, a nib will be sought with the same name as the view controller’s class.

Let’s try it. Rename *MyNib.xib* to *RootViewController.xib*, and change the code that instantiates and initializes our RootViewController like this:

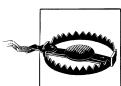
```
RootViewController* theRVC = [[RootViewController alloc] init];
self.window.rootViewController = theRVC; // retains, places view into interface
```

The project still works. But wait, there’s more! It seems ridiculous that we should end up with a nib that has “Controller” in its name merely because our view controller, as is so often the case, has “Controller” in *its* name. A nib, after all, is not a controller. Well, there’s an additional aspect to the shortcut: the runtime, in looking for a view controller’s corresponding nib, will in fact try stripping “Controller” off the end of the view controller class’s name. (This feature is undocumented, but it works reliably and I can’t believe it would ever be retracted.) Thus, we can name our nib file *RootView.xib* instead of *RootViewController.xib*, and it will *still* be properly associated with our RootViewController instance when we initialize that instance using `init`.



Recall from [Chapter 9](#) that when an image file is sought by name in the app's bundle, naming conventions allow different files to be loaded under different runtime conditions. The same is true for nib files. A nib file named `RootViewController~ipad.xib` will be loaded on an iPad when the name `@"RootViewController"` is specified, regardless of whether it is specified explicitly (as the first argument to `initWithNibName:bundle:`) or implicitly (because the view controller class is `RootViewController`). This principle will greatly simplify your life when you're writing a universal app.

When you create the files for a `UIViewController` subclass, the Xcode dialog has a checkbox (which we unchecked earlier) offering to create an eponymous `.xib` file at the same time ("With XIB for user interface"). If you accept that option, the nib is created with the File's Owner's class already set to the view controller's class and with its `view` outlet already hooked up to the view. This automatically created `.xib` file does *not* have "Controller" stripped off the end of its name; you can rename it manually later (I generally do) if the default name bothers you.



Not every built-in subclass of `UIViewController` obeys the convention that a nil nib name means a nib with the same name as the view controller's class. In particular, `UITableViewController` does not. I regard this as a bug; in any case, it has caught me by surprise several times (and can be difficult to track down, when the table view mysteriously fails to appear in the interface).

Earlier, I suggested deleting `viewDidLoad` from `RootViewController`'s code. This was mostly so that we couldn't confuse ourselves; I wanted you to see clearly that the view was being created and configured in the nib. In real life, however, it is perfectly acceptable to load a view controller's view from a nib file *and* proceed to further configurations and initializations in `viewDidLoad`. This is a common thing to do; it has the advantage that by the time `viewDidLoad` is called, we are guaranteed that the view has been loaded from the nib and that we can access it via `self.view`.

On the other hand, if a view controller's view is to come from a nib, you should *not* implement `loadView`. You'll just confuse yourself if you do. The truth is that `loadView` is *always* called when the view controller first decides that it needs its view. If we override `loadView`, we supply and set the `view` in code. If we don't override `loadView`, the default implementation is to load the view controller's nib and set the `view` through an outlet. (That is why, if we do override `loadView`, we must *not* call `super` — that would cause us to get *both* behaviors.) If we don't override `loadView` and there is no associated nib (because the nib name was nil in `initWithNibName:bundle:` and there is no nib whose name matches the name of the view controller class, in the way that I described a moment ago), the default implementation of `loadView` creates a generic `UIView` as discussed in the previous section.

When Is the View Loaded?

A UIViewController's `view` property is set "lazily": rather than setting it when the UIViewController itself is instantiated, the `view` property isn't set until it's needed — namely, when you or Cocoa tries to fetch its value for the first time. At that moment, `loadView` is called. This architecture has several advantages, but it can also trap you: I already mentioned (in [Chapter 11](#)) how I made the mistake of mentioning a UIViewController's `view` in its `awakeFromNib` and caused the view to be loaded twice. So don't mention a view controller's `view` until it's time to load it! The `isViewLoaded` method reports whether the view has in fact been loaded, without mentioning it in a way that will also cause it to be loaded accidentally.

It will not have escaped your attention that the architecture I've just been describing is the architecture used by most of the Xcode 4.2 project templates. Take, for example, the Single View Application template. You already know that, using this template, you can design the initial interface in a nib file and have it appear in the running app. You are now in a position to understand precisely why. In addition to the AppDelegate class, there's a ViewController class along with a nib file called `ViewController.xib`. The app delegate's `application:didFinishLaunchingWithOptions:` instantiates ViewController, associating it with its nib, and makes that instance the window's `rootViewController`:

```
self.window = [[UIWindow alloc] initWithFrame:[[UIScreen mainScreen] bounds]];
// Override point for customization after application launch.
self.viewController =
    [[ViewController alloc] initWithNibName:@"ViewController" bundle:nil];
self.window.rootViewController = self.viewController;
```

That code, however, can be considerably abbreviated, as you now know. There is no need to assign the view controller instance to a property, as it will be retained and available through the window's `rootViewController` property. And there is no need to specify the `nibName:` argument, because the nib file has the same name as the view controller. So we could have said this:

```
self.window = [[UIWindow alloc] initWithFrame:[[UIScreen mainScreen] bounds]];
// Override point for customization after application launch.
self.window.rootViewController = [[ViewController alloc] init];
```

Moreover, that code works even if we change the name of the nib file to `View.xib`.

Nib-Instantiated View Controller

Like any other object, a view controller can be represented by a nib object and be instantiated through the loading of the nib. In the nib editor, the Object library contains entries for View Controller (UIViewController) as well as for several built-in UIViewController subclasses. Any of these can be dragged into the nib (though not, of course, into a view, since they are not view objects). To retrieve a reference to a view controller

instantiated through the loading of a nib, you'd obviously need an outlet from a proxy object, probably the File's Owner.

To illustrate will require a certain temporary suspension of disbelief, as you've no immediate *need* to instantiate a view controller from a nib. Nevertheless, the ability to do so will prove its worth as soon as you have to do with multiple related view controllers, such as a UINavigatorController and its initial child view controller, a UITabBarController and its multiple child view controllers, or a storyboard. So let's modify our existing example so as to instantiate RootViewController from a nib, and you'll just have to trust me that this could be a useful thing to do, rather than (as it will first seem) an unnecessary complication.

Our first step will be to create an extra nib for no other purpose than to instantiate RootViewController:

1. Choose File → New → New File and specify an iOS User Interface Empty nib file. Click Next.
2. In the Device Family pop-up menu, specify iPhone.
3. Name the nib file *RVC* and create it in the appropriate fashion.
4. Edit the newly created *RVC.xib*. Drag a View Controller into the canvas. In the Attributes inspector, specify that this is a RootViewController instance.
5. We'll need an appropriate File's Owner object with an outlet pointing to our RootViewController instance. At the earliest stage of our application's launch process, there's really only one candidate — the app delegate. So select the File's Owner proxy and specify in the Attributes inspector that this is an AppDelegate.
6. We still need that outlet, and we can't make it without a corresponding property in AppDelegate. Option-click *AppDelegate.h* in the Project navigator so that *RVC.xib* is being edited in the main pane of the editor and *AppDelegate.h* is being edited in the assistant pane. Control-drag from the Root View Controller object in the nib into the @interface section in *AppDelegate.h*. You're offered the chance to create an outlet; call it *rvc*. You'll also need to type this line of code before the start of the interface section, to satisfy the compiler that RootViewController is a class name:

```
@class RootViewController;
```

Now we're ready to tell AppDelegate to load *RVC.xib* with itself as owner and extract the RootViewController instance and use it as the window's *rootViewController*. Return to *AppDelegate.m*, where you'll find that Xcode has already synthesized *rvc* for us. Change the start of *application:didFinishLaunchingWithOptions:* to look like this:

```
self.window = [[UIWindow alloc] initWithFrame:[[UIScreen mainScreen] bounds]];
// Override point for customization after application launch.
[[NSBundle mainBundle] loadNibNamed:@"RVC" owner:self options:nil];
self.window.rootViewController = self.rvc;
```

Build and run the app. It works, and the interface from *RootViewController.xib* appears (or from *RootView.xib* if you renamed it)! This makes perfect sense, since nothing has changed from our examples in the previous section except the way we instantiated *RootViewController*. It comes into existence from the loading of the nib *RVC.xib*, but the runtime then performs the very same search as before for a nib with the same name as the view controller class — and finds it.

But what if the nib had a different name? Edit *RVC.xib*, select the Root View Controller, and study its Attributes inspector. You'll find there's a NIB Name field. At the moment it's empty, signifying the equivalent of a nil `nibName`: argument in the initializer. But you could type (or use the combo box to choose) the name of a different nib file to associate that nib with this view controller as the view controller is instantiated by the loading of the nib. Thus, everything that was possible in the previous sections, where we instantiated the view controller in code, remains possible now that we're instantiating it from a nib file.

Indeed, it remains possible to supply the view controller's view in *any* of the other ways discussed earlier in this section. The mere fact that this view controller is instantiated from a nib rather than using code changes nothing. If the view controller in the nib has no contained view, and if no same-named view nib is found, `loadView` will be called, and you can implement it to supply a view, or delete it and let a generic `UIView` be supplied. No matter where the view comes from, you can configure it further or do any other initial tasks in `viewDidLoad`.

There is also, however, a completely new alternative: we can design the interface *right here in the same nib file* as the view controller (that is, in *RVC.xib*). In fact, we can design the interface *in the view controller itself*. You'll notice that the canvas representation of the view controller is the size of an iPhone screen, even though a view controller is not a view object. There's a reason for that: it's so that the view controller can *accommodate* a screen-sized view object, to serve as its `view`.

To see this, drag a generic View object from the Object library right into the Root View Controller object in the nib editor canvas. You can now proceed to design the interface within this view. For example, you can make its background color yellow (to distinguish it from all the other interfaces we've been designing) and drag a different label into it (perhaps you could make it say "Howdy, Universe" for a change).

Build and run the project. The yellow background and the "Howdy, Universe" label appear! This is because the presence of a view inside the view controller in the nib is taken as the equivalent of `loadView` in code: this view is instantiated and assigned to the view controller as its `view` property.

Take some time to look a little further at a view controller's Attributes inspector; it provides ways to set some options that would otherwise be set in code. For example, the Wants Full Screen checkbox is our friend the `wantsFullScreenLayout` property. The meanings of the other options will become evident as this chapter proceeds.



Like any other nib object, when a view controller is instantiated from a nib or storyboard, its designated initializer (`initWithNibName:bundle:`) is *not* called. If your `UIViewController` subclass needs very early access in code to the view controller instance, it can override `initWithCoder:` ([Chapter 7](#), [Chapter 11](#)).

Storyboard-Instantiated View Controller

iOS 5 introduces storyboards. A *storyboard* is, in effect, a single file representing a collection of nib files similar to *RVC.xib* in the previous section — that is, nib files containing view controller nib objects. A storyboard basically consists of multiple view controller nib objects arranged in a chain similar to the chains of child and presented view controllers diagrammed in figures earlier in this chapter. The value of a storyboard emerges not in its saving you from writing code (in my experience, use of a storyboard often entails writing *more* code than otherwise) but in its graphic expression of the view controller hierarchy. In a world where the possible states of your app’s overall interface can be best understood through a diagram of view controllers and the instructions for when and how to transition from one to the next, a storyboard can conveniently encapsulate these key aspects of your app’s operation, helping you as a developer create and maintain your app.

A storyboard, like a nib, is an actual file in your project (a `.Storyboard` file); it is compiled into your app’s bundle, and can be referenced by calling the `UIStoryboard` class method `storyboardWithName:bundle:`. A view controller object in a storyboard must still correspond to an actual view controller class in your project, either inherited from Cocoa or a `UIViewController` subclass of your own, just as with a nib. And editing a view controller in a storyboard is almost identical to editing a view controller in a nib.

However, because a storyboard can effectively combine what would have been separate nib files, the mechanism for how a view controller is instantiated from a storyboard is a little different from how it is instantiated from a nib. A storyboard has no owner, the way a nib file does, and so it contains no File’s Owner proxy. A view controller is instantiated from a storyboard in one of three ways:

- At most one view controller in the storyboard is designated the storyboard’s *initial view controller*. Typically, this will be the view controller to be used as the app’s root view controller at launch. The storyboard is sent `instantiateInitialViewController`, and returns an instance of the initial view controller’s class, configured in accordance with your edits in the storyboard.
- When the moment comes for one view controller’s view to transition to another view controller’s view, the second view controller is instantiated *automatically* (and configured in accordance with your edits in the storyboard) and made an appropriate child or presented view controller of the first. Such a transition is called a *segue*. (Storyboard-related terminology, including the word “storyboard,” is drawn

from the world of movie-making. View controllers themselves are referred to as “scenes.”)

- A view controller in a storyboard can be assigned an arbitrary string identifier using the storyboard editor. It is then possible to instantiate that view controller by sending `instantiateViewControllerWithIdentifier:` to the storyboard; an instance of the view controller’s class is returned, configured in accordance with your edits in the storyboard.

So far in this section of the chapter, we’ve been concerned with generating our app’s root view controller and its view. Let’s continue by seeing how to do that with a storyboard. Start all over again with a new project based on the Single View application template, but this time make sure that Use Storyboard is checked. The resulting project consists of an `AppDelegate` class and a `ViewController` class, and a storyboard (called `MainStoryboard.storyboard`) instead of a nib.

Look in `AppDelegate.m` and you’ll discover that `application:didFinishLaunchingWithOptions:` contains no code at all — not to load the storyboard file, nor even to generate the window and display it. That’s because `UIApplicationMain` does all the work behind the scenes. As I explained in [Chapter 6](#), `MainStoryboard.storyboard` is designated the main storyboard file in our `Info.plist`, using the key “Main storyboard file base name” (`UIStoryboardName`). The result, as I explained in [Chapter 14](#), is that after `UIApplicationMain` instantiates the app delegate class, it asks the app delegate for the value of its `window` property and then takes the resulting resulting window instance (or, if the `window` property was nil, an instance of `UIWindow` that it creates and assigns to the `window` property), loads the main storyboard, sends it `instantiateViewControllerWithIdentifier:` to get a root view controller instance, assigns that instance to the window’s `rootViewController` property, and calls `makeKeyAndVisible` on the window.

Now edit `MainStoryboard.storyboard`. It contains a single view controller object, with its class already specified as `ViewController`, already designated as the storyboard’s initial view controller, and already containing a view that will function as its `view` — exactly like the view controller in `RVC.xib` at the end of the preceding section.

So now you can give this view a background color, put a label into it, and build and run the project — and your view appears. *Be sure you understand why.* Storyboards are not magic, any more than nibs. A view controller in a storyboard is a view controller like any other. It is perfectly possible to get the view for a view controller instantiated from a storyboard in any of the other ways discussed earlier in this section. You could delete the View contained by the view controller in the storyboard. You could then provide the view by configuring a nib file with the same name as the view controller class. (However, there is no way to specify a different nib file name for this view controller in the storyboard editor; this feels like an oversight.) Alternatively, you could provide the view in the view controller class’s `loadView`, or by having no `loadView` and letting a default `UIView` be provided. No matter where the view comes from, you can configure it further or do any other initial tasks in `viewDidLoad`.

Take a moment to study the storyboard editing interface a little. Below the view controller in the canvas is a black bar containing two icons. The first is the First Responder proxy object, which is also present in a nib file. The second is the view controller itself: this is merely a convenient way to select the view controller (because otherwise you'd typically be working with its contained view); you could do the same thing using the view controller's listing in the dock. The dock itself is divided into scenes; at the moment there's just one scene. The black bar is effectively a collection of the very same top-level objects shown for this scene in the dock.

Rotation

A major part of a view controller's job is to know how to rotate the view. The user will experience this as rotation of the app itself: the top of the app shifts so that it is oriented against a different side of the device's display. There are two complementary uses for rotation:

Compensatory rotation

The app rotates to compensate for the orientation of the device, so that the app appears right way up with respect to how the user is holding the device. The challenge of compensatory rotation stems, quite simply, from the fact that the screen is not square. This means that if the app rotates 90 degrees, the interface no longer fits the screen, and must be changed to compensate.

Forced rotation

The app rotates when a particular view appears in the interface, or when the app launches, to indicate that the user needs to rotate the device in order to view the app the right way up. This is typically because the interface has been specifically designed, in the face of the fact that the screen is not square, to appear in one particular mode (portrait or landscape).

In the case of the iPhone, no law says that your app has to perform compensatory rotation. Most of my iPhone apps do not do so; indeed, I have no compunction about doing just the opposite, forcing the user to rotate the device differently depending on what view is being displayed. The iPhone is small and easily reoriented with a twist of the user's wrist, and it has a natural right way up, especially because it's a phone. (The iPod touch isn't a phone, but the same argument works by analogy.) On the other hand, Apple would prefer iPad apps to rotate to at least two opposed orientations (such as landscape with the button on the right and landscape with the button on the left), and preferably to all four possible orientations, so that the user isn't restricted in how the device is held and positioned.

It's fairly trivial to let your app rotate to two opposed orientations, because once the app is set up to work in one of them, it can work with no change in the other. But allowing an app to rotate between two orientations that are 90 degrees apart is trickier, because its dimensions must change — its height and width are swapped — and this

may require a change of layout and might even call for more substantial alterations to the interface, such as removal or addition of part of the interface. A good example is the behavior of Apple’s Mail app on the iPad: in landscape mode, the master pane and the detail pane appear side by side, but in portrait mode, the detail pane is removed and must be summoned using a button or by swiping, at which point the user can work only in the detail pane until the detail pane is dismissed.

The main thing your view controller must do in order to support rotation is to override `shouldAutorotateToInterfaceOrientation:`. The incoming parameter is the proposed orientation, and will be one of the following:

- `UIInterfaceOrientationPortrait`, with the home button at the bottom.
- `UIInterfaceOrientationPortraitUpsideDown`, with the home button at the top.
- `UIInterfaceOrientationLandscapeLeft`, with the home button at the left.
- `UIInterfaceOrientationLandscapeRight`, with the home button at the right.

You return YES for all permitted orientations and NO otherwise. The default, if you don’t override this method, is to return YES for `UIInterfaceOrientationPortrait` and NO otherwise; you *must* return YES for *some* orientation. (There is no `UIViewController` property that can be set to specify what orientations a view controller can adopt; overriding this method is the only way.)

The four interface orientations are matched by four device orientations with similar names (`UIDeviceOrientationPortrait` and so on), except that the two landscape orientations are reversed: `UIInterfaceOrientationLandscapeLeft` is the same as `UIDeviceOrientationLandscapeRight`, because if the user rotates the device 90 degrees left from an original portrait orientation, the interface must rotate 90 degrees right to compensate. The device actually has several more possible orientations, such as `UIDeviceOrientationFaceUp` (the device is lying on its back), but these will not, of themselves, trigger interface rotation.

In case your implementation of `shouldAutorotateToInterfaceOrientation:` wants to return YES to both of two opposed orientations, a pair of macros are provided: `UIInterfaceOrientationIsLandscape(io)` returns YES if `io` is either of the two landscape orientations, and `UIInterfaceOrientationIsPortrait(io)` returns YES if `io` is either of the two portrait orientations.

A new iOS 5 feature is the `UIViewController` class method `attemptRotationToDeviceOrientation`. The idea is that instead of implementing `shouldAutorotateToInterfaceOrientation:` so that it always returns the same thing, you might implement it in some dynamic way, such that at some moment in your app it starts returning YES or NO for different orientations than previously. At such a moment, you need a way to prod the runtime to call `shouldAutorotateToInterfaceOrientation:` again, right now, for all existing view controllers; this class method is how you do that.

What Rotates?

We say that your app rotates, and you'll think of it as rotating, but what really rotates is the status bar's position. When the device rotates, a `UIDeviceOrientationDidChangeNotification` is emitted by the `UIDevice`, and your app's root view controller is consulted with `shouldAutorotateToInterfaceOrientation:`; if the view controller returns YES for the proposed orientation, the `UIApplication` instance is sent the `setStatusOrientation:animated:` message. In a 90-degree rotation, the window's root view then has its width and height dimensions swapped, and a transform is applied so that it appears "right way up." Moreover, this is all accompanied by animation, so it really looks to the user as if the app is rotating. But the window *itself* doesn't budge; it remains "pinned" to the screen (window bounds are screen bounds), it is taller than it is wide, and its top is at the top of the device (away from the home button). As for the view, its bounds are wider than tall in a landscape orientation, but its frame remains taller than wide (though you really shouldn't be referring to the view's frame in this situation, because it has a transform applied; see [Chapter 14](#)).

Initial Orientation

Your app's initial orientation, as the user will see it when launching, is determined by the "Supported interface orientations" setting (`UISupportedInterfaceOrientations`) in the app's *Info.plist*. If there is only one orientation listed here, the user will see the app in that orientation initially. If there is more than one orientation listed, you are giving the system the option to rotate the app into whichever of the supported orientations is closest to the way the device is positioned at the time. In Xcode 4, you can specify this setting graphically by editing the target (in the Summary tab).



The orientation(s) specified by the "Supported interface orientations" setting must also be among the orientations to which the root view controller's `shouldAutorotateToInterfaceOrientation:` will return YES!

As I've already mentioned, on the iPhone, it is common for an app to launch into just one orientation, as it is no trouble for the user to twist the device to view the app correctly; but on the iPad, Apple would like you to list at least two opposed orientations (both portrait orientations or both landscape orientations), and preferably all four.

An app whose initial orientation is portrait mode with the button at the bottom has effectively no work to do, because this is the default orientation. An app with no "Supported interface orientations" setting adopts "Portrait (bottom home button)" by default (`UIInterfaceOrientationPortrait`), and the default `shouldAutorotateToInterfaceOrientation:` behavior is to return YES to `UIInterfaceOrientationPortrait` and NO to anything else. Supporting "Portrait (top home button)" (`UIInterfaceOrientationPortraitUpsideDown`) in addition is also very little work, because the interface that

worked for “Portrait (bottom home button)” will work here as well. (You would never write an app that can launch into portrait-upside-down orientation without also being able to launch into normal portrait mode.)

Now let’s talk about an app whose initial orientation is landscape mode. The cold, hard truth that you must understand is that in fact *all apps launch into portrait mode initially*. This is because the window goes only one way, with its top at the top of the device (away from the home button) — window bounds are screen bounds (see “[What Rotates?](#)” on page 519). No matter what you specify in the *Info.plist* “Supported interface orientations” setting, the app will launch into portrait orientation *first*, and will *then rotate* into your specified orientation. The user won’t necessarily *see* this initial rotation; it may have happened by the time the user sees the app’s actual interface. But it will happen.

Thus, an app whose initial orientation is landscape mode *must be configured to rotate from portrait to landscape* even if it doesn’t support rotation after that. The initial setup of such an app’s interface can be surprisingly tricky, because the interface takes on portrait dimensions before it takes on landscape dimensions; questions of the form “I’m having trouble getting my app to launch into landscape mode” are among the most frequently asked iOS questions on the Internet.

Let’s take the case of an iPhone app whose initial orientation is landscape, with the button on the right:

1. In *Info.plist*, set the “Supported interface orientations” to “Landscape (right home button).”
2. In the root view controller’s code, override `shouldAutorotateToInterfaceOrientation`: along these lines:

```
- (BOOL) shouldAutorotateToInterfaceOrientation:(UIInterfaceOrientation)io {
    return (io == UIInterfaceOrientationLandscapeRight);
}
```
3. If the root view interface requires further configuration in code, give the interface a chance to establish itself and to rotate into position *before* doing that configuration.

It is the neglect of that third step that causes beginners the most trouble (especially because the documentation doesn’t tell you about it). The problem is that when the root view is initially put into the window, its bounds are window bounds; only later, at the redraw moment ([Chapter 14](#)), does the view resize appropriately to the initial orientation you’ve requested. Thus, if you try to configure the interface too early (say, in the root view controller’s `viewDidLoad`), the width and height values of the view’s bounds will be reversed.

For example, let’s say that we’ve taken the first two steps but we’ve neglected the third, and that our root view controller’s `viewDidLoad` code looks like this:

```

- (void)viewDidLoad {
    [super viewDidLoad];
    UIView* square = [[UIView alloc] initWithFrame:CGRectMake(0,0,10,10)];
    square.backgroundColor = [UIColor blackColor];
    square.center = CGPointMake(CGRectGetMidX(self.view.bounds),5); // top center?
    [self.view addSubview:square];
}

```

The app apparently launches into a rotated landscape orientation; the user must hold the device with the home button at the right to see it correctly. That's good. But where's the little black square? *Not* at the top center of the screen! The square appears at the top of the screen, but only about a third of the way across. The trouble is that in order to calculate the x-coordinate of the square's center we examined the view's bounds *too soon*, at a time when the view's x-dimension (its width dimension) was still its shorter dimension.

A solution is to override one of the rotation events discussed in the next section and complete the configuration of your view there. For example, we could override `didRotateFromInterfaceOrientation:`. This method is called for the first time at launch *after* the initial rotation has been performed, so the dimensions of the root view are correct. There's always a chance that `didRotateFromInterfaceOrientation:` will be called again later in the app's lifetime, but we only want to perform our initial view configuration initially, so we need to take precautions against accidentally configuring the interface again; this is a good use of a static variable:

```

- (void) finishInitializingView {
    // static BOOL flag
    static BOOL done = NO;
    if (done)
        return;
    done = YES;
    // the static BOOL flag makes sure the following is performed exactly once
    UIView* square = [[UIView alloc] initWithFrame:CGRectMake(0,0,10,10)];
    square.backgroundColor = [UIColor blackColor];
    square.center = CGPointMake(CGRectGetMidX(self.view.bounds),5);
    [self.view addSubview:square];
}

- (void)didRotateFromInterfaceOrientation:
    (UIInterfaceOrientation)fromInterfaceOrientation {
    [self finishInitializingView];
}

```

When designing in the nib, if the interface is to appear in landscape mode, you can design in landscape mode; select your view and choose Landscape in the Orientation pop-up menu in the Simulated Metrics section of the Attributes inspector. However, this does not really rotate anything; you're merely swapping the view's height and width values, for convenience while editing the nib. The nib knows nothing of orientations, and the view, if it is placed into the window automatically as the app launches, will still have a narrower x-dimension initially. In practice this should cause no difficulty; if the

results at runtime are not what you expect, it may help to use autoresizing settings so that the bits of your interface settle into the correct position and dimensions.

Rotation Events

You can learn the current interface orientation through a view controller's `interfaceOrientation` property. Your `UIViewController` subclass can override any of the following methods (which are called in the order shown) to be alerted in connection with rotation:

`willRotateToInterfaceOrientation:duration:`

The first parameter is the new orientation; `self.interfaceOrientation` is the old orientation, and the view's bounds are the old bounds.

`willAnimateRotationToInterfaceOrientation:duration:`

The first parameter is the new orientation; `self.interfaceOrientation` is the new orientation, and the view's bounds are the new bounds. The call is wrapped by an animation block, so changes to animatable view properties are animated.

`didRotateFromInterfaceOrientation:`

The parameter is the old orientation; `self.interfaceOrientation` is the new orientation, and the view's bounds are the new bounds.

By judicious use of these rotation events, you can achieve some useful effects. For example, you might want your app's interface to look different in landscape mode from how it looks in portrait mode. In simple cases, the autoresizing settings of your view's subviews may suffice; but for more elaborate adjustments, you can perform layout in your implementation of these methods.

In this simple example, our app displays a black rectangle at the left side of the screen if the device is in landscape orientation, but not if the device is in portrait orientation. The example is marked by its defensive posture, quite typical of an event-driven architecture ([Chapter 11](#)) under which it is best to make as few assumptions as possible about what will really happen; and it is careful to work in concert with the rotation animation:

```
@synthesize blackRect = _blackRect;

- (UIView*) blackRect {
    if (!self->_blackRect) {
        if (UIInterfaceOrientationIsPortrait(self.interfaceOrientation))
            return nil;
        CGRect f = self.view.bounds;
        f.size.width /= 3.0;
        f.origin.x = -f.size.width;
        UIView* br = [[UIView alloc] initWithFrame:f];
        br.backgroundColor = [UIColor blackColor];
        self.blackRect = br;
    }
    return self->_blackRect;
}
```

```

}

-(void)willAnimateRotationToInterfaceOrientation:(UIInterfaceOrientation)io
    duration:(NSTimeInterval)duration {
    UIView* v = self.blackRect;
    if (UIInterfaceOrientationIsLandscape(io)) {
        if (!v.superview) {
            [self.view addSubview:v];
            CGRect f = v.frame;
            f.origin.x = 0;
            v.frame = f;
        }
    } else {
        if (v.superview) {
            CGRect f = v.frame;
            f.origin.x -= f.size.width;
            v.frame = f;
        }
    }
}

- (void) didRotateFromInterfaceOrientation:(UIInterfaceOrientation)io {
    if (UIInterfaceOrientationIsPortrait(self.interfaceOrientation))
        [self.blackRect removeFromSuperview];
}

```

We have a `UIView` property, `blackRect`, to retain the black rectangle; we implement its getter to create the black rectangle if it hasn't been created already, but only if we are in landscape orientation, since otherwise we cannot set the rectangle's dimensions properly. The implementation of `willAnimateRotationToInterfaceOrientation:duration:` slides the black rectangle in from the left as part of the rotation animation if we have ended up in a landscape orientation, but only if it isn't in the interface already; after all, the user might rotate the device 180 degrees, from one landscape orientation to the other. Similarly, it slides the black rectangle out to the left if we have ended up in a portrait orientation, but only if it *is* in the interface already. Finally, `didRotateFromInterfaceOrientation:`, called after the rotation animation is over, makes sure the rectangle is removed from its superview if we have ended up in a portrait orientation.

Presented View Controller

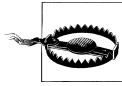
Aside from responding to device rotation, the chief purpose of view controllers to make views come and go coherently in the interface. By now, you're probably eager to make some views come and go. The rock-bottom simplest way of making views come and go is through a *presented view controller*.

In iOS 4 and before, the view provided by a presented view controller was called a *modal view*. To make a modal view appear, you'd send a view controller whose view occupies the interface the `presentModalViewController:animated:` message, handing it a second view controller. This second view controller was the modal view controller,

and its view would now appear as a modal view. The second view controller became the first view controller's `modalViewController`; the first view controller was the second view controller's `parentViewController`. This state of affairs persisted until `dismissModalViewControllerAnimated:` was called.

In iOS 5, although `presentModalViewController:animated:` and `dismissModalViewControllerAnimated:` are not officially deprecated, they are considered outmoded and are slated for official deprecation in the future. In this edition of the book, I won't use them, nor will I refer to modal view controllers. The term "modal" still plays a role because of legacy terminology, but, as much as possible, I will keep to the new style of terminology, which speaks of a *presenting view controller* and a *presented view controller*.

In iOS 5, then, to make a view controller present another view controller, you send the first view controller `presentViewController:animated:completion:`, handing it the second view controller, which you will probably instantiate for this very purpose. The first view controller is now the second view controller's `presentingViewController`; the second view controller is the first view controller's `presentedViewController`, and the first view controller is retaining the second view controller. This state of affairs persists until the first view controller is sent `dismissViewControllerAnimated:completion:`; afterwards, the first view controller releases the second view controller, which will typically go out of existence together with its view, its child view controllers and *their* views, and so on.



The fact that in iOS 5 a presenting view controller is *not* the `parentViewController` of the presented (modal) view controller may break your existing code from iOS 4 or before when you compile it with iOS 5 as the Base SDK.

When one view controller presents another, the view of the presenting view controller is automatically replaced in the interface with the view of the presented view controller, optionally with animation as the transition takes place ([Figure 19-6](#)). (On the iPad, the view of the presented view controller can instead optionally appear in front of the view of the presenting view controller; I'll talk about that later.) When the presented view controller is dismissed, the view of the presented view controller is automatically replaced in the interface with the view of the presenting view controller, optionally with animation. The `completion:` parameter in `presentViewController:animated:completion:` and `dismissViewControllerAnimated:completion:` lets you supply a block of code to be run after the view replacement has occurred.

View controller presentation and dismissal is therefore a way of swapping out interface *completely*. More often than not, this will be the *entire* interface; on the iPhone, it will *always* be the entire interface (in effect, the root view controller is always the presenting view controller on the iPhone).

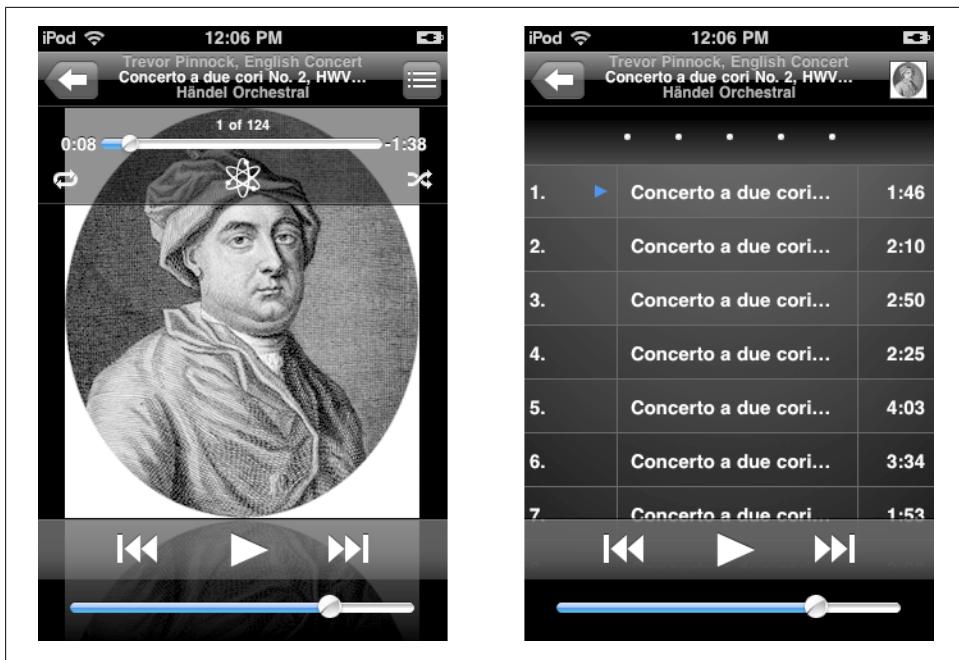


Figure 19-7. Two views that are equal partners

Because one view controller actively presents and the other is passively presented, there is a temptation to think of the presented view controller as secondary or temporary. This is often true, but such a conception misses the full power of a presented view controller, which is that it completely changes what the user sees. A presented view controller might never be dismissed; its view's interface might be complex; it might itself have child view controllers (Figure 19-6); it might itself present yet another view controller. You, the programmer, may be conscious that the presented view controller's view is in some sense a cover for the presenting view controller's view; but the user might not experience the interface that way. Perhaps it is better to think of view controller presentation as a simple, straightforward way of transitioning from one view (what a storyboard would call one *scene*) to another.

For example, in Apple's own Music app, the two alternating views that appear when you view the currently playing song are more like equal partners (Figure 19-7); there's no sense that one is secondary to the other. Yet it's likely that these are the views of a presenting view controller and a presented view controller.

Let's make one view controller present another. We already have an example project, from earlier in this chapter, containing an AppDelegate class and a RootViewController class. Let's modify it to add a second view controller class, and make RootViewController present it (don't use the project containing a storyboard; I'll talk about storyboards and presented view controllers later):

1. Choose File → New → New File and make a new Cocoa Touch UIViewController subclass. Click Next.
2. Name the class SecondViewController, make sure it is a subclass of UIViewController, and *check the checkbox* “With XIB for user interface” so that we can design this view controller’s view quickly and easily in a nib. Click Next.
3. Create the files.
4. Do something in *SecondViewController.xib* to make the view distinctive, so that you’ll know it when it appears.
5. We need a way to trigger the presenting of SecondViewController. Back in RootViewController’s view’s interface, add a button. Connect that button to an action method in *RootViewController.m*; let’s call it `doPresent:`.

Now we’ll write the code for `doPresent:`. First, import “*SecondViewController.h*” at the top of *RootViewController.m*, so that we can speak of SecondViewController. Here’s the code:

```
- (IBAction)doPresent:(id)sender {
    [self presentViewController:[[SecondViewController alloc] init]
                           animated:YES completion:nil];
}
```

Run the project. In RootViewController’s view, tap the button. SecondViewController’s view slides into place over RootViewController’s view.

In our lust for instant gratification, we have neglected to provide a way to dismiss the presented view controller. If you’d like to do that, put a button into SecondViewController’s view and connect it to an action method in *SecondViewController.m*:

```
- (IBAction)doDismiss:(id)sender {
    [self.presentingViewController dismissViewControllerAnimated:YES
                                                    completion:nil];
}
```

Run the project. You can now alternate between RootViewController’s view and SecondViewController’s view.



Instead of sending `dismissViewControllerAnimated:completion:` to the presenting view controller, our presented view controller could have sent it to `self`. The call is then automatically forwarded to the presenting view controller. However, this fact is not officially documented, so I’m not sure whether it’s safe to rely on it in future.

In real life, it is quite probable that both presentation and dismissal will be a little more involved than in that simple example. The presenting view controller will very likely have additional information to pass on to the presented view controller. Here’s a typical example from one of my apps (this is in fact the transition that engenders [Figure 19-6](#)):

```
DrillViewController* dvc = [[DrillViewController alloc] initWithData:drillTerms];
[self presentViewController:dvc animated:YES completion:nil];
```

As you can see, I've constructed `DrillViewController` to have a designated initializer `initWithData:` for this very purpose, so that the presenting view controller can pass it the data it will need to do its job while it exists.

The presented view controller, too, will very likely have additional information to pass back to the presenting view controller. The user is interacting with the presented view controller's view, so (as in our example above) it is the presented view controller that knows when it should be dismissed. It also knows what happened while it was in existence. So the presented view controller may need a way to turn to the presenting view controller and hand back some data as it tells the presenting view controller to dismiss it — and this might even involve calling some method in the presenting view controller. Moreover, identifying the presenting view controller from within the presented view controller might not be a simple task. What if there's a chain of presented view controllers?

The presented view controller may thus need a reference to the presenting view controller and a knowledge of at least some methods implemented by the presenting view controller's class. A standard architecture that solves this problem is for the presented view controller to define a protocol to which the presenting view controller conforms. The presenting view controller then hands the presented view controller a (weak!) reference to itself as it creates the presented view controller. In this way the presented view controller is the one that specifies what the communication callbacks will be, and it remains agnostic about the actual class of the presenting view controller. This is the architecture, exemplified by the Utility Application project template, that I discussed in [Chapter 10](#) (at the end of the Protocols section).

To implement this architecture in our existing example with `RootViewController` and `SecondViewController`, you'd modify `SecondViewController` to look like this:

```
// [in SecondViewController.h]
@protocol SecondViewControllerDelegate
- (void)dismissSecondViewControllerWithData: (id) data;
@end
@interface SecondViewController : UIViewController
@property (nonatomic, weak) id<SecondViewControllerDelegate> delegate;
@end

// [in SecondViewController.m]
@synthesize delegate;
- (IBAction)doDismiss:(id)sender {
    [self.delegate dismissSecondViewControllerWithData:nil];
}
```

`RootViewController` would then declare itself in `RootViewController.h` as adopting `SecondViewControllerDelegate`; observe that in order to do this, `RootViewController.h` must import "`SecondViewController.h`" (because that's where the protocol is de-

fined). RootViewController’s code would then present and dismiss SecondViewController like this:

```
- (IBAction)doPresent:(id)sender {
    SecondViewController* svc = [[SecondViewController alloc] init];
    svc.delegate = self;
    [self presentViewController:svc animated:YES completion:nil];
}
- (void)dismissSecondViewControllerWithData:(id)data {
    // ... do something with the data here ...
    [self dismissViewControllerAnimated:YES completion:nil];
}
```

To be sure, configuring this architecture involves considerable work, and I know from experience that there is a strong temptation to be lazy and avoid it. It may indeed be possible to get by with a simplified solution. For example, SecondViewController might know that it will be presented only by a RootViewController, and can thus import "RootViewController.h", cast its `presentingViewController` to a RootViewController, and call any RootViewController method. Or SecondViewController could post a notification for which RootViewController has registered. Nevertheless, a protocol is the fullest and most correct architecture for a presented view controller to communicate back to its presenter.

I am not saying that a simpler solution would never be possible or that it would never be right. One can imagine situations where a notification would arguably be a better solution. Perhaps there is a deep chain of presented views, and perhaps different user actions should cause dismissal at different levels up the chain; thus, the deepest view could not have a single delegate. Perhaps there is no point at which the presenter who must perform the dismissal is ever in direct contact with the view controller that handles the user action; in the previous example, they were in contact because the one created the other, but that might not be the case. Then, I think, a notification is perfectly reasonable. This all part of the larger topic of getting a reference, discussed in [Chapter 13](#). No one size fits all.

Presented View Animation

When a view is presented and later when it is dismissed, an animation can be performed, according to whether the `animated:` parameter of the corresponding method is YES. The possible animation styles (whose names preserve the legacy “modal” designation) are:

`UIModalTransitionStyleCoverVertical` (*the default*)

The presented view slides up from the bottom to cover the presenting view on presentation and down to reveal the presenting view on dismissal. (“Bottom” is defined differently depending on the orientation of the device and the orientations the view controllers support.)

UIModalTransitionStyleFlipHorizontal

The view flips on the vertical axis as if the two views were the front and back of a piece of paper. (The “vertical axis” is the device’s long axis, regardless of the app’s orientation.)

UIModalTransitionStyleCrossDissolve

The views remain stationary, and one fades into the other.

UIModalTransitionStylePartialCurl

The first view curls up like a page in a notepad to expose most of the second view, but remains covering the top left region of the second view. Thus there must not be any important interface in that region, as the user will not be able to see or touch it. This option seems more appropriate on the larger iPad screen, though it is legal on the iPhone.

You do not pass the animation style as a parameter when presenting or dismissing a view controller; rather, it is attached beforehand to a view controller as its `modalTransitionStyle` property. (It is legal, but not common, for the `modalTransitionStyle` value to differ at the time of dismissal from its value at the time of presentation; reversing on dismissal the same animation style that was used on presentation is a subtle cue to the user that we’re returning to a previous state.) The view controller that should have this `modalTransitionStyle` property set will generally be the *presented view controller* (I’ll talk about the exception to this rule in a moment).

Under this architecture, if the presenting view controller wants to dictate the transition style, it should set the presented view controller’s `modalTransitionStyle` property right after it instantiates the presented view controller. Otherwise, if the presented view controller wants to dictate its own transition style, it should set its own `modalTransitionStyle` property early in its lifetime; for example, it might override `initWithNibName:bundle:`. If the presented view controller is instantiated in a nib, there’s a Transition Style pop-up menu.

Presentation Styles

On the iPhone, the presented view controller’s view always replaces the presenting view controller’s view. On the iPad, there are additional options, in accordance with the iPad’s larger screen, the existence of popovers, and the extended ability to do things with custom parent controllers. These options are expressed through `UIViewController`’s `modalPresentationStyle` property.

The `modalPresentationStyle` property must be set on the presented view controller. As with `modalTransitionStyle` (discussed in the previous section), the presented view controller will want to set this value early in its lifetime if it isn’t being set by the presenting view controller.

Your choices (which display more legacy “modal” names) are:

UIModalPresentationFullScreen

The default. This is the only mode in which `UIModalTransitionStylePartialCurl` is legal. On the iPhone, although it is not illegal to set the `modalPresentationStyle` to another value, a presented view will always behave as if it were `UIModalPresentationFullScreen`; this amounts to another way of saying that on the iPhone the root view controller is always the presenting view controller.

UIModalPresentationPageSheet

In a portrait orientation, basically indistinguishable from fullscreen mode. But in a landscape orientation, the presented view has the width of the portrait-oriented screen, so the presenting view controller's view remains partially visible behind the presented view, but dimmed so that the user can't interact with it. Thus this mode is very like a modal dialog on Mac OS X.

UIModalPresentationFormSheet

Similar to `UIModalPresentationPageSheet`, but the presented view is smaller. As the name implies, this intended to allow the user to fill out a form (Apple describes this as “gathering structured information from the user”).

UIModalPresentationCurrentContext

The presented view replaces the view of the presenting view controller. This is useful with popovers, for instance (see [Chapter 22](#)), and with parent controllers where a child view controller occupies only a portion of the screen — in which case, the presented view can replace just that portion of the screen.

On the iPad, when the presented view controller's `modalPresentationStyle` is `UIModalPresentationCurrentContext`, a decision has to be made, when a presenting view controller is sent `presentViewController:animated:completion:`, as to what view controller's view should be replaced by the presented view controller's view. In iOS 5, the way this decision is made involves another `UIViewController` property, `definesPresentationContext` (a BOOL). Here's what happens: starting with the view controller to which `presentViewController:animated:completion:` was sent, we walk up the chain of parent view controllers, looking for one whose `definesPresentationContext` property is YES. If we find one, that is the one whose view will be replaced by the presented view controller's view. If we *don't* find one, the root view controller's view is replaced, just as if the presented view controller's `modalPresentationStyle` had been `UIModalPresentationFullScreen`.

Moreover, if, during the search just described, we *do* find a view controller whose `definesPresentationContext` property is YES, we look to see if that view controller's `providesPresentationContextTransitionStyle` property is *also* YES. If so, that view controller's `modalTransitionStyle` (defining the style of the transition animation when one view replaces the other, as discussed in the previous section) is used for this transition animation, instead of using the presented view controller's `modalTransitionStyle`.

To illustrate, I need a parent–child view controller pair to work with. This chapter hasn’t yet discussed any parent view controllers in detail, but the simplest is UITabBarController, which I discuss in the next section, and it’s easy to create a working app with a UITabBarController-based interface, so that’s the example I’ll use.

Start with a Universal version of the Tabbed Application project template. Make a new view controller class and an accompanying nib file to use as a presented view controller. Put a button in the first view controller’s view and connect it to an action method in the first view controller that summons the new view controller as a presented view controller. For example, if the new view controller’s class is ExtraViewController, you’d say this:

```
- (IBAction)doPresent:(id)sender {
    UIViewController* vc = [ExtraViewController new];
    [self presentViewController:vc animated:YES completion:nil];
}
```

Run the project. You’ll find that the presented view controller occupies the *entire* interface, covering even the tab bar. The presented view replaces the root view.

Now change the code to look like this:

```
- (IBAction)doPresent:(id)sender {
    UIViewController* vc = [ExtraViewController new];
    self.definesPresentationContext = YES;
    vc.modalPresentationStyle = UIModalPresentationCurrentContext;
    [self presentViewController:vc animated:YES completion:nil];
}
```

Run the project. This time, the presented view controller replaces only the first view controller’s view; the tab bar remains. That’s because the presented view controller’s `modalPresentationStyle` is `UIModalPresentationCurrentContext`, and when `presentViewController:animated:completion:` is sent to `self`, the `definesPresentationContext` property of `self` is YES. The search for a context stops, and the presented view replaces the first view controller’s view instead of the root view.

This is even clearer if we change the transition animation. Instead of changing the `modalTransitionStyle` of the presented view controller, `vc`, we can do this by changing the `modalTransitionStyle` of the presenting view controller, `self`. Add two more lines, like this:

```
- (IBAction)doPresent:(id)sender {
    UIViewController* vc = [ExtraViewController new];
    self.definesPresentationContext = YES;
    self.providesPresentationContextTransitionStyle = YES;
    self.modalTransitionStyle = UIModalTransitionStyleFlipHorizontal;
    vc.modalPresentationStyle = UIModalPresentationCurrentContext;
    [self presentViewController:vc animated:YES completion:nil];
}
```

Now the transition animation uses the flip horizontal animation; the presenting view controller is able to override the transition animation of the presented view controller.

It can do this only because all three conditions are met: the presented view controller's `modalPresentationStyle` is `UIModalPresentationCurrentContext`, the `definesPresentationContext` property of `self` is YES, and the `providesPresentationContextTransitionStyle` property of `self` is YES as well. The flip horizontal animation shows clearly what part of the interface is being replaced by the presented view.

Observe also that you can still switch between the first and second tabbed views, even while the presented view is occupying the place of the first tabbed view. Clearly, very powerful and interesting interface can be constructed using this technique.

It's helpful to experiment with the above code, commenting out individual lines to see what effect they have on the overall result. Also, run the code on the iPhone simulator and observe that none of this works; the `UIModalPresentationCurrentContext` presentation style, on which this entire behavior depends, is an iPad-only feature.

Presented Views and Rotation

No law requires that every "scene" of your interface should appear in the same rotation. On the iPhone especially, where the user can easily rotate the device while working with an app, it is reasonable and common for one scene to appear in portrait orientation and another to appear in landscape orientation.

One easy way to achieve this is to implement `shouldAutorotateToInterfaceOrientation:` differently for a presented view controller. When a view controller is presented, its `shouldAutorotateToInterfaceOrientation:` is consulted and used. If the device is not in an orientation that the presented view controller approves of, the app will rotate and the user will have to rotate the device to see it correctly.

For example, in my flashcard app pictured in [Figure 19-3](#), the flashcards are viewed only in landscape orientation. But there is also an option to display a list (a `UITableView`) of all flashcards. This list is far better viewed in portrait orientation, so as to accommodate the greatest possible number of items on the screen at once; therefore, it is permitted to assume portrait orientation only. The user must rotate the device with the hand holding the iPhone, but this is not objectionable; in fact, it quickly becomes automatic and subconscious.

This is achieved by implementing `shouldAutorotateToInterfaceOrientation:` this way in one view controller:

```
// [view controller A]
- (BOOL) shouldAutorotateToInterfaceOrientation:(UIInterfaceOrientation)io {
    return (io == UIInterfaceOrientationLandscapeRight);
}
```

Meanwhile, it is implemented this way in another view controller:

```
// [view controller B]
- (BOOL) shouldAutorotateToInterfaceOrientation:(UIInterfaceOrientation)io {
    return (io == UIInterfaceOrientationPortrait);
}
```

Then when view controller A's view appears, it will appear in landscape orientation. View controller A, on demand, presents view controller B; when view controller B's view appears, it will appear in portrait orientation.



If the presented view controller's presentation style is `UIModalPresentationPageSheet` or `UIModalPresentationFormSheet`, and if its `shouldAutorotateToInterfaceOrientation:` forces the user to rotate the device, the presenting view controller does *not* get any rotation events. This can result in a weird interface, so don't do that. The presented view controller's acceptable orientations should match those of the presenting view controller.

A presented view controller with a different preferred orientation can also be used as a way of rotating the interface in response to rotation of the device. Instead of performing complex layout on rotation, as in “[Rotation Events](#)” on page 522, we forbid the rotation of our view (in `shouldAutorotateToInterfaceOrientation:`), detect the rotation of the *device* instead, and replace our view with a presented view suited to the new orientation.

In this example, I assume that the root view controller (where this code is) has its `shouldAutorotateToInterfaceOrientation:` set to return `io==UIInterfaceOrientationPortrait`, while the alternate `LandscapeViewController` is set to return `UIInterfaceOrientationIsLandscape(io)`:

```
- (void)viewDidLoad {
    [super viewDidLoad];
    [[UIDevice currentDevice] beginGeneratingDeviceOrientationNotifications];
    [[NSNotificationCenter defaultCenter] addObserver:self
        selector:@selector(screenRotated:)
        name:UIDeviceOrientationDidChangeNotification
        object:nil];
}

- (void) screenRotated: (id) notif {
   NSUInteger rot = [[UIDevice currentDevice] orientation];
    if (UIDeviceOrientationIsLandscape(rot) && !self.presentedViewController) {
        [[UIApplication sharedApplication]
            statusBarOrientation:rot animated:YES];
        LandscapeViewController *c = [[LandscapeViewController alloc] init];
        c.modalTransitionStyle = UIModalTransitionStyleCrossDissolve;
        [self presentViewController:c animated:YES completion:nil];
    } else if (rot==UIDeviceOrientationPortrait && self.presentedViewController) {
        [[UIApplication sharedApplication]
            statusBarOrientation:rot animated:YES];
        [self dismissViewControllerAnimated:YES completion:nil];
    }
}
```

The judicious use of `setStatusOrientation:animated:` nets us the nice status bar animation exactly on those occasions when the system does not hand it to us, so that

we get it no matter how the user rotates the device. The result is basically indistinguishable from “normal” interface rotation.

Tab Bar Controllers

A *tab bar* (UITabBar, see also [Chapter 25](#)) is a horizontal bar containing items. Each item (a UITabBarItem) displays, by default, an image and a title. At all times, exactly one of these items is selected (highlighted); when the user taps an item, it becomes the selected item.

If there are too many items to fit on a tab bar, the excess items are automatically subsumed into a final More item. When the user taps the More item, a list of the excess items appears, and the user can select one; the user can also be permitted to edit the tab bar, determining which items appear in the tab bar itself and which ones spill over into the More list.

A tab bar is an independent interface object, but it is most commonly used in conjunction with a *tab bar controller* (UITabBarController, a subclass of UIViewController) to form a tab bar interface. The tab bar controller displays the tab bar at the bottom of its own view. From the user’s point of view, the tab bar items correspond to views; when the user selects a tab bar item, the corresponding view appears. These views are “scenes”; the user is employing the tab bar to choose an entire area of your app’s functionality. In reality, the UITabBarController is a parent view controller; you give it child view controllers, which the tab bar controller then contains, and the scenes are the views of those child view controllers.

A UITabBarController thus puts an organizational front end on the other interface that your app offers:

The tab bar is navigational

The user chooses an item from the tab bar to change scenes.

The tab bar is informative

The tab bar, possibly including the More item, tells the user what scenes are available, and the highlighting shows which scene is currently displayed.

Familiar examples of a tab bar interface on the iPhone are Apple’s Clock app, which has four tab bar items, and Apple’s Music app, which has four tab bar items plus a More item that reveals a list of six more.

On the small iPhone screen, you may feel that using a tab bar interface involves a certain degree of sacrifice. Some screen real estate is occupied by the tab bar, reducing the amount of space available for the views that it summons by about 50 pixels at the bottom, and the tab bar itself cannot readily be hidden. Don’t forget, however, about presented view controllers! A tab bar controller can present a view controller whose view occupies the whole screen, replacing the tab bar interface with a completely dif-

ferent interface. And a tab bar controller can itself be a presented view controller. Thus, there is no requirement that your app's *whole* interface should be a tab bar interface.

When using a tab bar interface by way of a UITabBarController, you do not interact (as a programmer) with the tab bar itself; you don't create it or set its delegate. You provide the UITabBarController with children, and it does the rest; when the UITabBarController's view is displayed, there's the tab bar along with the view of the selected item. Starting in iOS 5, however, you can customize the *look* of a tab bar (see [Chapter 25](#) for details); you're no longer confined to the black gradient look that characterized a tab bar in the past. You can get a reference to the tab bar controller's tab bar through its `tabBar` property.

As mentioned earlier in this chapter, you should not subclass UITabBarController (nor should you need to).

Since you don't subclass UITabBarController, you can't override its `shouldAutorotateToInterfaceOrientation:`. Instead, a tab bar interface's orientation is determined by the orientations of its child view controllers. If the device is rotated to a certain orientation, and all the child view controllers permit rotation to that orientation, the tab bar interface itself will rotate; similarly, if all the child view controllers agree in permitting rotation to a certain orientation, they can force the tab bar interface to assume that orientation.

Tab Bar Items

A tab bar item is a UITabBarItem; this is a subclass of UIBarItem, an abstract class that provides some of its most important properties, such as `title`, `image`, and `enabled`. For each view controller you assign as a tab bar controller's child, you're going to need a tab bar item, which will appear in the tab bar. This tab bar item will be your child view controller's `tabBarItem`.

There are two ways to make a tab bar item:

By borrowing it from the system

Instantiate UITabBarItem using `initWithTabBarSystemItem:tag:`, and assign the instance to your child view controller's `tabBarItem`. Consult the documentation for the list of available system items. Unfortunately you can't customize a system tab bar item's title; you must accept the title the system hands you. (You can't work around this annoying restriction by somehow copying a system tab bar item's image.)

By making your own

Instantiate UITabBarItem using `initWithTitle:image:tag:`. Alternatively, use the view controller's existing `tabBarItem` and set its `image` and `title`. Instead of setting the `title` of the `tabBarItem`, you can set the `title` property of the view controller itself; setting a view controller's `title` automatically sets the `title` of its current

`tabBarItem` (unless the tab bar item is a system tab bar item), but the converse is not true.

The `image` for a tab bar item should be a 30×30 PNG (if it is larger, it will be scaled down as needed). It should be a *transparency mask*; that is, it should consist of transparent pixels and opaque pixels (possibly including semiopaque pixels). Color is of no consequence and will be ignored; all that matters is the degree of transparency of each pixel. The runtime itself will color the image, adding a shine effect.

Starting in iOS 5, you can instead provide a normal image by calling `setFinishedSelectedImage:withFinishedUnselectedImage:`; the runtime will not modify this image in any way, and getting the size right is up to you. The selected and unselected image can be the same, but the runtime will not tint the selected image (as it does for an `image`), so you'll probably want two different images, to differentiate the two states. (See [Chapter 15](#) on how to use Core Image to tint or desaturate an image.)

You can also give a tab bar item a badge (see the documentation on the `badgeValue` property). Other ways in which you can customize the look of a tab bar item are discussed in [Chapter 25](#). For example, you can control the font and style of the title, or you can give it an empty title and offset the image.

Configuring a Tab Bar Controller

As I've already said, you configure a tab bar controller by handing it the view controllers that will be its children. To do so, collect those view controllers into an array and set the `UITabBarController`'s `viewControllers` property to that array. The view controllers in the array are now the tab bar controller's child view controllers; the tab bar controller is the `parentViewController` of the view controllers in the array. The tab bar controller is also the `tabBarController` of the view controllers in the array and of all their children; thus a child view controller at any depth can learn that it is contained by a tab bar controller and can get a reference to that tab bar controller. The tab bar controller retains the array, and the array retains the child view controllers.

The tab bar controller's tab bar will automatically display the `tabBarItem` of each child view controller. The order of the tab bar items is the order of the view controllers in the tab bar controller's `viewControllers` array. Thus, the child view controllers will probably want to configure their `tabBarItem` property very early in their lifetime, so that the `tabBarItem` is ready by the time the view controller is handed as a child to the tab bar controller. It is common to override `initWithNibName:bundle:` for this purpose.

Here's a simple example excerpted from the app delegate's `application:didFinishLaunchingWithOptions:` of one of my apps, in which I construct a tab bar interface and display it:

```
UITabBarController* tbc = [[UITabBarController alloc] init];
// create tabs
UIViewController* b = [[[GameBoardController alloc] init];
// some code omitted here... never mind what "s" is in the next line
```

```

UINavigationController* n =
    [[[UINavigationController alloc] initWithRootViewController:s];
// load up with tab views
tbc.viewControllers = [NSArray arrayWithObjects: b, n, nil];
// configure window
self.window.rootViewController = tbc;

```

You'll notice that I don't configure the contained view controllers' tab bar items. That's because those view controllers configure themselves early in their lifetimes. For example:

```

// [GameBoardController.m]

- (id) initWithNibName:(NSString *)NibName bundle:(NSBundle *)nibBundleOrNil {
    self = [super initWithNibName:nibNameOrNil bundle:nibBundleOrNil];
    if (self) {
        // we will be embedded in a tab bar interface, configure
        self.tabBarItem.image = [UIImage imageNamed:@"game.png"];
        self.title = @"Game";
    }
    return self;
}

```

If you change the tab bar controller's view controllers array later in its lifetime and you want the corresponding change in the tab bar items to be animated, call `setViewControllers:animated:`.

When a child view controller's view is displayed, it is resized to fit the region of the tab bar controller's view above the tab bar. Keep that fact in mind when designing your view. Autoresizing settings can help here, if you don't want an interface item near the bottom of your view to be left behind and disappear from the interface when the view is resized. Also, when editing your view in the nib editor, you can shrink the view to the size at which it will be displayed in the tab bar interface, thus helping you judiciously situate your interface items; to do so, choose Tab Bar in the Bottom Bar pop-up menu of the Simulated Metrics section of the Attributes inspector.

Initially, by default, the first child view controller's tab bar item is selected and its view is displayed. To tell the tab bar controller which tab bar item should be selected, you can couch your choice in terms of the contained view controller, by setting the `selectedViewController` property, or by index number in the array, by setting the `selectedIndex` property. The same properties also tell you what view controller's view the user has displayed by tapping in the tab bar.

You can also set the UITabBarController's delegate; the delegate should adopt the UITabBarControllerDelegate protocol. The delegate gets messages allowing it to prevent a given tab bar item from being selected, and notifying it when a tab bar item is selected and when the user is customizing the tab bar from the More item.

If the tab bar contains few enough items that it doesn't need a More item, there won't be one and the tab bar won't be user-customizable. If there *is* a More item, you can exclude some tab bar items from being customizable by setting the `customizableView-`

`Controllers` property to an array that lacks them; setting this property to nil means that the user can see the More list but can't rearrange the items. Setting the `viewControllers` property sets the `customizableViewControllers` property to the same value, so if you're going to set the `customizableViewControllers` property, do it *after* setting the `viewControllers` property. (If you do allow the user to rearrange items, you would presumably want to respond in the delegate to this rearrangement by saving the new arrangement in the NSUserDefaults and using it the next time the app runs.) The `moreNavigationController` property can be compared with the `selectedViewController` property to learn whether the user is currently viewing the More list; apart from this, the More interface is mostly out of your control, but I'll discuss some sneaky ways of customizing it in [Chapter 25](#).

You can also configure a UITabBarController in a nib or storyboard. The nib editor interface is quite clever about this. The UITabBarController's contained view controllers can be set directly in the nib. Moreover, each contained view controller contains a Tab Bar Item; you can select this and set its title and image directly in the nib. (If a view controller in a nib *doesn't* have a Tab Bar Item and you want to configure this view controller for use in a tab bar interface, drag a Tab Bar Item from the Object library onto the view controller.) The UITabBarController itself has a `delegate` outlet. Thus, it is possible to create a fully configured tab bar interface with essentially no code at all.

Navigation Controllers

A *navigation bar* (UINavigationBar, see also [xref:chap_id25](#)) is a horizontal bar displaying, at its simplest, a center title and a right button. When the user taps the right button, the navigation bar animates, sliding its interface out to the left and replacing it with a new interface that enters from the right, displaying a back button at the left side, and a new center title — and possibly a new right button. Thus the user can now either go further forward (to the right), tapping the right button to proceed to yet another center title, or else go back (to the left), tapping the back button to return to the first interface, with the first center title and the first right button.

There's a computer science name for the architecture I'm describing — a stack. Conceptually, a navigation bar represents a stack. Under the hood, a navigation bar really *is* a stack. A navigation bar holds an internal stack of navigation items (UINavigationItem). It starts out with one navigation item (the *root* or *bottom item*); you can then push another navigation item onto the stack, and from there you can either pop that navigation item to remove it from the stack or push yet *another* navigation item onto the stack.

At any moment, therefore, some navigation item is the *top item* on the stack, the most recently pushed item still present on the stack (the `topItem`). Furthermore, unless the top item is also the root item (because it is the only item in the stack), some navigation

item is the *back item* (the `backItem`), the item that would be top item if we were now to pop the top item.

The state of the stack is reflected in the navigation bar's interface. The navigation bar's center title comes automatically from the top item, and its back button comes from the back item. (See [Chapter 25](#) for a complete description.) Thus, typically, the center shows the user what item is current, and the left side is a button telling the user what item we would return to if the user were to tap that button. The animations reinforce this notion of directionality, giving the user a sense of position within a chain of items. When a navigation item is pushed onto the stack, the new interface slides in from the right; when an item is popped from the stack, the new interface slides in from the left.

A navigation bar is an independent interface object, but it is most commonly used in conjunction with a *navigation controller* (`UINavigationController`, a subclass of `UIViewController`) to form a navigation interface. To each navigation item in the navigation bar corresponds a “scene” — a view controller and its view. Just as there is a stack of navigation items in the navigation bar, there is a stack of view controllers in the navigation controller. These view controllers are the navigation controller's children, and the navigation items each belong a view controller; each navigation item is, in fact, a view controller's `navigationItem`. Whatever view controller is at the top of the navigation controller's stack, that is the view controller whose `navigationItem` is displayed in the navigation bar, and whose view is displayed in the interface as a whole. The animation in the navigation bar is matched by an animation of the interface as a whole: a view controller's view slides into the main interface from the left or right just as its navigation item slides into the navigation bar from the left or right.

If a tab bar interface is like offering the user a choice from a hand of cards, a navigation interface is like dealing one card at a time off the top of the deck. The user starts by looking at a single card, face up. When the user navigates to the right, this is like asking for another card, which now lies face up on top of the first card. At this point the user can perhaps ask for yet another card. When the user navigates to the left, this is like taking the topmost face-up card and returning it to the dealer; the user can repeat this until only the original card is left.

In real life, a navigation interface is far more flexible than what I have so far described. The user can indeed sometimes navigate to the right by tapping the right button in the navigation bar, but far more often the user will navigate to the right by tapping something inside the main interface, such as a listing in a table view. ([Figure 19-1](#) is a navigation interface that works this way.) It is also possible to allow the user to navigate back to the left by tapping something inside the main interface, though this is much rarer; the convention of going back by tapping whatever is leftmost in the navigation bar is ingrained in users. It is even possible for a navigation interface to hide the navigation bar, letting the user navigate to the right and later back to the left solely by means of interface objects within the main interface. (Conversely, it is not uncommon, especially on the iPhone, to put a view controller inside a navigation controller just to get

the convenience of the navigation bar, with its title and buttons, even when no actual push-and-pop navigation is going to take place.)

When using a navigation interface by way of a `UINavigationController`, you do not interact (as a programmer) with the navigation bar itself; you don't create it or set its delegate. You provide the `UINavigationController` with children, and it does the rest, handing each child view controller's `navigationItem` to the navigation bar for display and showing the child view controller's view each time navigation occurs. In iOS 5, however, you can customize the *look* of a navigation bar far more than in the past (see [Chapter 25](#) for details). You can get a reference to the navigation controller's navigation bar through its `navigationBar` property.



Prior to iOS 5, the features of a navigation item, and hence the possible contents of a navigation bar, were rather limited. You could have a back button or a left button, a center title or view, and right button, and that was pretty much all. In iOS 5, however, there can be multiple left buttons and multiple right buttons.

A navigation interface may also optionally display a toolbar at the bottom. A toolbar (`UIToolbar`) is a horizontal view displaying a row of items, any of which the user can tap. The tapped item may highlight momentarily but is not selected; it represents the initiation of an action, not a state or a mode, and should be thought of as (and may in fact look like) a button. In iOS 5, the look of the toolbar can be greatly customized ([Chapter 25](#)).

(A `UIToolbar` can be used independently, and often is. It then typically appears at the bottom on an iPhone — [Figure 19-3](#) has a toolbar at the bottom — but often appears at the top on an iPad, where it plays something of the role that the menu bar plays on the desktop. When a toolbar is displayed by a navigation controller, though, it always appears at the bottom.)

In a navigation interface, the contents of the toolbar are linked to the view controller that is currently the top item in the stack: they are its `toolbarItems`. The toolbar can also readily be hidden or shown as a certain view controller becomes the top item. Thus, the toolbar appears to the user to be associated with a particular view, not with the navigation interface as a whole.

A navigation controller is thus tremendously flexible:

- A view controller is not pushed onto the stack until it is needed for display. Thus, you get to decide in real time what the next view should be; typically, you won't even create the next view controller until the user asks to navigate to the right. This makes a navigation controller perfect for a master–detail interface, in which the user sees a list of possibilities and taps one to navigate to the detailed view of that thing ([Figure 19-1](#)).

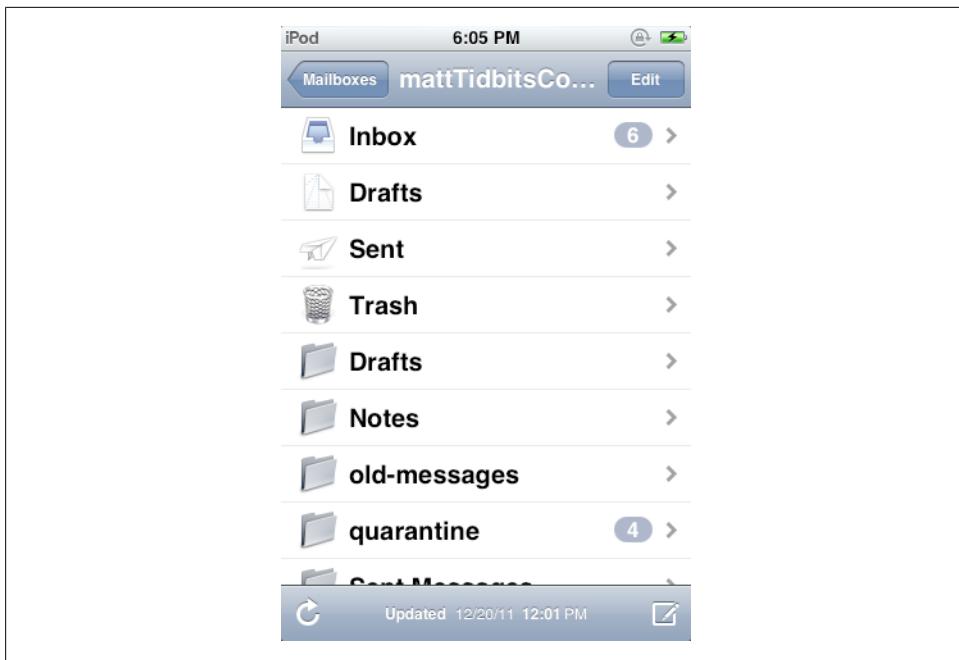


Figure 19-8. A familiar navigation interface

- Both the navigation bar and the toolbar are optional, in the sense that they can be hidden. Moreover, the look of both the navigation bar and the toolbar can be customized (and in iOS 5, they can *really* be customized). Thus, not all navigation interfaces look the same.
- The toolbar's contents and the navigation bar's contents respond automatically to the advent of a different view, and can also be changed in code.
- A navigation interface can readily be used in many ways: as the root view of your app, as a presented view, or as a view contained by a tab bar controller.

A familiar example of a navigation interface is Apple's Mail app (Figure 19-8), a master–detail interface with the navigation bar at the top and the toolbar displaying additional options and information at the bottom.

As mentioned earlier in this chapter, you should not subclass `UINavigationController` (nor should you need to).

Since you don't subclass `UINavigationController`, you can't override its `shouldAutorotateToInterfaceOrientation:`. Instead, a navigation interface's orientation is determined as follows:

- Initially, the navigation interface adopts the orientation rule of its root view controller.

- When a view controller is *pushed* onto the stack, the navigation interface *does not rotate*, even if the new view controller does not permit the current orientation. However, if the device is later rotated to an orientation that the new view controller permits, the navigation interface will rotate to it.
- When a view controller is *popped* from the stack, if the view controller that's now at the top of the stack permits the current orientation, the navigation interface stays in that orientation. Otherwise, the navigation interface rotates to a permitted orientation.

Bar Button Items

The buttons in a UIToolbar or a UINavigationBar are bar button items (UIBarButtonItem, a subclass of UIBarButtonItem). A bar button item can come in either of two broadly different flavors:

Basic bar button item

The bar button item looks and behaves like a simple button. A bar button item is not, however, the same as a UIButton; it has some button-like qualities, but it does not inherit from UIButton, from UIControl, or even from UIView. (It doesn't even have a frame.)

Custom view

The bar button item is assigned, and appears as, a `customView`. The `customView` is a UIView — *any* kind of UIView. Thus, a bar button item can put *any* sort of view into a toolbar or navigation bar, including a real UIButton or anything else (and implementing any button behavior would then be the responsibility of that view).

Let's start with the basic bar button item (no custom view). A bar button item, like a tab bar item, inherits from UIBarButtonItem the `title`, `image`, and `enabled` properties. A basic bar button item can have a title or an image, but generally not both; assigning an image removes the title if the bar button item is used in a navigation bar, but in a toolbar the title appears below the image. The image should usually be quite small (20×20 pixels is a good size).

A bar button item also has `target` and `action` properties. These give it its button-like behavior: tapping a bar button item can trigger an action method elsewhere ([Chapter 11](#)).

The overall look of a basic bar button item is determined by its `style` property; the choices are:

`UIBarButtonItemStyleBordered`

Looks like a button, with a round rectangular border around the image or title.

`UIBarButtonItemStylePlain`

In a toolbar, the bare title or image (or both) is displayed; if just a title, the text size is much larger by default than with a `UIBarButtonItemStyleBordered` button, but

in iOS 5 you can change that. Like a tab bar item's `image`, a bar button item's `image` must be a transparency mask when used with `UIBarButtonItemStylePlain` in a toolbar. In a navigation bar, `UIBarButtonItemStylePlain` is converted to `UIBarButtonItemStyleBordered`.

`UIBarButtonItemStyleDone`

Bordered, and with a blue fill. As the name implies, this is suitable for a Done button in a temporary view.

Starting in iOS 5, however, the look of a basic bar button item can be given a great deal of additional customization. It can have a tint color or a background image, and, as with a tab bar item, you can control the font and style of the title. Full details appear in [Chapter 25](#).

There are three ways to make a bar button item:

By borrowing it from the system

Instantiate `UIBarButtonItem` using `initWithBarButtonSystemItem:target:action:`. Consult the documentation for the list of available system items; they are not the same as for a tab bar item. You can't assign a title or change the image. (But you can change the tint color or assign a background image.)

By making your own basic bar button item

Instantiate `UIBarButtonItem` using `initWithTitle:style:target:action:` or `initWithImage:style:target:action:`.

In iOS 5, an additional method, `initWithImage:landscapeImagePhone:style:target:action:`, lets you supply two images, one for portrait orientation, the other for landscape orientation; this is because by default the bar's height might change when the interface is rotated.

By making a custom view bar button item

Instantiate `UIBarButtonItem` using `initWithCustomView:`. The tab bar item has no action and target; the `UIView` itself must somehow implement button behavior if that's what you want. For example, the `customView` might be a `UISegmentedButton`, but then it is the `UISegmentedButton`'s target and action that give it button behavior.

Bar button items in a toolbar are positioned automatically by the system. You can provide hints to help with this positioning. If you know that you'll be changing an item's title dynamically, you'll probably want its width to accommodate the longest possible title right from the start; to arrange that, set the `possibleTitles` property to an `NSSet` of strings that includes the longest title. Alternatively, you can supply an absolute width. Also, you can incorporate spacer system items into the toolbar; these have no visible appearance, and cannot be tapped, but serve to help distribute the other items. The `UIBarButtonItemStyleFlexibleSpace` is the one most frequently used; place these between the visible items to distribute the visible items equally across the width of the



Figure 19-9. A segmented control in the center of a navigation bar

toolbar. There is also a `UIBarButtonSystemItemFixedSpace` whose `width` lets you insert a space of defined size.

Navigation Items

What appears in a navigation bar (`UINavigationBar`) depends upon the navigation items (`UINavigationItem`) in its stack. In a navigation interface, the navigation controller will manage the navigation bar's stack for you, but you must still configure each navigation item, which you do by setting properties of the `navigationItem` of each child view controller. The properties are as follows (see also [Chapter 25](#)):

`title` or `titleView`

Determines what is to appear in the center of the navigation bar when this navigation item is at the top of the stack. The `title` is a string. Instead of setting `navigationItem.title`, however, you will usually set the view controller's `title` property instead; setting this property sets the `title` of the `navigationItem` automatically. The `titleView` can be any kind of `UIView`; if set, it will be displayed instead of the `title`. The `titleView` can implement further `UIView` functionality; for example, it can be tappable. Even if you are using a `titleView`, you should still give your view controller a `title`, as it will be needed for the back button when a view controller is pushed onto the stack on top of this one.

[Figure 19-1](#) shows the TidBITS News master view, with the navigation bar displaying a `titleView` which is a (tappable) image view; the master view's `title` is therefore not displayed. In the TidBITS News detail view controller, the `titleView` is a segmented control providing a Previous and Next button, and the back button displays the master view controller's `title` ([Figure 19-9](#)).

`prompt`

An optional string to appear centered above everything else in the navigation bar. The navigation bar's height will be increased to accommodate it.

`rightBarButtonItem` or `rightBarButtonItems`

A bar button item or, respectively, an array of bar button items to appear at the right side of the navigation bar. The option to provide multiple right bar button items (`rightBarButtonItems`) is new in iOS 5. In [Figure 19-1](#), the refresh button is a right bar button item; it has nothing to do with navigation, but is placed here merely because space is at a premium on the small iPhone screen. Similarly, in [Figure 19-9](#), the text size button is a right bar button item, placed here for the same reason.

`backBarButtonItem`

When a view controller is pushed on top of this view controller, the pushed view controller's navigation bar will display at its left a button pointing to the left, whose title is this view controller's `title`. That button is *this* view controller's back button item. That's right: the back button displayed in the navigation bar belongs, not to the top item (the `navigationItem` of the current view controller), but to the back item (the `navigationItem` of the view controller that is one level down in the stack). In [Figure 19-9](#), the back button in the detail view is the default back button, displaying the `title` of the master view.

The vast majority of the time, the default behavior is the behavior you'll want, and you'll leave the back button alone. If you wish, though, you can customize the back button by setting the `navigationItem.backBarButtonItem` of this view controller (the view controller that will provide the back item, not the one that will provide the top item), so that it contains an image, or a title differing from this view controller's `title`. The best technique is to provide a new `UIBarButtonItem` whose target and action are `nil` (its style doesn't matter); the runtime will provide a correct target and action, so as to create a working back button:

```
UIBarButtonItem* b = [[UIBarButtonItem alloc] initWithTitle:@"Go Back"
                                                       style:UIBarButtonItemStylePlain target:nil action:nil];
self.navigationItem.backBarButtonItem = b;
```

A `BOOL` property, `hidesBackButton`, allows the top navigation item to suppress display of the back item's back bar button item. Obviously, if you set this to `YES`, you'll need to consider providing some other means of letting the user navigate back.

`leftBarButtonItem` or `leftBarButtonItems`

A bar button item or, respectively, an array of bar button items to appear at the left side of the navigation bar. The option to provide multiple left bar button items (`leftBarButtonItems`) is new in iOS 5.

Prior to iOS 5, the presence of a left bar button item would cause the back button item not to appear; since this disabled the normal means of going back, use of a left bar button item was typically confined to the root view, where no back button is displayed anyway. In iOS 5, a new property, `leftItemsSupplementBackButton`, if set to `YES`, allows both the back button and one or more left bar button items to appear.

Here's the view controller code that configures its navigation item to generate the navigation bar shown in [Figure 19-1](#):

```
// title for back button in detail view
self.title = @"TidBITS";
// image to display in navigation bar
UIImageView* imv = [[UIImageView alloc] initWithImage:
    [UIImage imageNamed:@"tb_iphone_banner.png"]];
self.navigationItem.titleView = imv;
```



Figure 19-10. A highly dynamic navigation bar

```
// reload button for navigation bar
UIBarButtonItem* b = [[UIBarButtonItem alloc]
    initWithBarButtonSystemItem:UIBarButtonSystemItemRefresh
    target:self action:@selector(doRefresh:)];
self.navigationItem.rightBarButtonItem = b;
```

A view controller's navigation item can have its properties set at any time while being displayed in the navigation bar. This (and not direct manipulation of the navigation bar) is the way to change the navigation bar's contents dynamically. For example, in one of my apps, the visible right bar button should either be the system Play button, the system Pause button, or nothing, depending on whether music from the library is playing, paused, or stopped (Figure 19-10). So I have a timer that periodically checks the state of the music player:

```
int whichButton = -1;
if ([self->mp playbackState] == MPMusicPlaybackStatePlaying)
    whichButton = UIBarButtonItemSystemItemPause;
else if ([self->mp playbackState] == MPMusicPlaybackStatePaused)
    whichButton = UIBarButtonItemSystemItemPlay;
if (whichButton == -1)
    self.navigationItem.rightBarButtonItem = nil;
else {
    UIBarButtonItem* bb =
        [[UIBarButtonItem alloc] initWithBarButtonSystemItem:whichButton
            target:self action:@selector(doPlayPause:)];
    self.navigationItem.rightBarButtonItem = bb;
}
```

In that same app, and in the same navigation item, the `titleView` is a progress view (`UIProgressView`, Chapter 25). I treat this like any other progress view, constantly updating it (setting its `progress` value) and even making it visible or invisible (setting its `hidden` value) without regard for the fact that it's being displayed in a navigation bar (Figure 19-10).

Toolbar Items

Each view controller to be pushed onto the navigation controller's stack is responsible for supplying the items to appear in the navigation interface's toolbar, if there is one. This is done by setting the view controller's `toolbarItems` property to an array of `UIBarButtonItem` instances. You can change the toolbar items even while the view controller's view and current `toolbarItems` are showing, optionally with animation, by sending `setToolbarItems:animated:` to the view controller.

A view controller has the power to specify that the navigation interface's toolbar should be hidden whenever it (the view controller) is on the stack. To do so, set the view controller's `hidesBottomBarWhenPushed` property to YES. The trick is that you must do this early enough, namely before the view loads. (The view controller's `viewDidLoad` is too late; its designated initializer is a good place.) The bottom bar remains hidden from the time this view controller is pushed to the time it is popped, even if other view controllers are pushed and popped on top of it in the meantime. For more flexibility, you can call the `UINavigationController`'s `setToolbarHidden:animated:` at any time.

Configuring a Navigation Controller

You configure a navigation controller by manipulating its stack of view controllers. If a view controller is in the stack, it is a child view controller of the navigation controller; the navigation controller is its `parentViewController`, and it is also the `navigationController` of this view controller and its child view controllers at any depth. Thus a child view controller at any depth can learn that it is contained by a navigation controller and can get a reference to that navigation controller. The navigation controller retains the view controller as long as the view controller is on its stack; when the view controller is removed from the stack, the navigation releases the view controller, which is usually permitted to go out of existence at that point.

The normal way to manipulate a navigation controller's stack is one view controller at a time. When the navigation controller is instantiated, it is assigned a single root view controller, the view controller that goes at the bottom of the stack, whose view the navigation controller will initially display, by initializing it with `initWithRootViewController::`:

```
FirstViewController* fvc = [[FirstViewController alloc] init];
UINavigationController* nav =
    [[UINavigationController alloc] initWithRootViewController:fvc];
self.window.rootViewController = nav; // or whatever
```

Later, when the user asks to navigate to the right, you obtain the next view controller (typically by creating it) and push it onto the stack by calling `pushViewController:animated:` on the navigation controller, and the navigation controller displays its view:

```
// [in FirstViewController:]
SecondViewController* svc = [[SecondViewController alloc] init];
[self.navigationController pushViewController:svc animated:YES];
```

Typically, that's all there is to it! When the user taps the back button to navigate back to the left, the runtime will call `popViewControllerAnimated:` for you.

You can also set the `UINavigationController`'s delegate; the delegate should adopt the `UINavigationControllerDelegate` protocol. The delegate receives an event before and after the navigation controller changes what view controller's view is displayed.

You can manipulate the stack more directly if you wish. You can call `popViewControllerAnimated:` yourself; to pop multiple items so as to leave a particular view controller at

the top of the stack, call `popToViewController:animated:`, or to pop all the items down to the root view controller, call `popToRootViewControllerAnimated:`. All of these methods return the popped view controller (or view controllers, as an array), in case you want to do something with them.

To set the entire stack at once, call `setViewControllers:animated:`. You can access the stack through the `viewControllers` property. Manipulating the stack directly is the only way, for instance, to delete or insert a view controller in the middle of the stack.

The view controller at the top of the stack is the `topViewController`; the view controller whose view is displayed is the `visibleViewController`. Those will normally be the same view controller, but they needn't be, as the top view controller might be presenting a view controller. Other view controllers can be accessed through the `viewControllers` array by index number. The root item is the array's `objectAtIndex:0`; if the array's `count` is `c`, the back item is the array's `objectAtIndex:c-2`.

You'll notice that in the preceding code examples I didn't configure a view controller's `navigationItem` as I pushed it onto the stack. Sometimes, the code that creates a view controller may configure that view controller's `navigationItem`, but it is most common for a view controller to configure its own `navigationItem`. A view controller will be concerned to do this sufficiently early in its own lifetime. The earliest such point is an override of the designated initializer. Other possibilities are:

The view controller's `awakeFromNib`

Obviously, this choice is possible only in cases where the view controller instance comes from a nib.

The view controller's `viewDidLoad` (or `loadView`)

This seems an obvious choice, but Apple warns (in the `UIViewController` class reference, under `navigationItem`) that configuring a view controller's navigation item in conjunction with the creation of its view is not a good idea, because the circumstances under which the view is needed are not identical to the circumstances under which the navigation item is needed. However, Apple's own examples appear to violate this warning; they often do configure a view controller's navigation item in its `viewDidLoad` implementation.

When a child view controller's view is displayed, it is resized to fit the display area. The size of this area will depend on whether the navigation controller is showing a navigation bar or a toolbar or both (or neither). You should design your view in such a way as to be prepared for such resizing. Autoresizing settings can help here. Also, when editing your view in the nib editor, you can shrink the view to the size at which it will be displayed in the navigation interface, thus helping you judiciously situate your interface items; to do so, choose appropriately from the Top Bar and Bottom Bar pop-up menus of the Simulated Metrics section of the Attributes inspector.

A navigation controller's navigation bar can be hidden and shown with `setNavigationBarHidden:animated:`, and the toolbar can be hidden and shown with `setToolbar-`

`Hidden:animated:`. The current view controller's view will be resized then and there; be sure to design your view to accommodate such dynamic resizing, if needed.

You can also configure a UINavigationController or any view controller that is to serve in a navigation interface in a nib or storyboard. In a nib, a navigation view controller has a top bar and a bottom bar (in the Attributes inspector), and you can configure how they look and whether they are initially visible. The root view controller can be specified. Moreover, the root view controller has a Navigation Item where you can specify its title, its prompt, and the text of its back button. (If a view controller in a nib *doesn't* have a Navigation Item and you want to configure this view controller for use in a navigation interface, drag a Navigation Item from the Object library onto the view controller.) You can drag Bar Button Items into the navigation bar to set the left button and right button of a view controller's Navigation Item. Moreover, the Navigation Item has outlets, one of which permits you to set its `titleView`. Plus, you can give a view controller Bar Button Items that will appear in the toolbar. Thus the configuration of a navigation view controller, its root view controller, and any other view controllers that will be pushed onto its stack can be performed to a certain degree in a nib.

Page View Controller

A page view controller (`UIPageViewController`), new in iOS 5, has one or two child view controllers whose view(s) it displays within its own view. The user, by default, can then tap or swipe to navigate, revealing the view of a different view controller or pair of view controllers, analogously to the pages of a book. To reinforce the book analogy, there's a page curl animation, similar to the animation in Apple's iBooks.

Page view controllers are a great addition to the repertoire of built-in view controllers. Before iOS 5, I was accomplishing the same thing in my flashcard apps (without the page curl animation) by means of a scroll view ([Chapter 20](#)); the code was complex and tricky. With a page view controller, I was able to make my app's code far simpler, plus I get the page curl animation for free.

To create a `UIPageViewController`, initialize it with `initWithTransitionStyle:navigationOrientation:options:`. Here's what the parameters mean:

- The only legal transition style at present is `UIPageViewControllerTransitionStylePageCurl`.
- The possible navigation orientations are `UIPageViewControllerNavigationOrientationHorizontal` and `UIPageViewControllerNavigationOrientationVertical`.
- The options are a dictionary for which, at present, only one key is possible, `UIPageViewControllerOptionSpineLocationKey`, specifying the position of the spine (the pivot line around which page transitions rotate). The legal possible values at present are:

- `UIPageViewControllerSpineLocationMin` (left or top)
- `UIPageViewControllerSpineLocationMid` (middle; two pages are shown at once)
- `UIPageViewControllerSpineLocationMax` (right or bottom)

You then assign the page view controller a `dataSource`, which should conform to the `UIPageViewControllerDataSource` protocol, and assign the page view its initial content by handing it its initial child view controller(s). You do that by calling `setViewControllers:direction:animated:completion:`. Here are the parameters:

- The view controllers are an array of one or two view controllers — one if the spine is `min` or `max`, two if the spine is `mid`.
- The direction can be `UIPageViewControllerNavigationDirectionForward` or `UIPageViewControllerNavigationDirectionBackward`; which you specify probably won't matter when you're assigning the page view controller its initial content.
- When you're assigning the page view controller its initial content, you probably won't want any animation.

Here's a minimal example, from the app delegate's `application:didFinishLaunchingWithOptions::`:

```
// make a page view controller
UIPageViewController* pvc =
    [[UIPageViewController alloc]
        initWithTransitionStyle:UIPageViewControllerTransitionStylePageCurl
        navigationOrientation:UIPageViewControllerNavigationOrientationHorizontal
        options:nil];
// give it an initial page
Pep* page = [[Pep alloc] initWithPepBoy:[self.pep objectAtIndex:0]
                           nib: nil bundle: nil];
[pvc setViewControllers:[NSArray arrayWithObject:page]
                     direction:UIPageViewControllerNavigationDirectionForward
                     animated:NO completion:NULL];
// give it a data source
pvc.dataSource = self;
// stick it in the window, retain the page view controller
self.window.rootViewController = pvc;
```

From then on, when the user tries to navigate, the data source's `pageViewController:viewControllerAfterViewController:` or `pageViewController:viewControllerBeforeViewController:` will be called; its job is to return the requested view controller. You'll need a strategy for doing that; the strategy you devise will depend on your modal, that is, on how you're maintaining your data.

To illustrate, continuing from the preceding example, I've got a `UIViewController` subclass called `Pep` and an array (`self.pep`) of the names of the Pep Boys. After I initialize a `Pep` object by calling `initWithPepBoy:nib:bundle:`, `Pep`'s `viewDidLoad` fetches a corresponding image from the app bundle and assigns it as the image to a `UIImageView` within its own view. Thus, a page in the page view controller contains an image of a Pep Boy. So here's one of my data source methods:

```

-(UIViewController *)pageViewController:(UIPageViewController *)pageViewController
    viewControllerAfterViewController:(UIViewController *)viewController {
    NSString* boy = [(Pep*)viewController boy]; // string name of this Pep Boy
    NSUInteger ix = [self.pep indexOfObject:boy]; // find it in the data model
    ix++;
    if (ix >= [self.pep count])
        return nil; // there is no next page
    return [[Pep alloc] initWithPepBoy:[self.pep objectAtIndex:ix]
        nib: nil bundle: nil];
} // and "before" is similar

```

You can also call `setViewControllers:direction:animated:completion:` to change programmatically what page is being displayed, possibly with animation. I do so in my Latin flashcard app during drill mode ([Figure 19-5](#)), to advance to the next term in the current drill:

```

[self.terms shuffle];
Term* whichTerm = [self.terms objectAtIndex: 0];
CardController* cdc = [[CardController alloc] initWithTerm:whichTerm];
[self.pvc setViewControllers:[NSArray arrayWithObject:cdc]
    direction:UIPageViewControllerNavigationDirectionForward
    animated:YES completion:nil];

```

If you refer to `self` in the `completion` block of `setViewControllers:direction:animated:completion:`, ARC will warn of a possible retain cycle. I don't know why there would be a retain cycle, but I take no chances: I do the weak-strong dance described in [Chapter 13](#).

It is also possible to assign a page view controller a delegate, which adopts the `UIPageViewControllerDelegate` protocol. You get an event when the user finishes turning the page, and you get a chance to change the spine location dynamically in response to a change in device orientation.

A page view controller has a `gestureRecognizers` property that returns an array containing the gesture recognizers it uses to recognize a tap or swipe. You can do useful customizations with these gesture recognizers, such as attaching them to another view, or implementing their delegate. I do that in my Latin flashcard app in order to prevent the user's tap or swipe from working unless it is outside a text field. First, I make myself delegate of all the gesture recognizers:

```

for (UIGestureRecognizer* g in pvc.gestureRecognizers)
    g.delegate = self; // give me a chance to veto navigation

```

When my delegate method is called, I disallow all touches whose view isn't the image view that comprises the background of the card. I've assigned that image view its own `UIImageView` subclass for purposes of identification:

```

- (BOOL)gestureRecognizer:(UIGestureRecognizer *)gestureRecognizer
    shouldReceiveTouch:(UITouch *)touch {
    if (![touch.view isKindOfClass:[BackgroundImageView class]])
        return NO;
    return YES;
}

```

One further bit of configuration is performed through the `doubleSided` property. If it is YES, the next page occupies the back of the previous page. The default is NO, unless the spine is in the middle, in which case it's YES and can't be changed. Your only option here, therefore, is to set it to YES when the spine isn't in the middle, and in that case the back of each page would be a sort of throwaway page, glimpsed by the user during the page curl animation. For example, you might make every other page a solid white view.

A `UIPageViewController`'s policy on rotation is like a `UINavigationController`'s policy. Initially, the page view controller wants the interface to rotate to an orientation allowed by the initial page; later on, if a new page's view controller doesn't permit the current device orientation, there is no rotation, but if the device is then rotated to a permitted orientation, the interface rotates to match. It is probably best to use view controllers whose answers to `shouldAutorotateToInterfaceOrientation:` are the same.

Container View Controllers

Starting in iOS 5, the ability of built-in view controller subclasses such as `UITabBarController`, `UINavigationController`, and `UIPageViewController` to function as parents for other view controllers and put the views of those view controllers into their own view is generalized and extended to `UIViewController` as a whole. The consequence is that you are allowed to write your own parent view controllers, which Apple calls **container view controllers**. This should allow for some powerful and original new interfaces.

A `UIViewController` has a `childViewControllers` array. This array will be maintained for you. To act as a parent view controller, your `UIViewController` subclass must fulfill certain responsibilities.

- When a view controller is to become your parent view controller's child:
 - You call `addChildViewController:` on your parent view controller. The child is automatically added to the parent's `childViewControllers` array, which retains it.
 - You get the child view controller's into your interface, if that's what adding a child view controller means.
 - You send `didMoveToParentViewController:` to the child with the parent view controller as its argument.
- When a view controller is to cease being your parent view controller's child:
 - You send `willMoveToParentViewController:` to the child with a nil argument.
 - You remove the child view controller's view from your interface, if that's what removing a child view controller means.
 - You send `removeFromParentViewController` to the child. The child is automatically removed from the parent's `childViewControllers` array, which releases it.

This is a clumsy and rather poorly designed dance. The underlying reason for it is that a child view controller must always receive `willMoveToParentViewController:` followed by `didMoveToParentViewController:` (and your own child view controllers can take advantage of these events however you like). Well, it turns out that `addChildViewController:` sends `willMoveToParentViewController:` for you, and that `removeFromParentViewController` sends `didMoveToParentViewController:` for you; so in each case you must send the *other* message, the one that adding or removing a child view controller *doesn't* send for you—and of course you must send it so that everything happens in the correct order, as dictated by rules I just listed.

When you initialize your interface, with no replacement of one view by another, you'll simply set up the frame of the view of some child view controller and place it into your view. In this example, I do that in my parent view controller's `viewDidLoad`:

```
- (void)viewDidLoad
{
    [super viewDidLoad];
    UIViewController* vc = // whatever; this the initial view controller
    [self addChildViewController:vc]; // "will" is called for us
    vc.view.frame = // whatever
    [self.view addSubview: vc.view];
    // when we call "add", we must call "did" afterwards
    [vc didMoveToParentViewController:self];
}
```

Now comes the really interesting part. The purpose of view controllers, as you know, is to manage the replacement of one view with another in the interface. Thus, no matter what your parent view controller's particular functionality is, it will eventually want to replace the view of one child view controller with the view of another child view controller in the interface. To do so, it will send itself the message `transitionFromViewController:toViewController:duration:options:animations:completion:`. The `options` argument is a bitmask comprising the same possible options (whose names start with `UIViewControllerAnimatedOption...`) that apply to any block-based view transition ([“Block-Based View Animation” on page 423](#)). The `animations:` block is for *additional* view animations, as with any transition.

The `completion:` block will be important if this transition is part of removing or adding a child view controller. At the time `transitionFromViewController:...` is called, both view controllers involved must be children of the parent view controller; so if you're going to remove one of the view controllers as a child, you'll do it in the `completion:` block. Similarly, if you owe a new child view controller a `didMoveToParentViewController:` call, you'll use the `completion:` block to fulfill that debt.

In this simple example, our parent view controller has one child view controller at a time. It displays the view of that child view controller within its own view. When our parent view controller is handed a new child view controller, it substitutes that new child view controller for the old child view controller and replaces the old child view

controller's view with the new child view controller's view. The two view controllers are called `fromvc` and `tovc`:

```
// set up the new view controller's view's frame
tovc.view.frame = // ... whatever
// must have both as children before we can transition between them
[self addChildViewController:tovc]; // "will" is called for us
// when we call "remove", we must call "will" (with nil) beforehand
[fromvc willMoveToParentViewController:nil];
[self transitionFromViewController:fromvc
    toViewController:tovc
    duration:0.4
    options:UIViewAnimationOptionTransitionFlipFromLeft
    animations:nil
    completion:^(BOOL done){
        // we called "add"; we must call "did" afterwards
        [tovc didMoveToParentViewController:self];
        [fromvc removeFromParentViewController];
        // "did" is called for us
    }];
}];
```

Thus it is possible, using a custom parent view controller, to manage the regions of your app's interface just as a tabbed interface or a navigation interface does. This works even on an iPhone, because view transitions work on a iPhone and because subviews are just subviews. So, for example, in my Latin flashcard app ([Figure 19-3](#)), the interface comprises a `UIPageViewController`'s view and a toolbar (`UIToolbar`). Thus, I want to take charge of the frame of the `UIPageViewController`'s view and put it into a region of the interface myself. Prior to iOS 5, this would have been illegal: you weren't supposed to

There was (and actually still is, though it is presumably now outmoded) a warning in Apple's documentation, in the "View Controller Basics" chapter of the *View Controller Programming Guide for iOS*, saying not to do exactly what I want to do here:

If you want to divide a view hierarchy into multiple subareas and manage each one separately, use generic controller objects (custom objects descending from `NSObject`) instead of view controller objects to manage each subarea. Then use a single view controller object to manage the generic controller objects.

In iOS 5, though, I can write a parent view controller object that *does* divide my view hierarchy into multiple subareas, managing each subarea through its own child view controller ([Figure 19-4](#)). Another result of this generalized architecture is that if you don't like the way a built-in `UIViewController` subclass does its job, you can write your own replacement. For example, a `UITabBarController` really does nothing you couldn't now do yourself using a custom parent view controller.

Storyboards

The idea of a storyboard (["Storyboard-Instantiated View Controller" on page 515](#)) is to express compactly and diagrammatically the "story" of how the interface will change.

I imagine that the idea of storyboards came into existence from the complaints and expectations of programmers, especially beginners. When I first started programming iOS, I had great difficulty in understanding view controllers and in arriving at the explanation of them that I've presented in this chapter. When I discovered that view controllers could be in a nib file, and that in some cases (such as a navigation controller and its root view controller) multiple view controllers could be in a nib, I was disappointed to find that, to continue this way of obtaining view controllers, I'd need further nibs, one for each additional view controller to be pushed onto the stack. My feeling was: "If I'm going to be required to use view controllers, why can't they *all* be in one nib?" Storyboards seem to be a response to that sort of feeling.

A storyboard collects, in a single file, multiple view controllers — both parent and children, both presenter and presented. Each view controller is part of a *scene*, which is essentially an encapsulation of a nib containing that view controller. Connections are drawn between view controllers. If view controller A has view controller B as its child from the outset (as with a navigation controller and its root view controller, or a tab bar controller and its child view controllers), the connection is a *relationship* and emanates from view controller A. If view controller A's view is to be replaced later by view controller B's view, the connection is a *segue* and typically emanates from some interface item in view controller A's view that the user can tap to trigger the replacement. A segue is directional; it has a source view controller and a destination view controller. A segue can also be assigned a string identifier (in its Attributes inspector); you will just about always want to do this. A segue is a nib object, like any other; it is an instance of UIStoryboardSegue (or your custom subclass thereof).

Recall from "[Storyboard-Instantiated View Controller](#)" on page 515 that a storyboard is referenced in code by calling `storyboardWithNibName:`, and that a particular view controller can be instantiated from a storyboard by calling `instantiateInitialViewController` or `instantiateViewControllerWithIdentifier:`. Recall also from [Chapter 6](#) and [Chapter 14](#) that your app's main storyboard, if there is one, is automatically used to instantiate its initial view controller, which is made the window's root view controller.

Subsequent view controllers in a nib are usually instantiated *automatically* as they are needed. View controllers bound together by a relationship are instantiated together, just as they would be if they were together in a nib. But a view controller that is the destination view controller of a segue is instantiated *when the segue is triggered*, and its memory management is then handed over to its parent or presenting view controller, just as you would do in code; thus, for example, if a view controller is a presented view controller, it is instantiated automatically each time it is presented, and the instance goes out of existence each time it is dismissed.

A view controller instantiated from a storyboard has a non-nil `storyboard` property pointing to the storyboard that it comes from; this provides another way of getting a reference to a storyboard.

Here are the primary concerns in using and configuring a storyboard:

How to trigger a segue

Usually, you'll draw the segue as a connection from the interface object that is to trigger it. For example, when a button is to trigger the segue, you draw the connection from the button to the destination view controller; the runtime assumes that you want the segue to be triggered in response to the button's Touch Up Inside event. When a table view cell is to trigger the segue, you draw the connection from the cell to the destination view controller; the runtime assumes that you want the segue to be triggered in response to the user selecting the cell. If the default behavior doesn't cover your needs, you can trigger the segue yourself, in code, by calling `performSegueWithIdentifier:sender:` on the source view controller.

How to customize the transition performed by a segue

A segue's Attributes inspector contains pop-up menus that let you specify the segue as a presenting ("Modal") segue or a navigation controller ("Push") segue. (It's odd that the storyboard editor perpetuates the term "Modal" just when the rest of Cocoa is preparing to deprecate it.) If these actions don't cover your needs, you can make a custom `UIStoryboardSegue` subclass and implement its `perform` method; you'll need to call the appropriate method yourself, just as you would have done if you weren't using a storyboard. In the storyboard editor, you must set the segue's Style pop-up menu to Custom and type the name of your `UIStoryboardSegue` subclass in the Segue Class field. (You do *not* specify the class in the Identity inspector, and you do *not* get a pop-down menu or automatic name completion; this feels like a mistake in the design of the storyboard editor.) For example, for a custom transition when presenting a view controller, you'd implement `perform` to call `presentViewController:animated:completion:` yourself. Your code can work out what segue is being triggered by examining the segue's `identifier`, `sourceViewController`, and `destinationViewController` properties.

How to configure a view controller before a segue to it

Before a segue is performed, the source view controller is sent `prepareForSegue:sender:`. The view controller can work out what segue is being triggered by examining the segue's `identifier` and `destinationViewController` properties, and the `sender` is the interface object that was tapped to trigger the segue (or, if `performSegueWithIdentifier:sender:` was called in code, whatever object was supplied as the `sender` argument). This is the moment when the source view controller and the destination view controller meet; the source view controller can thus perform configurations on the destination view controller.

How to reverse the transition performed by a segue

This is easy, because the involvement of the storyboard makes no difference whatever. To pop a pushed view controller in a navigation interface, or to dismiss a presented view controller, you do in code whatever you would have done without the storyboard.

That's all there is to using a storyboard. You should experiment with the storyboard versions of the various Xcode 4.2 project templates, to see how they are constructed; you can then readily adapt them to your particular situation.

As an example, let's start with the storyboard version of the Utility Application template, which shows how a presenting–presented view controller segue is implemented in a storyboard. Both the presenting view controller (`MainViewController`) and the presented view controller (`FlipsideViewController`) are together in one storyboard file. The segue from the former to the latter emanates from a button in `MainViewController`'s view; its Attributes inspector specifies that its identifier is `@"showAlternate"`, and that it is a “Modal” segue (meaning we're to present a view) with a “Flip Horizontal” transition. This means that when the user taps the button, `presentViewController:animated:completion:` is called for us; the button needn't be connected to an action handler in our own code.

Suppose now that we want to trigger the segue in some other way. Perhaps we don't want to use the button's Touch Up Inside event, or perhaps we want to trigger the segue in response to some other gesture or event altogether. In that case, we delete the existing segue from the storyboard, and draw a new connection from the source view controller to the destination view controller, specifying that this is a “Modal” segue. We are now back to implementing some sort of action handler, just as we would have done if there were no storyboard; but instead of instantiating the destination view controller and presenting it, this action handler triggers the segue. It might look like this:

```
- (IBAction)doButton:(id)sender {
    [self performSegueWithIdentifier:@"showAlternate" sender:self];
}
```

Suppose next that we don't want any animation; the pop-up menus in the segue's Attributes inspector don't provide that option. Or suppose we have a `completion:` block we want performed. We'll need to subclass `UIStoryboardSegue` and override its `perform` method; this will be called instead of calling `presentViewController:animated:completion:`, and it is up to us to call `presentViewController:animated:completion:` ourselves, just as if we weren't using a storyboard.

So, we add a new `UIStoryboardSegue` subclass to our project; call it `MySegue`. Remember to specify in the storyboard that this segue is a `MySegue` (change the Style pop-up menu to Custom and type `MySegue` in the Segue Class field)! Then, in `MySegue.m`, implement `perform`:

```
-(void) perform {
    [self.sourceViewController
        presentViewController:self.destinationViewController
        animated:NO completion:nil];
}
```

If we needed to add a `completion:` block, we could do so.

Now let's suppose we need to configure the `FlipsideViewController` instance. For example, if we were instantiating `FlipsideViewController` in code, we might initialize that

instance then and there with a custom designated initializer such as `initWithData:`, handing it the data that it needs, as I demonstrated in an example from my own code earlier in this chapter when discussing presented view controllers:

```
DrillViewController* dvc = [[DrillViewController alloc] initWithData:drillTerms];
[self presentViewController:dvc animated:YES completion:nil];
```

With a storyboard, we can't do that; there's no way to tell the storyboard to call `initWithData:`. (To be fair, we'd have the same difficulty if the destination view controller were instantiated from a nib.) Instead, we must let the storyboard instantiate and initialize `FlipsideViewController`; the runtime then turns to `MainViewController` and calls `prepareForSegue:sender:`. `MainViewController` can identify the segue using its `identifier` property, and can get a reference to the `FlipsideViewController` instance as the segue's `destinationViewController` property. So now it can configure the `FlipsideViewController`:

```
- (void)prepareForSegue:(UIStoryboardSegue *)segue sender:(id)sender {
    if ([[segue identifier] isEqualToString:@"showAlternate"]) {
        FlipsideViewController* fvc =
            (FlipsideViewController*)segue.destinationViewController;
        fvc.theData = theData;
    }
}
```

Finally, what happens when we want to return from the `FlipsideViewController`'s view to the `MainViewController`'s view? Absolutely nothing! The storyboard gives us no assistance with this side of the equation. We might have a button in `FlipsideViewController`'s view with an action handler in `FlipsideViewController`, which calls (or causes to be called) `dismissModalViewControllerAnimated:` — just as if there were no storyboard at all.

I can't resist giving my opinion on storyboards in general, which is that they are probably well calculated to get beginners up and running, but that their architectural organization is likely to be advantageous only in the simplest situations. No accommodation is made for unknown parent-child relationships, so an app with a custom parent view controller can't define that parent view controller's relationships in a storyboard. A storyboard containing many scenes can quickly become unwieldy. The storyboard editor is not well designed; in particular, a segue's identifier is crucial, and yet it must be typed by hand in multiple places (an invitation to error) and is not displayed in the editor's canvas (you have to select the segue and examine its Attributes inspector).

Moreover, although storyboards assemble multiple view controllers and their views in a single editor, they don't make it any easier to get a reference from one view controller's world to another: you're not allowed to draw a connection from an object within one "scene" to an object within another "scene". You can't draw a connection between one *nib* and another, as that would make no sense: it wouldn't be clear what instances you're trying to connect, or even whether those instances will exist simultaneously at runtime (see [Chapter 7](#) and [Chapter 13](#)). But a *Storyboard* seems to promise that instances will come into being in a controlled and predictable manner, and in a known

relationship to one another. After all, a segue has both a `sourceViewController` and a `destinationViewController` at the same time; so if I could draw, say, an action connection from a button in a presented view controller's view to the presenting view controller in a storyboard, why isn't it clear that the particular instances in question are the ones that will be serving as the segue's `sourceViewController` and `destinationViewController` at the moment the segue is triggered? I may be missing some reason why connection across scenes is impossible, but this seems to me a major opportunity lost. The outcome, indeed, is that it often takes considerably *more* work, involving more fragile code, to establish the desired connections in your `prepareForSegue:sender:` implementation, than it would have been to create, configure and present the destination view controller in code in the first place.

My biggest problem with storyboards, though, is that, from an OOP perspective, everything happens in the wrong place. Consider the case of a `completion:` block in a custom `UIStoryboardSegue`'s `perform` implementation. If this code were in the source view controller, as it normally would be, we could access any variables in scope; but this code is *not* in the source view controller. Thus, the block is in the wrong place, sundered from its proper context. Also, the segue transition and its reversal are called in different places, which feels wrong. Thus, in our example above, `presentViewController:animated:completion:` is called in one place (either automatically or in a custom `UIStoryboardSegue` class), while the corresponding `dismissViewControllerAnimated:` is called somewhere else entirely (one of the two view controllers).

Similarly, `prepareForSegue:sender:` is an ugly bottleneck. If our source view controller is the source of several segues, its `prepareForSegue:sender:` becomes in effect a giant `switch` statement, doing different things depending on the identifier of the incoming segue; nothing could be more brittle, error-prone, and unmaintainable. Moreover, this way of having the source view controller and the destination view controller meet puts the wrong view controller in the driver's seat. If the destination view controller has a designated initializer, that initializer dictates to whoever instantiates it what the view controller wants ("I cannot come into existence without this `data:` argument"), and the view controller can then do what it likes with the data it is handed. By contrast, in `prepareForSegue:sender:`, the source view controller is responsible for knowing what data the destination view controller needs, and has to access its properties or call supplementary methods in order to provide that data.

This is not to say, however, that the notion of storyboards is fundamentally flawed, or that storyboards are completely unusable. As I said earlier, for simple apps that meet the basic expectations of the storyboard architecture, a storyboard could be a fine way to get yourself up and running quickly. And the encapsulation of what happens in connection with the transition from one view controller's view to another's as an object (a `UIStoryboardSegue`) is an intriguing insight, one that could ultimately lead to *better* code organization when and if storyboards ultimately mature. Plus, in [Chapter 21](#) it will turn out that the storyboard editor implements a couple of very useful features that you can't obtain any other way.

Personally, I avoid storyboards except in the simplest situations; outside of such situations, storyboards don't reduce my code, they scatter that code into inappropriate places, and the storyboard editor itself is more a hassle than a convenience. I spend more time satisfying the storyboard architecture's expectations than getting on with the real business of the app. You might justify reliance on storyboards as a crutch if you are a beginner and don't understand view controllers; but you've read this chapter and you *do* understand view controllers!

View Controller Lifetime Events

As views come and go, driven by view controllers and the actions of the user, events arrive that give your view controller the opportunity to respond to the various stages of its existence. By overriding these methods, your UIViewController subclass can perform appropriate tasks. Most commonly, you'll override `viewWillAppear:`, `viewDidAppear:`, `viewWillDisappear:`, or `viewDidDisappear:`. Note that *you must call super* in your override of any of these four methods.

Let's take the case of a UIViewController pushed onto the stack of a navigation controller. It receives, in this order, the following messages:

- `willMoveToParentViewController:`
- `viewWillAppear:`
- `viewWillLayoutSubviews`
- `viewDidLayoutSubviews`
- `viewDidAppear:`
- `didMoveToParentViewController:`

When this same UIViewController is popped off the stack of the navigation controller, it receives, in this order, the following messages:

- `willMoveToParentViewController:` (with argument nil)
- `viewWillDisappear:`
- `viewDidDisappear:`
- `didMoveToParentViewController:` (with argument nil)

In these names, the notions "appear" and "disappear" reflect the view's insertion into and removal from the interface. Disappearance can happen because the UIViewController itself is taken out of commission, but it can also happen because another UIViewController supersedes it. For example, let's take the case of the UIViewController functioning as the root view controller of a navigation controller. When another view controller is pushed on top of it, the root view controller gets these messages:

- `viewWillDisappear:`
- `viewDidDisappear:`

By the same token, appearance can happen because this `UIViewController` has been brought into play, but it can also happen because some other `UIViewController` is no longer superseding it. For example, when a view controller is popped from a navigation controller, the view controller that was below it in the stack receives these events:

- `viewWillAppear:`
- `viewWillLayoutSubviews`
- `viewDidLayoutSubviews`
- `viewDidAppear:`

Prior to iOS 5, it was rather difficult to distinguish these cases — that is, to find out precisely *why* your view was appearing or disappearing. In iOS 5, you can find out, from within the four “appear”/“disappear” methods, more about why they are being called, by calling these methods on `self`:

- `isBeingPresented`
- `isBeingDismissed`
- `isMovingToParentViewController`
- `isMovingFromParentViewController`

Here are some examples of how these events are used in my own apps:

- A certain view in a navigation controller needs the toolbar, whereas other views do not. As the view appears and disappears, the view controller calls the navigation controller to show and hide the toolbar.
- In a master–detail interface, the root view in a navigation interface displays a table. The data displayed by the root view’s table might change while the user is working in a detail view, so I reload the root view table’s data in the view controller’s `viewWillAppear:`.
- A view that can be pushed onto a navigation controller’s stack contains a progress view that is periodically updated through a timer. This timer needs to be in existence and running only when this view is in the interface. So I create the timer in the view controller’s `viewWillAppear:` and destroy it in `viewDidDisappear:`. Similarly, I commonly register for notifications in `viewWillAppear:` and unregister in `viewDidDisappear:`. In both cases, this architecture allows me to avoid the retain cycle that would result if I waited for a `dealloc` that might never come ([Chapter 12](#)).
- A certain view that can be shown by switching tab views must reflect the current state of certain user defaults. I refresh the view’s interface in its `viewWillAppear:`, so that whenever it does appear, it is current.
- In a master–detail interface, the detail is a long scrollable text. Whenever the user returns to a previously read detail view, I want to scroll it to wherever it was previously scrolled to. So I save the scroll position for this detail view into the user defaults in its `viewDidDisappear:`.

In the multitasking world of iOS 4 and later, `viewWillDisappear:` and `viewDidDisappear:` are *not* called when the app is suspended into the background. Moreover, once suspended, your app might never return to life; it could be terminated in the background. Some of your functionality performed in `viewWillDisappear:` and `viewDidDisappear:` may have to be duplicated in response to an application lifetime event (Chapter 11), such as `applicationDidEnterBackground:`, if you are to cover every case.

A custom parent view controller, as I explained earlier, must effectively pass `willMoveToParentViewController:` and `didMoveToParentViewController:` to its children manually. But other lifetime events, such as the `appear` events and rotation events, are normally passed along automatically. If you wish to take charge of passing any of these events to your view controller's children, you must take charge of them all:

- `viewWillAppear:`
- `viewDidAppear:`
- `viewWillDisappear:`
- `viewDidDisappear:`
- `willRotateToInterfaceOrientation:duration:`
- `willAnimateRotationToInterfaceOrientation:duration:`
- `didRotateFromInterfaceOrientation:`

To do so, override `automaticallyForwardAppearanceAndRotationMethodsToChildViewControllers` in the parent view controller to return NO.

View Controller Memory Management

Memory management works in a special way for view controllers. Memory is at a premium on a mobile device. A view controller itself is usually lightweight, but a view is memory-intensive, and a view controller might be holding onto a large block of data. A view controller can persist without its view being visible to the user — for example, because a presented view has replaced its view, or because it is in a tab interface but is not the currently selected, or because it is in a navigation interface but is not at the top of the stack. In such a situation, if memory is getting short, then even though the view controller itself persists, the runtime may nilify its view, thus releasing the view and its subviews. The view is effectively *unloaded*.

If a view controller's view is unloaded, then the next time that view is needed for display or mentioned in code, we'll go through the whole rigmarole of loading the view *again* — creating it in code if `loadView` is overridden, or loading the associated nib, and passing once again through `viewDidLoad`. In fact, that's the whole point of the way a view controller's view is loaded lazily in the first place. Memory may become tight, but a view controller's view needn't occupy memory unless it is actually needed, to appear in the interface.

(That is why it's generally better not to have the view controller and its view in the same nib file. In such case, the view can't be unloaded, because there would be no way to load it again. So such a view can't participate in this aspect of iOS app memory management. On the other hand, a view controller containing a view in a storyboard file *does* participate in this aspect of iOS app memory management, presumably because of the dynamic way in which things are loaded from a storyboard by the runtime; unlike a nib, a view controller loaded from a storyboard keeps a reference to that storyboard, and can return to the well to collect its view repeatedly. This could be a reason for using a storyboard.)

It comes as a surprise to beginners (and not-so-beginners) that the runtime can come along behind their backs and nilify an existing view; this seems to introduce a nasty element of indeterminacy into the app's behavior. However, look closely at your `loadView` and `viewDidLoad` overrides. If you've written them sensibly, doing within them only what needs to be done in connection with the view loading, there shouldn't be any problem if these methods are called multiple times. For example, if `viewDidLoad` sets a certain property, using appropriate memory management, what does it matter if it is called again later and sets the very same property to the very same value?

Moreover, view unloading shouldn't present any issues if you are obeying the dictates of model-view-controller ([Chapter 13](#)). The view controller is controller (hence the name). The view, on the other hand, should be just view; its temporary loss, and subsequent restoration, should not pose any special challenges to your code, because you aren't storing anything persistent in the view. The view must be configured, perhaps based on the model, when it comes into existence; as long as the view controller does this, the view should look and behave correctly no matter when and how often it may come into existence. Just don't write your code in such a way that you are *counting* on the view to come into existence just once or at a specific moment.

The object of releasing memory, in the multitasking world of iOS 4 and 5, is partly altruistic and partly selfish. You want to keep your memory usage as low as possible so that other apps can be launched and so that the user can switch between numerous backgrounded apps, bringing each one to the front and finding it in the state in which it was suspended. You also want to prevent your app from being terminated. If your app is backgrounded and is considered a memory hog, it may be terminated when memory runs short; hence you want to reduce your memory usage at the time the app goes into the background. If your app is warned that memory is running short and it doesn't take appropriate action to reduce its memory usage, your app may be killed even while running in the foreground!

When memory runs short, your view controller may receive this sequence of events:

`didReceiveMemoryWarning`

You are invited to release any data that you can do without. Do not release data that you can't readily and quickly recreate! Do not release the view, but do call `super`, which may or may not proceed to release the view. (This call will have been

preceded by a call to the app delegate's `applicationDidReceiveMemoryWarning:`, together with a `UIApplicationDidReceiveMemoryWarningNotification` posted to any registered objects.)

`viewWillUnload`

The view has not yet been unloaded, but it is about to be; `self.view` is still meaningful. You are invited to perform any cleanup related to the view and its subviews, while you still have a reference to it. This event is new in iOS 5.

`viewDidUnload`

The view has been unloaded; `self.view` is nil, and you must not speak of it in code, or you'll accidentally cause it to be loaded again, thus defeating the whole purpose of this exercise. You are invited to release any *other* interface objects that are associated with this view, that you are retaining, and that will be restored when the view loads again; typically, these will be objects assigned to an instance variable, possibly through an outlet, and under ARC you will set the instance variable (or its property) to nil.



The documentation on `viewDidUnload` does not say you have to call `super`. Some of Apple's examples do call `super`, but this is probably a (harmless) mistake.

If you're going to release data in `didReceiveMemoryWarning`, you must concern yourself with how you're going to get it back. You can't rely on `viewDidLoad` for this, because the data might be released without unloading the view, in which case `viewDidLoad` won't be called. The surest approach is to implement a getter that reconstructs or fetches the data if it is nil.

In this example, in `didReceiveMemoryWarning` we write `myBigData` out as a file to disk ([Chapter 36](#)) and release it from memory. At the same time, we override the synthesized accessors for `myBigData` (using the technique shown in [Example 12-10](#)) so that if we subsequently try to get `myBigData` and it's nil, we then try to fetch it from disk and, if we succeed, we delete it from disk (to prevent stale data) and set `myBigData` before returning it. The result is that `myBigData` is released when there's low memory, reducing our memory overhead until we actually *need* `myBigData`, at which time asking for its value (through the getter or property) restores it:

```
@synthesize myBigDataAlias=myBigData;

- (void) setMyBigData: (NSData*) data {
    self.myBigDataAlias = data;
}

- (NSData*) myBigData {
    NSFileManager* fm = [[NSFileManager alloc] init];
    NSString* f = [NSTemporaryDirectory() stringByAppendingPathComponent:@"myBigData"];
    BOOL fExists = [fm fileExistsAtPath:f];
```

```

if (!self.myBigDataAlias) {
    if (fExists) {
        NSData* data = [NSData dataWithContentsOfFile:f];
        self.myBigDataAlias = data;
        NSError* err = nil;
        BOOL ok = [fm removeItemAtPath:f error:&err];
        NSAssert(ok, @"Couldn't remove temp file");
    }
}
return self.myBigDataAlias;
}

- (void)didReceiveMemoryWarning {
    [super didReceiveMemoryWarning];
    if (self->myBigData) {
        NSString* f = [NSTemporaryDirectory() stringByAppendingPathComponent:@"myBigData"];
        [myBigData writeToFile:f atomically:NO];
        self.myBigData = nil;
    }
}

```

Xcode gives you a way to test low-memory circumstances artificially. Run your app in the Simulator; in the Simulator, choose Hardware → Simulate Memory Warning. I don’t believe this has any actual effect on memory, but a memory warning of sufficient severity is sent to your app, so you can see the results of triggering your low-memory response code, including the app delegate’s `applicationDidReceiveMemoryWarning:` and your view controller’s `didReceiveMemoryWarning` and (if the view isn’t showing) `viewDidUnload`.

On the device, the equivalent is call an undocumented method:

```
[[UIApplication sharedApplication]
    performSelector:@selector(_performMemoryWarning)];
```

That’s helpful if your app won’t run on the Simulator (because it uses device-only features), and you can use it in the Simulator as well; basically it’s the code equivalent of Hardware → Simulate Memory Warning. (And remember to remove this code when it is no longer needed, as the App Store won’t accept an app that calls an undocumented method.)

CHAPTER 20

Scroll Views

A scroll view (`UIScrollView`) is a view whose content is larger than its bounds. To reveal a desired area, the user can scroll the content by dragging or flicking, or you can reposition the content in code.

Think of the scroll view as consisting of two things:

The scroll view itself

The scroll view itself acts like a window (a window in a house, not a `UIWindow`). The scroll view's bounds size is the size of that window.

The content

The content is the scene viewed through the window (the scroll view). The content is presumably larger than the scroll view, because otherwise there would be nothing to scroll. By sliding the content, a desired portion of it can be positioned within the scroll view and thus made visible.

Although it is useful to *think* of the scroll view in this way, the truth is far simpler. The scroll view isn't really specially window-like; it's just a view (whose `clipsToBounds` is usually YES). And the content isn't really a separate thing; it's just a set of parameters for positioning the scroll view's subviews. When the scroll view scrolls, what's really changing is the scroll view's own bounds origin; the subviews are positioned with respect to the bounds origin, so they move with it. (See [Chapter 14](#).)

However, a scroll view does bring to the table some nontrivial additional abilities:

- It knows how to shift its bounds origin in response to the user's gestures.
- It provides scroll indicators whose size and position give the user a clue as to the content's size and position.
- It can optionally enforce paging, whereby the user can view only integral portions of the content.
- It can support zooming, so that the user can resize the apparent content by pinching.

Creating a Scroll View

To provide a scroll view with scrollable content, you tell the scroll view how large the content is, by setting its `contentSize`, and you populate the scroll view with subviews whose visibility will be managed as if they were sitting in a view the size of the `contentSize`. It may help to think of the content as a rectangle whose dimensions are given by `contentSize` and whose top left corner is `{0,0}`. When you populate the scroll view, you're effectively placing its subviews into that rectangle. Evidently, you have two choices about how to provide these subviews: you can supply them directly in code, or you can design the scroll view in a nib.

Here's an example of the first approach. Let's start with the Single View Application template so that we have a root view controller. In the view controller's `loadView` I'll create the root view in code (ignoring the view controller nib supplied by the template); I'll create the scroll view and make it the root view, and populate it with 30 `UILabels` whose text contains a sequential number so that we can see where we are when we scroll:

```
- (void) loadView {
    UIScrollView* sv = [[UIScrollView alloc] initWithFrame:
                           [[UIScreen mainScreen] applicationFrame]];
    sv.backgroundColor = [UIColor whiteColor];
    self.view = sv;
    CGFloat y = 10;
    for (int i=0; i<30; i++) {
        UILabel* lab = [[UILabel alloc] init];
        lab.text = [NSString stringWithFormat:@"This is label %i", i+1];
        [lab sizeToFit];
        CGRect f = lab.frame;
        f.origin = CGPointMake(10,y);
        lab.frame = f;
        [sv addSubview:lab];
        y += lab.bounds.size.height + 10;
    }
    CGSize sz = sv.bounds.size;
    sz.height = y;
    sv.contentSize = sz; // This is the crucial line
}
```

The crucial move, as the comment notes, is that we tell the scroll view how large its content is to be. If we omit this step, the scroll view won't be scrollable; the window will appear to consist of a static column of labels.

There is no rule about the order in which you perform the two operations of setting the `contentSize` and populating the scroll view with subviews. In this example, we set the `contentSize` afterward because it is more convenient to track the heights of the subviews as we add them than to calculate their total height in advance. Similarly, you can alter a scroll view's content (subviews) and `contentSize` dynamically as the app runs.

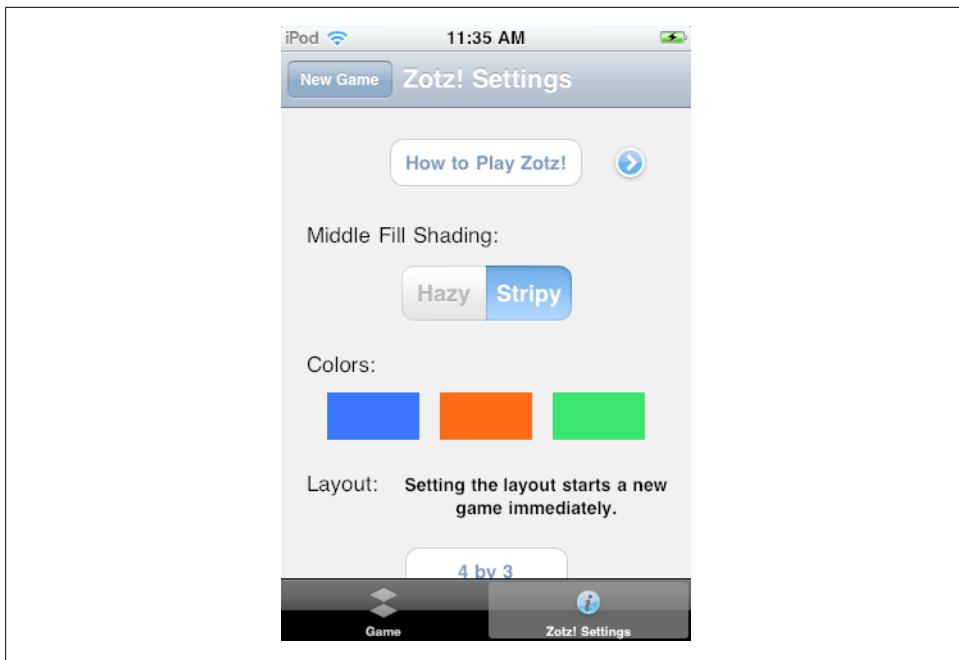


Figure 20-1. The Zotz! settings view

Any direct subviews of the scroll view may need to have their autoresizing set appropriately in case the scroll view is resized, as would happen, for instance, if our root view controller allowed autorotation. To see this, add these lines inside the for loop:

```
lab.backgroundColor = [UIColor redColor]; // make label bounds visible  
lab.autoresizingMask = UIViewAutoresizingFlexibleWidth;
```

Now implement `shouldAutorotateToInterfaceOrientation`: to return YES for all orientations, run the app, and rotate the device (or the Simulator). The labels are wider in portrait orientation because the scroll view itself is wider. Note that this has nothing to do with the `contentSize`! Resizing the `contentSize` has no effect on the size or position of the scroll view's subviews.

Populating a scroll view in code is a very common approach, especially because the content of the scroll view is often not known until runtime. However, sometimes a scroll view is just a way of presenting a fixed view that's larger than the available space; in that case, it is simpler to design the whole scroll view in the nib. The nib editor makes this easy to do.

For example, in my Zotz! app, the user specifies preference settings in a navigation interface inside a tab bar interface, and there isn't enough vertical space for the various interface objects ([Figure 20-1](#)). The obvious solution is that the view shown in this navigation interface should be a scroll view. To lay out this scroll view's subviews in code would be painful and unmaintainable; a nib-based solution is better.

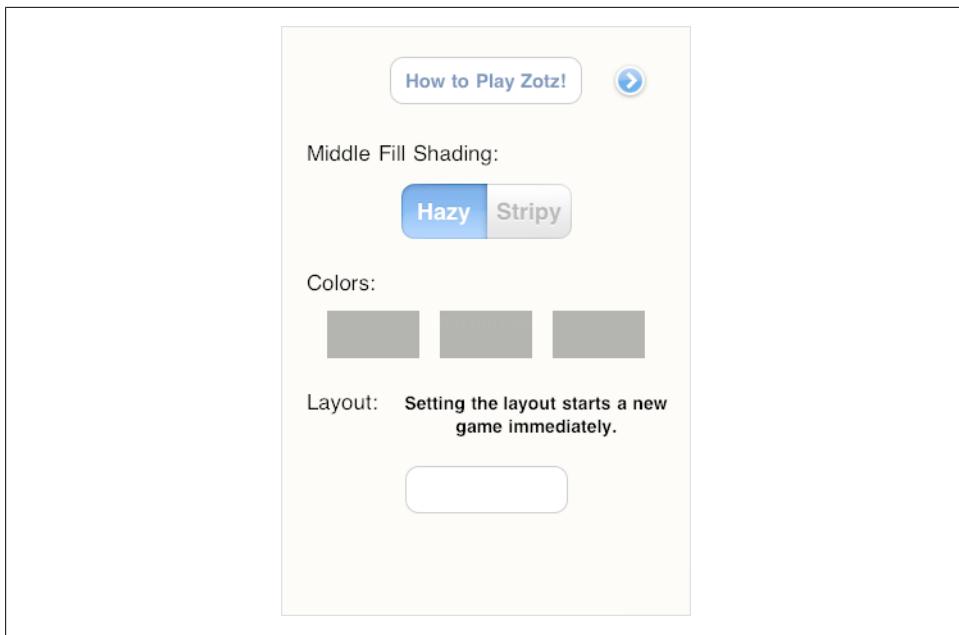


Figure 20-2. The Zotz! settings view, designed in the nib

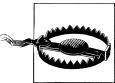
A view in a nib can be a UIScrollView, and this can be any size; its subviews will be the scroll view's subviews when the app runs, which means, as we've already seen, that they will constitute the scroll view's scrollable content. So you design the scroll view's content in the nib exactly as you'd design the content of any view (Figure 20-2). In this case, the scroll view itself is a view controller's view and will automatically be resized appropriately when it is placed into the interface.

Unfortunately, the nib editor provides no way to set the scroll view's `contentSize`. Thus it is up to your code to set an appropriate `contentSize`, large enough to embrace all the scroll view's subviews. But how will your code know what size that is? In the Zotz! app, where the scroll view's content needs to scroll vertically, I solve this problem by means of an outlet to the bottommost subview in the nib. When the view loads, I use the view controller's `viewDidLoad` to learn the y-position of this subview; I allow some additional space at the bottom and set the content size:

```
UIScrollView* sc = (UIScrollView*)self.view;
float width = sc.bounds.width;
// use lowest subview, "layout", as reference for content height
float height = self.layout.frame.origin.y + self.layout.frame.size.height + 20.0;
sc.contentSize = CGSizeMake(width, height);
```

A more elegant approach, perhaps, would have been to put all those views into a single container view that is itself inside the scroll view. The purpose of the container view is to show the scroll view what its `contentSize` should be. We no longer need an outlet because we know that the container view is the scroll view's first subview:

```
UIScrollView* sc = (UIScrollView*)self.view;
sc.contentSize = ((UIView*)[sc.subviews objectAtIndex:0]).bounds.size;
```



Do not assume that the subviews you add to a UIScrollView are its only subviews! The scroll indicators managed by the scroll view, discussed in the next section, are also subviews (they are actually UIImageViews).

Scrolling

For the most part, the purpose of a scroll view will be to let the user scroll. A number of properties affect the user experience with regard to scrolling:

`scrollEnabled`

If NO, the user can't scroll, but you can still scroll in code (as explained later in this section). You could put a UIScrollView to various creative purposes other than letting the user scroll; for example, scrolling in code to a different region of the content might be a way of replacing one piece of interface by another, possibly with animation.

`scrollsToTop`

If YES (the default), and assuming scrolling is enabled, the user can tap on the status bar as a way of making the scroll view scroll its content to the top. You can also override this setting dynamically through the scroll view's delegate (discussed later in this chapter).

`bounces`

If YES (the default), then when the user scrolls to a limit of the content, it is possible to scroll somewhat further (possibly revealing the scroll view's `backgroundColor` behind the content, if a subview was covering it); the content then snaps back into place when the user releases it. Otherwise, the user experiences the limit as a sudden inability to scroll further in that direction.

`alwaysBounceVertical`, `alwaysBounceHorizontal`

If YES, and assuming that `bounces` is YES, then even if the `contentSize` in the given dimension isn't larger than the scroll view (so that no scrolling is actually possible in that dimension), the user can nevertheless scroll somewhat and the content then snaps back into place when the user releases it; otherwise, the user experiences a simple inability to scroll in that dimension.

`directionalLockEnabled`

If YES, and if scrolling is possible in both dimensions (even if only because the appropriate `alwaysBounce...` is YES), then the user, having begun to scroll in one dimension, can't scroll in the other dimension without ending the gesture and starting over. In other words, the user is constrained to scroll vertically or horizontally but not both at once.

`decelerationRate`

The rate at which scrolling is damped out, and the content comes to a stop, after a flick gesture. As convenient examples, standard constants `UIScrollViewDecelerationRateNormal` (0.998) and `UIScrollViewDecelerationRateFast` (0.99) are provided. Lower values mean faster damping; experimentation suggests that values lower than 0.5 are viable but barely distinguishable from one another. You can also effectively override this value dynamically through the scroll view's delegate (discussed later in this chapter).

`showsHorizontalScrollIndicator`, `showsVerticalScrollIndicator`

The scroll indicators are bars that appear only while the user is scrolling in a scrollable dimension (where the content is larger than the scroll view), and serve to indicate both the size of the content in that dimension relative to the scroll view and where the user is within it. The default is YES for both.

Because the user cannot see the scroll indicators except when actively scrolling, there is normally no indication that the view is scrollable. I regard this as somewhat unfortunate, because it makes the possibility of scrolling less discoverable; I'd prefer an option to make the scroll indicators constantly visible. Apple suggests that you call `flashScrollIndicator`s when the scroll view appears, to make the scroll indicators visible momentarily.

`indicatorStyle`

The way the scroll indicators are drawn. Your choices are:

- `UIScrollViewIndicatorStyleDefault` (black with a white border)
- `UIScrollViewIndicatorStyleBlack` (black)
- `UIScrollViewIndicatorStyleWhite` (white)

`contentInset`

A `UIEdgeInsets` struct (four `CGFloats` in the order `top`, `left`, `bottom`, `right`) specifying margins around the content. A typical use for this would be that your scroll view underlaps an interface element, such as a translucent status bar, navigation bar, or toolbar, and you want your content to be visible even when scrolled to its limit.

For example, suppose that our app with the 30 labels has its `Info.plist` configured with the “Status bar style” key set to “Transparent black style,” and that our scroll view’s view controller sets its `wantsFullScreenLayout` to YES. The scroll view now underlaps the status bar. This looks cool while scrolling, but at launch time, and if scrolled all the way to the top, the first label is partly covered by the status bar. We can fix this by supplying a `contentInset` whose `top` matches the height of the status bar. We may also have to scroll the content into position at launch time in code so that it looks right:

```
CGFloat top = [[UIApplication sharedApplication] statusBarFrame].size.height;
sv.contentInset = UIEdgeInsetsMake(top,0,0,0);
[sv scrollRectToVisible:CGRectMake(0,0,1,1) animated:NO];
```

`scrollIndicatorInsets`

A `UIEdgeInsets` struct specifying a shift in the position of the scroll indicators. A typical use is to compensate for the `contentInset`. For example, returning to our scroll view that underlaps the translucent status bar, the content is no longer hidden under the status bar when scrolled to the top, but the top of the vertical scroll indicator is. We can fix this by setting the `scrollIndicatorInsets` to the same value as the `contentInset`.



Here's a trick I've sometimes used: by setting a `scrollIndicatorInsets` component to a negative number and setting the scroll view's `clipsToBounds` to NO, you can make the scroll indicators appear *outside* the scroll view. But because you've turned off `clipsToBounds`, you might have to impose some opaque views on top of the interface to mask off the edges of the scroll view, so that its content isn't visible outside its bounds.

You can scroll in code even if the user can't scroll. The content simply moves to the position you specify, with no bouncing and no exposure of the scroll indicators. You can specify the new position in two ways:

`contentOffset`

The point (`CGPoint`) of the content that is located at the scroll view's top left. The numbers will work correctly if, as I suggested earlier, you pretend that the content is a big rectangle whose own top left is its `{0,0}` point. You can get this property to learn the current scroll position, and set it to change the current scroll position. There is an obvious implication here that you could equally scroll by changing the scroll view's bounds origin; you are unlikely to do that, but the truth is that the bounds origin and the `contentOffset` are effectively the same thing.

To set the `contentOffset` with animation, call `setContentOffset:animated:`. The animation does not cause the scroll indicators to appear; it just slides the content to the desired position.

`scrollRectToVisible:animated:`

Adjusts the content so that the specified `CGRect` of the view is within the scroll view's bounds. This is less precise than setting the `contentOffset`, because you're not saying exactly what the resulting scroll position will be, but sometimes guaranteeing the visibility of a certain portion of the content is exactly what you're after.

If you call a method to scroll with animation and you need to know when the animation ends, implement `scrollViewDidEndScrollingAnimation:` in the scroll view's delegate.

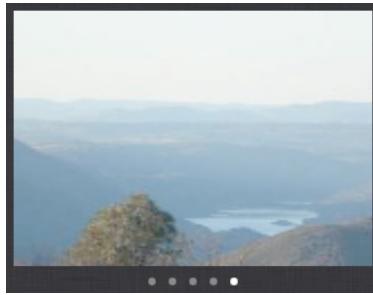


Figure 20-3. A scroll view coordinated with a page control

Paging

If its `pagingEnabled` property is YES, the scroll view doesn't let the user scroll freely; instead, the content is considered to consist of sections the size of the scroll view's bounds, and the user can scroll only in such a way as to move to an adjacent section.

For instance, one of Apple's examples consists of a scroll view containing image views. Each image view is the size of the scroll view. This is an appropriate use of `pagingEnabled`: the user can scroll to see the entire next image or the entire previous image.

The scroll indicator, if it appears, gives the user a sense of how many "pages" constitute the view. Alternatively, you could use delegate messages to coordinate with a UIPager-Controller ([Chapter 25](#)). [Figure 20-3](#) shows my modification of Apple's Scrolling example, where I've added a UIPagerControl below the paging scroll view. Here's the code that updates the page control (`pager`) when the user scrolls:

```
- (void)scrollViewDidEndDecelerating:(UIScrollView *)scrollView {
    CGFloat x = scrollView.contentOffset.x;
    CGFloat w = scrollView.bounds.size.width;
    self.pager.currentPage = x/w;
}
```

And here's the code that scrolls the scroll view (`sv`) when the user taps the page control:

```
- (void) userDidPage: (id) sender {
    NSInteger p = self.pager.currentPage;
    CGFloat w = self.sv.bounds.size.width;
    [self.sv setContentOffset:CGPointMake(p*w,0) animated:YES];
}
```

A useful trick is to have *no* scroll indicator and *no* page control, so that the user has *no* indication of how many "pages" there are, and then to supply pages dynamically as the user scrolls. The result is that the user gets a paging environment that just keeps going and going. In this way, you can display a huge number of pages without having to put them all into the scroll view at once. That, in fact, is how my flashcards apps worked in iOS 3 and iOS 4 (when iOS 5 came along, I replaced that implementation with a UIPagerViewController, [Chapter 19](#)); if you're curious about the technique I was

using, watch the Advanced Scroll View Techniques video from WWDC 2011, which describes something very similar (calling it “infinite scrolling”).

Tiling

Suppose we have some finite but really big content that we want to display in a scroll view, such as a very large image that the user can inspect, piecemeal, by scrolling. To hold the entire image in memory may be onerous or impossible.

Tiling is one solution to this kind of problem. It takes advantage of the insight that there’s really no need to hold the entire image in memory; all we really need at any given moment is the part of the image the user is looking at right now. Mentally, divide the content rectangle into a matrix of rectangles; these rectangles are the tiles. In reality, divide the huge image into corresponding rectangles. Then whenever the user scrolls, we look to see whether part of any empty tile has become visible, and if so, we supply its content. At the same time, we can release the content of all tiles that are completely offscreen. Thus, at any given moment, only the tiles that are showing have content. There is some latency associated with this approach (the user scrolls, then any empty newly visible tiles are filled in), but we will have to live with that.

There is actually a built-in CALayer subclass for helping us implement tiling — CATiledLayer. Its `tileSize` property sets the dimensions of a tile. Its `drawLayer:in-Context:` is called when content for an empty tile is needed; calling `CGContextGetClip-BoundingBox` on the context reveals the location of desired tile, and now we can supply that tile’s content.

To illustrate, we’ll use some tiles already created for us as part of Apple’s own PhotoScroller example. In particular, I’ll use the “Shed_1000” images. These all have names of the form `Shed_1000_x_y.png`, where `x` and `y` are integers corresponding to the picture’s position within the matrix. The images are 256×256 pixels (except for the ones on the extreme right and bottom edges of the matrix, which are shorter in one dimension).

Once again I’ll start with the Single View Application template, adding to it a TiledView class (a UIView subclass) and ignoring the nib file. And once again I’ll implement `loadView` to make the root view controller’s view a UIScrollView; our scroll view’s sole subview will be a TiledView, which exists purely to give our CATiledLayer a place to live. We have just one set of tile images and we want these to appear the same size regardless of the display resolution, so we’ll set the CATiledLayer’s tile size with respect to its native scale:

```
- (void)loadView {
    UIScrollView* sv = [[UIScrollView alloc] initWithFrame:
        [[UIScreen mainScreen] applicationFrame]];
    sv.backgroundColor = [UIColor whiteColor];
    self.view = sv;
    CGRect f = CGRectMake(0,0,9*256,13*256);
    TiledView* content = [[TiledView alloc] initWithFrame:f];
```

```

float tsz = 256 * content.layer.contentsScale;
[(CATiledLayer*)content.layer setTileSize: CGSizeMake(tsz, tsz)];
[self.view addSubview:content];
[sv setContentSize: f.size];
}

```

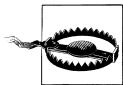
Here's the code for TiledView. The CATiledLayer is our underlying layer; therefore we are its delegate. This means that when `drawLayer:inContext:` is called, `drawRect:` is called, and the argument to `drawRect:` is the same as the result of calling `CGContextGetClipBoundingBox`, namely, it's the rect of the tile we are to draw. As Apple's code points out, we must fetch images with `imageWithContentsOfFile:` so as to avoid the automatic caching behavior of `imageNamed:`, because we're doing all this exactly to avoid using more memory than we have to:

```

+ (Class) layerClass {
    return [CATiledLayer class];
}

-(void)drawRect:(CGRect)r {
    CGRect tile = r;
    int x = tile.origin.x/256;
    int y = tile.origin.y/256;
    NSString *tileName = [NSString stringWithFormat:@"Shed_1000_%i_%i", x, y];
    NSString *path =
        [[NSBundle mainBundle] pathForResource:tileName ofType:@"png"];
    UIImage *image = [UIImage imageWithContentsOfFile:path];
    [image drawAtPoint:tile.origin];
    // uncomment the following to see the tile boundaries
    /*
    UIBezierPath* bp = [UIBezierPath bezierPathWithRect: r];
    [[UIColor whiteColor] setStroke];
    [bp stroke];
    */
}

```



This code can be dangerous on a system earlier than iOS 4. As explained in Apple's tech note QA1637, "CATiledLayer and UIKit graphics," a CATiledLayer draws itself on a background thread. Starting in iOS 4, accessing the current context and drawing to it with UIKit is safe in a background thread; before that, it wasn't.

There is no special call for invalidating an offscreen tile. You can call `setNeedsDisplay` or `setNeedsDisplayInRect:` on the TiledView, but this doesn't erase offscreen tiles. You're just supposed to trust that the CATiledLayer will eventually clear offscreen tiles if needed to conserve memory.

CATiledLayer has a class method `fadeDuration` that dictates the duration of the animation that fades a new tile into view. You can create a CATiledLayer subclass and override this method to return a value different from the default (0.25), but in general this is probably not worth doing, as the default value is a good one. Returning a smaller

value won't make tiles appear faster; it just replaces the nice fade-in with an annoying flash.

Zooming

To implement zooming of a scroll view's content, you set the scroll view's `minimumZoomScale` and `maximumZoomScale` so that at least one of them isn't 1 (the default). You also implement `viewForZoomingInScrollView:` in the scroll view's delegate to tell the scroll view which of its subviews is to be the scalable view. The scroll view then zooms by applying a scaling transform ([Chapter 14](#)) to this subview. Typically, you'll want the scroll view's entire content to be scalable, so you'll have one direct subview of the scroll view that acts as the scalable view, and anything else inside the scroll view will be a subview of the scalable view, so as to be scaled together with it.

To illustrate, let's return to the first example in this chapter, where we created a scroll view containing 30 labels. To make this scroll view zoomable, we'll need to modify the way we create it. As it stands, the scroll view's subviews are just the 30 labels; there is no single view that we would scale in order to scale all the labels together. This time, as we create the scroll view in our root view controller's `loadView` implementation, instead of making the 30 labels subviews of the scroll view, we'll make them subviews of a single scalable view and make the scalable view the subview of the scroll view:

```
UIScrollView* sv = [[UIScrollView alloc] initWithFrame:  
    [[UIScreen mainScreen] applicationFrame]];  
self.view = sv;  
UIView* v = [[UIView alloc] init];  
CGFloat y = 10;  
for (int i=0; i<30; i++) {  
    UILabel* lab = [[UILabel alloc] init];  
    lab.text = [NSString stringWithFormat:@"This is label %i", i+1];  
    [lab sizeToFit];  
    CGRect f = lab.frame;  
    f.origin = CGPointMake(10,y);  
    lab.frame = f;  
    [v addSubview:lab];  
    y += lab.bounds.size.height + 10;  
}  
CGSize sz = sv.bounds.size;  
sz.height = y;  
sv.contentSize = sz;  
v.frame = CGRectMake(0,0,sz.width,sz.height);  
[sv addSubview:v];
```

So far, nothing has changed; the scroll view works just as before, but it isn't zoomable. To make it zoomable, we add these lines:

```
v.tag = 999;  
sv.minimumZoomScale = 1.0;  
sv.maximumZoomScale = 2.0;  
sv.delegate = self;
```

We have assigned a tag to the view that is to be scaled, so we can find it later. We have set the scale limits for the scroll view. And we have made ourselves the scroll view's delegate. Now all we have to do is implement `viewForZoomingInScrollView:` and return the scalable view:

```
- (UIView *)viewForZoomingInScrollView:(UIScrollView *)scrollView {
    return [scrollView viewWithTag:999];
}
```

The scroll view now responds to pinch gestures by scaling appropriately! The user can actually scale considerably beyond the limits we set in both directions; when the gesture ends, the scale returns to the limit value. If we wish to confine scaling strictly to our defined limits, we can set the scroll view's `bouncesZoom` to NO; when the user reaches a limit, scaling will simply stop.



The scroll view zooms by applying a scaling transform to the scalable view; therefore *the frame of the scalable view is scaled as well*. Moreover, the scroll view is concerned to make scrolling continue to work correctly: the limits as the user scrolls should continue to match the limits of the content, and commands like `scrollRectToVisible:animated:` should continue to work the same way for the same values. Therefore, the scroll view *automatically scales its own contentSize* to match the current zoom scale. (You can actually detect this happening by overriding `setContentSize` in a UIScrollView subclass: you can see the scroll view adjusting its own content size as you zoom.)

If the `minimumZoomScale` is less than 1, then when the scalable view becomes smaller than the scroll view, it is pinned to the scroll view's top left. If you don't like this, you can change it by subclassing UIScrollView and overriding `layoutSubviews`, or by implementing the scroll view delegate method `scrollViewDidZoom:`. Here's a simple example (drawn from a WWDC 2010 video) demonstrating an override of `layoutSubviews` that keeps the scalable view centered when it becomes smaller than the scroll view:

```
-(void)layoutSubviews {
    [super layoutSubviews];
    UIView* v = [self.delegate viewForZoomingInScrollView:self];
    CGFloat svw = self.bounds.size.width;
    CGFloat svh = self.bounds.size.height;
    CGFloat vw = v.frame.size.width;
    CGFloat vh = v.frame.size.height;
    CGRect f = v.frame;
    if (vw < svw)
        f.origin.x = (svw - vw) / 2.0;
    else
        f.origin.x = 0;
    if (vh < svh)
        f.origin.y = (svh - vh) / 2.0;
```

```
        else
            f.origin.y = 0;
        v.frame = f;
    }
```

Zooming Programmatically

To zoom programmatically, you have two choices:

`setZoomScale:animated:`

Zooms in terms of scale value. The `contentOffset` is automatically adjusted to keep the current center centered and the content occupying the entire scroll view.

`zoomToRect:animated:`

Zooms so that the given rectangle of the content occupies as much as possible of the scroll view's bounds. The `contentOffset` is automatically adjusted to keep the content occupying the entire scroll view.

In this example, I implement double-tapping as a zoom gesture. Detecting the double-tap is easy thanks to a gesture recognizer attached to the scalable view ([Chapter 18](#)). In this implementation of the action handler for the double-tap UITapGestureRecognizer, a double-tap means to zoom to maximum scale, minimum scale, or actual size, depending on the current scale value:

```
- (void) tapped: (UIGestureRecognizer*) tap {
    UIView* v = tap.view;
    UIScrollView* sv = (UIScrollView*)v.superview;
    if (sv.zoomScale < 1) {
        [sv setZoomScale:1 animated:YES];
        CGPoint pt =
            CGPointMake((v.bounds.size.width - sv.bounds.size.width)/2.0,0);
        [sv setContentOffset:pt animated:NO];
    }
    else if (sv.zoomScale < sv.maximumZoomScale)
        [sv setZoomScale:sv.maximumZoomScale animated:YES];
    else
        [sv setZoomScale:sv.minimumZoomScale animated:YES];
}
```

Zooming with Detail

By default, when a scroll view zooms, it merely applies a scale transform to the scaled view. The scaled view's drawing is cached beforehand into its layer, and the bits of the resulting bitmap are drawn larger. This means that a zoomed-in scroll view's content may be fuzzy (pixellated). In some cases this might be acceptable, but in others you might like the content to be redrawn more sharply at its new size.



On a double-resolution device, this might not be such an issue. For example, if the user is allowed to scroll only up to double scale, you can draw at double scale right from the start; the results will look good at single scale, because the screen has double resolution, as well as at double scale, because that's the scale you drew at.

One solution is to take advantage of a CATiledLayer feature that I didn't mention earlier. It turns out that CATiledLayer is aware not only of scrolling but also of scaling: you can configure it to ask for tiles to be drawn when the layer is scaled to a new order of magnitude. This approach is extremely easy: your drawing routine is called and you simply draw, the graphics context itself having already been scaled appropriately. In fact, your drawing doesn't even have to involve multiple tiles! Of course it *can* involve tiles; for a large tiled image, you would be forearmed with multiple versions of the image broken into an identical quantity of tiles, each set having double the tile size of the previous set (as in Apple's PhotoScroller example). But you can also just draw directly.



This technique is also demonstrated in the “Basic Zooming Using the Pinch Gestures” chapter of Apple’s *ScrollView Programming Guide for iOS*. That chapter ends with a warning that “the UIKit drawing methods are not thread-safe,” but you can ignore that warning if you’re targeting iOS 4 and iOS 5, as it’s no longer true. As throughout this book, the code in this chapter makes no attempt to be compatible with iOS 3.

Besides its `tileSize`, you’ll need to set two additional CATiledLayer properties:

`levelsOfDetail`

The number of different resolutions at which you want to redraw, where each level has twice the resolution of the previous level. So, for example, with two levels of detail we can ask to redraw when zooming to double size (2x) and when zooming back to single size (1x).

`levelsOfDetailBias`

The number of levels of detail that are *larger* than single size (1x). For example, if `levelsOfDetail` is 2, then if we want to redraw when zooming to 2x and when zooming back to 1x, the `levelsOfDetailBias` is 1, because one of those levels is larger than 1x; if we were to leave `levelsOfDetailBias` at 0 (the default), we would be saying we want to redraw when zooming to 0.5x and back to 1x — we have two levels of detail but neither is larger than 1x, so one must be smaller than 1x.

So, just to hammer home the point, let’s say we want to redraw at .5x, 1x, 2x, and 4x. That’s four levels of detail, so `levelsOfDetail` is 4; and two of them are larger than 1x, so `levelsOfDetailBias` is 2.

The CATiledLayer will ask for a redraw at a higher resolution as soon as the view’s size becomes larger than the previous resolution. In other words, if there are two levels of

detail with a bias of 1, the layer will be redrawn at 2x as soon as it is zoomed even a little bit larger than 1x. This is an excellent approach, because although a level of detail would look blurry if scaled up, it looks pretty good scaled down. There's just one catch: your settings will need to depend on the screen resolution if you want them to work the same way on both single- and double-resolution devices.

To illustrate, I'll reuse our previous example, where the root view controller's view is a scroll view whose subview is a TiledView that hosts a CATiledLayer; but this time I'll draw our 30 labels into the CATiledLayer. The tiles are of no interest during scrolling — we're not going to run low on memory drawing a few labels — so we may as well set the tile size to the bounds size:

```
- (void)loadView {
    UIScrollView* sv = [[UIScrollView alloc] initWithFrame:
        [[UIScreen mainScreen] applicationFrame]];
    self.view = sv;
    CGRect f = CGRectMake(0,0,self.view.bounds.size.width,940);
    TiledView* content = [[TiledView alloc] initWithFrame:f];
    content.tag = 999;
    CATiledLayer* lay = (CATiledLayer*)content.layer;
    lay.tileSize = f.size;
    lay.levelsOfDetail = 2;
    lay.levelsOfDetailBias = 1;
    [self.view addSubview:content];
    [sv setContentSize: f.size];
    sv.minimumZoomScale = 1.0;
    sv.maximumZoomScale = 2.0;
    sv.delegate = self;
    // deal with double-resolution screen
    if ([[UIScreen mainScreen] scale] > 1.0) {
        f.size.width *= 4;
        f.size.height *= 4;
        lay.tileSize = f.size;
        lay.levelsOfDetailBias = 2;
    }
}
- (UIView *)viewForZoomingInScrollView:(UIScrollView *)scrollView {
    return [scrollView viewWithTag:999];
}
```

Note the compensation for the double-resolution screen. On a double-resolution device, the CATiledLayer starts life at the 2x level of detail, so instead of asking to be redrawn at 1x and 2x, we ask to be redrawn at 2x and 4x. Second, because we are at double resolution at the 2x level of detail, we must make our tiles *four* times bigger if a single tile is to be the size of the view initially.

Here's the code for TiledView. Its `drawRect:` essentially does the work of putting the labels into place that we were previously doing in `loadView`, except that now there are no labels: we're in `drawRect:` so we draw the text directly. The important thing to notice is that we do nothing whatever about zooming and scaling! The incoming `drawRect:` argument is the bounds of our view, and we just draw into it:

```

+ (Class) layerClass {
    return [CATiledLayer class];
}

-(void)drawRect:(CGRect)r {
    [[UIColor whiteColor] set];
    UIRectFill(self.bounds);
    [[UIColor blackColor] set];
    UIFont* f = [UIFont fontWithName:@"Helvetica" size:18];
    // height consists of 31 spacers with 30 texts between them
    CGFloat viewh = self.bounds.size.height;
    CGFloat spacerh = 10;
    CGFloat texth = (viewh - (31*spacerh))/30.0;
    CGFloat y = spacerh;
    for (int i = 0; i < 30; i++) {
        NSString* s = [NSString stringWithFormat:@"This is label %i", i];
        [s drawAtPoint:CGPointMake(10,y) withFont:f];
        y += texth + spacerh;
    }
    // uncomment the following to see the tiling
/*
    UIBezierPath* bp = [UIBezierPath bezierPathWithRect:r];
    [[UIColor redColor] setStroke];
    [bp stroke];
*/
}

```

An alternative and much simpler approach (which I didn’t know about when I wrote the first edition of this book — I got it from a WWDC 2011 video) is to make yourself the scroll view’s delegate so that you get an event when the zoom ends, and then change the scalable view’s `contentScaleFactor` to match the current zoom scale, compensating for the double-resolution screen at the same time:

```

- (void)scrollViewDidEndZooming:(UIScrollView *)scrollView
    withView:(UIView *)view
    atScale:(float)scale {
    view.contentScaleFactor = scale * [UIScreen mainScreen].scale;
}

```

That approach comes with a caveat, however: you mustn’t overdo it. If the zoom scale, screen resolution, and scalable view size are high, you will be asking for a very large graphics context to be maintained in memory, which could cause your app to run low on memory or even to be abruptly terminated by the system.

ScrollView Delegate

The scroll view’s delegate (adopting the `UIScrollViewDelegate` protocol) receives lots of messages that can help you track what the scroll view is up to:

`scrollViewDidScroll:`

If you scroll in code without animation, you will receive this message *once*. If the user drags or flicks, or uses the scroll-to-top feature, or if you scroll in code with

animation, you will receive this message *repeatedly* throughout the scroll, including during the time the scroll view is decelerating after the user's finger has lifted; there are other delegate messages that tell you, in those cases, when the scroll has really ended.

scrollViewDidEndScrollingAnimation:

If you scroll in code with animation, you will receive this message when the animation ends.

scrollViewWillBeginDragging:

scrollViewWillEndDragging:withVelocity:targetContentOffset:

scrollViewDidEndDragging:willDecelerate:

If the user scrolls by dragging or flicking, you will receive these messages at the start and end of the user's finger movement. If the user brings the scroll view to a stop before lifting the finger, `willDecelerate` is NO and the scroll is over. If the user lets go of the scroll view while the finger is moving, or if paging is turned on and the user has not paged perfectly already, `willDecelerate` is YES and we proceed to the delegate messages reporting deceleration.

The purpose of `scrollViewWillEndDragging:....`, new in iOS 5, is to let you customize the outcome of the content's deceleration. The third argument is a pointer to a `CGPoint`; thus you can use it to set a different `CGPoint`, specifying the `contentOffset` value the content should have when the deceleration is over.

scrollViewWillBeginDecelerating:

scrollViewDidEndDecelerating:

Sent once each after `scrollViewDidEndDragging:willDecelerate:` arrives with a value of YES. When `scrollViewDidEndDecelerating:` arrives, the scroll is over.

scrollViewShouldScrollToTop:

scrollViewDidScrollToTop:

You won't get either of these if `scrollsToTop` is NO, because the scroll-to-top feature is turned off in that case. The first lets you prevent the user from scrolling to the top on this occasion even if `scrollsToTop` is YES. The second tells you that the user has employed this feature and the scroll is over.

In addition, the scroll view has read-only properties reporting its state:

tracking

The user has touched the scroll view, but the scroll view hasn't decided whether this is a scroll or some kind of tap.

dragging

The user is dragging to scroll.

decelerating

The user has scrolled and has lifted the finger, and the scroll is continuing.

So, if you wanted to do something after a scroll ends completely regardless of how the scroll was performed, you'd need to implement many delegate methods:

- `scrollViewDidEndDragging:willDecelerate`: in case the user drags and stops (`willDecelerate` is NO).
- `scrollViewDidEndDecelerating`: in case the user drags and the scroll continues afterward.
- `scrollViewDidScrollToTop`: in case the user uses the scroll-to-top feature.
- `scrollViewDidEndScrollingAnimation`: in case you scroll in code with animation.

You don't need a delegate method to tell you when the scroll is over after you scroll in code without animation: it's over immediately, so if you have work to do after the scroll ends, you can do it in the next line of code.

There are also three delegate messages that report zooming:

`scrollViewWillBeginZooming:withView:`

If the user zooms or you zoom in code, you will receive this message as the zoom begins.

`scrollViewDidZoom:`

If you zoom in code, even with animation, you will receive this message *once*. If the user zooms, you will receive this message *repeatedly* as the zoom proceeds. (You will probably also receive `scrollViewDidScroll:`, possibly many times, as the zoom proceeds.)

`scrollViewDidEndZooming:withView:atScale:`

If the user zooms or you zoom in code, you will receive this message after the last `scrollViewDidZoom:`.

In addition, the scroll view has read-only properties reporting its state during a zoom:

`zooming`

The scroll view is zooming. It is possible for `dragging` to be true at the same time.

`zoomBouncing`

The scroll view is returning automatically from having been zoomed outside its minimum or maximum limit. As far as I can tell, you'll get only one `scrollViewDidZoom:` while the scroll view is in this state.

ScrollView Touches

Improvements in the scroll view implementation have eliminated most of the worry once associated with scroll view touches. A scroll view will interpret a drag or a pinch as a command to scroll or zoom, and any other gesture will fall through to the subviews; thus buttons and similar interface objects inside a scroll view work just fine.

You can even put a scroll view inside a scroll view, and this can be quite a useful thing to do, in contexts where you might not think of it at first. A WWDC 2010 presentation uses as an example Apple’s Photos app, where a single photo fills the screen: you can page-scroll from one photo to the next, and you can zoom the current photo with a pinch-out gesture. This, the presentation demonstrates, can be implemented with a scroll view inside a scroll view: the outer scroll view is for paging between images, and the inner scroll view contains the current image and is for zooming.

Gesture recognizers ([Chapter 18](#)) have also greatly simplified the task of adding custom gestures to a scroll view. For instance, some older code in Apple’s documentation, showing how to implement a double-tap to zoom in and a two-finger tap to zoom out, uses old-fashioned touch handling, but this is no longer necessary. Simply attach to your scroll view’s scalable subview any gesture recognizers for these sorts of gesture, and they will mediate automatically among the possibilities.

In the past, making something inside a scroll view draggable required setting the scroll view’s `canCancelContentTouches` property to NO. (The reason for the name is that the scroll view, when it realizes that a gesture is a drag or pinch gesture, normally sends `touchesCancelled:forEvent:` to a subview tracking touches, so that the scroll view and not the subview will be affected.) However, unless you’re implementing old-fashioned direct touch handling, you probably won’t have to concern yourself with this. Regardless of how `canCancelContentTouches` is set, a draggable control, such as a `UISlider`, remains draggable inside a scroll view.

On the other hand, something like a `UISlider` might prove more quickly responsive if you set the scroll view’s `delaysContentTouches` to NO. Without this, the user may have to hold a finger on the slider briefly before it becomes draggable. But even this will be a concern only if the scroll view is scrollable in the same dimension as the slider is oriented; a horizontal slider in a scroll view that can be scrolled only vertically is instantly draggable.

Here’s an example of a draggable object inside a scroll view implemented through a gesture recognizer. Suppose we have an image of a map, larger than the screen, and we want the user to be able to scroll it in the normal way to see any part of the map, but we also want the user to be able to drag a flag into a new location on the map. I’ll arrange this as I configure the scroll view in `loadView`, by attaching to the flag a `UIPanGestureRecognizer` with the same `dragging:` action handler developed in [Chapter 18](#):

```
- (void) loadView {
    UIScrollView* sv = [[UIScrollView alloc]
        initWithFrame:[[UIScreen mainScreen] applicationFrame]];
    self.view = sv;
    UIImageView* imv =
        [[UIImageView alloc] initWithImage: [UIImage imageNamed:@"map.jpg"]];
    [sv addSubview:imv];
    sv.contentSize = imv.bounds.size;
    UIImageView* flag = [[UIImageView alloc] initWithImage:
        [UIImage imageNamed:@"redflag.png"]];
    [sv addSubview: flag];
```

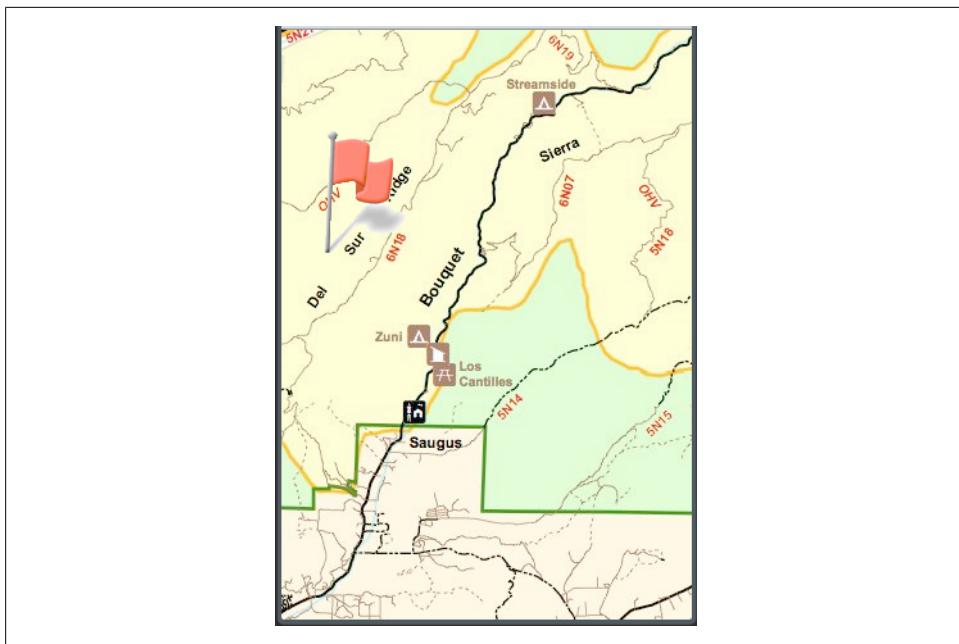


Figure 20-4. A scrollable map with a draggable flag

```

UIPanGestureRecognizer* pan = [[UIPanGestureRecognizer alloc]
    initWithTarget:self
    action:@selector(dragging:)];
[flag addGestureRecognizer:pan];
flag.userInteractionEnabled = YES;
}

- (void) dragging: (UIPanGestureRecognizer*) p {
    UIView* v = p.view;
    if (p.state == UIGestureRecognizerStateBegan ||
        p.state == UIGestureRecognizerStateChanged) {
        CGPoint delta = [p translationInView: v.superview];
        CGPoint c = v.center;
        c.x += delta.x; c.y += delta.y;
        v.center = c;
        [p setTranslation: CGPointMakeZero inView: v.superview];
    }
}
  
```

The user can now drag the map or the flag (Figure 20-4). The state of the scroll view's `canCancelContentTouches` is irrelevant, because the flag view isn't tracking the touches manually.

An interesting addition to that example would be to implement autoscrolling, meaning that the scroll view scrolls itself when the user drags the flag close to its edge. This, too, is greatly simplified by gesture recognizers; in fact, we can add autoscrolling code directly to the `dragging:` action handler:

```

- (void) dragging: (UIPanGestureRecognizer*) p {
    UIView* v = p.view;
    if (p.state == UIGestureRecognizerStateBegan || 
        p.state == UIGestureRecognizerStateChanged) {
        CGPoint delta = [p translationInView: v.superview];
        CGPoint c = v.center;
        c.x += delta.x; c.y += delta.y;
        v.center = c;
        [p setTranslation: CGPointMakeZero inView: v.superview];
    }
    // autoscroll
    if (p.state == UIGestureRecognizerStateChanged) {
        CGPoint loc = [p locationInView: self.view.superview];
        CGRect f = self.view.frame;
        UIScrollView* sv = (UIScrollView*)self.view;
        CGPoint off = sv.contentOffset;
        CGSize sz = sv.contentSize;
        CGPoint c = v.center;
        // to the right
        if (loc.x > CGRectGetMaxX(f) - 30) {
            CGFloat margin = sz.width - CGRectGetMaxX(sv.bounds);
            if (margin > 6) {
                off.x += 5;
                sv.contentOffset = off;
                c.x += 5;
                v.center = c;
                [self performSelector:@selector(dragging:)
                    withObject:p afterDelay:0.2];
            }
        }
        // to the left
        if (loc.x < f.origin.x + 30) {
            CGFloat margin = off.x;
            if (margin > 6) {
                // ... omitted ...
            }
        }
        // to the bottom
        if (loc.y > CGRectGetMaxY(f) - 30) {
            CGFloat margin = sz.height - CGRectGetMaxY(sv.bounds);
            if (margin > 6) {
                // ... omitted ...
            }
        }
        // to the top
        if (loc.y < f.origin.y + 30) {
            CGFloat margin = off.y;
            if (margin > 6) {
                // ... omitted ...
            }
        }
    }
}

```

The material marked as omitted in the second, third, and fourth cases is obviously parallel to the first case, and is left as an exercise for the reader.

Starting in iOS 5, a scroll view's touch handling is itself based on gesture recognizers attached to the scroll view — and these are available to your code through the scroll view's `panGestureRecognizer` and `pinchGestureRecognizer` properties. This means that if you want to customize a scroll view's touch handling, it's easy to add more gesture recognizers and have them interact with those already attached to the scroll view.

To illustrate, I'll build on the previous example. Suppose we want the flag to start out offscreen, and we'd like the user to be able to summon it with a rightward swipe. We can attach a `UISwipeGestureRecognizer` to our scroll view, but it will never recognize its gesture because the scroll view's own pan gesture recognizer will recognize first. But we have access to the scroll view's panning gesture recognizer, so we can compel it to yield to our swipe gesture recognizer by sending it `requireGestureRecognizerToFail::`:

```
- (void) loadView {
    UIScrollView* sv = [[UIScrollView alloc] initWithFrame:
        [[UIScreen mainScreen] applicationFrame]];
    self.view = sv;
    UIImageView* imv = [[UIImageView alloc] initWithImage:
        [UIImage imageNamed:@"map.jpg"]];
    [sv addSubview:imv];
    sv.contentSize = imv.bounds.size;
    UISwipeGestureRecognizer* swipe =
        [[UISwipeGestureRecognizer alloc] initWithTarget:self
            action:@selector(swiped:)];
    [sv addGestureRecognizer:swipe];
    [sv.panGestureRecognizer requireGestureRecognizerToFail:swipe];
    // no flag; user must summon it with a swipe
}
```

Here's my implementation of `swiped::`:

```
- (void) swiped: (UISwipeGestureRecognizer*) g {
    if (g.state == UIGestureRecognizerStateEnded ||
        g.state == UIGestureRecognizerStateCancelled) {
        UIImageView* flag = [[UIImageView alloc]
            initWithImage: [UIImage imageNamed:@"redflag.png"]];
        UIPanGestureRecognizer* pan = [[UIPanGestureRecognizer alloc]
            initWithTarget:self
            action:@selector(dragging:)];
        [flag addGestureRecognizer:pan];
        flag.userInteractionEnabled = YES;
        // put the flag offscreen and animate it onto the screen
        UIScrollView* sv = (UIScrollView*)self.view;
        CGPoint p = sv.contentOffset;
        CGRect f = flag.frame;
        f.origin = p;
        f.origin.x -= flag.bounds.size.width;
        flag.frame = f;
        [sv addSubview: flag];
        [UIView beginAnimations:nil context:NULL];
        f.origin.x = p.x;
```

```

    flag.frame = f;
    [UIView commitAnimations];
    // thanks for the flag, now stop operating altogether
    g.enabled = NO;
}
}

```

ScrollView Performance

At several points in earlier chapters I've mentioned performance problems and ways to increase drawing efficiency. Nowhere are you so likely to need these as in connection with a scroll view. As a scroll view scrolls, views must be drawn very rapidly as they appear on the screen. If the view-drawing system can't keep up with the speed of the scroll, the scrolling will visibly stutter.

Performance testing and optimization is a big subject, so I can't tell you exactly what to do if you encounter stuttering while scrolling. But certain general suggestions (mostly extracted from a really great presentation at the 2010 WWDC) should come in handy:

- Everything that can be opaque should be opaque: don't force the drawing system to composite transparency, and remember to tell it that an opaque view or layer *is* opaque by setting its `opaque` property to YES. If you really must composite transparency, keep the size of the nonopaque regions to a minimum; for example, if a large layer is transparent at its edges, break it into five layers — the large central layer, which is opaque, and the four edges, which are not.
- If you're drawing shadows, don't make the drawing system calculate the shadow shape for a layer: supply a `shadowPath`, or use Core Graphics to create the shadow with a drawing. Similarly, avoid making the drawing system composite the shadow as a transparency against another layer; for example, if the background layer is white, your opaque drawing can itself include a shadow already drawn on a white background.
- Don't make the drawing system scale images for you; supply the images at the target size for the correct resolution.
- In a pinch, you can just eliminate massive swatches of the rendering operation by setting a layer's `shouldRasterize` to YES. You could, for example, do this when scrolling starts and then set it back to NO when scrolling ends.

Apple's documentation also says that setting a view's `clearsContextBeforeDrawing` to NO may make a difference. I can't confirm or deny this; it may be true, but I haven't encountered a case that positively proves it.

As I've already mentioned, Xcode provides tools that will help you detect inefficiencies in the drawing system. In the Simulator, the Debug menu shows you blended layers (where transparency is being composited) and images that are being copied, misaligned, or rendered offscreen (rendering directly to the screen is always more efficient). On the device, the Core Animation module of Instruments provides the same functionality,

plus it tracks the frame rate for you, allowing you to scroll and measure performance objectively.

CHAPTER 21

Table Views

*I'm gonna ask you the three big questions. — Go ahead.
— Who made you? — You did. — Who owns the biggest
piece of you? — You do. — What would happen if I
dropped you? — I'd go right down the drain.*

—Dialogue by Garson Kanin and Ruth Gordon,
Pat and Mike

A table view (UITableView) is a scrolling interface (a vertically scrolling UIScrollView, [Chapter 20](#)) for presenting a single column of rectangular cells (UITableViewCell, a UIView subclass). It is a keystone of Apple’s strategy for making the small iPhone screen useful and powerful, and has three main purposes:

Presentation of information

The cells typically contain text, which the user can read. The cells are usually quite small, in order to maximize the number of them that appear on the screen at once, so this text is often condensed, truncated, or otherwise simplified.

Selection

A cell can be selected by tapping. A table view can thus be used to provide the user with a column of choices. The user chooses by tapping a cell, and the app responds appropriately to that choice.

Navigation

The appropriate response to the user’s choosing a cell is often navigation to another interface. This might be done, for example, through a presented view controller or a navigation interface ([Chapter 19](#)). An extremely common configuration is a master–detail interface, in which the master view is (or contains) a table view, often at the root of a navigation interface; the user taps a listing in the table to navigate to the details for that choice. This is one reason why truncation of text in a table view is acceptable: the detail view contains the full information.

In addition to its column of cells, a table view can be extended by a number of other features that make it even more useful and flexible:

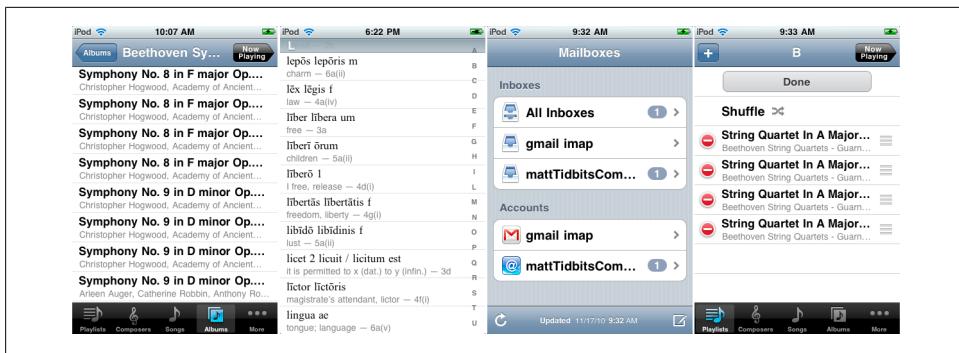


Figure 21-1. Four table view variations

- A table can start with a header view (at the top) and end with a footer view (at the bottom).
- The cells can be clumped into sections. Each section can have a header and footer, and these remain visible as long as the section itself occupies the screen, giving the user a clue as to where we are within the table. Moreover, a section index can be provided, in the form of an overlay column of abbreviated section titles, which the user can tap to jump to the start of a section (thus making a long table tractable).
- A table can have a “grouped” format. This is often used for presenting small numbers of related cells.
- Tables can be editable: the user can be permitted to insert, delete, and reorder cells.

Figure 21-1 illustrates four variations of the table view:

1. Apple’s Music app lists song titles and artists for a given album in truncated form in a table view within a navigation interface which is itself within a tab bar interface; tapping an album in a table of album titles summons this list of songs within that album, and tapping a song in this list plays it.
2. An app of mine lists Latin words and their definitions in alphabetical order, divided into sections by first letter, with section headers and a section index.
3. Apple’s Mail app lists inboxes and accounts in a grouped format, clumped into sections with headers.
4. Apple’s Music app allows a custom playlist to be edited, with interface for deleting and rearranging cells.

Table cells, too, can be extremely flexible. Some basic table cell formats are provided, such as a text label along with a small image view, but you are free to design your own table cell as you would any other view. There are also some standard interface items that are commonly used in a table cell, such as a checkmark to indicate selection or a right-pointing chevron to indicate that tapping the cell navigates to a detail view.

It would be difficult to overestimate the importance of table views. An iOS app without a table view somewhere in its interface would be a rare thing, especially on the small

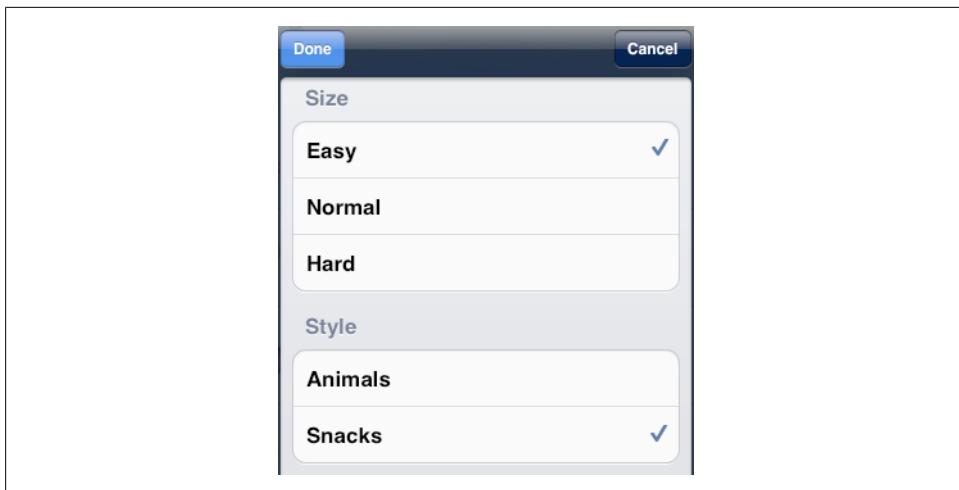


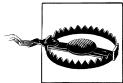
Figure 21-2. A grouped table view as an interface for choosing options

iPhone screen. I've written apps consisting almost entirely of table views. Indeed, it is not uncommon to use a table view even in situations where there is nothing particularly table-like about that situation, simply because it is so convenient. For example, in one of my apps I want the user to be able to choose between three levels of difficulty. In a desktop application I'd probably use radio buttons; but there are no radio buttons among the standard iOS interface objects. Instead, I use a grouped table so small that it doesn't even scroll. This gives me a section header, three tappable cells, and a checkmark indicating the current choice (Figure 21-2).

There is a UIViewController subclass, UITableViewController, dedicated to the presentation of a table view. It is important to stress, however, that you never really *need* to use a UITableViewController — it's just a convenience, and doesn't do anything that you couldn't do yourself by other means — and that if your table view does not constitute the *entire* view to be handled by a UIViewController, you *can't* use a UITableViewController.

Here's some of what using a UITableViewController gives you:

- UITableViewController's `initWithStyle:`: creates the table view with a plain or grouped format.
- The view controller is automatically made the table view's delegate and data source, unless you specify otherwise.
- The table view is made the view controller's `tableView`. It is also, of course, the view controller's `view`, but the `tableView` property is typed as a UITableView, so you can send table view messages to it without typecasting.



A UITableViewController does *not* obey the UIViewController rule about what happens if it is sent `initWithNibName:bundle:` with a nil nib name. A UITableViewController will *not* find its nib automatically by comparison with its own class name ([Chapter 19](#)); instead, a nil nib name will cause the view controller to generate its table view from scratch. I regard this as a bug. If it's not a bug, it's certainly annoying — more than annoying, it feels like a trap lying in wait for the unsuspecting programmer. I wish I had a nickel for every time I've fallen into it. If I'm loading a UITableViewController's table from a nib, I like to give the table view in the nib an unusual background color during development, just to prove to myself that the table I'm seeing when the app runs really was loaded from the nib; I suggest you do the same!

Table View Cells

Beginners may be surprised to learn that a table view's structure and contents are not configured in advance. Rather, you supply the table view with a data source and a delegate (which will often be the same object; see [Chapter 11](#)), and the table view turns to these in real time, as the app runs, whenever it needs a piece of information about its structure and contents.

This architecture is actually part of a brilliant strategy to conserve resources. Imagine a long table consisting of thousands of rows. It must appear, therefore, to consist of thousands of cells as the user scrolls. But a cell is a UIView and is memory-intensive; to maintain thousands of cells internally would put a terrible strain on memory. Therefore, the table typically maintains only as many cells as are showing simultaneously at any one moment (about ten, let's say). As the user scrolls, the table grabs a cell that is no longer showing on the screen and is therefore no longer needed, and hands it back to you and asks you to configure it as the cell that is about to be scrolled into view. Cells are thus *reused* to minimize the number of actual cells in existence at any one moment.

Therefore your code must be prepared, on demand, to supply the table with pieces of requested data. Of these, the most important is the table cell to be slotted into a given position. A position in the table is specified by means of an index path (NSIndexPath), a class used here to combine a section number with a row number, and is often referred to simply as a *row* of the table. Your data source object may at any moment be sent the message `tableView:cellForRowAtIndexPath:`, and must respond by returning the UITableViewCell to be displayed at that row of the table. And you must return it *fast*: the user is scrolling *now*, so the table needs the next cell *now*.

In this section, then, I'll discuss *what* you're going to be supplying — the table view cell. In the next section, I'll talk about *how* you supply it.



In iOS 5, a table view whose cell contents are known beforehand, such as the one shown in [Figure 21-2](#), can in fact be configured in advance, by designing the table's view controller in a storyboard. I'll discuss how to do that later in this chapter. (This is the unique and valuable use of a storyboard to which I referred in [Chapter 19](#); for reasons known only to Apple, this feature doesn't appear in the nib editor.)

Built-In Cell Styles

To create a cell using one of the built-in cell styles, call `initWithStyle:reuseIdentifier:`. The `reuseIdentifier` is what allows cells previously assigned to rows that are now longer showing to be reused for cells that are; it will usually be the same for all cells in a table. Your choices of cell style are:

`UITableViewCellStyleDefault`

The cell has a `UILabel` (its `textLabel`), with an optional `UIImageView` (its `imageView`) at the left. If there is no image, the label occupies the entire width of the cell.

`UITableViewCellStyleValue1`

The cell has two `UILabels` (its `textLabel` and its `detailTextLabel`), side by side, with an optional `UIImageView` (its `imageView`) at the left. The first label is left-aligned; the second label is right-aligned. If the first label's text is too long, the second label won't appear.

`UITableViewCellStyleValue2`

The cell has two `UILabels` (its `textLabel` and its `detailTextLabel`), side by side; the first label is small. No `UIImageView` will appear. The label sizes are fixed, and the text of either will be truncated if it's too long.

`UITableViewCellStyleSubtitle`

The cell has two `UILabels` (its `textLabel` and its `detailTextLabel`), one above the other, with an optional `UIImageView` (its `imageView`) at the left.

To experiment with the built-in cell styles, do this.

1. Make a new iPhone project from the Empty Application project template.
2. Choose File → New → New File and ask for a Cocoa Touch `UIViewController` subclass.
3. Call the subclass `RootViewController` and *use the combo box to specify that it is to be `UITableViewController` subclass*. (`UITableViewController` is a `UIViewController` subclass, to be discussed later in this chapter; its view is a `UITableView`, which will thus occupy our app's interface.) “With XIB for user interface” should be checked, so that Xcode will hook up the table view in the nib correctly.
4. Click Next and create the files.

To get our table view into the interface, import "RootViewController.h" into *AppDelegate.m*, and add this line to *AppDelegate*'s `application:didFinishLaunchingWithOptions:` (recall that we must specify the nib name explicitly):

```
self.window.rootViewController =
    [[RootViewController alloc] initWithNibName:@"RootViewController" bundle:nil];
```

Now modify the *RootViewController* class (which comes with a lot of templated code), as in [Example 21-1](#).

Example 21-1. The world's simplest table

```
- (NSInteger)numberOfSectionsInTableView:(UITableView *)tableView {
    return 1; ①
}

- (NSInteger)tableView:(UITableView *)tableView
    numberOfRowsInSection:(NSInteger)section {
    return 20; ②
}

- (UITableViewCell *)tableView:(UITableView *)tableView
    cellForRowAtIndexPath:(NSIndexPath *)indexPath {
    static NSString *CellIdentifier = @"Cell";
    UITableViewCell *cell =
        [tableView dequeueReusableCellWithIdentifier:CellIdentifier];
    if (cell == nil) {
        cell = [[UITableViewCell alloc]
            initWithStyle:UITableViewCellStyleDefault ③
            reuseIdentifier:CellIdentifier];
    } ④
    cell.textLabel.text = @"Howdy there"; ⑤
    return cell;
}
```

The idea is to start by generating a single cell in a built-in cell style and then to examine and experiment with its appearance by tweaking the code and running the app. The key parts of the code are:

- ① This code is unchanged from the template; our table will have one section.
- ② Our table will consist of 20 rows. We're going to make our cell without regard to what row it is slotted into; so all 20 rows will be identical. But having multiple rows will give us a sense of how our cell looks when placed next to other cells.
- ③ This is where you specify the built-in table cell style you want to experiment with. Change `UITableViewCellStyleDefault` to a different style as desired.
- ④ At this point in the code you can modify characteristics of the cell (`cell`) that are to be the same for *every* cell of the table. I'll give an example later.
- ⑤ We now have the cell to be used for *this* row of the table, so at this point in the code you can modify characteristics of the cell (`cell`) that are unique to this row. Of

course, that isn't what I've done in the example code; as I just said, all the cells will be identical for now. But that's just because we're only beginners. In real life they'd obviously be likely to have different text, so this is where we put the code that sets the text.

Build and run the app. Behold your table.

Now you can start experimenting, by changing the code at the point where we're now setting `cell.textLabel.text`, and by trying different cell styles other than `UITableViewCellStyleDefault`.

The flexibility of the built-in styles is based mostly on the flexibility of `UILabels` (see also [Chapter 23](#)). Not everything can be customized, because after you return the cell some further configuration takes place, which may override your settings. For example, the size and position of the cell's subviews are not up to you. (I'll explain how to get around that in the next section.) But you get a remarkable degree of freedom. Here are some of the `UILabel` properties you can try changing:

`text`

The string shown in the label.

`textColor`, `highlightedTextColor`

The color of the text. The `highlightedTextColor` applies when the cell is selected (tap on a cell to select it); if you don't set it, the label may choose its own variant of the `textColor` when the cell is highlighted.

`textAlignment`

How the text is aligned; your choices are `UITextAlignmentLeft`, `UITextAlignmentCenter`, and `UITextAlignmentRight`.

`numberOfLines`

The maximum number of lines of text to appear in the label. Text that is long but permitted to wrap, or that contains explicit linefeed characters, can appear completely in the label if the label is tall enough and the number of permitted lines is sufficient. `0` means there's no maximum.

`lineBreakMode`

The wrapping rule for text that is too long for the label's width. The default is `UILineBreakModeTailTruncation`, which means that text wraps at word ends and then, if the last permitted line is still too long for the label, an ellipsis mark is its last visible character.

`font`

The label's font. You could reduce the font size as a way of fitting more text into the label. A font name includes its style, so you could change a bold font to nonbold, or *vice versa*. For example:

```
cell.textLabel.font = [UIFont fontWithName:@"Helvetica-Bold" size:12.0];
```

`minimumFontSize`, `adjustsFontSizeToFitWidth`

If the `numberOfLines` is 1, setting these will allow the font size to shrink automatically in an attempt to fit the entire text into the label's width.

`shadowColor`, `shadowOffset`

The text shadow. Adding a little shadow can increase clarity and emphasis for large text.

The image view's frame can't be changed, but you can inset its apparent size by supplying a smaller image and setting the image view's `contentMode` to `UIViewContentModeCenter`. It's probably a good idea in any case, for performance reasons, to supply images at their drawn size and resolution rather than making the drawing system scale them for you (see the last section of [Chapter 20](#)). For example:

```
UIImage* im = [UIImage imageNamed:@"pic.png"];
UIGraphicsBeginImageContextWithOptions(CGSizeMake(36,36), YES, 0);
[im drawInRect:CGRectMake(0,0,36,36)];
UIImage* im2 = UIGraphicsGetImageFromCurrentImageContext();
UIGraphicsEndImageContext();
cell.imageView.image = im2;
cell.imageView.contentMode = UIViewContentModeCenter;
```

The cell itself also has some properties you can play with:

`accessoryType`

A built-in type of accessory view, which appears at the cell's right end. For example:

```
cell.accessoryType = UITableViewCellStyleDisclosureIndicator;
```

`accessoryView`

Your own `UIView`, which appears at the cell's right end (overriding the `accessoryType`). For example:

```
UIButton* b = [UIButton buttonWithType:UIButtonTypeRoundedRect];
[b setTitle:@"Tap Me" forState:UIControlStateNormal];
[b sizeToFit];
// ... also assign button a target and action ...
cell.accessoryView = b;
```

`indentationLevel`, `indentationWidth`

These properties give the cell a left margin, useful for suggesting a hierarchy among cells. You can also set a cell's indentation level in real time, with respect to the table row into which it is slotted, by implementing the delegate's `tableView:indentationLevelForRowAtIndexPath:` method.

`selectionStyle`

How the background looks when the cell is selected. The default is a blue gradient (`UITableViewCellStyleBlueBackground`), or you can choose `UITableViewCellStyleGrayBackground` (gray gradient) or `UITableViewCellStyleNone`.

```
backgroundColor  
backgroundView  
selectedBackgroundView
```

What's behind everything else drawn in the cell. The `selectedBackgroundView` is drawn in front of the `backgroundView` (if any) when the cell is selected, and will appear instead of whatever the `selectionStyle` dictates. The `backgroundColor` is behind the `backgroundView`. (Thus, if both the `selectedBackgroundView` and the `backgroundView` have some transparency, both of them and the `backgroundColor` can appear composited together when the cell is selected.)

`multipleSelectionBackgroundView`

If defined (not nil), and if the table's `allowsMultipleSelection` (or, if editing, `allowsMultipleSelectionDuringEditing`) is YES, used instead of the `selectedBackgroundView` when the cell is selected.

Applying a `backgroundView` or a `backgroundColor` can be tricky, because:

- The cell's default interface elements, such as the `textLabel`, automatically adopt the cell's background color as their own background color when the cell is not selected. Thus, they will appear to "punch a hole" through the `backgroundView`, revealing the background color behind it. (This problem doesn't arise for a selected cell, because when the cell is selected the cell's interface elements automatically switch to a transparent background, allowing the `selectionStyle` or `selectedBackgroundView` to show through.) The solution, if you want the `backgroundView` to appear behind the interface elements, is to set the `backgroundColor` of the interface elements to a color with some transparency, possibly `[UIColor clearColor]`.
- The cells automatically take on the same `backgroundColor` as the table, and getting them to stop doing this is not easy. The problem is that `tableView:cellForRowAtIndexPath:` is *too soon*; when you set the `backgroundColor` here, your command is obeyed, but then the background color reverts to the table's background color as the cell's own `setSelected:animated:` is called automatically and the cell does various things to its own appearance. The solution is to implement a delegate method, `tableView:willDisplayCell:forRowAtIndexPath:`, and set the `backgroundColor` there.

In this example, we set the `backgroundView` to display an image with some transparency at the outside edges, so that the `backgroundColor` shows behind it. We set the `selectedBackgroundView` to an almost transparent dark rectangle, to darken that image when the cell is selected. And we give the `textLabel` a clear background color so that the rest of our work shows through:

```
UIImageView* v = [[UIImageView alloc] initWithFrame:cell.bounds];  
v.contentMode = UIViewContentModeScaleToFill;  
v.image = [UIImage imageNamed:@"linen.png"];  
cell.backgroundView = v;
```

```
UIView* v2 = [[UIView alloc] initWithFrame:cell.bounds];
v2.backgroundColor = [UIColor colorWithRed:0.2 green:0.2 blue:0.2 alpha:0.1];
cell.selectedBackgroundView = v2;
cell.textLabel.backgroundColor = [UIColor clearColor];
```

I'd put that code in spot numbered 4 in [Example 21-1](#). These features are to be true of every cell ever displayed in the table, and they need to be configured just once for every cell as it first comes into existence. There's no need to waste time doing the same thing all over again when an existing cell is reused.

Finally, there are a few properties of the table view itself worth playing with:

`rowHeight`

The height of a cell. This is another way to deal with text that is too long; besides decreasing the font size, you can increase the cell size. You can also change this value in the nib file; the table view's row height appears in the Size inspector. The cell's subviews have their autoresizing set so as to compensate correctly. You can also set a cell's height in real time by implementing the delegate's `tableView:heightForRowAtIndexPath:` method; thus a table's cells may differ from one another in height (more about that later in this chapter).

`separatorColor`, `separatorStyle`

These can also be set in the nib. The choices of separator style are:

- `UITableViewCellSeparatorStyleNone` (plain style table only)
 - `UITableViewCellSeparatorStyleSingleLine`
 - `UITableViewCellSeparatorStyleSingleLineEtched` (grouped style table only)
- Oddly, the separator style names are associated with `UITableViewCell` even though the separator style itself is a `UITableView` property.

`backgroundColor`, `backgroundView`

What's behind all the cells of the table; this may be seen if the cells have transparency, or if the user bounces the cells beyond their limit. I've already mentioned that cells will, by default, take on the `backgroundColor` of their table. The `backgroundView` is drawn on top of the `backgroundColor`.

`tableHeaderView`, `tableFooterView`

Views to be shown before the first row and after the last row, respectively. Their background color is, by default, the background color of the table, but you can change that. The user can, if you like, interact with the views (and their subviews). You can keep a reference to these views and alter them dynamically during the lifetime of the app. If you change the height of one of these views dynamically, you must set the corresponding table view property afresh so that the table view learns what has happened.

A `UITableViewCell` is a `UIView` and has a `CALayer`, of course, so you can use all the view and layer properties to customize it still further. Here's a further example, showing heavy customization of a table view cell's shape and background. It's a `UITableView`-



Figure 21-3. A cell with a custom gradient background

`CellStyleDefault` cell. I shrink the image and center it in the image view, using the code shown earlier, and I set `textLabel.text`; here, I also supply a gray gradient as the cell's background, and put a round-rect border around the cell (Figure 21-3). Again, I'd put this code in the spot numbered 4 in [Example 21-1](#):

```
UIView* v = [[UIView alloc] initWithFrame:cell.frame];
v.backgroundColor = [UIColor blackColor];
CAGradientLayer* lay = [CAGradientLayer layer];
lay.colors = [NSArray arrayWithObjects:
    (id)[UIColor colorWithWhite:0.6 alpha:1].CGColor,
    [UIColor colorWithWhite:0.4 alpha:1].CGColor, nil];
lay.frame = v.layer.bounds;
[v.layer addSublayer:lay];
lay.borderWidth = 1;
lay.borderColor = [UIColor blackColor].CGColor;
lay.cornerRadius = 5;
cell.backgroundView = v;
cell.textLabel.font = [UIFont fontWithName:@"Helvetica-Bold" size:16];
cell.textLabel.lineBreakMode = UILineBreakModeWordWrap;
cell.textLabel.numberOfLines = 2;
cell.textLabel.textColor = [UIColor whiteColor];
cell.textLabel.backgroundColor = [UIColor clearColor];
```

Custom Cells

The built-in cell styles give the beginner a leg up in getting started with table views, but there is nothing sacred about them, and sooner or later you'll probably want to go beyond them and put yourself in charge of how a table's cells look and what subviews they contain. There are four possible approaches:

- Supply a `UITableViewCell` subclass and override `layoutSubviews` to alter the frames of the built-in subviews.
- In `tableView:cellForRowAtIndexPath:`, add subviews to each cell's `contentView` as the cell is created. The `contentView` is the superview for the cell's subviews, exclusive of things like the `accessoryView`; so by confining yourself to the `contentView`, you allow the cell to continue working correctly. This approach can be combined with the previous one, or you can ignore the built-in subviews and use your own exclusively. As long as the built-in subviews for a particular built-in cell style are not referenced, they are never created or inserted into the cell.
- Design the cell in a nib, and load that nib in `tableView:cellForRowAtIndexPath:` each time a cell needs to be created.
- Design the cell in a storyboard.

I'll illustrate each approach.

Overriding a cell's subview layout

You can't directly change the frame of a built-in cell style subview in `tableView:cellForRowAtIndexPath:` or `tableView:willDisplayCell:forRowAtIndexPath:`, because after your changes, the cell's `layoutSubviews` comes along and overrides them. The work-around is to override the cell's `layoutSubviews!` This is a straightforward solution if your main objection to a built-in style is the frame of a subview.

To illustrate, let's modify a `UITableViewCellStyleDefault` cell so that the image is at the right end instead of the left end. We'll make a `UITableViewCell` subclass, `MyCell`; Xcode makes this easy by supplying `UITableViewCell` as one of the default classes in the "Subclass of" pop-up menu when you make a new Cocoa Touch Objective-C Class file. Here is `MyCell`'s `layoutSubviews`:

```
- (void) layoutSubviews {
    [super layoutSubviews];
    CGRect cvb = self.contentView.bounds;
    CGRect imf = self.imageView.frame;
    imf.origin.x = cvb.size.width - imf.size.width;
    self.imageView.frame = imf;
    CGRect tf = self.textLabel.frame;
    tf.origin.x = 5;
    self.textLabel.frame = tf;
}
```

Now, in our table view's data source, we change the line where new cells are created, so as to generate an instance of `MyCell` rather than `UITableViewCell` (you'll need to import "`MyCell.h`" so as to be able to speak of `MyCell`):

```
cell = [[MyCell alloc]
        initWithStyle:UITableViewCellStyleDefault
        reuseIdentifier:CellIdentifier];
```

Adding subviews in code

Let's rewrite the previous example so that we don't need our own `UITableViewCell` subclass. Instead of modifying the existing `imageView` and `textLabel`, we'll add to each `UITableViewCell`'s content view a completely new `UIImageView` and `UILabel`, each of which can be assigned a frame that won't be changed by the runtime. Here are some things to keep in mind:

- The new views must be added when we instantiate a new cell, but not when we reuse a cell (because a reused cell already has them).
- We must never send `addSubview:` to the cell itself — only to its `contentView` (or some subview thereof).
- We should assign the new views an appropriate `autoresizingMask`, because the cell's content view might be resized.

- Each new view should be assigned a tag so that it can be referred to elsewhere.

Our implementation of `tableView:cellForRowAtIndexPath:` might look like this:

```

- (UITableViewCell *)tableView:(UITableView *)tableView
    cellForRowAtIndexPath:(NSIndexPath *)indexPath {
    static NSString *CellIdentifier = @"Cell";
    UITableViewCell *cell =
        [tableView dequeueReusableCellWithIdentifier:CellIdentifier];
    if (cell == nil) {
        cell = [[UITableViewCell alloc] initWithStyle:UITableViewCellStyleDefault
                                       reuseIdentifier:CellIdentifier];
        CGFloat side = cell.contentView.bounds.size.height;
        UIImageView* iv = [[UIImageView alloc] init];
        iv.frame =
            CGRectMake(cell.contentView.bounds.size.width - side, 0, side, side);
        iv.tag = 1;
        iv.autoresizingMask = (UIViewAutoresizingFlexibleHeight |
                               UIViewAutoresizingFlexibleLeftMargin);
        [cell.contentView addSubview:iv];

        UILabel* lab = [[UILabel alloc] init];
        lab.frame =
            CGRectMake(5, 0, cell.contentView.bounds.size.width - side - 10, side);
        lab.tag = 2;
        lab.autoresizingMask = (UIViewAutoresizingFlexibleHeight |
                               UIViewAutoresizingFlexibleRightMargin);
        [cell.contentView addSubview:lab];
    }
    UILabel* lab = (UILabel*)[cell viewWithTag: 2];
    // ... set up lab here ...
    UIImageView* iv = (UIImageView*)[cell viewWithTag: 1];
    // ... set up iv here ...
    return cell;
}

```

Note how we can refer to the label and the image view, even when we're handed an existing cell for reuse, because we had the foresight to give them tags.

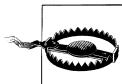
Using our own cell subviews instead of the built-in cell style subviews has some clear advantages; we no longer have to perform an elaborate dance to escape from the restrictions imposed by the runtime. Still, the verbosity of this code is somewhat overwhelming. We can avoid this by designing the cell in a nib.

Designing a cell in a nib

In designing a cell in a nib, we start by creating a nib file that will consist, in effect, solely of this one cell. In Xcode, we create a new iOS User Interface View nib file. Let's call it *MyCell.xib*. In the nib editor, delete the existing View and replace it with a Table View Cell from the Object library.

The cell's design window shows a standard-sized cell; you can resize it as desired, but bear in mind that the final size of the cell in the interface will be dictated by the table

view's `rowHeight`. The cell's style can be specified in the Style pop-up menu of the Attributes inspector, and this gives you the default subviews, locked in their standard positions; for example, if you choose Basic, the `textLabel` appears, and if you specify an image in the Image combo box, the `imageView` appears.



The Style pop-up menu is new in iOS 4.2, and using it in the nib editor may cause the compiler to complain. The solution is to switch, in the nib editor, to the File inspector and change the Development pop-up menu (under Interface Builder Document) to "Xcode 4.2." The fact that this is necessary feels like a bug.

For purposes of the example, however, set the Style pop-up menu to Custom and let's start with a blank slate. We'll implement, from scratch, the same subviews we've already implemented in the preceding two examples: a `UILabel` on the left side of the cell, and a `UIImageView` on the right side. Just as when we add subviews in code, we should set each subview's autoresizing behavior and give each subview a tag. The difference is that we now do both those tasks in the nib, not in code. (If you're following along hands-on, go ahead and do that now.)

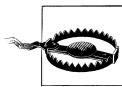
Now comes the most important step. Our goal is to load this nib and grab the cell. To do so, we need an existing instance to function as the File's Owner for this nib, and its class must have an outlet pointing to the cell ([Chapter 7](#)). I'll assume that we've still got a `UITableViewController` subclass called `RootViewController`:

1. Give `RootViewController` a `UITableViewCell` property outlet (and don't forget to `@synthesize` this property):

```
@property (nonatomic, strong) IBOutlet UITableViewCell* tvc;
```

2. In the nib, set the File's Owner class to `RootViewController`, and hook its `tvc` outlet to the cell. (Don't use its `view` outlet by mistake! We don't want to accidentally repoint this view controller's view to the cell.)
3. Assign to the cell in the nib, in the Attributes inspector, the same string Identifier that is used as a cell identifier in our code. We have not altered the template, so that value is `Cell`, in accordance with this line of code:

```
static NSString* CellIdentifier = @"Cell";
```



It is *unbelievably important* to get that last step right. If the cell's Identifier in the nib doesn't match the identifier string in code, you might not get an error, but cells will not be reused properly in the construction of the table. Omitting this step or performing it incorrectly is a common beginner mistake.

We are now ready to modify our implementation of `tableView:cellForRowIndexPath:`. Each time a new cell is needed, at the spot where we previously instantiated a cell in code, we instead load the nib with `self` as owner. To do this, you should use

the UINib method `instantiateWithOwner:options:`, as this will provide some extra speed and efficiency (possibly by caching the nib behind the scenes). As the nib loads, a new instance of the cell is assigned to the `tvc` instance variable ([Chapter 7](#)). We assign that to `cell`, which (under ARC) retains it; we then nilify the instance variable, to release it:

```
UINib* theCellNib = [UINib nibWithNibName:@"MyCell" bundle:nil];
[theCellNib instantiateWithOwner:self options:nil];
cell = self.tvc;
self.tvc = nil;
```

This gives us all the advantages of customized cell contents without any verbosity; the cell is created in two easy lines of code, and we proceed to configure it. Moreover, our configuration code (“set up `lab` here”) can perhaps be greatly reduced, because this subview can be configured in the nib. For example, as I was evolving this example from earlier examples, I had this configuration code:

```
lab.numberOfLines = 2;
lab.textAlignment = NSTextAlignmentLeft;
lab.adjustsFontSizeToFitWidth = YES;
lab.lineBreakMode = UILineBreakModeClip;
lab.minimumFontSize = 0.1;
lab.textColor = [UIColor whiteColor];
lab.highlightedTextColor = [UIColor blackColor];
lab.font = [UIFont fontWithName:@"Helvetica-Bold" size:12.0];
lab.text = @"This is a test";
```

But all of that code can now be deleted! Those are all aspects of the `UILabel` that can be configured in the nib instead. Moreover, a table view cell in a nib has outlets for its `accessoryView` (and `editingAccessoryView`) and its `backgroundView` and `selectedBackgroundView` (but not, oddly, its `multipleSelectionBackgroundView`; one really does wonder whether the various departments at Apple ever talk to one another). Thus we can set all of these features directly in the nib. For example, to set the background view to the custom gradient we developed earlier, I’d create a `UIImageView` subclass that draws the gradient and makes the result its own `image`, put an instance of this `UIImageView` subclass into the cell’s nib, and hook the cell’s `backgroundView` outlet to that `UIImageView` instance. For these and similar reasons, I am very much partial to the nib-loading approach to cell instantiation.

Moreover, starting in iOS 5, there’s yet another reason to instantiate your cell’s from a nib — a new way of loading the cells, as needed, *automatically*, so that you don’t need an outlet to the cell. Up to now, each time after the call to `dequeueReusableCellWithIdentifier:`, we have been checking to see whether the result is nil, and instantiating a cell manually if so:

```
UITableViewCell *cell =
    [tableView dequeueReusableCellWithIdentifier:CellIdentifier];
if (cell == nil) {
    // ... instantiate a cell, in code or by loading a nib
```

In iOS 5, if we've got a cell in a nib, we can eliminate that code. Here's how: before the table view needs any cells, call its method `registerNib:forCellReuseIdentifier:`. This effectively pairs the reuse identifier string with a nib name. From then on, when you call `dequeueReusableCellWithIdentifier:`, *a cell is always returned*. If there's a reusable cell available, it is returned as before; if there's no reusable cell available, the nib corresponding to the given identifier is loaded and the cell is returned.

The advantages of this approach are:

- We no longer have to specify the reuse identifier in the nib, because we've performed the pairing in code. (In fact, it is best to leave the Identifier in the nib blank; a wrong Identifier value will break things, but a blank Identifier won't.)
- There is no need to specify a File's Owner class in the nib (in fact, it should be left at `NSObject`), and there is no outlet pointing from the File's Owner to the cell; the runtime simply looks for a single top-level nib object. Similarly, there is no need for the class running all this code to have a property corresponding to the outlet; there is no outlet.
- There is no need for the test (`if cell == nil`); there will always be a cell. The entire section of the cell-retrieving code marked as step 4 in [Example 21-1](#) is eliminated. If you have configurations to perform on each cell just once, when the cell is freshly instantiated, make a new test examining something about its state that differentiates a freshly instantiated, unconfigured cell from a configured one.

Here's how our implementation of `tableView:cellForRowAtIndexPath:` is structured under this new regime:

```
- (UITableViewCell *)tableView:(UITableView *)tableView
    cellForRowAtIndexPath:(NSIndexPath *)indexPath {
    static NSString *CellIdentifier = @"Cell";
    [self.tableView registerNib:
        [UINib nibWithNibName:@"MyCell" bundle:nil]
        forCellReuseIdentifier:CellIdentifier];
    UITableViewCell *cell =
        [tableView dequeueReusableCellWithIdentifier:CellIdentifier];
    // that's it! we've got a cell
    UILabel* lab = (UILabel*)[cell viewWithTag: 2];
    // ... set up lab here ...
    UIImageView* iv = (UIImageView*)[cell viewWithTag: 1];
    // ... set up iv here ...
    return cell;
}
```

In reality we only need to call `registerNib:forCellReuseIdentifier:` once. There's no harm in calling it repeatedly, however; all we're doing is pairing two string names, the nib and the identifier, which takes no time at all. If the repeated call offends you, promote the `CellIdentifier` to an external global variable, so that code throughout this class can see it, and make the call once, earlier in the table view's life (such as the view controller's `viewDidLoad`).

In just one aspect, Apple seems to have designed this approach in a half-baked way: in a nib to be loaded automatically after calling `registerNib:forCellReuseIdentifier:`, there must be *just one top-level nib object*, the `UITableViewCell`. This seems an unnecessary restriction; why can't the nib-loader examine the top-level objects and discover the one that's a `UITableViewCell`? The unfortunate result is that, say, an object to be pointed to by the cell's `backgroundView` outlet can't be included in the nib as *another* top-level nib object. The workaround is to make such an object a subview of the cell; you might need to use the Size inspector to move the object out of the cell's bounds so that you can continue working in the cell (and the cell will have to use the Custom style, as you can't add any subviews to a built-in cell style in a nib).

We are still referring to the cell's subviews in code by way of `viewWithTag:`. If you would prefer to use names, simply provide a `UITableViewCell` subclass with outlet properties, and configure the nib file accordingly:

1. Create the files for a `UITableViewCell` subclass; let's call it `MyCell`. (If you still have the `MyCell` files from the earlier example, you can use them, but delete the `layoutSubviews` override.) Give the class two outlet properties (and remember to `@synthesize` them):

```
@property (nonatomic, strong) IBOutlet UILabel* theLabel;
@property (nonatomic, strong) IBOutlet UIImageView* theImageView;
```

2. In the table view cell nib, change the class of the cell to `MyCell`, and link up the outlets from the cell to the respective subviews.

In our implementation of `tableView:cellForRowAtIndexPath:`, once we've typed the cell as a `MyCell` (and imported "`MyCell.h`"), the compiler will let us use the property names to access the subviews:

```
MyCell* theCell = (MyCell*)cell;
UILabel* lab = theCell.theLabel;
// ... set up lab here ...
UIImageView* iv = theCell.theImageView;
// ... set up iv here ...
```

Designing a cell in a storyboard

In iOS 5, if we're using a `UITableViewController` subclass, its table view's cells can be designed in a storyboard. In the storyboard editor, the `UITableViewController` comes with a table view. In the Attributes inspector, you set the table view's Content pop-up menu to Dynamic Prototypes, and use the Prototype Cells field to say how many different cell types there are to be — that is, how many different cell identifiers your table view controller's code will be using. In our case (and in most cases) this is 1. The table view in the storyboard editor is shown with as many table view cells as the Prototype Cells field dictates. Again, in our case that means there's one table view cell. This effectively corresponds to the table view cell in the nib file from which we were previously instantiating it.

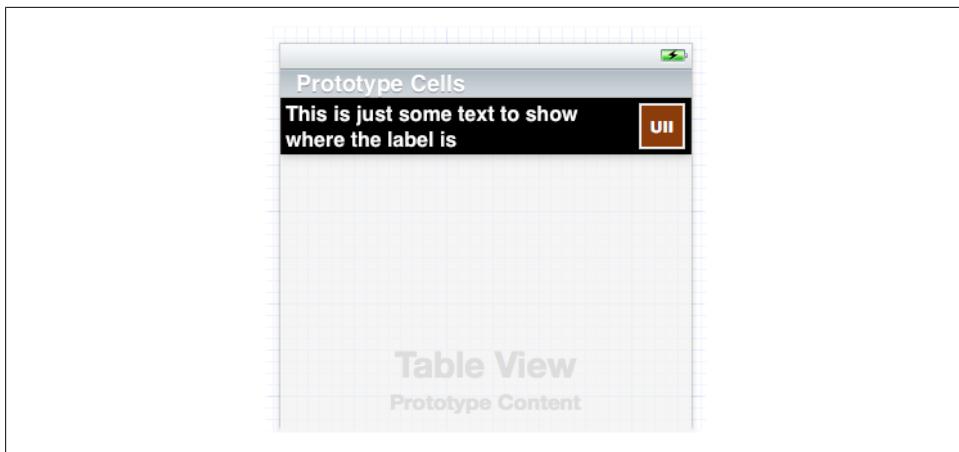


Figure 21-4. A table view with a cell prototype in a storyboard

Now you can design the table view cell graphically; for example, once again we could have a label on the left and an image view on the right (Figure 21-4). If you want to add views to be the target of any of the cell’s outlets (such as its `backgroundView`), you must make them subviews of the cell, as we did a moment ago.

I’ve already discussed (Chapter 19) the ways in which a view controller might be instantiated from a storyboard. This storyboard might be your app’s main storyboard, and the table view controller might be instantiated automatically. On the other hand, you might take advantage of a storyboard solely in order to use this feature (the ability to design a nib in a storyboard), even if your app has no main storyboard, in which case you would explicitly instantiate the table view controller from the storyboard in code.

The reason why you might use a storyboard just to design a table view cell is that this approach requires *even less code* than what we’ve been doing so far. In `tableView:cellForRowAtIndexPath:`, we still need to call `dequeueReusableCellWithIdentifier:`, but that’s *all* we need to do. If there is no reusable cell, there is no need for us to instantiate a cell in code, and there is no need to load a nib; the cell will be instantiated from the storyboard, which will be found automatically because the view controller keeps a pointer to it. There is no need to call `registerNib:forCellReuseIdentifier:`, because the runtime knows where to find the cell prototype—in the storyboard. The only thing you have to do is give the cell the correct Identifier value in its Attributes inspector in the storyboard; the rest happens automatically:

```
UITableViewCell *cell = [tableView dequeueReusableCellWithIdentifier:CellIdentifier];
```

If a new cell is needed, the runtime looks in the storyboard to find the cell prototype for this table view whose Identifier value matches `CellIdentifier`. It instantiates that cell and returns it. End of story!

As in the previous example, the cell in the storyboard can be an instance of a subclass of `UITableViewCell`. If our subclass and the outlets in the storyboard are set up as

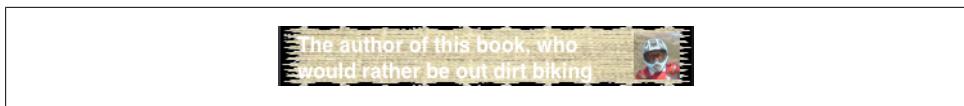


Figure 21-5. A table view cell instantiated from a storyboard

before, we can refer to the cell's custom subviews through property names. Here is the *entire* code used to generate Figure 21-5 (everything else is configured in the storyboard):

```
- (UITableViewCell *)tableView:(UITableView *)tableView
    cellForRowAtIndexPath:(NSIndexPath *)indexPath {
    static NSString *CellIdentifier = @"Cell";
    MyCell *cell =
        (MyCell*)[tableView dequeueReusableCellWithIdentifier:CellIdentifier];
    cell.theLabel.text =
        @"The author of this book, who would rather be out dirt biking";
    UIImage* im = [UIImage imageNamed:@"moi.png"];
    UIGraphicsBeginImageContextWithOptions(CGSizeMake(36,36), YES, 0.0);
    [im drawInRect:CGRectMake(0,0,36,36)];
    UIImage* im2 = UIGraphicsGetImageFromCurrentImageContext();
    UIGraphicsEndImageContext();
    cell.theImageView.image = im2;
    return cell;
}
```

Table View Data

The structure and content of the actual data portrayed in a table view comes from the data source, an object pointed to by the table view's `dataSource` property and adopting the `UITableViewDataSource` protocol. The data source is thus the heart and soul of the table. What surprises beginners is that the data source operates not by *setting* the table view's structure and content, but by *responding on demand*. The data source, *qua* data source, consists of a set of methods that the table view will call when it needs information. This architecture has important consequences for how you write your code, which can be summarized by these simple guidelines:

Be ready

Your data source cannot know *when* or *how often* any of these methods will be called, so it must be prepared to answer *any question at any time*.

Be fast

The table view is asking for data in real time; the user is probably scrolling through the table *right now*. So you mustn't gum up the works; you must be ready to supply responses just as fast as you possibly can. (If you can't supply a piece of data fast enough, you may have to skip it, supply a placeholder, and insert the data into the table later. This, however, may involve you in threading issues that I don't want to get into here. I'll give an example in [Chapter 37](#).)

Be consistent

There are multiple data source methods, and you cannot know *which* one will be called at a given moment. So you must make sure your responses are mutually consistent at *any* moment. For example, a common beginner error is forgetting to take into account, in your data source methods, the possibility that the data might not be ready yet.

This may sound daunting, but you'll be fine as long as you maintain an unswerving adherence to the principles of model–view–controller ([Chapter 13](#)). How and when you accumulate the actual data, and how that data is structured, is a *model* concern. Acting as a data source is a *controller* concern. So you can acquire and arrange your data whenever and however you like, just so long as when the table view actually turns to you and asks what to do, you can lay your hands on the relevant data rapidly and consistently. You'll want to design the model in such a way that the controller can access any desired piece of data more or less instantly.

Another source of confusion for beginners is that methods are rather oddly distributed between the data source and the delegate, an object pointed to by the table view's `delegate` property and adopting the `UITableViewDelegate` protocol; in some cases, one may seem to be doing the job of the other. This is not usually a cause of any real difficulty, because the object serving as data source will probably also be the object serving as delegate. Nevertheless, it is rather inconvenient when you're consulting the documentation; you'll probably want to keep the data source and delegate documentation pages open simultaneously as you work.



In iOS 5, if a table view's contents are known beforehand, you can design the entire table, *including the contents of individual cells*, in a storyboard.

This could be a reason for using a storyboard, even if your app has no main storyboard. I'll give an example later in this chapter.

The Three Big Questions

Like Katherine Hepburn in *Pat and Mike*, the basis of your success as a data source is your ability, at any time, to answer the Three Big Questions. The questions the table view will ask you are a little different from the questions Mike asks Pat, but the principle is the same: know the answers, and be able to recite them at any moment. Here they are:

How many sections does this table have?

The table will call `numberOfSectionsInTableView:`; respond with an integer. In theory you can omit this method, as the default response is `1`, which is often correct. However, I never omit it; for one thing, returning `0` is a good way to say that the table has no data, and will prevent the table view from asking any other questions.

How many rows does this section have?

The table will call `tableView:numberOfRowsInSection:`. The table supplies a section number — the first section is numbered 0 — and you respond with an integer. In a table with only one section, of course, there is probably no need to examine the incoming section number.

What cell goes in this row of this section?

The table will call `tableView:cellForRowAtIndexPath:`. The index path is expressed as an `NSIndexPath`; this is a sophisticated and powerful class, but you don't actually have to know anything about it, because `UITableView` provides a category on it that adds two read-only properties — `section` and `row`. Using these, you extract the requested section number and row number, and return a `UITableViewCell`. The first row of a section is numbered 0.

The strategy for implementing `tableView:cellForRowAtIndexPath:` is a little complicated, because you will probably want to keep memory usage at a minimum by reusing cells. The idea, as I've already mentioned, is that once a cell is no longer visible on the screen, it can be slotted into a row that *is* visible — with its portrayed data appropriately modified, of course! — so that no more than the number of simultaneously visible cells need to exist at any given moment. Luckily, a table view is ready to implement this strategy for you, and the template shows you how to write `tableView:cellForRowAtIndexPath:` accordingly. Let's examine the template code more closely ([Example 21-2](#)).

Example 21-2. The template code for the third big question

```
// Customize the appearance of table view cells.  
- (UITableViewCell *)tableView:(UITableView *)tableView  
    cellForRowAtIndexPath:(NSIndexPath *)indexPath {  
    static NSString *CellIdentifier = @"Cell"; ❶  
    UITableViewCell *cell = ❷  
        [tableView dequeueReusableCellWithIdentifier:CellIdentifier];  
    if (cell == nil) {  
        cell = [[UITableViewCell alloc] ❸  
            initWithStyle:UITableViewCellStyleDefault  
            reuseIdentifier:CellIdentifier];  
    } ❹  
    return cell;  
}
```

- ❶ The table view is ready to maintain a cache of reusable cells. In fact, it can maintain more than one such cache; this could be useful if your table view contains more than one type of cell (where the meaning of the concept “type of cell” is pretty much up to you). Therefore you must *name* each cache, by attaching a cell identifier string to any cell that can be reused. In most cases, there will be just one cache and therefore just one cell identifier string, and you will want every cell to be reusable, so you will attach that identifier string to every cell. In the template code, this string is main-

tained as a static NSString literal, ensuring that the memory for it will be set aside only once and maintained between calls to this method.

- ❷ To generate the cell being requested by the method call, you start by sending the table view `dequeueReusableCellWithIdentifier:`. This asks the table view to pull out of the named cache a currently unused cell, if there is one. There are two possibilities: either the table view will return a cell, in which case that's the cell you're going to be returning (with its portrayed data appropriately modified, of course!), or the table view will return nil (because there isn't an unused cell already in the cache), in which case a new cell will need to be instantiated.
- ❸ If `dequeueReusableCellWithIdentifier:` returned nil, one of several things must happen, as I explained in great detail in the preceding section of this chapter. I'll summarize again what they are:
 - If you're instantiating new cells in code, you'll do so, initializing the cell with `initWithStyle:reuseIdentifier:`. The style is one of the built-in styles; we have already seen what this implies, and how you can configure a cell yourself regardless of its style. The reuse identifier is the cell identifier that makes this cell reusable. If you were to supply nil here, the cell would not be reusable.
 - If you're instantiating new cells by loading a nib file, you might now load the nib file with yourself as owner, accessing the cell by means of an outlet.
 - Alternatively, if you've previously called `registerNib:forCellReuseIdentifier:`, or if this is a UITableViewController instantiated from a storyboard, the runtime will automatically load the nib or storyboard file, instantiate the cell, and return the new instance as the result of `dequeueReusableCellWithIdentifier:`. In that case, this part of the code won't exist.
- ❹ If `dequeueReusableCellWithIdentifier:` returned nil, you have just instantiated a new cell; you may now want to perform any other tasks appropriate to the initial configuration of a new cell. For example, earlier we saw that you can add subviews to the cell's content view; you would want to do this only when the cell is first created, so this is the place for it. If you're generating new cells by the third method I just listed, this part of the code won't exist (you won't check for `cell == nil`, because `cell` will *never* be nil); instead, you'll devise some other condition to differentiate newly instantiated unconfigured cells from reused cells that you've already configured.
- ❺ You now have a cell (here, called `cell`), from whatever source derived. No matter whether you just created it or you are reusing a previously used cell, you now configure the cell appropriately to the section and row into which this cell is now to be slotted. *This* is the thing we did *not* do in any of our earlier examples; we were concentrating on cell configuration itself, so we configured the cell the same way regardless of its row. But in real life, this is the most important step: this is the step that differentiates one cell from another, reflecting the model data into the view.

You do not know or care, when configuring the cell at the final stage, whether the cell is new or reused. Therefore, you should always configure *everything* about the cell that might need configuring. If you fail to do this, and if the cell is reused, you might be surprised when some aspect of the cell is left over from its previous use; on the other hand, if you fail to do this, and if the cell is new, you might be surprised when some aspect of the cell isn't configured at all.

I'm speaking from experience here. For example, in one of my apps that lists article titles in a table, there is a little loudspeaker icon that should appear in the cell if there is a recording associated with this article. So I initially wrote this code:

```
if (item.enclosures && [item.enclosures count])
    [cell viewWithTag: 5].hidden = NO;
```

This turned out to be a mistake, because when a cell was reused, it had a visible loudspeaker icon if, in a previous incarnation, it had *ever* had a visible loudspeaker icon. The solution was to rewrite the logic to express all possibilities, like this:

```
[cell viewWithTag: 5].hidden = !(item.enclosures && [item.enclosures count]);
```

To illustrate the efficiency of the cell-caching architecture, I'll use the nib-instantiated cell left over from our earlier examples (so that it will be possible to differentiate between cases when a cell is instantiated and cases when a cell is reused). Presume that we have a table of one section and some large number of rows - let's say there are 100 of them:

```
- (UITableViewCell *)tableView:(UITableView *)tableView
    cellForRowAtIndexPath:(NSIndexPath *)indexPath {
    static NSString *CellIdentifier = @"Cell";
    UITableViewCell *cell =
        [tableView dequeueReusableCellWithIdentifier:CellIdentifier];
    if (cell == nil) {
        [[NSBundle mainBundle] loadNibNamed:@"MyCell" owner:self options:nil];
        cell = self.tvc;
        NSLog(@"creating a new cell");
    }
    MyCell* theCell = (MyCell*)cell;
    UILabel* lab = theCell.theLabel;
    lab.text = [NSString stringWithFormat: @"This is row %i of section %i",
                indexPath.row, indexPath.section];
    return cell;
}
```

When we run this code and scroll through the table, every cell is numbered correctly, so there appear to be 100 cells. But the log messages show us that only 11 cells are actually created.

In real life, of course, you'd probably be setting `lab.text` by consulting a data store (the model) for the value appropriate to this row (or, if there is more than one section, this row of this section). A line from a typical real-life implementation of `tableView:cellForRowAtIndexPath:` consults an array (`titles`):

```
cell.textLabel.text = [titles objectAtIndex: [indexPath row]];
```

It may surprise you that that's all I have to say about how to get the data from your model. But it is! The fact is that I can have no earthly idea what your app does, what its data is about, or what your table is trying to portray. The important thing is to remember that you're going to be receiving an `NSIndexPath` specifying a section and a row, and you need to be lay your hands on the data corresponding to that slot *now* and configure the cell accordingly *now*. So construct your model, and your algorithm for consulting it in this last part of the Third Big Question, accordingly.

You do get a sort of second bite of the cherry: there's a delegate method, `tableView:willDisplayCell:forRowAtIndexPath:`, that is called for every cell just before it appears in the table. This is absolutely the last minute to configure a cell. But don't misuse this method. You're functioning as the delegate here, not the data source; you may set the final details of the cell's appearance — as I discussed earlier, this is a good place to set a cell's background color if you don't want it to come from the table's background color — but you shouldn't be consulting the data model at this point.

Table View Sections

Your table data can be expressed as divided into sections. You might clump your data into sections for various reasons (and doubtless there are other reasons beyond these):

- You want to supply section headers (or footers, or both). This can clarify the presentation of your data by dividing the rows into groups. Also, a section header or footer can be a custom `UIView`, so it's a place where you might put additional information or functional interface (such as a button the user can tap).
- You want to make navigation of the table easier by supplying an index down the right side. You can't have an index without sections.
- You want to facilitate programmatic rearrangement of the table. For example, it's very easy to hide or move an entire section at once, possibly with animation.

The number of sections is determined by your reply to `numberOfSectionsInTableView:`. For each section, the table view will consult your data source and delegate to learn whether this section has a header or a footer, or both, or neither (the default). You can supply headers and footers in two ways:

Header or footer string

You implement the data source method `tableView:titleForHeaderInSection:` or `tableView:titleForFooterInSection:` (or both). Return `nil` to indicate that the given section has no header (or footer). Return a string to use it as the section's header (or footer). You cannot change the style of the label or of the header or footer as a whole.

Header or footer view

You implement the delegate method `tableView:viewForHeaderInSection:` or `tableView:viewForFooterInSection:` (or both). The corresponding `titleFor...` data

source method, if implemented, is ignored. The view you supply is used as the entire header or footer and is automatically resized to the table's width and the section header or footer height. If the view you supply has subviews, be sure to set proper autoresizing behavior so that they'll be positioned and sized appropriately when the view itself is resized.

Supplying a header or a footer as a string is simpler, but supplying it as a view is vastly more powerful, because not only can you style a label, but you can also employ other kinds of interface, possibly interactive (such as a button). Moreover, a view can be customized in any of the ways that views can be customized: for example, you can set the transparency of a header (so that rows are visible as they scroll behind it), or give it a shadow (so that it casts its shadow on the row adjacent to it). Plus, you can keep a reference to a header or footer view and change that view's content later, dynamically, in code, possibly with animation, in the course of the app's lifetime — without reloading the table's data. In other words, once you've handed the table a section header or footer view, the table leaves it alone; it knows all it needs to know, and the view is yours to command.

What the table *doesn't* know, and what you *can't* change arbitrarily during the lifetime of the app, is a header or footer view's height; you can change the height, but you must do it in a controlled manner that the table knows about. Obviously, the table needs to know the height of every header or footer so that it can slot it in among the rows correctly; if you want to change a header or footer's height, you must ask the table to reload its data (discussed in the next section) so that can lay out its content afresh.

No matter which way you supply the header or footer, its height is determined as follows. The default height comes from the table itself, which has a `sectionHeaderHeight` property and a `sectionFooterHeight` property. But you can override this to supply header and footer heights on an individual basis by implementing the delegate method `tableView:heightForHeaderInSection:` or `tableView:heightForFooterInSection:`; if you implement one of these methods, it must always return an actual value, even if the header or footer is a nil view (that requirement is new in iOS 5). As a convenience, if the header or footer is just a title (that is, you implemented `titleFor...` and not `viewFor...`), you can have the table supply the height automatically, just as it would if you hadn't implemented the `heightFor...` method in the first place, by returning `UITableViewAutomaticDimension` (new in iOS 5).



Don't confuse the section headers and footers with the header and footer of the table as a whole. The latter are view properties of the table view itself and are set through its properties `tableHeaderView` and `tableFooterView`, discussed earlier in this chapter.

If the table view has the plain style, you can add an index down the right side of the table, which the user can tap to jump to the start of a section. This is very helpful for navigating long tables. To generate the index, implement the data source method

`sectionIndexTitlesForTableView:`, returning an `NSArray` of string titles to appear as entries in the index. This works even if there are no section headers. The index will appear only if the number of rows exceeds the table view's `sectionIndexMinimumDisplayRowCount` property value; the default is 0 (not `NSIntegerMax` as claimed by the documentation), so the index is always displayed by default. You will want the index entries to be short — preferably just one character — because they will be partially obscuring the right edge of the table; plus, the cell's content view will shrink to match, so you're sacrificing some cell real estate. Unfortunately, there is no official way to modify the index's appearance (such as the color of its entries).

Normally, there will be a one-to-one correspondence between the index entries and the sections; when the user taps an index entry, the table jumps to the start of the corresponding section. However, under certain circumstances you may want to customize this correspondence. For example, suppose there are 40 sections, but there isn't room to display 40 index entries comfortably on the iPhone. The index will automatically curtail itself, omitting some index entries and inserting bullets to suggest this, but you might prefer to take charge of the situation by supplying a shorter index. In such a case, implement the data source method `tableView:sectionForSectionIndexTitleAtIndex:`, returning the index of the section to jump to for this section index. Both the section index title and its index are passed in, so you can use whichever is convenient.

Apple's documentation elaborates heavily on the details of implementing the model behind a table with an index and suggests that you rely on a class called `UILocalizedIndexedCollation`. This class is effectively a way of generating an ordered list of letters of the alphabet, with methods for helping to sort an array of strings and separate it into sections. This might be useful if you need your app to be localized, because the notion of the alphabet and its order changes automatically depending on the user's preferred language. But this notion is also fixed; you can't readily use a `UILocalizedIndexCollation` to implement your own sort order. For example, `UILocalizedIndexCollation` was of no use to me in writing my Greek and Latin vocabulary apps, in which the Greek words must be sorted, sectioned, and indexed according to the Greek alphabet, and the Latin words use a reduced version of the English alphabet (no initial J, K, or V through Z).

I'll demonstrate a technique for implementing a sectioned indexed table by describing the approach I take in my Latin vocabulary app (this is the second app in [Figure 21-1](#)). Everything depends on preparing the data in advance, storing it in appropriate structures. The data, consisting of Latin words and their English definitions, start out in an array (`data`) of `Term` objects; in particular, each Latin word is stored in its `Term` under the `latin` property, which is key-value coding compliant ([Chapter 12](#)). So we begin by sorting the array on that key's value:

```
NSArray* terms =
    [data sortedArrayUsingDescriptors:
        [NSArray arrayWithObject:
            [NSSortDescriptor sortDescriptorWithKey:@"latin" ascending:YES]]];
```

We proceed now to make two arrays:

1. A list of section names (`sectionNames`), consisting of the unique capitalized first letters of the Latin words. The trick here is the word “unique”; as we cycle through the sorted Latin words, we add a letter to this list only if it isn’t the same as the last letter we added.
2. An array of arrays (`sectionData`) of all words starting with each letter.

```
self.sectionNames = [NSMutableArray array];
self.sectionData = [NSMutableArray array];
NSString* previousSection = @"";
for (Term* term in terms) {
    NSString* currentSection = [[term.latin substringToIndex:1] uppercaseString];
    if (![currentSection isEqualToString: previousSection]) {
        previousSection = currentSection;
        [sectionNames addObject: currentSection];
        [sectionData addObject: [NSMutableArray array]];
    }
    // in every case, add this term to the last section array
    [[sectionData lastObject] addObject: term];
}
```

Now that we have our two arrays, the business of actually describing the table view’s structure based on them is trivial; the methods are all one-liners, demonstrating that design and preparation of the data model is everything:

```
- (NSInteger)numberOfSectionsInTableView:(UITableView *)tableView {
    return [sectionNames count];
}

- (NSInteger)tableView:(UITableView *)tableView
    numberOfRowsInSection:(NSInteger)section {
    return [[sectionData objectAtIndex: section] count];
}

- (NSString *)tableView:(UITableView *)tableView
    titleForHeaderInSection:(NSInteger)section {
    return [sectionNames objectAtIndex: section];
}

- (NSArray *)sectionIndexTitlesForTableView:(UITableView *)tableView {
    return sectionNames;
}
```

The implementation of `tableView:cellForRowIndexPath:` requires consideration of both the `section` and the `row` of the incoming `indexPath`, but here too we can lay our hands on the correct `Term` for this table row instantly:

```
Term* term = [[sectionData objectAtIndex: [indexPath section]]
    objectAtIndex: [indexPath row]];
```

Refreshing Table View Data

The table view (view) has no direct connection to the underlying data (model). If you want the table view display to change because the underlying data have changed, you have to cause the table view to refresh itself; basically, you're requesting that the Three Big Questions be asked all over again (controller). At first blush, this seems inefficient ("regenerate *all* the data??"); but it isn't. Remember, in a table that caches reusable cells, there are no cells of interest other than those actually showing in the table at this moment. Thus, having worked out the layout of the table through the section header and footer heights and row heights, the table has to regenerate only those cells that are actually visible.

You can cause the table data to be refreshed using any of several methods:

`reloadData`

The table view will ask the Three Big Questions all over again, including section headers and footers, and index entries.

`reloadRowsAtIndexPaths:withRowAnimation:`

The table view will ask the Three Big Questions all over again, including section headers and footers, but not index entries. The first parameter is an array of index paths; to form an index path, use the `NSIndexPath` class method `indexPathForRow:inSection:`.

`reloadSections:withRowAnimation:`

The table view will ask the Three Big Questions all over again, including section headers and footers, but not index entries. The first parameter is an `NSIndexSet` (see [Chapter 10](#)).

The second two methods can perform animations that cue the user as to what's changing. The `withRowAnimation:` parameter is one of the following:

`UITableViewRowAnimationFade`

The old fades into the new.

`UITableViewRowAnimationRight`

`UITableViewRowAnimationLeft`

`UITableViewRowAnimationTop`

`UITableViewRowAnimationBottom`

The old slides out in the stated direction, and is replaced from the opposite direction.

`UITableViewRowAnimationNone`

No animation.

`UITableViewRowAnimationMiddle`

Hard to describe; it's a sort of venetian blind effect on each cell individually.

UITableViewRowAnimationAutomatic

The table view just “does the right thing” (new in iOS 5). This is especially useful for grouped style tables, because if you pick the wrong animation the display can look very funny as it proceeds.

If all you need to do is to refresh the index, call `reloadSectionIndexTitles`; this calls the data source’s `sectionIndexTitlesForTableView:`. (This is new in iOS 5; the method itself isn’t new, but prior to iOS 5, calling it had no effect.)

It is also possible to access and alter a table’s individual cells directly. This can be a far more lightweight approach to refreshing the table, plus you can supply your own animation within the cell as it alters its appearance. To do this, you need direct access to the cell you want to change. You’ll probably want to make sure the cell is visible within the table view’s bounds; if you’re taking proper advantage of the table’s reusable cell caching mechanism, nonvisible cells don’t really exist (except as potential cells waiting in the reuse cache), and there’s no point changing them, as they’ll be changed when they are scrolled into view, through the usual call to `tableView:cellForRowIndexPath:`. Here are some UITableView methods that mediate between cells, rows, and visibility:

visibleCells

An array of the cells actually showing within the table’s bounds.

indexPathsForVisibleRows

An array of the rows actually showing within the table’s bounds.

cellForRowIndexPath:

Returns a UITableViewCell if the table is maintaining a cell for the given row (typically because this is a visible row); otherwise, returns nil.

indexPathForCell:

Given a cell obtained from the table view, returns the row into which it is slotted.

It is important to bear in mind that the cells are not the data (view is not model). If you change the content of a cell manually, make sure that you have also changed the model corresponding to it, so that the row will appear correctly if its data is reloaded later.

Variable Row Heights

Most tables have rows that are all the same height, as set by the table view’s `rowHeight`. However, the delegate’s `tableView:heightForRowAtIndexPath:` can be used to make different rows different heights. You can see this in the TidBITS News app; look at [Figure 19-1](#), where the first cell is shorter than the second cell (because the headline is one line instead of two).

Here are some things to remember when implementing a table whose rows can have different heights:

Avoid performance limits

Variable row heights work best if the table is short and simple (not too many rows). The table view must effectively lay out the entire table in order to load the data and in order at any moment to know the size and offset of the scrolling content. With a table consisting of a large number of rows, this can become too much information for the table to manipulate fast enough as the user scrolls, and the whole mechanism will break down.

Lay out subviews correctly

As a cell is reused, its height may be changed, because the new row into which it is to be slotted is a different height from the old row. Similarly, if the cell comes from a nib, its height in the table view may be changed from its height in the nib. This will expose any weaknesses in your practice for laying out subviews. For example, a mistake in the `autoresizingMask` value of subviews can result in display errors that would not have been exposed if all the rows were the same height. You may have to resort to manual layout (implementing `layoutSubviews` in a `UITableViewCell` subclass).

Plan ahead

You (the delegate) are going to be asked for all the heights of all the rows well before you (the data source) are called upon to provide the data for any individual rows. You will want to provide this information quickly and accurately. So you will have to plan how the data will appear in *every* row before actually causing the data to appear in *any* row.

For example, here's how I implemented the variable-height rows in the TidBITS News app. The cells in its table view contain two labels. In order to determine the height of a cell, I need to know how much vertical space each of its labels should occupy. For example, the first label, the article headline, permits two lines of text (`numberOfLines`), but if the headline is short it might need only one. So I supply a utility method, `labelHeightsForItem:`, that calls `NSString's sizeWithFont:constrainedToSize:lineBreakMode:` to calculate the heights of both labels given their text and font. Note that this utility method must be able to consult the data model in order to learn what the text will be for each label:

```
- (NSArray*) labelHeightsForRow: (NSInteger) row {
    NSString* title = [self->titles objectAtIndex: row];
    NSString* artist = [self->artists objectAtIndex: row];
    // values used in next two lines have been precached as ivars at load time
    CGSize tsz =
        [title sizeWithFont:self->titleFont
                  constrainedToSize:CGSizeMake(tw, 4000)];
    CGSize asz =
        [artist sizeWithFont:self->artistFont
                  constrainedToSize:CGSizeMake(aw, 4000)];
    NSArray* result = [NSArray arrayWithObjects:
                           [NSNumber numberWithFloat: tsz.height],
```

```

        [NSNumber numberWithFloat:asz.height],
        nil];
    return result;
}

```

My `tableView:heightForRowAtIndexPath:` implementation can then call `labelHeightsForItem:` and use those heights, along with some #defined spacer values, to work out the total height of the cell, while restricting it to the maximum height I'm willing to allow:

```

- (CGFloat)tableView:(UITableView *)tableView
    heightForRowAtIndexPath:(NSIndexPath *)indexPath {
    id item = [self.parsedData.items objectAtIndex:[indexPath row]];
    NSArray* arr = [self labelHeightsForItem: item]; // call our utility method
    // label heights are now in a two-element array of NSNumbers
    CGFloat proposedHeight =
        [[arr objectAtIndex:0] floatValue] + [[arr objectAtIndex:1] floatValue]
        + _topspace + _midspace + _bottomspace;
    CGFloat result =
        (proposedHeight >= tableView.rowHeight) ?
            tableView.rowHeight : proposedHeight;
    return result;
}

```

My `tableView:willDisplayCell:forRowAtIndexPath:` implementation calls `labelHeightsForItem:` again, using that same information and the same #defined spacer values again, this time setting the frames of the two labels within the cell's content view:

```

- (void)tableView:(UITableView *)tableView willDisplayCell:(UITableViewCell *)cell
    forRowAtIndexPath:(NSIndexPath *)indexPath {
    FPIItem* item = [self.parsedData.items objectAtIndex:[indexPath row]];
    NSArray* arr = [self labelHeightsForItem: item];
    CGRect f1 = [cell viewWithTag: 1].frame;
    f1.size.height = [[arr objectAtIndex: 0] floatValue];
    f1.origin.y = _topspace;
    [cell viewWithTag: 1].frame = f1;
    // ... and similarly for the second label ...
}

```

Table View Selection

A table view cell has a normal (deselected) state and a selected state, according to its `selected` property. It is possible to change a cell's `selected` property directly (possibly with animation, using `setSelected:animated:`), but you are more likely to manage selection through the table view, letting the table view deal with the `selected` property of its cells.

One of the chief purposes of your table view is likely to be to let the user select a cell. This will be possible, provided you have not set the value of the table view's `allowsSelection` property to NO. The user taps a normal cell, and the cell switches to its selected state. As we've already seen, this will usually mean that the cell is redrawn with

a blue (or gray) background view, but you can change this. If the user taps an already selected cell, by default it stays selected.

Starting in iOS 5, table views can permit the user to select multiple cells simultaneously. (This was possible previously by deliberately misusing some delegate methods, but the multiple selection was effectively happening behind the table view's back, which had some unfortunate side effects.) Set the table view's `allowsMultipleSelection` property to YES. If the user taps an already selected cell, by default it is deselected.

Your code can also learn and manage the selection through these `UITableView` instance methods:

```
indexPathForSelectedRow  
indexPathsForSelectedRows
```

These methods report the currently selected row(s), or nil if there is no selection. Don't accidentally call the wrong one. For example, calling `indexPathForSelectedRow` when the table view allows multiple selection give a result that will have you scratching your head in confusion. (As usual, I speak from experience.)

`selectRowAtIndexPath:animated:scrollPosition:`

The animation involves fading in the selection, but the user may not see this unless the selected row is already visible. The last parameter dictates whether and how the table view should scroll to reveal the newly selected row:

- `UITableViewScrollPositionTop`
- `UITableViewScrollPositionMiddle`
- `UITableViewScrollPositionBottom`
- `UITableViewScrollPositionNone`

For the first three options, the table view scrolls (with animation, if the second parameter is YES) so that the selected row is at the specified position among the visible cells. For `UITableViewScrollPositionNone`, the table view does not scroll; if the selected row is not already visible, it does not become visible. You can make it visible by calling `scrollToRowAtIndexPath:atScrollPosition:` (discussed later in this chapter).

`deselectRowAtIndexPath:animated:`

Deselects the given row (if it is selected); the optional animation involves fading out the selection. No automatic scrolling takes place. To deselect all currently selected rows, call `selectRowAtIndexPath:animated:scrollPosition:` with a nil index path.

When a table view changes a cell's `selected` state, it also changes its `highlighted` state. This causes the cell to propagate the highlighted state down through its subviews by setting each subview's `highlighted` property if it has one. That is why a `UILabel`'s `highlightedTextColor` applies when the cell is selected. Similarly, a `UIImageView` (such as the cell's `imageView`) can have a `highlightedImage` that is shown when the cell is selected, and a `UIControl` (such as a `UIButton`) takes on its `highlighted` state when the

cell is selected. You can set a cell’s `highlighted` state directly, with the `highlighted` property or `setHighlighted:animated:`; but you are unlikely to do so, instead leaving the table view to manage selection and highlighting together.



If you set a cell’s `selected` or `highlighted` property directly, you are effectively acting behind the table view’s back; it has no knowledge of what you’re doing, so it won’t track selection properly. This is another reason why it’s better to let the user tap to select, or call the table view selection methods.

Response to user selection is through the table view’s delegate. Despite their names, the “will” methods are actually “should” methods: return `nil` to prevent the selection (or deselection) from taking place; return the index path handed in as argument to permit the selection (or deselection), or a different index path to cause a different cell to be selected (or deselected):

- `tableView:willSelectRowAtIndexPath:`
- `tableView:didSelectRowAtIndexPath:`
- `tableView:willDeselectRowAtIndexPath:`
- `tableView:didDeselectRowAtIndexPath:`

When `tableView:willSelectRowAtIndexPath:` is called because the user taps a cell, and if this table view permits only single cell selection, `tableView:willDeselectRowAtIndexPath:` will be called subsequently for any previously selected cells.

Here’s an example of implementing `tableView:willSelectRowAtIndexPath:`. The default behavior for `allowsSelection` (not multiple selection) is that the user can select by tapping; if the user taps a selected row, the selection does not change. We can alter this so that tapping a selected row deselects it:

```
- (NSIndexPath*) tableView:(UITableView*)tv  
    willSelectRowAtIndexPath:(NSIndexPath*)ip {  
    if ([tv cellForRowAtIndexPath:indexPath:ip].selected) {  
        [tv deselectRowAtIndexPath:indexPath animated:NO];  
        return nil;  
    }  
    return ip;  
}
```

An extremely common response to user selection is navigation. A master–detail architecture is typical: the table view lists things the user can see in more detail, and a tap replaces the table view with the detailed view of the selected thing. Very often the table view will be in a navigation interface, and you will respond to user selection by creating the detail view and pushing it onto the navigation controller’s stack. This interface is so common that Xcode’s Master–Detail Application project template implements it for you.

For example, here's the code from my TidBITS News app that navigates from the list of articles to the actual article on which the user has tapped:

```
- (void)tableView:(UITableView *)tableView  
    didSelectRowAtIndexPath:(NSIndexPath *)indexPath {  
    id item = [self.parsedData.items objectAtIndex: [indexPath row]];  
    UIViewController* svc = [[StoryViewController alloc] initWithItem: item];  
    [self.navigationController pushViewController:svc animated:YES];  
}
```

If you're using a UITableViewController, then by default, whenever the table view appears, the selection is cleared automatically in `viewWillAppear:` (unless you disable this by setting the table view controller's `clearsSelectionOnViewWillAppear` to NO), and the scroll indicators are flashed in `viewDidAppear:`. I sometimes prefer to set `clearsSelectionOnViewWillAppear` to NO and implement deselection in `viewDidAppear:`; the effect is that when the user returns to the table, the row is still momentarily selected before it deselects itself:

```
- (void) viewDidAppear:(BOOL)animated {  
    // deselect selected row  
    [tableView selectRowAtIndexPath:nil animated:NO  
        scrollPosition:UITableViewScrollPositionNone];  
    [super viewDidAppear:animated];  
}
```

By convention, if selecting a table view cell causes navigation, the cell should be given an `accessoryType` of `UITableViewCellAccessoryDisclosureIndicator`. This is a plain gray right-pointing chevron at the right end of the cell. The chevron doesn't of itself respond to user interaction; it's just a visual cue that we'll "move to the right" if the user taps the cell.



Recall from [Chapter 19](#) that in a storyboard, if a segue emanates from a UITableViewCell, the storyboard assumes that you want the segue to be triggered when the user taps a cell.

An alternative `accessoryType` is `UITableViewCellAccessoryDetailDisclosureButton`. It is a button and *does* respond to user interaction, through your implementation of the table view delegate's `tableView:accessoryButtonTappedForRowWithIndexPath:`. The button has a right-pointing chevron, so once again you'd be likely to respond by navigating; in this case, however, you would probably use the button *instead* of selection as a way of letting the user navigate. A common convention is that selecting the cell as a whole does one thing and tapping the disclosure button does something else (involving navigation to the right). For example, in Apple's Phone app, tapping a contact's listing in the Recents table places a call to that contact, but tapping the disclosure button switches to that contact's detail view.

Another use of cell selection is to implement a choice among cells, where a section of a table effectively functions as an iOS alternative to Mac OS X radio buttons. The table

view usually has the grouped format. An `accessoryType` of `UITableViewCellAccessoryCheckmark` is typically used to indicate the current choice. Implementing radio-button behavior is up to you.

For example, here's how [Figure 21-2](#) is implemented. The table view is created in the grouped style, in code, as part of the view controller's `loadView` implementation:

```
UITableView* tv = [[UITableView alloc] initWithFrame:CGRectMake(0,0,320,310)
                                              style:UITableViewStyleGrouped];
[v addSubview:tv];
tv.dataSource = self;
tv.delegate = self;
tv.scrollEnabled = NO;
self.tableView = tv;
```

As data source, we supply the structure and content of the table. The user defaults are storing the current choice in each of the two categories, so we use them to decide where the checkmarks go:

```
- (NSInteger)numberOfSectionsInTableView:(UITableView *)tableView {
    return 2;
}

- (NSString *)tableView:(UITableView *)tableView
    titleForHeaderInSection:(NSInteger)section {
    if (section == 0)
        return @"Size";
    return @"Style";
}

- (NSInteger)tableView:(UITableView *)tableView
    numberOfRowsInSection:(NSInteger)section {
    if (section == 0)
        return 3;
    return 2;
}

- (UITableViewCell *)tableView:(UITableView *)tv
    cellForRowAtIndexPath:(NSIndexPath *)indexPath {
    static NSString *CellIdentifier = @"Cell";
    UITableViewCell *cell = [tv dequeueReusableCellWithIdentifier:CellIdentifier];
    if (cell == nil) {
        cell = [[UITableViewCell alloc]
            initWithStyle:UITableViewCellStyleDefault
            reuseIdentifier:CellIdentifier];
    }
    NSUInteger section = [indexPath section];
    NSUInteger row = [indexPath row];
    NSUserDefaults* ud = [NSUserDefaults standardUserDefaults];
    if (section == 0) {
        if (row == 0) {
            cell.textLabel.text = @"Easy";
        } else if (row == 1) {
            cell.textLabel.text = @"Normal";
        } else if (row == 2) {
            cell.textLabel.text = @"Large";
        }
    }
}
```

```

        cell.textLabel.text = @"Hard";
    }
} else if (section == 1) {
    if (row == 0) {
        cell.textLabel.text = @"Animals";
    } else if (row == 1) {
        cell.textLabel.text = @"Snacks";
    }
}
cell.accessoryType = UITableViewCellAccessoryNone;
if ([[ud valueForKey:@"Style"] isEqualToString:cell.textLabel.text] ||
    [[ud valueForKey:@"Size"] isEqualToString:cell.textLabel.text])
    cell.accessoryType = UITableViewCellAccessoryCheckmark;
return cell;
}

```

As delegate, we are called when the user taps a cell. We store the user's selection into the user defaults; then we reload the table data, which deselects the current selection and reassigned the checkmark:

```

- (void)tableView:(UITableView *)tv
    didSelectRowAtIndexPath:(NSIndexPath *)indexPath {
   NSUserDefaults* ud = [NSUserDefaults standardUserDefaults];
    NSString* setting = [tv cellForRowAtIndexPath:indexPath].textLabel.text;
    [ud setValue:setting forKey:
        [self tableView:tv titleForHeaderInSection:indexPath.section]];
    [self.tableView reloadData];
}

```

In iOS 5, I could rewrite that table view to be a *static table*. This is a new feature in iOS 5, and is available only through the storyboard editor, so I'd need a UITableViewController subclass instantiated from a storyboard. In the storyboard editor, you select the table and set its Content pop-up menu in the Attributes inspector to Static Cells. Then you can design the entire table, including section header and footer text, and the content of each cell ([Figure 21-6](#)).

In our code, we can then eliminate our implementations of `numberOfSectionsInTableView:`, `tableView:titleForHeaderInSection:`, and `tableView:numberOfRowsInSection:`; all of that information comes from the table designed in the storyboard. We still need to deal somehow with the display of the checkmarks. I can think of two possibilities (and there may be others):

- We can still implement `tableView:cellForRowIndexPath:`. We start by calling `super` to obtain the cell, completely configured according to the storyboard; we can then add or remove the checkmark:

```

// we can still modify the cell as long as we fetch it from super
UITableViewCell* cell = [super tableView:tv cellForRowAtIndexPath:indexPath];
NSUserDefaults* ud = [NSUserDefaults standardUserDefaults];
cell.accessoryType = UITableViewCellAccessoryNone;

```

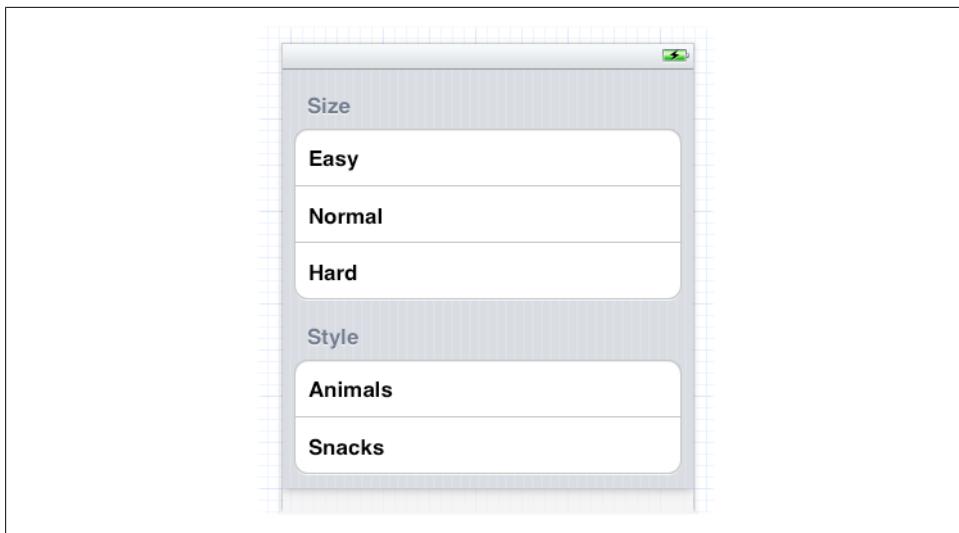


Figure 21-6. Designing a static table in the storyboard editor

```
if ([[ud valueForKey:@"Style"] isEqualToString:cell.textLabel.text] ||
    [[ud valueForKey:@"Size"] isEqualToString:cell.textLabel.text])
    cell.accessoryType = UITableViewCellAccessoryCheckmark;
return cell;
```

- We can omit `tableView:cellForRowAtIndexPath:` and implement `tableView:willDisplayCell:forRowAtIndexPath:` instead:

```
NSUserDefaults* ud = [NSUserDefaults standardUserDefaults];
cell.accessoryType = UITableViewCellAccessoryNone;
if ([[ud valueForKey:@"Style"] isEqualToString:cell.textLabel.text] ||
    [[ud valueForKey:@"Size"] isEqualToString:cell.textLabel.text])
    cell.accessoryType = UITableViewCellAccessoryCheckmark;
```

Table View Scrolling and Layout

A UITableView is a UIScrollView, so everything you already know about scroll views is applicable ([Chapter 20](#)). In addition, a table view supplies two convenience scrolling methods:

- `scrollToRowAtIndexPath:atScrollPosition:animated:`
- `scrollToNearestSelectedRowAtIndexPath:atScrollPosition:animated:`

The `scrollPosition` parameter is as for `selectRowAtIndexPath:...,` discussed earlier in this chapter.

The following UITableView methods mediate between the table's bounds coordinates on the one hand and table structure on the other:

- `indexPathForRowAtPoint:`

- `indexPathsForRowsInRect:`
- `rectForSection:`
- `rectForRowAtIndexPath:`
- `rectForFooterInSection:`
- `rectForHeaderInSection:`

The table's header and footer are views, so their coordinates are given by their frames.

Table View Searching

A table view is a common way to present the results of a search performed through a search field (a UISearchBar; see [Chapter 25](#)). This is such a standard interface, in fact, that a class is provided, UISearchDisplayController, to mediate between the search field where the user enters a search term and the table view listing the results of the search. The UISearchDisplayController needs the following things:

A search bar

A UISearchBar in the interface. This will be the UISearchDisplayController's `searchBar`.

A view controller

The view controller managing the view in the interface over which the search results are to appear. This will be the UISearchDisplayController's `searchContentsController`.

A results table view

The table view in which the search results will be presented. This will be the UISearchDisplayController's `searchResultsTableView`. It can already exist, or the UISearchDisplayController will create it.

A data source and delegate for the results table view

The UISearchDisplayController's `searchResultsDataSource` and `searchResultsDelegate`. They will control the data and structure of the search results table. They are commonly the same object, as for any table view; moreover, they are commonly the view controller.

A delegate

An optional object adopting the UISearchDisplayDelegate protocol. It will be notified of events relating to the display of results. It, too, is commonly the view controller.

Moreover, the UISearchBar itself can also have a delegate, and this, too, is commonly the view controller.

A UISearchDisplayController's `searchContentsController` needn't be a UITableView-Controller, and the data that the user is searching needn't be the content of an existing

table view. But they frequently are! That's because the mental connection between a table and a search is a natural one; when the search results are presented as a table view, the user feels that the search field is effectively filtering the contents of the original table view. A single object may thus be playing all of the following roles:

- The searchable table view's view controller
- The searchable table view's data source
- The searchable table view's delegate
- The view controller for the view over which the search results will appear
- The search results table view's data source
- The search results table view's delegate
- The UISearchDisplayController's delegate
- The UISearchBar's delegate

To illustrate, we will implement a table view that is searchable through a UISearchBar and that displays the results of that search in a second table view managed by a UISearchDisplayController.

The first question is how to make the search field appear along with the table view. Apple's own apps, such as the Contacts app, have popularized an interface in which the search field is the table view's header view. Indeed, this is such a common arrangement that the nib editor's Object library contains an object called Search Bar and Search Display Controller; if you drag this onto a UITableView in the nib editor, the search field becomes the table's header view and a UISearchDisplayController is created for you automatically, with all properties hooked up appropriately through outlets, much as I just described. In our example, however, we'll create the UISearchDisplayController and the UISearchBar in code.

Another feature of Apple's standard interface is that the search field isn't initially showing. To implement this, we'll scroll to the first actual row of data when the table view appears.

We're going to start with a table managed by a UITableViewController. In this view controller's `viewDidLoad`, we create the search bar and slot it in as the table's header view; we then load the data and scroll the header view out of sight. We also create the UISearchDisplayController and tie it to the search bar — and to ourselves (the UITableViewController) as the UISearchDisplayController's controller, delegate, search table data source, and search table delegate, as well as making ourselves the UISearchBar delegate. We also retain the UISearchDisplayController by assigning it to a property, so that it doesn't vanish in a puff of smoke before we can use it:

```
UISearchBar* b = [[UISearchBar alloc] init];
[b sizeToFit];
b.delegate = self;
[self.tableView setTableHeaderView:b];
[self.tableView reloadData];
```

```

[self.tableView
    scrollToRowAtIndexPath:[NSIndexPath indexPathForRow:0 inSection:0]
    atScrollPosition:UITableViewScrollPositionTop animated:NO];
UISearchDisplayController* c =
    [[UISearchDisplayController alloc] initWithSearchBar:b
                                                contentsController:self];
self.sbc = c; // retain the UISearchDisplayController
c.delegate = self;
c.searchResultsDataSource = self;
c.searchResultsDelegate = self;

```

When the user initially taps in the search field, the `UISearchDisplayController` automatically constructs a new interface along with a nice animation. This indicates to the user that the search field is ready to receive input; when the user proceeds to enter characters into the search field, the `UISearchDisplayController` is ready to display its own search results table view in this interface. The `UISearchBar` has a `Cancel` button that the user can tap to dismiss the interface created by the `UISearchDisplayController`.

Populating the search results table in response to what the user does in the `UISearchBar` is up to us. The `UITableViewContrroller` is both data source and delegate for the original table view, as well as data source and delegate for the search results table. This means that our search is already almost working, because the search results table will automatically have the same data and structure as the original table! Our only additional task, beyond what our code already does, is to check whether the table view that's talking to us is the search results table view (this will be the `UISearchDisplayController`'s `searchResultsTableView`) and, if it is, to limit our returned data with respect to the search bar's text. The strategy for doing this should be fairly obvious if we are maintaining our source data in a sensible model.

Let's say, for the sake of simplicity, that our original table is displaying the names of the 50 United States, which it is getting from an array of strings called `states`:

```

- (NSInteger)numberOfSectionsInTableView:(UITableView *)tableView {
    return 1;
}

- (NSInteger)tableView:(UITableView *)tableView
    numberOfRowsInSection:(NSInteger)section {
    NSArray* model = self.states;
    return [model count];
}

// Customize the appearance of table view cells.
- (UITableViewCell *)tableView:(UITableView *)tableView
    cellForRowAtIndexPath:(NSIndexPath *)indexPath {
    static NSString *CellIdentifier = @"Cell";
    UITableViewCell *cell =
        [tableView dequeueReusableCellWithIdentifier:CellIdentifier];
    if (cell == nil) {
        cell = [[UITableViewCell alloc]
            initWithStyle:UITableViewCellStyleDefault
            reuseIdentifier:CellIdentifier];
    }
    cell.textLabel.text = [self.states objectAtIndex:indexPath.row];
    return cell;
}

```

```

    }
    NSArray* model = self.states;
    cell.textLabel.text = [model objectAtIndex: indexPath.row];
    return cell;
}

```

To make this work with a UISearchDisplayController, the only needed change is this: Each time we speak of the NSArray called `model`, we must decide whether it should be `self.states`, as now, or whether it should be a *different* array that is filtered with respect to the current search — let's call it `self.filteredStates`. There are two occurrences of this line:

```
NSArray* model = self.states;
```

They are now to be replaced by this:

```
NSArray* model =
    (tableView == sbc.searchResultsTableView) ? self.filteredStates : self.states;
```

The only remaining question is when and how this `filteredStates` array should be calculated. One approach is to ignore the user typing into the search field and calculate the filtered array only when the user taps the Search button in the keyboard. We are the UISearchBar delegate, so we can hear about the user tapping the Search button by implementing `searchBarSearchButtonClicked:`. We create the filtered array and reload the UISearchDisplayController's search results table view:

```

- (void)searchBarSearchButtonClicked:(UISearchBar *)searchBar {
    NSPredicate* p = [NSPredicate predicateWithBlock:
        ^BOOL(id obj, NSDictionary *d) {
            NSString* s = obj;
            return ([s rangeOfString:searchBar.text
                options:NSCaseInsensitiveSearch].location != NSNotFound);
        }];
    self.filteredStates = [states filteredArrayUsingPredicate:p];
    [self.sbc.searchResultsTableView reloadData];
}

```

This works fine. However, as soon as the user starts typing into the search field, the results table becomes empty, and sits there empty as the user types. This is because the results table reloads whenever the user changes the contents of the search field and is finding that it has no data (because `filteredStates` is nil).

One solution is to set the `filteredStates` array to the `states` array before the search begins. We can know when this will be through a UISearchDisplayController delegate method:

```

- (void)searchDisplayControllerWillBeginSearch:(UISearchDisplayController *)c {
    self.filteredStates = self.states;
}

```

Alternatively, we can generate a new set of search results every time the user types in the search field, effectively implementing a “live” search ([Figure 21-7](#)). This is perfectly reasonable for our extremely small data set of 50 states, though of course it mightn't

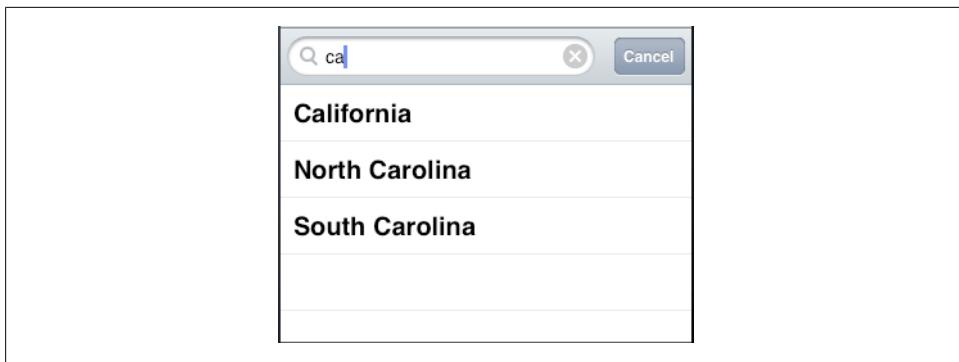


Figure 21-7. Filtering a table with a search bar

work well if the data set were very large, or if there were for some other reason a delay in filtering it. To implement live search, we turn our implementation of `searchBarSearchButtonClicked:` into an implementation of a different `UISearchBar` delegate method, `searchBar:textDidChange:`. Everything else stays the same, except that now there is no need to reload the search results table's data, as by default the `UISearchDisplayController` will do that automatically:

```
- (void)searchBar:(UISearchBar *)searchBar textDidChange:(NSString *)searchText {
    NSPredicate* p = [NSPredicate predicateWithBlock:
        ^BOOL(id obj, NSDictionary *d) {
            NSString* s = obj;
            return ([s rangeOfString:searchText
                options:NSCaseInsensitiveSearch].location != NSNotFound);
        }];
    self.filteredStates = [states filteredArrayUsingPredicate:p];
}
```

A `UISearchBar` can also display scope buttons, letting the user alter the meaning of the search. If you add these, then of course you must take them into account when filtering the model data. For example, let's have two scope buttons, "Starts With" and "Contains":

```
UISearchBar* b = [[UISearchBar alloc] init];
[b sizeToFit];
b.scopeButtonTitles = [NSArray arrayWithObjects: @"Starts With", @"Contains", nil];
// ...
```

Our filtering routine must now take the state of the scope buttons into account. Moreover, the search results table view will reload when the user changes the scope (which we can detect in another `UISearchBar` delegate method, `searchBar:selectedScopeButtonIndexDidChange:`), so if we're doing a live search, we must respond by filtering the data then as well. To prevent repetition, we'll abstract the filtering routine into a method of its own:

```

- (void)filterData {
    NSPredicate* p = [NSPredicate predicateWithBlock:
        ^BOOL(id obj, NSDictionary *d) {
            NSString* s = obj;
            NSStringCompareOptions options = NSCaseInsensitiveSearch;
            if (sbc.searchBar.selectedScopeButtonIndex == 0)
                options |= NSAnchoredSearch;
            return ([s rangeOfString:sbc.searchBar.text
                options:options].location != NSNotFound);
        }];
    self.filteredStates = [states filteredArrayUsingPredicate:p];
}

- (void)searchBar:(UISearchBar *)searchBar textDidChange:(NSString *)searchText {
    [self filterData];
}

- (void)searchBar:(UISearchBar *)searchBar
    selectedScopeButtonIndexDidChange:(NSInteger)selectedScope {
    [self filterData];
}

```

In an indexed list — one with sections and an index running down the right side — a “magnifying glass” search symbol can be made to appear in the index by including `UITableViewIndexSearch` (usually as the first item) in the string array returned from `sectionIndexTitlesForTableView:`. For example, suppose that as in our earlier example, the section names are to be used as index entries and are in an array called `sectionNames`:

```

- (NSArray *)sectionIndexTitlesForTableView:(UITableView *)tableView {
    return [[NSArray arrayWithObject: UITableViewIndexSearch]
        arrayByAddingObjectsFromArray:sectionNames];
}

```

You’ll also need to implement `tableView:sectionForSectionIndexTitleAtIndex:`, because now the correspondence between index entries and sections is off by one. If the user taps the magnifying glass in the index, you scroll to reveal the search field (and you’ll also have to return a bogus section number, but there is no penalty for that):

```

- (NSInteger)tableView:(UITableView *)tableView
    sectionForSectionIndexTitle:(NSString *)title
    atIndex:(NSInteger)index {
    if (index == 0)
        [tableView scrollRectToVisible:tableView.tableHeaderView.frame
            animated:NO];
    return index-1;
}

```

Whenever the search results table becomes empty (because the search bar is nonempty and `filteredStates` is nil), the words “No Results” appear superimposed on it. I find this incredibly obnoxious, and I can’t believe that after all these years Apple still hasn’t granted programmers an official way to remove or customize it. Here’s an unofficial way:

```

-(BOOL)searchDisplayController:(UISearchDisplayController *)controller
    shouldReloadTableForSearchString:(NSString *)searchString {
    dispatch_time_t popTime = dispatch_time(DISPATCH_TIME_NOW, 1);
    dispatch_after(popTime, dispatch_get_main_queue(), ^{
        for (UIView* v in self.sbc.searchResultsTableView.subviews) {
            if ([v isKindOfClass: [UILabel class]] &&
                [[(UILabel*)v text] isEqualToString:@"No Results"]) {
                [(UILabel*)v setText: @""];
                break;
            }
        }
    });
    return YES;
}

```



A UISearchBar has many properties through which its appearance can be configured; I'll discuss them in [Chapter 25](#). Both the UISearchBar and UISearchDisplayController send their delegate numerous messages that you can take advantage of to customize behavior; consult the documentation. A UISearchBar in a UIToolbar on the iPad can display its results in a popover; I'll talk about that in [Chapter 22](#).

Table View Editing

A table view cell has a normal state and an editing state, according to its `editing` property. The editing state is typically indicated visually by one or more of the following:

Editing controls

At least one editing control will usually appear, such as a minus button (for deletion) at the left side.

Shrinkage

The content of the cell will usually shrink to allow room for an editing control. You can prevent a cell in a grouped-style table from shifting its left end rightward in editing mode by setting its `shouldIndentWhileEditing` to NO, or with the table delegate's `tableView:shouldIndentWhileEditingRowAtIndexPath:`.

Changing accessory view

The cell's accessory view will change automatically in accordance with its `editingAccessoryType` or `editingAccessoryView`. If you assign neither, so that they are nil, the cell's accessory view will vanish when in editing mode.

As with selection, you can set a cell's `editing` property directly (or use `setEditing:animated:` to get animation), but you are more likely to let the table view manage editability. Table view editability is controlled through the table view's `editing` property, usually by sending the table the `setEditing:animated:` message. The table is then responsible for putting its cells into edit mode.



A cell in edit mode can also be selected by the user if the table view's `allowsSelectionDuringEditing` (or, in iOS 5, `allowsMultipleSelectionDuringEditing`) is YES. But this would be unusual.

Putting the table into edit mode is usually left up to the user. A typical interface would be an Edit button that the user can tap. In a navigation interface, we might have our view controller supply the button as the navigation item's right button:

```
UIBarButtonItem* bbi =
    [[UIBarButtonItem alloc] initWithBarButtonSystemItem:UIBarButtonSystemItemEdit
                                                target:self action:@selector(doEdit:)];
self.navigationItem.rightBarButtonItem = bbi;
```

Our action handler will be responsible for putting the table into edit mode, so in its simplest form it might look like this:

```
- (void) doEdit: (id) sender {
    [self.tableView setEditing:YES animated:YES];
}
```

But that does not solve the problem of getting *out* of editing mode. The standard solution is to have the Edit button replace itself by a Done button:

```
- (void) doEdit: (id) sender {
    int which;
    if (![self.tableView isEditing]) {
        [self.tableView setEditing:YES animated:YES];
        which = UIBarButtonItemStyleDone;
    } else {
        [self.tableView setEditing:NO animated:YES];
        which = UIBarButtonItemStyleEdit;
    }
    UIBarButtonItem* bbi = [[UIBarButtonItem alloc]
                           initWithBarButtonSystemItem:which target:self action:@selector(doEdit:)];
    self.navigationItem.rightBarButtonItem = bbi;
}
```

However, it turns out that all of this is completely unnecessary if we want standard behavior, as it is already implemented for us! A `UIViewController` supplies an `editButtonItem` that calls the `UIViewController`'s `setEditing:animated:` when tapped, tracks whether we're in edit mode with the `UIViewController`'s `editing` property, and changes its own title accordingly. Moreover, a `UITableViewController`'s implementation of `setEditing:animated:` is to call `setEditing:animated:` on its table view. Thus, if we're using a `UITableViewController`, we get all of that behavior for free just by inserting the `editButtonItem` into our interface:

```
self.navigationItem.rightBarButtonItem = self.editButtonItem;
```

When the table view enters edit mode, it consults its data source and delegate about the editability of individual rows:

tableView:canEditRowAtIndexPath: *to the data source*

The default is YES. The data source can return NO to prevent the given row from entering edit mode.

tableView:editingStyleForRowAtIndexPath: *to the delegate*

Each standard editing style corresponds to a control that will appear in the cell. The choices are:

UITableViewCellStyleDelete

The cell shows a minus button at its left end. The user can tap this to summon a Delete button, which the user can then tap to confirm the deletion. This is the default.

UITableViewCellStyleInsert

The cell shows a plus button at its left end; this is usually taken to be an insert button.

UITableViewCellStyleNone

No editing control appears.

If the user taps an insert button (the plus button) or a delete button (the Delete button that appears after the user taps the minus button), the data source is sent the **tableView:commitEditingStyle:forRowAtIndexPath:** message and is responsible for obeying it. In your response, you will probably want to alter the structure of the table, and UITableView methods for doing this are provided:

- **insertRowsAtIndexPaths:withRowAnimation:**
- **deleteRowsAtIndexPaths:withRowAnimation:**
- **insertSections:withRowAnimation:**
- **deleteSections:withRowAnimation:**

The row animations here are effectively the same ones discussed earlier in connection with refreshing table data; “left” for an insertion means to slide in from the left, and for a deletion it means to slide out to the left, and so on.

New in iOS 5 are two more methods for rearranging the structure of the table; there is animation, but there’s no provision for customizing it:

- **moveSection:toSection:**
- **moveRowAtIndexPath:toIndexPath:**

If you’re issuing more than one of these commands, you can combine them by surrounding them with **beginUpdates** and **endUpdates**, forming an *updates block*. An updates block combines not just the animations but the requested changes themselves. This relieves you from having to worry about how a command is affected by earlier commands in the same updates block; indeed, order of commands within an updates block doesn’t really matter.

For example, if you delete row 1 of a certain section and then (in a separate command) delete row 2 of the same section, you delete two successive rows, just as you would expect; the notion “2” does not change its meaning because you deleted an earlier row first, because you *didn’t* delete an earlier row first — the updates block combines the commands for you, interpreting both index paths with respect to the state of the table before any changes are made. If you perform insertions and deletions together in one animation, the deletions are performed first, regardless of the order of your commands, and the insertion row and section numbers refer to the state of the table after the deletions.

An updates block can also include `reloadRows...` and `reloadSections...` commands (but not `reloadData`).

I need hardly emphasize once again (but I will anyway) that view is not model. It is one thing to rearrange the appearance of the table, another to alter the underlying data. It is up to you to make certain you do both together. Do not, even for a moment, permit the data and the view to get out of synch with each other. If you delete a row, remove from the model the datum that it represents. The runtime will try to help you with error messages if you forget to do this, but in the end the responsibility is yours. I’ll give examples as we proceed.



Here’s an interesting trick: an empty updates block lays out the table view, fetching the section header and footer titles or views, their heights, and the row heights, *without reloading any cells*. Apple takes advantage of this in the Table View Animations and Gestures example, in which a pinch gesture is used to change a table’s row height in real time.

Deleting Table Items

Deletion of table items is the default, so there’s not much for us to do in order to implement it. If our view controller is a `UITableViewController` and we’ve displayed the Edit button as its navigation item’s right button, everything happens automatically: the user taps the Edit button, the view controller’s `setEditing:animated:` is called, the table view’s `setEditing:animated:` is called, and the cells all show the minus button at the left end. The user can then tap a minus button; a Delete button appears at the cell’s right end. You can customize the Delete button’s title with the table delegate method `tableView:titleForDeleteConfirmationButtonForRowAtIndexPath:`.

What is *not* automatic is the actual response to the Delete button. For that, we need to implement `tableView:commitEditingStyle:forRowAtIndexPath:`. Typically, you’ll remove the corresponding entry from the underlying model data, and you’ll call `deleteRowsAtIndexPaths:withRowAnimation:` or `deleteSections:withRowAnimation:` to update the appearance of the table. As I said a moment ago, you must delete the row or section in such a way as to keep the table display coordinated with the model’s structure. Otherwise, the app may crash (but with an extremely helpful error message).

To illustrate, let's suppose that the underlying model is an array of arrays, maintained as `sectionNames` and `sectionData`, as in our earlier example. These arrays must now be mutable. Our approach will be in two stages:

1. Deal with the model data. We'll delete the requested row; if this empties the section array, we'll also delete that section array and the corresponding section name.
2. Deal with the table's appearance. If we deleted the section array, we'll call `deleteSections:withRowAnimation:`; otherwise, we'll call `deleteRowsAtIndexPaths:withRowAnimation:`. In the former case, if a section index appears in the table, we will also reload the section index:

```
- (void)tableView:(UITableView *)tableView
    commitEditingStyle:(UITableViewCellEditingStyle)editingStyle
        forRowAtIndexPath:(NSIndexPath *)indexPath {
    [[self.sectionData objectAtIndex:indexPath.section]
        removeObjectAtIndex:indexPath.row];
    if ([[self.sectionData objectAtIndex:indexPath.section] count] == 0) {
        [self.sectionData removeObjectAtIndex:indexPath.section];
        [self.sectionNames removeObjectAtIndex:indexPath.section];
        [tableView deleteSections:[NSSet indexSetWithIndex: indexPath.section]
            withRowAnimation:UITableViewRowAnimationLeft];
        [tableView reloadSectionIndexTitles];
    } else {
        [tableView deleteRowsAtIndexPaths:[NSArray arrayWithObject:indexPath]
            withRowAnimation:UITableViewRowAnimationLeft];
    }
}
```

The user can also delete a row by swiping it to summon its Delete button *without* having explicitly entered edit mode; no other row is editable, and no other editing controls are shown. This feature is implemented “for free” by virtue of our having supplied an implementation of `tableView:commitEditingStyle:forRowAtIndexPath:`. If you're like me, your first response will be: “Thanks for the free functionality, Apple, and now how do I turn this off?” Because the Edit button is already using the `UIViewController`'s `editing` property to track edit mode, we can take advantage of this and refuse to let any cells be edited unless the view controller *is* in edit mode:

```
- (UITableViewCellEditingStyle)tableView:(UITableView *)aTableView
    editingStyleForRowAtIndexPath:(NSIndexPath *)indexPath {
    return self.editing ?
        UITableViewCellEditingStyleDelete : UITableViewCellEditingStyleNone;
}
```

Editable Content in Table Items

A table item might have content that the user can edit directly, such as a `UITextField` ([Chapter 23](#)). Because the user is working in the view, you need a way to reflect the user's changes into the model. This will probably involve putting yourself in contact with the interface objects where the user does the editing.

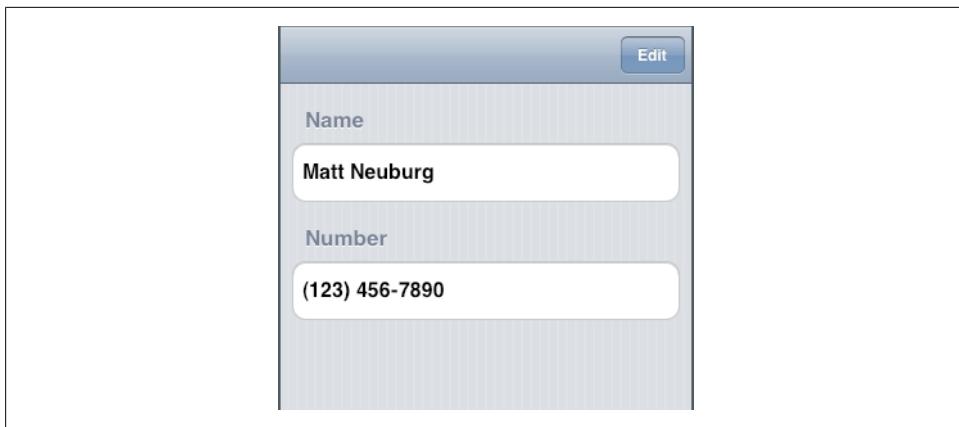


Figure 21-8. A simple phone directory app

To illustrate, I'll implement a table view cell with a text field that is editable when the cell is in editing mode. Imagine an app that maintains a list of names and phone numbers. A name and phone number are displayed as a grouped-style table, and they become editable when the user taps the Edit button (Figure 21-8).

A UITextField is editable if its `enabled` is YES. To tie this to the cell's `editing` state, it is probably simplest to implement a custom UITableViewCell class. I'll call it MyCell, and I'll design it in the nib, giving it a single UITextField that's pointed to through a property called `textField`. In the code for MyCell, we override `didTransitionToState:`, as follows:

```
- (void) didTransitionToState:(UITableViewCellStateMask)state {
    [super didTransitionToState:state];
    if (state == UITableViewCellStateEditingMask) {
        self.textField.enabled = YES;
    }
    if (state == UITableViewCellStateDefaultMask) {
        self.textField.enabled = NO;
    }
}
```

In the table's data source, we make ourselves the text field's delegate when we create and configure the cell (the nib is previously registered with `registerNib:forCellReuseIdentifier:`, and `CellIdentifier` is an external global, as I suggested earlier):

```
- (UITableViewCell *)tableView:(UITableView *)tableView
    cellForRowAtIndexPath:(NSIndexPath *)indexPath {
    MyCell* cell =
        (MyCell*)[tableView dequeueReusableCellWithIdentifier:CellIdentifier];
    if (indexPath.section == 0)
        cell.textField.text = self.name;
    if (indexPath.section == 1) {
        cell.textField.text = [self.numbers objectAtIndex: indexPath.row];
        cell.textField.keyboardType = UIKeyboardTypeNumbersAndPunctuation;
```

```

    }
    cell.textField.delegate = self;
    return cell;
}

```

We are the UITextField's delegate, so we are responsible for implementing the Return button in the keyboard to dismiss the keyboard:

```

- (BOOL)textFieldShouldReturn:(UITextField *)tf {
    [tf endEditing:YES];
    return NO;
}

```

Now comes the interesting part. When a text field stops editing, we are its delegate, so we can hear about it in `textFieldDidEndEditing:`. We work out which cell it belongs to, and update the model accordingly:

```

- (void)textFieldDidEndEditing:(UITextField *)tf {
    // some cell's text field has finished editing; which cell?
    UIView* v = tf;
    do {
        v = v.superview;
    } while (![v isKindOfClass: [UITableViewCell class]]);
    MyCell* cell = (MyCell*)v;
    // update data model to match
    NSIndexPath* ip = [self.tableView indexPathForCell:cell];
    if (ip.section == 1)
        [self.numbers replaceObjectAtIndex:ip.row withObject:cell.textField.text];
    else if (ip.section == 0)
        self.name = cell.textField.text;
}

```

Inserting Table Items

You are unlikely to attach a plus (insert) button to every row. A more likely interface is that when a table is edited, every row has a minus button except the last row, which has a plus button; this shows the user that a new row can be inserted at the end of the table.

Let's implement this for phone numbers in our name-and-phone-number app, allowing the user to give a person any quantity of phone numbers ([Figure 21-9](#)):

```

- (UITableViewCellEditingStyle)tableView:(UITableView *)tableView
    editingStyleForRowAtIndexPath:(NSIndexPath *)indexPath {
    if (indexPath.section == 1) {
        NSInteger ct =
            [self.tableView tableView numberOfRowsInSection:indexPath.section];
        if (ct-1 == indexPath.row)
            return UITableViewCellEditingStyleInsert;
        return UITableViewCellEditingStyleDelete;
    }
    return UITableViewCellEditingStyleNone;
}

```



Figure 21-9. Phone directory app in editing mode

The person's name has no editing control (a person must have exactly one name), so we prevent it from indenting in edit mode:

```
- (BOOL)tableView:(UITableView *)tableView
    shouldIndentWhileEditingRowAtIndexPath:(NSIndexPath *)indexPath {
    if (indexPath.section == 1)
        return YES;
    return NO;
}
```

When the user taps an editing control, we must respond. We immediately force our text fields to cease editing: the user may have tapped the editing control while editing, and we want our model to contain the very latest changes, so this is effectively a way of causing our `textFieldDidEndEditing:` to be called. The model for our phone numbers is a mutable array of strings, `numbers`. We already know what to do when the tapped control is a delete button; things are similar when it's an insert button, but we've a little more work to do. The new row will be empty, and it will be at the end of the table; so we append an empty string to the `numbers` model array, and then we insert a corresponding row at the end of the view. But now two successive rows have a plus button; the way to fix that is to reload the first of those rows. Finally, we also show the keyboard for the new, empty phone number, so that the user can start editing it immediately; we do that outside the update block:

```

- (void) tableView:(UITableView *)tableView
commitEditingStyle:(UITableViewCellEditingStyle)editingStyle
forRowAtIndexPath:(NSIndexPath *)indexPath {
    [tableView endEditing:YES]; // added, because user can click minus/plus while still editing
                                // so we must force saving to the model
    if (editingStyle == UITableViewCellEditingStyleInsert) {
        [self.numbers addObject: @" "];
        NSInteger ct = [self.numbers count];
        [tableView beginUpdates];
        [tableView insertRowsAtIndexPaths:
            [NSArray arrayWithObject: [NSIndexPath indexPathForRow: ct-1 inSection:1]]
            withRowAnimation:UITableViewRowAnimationAutomatic];
        [self.tableView reloadRowsAtIndexPaths:
            [NSArray arrayWithObject:[NSIndexPath indexPathForRow:ct-2 inSection:1]]
            withRowAnimation:UITableViewRowAnimationAutomatic];
        [tableView endUpdates];
        // crucial that this next bit be *outside* the update block
        UITableViewCell* cell = [self.tableView cellForRowAtIndexPath:
            [NSIndexPath indexPathForRow:ct-1 inSection:1]];
        [((MyCell*)cell).textField becomeFirstResponder];
    }
    if (editingStyle == UITableViewCellEditingStyleDelete) {
        [self.numbers removeObjectAtIndex:indexPath.row];
        [tableView beginUpdates];
        // update block makes animations work together
        [tableView deleteRowsAtIndexPaths:[NSArray arrayWithObject: indexPath]
            withRowAnimation:UITableViewRowAnimationAutomatic];
        [tableView reloadSections:[NSIndexSet indexSetWithIndex:1]
            withRowAnimation:UITableViewRowAnimationAutomatic];
        [tableView endUpdates];
    }
}

```

Rearranging Table Items

If the data source implements `tableView:moveRowAtIndexPath:toIndexPath:`, the table displays a reordering control at the right end of each row in editing mode, and the user can drag it to rearrange table items. The reordering control can be prevented for individual table items by implementing `tableView:canMoveRowAtIndexPath:`. The user is free to move rows that display a reordering control, but the delegate can limit where a row can be moved to by implementing `tableView:targetIndexPathForMoveFromRowAtIndexPath:toProposedIndexPath:`.

To illustrate, we'll add to our name-and-phone-number app the ability to rearrange phone numbers. There must be multiple phone numbers to rearrange:

```

- (BOOL)tableView:(UITableView *)tableView
canMoveRowAtIndexPath:(NSIndexPath *)indexPath {
    if (indexPath.section == 1 && [self.numbers count] > 1)
        return YES;
    return NO;
}

```

In our example, a phone number must not be moved out of its section, so we implement the delegate method to prevent this. We also take this opportunity to dismiss the keyboard if it is showing.

```
- (NSIndexPath *)tableView:(UITableView *)tableView
    targetIndexPathForMoveFromRowAtIndexPath:(NSIndexPath *)sourceIndexPath
        toProposedIndexPath:(NSIndexPath *)proposedDestinationIndexPath {
    [tableView endEditing:YES];
    if (proposedDestinationIndexPath.section == 0)
        return [NSIndexPath indexPathForRow:0 inSection:1];
    return proposedDestinationIndexPath;
}
```

After the user moves an item, `tableView:moveRowAtIndexPath:toIndexPath:` is called, and we trivially update the model to match. We also reload the table, to fix the editing controls:

```
- (void)tableView:(UITableView *)tableView
    moveRowAtIndexPath:(NSIndexPath *)fromIndexPath
        toIndexPath:(NSIndexPath *)toIndexPath {
    NSString* s = [self.numbers objectAtIndex: fromIndexPath.row];
    [self.numbers removeObjectAtIndex: fromIndexPath.row];
    [self.numbers insertObject:s atIndex: toIndexPath.row];
    [tableView reloadData];
}
```

Dynamic Table Content

We can rearrange a table not just in response to the user working in edit mode, but for some other reason entirely. In this way, many interesting and original interfaces are possible. In this example (the idea for which is shamelessly stolen from a WWDC 2010 video), we permit the user to double-tap on a section header as a way of collapsing or expanding the section — that is, we'll suppress or permit the display of the rows of the section, with a nice animation when we switch from one to the other.

Presume that our data model once again consists of the two arrays, `sectionNames` and `sectionData`. I've also got an `NSMutableSet`, `hiddenSections`, in which I'll list the sections aren't showing their rows. That list is all I'll need, since either a section is showing all its rows or it's showing none of them:

```
- (NSInteger)tableView:(UITableView *)tableView
    numberOfRowsInSection:(NSInteger)section {
    if ([self.hiddenSections containsObject:
        [NSNumber numberWithUnsignedInteger:section]])
        return 0;
    return [[sectionData objectAtIndex: section] count];
}
```

The section headers are a custom `UILabel` to which a `UITapGestureRecognizer` is attached, so we can detect a double-tap. Here's how we respond to a double-tap. We examine the tapped label to learn what section this is, and find out how many rows it has, as we'll need to know that later whether we're about to show or hide rows. Then

we look for the section number in our `hiddenSections` set. If it's there, we're about to display the rows, so we remove that number from `hiddenSections`; now we work out the index paths of the rows we're about to insert, and we insert them. If it's not there, we're about to hide the rows, so we insert that number into `hiddenSections`; again, we work out the index paths of the rows we're about to delete, and we delete them.

```
- (void) tap: (UIGestureRecognizer*) g {
    UILabel* lab = (UILabel*)g.view;
    NSString* s = lab.text;
    NSUInteger sec = [self.sectionNames indexOfObject:s];
    NSUInteger ct = [(NSArray*)[self.sectionData objectAtIndex:sec] count];
    NSNumber* secnum = [NSNumber numberWithUnsignedInteger:sec];

    if ([self.hiddenSections containsObject:secnum]) {
        [self.hiddenSections removeObject:secnum];
        [self.tableView beginUpdates];
        NSMutableArray* arr = [NSMutableArray array];
        for (int ix = 0; ix < ct; ix++) {
            NSIndexPath* ip = [NSIndexPath indexPathForRow:ix inSection:sec];
            [arr addObject: ip];
        }
        [self.tableView insertRowsAtIndexPaths:arr
            withRowAnimation:UITableViewRowAnimationAutomatic];
        [self.tableView endUpdates];
        // if necessary, scroll to make insertion visible
        [self.tableView scrollToRowAtIndexPath:[arr lastObject]
            atScrollPosition:UITableViewScrollPositionNone animated:YES];
    } else {
        [self.hiddenSections addObject:secnum];
        [self.tableView beginUpdates];
        NSMutableArray* arr = [NSMutableArray array];
        for (int ix = 0; ix < ct; ix++) {
            NSIndexPath* ip = [NSIndexPath indexPathForRow:ix inSection:sec];
            [arr addObject: ip];
        }
        [self.tableView deleteRowsAtIndexPaths:arr
            withRowAnimation:UITableViewRowAnimationAutomatic];
        [self.tableView endUpdates];
    }
}
```

Table View Menus

Starting in iOS 5, it is possible to display a menu from a table view cell by performing a long press on the cell. A menu, in iOS, is a sort of balloon containing tappable words such as Copy, Cut, and Paste. And as far as I can tell, those are the only words you'll be including in a table view cell's menu; I tried to customize the menu to include other terms, but I failed.

To allow the user to display a menu from a table view's cells, you implement three delegate methods:

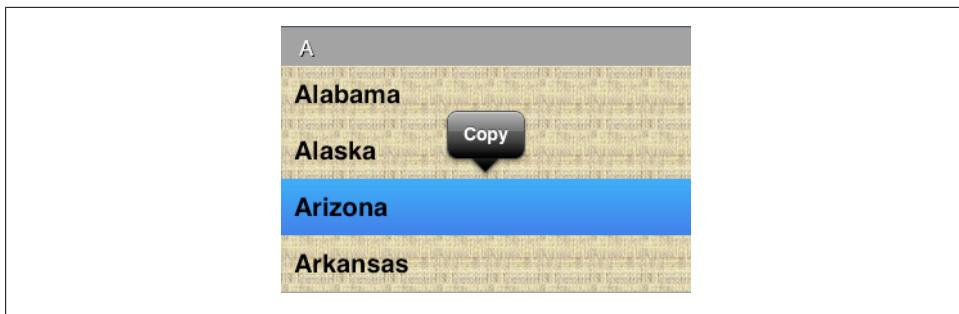


Figure 21-10. A table view cell with a menu

tableView:shouldShowMenuForRowAtIndexPath:

Return YES if the user is to be permitted to summon a menu by performing a long press on this cell.

tableView:canPerformAction:forRowAtIndexPath:withSender:

You'll be called repeatedly with a bunch of selectors for various actions that the system knows about, but as far as I can tell, the only ones worth responding YES to are `cut:`, `copy:`, and `paste:`. Whichever ones you respond YES to will appear in the menu; returning YES, regardless, causes all three menu items to appear in the menu. The menu will now appear unless you return NO to all three actions. The sender is the shared `UIMenuController`, which I'll discuss more in [Chapter 23](#) and [Chapter 39](#).

tableView:performAction:forRowAtIndexPath:withSender:

The user has tapped one of the menu items; your job is to respond to it somehow.

Here's an example where the user can summon a Copy menu item from any cell ([Figure 21-10](#)):

```
- (BOOL)tableView:(UITableView *)tableView  
    shouldShowMenuForRowAtIndexPath:(NSIndexPath *)indexPath {  
    return YES;  
}  
  
- (BOOL)tableView:(UITableView *)tableView canPerformAction:(SEL)action  
    forRowAtIndexPath:(NSIndexPath *)indexPath withSender:(id)sender {  
    return (action == @selector(copy:));  
}  
  
- (void)tableView:(UITableView *)tableView performAction:(SEL)action  
    forRowAtIndexPath:(NSIndexPath *)indexPath withSender:(id)sender {  
    NSString* s = [[self.sectionData objectAtIndex:indexPath.section]  
        objectAtIndex:indexPath.row];  
    if (action == @selector(copy:)) {  
        // ... do whatever copying consists of ...  
    }  
}
```

As Figure 21-10 shows, the long press gesture used to summon a menu also causes the pressed cell to assume its selected state—and hence its selected appearance. Moreover, tapping a menu item to choose it deselects the cell, even if it was previously selected; and tapping elsewhere, to dismiss the menu without choosing any menu item, may then select the cell under that tap.

This interweaving of the ability to summon a menu with the cell selection mechanism is most unfortunate, especially since Apple has not also provided new properties for detecting that menu display is occurring or for customizing what happens when it is. You cannot prevent the long-pressed cell from being selected by any of the normal means—for example, by returning nil from `tableView:willSelectRowAtIndexPath:` or by setting the table view's `allowsSelection` to NO.

If cell selection is not permitted in general, and if you want to prevent the cell that is showing a menu to appear selected, you can use properties like a label's `highlightedTextColor` and a cell's `selectedBackgroundView` (or an implementation of `setSelected:animated:` in a custom `UITableViewCell` subclass) to suppress the selected appearance. But if cell selection *is* permitted in general, making a cell that's displaying a menu look different from a normally selected cell, or preventing the cell's deselection after menu display if the cell was previously selected, is extremely difficult, because your code can't readily obtain enough state information to interpret what's happening. I have experimented with workarounds, but the results are not sufficiently convincing to be worth printing here. Until there's an iOS update that attends to these issues, your use of table cell menu display will probably be fairly limited.

Popovers and Split Views

Popovers and split views are forms of interface that exist only on the iPad.

A *popover* (managed by a `UIPopoverController`) is a sort of secondary window or dialog: it displays a view layered on top of the main interface. It is associated, through a sort of arrow, with a view in the main interface, such as the button the user tapped to summon the popover. It does not dim out the rest of the screen, like a presented view whose presentation mode is `UIModalPresentationPageSheet` or `UIModalPresentationFormSheet` (see [Chapter 19](#)). It can be effectively modal, preventing the user from working in the rest of the interface; or it can vanish if the user taps outside it; or you can allow the user to tap some or all of the interface outside it without dismissing the popover.

A popover, in effect, superimposes a roughly iPhone-sized screen on top of the iPad screen, and is useful in part precisely because it brings to the larger iPad the smaller, more lightweight flavor of the iPhone. For example, in my LinkSame app, both the settings view (where the user configures and begins a new game) and the help view (which describes how to play the game) are popovers ([Figure 22-1](#)). On the iPhone, both these views would occupy the entire screen; for each, we'd need a way to navigate to it, and then to return to the main interface when the user is finished with it — both would probably involve presented view controllers. But with the larger iPad screen, that would make no sense; neither view is large enough, or important enough, to occupy the entire screen exclusively. As popovers, these views are shown as what they are: smaller, secondary views which the user summons temporarily and then dismisses.

A *split view* (managed by a `UISplitViewController`) is a combination of two views, the first having the width of an iPhone screen in portrait orientation. When the iPad is in landscape orientation, the two views appear side by side. When the iPad is in portrait orientation, there are two possibilities:

- Only the second view appears, with an option to summon the first view as a popover by tapping a `UIBarButtonItem`.

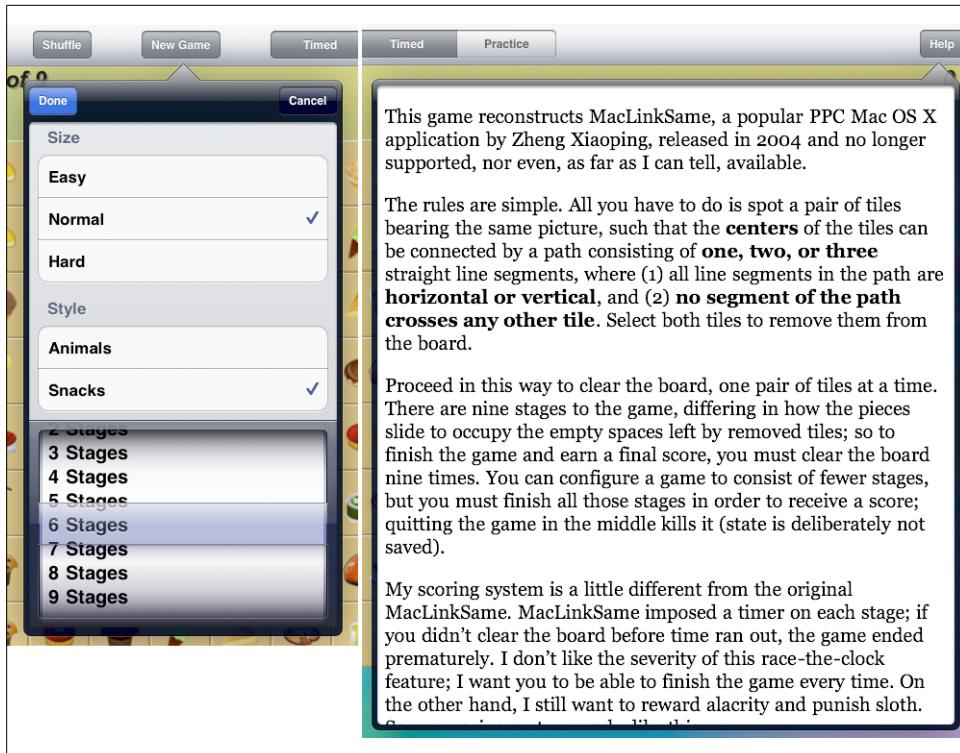


Figure 22-1. Two popovers

- Both views continue to appear side by side; the second view is narrower, because the screen is narrower. (This option is new in iOS 5.)

Like popovers, a split view may be regarded as an evolutionary link between the smaller iPhone interface and the larger iPad interface. On the iPhone, you might have a master–detail architecture in a navigation interface, where the master view is a table view ([Chapter 21](#)). On the iPad, the large screen can accommodate the master view and the detail view *simultaneously*; the split view is a built-in way to do that (and it is no coincidence that its first view is sized to hold the master table that occupied the entire screen on the iPhone).

Split views were important before iOS 5, because `UISplitViewController` was the only legal way in which a single view controller could display the views of two child view controllers side by side. In iOS 5, however, where you are free to design your own custom parent view controllers, `UISplitViewController` will probably diminish in value. Apple's Mail app on the iPad used a `UISplitViewController` in iOS 3.2 and iOS 4, but in iOS 5, it uses a different interface (which I'll try to reverse-engineer later in this chapter).

Configuring and Displaying a Popover

To display a popover, you'll need a UIPopoverController, along with a view controller (UIViewController) whose view the popover will contain. UIPopoverController is not itself a UIViewController subclass. The view controller is the UIPopoverController's `contentViewController`. You'll set this property initially through UIPopoverController's designated initializer, `initWithContentViewController:`. Later, you can swap out a popover controller's view controller (and hence its contained view) by calling `setContentViewController:animated:.`)

Here's how the UIPopoverController for the first popover in [Figure 22-1](#) is initialized. I have a UIViewController subclass, NewGameController. NewGameController's view contains a table ([Figure 21-2](#)) and a UIPickerView ([Chapter 11](#)), and is itself the data source and delegate for both. I instantiate NewGameController and use this instance as the root view controller of a UINavigationController, giving its `navigationItem` a `leftBarButtonItem` (Done) and a `rightBarButtonItem` (Cancel). I don't really intend to do any navigation; I just want the two buttons, and this is an easy way of getting them into the interface. That UINavigationController then becomes a UIPopoverController's view controller:

```
NewGameController* dlg = [[NewGameController alloc] init];
UIBarButtonItem* b = [[UIBarButtonItem alloc]
    initWithBarButtonSystemItem: UIBarButtonItemSystemItemCancel
        target: self
        action: @selector(cancelNewGame:)];
dlg.navigationItem.rightBarButtonItem = b;
b = [[UIBarButtonItem alloc]
    initWithBarButtonSystemItem: UIBarButtonItemSystemItemDone
        target: self
        action: @selector(saveNewGame:)];
dlg.navigationItem.leftBarButtonItem = b;
UINavigationController* nav =
    [[UINavigationController alloc] initWithRootViewController:dlg];
UIPopoverController* pop =
    [[UIPopoverController alloc] initWithContentViewController:nav];
```

Note that that code doesn't cause the popover to appear on the screen! I'll come to that in a moment.

The popover controller should also be told the size of the view it is to display, which will be the size of the popover. The default popover size is {320, 1100}; Apple would like you to stick to the default width of 320 (the width of an iPhone screen), but a maximum width of 600 is permitted, and the second popover in [Figure 22-1](#) uses it. The popover's height might be shorter than requested if there isn't enough vertical space, so the view to be displayed needs to be prepared for the possibility that it might be resized.

You can provide the popover size in one of two ways:

UIPopoverController's popoverContentSize property

This property can be set before the popover appears; it can also be changed while the popover is showing, with `setPopoverContentSize:animated:`.

UIViewController's contentSizeForViewInPopover property

The UIViewController is the UIPopoverController's `contentViewController` (or is contained by that view controller, as in a tab bar interface or navigation interface). This approach often makes more sense, because a UIViewController will generally know its own view's ideal size. If a view controller is instantiated in a nib or storyboard, this value can be set in the nib or storyboard.

In the case of the first popover in [Figure 22-1](#), the NewGameController sets its own `contentSizeForViewInPopover` in `viewDidLoad`; its popover size is simply the size of its view:

```
self.contentSizeForViewInPopover = self.view.bounds.size;
```

The popover itself, however, will need to be somewhat taller, because the NewGameController is embedded in a UINavigationController, whose navigation bar occupies additional vertical space. Delightfully, the UINavigationController takes care of that automatically; its own `contentSizeForViewInPopover` adds the necessary height to that of its root view controller.

In case of a conflict, the rule seems to be that if the UIPopoverController and the UIViewController have different settings for their respective content size properties at the time the popover is initially displayed, the UIPopoverController's setting wins. But once the popover is visible, if *either* property is changed, the change is obeyed; specifically, my experiments suggest that if the UIViewController's `contentSizeForViewInPopover` is changed (not merely set to the value it already has), the UIPopoverController adopts that value as its `popoverContentSize` and the popover's size is adjusted accordingly.

If a popover's `contentViewController` is a UINavigationController, and a view controller is pushed onto or popped off of its stack, then if the current view controller's `contentSizeForViewInPopover` differs from that of the previously displayed view controller, my experiments suggest that the popover's *width* will change to match the new width, but the popover's *height* will change only if the new height is *taller*. This feels like a bug. A workaround is to implement the UINavigationController's delegate method `navigationController:didShowViewController:animated:`, so as to set the navigation controller's `contentSizeForViewInPopover` explicitly:

```
- (void)navigationController:(UINavigationController *)navigationController  
    didShowViewController:(UIViewController *)viewController  
        animated:(BOOL)animated {  
    navigationController.contentSizeForViewInPopover =  
        viewController.contentSizeForViewInPopover;  
}
```

(That workaround is not entirely satisfactory from a visual standpoint, as two animations succeed one another, but I tried implementing `willShowViewController...` instead and liked the results even less.)

The popover is made to appear on screen by sending the UIPopoverController one of the following messages (and the UIPopoverController's `popoverVisible` property then becomes YES):

- `presentPopoverFromRect:inView:permittedArrowDirections:animated:`
- `presentPopoverFromBarButtonItem:permittedArrowDirections:animated:`

The popover has a sort of triangular bulge (called its *arrow*) on one edge, pointing to some region of the existing interface, from which the popover thus appears to emanate and to which it seems to be related. The difference between the two methods lies only in how this region is specified. With the first method, you can provide any CGRect with respect to any visible UIView's coordinate system; for example, to make the popover emanate from a UIButton, you could provide the UIButton's frame with respect to its superview, or (better) the UIButton's bounds with respect to itself. But you can't do that with a UIBarButtonItem, because a UIBarButtonItem isn't a UIView and doesn't have a frame or bounds; hence the second method is provided.

The permitted arrow directions restrict which sides of the popover the arrow can appear on. It's a bitmask, and your choices are:

- `UIPopoverArrowDirectionUp`
- `UIPopoverArrowDirectionDown`
- `UIPopoverArrowDirectionLeft`
- `UIPopoverArrowDirectionRight`
- `UIPopoverArrowDirectionAny`

Usually, you'd specify `UIPopoverArrowDirectionAny`, allowing the runtime to put the arrow on whatever side it feels is appropriate.

Even if you specify a particular arrow direction, you still have no precise control over a popover's location. Starting in iOS 5, however, you do get some veto power: set the UIPopoverController's `popoverLayoutMargins` to a UIEdgeInsets stating the margins, with respect to the root view bounds, within which the popover must appear. If an inset that you give is so large that the arrow can no longer touch the presenting rect, it may be ignored, or the arrow may become disconnected from its presenting rect; you probably should try not to do that.

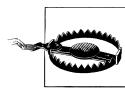
Starting in iOS 5, you can customize much of a popover's appearance. For example, the first popover in [Figure 22-1](#) has a dark navigation bar even though no such thing was requested when the UINavigationController was created. This is because a popover whose content view controller is a navigation controller likes to take control of its navigation bar's `barStyle` and set it to a special undocumented style, evidently to make



Figure 22-2. A very silly popover

it harmonize with the popover’s border. There isn’t much you can do about this: setting the navigation bar’s `tintColor` has no effect, and setting its `backgroundColor` is effective but ugly. But in iOS 5, you can take charge of a `UINavigationBar`’s background image and you can customize the position and appearance of its bar button items ([Chapter 25](#)), so you can achieve some pleasing results.

Moreover, in iOS 5 you can customize the outside of the popover—that is, the “frame” and the arrow. To do so, you set the `UIPopoverController`’s `popoverBackgroundViewClass` to your subclass of `UIPopoverBackgroundView` (a `UIView` subclass)—at which point you can achieve just about anything you want, including the very silly popover shown in [Figure 22-2](#). You get to dictate (by implementing `contentViewInset`) the thickness of all four sides of the “frame”.



In my experiments, I found it necessary to import `<UIKit/UIPopoverBackgroundView.h>` into the header file of my `UIPopoverBackgroundView` subclass. This feels like a bug; to subclass a class that is part of `UIKit`, importing `UIKit` should be sufficient.

Configuring your `UIPopoverBackgroundView` subclass is a bit tricky, because this single view is responsible for drawing both the arrow and the “frame”. Thus, in a complete and correct implementation, you’ll have to draw differently depending on the arrow direction, which you can learn from the `UIPopoverBackgroundView`’s `arrowDirection` property. I’ll give a simplified example in which I assume that the arrow direction will be `UIPopoverArrowDirectionUp`. Then drawing the “frame” is easy: here,

I divide the view's overall rect into two areas, the arrow area on top (its height is a #defined constant, ARHEIGHT) and the “frame” area on the bottom, and draw the “frame” into the bottom area as a resizable image ([Chapter 15](#)):

```
UIImage* linOrig = [UIImage imageNamed: @"linen.png"];
CGFloat capw = linOrig.size.width / 2.0 - 1;
CGFloat caph = linOrig.size.height / 2.0 - 1;
UIImage* lin = [linOrig resizableImageWithCapInsets:
    UIEdgeInsetsMake(caph, capw, caph, capw)];
CGRect arrow;
CGRect body;
CGRectDivide(rect, &arrow, &body, ARHEIGHT, CGRectGetMinYEdge);
[lin drawInRect:body];
```

The documentation claims that the popover will be awarded a shadow automatically, but in my tests this was not the case. So I also apply my own shadow around the “frame” shape:

```
self.layer.shadowPath = CGPathCreateWithRect(body, NULL);
self.layer.shadowColor = [UIColor grayColor].CGColor;
self.layer.shadowRadius = 20;
self.layer.shadowOpacity = 0.4;
```

Now for the arrow. The UIPopoverBackgroundView has an `arrowHeight` property and an `arrowBase` property that you've set to describe the arrow dimensions to the runtime. (In my code, their values are provided by two #defined constants, ARHEIGHT and ARBASE.) My arrow will consist simply of a texture-filled isosceles triangle, with an excess base consisting of a rectangle to make sure it's attached to the “frame”. (In my real code I draw the arrow first, so that its excess is behind the “frame”.) The UIPopoverBackgroundView also has an `arrowOffset` property that the runtime has set to tell you where to draw the arrow: this offset measures the positive distance between the center of the view and the center of the arrow. However, the runtime will have no hesitation in setting the `arrowOffset` all the way at the edge of view, or even beyond its bounds (in which case it won't be drawn); to prevent this, I provide a maximum offset limit:

```
CGContextRef con = UIGraphicsGetCurrentContext();
CGContextSaveGState(con);
CGFloat proposedX = self.arrowOffset;
CGFloat limit = 22.0;
CGFloat maxX = rect.size.width/2.0 - limit;
if (proposedX > maxX)
    proposedX = maxX;
if (proposedX < limit)
    proposedX = limit;
CGContextTranslateCTM(con, rect.size.width/2.0 + proposedX - ARBASE/2.0, 0);
CGContextMoveToPoint(con, 0, ARHEIGHT);
CGContextAddLineToPoint(con, ARBASE / 2.0, 0);
CGContextAddLineToPoint(con, ARBASE, ARHEIGHT);
CGContextClosePath(con);
CGContextAddRect(con, CGRectMake(0,ARHEIGHT,ARBASE,15));
CGContextClip(con);
[lin drawAtPoint:CGPointMake(-40,-40)];
CGContextRestoreGState(con);
```

Managing a Popover

Unlike a presented view controller or a child view controller, a UIPopoverController instance is not automatically retained for you by some presenting view controller or parent view controller; you must retain it yourself. If you fail to do this, then if the UIPopoverController goes out of existence while its popover is on the screen, your app will crash (with a helpful message: “-[UIPopoverController dealloc] reached while popover is still visible”). Also, you might need the retained reference to the UIPopoverController later, when the time comes to dismiss the popover.

There are actually two ways in which a popover can be dismissed: the user can tap outside the popover, or you can explicitly dismiss the popover (as I do with the first popover in [Figure 22-1](#) when the user taps the Done button or the Cancel button). In order to dismiss the popover explicitly, you send its UIPopoverController the `dismissPopoverAnimated:` message. Obviously, then, you need a reference to the UIPopoverController.

Even if a popover is normally dismissed automatically by the user tapping outside it, you *still* might want to dismiss it explicitly on certain occasions — so you still might need a reference to the popover controller. For example, in keeping with the transient nature of popovers, I like to dismiss the current popover programmatically when the application undergoes certain strong transitions, such as going into the background or being rotated. (See also Apple’s technical note on what to do when the interface rotates while a popover is showing, QA1694, “Handling Popover Controllers During Orientation Changes.”) You can listen for the former by registering for `UIApplicationDidEnterBackgroundNotification`, and for the latter by implementing `willRotateToInterfaceOrientation:duration:`. This policy is not merely aesthetic; some view controllers, especially certain built-in specialized view controllers, recover badly from such transitions when displayed in a popover.

The obvious solution is an instance variable or property with a strong (retain) policy. The question then is how many such instance variables to use if the app is going to be displaying more than one popover. We could have one instance variable for *each* popover controller. On the other hand, a well-behaved app, in accordance with Apple’s interface guidelines, is probably never going to display more than one popover simultaneously; so a *single* UIPopoverController instance variable (we might call it `currentPop`) should suffice. This one instance variable could be handed a reference to the current popover controller each time we present a popover; using that reference, we would be able later to dismiss the current popover and release its controller.

Dismissing a Popover

An important feature of a popover’s configuration is whether and to what extent the user can operate outside it without automatically dismissing it. There are two aspects to this configuration:

UIViewController's `modalInPopover` property

If this is YES for the popover controller's view controller (or for its current child view controller, as in a tab bar interface or navigation interface), the popover is absolutely modal; any tap outside it will be ignored and won't have any effect at all, not even to dismiss the popover. The default is NO.

UIPopoverController's `passthroughViews` property

This matters only if `modalInPopover` is NO. It is an array of views in the interface behind the popover; the user can interact with these views, but a tap anywhere else outside the popover will dismiss it (with no effect on the thing tapped). If `passthroughViews` is nil, a tap *anywhere* outside the popover will dismiss it.



Setting a UIPopoverController's `passthroughViews` might not have any effect unless the popover is already showing (the UIPopoverController has been sent `presentPopover...`).

A popover can present a view controller internally; you'll specify a `modalPresentationStyle` of `UIModalPresentationCurrentContext`, because otherwise the presented view will be fullscreen by default. You'll also specify a transition style of `UIModalTransitionStyleCoverVertical` — with any other transition style, your app will crash with this message: “Application tried to present inside popover with transition style other than `UIModalTransitionStyleCoverVertical`.“ The presented view controller's `modalInPopover` is automatically set to YES: that is, while the presented view controller is being presented within the popover, the user can't make anything happen by tapping outside the popover, not even to dismiss the popover. (You can subvert this by setting the presented view controller's `modalInPopover` to NO *after* it is presented, but you probably shouldn't.)

If `modalInPopover` is NO, you should pay attention to the `passthroughViews`, as the default behavior may be undesirable. For example, if a popover is summoned by the user tapping a UIBarButtonItem item in a toolbar using `presentPopoverFromBarButtonItem:...`, the entire toolbar is a passthrough view; this means that the user can tap any button in the toolbar, *including the button that summoned the popover*. The user can thus by default summon the popover again while it is still showing, which is certainly not what you want. I like to set the `passthroughViews` to nil; at the very least, while the popover is showing, you should probably disable the UIBarButtonItem that summoned the popover.

We are now ready for a rigorous specification of the two ways in which a popover can be dismissed:

- The popover controller's view controller's `modalInPopover` is NO, and the user taps outside the popover on a view not listed in the popover controller's `passthroughViews`. The UIPopoverController's delegate (adopting the UIPopoverController-Delegate protocol) is sent `popoverControllerShouldDismissPopover:`; if it doesn't

- return NO (which might be because it doesn't implement this method), the popover is dismissed, and the delegate is sent `popoverControllerDidDismissPopover:`.
- The UIPopoverController is sent `dismissPopoverAnimated:` by your code; the delegate methods are *not* sent in that case. Typically this would be because you've included some interface item *inside* the popover that the user can tap to dismiss the popover (like the Done and Cancel buttons in the first popover in [Figure 22-1](#)).

Because a popover can be dismissed in two different ways, if you have a cleanup task to perform as the popover vanishes, you may have to see to it that this task is performed under two different circumstances. That can get tricky.

To illustrate, I'll describe what happens when the first popover in [Figure 22-1](#) is dismissed. Within this popover, the user is interacting with several settings in the user defaults. But if the user cancels, or if the user taps outside the popover (which I take to be equivalent to canceling), I want to revert those defaults to the way they were before the popover was summoned. So, as I initially present the popover, I preserve the relevant current user defaults as an ivar:

```
// save defaults so we can restore them later if user cancels
self.oldDefs = [[NSUserDefaults standardUserDefaults] dictionaryWithValuesForKeys:
    [NSArray arrayWithObjects:@"Style", @"Size", @"Stages", nil]];
```

Now, if the user taps Save, the user's settings within the popover have *already* been saved (in the user defaults), so I explicitly dismiss the popover and proceed to initiate the new game that the user has asked for. On the other hand, if the user taps Cancel, I must revert the user defaults as I dismiss the popover:

```
- (void)cancelNewGame: (id) sender { // cancel button in New Game popover
    [self.currentPop dismissPopoverAnimated:YES];
    self.currentPop = nil;
    [[NSUserDefaults standardUserDefaults]
        setValuesForKeysWithDictionary:self.oldDefs];
}
```

But I must also do the same thing if the user taps outside the popover. Therefore I implement the delegate method and revert the user defaults there as well:

```
- (void)popoverControllerDidDismissPopover:(UIPopoverController *)pc {
    [[NSUserDefaults standardUserDefaults]
        setValuesForKeysWithDictionary:self.oldDefs];
    self.currentPop = nil;
}
```

There is a problem with the foregoing implementation, however. My app, you may remember, has *another* popover (the second popover in [Figure 22-1](#)). This popover, too, can be dismissed by the user tapping outside it; in fact, that's the only way the user can dismiss it. This means that `popoverControllerDidDismissPopover:` will be called. But now we *don't* want to call `setValuesForKeysWithDictionary:`; it's the wrong popover, and we have no preserved defaults to revert. This means that I must somehow test for *which* popover controller is being passed in as the parameter to `popoverController-`

`DidDismissPopover`: But how can I distinguish one popover controller from another? Luckily, my popover controllers have different types of view controller:

```
- (void)popoverControllerDidDismissPopover:(UIPopoverController *)pc {
    if ([pc.contentViewController isKindOfClass: [UINavigationController class]])
        [[NSUserDefaults standardUserDefaults]
            setValuesForKeysWithDictionary:self.oldDefs];
    self.currentPop = nil;
}
```

If this were not the case — for example, if I had two different popovers each of which had a UINavigationController as its view controller — I'd need some other way of distinguishing them. This is rather a knotty problem, and in the past I've resorted to various desperate measures to resolve it, such as subclassing UIPopoverController.

Earlier, I mentioned that I also want to dismiss any currently displayed popover if the interface rotates, or if the app goes into the background. That means that I must perform these same tests *again* in response to the appropriate notifications:

```
- (void)willRotateToInterfaceOrientation:(UIInterfaceOrientation)io
    duration:(NSTimeInterval)duration {
    UIPopoverController* pc = self.currentPop;
    if (pc) {
        if ([pc.contentViewController isKindOfClass:[UINavigationController class]])
            [[NSUserDefaults standardUserDefaults]
                setValuesForKeysWithDictionary:self.oldDefs];
        [pc dismissPopoverAnimated:NO];
        self.currentPop = nil; // wrong in previous version
    }
}

-(void)backgrounding:(id)dummy {
    UIPopoverController* pc = self.currentPop;
    if (pc) {
        if ([pc.contentViewController isKindOfClass:[UINavigationController class]])
            [[NSUserDefaults standardUserDefaults]
                setValuesForKeysWithDictionary:self.oldDefs];
        [pc dismissPopoverAnimated:NO];
        self.currentPop = nil; // wrong in previous version
    }
}
```

In my view, the need for all this work, all this testing and caution and duplication of functionality, just to display a couple of popovers, demonstrates that the framework's implementation of popover controllers is deeply flawed. You don't get sufficient help with getting a reference to a UIPopoverController from a view currently being displayed within it, with managing a popover controller's memory, or with distinguishing one popover controller from another. I've shown how you can work around these shortcomings, but in a better world such workarounds wouldn't be necessary. (For example, because only one popover is supposed to be showing at a time, the framework could just maintain a reference to its controller for you.) Popovers were invented for iOS 3.2;

frankly, I'm astonished that we've reached iOS 5.0 with no improvement in this area of the API.

Popover Segues

In an iPad storyboard, a segue can be designated a popover segue, by choosing Popover from the Style pop-up menu in the Attributes inspector. The consequences of doing so are:

- When the segue is triggered, a popover is displayed. The runtime constructs a UIPopoverController and makes the segue's destination view controller the UIPopoverController's view controller. The popover's "anchor" (the view or bar button item to which its arrow points) is the source object from which you control-drag to form the segue, or it can be set in the Attributes inspector.
- The segue is a UIStoryboardPopoverSegue, a UIStoryboardSegue subclass that adds a single read-only property, `popoverController`. You can use this, for instance in `prepareForSegue:sender:`, to customize the popover controller.

The UIPopoverController created by the triggering of a popover segue is retained behind the scenes; the app does not crash if you fail to retain it explicitly yourself. (I don't know what object is retaining it, but I think it's the storyboard.) Nevertheless, you will *still* probably want to retain your own reference to the popover controller, because you're still likely to need that reference for the reasons discussed in the preceding section. You'll probably obtain that reference in your `prepareForSegue:sender:` implementation.

Popover segues sound tempting, but they do not appreciably reduce the amount of code required to configure and manage a popover. On the contrary, in my experience they tend to increase it. Consider, for example, the code I cited earlier for creating a popover controller whose view controller is a navigation view controller:

```
NewGameController* dlg = [[NewGameController alloc] init];
UIButtonBarItem* b = [[UIButtonBarItem alloc]
    initWithTitleBarButtonSystemItem: UIBarButtonSystemItemCancel
        target: self
        action: @selector(cancelNewGame:)];
dlg.navigationItem.rightBarButtonItem = b;
b = [[UIButtonBarItem alloc]
    initWithTitleBarButtonSystemItem: UIBarButtonSystemItemDone
        target: self
        action: @selector(saveNewGame:)];
dlg.navigationItem.leftBarButtonItem = b;
UINavigationController* nav =
    [[UINavigationController alloc] initWithRootViewController:dlg];
UIPopoverController* pop =
    [[UIPopoverController alloc] initWithContentViewController:nav];
```

In a storyboard, you could create a popover segue to a navigation controller whose root view is a NewGameController. But now how would you hook up the NewGameCon-

troller's bar button items to make their actions call the correct methods in `self`? In `prepareForSegue:sender:`, you'd have to work your way down from the navigation controller to the NewGameController; then you'd need either to create the bar button items in code, exactly as before, or to hook up their target and action if they were instantiated from the storyboard:

```
UINavigationController* nav = segue.destinationViewController;
UIViewController* vc = [nav.childViewControllers objectAtIndex:0];
vc.navigationItem.leftBarButtonItem.target = self;
vc.navigationItem.leftBarButtonItem.action = @selector(savePop1:);
vc.navigationItem.rightBarButtonItem.target = self;
vc.navigationItem.rightBarButtonItem.action = @selector(cancelPop1:);
```

Similarly, how would you set a popover controller's `passthroughViews` to `nil` when the popover controller is generated by a popover segue? You can't do that in the storyboard, so you'd have to do it in `prepareForSegue:sender:`. But, as I've already said, `prepareForSegue:sender:` is too soon; you need to do it *after* the popover has been displayed. So you'd have to use delayed performance:

```
UIStoryboardPopoverSegue* seg = (id)segue;
UIPopoverController* pop = seg.popoverController;
dispatch_time_t popTime = dispatch_time(DISPATCH_TIME_NOW, 1);
dispatch_after(popTime, dispatch_get_main_queue(), ^{
    pop.passthroughViews = nil;
});
```

Finally, as I mentioned earlier, you'd still need to maintain a reference to the UIPopoverController (`pop` in the above code), set yourself as its delegate, and configure any other features of the popover controller such as its `popoverLayoutMargins` and `popoverBackgroundViewClass`, just as you would have done if there were no storyboard at all. I find the code needed to deal with a UIStoryboardPopoverSegue's popover controller confusing and fragile; I'd rather create the popover controller in code to begin with.

Automatic Popovers

In a few situations, the framework will automatically create and display a popover for you. I am not personally fond of this behavior; the advantages of the automatic behavior (such as the fact that you do not have to provide code to create the popover) are counterbalanced, in my view, by the disadvantages — in particular, you can't get access to the UIPopoverController, so the behavior and appearance of the popover is completely out of your hands.

One such situation is what happens when a search bar (a UISearchBar) tied to a search display controller (UISearchDisplayController) appears in a toolbar (UIToolbar) on the iPad. Recall the example of search display controller code from [Chapter 21](#), where we search a list of the 50 United States; I'll modify that example to demonstrate. In the nib editor, start with a toolbar at the top of the root view, and drag into it the combined Search Bar and Search Display Controller object from the Object library. As before,

this simple move causes a whole bunch of outlets to be configured automatically. In particular:

- The search bar's delegate is the File's Owner.
- The File's Owner's `searchDisplayController` is the search display controller. This is a `UIViewController` property that I didn't mention in [Chapter 21](#), because its worth is not clear when a `UISearchDisplayController` is created and configured in code. When a `UISearchDisplayController` is instantiated from a nib, however, this property is an outlet that retains the search display controller, as well as providing access to it in code.
- The search display controller's search bar is the search bar.
- The search display controller's delegate, `searchContentsController`, `searchResultsDataSource`, and `searchResultsDelegate` is the File's Owner. Of these, only the latter two appear to be of importance in this example.

Now for the code. When our view controller loads its view, we also load the model (the list of states) into an `NSArray` property called `states`. We also have an `NSArray` property called `filteredStates`. Here is the code for dealing with the search bar and the search display controller's results table:

```
- (NSInteger)numberOfSectionsInTableView:(UITableView *)tableView {
    return 1;
}

- (NSInteger)tableView:(UITableView *)tableView
    numberOfRowsInSection:(NSInteger)section {
    return [self.filteredStates count];
}

- (UITableViewCell *)tableView:(UITableView *)tableView
    cellForRowAtIndexPath:(NSIndexPath *)indexPath {
    static NSString *CellIdentifier = @"Cell";
    UITableViewCell *cell =
        [tableView dequeueReusableCellWithIdentifier:CellIdentifier];
    if (cell == nil)
        cell = [[UITableViewCell alloc] initWithStyle:UITableViewCellStyleDefault
                                       reuseIdentifier:CellIdentifier];
    cell.textLabel.text = [self.filteredStates objectAtIndex:indexPath.row];
    return cell;
}

- (void)filterData {
    NSString* target = self.searchDisplayController.searchBar.text;
    NSPredicate* p = [NSPredicate predicateWithBlock:
        ^(id obj, NSDictionary *d) {
            NSString* s = obj;
            NSStringCompareOptions options = NSCaseInsensitiveSearch;
            BOOL b = [s rangeOfString:target options:options].location != NSNotFound;
            return b;
        }];
    self.filteredStates = [states filteredArrayUsingPredicate:p];
```

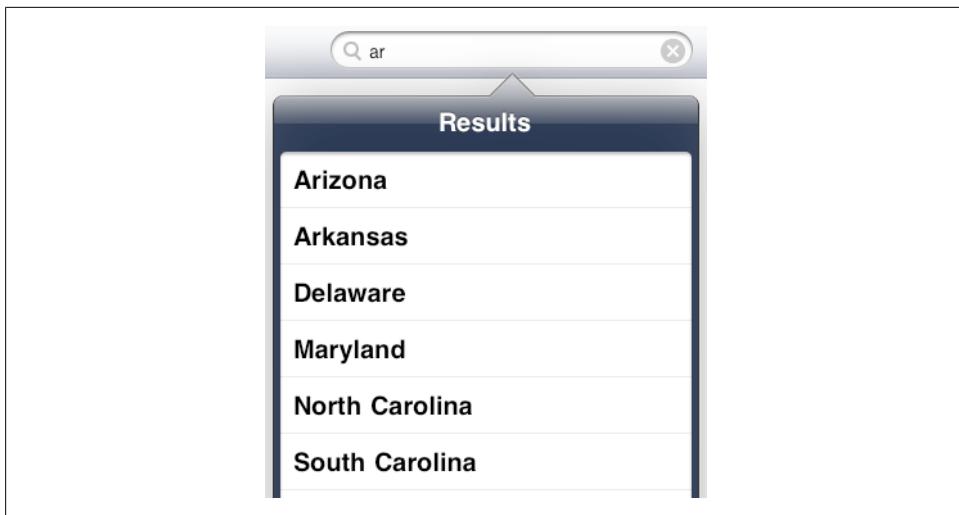


Figure 22-3. An automatically created search results popover

```
}

- (void)searchBar:(UISearchBar *)searchBar textDidChange:(NSString *)searchText {
    [self filterData];
}
```

That's all. There is no creation of a UITableView and no mention anywhere of a UIPopoverController. Nevertheless, when the user enters text in the search bar, a popover appears, containing a table of search results (Figure 22-3). Note the "Results" title at the top of the popover; starting in iOS 5, you can customize that title by setting the UISearchDisplayController's `searchResultsTitle` property (you can do this in the nib editor). However, as I mentioned before, you get no official access to this popover's controller, so you can't set its passthrough views and so on.

Another example of an automatic popover on the iPad is the alert sheet, discussed in [Chapter 26](#).

Split Views

A split view is implemented through a UISplitViewController (a UIViewController subclass) whose children are the two UIViewControllers whose views are to be displayed in the two regions of the split view. You provide the children through the UISplitViewController's `viewControllers` property (an NSArray); it can be configured in code or in a nib. A UIViewController that is a child, at any depth, of a UISplitViewController has a reference to the UISplitViewController through its `splitViewController` property.

There is very little work for you to do with regard to a split view. A split view controller has no further properties or methods beyond those inherited from `UIViewController`. You can hear about what the split view is doing through its delegate (adopting the `UISplitViewControllerDelegate` protocol), which receives these messages:

`splitViewController:willHideViewController:withBarButtonItem:forPopoverController:`

The split view is rotating to portrait orientation, so it's hiding the first view. The split view creates a `UIBarButtonItem` and hands it to you as the third parameter. Your mission, should you decide to accept it, is to put that bar button item into the interface, typically in a toolbar at the top of the root view. You are free to configure the bar button item's title and image as well.

What you do with the other parameters is up to you, but none of them are needed; the split view has already set things up so that if you do put this bar button item into the interface, then if the user taps it, a popover will be presented through the popover controller (fourth parameter) displaying the view of the view controller (second parameter). It's common practice to keep a reference to the popover controller, in case you need it in order to dismiss the popover later (but don't set its delegate).

`splitViewController:popoverController:willPresentViewController:`

The user has tapped the bar button item you were handed in the first delegate method, and the popover is about to appear. You probably won't need to implement this method.

`splitViewController:willShowViewController:invalidatingBarButtonItem:`

This is the opposite of the first delegate method: the split view is rotating to landscape orientation, so it's going to break the connection between the bar button item and the popover controller and is going to put the first view back into the interface. You should remove the bar button item from the interface (the split view will not do that for you).

As an app with a split view interface launches, the `willHide` delegate method is called; if the device is being held in landscape orientation at launch time, and if the split view is free to rotate, it then immediately rotates to landscape orientation and the `willShow` delegate method is called.

The only even slightly tricky part of all this lies hidden in the words "put that `UIBarButtonItem` into the interface" and "remove the `UIBarButtonItem` from the interface." How you do this depends on your interface. Let's take the case of a toolbar. `UIToolbar` has no method for inserting or removing a single bar button item; you have to set its entire array of items at once. So, for example, to put a `UIBarButtonItem` at the left end of a toolbar, you'll set the toolbar's `items` array to an array composed of the `UIBarButtonItem` and the items from the existing `items` array. You will probably set the toolbar's `items` array by calling `setItems:animated:` in order to get the animation. Removing a `UIBarButtonItem` from a toolbar is similar. So:

```

NSArray* arr = [[NSArray arrayWithObject: barButtonItem]
    arrayByAddingObjectsFromArray: self.toolbar.items];
[self.toolbar setItems:arr animated:YES];

```

On the other hand, the Xcode 4.2 version of the iPad Master–Detail Application project template, which demonstrates a split view interface, uses a navigation bar instead of a toolbar:

```

DetailViewController *detailViewController =
    [[DetailViewController alloc] initWithNibName:@"DetailViewController" bundle:nil];
UINavigationController *detailNavController =
    [[UINavigationController alloc] initWithRootViewController:detailViewController];

```

As a result, the DetailViewController instance, functioning as the UISplitViewController’s delegate, has an easy time getting the bar button item into and out of the interface: it sets its own `navigationItem.leftBarButtonItem`.

If you ask for a main storyboard as you generate your iPad project from the Master–Detail Application project template, you can see how little code is required to implement a working split view. The split view controller is hooked to its child view controllers by relationships in the storyboard; those child view controllers are both navigation controllers (mostly in order to get the navigation bar), and they are hooked to their root view controllers by relationships in the storyboard as well. Thus, all five view controllers are instantiated together, automatically, as the app launches.

The only thing that can’t be set in the storyboard is the delegate relationship between the UISplitViewController and the DetailViewController (they are both in the storyboard together, but they’re in different scenes, so no outlet can be drawn between them; see my rant on this topic in [Chapter 19](#)). To solve this, the template includes the following code in the app delegate’s `application:didFinishLaunchingWithOptions:`:

```

UISplitViewController *splitViewController =
    (UISplitViewController *)self.window.rootViewController;
UINavigationController *navigationController =
    [splitViewController.viewControllers lastObject];
splitViewController.delegate = (id)navigationController.topViewController;

```

By the time that code runs, the window’s `rootViewController` has been set and all the view controllers are in place, in their various parent–child relationships; the code is thus able to work its way through the parent–child hierarchy ([Chapter 13](#)) to get references to the two desired view controllers and can make the one the delegate of the other.

The only other code of interest is as follows:

- Both the master view controller and the detail view controller permit rotation, returning YES from `shouldAutorotateToInterfaceOrientation:`. Otherwise, the split view itself won’t permit rotation.
- The DetailViewController implements the two delegate methods to add and remove the UIBarButtonItem in the toolbar. It also retains a reference to the popover when it’s showing.

```

- (void)splitViewController:(UISplitViewController *)splitController
    willHideViewController:(UIViewController *)viewController
    withBarButtonItem:(UIBarButtonItem *)barButtonItem
    forPopoverController:(UIPopoverController *)popoverController {
    [self.navigationItem setLeftBarButtonItem:barButtonItem animated:YES];
    self.masterPopoverController = popoverController;
}

- (void)splitViewController:(UISplitViewController *)splitController
    willShowViewController:(UIViewController *)viewController
    invalidatingBarButtonItem:(UIBarButtonItem *)barButtonItem {
    [self.navigationItem setLeftBarButtonItem:nil animated:YES];
    self.masterPopoverController = nil;
}

```

The split view interface is now working. In landscape orientation, the master view (a table view) appears on the left, with its navigation bar at the top, and the detail view appears on the right, with its navigation bar at the top. In portrait orientation, the master view vanishes; the detail view, with its navigation bar, occupies the entire screen. A UIBarButtonItem appears at the navigation bar's left end; tapping it summons the master view in a popover.

Starting in iOS 5, it is also possible for a split view interface *not* to hide the left view in portrait orientation. Instead, the left and right view both appear in both orientations; the left view's width is unchanged, while the right view is resized appropriately. (Compare Apple's Settings app on the iPad.) To get that behavior, implement this delegate method:

`splitViewController:shouldHideViewController:inOrientation:`

Allows the left view controller to be hidden (return YES) or not (return NO) as the interface rotates to the given orientation. If you return NO, the other delegate methods won't be called in this orientation; there will be no bar button item and no popover.

Alternatively, seeing that iOS 5 permits it, you might dispense with UISplitViewController altogether and write your own parent view controller class. It isn't difficult to write a class that does what UISplitViewController does, and it isn't difficult to write a class that does something similar but preferable.

For example, consider Apple's Mail app in iOS 5. It's a split view interface, in the sense that there's a narrow master view and a wider detail view that occupy the screen together in landscape orientation. But in portrait orientation, there's no popover. Instead, the user taps a bar button item or swipes the detail view, and the master view slides in from the left, sitting on top of the detail view, and behaves modally — that is, the user can work only in the master view until either it is explicitly dismissed by tapping elsewhere or the device is rotated to landscape orientation so that the master view resumes its place beside the detail view.

Here's some code that implements such an interface. In the parent view controller, we have references to the two child view controllers as `vc1` and `vc2`. Our `loadView` imple-

mentation is called in portrait orientation ([Chapter 19](#)), so we position `vc1`'s view off-screen and make `vc2`'s view occupy the entire interface:

```
UIView* v =
    [[UIView alloc] initWithFrame: [[UIScreen mainScreen] applicationFrame]];
self.view = v;
v.backgroundColor = [UIColor blackColor];
UIView* v1 = self.vc1.view;
UIView* v2 = self.vc2.view;
CGRect frame1 = CGRectMake(-320, 0, 320, self.view.bounds.size.height);
CGRect frame2 = self.view.bounds;
// ... some code omitted ...
[v addSubview: leftView];
[v addSubview: rightView];
```

In the omitted code, I embed `v1` in `leftView` and `v2` in `rightView` (which are instances of my own `UIView` subclasses) so as to get the rounded corners and shadow-casting behavior of the Mail app; I also tag `leftView` and `rightView` with values 1 and 2 so that I can find them in the interface later. When the user rotates the device, I respond by rearranging the interface as necessary, with animation:

```
-(void)willAnimateRotationToInterfaceOrientation:(UIInterfaceOrientation)io
                                         duration:(NSTimeInterval)d
{
    [UIView animateWithDuration:d animations:^{
        if (UIInterfaceOrientationIsPortrait(io)) {
            [self.view viewWithTag: 1].frame =
                CGRectMake(-320, 0, 320, self.view.bounds.size.height);
            [self.view viewWithTag: 2].frame = self.view.bounds;
        } else {
            CGRect frame1;
            CGRect frame2;
            CGRectDivide(self.view.bounds, &frame1, &frame2, 320, CGRectGetMinXEdge);
            [self.view viewWithTag: 1].frame = frame1;
            [self.view viewWithTag: 2].frame = frame2;
        }
    } completion:^(BOOL b) {
        ((ShadowView*)[self.view viewWithTag: 1]).showsShadow = NO;
    }];
}
```

The detail view controller, which sits inside a navigation controller in order to get the navigation bar, responds to rotation by showing or hiding a bar button item whose action method calls the split view controller's `showView1` method:

```
- (void)viewDidLoad {
    [super viewDidLoad];
    UIBarButtonItem* b = [[UIBarButtonItem alloc] initWithTitle:@"Master"
        style:UIBarButtonItemStyleBordered
        target:self
        action:@selector(doMasterButton:)];
    self.btn = b;
    [self.navigationItem setLeftBarButtonItem:self.btn animated:NO];
}
```

```

- (BOOL)shouldAutorotateToInterfaceOrientation:(UIInterfaceOrientation)io {
    // Return YES for supported orientations
    return YES;
}

- (void)willAnimateRotationToInterfaceOrientation:(UIInterfaceOrientation)io
                                         duration:(NSTimeInterval)duration {
    if (UIInterfaceOrientationIsPortrait(io))
        [self.navigationItem setLeftBarButtonItem:self.btnExit animated:NO];
    else
        [self.navigationItem setLeftBarButtonItem:nil animated:NO];
}

- (void) doMasterButton: (id) sender {
    UIViewController* vc = self.parentViewController;
    while (![vc isKindOfClass: [MySplitViewController class]])
        vc = vc.parentViewController;
    [(MySplitViewController*)vc showView1];
}

```

Now we come to the really interesting part. When the user taps the bar button item configured by the detail view controller, the split view controller animates the movement of the master view onto the screen, and puts a screen-sized invisible view with a tap gesture recognizer right behind it, to enforce the modal behavior of the master view:

```

- (void) showView1 {
    UIView* v = [self.view viewWithTag:1];
    [self.view bringSubviewToFront:v];
    // insert modal-enforcing tap-detecting invisible view
    UIView* v2 = [[UIView alloc] initWithFrame:self.view.bounds];
    v2.tag = 3;
    UITapGestureRecognizer* t = [[UITapGestureRecognizer alloc]
                                 initWithTarget:self action:@selector(hideView1:)];
    [v2 addGestureRecognizer:t];
    [self.view insertSubview:v2 belowSubview:v];
    [UIView animateWithDuration:.2 animations:^{
        v.frame = CGRectMake(0,0,320,self.view.bounds.size.height);
    }];
}

```

If, while the master view is showing, the user taps anywhere outside the master view, the tap gesture recognizer detects the tap. In response, we remove the invisible view and animate the master view back off the screen:

```

- (void) hideView1: (UIGestureRecognizer*) g {
    [g.view removeFromSuperview]; // thank you very much
    UIView* v = [self.view viewWithTag:1];
    [UIView animateWithDuration:.2 animations:^{
        v.frame = CGRectMake(-320,0,320,self.view.bounds.size.height);
    } completion:^(BOOL finished) {
    }];
}

```

There's a little more to my code, to imitate the shadow-casting behavior of the master view in the Mail app, but including it here would have cluttered the discussion.

Text

Text can be displayed in various ways:

UILabel

Displays text, possibly consisting of multiple lines, in a single font and size, with color (and highlighted color), alignment, and wrapping and truncation. Discussed in this chapter.

UITextField

Displays a single line of editable text, in a single font and size, with color and alignment; may have a border, may have a background image, and overlay views may appear at its right and left end. Discussed in this chapter; a UITextField is a UIControl subclass, so see also [Chapter 25](#).

UITextView

Displays scrollable text, possibly editable, in a single font and size, with color and alignment; can use data detectors to display tappable links. Discussed in this chapter.

UIWebView

A scrollable view displaying rendered HTML. Because HTML can express text attribute spans, this is a good way to show text in multiple fonts, sizes, colors, alignments, and so on, and to include images and tappable links. Can also display various additional document types, such as PDF, RTF, and .doc. Discussed in the next chapter.

Drawing

There are three main ways to draw strings directly:

Core Graphics

Low-level methods for drawing text (not NSStrings). For drawing in general, see [Chapter 15](#).

NSString

At a high level, the `UIStringDrawing` category on `NSString` endows strings with the ability to draw themselves, along with metrics methods (the `sizeWithFont...` methods) for learning the dimensions at which a given string will be drawn. Examples appeared in [Chapter 12](#) and [Chapter 20](#).

Core Text

The only way in iOS to draw strings with multiple fonts and styles. Core Text also provides access to advanced font typographical features. Discussed in this chapter.

An app can include fonts within its bundle; these will be loaded at launch time if the app lists them in its `Info.plist` under the “Fonts provided by application” key (`UIAppFonts`). In this way, your app can use fonts not present by default on the device.

UILabel

`UILabel` was introduced in [Chapter 7](#); its main properties are discussed in [“Built-In Cell Styles” on page 595](#). A label’s most important feature, its text, is its `text` property; you are likely also to set its `font`, `textColor`, and `textAlignment` properties, and possibly its `shadowColor` and `shadowOffset` properties. The label’s text can have an alternate `highlightedTextColor`, to be used when its `highlighted` property is YES (as happens, for example, when the label is in a selected cell of a table view).

If a `UILabel` consists of only one line of text (`numberOfLines` is 1, the default), then you can set `adjustsFontSizeToFitWidth` to YES and provide a `minimumFontSize` if you want the label to shrink the font size smaller than its `font` setting in an attempt to display as much of the `text` as possible. How the text is repositioned when this happens is determined by the label’s `baselineAdjustment` property.

A `UILabel` may alternatively consist of multiple lines of text (`numberOfLines` is greater than 1), but in that case `adjustsFontSizeToFitWidth` is ignored; the font size set in `font` is used even if not all of the text will fit.

If `numberOfLines` is 1, any line breaks in the `text` are treated as spaces. Further line breaking (wrapping) and truncation behavior, which applies to both single-line and multiline labels, is determined by its `lineBreakMode`. You can get a feel for this behavior by experimenting in the nib. Your options are:

`UILineBreakModeWordWrap`

Lines break at word-end. This is the default.

`UILineBreakModeClip`

Lines break at word-end, but the last line can break in the middle of a word.

`UILineBreakModeCharacterWrap`

All lines can break in the middle of a word.

```
UILineBreakModeHeadTruncation  
UILineBreakModeMiddleTruncation  
UILineBreakModeTailTruncation
```

Lines break at word-end. But now suppose the text is too long for the label. (This might be because a single-line label isn't wide enough, or because the `numberOfLines` is insufficient given the label's width, or because the label isn't tall enough to display the `numberOfLines`.) Then the last line displays an ellipsis at the start, middle, or end respectively, and the omitted material, which would normally be the end of the text, is instead at the point of the ellipsis. Thus, if the `lineBreakMode` is `UILineBreakModeHeadTruncation`, the last line is always displayed, and if the `lineBreakMode` is `UILineBreakModeMiddleTruncation`, the last word(s) are displayed at the end of the last line — preceded, in both cases, by everything that will fit from the start of the text to the ellipsis.

If `numberOfLines` is larger than the number of lines actually needed, the text is vertically centered in the label. This may be undesirable; you might prefer to shrink (or grow) the label to fit its text (and then perhaps to reposition the label). On the face of it, you can't use `sizeToFit` to do this, because the default `UILabel` implementation of `sizeToFit` is to make the label the right width to contain all its text on a *single* line. However, you can modify that response by overriding `UILabel`'s `textRectForBounds:limitedToNumberOfLines:` in a subclass. Here, we create a `UILabel` subclass such that the label responds to `sizeToFit` by making itself the right *height* to contain all of its text on *multiple* lines, without changing its width:

```
- (CGRect)textRectForBounds:(CGRect)bounds  
    limitedToNumberOfLines:(NSInteger)numberOfLines {  
    CGSize sz = [self.text sizeWithFont:self.font  
        constrainedToSize:CGSizeMake(self.bounds.size.width, 10000)  
        lineBreakMode:self.lineBreakMode];  
    return (CGRect){bounds.origin, CGSizeMake(self.bounds.size.width, sz.height)};  
}
```

We use `NSString`'s `sizeWithFont:constrainedToSize:lineBreakMode:` to work out the actual height needed for a given text. 10,000 is just an arbitrarily big number, which we assume the label's real height would never reach. Of course, the lines of the label won't all be displayed unless its `numberOfLines` is also sufficiently large.

The other `UILabel` method that you can override in a subclass is `drawTextInRect:`. This is the equivalent of `drawRect:` — that is, it's your chance to modify the overall drawing of the label. An example appears in [Chapter 10](#).

UITextField

A text field has many of the same properties as a label, without the ability to contain multiple lines. So it has a `text`, `font`, `textColor`, and `textAlignment`. It also has `adjustsFontSizeToFitWidth` and `minimumFontSize` properties, although these don't work exactly like a label; a text field won't allow its font size to shrink automatically as small as a

label will. If the user enters text that is too long for the width of the field, the text moves (scrolls horizontally) to show the insertion point, but when the field is no longer being edited, text that is too long is displayed with an ellipsis at the end.

A text field also has a `placeholder` property, which is the text that appears faded within the text field when it has no text; the idea is that you can use this to suggest to the user what the text field is for. If its `clearsOnBeginEditing` property is YES, the text field automatically deletes its existing text when the user begins editing within it.

A text field's border drawing is determined by its `borderStyle` property. Your options are:

UITextBorderStyleNone

No border.

UITextBorderStyleLine

A plain rectangle.

UITextBorderStyleBezel

A slightly bezeled rectangle: the top and left sides have a very slight, thin shadow.

UITextBorderStyleRoundedRect

A rounded rectangle; the top and left sides have a stronger shadow, so that the text appears markedly recessed behind the border.

A text field can have a background color (because it is a `UIView`) or a background image (`background`), possibly along with a second image (`disabledBackground`) to be displayed when the text field's `enabled` property (inherited from `UIControl`) is NO. The user can't interact with a disabled text field, but without a `disabledBackground` image, the user may lack any visual clue to this fact.

The text background of a `UITextBorderStyleRoundedRect` text field is always white; its background image is ignored. But its background color is visible at its corners, outside the rounded border, and therefore, to look good, should match what's behind the text field or should be `clearColor`.

A text field may contain up to two ancillary overlay views (such as a magnifying glass icon to suggest that the field initiates a search), its `leftView` and `rightView`, and a Clear button (a gray circle with a white "x"). The automatic visibility of each of these is determined by the `leftViewMode`, `rightViewMode`, and `clearViewMode`, respectively. The view mode values are:

UITextFieldViewModeNever

The view never appears.

UITextFieldViewModeWhileEditing

A Clear button appears if there is text in the field and the user is editing. A left or right view appears if there is *no* text in the field and the user is editing.

UITextFieldViewModeUnlessEditing

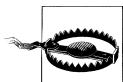
A Clear button appears if there is text in the field and the user is not editing. A left or right view appears if the user is not editing, or if the user is editing but there is no text in the field.

UITextFieldViewModeAlways

A left or right view always appears; a Clear button appears if there is text in the field.

Depending on what sort of view you use, your `leftView` and `rightView` may have to be sized manually so as not to overwhelm the text view contents. A right view and a Clear button can conflict, trying to appear at the same time (for example, if they both have `UITextFieldViewModeAlways` and there is text in the field); in this case, the right view may cover the Clear button unless you reposition it. The positions and sizes of *any* of the components of the text field can be set in relation to the text field's bounds by overriding the appropriate method in a subclass:

- `clearButtonRectForBounds`:
- `leftViewRectForBounds`:
- `rightViewRectForBounds`:
- `borderRectForBounds`:
- `textRectForBounds`:
- `placeholderRectForBounds`:
- `editingRectForBounds`:



You should make no assumptions about how often or how frequently these methods will be called; the same method might be called several times in quick succession. Also, these methods should all be called with a parameter that is the bounds of the text field, but some are called with a 100×100 bounds; this feels like a bug.

You can also override the following:

drawTextInRect:

Called when the text has changed and the user is not editing or ends editing. You should either draw the text or call `super` to draw it; if you do neither, the text will become blank. Observe that you get no method to customize the drawing of the text while it is being edited.

drawPlaceholderInRect:

Called when the placeholder text is about to appear. You should either draw the placeholder text or call `super` to draw it; if you do neither, the placeholder will become blank.

Both these methods are called with a parameter whose size is the dimensions of the text field's text area, but whose origin is `{0, 0}`. In effect what you've got is a graphics context for just the text area; any drawing you do outside the given rectangle will be clipped.

Editing and the Keyboard

Recall from [Chapter 11](#) the UIResponder-related notion of the *first responder*. A text field's editing status, as well as the presence or absence of the onscreen simulated keyboard, is intimately tied to this notion:

- When a text field is first responder, it is being edited and the keyboard is present.
- When a text field is no longer first responder, it is no longer being edited, and if no other text field (or text view) becomes first responder, the keyboard is not present. If the keyboard is present because one text field is first responder, and another text field becomes first responder (for example, because the user taps in it), the keyboard is not dismissed and brought back; it just remains onscreen.

You can programmatically control a text field's editing status, as well as the presence or absence of the keyboard, by way of the text field's first responder status. To make the insertion point appear within a text field and to cause the keyboard to appear, you send `becomeFirstResponder` to that text field; to make a text field stop being edited and to cause the keyboard to disappear, you send `resignFirstResponder` to that text field. Actually, `resignFirstResponder` returns a BOOL, because a responder might return NO to indicate that for some reason it refuses to obey this command. Note also the `UIView endEditing:` method, which can be sent to the first responder or any superview (including the window) to ask or compel the first responder to resign first responder status.

In a view presented in the `UIModalPresentationFormSheet` style on the iPad ([Chapter 19](#)), the keyboard, by default, does *not* disappear when a text field resigns first responder status. This is apparently because a form sheet is intended primarily for text input, so the keyboard is felt as accompanying the form as a whole, not individual text fields. Starting in iOS 4.3, you can prevent this exceptional behavior: in your `UIViewController` subclass, override `disablesAutomaticKeyboardDismissal` to return NO.



There is no simple way to learn what view is first responder! This is very odd, because a window surely knows what its first responder is — but it won't tell you. There's a method `isFirstResponder`, but you'd have to send it to every view in a window until you find the first responder. One workaround is to store a reference to the first responder yourself, typically in your implementation of the text field delegate's `textFieldDidBeginEditing:`. *Do not name this reference `firstResponder`!* This name is apparently already in use by Cocoa, and a name collision can cause your app to misbehave.

Dismissing the keyboard

Once the user has tapped in a text field and the keyboard has automatically appeared, how is the user supposed to get rid of it? This is unlikely to be a problem on the iPad, where the keyboard typically contains a special button that dismisses the keyboard. But on the iPhone, it's an oddly tricky issue. You would think that the "return" button in the keyboard would dismiss the keyboard; but, of itself, it doesn't.

One solution is to be the text field's delegate and to implement a text field delegate method, `textFieldShouldReturn:`. When the user taps the Return key in the keyboard, we hear about it through this method, and we tell the text field to resign its first responder status, which dismisses the keyboard:

```
- (BOOL)textFieldShouldReturn: (UITextField*) tf {  
    [tf resignFirstResponder];  
    return YES;  
}
```

I'll provide a more self-contained, automatic solution later in this chapter.

Keyboard covers text field

The keyboard has a position "docked" at the bottom of the screen. This may cover the text field in which the user wants to type, even if it is first responder. On the iPad with iOS 5, this may not be an issue, because the user can "undock" the keyboard (possibly also splitting and shrinking it) and slide it up and down the screen freely. And some scrolling views, such as a table view, may help out by scrolling automatically. Otherwise, you'll probably want to do something to reveal the text field.

To help with this, you can register for any of four keyboard-related notifications:

- `UIKeyboardWillShowNotification`
- `UIKeyboardDidShowNotification`
- `UIKeyboardWillHideNotification`
- `UIKeyboardDidHideNotification`

Those notifications all have to do with the *docked* position of the keyboard. On the iPhone, keyboard docking and keyboard visibility are equivalent: the keyboard is visible if and only if it is docked. On the iPad under iOS 5, the keyboard is said to "show" if it is being docked, whether that's because it is appearing from offscreen or because the user is docking it; and it is said to "hide" if it is undocked, whether that's because it is moving offscreen or because the user is undocking it.

Alternatively (or also), in iOS 5 you can register for two more notifications, which are sent both when the keyboard enters and leaves the screen and (on the iPad) when the user drags it, splits or unsplits it, and docks or undocks it:

- `UIKeyboardWillChangeFrameNotification`
- `UIKeyboardDidChangeFrameNotification`

The notification's `userInfo` dictionary contains information about the keyboard describing what it will do or has done, under these keys:

- `UIKeyboardFrameBeginUserInfoKey`
- `UIKeyboardFrameEndUserInfoKey`
- `UIKeyboardAnimationDurationUserInfoKey`
- `UIKeyboardAnimationCurveUserInfoKey`

Thus, to a large extent, you can coordinate your actions with those of the keyboard. In particular, by looking at the `UIKeyboardFrameEndUserInfoKey`, you know what position the keyboard is moving to; if necessary, you can compare this with the screen bounds to learn whether the keyboard will now be on or off the screen.

(In the case of `UIKeyboardWillChangeFrameNotification`, however, there won't be any `UIKeyboardFrameEndUserInfoKey` when the user starts dragging the keyboard on the iPad, because the runtime doesn't know where the user will drag the keyboard to. The frame value for the split keyboard on the iPad is the size of the *entire* keyboard, as if it weren't split; but its height is shorter than the height of the normal keyboard, so you may be able to deduce that it is split, if you really need to know that.)

Finding a strategy for dealing with the keyboard's presence depends on the needs of your particular app. It may well be that even on the iPad you can ignore `UIKeyboardDidChangeFrameNotification` and concern yourself only with the docked position of the keyboard, because, as I mentioned a moment ago, if the keyboard isn't docked, the user is free to move and split it. To illustrate, therefore, I'll concentrate on the most universal case, where the keyboard moves into and out of docked position and we detect this with `UIKeyboardWillShowNotification` and `UIKeyboardWillHideNotification`. What should we do if, when the keyboard appears, it covers the text field being edited?

One simple approach is to start with your interface embedded in a scroll view. The user need not be aware of the scroll view, and the scroll view need not be scrollable by the user; its purpose is then merely so that *you* can scroll the interface. (You might object that in that case the scroll view is unnecessary, because we can effectively scroll the interface by setting our containing view's bounds origin. This is true, but using a scroll view provides more flexibility, and in any case I'll need the scroll view as I develop the example.)

Let's assume, then, that the whole interface is in a scroll view. As our view controller's view loads (`viewDidLoad`), we configure the scroll view's content size, and register for two of the keyboard notifications:

```
[super viewDidLoad];
CGSize sz = self.scrollView.bounds.size;
sz.height *= 2;
self.scrollView.contentSize = sz;
self.scrollView.scrollEnabled = NO;
[[NSNotificationCenter defaultCenter] addObserver:self
                                         selector:@selector(keyboardShow:)]
```

```

        name:UIKeyboardWillShowNotification
        object:nil];
[[NSNotificationCenter defaultCenter] addObserver:self
                                         selector:@selector(keyboardHide:)
                                         name:UIKeyboardWillHideNotification
                                         object:nil];

```

As the keyboard appears, we store the old content offset in an instance variable and scroll to keep the text field visible. (Observe that the keyboard's frame comes to us in window coordinates, so it is necessary to convert it to our scroll view's coordinates in order to make sense of it.) When the keyboard is dismissed, we scroll back to our previous position:

```

- (void) keyboardShow: (NSNotification*) n {
    self->oldOffset = self.scrollView.contentOffset;
    NSDictionary* d = [n userInfo];
    CGRect r = [[d objectForKey:UIKeyboardFrameEndUserInfoKey] CGRectValue];
    r = [self.scrollView convertRect:r fromView:nil];
    CGRect f = self.fr.frame;
    CGFloat y =
        CGRectGetMaxY(f) + r.size.height - self.scrollView.bounds.size.height + 5;
    if (r.origin.y < CGRectGetMaxY(f))
        [self.scrollView setContentOffset:CGPointMake(0, y) animated:YES];
}

- (void) keyboardHide: (NSNotification*) n {
    [self.scrollView setContentOffset:self->oldOffset animated:YES];
}

```

The heart of that code is, of course, the determination of the value *y*, the vertical value to which we're going to scroll the scroll view. I've elected to scroll just enough to keep the entire text field above the keyboard, with a space of five pixels between them. The decision involves both aesthetics and functionality; a completely different decision could be equally valid. For example, if my interface contains exactly three text fields, it might make sense to scroll in such a way as to make them all visible, so that the user can work in any of them without dismissing the keyboard.

Let's now extend the example to cover the situation where the interface is a scroll view that the user *can* normally scroll. In that case, we shouldn't change the scroll view's content size. Instead, we should change the behavior of the scroll view so that it operates coherently within the reduced space left by the keyboard. This is a job for **contentInset**, whose purpose, you will recall ([Chapter 20](#)), is precisely to make it possible for the user to view all of the scroll view's content even though part of the scroll view is being covered by something.

So, in `viewDidLoad`, we still register for keyboard notifications, but we won't touch the scroll view. When the keyboard appears, we store not only the current content offset but the current content inset and scroll indicator inset as well; then we alter them:

```

- (void) keyboardShow: (NSNotification*) n {
    self->oldContentInset = self.scrollView.contentInset;
    self->oldIndicatorInset = self.scrollView.scrollIndicatorInsets;
    self->oldOffset = self.scrollView.contentOffset;
    NSDictionary* d = [n userInfo];
    CGRect r = [[d objectForKey:UIKeyboardFrameEndUserInfoKey] CGRectValue];
    r = [self.scrollView convertRect:r fromView:nil];
    CGRect f = self.fr.frame;
    CGFloat y =
        CGRectGetMaxY(f) + r.size.height - self.scrollView.bounds.size.height + 5;
    if (r.origin.y < CGRectGetMaxY(f))
        [self.scrollView setContentOffset:CGPointMake(0, y) animated:YES];
    UIEdgeInsets insets;
    insets = self.scrollView.contentInset;
    insets.bottom = r.size.height;
    self.scrollView.contentInset = insets;
    insets = self.scrollView.scrollIndicatorInsets;
    insets.bottom = r.size.height;
    self.scrollView.scrollIndicatorInsets = insets;
}

```

When the keyboard disappears, we restore not only the content offset but the insets as well; this works best if the insets are restored using delayed performance:

```

- (void) keyboardHide: (NSNotification*) n {
    [self.scrollView setContentOffset:self->oldOffset animated:YES];
    dispatch_time_t popTime = dispatch_time(DISPATCH_TIME_NOW, 0.5 * NSEC_PER_SEC);
    dispatch_after(popTime, dispatch_get_main_queue(), ^{
        self.scrollView.scrollIndicatorInsets = self->oldIndicatorInset;
        self.scrollView.contentInset = self->oldContentInset;
    });
}

```

This second approach works equally well even if the scroll view was not originally user-scrollable (because its content size is the same as its bounds size); the interface becomes user-scrollable only when the keyboard is present, allowing the user to see any part of it. A nice byproduct is that when the keyboard is present, the scroll view scrolls automatically if the user enters characters into a text field that has been scrolled out of sight or taps in a text field that's partially hidden.

Configuring the Keyboard

A UITextField implements the UITextInputTraits protocol, which defines properties on the UITextField that you can set to determine how the keyboard will look and how typing in the text field will behave. (These properties can also be set in the nib.) For example, you can set the `keyboardType` to `UIKeyboardTypePhonePad` to make the keyboard for this text field consist of digits only. You can set the `returnKeyType` to determine the text of the Return key (if the keyboard is of a type that has one). You can even supply your own keyboard by setting the text field's `inputView`. You can turn off autocapitalization (`autocapitalizationType`) or autocorrection (`autocorrectionType`), make the

Return key disable itself if the text field has no content (`enablesReturnKeyAutomatically`), and make the text field a password field (`secureTextEntry`).



The user's choices in the Settings app with regard to certain text input features, such as autocapitalization or autocorrection, take priority over your configuration of these same features for a particular text field.

You can attach an accessory view to the top of the keyboard by setting the text field's `inputAccessoryView`. In this example, the accessory view is a `UIButton` configured in the nib and accessed through an outlet property, `buttonView`. When editing starts, the text field delegate's `textFieldDidBeginEditing:` arrives before the `UIKeyboardWillShowNotification`, so we can use it to configure the keyboard before it appears. We also store a reference to the text field so that we know the first responder later:

```
- (void)textFieldDidBeginEditing:(UITextField *)tf {
    self.fr = tf; // keep track of first responder
    tf.inputAccessoryView = self.buttonView;
}
```

The button is a Next button. When the user taps it, we use it to move editing to the next text field:

```
- (IBAction)doNextField:(id)sender {
    UITextField* nextField = // ... figure it out, based on self.fr
    [nextField becomeFirstResponder];
}
```

The user can control the localization of the keyboard character set in the Settings app, either through a choice of the system's base language or by enabling additional "international keyboards." In the latter case, the user can switch among keyboard character sets while the keyboard is showing. But, as far as I can tell, your code can't make this choice, so you can't, for example, have a Russian-teaching app in which a certain text field automatically shows the Cyrillic keyboard. You can ask the user to switch keyboards manually, but if you really want a particular keyboard to appear regardless of the user's settings and behavior, you'll have to create it yourself and provide it as the `inputView`.

Text Field Delegate and Control Event Messages

As editing begins and proceeds in a text field, a sequence of messages is sent to the text field's delegate (some of which are also available to other objects as notifications). Using these, you can customize the text field's behavior during editing:

`textFieldShouldBeginEditing:`

Return NO to prevent the text field from becoming first responder.

textFieldDidBeginEditing: (*and UITextFieldTextDidBeginEditingNotification*)

The text field has become first responder.

textFieldShouldClear:

Return NO to prevent the operation of the Clear button or of automatic clearing on entry (*clearsOnBeginEditing*).

textFieldShouldReturn:

The user has tapped the Return button in the keyboard. We have already seen that this can be used as way of dismissing the keyboard.

textField:shouldChangeCharactersInRange:replacementString:

Sent when the user changes the text in the field by typing or pasting, or by backspacing or cutting (in which case the replacement string will have zero length). Return NO to prevent the proposed change; you can substitute text by changing the text field's `text` directly (there is no circularity, as this delegate method is not called when you do that). In this example, the user can enter only lowercase characters:

```
- (BOOL)textField:(UITextField *)textField  
    shouldChangeCharactersInRange:(NSRange)range  
    replacementString:(NSString *)string {  
    NSString* lc = [string lowercaseString];  
    if ([string isEqualToString:lc])  
        return YES;  
    textField.text = [textField.text stringByReplacingCharactersInRange:range  
                                              withString:lc];  
    return NO;  
}
```

It is common practice to implement this delegate method as a way of learning that the text has been changed, even if you then always return YES. The `UITextFieldTextDidChangeNotification` corresponds loosely.

textFieldShouldEndEditing:

Return NO to prevent the text field from resigning first responder (even if you just sent `resignFirstResponder` to it). You might do this, for example, because the text is invalid or unacceptable in some way. The user will not know why the text field is refusing to end editing, so the usual thing is to put up an alert explaining the problem.

textFieldDidEndEditing: (*and UITextFieldTextDidEndEditingNotification*)

The text field has resigned first responder. See [Chapter 21](#) for an example of using `textFieldDidEndEditing:` to fetch the text field's current text and store it in the model.

A text field is also a control. This means you can attach a target–action pair to any of the events that it reports in order to receive a message when that event occurs (see [Chapter 11](#)):

- The user can touch and drag, triggering Touch Down and the various Touch Drag events.
- If the user touches in such a way that the text field enters editing mode, Editing Did Begin and Touch Cancel are triggered; if the user causes the text field to enter editing mode in some other way (such as by tabbing into it), Editing Did Begin is triggered without any Touch events.
- As the user edits, Editing Changed is triggered. If the user taps while in editing mode, Touch Down (and possibly Touch Down Repeat) and Touch Cancel are triggered.
- Finally, when editing ends, Editing Did End is triggered; if the user stops editing by tapping Return in the keyboard, Did End on Exit is triggered first.

In general, you’re more likely to treat a text field as a text field (through its delegate messages) than as a control (through its control events). However, the Did End on Exit event message has an interesting property: it provides an alternative way to dismiss the keyboard when the user taps a text field keyboard’s Return button. If there is a Did End on Exit target–action pair for this text field, then if the text field’s delegate does not return NO from `textFieldShouldReturn:`, the keyboard will be dismissed *automatically* when the user taps the Return key. (The action handler for Did End on Exit doesn’t actually have to *do* anything.)

This suggests the following trick for getting automatic keyboard dismissal *with no code at all*. In the nib, edit the First Responder proxy object in the Attributes inspector, adding a new First Responder Action; let’s call it `dummy:`. Now hook the Did End on Exit event of the text field to the `dummy:` action of the First Responder proxy object. That’s it! Because the text field’s Did End on Exit event now has a target–action pair, the text field automatically dismisses its keyboard when the user taps Return; because there is no penalty for not finding a handler for a message sent up the responder chain, the app doesn’t crash even though there is no implementation of `dummy:` anywhere.

Of course, you can implement that trick in code instead:

```
[textField addTarget:nil action:@selector(dummy:)
    forControlEvents:UIControlEventEditingDidEndOnExit];
```

A disabled text field emits no delegate messages or control events.

The Text Field Menu

When the user double-taps or long-presses in a text field, the menu appears. It contains menu items such as Select, Select All, Paste, Copy, Cut, and Replace; which menu items appear depends on the circumstances.

The menu can be customized, but you are unlikely to do this with respect to a text field, because you don’t get any access to information about the text field’s selection, making it difficult to decide intelligently what menu items should appear or what they should

do when chosen. Thus, for the most part, it is best not to alter a text field’s menu or to interfere with its behavior.

If you do want to alter the menu, the key facts you need to know are these:

- You can add menu items to the menu through the singleton global shared `UIMenuController` object. Its `menuItems` property is an array of *custom* menu items — that is, menu items that may appear *in addition* to those that the system puts there. A menu item is a `UIMenuItem`, which is simply a title (which appears in the menu) plus an action selector. The action will be called, nil-targeted, thus sending it up the responder chain, when the user taps the menu item (and, by default, the menu will be dismissed).
- The actions for the standard menu items are nil-targeted, so they percolate up the responder chain, and you can interfere with their behavior by implementing their actions. Their selectors are listed in the `UIResponderStandardEditActions` informal protocol (except for Replace, which is implemented through an undocumented selector `promptForReplace:`).
- You govern the presence or absence of *any* menu item by implementing the `UIResponder` method `canPerformAction:withSender:` in the responder chain.

As an example, we’ll devise a text field in which the standard menu is completely replaced by our own menu, which contains a single menu item, Expand. I’m imagining here, for instance, a text field where the user can type a U.S. state two-letter abbreviation (such as “CA”) and can then summon the menu and tap Expand to get the state’s full name (such as “California”). We’ll implement this by means of a `UITextField` subclass.

At some point before the user can tap in an instance of our `UITextField` subclass, we modify the global menu; we could do this in the app delegate as the app starts up, for example:

```
UIMenuItem *mi = [[UIMenuItem alloc] initWithTitle:@"Expand"
                                              action:@selector(expand:)];
UIMenuController *mc = [UIMenuController sharedMenuController];
mc.menuItems = [NSArray arrayWithObject:mi];
```

In our `UITextField` subclass, we implement `canPerformAction:withSender:` to govern the contents of the menu. The placement of this implementation is crucial. By putting it here, we guarantee that this implementation will be called when an instance of this subclass is first responder, but at no other time. Therefore, every other text field (or any other object that displays a menu) will behave normally, displaying Cut or Select All or whatever’s appropriate; only an instance of our subclass will have the special menu, displaying only Expand.

```
- (BOOL) canPerformAction:(SEL)action withSender: (id) sender {
    if (action == @selector(expand:))
        return ([self.text length] == 2); // could be more intelligent here
    return NO;
}
```

When the user chooses the Expand menu item, the `expand:` message is sent up the responder chain. We catch it in our `UITextField` subclass and obey it. Proceeding to match abbreviations with state names is left as an exercise for the reader:

```
- (void) expand: (id) sender {
    NSString* s = self.text;
    // ... alter s here ...
    self.text = s;
}
```

By setting the `text` property at the end of the `expand:` method, we cause the selection handles to vanish if there are any.

To demonstrate interference with the standard menu items, we'll modify the example to allow the Copy menu item to appear if it wants to:

```
- (BOOL) canPerformAction:(SEL)action withSender:(id)sender {
    if (action == @selector(expand:))
        return ([self.text length] == 2);
    if (action == @selector(copy:))
        return [super canPerformAction:action withSender:sender];
    return NO;
}
```

Now we'll implement `copy:` and modify its behavior. First we call `super` to get standard copying behavior; then we modify what's now on the pasteboard:

```
- (void) copy: (id) sender {
    [super copy: sender];
    UIPasteboard* pb = [UIPasteboard generalPasteboard];
    NSString* s = pb.string;
    // ... alter s here ....
    pb.string = s;
}
```

UITextView

A text view (`UITextView`) is sort of a scrollable, multiline version of a text field (`UITextField`, with which it should not be confused). It is a scroll view subclass (`UIScrollView`, [Chapter 20](#)), and thus has (by default) no border; it is *not* a control. Nevertheless, it has many close similarities to a text field. It has `text`, `font`, `textColor`, and `textAlignment` properties; it can be editable or not, according to its `editable` property. (You might use a scrollable noneditable text view instead of a `UILabel`, so as not to be limited to a fixed number of lines of text.) An editable text view governs its keyboard just as a text field does: when it is first responder, it is being edited and shows the keyboard, and it implements the `UITextInput` protocol and has `inputView` and `inputAccessoryView` properties. Its menu works the same way as a text field's as well.

One big difference, from the programmer's point of view, between a text view and a text field is that a text view gives you information about, and control of, its selection: it has a `selectedRange` property which you can get and set, and it adds a `scrollRangeTo-`

Visible: method so that you can scroll in terms of a range of its text. The `selectedRange` is useful especially if the text view is first responder, because the selection is then meaningful and visible, but it does work (invisibly) even if the text view is not first responder.

A text view also has a `dataDetectorTypes` property that, if the text view is not editable, allows text of certain types (presumably located using `NSDataDetector`, [Chapter 10](#)) to be rendered as tappable links.

A text view's delegate messages (`UITextViewDelegate` protocol) and notifications are quite parallel to those of a text field. The big differences are:

- There's a `textViewDidChange:` delegate message (and an accompanying `UITextViewTextDidChangeNotification`), whereas a text field has its `Editing Changed` control event (and notification).
- There's a `textViewDidChangeSelection:` delegate message, whereas a text field is uninformative about the selection.

A text view's `contentSize` is maintained for you, automatically, as the text changes. You can track changes to the content size (in `textViewDidChange:`, for example), but you probably shouldn't try to change it. A common use of content size tracking is to implement a *self-sizing* text view, that is, a text view that adjusts its size automatically to embrace the amount of text it contains:

```
- (void) adjust {
    CGSize sz = self->tv.contentSize;
    CGRect f = self->tv.frame;
    f.size = sz;
    self->tv.frame = f;
}

- (void)textViewDidChange:(UITextView *)textView {
    [self adjust];
}
```

A self-sizing text view works best if the text view is not user-scrollable (`scrollEnabled` is NO). If it *is* user-scrollable, it might scroll itself as the user enters text, and you might then have to struggle to prevent it from doing so:

```
- (void)scrollViewDidScroll:(UIScrollView *)scrollView {
    scrollView.contentOffset = CGPointMakeZero;
}
```

Dismissing the keyboard for a text view works differently than for a text field. Because a text view is multiline, the Return key is meaningful for character entry; you aren't likely to want to misuse it as a way of dismissing the keyboard, and you don't get a special delegate message for it. On the iPad, the virtual keyboard may contain a button that dismisses the keyboard. On the iPhone, the interface might well consist of a text view and the keyboard, so that instead of dismissing the keyboard, the user dismisses the entire interface. For example, in the Mail app on the iPhone, when the user is

composing a message, the keyboard is present the whole time. On the other hand, in the Notes app, a note alternates between being read fullscreen and being edited with the keyboard present; in the latter case, a Done button is provided to dismiss the keyboard. If there's no good place to put a Done button in the interface, you could attach an accessory view to the keyboard itself.

The problem of having part of a text view be covered by the virtual keyboard can't be solved by meddling with its `contentInset`, because for some undocumented reason a text view doesn't accept changes to the `bottom` of its `contentInset` — the value is always reset to 32. The solution is therefore to shrink the text view itself.

In this example, we imagine a simple interface containing a single text view (`tv`) and no text fields; thus, if the keyboard shows, it must be because our text view is being edited. As our view controller loads its view, we register for keyboard notifications `UIKeyboardWillShowNotification` and `UIKeyboardWillHideNotification` as we did in our text field example. When the keyboard appears, we store the text view's frame and shrink the text view; when the keyboard hides, we restore the text view's frame:

```
- (void) keyboardShow: (NSNotification*) n {
    NSDictionary* d = [n userInfo];
    CGRect r = [[d objectForKey:UIKeyboardFrameEndUserInfoKey] CGRectValue];
    r = [self.view convertRect:r fromView:nil];
    CGRect f = self.tv.frame;
    self->oldFrame = f;
    f.size.height = self.view.frame.size.height - f.origin.y - r.size.height;
    self.tv.frame = f;
}

- (void) keyboardHide: (NSNotification*) n {
    self.tv.frame = self->oldFrame;
}
```

This approach works, but the scroll view's change of frame lacks animation to match that of the keyboard. We can easily fix that, because the notification's `userInfo` dictionary describes that animation in detail:

```
- (void) keyboardShow: (NSNotification*) n {
    NSDictionary* d = [n userInfo];
    CGRect r = [[d objectForKey:UIKeyboardFrameEndUserInfoKey] CGRectValue];
    NSNumber* curve = [d objectForKey:UIKeyboardAnimationCurveUserInfoKey];
    NSNumber* duration = [d objectForKey:UIKeyboardAnimationDurationUserInfoKey];
    r = [self.view convertRect:r fromView:nil];
    CGRect f = self.tv.frame;
    self->oldFrame = f;
    f.size.height = self.view.frame.size.height - f.origin.y - r.size.height;
    [UIView beginAnimations:nil context:NULL];
    [UIView setAnimationDuration:[duration floatValue]];
    [UIView setAnimationCurve:[curve intValue]];
    self.tv.frame = f;
    [UIView commitAnimations];
}

- (void) keyboardHide: (NSNotification*) n {
```

```

NSDictionary* d = [n userInfo];
NSNumber* curve = [d objectForKey:UIKeyboardAnimationCurveUserInfoKey];
NSNumber* duration = [d objectForKey:UIKeyboardAnimationDurationUserInfoKey];
[UIView beginAnimations:nil context:NULL];
[UIView setAnimationDuration:[duration floatValue]];
[UIView setAnimationCurve:[curve intValue]];
self.tv.frame = self->oldFrame;
[UIView commitAnimations];
}

```

Apple's own apps display some interesting uses of text views. How, for instance, is the iPhone Notes app interface for reading and editing a note actually achieved? Its key features are: the text appears to be on lined paper; the first line of the text has some space above it, where the date appears; and the text has a wide left margin.

The text is presumably a text view; it must have a background color of `clearColor` to allow the paper to show through. My guess is that the space at the start of the text is achieved with the text view's `contentInset`, whose `top` value, at least, is obeyed; the wide left margin, on the other hand, suggests that the left edge of the text view is inset from the left edge of the "paper." The paper itself might be an image view; by tracking the text view's `contentSize` (in `textViewDidChange:`, for example), the app can make sure that the paper image is always sufficiently tall, and by tracking the scroll position (in `scrollViewDidScroll`), the image view's transform can be adjusted to keep the lines of text coordinated with the lines of the paper:

```

- (void)scrollViewDidScroll:(UIScrollView *)scrollView {
    CGFloat yoff = scrollView.contentOffset.y + scrollView.contentInset.top;
    paperMiddle.transform = CGAffineTransformMakeTranslation(0, -yoff);
}

```

Another interesting interface is the Mail app's screen for composing a message. At the top are fields for addresses and subject; then comes the body of the message, which gets longer as the user types into it. The message body is evidently a text view, but it lacks scroll indicators and cannot itself be scrolled; what scrolls is the interface as a whole. So this must be a self-sizing text view; presumably the text view's `contentSize` is tracked, and both the text view's size and the enclosing scroll view are adjusted accordingly. (But the Mail app's screen for *replying* to a message uses styled text, so it is presumably drawn with Core Text, discussed in the next section.)

Core Text

Core Text allows strings to be drawn with multiple fonts and styles. It is implemented by the Core Text framework; to utilize it, your app must link to `CoreText.framework`, and your code must import `<CoreText/CoreText.h>`. It uses C, not Objective-C, and it's rather verbose, but getting started with it is not difficult.

A typical simple Core Text drawing operation begins with an *attributed string*. This is an `NSAttributedString` (or `CFAAttributedString`; they are toll-free bridged), which is a string accompanied by attributes (such as font, size, and style) applied over ranges.

Each attribute is described as a name–value pair. The names of the attributes are listed in Apple’s *Core Text String Attributes Reference*, along with their value types. The most commonly used attribute is probably `kCTFontAttributeName`, which determines the font and size of a stretch of text; its value is a `CTFontRef`, a Core Text type which is *not* bridged to `UIFont`. You’ll typically supply attributes as a dictionary of name–value pairs.

For example, imagine that we have a `UIView` subclass called `StyledText`, which has a public `text` property that is an attributed string. Its job will be to draw that attributed string into itself:

```
@interface StyledText : UIView
@property (nonatomic, copy) NSAttributedString* text;
@end
```

Imagine further that an instance of `StyledText` appears in the interface and that we have a reference to it as an instance variable called `styler`. How, then, might we create an `NSAttributedString` and assign it to `styler.text`? Let’s start with a mutable attributed string:

```
NSString* s = @"Yo ho ho and a bottle of rum!";
NSMutableAttributedString* mas =
    [[NSMutableAttributedString alloc] initWithString:s];
```

Now I’ll apply some attributes. I’ll cycle through the words of the string; to each word I’ll apply a slightly larger size of the same font. My base font will be `Baskerville 18`. Note that the name supplied when creating a `CTFont` must be a PostScript name; a free app, `Typefaces`, is helpful for learning all the fonts on a device along with their PostScript names:

```
block CGFloat f = 18.0;
CTFontRef basefont = CTFontCreateWithName((CFStringRef)@"Baskerville", f, NULL);
[s enumerateSubstringsInRange:NSMakeRange(0, [s length])
    options:NSStringEnumerationByWords
    usingBlock:
^(NSString *substring, NSRange substringRange, NSRange encRange, BOOL *stop) {
    f += 3.5;
    CTFontRef font2 = CTFontCreateCopyWithAttributes(basefont, f, NULL, NULL);
    NSDictionary* d2 = [[NSDictionary alloc] initWithObjectsAndKeys:
        CFBridgingRelease(font2),
        (NSString*)kCTFontAttributeName, nil];
    [mas addAttributes:d2 range:encRange];
}];
```

Finally, I’ll make the last word bold. The easiest way to obtain the range of the last word is to cycle through the words backward and stop after the first one. Boldness is a font trait; we must obtain a bold variant of the original font. The font we started with, `Baskerville`, has such a variant, so this will work:

```
[s enumerateSubstringsInRange:NSMakeRange(0, [s length])
    options: (NSStringEnumerationByWords |
        NSStringEnumerationReverse)
    usingBlock:
```

```

^(NSString *substring, NSRange substringRange, NSRange encRange, BOOL *stop) {
    CTFontRef font2 =
        CTFontCreateCopyWithSymbolicTraits (
            basefont, f, NULL, kCTFontBoldTrait, kCTFontBoldTrait);
    NSDictionary* d2 = [[NSDictionary alloc] initWithObjectsAndKeys:
        CFStringRelease(font2),
        (NSString*)kCTFontAttributeName, nil];
    [mas addAttributes:d2 range:encRange];
    *stop = YES; // do just once, last word
}];
```

You’re probably wondering why I seem to ask for the bold variant (`kCTFontBoldTrait`) twice. The first time (the fourth argument in the call to `CTFontCreateCopyWithSymbolicTraits`) I’m providing a bitmask. The second time (the fifth argument) I’m providing a second bitmask that says which bits of the first bitmask are meaningful. For example, suppose I’m starting with a font that might or might not be italic, and I want to obtain its bold variant — meaning that if it *is* italic, I want a bold italic font. It isn’t enough to supply a bitmask whose value is `kCTFontBoldTrait`, because this appears to switch boldness on and everything else off. Thus, the second bitmask says, “Only this one bit is important; leave all other attributes alone.” By the same token, to get a nonbold variant of a font that might be bold, you’d supply `0` as the fourth argument and `kCTFontBoldTrait` as the fifth argument.

Finally, I’ll hand the attributed string over to our self-drawing interface object and complete our memory management:

```

self.styler.text = mas;
[self.styler setNeedsDisplay];
CFRelease(basefont);
```

(And did you notice the cool use of ARC’s `CFBridgingRelease` earlier? Recall from [Chapter 12](#) that this is a way of crossing the bridge from a `CFTyperef` to an Objective-C object and, at the same time, giving ARC responsibility for completing the memory management that we started by calling a `Create` function.)

We have now generated an `NSAttributedString` and handed it over to our `StyledText`. How will the `StyledText` draw itself? There are two main ways: a `CATextLayer` and direct drawing with Core Text itself.

Let’s start by using a `CATextLayer` ([Chapter 16](#)). Because this `UIView` subclass will be instantiated from a nib, I’ll give it a `CATextLayer` in `awakeFromNib`, retaining a reference to it as an instance variable, `textLayer`. I’ll also implement `layoutSublayersOfLayer`: so that the `CATextLayer` always has the bounds of the view as a whole:

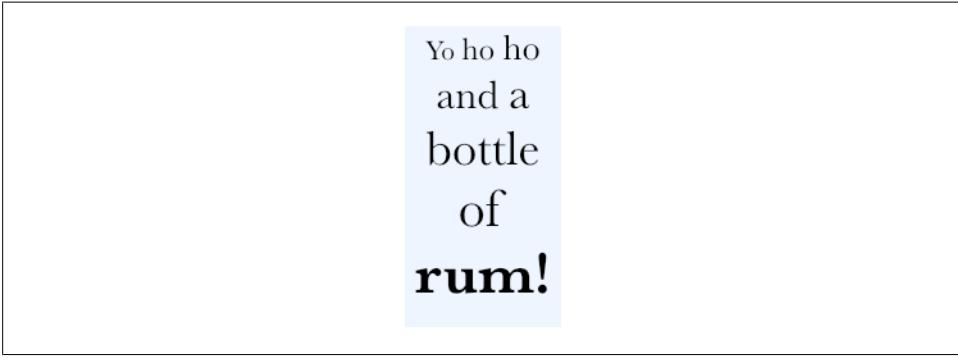
```

- (void) awakeFromNib {
    CATextLayer* lay = [[CATextLayer alloc] init];
    lay.frame = self.layer.bounds;
    [self.layer addSublayer:lay];
    self.textLayer = lay;
}
```



Yo ho ho and a bottle of **rum!**

Figure 23-1. Text whose size increases word by word



Yo ho ho
and a
bottle
of
rum!

Figure 23-2. The same text wrapped and centered

```
- (void) layoutSublayersOfLayer:(CALayer *)layer {
    [[layer.sublayers objectAtIndex:0] setFrame:layer.bounds];
}
```

Our `drawRect:` implementation is now trivial; we simply set the `CATextLayer`'s `string` property to our attributed string:

```
- (void)drawRect:(CGRect)rect {
    if (!self.text)
        return;
    self.textLayer.string = self.text;
}
```

Sure enough, our attributed string is drawn ([Figure 23-1](#))!

`CATextLayer` has some additional useful properties. If the width of the layer is insufficient to display the entire string, we can get truncation behavior with the `truncationMode` property. If the `wrapped` property is set to YES, the string will wrap. We can also set the alignment with the `alignmentMode` property. Here, I'll wrap and center the text, and I'll narrow the view so that the results are clear ([Figure 23-2](#)):

```
lay.wrapped = YES;
lay.alignmentMode = kCAAlignmentCenter;
self.autoresizingMask = UIViewAutoresizingFlexibleHeight;
CGRect f = self.frame;
f.size.width = 130;
self.frame = f;
```

The second way to display an attributed string is to draw it directly into a graphics context with Core Text. The text will be drawn upside-down unless we flip the graphics

context's coordinate system. If the string is a single line we can draw it directly into a graphics context with a CTLineRef. Positioning the drawing is up to us.

The following code results in a drawing that looks like [Figure 23-1](#):

```
- (void)drawRect:(CGRect)rect {
    if (!self.text)
        return;
    CGContextRef ctx = UIGraphicsGetCurrentContext();
    // flip context
    CGContextSaveGState(ctx);
    CGContextTranslateCTM(ctx, 0, self.bounds.size.height);
    CGContextScaleCTM(ctx, 1.0, -1.0);
    CTLineRef line =
        CTLineCreateWithAttributedString(
            (__bridge CFAttributedStringRef)self.text);
    CGContextSetTextPosition(ctx, 1, 3);
    CTLineDraw(line, ctx);
    CFRelease(line);
    CGContextRestoreGState(ctx);
}
```

If we want our string to be drawn wrapped, we must use a CTFramesetter. The framesetter requires a frame into which to draw; this is expressed as a CGPath, but don't get all excited about the possibility of drawing wrapped into some interesting shape, such as an ellipse, because on iOS the path must describe a rectangle:

```
- (void)drawRect:(CGRect)rect {
    if (!self.text)
        return;
    CGContextRef ctx = UIGraphicsGetCurrentContext();
    // flip context
    CGContextSaveGState(ctx);
    CGContextTranslateCTM(ctx, 0, self.bounds.size.height);
    CGContextScaleCTM(ctx, 1.0, -1.0);
    CTFramesetterRef fs =
        CTFramesetterCreateWithAttributedString(
            (__bridge CFAttributedStringRef)self.text);
    CGMutablePathRef path = CGPathCreateMutable();
    CGPathAddRect(path, NULL, rect);
    // range (0,0) means "the whole string"
    CTFrameRef f = CTFramesetterCreateFrame(fs, CFRRangeMake(0, 0), path, NULL);
    CTFrameDraw(f, ctx);
    CGPathRelease(path);
    CFRelease(f);
    CFRelease(fs);
    CGContextRestoreGState(ctx);
}
```

That code wraps, but the text is left-aligned. With a CTFramesetter, drawing behaviors such as alignment and truncation can be expressed as part of the attributed string itself by applying a CTParagraphStyle. Paragraph styles can also include first-line indent, tab stops, line height, spacing, and more. To center our text (which results in a drawing that looks like [Figure 23-2](#)), we apply a centered style to our text before drawing it:

ALABAMA	MONTANA
ALASKA	NEBRASKA
ARIZONA	NEVADA
ARKANSAS	NEW HAMPSHIRE
CALIFORNIA	NEW JERSEY
COLORADO	NEW MEXICO
CONNECTICUT	NEW YORK
DELAWARE	NORTH CAROLINA
FLORIDA	NORTH DAKOTA
GEORGIA	OHIO
HAWAII	OKLAHOMA
IDAHO	OREGON
ILLINOIS	PENNSYLVANIA
INDIANA	RHODE ISLAND
IOWA	SOUTH CAROLINA
KANSAS	SOUTH DAKOTA
KENTUCKY	TENNESSEE
LOUISIANA	TEXAS
MAINE	UTAH
MARYLAND	VERMONT
MASSACHUSETTS	VIRGINIA
MICHIGAN	WASHINGTON
MINNESOTA	WEST VIRGINIA
MISSISSIPPI	WISCONSIN
MISSOURI	WYOMING

Figure 23-3. Two-column text in small caps

```

NSMutableAttributedString* mas = [self.text mutableCopy];
NSString* s = [mas string];
CTTextAlignment centerValue = kCTCenterTextAlignment;
CTParagraphStyleSetting center =
    {kCTParagraphStyleSpecifierAlignment, sizeof(centerValue), &centerValue};
CTParagraphStyleSetting pss[1] = {center};
CTParagraphStyleRef ps = CTParagraphStyleCreate(pss, 1);
[mas addAttribute:(NSString*)kCTParagraphStyleAttributeName
    value:CFBridgingRelease(ps)
    range:NSMakeRange(0, [s length])];
self.text = mas;

```

Core Text can also access font typographical features that can't be accessed in any other way, such as the built-in ability of Didot and Hoefler Text to render themselves in small caps. As an example, we'll draw the names of the 50 U.S. states in small caps, centered, in two columns on an iPad (Figure 23-3).

As we create the `NSAttributedString`, we use a convenience function, `CTFontDescriptorCreateCopyWithFeature`, to access Didot's small caps variant. I had to log the result of `CTFontCopyFeatures` to learn the "magic numbers" for this variant of this font (there is also old documentation of font features at <http://developer.apple.com/fonts/registry>). We apply a centered style, just as before, and set our `styler` object's text:

```

NSString* path = [[NSBundle mainBundle] pathForResource:@"states" ofType:@"txt"];
NSString* s = [NSString stringWithContentsOfFile:path
    encoding:NSUTF8StringEncoding error:nil];
CTFontRef font = CTFontCreateWithName((CFStringRef)@"Didot", 18, NULL);
CTFontDescriptorRef fontdesc1 = CTFontCopyFontDescriptor(font);

```

```

CTFontDescriptorRef fontdesc2 =
CTFontDescriptorCreateCopyWithFeature(fontdesc1,
    (__bridge CFNumberRef)[NSNumber numberWithInt:3],
    (__bridge CFNumberRef)[NSNumber numberWithInt:3]);
CTFontRef basefont = CTFontCreateWithFontDescriptor(fontdesc2, 0, NULL);
NSDictionary* d = [[NSDictionary alloc] initWithObjectsAndKeys:
    CFBridgingRelease(basefont),
    (NSString*)kCTFontAttributeName, nil];
NSMutableAttributedString* mas =
    [[NSMutableAttributedString alloc] initWithString:s attributes:d];
CTTextAlignment centerValue = kCTCenterTextAlignment;
CTParagraphStyleSetting center =
    {kCTParagraphStyleSpecifierAlignment, sizeof(centerValue), &centerValue};
CTParagraphStyleSetting pss[1] = {center};
CTParagraphStyleRef ps = CTParagraphStyleCreate(pss, 1);
[mas addAttribute:(NSString*)kCTParagraphStyleAttributeName
    value:CFBridgingRelease(ps)
    range:NSMakeRange(0, [s length])];
self.styler.text = mas;
CFRelease(fontdesc2); CFRelease(fontdesc1); CFRelease(font);

```

The two-column arrangement is achieved by drawing into two frames. In our `drawRect` code, after flipping the context as before (not shown), we draw the entire text into the first frame and then use `CTFrameGetVisibleStringRange` to learn how much of the text actually fits into it; this tells us where in the attributed string to start drawing into the second frame:

```

CGRect r1 = rect;
r1.size.width /= 2.0; // column 1
CGRect r2 = r1;
r2.origin.x += r2.size.width; // column 2
CTFramesetterRef fs =
    CTFramesetterCreateWithAttributedString(
        (__bridge CFAAttributedStringRef)self.text);
// draw column 1
CGMutablePathRef path = CGPathCreateMutable();
CGPathAddRect(path, NULL, r1);
CTFrameRef f = CTFramesetterCreateFrame(fs, CFRangeMake(0, 0), path, NULL);
CTFrameDraw(f, ctx);
CGPathRelease(path);
CFRange drawnRange = CTFrameGetVisibleStringRange(f);
CFRelease(f);
// draw column 2
path = CGPathCreateMutable();
CGPathAddRect(path, NULL, r2);
f = CTFramesetterCreateFrame(fs,
    CFRangeMake(drawnRange.location + drawnRange.length, 0), path, NULL);
CTFrameDraw(f, ctx);
CGPathRelease(path);
CFRelease(f);
CFRelease(fs);

```

The result is [Figure 23-3](#). But now let's go further. A frame is itself composed of CTLines describing how each line of text was laid out. To demonstrate, let's turn our two-



Figure 23-4. The user has tapped on California

column list of states into an interactive interface: when the user taps the name of a state, we'll fetch that name, and we'll briefly draw a rectangle around the name to provide feedback (Figure 23-4).

We have two NSMutableArray properties, `theLines` and `theBounds`. We initialize them to empty arrays at the start of our `drawRect:`, and each time we call `CTFrameDraw` we also call a utility method:

```
[self appendLinesAndBoundsOfFrame:f context:ctx];
```

In `appendLinesAndBoundsOfFrame:context:` we save the CTLines of the frame into `theLines`; we also calculate the drawn bounds of each line and save it into `theBounds`:

```
- (void) appendLinesAndBoundsOfFrame:(CTFrameRef)f context:(CGContextRef)ctx{
    CGAffineTransform t1 =
    CGAffineTransformMakeTranslation(0, self.bounds.size.height);
    CGAffineTransform t2 = CGAffineTransformMakeScale(1, -1);
    CGAffineTransform t = CGAffineTransformConcat(t2, t1);
    CGPathRef p = CTFrameGetPath(f);
    CGRect r = CGPathGetBoundingBox(p); // this is the frame bounds
    NSArray* lines = (__bridge NSArray*)CTFrameGetLines(f);
    [self.theLines addObject:lines];
    CGPoint origins[[lines count]];
    CTFrameGetLineOrigins(f, CFRangeMake(0,0), origins);
    for (int i = 0; i < [lines count]; i++) {
        CTLineRef aLine = (__bridge CTLineRef)[lines objectAtIndex:i];
        CGRect b = CTLineGetImageBounds((CTLineRef)aLine, ctx);
        // the line origin plus the image bounds size is the bounds we want
        CGRect b2 = { origins[i], b.size };
        // but it is expressed in terms of the frame, so we must compensate
        b2.origin.x += r.origin.x;
        b2.origin.y += r.origin.y;
        // we must also compensate for the flippedness of the graphics context
        b2 = CGRectMakeApplyAffineTransform(b2, t);
        [self.theBounds addObject: [NSValue valueWithCGRect:b2]];
    }
}
```

We have attached a UITapGestureRecognizer to our view; when the user taps, we cycle through the saved bounds to see if any of them contains the tap point. If it does, we fetch the name of the state, and we draw a rectangle around it:

UITextInput

The UITextInput protocol, along with related protocols, allows you to design your own interface object that interacts with the text input system to receive keypresses from the keyboard and draw the corresponding text (with Core Text); see the chapter “Drawing and Managing Text” in Apple’s *Text, Web, and Editing Programming Guide for iOS*, as well as Apple’s SimpleTextInput example code. This is an involved subject and beyond the scope of this book, not least because I’ve no experience with it.

```
- (void) tapped: (UITapGestureRecognizer*) tap {
    CGPoint loc = [tap locationInView:self];
    for (int i = 0; i < [self.theBounds count]; i++) {
        CGRect rect = [[self.theBounds objectAtIndex: i] CGRectValue];
        if (CGRectContainsPoint(rect, loc)) {
            // draw rectangle for feedback
            CALayer* lay = [CALayer layer];
            lay.frame = CGRectInset(rect, -5, -5);
            lay.borderWidth = 2;
            [self.layer addSublayer: lay];
            [lay performSelector:@selector(removeFromSuperlayer)
                withObject:nil afterDelay:0.3];
            // fetch the drawn string tapped on
            CTLineRef theLine =
                (__bridge CTLineRef)[self.theLines objectAtIndex:i];
            CFRange range = CTLineGetStringRange(theLine);
            CFStringRef s = CFStringCreateWithSubstring(
                NULL, (__bridge CFStringRef)[self.text string], range);
            // ... could do something useful with string here ...
            NSLog(@"%@", s);
            CFRelease(s);
            break;
        }
    }
}
```

If we needed to, we could even learn what character the user tapped by going down to the level of glyph runs (CTRun) and glyphs (CTGlyph). We have barely scratched the surface of what Core Text can do. Read Apple’s *Core Text Programming Guide* for further information.

Web Views

A web view (`UIWebView`) is a `UIView` subclass that acts as a versatile renderer of text in various formats, including:

- HTML
- PDF
- RTF, including `.rtfd` (which must be supplied in a zipped format, `.rtfd.zip`)
- Microsoft Word (`.doc`), Excel (`.xls`), and PowerPoint (`.ppt`)
- Pages, Numbers, and Keynote; before iWork 2009, these must be zipped (e.g. `.key.zip`), but starting with iWork 2009 they must *not* be zipped.

In addition to displaying rendered text, a web view is, by default, a web browser. This means that if you ask a web view to display HTML that makes reference to an object available on disk or over the Internet, such as an image to be shown as the source of an `img` tag, the web view by default will attempt to fetch it and display it. Similarly, if the user taps, within the web view, on a link that leads to content on disk or over the Internet that the web view can render, the web view by default will attempt to fetch that content and display it. Indeed, a web view is, in effect, a front end for WebKit, the same rendering engine used by Mobile Safari (and, for that matter, by Safari on Mac OS X). A web view can display non-HTML file formats such as PDF, RTF, and so on, precisely because WebKit can display them.

As the user taps links and displays web pages, the web view keeps back-and-forward lists, just like a web browser. Two properties, `canGoBack` and `canGoForward`, and two methods, `goBack` and `goForward`, let you interact with this list. Your interface could thus contain Back and Forward buttons, like a miniature web browser.

A web view is scrollable, but `UIWebView` is *not* a `UIScrollView` subclass; it *has* a scroll view, rather than *being* a scroll view. Starting in iOS 5, however, you can access a web view's scroll view as its `scrollView` property.

A web view is zoomable if its `scalesToFit` property is YES; in that case, it initially scales its content to fit, and the user can zoom the content (this includes use of the gesture,

familiar from Mobile Safari, whereby double-tapping part of a web page zooms to that region of the page). Like a text view, its `dataDetectorTypes` property lets you set certain types of data to be automatically converted to clickable links. An obvious difference from a text view is that the target of a web page link may be displayed right there in the web view, rather than switching to Mobile Safari.

`UIWebView` is not intended for subclassing.

It is possible to design an entire app that is effectively nothing but a `UIWebView` — especially if you have control of the server with which the user is interacting. Indeed, before the advent of iOS, an iPhone app *was* a web application. There are still iPhone apps that work this way, but such an approach to app design is outside the scope of this book. (See Apple's *Mobile Safari Web Application Tutorial* if you're curious.)

A web view's most important task is to render HTML content; like any browser, a web view understands HTML, CSS, and JavaScript. In order to construct content for a web view, *you* must know HTML, CSS, and JavaScript. Discussion of those languages is beyond the scope of this book; each would require a book (at least) of its own.

Loading Content

To load a web view with content initially, you're going to need one of three things:

An `NSURLRequest`

Construct an `NSURLRequest` and call `loadRequest:`. An `NSURLRequest` might involve a file URL referring to a file on disk (within your app's bundle, for instance); the web view will deduce the file's type from its extension. But it might also involve the URL of a resource to be fetched across the Internet, in which case you can configure various additional aspects of the request (for example, you can form a POST request). This is the only form of loading that works with `goBack` (because in the other two forms, there is no URL to go back to).

An HTML string

Construct an `NSString` consisting of valid HTML and call `loadHTMLString:baseURL:`. The `baseURL:` will be used to fetch any resources referred to by a partial (relative) URL in the string. For example, you could use this to cause partial URLs to refer to resources inside your app's bundle.

Data and a MIME type

Obtain an `NSData` object and call `loadData:MIMEType:textEncodingName:baseURL:`. Obviously, this requires that you know the appropriate MIME type, and that you obtain the content as `NSData` (or convert it to `NSData`). Typically, this will be because the content was itself obtained by fetching it from the Internet ([Chapter 37](#)).

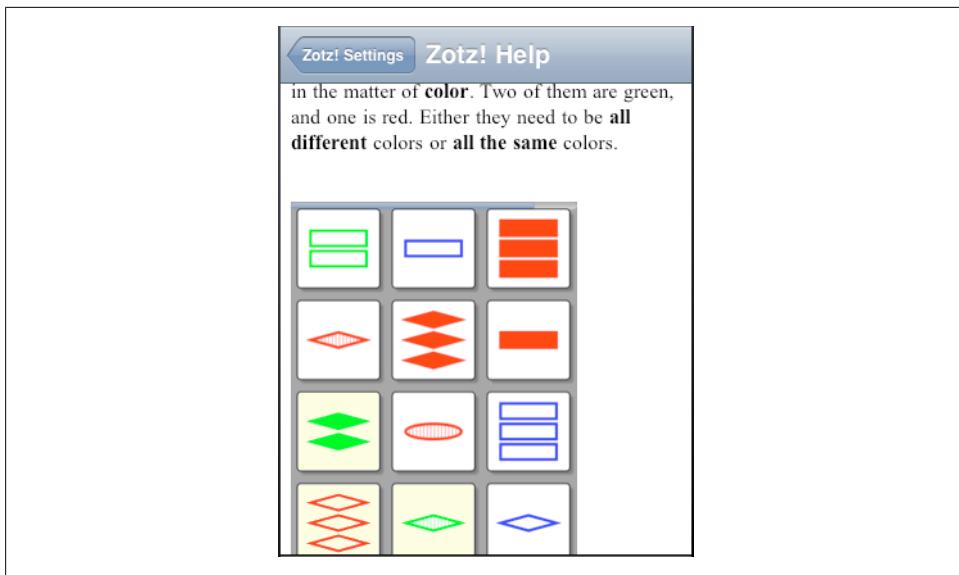


Figure 24-1. A Help screen that's a web view

There is often more than one way to load a given piece of content. For instance, one of Apple's own examples suggests that you display a PDF file in your app's bundle by loading it as data, along these lines:

```
NSString *thePath = [[NSBundle mainBundle] pathForResource:@"MyPDF" ofType:@"pdf"];
NSData *pdfData = [NSData dataWithContentsOfFile:thePath];
[self.wv loadData:pdfData MIMEType:@"application/pdf"
    textEncodingName:@"utf-8" baseURL:nil];
```

But the same thing can be done with a file URL and `loadRequest:`, like this:

```
NSURL* url = [[NSBundle mainBundle] URLForResource:@"MyPDF" withExtension:@"pdf"];
NSURLRequest* req = [[NSURLRequest alloc] initWithURL:url];
[self.wv loadRequest:req];
[req release];
```

Similarly, in one of my apps, where the Help screen is a web view ([Figure 24-1](#)), the content is an HTML file along with some referenced image files, and I load it like this:

```
NSString* path = [[NSBundle mainBundle] pathForResource:@"help" ofType:@"html"];
NSURL* url = [NSURL fileURLWithPath:path];
NSError* err = nil;
NSString* s = [NSString stringWithContentsOfURL:url
    encoding:NSUTF8StringEncoding error:&err];
// error-checking omitted
[view loadHTMLString:s baseURL:url];
```

Observe that I supply both the string contents of the HTML file and the URL reference to the same file, the latter to act as a base URL so that the relative references to the images will work properly. (At the time I wrote that code, the `NSBundle` method `URLForResource:withExtension:` didn't yet exist, so I had to form a pathname reference to the

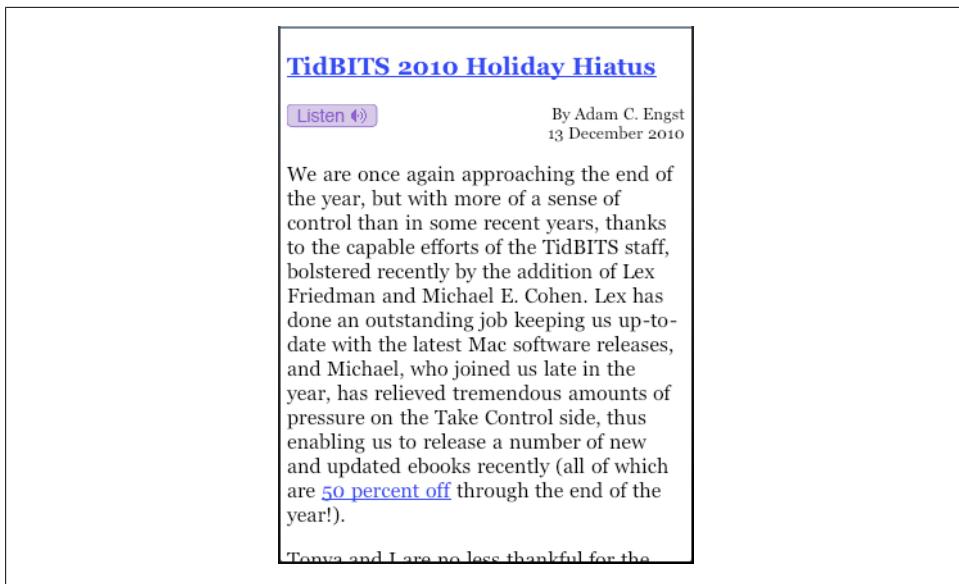


Figure 24-2. A web view with dynamically formed content

file and convert it to a URL.) In this instance, I could have used `loadRequest:` and the file URL:

```
NSString* path = [[NSBundle mainBundle] pathForResource:@"help" ofType:@"html"];
NSURL* url = [NSURL fileURLWithPath:path];
NSURLRequest* req = [[NSURLRequest alloc] initWithURL:url];
[view loadRequest: req];
[req release];
```

You can use `loadHTMLString:baseURL:` to form your own web view content dynamically. For example, in the TidBITS News app, the content of an article is displayed in a web view that is loaded using `loadHTMLString:baseURL:`. The body of the article comes from an RSS feed, but it is wrapped in programmatically supplied material. Thus, in [Figure 24-2](#), the title of the article and the fact that it is a link, the right-aligned author byline and publication date, and the Listen button, along with the overall formatting of the text (including the font size), are imposed as the web view appears.

There are many possible strategies for doing this. In the case of the TidBITS News app, I start with a template loaded from disk:

```
<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 4.01 Transitional//EN"
"http://www.w3.org/TR/html4/loose.dtd">

<style type="text/css">
```

```

    p.inflow_image {
        text-align:center;
    }
    img {
        width:100%;
        max-width:%ipx;
        height:auto
    }

```

</style>

<!-- preload alt image -->

```

<script type="text/javascript" language="javascript">
    (new Image()).src='tb_iphone_listen_pressed_02.png';
</script>
```

<title>no title</title>

</head>

<body style="font-size:%ipx; font-family:Georgia; margin:1px %ipx">

<!-- title, which is a link to original article at our site -->

```

<div style="margin-top: 0px; margin-bottom: 15px">
    <h3><a href="#">%@</a></h3>

```

</div>

<!-- playbutton or nothing; author and date -->

```

<div style="width:100%">
    <span style="float:left; display:block; vertical-align:middle">%@</span>
    <span style="float:right; margin-bottom: 15px; display:block;
    text-align:right; font-size:80%;">
        By %@<br>%@
    </span>

```

</div>

<!-- body, from feed -->

```

<div style="clear:both; margin:30px 0px;">
    %@
</div>

```

</body>

</html>

As you can see, the template defines the structure of a valid HTML document — the opening and closing tags, the head area (including some CSS styling and a little JavaScript), and a body consisting of `divs` laying out the parts of the page. The template is designed to be used as the format string in a `stringWithFormat:` method call; hence the various format specifiers scattered throughout it (and literal percent signs are escaped by doubling them). When the web view is to be loaded, the template is read from disk and handed over to `stringWithFormat:`, with every format specifier matched by an argument:

```

NSError* err = nil;
NSString* template =
    [NSString stringWithContentsOfFile:
        [[NSBundle mainBundle] pathForResource:@"htmltemplate" ofType:@"txt"]
            encoding: NSUTF8StringEncoding error:&err];
// error-checking omitted
NSString* s = [NSString stringWithFormat: template,
    maxImageWidth,
    [fontsize intValue],
    margin,

```

```

anitem.guid,
anitem.title,
(canPlay ? playbutton : @""),
[anitem authorOfItem],
date,
anitem.content
];

```

Some of these arguments (such as `anitem.title`, `date`, `anitem.content`) slot values more or less directly from the app’s model into the web view. Others are derived from the current circumstances. For example, `maxImageWidth` and `margin` have been set depending on whether the app is running on the iPhone or on the iPad; `fontsize` comes from the user defaults, because the user is allowed to determine how large the text should be. The result is an HTML string ready for `loadHTMLString:baseURL:`.

Web view content is loaded *asynchronously* (gradually, in a thread of its own), and it might not be loaded at all (because the user might not be connected to the Internet, the server might not respond properly, and so on). If you’re loading a resource directly from disk, loading is quick and nothing is going to go wrong; even then, though, rendering the content can take time, and even a resource loaded from disk, or content formed directly as an HTML string, might itself refer to material out on the Internet that takes time to fetch.

Your app’s interface is not blocked or frozen while the content is loading. On the contrary, it remains accessible and operative; that’s what “asynchronous” means. The web view, in fetching a web page and its linked components, is doing something quite complex, involving both threading and network interaction, but it shields you from this complexity. Your own interaction with the web view stays on the main thread and is straightforward. You ask the web view to load some content, and then you just sit back and let it worry about the details.

Indeed, there’s very little you *can* do once you’ve asked a web view to load content. Your main concerns will probably be to know when loading really starts, when it has finished, and whether it succeeded. To help you with this, a `UIWebView`’s delegate (adopting the `UIWebViewDelegate` protocol) gets three messages:

- `webViewDidStartLoad:`
- `webViewDidFinishLoad:`
- `webView:didFailLoadWithError:`

In this example from the TidBITS News app, I mask the delay while the content loads by displaying an activity indicator (a `UIActivityIndicatorView`, referred to by a property, `activity`) at the center of the web view:

```

- (void)webViewDidStartLoad:(UIWebView *)wv {
    self.activity.center =
        CGPointMake(CGRectGetMidX(wv.bounds), CGRectGetMidY(wv.bounds));
    [self.activity startAnimating];
}

```

```

- (void)webViewDidFinishLoad:(UIWebView *)wv {
    [self.activity stopAnimating];
}

- (void)webView:(UIWebView *)wv didFailLoadWithError:(NSError *)error {
    [self.activity stopAnimating];
}

```

Before designing the HTML to be displayed in a web view, you might want to read up on the brand of HTML native to the mobile WebKit engine. Of course a web view *can* display any valid HTML you throw at it, but the mobile WebKit has certain limitations. For example, mobile WebKit notoriously doesn't use plug-ins, such as Flash; it doesn't implement scrollable frames within framesets; and it imposes limits on the size of resources (such as images) that it can display. On the plus side, it has many special abilities and specifications that you'll want to take advantage of; for example, WebKit is in the forefront of the march towards HTML 5.

A good place to start is Apple's *Safari Web Content Guide* (<http://developer.apple.com/library/safari/documentation/AppleApplications/Reference/SafariWebContent>). It contains links to all the other relevant documentation, such as the *Safari CSS Visual Effects Guide* (<http://developer.apple.com/library/safari/documentation/InternetWeb/Conceptual/SafariVisualEffectsProgGuide>), which describes some things you can do with WebKit's implementation of CSS3 (like animations), and the *Safari HTML5 Audio and Video Guide* (http://developer.apple.com/library/safari/documentation/AudioVideo/Conceptual/Using_HTML5_Audio_Video), which describes WebKit's audio and video player support.

If nothing else, you'll definitely want to be aware of one important aspect of web page content — the *viewport*. You'll notice that the TidBITS News HTML template I showed a moment ago contains this line within its `<head>` area:

```
<meta name="viewport" content="initial-scale=1.0" />
```

Without that line, the HTML string is laid out incorrectly when it is rendered. This is noticeable especially with the iPad version of TidBITS News, where the web view can be rotated when the device is rotated, causing its width to change: in one orientation or the other, the text will be too wide for the web view, and the user has to scroll horizontally in order to read it all. The *Safari Web Content Guide* explains why: if no *viewport* is specified, the *viewport* can change when the app rotates. Setting the *initial-scale* causes the *viewport* size to adopt correct values in both orientations.

A web view's *loading* property tells you whether it is in the process of loading a request. If, at the time a web view is to be destroyed, its *loading* is YES, it is up to you to cancel the request by sending it the *stopLoading* message first; actually, it does no harm to send the web view *stopLoading* in any case. In addition, UIWebView is one of those weird classes I warned you about (Chapter 12) whose memory management behavior is odd: Apple's documentation warns that if you assign a UIWebView a delegate, you

must nilify its `delegate` property before releasing the web view. Thus, in a controller class that retains a web view, I do an extra little dance in `dealloc`:

```
- (void) dealloc {
    [wv stopLoading];
    wv.delegate = nil;
}
```

Communicating with a Web View

Having loaded a web view with content, you don't so much configure or command the web view as communicate with it. There are two modes of communication with a web view and its content:

Load requests

When a web view is asked to load content, possibly because the user has tapped a link within it, its delegate is sent the message `webView:shouldStartLoadWithRequest:navigationType:`. This is your opportunity to interfere with the web view's loading behavior; if you return NO, the content won't load.

The second argument is an `NSURLRequest`, whose `URL` property you can analyze (very easily, because it's an `NSURL`). The third argument is a constant describing the type of navigation involved, whose value will be one of the following:

- `UIWebViewNavigationTypeLinkClicked`
- `UIWebViewNavigationTypeFormSubmitted`
- `UIWebViewNavigationTypeBackForward`
- `UIWebViewNavigationTypeReload`
- `UIWebViewNavigationTypeFormResubmitted`
- `UIWebViewNavigationTypeOther` (includes loading the web view with content initially)

JavaScript execution

You can speak JavaScript to a web view's content by sending it the `stringByEvaluatingJavaScriptFromString:` message. Thus you can enquire as to the nature and details of that content, and you can alter the content dynamically.

The TidBITS News app uses `webView:shouldStartLoadWithRequest:navigationType:` to distinguish between the user tapping an ordinary link and tapping the Listen button (shown in [Figure 24-2](#)). The `onClick` script for the `<a>` tag surrounding the Listen button image executes this JavaScript code:

```
document.location='play:me'
```

This causes the web view to attempt to load an `NSURLRequest` whose `URL` is `play:me`, which is totally bogus; it's merely an internal signal to ourselves. In the web view's delegate, we intercept the attempt to load this request, examine the `NSURLRequest`,

observe that its URL has a scheme called @"play", and prevent the loading from taking place; instead, we head back to the Internet to start playing the online podcast recording associated with this article. Any other load request caused by tapping a link is also prevented and redirected instead to Mobile Safari, because we don't want our web view used as an all-purpose browser. But we do let our web view load a request in the general case, because otherwise it wouldn't even respond to our attempt to load it with HTML content in the first place:

```
- (BOOL)webView:(UIWebView *)webView shouldStartLoadWithRequest:(NSURLRequest *)request navigationType:(UIWebViewNavigationType)nt {
    if ([r.URL.scheme isEqualToString: @"play"]) {
        [self doPlay:nil];
        return NO;
    }
    if (nt == UIWebViewNavigationTypeLinkClicked) {
        [[UIApplication sharedApplication] openURL:r.URL];
        return NO;
    }
    return YES;
}
```

The TidBITS News app uses JavaScript in several ways: I'll describe one. If the user reads an article, then leaves that screen to examine the list of articles (or terminates the app), but then returns to the same article, we'd like to display the article vertically scrolled to the same position where it was before. A web view is not a UIScrollView, and prior to iOS 5 there was no access to its scroll view, so there was no way to learn or control its scroll position. However, a web view is a browser, so we can learn the scroll position of its content using JavaScript (`wv` here is the web view):

```
NSString* scrollY = [wv stringByEvaluatingJavaScriptFromString: @"scrollY"];
```

Later, we can restore the scroll position by using the converse:

```
[wv stringByEvaluatingJavaScriptFromString:
    [NSString stringWithFormat: @"window.scrollTo(0, %@);", scrollY]];
```

In iOS 5, however, I would probably accomplish that by referring to the `contentOffset` of the web view's `scrollView`.

JavaScript and the document object model (DOM) are quite powerful. Here's some additional documentation you may find helpful:

- WebKit DOM Programming Topics (<http://developer.apple.com/library/safari/#documentation/AppleApplications/Conceptual/SafariJSPProgTopics/WebKitJavaScript.html>)
- WebKit DOM Reference (<http://developer.apple.com/library/safari/#documentation/AppleApplications/Reference/WebKitDOMRef>)
- Safari DOM Additions Reference (<http://developer.apple.com/library/safari/#documentation/AppleApplications/Reference/SafariJSRef>).

Controls and Other Views

This chapter discusses all `UIView` subclasses provided by `UIKit` that haven't been discussed already (except for the two modal dialog classes, which are described in the next chapter). It's remarkable how few of them there are; `UIKit` exhibits a noteworthy economy of means in this regard.

Additional `UIView` subclasses are provided by other frameworks. For example, the Map Kit framework provides the `MKMapView` ([Chapter 34](#)). Also, additional `UIViewController` subclasses are provided by other frameworks as a way of creating interface. For example, the MessageUI framework provides `MFMailComposeViewController`, which provides your app with interface for letting the user compose and send a mail message ([Chapter 33](#)). There will be lots of examples in [Part VI](#).

`UIActivityIndicatorView`

An activity indicator (`UIActivityIndicatorView`) appears as the spokes of a small wheel. You set the spokes spinning with `startAnimating`, giving the user a sense that some time-consuming process is taking place. You stop the spinning with `stopAnimating`. If the activity indicator's `hidesWhenStopped` is YES (the default), it is visible only while spinning.

An activity indicator comes in a style, its `activityIndicatorViewStyle`; if it is created in code, you'll set its style with `initWithActivityIndicatorStyle:`. Your choices are:

- `UIActivityIndicatorViewStyleWhiteLarge`
- `UIActivityIndicatorViewStyleWhite`
- `UIActivityIndicatorViewStyleGray`

In iOS 5 you can also assign an activity indicator a `color`; this overrides the color assigned through the style. An activity indicator is a `UIView`, so you can set its `backgroundColor`; a nice effect is to give an activity indicator a contrasting translucent background color and round the corners of that background by accessing its layer ([Figure 25-1](#)):



Figure 25-1. A large activity indicator

```
v.color = [UIColor yellowColor];
v.backgroundColor = [UIColor colorWithWhite:0.2 alpha:0.4];
```

An activity indicator has a standard size, which depends on its style. Changing its size in code changes the size of the region painted by its `backgroundColor`, but not the size of the spinning spokes.

Here's some code from a `UITableViewCell` subclass in one of my apps. In this app, it takes some time, after the user taps a cell to select it, for me to construct the next view and navigate to it, so to cover the delay, I show a spinning activity indicator in the center of a cell while it's selected:

```
- (void)setSelected:(BOOL)selected animated:(BOOL)animated {
    if (selected) {
        UIActivityIndicatorView* v =
            [[UIActivityIndicatorView alloc]
             initWithActivityIndicatorStyle:UIActivityIndicatorViewStyleWhiteLarge];
        v.center =
            CGPointMake(self.bounds.size.width/2.0, self.bounds.size.height/2.0);
        v.frame = CGRectMake(v.frame);
        v.tag = 1001;
        [self.contentView addSubview:v];
        [v startAnimating];
    } else {
        [[self.contentView viewWithTag:1001] removeFromSuperview];
        // no harm if nonexistent
    }
    [super setSelected:selected animated:animated];
}
```

If activity involves the network, you might want to set `UIApplication`'s `networkActivityIndicatorVisible` to YES. This displays a small spinning activity indicator in the status bar. The indicator is not reflecting actual network activity; if it's visible, it's spinning. Be sure to set it back to NO when the activity is over.

An activity indicator is simple and standard, but you can't change the way it's drawn. If you want your own custom activity indicator, though, it's easy to make one. One obvious way is to use a `UIImageView` with a sequence of custom images forming an animation (`animationImages`), as described in [Chapter 17](#).

UIProgressView

A progress view (`UIProgressView`) is a “thermometer,” graphically displaying a percentage. It is often used to represent a time-consuming process whose percentage of completion is known (if the percentage of completion is unknown, you're more likely



Figure 25-2. A progress view

to use an activity indicator), but it might also be used to represent a fairly static percentage. For example, in one of my apps, I use a progress view to show the current position within the song being played by the built-in music player; in another app, which is a card game, I use a progress view to show how many cards are left in the deck.

A progress view comes in a style, its `progressViewStyle`; if the progress view is created in code, you'll set its style with `initWithProgressViewStyle:`. Your choices are:

- `UIProgressViewStyleDefault`
- `UIProgressViewStyleBar`

The latter is intended for use in a `UIBarButtonItem`, as the title view of a navigation item, and so on.

The height (the narrow dimension) of a progress view is generally not up to you; it's determined by the progress view's style. Changing a progress view's height has no visible effect on how the thermometer is drawn and is not a good idea.

The fullness of the thermometer is the progress view's `progress` property. This is a value between 0 and 1, inclusive; obviously, you'll need to do some elementary arithmetic in order to convert from the actual value you're reflecting to a value within that range. For example, to reflect the number of cards remaining in a deck of 52 cards:

```
prog.progress = [[deck cards] count] / 52.0;
```

Starting in iOS 5, you can animate the change from one `progress` value to another by calling `setProgress:animated:`.

Also new in iOS 5, you can customize the colors or images of the parts of the progress view. To customize the colors, set the progress view's `progressTintColor` and `trackTintColor` (the track is the unfilled part of the progress view). To customize the images, set the progress view's `progressView` and `trackView`; these, if set, override the tint colors. The views will be squashed and stretched to fill the bounds of the track and the progress "thermometer," so in all probability you'll use a resizable image whose height is the progress view's standard height (9 points). The "thermometer" image is not clipped to the track image, so seeing that it fits inside is up to you. In this simple example, the track image and progress image are squares rotated 45 degrees, so that they both have pointed ends:

```
UIGraphicsBeginImageContextWithOptions(CGSizeMake(9,9), NO, 0);
CGContextRef con = UIGraphicsGetCurrentContext();
CGContextSetFillColorWithColor(con, [UIColor blackColor].CGColor);
CGContextMoveToPoint(con, 0, 4.5);
CGContextAddLineToPoint(con, 4.5, 9);
CGContextAddLineToPoint(con, 9, 4.5);
CGContextAddLineToPoint(con, 4.5, 0);
```



Figure 25-3. A custom progress view

```
CGContextClosePath(con);
CGPathRef p = CGContextCopyPath(con);
CGContextFillPath(con);
UIImage* im = UIGraphicsGetImageFromCurrentImageContext();
CGContextSetFillColorWithColor(con, [UIColor whiteColor].CGColor);
CGContextAddPath(con, p);
CGContextFillPath(con);
UIImage* im2 = UIGraphicsGetImageFromCurrentImageContext();
CGPathRelease(p);
UIGraphicsEndImageContext();
im = [im resizableImageWithCapInsets:UIEdgeInsetsMake(4, 4, 4, 4)];
im2 = [im2 resizableImageWithCapInsets:UIEdgeInsetsMake(4, 4, 4, 4)];
prog.trackImage = im;
prog.progressImage = im2;
```

Despite its new customizability, you can't really change the way a progress view is drawn; for example, you can't make it taller. If you want your own custom progress view, it's easy to make one; all you need is a custom UIView subclass that draws something similar to a thermometer. Figure 25-3 shows a simple custom thermometer view; it has a `value` property, and you set this to something between 0 and 1 and then call `setNeedsDisplay` to get the view to redraw itself. Here's its `drawRect:` code:

```
- (void)drawRect:(CGRect)rect {
    CGContextRef c = UIGraphicsGetCurrentContext();
    [[UIColor whiteColor] set];
    CGFloat ins = 2.0;
    CGRect r = CGRectMakeInset(self.bounds, ins, ins);
    CGFloat radius = r.size.height / 2.0;
    CGMutablePathRef path = CGPathCreateMutable();
    CGPathMoveToPoint(path, NULL, CGRectGetMaxX(r) - radius, ins);
    CGPathAddArc(path, NULL,
                 radius-ins, radius+ins, radius, -M_PI/2.0, M_PI/2.0, true);
    CGPathAddArc(path, NULL,
                 CGRectGetMaxX(r) - radius, radius+ins, radius, M_PI/2.0, -M_PI/2.0, true);
    CGPathCloseSubpath(path);
    CGContextAddPath(c, path);
    CGContextSetLineWidth(c, 2);
    CGContextStrokePath(c);
    CGContextAddPath(c, path);
    CGContextClip(c);
    CGContextFillRect(c, CGRectMake(
        r.origin.x, r.origin.y, r.size.width * self.value, r.size.height));
}
```

UIPickerView

A UIPickerView displays selectable choices using a rotating drum metaphor. It has a standard legal range of possible heights, which is undocumented and must be discovered by trial and error (attempting to set the height outside this range will fail with a warning in the console); its width is largely up to you. Each drum, or column, is called a *component*.

Your code configures the UIPickerView's content through its data source (UIPickerViewDataSource) and delegate (UIPickerViewDelegate), which are usually the same object (see also [Chapter 11](#)). Your data source and delegate must answer questions similar to those posed by a UITableView ([Chapter 21](#)):

`numberOfComponentsInPickerView: (data source)`

How many components (drums) does this picker view have?

`pickerView:numberOfRowsInComponent: (data source)`

How many rows does this component have? The first component is numbered 0.

`pickerView:titleForRow:forComponent:`

`pickerView:viewForRow:forComponent:reusingView: (delegate)`

What should this row of this component display? The first row is numbered 0. You can supply either a simple title string or an entire view such as a UILabel, giving you more control over formatting, but you must supply every row of every component the same way, because if `viewForRow` is implemented, `titleForRow` isn't called. The `reusingView` parameter, if not nil, is a view that you supplied for a row now no longer visible, giving you a chance to reuse it, much as cells are reused in a table view.

Here's the code for a UIPickerView ([Figure 25-4](#)) that displays the names of the 50 U.S. states, obtained from a text file. We implement `pickerView:viewForRow:forComponent:reusingView:` just because it's the more interesting case; as our views, we supply UILabel instances. The state names, drawn from an NSArray property `states`, are drawn centered because the labels are themselves centered within the picker view:

```
- (NSInteger)numberOfComponentsInPickerView:(UIPickerView *)pickerView {
    return 1;
}

- (NSInteger)pickerView:(UIPickerView *)pickerView
    numberOfRowsInComponent:(NSInteger)component {
    return 50;
}

- (UIView *)pickerView:(UIPickerView *)pickerView viewForRow:(NSInteger)row
    forComponent:(NSInteger)component reusingView:(UIView *)view {
    UILabel* lab;
    if (view)
        lab = (UILabel*)view; // reuse it
    else
```



Figure 25-4. A picker view

```
lab = [[UILabel alloc] init];
lab.text = [self.states objectAtIndex:row];
lab.backgroundColor = [UIColor clearColor];
[lab sizeToFit];
return lab;
}
```

The delegate may further configure the UIPickerView's physical appearance by means of these methods:

- `pickerView:rowHeightForComponent:`
- `pickerView:widthForComponent:`

The delegate may implement `pickerView:didSelectRow:inComponent:` to be notified each time the user spins a drum to a new position. You can also query the picker view directly by sending it `selectedRowInComponent:`.

You can set the value to which any drum is turned using `selectRow:inComponent:animated:`. Other handy picker view methods allow you to request that the data be reloaded, and there are properties and methods to query the picker view's contents (though of course they do not relieve you of responsibility for knowing the data model from which the picker view's contents are supplied):

- `reloadComponent:`
- `reloadAllComponents`
- `numberOfComponents`
- `numberOfRowsInComponent:`
- `viewForRow:forComponent:`

By implementing `pickerView:didSelectRow:inComponent:` and using `reloadComponent:` you can make a picker view where the values displayed by one drum depend dynamically on what is selected in another. For example, one can imagine expanding our U.S. states example to include a second drum listing major cities in each state; when the user switches to a different state in the first drum, a different set of major cities appears in the second drum.



Figure 25-5. A search bar with a search results button

UIearchBar

A search bar (UIearchBar) is essentially a variety of text field, though it is not in fact a UITextField subclass. It is displayed by default as a rounded rectangle containing a magnifying glass icon, where the user can enter text ([Figure 25-5](#)). It does not, of itself, do any searching or display the results of a search; a common interface involves displaying the results of a search as a table, and the UISearchDisplayController class makes this easy to do (see [Chapter 21](#)).

A search bar's current text is its `text` property. It can have a `placeholder`, which appears when there is no text. A `prompt` can be displayed above the search bar to explain its purpose. Delegate methods (UIearchBarDelegate) notify you of editing events; for their use, compare the text field and text view delegate methods discussed in [Chapter 23](#):

- `searchBarShouldBeginEditing:`
- `searchBarTextDidBeginEditing:`
- `searchBar:textDidChange:`
- `searchBar:shouldChangeTextInRange:replacementText:`
- `searchBarShouldEndEditing:`
- `searchBarTextDidEndEditing:`

A search bar has a `barStyle`, for which your choices are `UIBarStyleDefault` or `UIBarStyleBlack`, and either `translucent` or not. Alternatively, the search bar may have a `tintColor`. The bar style and tint color are the same as for a navigation bar or toolbar, and are drawn the same way a navigation bar or toolbar would draw them (see later in this chapter); because of this, a search bar looks good at the top of the screen, where a navigation bar or toolbar might go — in effect, it is drawn as if it *were* a navigation bar or toolbar.

Starting in iOS 5, a search bar can have a custom `backgroundImage`; this will be treated as a resizable image (that is, it will be either stretched or tiled; see [Chapter 15](#)), and overrides the bar style or tint color.

The search field area where the user enters text can be offset with respect to its background with the `searchFieldBackgroundPositionAdjustment` property; you might use this, for example, if you had enlarged the search bar's height and wanted to position the field area within that height. The text can be offset within the field area with the `searchTextPositionAdjustment` property.

You can also replace the image of the search field itself; this is the image that is normally a rounded rectangle. The image you supply will be drawn in front of the background and behind the contents of the search field (such as the text); its width will be adjusted for you, but its height will not be — instead, the image is placed vertically centered where the search field needs to go, and choosing an appropriate height, as well as ensuring a light-colored area in the middle so the user can read the text, is up to you.

A search bar displays an internal cancel button automatically (normally an “x” in a circle) if there is text in the search field. Internally, at its right end, a search bar may display a search results button (`showsSearchResultsButton`), which may be selected or not (`searchResultsButtonSelected`), or a bookmark button (`showsBookmarkButton`); if you ask to display both, you’ll get the search results button. These buttons vanish if text is entered in the search bar so that the cancel button can be displayed. There is also an option to display a Cancel button externally (`showsCancelButton`, or call `setShowsCancelButton:animated:`). The internal cancel button works automatically to remove whatever text is in the field; the other buttons do nothing, but delegate methods notify you when they are tapped:

- `searchBarResultsListButtonClicked:`
- `searchBarBookmarkButtonClicked:`
- `searchBarCancelButtonClicked:`

Starting in iOS 5, you can customize the images used for the internal left icon (the magnifying glass, by default) and any of the internal right icons (the cancel button, the search results button, and the bookmark button) with `setImage:forSearchBarIcon:state:`. The image you supply may be sized down as required, but sometimes it isn’t, so perhaps it is best to size the image first; about 20×20 seems to be a good size. The `state:` value can be `UIControlStateNormal` or `UIControlStateDisabled`; if you don’t supply an image for a state, the default image for that state is used. The icons are specified with constants:

- `UISearchBarIconSearch`
- `UISearchBarIconClear`
- `UISearchBarIconBookmark`
- `UISearchBarIconResultsList`

The position of an icon can be adjusted with `setPositionAdjustment:forSearchBarIcon:`.

A search bar may also display scope buttons (see the example in [Chapter 21](#)). These are intended to let the user alter the meaning of the search; precisely how you use them is up to you. To make the scope buttons appear, use the `showsScopeBar` property; the button titles are the `scopeButtonTitles` property, and the currently selected scope button is the `selectedScopeButtonIndex` property. The delegate is notified when the user taps a different scope button:

- `searchBar:selectedScopeButtonIndexDidChange:`

Starting in iOS 5, the look of the scope bar can be heavily customized. Its background is the `scopeBarBackgroundImage`, which will be stretched or tiled as needed. To set the background of the smaller area constituting the actual buttons, call `setScopeBarButtonBackgroundImage:forState:`; the states are `UIControlStateNormal` and `UIControlStateSelected`. If you don't supply a separate selected image, a darkened version of the normal image is used. If you don't supply a resizable image, the image will be made resizable for you; the runtime decides what region of the image will be stretched behind each button.

The dividers between the buttons are normally vertical lines, but you can customize them as well: call `setScopeBarButtonDividerImage:forLeftSegmentState:rightSegmentState:`. A full complement of dividers consists of three images, one when the buttons on both sides of the divider are normal (unselected) and one each when a button on one side or the other is selected; if you supply an image for just one state combination, it is used for the other two state combinations. The height of the divider image is adjusted for you, but the width is not; you'll normally use an image just a few pixels wide.

The font attributes of the titles of the scope buttons can be customized with respect to their font, color, shadow color, and shadow offset; this is done by calling `setScopeBarButtonTitleTextAttributes:forState:`. The `attributes:` argument is a dictionary whose possible keys will be used for several interface objects in this chapter, so I'll call it a *text attributes dictionary*:

- `UITextAttributeFont`, a `UIFont`; a zero size means “the default size”
- `UITextAttributeTextColor`, a `UIColor`
- `UITextAttributeTextShadowColor`, a `UIColor`
- `UITextAttributeTextShadowOffset`, a `UIOffset` wrapped up as an `NSValue`

(All the customizing `set...` methods I've mentioned have a corresponding getter, whose name is the same without the “set” prefix.)

It may appear that there is no way to customize the external Cancel button, but in fact, although you've no direct access to it through the search bar, the Cancel button is a `UIBarButtonItem` and you can customize it using the `UIBarButtonItem` appearance proxy, discussed later in this chapter.

By combining the various customization possibilities, a completely unrecognizable search bar of inconceivable ugliness can easily be achieved. Let's be careful out there.

The problem of allowing the keyboard to appear without hiding the search bar is exactly as for a text field ([Chapter 23](#)). Text input properties of the search bar configure its keyboard and typing behavior like a text field as well: `keyboardType`, `autocapitalizationType`, and `autocorrectionType` (and, new in iOS, its `spellCheckingType`, which lets you turn spell checking on and off). When the user taps the Search key in the keyboard, the

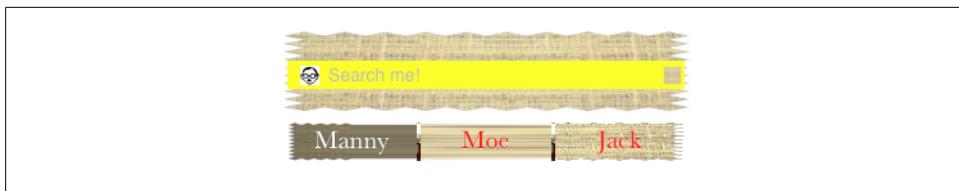


Figure 25-6. A horrible search bar

delegate is notified, and it is then up to you to dismiss the keyboard (`resignFirstResponder`) and perform the search:

- `searchBarSearchButtonClicked`:

A common interface on the iPad is to embed a search bar as a bar button item's view in a toolbar at the top of the screen. This approach has its pitfalls; for example, there is no room for a prompt, and scope buttons or an external Cancel button may not appear either. One rather slimy workaround is to layer the search bar over the toolbar rather than having it genuinely live in the toolbar. Another is to have the search bar itself occupy the position of the toolbar at the top of the screen. On the other hand, a search bar in a toolbar that is managed by a `UISearchDisplayController` will automatically display search results in a popover, which can be a considerable savings of time and effort (though the popover controller is unfortunately out of your hands); see [Chapter 22](#) for an example. An interesting thing about that example, which I didn't mention before, is that the search bar contains a results list button that summons the popover when tapped, and in that case the popover's top bar contains a Clear button that empties the search bar and dismisses the popover; that behavior is apparently entirely automatic and due to the search display controller.

UIControl

`UIControl` is a subclass of `UIView` whose chief purpose is to be the superclass of several further built-in classes and to endow them with common behavior. These are classes representing views with which the user can interact (controls).

The most important thing that controls have in common is that they automatically track and analyze touch events ([Chapter 18](#)) and report them to your code as significant control events by way of action messages. Each control implements some subset of the possible control events; see [Chapter 11](#) for a list of which control events are implemented by which controls. For each control event that you want to hear about automatically, you attach to the control one or more target-action pairs. You can do this in the nib ([Chapter 7](#)) or in code ([Chapter 11](#)).

For any given control, each control event and its target-action pairs form a dispatch table. The following methods permit you to manipulate and query the dispatch table:

- `addTarget:action:forControlEvents:`:

- `removeTarget:action:forControlEvents:`
- `actionsForTarget:forControlEvents:`
- `allTargets`
- `allControlEvents` (a bitmask of control events to which a target–action pair is attached)

An action selector may adopt any of three signatures, whose parameters are (see [Chapter 11](#)):

- The control and the UIEvent
- The control only
- No parameters

To make a control emit its action message to a particular control event right now, in code, call its `sendActionsForControlEvents:` method (which is never called automatically by the framework). For example, suppose you tell a UISwitch programmatically to change its setting from Off to On. This doesn’t cause the switch to report a control event, as it would if the user had slid the switch from off to on; if you wanted it to do so, you could use `sendActionsForControlEvents:`, like this:

```
[switch setOn: YES animated: YES];
[switch sendActionsForControlEvents: UIControlEventValueChanged];
```

You might also use `sendActionsForControlEvents:` in a subclass to customize the circumstances under which a control reports control events.

A control has `enabled`, `selected`, and `highlighted` properties; any of these can be YES or NO independently of the others. These correspond to its `state`, which is reported as a bitmask of three possible values:

- `UIControlStateHighlighted`
- `UIControlStateDisabled`
- `UIControlStateSelected`

A fourth state, `UIControlStateNormal`, corresponding to a zero `state` bitmask, means that `enabled`, `selected`, and `highlighted` are all NO.

A control that is not enabled does not respond to user interaction; whether the control also portrays itself differently, to cue the user to this fact, depends upon the control. For example, a disabled UISwitch is faded. But a round rect text field, unless you explicitly configure it to display a different background image when disabled ([Chapter 23](#)), gives the user no cue that it is disabled. The visual nature of control selection and highlighting, too, depends on the control. Neither highlighting nor selection make any difference to the appearance of a UISwitch, but a highlighted UIButton usually looks quite different from a nonhighlighted UIButton.



Figure 25-7. A switch

A control has `contentHorizontalAlignment` and `contentVerticalAlignment` properties. Again, these matter only if the control has content that can be aligned. You are most likely to use these properties in connection with a `UIButton` to position its title and internal image.

A text field (`UITextField`) is a control; see [Chapter 23](#). The remaining controls are covered here, and then I'll give a simple example of writing your own custom control.

UISwitch

A `UISwitch` portrays a `BOOL` value: it looks like a sliding switch whose positions are labeled ON and OFF, and its `on` property is either YES or NO. The user can slide or tap to toggle the switch's position. When the user changes the switch's position, the switch reports a Value Changed control event. To change the `on` property's value with accompanying animation, call `setOn:animated:`.

A switch has only one size (apparently 79×27); any attempt to set its size will be ignored. In iOS 5 you can set its `onTintColor`, but that's the only customization you can perform. You can't customize the ON and OFF labels; the only solution is to roll your own switch-like interface widget (several third-party implementations are available).



Don't name a `UISwitch` instance variable or property `switch`, as this is a reserved word in C.

UIStepper

A `UIStepper` (new in iOS 5, [Figure 25-8](#)) lets the user increase or decrease a numeric value: it looks like two buttons side by side, one labeled with a minus sign, the other with a plus sign. The user can slide a finger from one button to the other as part of the same interaction with the stepper. It has only one size (apparently 94×27). It maintains a numeric value, which is its `value`. Each time the user increments or decrements the value, it changes by the stepper's `stepValue`. If the `minimumValue` or `maximumValue` is reached, the user can go no further in that direction, and to show this, the corresponding button is disabled — unless the stepper's `wraps` property is YES, in which case the value goes beyond the maximum by starting again at the minimum, and *vice versa*.

As the user changes the stepper's `value`, a Value Changed control event is reported. Portraying the numeric value itself is up to you; you might, for example, use a label or (as in this example) a progress view:



Figure 25-8. A stepper

```
- (IBAction)doStep:(UIStepper*)step {  
    self->prog.progress = step.value / (step.maximumValue - step.minimumValue);  
}
```

If a stepper's `continuous` is YES (the default), a long touch on one of the buttons will update the value repeatedly; the updates start slowly and get faster. If the stepper's `autorepeat` is NO, the updated value is not reported as a Value Changed control event until the entire interaction with the stepper ends; the default is YES.

UIPageControl

A UIPageControl is a row of dots; each dot is called a *page*, because it is intended to be used in conjunction with some other interface that portrays something analogous to pages, such as a UIScrollView with its `pagingEnabled` set to YES. Coordinating the page control with this other interface is up to you. (See [Chapter 20](#) for an example and [Figure 20-3](#) for an illustration.)

The number of dots is the page control's `numberOfPages`; this should be small, as the dots need to fit within the page control's bounds. The current page, its `currentPage`, is portrayed as a solid dot; the others are slightly transparent. The user can tap to one side or the other of the current page's dot to increment or decrement the current page; the page control then reports a Value Changed control event.

You can make the page control wider than the dots to increase the target region on which the user can tap. You can make the page control taller as well, but only the horizontal component of a tap is taken into account, so this would probably be pointless as well as confusing to the user. To learn the minimum size required for a given number of pages, call `sizeForNumberOfPages:`.

If a page control's `hidesForSinglePage` is YES, the page control becomes invisible when its `numberOfPages` changes to 1.

If a page control's `defersCurrentPageDisplay` is YES, then when the user taps to increment or decrement the page control's value, the display of the current page is not changed. A Value Changed control event is reported, but it is up to your code to handle this action and call `updateCurrentPageDisplay`. A case in point might be if the user's changing the current page starts an animation, but you don't want the current page dot to change until the animation ends.

UIDatePicker

A UIDatePicker looks like a UIPickerView (discussed earlier in this chapter), but it is *not* a UIPickerView subclass (it *uses* a UIPickerView to draw itself, but it provides no official access to that picker view). Its purpose is to express the notion of a date and time, taking care of the calendrical and numerical complexities so that you don't have to. When the user changes its setting, the date picker reports a Value Changed control event.

A UIDatePicker has one of four modes (`datePickerMode`), determining how it is drawn:

UIDatePickerModeTime

The date picker displays a time; for example, it has an hour component and a minutes component.

UIDatePickerModeDate

The date picker displays a date; for example, it has a month component, a day component, and a year component.

UIDatePickerModeDateAndTime

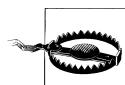
The date picker displays a date and time; for example, it has a component showing day of the week, month, and day, plus an hour component and a minutes component.

UIDatePickerModeCountDownTimer

The date picker displays a number of hours and minutes; for example, it has an hours component and a minutes component.

Exactly what components a date picker displays, and what values they contain, depends by default upon system settings. For example, a U.S. time displays an hour (numbered 1 through 12), minutes, and AM or PM, but a British time displays an hour (numbered 1 through 24) and minutes. If your app contains a date picker displaying a time, and the user changes the system region format on the device from United States to United Kingdom, the date picker's display will change immediately, eliminating the AM/PM component and changing the hour numbers to run from 1 to 24.

A date picker has `calendar` and `timeZone` properties, respectively an `NSCalendar` and an `NSTimeZone`; these are nil by default, meaning that the date picker responds to the user's system-level settings. You can also change these values manually; for example, if you live in California and you set a date picker's `timeZone` to GMT, the displayed time is shifted forward by 8 hours, so that 11 AM is displayed as 7 PM (if it is winter).

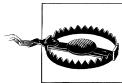


Don't change the `timeZone` of a `UIDatePickerModeCountDownTimer` date picker, or the displayed value will be shifted and you will confuse the heck out of yourself and your users.

The minutes component, if there is one, defaults to showing every minute, but you can change this with the `minuteInterval` property. The maximum value is 30, in which case the minutes component values are 0 and 30.

The maximum and minimum values enabled in the date picker are determined by its `maximumDate` and `minimumDate` properties. Values outside this range may appear disabled. There isn't really any practical limit on the range that a date picker can display, because the “drums” representing its components are not physical, and values are added dynamically as the user spins them. In this example, we set the initial minimum and maximum dates of a date picker (`dp`) to the beginning and end of 1954. We also set the actual date, because otherwise the date picker will appear initially set to now, which will be disabled because it isn't within the minimum–maximum range:

```
NSDateComponents* dc = [[NSDateComponents alloc] init];
[dc setYear:1954];
[dc setMonth:1];
[dc setDay:1];
NSCalendar* c = [[NSCalendar alloc] initWithCalendarIdentifier:NSGregorianCalendar];
NSDate* d = [c dateFromComponents:dc];
dp.minimumDate = d;
dp.date = d;
[dc setYear:1955];
d = [c dateFromComponents:dc];
dp.maximumDate = d;
```



Don't set the `maximumDate` and `minimumDate` properties values for a `UIDatePickerModeCountDownTimer` date picker, or you might cause a crash with an out-of-range exception.

The date represented by a date picker (unless its mode is `UIDatePickerModeCountDown-Timer`) is its `date` property, an `NSDate`. The default date is now, at the time the date picker is instantiated. For a `UIDatePickerModeDate` date picker, the time by default is 12 AM (midnight), local time; for a `UIDatePickerModeTime` date picker, the date by default is today. The internal value is reckoned in the local time zone, so it may be different from the displayed value, if you have changed the date picker's `timeZone`.

The value represented by a `UIDatePickerModeCountDownTimer` date picker is its `countDownDuration`. The date picker does not actually do any counting down; changing its `countDownDuration` at appropriate intervals, if desired, is up to you, though you are more likely to use some other interface to display the countdown, especially because the date picker doesn't display seconds. The Timer tab of Apple's Clock app shows a typical interface; the user configures the date picker to set the `countDownDuration` initially, but once the counting starts, the date picker is hidden and a label displays the remaining time. The `countDownDuration` is an `NSTimeInterval`, which is a double representing a number of seconds; dividing by 60 to convert to minutes, and again to convert to hours, is up to you — or you could use the built-in calendrical classes:

```

NSTimeInterval t = [datePicker countDownDuration];
NSDate* d = [NSDate dateWithTimeIntervalSinceReferenceDate:t];
NSCalendar* c = [[NSCalendar alloc] initWithCalendarIdentifier:NSGregorianCalendar];
[c setTimeZone: [NSTimeZone timeZoneForSecondsFromGMT:0]]; // normalize
NSUInteger units = NSHourCalendarUnit | NSMinuteCalendarUnit;
NSDateComponents* dc = [c components:units fromDate:d];
NSLog(@"%@", [dc hour], [dc minute]);

```

Similarly, to convert between an NSDate and a string, you'll need an NSDateFormatter (see [Chapter 10](#), and Apple's *Date and Time Programming Guide*):

```

NSDate* d = [datePicker date];
NSDateFormatter* df = [[NSDateFormatter alloc] init];
[df setTimeStyle:kCFDateFormatFullStyle];
[df setDateStyle:kCFDateFormatFullStyle];
NSLog(@"%@", [df stringFromDate:d]);
// "Wednesday, August 10, 2011 3:16:25 AM Pacific Daylight Time"

```

UISlider

A slider (UISlider) is an expression of a continuously settable value (its `value`) between some minimum and maximum (its `minimumValue` and `maximumValue`; they are 0 and 1 by default). It is portrayed as an object, the *thumb*, positioned along a *track*. As the user changes the thumb's position, the slider reports a Value Changed control event; it may do this continuously as the user presses and drags the thumb (if the slider's `continuous` is YES, the default) or only when the user releases the thumb (if its `continuous` is NO). While the user is pressing on the thumb, the slider is in the `highlighted` state. To change the slider's value with animation, call `setValue:animated:`.

A commonly expressed desire is to modify a slider's behavior so that if the user taps on its track, the slider moves to the spot where the user tapped. Unfortunately, a slider does not, of itself, respond to taps on its track; such a tap doesn't even cause it to report a Touch Up Inside. However, with a gesture recognizer, most things are possible; here's the action handler for a UITapGestureRecognizer attached to a UISlider:

```

- (void) tapped: (UITapGestureRecognizer*) g {
    UISlider* s = (UISlider*)g.view;
    if (s.highlighted)
        return; // tap on thumb, let slider deal with it
    CGPoint pt = [g locationInView: s];
    CGFloat percentage = pt.x / s.bounds.size.width;
    CGFloat delta = percentage * (s.maximumValue - s.minimumValue);
    CGFloat value = s.minimumValue + delta;
    [s setValue:value animated:YES];
}

```

A slider's appearance is extremely customizable. You can change the color of the thumb and the track on either side of it (`thumbTintColor`, `minimumTrackTintColor`, and `maximumTintColor`, all new in iOS 5), or you can go even further and provide your own thumb image and your own track image, along with images to appear at each end of the track, and you can override in a subclass the methods that position these.

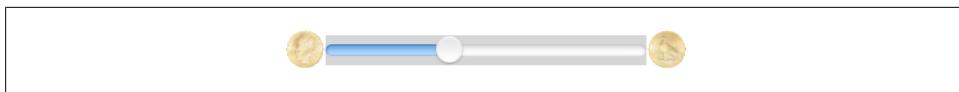


Figure 25-9. Repositioning a slider's images and track

The images at the ends of the track are the slider's `minimumValueImage` and `maximumValueImage`, and they are `nil` by default. If you set them to actual images (which can also be done in the nib), the slider will attempt to position them within its own bounds, shrinking the drawing of the track to compensate. The slider does not clip its subviews by default, so the images can extend outside the slider's bounds.

For example, suppose the slider's dimensions are 250×23 (the standard height), and suppose the images are 30×30 . Then the minimum image is drawn with its origin at $\{0, -4\}$ — its left edge matches the slider's left edge, and its top is raised so that the center of its height matches the center of the slider's height — and the maximum image is drawn with its origin at $\{220, -4\}$. But the track is drawn with a width of only 164 pixels, instead of the normal 246; that is, instead of being nearly the full width of the slider, the track is contracted to allow room for the images. (This, by the way, wrecks the behavior of the `tapped:` handler in the previous example, which relies on the actual bounds of the slider, not the apparent width and position of the track. Fixing the example is left as an exercise for the reader.)

You can change these dimensions by overriding `minimumValueImageRectForBounds:`, `maximumValueImageRectForBounds:`, and `trackRectForBounds:` in a subclass. The bounds passed in are the slider's bounds. In this example, we expand the track width to the full width of the slider, and draw the images outside the slider's bounds (Figure 25-9; I've given the slider a gray background color so you can see how the track and images are related to its bounds):

```
- (CGRect)maximumValueImageRectForBounds:(CGRect)bounds {
    CGRect result = [super maximumValueImageRectForBounds:bounds];
    result = CGRectMakeOffset(result, 31, 0);
    return result;
}

- (CGRect)minimumValueImageRectForBounds:(CGRect)bounds {
    CGRect result = [super minimumValueImageRectForBounds:bounds];
    result = CGRectMakeOffset(result, -31, 0);
    return result;
}

- (CGRect)trackRectForBounds:(CGRect)bounds {
    CGRect result = [super trackRectForBounds:bounds];
    result.origin.x = 0;
    result.size.width = bounds.size.width;
    return result;
}
```

The thumb is also an image, and you set it with `setThumbImage:forState:`. There are two chiefly relevant states, `UIControlStateNormal` (not highlighted) and `UIControlState-`



Figure 25-10. Replacing a slider's thumb

Highlighted, so if you supply images for both, the thumb will change automatically while the user is dragging it. If you supply just an image for the normal state, the thumb image won't change while the user is dragging it. By default, the image will be centered in the track at the point represented by the slider's current value; you can shift this position by overriding `thumbRectForBounds:trackRect:value:` in a subclass. In this example, the image is repositioned upward slightly (Figure 25-10):

```
- (CGRect)thumbRectForBounds:(CGRect)bounds
    trackRect:(CGRect)rect value:(float)value {
    CGRect result = [super thumbRectForBounds:bounds trackRect:rect value:value];
    result = CGRectMakeOffset(result, 0, -7);
    return result;
}
```

Enlarging a slider's thumb can mislead the user as to the area on which it can be tapped to drag it. The slider, not the thumb, is the touchable UIControl; so if the slider's height is 23 pixels, only the part of the thumb that intersects that 23-pixel height will be draggable. The user may try to drag the part of the thumb that is drawn outside the slider's bounds, and will fail (and be confused). The solution is to increase the slider's height; you can't do this in the nib editor, but you can do it in code.

The track is two images, one appearing to the left of the thumb, the other to its right. They are set with `setMinimumTrackImage:forState:` and `setMaximumTrackImage:forState:`. If you supply images both for normal state and for highlighted state, the images will change while the user is dragging the thumb.

The images should be resizable (Chapter 15), because that's how the slider performs the clever illusion of making it look like the user is dragging the thumb along a single static track. In reality, there are two images; as the user drags the thumb, one image grows horizontally and the other shrinks horizontally. For the left track image, the right end cap inset will be partially or entirely hidden under the thumb; for the right track image, the left end cap inset will be partially or entirely hidden under the thumb. Figure 25-11 shows a track derived from a single 15×15 image of a circular object (a coin):

```
UIImage* coin = [UIImage imageNamed: @"coin.png"];
UIImage* coinEnd = [coin resizableImageWithCapInsets:UIEdgeInsetsMake(0,7,0,7)];
[slider setMinimumTrackImage:coinEnd forState:UIControlStateNormal];
[slider setMaximumTrackImage:coinEnd forState:UIControlStateNormal];
```

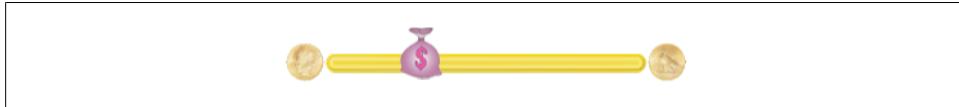


Figure 25-11. Replacing a slider's track



Figure 25-12. A segmented control

UISegmentedControl

A segmented control (`UISegmentedControl`) is a row of tappable segments; a segment is rather like a button. This provides a way for the user to choose among several related options. By default (`momentary` is NO), the most recently tapped segment remains selected; alternatively (`momentary` is YES), the tapped segment is shown as selected momentarily, but then no segment selection is displayed, though internally the tapped segment remains the selected segment. The selected segment can be retrieved with the `selectedSegmentIndex` property; it can also be set with the `selectedSegmentIndex` property, and remains visibly selected (even for a `momentary` segmented control). A `selectedSegmentIndex` value of `UISegmentedControlNoSegment` means no segment is selected. When the user taps a segment that is not already visibly selected, the segmented control reports a Value Changed event.

A segment can be separately enabled or disabled with `setEnabled:forSegmentAtIndex:`, and its enabled state can be retrieved with `isEnabledForSegmentAtIndex:`. A disabled segment is drawn faded, and the user can't tap it, but it can still be selected in code.

A segment has either a title or an image; when one is set, the other becomes nil. The methods for setting and fetching the title and image for existing segments are:

- `setTitle:forSegmentAtIndex:`
- `setImage:forSegmentAtIndex:`
- `titleForSegmentAtIndex:`
- `imageForSegmentAtIndex:`

After changing an existing segment's title or image, you might want to call `sizeToFit` to resize the segments automatically.

You will also want to set the title or image when creating the segment. You can do this in code if you're creating the segmented control from scratch, with `initWithItems:`, which takes an array each item of which is either a string or an image.

Methods for managing segments dynamically are:

- `insertSegmentWithTitle:atIndex:animated:`
- `insertSegmentWithImage:atIndex:animated:`
- `removeSegmentAtIndex:animated:`
- `removeAllSegments`

The number of segments can be retrieved with the read-only `numberOfSegments` property.

A segment's width is adjusted automatically when you create it or call `sizeToFit`, or you can set it manually with `setWidth:forSegmentAtIndex:` (and retrieve it with `widthForSegmentAtIndex:`). Alternatively, if you set a segment's width to 0, the system will adjust the width for you if the segmented control's `apportionsSegmentWidthsByContent` property is YES (new in iOS 5).

You can also change the position of the content (title or image) within a segment. To set this position in code, call `setContentOffset:forSegmentAtIndex:` (and retrieve it with `contentOffsetForSegmentAtIndex:`), where the offset is expressed as a `CGSize` describing how much to move the content from its default centered position.

A segmented control's height is standard in accordance with its style. You can change a segmented control's height in code, but of course if you later call `sizeToFit`, it will resume its standard height.

A segmented control comes in a choice of styles (its `segmentedControlStyle`):

`UISegmentedControlStylePlain`

Large default height (44 pixels) and large titles. Deselected segments are gray; the selected segment is blue and has a depressed look.

`UISegmentedControlStyleBordered`

Just like `UISegmentedControlStylePlain`, but a dark border emphasizes the segmented control's outline.

`UISegmentedControlStyleBar`

Small default height (30 pixels) and small titles. All segments are blue, but you can change this by setting the `tintColor`; the selected segment is slightly darker.

`UISegmentedControlStyleBezeled`

Large default height (44 pixels) and small titles. Similar to `UISegmentedControlStyleBar`. All segments are blue, but you can change this by setting the `tintColor`; the selected segment is brighter.

Starting in iOS 5, you can further customize a segmented control's appearance. The methods involved here are parallel to those for setting the look of the scope bar portion of a search bar, described earlier in this chapter. You can set the overall background, the divider image, the text attributes for the segment titles, and the position of segment contents:

- `setBackgroundImage:forState:barMetrics:`

- `setDividerImage:forLeftSegmentState:rightSegmentState:barMetrics:`
- `setTitleTextAttributes:forState:`
- `setContentPositionAdjustment:forSegmentType:barMetrics:`

The `segmentType:` parameter in `setContentPositionAdjustment:forSegmentType:barMetrics:` is needed because the segments at the two extremes have rounded ends (and, if a segment is the lone segment, both its ends are rounded). The argument allows you distinguish between the various possibilities:

- `UISegmentedControlSegmentAny`
- `UISegmentedControlSegmentLeft`
- `UISegmentedControlSegmentCenter`
- `UISegmentedControlSegmentRight`
- `UISegmentedControlSegmentAlone`

The `barMetrics:` parameter will recur later in this chapter, in the discussion of navigation bars, toolbars, and bar button items; for a segmented control, its value matters only if the segmented control has `UISegmentedControlStyleBar` and only if it is being used inside a navigation bar or toolbar. Otherwise, use `UIBarMetricsDefault`. See “[Landscape iPhone Bar Metrics](#)” on page 732 for more information.

UIButton

A button (`UIButton`) is a fundamental tappable control; its appearance is extremely flexible. It is endowed at creation with a type. The code creation method is a class method, `buttonWithType::`. The types are:

`UIButtonTypeCustom`

Could be completely invisible, if the `backgroundColor` is `clearColor` and there's no title or other content. If a `backgroundColor` is supplied, a thin, subtle rectangular border is also present; you can add more of a border by modifying the button's `layer`. Alternatively, you can provide a background image, thus making the button appear to be any shape you like (though this does not automatically affect its tapable region).

`UIButtonTypeDetailDisclosure`

`UIButtonTypeContactAdd`

`UIButtonTypeInfoLight`

`UIButtonTypeInfoDark`

Basically, these are all `UIButtonTypeCustom` buttons whose image is set automatically to standard button images: a right-pointing chevron, a plus sign, a light letter “i,” and a dark letter “i,” respectively.

`UIButtonTypeRoundedRect`

A rounded rectangle with a white background and an antialiased gray border. However, supplying a rectangular opaque background image results in a rectangle similar to a `UIButtonTypeCustom` button. (A rounded rect button is actually an instance of a `UIButton` subclass, `UIRoundedRectButton`, but you're probably not supposed to know that.)

New in iOS 5, a rounded rect button can have a `tintColor` which fills the button only while the button is highlighted.

A button has a title, a title color, a title shadow color, an image, and a background image. The background image, if any, is stretched to fit the button's bounds. The image, on the other hand, if smaller than the button, is not resized, and is thus shown internally within the button. The button can have both a title and an image, if the image is small enough; in that case, the image is shown to the left of the title by default.

These five features (title, title color, title shadow color, image, and background image) can all be made to vary depending on the button's current state: `UIControlStateHighlighted`, `UIControlStateSelected`, `UIControlStateDisabled`, and `UIControlStateNormal`. The button can be in more than one state at once, except for `UIControlStateNormal` which effectively means “none of the other states”. A state change, whether automatic (the button is highlighted while the user is tapping it) or programmatically imposed, will thus in and of itself alter a button's appearance. To make this possible, the methods for setting these button features all involve specifying a corresponding state — or multiple states, using a bitmask:

- `setTitle:forState:`
- `setTitleColor:forState:` (by default, the title color is white when the button is highlighted)
- `setTitleShadowColor:forState:`
- `setImage:forState:`
- `setBackgroundImage:forState:`

Similarly, when getting these button features, you must either use a method to specify a single state you're interested in or use a property to ask about the feature as currently displayed:

- `titleForState:`
- `titleColorForState:`
- `titleShadowColorForState:`
- `imageForState:`
- `backgroundImageForState:`
- `currentTitle`
- `currentTitleColor`

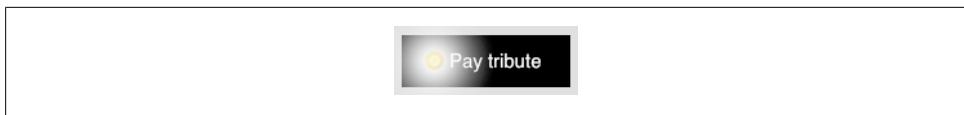


Figure 25-13. A button with highlighted glow

- `currentTitleShadowColor`
- `currentImage`
- `currentBackgroundImage`

If you don't specify a feature for a particular state, or if the button adopts more than one state at once, an internal heuristic is used to determine what to display. I can't describe all possible combinations, but here are some general observations:

- If you specify a feature for a particular state (highlighted, selected, or disabled), and the button is in only that state, that feature will be used.
- If you *don't* specify a feature for a particular state (highlighted, selected, or disabled), and the button is in only that state, the normal version of that feature will be used as fallback. (That's why many examples earlier in this book have assigned a title for `UIControlStateNormal` only; this is sufficient to give the button a title in every state.)
- Combinations of states often cause the button to fall back on the feature for normal state. For example, if a button is both highlighted and selected, the button will display its normal title, even if it has a highlighted title, a selected title, or both.

In addition, a `UIButton` has some properties determining how it draws itself in various states, which can save you the trouble of specifying different images for different states:

`showsTouchWhenHighlighted`

If YES, then the button projects a circular white glow when highlighted. If the button has an internal image, the glow is centered behind it (Figure 25-13); thus, this feature is suitable particularly if the button image is small and circular; for example, it's the default behavior for a `UIButtonTypeInfoLight` or `UIButtonTypeInfoDark` button. (If the button has no internal image, the glow is centered at the button's center.) The glow is drawn on top of the background image or color, if any.

`adjustsImageWhenHighlighted`

If YES (the default), then if there is no separate highlighted image (and if `showsTouchWhenHighlighted` is NO), the normal image is darkened when the button is highlighted. This applies equally to the internal image and the background image.

`adjustsImageWhenDisabled`

If YES (the default), then if there is no separate disabled image, the normal image is lightened (faded) when the button is disabled. This applies equally to the internal image and the background image.

The title is a `UILabel` ([Chapter 23](#)), and the label features of the title can be accessed through the button's `titleLabel`. Thus, for example, you can set the title's `font`, `lineBreakMode`, and `shadowOffset`. If the `shadowOffset` is not `{0,0}`, then the title has a shadow, and the title shadow color feature comes into play; the button's `reversesTitleShadowWhenHighlighted` property also applies: if YES, the `shadowOffset` values are replaced with their additive inverses when the button is highlighted.

An easy way to make a button's title consist of multiple lines is to set the button's `titleLabel.lineBreakMode` to `UILineBreakModeWordWrap` and put manual line breaks into the button's title: `@"This is a line\nand this is a line"`. (To insert a line break in the nib editor, type Option-Return.)

The internal image is drawn by a `UIImageView` ([Chapter 15](#)) whose features can be accessed through the button's `imageView`. Thus, for example, you can change the internal image view's `alpha` to make the image more transparent.

The internal position of the image and title as a whole are governed by the button's `contentVerticalAlignment` and `contentHorizontalAlignment` (recall that these properties are inherited from `UIControl`). You can also tweak the position of the image and title, together or separately, by setting the button's `contentEdgeInsets`, `titleEdgeInsets`, or `imageEdgeInsets`. Increasing an inset component increases that margin; thus, for example, a positive `top` component makes the distance between that object and the top of the button larger than normal (where "normal" is where the object would be according to the alignment settings). The `titleEdgeInsets` or `imageEdgeInsets` values are added to the overall `contentEdgeInsets` values. So, for example, if you really wanted to, you could make the internal image appear to the right of the title by decreasing the left `titleEdgeInsets` and increasing the left `imageEdgeInsets`.

Four methods also provide access to the button's positioning of its elements:

- `titleRectForContentRect`:
- `imageRectForContentRect`:
- `contentRectForBounds`:
- `backgroundRectForBounds`:

These methods are called whenever the button is redrawn, including every time it changes state. The content rect is the area in which the title and image are placed. By default, `contentRectForBounds`: and `backgroundRectForBounds`: yield the same result.

You can override these methods in a subclass to change the way the button's elements are positioned. In this example, we shrink the button slightly when highlighted as a way of providing feedback:



Figure 25-14. A button with a stretched background image

```
- (CGRect)backgroundRectForBounds:(CGRect)bounds {
    CGRect result = [super backgroundRectForBounds:bounds];
    if (self.highlighted)
        result = CGRectInset(result, 3, 3);
    return result;
}
```

A button's background image is stretched if the image is smaller, in both dimensions, than the button's `backgroundRectForBounds:`. You can take advantage of this stretching, for example, to construct a rounded rectangle background for the button by supplying a resizable image. In this example (Figure 25-14), both the internal image and the background image are generated from the same image (which is in fact the same image used to generate the track in Figure 25-11):

```
UIImage* im = [UIImage imageNamed: @"coin.png"];
CGSize sz = [im size];
UIImage* im2 = [im resizableImageWithCapInsets:UIEdgeInsetsMake(
    sz.height/2.0, sz.width/2.0, sz.height/2.0, sz.width/2.0)];
[button setBackgroundImage: im2 forState: UIControlStateNormal];
button.backgroundColor = [UIColor clearColor];
```

Custom Controls

The `UIControl` class implements several touch-tracking methods that you might override in order to customize a built-in `UIControl` type or to create your own `UIControl` subclass, along with properties that tell you whether touch tracking is going on:

- `beginTrackingWithTouch:withEvent:`
- `continueTrackingWithTouch:withEvent:`
- `endTrackingWithTouch:withEvent:`
- `cancelTrackingWithEvent:`
- `tracking` (property)
- `touchInside` (property)

With the advent of gesture recognizers (Chapter 18), such direct involvement with touch tracking is probably less needed than it used to be, especially if your purpose is to modify the behavior of a built-in `UIControl` subclass. So, to illustrate their use, I'll give a simple example of creating a custom control. The main reason for doing this (rather than using, say, a `UIView` and gesture recognizers) would probably be to obtain the convenience of control events. Also, the touch-tracking methods, though of course nowhere near as high-level as gesture recognizers, are at least a level up from the



Figure 25-15. A custom control

UIResponder methods `touchesBegan:withEvent:` and so forth ([Chapter 18](#)): they track a single touch, and both `beginTracking...` and `continueTracking...` return a BOOL, giving you a chance to stop tracking the current touch.

We'll build a simplified knob control ([Figure 25-15](#)). The control starts life at its minimum position, with an internal angle value of 0; it can be rotated clockwise with a single finger as far as its maximum position, with an internal angle value of 5 (radians). To keep things simple, the words "Min" and "Max" appearing in the interface are actually labels; the control just draws the knob, and to rotate it we'll simply apply a rotation transform.

Our control is a UIControl subclass, MyKnob. It has a CGFloat property `angle`, and a CGFloat instance variable `initialAngle` that we'll use internally during rotation. Because a UIControl is a UIView, it can draw itself, which it does with a UIImage included in our app bundle:

```
- (void) drawRect:(CGRect)rect {
    UIImage* knob = [UIImage imageNamed:@"knob.png"];
    [knob drawInRect:rect];
}
```

We'll need a utility function for transforming a touch's Cartesian coordinates into polar coordinates, giving us the angle to be applied as a rotation to the view:

```
static CGFloat pToA (UITouch* touch, UIView* self) {
    CGPoint loc = [touch locationInView: self];
    CGPoint c = CGPointMake(CGRectGetMidX(self.bounds),
                           CGRectGetMidY(self.bounds));
    return atan2(loc.y - c.y, loc.x - c.x);
}
```

Now we're ready to override the tracking methods. `beginTrackingWithTouch:withEvent:` simply notes down the angle of the initial touch location. `continueTrackingWithTouch:withEvent:` uses the difference between the current touch location's angle and the initial touch location's angle to apply a transform to the view, and updates the `angle` property. `endTrackingWithTouch:withEvent:` triggers the Value Changed control event. So our first draft looks like this:

```

- (BOOL) beginTrackingWithTouch:(UITouch *)touch withEvent:(UIEvent *)event {
    self->initialAngle = pToA(touch, self);
    return YES;
}

- (BOOL) continueTrackingWithTouch:(UITouch *)touch withEvent:(UIEvent *)event {
    CGFloat ang = pToA(touch, self);
    ang -= self->initialAngle;
    CGFloat absoluteAngle = self->angle + ang;
    self.transform = CGAffineTransformRotate(self.transform, ang);
    self->angle = absoluteAngle;
    return YES;
}

- (void) endTrackingWithTouch:(UITouch *)touch withEvent:(UIEvent *)event {
    [self sendActionsForControlEvents:UIControlEventValueChanged];
}

```

This works: we can put a MyKnob into the interface and hook up its Value Changed control event (this can be done in the nib editor), and sure enough, when we run the app, we can rotate the knob and, when our finger lifts from the knob, the Value Changed action handler is called. However, `continueTrackingWithTouch:withEvent:` needs modification.

First, we need to peg the minimum and maximum rotation at 0 and 5, respectively. For simplicity, we'll just stop tracking, by returning NO, if the rotation goes below 0 or above 5, fixing the angle at the exceeded limit. However, because we're no longer tracking, `endTracking...` will never be called, so we also need to trigger the Value Changed control event. (Doubtless you can come up with a more sophisticated way of pegging the knob at its minimum and maximum, but remember, this is only a simple example.) Second, it might be nice to give the programmer the option to have the Value Changed control event reported continuously as `continueTracking...` is called repeatedly. So we'll add a `continuous` BOOL property and obey it.

Here, then, is our revised `continueTracking...` implementation:

```

- (BOOL) continueTrackingWithTouch:(UITouch *)touch withEvent:(UIEvent *)event {
    CGFloat ang = pToA(touch, self);
    ang -= self->initialAngle;
    CGFloat absoluteAngle = self->angle + ang;
    if (absoluteAngle < 0) {
        self.transform = CGAffineTransformIdentity;
        self->angle = 0;
        [self sendActionsForControlEvents:UIControlEventValueChanged];
        return NO;
    }
    if (absoluteAngle > 5) {
        self.transform = CGAffineTransformMakeRotation(5);
        self->angle = 5;
        [self sendActionsForControlEvents:UIControlEventValueChanged];
        return NO;
    }
    self.transform = CGAffineTransformRotate(self.transform, ang);
}

```

```

    self->angle = absoluteAngle;
    if (self->continuous)
        [self sendActionsForControlEvents:UIControlEventValueChanged];
    return YES;
}

```

Finally, we'll probably want to be able to set the angle programmatically as a way of rotating the knob:

```

- (void) setAngle: (CGFloat) ang {
    if (ang < 0)
        ang = 0;
    if (ang > 5)
        ang = 5;
    self.transform = CGAffineTransformMakeRotation(ang);
    self->angle = ang;
}

```

This is more work than using a gesture recognizer (which is left as an exercise for the reader), but not much, and it gives a sense of what's involved in creating a custom control.

Bars

As you saw in [Chapter 19](#), the three bar types — UINavigationBar, UIToolbar, and UITabBar — are often used in conjunction with a built-in view controller:

- A UINavigationController has a UINavigationBar, which it can display at the top of its view.
- A UINavigationController has a UIToolbar, which it can display at the bottom of its view.
- A UITabBarController has a UITabBar, which it displays at the bottom of its view.

You can also use these bar types independently, though you're not very likely to do that except with a UIToolbar, which is often used on the iPad as a top bar, adopting a role analogous to a menu bar on Mac OS X. That's such a common interface, in fact, that certain special automatic behaviors are associated with it; for example, a UISearchBar in a UIToolbar and managed by a UISearchDisplayController will automatically display its search results table in a popover ([Chapter 22](#)), which is different from what happens if the UISearchBar is *not* in a UIToolbar.

This section summarizes the facts about the three bar types and the items that populate them.

UINavigationBar

A UINavigationBar is populated by UINavigationItems. The UINavigationBar maintains a stack; UINavigationItems are pushed onto and popped off of this stack. Whatever UINavigationItem is currently topmost in the stack (the UINavigationBar's top-

`Item`), in combination with the `UINavigationItem` just beneath it in the stack (the `UINavigationBar`'s `backItem`), determines what appears in the navigation bar:

- The `title` (string) or `titleView` (`UIView`) of the `topItem` appears in the center of the navigation bar.
- The `prompt` (string) of the `topItem` appears at the top of the navigation bar.
- The `rightBarButtonItem` and `leftBarButtonItem` appear at the right and left ends of the navigation bar. These are `UIBarButtonItem`s. A `UIBarButtonItem` can be a system button, a titled button, an image button, or a container for a `UIView`. A `UIBarButtonItem` is not itself a `UIView`, however. I'll discuss bar button items further in a moment (and refer also to the discussion in [Chapter 19](#)).

In iOS 5, a `UINavigationItem` can have multiple right bar button items and multiple left bar button items; its `rightBarButtonItems` and `leftBarButtonItems` properties are arrays (of `UIBarButtonItem`s). The bar button items are displayed from the outside in: that is, the first item in the `leftBarButtonItems` is leftmost, while the first item in the `rightBarButtonItems` is rightmost. Even if there are multiple buttons on a side, you can still speak of that button in the singular: the `rightBarButtonItem` is the first item of the `rightBarButtonItems` array, and the `leftBarButtonItem` is the first item of the `leftBarButtonItems` array.

- The `backBarButtonItem` of the `backItem` appears at the left end of the navigation bar. It typically points to the left, and is automatically configured so that, when tapped, the `topItem` is popped off the stack. If the `backItem` has no `backBarButtonItem`, then there is *still* a back button at the left end of the navigation bar, taking its title from the `title` of the `backItem`. However, if the `topItem` has its `hidesBackButton` set to YES, the back button is suppressed. Also, unless the `topItem` has its `leftItemsSupplementBackButton` set to YES (new in iOS 5), the back button is suppressed if the `topItem` has a `leftBarButtonItem`.

Changes to the navigation bar's buttons can be animated by sending its `topItem` any of these messages:

- `setRightBarButtonItem:animated:`
- `setLeftBarButtonItem:animated:`
- `setRightBarButtonItems:animated:`
- `setLeftBarButtonItems:animated:`
- `setHidesBackButton:animated:`

`UINavigationItems` are pushed and popped with `pushNavigationItem:animated:` and `popNavigationItemAnimated:`, or you can set all items on the stack at once with `setItems:animated:` or by directly setting the `items`.

A `UINavigationBar` can be styled using its `barStyle`, `translucent`, and `tintColor` properties. Possible `barStyle` values are `UIBarStyleDefault` and `UIBarStyleBlack`; setting a `tintColor` overrides the `barStyle`. In iOS 5 you can go further, providing a background

Landscape iPhone Bar Metrics

Several methods that customize the look of a navigation bar, a toolbar, a bar button item, or a segmented control depend upon its bar metrics (`barMetrics:`), which can be `UIBarMetricsDefault` or `UIBarMetricsLandscapePhone`. This is intended to cover the case of a navigation bar or toolbar belonging to a `UINavigationController`. On the iPhone (not on the iPad), the `UINavigationController` changes the height of the bar when its interface rotates; the standard heights are 44 (portrait) and 32 (landscape). The idea is that you might like the look of whatever you’re customizing to change as well. The same distinction is drawn by a bar button item’s `landscapeImagePhone...` properties.

Under the hood, the navigation bar or toolbar is simply responding to the change in its height, no matter how that change is caused; thus you could in fact trigger the same response in an independent bar (not belonging to a navigation controller). My experiments suggest that a height of 32 or less is considered `UIBarMetricsLandscapePhone`, while a height of 34 or more is considered `UIBarMetricsDefault`.

`image (setBackgroundImage:forBarMetrics:)` and customizing the title’s text attributes (`titleTextAttributes`). You can also shift the title’s vertical position by calling `setTitleVerticalPositionAdjustment:forBarMetrics:`.

Recall from [Chapter 19](#) that a bar button item may be instantiated with any of five methods:

- `initWithBarButtonSystemItem:target:action:`
- `initWithTitle:style:target:action:`
- `initWithImage:style:target:action:`
- `initWithImage:landscapeImagePhone:style:target:action:`
- `initWithCustomView:`

The styles are:

- `UIBarButtonItemStyleBordered`
- `UIBarButtonItemStylePlain` (converted to `UIBarButtonItemStyleBordered` in a navigation bar)
- `UIBarButtonItemStyleDone` (only in a navigation bar)

In addition to its `title` and `image` (and, on iOS 5, its `landscapeImagePhone`), a bar button item inherits from `UIBarButtonItem` the ability to adjust the image position with `imageInsets` (and `landscapeImagePhoneInsets`), plus the `enabled` and `tag` properties. Recall from [Chapter 19](#) that you can also set a bar button item’s `possibleTitles` and `width` properties.

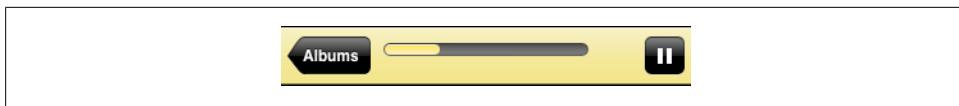


Figure 25-16. A colorful navigation bar

You can also customize the look of bar button items in iOS 5. It has a `tintColor` property, or you can give it a background image; and you can apply text attributes to its title. These are the customization methods:

- `setTitleTextAttributes:forState:` (inherited from `UIBarButtonItem`)
- `setTitlePositionAdjustment:forBarMetrics:`
- `setBackgroundImage:forState:barMetrics:`
- `setBackgroundVerticalPositionAdjustment:forBarMetrics:`

In addition, these methods apply only if the bar button item is being used as a back button item:

- `setBackButtonTitlePositionAdjustment:forBarMetrics:`
- `setBackButtonBackgroundImage:forState:barMetrics:`
- `setBackButtonBackgroundVerticalPositionAdjustment:forBarMetrics:`

[Figure 19-10](#) shows how the navigation bar of my Alumen app used to look. When iOS 5 came along, I jazzed it up a little with some custom colors ([Figure 25-16](#)).

When you use a `UINavigationBar` implicitly as part of a `UINavigationController` interface, the controller is the navigation bar's delegate. If you were to use a `UINavigationBar` on its own, you might want to supply your own delegate. The delegate methods are:

- `navigationBar:shouldPushItem:`
- `navigationBar:didPushItem:`
- `navigationBar:shouldPopItem:`
- `navigationBar:didPopItem:`

This simple (and silly) example of a stand-alone `UINavigationBar` ([Figure 25-17](#)) implements the legendary baseball combination trio of Tinker to Evers to Chance (see the relevant Wikipedia article if you don't know about them):

```
- (void)viewDidLoad {
    [super viewDidLoad];
    UINavigationItem* ni = [[UINavigationItem alloc] initWithTitle:@"Tinker"];
    UIBarButtonItem* b = [[UIBarButtonItem alloc] initWithTitle:@"Evers"
        style:UIBarButtonItemStyleBordered
        target:self action:@selector(pushNext:)];
    ni.rightBarButtonItem = b;
    nav.items = [NSArray arrayWithObject: ni]; // nav is the UINavigationBar
}
```



Figure 25-17. A navigation bar

```
- (void) pushNext: (id) sender {
    UIBarButtonItem* oldb = sender;
    NSString* s = oldb.title;
    UINavigationItem* ni = [[UINavigationItem alloc] initWithTitle:s];
    if ([s isEqualToString:@"Evers"]) {
        UIBarButtonItem* b = [[UIBarButtonItem alloc] initWithTitle:@"Chance"
            style:UIBarButtonItemStyleBordered
            target:self action:@selector(pushNext:)];
        ni.rightBarButtonItem = b;
    }
    [nav pushNavigationItem:ni animated:YES];
}
```

UIToolbar

A UIToolbar displays a row of UIBarButtonItem objects, which are its items. The items are displayed from left to right in the order in which they appear in the items array. You can set the items with animation by calling `setItems:animated:`. You can use the system bar button items `UIBarButtonItemSystemItemFlexibleSpace` and `UIBarButtonItemSystemItemFixedSpace`, along with the `UIBarButtonItem width` property, to position the items within the toolbar.

See the previous section and [Chapter 19](#) for more about creation and customization of UIBarButtonItem objects. A bar button item's image, to be used with `UIBarButtonItemStylePlain` in a toolbar, must be a transparency mask; colors will be ignored — all that matters is the transparency of the various parts of the image. The color for the image will be supplied by default, or you can customize it with the bar button item's `tintColor`. For `UIBarButtonItemStyleBordered`, on the other hand, the image color does matter, and it is the background of the button that will be colored by the `tintColor`.

For examples of tinted bar button items in a toolbar, see [Figure 19-5](#); the toolbar has a `UIBarStyleBlack` style, its height is taller than normal to accommodate larger bar button items, and it is populated with three tinted bar button items — a `UIBarButtonItemSystemItemCancel` bar button item (the tint color tints the button background) and two `UIBarButtonItemStylePlain` bar button items with transparency mask images (the tint color tints the images).

A toolbar can be styled using its `barStyle`, `translucent`, and `tintColor` properties. Possible `barStyle` values are `UIBarStyleDefault` and `UIBarStyleBlack`; setting a `tintColor` overrides the `barStyle`. In iOS 5 you can go further, providing a background image with `setBackgroundImage:forToolbarPosition:barMetrics:`; the toolbar positions are:

- `UIToolbarPositionAny`

- `UIToolbarPositionBottom`
- `UIToolbarPositionTop`

UITabBar

A UITabBar displays UITabBarItems (its `items`), each consisting of an image and a name, and maintains a current selection among those items (its `selectedItem`, which is a UITabBarItem, not an index number). To hear about a change of selection, implement `tabBar:didSelectItem:` in the delegate (UITabBarDelegate). To change the items in an animated fashion, call `setItems:animated:`.

In iOS 5, the look of a tab bar can be customized. You can set its `tintColor` and `backgroundImage`. The `tintColor` is used to color a tab bar item's image when it is not selected (even if you also set the `backgroundImage`). The tint color used to color a tab bar item when it is selected is the tab bar's `selectedImageTintColor`. You can also set the image drawn behind the selected tab bar item to indicate that it's selected, the `selectionIndicatorImage`.

A UITabBarItem is created with one of these two methods:

- `initWithTabBarSystemItem:tag:`
- `initWithTitle:image:tag:`

UITabBarItem is a subclass of UIBarButtonItem, so in addition to its `title` and `image` it inherits the ability to adjust the image position with `imageInsets`, plus the `enabled` and `tag` properties.

A tab bar item's `image` must be a transparency mask; its colors are ignored — only the transparency matters, with tinting applied to the nontransparent areas of the image (the tab bar's `tintColor` and `selectedImageTintColor`). In iOS 5, however, you can instead call `setFinishedSelectedImage:withFinishedUnselectedImage:` to supply normal images to be shown when the tab bar item is selected and unselected respectively.

You can also customize the look of a tab bar item's title in iOS 5. Call `setTitleTextAttributes:forState:` to apply a text attributes dictionary; and you can adjust the title's position with the `titlePositionAdjustment` property.

The user can be permitted to alter the contents of the tab bar, setting its tab bar items from among a larger repertory of tab bar items. To summon the interface that lets the user do this, call `beginCustomizingItems:`, passing an array of UITabBarItems that may or may not appear in the tab bar. (To prevent the user from removing an item from the tab bar, include it in the tab bar's `items` and *don't* include it in the argument passed to `beginCustomizingItems:`.) A presented view with a Done button appears, behind the tab bar but in front of everything else, displaying the customizable items. The user can then drag an item into the tab bar, replacing an item that's already there. To hear about the customizing view appearing and disappearing, implement delegate methods:

- `tabBar:willBeginCustomizingItems:`
- `tabBar:didBeginCustomizingItems:`
- `tabBar:willEndCustomizingItems:changed:`
- `tabBar:didEndCustomizingItems:changed:`

A UITabBar on its own (outside a UITabBarController) does not provide any automatic access to the user customization interface; it's up to you. In this (silly) example, we populate a UITabBar with four system tab bar items and a More item; we also populate an instance variable array with those same four system tab bar items, plus four more. When the user taps the More item, we show the user customization interface with all eight tab bar items:

```

- (void)viewDidLoad {
    [super viewDidLoad];
    NSMutableArray* arr = [NSMutableArray array];
    for (int ix = 1; ix < 8; ix++) {
        UITabBarItem* tbi =
            [[UITabBarItem alloc] initWithTitleTabBarSystemItem:ix tag:ix];
        [arr addObject: tbi];
    }
    self.items = arr; // copy policy
    [arr removeAllObjects];
    [arr addObjectsFromArray: [self.items subarrayWithRange:NSMakeRange(0,4)]];
    UITabBarItem* tbi = [[UITabBarItem alloc] initWithTitleTabBarSystemItem:0 tag:0];
    [arr addObject: tbi]; // More button
    tb.items = arr; // tb is the UITabBar
}

- (void)tabBar:(UITabBar *)tabBar didSelectItem:(UITabBarItem *)item {
    NSLog(@"did select item with tag %i", item.tag);
    if (item.tag == 0) {
        // More button
        tabBar.selectedItem = nil;
        [tabBar beginCustomizingItems:self.items];
    }
}

```

When used in conjunction with a UITabBarController, the customization interface is provided automatically, in an elaborate way. If there are a lot of items, a More item is automatically present, and can be used to access the remaining items in a table view. Here, the user can select any of the excess items, navigating to the corresponding view. Or, the user can switch to the customization interface by tapping the Edit button. (See the iPhone Music app for a familiar example.) [Figure 25-18](#) shows how a More list looks by default:

The way this works is that the automatically provided More item corresponds to a UINavigationController with a root view controller (UIViewController) whose `view` is a UITableView. Thus, a navigation interface containing this UITableView appears as a tab view when the user taps the More button. When the user selects an item in the

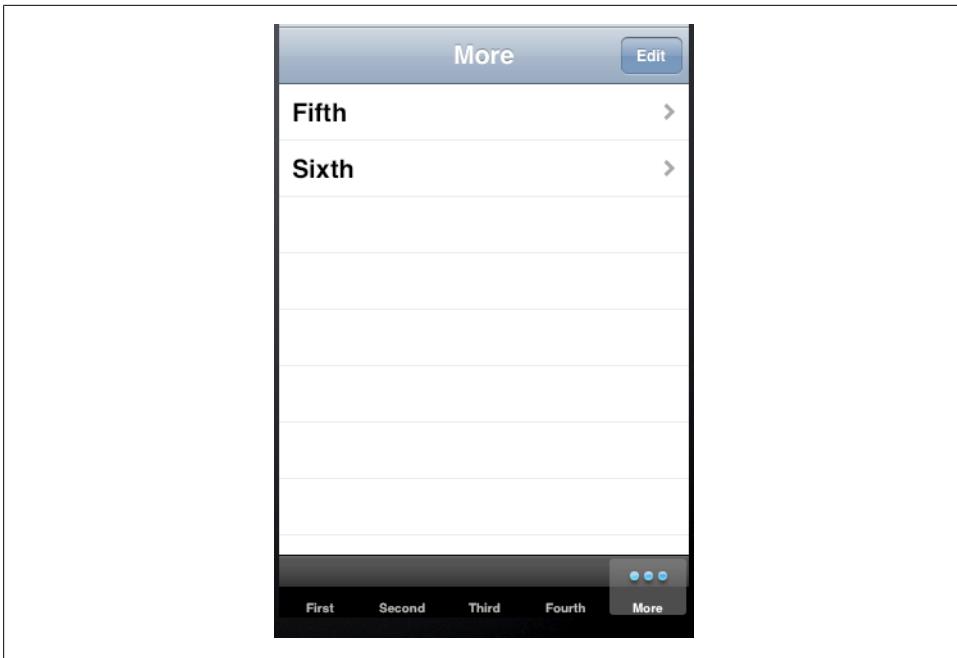


Figure 25-18. Automatically generated More list

table, the corresponding UIViewController is pushed onto the UINavigationController's stack.

You can access this UINavigationController: it is the UITabBarController's `moreNavigationController`. Through it, you can access the root view controller: it is the first item in the UINavigationController's `viewControllers` array. And through that, you can access the table view: it is the root view controller's `view`. This means you can customize what appears when the user taps the More button! For example, let's make the navigation bar black, and let's remove the word More from its title:

```
UINavigationController* more = self.tabBarController.moreNavigationController;
UIViewController* list = [more.viewControllers objectAtIndex:0];
list.title = @"";
UIBarButtonItem* b = [[UIBarButtonItem alloc] init];
b.title = @"Back";
list.navigationItem.backBarButtonItem = b; // so user can navigate back
more.navigationBar.barStyle = UIBarStyleBlack;
```

We can go even further by supplementing the table view's data source with a data source of our own, thus proceeding to customize the table itself. This is tricky because we have no internal access to the actual data source, and we mustn't accidentally disable it from populating the table. Still, it can be done. I'll start by replacing the table view's data source with an instance of my own `MyDataSource`, storing a reference to the original data source object in an instance variable of `MyDataSource`:

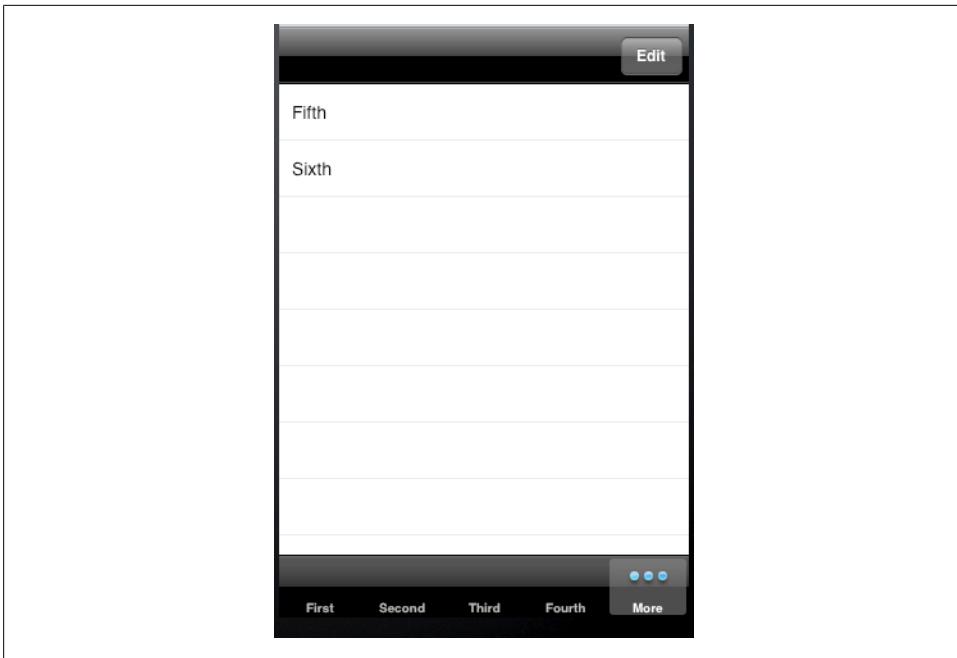


Figure 25-19. Customized More list

```
UITableView* tv = (UITableView*)list.view;
MyDataSource* mds = [[MyDataSource alloc] init];
self.myDataSource = mds; // retain policy
self.myDataSource.originalDataSource = tv.dataSource;
tv.dataSource = self.myDataSource;
```

Next, I'll use Objective-C's automatic message forwarding mechanism (see the *Objective-C Runtime Programming Guide*) so that MyDataSource acts as a front end for `originalDataSource`. MyDataSource will magically appear to respond to any message that `originalDataSource` responds to, and any message that arrives that MyDataSource can't handle will be magically forwarded to `originalDataSource`. This way, the insertion of the MyDataSource instance as data source doesn't break whatever the original data source does:

```
- (id)forwardingTargetForSelector:(SEL)aSelector {
    if ([self.originalDataSource respondsToSelector: aSelector])
        return self.originalDataSource;
    return [super forwardingTargetForSelector:aSelector];
}
```

Finally, we'll implement the two Big Questions required by the `UITableViewDataSource` protocol, to quiet the compiler. In both cases, we first pass the message along to `originalDataSource` (somewhat analogous to calling `super`); then we add our own customizations as desired. Here, I'll remove each cell's disclosure indicator and change its text font. The outcome is shown in [Figure 25-19](#):

```

- (NSInteger)tableView:(UITableView *)tv numberOfRowsInSection:(NSInteger)section {
    // this is just to quiet the compiler
    return [self.originalDataSource tableView:tv numberOfRowsInSection:section];
}

- (UITableViewCell *)tableView:(UITableView *)tv
    cellForRowAtIndexPath:(NSIndexPath *)indexPath {
    UITableViewCell *cell =
        [self.originalDataSource tableView:tv cellForRowAtIndexPath:indexPath];
    cell.accessoryType = UITableViewCellAccessoryNone;
    cell.textLabel.font = [UIFont systemFontOfSize:14];
    return cell;
}

```

Apearance Proxy

In iOS 5, instead of sending messages that customize the look of an interface object to the object itself, you can send them instead to an *appearance proxy* for that object's class. The appearance proxy then passes that same message along to the actual instances of that class.

Thus, for example, instead of sending `setTitleTextAttributes:forState:` to a particular UIBarButtonItem, you could instead send it to a UIBarButtonItem appearance proxy. Any actual UIBarButtonItems from then on would have the text attributes you specified.

This architecture has two chief uses:

- It simplifies the task of giving your app a consistent overall appearance. Suppose you want all UIBarButtonItems to have a certain title font. Instead of having to remember to send `setTitleTextAttributes:forState:` to every UIBarButtonItem your app ever instantiates, you send it once to the appearance proxy and it is sent to those UIBarButtonItems for you.
- It provides access to interface objects that might otherwise be difficult to refer to. For example, you don't get direct access to a search bar's external Cancel button, but it is a UIBarButtonItem and you can customize it through the UIBarButtonItem appearance proxy.

There are two class methods for obtaining an appearance proxy:

`appearance`

Returns a general appearance proxy for that class.

`appearanceWhenContainedIn:`

The argument is a nil-terminated list (not an array) of classes, arranged in order of containment from inner to out. The method you send to the appearance proxy returned from this call will be passed on only to instances of the target class that are actually contained in the way you describe. The notion of what "contained"

means is deliberately left vague; basically, it works the way you intuitively expect it to work.

When configuring appearance proxy objects, *specificity trumps generality*. Thus, you could call `appearance` to say what should happen for *most* instances of some class, and call `appearanceWhenContainedIn:` to say what should happen *instead* for *certain* instances of that class. Similarly, longer `appearanceWhenContainedIn:` chains are more specific than shorter ones.

For example, here's some code from my Latin flashcard app (`myGolden` and `myPaler` are methods defined by a category on `UIColor`):

```
[[UIBarButtonItem appearance]
    setTintColor: [UIColor myGolden]];
[[UIBarButtonItem appearanceWhenContainedIn:
    [UIToolbar class], nil]
    setTintColor: [UIColor myPaler]];
[[UIBarButtonItem appearanceWhenContainedIn:
    [UIToolbar class], [DrillViewController class], nil]
    setTintColor: [UIColor myGolden]];
```

That means:

1. In general, bar button items should be tinted golden.
2. But bar button items in a toolbar are an exception: they should be paler.
3. But bar button items in a toolbar in `DrillViewController`'s view are an exception to *that*: they should be golden.

(If you're looking at this book's figures in color, you can see this difference made manifest in [Figure 19-3](#) and [Figure 19-5](#).)

Sometimes, in order to express sufficient specificity, I find myself defining subclasses for no other purpose than to refer to them when obtaining an appearance proxy. For example, here's some more code from my Latin flashcard app:

```
[[UINavigationBar appearance] setBackgroundImage:marble2
    forBarMetrics:UIBarMetricsDefault];
// counteract the above for the black navigation bar
[[BlackNavigationBar appearance] setBackgroundImage:nil
    forBarMetrics:UIBarMetricsDefault];
```

In that code, `BlackNavigationBar` is a `UINavigationBar` subclass that does nothing whatever. Its sole purpose is to tag one navigation bar in my interface so that I can refer to it in that code! Thus, I'm able to say, in effect, "All navigation bars in this app should have `marble2` as their background image, unless they are instances of `BlackNavigationBar`."

The ultimate in specificity is, of course, to customize the look of an instance directly. Thus, for example, if you set one particular `UIBarButtonItem`'s `tintColor` property, then `setTintColor:` sent to a `UIBarButtonItem` appearance proxy will have no effect on that particular bar button item.

You'll want to know which messages can be sent to the appearance proxies for which classes. The best way to find out is to look in the header for that class (or a superclass); any appropriate property or method will be tagged `UI_APPEARANCE_SELECTOR`. For example, here's how the `tintColor` property is declared in `UIBarButtonItem.h`:

```
@property(nonatomic,retain) UIColor *tintColor  
    OSX_AVAILABLE_STARTING(__MAC_NA,__IPHONE_5_0)  
    UI_APPEARANCE_SELECTOR;
```

You may also be able to deduce this information from the classification of properties and methods in the documentation, but I find the header to be far more reliable and explicit.

The appearance proxy is an `id`. Therefore, it can be sent *any* message for which a method signature can be found; if you send it a message that isn't tagged `UI_APPEARANCE_SELECTOR` for the class that the proxy represents, the compiler can't stop you, but you'll crash at runtime when the message is actually sent. An `id` has no properties; that's why we must call `setTintColor:` in order to set the `UIBarButtonItem`'s `tintColor` property.

I've been deliberately vague as to *when* instances are affected by the appearance proxy. The rule seems to be that sending a message to an appearance proxy will affect *future* instances of that class, but it won't spontaneously alter the look of an object that is *currently visible* in the interface. The look of an existing object will be changed, however, if that object is refreshed somehow.

For example, in a tab bar interface where the tab bar is visible, this code has no immediate effect:

```
[[UITabBarItem appearance] setTitleTextAttributes:  
    [NSDictionary dictionaryWithObjectsAndKeys:  
        [UIFont fontWithName:@"Baskerville" size:18],  
        UITextAttributeFont,  
        nil]  
    forState:UIControlStateNormal];
```

But if a presented view controller's view then covers the whole interface, and is later dismissed, the tab bar items will have adopted their new text attributes. Or, if this code is in a child view controller of the `UITabBarController`, changing the value of `self.tabBarItem.title` will cause the text attributes for this tab bar item to be refreshed.

In general, though, such trickery is probably not appropriate. You'll usually configure your appearance proxies very early in the lifetime of the app, and never again. The app delegate's `application:didFinishLaunchingWithOptions:`, before the app's window has been displayed, is the most obvious and common location.

Modal Dialogs

A modal dialog demands attention; while it is present, the user can do nothing other than work within it or dismiss it. You might need to put up a simple modal dialog in order to give the user some information or to ask the user how to proceed. Two UIView subclasses construct and present rudimentary modal dialogs:

UIAlertView

A UIAlertView pops up unexpectedly with an elaborate animation and may be thought of as an attention-getting interruption. An alert is displayed in the center of the screen; it contains a title, a message, and an indefinite number of additional buttons, one of which may be the cancel button, meaning that it does nothing but dismiss the alert. Often there is *only* a cancel button, the primary purpose of the alert being to show the user the message (“You won the game”); the additional buttons may be used to give the user a choice of how to proceed (“You won the game; would you like to play another?” “Yes,” “No,” “Replay”).

Starting in iOS 5, an alert view may contain a text field, a password field, or both.

UIActionSheet

A UIActionSheet may be considered the iOS equivalent of a Mac OS X menu. An action sheet is displayed arising from the interface: on the iPhone, it slides up from the bottom of the screen; on the iPad, it is typically shown in a popover. It consists of an indefinite number of buttons (there can be a title, optionally, but there usually isn’t); one may be the cancel button, which appears last (though on the iPad, for a popover, this may not be needed), and one may be a “destructive” button, which appears first in red, emphasizing the severity of that option.

Where a UIAlertView is an interruption, a UIActionSheet is a logical branching of what the user is already doing: it typically divides a single piece of interface into multiple possible courses of action. For example, in Mobile Safari a single “More” button summons an action sheet that lets the user add the current page as a bookmark, add it to the home screen, mail a link to it, or print it (or cancel and do nothing).



Figure 26-1. An alert view (UIAlertView)



One occasionally sees a misuse of the built-in dialogs to include additional interface. For example, a UIActionSheet is a UIView, so in theory you can add a subview to it. I cannot recommend such behavior; it clearly isn't intended, and there's no need for it. If what you need isn't what a built-in dialog normally does, don't use a built-in dialog.

Alert View

The basic method for constructing an alert view (UIAlertView) is `initWithTitle:message:delegate:cancelButtonTitle:otherButtonTitles:`. The method for making a constructed alert view appear onscreen is `show`. The alert is automatically dismissed as soon as the user taps any button. Here's an example ([Figure 26-1](#)):

```
UIAlertView* alert = [[UIAlertView alloc] initWithTitle:@"Not So Fast!"
                                              message:@"Do you really want to do this tremendously destructive thing?"
                                             delegate:self cancelButtonTitle:@"Yes" otherButtonTitles:@"No", @"Maybe", nil];
[alert show];
```

The `otherButtonTitles` parameter is of indefinite length, so it must be either `nil` or a `nil`-terminated list (not an array) of strings. The cancel button needn't involve canceling anything; it is drawn darker than the other buttons and comes last in a column of buttons, as you can see from [Figure 26-1](#), but if there were three `otherButtonTitles` and a `nil` `cancelButtonTitle`, the alert dialog would look exactly the same.

Starting in iOS 5, before calling `show` you can modify the alert view's `alertViewStyle` to add text fields. Your choices are:

- `UIAlertViewStyleDefault`, the default (no text fields)
- `UIAlertViewStylePlainTextInput`, one normal text field
- `UIAlertViewStyleSecureTextInput`, one `secureTextEntry` text field
- `UIAlertViewStyleLoginAndPasswordInput`, one normal text field and one `secureTextEntry` text field

If an alert view is to contain a text field, it probably should have at most one or two buttons, with short titles such as “OK” and “Cancel”. Otherwise, there might not be room on the screen for the alert view and the keyboard. You can retrieve the text fields with `textFieldAtIndex:`; possible parameters are `0` and `1` (where `1` is the password field when the style is `UIAlertViewStyleLoginAndPasswordInput`). You can treat the text fields as you would any text field; for example, you can set the text field’s delegate, arrange to receive control events from the text field, determine the text field’s keyboard type, and so on:

```
UIAlertView* alert = [[UIAlertView alloc] initWithTitle:@"Enter a number:"
                                              message:nil delegate:self cancelButtonTitle:@"Cancel"
                                              otherButtonTitles:@"OK", nil];
alert.alertViewStyle = UIAlertViewStylePlainTextInput;
UITextField* tf = [alert textFieldAtIndex:0];
tf.keyboardType = UIKeyboardTypeNumberPad;
[tf addTarget:self action:@selector(textChanged:)
    forControlEvents:UIControlEventEditingChanged];
[alert show];
```

The alert dialog is modal, but the code that presents it is not: after the alert is shown, your code continues to run. If an alert consists of a single button (the cancel button), you might show it and forget about it, secure in the knowledge that the user must dismiss it sooner or later and that nothing can happen until then. But if you want to respond at the time the user dismisses the alert, or if there are several buttons and you want to know which one the user tapped to dismiss the alert, you’ll need to implement at least one of these delegate methods (`UIAlertViewDelegate`):

- `alertView:clickedButtonAtIndex:`
- `alertView:willDismissWithButtonIndex:`
- `alertView:didDismissWithButtonIndex:`

The cancel button index is usually `0`, with the remaining button indexes increasing in the order in which they were defined. If you’re in any doubt, or if you need the button title for any other reason, you can call `buttonTitleAtIndex:`. Properties allow you to work out the correspondence between indexes and buttons without making any assumptions:

- `cancelButtonIndex` (-1 if none)
- `firstOtherButtonIndex` (-1 if none)
- `numberOfButtons` (including the cancel button)

You can also dismiss an alert view programmatically, with `dismissWithClickedButtonIndex:animated:`. When an alert view is dismissed programmatically, the delegate method `alertView:clickedButtonAtIndex:` is *not* called, because no button was actually clicked by the user. But the button index you specify is still passed along to the two `dismiss` delegate methods. The button index you specify needn’t correspond to any

existing button; thus, you could use it as a way of telling your delegate method that your code, and not the user, dismissed the alert.

Two additional delegate methods notify you when the alert is initially shown:

- `willPresentAlertView:`
- `didPresentAlertView:`

In iOS 5, a new delegate method asks whether the first “other button” should be enabled:

- `alertViewShouldEnableFirstOtherButton:`

The delegate receives that message each time the state of things changes in the alert — in particular, when the alert appears and when the text in a text field changes. In this example, there’s a text field, my cancel button says “Cancel”, and my other button says “OK”; I enable the OK button only if there is text in the text field:

```
- (BOOL)alertViewShouldEnableFirstOtherButton:(UIAlertView *)alertView {
    UITextField* tf = [alertView textFieldAtIndex:0];
    return [tf.text length] > 0;
}
```

One further delegate method notifies you if the alert is dismissed by the system:

- `alertViewCancel:`

Before iOS 4.0, this could happen because the user quit the app with the alert showing; the system dismissed the alert, and your code had a chance to respond before actually terminating. But iOS 4.0 introduced multitasking; if the user clicks the Home button, your app is backgrounded without the system dismissing the alert, and `alertViewCancel:` may be a dead letter. It would thus be up to your code, as the app is backgrounded, whether to leave the alert there or to dismiss the alert and perhaps take some default action.

Action Sheet

The basic method for constructing an action sheet (`UIActionSheet`) is `initWithTitle:delegate:cancelButtonTitle:destructiveButtonTitle:otherButtonTitles:`. There are various methods for summoning the actual sheet, depending on what part of the interface you want the sheet to arise from. The following are appropriate on the iPhone, where the sheet typically rises from the bottom of the screen:

`showInView:`

On the iPhone, far and away the most commonly used method. You will usually specify the root view controller’s view. Don’t specify a view whose view controller is contained by a view controller that hides the bottom of the interface, such as a tab bar controller or a navigation controller with a toolbar; if you do, some of the buttons may not function (and you get a helpful warning in the console: “Present-



Figure 26-2. An action sheet on the iPhone

ing action sheet clipped by its superview”). Instead, specify the tab bar controller’s view itself, or the navigation controller’s view itself, or use one of the other methods. For example, in my Zotz app, which has a tab bar interface, the settings view controller summons an action sheet like this ([Figure 26-2](#)):

```
[sheet showInView: self.tabBarController.view];
```

`showFromTabBar:, showFromToolbar:`

On the iPhone, these cause the sheet to rise from the bottom of the screen, just like `showInView:`, because the tab bar or toolbar is at the bottom of the screen; however, they avoid the problem with `showInView:` described earlier.

On the iPad, you are more likely to use one of the following methods, which resemble the methods for presenting a popover ([Chapter 22](#)); they do in fact present the action sheet as a popover, with its arrow pointing to the specified part of the interface ([Figure 26-3](#)):

- `showFromRect:inView:animated:`
- `showFromBarButtonItem:animated:`

(On the iPhone, those methods should be avoided; they don’t cause an error, and they do work — the sheet still ends up at the bottom at the screen — but they can do messy things to the interface.)



Figure 26-3. An action sheet presented as a popover



Figure 26-4. An action sheet presented inside a popover

On the iPad, there is usually no point including a cancel button title: if the alert sheet is shown as a popover, no cancel button will appear. This is because the popover is configured to be dismissed when the user taps outside it, which is the same as canceling it.

However, it is also possible on the iPad to show an alert sheet *inside an existing popover*. In this scenario, we are already presenting the popover, and then we summon an action sheet within the popover's view. In that case, the action sheet behaves as if the popover were an iPhone: you summon it with `showInView:`, it slides up from the bottom of the popover, and the cancel button appears (Figure 26-4). Moreover, the action sheet is then modal: the user can't dismiss the popover, or do anything else, without dismissing the action sheet first.

An action sheet also has a style, its `actionSheetStyle`:

- `UIActionSheetStyleAutomatic`
- `UIActionSheetStyleDefault`
- `UIActionSheetStyleBlackTranslucent`
- `UIActionSheetStyleBlackOpaque`

These values are closely related to the possible styles (`barStyle`) of a UIToolbar ([Chapter 25](#)). However, an action sheet's style depends also on the mode of presentation; experimentation suggests, for example, that setting the `actionSheetStyle` of an action sheet that appears as a popover may make no perceptible difference, and that an action sheet that is shown from a tab bar will always be black opaque.

In other respects an action sheet is managed in a manner completely parallel to an alert view. When one of its buttons is tapped, the sheet is dismissed automatically, but you'll probably want to implement a delegate method (`UIActionSheetDelegate`) in order to learn which button it was:

- `actionSheet:clickedButtonAtIndex:`
- `actionSheet:willDismissWithButtonIndex:`
- `actionSheet:didDismissWithButtonIndex:`

If the action sheet is shown as a popover on the iPad, and if the popover is dismissed by the user tapping outside it, the button index is `-1`.

To respond appropriately to the delegate methods without making assumptions about how the indexes correspond to the buttons, you can use the `buttonTitleAtIndex:` method, and these properties:

- `cancelButtonIndex`
- `destructiveButtonIndex`
- `firstOtherButtonIndex`
- `numberOfButtons`

You can dismiss an action sheet programmatically with `dismissWithClickedButtonIndex:animated:`, in which case `actionSheet:clickedButtonAtIndex:` is not called, but the two `dismiss` delegate methods are. Two additional delegate methods notify you when the sheet is initially shown:

- `willPresentActionSheet:`
- `didPresentActionSheet:`

A further delegate method, `actionSheetCancel:`, notifies you if the sheet is dismissed by the system, though in iOS 4.0 or later this is unlikely to happen; if your app is backgrounded with an action sheet showing, it's up to you to decide how to proceed.

Here's the code that presents the action sheet shown in [Figure 26-2](#), along with the code that responds to its dismissal:

```
- (void)chooseLayout: (id) sender {
    UIActionSheet* sheet =
        [[UIActionSheet alloc] initWithTitle:@"Choose New Layout" delegate:self
            cancelButtonTitle:(NSString *)@"Cancel" destructiveButtonTitle:nil
            otherButtonTitles:@[@"3 by 3", @"4 by 3", @"4 by 4", @"5 by 4", @"5 by 5",
            nil];
    [sheet showInView: self.tabBarController.view];
}

- (void)actionSheet:(UIActionSheet *)as clickedButtonAtIndex:(NSInteger)ix {
    if (ix == as.cancelButtonIndex)
        return;
    NSString* s = [as buttonTitleAtIndex:ix];
    // ...
}
```

On the iPad, if an action sheet is shown as a popover from a bar button item in a toolbar, the toolbar becomes a passthrough view for the popover. This behavior is troublesome, for the same reasons I gave in [Chapter 22](#): the user can now tap another bar button item without causing the action sheet's popover to be dismissed (and possibly even summoning another popover—perhaps even another instance of the same action sheet—simultaneously). Preventing this sort of conflict is entirely up to your code. You can't solve the problem by adjusting the popover controller's `passthroughViews`, because you can't get access to the popover controller! This situation seems like a massive bug. The best solution seems to be to implement the delegate methods to toggle user interaction in the toolbar:

```
- (IBAction)doButton:(id)sender {
    UIActionSheet* act = [[UIActionSheet alloc]
        initWithTitle:nil delegate:self cancelButtonTitle:nil
        destructiveButtonTitle:nil otherButtonTitles:@[@"Hey", @"Ho", nil];
    [act showFromBarButtonItem:sender animated:YES];
}

- (void)didPresentActionSheet:(UIActionSheet *)actionSheet {
    [self.toolbar setUserInteractionEnabled:NO];
}

- (void)actionSheet:(UIActionSheet *)actionSheet
    didDismissWithButtonIndex:(NSInteger)buttonIndex {
    [self.toolbar setUserInteractionEnabled:YES];
}
```

Dialog Alternatives

Alert views and actions sheets are limited, inflexible, and inappropriate to any but the simplest cases. In more complex situations, it really isn't that much work to implement an alternative.

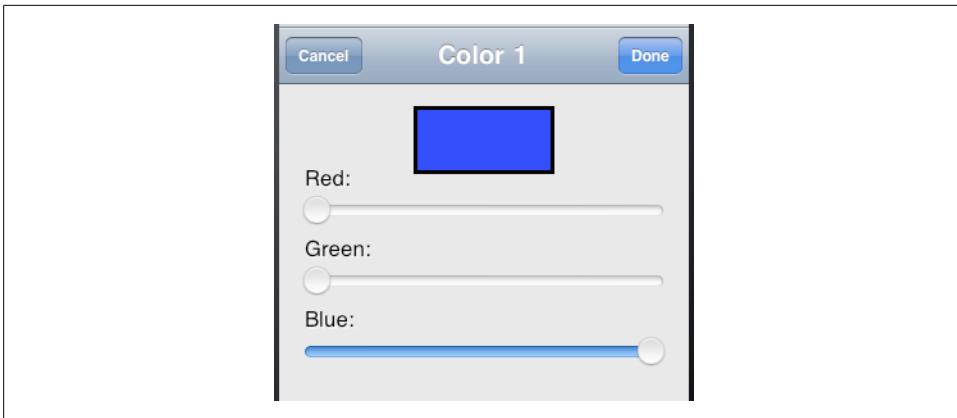


Figure 26-5. A presented view functioning as a modal dialog

On the iPhone, the main alternative is to navigate to a new screenful of interface. This might be by way of a navigation interface, or using a presented view. For example, in the Zotz app, in the Settings view, when the user taps a color, I summon a presented view, using a `UIViewController` subclass of my own, `ColorPickerController` (Figure 26-5).

On Mac OS X, the color picker in Figure 26-5 might be presented as a secondary window acting as a dialog. On the small iPhone screen, where there are no secondary windows, the presented view *is* the equivalent of a dialog. Indeed, one might argue that the action sheet shown in Figure 26-2 is not a very appropriate use of an action sheet, that it's too intrusive and has too many buttons. It might have been better if I'd designed my own presented view; I probably picked an action sheet because it required just a few lines of code — basically, I was being lazy.

On the iPad, a popover is virtually a secondary window, and can be truly modal. An action sheet is usually presented as a popover, but it's limited, and you don't get access to the popover controller; in many cases, you'll probably be better off designing your own view to be presented in a popover. The popovers in Figure 22-1, for example, are effectively modal dialogs. A popover can internally present a secondary presented view or even an action sheet, as we've already seen. Also on the iPad, a presented view can use the `UIModalPresentationFormSheet` presentation style, which is effectively a dialog window smaller than the screen.

Local Notifications

A *local notification* is an alert to the user that can appear even if your app is not running. In one of its manifestations, it appears as a dialog on top of whatever the user is doing at that moment, which is why it is treated in this chapter. (If a local notification from some other app were to appear as a dialog while *your* app is frontmost, *your* app would

become inactive; see [Chapter 11](#) and the `applicationWillResignActive:` app delegate message.) Local notifications were introduced in iOS 4.



This use of the term *notification* has nothing to do with `NSNotification` ([Chapter 11](#)). The ambiguity is unfortunate.

Your app does not present a local notification alert: indeed, your app *can't* present a local notification alert, because if your app's local notification alert appears, your app *ex hypothesi* isn't frontmost. Rather, your app hands a local notification to the system along with instructions about when the local notification should *fire*. When the specified time arrives, if your app isn't frontmost, the system presents the notification on your behalf.

In iOS 5, the user has several choices as to how a notification from your app should be presented — probably because users complained that the old way was too intrusive. These choices appear in the Settings app, under Notifications. iOS 5 introduces the Notification Center, which appears when the user swipes downward from the very top of the screen, and the user can enable or disable your app's notifications appearing there. A notification from your app can also appear as an alert (the old style from iOS 4), as a temporary banner at the top of the screen, or not at all. The user can also prohibit your app's alerts from appearing in the lock screen. It is thus perfectly possible for the user to suppress your app's alerts altogether!

No matter the interface whereby a notification presents itself, it generally provides some way for the user to summon your app in response. If the user does so, your app will be brought to the front, launching it if it isn't already suspended in the background. Your app may need to be concerned with detecting that it has been brought to the front under these special circumstances, and I'll talk in a moment about how it can do that.

To create a local notification, you configure a `UILocalNotification` object and hand it to the system with `UIApplication's scheduleLocalNotification:`. The `UILocalNotification` object has properties describing how the dialog should look and behave and when you want it to appear:

`alertBody`

The message displayed in the notification.

`alertAction`

This matters only if your notification is displayed as an alert; in that case, this is the text of the action button. If you don't set `alertAction`, the text of the action button will be "Launch."

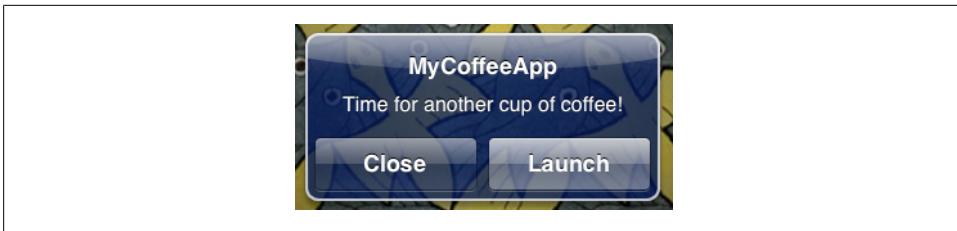


Figure 26-6. An alert posted by the system when a local notification fires



According to the documentation, you should be able to set `hasAction` to NO to suppress the action button altogether, but in my testing, doing so had no effect.

`soundName`

The name of a sound file at the top level of your app bundle, to be played when the alert appears. This should be an uncompressed sound (AIFF or WAV). Alternatively, you can specify the default sound, `UILocalNotificationDefaultSoundName`. If you don't set this property, there won't be a sound. Regardless of the value you supply here, the user can prevent your app's notifications from emitting a sound.

`userInfo`

An optional NSDictionary whose contents are up to you. As you would expect, this is so that your app can identify the local notification object when it fires.

`fireDate, timeZone`

When you want the local notification to fire. The `fireDate` is an NSDate (see [Chapter 10](#) and [Chapter 25](#) for examples of date manipulation). If you don't include a `timeZone`, the date is measured against universal time; if you *do* include a `timeZone`, the date is measured against the user's local time zone, even if that time zone changes (because the user travels, for instance).

`repeatInterval, repeatCalendar`

If set, the local notification will recur.

As I've already mentioned, you hand a configured local notification to the system with UIApplication's `scheduleLocalNotification:`. Additional UIApplication methods let you manipulate the list of local notifications you've already scheduled. You can cancel one or all scheduled local notifications (`cancelLocalNotification:`, `cancelAllLocalNotifications:`); starting in iOS 4.2, you can also manipulate the list directly by setting UIApplication's `scheduledLocalNotifications`, an NSArray property (previously, this property was read-only).

[Figure 26-6](#) shows an alert generated by the firing of a local notification. Here's a simple example of creating and scheduling the local notification that resulted in that alert:

```
UILocalNotification* ln = [[UILocalNotification alloc] init];
ln.alertBody = @"Time for another cup of coffee!";
ln.fireDate = [NSDate dateWithTimeIntervalSinceNow:15];
ln.soundName = UILocalNotificationDefaultSoundName;
[[UIApplication sharedApplication] scheduleLocalNotification:ln];
```

Now let's talk about what happens when one of your scheduled local notifications fires. There are three possibilities, depending on the state of your app at that moment:

Your app is suspended in the background

If the user summons your app from a notification, your app is brought to the front; your app delegate will then receive `application:didReceiveLocalNotification:`, where the second parameter is the `UILocalNotification`, and your application's `applicationState` will be `UIApplicationStateInactive`.

Your app is frontmost

The user won't be informed by the system that the notification has fired (though the notification may be stored in the Notification Center). Your app delegate will receive `application:didReceiveLocalNotification:`, where the second parameter is the `UILocalNotification`, and your application's `applicationState` will be `UIApplicationStateActive`. The idea is that if your app wants to let the user know that something special is happening, that's your app's business and it can do it in its own way.

Your app isn't running

If the user summons your app from a notification, your app is launched; your app delegate will then receive, *not* `application:didReceiveLocalNotification:`, but rather `application:didFinishLaunchingWithOptions:` with an `NSDictionary` parameter that includes the `UIApplicationLaunchOptionsLocalNotificationKey`, whose value is the `UILocalNotification`.

Thus, you should implement `application:didReceiveLocalNotification:` to check the `UIApplication`'s `applicationState`, and you should implement `application:didFinishLaunchingWithOptions:` to check its second parameter to see whether we are launching in response to a local notification. In this way, you will be able to distinguish the three different possibilities, and you can respond appropriately.

In the first and third cases (your app is suspended in the background, or your app isn't running), you may want to show the user some interface appropriate to the local notification's situation. For example, you might want to push a particular view controller onto your navigation interface or present a particular view controller's view. However, when your app is launched from scratch, the first thing the user sees is its launch image ([Chapter 9](#)); and when your app is activated from a suspended state, the first thing the user sees is a screenshot image of your app, taken by the system when your app was suspended. Thus, there may be a mismatch between these images and the interface that you're about to show the user in this special situation. The user will thus see an odd flash as the image is removed to reveal your app's actual interface. To prevent this flash, you can include in the original `UILocalNotification` an `alertLaunchImage` to be presen-

ted instead of these images. The idea is that this `alertLaunchImage` should be a better match for the interface the user will actually see.

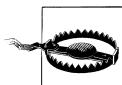


There is actually a fourth possibility for when a local notification fires. Under some special circumstances (addressed, for example, in [Chapter 27](#) and [Chapter 35](#)), your app might be running, *not* suspended, in the background. In this case, the situation is similar to what happens when your app *is* suspended: the user may be notified, and can summon your app to the front. Your running-in-the-background app can even schedule a notification to fire immediately with the convenience method `presentLocalNotificationNow:`.

Some Frameworks

In addition to the basic UIKit and Foundation frameworks, which supply the fundamental interface and utility classes for all apps, Cocoa supplies numerous optional frameworks that you can use if your app has special needs. This part of the book introduces some of these frameworks and their related topics. At the same time, it necessarily exercises some restraint. To explore *all* of the additional iOS frameworks in *full* depth would more than double the size of this book! So this part of the book fully explains the basics, but then stops and leaves you to go further on your own if you need to; it teaches you what you need to know to get started, and it trains you to understand and explore these and related frameworks independently if your app requires a further level of depth and detail.

- [Chapter 27](#) introduces the various iOS means for playing sound files, including audio sessions and playing sounds in the background.
- [Chapter 28](#) describes some basic ways of playing video (movies), along with an introduction to the powerful new AV Foundation framework.
- [Chapter 29](#) is about how an app can access the user's music library.
- [Chapter 30](#) is about how an app can access the user's photo library, along with the ability to take photos and capture movies.
- [Chapter 31](#) discusses how an app can access the user's address book.
- [Chapter 32](#) talks about how an app can access the user's calendar data.
- [Chapter 33](#) describes how an app can allow the user to compose and send email and SMS messages.
- [Chapter 34](#) explains how an app can display a Google map, along with custom annotations and overlays.
- [Chapter 35](#) is about how an app can learn where the device is located, how it is moving, and how it is oriented.



This chapter not yet revised for iOS 5.

iOS provides various means and technologies for allowing your app to produce sound (and even to input it). The topic is a large one, so this chapter can only introduce it. You'll want to read Apple's *Multimedia Programming Guide* and *Core Audio Overview*.



None of the classes discussed in this chapter provide any user interface within your application for allowing the user to stop and start playback of sound (though you can create your own interface and use it with them). However, a web view ([Chapter 24](#)) supports the HTML 5 `<audio>` tag; this can be a simple, lightweight way to play audio and to allow the user to control playback. Or treat the sound as a movie and use the `MPMoviePlayerController` class discussed in [Chapter 28](#); this can also be a good way to play a sound file located remotely over the Internet.

System Sounds

The simplest form of sound is *system sound*, which is the iOS equivalent of the basic computer “beep.” This is implemented through System Sound Services; you’ll need to import `<AudioToolbox/AudioToolbox.h>` and link to `AudioToolbox.framework`. You’ll be calling one of two C functions, which behave very similarly to one another:

`AudioServicesPlayAlertSound`

Plays a sound and, on an iPhone, may also vibrate the device, depending on the user’s settings. On the original iPod touch, plays only a built-in alert sound.

`AudioServicesPlaySystemSound`

Plays a short sound of your choice. On an iPhone, there won't be an accompanying vibration, but you can specifically elect to have this "sound" *be* a device vibration.

The sound needs to be an uncompressed AIFF or WAV file (or an Apple CAF file wrapping one of these). To hand the sound to these functions, you'll need a SystemSoundID, which you obtain by calling `AudioServicesCreateSystemSoundID` with a CFURLRef (or NSURL) that points to a sound file. In this example, the sound file is in our app bundle:

```
NSURL* sndurl = [[NSBundle mainBundle] URLForResource:@"test" withExtension:@"aif"];
SystemSoundID snd;
AudioServicesCreateSystemSoundID ((CFURLRef)sndurl, &snd);
AudioServicesPlaySystemSound(snd);
```

However, there's a problem with that code: we have failed to exercise proper memory management. We need to call `AudioServicesDisposeSystemSoundID` to release our SystemSoundID. But when shall we do this? `AudioServicesPlaySystemSound` executes asynchronously. So the solution can't be to call `AudioServicesDisposeSystemSoundID` in the next line of the same snippet, because this would release our sound just as it is about to start playing, resulting in silence. A solution that works is to implement a sound completion handler, a function that is called when the sound has finished playing. So, our sound-playing snippet now looks like this:

```
NSURL* sndurl = [[NSBundle mainBundle] URLForResource:@"test" withExtension:@"aif"];
SystemSoundID snd;
AudioServicesCreateSystemSoundID((CFURLRef)sndurl, &snd);
AudioServicesAddSystemSoundCompletion(snd, NULL, NULL, &SoundFinished, NULL);
AudioServicesPlaySystemSound(snd);
```

And here is our sound completion handler, the `SoundFinished` function referred to in the previous snippet:

```
void SoundFinished (SystemSoundID snd, void* context) {
    AudioServicesRemoveSystemSoundCompletion(snd);
    AudioServicesDisposeSystemSoundID(snd);
}
```

Note that because we are about to release the sound, we first release the sound completion handler information applied to it. The last argument passed to `AudioServicesAddSystemSoundCompletion` is a pointer-to-void that comes back as the second parameter of our sound completion handler function; you can use this parameter in any way you like, such as to help identify the sound.

Audio Session

If your app is going to use a more sophisticated way of producing sound, such as an audio player (discussed in the next section), it must specify a *policy* regarding that sound. This policy will answer such questions as: Should sound stop when the screen

is locked? Should sound interrupt existing sound (being played, for example, by the Music app) or should it be layered on top of it?

Your policy is declared in an *audio session*, which is a singleton AVAudioSession instance created automatically as your app launches. You can configure this AVAudioSession instance once at launch time (or, at any rate, before producing any sound), or you can change its configuration dynamically while your app runs. You can talk to the AVAudioSession instance in Objective-C (see the AVAudioSession class reference) or in C (see the Audio Session Services reference), or both.

To use the Objective-C API, you'll need to link to *AVFoundation.framework* and import `<AVFoundation/AVFoundation.h>`. You'll refer to your app's AVAudioSession by way of the class method `sharedInstance`.

To use the C API, you'll need to link to *AudioToolbox.framework* and import `<AudioToolbox/AudioToolbox.h>`. The `AudioSession...` functions don't require a reference to an audio session. You must explicitly initialize your audio session with `AudioSessionInitialize` before talking to it with the C API, unless you have already talked to it with the Objective-C API.

The basic policies for audio playback are:

Ambient

Your app's audio plays even while iPod music or other background audio is playing, and is silenced by the Silent switch and screen locking.

Solo Ambient

Your app stops iPod music or other background audio from playing, and is silenced by the Silent switch and screen locking (the default).

Playback

Your app stops iPod music or other background audio from playing, and is not silenced by the Silent switch or screen locking.

In addition, you can modify these policies, using the `AudioSessionSetProperty` function of the C API. For example:

- You can override the Playback policy so as to allow iPod music or other background audio to play (`kAudioSessionProperty_OverrideCategoryMixWithOthers`).
- You can override a policy that allows iPod music or other background audio to play, so as to *duck* (diminish the volume of) that background audio (`kAudioSessionProperty_OtherMIXableAudioShouldDuck`). Ducking does *not* depend automatically on whether your app is actively producing any sound; rather, it starts as soon as you turn this override on and remains in place until your audio session is deactivated.

Your audio session policy is not in effect unless your audio session is also *active*. By default, it isn't. Thus, asserting your audio session policy is done by a combination of configuring the audio session and activating the audio session. This is a little tricky

because of multitasking on iOS 4. Your audio session can be deactivated automatically if your app is no longer active. So if you want your policy to be obeyed under all circumstances, you must explicitly activate your audio session each time your app becomes active. (See [Chapter 11](#) for how an app resigns and resumes active status.)

For example, an audio session configured with a Playback policy that is activated as the app launches, in the app delegate's `application:didFinishLaunchingWithOptions:`, will silence background audio at launch time, but if the user then switches away to the Music app and starts playing some music, and then switches back to your app, the iPod music won't be silenced, even though it is now in the background, because your audio session has been deactivated. You might try to work around this by activating your audio session in the app delegate's `applicationWillEnterForeground:` as well. But then the user might double-click the Home button while your app is frontmost, and use the app switcher playback controls ([Figure 27-1](#)) to start playing iPod music; when the user leaves the app switcher interface, that iPod music continues to play, even though it is now in the background. To prevent *that*, you must activate your audio session in the app delegate's `applicationDidBecomeActive:`. It turns out that *only* `applicationDidBecomeActive:` is called in all three situations, so that is the place to activate your audio session so as to keep it active when running under iOS 4.

Here are the main methods you'll need:

setCategory:error:

How you configure your audio session policy. You can do the same thing with the C API and `AudioSessionSetProperty`.

setActive:error:

How you bring your audio session policy into force. You'll need to call this any time your audio session has been deactivated. At a minimum, calling it in `applicationDidBecomeActive:` keeps your audio session policy in force in the multitasking world of iOS 4.

If you deactivate your audio session using `setActive:withFlags:error:`, passing a flag of `AVAudioSessionSetActiveFlags_NotifyOthersOnDeactivation`, you tell the system to allow any audio suspended by the activation of your audio session to resume. Thus you can activate and deactivate a Playback-policy audio session as your app runs as a way of pausing and resuming background audio (if the app providing the background audio responds correctly to this flag). I'll give an example later in this chapter.

You can do the same thing with the C API and `AudioSessionSetActive` or `AudioSessionSetActiveWithFlags`.

Here's an example from an app where we want background sound such as Music app songs to continue playing while our app runs. We configure our audio session to use the Ambient policy in `application:didFinishLaunchingWithOptions:`, as follows:

```
[[AVAudioSession sharedInstance] setCategory:  
    AVAudioSessionCategoryAmbient error: NULL];
```

Or, using the C API:

```
AudioSessionInitialize (NULL, NULL, NULL, NULL);  
UInt32 ambi = kAudioSessionCategory_AmbientSound;  
AudioSessionSetProperty(kAudioSessionProperty_AudioCategory, sizeof(ambi), &ambi);
```

We activate our audio session every time our app becomes active, no matter how, in `applicationDidBecomeActive:`, like this:

```
[[AVAudioSession sharedInstance] setActive: YES error: NULL];
```

Or, using the C API:

```
AudioSessionSetActive(true);
```

That's all it takes to set and enforce your app's overall audio session policy. Now let's make the example more interesting by introducing ducking. Just before we're about to play a sound, we duck any external sound:

```
UInt32 duck = 1;  
AudioSessionSetProperty(kAudioSessionProperty_OtherMixableAudioShouldDuck,  
    sizeof(duck), &duck);
```

When we finish playing a sound, we turn off ducking. This is the tricky part. Not only must we remove the ducking property from our audio session policy, but we must also deactivate our audio session to make the change take effect immediately and bring the external sound back to its original level; there is then no harm in reactivating our audio session:

```
UInt32 duck = 0;  
AudioSessionSetProperty(kAudioSessionProperty_OtherMixableAudioShouldDuck,  
    sizeof(duck), &duck);  
AudioSessionSetActive(false);  
AudioSessionSetActive(true);
```

Your audio session can be *interrupted*. This could mean that some other app deactivates it: for example, on an iPhone a phone call can arrive. In the multitasking world of iOS 4, it could mean that another app asserts its audio session over yours, or simply that your app is no longer in the foreground. If you initially set your audio session's category using Objective-C and `AVAudioSession`, you can handle interruptions in the `AVAudioSession`'s delegate. If you use the C API and `AudioSessionInitialize`, you can set up an interruption handler there (for example code, see the "Audio Session Cookbook" section of Apple's *Audio Session Programming Guide*).

Handling an interruption can be as simple as reactivating your audio session when the interruption ends. So, in `application:didFinishLaunchingWithOptions:`, when we set our audio session's category, we can also set ourselves as its delegate, and implement the delegate method `endInterruption`:

```
- (void)endInterruption {
    AudioSessionSetActive(true);
}
```

Even this may not be necessary; if you use an audio player (`AVAudioPlayer`, discussed in the next section), it provides its own delegate methods for notifying you of interruptions and activates your audio session for you when an interruption ends.

In the multitasking world of iOS 4, when your app switches to the background, your audio session is interrupted (unless your app plays audio in the background, as discussed later in this chapter). But because your app is now suspended in the background, you won't find out about this until the user brings your app back to the front. At that point, you'll get your audio session's `beginInterruption` and `endInterruption` delegate messages, plus your app delegate's `applicationDidBecomeActive:`. In general this shouldn't cause any difficulty, but it's something to be aware of: make sure your implementations of these methods don't step on one another in some undesirable way.

You should look over the documentation for both the Objective-C API and the C API. The C API is more powerful. With the C API, as we've already seen, you can implement ducking, and, for example, you can arrange to be notified when a property of your audio session changes, you can learn whether background audio is playing, and you can track and manipulate how audio is routed. Read Apple's *Audio Session Programming Guide* for a full overview of how you can use your app's audio session.

Audio Player

An *audio player* is an instance of the `AVAudioPlayer` class. This is the easiest way to play sounds with any degree of sophistication. A wide range of sound types is acceptable, including MP3, AAC, and ALAC, as well as AIFF and WAV. You can set a sound's volume and stereo pan features, loop a sound, synchronize the playing of multiple sounds simultaneously, and set playback to begin somewhere in the middle of a sound.

To use an audio player, you'll need to link to `AVFoundation.framework` and import `<AVFoundation/AVFoundation.h>`. An audio player should always be used in conjunction with an audio session; see the previous section.



Not every device type can play a compressed sound format in every degree of compression, and the limits can be difficult or impossible to learn except by experimentation. I encountered this issue when an app of mine worked correctly on an iPod touch 32GB but failed to play its sounds on an iPod touch 8GB (even though the latter was newer). Even more frustrating, the files played just fine in the Music app on *both* devices. The problem appears to be that the compression bit rate of my sound files was too low for `AVAudioPlayer` on the 8GB device, but not on the 32GB device. But there is no documentation of any such limit.

An audio player can possess and play only one sound, but you can have multiple audio players playing simultaneously. Devising a strategy for instantiating, retaining, and releasing your audio players is up to you. An audio player is initialized with its sound, using a local file URL or NSData. To play the sound, first tell the audio player to `prepareToPlay`, causing it to load buffers and initialize hardware; then tell it to `play`. The audio player's delegate is notified when the sound finishes playing. Other useful methods include `pause` and `stop`; the chief difference between them is that `pause` doesn't release the buffers and hardware set up by `prepareToPlay`, but `stop` does (so you'd want to call `prepareToPlay` again before resuming play).

For example, one of my apps uses a class called Player, which implements a `play:` method expecting a string path to a sound file in the app bundle. This method creates a new audio player, stores it as an instance variable, and tells it to play the sound file; it also sets itself up as that audio player's delegate, and emits a notification when the sound finishes playing. In this way, by maintaining a single Player instance, I can play different sounds in succession:

```
- (void) play: (NSString*) path {
    NSURL *fileURL = [[NSURL alloc] initFileURLWithPath: path];
    NSError* err = nil;
    AVAudioPlayer *newPlayer =
        [[AVAudioPlayer alloc] initWithContentsOfURL: fileURL error: &err];
    // error-checking omitted
    [fileURL release];
    self.player = newPlayer; // retain policy
    [newPlayer release];
    [self.player prepareToPlay];
    [self.player setDelegate: self];
    [self.player play];
}

- (void)audioPlayerDidFinishPlaying:(AVAudioPlayer *)player // delegate method
    successfully:(BOOL)flag {
    [[NSNotificationCenter defaultCenter]
        postNotificationName:@"soundFinished" object:nil];
}
```

Here are some useful audio player properties:

`pan`, `volume`

Stereo positioning and loudness, respectively.

`numberOfLoops`

How many times the sound should repeat after it finishes playing; thus, 0 (the default) means it doesn't repeat. A negative value causes the sound to repeat indefinitely (until told to `stop`).

`duration`

The length of the sound.



Figure 27-1. The software remote controls in the iOS 4 app switcher

currentTime

The playhead position within the sound.

Starting in iOS 4.0, the `playAtTime:` method allows playing to be scheduled to start at a certain time. The time should be described in terms of the audio player's `device-CurrentTime` property.

As I mentioned in the previous section, an audio player handles sound interruptions automatically. In particular, your audio session is reactivated for you when the interruption ends. You can implement the delegate methods `audioPlayerBegin-Interruption:` and `audioPlayerEndInterruption:` (or, in iOS 4, `audioPlayerEnd-Interruption:withFlags:`) to add functionality; for example, you might respond by updating your interface, or you might want to resume play when the interruption ends (by calling `play`). What you do will depends on what makes sense for your particular application. An example appears later in this chapter.

Remote Control of Your Sound

Various sorts of signal constitute *remote control*. There is hardware remote control; the user might be using earbuds with buttons, for example. There is also software remote control; for example, in iOS 4, the playback controls that you see when you double-click the Home button to view the fast app switcher and then swipe to the right (Figure 27-1) are a form of software remote control. Similarly, the buttons that appear if you double-click the Home button when the screen is locked and sound is playing are a form of software remote control (Figure 27-2).

Your app can arrange to be targeted by *remote control events* reporting that the user has tapped a remote control. This is particularly appropriate in an app that plays sound. Your sound-playing app can respond to the remote play/pause button, for example, by playing or pausing its sound.

Remote control events are a form of `UIEvent`, and they are sent initially to the first responder. (See Chapter 11 and Chapter 18 on `UIResponders` and the responder chain.) To arrange to be a recipient of remote control events:

- Your app must contain a `UIResponder` in its responder chain that returns YES from `canBecomeFirstResponder`, and that responder must actually be first responder.
- Some `UIResponder` in the responder chain, at or above the first responder, must implement `remoteControlReceivedWithEvent:`.

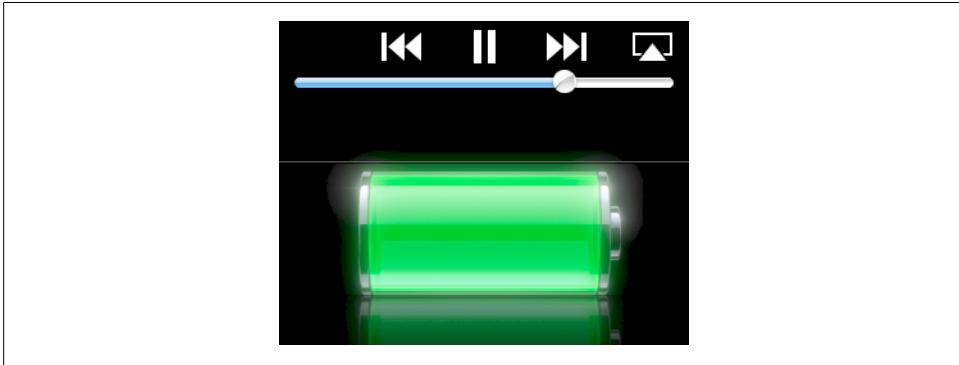


Figure 27-2. The software remote controls on the locked screen

- Your app must call the `UIApplication` instance method `beginReceivingRemoteControlEvents`.
- Your app audio session's policy must be Playback.
- Your app must emit some sound. The rule is that whatever running app capable of receiving remote control events last produced sound is the target of remote events. The user can tell what app this is because the icon at the right of [Figure 27-1](#) is the icon of that app. The remote control event target app defaults to the Music app if no other app takes precedence by this rule.

A typical place to put all of this is in your view controller, which is, after all, a `UIResponder`:

```
- (BOOL)canBecomeFirstResponder {
    return YES;
}

- (void)viewDidAppear:(BOOL)animated {
    [super viewDidAppear: animated];
    [self becomeFirstResponder];
    [[UIApplication sharedApplication] beginReceivingRemoteControlEvents];
}

- (void)remoteControlReceivedWithEvent:(UIEvent *)event {
    // ...
}
```

That's just a sketch, but it does work: when this app is running, the user can employ remote controls, such as the buttons in [Figure 27-1](#) or a physical button on earbuds, to cause `remoteControlReceivedWithEvent:` to be called.

The question then is how to implement `remoteControlReceivedWithEvent:`. Your implementation will examine the `subtype` of the incoming `UIEvent` in order to decide what to do. There are many possible `subtype` values, listed under `UISubtype` in the `UIEvent` class documentation; they have names like `UIEventSubtypeRemoteControlPlay`. A min-

imal implementation will respond to `UIEventSubtypeRemoteControlTogglePlayPause`. Here's an example in an app where sound is produced by an `AVAudioPlayer`:

```
- (void)remoteControlReceivedWithEvent:(UIEvent *)event {
    UIEventSubtype rc = event.subtype;
    if (rc == UIEventSubtypeRemoteControlTogglePlayPause) {
        if ([player isPlaying])
            [player pause];
        else
            [player play];
    }
}
```

Playing Sound in the Background

In the multitasking world of iOS 4, when the user switches away from your app to another app, by default, your audio session is interrupted and your app is suspended. But if the business of your app is to play sound, you might like your app to continue playing sound in the background. In earlier sections of this chapter, I've spoken about how your app, in the foreground, relates its sound production to background sound such as the Music app. Now we're talking about how your app can *be* that background sound, possibly playing sound while some other app is in the foreground.

To play sound in the background, your app must do two things:

- In your `Info.plist`, you must include the “Required background modes” key (`UIBackgroundModes`) with a value that includes “App plays audio” (`audio`).
- Your audio session’s policy must be Playback (and must be active, of course).

That's actually all it takes! If those two things are true, then if your app is producing sound, that sound will go right on playing when the user clicks the Home button and dismisses your application or switches to another app.

An extremely cool feature of playing sound in the background is that remote control events continue to work. Even if your app was not actively playing at the time it was put into the background, if it is the remote control target, then if the user causes a remote control event to be sent, your app will be woken up in the background in order to receive it and can begin playing sound. However, the rules for interruptions still apply; another app can interrupt your app's audio session while your app is in the background, and if that app receives remote control events, then your app is no longer the remote control target.

If your app is the remote control target in the background, then another app can interrupt your app's audio, play some audio of its own, and then deactivate its own audio session with the flag telling your app to resume playing. I'll give a minimal example of how this works with an `AVAudioPlayer`.

Let's call the two apps `BackgroundPlayer` and `Interrupter`. Suppose `Interrupter` has an audio session policy of `Ambient`. This means that when it comes to the front, back-

ground audio doesn't stop. But now Interrupter wants to play a sound of its own, temporarily stopping background audio. To pause the background audio, it sets its own audio session to Playback:

```
[[AVAudioSession sharedInstance]
    setCategory:AVAudioSessionCategoryPlayback error:NULL];
[[AVAudioSession sharedInstance] setActive:YES error:NULL];
[player setDelegate: self];
[player prepareToPlay];
[player play];
```

When Interrupter's sound finishes playing, the AVAudioPlayer's delegate is notified. In response, Interrupter deactivates its audio session with the flag; then it's fine for it to switch its audio session policy back to Ambient and activate it once again:

```
[[AVAudioSession sharedInstance] setActive:NO
    withFlags:AVAudioSessionSetActiveFlags_NotifyOthersOnDeactivation error:NULL];
[[AVAudioSession sharedInstance]
    setCategory:AVAudioSessionCategoryAmbient error:NULL];
[[AVAudioSession sharedInstance] setActive:YES error:NULL];
```

So much for Interrupter. Now let's turn to BackgroundPlayer. Remember, it is configured to play in the background and to receive remote control events. When Interrupter changes its own policy to Playback, BackgroundPlayer is interrupted. When Interrupter deactivates its audio session, BackgroundPlayer's AVAudioPlayer delegate is notified that the interruption has ended. It tests for the resume flag and, if it is set, starts playing again:

```
- (void)audioPlayerEndInterruption:(AVAudioPlayer *)p withFlags:(NSUInteger)flags {
    if (flags & AVAudioSessionInterruptionFlags_ShouldResume) {
        [p prepareToPlay];
        [p play];
    }
}
```



In [Chapter 11](#), I said that your app delegate will probably never receive the `applicationWillTerminate:` message, because by the time the app terminates, it will already have been suspended and incapable of receiving any events. However, an app that is playing sound in the background is *not* suspended, even though it is in the background. If it is terminated while playing sound in the background, it will receive `applicationDidEnterBackground:`, even though it has *already* received this previously when it was moved into the background, and then it *will* receive `applicationWillTerminate:`.

Further Topics in Sound

iOS is a powerful milieu for production and processing of sound; its sound-related technologies are extensive. Here are some further topics, beyond the scope of this book:

Recording sound

To record sound simply, use AVAudioRecorder. Your audio session policy will need to adopt a Record policy before recording begins.

Audio queues

Audio queues implement sound playing and recording through a C API with more granularity than the Objective-C AVAudioPlayer and AVAudioRecorder (though it is still regarded as a high-level API), giving you access to the buffers used to move chunks of sound data between stored sound data and sound hardware.

Extended Audio File Services

A C API for reading and writing sound files in chunks. It is useful in connection with technologies such as audio queues.

Audio Converter Services

A C API for converting sound files between formats.

Streaming audio

Audio streamed in real time over the network, such as an Internet radio station, can be played with Audio File Stream Services, in connection with audio queues.

OpenAL

An advanced technology for playing sound with fine control over its stereo stage and directionality.

Audio units

Plug-ins that filter and modify the nature and quality of a sound as it passes through them. See the *Audio Unit Hosting Guide for iOS*.

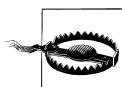
CoreMIDI

The CoreMIDI framework was introduced into iOS in version 4.2. It allows interaction with MIDI devices (but not, unfortunately, direct internal playback of MIDI files).

See also [Chapter 29](#) on accessing sound files in the user's music library.

CHAPTER 28

Video



This chapter not yet revised for iOS 5.

Basic video playback is performed in a view owned by a `MPMoviePlayerController`. You'll need to link to `MediaPlayer.framework` and import `<MediaPlayer/MediaPlayer.h>`. There are two relevant classes supplied by the Media Player framework:

`MPMoviePlayerController`

Vends and controls a view that plays a movie.

The behavior of this class has changed very greatly since it was introduced in iOS 2.0. It is difficult to use it compatibly with multiple system versions. In this chapter, I describe only its current behavior (in iOS 4.2 or later), with no attempt to discuss earlier differences or to advise you on backward compatibility.

`MPMoviePlayerViewController`

Owns an `MPMoviePlayerController`, and presents its view as a fullscreen view. Introduced in iOS 3.2.

A simple interface for trimming video (`UIVideoEditorController`) is also supplied. Starting in iOS 4, sophisticated video editing can be performed through AV Foundation, which this book does not discuss in depth.

A mobile device does not have unlimited power for decoding and presenting video in real time. A video that plays on your computer might not play at all on an iOS device. See the “Media Layer” chapter of Apple’s *iOS Technology Overview* for a list of specifications and limits within which video is eligible for playing.



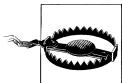
A web view ([Chapter 24](#)) supports the HTML 5 `<video>` tag. This can be a simple lightweight way to present video and to allow the user to control playback. Starting in iOS 4.3, both web view video and `MPMoviePlayerController` support AirPlay (streaming from the device to an Apple TV), but I have not tested this.

MPMoviePlayerController

An `MPMoviePlayerController` vends and controls a view, its `view` property, and a movie described by a URL, its `contentURL`, which it will present in that view. You are responsible for instantiating and retaining the `MPMoviePlayerController`, and for placing its `view` into your interface. No law says you *have* to put the `MPMoviePlayerController`'s view into your interface, but if you don't, the user won't be able to see the movie or the controls that accompany it by default. The view is a real view; you can set its `frame`, its `autoresizingMask`, and so forth, and you can give it subviews.

The movie URL can be a local file URL, so that the player can show, for example, a movie stored as a file in the app's bundle, or obtained from the Camera Roll / Saved Photos group in the user's photo library ([see Chapter 30](#)); or it can be a resource (possibly streamed) to be fetched over the Internet, in which case the `MPMoviePlayerController` initiates the download as soon as the `MPMoviePlayerController` has the `contentURL`.

Things happen slowly with a movie. Even when a movie is a local file, a certain amount of it has to load before the `MPMoviePlayer` actually has the movie and its specifications and can start playing it. In the case of a remote resource, this will take even longer. If `shouldAutoplay` is YES, play will begin as soon as it *is* possible.



By default, an `MPMoviePlayerController`'s `shouldAutoplay` is YES. This means that the movie will start loading and playing as soon as the `MPMoviePlayerController` has a `contentURL` — which, if you create it using `initWithContentURL:`, will be the moment it comes into existence. This will happen even if you don't put the `MPMoviePlayerController`'s view into your interface. If the movie has sound, the user will then hear it without being able to see it, which could be confusing. To prevent this, put the view into your interface, or set `shouldAutoplay` to NO (or both).

A movie file can be in a standard movie format, such as `.mov` or `.mp4`, but it can also be a sound file. An `MPMoviePlayerController` is thus an easy way to play a sound file, including a sound file obtained in real time over the Internet, along with standard controls for pausing the sound and moving the playhead.

In this example, we create an `MPMoviePlayerController`, give it a reference to a movie from our app bundle, retain it through a property, and put its view into our interface:

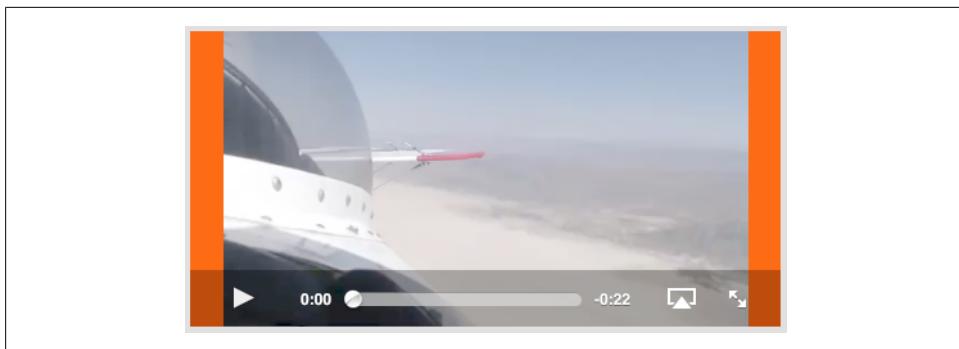


Figure 28-1. A movie player with controls



Figure 28-2. A movie player without controls

```
NSURL* m = [[NSBundle mainBundle] URLForResource:@"Movie" withExtension:@"m4v"];
MPMoviePlayerController* mp = [[MPMoviePlayerController alloc] initWithContentURL:m];
self.mpc = mp; // retain policy
[mp release];
self.mpc.shouldAutoplay = NO;
self.mpc.view.frame = CGRectMake(10, 10, 300, 230);
[self.view addSubview:self.mpc.view];
```

The controls (`controlStyle` is `MPMovieControlStyleEmbedded`) include a play/pause button, a slider for changing the current frame of the movie, and a fullscreen button; there may also be an AirPlay route button, if an appropriate device is found on the network (Figure 28-1).

The user can tap the view to show or hide the controls at the bottom; the controls may also disappear automatically after play begins (Figure 28-2).

If the user taps the fullscreen button (or pinches outwards) to enter fullscreen mode, the controls (`controlStyle` is `MPMovieControlStyleFullscreen`) at the top include a Done button, a slider, and an increased fullscreen button, and a second set of controls appears at the bottom with a play/pause button and rewind and fast-forward buttons, plus possibly a volume slider and an AirPlay route button. The user can tap to dismiss or

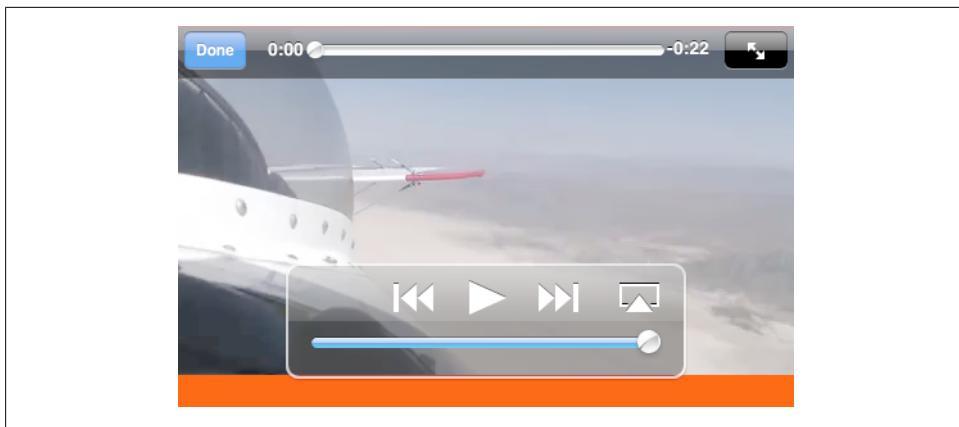


Figure 28-3. A movie player in fullscreen mode, with controls

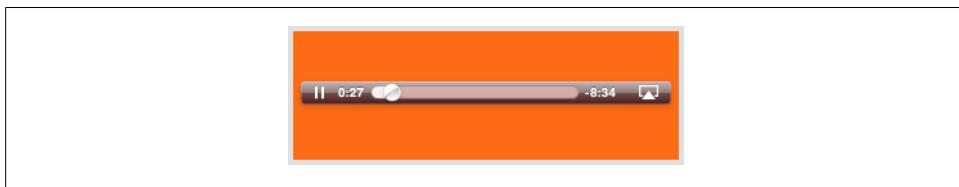


Figure 28-4. A movie player when the movie is a sound file

summon the controls, can double-tap to toggle increased fullscreen mode, and can tap Done to stop play and leave fullscreen mode ([Figure 28-3](#)).

You can also set the style of the controls (`controlStyle`) manually, though this would be an odd thing to do, because each style of control goes with a display mode (fullscreen or otherwise); you are most likely to use this feature to make it impossible for the user to summon the controls at all (`MPMovieControlStyleNone`).

If the movie is actually a sound file, the controls are drawn differently: there is a start/pause button, a slider, and possibly an AirPlay route button, and that's all ([Figure 28-4](#)).

The movie is scaled to fill the size of the view in accordance with the `MPMoviePlayerController scalingMode`; the default is `MPMovieScalingModeAspectFit`, which scales to fit, keeping the correct aspect ratio, and fills the unfilled dimension with the color of the `MPMoviePlayerController`'s `backgroundView`.

You might like to learn the actual size and aspect ratio of the movie, perhaps so as to eliminate the excess unfilled dimension. You can get the `MPMoviePlayerController`'s `naturalSize`, but, as I mentioned earlier, it takes time after the content URL is set for this to be determined. I'll show an example in a moment.

The fullscreen rendering can be rotated if the view in which the movie player subview is embedded can be rotated (because it is, or is in, a view controlled by a view controller

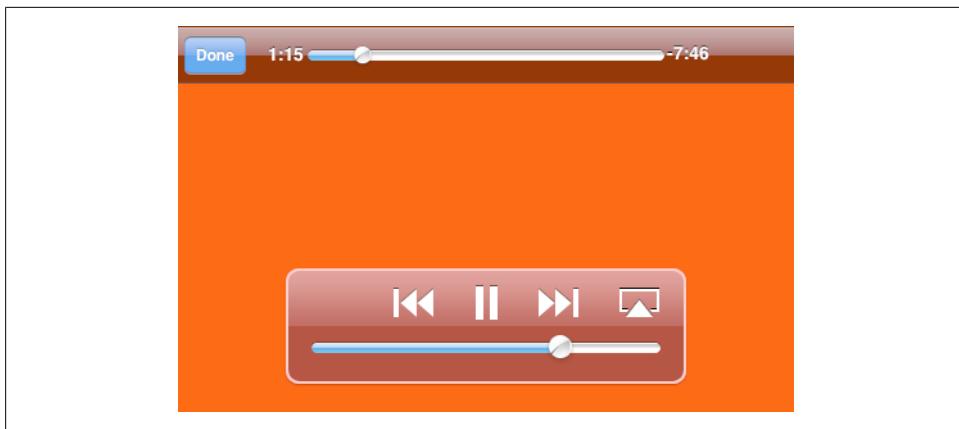


Figure 28-5. A fullscreen movie player when the movie is a sound file

that permits this). You can programmatically toggle between fullscreen and not, with `setFullscreen:animated:`. It is possible to toggle the view for a sound file into fullscreen mode programmatically even though it lacks a fullscreen control that the user can tap (Figure 28-5).

The movie can be made to repeat automatically (`repeatMode`) when it reaches its end. You can get the movie’s `duration`. You can change its `initialPlaybackTime` and `endPlaybackTime` (effectively trimming the start and end off the movie). Further programmatic control over the actual playing of the movie is obtained through the `MPMediaPlayback` protocol, which `MPMoviePlayerController` adopts. This gives you the expected `play`, `pause`, and `stop` methods, as well as commands for seeking quickly forward and backward, and you can get and set the `currentPlaybackTime` to position the playhead. Note the `prepareToPlay` method; it’s often a good idea to send this to the `MPMoviePlayerController` just after changing its content URL. You can also set the `currentPlaybackRate`, making the movie play slower or faster than normal, and even backward (though in my experience backward play doesn’t always work very well).

An `MPMoviePlayerController` doesn’t have a delegate. Instead, to learn of events as they happen, you must register for notifications. These notifications are how you know when, after assigning a content URL, it is safe for you to query properties of the movie such its `naturalSize` and `duration`. In this example, I’ll use a notification to embed the movie view into the interface, at the correct aspect ratio, as soon as the `naturalSize` is known:

```
- (void)setUpMPC {
    NSURL* m = [[NSBundle mainBundle] URLForResource:@"Movie" withExtension:@"m4v"];
    // ... the rest as before; do NOT add to view yet
    // [self.view addSubview:self.mpc.view];
    [[NSNotificationCenter defaultCenter] addObserver:self
        selector:@selector(finishSetup:)
        name:MPMovieNaturalSizeAvailableNotification
        object:self.mpc];
```

```

    }

    - (void) finishSetup: (id) n {
        CGRect f = self.mpc.view.bounds;
        f.size = self.mpc.naturalSize;
        // make width 300, keep ratio
        CGFloat ratio = 300/f.size.width;
        f.size.width *= ratio;
        f.size.height *= ratio;
        self.mpc.view.bounds = f;
        [self.view addSubview:self.mpc.view];
    }
}

```

Additional notifications tell such things as when fullscreen mode is entered and exited, and when the movie finishes playing. One of the most important notifications is `MPMoviePlayerPlaybackStateDidChangeNotification`; to learn the actual playback state, query the `MPMoviePlayerController`'s `playbackState`, which will be one of these:

- `MPMoviePlaybackStateStopped`
- `MPMoviePlaybackStatePlaying`
- `MPMoviePlaybackStatePaused`
- `MPMoviePlaybackStateInterrupted`
- `MPMoviePlaybackStateSeekingForward`
- `MPMoviePlaybackStateSeekingBackward`

If the content comes from the Internet, there is of course many a slip possible. Things take time; the Internet might slow down, or go away completely; the resource to be fetched might not exist. You'll want to register for notifications that tell you when things happen, and especially when things go wrong.

In this example, we've registered for the `MPMoviePlayerPlaybackDidFinishNotification`. There are two ways to detect an error by looking in the notification's `userInfo` dictionary: we can examine its key called `MPMoviePlayerPlaybackDidFinishReasonUserInfoKey`, which will be `MPMovieFinishReasonPlaybackError`; or, we can look to see if it has a key called `@"error"`, which will be an `NSError`:

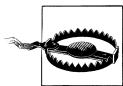
```

- (void) didFinish: (NSNotification*) n {
    // first way
    NSNumber* num =
        [[n userInfo] objectForKey:MPMoviePlayerPlaybackDidFinishReasonUserInfoKey];
    int reason = [num intValue];
    if (reason == MPMovieFinishReasonPlaybackError)
        NSLog(@"there was an error of some sort!");
    // second way
    NSError* err = [[n userInfo] objectForKey: @"error"];
    if (err)
        NSLog(@"%@", [err localizedDescription]);
}

```

However, the `MPMoviePlayerPlaybackDidFinishNotification` will not be sent if the movie starts downloading and playing, but the download is then cut off. To detect this, we register for `MPMoviePlayerLoadStateDidChangeNotification` and check for whether the `MPMoviePlayerController`'s `loadState` (a bitmask) has the `MPMovieLoadStateStalled` bit set. If so, we're in trouble, but play will not automatically stop; the `MPMoviePlayerController` will keep trying to obtain data. If we want to prevent that, we have to stop it manually:

```
- (void) loadStateChanged: (id) n {
    int ls = self.mpc.loadState;
    if (ls & MPMovieLoadStateStalled) {
        [self.mpc stop];
        NSLog(@"The download seems to have stalled out.");
    }
}
```



Only one `MPMoviePlayerController` can display a movie in your interface. To avoid confusion about why one of your `MPMoviePlayerController`s is not playing its movie successfully, restrict your interface so that it contains only one `MPMoviePlayerController` in the first place.

MPMoviePlayerViewController

An `MPMoviePlayerViewController` is, as its name implies, a view controller (a `UIViewController` subclass). It manages an `MPMoviePlayerController` (`moviePlayer`) and automatically provides a fullscreen presentation of the `MPMoviePlayerController`'s view. Thus, an `MPMoviePlayerViewController` has some strong advantages of simplicity. You don't have to put the `MPMoviePlayerController`'s view into your interface, and you don't have to worry about the user toggling between fullscreen and non-fullscreen modes. The view is either present, occupying the entire screen, or it isn't.

You can use an `MPMoviePlayerViewController` wherever you would use a `UIViewController`, pushing it onto a navigation stack, for example, or making it part of a tab bar interface, or presenting it. Here's a simple example:

```
NSURL* m = [[NSBundle mainBundle] URLForResource:@"Movie" withExtension:@"m4v"];
MPMoviePlayerViewController* mpvc =
    [[MPMoviePlayerViewController alloc] initWithContentURL: m];
[self presentModalViewController:mpvc animated:YES];
[mpvc release];
```

An alternative method for presenting the view is `presentMoviePlayerViewControllerAnimated:`. It uses a style of animation otherwise unavailable, in which the current view slides out to reveal the movie view.

If the `MPMoviePlayerViewController` is a presented view controller, it is dismissed automatically when the user taps the Done button. If you use the `MPMoviePlayerViewController` in some other way, the Done button stops play but that's all. You can

detect the tapping of the Done button by registering for the `MPMoviePlayerPlaybackDidFinishNotification`.

`MPMoviePlayerViewController` overrides `shouldAutorotateToInterfaceOrientation`: to return YES for the primary portrait orientation and both landscape orientations. This means the view is rotatable without your having to subclass `MPMoviePlayerViewController` and override `shouldAutorotateToInterfaceOrientation`: yourself.



When an `MPMoviePlayerViewController`'s view is showing, it becomes a recipient of remote control events (see [Chapter 27](#)). This feature is convenient, but if it's not what you want, it is not easily overcome; there is no property for turning it off. The best way to avoid it is to use an `MPMoviePlayerController` instead.

That's all there is to an `MPMoviePlayerViewController`; the rest of your interaction with it is through its `MPMoviePlayerController` (`moviePlayer`), including the latter's notifications.

UIVideoEditorController

`UIVideoEditorController` is a view controller that presents an interface for trimming video. Its view and internal behavior are outside your control. You are expected to show the view controller's view, probably as a presented view (or, on the iPad, in a popover, though in fact I haven't been able to get this view controller to work properly on the iPad no matter how I present it), and respond by way of its delegate.

Before summoning a `UIVideoEditorController`, be sure to call its class method `canEditVideoAtPath:`. Not every video format is editable, and not every device supports video editing. If this call returns NO, you can't use an instance of this class. (This call can take some noticeable time to return.) You must also set the `UIVideoEditorController` instance's delegate and `videoPath` before presenting it:

```
NSString* path = [[NSBundle mainBundle] pathForResource:@"movie" ofType:@"mov"];
BOOL can = [UIVideoEditorController canEditVideoAtPath:path];
if (!can) {
    NSLog(@"can't edit this video");
    return;
}
UIVideoEditorController* vc = [[UIVideoEditorController alloc] init];
vc.delegate = self;
vc.videoPath = path;
[self presentViewController:vc animated:YES];
[vc release];
```

The view's interface contains Cancel and Save buttons, a trimming box displaying thumbnails from the movie, a Play/Pause button, and the movie itself. The user slides the ends of the trimming box to set the beginning and end of the saved movie. The

Cancel and Save buttons do *not* dismiss the presented view; you must do that in your implementation of the delegate methods. There are three of them, and you should implement all three and dismiss the presented view in all of them:

- `videoEditorController:didSaveEditedVideoToPath:`
- `videoEditorControllerDidCancel:`
- `videoEditorController:didFailWithError:`

It's important to implement the `didFail...` method, because things can go wrong even at this stage.

By the time `videoEditorController:didSaveEditedVideoToPath:` is called, the video has already been saved to a file in your app's temporary directory (the same directory returned from a call to `NSTemporaryDirectory`). Doing something useful with the saved file is up to you; if you merely leave it in the temporary directory, you can't rely on it to persist. In this example, I copy the edited movie into the user's Camera Roll photo album (called Saved Photos if the device has no camera):

```
- (void) videoEditorController: (UIVideoEditorController*) editor
    didSaveEditedVideoToPath: (NSString*) editedVideoPath {
    if (UIVideoAtPathIsCompatibleWithSavedPhotosAlbum(editedVideoPath))
        UISaveVideoAtPathToSavedPhotosAlbum(editedVideoPath, nil, nil, NULL);
    else
        NSLog(@"need to think of something else to do with it");
    [self dismissModalViewControllerAnimated:YES];
}
```

Further Topics in Video

A large suite of AV Foundation classes was introduced in iOS 4 that provide detailed access to media components, analogous to QuickTime on Mac OS X. To access these, you'll need to link to `AVFoundation.framework` (and probably `CoreMedia.framework` as well), and import `<AVFoundation/AVFoundation.h>`. For a list of classes, see the *AV Foundation Framework Reference*.

The AV Foundation class that performs actual playing of media is `AVPlayer`. An `AVPlayer` has an `AVPlayerItem`; this is its media. An `AVPlayerItem` comprises tracks (`AVPlayerItemTrack`), which can be individually enabled or disabled. It gets these from its underlying `AVAsset`; this is the basic media unit, as it were, providing you with access to tracks and metadata.

Starting in iOS 4.1, an `AVPlayer` can be an `AVQueuePlayer`, a subclass that allows multiple `AVPlayerItems` to be loaded up and then played in sequence. `AVQueuePlayer` also has an `advanceToNextItem` method, and its list of items can be changed dynamically, so you could use it to give the user access to a set of "chapters."

To display an `AVPlayer`'s movie, you need an `AVPlayerLayer` (a `CALayer` subclass). You are unlikely to take this approach unless you need either the extended powers of

AV Foundation or the sequential playing power of AVQueuePlayer or the flexibility of working directly with a layer and Core Animation. The AVPlayerLayer doesn't even come with controls for playing a movie and visualizing its progress; you have to create these yourself. Nevertheless, simply displaying a movie in this way is quite easy:

```
NSURL* m = [[NSBundle mainBundle] URLForResource:@"movie" withExtension:@"mov"];
AVPlayer* p = [AVPlayer playerWithURL:m];
self.player = p; // might need a reference later
AVPlayerLayer* lay = [AVPlayerLayer playerLayerWithPlayer:p];
lay.frame = CGRectMake(10,10,300,200);
[self.view.layer addSublayer:lay];
```

To let the user choose to play the movie, we might provide a Play button. In this example, the button toggles the playing status of the movie by changing its rate:

```
- (IBAction) doButton: (id) sender {
    CGFloat rate = self.player.rate;
    if (rate < 0.01)
        self.player.rate = 1;
    else
        self.player.rate = 0;
}
```

Another intriguing feature of an AVPlayer is that you can coordinate animation in your interface ([Chapter 17](#)) with the playing of the movie. In other words, you attach an animation to a layer in more or less the usual way, but the animation takes place in movie playback time: if the movie is stopped, the animation is stopped, and if the movie is run at double rate, the animation runs at double rate. This is done by embedding the layer to be animated in an AVSynchronizedLayer, which is coupled with an AVPlayerItem.

To demonstrate, I'll extend the previous example; after we insert our AVPlayerLayer into the interface, we also create and insert an AVSynchronizedLayer:

```
// create synch layer, put it in the interface
AVPlayerItem* item = p.currentItem;
AVSynchronizedLayer* syncLayer =
    [AVSynchronizedLayer synchronizedLayerWithPlayerItem:item];
syncLayer.frame = CGRectMake(10,220,300,10);
syncLayer.backgroundColor = [[UIColor whiteColor] CGColor];
[self.view.layer addSublayer:syncLayer];
// give synch layer a sublayer
CALayer* subLayer = [CALayer layer];
subLayer.backgroundColor = [[UIColor blackColor] CGColor];
subLayer.frame = CGRectMake(0,0,10,10);
[syncLayer addSublayer:subLayer];
// animate the sublayer
CABasicAnimation* anim = [CABasicAnimation animationWithKeyPath:@"position"];
anim.fromValue = [NSValue valueWithCGPoint: subLayer.position];
anim.toValue = [NSValue valueWithCGPoint: CGPointMake(295,5)];
anim.removedOnCompletion = NO;
anim.beginTime = AVCoreAnimationBeginTimeAtZero; // important trick
anim.duration = CMTimeGetSeconds(item.asset.duration);
[subLayer addAnimation:anim forKey:nil];
```



Figure 28-6. The black square's position is synchronized to the movie

The result is shown in Figure 28-6. The white rectangle is the `AVSynchronizedLayer`, tied to our movie. The little black square inside it is its sublayer; if we animate the black square, that animation will be synchronized to the movie. We do animate the black square, changing its position from the left end of the white rectangle to the right end, starting at the beginning of the movie and with the same duration as the movie, but although we attach this animation to the black square layer in the usual way, the black square *doesn't move* until we tap the button to call `doButton:` and start the movie playing. Moreover, if we tap the button again to pause the movie, the black square stops. The black square is thus *automatically* representing the current play position within the movie!

Despite the simplicity of those examples, writing real-life AV Foundation code usually presents certain sophisticated challenges. For instance, as we've already seen, it takes time for media values to become available. Even with `MPMoviePlayerController` we couldn't fetch a movie's `naturalSize` immediately; we had to wait for an `MPMovieNaturalSizeAvailableNotification`. This is even more true with AV Foundation, where you use Objective-C blocks to call back into your code when values are ready or tasks are finished; moreover, such blocks may run on a background thread. AV Foundation also prefers key-value observing (Chapter 13) rather than notifications for letting you track value changes in your media (such as the progress of a movie as it plays).

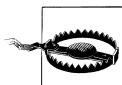
The mighty powers with which AV Foundation endows you are beyond the scope of this book. Here's a quick survey to whet your appetite. With AV Foundation, you can:

- Construct your own media asset (`AVComposition`, an `AVAsset` subclass). For example, you might combine part of the sound from one asset and part of the video from another into a single movie.
- Apply audio volume changes, and video opacity and transform changes, to the playback of individual tracks.
- Extract single images (“thumbnails”) from a movie (`AVAssetImageGenerator`).
- Export a movie in a different format (`AVAssetExportSession`).

- Capture audio, video, and stills, on a device that supports it (such as an iPhone, or another device connected to external hardware), including capturing video frames as still images (see Technical Q&A QA1702).

It should be evident from even so brief a summary that you could use AV Foundation to write a movie editor or a sound mixer. To learn more, you'll want to read the *AV Foundation Programming Guide*.

Music Library



This chapter not yet revised for iOS 5.

An iOS device, in addition to running apps, can also be used for the same purpose as the original iPod — to hold and play music and podcasts. These items, usually moved onto the device by way of iTunes on a computer, constitute the device’s *music library*; the user can play them with the Music app (formerly called the iPod app on some devices). Since version 3, iOS has provided the programmer with access to the device’s music library; this access has improved over time, and some aspects are as new as iOS 4.2. You can:

- Explore the music library
- Play an item from the music library
- Learn and control what the Music app’s music player is doing
- Present a standard interface for allowing the user to select a music library item

These abilities are provided by the Media Player framework. You’ll need to link to *MediaPlayer.framework* and import `<MediaPlayer/MediaPlayer.h>`.

Exploring the Music Library

Everything in the music library, as seen by your code, is an `MPMediaEntity`. This is an abstract class that endows its subclasses with the ability to describe themselves through key–value pairs called *properties*. (This use of the word “properties” has nothing to do with Objective-C properties, explained in [Chapter 12](#); these properties are more like entries in an `NSDictionary`, except that the keys are not objects.) The repertoire of properties depends on the sort of entity you’re looking at; many of them will be intuitively familiar from your use of iTunes. For example, a media item has a title, an album title, a track number, an artist, a composer, and so on; a playlist has a title, a flag

indicating whether it is a “smart” playlist, and so on. The property keys have names like `MPMediaItemPropertyTitle`.

To fetch a property’s value, call `valueForProperty:` with its key. Starting in iOS 4, you can fetch multiple properties with `enumerateValuesForProperties:usingBlock:`.

An individual item in the music library is an `MPMediaItem`, an `MPMediaEntity` subclass. It has a type, according to the value of its `MPMediaItemPropertyMediaType` property: it might be music, a podcast, or an audiobook. Different types of item have slightly different properties; for example, a podcast has a podcast title (in addition to its normal title). An item’s artwork image is an instance of the `MPMediaItemArtwork` class, from which you can get the image itself scaled to a specified size.

A playlist is an `MPMediaPlaylist`. As you would expect, it has `items` and a `count` of those items. It inherits those properties from its superclass, `MPMediaItemCollection`, which is the other `MPMediaEntity` subclass. I’ll talk more about `MPMediaItemCollection` in a moment.

Obtaining actual information from the music library requires a *query*, an `MPMediaQuery`. First, you *form* the query. There are two main ways to do this:

With a convenience constructor

`MPMediaQuery` provides several class methods that form a query ready to ask the music library for all of its songs, or all of its podcasts, and so on. Here’s the complete list:

- `songsQuery`
- `podcastsQuery`
- `audiobooksQuery`
- `playlistsQuery`
- `albumsQuery`
- `artistsQuery`
- `composersQuery`
- `genresQuery`
- `compilationsQuery`

With filter predicates

You can attach to the query one or more `MPMediaPropertyPredicate` instances, forming a set (`NSSet`) of predicates. These predicates filter the music library according to criteria you specify; to be included in the result, a media item must successfully pass through all the filters (in other words, the predicates are combined using logical-and). A predicate is a simple comparison. It has two, or possibly three, aspects:

A property

The key to the property you want to compare against. Not every property can be used in a filter predicate; the documentation makes the distinction clear (and starting in iOS 4.2, so does an `MPMediaEntity` class method, `canFilterByProperty:`).

A value

The value that the specified property must have in order to pass through the filter.

A comparison type (optional)

In order to pass through the filter, a media item's property value can either *match* the value you provide (`MPMediaPredicateComparisonEqualTo`, the default) or *contain* the value you provide (`MPMediaPredicateComparisonContains`).

These two ways of forming a query are actually the same; a convenience constructor is just a quick way of supplying a query already endowed with a filter predicate.

A query also *groups* its results, according to its `groupingType`. Your choices are:

- `MPMediaGroupingTitle`
- `MPMediaGroupingAlbum`
- `MPMediaGroupingArtist`
- `MPMediaGroupingAlbumArtist`
- `MPMediaGroupingComposer`
- `MPMediaGroupingGenre`
- `MPMediaGroupingPlaylist`
- `MPMediaGroupingPodcastTitle`

The query convenience constructors all supply a `groupingType` in addition to a filter predicate. Indeed, the grouping is often the salient aspect of the query. For example, an `albumsQuery` is in fact a `songsQuery` with the added feature that its results are grouped by album.

The groups resulting from a query are *collections*; that is, each is an `MPMediaItemCollection`. This class, you will recall, is the superclass of `MPMediaPlaylist`, and is an `MPMediaEntity` subclass. So, a collection has properties; it also has items and a count. It also has a `representativeItem` property, which gives you just one item from the collection. The reason you need this is that properties of a collection are often embodied in its items rather than in the collection itself. For example, an album has no title; rather, its items have album titles that are all the same. So to learn the title of an album, you ask for the album title of a representative item.

After you form the query, you *perform* the query. You do this simply by asking for the query's results. You can ask either for its `collections` (if you care about the groups returned from the query) or for its `items`. Here, I'll discover the titles of all the albums:

```

MPMediaQuery* query = [MPMediaQuery albumsQuery];
NSArray* result = [query collections];
// prove we've performed the query, by logging the album titles
for (MPMediaItemCollection* album in result)
    NSLog(@"%@", [[album representativeItem]
                  valueForProperty:MPMediaItemPropertyAlbumTitle]);

/*
Output starts like this on my device:
Bach Instrumental
Bach Keyboard
Bach Masses, Oratorios, Passions
Bach, CPE, Misc
Bach, CPE, Trio Sonatas
Beethoven Canons
Beethoven Choral / Vocal
...
*/

```

Now let's make our query more elaborate; we'll get the titles of all the albums whose name contains "Bach". Observe that what we really do is to ask for all songs whose album title contains "Bach", grouped by album:

```

MPMediaQuery* query = [MPMediaQuery albumsQuery];
MPMediaPropertyPredicate* hasBach =
    [MPMediaPropertyPredicate predicateWithValue:@"Bach"
        forProperty:MPMediaItemPropertyAlbumTitle
        comparisonType:MPMediaPredicateComparisonContains];
[query addFilterPredicate:hasBach];
NSArray* result = [query collections];
for (MPMediaItemCollection* album in result)
    NSLog(@"%@", [[album representativeItem]
                  valueForProperty:MPMediaItemPropertyAlbumTitle]);
/*
Complete output on my device:
Bach Instrumental
Bach Keyboard
Bach Masses, Oratorios, Passions
Bach, CPE, Misc
Bach, CPE, Trio Sonatas
*/

```

Because the results of that query are actually songs (MPMediaItems), we can immediately access any song in any of those albums. Let's modify our output from that query to print the titles of all the songs in the first album returned, which happens to be the Bach Instrumental album. We don't have to change our query, so I'll start at the point where we perform it:

```

// ... same as before ...
NSArray* result = [query collections];
MPMediaItemCollection* album = [result objectAtIndex: 0];
for (MPMediaItem* song in album.items)
    NSLog(@"%@", [song valueForProperty:MPMediaItemPropertyTitle]);
/*
Output starts like this:

```

```
Lute Suite In G Minor BWV 995 1. Prélude  
Lute Suite In G Minor BWV 995 2. Allemande  
Lute Suite In G Minor BWV 995 3. Courante  
Lute Suite In G Minor BWV 995 4. Sarabande  
...  
*/
```

One of the properties of an `MPMediaEntity` is its *persistent ID* (`MPMediaItemProperty-PersistentID` and `MPMediaPlaylistPropertyPersistentID`). This is important, as it uniquely identifies this song or playlist. No other means of identification is guaranteed unique; two songs or two playlists can have the same title, for example. Using the persistent ID, you can retrieve again at a later time the same song or playlist you retrieved earlier, even across launches of your app. Starting in iOS 4.2, the repertoire of available persistent IDs is extended to entities in general (`MPMediaEntityProperty-PersistentID`), album, artist, composer, and more.

While you are maintaining the results of a search, the contents of the music library may themselves change. For example, the user might connect the device to a computer and add or delete music with iTunes. This can put your results out of date. For this reason, the library's own modified state is available through the `MPMediaLibrary` class. Call the class method `defaultMediaLibrary` to get the actual library instance; now you can ask it for its `lastModifiedDate`. You can also register to receive a notification, `MPMediaLibraryDidChangeNotification`, when the music library is modified; this notification is not emitted unless you first send the library `beginGeneratingLibraryChangeNotifications`. You should eventually balance this with `endGeneratingLibraryChangeNotifications`.

The Music Player

The Media Player framework class for playing an `MPMediaItem` is `MPMusicPlayerController`. It comes in two flavors, depending on which class method you use to get an instance:

`applicationMusicPlayer`

Plays an `MPMediaItem` from the music library within your application. The song being played by the `applicationMusicPlayer` can be different from the Music app's current song. This player stops when your app is not in the foreground.

`iPodMusicPlayer`

The global music player—the very same player used by the Music app. This might already be playing an item, or be paused with a current item, at any time while your app runs; you can learn what item this is, and play music with this player. It continues playing independently of the state of your app. The user can at any time completely change what this player is doing.



An `applicationMusicPlayer` is not really inside your app. It is actually the global music player behaving differently. It has its own audio session; if you activate an audio session with a Solo Ambient or Playback policy ([Chapter 27](#)), the player will stop. You cannot play its audio when your app is in the background. You cannot make it the target of remote control events. If these limitations prove troublesome, use the `iPodMusicPlayer` (or `AVPlayer`, as discussed later in this chapter).

A music player doesn't merely play an item; it plays from a *queue* of items. This behavior is familiar from iTunes and the Music app, though you might not have considered it explicitly. For example, in iTunes, when you switch to a playlist and double-click the first song to start playing, when iTunes comes to the end of that song, it proceeds to the next song in the playlist. So at that moment, the totality of songs in the playlist is its queue. The music player behaves the same way; when it reaches the end of a song, it proceeds to the next song in its queue.

Your methods for controlling playback also reflect this queue-based orientation. In addition to the expected `play`, `pause`, and `stop` commands, there's a `skipToNextItem` and `skipToPreviousItem` command. Anyone who has ever used iTunes or the Music app (or, for that matter, an old-fashioned iPod) will have an intuitive grasp of this and everything else a music player does. For example, you can also set a music player's `repeatMode` and `shuffleMode`, just as in iTunes and so forth.

You provide a music player with its queue in one of two ways:

With a query

You hand the music player an `MPMediaQuery`. The query's `items` are the items of the queue.

With a collection

You hand the music player an `MPMediaItemCollection`. This might be obtained from a query you performed, but you can also assemble your own collection of `MPMediaItems` in any way you like, putting them into an array and calling `collectionWithItems:` or `initWithItems:`.

In this example, we collect all songs in the library shorter than 30 seconds into a queue and set the queue playing in random order using the application-internal music player:

```
MPMediaQuery* query = [MPMediaQuery songsQuery];
NSMutableArray* marr = [NSMutableArray array];
MPMediaItemCollection* queue = nil;
for (MPMediaItem* song in query.items) {
    CGFloat dur =
        [[song valueForProperty:MPMediaItemPropertyPlaybackDuration] floatValue];
    if (dur < 30)
        [marr addObject: song];
}
if ([marr count] == 0)
    NSLog(@"No songs that short!");
else
```

```

queue = [MPMediaItemCollection collectionWithItems:marr];
if (queue) {
    MPMusicPlayerController* player =
        [MPMusicPlayerController applicationMusicPlayer];
    [player setQueueWithItemCollection:queue];
    player.shuffleMode = MPMusicShuffleModeSongs;
    [player play];
}

```

If a music player is currently playing, setting its queue will stop it; restarting play appropriately is up to your code, if desired. Unfortunately, you can't query a music player as to its queue. You can keep your own copy of the array constituting the queue when you set the queue, but the user can completely change the queue of an `iPodMusicPlayer`, so if control over the queue is important to you, you'll have to use the `applicationMusicPlayer`.

A music player has a `playbackState` that you can query to learn what it's doing (whether it is playing, paused, stopped, or seeking). It also emits notifications so you can hear about changes in its state:

- `MPMusicPlayerControllerPlaybackStateDidChangeNotification`
- `MPMusicPlayerControllerNowPlayingItemDidChangeNotification`
- `MPMusicPlayerControllerVolumeDidChangeNotification`

These notifications are not emitted, however, until you tell the music player to `beginGeneratingPlaybackNotifications`. This is an instance method, so you can arrange to receive notifications from just one particular music player if you like. If you do receive notifications from both, you can distinguish them by examining the `NSNotification`'s `object` and comparing it to each player. You should eventually balance this call with `endGeneratingPlaybackNotifications`.

To illustrate, I'll extend the previous example to set a `UILabel` in our interface every time a different song starts playing. Before we set the player playing, we insert these lines to generate the notifications:

```

[player beginGeneratingPlaybackNotifications];
[[NSNotificationCenter defaultCenter] addObserver:self
    selector:@selector(changed:)
    name:MPMusicPlayerControllerNowPlayingItemDidChangeNotification
    object:nil];

```

And here's how we respond to those notifications:

```

- (void) changed: (NSNotification*) n {
    MPMusicPlayerController* player =
        [MPMusicPlayerController applicationMusicPlayer];
    if ([n object] == player) { // just playing safe
        NSString* title =
            [player.nowPlayingItem valueForProperty:MPMediaItemPropertyTitle];
        [self->label setText: title];
    }
}

```

There's no periodic notification as a song plays and the current playhead position advances. To get this information, you'll have to resort to polling. This is not objectionable as long as your polling interval is reasonably sparse; your display may occasionally fall a little behind reality, but this won't usually matter. For example, in one of my apps I use a `UIProgressView` (`p`) to show the current percentage of the current song played by the global player. There's no notification, so I use an `NSTimer` and poll the state of the player every 2 seconds. (I described this architecture in [Chapter 11](#), and showed some of the code triggered by the firing of this timer in [Chapter 19](#); [Figure 19-10](#) is a screenshot containing this `UIProgressView`.) If we are playing or paused in a song, I show the proportion played; otherwise, I hide the `UIProgressView` entirely:

```
MPMusicPlayerController* mp = [MPMusicPlayerController iPodMusicPlayer];
if ([mp playbackState] == MPMusicPlaybackStatePlaying ||
    [mp playbackState] == MPMusicPlaybackStatePaused) {
    p.hidden = NO;
    MPMediaItem* item = mp.nowPlayingItem;
    NSTimeInterval current = mp.currentPlaybackTime;
    NSTimeInterval total =
        [[item valueForProperty:MPMediaItemPropertyPlaybackDuration] doubleValue];
    p.progress = current / total;
} else {
    p.hidden = YES;
}
```

An `MPMusicPlayerController` has no user interface; if you want your app to provide the user with controls for playing and stopping a song, you'll have to provide those controls yourself. The `iPodMusicPlayer` has its own natural interface already, of course — namely, the Music app and the remote playback controls ([Figure 27-1](#)). The Media Player framework does offer a slider for setting the system output volume, along with an AirPlay route button if appropriate; this is an `MPVolumeView`.

`MPMusicPlayerController` is convenient and simple, but it's also simple-minded. As we've seen, its audio session isn't your audio session; the music player doesn't really belong to you. Starting in iOS 4.0, an `MPMediaItem` has an `MPMediaItemPropertyAssetURL` key, whose value is a URL suitable for forming an `AVAsset`. Thus, another way to play an `MPMediaItem` is through AV Foundation ([Chapter 28](#)). This approach gives you independence from the Music player; it puts playback of the song into your app's audio session and allows you to control it in response to remote control events and to play it while your app is in the background. (Of course, you can do a lot more with AV Foundation than merely to *play* a song from the music library. For example, you could incorporate a song, or part of a song, as the sound track to a movie.) In this simple example, we start with an array of `MPMediaItems` and initiate play of those items in an `AVQueuePlayer`:

```
NSArray* arr = // array of MPMediaItem;
NSMutableArray* assets = [NSMutableArray array];
for (MPMediaItem* item in arr) {
    AVPlayerItem* pi = [[AVPlayerItem alloc] initWithURL:
        [item valueForProperty:MPMediaItemPropertyAssetURL]];
```

```

[assets addObject:pi];
[pi release];
}
self.qp = [AVQueuePlayer queuePlayerWithItems:assets];
[self.qp play];

```

The Music Picker

The music picker (`MPMediaPickerController`) is a view controller (`UIViewController`) whose view is an self-contained navigation interface in which the user can select a media item. This interface looks very much like the Music app. You have no access to the actual view; you are expected to present the view controller (or, on the iPad, to use a popover).

You can limit the type of media items displayed by creating the controller using `initWithMediaTypes:`. You can make a prompt appear at the top of the navigation bar (`prompt`). And you can govern whether the user can choose multiple media items or just one, with the `allowsPickingMultipleItems` property. That's all there is to it.

While the view is showing, you learn what the user is doing through two delegate methods (`MPMediaPickerController`):

- `mediaPickerController:didPickMediaItems:`
- `mediaPickerControllerDidCancel:`

How you use these depends on the value of the controller's `allowsPickingMultipleItems`:

The controller's `allowsPickingMultipleItems` is NO (the default)

Every time the user taps a media item, your `mediaPickerController:didPickMediaItems:` is called, handing you an `MPMediaItemCollection` consisting of all items the user has tapped so far (including the same item multiple times if the user taps the same item more than once). When the user taps Cancel, your `mediaPickerControllerDidCancel:` is called.

The controller's `allowsPickingMultipleItems` is YES

The interface has Plus buttons at the right end of every media item, similar to the Music app interface for creating a playlist. When the user taps Done, `mediaPickerController:didPickMediaItems:` is called, handing you an `MPMediaItemCollection` consisting of all items for which the user has tapped the Plus button (including the same item multiple times if the user taps the same item's Plus button more than once). Your `mediaPickerControllerDidCancel:` is *never* called.

The view is *not* automatically dismissed; it is up to you to dismiss the presented view controller. The standard behavior, in order to manage the interface sensibly, would be for you to dismiss the presented view controller as soon as you get either delegate message.

In this example, we put up the music picker, allowing the user to choose one media item; we then play that media item with the application’s music player:

```
- (void) presentPicker {
    MPMediaPickerController* picker = [[MPMediaPickerController alloc] init];
    picker.delegate = self;
    [self presentModalViewController:picker animated:YES];
    [picker release];
}

- (void) mediaPickerController:(MPMediaPickerController*) mediaPickerController didPickMediaItems:(MPMediaItemCollection*) mediaItemCollection {
    MPMusicPlayerController* player =
        [MPMusicPlayerController applicationMusicPlayer];
    [player setQueueWithItemCollection:mediaItemCollection];
    [player play];
    [self dismissModalViewControllerAnimated:YES];
}

- (void) mediaPickerControllerDidCancel:(MPMediaPickerController*) mediaPickerController {
    [self dismissModalViewControllerAnimated:YES];
}
```

On the iPad, the music picker should be presented in a popover. (You can show the music picker as a presented view on the iPad, but my experience is that it then doesn’t work correctly. I regard as a bug the fact that you can present it at all; if that isn’t going to work, it would be better if the framework threw an exception.) On the iPad, the popover has no Cancel button, but this is no trouble because the user can tap outside the popover to dismiss it.

[Example 29-1](#) rewrites the preceding code so that on the iPhone it shows the picker as a presented view but on the iPad it shows the picker in a popover. The `presentPicker` method is now a button’s control event action, so that we can point the popover’s arrow to the button. How we summon the picker depends on the device (we use `UI_USER_INTERFACE_IDIOM` to distinguish the two cases); if it’s an iPad, we create a popover and set an instance variable to retain it (as discussed in [Chapter 22](#)). How we dismiss the picker depends on how it is being presented.

Example 29-1. A presented view on the iPhone, a popover on the iPad

```
- (void) presentPicker: (id) sender {
    MPMediaPickerController* picker =
        [[[MPMediaPickerController alloc] init] autorelease];
    picker.delegate = self;
    if (UI_USER_INTERFACE_IDIOM() == UIUserInterfaceIdiomPhone)
        [self presentModalViewController:picker animated:YES];
    else {
        UIPopoverController* pop =
            [[UIPopoverController alloc] initWithContentViewController:picker];
        self.currentPop = pop;
        [pop presentPopoverFromRect:[sender bounds] inView:sender
            permittedArrowDirections:UIPopoverArrowDirectionAny animated:YES];
        [pop release];
    }
}
```

```

        }
    }

- (void) dismissPicker: (MPMediaPickerController*) mediaPicker {
    if (self.currentPop && self.currentPop.popoverVisible) {
        [self.currentPop dismissPopoverAnimated:YES];
    } else {
        [self dismissModalViewControllerAnimated:YES];
    }
}

- (void)mediaPicker: (MPMediaPickerController *)mediaPicker
    didPickMediaItems:(MPMediaItemCollection *)mediaItemCollection {
    MPMusicPlayerController* player = [MPMusicPlayerController applicationMusicPlayer];
    [player setQueueWithItemCollection:mediaItemCollection];
    [player play];
    [self dismissPicker: mediaPicker];
}

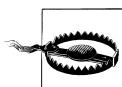
- (void)mediaPickerDidCancel:(MPMediaPickerController *)mediaPicker {
    [self dismissPicker: mediaPicker];
}

```



MPMediaPickerController is a good example of a view controller that behaves badly in a popover if the app is suspended to the background and then comes back to the foreground with the popover still visible (see my warning about this in [Chapter 22](#)).

Photo Library



This chapter not yet revised for iOS 5.

The still photos and movies accessed by the user through the Photos app constitute the *photo library*. Your app can provide an interface for exploring this library through the `UIImagePickerController` class. The same class can also be used to take photos and videos on devices with the necessary hardware.

Also, starting in iOS 4, the Assets Library framework lets you access the photo library and its contents programmatically.

You can write files into the Camera Roll / Saved Photos album. The ability to save an image file into the Camera Roll / Saved Photos album, by calling `UIImageWriteToSavedPhotosAlbum`, has existed since the early days of iOS 2.0. Some kinds of video file can also be saved into the Camera Roll / Saved Photos album; in an example in [Chapter 28](#), I checked whether this was true of a certain video file by calling `UIVideoAtPathIsCompatibleWithSavedPhotosAlbum`, and saved the file by calling `UISaveVideoAtPathToSavedPhotosAlbum`, functions that have existed since iOS 3.1. The Assets Library framework adds further options.

To use constants such as `kUTTypeImage`, referred to in this chapter, your app must link to `MobileCoreServices.framework` and import `<MobileCoreServices/MobileCoreServices.h>`.

`UIImagePickerController`

`UIImagePickerController` is a view controller (`UINavigationController`) whose view provides a navigation interface, similar to the Photos app, in which the user can choose an item from the photo library. Alternatively, it can provide an interface for taking a video or still photo if the necessary hardware is present. You will typically present the

view controller (or, on the iPad, show it in a popover; for code that does both, see [Example 29-1](#)).

Choosing from the Photo Library

To let the user choose an item from the photo library, instantiate UIImagePickerController and assign its `sourceType` one of these values:

- `UIImagePickerControllerSourceTypeSavedPhotosAlbum`
- `UIImagePickerControllerSourceTypePhotoLibrary`

You should call the class method `isSourceTypeAvailable:` beforehand; if it doesn't return YES, don't present the controller with that source type.

You'll probably want to specify an array of `mediaTypes` you're interested in. This array will usually contain `kUTTypeImage`, `kUTTypeMovie`, or both; or you can specify all available types by calling the class method `availableMediaTypesForSourceType:`.

After doing all of that, and having supplied a delegate, present the view controller:

```
UIImagePickerControllerSourceType type =
    UIImagePickerControllerSourceTypePhotoLibrary;
BOOL ok = [UIImagePickerController isSourceTypeAvailable:type];
if (!ok) {
    NSLog(@"alas");
    return;
}
UIImagePickerController* picker = [[UIImagePickerController alloc] init];
picker.sourceType = type;
picker.mediaTypes =
    [UIImagePickerController availableMediaTypesForSourceType:type];
picker.delegate = self;
[self presentModalViewController:picker animated:YES];
[picker release];
```

The delegate (`UIImagePickerControllerDelegate`) will receive one of these messages:

- `imagePickerController:didFinishPickingMediaWithInfo:`
- `imagePickerControllerDidCancel:`

If a delegate method is not implemented, the view controller is dismissed automatically, but rather than relying on this, you should implement both delegate methods and dismiss the view controller yourself in both. The `didFinish...` method is handed a dictionary of information about the chosen item. The keys in this dictionary depend on the media type.

An image

The keys are:

`UIImagePickerControllerMediaType`

A UTI; probably @"public.image", which is the same as `kUTTypeImage`.

UIImagePickerControllerOriginalImage

A UIImage.

UIImagePickerControllerReferenceURL

An ALAsset URL (discussed later in this chapter).

A movie

The keys are:

UIImagePickerControllerMediaType

A UTI; probably @"public.movie", which is the same as kUTTypeMovie.

UIImagePickerControllerMediaURL

A file URL to a copy of the movie saved into a temporary directory. This would be suitable, for example, to display the movie with an MPMoviePlayerController ([Chapter 28](#)).

UIImagePickerControllerReferenceURL

An ALAsset URL (discussed later in this chapter).

Optionally, you can set the view controller's `allowsEditing` to YES. In the case of an image, the interface then allows the user to scale the image up and to move it so as to be cropped by a preset square; the dictionary will include two additional keys:

UIImagePickerControllerCropRect

An NSValue wrapping a CGRect.

UIImagePickerControllerEditedImage

A UIImage.

In the case of a movie, the user can trim the movie just as with a UIVideoEditorController ([Chapter 28](#)). The dictionary keys are the same as before, but the file URL points to the trimmed copy in the temporary directory.



Because of restrictions on how many movies can play at once, if you use a UIImagePickerController to let the user choose a movie and you then want to play that movie in an MPMoviePlayerController, you must destroy the UIImagePickerController first. How you do this depends on how you created and presented the UIImagePickerController. If you're using a presented view controller, you can use delayed performance to ensure that the MPMoviePlayerController isn't configured until after the animation dismissing the presented view. If you're using a popover on the iPad, you can release the UIPopoverController (probably by nilifying the instance variable that's retaining it).

Using the Camera

To prompt the user to take a photo or video, instantiate UIImagePickerController and set its source type to `UIImagePickerControllerSourceTypeCamera`. Be sure to check `isSourceTypeAvailable`: beforehand; it will be NO if the user's device has no camera or

the camera is unavailable. If it is YES, call `availableMediaTypesForSourceType`: to learn whether the user can take a still photo (`kUTTypeImage`), a video (`kUTTypeMovie`), or both. The result will guide your `mediaTypes` setting. Set a delegate, and present the view controller.

Starting in iOS 3.1 (when video recording was introduced), you can also specify the `videoQuality` and `videoMaximumDuration`.

Starting in iOS 4.0, additional properties and class methods allow you to determine the camera capabilities:

isCameraDeviceAvailable:

Checks to see whether the front or rear camera is available, using one of these parameters:

- `UIImagePickerControllerCameraDeviceFront`
- `UIImagePickerControllerCameraDeviceRear`

cameraDevice

Lets you learn and set which camera is being used.

availableCaptureModesForCameraDevice:

Checks whether the given camera can capture still images, video, or both. You specify the front or rear camera; returns an NSArray of NSNumbers, from which you can extract the integer value. Possible modes are:

- `UIImagePickerControllerCameraCaptureModePhoto`
- `UIImagePickerControllerCameraCaptureModeVideo`

cameraCaptureMode

Lets you learn and set the capture mode (still or video).

isFlashAvailableForCameraDevice:

Checks whether flash is available.

cameraFlashMode

Lets you learn and set the flash mode (or, for a movie, toggles the “video torch”).

Your choices are:

- `UIImagePickerControllerCameraFlashModeOff`
- `UIImagePickerControllerCameraFlashModeAuto`
- `UIImagePickerControllerCameraFlashModeOn`

When the view controller appears, the user will see the interface for taking a picture, familiar from the Camera app, possibly including flash button, camera selection button, and digital zoom (if the hardware supports these), still/video switch (if your `mediaTypes` setting allows both), and Cancel and Shutter buttons.

Alternatively, you can hide the standard controls by setting `showsCameraControls` to NO, replacing them with your own overlay view, which you supply as the value of the

`cameraOverlayView`. In this case, you’re probably going to want a button in your overlay view to make the camera to take a picture! You can do that through these methods:

- `takePicture`
- `startVideoCapture`
- `stopVideoCapture`

You can supply a `cameraOverlayView` even if you don’t set `showsCameraControls` to NO, but you’ll need to work out a sensible frame if you don’t want your view to obscure existing controls.

You can also zoom or otherwise transform the preview image by setting the `cameraViewTransform` property. This can be tricky, not least because the default transform is not the identity transform, and also because different versions of iOS apply your transform differently; in iOS 4 and later, it is applied from the center, but before that it is applied from the top.

If you set `showsCameraControls` to NO, you may notice a blank area the size of the toolbar at the bottom of the screen. You can set the `cameraViewTransform` property so that this area is filled in, but knowing just what values to use can be tricky. Experimentation suggests that in iOS 4 this might be a good approach:

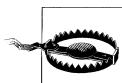
```
CGAffineTransform translate = CGAffineTransformMakeTranslation(0.0, 27.0);
CGAffineTransform scale = CGAffineTransformMakeScale(1.125, 1.125);
picker.cameraViewTransform = CGAffineTransformConcat(translate, scale);
```

I’m not very happy with that code, as it relies so heavily on “magic numbers,” but at least it provides a starting point.

Allowing the user to edit the captured image or movie, and handling the outcome with the delegate messages, is just as described in the previous section. In this example, having presented the interface for the user to take a still photo, we stick the resulting photo into a `UIImageView` in our interface:

```
- (void)imagePickerController:(UIImagePickerController *)picker
    didFinishPickingMediaWithInfo:(NSDictionary *)info {
    UIImage* im = [info objectForKey:UIImagePickerControllerOriginalImage];
    [self->iv setImage:im];
    [self dismissModalViewControllerAnimated:YES];
}
```

There won’t be any `UIImagePickerControllerReferenceURL` key in the dictionary delivered to the delegate because the image isn’t in the photo library. Also, starting in iOS 4.1, a still image might report a `UIImagePickerControllerMediaMetadata` key containing the metadata for the photo.



Setting camera-related properties such as `cameraViewTransform` when there is no camera or when the `UIImagePickerController` is not set to camera mode can crash your app.

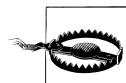


Figure 30-1. The user might see this

Instead of using UIImagePickerController, you can control the camera and capture images using the AV Foundation framework; see the “Media Capture” chapter of the *AV Foundation Programming Guide*.

The Assets Library Framework

The Assets Library framework, introduced in iOS 4, does for the photo library roughly what the Media Player framework does for the music library ([Chapter 29](#)), letting your code explore the library’s contents. You’ll need to link to `AssetsLibrary.framework` and import `<AssetsLibrary/AssetsLibrary.h>`. One obvious use of the Assets Library framework might be to implement your own interface for letting the user choose an image in a way that transcends the limitations of UIImagePickerController.



You might not be able to explore the photo library in this way; instead, you might get a “Global denied access” error, because the user must explicitly grant permission to access the library. If the user has never done that, the interface for giving such permission ([Figure 30-1](#)) requires that Location Services be turned on. So if the user has never given permission and Location Services is off, or if Location Services is on and the user refuses permission, you are blocked; your code must always be prepared for this possibility. (In my tests, even when I tapped OK in the [Figure 30-1](#) interface, my code seemed to halt; I didn’t get an error notification, but my code didn’t access the photo library either. Things worked fine the *next* time my code ran. This feels like a bug.)

A photo or video in the photo library is an ALAsset. Like a media entity ([Chapter 29](#)), an ALAsset can describe itself through key–value pairs called *properties*. (This use of the word “properties” has nothing to do with Objective-C properties, as explained in [Chapter 12](#).) For example, it can report its type (photo or video), its creation date, its orientation if it is a photo whose metadata reports this, and its duration if it is a video. You fetch a property value with `valueForProperty:`. The properties have names like `ALAssetPropertyName`.

A complicating factor in the case of photos is that a photo can provide multiple *representations* (roughly, image file formats). A given photo ALAsset lists these representations as one its properties, `ALAssetPropertyRepresentations`, an array of strings giving the UTIs identifying the file format; a typical UTI might be @"public.jpeg". A representation is an ALAssetRepresentation. You can get a photo's `defaultRepresentation`, or ask for a particular representation by submitting a file format's UTI to `representationForUTI:`.

Once you have an ALAssetRepresentation, you can interrogate it to get the actual image, either as raw data or as a CGImage (see [Chapter 15](#)). The simplest way is to ask for its `fullResolutionImage`; you may then want to derive a UIImage from this using `imageWithCGImage:scale:orientation:`. The original scale and orientation of the image are available as the ALAssetRepresentation's `scale` and `orientation`. An ALAssetRepresentation also has a `url`, which is important, because it's the unique identifier for the ALAsset (similar to the persistent ID of a song in the music library).

The photo library itself is an ALAssetsLibrary instance. It is divided into groups, which have types. For example, the user might have multiple albums; each of these is a group of type `ALAssetsGroupAlbum`. To fetch assets from the library, you either fetch one specific asset by providing its URL, or you enumerate its groups of a certain type. If you take the second approach, you are handed each group as an ALAssetsGroup; you can then enumerate the group's assets. Before doing so, you may optionally filter the group using a simple ALAssetsFilter; this limits any subsequent enumeration to photos only, videos only, or both. An ALAssetsGroup also has properties, such as a name, which you can fetch with `valueForProperty:`.

The Assets Library framework uses Objective-C blocks for fetching and enumerating assets and groups. These blocks behave rather oddly: at the end of the enumeration, they are called one extra time with a nil first parameter. Thus, you must code your block defensively to avoid treating the first parameter as real on that final call.

We now know enough for an example! I'll fetch the first photo from the album named "mattBestVertical" in my photo library and stick it into a UIImageView in the interface. For readability, I've set up the blocks in my code separately as variables *before* they are used, so it will help to read backward: we enumerate (at the end of the code) using the `getGroups` block (previously defined), which itself enumerates using the `getPix` block (defined before that). Here we go:

```
// what I'll do with the assets from the group
ALAssetsGroupEnumerationResultsBlock getPix =
^ (ALAsset *result, NSUInteger index, BOOL *stop) {
    if (!result)
        return;
    ALAssetRepresentation* rep = [result defaultRepresentation];
    CGImageRef im = [rep fullResolutionImage];
    UIImage* im2 = [UIImage imageWithCGImage:im];
    [self->iv setImage:im2]; // put image into our UIImageView
    *stop = YES; // got first image, all done
};
```

```

// what I'll do with the groups from the library
ALAssetsLibraryGroupsEnumerationResultsBlock getGroups =
    ^ (ALAssetsGroup *group, BOOL *stop) {
        if (!group)
            return;
        NSString* title = [group valueForProperty: ALAssetsGroupPropertyName];
        if ([title isEqualToString: @"mattBestVertical"]) {
            [group enumerateAssetsUsingBlock:getPix];
            *stop = YES; // got target group, all done
        }
    };
// might not be able to access library at all
ALAssetsLibraryAccessFailureBlock oops = ^ (NSError *error) {
    NSLog(@"oops! %@", [error localizedDescription]);
    // e.g. "Global denied access"
};
// and here we go with the actual enumeration!
ALAssetsLibrary* library = [[ALAssetsLibrary alloc] init];
[library enumerateGroupsWithTypes: ALAssetsGroupAlbum
                           usingBlock: getGroups
                           failureBlock: oops];
[library release];

```

As I mentioned at the start of this chapter, writing an image file into the Camera Roll / Saved Photos album has been supported since iOS 2.0, with the ability to write a movie file added in iOS 3.1. In iOS 4.0 and iOS 4.1, these functionalities were extended through the `ALAssetsLibrary` class, which offers these five methods:

`writeImageToSavedPhotosAlbum:orientation:completionBlock:`

Takes a `CGImageRef` and orientation.

`writeImageToSavedPhotosAlbum:metadata:completionBlock:`

Takes a `CGImageRef` and optional metadata dictionary (such as might arrive through the `UIImagePickerControllerMediaMetadata` key when the user takes a picture using `UIImagePickerController`).

`writeImageDataToSavedPhotosAlbum:metadata:completionBlock:`

Takes raw image data (`NSData`) and optional metadata.

`videoAtPathIsCompatibleWithSavedPhotosAlbum:`

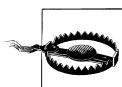
Takes a file path string. Returns a boolean.

`writeVideoAtPathToSavedPhotosAlbum:completionBlock:`

Takes a file path string.

Saving takes time, so a completion block allows you to be notified when it's over. The completion block supplies two parameters: an `NSURL` and an `NSError`. If the first parameter is not nil, the write succeeded, and this is the URL of the resulting `ALAsset`. If the first parameter is nil, the write failed, and the second parameter describes the error.

Address Book



This chapter not yet revised for iOS 5.

The user's address book, which the user sees through the Contacts app, is effectively a database that can be accessed directly through a C API provided by the Address Book framework. You'll link to *AddressBook.framework* and import `<AddressBook/AddressBook.h>`.

A user interface for interacting with the address book is also provided, through Objective-C classes, by the Address Book UI framework. You'll link to *AddressBookUI.framework* and import `<AddressBookUI/AddressBookUI.h>`.

Address Book Database

The address book is an `ABAddressBookRef` obtained by calling `ABAddressBookCreate`. This method's name contains "Create," so you must `CFRelease` the `ABAddressBookRef` when you're finished with it. The address book's data starts out exactly the same as the user's Contacts data. If you make any changes to the data, they are not written through to the user's real address book until you call `ABAddressBookSave`.

The primary constituent record of the address book database is the `ABPerson`. You'll typically extract persons from the address book by using these functions:

- `ABAddressBookGetPersonCount`
- `ABAddressBookGetPersonWithRecordID`
- `ABAddressBookCopyPeopleWithName`
- `ABAddressBookCopyArrayOfAllPeople`

The result of the latter two is a CFArrayRef. Their names contain “Copy,” so you must `CFRelease` the array when you’re finished with it. (I’m going to stop reminding you about memory management from here on.)

An ABPerson doesn’t formally exist as a type; it is actually an ABRecord (ABRecordRef), and by virtue of this has an ID, a type, and properties with values. To fetch the value of a property, you’ll call `ABRecordCopyValue`, supplying a property ID to specify the property that interests you. ABPerson properties, as you might expect, include things like first name, last name, and email.

Working with a property value is a little tricky because the way you treat it depends on what type of value it is. (You can find this out dynamically by calling `ABPersonGetTypeOfProperty`, but usually you’ll know in advance.) Some values are simple, but some are not. For example, a last name is a string, which is straightforward. But a person can have more than one email, so an email value is a “multistring.” To work with it, you’ll treat it as an ABMultiValue (ABMultiValueRef). This is like an array of values in which each item has, in addition to a value, a label and an identifier. The label categorizes (for example, a Home email as opposed to a Work email) but is not a unique specifier (because a person might have, say, two or more Work emails); the identifier is the unique specifier.

A person’s address is even more involved because not only is it an ABMultiValue (a person can have more than one address), but also a particular address is itself a dictionary (a CFDictinary). Each dictionary may have a key for street, city, state, country, and so on.

There is a lot more than this to parsing address book information, but this is enough to get you started and to illustrate by an example. I’ll fetch my own record out of the address book database on my device and detect that I’ve got two email addresses:

```
ABAddressBookRef adb = ABAddressBookCreate();
ABRecordRef moi = NULL;
CFArrayRef matts = ABAddressBookCopyPeopleWithName(adb, @"Matt");
// might be multiple matts, but let's find the one with last name Neuburg
for (CFIndex ix = 0; ix < CFArrayGetCount(matts); ix++) {
    ABRecordRef matt = CFArrayGetValueAtIndex(matts, ix);
    CFStringRef last = ABRecordCopyValue(matt, kABPersonLastNameProperty);
    if (last && CFStringCompare(last, (CFStringRef)@"Neuburg", 0) == 0)
        moi = matt;
    if (last)
        CFRelease(last);
}
if (NULL == moi) {
    NSLog(@"Couldn't find myself");
    CFRelease(matts);
    CFRelease(adb);
    return;
}
// parse my emails
ABMultiValueRef emails = ABRecordCopyValue(moi, kABPersonEmailProperty);
for (CFIndex ix = 0; ix < ABMultiValueGetCount(emails); ix++) {
```



Figure 31-1. A contact created programmatically

```
CFStringRef label = ABMultiValueCopyLabelAtIndex(emails, ix);
CFStringRef value = ABMultiValueCopyValueAtIndex(emails, ix);
NSLog(@"I have a %@ address: %@", (NSString*)label, (NSString*)value);
CFRelease(label);
CFRelease(value);
}
CFRelease(emails);
CFRelease(matts);
CFRelease(adbk);
/*
output:
I have a _$!<Home>!$_ address: matt@tidbits.com
I have a _$!<Work>!$_ address: mattworking@tidbits.com
*/
```

You can also modify an existing record, add a new record (`ABAddressBookAddRecord`), and delete a record (`ABAddressBookRemoveRecord`). In this example, I'll create a person, add him to the database, and save the database:

```
ABAddressBookRef adb = ABAddressBookCreate();
ABRecordRef snidely = ABPersonCreate();
ABRecordSetValue(snidely, kABPersonFirstNameProperty, @"Snidely", NULL);
ABRecordSetValue(snidely, kABPersonLastNameProperty, @"Whiplash", NULL);
ABMutableMultiValueRef addr = ABMultiValueCreateMutable(kABStringPropertyType);
ABMultiValueAddValueAndLabel(addr, @"snidely@villains.com", kABHomeLabel, NULL);
ABRecordSetValue(snidely, kABPersonEmailProperty, addr, NULL);
ABAddressBookAddRecord(adbk, snidely, NULL);
ABAddressBookSave(adbk, NULL);
CFRelease(addr);
CFRelease(snidely);
CFRelease(adbk);
```

Sure enough, if we then check the state of the database through the Contacts app, the new person exists ([Figure 31-1](#)).

Address Book Interface

The Address Book UI framework puts a user interface in front of common tasks involving the address book database and its manipulation by means of the functions and

Library Access Inconsistencies

You need special permission even to *look* in the photo library with the Assets Library framework ([Chapter 30](#)), and you can't modify the music library ([Chapter 29](#)), but you can freely read the user's address book, altering a person's details, adding new people, and even possibly deleting every record in the database, without the user's knowledge or permission. Don't ask me to justify that inconsistency; it makes no sense to me. Indeed, as Apple's documentation points out, your code has *more* power over the address book database than the user does, because you can manipulate groups of persons (ABGroup, not discussed here), but the user can't; the user acquires groups only by syncing with a Mac with groups in its Address Book.

data types discussed in the preceding section. This is a great help, because designing your own interface to do the same thing would be tedious and involved, especially given properties with multiple values and the added complexity of addresses. The framework provides four UIViewController subclasses:

ABPeoplePickerController

Presents a navigation interface, effectively the same as the Contacts app but without an Edit button: it lists the people in the database and allows the user to pick one and view the details.

ABPersonViewController

ABNewPersonViewController

ABUnknownPersonViewController

Presents an interface showing, respectively:

- The properties of a specific person, possibly editable
- The editable properties of a new person
- A proposed person with a partial set of noneditable properties

ABPeoplePickerController

An ABPeoplePickerController is a UINavigationController. Presenting it can be as simple as instantiating it, assigning it a delegate, and showing it as a presented view controller. (On the iPad, you'll probably use a popover; presenting the controller does work, but a popover looks better. For code that does both, see [Example 29-1](#).) The user can survey groups and the names of all contacts in each:

```
ABPeoplePickerController* picker =
    [[ABPeoplePickerController alloc] init];
picker.peoplePickerDelegate = self; // note: not merely "delegate"
[self presentViewController:picker animated:YES];
[picker release];
```

You should certainly provide a delegate, because without it the presented view will never be dismissed. This delegate is *not* the controller's `delegate` property! It is the controller's `peoplePickerDelegate` property. You should implement all three delegate methods:

`peoplePickerNavigationController:shouldContinueAfterSelectingPerson:`

The user has tapped a person in the contacts list, provided to you as an `ABRecordRef`. You have two options:

- Return NO. The user has chosen a person and that's all you wanted done. The selected person remains selected unless the user chooses another person. You are likely to dismiss the picker at this point.
- Return YES (and don't dismiss the picker). The view will navigate to a view of the person's properties. You can limit the set of properties the user will see at this point by setting the `ABPeoplePickerController`'s `displayedItems`. This is an array of `NSNumbers` wrapping the property identifiers such as `kABPersonEmailProperty`.

`peoplePickerNavigationController:shouldContinueAfterSelectingPerson:property:identifier:`

The user is viewing a person's properties and has tapped a property. Note that you are not handed the value of this property! You can fetch that yourself if desired, because you have the person and the property; plus, if the property has multiple values, you are handed an identifier so you can pick the correct one out of the array of values by calling `ABMultiValueGetIndexForIdentifier` and fetching the value at that index. You have two options:

- Return NO. The view is now still sitting there, displaying the person's properties. You are likely to dismiss the picker at this point.
- Return YES. This means that if the property is one that can be displayed in some other app, we will switch to that app. For example, if the user taps an address, it will be displayed in the Maps app; if the user taps an email, we will switch to the Mail app and compose a message addressed to that email.

`peoplePickerNavigationControllerDidCancel:`

The user has cancelled; you should dismiss the picker.

In this example, we want the user to pick an email. We have limited the display of properties to emails only:

```
picker.displayedProperties =  
    [NSArray arrayWithObject: [NSNumber numberWithInt: kABPersonEmailProperty]];
```

We return YES from the first delegate method. The second delegate method fetches the value of the tapped email and dismisses the picker:

```
- (BOOL)peoplePickerNavigationController:  
    (ABPeoplePickerController *)peoplePicker  
shouldContinueAfterSelectingPerson:(ABRecordRef)person  
property:(ABPropertyID)property
```

```

        identifier:(ABMultiValueIdentifier)identifier {
    ABMultiValueRef emails = ABRecordCopyValue(person, property);
    CFIndex ix = ABMultiValueGetIndexForIdentifier(emails, identifier);
    CFStringRef email = ABMultiValueCopyValueAtIndex(emails, ix);
    NSLog(@"%@", email); // do something with the email here
    CFRelease(email);
    CFRelease(emails);
    [self dismissModalViewControllerAnimated:YES];
    return NO;
}

```

ABPersonViewController

An `ABPersonViewController` is a `UIViewController`. To use it, instantiate it, set its `displayedPerson` and `personViewDelegate` (*not delegate*), and push it onto an existing navigation controller's stack. The user's only way out of the resulting interface will be through the Back button. You can limit the properties to be displayed, as with `ABPeoplePickerController`, by setting the `displayedProperties`. You can highlight a property with `setHighlightedItemForProperty:withIdentifier:`.

The delegate is notified when the user taps a property, similar to `ABPeoplePickerController`'s second delegate method illustrated in the code just above. As with `ABPeoplePickerController`'s delegate, return YES to allow some other app, such as Maps or Mail, to open the tapped value; return NO to prevent this.

If `ABPersonViewController`'s `allowsEditing` is YES, the right bar button is an Edit button. If the user taps this, the interface is transformed into the same sort of editing interface as `ABNewPersonViewController`. The user can tap Done or Cancel; if Done, the edits are automatically saved into the database. Either way, the user returns to the original display of the person's properties.

On the iPad, the same interface works, or alternatively you can use a popover. In that case you'll probably make the `ABPersonViewController` the root view of a `UINavigationController` created on the fly, especially if you intend to set `allowsEditing` to YES, since without the navigation interface the Edit button won't appear. No Back button is present or needed, because the user can dismiss the popover by tapping outside it.

ABNewPersonViewController

An `ABNewPersonController` is a `UIViewController`. To use it, instantiate it, set its `newPersonViewDelegate` (*not delegate*), instantiate a `UINavigationController` with the `ABNewPersonController` as its root view, and present the navigation controller:

```

ABNewPersonViewController* npvc = [[ABNewPersonViewController alloc] init];
npvc.newPersonViewDelegate = self;
UINavigationController* nc =
    [[UINavigationController alloc] initWithRootViewController:npvc];
[self presentModalViewController:nc animated:YES];
[nc release];
[npvc release];

```

The interface allows the user to fill in all properties of a new contact. You cannot limit the properties displayed. You can provide properties with default values by creating a fresh ABRecordRef representing an ABPerson with `ABPersonCreate`, giving it any property values you like, and assigning it to the `displayedPerson` property.

The delegate has one method, `newPersonViewController:didCompleteWithNewPerson:`, which is responsible for dismissing the presented view. If the new person is NULL, the user tapped Cancel. Otherwise, the user tapped Done; the new person is an ABRecordRef and has already been saved into the database.

But what if you don't want the new person saved into the database? What if you were presenting this interface merely because it's such a convenient way of letting the user fill in the property values of an ABPerson? Then simply remove the newly created person from the database, like this:

```

- (void)newPersonViewController:(ABNewPersonViewController*)newPersonViewController
    didCompleteWithNewPerson:(ABRecordRef)person {
    if (NULL != person) {
        ABAddressBookRef adb = ABAddressBookCreate();
        ABAddressBookRemoveRecord(adb, person, NULL);
        ABAddressBookSave(adb, NULL);
        CFStringRef name = ABRecordCopyCompositeName(person);
        NSLog(@"I have a person named %@", name); // do something with new person
        CFRelease(name);
        CFRelease(adb);
    }
    [self dismissModalViewControllerAnimated:YES];
}

```

The presented display works on the iPad as well. Alternatively, you can display the UINavigationController in a popover; the resulting popover is effectively modal, so it will be dismissed only through your implementation of `newPersonViewController:didCompleteWithNewPerson:`.

ABUnknownPersonViewController

An ABUnknownPersonViewController is a UIViewController. It presents, as it were, a proposed partial person. You can set the first and last name displayed as the controller's `alternateName` property and text below this as the controller's `message` property. You'll add other property values just as for an ABNewPersonViewController, namely, by creating a fresh ABRecordRef representing an ABPerson with `ABPersonCreate`, giving it some property values, and assigning it to the `displayedPerson` property.

To use ABUnknownPersonViewController, instantiate it, set the properties listed in the foregoing paragraph, set its `unknownPersonViewDelegate` (*not delegate*), and push it onto the stack of an existing navigation controller. The user's only way out of the resulting interface will be through the Back button.

On the iPad, make the ABUnknownPersonViewController the root view of a UINavigationController and present the navigation controller as a popover. No Back button is present or needed, because the user can dismiss the popover by tapping outside it.

What the user can do here depends on two other properties:

`allowsAddingToAddressBook`

If Yes, a Create New Contact button and an Add to Existing Contact button appear:

- If the user taps Create New Contact, the editing interface appears (as in ABNewPersonViewController and an editable ABPersonViewController). It is filled in with the property values of the `displayedPerson`. If the user taps Done, the person is saved into the database.
- If the user taps Add to Existing Contact, a list of all contacts appears (as in the first screen of ABPersonViewController). The user can Cancel or tap a person. If the user taps a person, the properties from the `displayedPerson` are merged into that person's record.

`allowsActions`

If YES, a Share Contact button appears. The user can tap this to edit and send a mail message containing a `.vcf` file embodying this contact.

The delegate has two methods, the first of which is required:

`unknownPersonViewController:didResolveToPerson:`

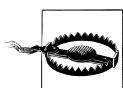
Called if `allowsAddingToAddressBook` is YES and the user finishes working in a presented editing view. The editing view has already been dismissed and the user has either cancelled (the second parameter is NULL) or has tapped Done (the second parameter is the ABPerson already saved into the database).

`unknownPersonViewController:shouldPerformDefaultActionForPerson:property:identifier:`

Return NO, as with ABPeoplePickerNavigationController, to prevent a tap on a property value from navigating to another app.

CHAPTER 32

Calendar



This chapter not yet revised for iOS 5.

The user's calendar information, which the user sees through the Calendar app, is effectively a database. Starting in iOS 4.0, this database can be accessed directly through the Event Kit framework. You'll link to *EventKit.framework* and import `<EventKit/EventKit.h>`.

A user interface for interacting with the calendar is also provided, by the Event Kit UI framework. This interface basically replicates part of the Calendar app. You'll link to *EventKitUI.framework* and import `<EventKitUI/EventKitUI.h>`.

Calendar Database

The calendar database is accessed as an instance of the `EKEventStore` class. Starting with this instance, you can obtain two kinds of object:

A calendar

A calendar is a collection of events, usually categorized for some purpose, such as Work or Home. It is an instance of `EKCalendar`. You can fetch all calendars with the `calendars` property, or the default calendar with the `defaultCalendarForNewEvents` property. Calendars have various types (`type`), reflecting the nature of their origin: a calendar can be created and maintained by the user locally (`EKCalendarTypeLocal`), but it might also live remotely on the network (`EKCalendarTypeCalDAV`, `EKCalendarTypeExchange`), possibly being updated by subscription (`EKCalendarTypeSubscription`); the Birthday calendar (`EKCalendarTypeBirthday`) is generated automatically from information in the address book. You can't create a calendar.

An event

An event is a memorandum describing when something happens. It is an instance of `EKEvent`. An event is associated with a calendar, its `calendar`, and you must

specify a calendar when creating an event. However, not every calendar permits modification or creation of events; if in doubt, check the calendar’s `allowsContentModifications` property. Given a calendar, you can’t fetch any events; rather, you fetch an event from the database as a whole. You do this either by unique identifier or by a date range.

I’ll give an example of fetching events from the calendar database in a moment, but first let’s focus on events. An event is an instance of `EKEvent`, and is where all the real action is. You fetch events out of the calendar database; you can modify an existing event, or make a new one; you can save an event into the calendar database with `saveEvent:span:error:`, and delete it from the database with `removeEvent:span:error:`. (You never save the database itself.)

If you’ve ever used the Calendar app, or iCal on the Mac, you have a sense for how an `EKEvent` can be configured. It has a `title` and optional `notes`. It is associated with a `calendar`, as I’ve already said. Most important, it has a `startDate` and an `endDate`; these are `NSDate`s and involve both date and time. If the event’s `allDay` property is YES, the time aspect of its dates is ignored; the event is associated with a day or a stretch of days as a whole. If the event’s `allDay` property is NO, the time aspect of its dates matters; a typical event will then usually be bounded by two times on the same day.

An event can be a recurring event, repeating at intervals according to some rule. That rule is the event’s `recurrenceRule`, and is an `EKRecurrenceRule`. A simple `EKRecurrenceRule` is described by three properties:

Frequency

By day, by week, by month, or by year.

Interval

Fine-tunes the notion “by” in the frequency. A value of 1 means “every.” A value of 2 means “every other.” And so on.

End

Optional, because the event might recur forever. It is an `EKRecurrenceEnd` instance, describing the limit of the event’s recurrence either as an end date or as a maximum number of occurrences.

The options for describing a more complex `EKRecurrenceRule` are best summarized by its initializer:

```
- (id)initRecurrenceWithFrequency:(EKRecurrenceFrequency)type
                           interval:(NSInteger)interval
                           daysOfTheWeek:(NSArray *)days
                           daysOfTheMonth:(NSArray *)monthDays
                           monthsOfTheYear:(NSArray *)months
                           weeksOfTheYear:(NSArray *)weeksOfTheYear
                           daysOfTheYear:(NSArray *)daysOfTheYear
                           setPositions:(NSArray *)setPositions
                           end:(EKRecurrenceEnd *)end
```

The meanings of all these parameters are mostly obvious from their names. The arrays are mostly of `NSNumber`, except for `daysOfTheWeek`, which is an array of `EKRecurrenceDayOfWeek`, a class that allows specification of a week number as well as a day number so that you can say things like “the fourth Thursday of the month.” The `setPositions` parameter is an array of numbers filtering the occurrences defined by the rest of the specification against the interval; however, it doesn’t seem to be always obeyed correctly. You can use any valid combination of parameters; the penalty for an invalid combination that can’t be resolved is a return value of `nil`.

Unfortunately, an `EKRecurrenceRule` has rather severe limitations. In theory, it is intended to reflect the `RRULE` event component in the iCalendar standard specification (original published as RFC 2445 and recently superseded by RFC 5545, <http://data-tracker.ietf.org/doc/rfc5545>). According to the standard, you should be allowed to specify both month and day in a yearly recurring event. The standard even gives an explicit example:

```
RRULE:FREQ=YEARLY;INTERVAL=2;BYMONTH=1;BYDAY=SU
```

This means “every Sunday in January, every other year.” But the documentation for `EKRecurrenceRule` says that `daysOfTheMonth` is valid only if the frequency is `EKRecurrenceFrequencyMonthly`, so you can’t use `EKRecurrenceRule` to form this rule.

For example, let’s consider how to express a single recurring event that marks when I have to pay my quarterly estimated taxes. The deadlines are the 15th of January, April, June, and September. This is not a regular interval (it isn’t every three months). The solution would seem to be to describe this as an annual recurring event, listing the relevant months explicitly as the `monthsOfTheYear` parameter, and putting the 15th of the month as the `daysOfTheMonth` parameter, like this:

```
EKRecurrenceRule* recur =
[[EKRecurrenceRule alloc]
initRecurrenceWithFrequency:EKRecurrenceFrequencyYearly
interval:1
daysOfTheWeek:nil
daysOfTheMonth:[NSArray arrayWithObject:[NSNumber numberWithInt:15]]
monthsOfTheYear:[NSArray arrayWithObjects:
    [NSNumber numberWithInt:1],
    [NSNumber numberWithInt: 4],
    [NSNumber numberWithInt: 6],
    [NSNumber numberWithInt: 9],
    nil]
weeksOfTheYear:nil
daysOfTheYear:nil
setPositions: nil
end:nil];
```

However, if we inspect the resulting recurrence rule, the `daysOfTheMonth` component has been thrown away, presumably because, as I just said, `daysOfTheMonth` is valid only if the frequency is `EKRecurrenceFrequencyMonthly`:

```
NSLog(@"%@", recur); // RRULE FREQ=YEARLY;INTERVAL=1;BYMONTH=1,4,6,9
```

So EKRecurrenceRule falls short of the standard. Nevertheless, it turns out that in this case, by also assigning the event itself a start and end date of the 15th of the month, we can obtain the desired event. I'll create the event and save it into the calendar database:

```
EKRecurrenceRule* recur =
[[EKRecurrenceRule alloc]
initRecurrenceWithFrequency:EKRecurrenceFrequencyYearly
interval:1
daysOfTheWeek:nil
daysOfTheMonth:nil
monthsOfTheYear:[NSArray arrayWithObjects:
[NSNumber numberWithInt: 1],
[NSNumber numberWithInt: 4],
[NSNumber numberWithInt: 6],
[NSNumber numberWithInt: 9],
nil]
weeksOfTheYear:nil
daysOfTheYear:nil
setPositions: nil
end:nil];
EKEventStore* database = [[EKEventStore alloc] init];
EKEvent* taxes = [EKEvent eventWithEventStore:database];
taxes.title = @"Estimated tax payment due";
taxes.recurrenceRule = recur;
NSCalendar* greg =
[[NSCalendar alloc] initWithCalendarIdentifier:NSGregorianCalendar];
NSDateComponents* comp = [[NSDateComponents alloc] init];
[comp setYear:2011];
[comp setMonth:4];
[comp setDay:15];
NSDate* date = [greg dateFromComponents:comp];
taxes.calendar = [database defaultCalendarForNewEvents];
taxes.startDate = date;
taxes.endDate = date;
taxes.allDay = YES;
NSError* err = nil;
BOOL ok = [database saveEvent:taxes span: EKSpanFutureEvents error:&err];
if (ok)
    NSLog(@"ok!");
else
    NSLog(@"error: %@", [err localizedDescription]);
[database release];
[greg release];
[comp release];
[recur release];
```

In that code, the event we save into the database is a recurring event. When we save or delete a recurring event, we must specify its `span`. This is either `EKSpanThisEvent` or `EKSpanFutureEvents`, and corresponds exactly to the two buttons the user sees in the Calendar interface when saving or deleting a recurring event ([Figure 32-1](#), and there is a similar choice on the Mac in iCal). The buttons and the span types reflect their meaning exactly: either the change affects this event alone, or this event plus all *future* (not

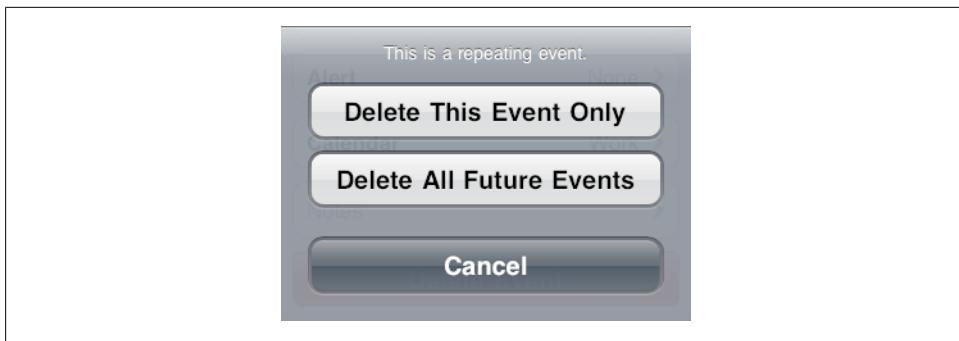


Figure 32-1. The user specifies a span

past) recurrences. This choice determines not only how this and future recurrences of the event are affected now, but also how they relate to one another from now on.

An EKEvent can have alarms. An alarm is an EKAlarm and can be set to fire either at an absolute date or at a relative offset from the event time. On an iOS device, an alarm fires through a local notification ([Chapter 26](#)).

An EKEvent can also be used to embody a meeting, with `attendees` (EKParticipant) and an `organizer`, but that is not a feature of an event that you can set.

Now let's return to the database as a whole. Once you've instantiated EKEventStore, you should maintain that instance for as long as your app will need access to the calendar database. This is partly because opening access to the database is somewhat expensive and time-consuming, and partly because an EKEvent belongs only to a single instance of EKEventStore.

An event in the database can be retrieved across instantiations of EKEventStore by its unique identifier (`eventIdentifier`), by calling `eventWithIdentifier:`; however, even this unique identifier might not survive changes in a calendar between launches of your app.

You can also extract events from the database by matching a predicate (NSPredicate). To form this predicate, you specify a start and end date and call the EKEventStore method `predicateForEventsWithStartDate:endDate:calendars:`. Because the date range is the only way to fetch events, any further filtering of events is up to you. In this example, I'll gather all events with "insurance" in their title; because I have to specify a date range, I ask for events occurring over the next year:

```
EKEventStore* database = [[EKEventStore alloc] init];
NSDate* d1 = [NSDate date];
NSDate* d2 = [NSDate dateWithTimeInterval:60*60*24*365 sinceDate:d1];
NSPredicate* pred =
    [database predicateForEventsWithStartDate:d1 endDate:d2
                                         calendars:database.calendars];
NSMutableArray* marr = [NSMutableArray array];
[database enumerateEventsMatchingPredicate:pred usingBlock:
    ^(EKEvent *event, BOOL *stop) {
```

```

NSRange r = [event.title rangeOfString:@"insurance"
                                options:NSCaseInsensitiveSearch];
if (r.location != NSNotFound)
    [marr addObject: event];
}];
```

After that, `marr` contains three events, because I pay my motorcycle insurance annually and my car insurance biannually. Those events are in no particular order; a convenience method on `EKEvent`, `compareStartDateWithEvent:`, is suitable as a sort selector:

```
[marr sortUsingSelector:@selector(compareStartDateWithEvent:)];
```

When you extract events from the database, event recurrences are treated as separate events. Thus, for example, although we created the tax payment event as a single recurring event, we might retrieve as many as four `EKEvent` instances of it over the next year. These `EKEvents` will have different start and end dates but the same `eventIdentifier`. When you fetch an event by calling `eventWithIdentifier:` you get the *earliest* event with that identifier. This makes sense, because if you’re going to make a change affecting this and future recurrences of the event, you need the option to start with the earliest possible recurrence (so that “future” means “all”).

The calendar database is an odd sort of database, because calendars can be maintained in so many ways and places. A calendar can change while your app is running (the user might sync, for example), which can put your information out of date. You can register for a single `EKEventStore` notification, `EKEventStoreChangedNotification`; if you receive it, you should assume that any calendar-related instances you’re holding are invalid. If you are in the middle of editing an event, you can reload its properties by calling `refresh`, but in general the database, and all events and calendars you’re working with, need to be fetched all over again.

Calendar Interface

The graphical interface consists of two views for letting the user work with an event:

`EKEventViewController`

Shows the description of a single event, possibly editable.

`EKEventEditViewController`

Allows the user to create or edit an event.

`EKEventViewController` simply shows the little rounded rectangle containing the event’s title, date, and time, familiar from the Calendar app, possibly with additional rounded rectangles describing alarms and notes. The user can’t tap these to do anything. To use `EKEventViewController`, instantiate it, give it an event, and push it onto the stack of an existing `UINavigationController`. The user’s only way out will be the Back button. So, for example:

```
EKEEventViewController* evc = [[EKEEventViewController alloc] init];
evc.event = [marr objectAtIndex:0];
evc.allowsEditing = NO;
[self.navigationController pushViewController:evc animated:YES];
[evc release];
```

The documentation says that `allowsEditing` is NO by default, but in my testing the default was YES; perhaps you'd best play safe and set it regardless. If it is YES, an Edit button appears as the right bar button item in the navigation bar, and by tapping this, the user can edit the various aspects of an event in the same navigation interface that should be familiar from the Calendar app, including the large red Delete button at the bottom. If the user ultimately deletes the event or edits it and taps Done, the change is saved into the database.

Starting in iOS 4.2, you can assign the EKEEventViewController a delegate in order to hear about what the user did. The delegate method `eventViewController:didCompleteWithAction:` is called only if the user deletes an event or accepts an invitation.

On the iPad, you use the EKEEventViewController as the root view of a navigation controller created on the fly and use the navigation controller as a popover's view controller. By default, the EKEEventViewController's `modalInPopover` is YES and a Done button appears as the right bar button; the delegate method `eventViewController:didCompleteWithAction:` is called if the user taps the Done button, and you'll need to dismiss the popover there. However, if `allowsEditing` is YES, you must set `modalInPopover` to NO, so to replace the Done button with the Edit button. Here's code that works both on the iPhone and on the iPad:

```
EKEEventViewController* evc = [[EKEEventViewController alloc] init];
evc.delegate = self;
evc.event = // ... whatever ...
evc.allowsEditing = NO;
// on iPhone, push onto existing navigation interface
if (UI_USER_INTERFACE_IDIOM() == UIUserInterfaceIdiomPhone)
    [self.navigationController pushViewController:evc animated:YES];
// on iPad, create navigation interface in popover
else {
    UINavigationViewController* nc =
        [[UINavigationController alloc] initWithRootViewController:evc];
    evc.modalInPopover = NO;
    UIPopoverController* pop =
        [[UIPopoverController alloc] initWithContentViewController:nc];
    self.currentPop = pop;
    [pop presentPopoverFromRect:[sender bounds] inView:sender
        permittedArrowDirections:UIPopoverArrowDirectionAny animated:YES];
    [nc release];
    [pop release];
}
[evc release];
```

EKEEventEditViewController (a UINavigationController) presents the interface for editing an event. To use it, set its `eventStore` and `editViewDelegate` (*not* delegate), and optionally its `event`, and present it as a presented view controller (or, on the iPad, in a

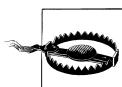
popover). The event can be nil for a completely empty new event; it can be an event you've just created (and possibly partially configured) and not stored in the database, or it can be an existing event from the database.

The delegate method `eventEditViewControllerDefaultCalendarForNewEvents:` may be implemented to specify what calendar a completely new event should be assigned. If you're partially constructing a new event, you can assign it a calendar then, and of course an event from the database already has a calendar.

You must implement the delegate method `eventEditViewController:didCompleteWithAction:` so that you can dismiss the presented view. Possible actions are that the user cancelled, saved the edited event, or deleted an already existing event from the database.

On the iPad, the presented view works, or you can present the EKEventEditViewController as a popover. By default, its `modalInPopover` is YES, which means you'll use `eventEditViewController:didCompleteWithAction:` to dismiss the popover.

Both view controllers automatically listen for changes in the database and, if needed, will automatically call `refresh` on the event being edited, updating their display to match. If the event was in the database and has been deleted while the user is viewing it, the delegate will get the same notification as if the user had deleted it.



This chapter not yet revised for iOS 5.

Your app can present interface allowing the user to edit and send a mail message (and, starting in iOS 4, an SMS message). Two view controller classes are provided by the Message UI framework; your app will link to *MessageUI.framework* and import `<MessageUI/MessageUI.h>`. The classes are:

MFMailComposeViewController

Allows composition and sending of a mail message.

MFMessageComposeViewController

Allows composition and sending of an SMS message.

Mail Message

The MFMailComposeViewController class, a UINavigationController, allows the user to edit a mail message. The user can attempt to send the message there and then or can cancel but save a draft, or can cancel completely. Before using this class to present a view, call `canSendMail`; if the result is NO, go no further. A negative result means that the device is not configured for sending mail. A positive result does not mean that the device is connected to the network and can send mail right now, only that sending mail is generally possible with this device.

To use MFMailComposeViewController, instantiate it, provide a `mailComposeDelegate` (*not delegate*), and configure the message to any desired extent. The user can alter your preset configurations. Configuration methods are:

- `setSubject:`
- `setToRecipients:`
- `setCcRecipients:`

- `setBccRecipients:`
- `setMessageBody:isHTML:`
- `addAttachmentData:mimeType:fileName:`

The delegate method `mailComposeController:didFinishWithResult:error:` is called with the user's final action, which might be any of these:

- `MFMailComposeResultCancelled`
- `MFMailComposeResultSaved`
- `MFMailComposeResultSent`
- `MFMailComposeResultFailed`

Typically, you'll show the `MFMailComposeViewController` as a presented view controller. Dismissing the presented view is up to you, in the delegate method. You don't *have* to use it as a presented view controller, but it's a sensible interface and works well even on the iPad (use `UIModalPresentationFormSheet` if a full-screen presentation feels too overwhelming).

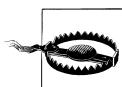
SMS Message

The `MFMessageComposeViewController` class, a `UINavigationController`, is even simpler. Instantiate the class, give it a `messageComposeDelegate`, configure it as desired, and show it as a presented view controller. Configuration is through the `recipients` and `body` properties.

The delegate method `messageComposeViewController:didFinishWithResult:` is called with the same possible results as for a mail message. Dismissing the presented view is up to you, in the delegate method.

CHAPTER 34

Maps



This chapter not yet revised for iOS 5.

Your app can imitate the Maps app, communicating with Google Maps to present a map interface and placing annotations and overlays on the map. UIView subclasses for displaying the map, along with the programming API, are provided by the Map Kit framework. You'll link to *MapKit.framework* and import `<MapKit/MapKit.h>`. You might also need the Core Location framework to express locations by latitude and longitude; you'll link to *CoreLocation.framework* and import `<CoreLocation/CoreLocation.h>`.

Presenting a Map

A map is presented through a UIView subclass, an MKMapView. The map is potentially a map of the entire world; the map view is usually configured to display a particular area. An MKMapView instance can be created in code or through the nib editor. A map has a type, which is one of the following:

- `MKMapViewTypeStandard`
- `MKMapViewTypeSatellite`
- `MKMapViewTypeHybrid`

The area displayed on the map is its `region`, an MKCoordinateRegion. This is a struct comprising a location (a `CLLocationCoordinate2D`), describing the latitude and longitude of the point at the center of the region, along with a span (an `MKCoordinateSpan`), describing the quantity of latitude and longitude embraced by the region. Convenience functions help you construct an MKCoordinateRegion.

In this example, I'll initialize the display of an MKMapView to show a place where I like to go dirt biking. The MKMapView is placed into the interface through the nib

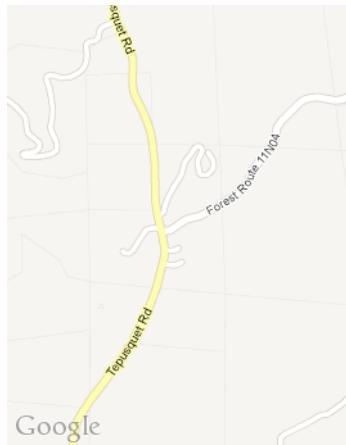


Figure 34-1. A map view showing a happy place

editor and is initially hidden so that the user doesn't see the default map of the world. I provide the region by setting the map view's `region` property, and show the view (Figure 34-1):

```
CLLocationCoordinate2D loc = CLLocationCoordinate2DMake(34.923964, -120.219558);
MKCoordinateRegion reg = MKCoordinateRegionMakeWithDistance(loc, 1000, 1000);
self->map.region = reg;
self->map.hidden = NO;
```

By default, the user can zoom and scroll the map with the usual gestures; you can turn this off by setting the map view's `zoomEnabled` and `scrollEnabled` to NO. Usually you will set them both to YES or both to NO. For example, if your aim is to prevent the user from changing the center coordinate, setting `scrollEnabled` to NO is insufficient, because the user can still zoom, and zooming includes double-tapping, which can change the center coordinate. (Workarounds to allow the user to zoom while keeping the center coordinate fixed, perhaps by implementing the delegate method `mapView:regionDidChangeAnimated:`, can quickly become ugly and difficult to maintain, involving raising and lowering of boolean flags and so forth.)

You can change the region displayed, optionally with animation, by calling these methods:

- `setRegion:animated:`
- `setCenterCoordinate:animated:`
- `setVisibleMapRect:animated:`
- `setVisibleMapRect:edgePadding:animated:`

`MKMapRect` (used in `setVisibleMapRect:animated:`) was introduced in iOS 4. Along with `MKMapPoint` and `MKMapSize`, it can simplify calculations by projecting the world's surface onto a flat rectangle. I'll give an example later in this chapter.

Annotations

An *annotation* is a marker associated with a location on a map. To make an annotation appear on a map, two objects are needed:

The object attached to the MKMapView

The annotation itself is attached to the MKMapView. It consists of any instance whose class adopts the MKAnnotation protocol, which specifies a coordinate, a title, and a subtitle for the annotation. You might have reason to define your own class to handle this task, or you can use the simple built-in MKPointAnnotation class. The annotation's `coordinate` is its most important property; this says where on earth the annotation should be drawn. The title and subtitle are optional, to be displayed in a callout.

The object that draws the annotation

An annotation is drawn by an MKAnnotationView, a UIView subclass. This can be extremely simple. In fact, even a nil MKAnnotationView might be perfectly satisfactory: it draws a red pin. If red is not your favorite color, a built-in MKAnnotationView subclass, MKPinAnnotationView, displays a pin in red, green, or purple; by convention you are supposed to use these colors for different purposes (destination points, starting points, and user-specified points, respectively). For more flexibility, you can provide your own UIImage as the MKAnnotationView's `image` property. And for even *more* flexibility, you can take over the drawing of an MKAnnotationView by overriding `drawRect:` in a subclass.

Not only does an annotation require two separate objects, but in fact those objects do not initially exist together. An annotation object has no pointer to the annotation view object that will draw it. Rather, it is up to you to supply the annotation view object in real time, on demand, in the MKMapView's delegate. This architecture may sound confusing, but in fact it's a very clever way of reducing the amount of resources needed at any given moment. Think of it this way: an annotation itself is merely a lightweight object that a map can always possess; the corresponding annotation view is a heavyweight object that is needed only so long as that annotation's coordinates are within the visible portion of the map.

To illustrate the simplest possible case, let's return to the code where we initially configured our map. Here's the same code again, but this time I'll add an annotation:

```
CLLocationCoordinate2D loc = CLLocationCoordinate2DMake(34.923964, -120.219558);
MKCoordinateRegion reg = MKCoordinateRegionMakeWithDistance(loc, 1000, 1000);
self->map.region = reg;
MKPointAnnotation* ann = [[MKPointAnnotation alloc] init];
ann.coordinate = loc;
ann.title = @"Park here";
ann.subtitle = @"Fun awaits down the road!";
[self->map addAnnotation:ann];
[ann release];
self->map.hidden = NO;
```

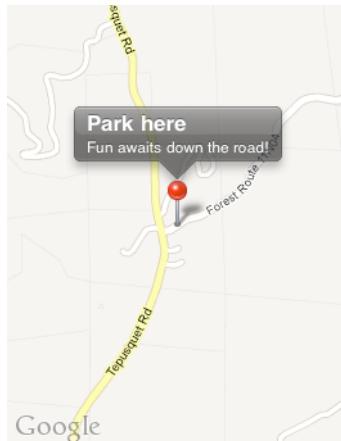


Figure 34-2. A simple annotation

That code is sufficient to produce Figure 34-2. I didn't implement any MKMapView delegate methods, so the MKAnnotationView is nil. But a nil MKAnnotationView, as I've already said, produces a red pin. I've also tapped the annotation, to display its callout, containing the annotation's title and subtitle.

This location, however, is a starting point, so by convention the pin should be green. We can create such a pin using MKPinAnnotationView, which has a `pinColor` property. We supply the annotation view in the map view's delegate (MKMapViewDelegate), by implementing `mapView:viewForAnnotation:`.

The structure of `mapView:viewForAnnotation:` is rather similar to the structure of `tableView:cellForRowAtIndexPath:` (Chapter 21), because they both do the same sort of thing. Recall that the goal of `tableView:cellForRowAtIndexPath:` is to allow the table view to reuse cells, so that at any given moment only as many cells are needed as are *visible* in the table view, regardless of how many rows the table as a whole may consist of. The same thing holds for a map and its annotation views. The map may have a huge number of annotations, but it needs to display annotation views for only those annotations that are within its current `region`. Any extra annotation views that have been scrolled out of view can thus be reused and are held for us by the map view in a cache for exactly this purpose.

So, in `mapView:viewForAnnotation:`, we start by calling `dequeueReusableAnnotationViewWithIdentifier:` to see whether there's an already existing annotation view that's not currently being displayed and that we might be able to reuse. If there isn't, we create one, attaching to it an appropriate reuse identifier.

Here's our implementation of `mapView:viewForAnnotation:`. We examine the incoming annotation to see whether the annotation view that draws it might be of a type susceptible to reuse. How we categorize views for reuse is up to us. In this case, let's say

that one category is our stock of green pins. We look to see whether this annotation is one that takes a green pin; in this case, I use the annotation's title to determine this. If so, we either create a green pin or reuse one from the map view's cache. Observe that in creating our green pin, we must explicitly set its `canShowCallout` to YES, as this is not the default:

```
- (MKAnnotationView *)mapView:(MKMapView *)mapView
    viewForAnnotation:(id <MKAnnotation>)annotation {
    MKAnnotationView* v = nil;
    if ([annotation.title isEqualToString:@"Park here"]) { ❶
        static NSString* ident = @"greenPin"; ❷
        v = [mapView dequeueReusableCellWithIdentifier:ident];
        if (v == nil) {
            v = [[[MKPinAnnotationView alloc] initWithAnnotation:annotation
                                                reuseIdentifier:ident]
                  autorelease];
            ((MKPinAnnotationView*)v).pinColor = MKPinAnnotationColorGreen;
            v.canShowCallout = YES;
        } else {
            v.annotation = annotation; ❸
        }
    }
    return v;
}
```

The structure of this implementation of `mapView:viewForAnnotation:` is typical (even though much of it seems pointlessly elaborate when we have only one annotation in our map):

- ❶ We might have more than one reusable type of annotation view. (A view can perhaps be reconfigured and thus reused, but cannot be magically converted into a view of a different type.) Here, some of our annotations might be marked with green pins, and other annotations might be marked by a different sort of annotation view altogether. So we must first somehow distinguish these cases, based on something about the incoming annotation.
- ❷ After that, for each reusable type, we proceed much as with table view cells. We have an identifier that categorizes this sort of reusable view. We try to dequeue an unused annotation view of the appropriate type, and if we can't, we create one and configure it.
- ❸ Even if we *can* dequeue an unused annotation view, and even if we have no other configuration to perform, we must associate the annotation view with the incoming annotation by assigning the annotation to this view's `annotation` property.

Now let's go further. Instead of a green pin, let's substitute our own artwork. Here, instead of creating an `MKPinAnnotationView`, I create an `MKAnnotationView` and give it a custom image. The image is too large, so I shrink the view's bounds before returning it; I also move the view up a bit, so that the bottom of the image is at the coordinates on the map ([Figure 34-3](#)):

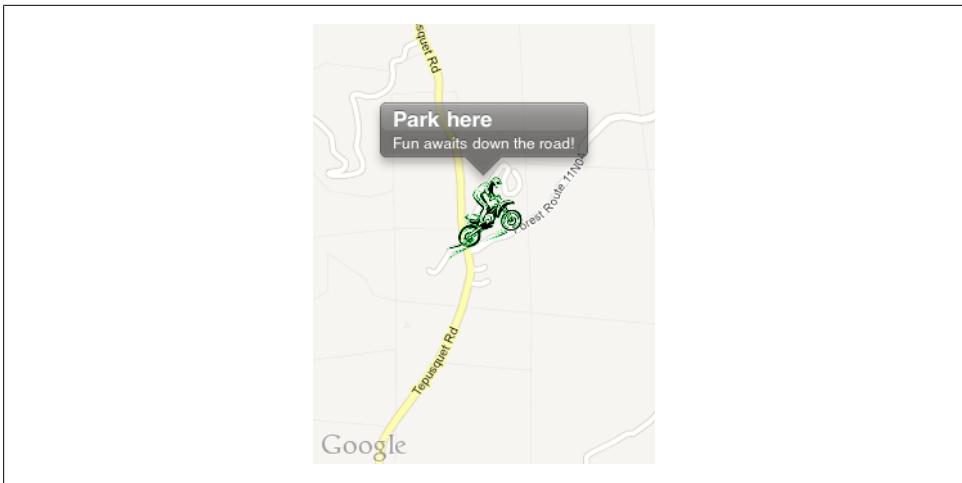


Figure 34-3. A custom annotation image

```
v = [[[MKAnnotationView alloc] initWithAnnotation:annotation
                                         reuseIdentifier:ident] autorelease];
v.image = [UIImage imageNamed:@"clipartdirtbike.gif"];
CGRect f = v.bounds;
f.size.height /= 3.0;
f.size.width /= 3.0;
v.bounds = f;
v.centerOffset = CGPointMake(0,-20);
v.canShowCallout = YES;
```

For more flexibility, we can create our own MKAnnotationView subclass and endow it with the ability to draw itself. At a minimum, such a subclass should override the initializer and assign itself a frame, and should implement `drawRect:`. Here's the implementation for a class `MyAnnotationView` that draws a dirt bike:

```
- (id)initWithAnnotation:(id <MKAnnotation>)annotation
    reuseIdentifier:(NSString *)reuseIdentifier {
    self = [super initWithAnnotation:annotation reuseIdentifier:reuseIdentifier];
    if (self) {
        UIImage* im = [UIImage imageNamed:@"clipartdirtbike.gif"];
        self.frame =
            CGRectMake(0, 0, im.size.width / 3.0 + 5, im.size.height / 3.0 + 5);
        self.centerOffset = CGPointMake(0,-20);
        self.opaque = NO;
    }
    return self;
}

- (void)drawRect: (CGRect) rect {
    UIImage* im = [UIImage imageNamed:@"clipartdirtbike.gif"];
    [im drawInRect:CGRectMakeInset(self.bounds, 5, 5)];
}
```

The corresponding implementation of `mapView:viewForAnnotation:` now has much less work to do:

```
- (MKAnnotationView *)mapView:(MKMapView *)mapView
    viewForAnnotation:(id <MKAnnotation>)annotation {
    MKAnnotationView* v = nil;
    if ([annotation.title isEqualToString:@"Park here"]) {
        static NSString* ident = @"bike";
        v = [mapView dequeueReusableAnnotationViewWithIdentifier:ident];
        if (v == nil) {
            v = [[[MyAnnotationView alloc] initWithAnnotation:annotation
                                                reuseIdentifier:ident]
                  autorelease];
            v.canShowCallout = YES;
        } else {
            v.annotation = annotation;
        }
    }
    return v;
}
```

For ultimate flexibility, we should provide our own annotation class as well. A minimal annotation class will look like this:

```
@interface MyAnnotation : NSObject <MKAnnotation> {
}
@property (nonatomic, readonly) CLLocationCoordinate2D coordinate;
@property (nonatomic, copy) NSString *title, *subtitle;
- (id)initWithLocation:(CLLocationCoordinate2D)coord;
@end

@implementation MyAnnotation
@synthesize coordinate, title, subtitle;
- (void)dealloc {
    [title release];
    [subtitle release];
    [super dealloc];
}

- (id)initWithLocation: (CLLocationCoordinate2D) coord {
    self = [super init];
    if (self) {
        self->coordinate = coord;
    }
    return self;
}
@end
```

Now when we create our annotation and add it to our map, our code looks like this:

```
MyAnnotation* ann = [[MyAnnotation alloc] initWithLocation:loc];
ann.title = @"Park here";
ann.subtitle = @"Fun awaits down the road!";
[self->map addAnnotation:ann];
[ann release];
```

A major advantage of this change appears in our implementation of `mapView:viewForAnnotation:`, where we test for the annotation type. Formerly, it wasn't easy to distinguish those annotations that needed to be drawn as a dirt bike; we were rather artificially examining the title:

```
if ([annotation.title isEqualToString:@"Park here"]) {
```

Now, however, we can just look at the class:

```
if ([annotation isKindOfClass:[MyAnnotation class]]) {
```

A further advantage of supplying our own annotation class is that gives our implementation room to grow. For example, at the moment, every `MyAnnotation` is drawn as a bike, but we could now add another property to `MyAnnotation` that tells us what drawing to use. We could also give `MyAnnotation` further properties saying such things as which way the bike should face, what angle it should be drawn at, and so on. Our implementation of `mapView:viewForAnnotation:` would read these properties and pass them on to `MyAnnotationView`, which would then draw itself appropriately.

The callout is visible in [Figure 34-2](#) and [Figure 34-3](#) because before taking the screenshot I tapped on the annotation, thus *selecting* it. `MKMapView` has methods allowing annotations to be selected or deselected programmatically, thus causing their callouts to appear or disappear; the delegate has methods notifying you when the user selects or deselects an annotation.

A callout can contain left and right accessory views; these are the `MKAnnotationView`'s `leftCalloutAccessoryView` and `rightCalloutAccessoryView`. These are `UIViews`, and should be small (less than 32 pixels in height). You can respond to taps on these views as you would any view or control.

`MKMapView` has extensive support for adding and removing annotations. Annotation views don't change size as the map is zoomed in and out, so if there are several annotations and they are brought close together by the user zooming out, the display can become crowded. The only way to prevent this is to respond to zooming by removing and adding annotations dynamically.

Overlays

An overlay differs from an annotation in being drawn *entirely* with respect to points on the surface of the earth. Thus, although an annotation's size is always the same, an overlay's size is tied to the zoom of the map view. (Overlays were introduced in iOS 4.)

Overlays are implemented much like annotations. You provide an object that adopts the `MKOverlay` protocol (which itself conforms to the `MKAnnotation` protocol) and add it to the map view. When the map view delegate method `mapView:viewForOverlay:` is called, you provide an `MKOverlayView` that actually draws the overlay. As with annotations, this architecture means that the overlay itself is a lightweight object,

and the overlay view is needed only if the part of the earth that the overlay covers is actually being displayed in the map view.

Some built-in MKShape subclasses adopt the MKOverlay protocol: MKCircle, MKPolygon, and MKPolyline. MKOverlayView has subclasses MKCircleView, MKPolygonView, and MKPolylineView, ready to draw the corresponding shapes. Thus, as with annotations, you can base your overlay entirely on the power of existing classes.

In this example, I'll use MKPolygonView to draw an overlay triangle pointing up the road from the parking place annotated in our earlier examples ([Figure 34-4](#)). We add the MKPolygon as an overlay to our map view, and derive the MKPolygonView from it in our implementation of `mapView:viewForOverlay:`. First, the MKPolygon overlay:

```
CLLocationCoordinate2D loc = self->map.region.center;
CGFloat lat = loc.latitude;
CLLocationDistance metersPerPoint = MKMetersPerMapPointAtLatitude(lat);
MKMapPoint c = MKMapPointForCoordinate(loc);
c.x += 150/metersPerPoint;
c.y -= 50/metersPerPoint;
MKMapPoint p1 = MKMapPointMake(c.x, c.y);
p1.y -= 100/metersPerPoint;
MKMapPoint p2 = MKMapPointMake(c.x, c.y);
p2.x += 100/metersPerPoint;
MKMapPoint p3 = MKMapPointMake(c.x, c.y);
p3.x += 300/metersPerPoint;
p3.y -= 400/metersPerPoint;
MKMapPoint pts[3] = {
    p1, p2, p3
};
MKPolygon* tri = [MKPolygon polygonWithPoints:pts count:3];
[self->map addOverlay:tri];
```

Second, the delegate method, where we provide the MKPolygonView:

```
- (MKOverlayView *)mapView:(MKMapView *)mapView
    viewForOverlay:(id <MKOverlay>)overlay {
    MKPolygonView* v = nil;
    if ([overlay isKindOfClass:[MKPolygon class]]) {
        v = [[[MKPolygonView alloc] initWithPolygon:(MKPolygon*)overlay]
              autorelease];
        v.fillColor = [[UIColor redColor] colorWithAlphaComponent:0.1];
        v.strokeColor = [[UIColor redColor] colorWithAlphaComponent:0.8];
        v.lineWidth = 2;
    }
    return v;
}
```

Now let's go further. The triangle in [Figure 34-4](#) is rather crude; I could draw a better arrow shape using a CGPath ([Chapter 15](#)). The built-in MKOverlayView subclass that lets me do that is MKOverlayPathView. To structure my use of MKOverlayView similarly to the preceding example, I'll supply the CGPath when I add the overlay instance to the map view. No built-in class lets me do that, so I'll use a custom class that implements the MKOverlay protocol.

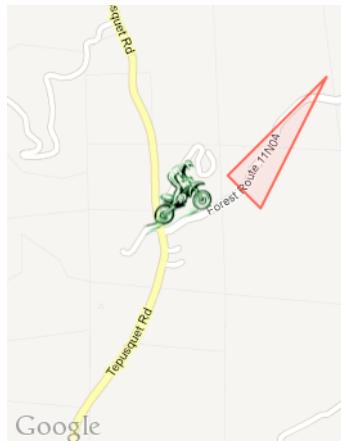


Figure 34-4. An overlay view

A minimal overlay class looks like this:

```
@interface MyOverlay : NSObject <MKOverlay> {
}
@property (nonatomic, readonly) MKMapRect boundingMapRect;
- (id) initWithRect: (MKMapRect) rect;
@end

@implementation MyOverlay
@synthesize boundingMapRect, coordinate;
- (id) initWithRect: (MKMapRect) rect {
    self = [super init];
    if (self) {
        self->boundingMapRect = rect;
    }
    return self;
}
@end
```

Our actual MyOverlay class will also have a `path` property, a UIBezierPath, that holds our CGPath and supplies it to the MKOverlayView.

Just as the `coordinate` property of an annotation tells the map view where on earth the annotation is to be drawn, the `boundingMapRect` property of an overlay tells the map view where on earth the overlay is to be drawn. Whenever any part of the `boundingMapRect` is displayed within the map view's bounds, the map view will have to concern itself with drawing the overlay. With MKPolygon, we supplied the points of the polygon in earth coordinates and the `boundingMapRect` was calculated for us. With our custom overlay class, we must supply or calculate it ourselves.

At first it may appear that there is a typological impedance mismatch: the `boundingMapRect` is an MKMapRect, whereas a CGPath is defined by CGPoints. However, it turns out that these units are interchangeable: the CGPoints of our CGPath will be translated

for us directly into MKMapPoints on the same scale — that is, the *distance* between any two CGPoints will be the distance between the two corresponding MKMapPoints. However, the *origins* are different: the CGPath must be described relative to the top left corner of the `boundingMapRect` — that is, the `boundingMapRect` is described in earth coordinates, but its top left corner is `{0,0}` as far as the CGPath is concerned. (You might think of this difference as analogous to the difference between a UIView's frame and its bounds.)

To make life simple, I'll think in meters; actually, I'll think in chunks of 75 meters, because this turned out to be a good unit for positioning and laying out the arrow. In other words, a line one `unit` long would in fact be 75 meters long if I were to arrive at this actual spot on the earth and discover the overlay literally drawn on the ground. Having derived this chunk (`unit`), I use it to lay out the `boundingMapRect`, four units on a side and positioned slightly east and north of the annotation point (because that's where the road is). Then I simply construct the arrow shape within the 4×4-unit square, rotating it so that it points in roughly the same direction as the road:

```
// start with our position and derive a nice unit for drawing
CLLocationCoordinate2D loc = self->map.region.center;
CGFloat lat = loc.latitude;
CLLocationDistance metersPerPoint = MKMetersPerMapPointAtLatitude(lat);
MKMapPoint c = MKMapPointForCoordinate(loc);
CGFloat unit = 75.0/metersPerPoint;
// size and position the overlay bounds on the earth
CGSize sz = CGSizeMake(4*unit, 4*unit);
MKMapRect mr = MKMapRectMake(c.x + 2*unit, c.y - 4.5*unit, sz.width, sz.height);
// describe the arrow as a CGPath
CGMutablePathRef p = CGPathCreateMutable();
CGPoint start = CGPointMake(0, unit*1.5);
CGPoint p1 = CGPointMake(start.x+2*unit, start.y);
CGPoint p2 = CGPointMake(p1.x, p1.y-unit);
CGPoint p3 = CGPointMake(p2.x+unit*2, p2.y+unit*1.5);
CGPoint p4 = CGPointMake(p2.x, p2.y+unit*3);
CGPoint p5 = CGPointMake(p4.x, p4.y-unit);
CGPoint p6 = CGPointMake(p5.x-2*unit, p5.y);
CGPoint points[] = {
    start, p1, p2, p3, p4, p5, p6
};
// rotate the arrow around its center
CGAffineTransform t1 = CGAffineTransformMakeTranslation(unit*2, unit*2);
CGAffineTransform t2 = CGAffineTransformRotate(t1, -M_PI/3.5);
CGAffineTransform t3 = CGAffineTransformTranslate(t2, -unit*2, -unit*2);
CGPathAddLines(p, &t3, points, 7);
CGPathCloseSubpath(p);
// create the overlay and give it the path
MyOverlay* over = [[[MyOverlay alloc] initWithRect:mr] autorelease];
over.path = [UIBezierPath bezierPathWithCGPath:p];
CGPathRelease(p);
// add the overlay to the map
[self->map addOverlay:over];
```

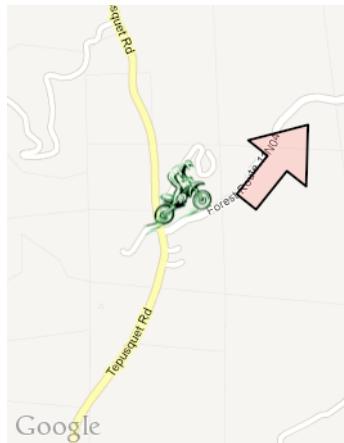


Figure 34-5. A nicer overlay view

The delegate method, where we provide the MKOverlayPathView, is simple. We pull the CGPath out of the MyOverlay instance and hand it to the MKOverlayPathView, also telling the MKOverlayPathView how to stroke and fill that path:

```
- (MKOverlayView*)mapView:(MKMapView*)mapView  
    viewForOverlay:(id <MKOverlay>)overlay {  
    MKOverlayView* v = nil;  
    if ([overlay isKindOfClass: [MyOverlay class]]) {  
        v = [[[MKOverlayPathView alloc] initWithOverlay:overlay] autorelease];  
        MKOverlayPathView* vv = (MKOverlayPathView*)v; // typecast for simplicity  
        vv.path = ((MyOverlay*)overlay).path.CGPath;  
        vv.strokeColor = [UIColor blackColor];  
        vv.fillColor = [[UIColor redColor] colorWithAlphaComponent:0.2];  
        vv.lineWidth = 2;  
    }  
    return v;  
}
```

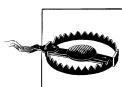
The result is a much nicer arrow (Figure 34-5), and of course this technique can be generalized to draw an overlay from any CGPath we like.

For full generality, you could define your own MKOverlayView subclass; your subclass must override and implement `drawMapRect:zoomScale:inContext:`. An example appears in the “Annotating Maps” chapter of Apple’s *Location Awareness Programming Guide*.

Overlays are maintained by the map view as an array and are drawn from back to front starting at the beginning of the array. MKMapView has extensive support for adding and removing overlays, and for managing their layering order.

CHAPTER 35

Sensors



This chapter not yet revised for iOS 5.

A device may contain hardware for sensing the world around itself—where it is located, how it is oriented, how it is moving.

Information about the device’s current location, orientation, and motion using its Wi-Fi and cellular networking, GPS, and magnetometer is provided through the Core Location framework. You’ll link to *CoreLocation.framework* and import `<CoreLocation/CoreLocation.h>`.

Information about the device’s change in speed and orientation using its accelerometer is provided through the *UIEvent* and *UIAccelerometer* classes. In iOS 4, this information is supplemented by the device’s gyroscope and is accessed through the Core Motion framework; you’ll link to *CoreMotion.framework* and import `<CoreMotion/CoreMotion.h>`.

Not all devices have all of this hardware. As of this writing, only an iPad2, an iPhone 4, or a fourth-generation iPod touch has a gyroscope. An iPod touch has no magnetometer. A device with only Wi-Fi (no cellular networking) cannot detect cell towers and also lacks a built-in GPS. And so forth.

Location

Core Location provides facilities for the device to determine and report its location (*location services*). Even a device without GPS or cellular capabilities might be able to do this, by scanning for nearby Wi-Fi devices and comparing these against an online database. Core Location will automatically use whatever facilities the device does have; all you have to do is ask for the device’s location.

Asking a map view (MKMapView, [Chapter 34](#)) to display the device’s location can be as simple as setting its `showsUserLocation` property to YES. The map automatically puts an annotation at that location, but displaying the appropriate region of the map is up to you. You can use the map delegate’s `mapView:didUpdateUserLocation:` to detect when it’s time to do this:

```
- (void)mapView:(MKMapView *)mapView
    didUpdateUserLocation:(MKUserLocation *)userLocation {
    CLLocationCoordinate2D coordinate = userLocation.location.coordinate;
    MKCoordinateRegion reg =
        MKCoordinateRegionMakeWithDistance(coordinate, 600, 600);
    mapView.region = reg;
}
```

This approach, however, is extremely coarse-grained. By setting `showsUserLocation` to YES, you are turning on location services and leaving them on, which can represent a significant power drain. You can turn location services off by setting `showsUserLocation` to NO, but then the annotation vanishes.

Instead, let’s turn on location services ourselves, just long enough to see if we can determine our position. If we can, we’ll turn location services back off, and display the location in our map with an annotation manually. We begin by ascertaining that location services are in fact available. If they are, we instantiate CLLocationManager, set ourselves as the delegate, and call `startUpdatingLocation` to turn on location services:

```
BOOL ok = [CLLocationManager locationServicesEnabled];
if (!ok) {
    NSLog(@"oh well");
    return;
}
CLLocationManager* lm = [[CLLocationManager alloc] init];
self.locman = lm;
[lm release];
self.locman.delegate = self;
[self.locman startUpdatingLocation];
```

(If we were going to track our location continually, we would also set the CLLocationManager’s `desiredAccuracy` and `distanceFilter`, but because we just want a single hit, we are skipping that step.) The delegate will eventually receive one of two messages. Something might go wrong (including that the user might refuse permission, as in [Figure 30-1](#)), in which case we’ll just turn location services back off:

```
- (void)locationManager:(CLLocationManager *)manager
    didFailWithError:(NSError *)error {
    NSLog(@"error: %@", [error localizedDescription]);
    [manager stopUpdatingLocation];
}
```

If things *don’t* go wrong, we’ll be handed our location as soon as it is determined. We’ll turn off location services and display our location on the map, along with a simple annotation:

```

- (void)locationManager:(CLLocationManager *)manager
    didUpdateToLocation:(CLLocation *)newLocation
    fromLocation:(CLLocation *)oldLocation {
    [manager stopUpdatingLocation];
    CLLocationCoordinate2D coordinate = newLocation.coordinate;
    MKCoordinateRegion reg =
        MKCoordinateRegionMakeWithDistance(coordinate, 600, 600);
    self->map.region = reg;
    MKPointAnnotation* ann = [[MKPointAnnotation alloc] init];
    ann.coordinate = coordinate;
    ann.title = @"You are here";
    [self->map addAnnotation:ann];
    [ann release];
}

```

Your app can track the device's location even when the app is not frontmost. As with sound ([Chapter 27](#)), you can set the `UIBackgroundModes` key of your app's `Info.plist`, giving it a value of `location`. This tells the system that if you have turned on location services and the user clicks the Home button, your app should not be suspended, the use of location services should still continue, and your app should still be notified if an event arrives. Background use of location services can cause a power drain, but if you want your app to function as a positional data logger, for instance, it may be the only way; you can also help conserve power by making judicious choices, such as setting a coarse `distanceFilter` value and not asking for high accuracy.

Starting in iOS 4.0, on a device with cellular capabilities, you can call `startMonitoringSignificantLocationChanges` instead of `startUpdatingLocation`. This technology (*significant location monitoring*) uses cell tower positioning and thus requires much less power than leaving the GPS turned on continuously.

Also starting in iOS 4.0, you can supply a circular `CLRegion` and turn on region monitoring (`startMonitoringForRegion:desiredAccuracy:`). When the device enters or exits the specified region, you get a notification. This technology (*region monitoring*) also uses cell tower positioning, and thus requires little power.

Both significant location monitoring and region monitoring automatically notify your app even if it is not frontmost; you do *not* have to set the `UIBackgroundModes` key. Precisely what happens depends on the state of your app when an event arrives:

Your app is suspended in the background

Your app is woken up long enough to receive the notification and do something with it.

Your app is not running at all

Your app is relaunched (remaining in the background), and your app delegate will be sent `application:didFinishLaunchingWithOptions:` with an `NSDictionary` containing `UIApplicationLaunchOptionsLocationKey`, thus allowing it to discern the special nature of the situation.

Heading and Course

For appropriately equipped devices, Core Location also supports use of the magnetometer to determine which way the device is facing (its *heading*) and the GPS to determine which way it is moving (its *course*).

In this example, I'll take advantage of the magnetometer and use the device as a compass. The `headingFilter` setting is to prevent us from being bombarded constantly with readings. For best results, the device should be held level (like a tabletop, or a compass); I have not changed the `headingOrientation` property, so the reported heading will then be the direction that the top of the device (the end away from the Home button) is pointing:

```
CLLocationManager* lm = [[CLLocationManager alloc] init];
self.locman = lm;
[lm release];
self.locman.delegate = self;
self.locman.headingFilter = 3;
[self.locman startUpdatingHeading];
```

Readings arrive as messages to the delegate. I'll simply log our magnetic heading along with a rough corresponding cardinal direction. I choose to use the magnetic heading (`magneticHeading`) rather than the true heading (`trueHeading`) because, as the documentation explains, the latter can be calculated correctly only if we are getting location updates as well as heading updates:

```
- (void)locationManager:(CLLocationManager *)manager
    didUpdateHeading:(CLHeading *)newHeading {
    CGFloat h = newHeading.magneticHeading;
    NSString* dir = @"N";
    NSArray* cards = [NSArray arrayWithObjects: @"N", @"NE", @"E", @"SE",
                      @"S", @"SW", @"W", @"NW", nil];
    for (int i = 0; i < 8; i++)
        if (h < 45.0/2.0 + 45*i) {
            dir = [cards objectAtIndex: i];
            break;
        }
    NSLog(@"%@", h, dir);
}
```

(Combining the magnetometer with the compass interface we developed in [Chapter 16](#) and [Chapter 17](#), so as to simulate a physical compass, is left as an exercise for the reader.)

GPS-based course information, if available, is returned automatically as part of a `CLLocation` object in `locationManager:didUpdateToLocation:fromLocation:`, through its `speed` and `course` properties.

Acceleration

Acceleration information can arrive in three ways:

As a prepackaged UIEvent

You can receive a UIEvent notifying you of a predefined gesture performed by accelerating the device. At present, the only such gesture is the user shaking the device.

From the shared UIAccelerometer

You can set yourself as the shared UIAccelerometer's delegate to receive acceleration notifications in `accelerometer:didAccelerate:`.

With the Core Motion framework

You instantiate CMMotionManager and then obtain information of a desired type. You can ask for accelerometer information, gyroscope information, or device motion information; device motion is a combination of accelerometer and gyroscope information that accurately describes the device's orientation in space, and is what you're most likely to want.

Shake Events

A shake event is a UIEvent ([Chapter 18](#)). Receiving shake events is rather like receiving remote events ([Chapter 27](#)), involving the notion of the first responder. To receive shake events, your app must contain a UIResponder which:

- Returns YES from `canBecomeFirstResponder`
- Is in fact first responder

This responder, or a UIResponder further up the responder chain, should implement some or all of these methods:

`motionBegan:withEvent:`

Something has started to happen that might or might not turn out to be a shake.

`motionEnded:withEvent:`

The motion reported in `motionBegan:withEvent:` is over and has turned out to be a shake.

`motionCancelled:withEvent:`

The motion reported in `motionBegan:withEvent:` wasn't a shake after all.

Thus, it might be sufficient to implement `motionEnded:withEvent:`, because this arrives if and only if the user performs a shake gesture. The first parameter will be the event subtype, but at present this is guaranteed to be `UIEventSubtypeMotionShake`, so testing it is pointless.

The view controller in charge of the current view is a good candidate to receive shake events. Thus, a minimal implementation might look like this:

```

- (BOOL) canBecomeFirstResponder {
    return YES;
}

- (void) viewDidAppear: (BOOL) animated {
    [super viewDidAppear: animated];
    [self becomeFirstResponder];
}

- (void)motionEnded:(UIEventSubtype)motion withEvent:(UIEvent *)event {
    NSLog(@"hey, you shook me!");
}

```

However, if the first responder is of a type that supports undo (such as an NSTextField), and if `motionBegan:withEvent:` arrives at the end of the responder chain, and if you have not set the shared UIApplication’s `applicationSupportsShakeToEdit` property to NO, a shake will be handled through an Undo or Redo alert. Your view controller might not want to rob any responders in its view of this capability. A simple way to prevent this is to test whether the view controller is, in fact, first responder; if it isn’t, we call `super` to pass the event on up the responder chain:

```

- (void)motionEnded:(UIEventSubtype)motion withEvent:(UIEvent *)event {
    if ([self isFirstResponder])
        NSLog(@"hey, you shook me!");
    else
        [super motionEnded:motion withEvent:event];
}

```

UIAccelerometer

To use the shared UIAccelerometer, call the class method `sharedAccelerometer` to get the global shared instance, set its `updateInterval` to prevent being swamped by notifications, and set its `delegate`. The delegate will immediately start receiving `accelerometer:didAccelerate:` calls; to turn these off, set the shared accelerometer’s delegate to nil.

The second parameter of `accelerometer:didAccelerate:` is a UIAcceleration, a simple class — nothing more than a struct, really — consisting of a timestamp and three acceleration values (UIAccelerationValue, equivalent to a double), one for each axis of the device. The positive x-axis points to the right of the device. The positive y-axis points toward the top of the device, away from the Home button. The positive z-axis points out of the screen toward the user. Acceleration values are measured in Gs. These values are only approximate — in fact, they are noisy and possibly quantized — so you have to allow room for this.

Even if the device is completely motionless, its acceleration values will constitute a vector of approximately 1 pointing toward the center of the earth, popularly known as *gravity*. The accelerometer is thus constantly reporting a combination of gravity and user-induced acceleration. This is good and bad. It’s good because it means you can

use the accelerometer to detect the device's orientation in space. It's bad because gravity values and user-induced acceleration values are mixed together.

You can attempt to separate these values, at the expense of some latency, with a low-pass filter to detect gravity only, or with a high-pass filter to eliminate the effect of gravity and detect user acceleration only, reporting a motionless device as having zero acceleration; often you'll do both (and a common technique is to run the output of the high-pass filter through an additional low-pass filter to reduce noise and small twitches). Apple provides some very nice sample code for implementing a low-pass or a high-pass filter (see especially the AccelerometerGraph example, which is also very helpful for learning how the accelerometer behaves); it's pointless to try to improve on this, so I won't bother to reproduce or discuss it.

In this example, I will simply report whether the device is lying flat on its back. To start with, I won't bother to use a filter. The two axes orthogonal to gravity, which in this position are the x and y axes, are much more accurate and sensitive to small variation than the axis pointing toward or away from gravity. So our approach is to ask first whether the x and y values are close to zero; only then do we use the z value to learn whether the device is on its back or on its face. To keep from updating our interface constantly, we implement a crude state machine; the state (an instance variable) starts out at -1, and then switches between 0 (device on its back) and 1 (device not on its back), and we update the interface only when there is a state change:

```
- (void)accelerometer:(UIAccelerometer *)accelerometer  
    didAccelerate:(UIAcceleration *)acceleration {  
    CGFloat x = acceleration.x;  
    CGFloat y = acceleration.y;  
    CGFloat z = acceleration.z;  
    CGFloat accu = 0.08;  
    if (fabs(x) < accu && fabs(y) < accu && z < -0.5) {  
        if (state == -1 || state == 1) {  
            state = 0;  
            self->label.text = @"I'm lying on my back... ahhh...";  
        }  
    } else {  
        if (state == -1 || state == 0) {  
            state = 1;  
            self->label.text = @"Hey, put me back down on the table!";  
        }  
    }  
}
```

This works, but it's sensitive to small motions of the device on the table. To damp this sensitivity, we can run our input through a low-pass filter. The low-pass filter code comes straight from Apple's own examples, and involves maintaining the previously filtered reading as a set of instance variables:

```
-(void)addAcceleration:(UIAcceleration*)accel {  
    double alpha = 0.1;  
    self->oldX = accel.x * alpha + self->oldX * (1.0 - alpha);  
    self->oldY = accel.y * alpha + self->oldY * (1.0 - alpha);
```

```

        self->oldZ = accel.z * alpha + self->oldZ * (1.0 - alpha);
    }

- (void)accelerometer:(UIAccelerometer *)accelerometer
    didAccelerate:(UIAcceleration *)acceleration {
    [self addAcceleration: acceleration];
    CGFloat x = self->oldX;
    CGFloat y = self->oldY;
    CGFloat z = self->oldZ;
    CGFloat accu = 0.08;
    if (fabs(x) < accu && fabs(y) < accu && z < -0.5) {
        // ... and the rest is as before ...
    }
}

```

In this next example, the user is allowed to slap the side of the device against an open hand as a way of telling it to go the next or previous image or whatever it is we're displaying. We pass the acceleration input through a high-pass filter to eliminate gravity (again, the filter code comes straight from Apple's examples). What we're looking for is a high positive or negative *x* value. A single slap is likely to consist of several consecutive readings above our threshold, but we want to report each slap only once, so we maintain the timestamp of our previous high reading as an instance variable and ignore readings that are too close to one another in time:

```

-(void)addAcceleration:(UIAcceleration*)accel {
    double alpha = 0.1;
    self->oldX = accel.x - ((accel.x * alpha) + (self->oldX * (1.0 - alpha)));
    self->oldY = accel.y - ((accel.y * alpha) + (self->oldY * (1.0 - alpha)));
    self->oldZ = accel.z - ((accel.z * alpha) + (self->oldZ * (1.0 - alpha)));
}

- (void)accelerometer:(UIAccelerometer *)accelerometer
    didAccelerate:(UIAcceleration *)acceleration {
    [self addAcceleration: acceleration];
    CGFloat x = self->oldX;
    // CGFloat y = self->oldY;
    // CGFloat z = self->oldZ;
    CGFloat thresh = 1.0;
    if (acceleration.timestamp - self->oldTime < 0.5)
        return;
    if (x < -thresh) {
        NSLog(@"left");
        self->oldTime = acceleration.timestamp;
    }
    if (x > thresh) {
        NSLog(@"right");
        self->oldTime = acceleration.timestamp;
    }
}

```

This works, but there's a problem. A sudden jerk involves both an acceleration (as the user starts the device moving) and a deceleration (as the device stops moving). Thus a left slap might be preceded by a high value in the opposite direction, which we will

interpret wrongly as a right slap. We can compensate crudely, at the expense of some latency, with delayed performance:

```
CGFloat thresh = 1.0;
if (x < -thresh) {
    if (acceleration.timestamp - self->oldTime > 0.5 || self->lastSlap == 1) {
        self->oldTime = acceleration.timestamp;
        self->lastSlap = -1;
        [NSObject cancelPreviousPerformRequestsWithTarget:self];
        [self performSelector:@selector(report:) withObject:@"left" afterDelay:0.5];
    }
}
if (x > thresh) {
    if (acceleration.timestamp - self->oldTime > 0.5 || self->lastSlap == -1) {
        self->oldTime = acceleration.timestamp;
        self->lastSlap = 1;
        [NSObject cancelPreviousPerformRequestsWithTarget:self];
        [self performSelector:@selector(report:) withObject:@"right" afterDelay:0.5];
    }
}
```

A more sophisticated analysis might involve storing a stream of all the most recent UIAcceleration objects and studying the entire stream to work out the overall trend.

Core Motion

Core Motion, introduced in iOS 4.0, takes advantage of the device's gyroscope if there is one. This means that instead of using the accelerometer alone and suffering the disadvantages and inaccuracies of filtering, you can have the device do the math, interpreting the accelerometer and gyroscope together for you and distinguishing gravity from user acceleration with negligible latency. To do so, you'll elect to receive CMDeviceMotion instances, consisting of the following properties:

gravity

A vector with value 1 pointing to the center of the earth.

userAcceleration

A vector describing user-induced acceleration, with no gravity component.

rotationRate

A vector describing how the device is rotating around its center. This sort of motion was often difficult or impossible to detect before the addition of the gyroscope.

attitude

A description of the device's instantaneous orientation in space. It is described with respect to an initial frame of reference in which the negative z-axis points at the center of the earth, but the x-axis and y-axis, though orthogonal to the other axes, could be pointing anywhere. Thus, it is of interest primarily as a change relative to an earlier attitude. To convert the current attitude to a description of this change,

you store an earlier attitude, and then send the `multiplyByInverseOfAttitude:` message to the current attitude, passing the stored attitude as the argument.

Further detail about Core Motion is beyond the scope of this book. The important thing to understand is that the presence of the gyroscope, combined with the `CMDeviceMotion` class, compensates for a number of shortcomings of reading the accelerometer alone. (The device may use the gyroscope in other ways as well; for example, if you ask Core Location for updates with a `desiredAccuracy` of `kCLLocationAccuracyBestForNavigation`, it uses the gyroscope to supplement the calculation of location changes.)

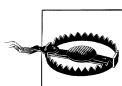
PART VII

Final Topics

This part of the book is a miscellany of topics that didn't fit easily into any of the preceding chapters.

- [Chapter 36](#) is about files. It explains how your app can store data on disk to be retrieved the next time the app runs (including both standalone files and user defaults). It also discusses sharing files with the user through iTunes and with other apps and concludes with a survey of how iOS can work with some common file formats (XML, SQLite, and image files).
- [Chapter 37](#) introduces networking, with an emphasis on HTTP downloading of data, and giving a nod to other aspects of networks (such as Bonjour and push notifications) that you can explore independently if your app requires them.
- [Chapter 38](#) is about threads. Making your app multithreaded (beyond the automatic threading support provided by the built-in interface widgets and their supporting frameworks) can introduce great complexity and is not a beginner topic, but you still might need to understand the basic concepts of multithreading, either in order to prevent a lengthy task from blocking user interaction with your app, or because some framework explicitly relies on it. Special attention is paid to the advantages of NSOperation and (especially) Grand Central Dispatch.
- [Chapter 39](#) describes how iOS supports Undo in your app.
- You are now a proud graduate of this book's school of iOS programming fundamentals. You are fully prepared to proceed independently. [Chapter 40](#) lists additional frameworks and facilities that were found to be beyond the scope of this book. Your mission, should you decide to accept it, is to explore these if and when you need them. iOS is huge; you'll never stop learning and experimenting. Good hunting!

Persistent Storage



This chapter not yet revised for iOS 5.

The device on which your app runs contains flash memory that functions as the equivalent of a hard disk, holding files that survive the device's being powered down (*persistent storage*). Apps can store files to, and retrieve them from, this virtual hard disk. Apps can also define document types in which they specialize and can hand such documents to one another.

The Sandbox

The hard disk as a whole is not open to your app's view. A limited portion of the hard disk is dedicated to your app alone: this is your app's *sandbox*. The idea is that every app, seeing only its own sandbox, is hindered from spying or impinging on the files belonging to other apps. Your app can also see some higher-level directories owned by the system as a whole, but cannot write to them.

You can create directories (folders) within the sandbox. In addition, the sandbox contains some standard directories. For example, suppose you want a reference to the Documents directory. Here's how to access it:

```
NSString* docs = [NSSearchPathForDirectoriesInDomains(
    NSDocumentDirectory, NSUserDomainMask, YES) lastObject];
```

That code returns a path string for the Documents directory. Starting in iOS 4, the preferred way to refer to a file or directory is with a URL. You obtain this from an `NSFileManager` instance:

```
NSFileManager* fm = [[NSFileManager alloc] init];
// ...
NSError* err = nil;
NSURL* docsurl = [fm URLForDirectory:NSDocumentDirectory
```

```

    inDomain:NSUserDomainMask appropriateForURL:nil create:YES error:&err];
// error-checking omitted
// ...
[fm release];

```

A question that will immediately occur to you is: where should I put secondary files and folders that I want to save now and read later? Once upon a time, the Documents directory would have been a good place. But in iOS 3.2, file sharing was introduced, meaning that your app can be configured so that the user can see and modify your app's Documents directory through iTunes. So you might not want to put things there that the user isn't supposed to see and change.

Settling upon an alternative location for your app's files is up to you. Personally, I favor the Application Support directory. On a Mac, this directory is shared by multiple applications, but on iOS each app has its own private Application Support directory in its own sandbox, so you can safely put files directly into it. This directory may not exist initially, so you can obtain it and create it at the same time:

```

NSURL* suppurl = [fm URLForDirectory:NSApplicationSupportDirectory
    inDomain:NSUserDomainMask appropriateForURL:nil create:YES error:&err];

```

After that, if you need a file path reference (an NSString), just ask for [suppurl path].

Basic File Operations

Let's say we intend to create folder *MyFolder* inside the Documents directory. Starting with the path string *docs* pointing at the Documents directory (as obtained in the previous section), we can generate a reference to *MyFolder*, using one of the many NSString methods specifically aimed at manipulating path strings. (There is no way to create a folder specified by a URL.) We can then use an NSFileManager instance (*fm*) to learn whether our target folder exists, and to create it if it doesn't:

```

NSString* myfolder = [docs stringByAppendingPathComponent:@"MyFolder"];
BOOL exists = [fm fileExistsAtPath:myfolder];
if (!exists) {
    NSError* err = nil;
    [fm createDirectoryAtPath:myfolder withIntermediateDirectories:NO
        attributes:nil error:&err];
    // error-checking omitted
}

```

To learn what files and folders exist within a directory, you can ask for an array of the directory's contents:

```

NSError* err = nil;
NSArray* arr = [fm contentsOfDirectoryAtPath:docs error:&err];
// error-checking omitted
/*
MyFolder
*/

```

That array is shallow, showing only the directory's immediate contents. For a deep array, ask for the directory's subpaths:

```
NSError* err = nil;
NSArray* arr = [fm subpathsOfDirectoryAtPath:docs error:&err];
// error-checking omitted
/*
MyFolder
MyFolder/moi.txt
*/
```

A deep array might be very big. If you're looking for something in particular, you might prefer to enumerate the directory, so that you are handed only one file reference at a time:

```
NSDirectoryEnumerator* dir = [fm enumeratorAtPath:docs];
for (NSString* file in dir) {
    // do something with each string
}
```

A directory enumerator also permits you to decline to dive into a particular subdirectory, so you can make your traversal even more efficient.

Note that in all those cases you are handed a relative pathname. To make it usable as an actual file reference, you can either append it to the original directory pathname or make the original directory the current directory:

```
NSDirectoryEnumerator* dir = [fm enumeratorAtPath:docs];
[fm changeCurrentDirectoryPath: docs];
for (NSString* f in dir)
    if ([[f pathExtension] isEqualToString: @"txt"])
        // f is valid for referring to this file
```

In that example, `f` is a valid pathname even though it's a relative (partial) pathname, because its containing folder is the current directory. Similarly, if you convert `f` to a URL (using `fileURLWithPath:`), the resulting NSURL object has an appropriate `baseURL` and so is a valid reference.

Consult the NSFileManager class documentation for more about what you can do with files, and see also Apple's *Low-Level File Management Programming Topics*.

Saving and Reading Files

To save or read a file, you are most likely to use one of the convenience methods for the class appropriate to the file's contents. NSString, NSData, NSArray, and NSDictionary provide `writeToFile...` and `initWithContentsOfFile...` methods (as well as `writeToURL...` and `initWithContentsOfURL...`). Recall that NSArray and NSDictionary files are actually property lists ([Chapter 10](#)) and work only if all the contents of the array or dictionary are property list types (NSString, NSData, NSDate, NSNumber, NSArray, and NSDictionary).

If an object's class adopts the NSCoding protocol, you can convert it to an NSData and back again using NSKeyedArchiver and NSKeyedUnarchiver. An NSData can be saved as a file or in a property list. Thus, NSCoding provides a way to save an object to disk. An example of doing this with a UIColor object appears in [Chapter 10](#).

You can make your own class adopt the NSCoding protocol. This can become somewhat complicated because an object can refer (through an instance variable) to another object, which may also adopt the NSCoding protocol, and thus you can end up saving an entire *graph* of interconnected objects if you wish. However, I'll confine myself to illustrating a simple case (and you can read the *Archives and Serializations Programming Guide* for more information).

Let's say, then, that we have a simple Person class with a `firstName` property and a `lastName` property. We'll declare that it adopts the NSCoding protocol:

```
@interface Person : NSObject <NSCoding> {
```

To make this class actually conform to NSCoding, we must implement `encodeWithCoder:` (to archive the object) and `initWithCoder:` (to unarchive the object). In `encodeWithCoder:`, we must first call `super` if the superclass adopts NSCoding, and then call the appropriate `encode...` method for each instance variable we want preserved:

```
- (void)encodeWithCoder:(NSCoder *)encoder {
    // [super encodeWithCoder: encoder]; // not in this case
    [encoder encodeObject:self->lastName forKey:@"last"];
    [encoder encodeObject:self->firstName forKey:@"first"];
}
```

In `initWithCoder`, we must call `super`, using either `initWithCoder:` if the superclass adopts the NSCoding protocol or the designated initializer if not, and then call the appropriate `decode...` method for each instance variable stored earlier, finally returning `self`; memory management is up to us:

```
- (id) initWithCoder:(NSCoder *)decoder {
    // self = [super initWithCoder: decoder]; // not in this case
    self = [super init];
    self->lastName = [decoder decodeObjectForKey:@"last"];
    [self->lastName retain];
    self->firstName = [decoder decodeObjectForKey:@"first"];
    [self->firstName retain];
    return self;
}
```

Now we'll test this by creating, configuring, and saving a Person instance as a file:

```
Person* moi = [[Person alloc] init];
moi.firstName = @"Matt";
moi.lastName = @"Neuburg";
NSData* moidata = [NSKeyedArchiver archivedDataWithRootObject:moi];
NSString* moofile = [myfolder stringByAppendingPathComponent:@"moi.txt"];
[moidata writeToFile:moofile atomically:NO];
[moi release];
```

Now we should be able to retrieve the saved Person at a later time:

```
NSData* persondata = [[NSData alloc] initWithContentsOfFile:moifile];
Person* person = [NSKeyedUnarchiver unarchiveObjectWithData:persondata];
[persondata release];
NSLog(@"%@", person.firstName, person.lastName); // Matt Neuburg
```

If the NSData object is itself the entire content of the file, as here, then instead of using `archivedDataWithRootObject:` and `unarchiveObjectWithData:`, you can skip the intermediate NSData object altogether and use `archiveRootObject:toFile:` and `unarchiveObjectWithFile:`.

Saving a single Person as an archive may seem like overkill; why didn't we just make a text file consisting of the first and last names? But imagine that a Person has a lot more properties, or that we have an array of hundreds of Persons, or an array of hundreds of dictionaries where one value in each dictionary is a Person; now all of a sudden the power of an archivable Person becomes clear. Even though Person now adopts the NSCoding protocol, an NSArray containing a Person object still cannot be written to disk using NSArray's `writeToFile...` or `writeToURL...`, because Person is still not a property list type. But the array can be archived and written to disk with NSKeyedArchiver.

User Defaults

User defaults, which have been referred to often already in this book (see especially [Chapter 10](#) and [Chapter 13](#)), are intended as the persistent storage of the user's preferences, as well as for maintaining state when your app quits so that you can restore the situation the next time the app launches. They are little more, really, than a special case of an NSDictionary property list file. You talk to the `NSUserDefaults` `standardUserDefaults` object much as if it were a dictionary; it has keys and values. And the only legal values are property list values (see the preceding section). Thus, to store a Person in user defaults, you'd have to archive it first to an NSData object. Unlike NSDictionary, `NSUserDefaults` provides convenience methods for converting between a simple data type such as a float or a BOOL and the object that is stored in the defaults (`setFloat:forKey:`, `floatForKey:`, and so forth). But the defaults themselves are still a dictionary.

Meanwhile, somewhere on disk, this dictionary is being saved for you automatically as a property list file — though you don't concern yourself with that. You simply set or retrieve values from the dictionary by way of their keys, secure in the knowledge that the file is being read into memory or written to disk as needed. Your chief concern is to make sure that you've written everything needful into user defaults before your app terminates; as we saw in [Chapter 11](#), in a multitasking world this will usually mean when the app delegate receives `applicationDidEnterBackground:` at the latest. If you're worried that your app might crash, you can tell the `standardUserDefaults` object to `synchronize`, but this is rarely necessary.

To provide the value for a key before the user has had a chance to do so — the default default, as it were — use `registerDefaults:`. What you're supplying here is a dictionary

whose key–value pairs will each be written into the defaults, but only if there is no such key already. Recall this example from [Chapter 10](#):

```
[[NSUserDefaults standardUserDefaults] registerDefaults:  
    [NSDictionary dictionaryWithObjectsAndKeys:  
        [NSNumber numberWithInt: 4],  
        @"cardMatrixRows",  
        [NSNumber numberWithInt: 3],  
        @"cardMatrixColumns",  
        nil]];
```

The idea is that we call `registerDefaults:` extremely early as the app launches. Either the app has run at some time previously and the user has set these preferences, in which case this call has no effect and does no harm, or not, in which case we now have initial values for these preferences with which to get started. So, in the game app from which that code comes, we start out with a 4×3 game layout, but the user can change this at any time.

This leaves only the question of how the user is to interact with the defaults. One way is that your app provides some kind of interface. For example, the game app from which the previous code comes has a tabbed interface; the second tab is where the user sets preferences ([Figure 20-1](#)). In the TidBITS News app, there’s a single button for setting the size of text, and that’s the only preference with which the user ever interacts directly.

(Both apps also store state information in the user defaults, but without the user’s knowledge or direct participation, and not with keys that the user has any way of accessing. For example, the game app records the state of the game board and the card deck into user defaults every time these change, so that if the app is terminated we can restore the game, the next time the app is launched, as it was when the user left off.)

Alternatively, you can provide a *settings bundle*, consisting mostly of one or more property list files describing an interface and the corresponding user default keys and their initial values; the Settings app is then responsible for translating your instructions into an actual interface, and for presenting it to the user.

Using a settings bundle has some obvious disadvantages: the user may not think to look in the Settings app, for example; the user has to leave your app to access preferences; and you don’t get the kind of control over the interface that you have within your own app. Also, in a multitasking world, this means that the user can set preferences while your app is backgrounded; you’ll need to register for `NSUserDefaultsDidChangeNotification` in order to hear about this.

In some situations, though, a settings bundle has some clear advantages. Keeping the preferences interface out of your app can make your app’s own interface cleaner and simpler. You don’t have to write any of the “glue” code that coordinates the preferences interface with the user default values. And it can be nice for the user to be able to set preferences for your app even when your app isn’t running.

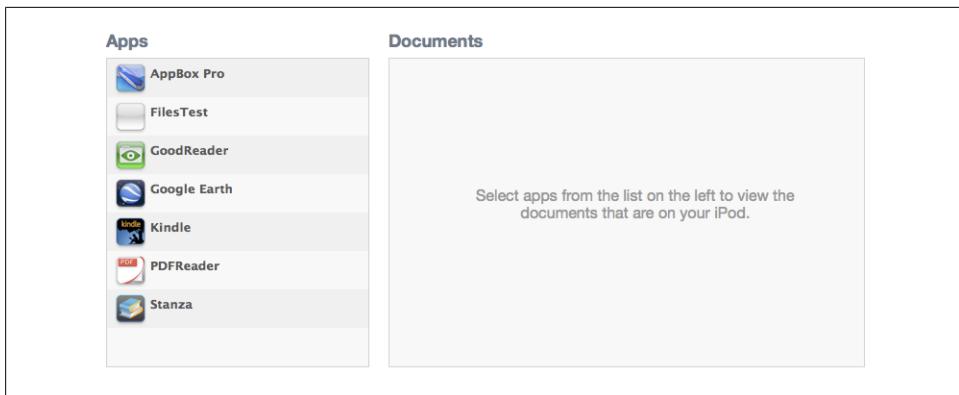


Figure 36-1. The iTunes file sharing interface

Writing a settings bundle is described in the “Implementing Application Preferences” chapter of Apple’s *iOS Application Programming Guide*, along with the *Settings Application Schema Reference*.

File Sharing

If your app supports file sharing, its Documents directory becomes available to the user through iTunes (Figure 36-1). The user can add files to your app’s Documents directory, and can save files and folders from your app’s Documents directory to the computer, as well as renaming and deleting files and folders. This could be appropriate, for example, if the purpose of your app is to display some common file type that the user might obtain elsewhere, such as PDFs or JPEGs.

To support file sharing, set the *Info.plist* key “Application supports iTunes file sharing” (`UIFileSharingEnabled`).

Once your entire Documents directory is exposed to the user this way, you are suddenly not so likely to use the Documents directory to store private files. As I mentioned earlier, I like to use the Application Support directory instead.

Document Types

Your app can declare itself willing to open documents of a certain type. In this way, if another app obtains a document of this type, it can propose to hand the document off to your app. For example, the user might download the document with Mobile Safari, or receive it in a mail message with the Mail app; now we need a way to get it from Safari or Mail to you.

To let the system know that your app is a candidate for opening a certain kind of document, you will configure the `CFBundleDocumentTypes` key in your *Info.plist*. This is an array, where each entry will be a dictionary specifying a document type by using

keys such as `LSItemContentTypes`, `CFBundleTypeName`, `CFBundleTypeIconFiles`, and `LSHandlerRank`.

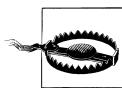
For example, suppose I want to declare that my app opens PDFs. My `Info.plist` could contain this simple entry (as seen in the standard editor):

```
Document types           (1 item)
  Item 0                (1 item)
    Document Content Type UTIs (1 item)
      Item 0              com.adobe.pdf
```

In Xcode 4, you can also specify document types by editing the target; switch to the Info tab. This same `Info.plist` entry would appear here as an untitled document type with `com.adobe.pdf` in the Types field.

Now suppose the user receives a PDF in an email message. The Mail app can display this PDF, but the user can also tap the Action button to bring up an action sheet containing two Open In buttons. The first button might actually specify my app as the default, but even if it doesn't, tapping the second button will bring up a second action sheet where my app appears as a button. (The interface will look like [Figure 36-2](#), except that my app will be listed as one of the buttons.)

But now suppose the user actually *taps* the button that hands the PDF off to my app. For this to work, my app delegate must implement `application:handleOpenURL:`. At this point, my app has been brought to the front, either by launching it from scratch or by reviving it from background suspension; its job is now to handle the opening of the document whose URL has arrived as the second parameter. To prevent me from peeking into another app's sandbox, the system has already copied the document into my sandbox, into the Inbox directory, which is created for exactly this purpose.



Unfortunately, the Inbox directory is currently created in your Documents folder. Thus, if your app implements file sharing, the user can see the Inbox folder; you may wish to delete the Inbox folder, therefore, as soon as you're done retrieving files from it.

In this simple example, my app has just one view controller, which has an outlet to a `UIWebView` where we will display any PDFs that arrive in this fashion. So my app delegate contains this code:

```
- (BOOL)application:(UIApplication *)application handleOpenURL:(NSURL *)url {
  [viewController displayPDF:url];
  return YES;
}
```

And my view controller contains this code:

```
- (void) displayPDF: (NSURL*) url {
  NSURLRequest* req = [NSURLRequest requestWithURL:url];
  [self->wv loadRequest:req];
}
```

In real life, things might be more complicated. Our implementation of `application:handleOpenURL:` might check to see whether this really *is* a PDF, and return NO if it isn't. Also, our app might be in the middle of something else, possibly displaying a completely different view controller's view; because `application:handleOpenURL:` can arrive at any time, we may have to be prepared to drop whatever we were doing and showing previously and display the incoming document instead.

If our app is launched from scratch by the arrival of this URL, `application:didFinishLaunchingWithOptions:` will be sent to our app delegate as usual. The options dictionary (the second parameter) will contain the `UIApplicationLaunchOptionsURLKey`, and we can take into account, if we like, the fact that we are being launched specifically to open a document. The usual thing, however, is to ignore this key and launch in the normal way; `application:handleOpenURL:` will then arrive in good order after our interface has been set up, and we can handle it just as we would if we had already been running.

Starting in iOS 4.2, your app delegate can implement `application:openURL:sourceApplication:annotation:` in order to receive more information about the incoming URL. If implemented, this will be called in preference to `application:handleOpenURL:`, and it won't be called at all on a device running an earlier system, so there is no penalty for implementing both methods.

Handing Off a Document

The converse of the situation discussed in the previous section is this: your app has somehow acquired a document that it wants to hand off to whatever app can deal with it. This is done through the `UIDocumentInteractionController` class. This class operates asynchronously, so retaining an instance of it is up to you; typically, you'll store it in an instance variable with a retain setter policy.

For example, let's say our app has a PDF sitting in its Documents directory. Assuming we have an `NSURL` pointing to this document, presenting the interface for handing the document off to some other application ([Figure 36-2](#)) could be as simple as this (`sender` is a button that the user has just tapped):

```
self.dic = [UIDocumentInteractionController interactionControllerWithURL:url];
BOOL y =
[dic presentOptionsMenuFromRect:[sender bounds] inView:sender animated:YES];
```

There are actually two action sheets available. The first action sheet in [Figure 36-2](#), the Options action sheet, is summoned by `presentOptionsMenu...`; the second action sheet in [Figure 36-2](#), the Open In action sheet, is summoned by `presentOpenInMenu...`, but it can also be summoned by one of the buttons in the first action sheet. These methods are cleverly designed to work on both iPhone and iPad interfaces; on the iPad, the buttons appear in a popover.

Your app can't learn *which* other applications are capable of accepting the document! Indeed, it can't even learn in advance whether *any* other applications are capable of

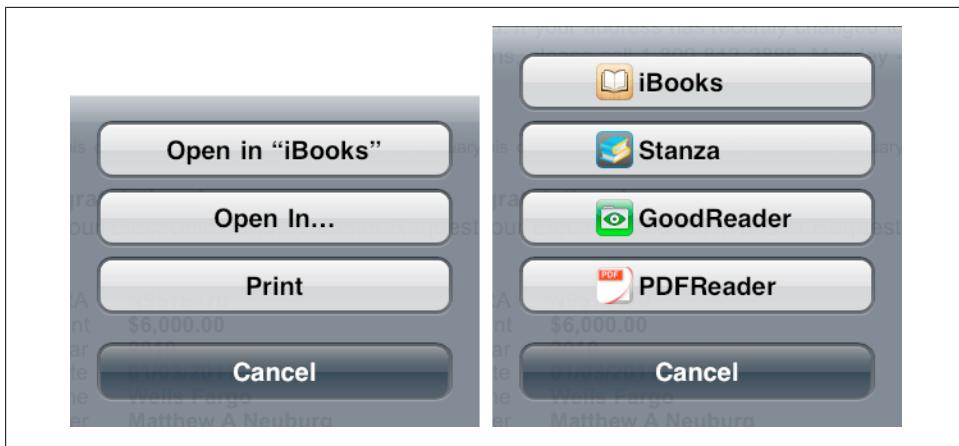


Figure 36-2. The document Options action sheet and Open In action sheet

accepting the document; your only clue is that the returned BOOL value afterward will be NO if UIDocumentInteractionController couldn't present the interface you requested.

UIDocumentInteractionController can, however, be interrogated for *some* information about the document type. In this example, we place a button into our interface whose image is the icon of the document type (the idea, perhaps, is that the user would then tap this button to do something with the document):

```
self.dic = [UIDocumentInteractionController interactionControllerWithURL:url];
UIImage* icon = [[self.dic icons] lastObject];
UIButton* b = [UIButton buttonWithType:UIButtonTypeRoundedRect];
[b setImage:icon forState:UIControlStateNormal];
[b sizeToFit]; // ... and probably also set frame origin here ...
[self.view addSubview: b];
```

A UIDocumentInteractionController can also present a preview of the document, if the document is of a type for which preview is enabled. You must give the UIDocumentInteractionController a delegate, and the delegate must implement `documentInteractionControllerViewControllerForPreview:`, returning an existing view controller that will contain the preview's view controller. So, here we ask for the preview:

```
self.dic = [UIDocumentInteractionController interactionControllerWithURL:url];
self.dic.delegate = self;
[self.dic presentPreviewAnimated:YES];
```

Here we supply the view controller:

```
- (UIViewController *) documentInteractionControllerViewControllerForPreview:
    (UIDocumentInteractionController *) controller {
    return self;
}
```

If the view controller returned were a UINavigationController, the preview's view controller would be pushed onto it. In this case it isn't, so the preview's view controller is

a presented view controller. The preview interface also contains an Action button that lets the user summon the Options action sheet. In fact, this preview interface is exactly the same interface already familiar from the Mail app.

Delegate methods allow you to track what's happening in the interface presented by the UIDocumentInteractionController. Probably the most important of these are the ones that inform you that key stages of the interaction are ending:

- `documentInteractionControllerDidDismissOptionsMenu:`
- `documentInteractionControllerDidDismissOpenInMenu:`
- `documentInteractionControllerDidEndPreview:`
- `documentInteractionController:didEndSendingToApplication:`

Previews are actually provided through the Quick Look framework, and you can skip the UIDocumentInteractionController altogether and present the preview yourself through a QLPreviewController (link to *QuickLook.framework* and import `<QuickLook/QuickLook.h>`). It's a view controller, so to display the preview you show it as a presented view controller or push it onto a navigation controller's stack (just as UIDocumentInteractionController would have done). A nice feature of QLPreviewController is that you can give it more than one document to preview; the user can move between these, within the preview, using arrow buttons that appear at the bottom of the interface. Plus, if a document can be opened in another app, the interface includes the action button that summons UIDocumentInteractionController's Options or Open In action sheet.

In this example, I have in my Documents directory several PDF documents. I acquire a list of their URLs and present a preview for them:

```
// obtain URLs of PDFs as an array
NSString* docsdir = [NSSearchPathForDirectoriesInDomains(
    NSDocumentDirectory, NSUserDomainMask, YES) lastObject];
NSFileManager* fm = [[NSFileManager alloc] init];
NSDirectoryEnumerator* direnum = [fm enumeratorAtPath:docsdir];
[fm changeCurrentDirectoryPath: docsdir];
NSMutableArray* marr = [NSMutableArray array];
for (NSString* file in direnum) {
    [direnum skipDescendants];
    if ([[file pathExtension] isEqualToString: @"pdf"]) {
        NSURL* url = [NSURL fileURLWithPathComponents:
            [NSArray arrayWithObjects: docsdir, file, nil]];
        [marr addObject: url];
    }
}
self.pdfs = marr; // retain policy
[fm release];
// show preview interface
QLPreviewController* preview = [[QLPreviewController alloc] init];
preview.dataSource = self;
[self presentModalViewController:preview animated:YES];
[preview release];
```

You'll notice that I haven't told the QLPreviewController what documents to preview. That is the job of QLPreviewController's data source. I am the data source! I simply fetch the requested information from the list of URLs:

```
- (NSInteger) numberOfPreviewItemsInPreviewController:  
    (QLPreviewController *) controller {  
    return [self.pdfs count];  
}  
  
- (id <QLPreviewItem>) previewController: (QLPreviewController *) controller  
    previewItemAtIndex: (NSInteger) index {  
    return [self.pdfs objectAtIndex:index];  
}
```

The second data source method requires me to return an object that adopts the QLPreviewItem protocol. By a wildly improbable coincidence, NSURL *does* adopt this protocol, so the example works.

XML

XML is a highly flexible and widely used general-purpose text file format for storage and retrieval of structured data. You might use it yourself to store data that you'll need to retrieve later, or you could encounter it when obtaining information from elsewhere, such as the Internet.

Mac OS X Cocoa provides a set of classes (NSXMLDocument and so forth) for reading, parsing, maintaining, searching, and modifying XML data in a completely general way, but iOS does *not* include these. I think the reason must be that their tree-based approach is too memory-intensive. Instead, iOS provides NSXMLParser, a much simpler class that walks through an XML document, sending delegate messages as it encounters elements. With this, you can parse an XML document once, but what you do with the pieces as they arrive is up to you. The general assumption here is that you know in advance the structure of the particular XML data you intend to read and that you have provided classes for storage of the same data in object form and for transforming the XML pieces into that storage.

To illustrate, let's return to our earlier example of a Person class with a `firstName` and a `lastName` property. Imagine that as our app starts up, we would like to populate it with Person objects, and that we've stored the data describing these objects as an XML file in our app bundle, like this:

```
<?xml version="1.0" encoding="utf-8"?>  
<people>  
    <person>  
        <firstName>Matt</firstName>  
        <lastName>Neuburg</lastName>  
    </person>  
    <person>  
        <firstName>Snidely</firstName>  
        <lastName>Whiplash</lastName>
```

```
</person>
<person>
    <firstName>Dudley</firstName>
    <lastName>Doright</lastName>
</person>
</people>
```

This data could be mapped to an array of Person objects, each with its `firstName` and `lastName` properties appropriately set. (This is a deliberately easy example, of course; not all XML is so easily or obviously expressed as objects.) Let's consider how we might do that.

Using NSXMLParser is not difficult in theory. You create the NSXMLParser, handing it the URL of a local XML file (or an NSData, perhaps downloaded from the Internet), set its delegate, and tell it to `parse`. The delegate starts receiving delegate messages. For simple XML like ours, there are only three delegate messages of interest:

parser:didStartElement:namespaceURI:qualifiedName:attributes:

The parser has encountered an opening element tag. In our document, this would be `<people>`, `<person>`, `<firstName>`, or `<lastName>`.

parser:didEndElement:namespaceURI:qualifiedName:

The parser has encountered the corresponding closing element tag. In our document this would be `</people>`, `</person>`, `</firstName>`, or `</lastName>`.

parser:foundCharacters:

The parser has encountered some text between the starting and closing tags for the current element. In our document this would be, for example, "Matt" or "Neuburg" and so on.

In practice, responding to these delegate messages poses challenges of maintaining state. If there is just one delegate, it will have to bear in mind at every moment what element it is currently encountering; this could make for a lot of instance variables and a lot of if-statements in the implementation of the delegate methods. To aggravate the issue, `parser:foundCharacters:` can arrive multiple times for a single stretch of text; that is, the text may arrive in pieces, so we have to accumulate it into an instance variable, which is yet another case of maintaining state.

An elegant way to meet these challenges is by resetting the NSXMLParser's delegate to different objects at different stages of the parsing process. We make each delegate responsible for parsing one element; when a child of that element is encountered, we make a new object and make *it* the delegate. The child element delegate is then responsible for making us, the parent, the delegate once again when it finishes parsing its own element. This is slightly counterintuitive because it means `parser:didStartElement...` and `parser:didEndElement...` for the same element are arriving at two different objects. Imagine, for example, what the job of our `<people>` parser will be:

- When `parser:didStartElement...` arrives, the `<people>` parser looks to see if this is a `<person>`. If so, it creates an object that knows how to deal with a `<person>`,

handing that object a reference to itself (the `<people>` parser), and makes it the delegate.

- Delegate messages now arrive at this newly created `<person>` parser. If any text is encountered, `parser:foundCharacters:` will be called, and the text must be accumulated into an instance variable.
- Eventually, `parser:didEndElement...` arrives. The `<person>` parser now uses its reference to make the `<people>` parser the delegate once again. Thus, the `<people>` parser is in charge once again, ready if another `<person>` element is encountered (and the old `<person>` parser might now go quietly out of existence).

With this in mind, we can design a simple all-purpose base class for parsing an element (simple especially because we are taking no account of namespaces, attributes, and other complications):

```
@interface MyXMLParserDelegate : NSObject <NSXMLParserDelegate> {
}
@property (nonatomic, copy) NSString* name;
@property (nonatomic, retain) NSMutableString* text;
@property (nonatomic, assign) MyXMLParserDelegate* parent;
@property (nonatomic, retain) MyXMLParserDelegate* child;
- (void) start: (NSString*) elementName parent: (id) parent;
- (void) makeChild: (Class) class
    elementName: (NSString*) elementName
    parser: (NSXMLParser*) parser;
- (void) finishedChild: (NSString*) s;

@end
```

Here's how these properties and methods are intended to work:

name

The name of the element we are parsing now.

text

A place for any characters to accumulate as we parse our element.

parent

The `MyXMLParserDelegate` who created us and whose child we are.

child

If we encounter a child element, we'll create a `MyXMLParserDelegate` and retain it here, making it the delegate.

start:parent:

When we create a child parser, we'll call this method on the child so that it knows who its parent is. The first parameter is the name of the element the child will be parsing; we know this because we, not the child, received `parser:didStartElement....` (In a fuller implementation, this method would be more elaborate and we'd hand the child *all* the information we got with `parser:didStartElement....`)

makeChild:elementName:parser:

If we encounter a child element, there's a standard dance to do: instantiate some subclass of MyXMLParserDelegate, make it our `child`, make it the parser's delegate, and send it `start:parent:.` This is a utility method that embodies that dance.

finishedChild:

When a child receives `parser:didEndElement...,` it sends this message to its parent before making its parent the delegate. The parameter is the `text`, but the parent can use this signal to obtain any information it expects from the child before the child goes out of existence.

Now we can sketch in the default implementation for MyXMLParserDelegate:

```
- (void) start: (NSString*) el parent: (id) p {
    self.name = el;
    self.parent = p;
    self.text = [NSMutableString string];
}

- (void) makeChild: (Class) class
    NSString* elementName
    NSXMLParser* parser {
    MyXMLParserDelegate* del = [[class alloc] init];
    self.child = del;
    parser.delegate = del;
    [del start: elementName parent: self];
    [del release];
}

- (void) finishedChild: (NSString*) s { // subclass implements as desired
}

- (void)parser:(NSXMLParser *)parser foundCharacters:(NSString *)string {
    [self.text appendString:string];
}

- (void)parser:(NSXMLParser *)parser didEndElement:(NSString *)elementName
    NSString* namespaceURI: NSString* namespaceURI
    NSString* qualifiedName: NSString* qName {
    if (parent) {
        [parent finishedChild: [[self.text copy] autorelease]];
        parser.delegate = self.parent;
    }
}

- (void) dealloc {
    [text release];
    [child release];
    [name release];
    [super dealloc];
}
```

We can now create subclasses of MyXMLParserDelegate: one for each kind of element we expect to parse. The chief responsibility of such a subclass, if it encounters a child

element in `parser:didStartElement...`, is to create an instance of the appropriate `MyXMLParserDelegate` subclass, send it `start:parent:`, and make it the delegate; we have already embodied this in the utility method `makeChild:elementName:parser:`. The reverse process is already built into the default implementation of `parser:didEndElement...`: we call the parent's `finishedChild:` and make the parent the delegate.

We can now parse our sample XML into an array of Person objects very easily. We start by obtaining the URL of the XML file, handing it to an `NSXMLParser`, creating our first delegate parser and making it the delegate, and telling the `NSXMLParser` to start:

```
NSURL* url = [[NSBundle mainBundle] URLForResource:@"folks" withExtension:@"xml"];
NSXMLParser* parser = [[NSXMLParser alloc] initWithContentsOfURL:url];
MyPeopleParser* people = [[MyPeopleParser alloc] init];
[parser setDelegate: people];
[parser parse];
// ... do something with people.people ...
[people release];
[parser release];
```

Here is `MyPeopleParser`. It is the top-level parser so it has some extra work to do: when it encounters the `<people>` element, which is the first thing that should happen, it creates the `people` array that will hold the Person objects; this array will be the final result of the entire parsing operation. If it encounters a `<person>` element, it does the standard dance I described earlier, creating a `<person>` parser (`MyPersonParser`) as its child and making it the delegate; when the `<person>` parser calls back to tell us it's finished, `MyPeopleParser` expects the `<person>` parser to supply a Person through its `person` property:

```
- (void)parser:(NSXMLParser *)parser didStartElement:(NSString *)elementName
          namespaceURI:(NSString *)namespaceURI
          qualifiedName:(NSString *)qualifiedName
          attributes:(NSDictionary *)attributeDict
{
    if ([elementName isEqualToString: @"people"])
        self.people = [NSMutableArray array];
    if ([elementName isEqualToString: @"person"])
        [self makeChild:[MyPersonParser class] elementName:elementName
                    parser:parser];
}

- (void) finishedChild: (NSString*) s {
    [people addObject: [(MyPersonParser*)self.child person]];
}
```

`MyPersonParser` does the same child-making dance when it encounters a `<firstName>` or a `<lastName>` element; it uses a plain vanilla `MyXMLParserDelegate` to parse these children, because the built-in ability to accumulate text and hand it back is all that's needed. In `finishedChild:`, it makes sure it has a Person object ready to hand back to its parent through its `person` property; key-value coding is elegantly used to match the name of the element with the name of the `Person` property to be set:

```

- (void)parser:(NSXMLParser *)parser didStartElement:(NSString *)elementName
    namespaceURI:(NSString *)namespaceURI
    qualifiedName:(NSString *)qualifiedName
    attributes:(NSDictionary *)attributeDict {
    [self makeChild:[MyXMLParserDelegate class] elementName:elementName
        parser:parser];
}

- (void) finishedChild:(NSString *)s {
    if (!self.person) {
        Person* p = [[Person alloc] init];
        self.person = p; // retain policy
        [p release];
    }
    [self.person setValue: s forKey: self.child.name];
}

- (void) dealloc {
    [person release];
    [super dealloc];
}

```

This may seem like a lot of work to parse such a simple bit of XML, but it is neatly object-oriented and requires very little new code once we've established the MyXMLParserDelegate superclass, which is of course reusable in many other situations.

On the other hand, if you really want tree-based XML parsing along with XPath and so forth, you can have it, because the `libxml2` library is present in the SDK (and on the device). This is a C *dylib* (short for “dynamic library,” extension `.dylib`), and Xcode doesn't automatically know during the build process where to find its headers (even though it's part of the SDK), so the instructions for accessing it in your project are a tiny bit more involved than linking to an Objective-C framework:

1. In Xcode, add `libxml2.dylib` to the Link Binary With Libraries build phase for your target, just as you would do with a framework.
2. Now comes the extra step that differs from using a framework; it is needed because, although the Xcode build process automatically looks inside `iphoneos/usr/include/` for headers, it doesn't automatically recurse down into folders, so it won't look inside the `libxml2` folder unless you tell it to. Edit the target's build settings and set the Header Search Paths build setting to `$SDKROOT/usr/include/libxml2`. When you close the dialog for adding a search path, this will transform itself into `iphoneos/usr/include/libxml2`.
3. In your code, import `<libxml/tree.h>`.

You now have to talk to `libxml2` using C. This is no trivial task. Here's an example proving we can do it; we read our XML file, parse it into a tree, and traverse all its elements:

```

NSURL* url = [[NSBundle mainBundle] URLForResource:@"folks" withExtension:@"xml"];
NSString* path = [url absoluteString];
const char* filename = [path UTF8String];
 xmlDocPtr doc = NULL;
 xmlNode *root_element = NULL;
 doc = xmlReadFile(filename, NULL, 0);
 root_element = xmlDocGetRootElement(doc);
 traverse_elements(root_element); // must be previously defined
 xmlFreeDoc(doc);
 xmlCleanupParser();

```

Here's our definition for `traverse_elements`; it logs each person and the person's first and last name, just to prove we are traversing successfully:

```

void traverse_elements(xmlNode * a_node) {
    xmlNode *cur_node = NULL;
    for (cur_node = a_node; cur_node; cur_node = cur_node->next) {
        if (cur_node->type == XML_ELEMENT_NODE) {
            if (strcmp(cur_node->name, "person") == 0)
                NSLog(@"found a person");
            if (strcmp(cur_node->name, "firstName") == 0)
                NSLog(@"First name: %s", cur_node->children->content);
            if (strcmp(cur_node->name, "lastName") == 0)
                NSLog(@"Last name: %s", cur_node->children->content);
        }
        traverse_elements(cur_node->children);
    }
}

```

If talking C to `libxml2` is too daunting, you can interpose an Objective-C front end by taking advantage of a third-party library. See, for example, <https://github.com/TouchCode/TouchXML>.

Keep in mind, however, that you're really not supposed to do what I just did. Even if you use `libxml2`, you're supposed to use stream-based parsing, not tree-based parsing. See Apple's `XMLPerformance` example code.

SQLite

SQLite (<http://www.sqlite.org/docs.html>) is a lightweight, full-featured relational database that you can talk to using SQL, the universal language of databases. This can be an appropriate storage format when your data comes in rows and columns (records and fields) and needs to be rapidly searchable.

In the same way as you can link to `libxml2.dylib`, you can link to `libsqLite3.dylib` (and import `<sqlite3.h>`) to access the power of SQLite. As with `libxml2`, talking C to `sqlite3` may prove annoying. There are a number of lightweight Objective-C front ends. In this example, I use `fmdb` (<https://github.com/ccgus/fmdb>) to read the names of people out of a previously created database:

```

NSString* docsdir = [NSSearchPathForDirectoriesInDomains(
    NSDocumentDirectory, NSUserDomainMask, YES) lastObject];
NSString* dbpath = [docsdir stringByAppendingPathComponent:@"people.db"];
FMDatabase* db = [FMDatabase databaseWithPath:dbpath];
if (![db open]) {
    NSLog(@"Ooops");
    return;
}
FMResultSet *rs = [db executeQuery:@"select * from people"];
while ([rs next]) {
    NSLog(@"%@", [
        [rs stringForColumn:@"firstname"],
        [rs stringForColumn:@"lastname"]]);
}
[db close];
/* output:
Matt Neuburg
Snidely Whiplash
Dudley Doright
*/

```

You can include a previously constructed SQLite file in your app bundle, but you can't write to it there; the solution is to copy it from your app bundle into another location, such as the Documents directory, before you start working with it.

The Core Data framework also uses SQLite as a storage format (or, alternatively, it can use XML). Core Data is a generalized way of dealing with objects and properties; it is appropriate particularly when these form a complex relational graph. For example, a person might have not only multiple addresses but also multiple friends who are also persons; expressing persons and addresses as explicit object types, working out how to link them and how to translate between objects in memory and data in storage, and tracking the effects of changes, such as when a person is deleted from the data, can be tedious. Core Data can help, but it is *not* a beginner-level technology, nor should it be seen as a substitute for a true relational database. Core Data is beyond the scope of this book; entire books can be written about Core Data alone (and have been). See the *Core Data Programming Guide* and the other resources referred to there.

Image File Formats

Starting in iOS 4, the Image I/O framework provides a simple, unified way to open image files (from disk or downloaded from the network, as described in [Chapter 37](#)), to save image files, to convert between image file formats, and to read metadata from standard image file formats, including EXIF and GPS information from a digital camera. You'll need to link to *ImageIO.framework* and import `<ImageIO/ImageIO.h>`.

Obviously, such features were not entirely missing before iOS 4. `UIImage` can read the data from most standard image formats, and you can convert formats with functions such as `UIImageJPEGRepresentation` and `UIImagePNGRepresentation`. But you could not, for example, save an image as TIFF without the Image I/O framework.

The Image I/O framework introduces the notion of an *image source* (`CGImageSourceRef`). This can be created from the URL of a file on disk or from `NSData` (actually `CFDataRef`, to which `NSData` is toll-free bridged) obtained or generated in some way. You can use this to obtain a `CGImage` of the source's image (or, if the source format contains multiple images, a particular image). But you can also obtain metadata from the source *without* transforming the source into a `CGImage`, thus conserving memory. For example:

```
NSURL* url = [[NSBundle mainBundle] URLForResource:@"colson" withExtension:@"jpg"];
CGImageSourceRef src = CGImageSourceCreateWithURL((CFURLRef)url, NULL);
NSDictionary* result = (id)CGImageSourceCopyPropertiesAtIndex(src, 0, NULL);
// ... do something with result ...
CFRelease(src); [result release];
```

Without having opened the image file as an image, we now have a dictionary full of information about it, including its pixel dimensions (`kCGImagePropertyPixelWidth`, `kCGImagePropertyPixelHeight`), its resolution, its color model, its color depth, and its orientation — plus, because this picture originally comes from a digital camera, the EXIF data such as the aperture and exposure at which it was taken, plus the make and model of the camera.

We can obtain the image as a `CGImage`, with `CGImageSourceCreateImageAtIndex`. Alternatively, we can request a thumbnail version of the image. I'm afraid that Apple's documentation fails to impress sufficiently on the reader the value of the thumbnail. If your purpose in opening this image is to display it in your interface, you don't care about the original image data; a thumbnail is precisely what you want, especially because you can specify any size for this "thumbnail" all the way up to the original size of the image! This is tremendously convenient, because to assign a small `UIImageView` a large image wastes all the memory reflected by the size difference.

To generate a thumbnail at a given size, you start with a dictionary specifying the size along with other instructions, and pass that, together with the image source, to `CGImageSourceCreateThumbnailAtIndex`. The only pitfall is that, because we are working with a `CGImage` and specifying actual pixels, we must remember to take account of the scale of our device's screen. So, for example, let's say we want to scale our image so that its largest dimension is no more than 100 points:

```
NSURL* url = [[NSBundle mainBundle] URLForResource:@"colson" withExtension:@"jpg"];
CGImageSourceRef src = CGImageSourceCreateWithURL((CFURLRef)url, NULL);
CGFloat scale = [UIScreen mainScreen].scale;
NSDictionary* d =
    [NSDictionary dictionaryWithObjectsAndKeys:
        (id)kCFBooleanTrue, kCGImageSourceShouldAllowFloat,
        (id)kCFBooleanTrue, kCGImageSourceCreateThumbnailWithTransform,
        (id)kCFBooleanTrue, kCGImageSourceCreateThumbnailFromImageAlways,
        [NSNumber numberWithInt:100*scale], kCGImageSourceThumbnailMaxPixelSize,
        nil];
CGImageRef imref = CGImageSourceCreateThumbnailAtIndex(src, 0, (CFDictionaryRef)d);
```

```

UIImage* im =
    [UIImage imageWithCGImage:imref scale:scale orientation:UIImageOrientationUp];
self->iv.image = im; // assign image to UIImageView
CFRelease(imref); CFRelease(src);

```

The Image I/O framework also introduces the notion of an *image destination*, used for saving an image into a specified file format. As a final example, I'll show how to save our image as a TIFF, which, as I mentioned before, was impossible previously. Notice that in this case we never even need to open the image as an image: we save directly from the image source to the image destination:

```

NSURL* url = [[NSBundle mainBundle] URLForResource:@"colson" withExtension:@"jpg"];
CGImageSourceRef src = CGImageSourceCreateWithURL((CFURLRef)url, NULL);
NSFileManager* fm = [[NSFileManager alloc] init];
NSURL* suppurl = [fm URLForDirectory:NSApplicationSupportDirectory
                           inDomain:NSUserDomainMask
                           appropriateForURL:nil
                           create:YES error:NULL];
NSURL* tiff = [suppurl URLByAppendingPathComponent:@"mytiff.tiff"];
CGImageDestinationRef dest =
    CGImageDestinationCreateWithURL((CFURLRef)tiff,
    (CFStringRef)@"public.tiff", 1, NULL);
CGImageDestinationAddImageFromSource(dest, src, 0, NULL);
bool ok = CGImageDestinationFinalize(dest);
// error-checking omitted
[fm release]; CFRelease(src); CFRelease(dest);

```


Basic Networking



This chapter not yet revised for iOS 5.

Networking is difficult and complicated, chiefly because it's ultimately out of your control. My favorite phrase with regard to the network is, "There's many a slip 'twixt the cup and the lip." You can ask for a resource from across the network, but at that point anything can happen: the resource might not be found (the server is down, perhaps), it might take a while to arrive, it might never arrive, the network itself might vanish after the resource has partially arrived. iOS, however, makes at least the *basics* of networking very easy, so that's all that this chapter will deal with.

Many earlier chapters have dealt with interface and frameworks that network for you automatically. Put a UIWebView in your interface ([Chapter 24](#)) and poof, you're networking; the UIWebView does all the grunt work, and it does it a lot better than you'd be likely to do it from scratch. The same is true of MPMovieViewController ([Chapter 28](#)), MFMailComposeViewController ([Chapter 33](#)), and MKMapView ([Chapter 37](#)).

HTTP Requests

A simple HTTP request is made through an NSURLConnection object. You hand it an NSURLRequest describing what you'd like to do, along with a delegate; the download begins automatically (unless you specify otherwise). Then you stand back and let delegate messages arrive. The actual network operations happen asynchronously (unless you specifically demand that they happen synchronously, which you'd never do); the NSURLConnection object does all its work in the background and sends you delegate messages when something occurs that you need to know about.

Data received from the network in response to your request will arrive as an NSData object. It will arrive piecemeal, so you have to maintain state; in particular, you'll have

an NSMutableData object to which you'll keep appending each new bit of NSData until you're told that the entire data has arrived — or that the request has failed. (The whole process is somewhat reminiscent of what we did with an NSXMLParser in [Chapter 36](#).)

All the real work happens in four delegate methods:

connection:didReceiveResponse:

The server is responding. We can now hope that our data will start to arrive, so get ready. If you like, you can interrogate the NSURLResponse object that is handed to you, to learn things from the response headers such as the data's expected size and MIME type.

connection:didReceiveData:

Some data has arrived. Append it to the NSMutableData object.

connectionDidFinishLoading:

All of the data has arrived; the NSMutableData object presumably contains it. Clean up as needed.

connection:didFailWithError:

Something went wrong. Clean up as needed.

Here's an example of initiating a download of a JPEG image file:

```
self.receivedData = [NSMutableData data];
NSString* s = @"http://www.someserver.com/somefolder/someimage.jpg";
NSURL* url = [NSURL URLWithString:s];
NSURLRequest* req = [NSURLRequest requestWithURL:url];
NSURLConnection* conn = [NSURLConnection connectionWithRequest:req delegate:self];
[conn retain];
```

Here are the corresponding delegate method implementations:

```
- (void) connection:(NSURLConnection *)connection
    didReceiveResponse:(NSURLResponse *)response {
    // connection is starting, clear buffer
    [receivedData setLength:0];
}

- (void) connection:(NSURLConnection *)connection didReceiveData:(NSData *)data {
    // data is arriving, add it to the buffer
    [receivedData appendData:data];
}

- (void) connection:(NSURLConnection *)connection didFailWithError:(NSError *)error {
    // something went wrong, release connection
    [connection release];
    // clean up interface as needed
}

- (void) connectionDidFinishLoading:(NSURLConnection *)connection {
```

```

    // all done, release connection, we are ready to rock and roll
    [connection release];
    // do something with receivedData
}

```

You may have noticed that, in creating the `NSURLConnection` initially, we retain it without releasing it. The release takes place in our delegate methods. This style of memory management is copied from the “Using `NSURLConnection`” chapter of Apple’s *URL Loading System Programming Guide*. The `NSURLConnection` instance is allowed to float in space, as it were, retained (so that it won’t vanish in a puff of smoke) but with no reference to it, until the final delegate method is called.

An alternative approach would be to keep the `NSURLConnection` in an instance variable and balance the retain and release normally. In that case we would probably wrap the entire connection process in a dedicated object to hold this instance variable, because otherwise keeping track of multiple simultaneous `NSURLConnections` would be a nightmare. Here’s the complete implementation for such a wrapper object, MyDownloader:

```

- (id) initWithRequest: (NSURLRequest*) req {
    self = [super init];
    if (self) {
        self->request = [req copy];
        self->connection = [[NSURLConnection alloc]
            initWithRequest:req delegate:self startImmediately:NO];
        self->receivedData = [[NSMutableData alloc] init];
    }
    return self;
}

- (void) dealloc {
    [request release];
    [connection release];
    [receivedData release];
    [super dealloc];
}

- (void) connection:(NSURLConnection *)connection
    didReceiveResponse:(NSURLResponse *)response {
    [receivedData setLength:0];
}

- (void) connection:(NSURLConnection *)connection didReceiveData:(NSData *)data {
    [receivedData appendData:data];
}

- (void)connection:(NSURLConnection *)connection didFailWithError:(NSError *)err {
    [[NSNotificationCenter defaultCenter]
        postNotificationName:@"connectionFinished" object:self
        userInfo:[NSDictionary dictionaryWithObject:err forKey:@"error"]];
}

```

```

- (void)connectionDidFinishLoading:(NSURLConnection *)connection {
    [[NSNotificationCenter defaultCenter]
        postNotificationName:@"connectionFinished" object:self];
}

```

In the line that creates the `NSURLConnection`, we have added the `startImmediately:` parameter, with a value of NO. Thus, a `MyDownloader` object can exist before doing any actual downloading. To set the download into motion, we tell `MyDownloader`'s `connection` to `start`. (Sending `start` to an `NSURLConnection` that is already downloading has no effect.) In the past, there have been complaints that sending `start` to an `NSURLConnection` that does not start immediately can cause a crash. I have not seen this myself, so perhaps it has been fixed in more recent iOS versions, but the solution is to schedule the connection on a run loop explicitly just before starting it:

```

[connection scheduleInRunLoop:[NSRunLoop currentRunLoop]
    forMode:NSDefaultRunLoopMode];
[connection start];

```

How would we use `MyDownloader` if we have several objects to download? We might, for example, keep a mutable array of `MyDownloader` objects. To initiate a download, we create a `MyDownloader` object, register for its @"`connectionFinished`" notification, stuff it into the array, and set its connection going:

```

if (!self.connections)
    self.connections = [NSMutableArray array];
NSString* s = @"http://www.someserver.com/somefolder/someimage.jpg";
NSURL* url = [NSURL URLWithString:s];
NSURLRequest* req = [NSURLRequest requestWithURL:url];
MyDownloader* d = [[MyDownloader alloc] initWithRequest:req];
[self.connections addObject:d];
[[NSNotificationCenter defaultCenter] addObserver:self
    selector:@selector(finished:) name:@"connectionFinished" object:d];
[d.connection start];
[d release];

```

When the notification arrives, either we've failed with an error or we've finished in good order. In the latter case, we grab the received data and retain it; either way, we remove the `MyDownloader` from the array, thus releasing it, its connection, and its data:

```

- (void) finished: (NSNotification*) n {
    MyDownloader* d = [n object];
    NSData* data = nil;
    if ([n userInfo]) {
        // ... error of some kind! ...
    } else {
        data = [d receivedData];
        [data retain]; // about to go out of existence otherwise
        // ... and do something with the data right now ...
    }
    [self.connections removeObject:d];
}

```

In real life, you might also query the `MyDownloader`'s `request`, or perhaps some other instance variable in a `MyDownloader` subclass, to identify what material this data rep-

resents. Also, you're not so likely to use a temporary array of downloaders; rather, you'll incorporate downloaders directly into your application's model, letting them fetch the data on demand.

Suppose, for example, you need to download images to serve as thumbnails in the cells of a UITableView. Let's consider how these images can be supplied lazily on demand. The model, as we saw in [Chapter 21](#), might be an array of dictionaries. In this case, the dictionary might contain some text and a downloader whose job is to supply the image. So what I'm proposing is a model like this:

```
array
  dictionary
    text: @"Manny"
    pic: Downloader whose job is to supply an image of Manny
  dictionary
    text: @"Moe"
    pic: Downloader whose job is to supply an image of Moe
  dictionary
    text: @"Jack"
    pic: Downloader whose job is to supply an image of Jack
  ....
```

When the table turns to the data source for data, the data source will turn to the dictionary for the requested row and ask the downloader for its image. At that point, either the downloader has an image, in which case it supplies it, or it hasn't, in which case it returns nil (or some placeholder) and begins the download.

Here's the key point. When a downloader succeeds in downloading its image, it notifies the data source. If the corresponding row is visible, the data source immediately tells the table to reload the corresponding row; the table asks the data source for the data, the data source turns to the dictionary for the requested row, and this time it obtains the image! Moreover, once an image is downloaded, the downloader continues to hold on to it and to supply it on request, so as the user scrolls, previously downloaded images just appear as part of the table.

The downloader we're imagining here is a MyDownloader subclass, MyImageDownloader, with an `image` property so that the data source can request the image. MyImageDownloader's implementation is straightforward:

```
- (UIImage*) image {
  if (image)
    return image;
  [self.connection start];
  return nil; // or a placeholder
}

- (void)connectionDidFinishLoading:(NSURLConnection *)connection {
  UIImage* im = [UIImage imageWithData:self->receivedData];
  if (im) {
    self.image = im;
}
```

```

        [[NSNotificationCenter defaultCenter]
            postNotificationName:@"imageDownloaded" object:self];
    }
}

```

The data source looks perfectly normal:

```

- (UITableViewCell *)tableView:(UITableView *)tableView
    cellForRowAtIndexPath:(NSIndexPath *)indexPath {
    static NSString *CellIdentifier = @"Cell";
    UITableViewCell *cell =
        [tableView dequeueReusableCellWithIdentifier:CellIdentifier];
    if (cell == nil) {
        cell = [[[UITableViewCell alloc] initWithStyle:UITableViewCellStyleDefault
                                         reuseIdentifier:CellIdentifier] autorelease];
    NSDictionary* d = [self.model objectAtIndex: indexPath.row];
    cell.textLabel.text = [d objectForKey:@"text"];
    MyImageDownloader* imd = [d objectForKey:@"pic"];
    cell.imageView.image = imd.image;
    return cell;
}

```

Now for the key point. The data source is also registered for an @"imageDownloaded" notification. When such a notification arrives, it works out the table row corresponding to the MyImageDownloader that posted the notification and reloads that row:

```

- (void) imageDownloaded: (NSNotification*) n {
    MyImageDownloader* d = [n object];
    NSUInteger row = [self.model indexOfObjectPassingTest:
        ^BOOL(id obj, NSUInteger idx, BOOL *stop) {
            return ((NSDictionary*)obj objectForKey:@"pic"] == d);
    }];
    if (row == NSNotFound) return; // shouldn't happen
    NSIndexPath* ip = [NSIndexPath indexPathForRow:row inSection:0];
    NSArray* ips = [self.tableView indexPathsForVisibleRows];
    if ([ips indexOfObject:ip] != NSNotFound) {
        [self.tableView reloadRowsAtIndexPaths:[NSArray arrayWithObject: ip]
                                         withRowAnimation:UITableViewRowAnimationFade];
    }
}

```

This works, and demonstrates the basic technique for doing one of the most-desired network-related operations, namely obtaining lazily, on demand, from the network, a piece of data and updating the interface accordingly. What's missing from the example is robustness with regard to failure. The trouble is that once an NSURLConnection has failed, it's dead; you can't use the same NSURLConnection to try again later. We can rectify that in MyImageDownloader by replacing the NSURLConnection on failure:

```

- (void)connection:(NSURLConnection *)connection didFailWithError:(NSError *)err {
    self.connection = [[[NSURLConnection alloc] initWithRequest:self.request
                                         delegate:self startImmediately:NO] autorelease];
}

```

This partially solves the problem: when the user scrolls a failed cell out of view and later scrolls it back into view, the table will ask the data source for its data and the

MyImageDownloader will try again to download its image. But that won't happen for a failed cell that's never scrolled out of view. How you deal with this is up to you; it's a matter of providing the best user experience without having an undue impact upon performance, battery, and so forth. In this instance, because these images are fairly unimportant, I might arrange that when an NSTimer with a fairly large interval fires (every 60 seconds, say), we reload the visible rows; this will cause any failed MyImageDownloader whose corresponding row is visible to try again.

In planning your interface, it is useful to draw a distinction as to whether the user will experience a certain bit of networking explicitly or implicitly. This changes nothing about *how* you network; it's a matter of presentation. The earlier example of downloading images to be slotted into the cells of an existing table view would be implicit networking: it happens regardless of whether the user wants it, and it doesn't seriously affect overall functionality, even if some or all of the images fail to arrive. In the TidBITS News app, on the other hand, everything displayed comes from a downloaded RSS feed: no feed, no data. The app saves the previously downloaded feed (in user defaults, see [Chapter 36](#)), so the user has something to read even in the absence of the network, but the feed is explicitly refreshed at launch or if the user taps a button (along with the spinning network activity indicator, [Chapter 25](#)), and if the download fails, we put up an alert.

Bonjour

Bonjour is the ingenious technology, originated at Apple and now becoming a universal standard, for allowing network devices to advertise services they provide and to discover dynamically other devices offering such services. Once an appropriate service is detected, a client device can resolve it to get a network address and can then begin communicating with the server device. Actually communicating is outside the scope of this book, but device discovery via Bonjour is easy.

In this example, we'll look to see whether any device, such as a Mac, is running iTunes with library sharing turned on. We can search for domains or for a particular service; here, we'll pass the empty string as the domain to signify "any domain," and concentrate on the service, which is @"_daap._tcp". We maintain two instance variables, the NSNetServiceBrowser that will look for devices, and a mutable array in which to store any services it discovers:

```
self.services = [NSMutableArray array];
NSNetServiceBrowser* browser = [[NSNetServiceBrowser alloc] init];
self.nsb = browser;
[browser release];
self.nsb.delegate = self;
[self.nsb searchForServicesOfType:@"_daap._tcp" inDomain:@""];
```

The NSNetServiceBrowser is now searching for devices advertising iTunes sharing and will keep doing so until we destroy it or tell it to stop. It is common to leave the service browser running, because devices can come and go very readily. As they do, the service

browser's delegate will be informed. For purposes of this example, I'll simply maintain a list of services, and update the app's interface when the situation changes:

```
- (void)netServiceBrowser:(NSNetServiceBrowser *)netServiceBrowser
    didFindService:(NSNetService *)netService
    moreComing:(BOOL)moreServicesComing {
    [self.services addObject:netService];
    if (!moreServicesComing)
        [self updateInterface];
}

- (void)netServiceBrowser:(NSNetServiceBrowser *)netServiceBrowser
    didRemoveService:(NSNetService *)netService
    moreComing:(BOOL)moreServicesComing {
    [self.services removeObject:netService];
    if (!moreServicesComing)
        [self updateInterface];
}
```

The delegate messages very kindly tell me when they have completed the task of informing me of a series of changes, so I can wait to update the interface until after a full batch of changes has ended. In this example, I don't really have any interface to update; I'll just log the list of services, each of which is an NSNetService instance:

```
- (void) updateInterface {
    for (NSNetService* service in self.services) {
        if (service.port == -1) {
            NSLog(@"service %@ of type %@", service.name, service.type);
        }
    }
}
```

To connect to a service, we would first need to *resolve* it, thus obtaining an address and other useful information. An unresolved service has port **-1**, as shown in the previous code. To resolve a service, you tell it to resolve; you will probably also set a delegate on the service, so as to be notified when the resolution succeeds (or fails). Here, I'll have the delegate call my `updateInterface` method again if a resolution succeeds, and I'll extend `updateInterface` to show the port number for any resolved services:

```
- (void) updateInterface {
    for (NSNetService* service in self.services) {
        if (service.port == -1) {
            NSLog(@"service %@ of type %@", service.name, service.type);
            [service setDelegate:self];
            [service resolveWithTimeout:10];
        } else {
            NSLog(@"service %@ of type %@", service.name, service.type, service.port, service.addresses);
        }
    }
}
```

```
- (void)netServiceDidResolveAddress:(NSNetService *)sender {
    [self updateInterface];
}
```

The `addresses` of a resolved service constitute an array of `NSData`. Logging an address like this is largely pointless, as it is not human-readable, but it's useful for handing to a `CFSocket`. In general you'll call the service's `getInputStream:outputStream:` to start talking over the connection; that's outside the scope of this discussion. See Apple's WiTap example for more.

Push Notifications

If your app uses a server on the network that's under your control, you can arrange for the user to be notified when a significant event takes place on the server. This is called a *push notification* (or *remote notification*). The user interface for a push notification is the same as for a local notification ([Chapter 26](#)): the user is shown an alert in front of whatever is happening at that moment and can use this alert to launch your app.

For example, the TidBITS News app is about news stories on the TidBITS website. The app's data comes from an RSS feed, which is refreshed on the server side whenever something changes on the site, such as a new news story being posted. It might be appropriate (and cool) if we were to add push notifications to the server code that refreshes the RSS feed, so that users could be alerted to the fact that they might like to launch TidBITS News and read a newly posted story.

Implementing push notifications is not trivial, and requires cooperation across the network between your app and your server, and between your server and Apple's push notification server. I've never actually tried this, so I'm just describing what the architecture is like; for details, read Apple's *Local and Push Notification Programming Guide*.

When developing your app, you obtain from the iOS Provisioning Portal ([Chapter 9](#)) credentials identifying your app, and allowing communication between your server and Apple's push notification server, and between Apple's push notification server and your app running on the user's device. When your app launches, it calls the `UIApplication` method `registerForRemoteNotificationTypes:`, which communicates asynchronously with Apple's push notification server to obtain a token identifying this instance of your app. If successful, the token comes back in the app delegate method `application:didRegisterForRemoteNotificationsWithDeviceToken:`. At that point, your app must communicate with your server to provide it with this token.

The server is now maintaining two pieces of information: its credentials and a list of tokens effectively representing users. When an event occurs at your server for which the server wishes to push a notification out to users, the server uses its credentials to connect with Apple's push notification server and — for *every individual user* whom the server wishes to notify — streams a message to Apple's push notification server, providing the user token plus a “payload” that describes the notification, much as a

`UILocalNotification` does ([Chapter 26](#)). The payload is written in JSON, which is a lightweight dictionary-like structure.

Meanwhile, the user's device, if it is still on, is (with luck) connected to the network in a low-power mode that allows it to hear from Apple's push notification server. The push notification server sends the message to the user's device, where the system treats it much like a local notification: if your app isn't frontmost, the user sees an alert, and if the user taps the action button in the alert, your app is brought to the front (launching if necessary). Either way, your app can then learn what has happened through either the app delegate message `application:didReceiveRemoteNotification:` or (if the app had to be launched from scratch) through `application:didFinishLaunchingWithOptions:`, whose dictionary will contain `UIApplicationLaunchOptionsRemoteNotificationKey`. The notification itself, instead of being a `UILocalNotification` object, is an `NSDictionary` corresponding to the original JSON payload.

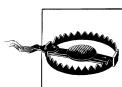
Beyond Basic Networking

An `NSURLRequest` has a cache policy, which you can set to determine whether the request might be satisfied without freshly downloading previously downloaded data. An `NSURLRequest` to be handed to an `NSURLConnection` can specify the FTP, HTTP, or HTTPS scheme, including POST requests. An `NSURLConnection` can handle redirects and authentication. See the *URL Loading System Programming Guide*. You can also get as deep into the details of networking as you like; see in particular the *CFNetwork Programming Guide*.

Apple provides a generous amount of sample code. See in particular `SimpleURLConnections`, `AdvancedURLConnections`, `SimpleNetworkStreams`, `SimpleFTPSample`, and `MVCNetworking`.

CHAPTER 38

Threads



This chapter not yet revised for iOS 5.

A *thread* is, simply put, a subprocess of your app that can execute even while other such subprocesses are also executing. Such simultaneous execution is called *concurrency*. The iOS frameworks use threads all the time; if they didn’t, your app would be less responsive to the user — perhaps even completely unresponsive. The genius of the frameworks, though, is that they use threads precisely so that you don’t have to.

For example, suppose your app is downloading something from the network ([Chapter 37](#)). This download doesn’t happen all by itself; somewhere, someone is running code that interacts with the network and obtains data. Yet a long download doesn’t prevent *your* code from running, nor does it prevent the user from tapping and swiping things in your interface. That’s concurrency in action.

It is a testament to the ingenuity of the iOS frameworks that this book has proceeded so far with no explicit discussion of threads. Indeed, it would have been nice to avoid the topic altogether. Threads are difficult and dangerous, and if at all possible you should avoid them. But sometimes that *isn’t* possible. So this chapter introduces threads, along with a warning: here be dragons. There is much more to threading, and especially to making your threaded code safe, than this chapter can possibly touch on. For detailed information about the topics introduced in this chapter, read Apple’s *Concurrency Programming Guide* and *Threading Programming Guide*.

The Main Thread

You are always using *some* thread. All your code must run somewhere; “somewhere” means a thread. When code calls a method, that method normally runs on the same thread as the code that called it. Your code is called through events ([Chapter 11](#)); those

What Is Concurrency, Really?

Part of the power of threads is that they implement concurrency without your having to worry about precisely how they do it. On a Mac Pro with eight cores, you could theoretically run eight threads truly *simultaneously*, one on each core. An iPhone device, however, has only one core. Nevertheless, an iOS app is multithreaded. How can this be? Basically, the processor performs a little code from one thread, then a little code from another, then a little code from yet another, and so on; it switches its attention between threads so quickly that they seem to run at the same time. But this is still concurrency. Concurrency is every bit as meaningful a reality on an iPhone as on a Mac Pro; how it is implemented at the level of the kernel and the processor is, mercifully, no concern of yours.

events normally call your code on the *main thread*. The main thread has certain special properties:

The main thread automatically has a run loop

A *run loop* is a recipient of events. It is how your code is notified that something is happening; without a run loop, a thread can't receive events. Cocoa events normally arrive on the main thread's run loop; that's why your code, called by those events, executes on the main thread.

The main thread is the interface thread

When the user interacts with the interface, those interactions are reported on the main thread. When your code interacts with the interface, it must do so on the main thread. Of course that will normally happen automatically, because your code normally runs on the main thread.

The main thread thus has a very great deal of work to do. Here's how life goes in your app:

1. An event arrives on the main thread; the user has tapped a button, for example, and now its control action handler is being called.
2. The event calls your code on the main thread. Your code now runs on the main thread. While this happens, nothing else can happen on the main thread. Your code might command some changes in the interface; this is safe, because your code is running on the main thread.
3. Your code finishes. The main thread's run loop is now free to report more events, and the user is free to interact with the interface once again.

The bottleneck here is obviously step 2, the running of your code. Your code runs on the main thread. That means the main thread can't do anything else while your code is running. No events can arrive while your code is running. The user can't interact with the interface while your code is running. But this is usually no problem, because:

- Your code executes really fast. It's true that the user can't interact with the interface while your code runs, but this is such a tiny interval of time that the user will probably never even notice.
- Your code, as it runs, blocks the user from interacting with the interface. But that's not bad: it's good! Your code, in response to what the user does, might update the interface; it would be insane if the user could do something else in the interface while you're in the middle of updating it.

On the other hand, as I've already mentioned, the frameworks operate in secondary threads all the time. The reason this doesn't affect you is that they talk to *your* code on the *main* thread. You have seen many examples of this in the preceding chapters. For example:

- You saw in [Chapter 17](#) that during an animation, the interface remains responsive to the user, and it is possible for your code to run. The Core Animation framework is running the animation and updating the presentation layer on a background thread. But any delegate methods or completion blocks are called on the main thread.
- You saw in [Chapter 24](#) that a UIWebView's fetching and loading of its content is asynchronous; that means the work is done in a background thread. But delegate methods are called on the main thread. The same is true of downloading a resource from the network with NSURLConnection ([Chapter 37](#)).
- You saw in [Chapter 27](#) that sounds are played asynchronously. But delegate methods are called on the main thread. Obviously, the same is true of music players ([Chapter 29](#)). Similarly, you saw in [Chapter 28](#) that movie loading, preparation, and playing happens asynchronously. But delegate methods are called on the main thread.
- You saw in [Chapter 28](#) and [Chapter 30](#) that saving a movie file takes time. So the saving takes place on a background thread. But delegate methods or completion blocks are called on the main thread.

Thus, you can (and should) usually ignore threads and just keep plugging away on the main thread. However, there are two kinds of situation in which your code will need to be explicitly aware of threading issues:

Your code is not called back on the main thread.

Some frameworks explicitly inform you that callbacks are not guaranteed to take place on the main thread. For example, CATiledLayer ([Chapter 20](#)) warns that `drawLayer:inContext:` is called in a background thread. As I explained, this means that in iOS 3.x and earlier, you can't draw here using UIKit, because UIKit classes are generally not thread-safe. In iOS 4, not only are the UIKit drawing-related classes thread-safe, but so is accessing the current context, so you can draw with `drawRect:` instead. Nevertheless, even in iOS 4 you need to be aware of the fact that

`drawRect:`, when triggered by a CATiledLayer, might be running in a background thread.

Similarly, AV Foundation ([Chapter 28](#)) warns that its blocks and notifications can arrive on a background thread. So if you intend to update the user interface, or use a value that might also be used by your main-thread code, you'll need to be thread-conscious.

Your code takes significant time.

If your code takes a long time to run and if running it on the main thread would prevent user interaction, you'll need to run that code on a background thread instead. This isn't just a matter of aesthetics; the system will actually kill your app if it discovers that its main thread is blocked for too long.

Why Threading Is Hard

The one certain thing about computer code is that it just clunks along the path of execution, one statement at a time. Lines of code, in effect, are performed in the order in which they appear. With threading, that certainty goes right out the window. If you have code that can be performed on a background thread, then you don't know when it will be performed in relation to your main-thread code. Any line of your background-thread code could be interleaved between any two lines of your main-thread code.

You also might not know *how many times* a piece of your background-thread code might be running simultaneously. Unless you take steps to prevent it, the same code could be spawned off as a thread even while it's already running in a thread. So any line of your background-thread code could be interleaved between any two lines of *itself*.

This situation is particularly threatening with regard to *shared data*. Suppose two threads were to get hold of the same object and change it. Who knows what horrors might result? Objects in general have state, adding up to the state of your app as a whole. If multiple threads are permitted to access your objects, they and your entire app can be put into an indeterminate or nonsensical state.

This problem cannot be solved by simple logic. For example, suppose you try to make data access safe with a condition, as in this pseudo-code:

```
if (no other thread is touching this data)
    do something to the data...
```

Such logic cannot succeed. Suppose the condition succeeds; no other thread is touching this data. But between the time when that condition is executed and the time when the next line executes and you start to do something to the data, another thread can come along and start touching the data!

It is possible to request assistance at a deeper level to ensure that a section of code is not run by two threads simultaneously. For example, you can implement a *lock* around

a section of code. But locks generate an entirely new level of potential pitfalls. In general, a lock is an invitation to forget to use the lock, or to forget to remove the lock after you've set it. And threads can end up contending for a lock in a way that permits neither thread to proceed.

Another problem is that the lifetime of a thread is independent of the lifetimes of other objects in your app. When an object is about to go out of existence and its `dealloc` is called, you are guaranteed that after `dealloc` calls `super`, none of your code in that object will ever run again. But a thread might still be running, and might try to talk to your object, even after your object has gone out of existence. You cannot solve this problem by having the thread retain your object, because then there is the danger that the thread might be the *last* code retaining your object, so that when the thread releases your object, its `dealloc` is called on that thread rather than the main thread, which could be a disaster.

Not only is threaded code hard to get right; it's also hard to test and hard to debug. It introduces indeterminacy, so you can easily make a mistake that never appears in your testing, but that does appear for some user. The real danger is that the user's experience will consist only of distant consequences of your mistake, long after the point where you made it, making the real cause of the problem extraordinarily difficult to track down.

All of this is meant to scare you away from using threads if you can possibly avoid it. For an excellent (and suitably frightening) account of some of the dangers and considerations that threading involves, see Apple's tech note TN2109. If terms like *race condition* and *deadlock* don't strike fear into your veins, look them up on Wikipedia.

Three Ways of Threading

Without pretending to completeness or even safety, this section will illustrate three approaches to threading. To give the examples a common base, we envision an app that draws the Mandelbrot set. (The actual code, not all of which is shown here, is adapted from a small open source project I downloaded from the Internet.) All it does is draw the basic Mandelbrot set in black and white, but that's enough number-crunching to introduce a significant delay. The idea is then to see how we can get that delay off the main thread.

The app contains a `UIView` subclass, `MyMandelbrotView`, which has one instance variable:

```
@interface MyMandelbrotView : UIView {  
    CGContextRef bitmapContext;  
}  
// ...  
@end
```

Here's the structure of its implementation:

```

// jumping-off point: draw the Mandelbrot set
- (void) drawThatPuppy {
    [self makeBitmapContext: self.bounds.size];
    CGPoint center =
        CGPointMake(CGRectGetMidX(self.bounds), CGRectGetMidY(self.bounds));
    [self drawAtCenter: center zoom: 1];
    [self setNeedsDisplay];
}

// create (and memory manage) instance variable
- (void) makeBitmapContext:(CGSize)size {
    if (self->bitmapContext)
        CGContextRelease(self->bitmapContext);
    // ... configure arguments ...
    CGContextRef context = CGBitmapContextCreate(NULL, /* ... */);
    self->bitmapContext = context;
}

// draw pixels of self->bitmapContext
- (void) drawAtCenter:(CGPoint)center zoom:(CGFloat)zoom {
    // .... do stuff to self->bitmapContext
}

// turn pixels of self->bitmapContext into CGImage, draw into ourselves
- (void) drawRect:(CGRect)rect {
    CGContextRef context = UIGraphicsGetCurrentContext();
    CGImageRef im = CGBitmapContextCreateImage(self->bitmapContext);
    CGContextDrawImage(context, self.bounds, im);
    CGImageRelease(im);
}

// final memory management
- (void) dealloc {
    if (self->bitmapContext)
        CGContextRelease(bitmapContext);
    [super dealloc];
}

```

(I haven't discussed creating a bitmap context from scratch; see "Graphics Contexts" in the *Quartz 2D Programming Guide* for example code. In this case, we take advantage of an iOS 4 feature that lets us pass NULL as the first argument to `CGBitmapContextCreate`, which relieves us of the responsibility for creating and memory-managing a data buffer associated with the graphics context.)

The `drawAtCenter:zoom:` method, which calculates the pixels of the instance variable `bitmapContext`, is time-consuming. We will consider three ways of moving this work off onto a background thread: with an old-fashioned manual thread, with NSOperation, and with Grand Central Dispatch.

Manual Threads

The simple way to make a thread manually is to send `performSelectorInBackground:withObject:` to some object containing a method to be performed on a background thread. Even with this simple approach, there is additional work to do:

Pack the arguments.

The method designated by the first argument to `performSelectorInBackground:withObject:` can take only one parameter, which you supply as the second argument to `performSelectorInBackground:withObject:`. So, if you want to pass more than one piece of information into the thread, or if the information you want to pass isn't an object, you'll need to pack it into a single object. Typically, this will be an `NSDictionary`. You should create this dictionary with `init...` and release it after passing it, because passing an autoreleased argument is unreliable.

Set up an autorelease pool.

Secondary threads don't participate in the global autorelease pool. So the first thing you must do in your threaded code is to create an `NSAutoreleasePool`, and the last thing you must do in your threaded code is to release it. Otherwise, you'll probably leak memory as autoreleased objects are created behind the scenes and are never released.

We'll rewrite `MyMandelbrotView` to use manual threading. Our `drawAtCenter:zoom:` method takes two parameters (and neither is an object), so we'll have to pack the argument that we pass into the thread, as a dictionary. Once inside the thread, we'll set up our autorelease pool and unpack the dictionary. This will all be made much easier if we interpose a trampoline method between `drawThatPuppy` and `drawAtCenter:zoom:`. So our implementation now looks like this (ignoring the parts that haven't changed):

```
- (void) drawThatPuppy {
    [self makeBitmapContext: self.bounds.size];
    CGPoint center =
        CGPointMake(CGRectGetMidX(self.bounds), CGRectGetMidY(self.bounds));
    NSDictionary* d = [[NSDictionary alloc] initWithObjectsAndKeys:
        [NSNumber numberWithCGPoint:center], @"center",
        [NSNumber numberWithInt: 1], @"zoom",
        nil];
    [self performSelectorInBackground:@selector(reallyDraw:) withObject:d];
    [d release];
    // [self setNeedsDisplay];
}

// trampoline, background thread entry point
- (void) reallyDraw: (NSDictionary*) d {
    NSAutoreleasePool* pool = [[NSAutoreleasePool alloc] init];
    [self drawAtCenter: [[d objectForKey:@"center"] CGPointValue]
                  zoom: [[d objectForKey:@"zoom"] intValue]];
    [pool release];
}
```

So far so good, but we haven't yet figured out how to draw our view. We have commented out the call to `setNeedsDisplay` in `drawThatPuppy`, because it's too soon; the call to `performSelectorInBackground:withObject:` launches the thread and returns immediately, so our `bitmapContext` instance variable isn't ready yet. Clearly, we need to call `setNeedsDisplay` after `drawAtCenter:zoom:` finishes generating the pixels of the graphics context. We can do this at the end of our trampoline method `reallyDraw:`, but we must remember that we're now in a background thread. Because `setNeedsDisplay` is a form of communication with the interface, we should call it on the main thread. We can do that with easily with `performSelectorOnMainThread:withObject:waitUntilDone:`. For maximum flexibility, it will probably be best to implement a second trampoline method:

```
// trampoline, background thread entry point
- (void) reallyDraw: (NSDictionary*) d {
    NSAutoreleasePool* pool = [[NSAutoreleasePool alloc] init];
    [self drawAtCenter: [[d objectForKey:@"center"] CGPointValue]
                  zoom: [[d objectForKey:@"zoom"] intValue]];
    [self performSelectorOnMainThread:@selector(allDone)
                           withObject:nil waitUntilDone:NO];
    [pool release];
}

// called on main thread! background thread exit point
- (void) allDone {
    [self setNeedsDisplay];
}
```

This code is specious; the seeds of nightmare are already sown. We now have a single object, `MyMandelbrotView`, some of whose methods are to be called on the main thread and some on a background thread; this invites us to become confused at some later time. Even worse, the main thread and the background thread are constantly sharing a piece of data, the instance variable `bitmapContext`; what's to stop some other code from coming along and triggering `drawRect:` while `drawAtCenter:zoom:` is in the middle of filling `bitmapContext`?

To solve these problems, we might need to use locks, and we would probably have to manage the thread more explicitly. For instance, we might use the `NSThread` class, which lets us retain our thread as an instance and query it from outside (with `isExecuting` and similar). For example code, see Apple's Metronome example, in *MetronomeView.m*. You'll observe immediately that the Metronome thread-handling code is quite elaborate and difficult to understand, yet it is an extremely basic implementation. In any case, it will be easier at this point to use `NSOperation`, the subject of the next threading approach.

NSOperation

The essence of `NSOperation` is that it encapsulates a task, not a thread. The operation described by an `NSOperation` object may be performed on a background thread, but

you don't have to concern yourself with that directly. You describe the operation and add the NSOperation to an NSOperationQueue to set it going. When the operation finishes, you are notified, typically by the NSOperation posting a notification. You can query both the queue and its operations from outside with regard to their state.

We'll rewrite MyMandelbrotView to use NSOperation. We need a new instance variable, an NSOperationQueue; we'll call it `queue`. And we have a new class, `MyMandelbrotOperation`, an NSOperation subclass. It is possible to take advantage of a built-in NSOperation subclass such as `NSInvocationOperation` or (in iOS 4) `NSBlockOperation`, but I'm deliberately illustrating the more general case by subclassing NSOperation itself.

Our implementation of `drawThatPuppy` makes sure that the queue exists; it then creates an instance of `MyMandelbrotOperation`, configures it, registers for its notification, and adds it to the queue:

```
- (void) drawThatPuppy {
    CGPoint center =
    CGPointMake(CGRectGetMidX(self.bounds), CGRectGetMidY(self.bounds));
    if (!self.queue) {
        NSOperationQueue* q = [[NSOperationQueue alloc] init];
        self.queue = q; // retain policy
        [q release];
    }
    NSOperation* op =
        [[MyMandelbrotOperation alloc] initWithSize:self.bounds.size
                                         center:center zoom:1];
    [[NSNotificationCenter defaultCenter] addObserver:self
                                             selector:@selector(operationFinished:)
                                               name:@"MyMandelbrotOperationFinished"
                                             object:op];
    [self.queue addOperation:op];
    [op release];
}
```

Our time-consuming calculations are performed by `MyMandelbrotOperation`. An NSOperation subclass, such as `MyMandelbrotOperation`, will typically have at least two methods:

A designated initializer

The NSOperation may need some configuration data. Once the NSOperation is added to a queue, it's too late to talk to it, so you'll usually hand it this configuration data as you create it, in its designated initializer.

A main method

This method will be called (with no parameters) automatically by the `NSOperationQueue` when it's time for the NSOperation to start.

Here's the interface for `MyMandelbrotOperation`:

```

@interface MyMandelbrotOperation : NSOperation {
    CGSize size;
    CGPoint center;
    CGFloat zoom;
    CGContextRef bitmapContext;
}
- (id) initWithSize: (CGSize) sz center: (CGPoint) c zoom: (CGFloat) z;
- (CGContextRef) bitmapContext;
@end

```

We have provided three instance variables for configuration, to be set in the initializer. Because `MyMandelbrotOperation` is completely separate from `MyMandelbrotView`, it must be told `MyMandelbrotView`'s size explicitly in the initializer. `MyMandelbrotOperation` also has its own `CGContextRef` instance variable, `bitmapContext`, along with an accessor so `MyMandelbrotView` can retrieve a reference to this graphics context when the operation has finished. Note that this is different from `MyMandelbrotView`'s `bitmapContext`; one of the benefits of using `NSOperation` is that we are no longer sharing data so promiscuously between threads.

Here's the implementation for `MyMandelbrotOperation`. All the calculation work has been transferred from `MyMandelbrotView` to `MyMandelbrotOperation` without change; the only difference is that `bitmapContext` now means `MyMandelbrotOperation`'s instance variable:

```

- (id) initWithSize: (CGSize) sz center: (CGPoint) c zoom: (CGFloat) z {
    self = [super init];
    if (self) {
        self->size = sz;
        self->center = c;
        self->zoom = z;
    }
    return self;
}

- (void) dealloc {
    if (self->bitmapContext)
        CGContextRelease(self->bitmapContext);
    [super dealloc];
}

- (CGContextRef) bitmapContext {
    return self->bitmapContext;
}

- (void) makeBitmapContext:(CGSize) size {
    // ... same as before ...
}

- (void) drawAtCenter:(CGPoint) center zoom:(CGFloat) zoom {
    // ... same as before ...
}

- (void) main {

```

```

if ([self isCancelled])
    return;
[self makeBitmapContext: self->size];
[self drawAtCenter: self->center zoom: self->zoom];
if (![self isCancelled])
    [[NSNotificationCenter defaultCenter]
        postNotificationName:@"MyMandelbrotOperationFinished" object:self];
}

```

The only method of interest is `main`. First, we test through the `NSOperation` method `isCancelled` to make sure we haven't been cancelled while sitting in the queue; this is good practice. Then, we do exactly what `drawThatPuppy` used to do, initializing our graphics context and drawing into its pixels.

When the operation is over, we need to notify `MyMandelbrotView` to come and fetch our data. There are two ways to do this; either `main` can post a notification through the `NSNotificationCenter`, or `MyMandelbrotView` can use key-value observing ([Chapter 13](#)) to be notified when our `isFinished` key path changes. We've chosen the former approach; observe that we check one more time to make sure we haven't been cancelled.

Now we are back in `MyMandelbrotView`, hearing that `MyMandelbrotOperation` has finished. We must immediately pick up any required data, because the `NSOperationQueue` is about to release this `NSOperation`. However, we must be careful; the notification may have been posted on a background thread, in which case our method for responding to it will also be called on a background thread. We are about to set our own graphics context and tell ourselves to redraw; those are things we want to do on the main thread. So we immediately step out to the main thread:

```

// warning! called on background thread
- (void) operationFinished: (NSNotification*) n {
    [self performSelectorOnMainThread:@selector(redrawWithOperation:)
        withObject:[n object] waitUntilDone:NO];
}

```

As we set `MyMandelbrotView`'s `bitmapContext` by reading `MyMandelbrotOperation`'s `bitmapContext`, we must concern ourselves with the memory management of a `CGContext` obtained from an object that may be about to release that context:

```

// now we're back on the main thread
- (void) redrawWithOperation: (MyMandelbrotOperation*) op {
    [[NSNotificationCenter defaultCenter]
        removeObserver:self
        name:@"MyMandelbrotOperationFinished"
        object:op];
    CGContextRef context = [op bitmapContext];
    if (self->bitmapContext)
        CGContextRelease(self->bitmapContext);
    self->bitmapContext = (CGContextRef) context;
    CGContextRetain(self->bitmapContext);
    [self setNeedsDisplay];
}

```

Using NSOperation instead of manual threading may not seem like any reduction in work, but it is a tremendous reduction in headaches:

The operation is encapsulated.

Because MyMandelbrotOperation is an object, we've been able to move all the code having to do with drawing the pixels of the Mandelbrot set into it. No longer does MyMandelbrotView contain some code to be called on the main thread and some code to be called on a background thread. The *only* MyMandelbrotView method that can be called in the background is `operationFinished:`, and that's a method we'd never call explicitly ourselves, so we won't misuse it accidentally.

The data sharing is rationalized.

Because MyMandelbrotOperation is an object, it has its own `bitmapContext`. The only moment of data sharing comes in `redrawWithOperation:`, when we must set MyMandelbrotView's `bitmapContext` to MyMandelbrotOperation's `bitmapContext`. Even if multiple MyMandelbrotOperation objects are added to queue, the moments when we set MyMandelbrotView's `bitmapContext` all occur on the main thread, so they cannot conflict with one another.

The coherence of MyMandelbrotView's `bitmapContext` does depend upon our obedience to an implicit contract not to set it or write into it anywhere except a few specific moments in MyMandelbrotView's code. But this is always a problem with data sharing in a multithreaded world, and we have done all we can to simplify the situation.

If we are concerned with the possibility that more than one instance of MyMandelbrotOperation might be added to the queue and executed concurrently, we have a further defense — we can set the NSOperationQueue's maximum concurrency level to 1:

```
NSOperationQueue* q = [[NSOperationQueue alloc] init];
[q setMaxConcurrentOperationCount:1];
self.queue = q;
[q release];
```

This turns the NSOperationQueue into a true serial queue; every operation on the queue must be completely executed before the next can begin. This might cause an operation added to the queue to take longer to execute, if it must wait for another operation to finish before it can even get started; however, this delay might not be important. What *is* important is that by executing the operations on this queue completely separately, we guarantee that only one operation at a time can do any data sharing. A serial queue is thus a form of data locking.

Because MyMandelbrotView can be destroyed (if, for example, its view controller is destroyed), there is still a risk that it will create an operation that will outlive it and will try to access it after it has been destroyed. We can reduce that risk by canceling all operations in our queue before releasing it:

```

- (void)dealloc {
    // release the bitmap context
    if (self->bitmapContext)
        CGContextRelease(bitmapContext);
    [self->queue cancelAllOperations];
    [self->queue release];
    [super dealloc];
}

```

Grand Central Dispatch

Grand Central Dispatch, or *GCD*, introduced in iOS 4, is a sort of low-level analogue to *NSOperation* and *NSOperationQueue*. When I say low-level, I'm not kidding; GCD is baked into the operating system kernel. Thus it can be used by any code whatsoever and is tremendously efficient.

Using GCD is like a mixture of the manual threading approach with the *NSOperationQueue* approach. It's like the manual threading approach because code to be executed on one thread appears together with code to be executed on another; however, you have a better chance of keeping the threads and data management straight, because GCD uses Objective-C blocks. It's like the *NSOperationQueue* approach because it uses queues; you express a task and add it to a queue, and the task is executed on a thread as needed. Moreover, by default these queues are serial queues, with each task on a queue finishing before the next is started, which, as we've already seen, is a form of data locking.

We'll rewrite *MyMandelbrotView* to use GCD. The structure of its interface is very slightly changed from the original, nonthreaded version. We have a new instance variable to hold our GCD queue; *makeBitmapContext*: now returns a graphics context rather than setting an instance variable directly; and *drawAtCenter:zoom:* now takes an additional parameter, the graphics context to draw into:

```

@interface MyMandelbrotView : UIView {
    CGContextRef bitmapContext;
    dispatch_queue_t draw_queue;
}
- (void)drawAtCenter:(CGPoint)center zoom:(CGFloat)zoom context:(CGContextRef)c;
- (CGContextRef)makeBitmapContext:(CGSize)size;
- (void) drawThatPuppy;
@end

```

In *MyMandelbrotView*'s implementation, we add management to create our GCD queue as the view is created and to tear it down as the view is destroyed:

```

- (id)initWithCoder:(NSCoder *)aDecoder {
    self = [super initWithCoder: aDecoder];
    if (self) {
        self->draw_queue = dispatch_queue_create("com.neuburg.mandeldraw", NULL);
    }
    return self;
}

```

```

- (void) dealloc {
    if (bitmapContext)
        CGContextRelease(bitmapContext);
    dispatch_release(draw_queue);
    [super dealloc];
}

```

Now for the implementation of `drawThatPuppy`. Here it is:

```

- (void) drawThatPuppy {
    CGPoint center = ❶
    CGPointMake(CGRectGetMidX(self.bounds), CGRectGetMidY(self.bounds));
    dispatch_async(draw_queue, ^{
        CGContextRef bitmap = [self makeBitmapContext: self.bounds.size];
        [self drawAtCenter: center zoom: 1 context:bitmap];
        dispatch_async(dispatch_get_main_queue(), ^{
            ❷
            if (self->bitmapContext)
                CGContextRelease(self->bitmapContext);
            self->bitmapContext = bitmap;
            [self setNeedsDisplay];
        });
    });
}

```

That's all there is to it. No trampoline methods. No `performSelector`.... No packing arguments into a dictionary. No autorelease pools. No instance variables. And effectively no sharing of data across threads. That's the beauty of blocks.

- ❶ We begin by calculating our `center`, as before. This value will be visible within the blocks, because blocks can see their surrounding context.
- ❷ Now comes our task to be performed in a background thread on our queue, `draw_queue`. We specify this task with the `dispatch_async` function. GCD has a lot of functions, but this is the one you'll use 99 percent of the time; it's the most important thing you need to know about GCD. We specify a queue and we provide a block saying what we'd like to do. Thanks to the block, we don't need any trampoline methods. In the block, we begin by declaring `bitmap` as a variable *local to the block*. We then call `makeBitmapContext:` to create the graphics context `bitmap`, and `drawAtCenter:zoom:context:` to set its pixels; we make these calls *directly*, just as we would do if we weren't threading in the first place.
- ❸ Now we need to get back onto the main thread. How do we do that? With `dispatch_async` again! We specify the main queue (which is effectively the main thread) with a function provided for this purpose and describe what we want to do in *another* block. This second block is nested inside the first, so it isn't performed until the preceding commands in the first block have finished; moreover, because the first block is part of its surrounding context, the second block can see our block-local `bitmap` variable! We set our `bitmapContext` instance variable (with no need for further memory management, because `makeBitmapContext` has returned a retained graphics context), and call `setNeedsDisplay`.

The benefits and elegance of GCD are stunning. The `bitmap` variable is not shared; it is local to each specific call to `drawThatPuppy`. The nested blocks are executed in succession, so any instance of `bitmap` must be completely filled with pixels before being used to set the `bitmapContext` instance variable. Moreover, `bitmapContext` is set on the main thread, plus the *entire* operation is performed on a serial queue; thus data sharing is reduced to a minimum, with no possibility of conflict. Our code is also highly maintainable, because the entire task on all threads is expressed within the single `drawThatPuppy` method, thanks to the use of blocks; indeed, the code is only very slightly modified from the original, nonthreaded version.

Threads and App Backgrounding

When your app is backgrounded and suspended ([Chapter 11](#)), a problem arises if your code is running. The system doesn't want to kill your code while it's executing; on the other hand, some other app may need to be given the bulk of the device's resources now. So as your app goes into the background, the system waits a short time for your app to finish doing whatever it may be doing, but it then suspends your app and stops it by force.

This shouldn't be a problem from your main thread's point of view, because your app shouldn't have any time-consuming code on the main thread in the first place; you now know that you can avoid this by using a background thread. On the other hand, it could be a problem for lengthy background operations, including asynchronous tasks performed by the frameworks. You can request time to complete a lengthy task (or at least abort it yourself, coherently) in case your app is backgrounded, by wrapping it in calls to `UIApplication`'s `beginBackgroundTaskWithExpirationHandler:` and `endBackgroundTask:`.

You call `beginBackgroundTaskWithExpirationHandler:` to announce that a lengthy task is beginning; it returns an identification number. At the end of your lengthy task, you call `endBackgroundTask:,` passing in that same identification number. This tells the application that your lengthy task is over and that, if your app has been backgrounded while the task was in progress, it is now okay to suspend you.

The argument to `beginBackgroundTaskWithExpirationHandler:` is a block, but this block does *not* express the lengthy task. It expresses what you will do *if your extra time expires* before you finish your lengthy task. At the very least, your expiration handler must call `endBackgroundTask:,` just as your lengthy task would have done; otherwise, your app won't just be suspended — it will be killed.



If your expiration handler block *is* called, you should make no assumptions about what thread it is running on.

Let's use `MyMandelbrotView`, from the preceding section, as an example. Let's say that if `drawThatPuppy` is started, we'd like it to be allowed to finish, even if the app is suspended in the middle of it, so that our `bitmapContext` instance variable is updated as requested. To try to ensure this, we call `beginBackgroundTaskWithExpirationHandler:` beforehand and call `endBackgroundTask:` at the end of the innermost block:

```
- (void) drawThatPuppy {
    CGPoint center =
        CGPointMake(CGRectGetMidX(self.bounds), CGRectGetMidY(self.bounds));
    UIBackgroundTaskIdentifier bti = [[UIApplication sharedApplication]
        beginBackgroundTaskWithExpirationHandler: ^{
            [[UIApplication sharedApplication] endBackgroundTask:bti];
        }];
    dispatch_async(draw_queue, ^{
        CGContextRef bitmap = [self makeBitmapContext: self.bounds.size];
        [self drawAtCenter: center zoom: 1 context:bitmap];
        dispatch_async(dispatch_get_main_queue(), ^{
            if (self->bitmapContext)
                CGContextRelease(self->bitmapContext);
            self->bitmapContext = bitmap;
            [self setNeedsDisplay];
            [[UIApplication sharedApplication] endBackgroundTask:bti];
        });
    });
}
```

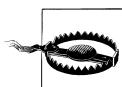
If our app is backgrounded while `drawThatPuppy` is in progress, it will (we hope) be given enough time to live that it can run all the way to the end. Thus, the instance variable `bitmapContext` will be updated, and `setNeedsDisplay` will be called, before we are actually suspended. Our `drawRect:` will not be called until our app is brought back to the front, but there's nothing wrong with that.

It's actually pretty good policy to use a similar technique when you're notified that your app is being backgrounded. It's common practice to respond to the app delegate message `applicationDidEnterBackground:` (or the corresponding `UIApplicationDidEnterBackgroundNotification`) by saving state and reducing memory usage, but this can take time, whereas what you'd like to do is return from `applicationDidEnterBackground:` as quickly as possible. A reasonable solution is to implement `applicationDidEnterBackground:` very much like `drawThatPuppy` earlier: call `beginBackgroundTaskWithExpirationHandler:` and then call `dispatch_async` to get off the main thread, and do your state-saving and so forth in its block. In this case, there's no point creating your own queue, because you don't care about serial order; you can use `dispatch_get_global_queue(0)`, the default-priority global background queue.

What about lengthy asynchronous operations such as networking ([Chapter 37](#))? As far as I can tell, it might not strictly be necessary to use `beginBackgroundTaskWithExpirationHandler:` in connection with `NSURLConnection`; it appears that `NSURLConnection` has the ability to resume automatically after an interruption when your app is suspended. Still, it might be better not to rely on that behavior (or on an assumption that, just because the network is present now, it will be present when the app awakes from

suspension), so you might like to integrate `beginBackgroundTaskWithExpirationHandler:` into your use of `NSURLConnection`.

Such integration can be just a little tricky, because `beginBackgroundTaskWithExpirationHandler:` and `endBackgroundTask:` rely on a shared piece of information, the `UIBackgroundTaskIdentifier`—but the downloading operation begins in one place (when the `NSURLConnection` is created, or when it is told to `start`) and ends in one of two other places (the `NSURLConnection`'s delegate is informed that the download has failed or succeeded), so information is not so easily shared. However, with something like our `MyDownloader` class, an entire single downloading operation is encapsulated, and we can give the class a `UIBackgroundTaskIdentifier` instance variable. So, we would set this instance variable with a call to `beginBackgroundTaskWithExpirationHandler:` just before telling the connection to `start`, and then both `connection:didFailWithError:` and `connectionDidFinishLoading:` would use the value stored in that instance variable to call `endBackgroundTask:` as their last action.



This chapter not yet revised for iOS 5.

The ability to undo the most recent action is familiar from Mac OS X. Typically, a Mac application will maintain an internal stack of undoable actions; choosing Edit → Undo or pressing Command-Z will reverse the action at the top of the stack, and will also make that action available for Redo so that the user can undo the most recent Undo. The idea is that, provided the user realizes soon enough that a mistake has been made, that mistake can be reversed.

A pervasive, extensive implementation of Undo makes sense on Mac OS X, especially when real-life objects are involved. For example, a window may represent an actual document, a file on disk; it would be terrible if every word typed or pasted or cut represented a permanent, irreversible change to that document. Given the transient, visual nature of the iOS interface, however, and the sorts of thing for which iOS apps are typically intended, users do not generally expect Undo at all.

Nevertheless, some iOS apps may benefit from at least a limited version of this facility. Not every action needs to be undoable, and the ability to undo needn't persist for very long. And limited Undo is not difficult to implement. Some built-in views — in particular, those that involve text entry, UITextField and UITextView ([Chapter 23](#)) — implement Undo already. And you can add it in other areas of your app.

The Undo Manager

Undo is provided through an instance of NSUndoManager, which basically just maintains a stack of undoable actions, along with a secondary stack of redoable actions. The goal in general is to work with the NSUndoManager so as to take care of handling both Undo and Redo in the standard manner: When the user chooses to undo the most

recent action, the action at the top of the Undo stack is popped off and reversed and is pushed onto the top of the Redo stack.

To illustrate, I'll use an artificially simple app in which the user can drag a small square around the screen. We'll start with an instance of a `UIView` subclass, `MyView`, to which has been attached a `UIPanGestureRecognizer` to make it draggable, as described in [Chapter 18](#). The gesture recognizer's action target is the `MyView` instance itself:

```
- (void) dragging: (UIPanGestureRecognizer*) p {
    if (p.state == UIGestureRecognizerStateBegan || 
        p.state == UIGestureRecognizerStateChanged) {
        CGPoint delta = [p translationInView: self.superview];
        CGPoint c = self.center;
        c.x += delta.x; c.y += delta.y;
        self.center = c;
        [p setTranslation: CGPointZero inView: self.superview];
    }
}
```

To make dragging of this view undoable, we need an `NSUndoManager` instance. Let's store this in an instance variable of `MyView` itself, accessible through a property, `undoer`.

There are two ways to register an action as undoable. One involves the `NSUndoManager` method `registerUndoWithTarget:selector:object:..`. This method uses a target-action architecture: you provide a target, a selector for a method that takes one parameter, and the object value to be passed as argument when the method is called. Then, later, if the `NSUndoManager` is sent the `undo` message, it simply sends that action with that argument to that target. What we want to undo here is the setting of our `center` property. This can't be expressed directly using a target-action architecture: we can call `setCenter:..`, but its parameter needs to be a `CGPoint`, which isn't an object. This means we're going to have to provide a secondary method that *does* take an object parameter. This is neither bad nor unusual; it is quite common for actions to have a special representation just for the purpose of making them undoable.

So, in our `dragging:..` method, instead of setting `self.center` to `c` directly, we now call a secondary method (let's call it `setCenterUndoably:..`):

```
[self setCenterUndoably: [NSValue valueWithCGPoint:c]];
```

At a minimum, `setCenterUndoably:..` should do the job that setting `self.center` used to do:

```
- (void) setCenterUndoably: (NSValue*) newCenter {
    self.center = [newCenter CGPointValue];
}
```

This works in the sense that the view is draggable exactly as before, but we have not yet made this action undoable. To do so, we must ask ourselves what message the `NSUndoManager` would need to send in order to undo the action we are about to perform. We would want the `NSUndoManager` to set `self.center` back to the value it has *now*, before we change it as we are about to do. And what method would `NSUn-`

doManager call in order to do that? It would call `setCenterUndoably:`, the very method we are implementing; that's *why* we are implementing it. So:

```
- (void) setCenterUndoably: (NSValue*) newCenter {
    [self.undoer registerUndoWithTarget:self
        selector:@selector(setCenterUndoably:)
        object:[NSValue valueWithCGPoint:self.center]];
    self.center = [newCenter CGPointValue];
}
```

This not only makes our action undoable, it also makes it redoable. Why? Because when we send `undo` to the `NSUndoManager`, it calls `setCenterUndoably:` and is immediately sent `registerUndo...` — and there's a rule that, if the `NSUndoManager` is sent this message *while it is undoing*, it puts the target–action information on the Redo stack instead of the Undo stack (because Redo *is* the Undo of an Undo, if you see what I mean). That's one of the chief tricks to working with an `NSUndoManager`: it will respond differently to `registerUndo...` depending on its state.

So far, so good. But our implementation of undoing is very annoying, because we are adding a single object to the Undo stack every time `dragging:` is called — and it is called many times during the course of a single drag. Thus, undoing merely undoes the tiny increment corresponding to one particular `dragging:` call. What we'd like, surely, is for undoing to undo an *entire* dragging gesture. We can implement this through *undo grouping*. As the gesture begins, we start a group; when the gesture ends, we end the group:

```
- (void) dragging: (UIPanGestureRecognizer*) p {
    if (p.state == UIGestureRecognizerStateBegan)
        [self.undoer beginUndoGrouping];
    if (p.state == UIGestureRecognizerStateBegan ||
        p.state == UIGestureRecognizerStateChanged) {
        CGPoint delta = [p translationInView: self.superview];
        CGPoint c = self.center;
        c.x += delta.x; c.y += delta.y;
        [self setCenterUndoably: [NSValue valueWithCGPoint:c]];
        [p setTranslation: CGPointZero inView: self.superview];
    }
    if (p.state == UIGestureRecognizerStateChanged ||
        p.state == UIGestureRecognizerStateCancelled)
        [self.undoer endUndoGrouping];
}
```

This works: each complete gesture of dragging `MyView`, from the time the user's finger contacts the view to the time it leaves, is now undoable (and then redoable) as a single unit.

Earlier I said that `registerUndo...` was one of two ways to register an action as undoable. The other is to use `prepareWithInvocationTarget:`. You provide the target and, *in the same line of code*, send to the object returned from this call the message and arguments you want sent when the `NSUndoManager` is sent `undo` (or, if we are undoing now, `redo`). So, in our example, instead of this line:

```
[self.undoer registerUndoWithTarget:self  
    selector:@selector(setCenterUndoably:)  
    object:[NSValue valueWithCGPoint:self.center]];
```

You'd say this:

```
[[self.undoer prepareWithInvocationTarget:self]  
    setCenterUndoably: [NSValue valueWithCGPoint:self.center]];
```

This code seems impossible: how can we send `setCenterUndoably:` without *calling* `setCenterUndoably:`? Either we are sending it to `self`, in which case it should actually be called at this moment, or we are sending it to some other object that doesn't implement `setCenterUndoably:`, in which case our app should crash. However, under the hood, the `NSUndoManager` is cleverly using Objective-C's dynamism (similarly to the message-forwarding example at the end of [Chapter 25](#)) to capture this call as an `NSInvocation` object, which it can use later to send the same message with the same arguments to the specified target.

In our example, `prepareWithInvocationTarget:` provides no particular advantage over `registerUndo....` In general, the advantage of `prepareWithInvocationTarget:` is that it lets you specify a method with any number of parameters, and those parameters needn't be objects.

The Undo Interface

We must now decide how to let the user request Undo and Redo. In testing the code from the preceding section, I used two buttons, an Undo button that sent `undo` to the `NSUndoManager` and a Redo button that sent `redo` to the `NSUndoManager`. This can be a perfectly reasonable interface, but let's talk about some others.

By default, your application supports *shake-to-edit*. This means the user can shake the device to bring up an undo/redo interface. We discussed this briefly in [Chapter 35](#). If you don't turn off this feature by setting the shared `UIApplication`'s `application-SupportsShakeToEdit` property to `NO`, then when the user shakes the device, the framework walks up the responder chain, starting with the first responder, looking for a responder whose `undoManager` property returns an actual `NSUndoManager` instance. If it finds one, it puts up the undo/redo interface and communicates appropriately with that `NSUndoManager`, depending on the user's choice in that interface.

You will recall what it takes for a `UIResponder` to be first responder in this sense: it must return `YES` from `canBecomeFirstResponder`, and it must actually be made first responder through a call to `becomeFirstResponder`. Let's suppose that `MyView` satisfies these requirements (for example, we might call `becomeFirstResponder` at the start of `dragging:`). Then, to make shake-to-edit work, we must also set its `undoManager` property instead of its `undoer` property as we are now doing. However, the inherited `undoManager` property is read-only, so in order to set it, we must override its property declaration, which we can do through a class extension. For example:



Figure 39-1. The shake-to-edit undo/redo interface

```
@interface MyView ()  
@property (nonatomic, retain) NSUndoManager *undoManager;  
@end  
  
@implementation MyView  
@synthesize undoManager;  
  
- (id)initWithCoder:(NSCoder *)aDecoder {  
    self = [super initWithCoder:aDecoder];  
    NSUndoManager* u = [[NSUndoManager alloc] init];  
    self.undoManager = u; // retain policy  
    [u release];  
    return self;  
}  
  
- (BOOL) canBecomeFirstResponder {  
    return YES;  
}  
  
- (void) dragging: (UIPanGestureRecognizer*) p {  
    [self becomeFirstResponder];  
    // ... the rest as before ...  
}  
//...  
@end
```

This works: shaking the device now brings up the undo/redo interface, and its buttons work correctly. However, I don't like the way the buttons are labeled; they just say Undo and Redo. To make them more expressive, we should provide a string describing each undoable action by calling `setActionName:`. We can appropriately and conveniently do this in `setCenterUndoably:`, as follows:

```
- (void) setCenterUndoably: (NSValue*) newCenter {  
    [self.undoManager registerUndoWithTarget:self  
                                         selector:@selector(setCenterUndoably:)  
                                         object:[NSValue valueWithCGPoint:self.center]];  
    [self.undoManager setActionName: @"Move"];  
    self.center = [newCenter CGPointValue];  
}
```

Now the buttons say Undo Move and Redo Move, which is a nice touch (Figure 39-1).

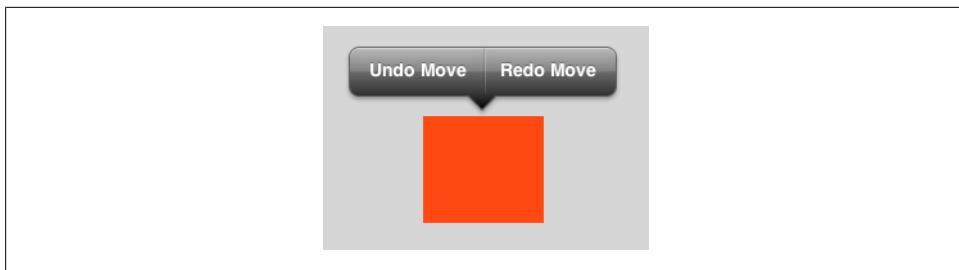


Figure 39-2. The shared menu as an undo/redo interface

Another possible interface is through a menu (Figure 39-2). Personally, I prefer this approach, as I am not fond of shake-to-edit (it seems both violent and unreliable). This is the same menu used by a UITextField or UITextView for displaying the Copy and Paste menu items (Chapter 23). The requirements for summoning this menu are effectively the same as those for shake-to-edit: we need a responder chain with a first responder at the bottom of it. So the code we've just supplied for making MyView first responder remains applicable.

We can make a menu appear, for example, in response to a long press on our MyView instance. So let's suppose we've attached another gesture recognizer to MyView. This will be a UILongPressGestureRecognizer, whose action handler is called `longPress:`. Recall from Chapter 23 how to implement the menu: we get the singleton global `UIMenuController` object and specify an array of custom `UIMenuItems` as its `menuItems` property. But a *particular* menu item will appear only if we return YES from `canPerformAction:withSender:` for that menu item's action:

```
- (void) longPress: (id) g {
    UIMenuController *m = [UIMenuController sharedMenuController];
    [m setTargetRect:self.bounds inView:self];
    UIMenuItem *mi1 =
        [[UIMenuItem alloc] initWithTitle:[self.undoManager undoMenuItemTitle]
                                      action:@selector(undo:)];
    UIMenuItem *mi2 =
        [[UIMenuItem alloc] initWithTitle:[self.undoManager redoMenuItemTitle]
                                      action:@selector(redo:)];
    [m setMenuItems:[NSArray arrayWithObjects: mi1, mi2, nil]];
    [mi1 release]; [mi2 release];
    [m setMenuVisible:YES animated:YES];
}

- (BOOL)canPerformAction:(SEL)action withSender:(id)sender {
    if (action == @selector(undo:))
        return [self.undoManager canUndo];
    if (action == @selector(redo:))
        return [self.undoManager canRedo];
    return [super canPerformAction:action withSender:sender];
}

- (void) undo: (id) dummy {
    [self.undoManager undo];
```

```
}

- (void) redo: (id) dummy {
    [self.undoManager redo];
}
```

Observe how we consult our NSUndoManager throughout. We get the titles for our custom menu items from the NSUndoManager (there might, after all, be more than one undoable kind of action, and therefore more than one title), and we know whether to display the Undo menu item or the Redo menu item by calling our NSUndoManager's `canUndo` and `canRedo`, which essentially asks whether there's anything on the respective stack.

The Undo Architecture

Implementing basic Undo is not particularly difficult. But maintaining an appropriate Undo stack at the right point (or points) in your responder hierarchy, so that the right thing happens at every moment, can require some work and some thought. Many questions can arise, and there are no simple answers.

In general, your chief concern will be maintaining a consistent state in your app and in the Undo and Redo stacks of any NSUndoManager instances. You don't want an Undo stack to contain a method call that, if actually sent, would be impossible to obey, or, if obeyed, would make nonsense of your app's state, because of things that have happened in the meantime. In order to prevent this, you have to make sure you are not implementing Undo only partially. Suppose, for example, your app presents a To-Do list in which the user can add items, edit items, and so forth. And suppose you implemented Undo and Redo for inserting an item but not for editing an item. Then if the user inserted an item and then edited it, and then did an Undo of an item insertion followed by a Redo of that item insertion, this would fail to restore the state of the app, because the editing has been omitted from the Redo.

This is why you typically want each undoable action to pass consistently through a bottleneck method that will register this action with the NSUndoManager. And you will usually want this bottleneck method to be the same method that is registered with the NSUndoManager, so that the Undo and Redo stacks are kept synchronized properly. The sole exception involves independent constructive and destructive actions, such as insertion into a list and deletion from that list; in that case, the Undo method for insertion will be the deletion method, and the Undo method for deletion will be the insertion method. You can customize the arrangement of bottlenecks further and in more complex ways, but it's easy to become confused, so you probably won't want to.

Not all aspects of communication with an NSUndoManager need to be performed in the same place, however. We already saw this in the examples earlier in this chapter: `setCenterUndoably:`, the bottleneck method, knows what method to register with the NSUndoManager, but `dragging:` knows what a complete gesture is and therefore

knows where to place the boundaries of a group. Similarly, it happens that our bottleneck method is the one that called `setActionName:`, but in real life it will often be some other method that knows best what name should be attached to a particular action. You will thus end up with a single `NSUndoManager` being bombarded with messages from various places in your code. Indeed, `NSUndoManager` accommodates exactly this sort of design; for example, it emits many notifications for which you can register, to help tie together operations that are performed at disparate locations in your code.

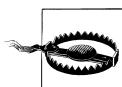
Then there are the larger architectural questions of how many `NSUndoManager` objects your app needs and how long each one needs to live. There's typically nothing wrong with an iOS app having occasional short-lived, short-depth Undo stacks and no Undo the rest of the time. Apple's `SimpleUndo` example constructs an app with an Edit interface, where the user makes changes and then taps either Cancel or Save, returning to the main interface. Here, the user can shake to undo what happened during that edit session. And that's all that's undoable within this app. If the user taps Edit again, one imagines that it would make sense to clear the existing Undo stack; there's no point in letting the user return to an earlier Edit session's state. If the user switches away to a different view controller altogether, one imagines that it would make sense to release the `NSUndoManager` completely and start with a clean slate when we come back; if the user had any intention of undoing, the time to do so was before abandoning this part of the interface.

Your architectural decisions will often be closely tied to the actual functionality and nature of your app. For example, consider again the `MyView` instance that the user can move, and whose movements the user can undo. Suppose our app has *two* `MyView` instances in the same window. In our earlier examples, we've implemented Undo at the level of the individual `MyView` instance. Is this right when there are multiple `MyView` instances, or should we move the implementation to a higher point in the responder chain that effectively contains them both — for example, to the view controller of whose view they are subviews? The answer is that there's no right answer. It depends on what makes sense for what our app actually does. If these are fairly independent objects, in terms of the app's functionality and the mental world it creates, then it might make sense to be able to undo a move of either view, independently of the other. But if these are, say, two playing cards in a deck, then obviously it isn't up to an individual card whether it can be put back into the place it was before; the only undoable card is the most recently moved of *all* cards.

For more about the `NSUndoManager` class and how to use it, read Apple's *Undo Architecture* as well as the documentation for the class itself.

CHAPTER 40

Epilogue



This chapter not yet revised for iOS 5.

You may go, for you're at liberty.

—W. S. Gilbert, *The Pirates of Penzance*

This book must come to an end, but your exploration of iOS will go on and on. There's much more to know and to explore. A single book that described completely, or even introduced, every aspect of iOS programming would be immense — many times the size of this one. Inevitably, severe limits have had to be set. Having read this book, you are now in a position to investigate many further areas of iOS that this book hasn't explored in any depth. Some of these areas have been mentioned in individual chapters; here are a few others:

OpenGL

An open source C library for drawing, including 3D drawing, that takes full advantage of graphics hardware. This is often the most efficient way to draw, especially when animation is involved. iOS incorporates a simplified version of OpenGL called OpenGL ES. See the *OpenGL Programming Guide for iOS*. Also, some forms of animated display (chiefly, but not exclusively, those using OpenGL) will benefit from CADisplayLink, a timer object that calls a method repeatedly based on the refresh rate of the screen's physical display.

Accelerate framework

Certain computation-intensive processes will benefit from the vector-based Accelerate framework, added in iOS 4. See the *vDSP Programming Guide*.

Game Kit

The Game Kit framework covers three areas that can enhance your user's game experience: Wireless or Bluetooth communication directly between devices (peer-to-peer); voice communication across an existing network connection; and Game

Center, a networking facility introduced in iOS 4.1 that facilitates these and many other aspects of interplayer communication, such as posting and viewing high scores and setting up combinations of players who wish to compete. See the *Game Kit Programming Guide*.

Advertising

The iAD framework, added in iOS 4, lets your free app attempt to make money by displaying advertisements provided by Apple. See the *iAD Programming Guide*.

Purchases

Your app can allow the user to buy something, using Apple's App Store to process payments. See the *In App Purchase Programming Guide*.

Printing

Printing was added to iOS in version 4.2. See the “Printing” chapter of the *Drawing and Printing Guide for iOS*.

Security

This book has not discussed security topics such as keychains, certificates, and encryption. See the *Security Overview* and the Security framework.

Accessibility

VoiceOver assists visually impaired users by describing the interface aloud. To participate, views must be configured to describe themselves usefully. Built-in views already do this to a large extent, and you can add to this functionality. See the *Accessibility Programming Guide for iOS*.

Telephone

The Core Telephony framework, introduced in iOS 4, lets your app get information about a particular cellular carrier and call.

External accessories

The user can attach an external accessory to the device, either directly via USB or wirelessly via Bluetooth. Your app can communicate with such an accessory. See *External Accessory Programming Topics*.